

The Sequential and Complete Transitional Course from PostgreSQL to ORACLE SQL §

Module 1: Core Oracle SQL Differences & Syntax Foundations §
*Key Differences, Data Types, DUAL, NULLs, Conditionals, ROWNUM,
Comments*

AS/Transitional SQL

May 25, 2025

Contents

1	Introduction: Charting a New Course to Oracle SQL	3
2	Oracle Data Types: The Building Blocks of Information	3
2.1	Character Data Types: VARCHAR2 and NVARCHAR2	3
2.2	Numeric Data Type: NUMBER	4
2.3	Date and Time Data Types: DATE and TIMESTAMP Variations	6
2.3.1	DATE	6
2.3.2	TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIME ZONE	7
2.4	Large Object Types: CLOB and BLOB (Brief Introduction)	8
3	DUAL Table: Oracle's Little Helper	9
4	NULL Handling: Embracing the Void (Oracle Style)	10
4.1	NVL(expr1, expr2)	11
4.2	NVL2(expr1, expr2, expr3)	12
4.3	COALESCE(expr1, expr2, ..., exprN)	12
5	Conditional Expressions: Making Choices in SQL	13
5.1	DECODE(expr, search1, result1, [search2, result2, ...], [default])	13
5.2	CASE Expression	15
6	ROWNUM Pseudo-column: Numbering Your Rows	16
7	Comments: Annotating Your Code	17
8	Bridging from PostgreSQL: Core Syntax & Foundational Concepts Summary	18

1. Introduction: Charting a New Course to Oracle SQL

WELCOME to your transitional journey from the familiar waters of PostgreSQL to the powerful realm of Oracle SQL! This module lays the foundational stones, focusing on key syntactical differences and core Oracle-specific concepts. While your PostgreSQL knowledge provides a robust compass, Oracle has its own unique landmarks and currents. Think of it as learning a new dialect of a language you already speak –many concepts are similar, but the “accent” and some “local phrases” are distinct.

We’ll cover essential building blocks: Oracle’s data types (some familiar, some new twists), the curious DUAL table, how Oracle prefers to handle NULL values, its conditional expressions (including the vintage DECODE), the often-misunderstood ROWNUM, and the universal language of comments.

Remember, SQL is SQL, but vendors add their spice. Oracle’s spice is often potent! Let’s get started.

2. Oracle Data Types: The Building Blocks of Information

Oracle offers a rich set of data types. Many will feel familiar from PostgreSQL, but Oracle often has its own flavor or specific type names. Understanding these is crucial for defining table structures and variables correctly.

2.1. Character Data Types: VARCHAR2 and NVARCHAR2

What Are They? (Meanings & Values in Oracle) `VARCHAR2(size [BYTE | CHAR])`¹ is Oracle’s primary variable-length character string type. It stores alphanumeric data. The ‘size’ specifies the maximum length. ‘BYTE’ (default) means the size is in bytes, while ‘CHAR’ means the size is in characters, which is important for multi-byte character sets. `NVARCHAR2(size)`² is similar to `VARCHAR2` but stores strings using a national character set, typically Unicode, making it ideal for multilingual data (e.g., storing names like ‘René Müller’, text in Japanese like ‘こんにちは’, or Chinese ‘你好’). A key Oracle peculiarity: An empty string (‘ ’) inserted into a `VARCHAR2` or `NVARCHAR2` column is treated as NULL!

PostgreSQL Parallel

In PostgreSQL, you’d use `VARCHAR(n)` or `TEXT`. PG’s `VARCHAR(n)` is analogous, and `TEXT` for longer strings. PostgreSQL treats an empty string (‘ ’) as distinct from NULL. This difference is a common “gotcha” when migrating. Also, PostgreSQL database encoding (e.g., UTF-8) typically handles multilingual data within standard `VARCHAR/TEXT` types, whereas Oracle has the explicit `NVARCHAR2` for this using a potentially different character set than the database default.

Relations: How They Play with Others (in Oracle) `VARCHAR2` and `NVARCHAR2` are fundamental. They are used in table definitions, PL/SQL variable declarations, func-

¹Refer to Oracle SQL Language Reference for `VARCHAR2` for complete details on size limits (e.g., 4000 bytes or 32767 bytes with `MAX_STRING_SIZE=EXTENDED`) and semantics.

²See Oracle SQL Language Reference for `NVARCHAR2`. Uses national character set (often AL16UTF16 or UTF8).

tion parameters, and return types. They often interact with string manipulation functions (covered later) and conditional expressions like CASE or DECODE. Since this is the first Oracle-specific content module, there are no *preceding* Oracle concepts from this transitional course to relate them to, beyond the general SQL context you already know.

How to Use Them: Structures & Syntax (in Oracle) Primarily used in CREATE TABLE statements and PL/SQL declarations.

```
1 CREATE TABLE EmployeeProfiles (  
2     employeeId NUMBER PRIMARY KEY,  
3     employeeName VARCHAR2(100 CHAR), -- 100 characters, good for multi-byte  
4     departmentCode VARCHAR2(10 BYTE), -- Max 10 bytes  
5     employeeBio NVARCHAR2(500) -- For multilingual bio, e.g., N' スティーブン (Steven)' or N' 李 '  
6 );  
7  
8 -- In PL/SQL (covered in detail later)  
9 DECLARE  
10     vFirstName VARCHAR2(50);  
11     vNotes NVARCHAR2(1000);  
12 BEGIN  
13     vFirstName := 'Oracleia';  
14     vNotes := N' こんにちは世界 - Hello World - Olá Mundo; -- N prefix for NVARCHAR2 literal  
15     -- ...  
16 END;  
17 /
```

Listing 1: Defining VARCHAR2 and NVARCHAR2 columns

Why Use Them? (Advantages in Oracle) VARCHAR2 is the workhorse for character data in Oracle. Using 'CHAR' semantics (VARCHAR2(n CHAR)) is advantageous for multi-byte character sets ensuring you can store 'n' characters regardless of their byte size. NVARCHAR2 provides robust support for internationalization by using a dedicated national character set, often Unicode. "For strings of text, both near and far, VARCHAR2 is Oracle's shining star. But if worldly scripts (e.g., '你好', 'こんにちは') you embrace, NVARCHAR2 wins the character race."

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Empty String is NULL:** This is a major difference from PostgreSQL and the SQL standard. '' IS NULL is true in Oracle. This can break logic migrated from systems where empty strings are distinct from NULL.
- **BYTE vs. CHAR Semantics:** If you define VARCHAR2(10 BYTE) in a UTF-8 database, you might only store 3 characters if they are 3 bytes each (e.g., some Chinese characters like 全球配件). Using VARCHAR2(10 CHAR) is generally safer for new development with multi-byte character sets. Check NLS_LENGTH_SEMANTICS instance/session parameter.
- **VARCHAR vs. VARCHAR2:** Oracle has a VARCHAR type, but its use is discouraged. VARCHAR2 is the standard. VARCHAR might have its behavior changed by Oracle in future releases to align with ANSI standards (where empty string is not NULL), so stick to VARCHAR2.

2.2. Numeric Data Type: NUMBER

What Are They? (Meanings & Values in Oracle) NUMBER(p,s)³ is Oracle's highly versatile numeric type. It stores positive, negative, fixed, and floating-point numbers, including zero.

³Consult Oracle SQL Language Reference for NUMBER. 'p' is precision (total digits), 's' is scale (digits after decimal).

- **p (precision):** Total number of digits (1 to 38).
- **s (scale):** Number of digits to the right of the decimal point (-84 to 127).

If 's' is omitted, it defaults to 0 (integer). If 'p' and 's' are omitted (NUMBER), it stores numbers as provided, up to Oracle's maximum precision and scale. Examples:

- **NUMBER(5):** Integer up to 99999.
- **NUMBER(7,2):** Number like 12345.67.
- **NUMBER(3,-2):** Number rounded to nearest 100 (e.g., 123 becomes 100).
- **NUMBER:** Floating point with max precision.

PostgreSQL Parallel

PostgreSQL has a richer set of distinct numeric types: SMALLINT, INTEGER, BIGINT for integers; NUMERIC(p,s) or DECIMAL(p,s) for exact fixed-point numbers; and REAL, DOUBLE PRECISION for floating-point. Oracle's NUMBER type aims to cover all these use cases with one type, configured by precision and scale. So, NUMBER(10,0) or NUMBER(10) is like PG's INTEGER or BIGINT (depending on range), and NUMBER(p,s) is like PG's NUMERIC(p,s).

Relations: How They Play with Others (in Oracle) NUMBER is fundamental for storing identifiers, counts, measurements, monetary values, etc. It's used in arithmetic operations, aggregate functions (like SUM, AVG - to be covered), and often as the data type for pseudo-columns like ROWNUM.

How to Use Them: Structures & Syntax (in Oracle)

```

1 CREATE TABLE ProductInventory (
2     productId NUMBER(10) PRIMARY KEY, -- Like an INTEGER or BIGINT
3     productName VARCHAR2(100),
4     quantityOnHand NUMBER(5),          -- Integer up to 99999
5     unitPrice NUMBER(8,2),             -- e.g., 123456.78
6     discountRate NUMBER(3,2)          -- e.g., 0.15 (15%)
7 );
8
9 -- In PL/SQL
10 DECLARE
11     vCounter NUMBER := 0;
12     vTotalAmount NUMBER(12,2);
13 BEGIN
14     vTotalAmount := 199.99;
15     -- ...
16 END;
17 /

```

Listing 2: Defining NUMBER columns

Why Use Them? (Advantages in Oracle) The main advantage is versatility. A single NUMBER type can store various kinds of numeric data, simplifying type choices compared to PostgreSQL's more specialized set. "NUMBER' is quite the flexible chap, for integers, decimals, it fills the gap. PostgreSQL has many, a numeric zoo, Oracle's 'NUMBER' says 'one will do!'" It also handles very large or very small numbers with high precision if needed.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Unspecified NUMBER:** Using just NUMBER without precision/scale provides maximum flexibility but can lead to storing numbers with widely varying scales in the same column. This might be problematic for applications expecting consistent formatting or for performance in some cases. For financial data, always specify precision and scale.
- **Rounding:** Oracle rounds values if they exceed the specified scale. E.g., inserting 12.348 into NUMBER(4,2) stores 12.35.
- **Portability/Understanding:** While versatile, developers from systems with distinct integer/float types might need adjustment. The same 'NUMBER' type can behave like an integer or a float depending on its definition.

2.3. Date and Time Data Types: DATE and TIMESTAMP Variations

2.3.1. DATE

What Is It? (Meanings & Values in Oracle) Oracle's DATE⁴ type stores date and time information. It includes year, month, day, hour, minute, and second components. It does not store fractional seconds or time zone information.

PostgreSQL Parallel

This is a critical difference! In PostgreSQL, DATE stores *only* the date (year, month, day). The PostgreSQL equivalent of Oracle's DATE is actually `TIMESTAMP WITHOUT TIME ZONE` or `TIMESTAMP(0) WITHOUT TIME ZONE` (since Oracle DATE has no fractional seconds). This distinction can lead to subtle bugs if not handled carefully during migration or when writing cross-database queries.

Relations: How They Play with Others (in Oracle) DATE is used for storing event timestamps, birth dates (though time part might be 00:00:00), etc. It interacts heavily with date functions (SYSDATE, TO_DATE, TO_CHAR, MONTHS_BETWEEN, etc. - many covered in the next module) and date arithmetic.

How to Use It: Structures & Syntax (in Oracle)

```

1 CREATE TABLE EventSchedule (
2     eventId NUMBER PRIMARY KEY,
3     eventName VARCHAR2(200),
4     eventStartDateTime DATE,
5     eventEndDateTime DATE
6 );
7
8 INSERT INTO EventSchedule (eventId, eventName, eventStartDateTime, eventEndDateTime)
9 VALUES (1, 'Team Meeting', TO_DATE('2024-07-15 10:00:00', 'YYYY-MM-DD HH24:MI:SS'), SYSDATE + 1/24);
10      -- SYSDATE plus 1 hour
11
12 SELECT eventName, TO_CHAR(eventStartDateTime, 'DD-Mon-YYYY HH:MI AM') AS formattedStart
13 FROM EventSchedule;
```

Listing 3: Defining and using DATE columns

Note: 'SYSDATE' is a function returning the current database server date and time as a 'DATE'. Adding '1' to a 'DATE' adds one day. '1/24' adds an hour.

Why Use It? (Advantages in Oracle) It's the traditional Oracle way to store date and time. Simpler than TIMESTAMP if fractional seconds and time zones are not needed.

⁴See Oracle SQL Language Reference for DATE. It always stores year, month, day, hour, minute, and second.

It's widely used in legacy Oracle systems. "Oracle's 'DATE', a tricky friend, holds date *and* time until the end."

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Time Component is Always Present:** Even if you insert only a date part (e.g., via `TO_DATE('2024-07-15', 'YYYY-MM-DD')`), the time component defaults to midnight (00:00:00). This means 'WHERE myDateColumn = TO_DATE('2024-07-15', 'YYYY-MM-DD')' might not behave as expected if 'myDateColumn' has a non-midnight time. Use `TRUNC(myDateColumn) = TO_DATE('2024-07-15', 'YYYY-MM-DD')` for date-only comparisons.
- **No Fractional Seconds/Time Zone:** If these are needed, use a `TIMESTAMP` type.
- **PG Confusion:** "PG's 'DATE' is simpler, it's true, so 'TRUNC' your 'DATE's, or you'll be blue!"

2.3.2. `TIMESTAMP`, `TIMESTAMP WITH TIME ZONE`, `TIMESTAMP WITH LOCAL TIME ZONE`

What Are They? (Meanings & Values in Oracle) These types extend `DATE` functionality:

- `TIMESTAMP(fractional_seconds_precision)`⁵: Stores year, month, day, hour, minute, second, and fractional seconds. Does NOT store time zone information. The precision (0-9, default 6) determines the number of digits in the fractional part of the seconds.
- `TIMESTAMP(p) WITH TIME ZONE`: Stores all that `TIMESTAMP` does, PLUS time zone information (either a region name or a UTC offset). The stored value includes the original time zone.
- `TIMESTAMP(p) WITH LOCAL TIME ZONE`: This is a clever one! When data is inserted, it's converted from the client's session time zone to the database's time zone and stored. When queried, it's converted from the database time zone back to the client's session time zone for display. The actual time zone offset is not stored with each value, but normalization occurs.

PostgreSQL Parallel

PostgreSQL's `TIMESTAMP WITHOUT TIME ZONE` (or just `TIMESTAMP`) is analogous to Oracle's `TIMESTAMP`. PostgreSQL's `TIMESTAMP WITH TIME ZONE` (or `TIMESTAMPTZ`) is analogous to Oracle's `TIMESTAMP WITH TIME ZONE` in concept. However, PG's `TIMESTAMPTZ` always normalizes input to UTC for storage and converts to the session's time zone on output. Oracle's `TIMESTAMP WITH TIME ZONE` stores the original explicit time zone. Oracle's `TIMESTAMP WITH LOCAL TIME ZONE` behaves more like PG's `TIMESTAMPTZ` in terms of normalization and session-dependent display, but the normalization target is the database time zone, not necessarily UTC.

Relations: How They Play with Others (in Oracle) These are used when high precision or time zone awareness is critical. They interact with specific timestamp functions (`SYSTIMESTAMP`, `CURRENT_TIMESTAMP`, `FROM_TZ`, `TO_TIMESTAMP_TZ`, etc. - covered in next module).

⁵See Oracle SQL Language Reference for `TIMESTAMP` family. Precision 0-9, default 6.

How to Use Them: Structures & Syntax (in Oracle)

```
1 CREATE TABLE AppLogs (  
2     logId NUMBER GENERATED ALWAYS AS IDENTITY,  
3     eventOccurred TIMESTAMP(3), -- Precision of 3 fractional seconds  
4     serverReportedTime TIMESTAMP(6) WITH TIME ZONE,  
5     userLocalEntryTime TIMESTAMP WITH LOCAL TIME ZONE  
6 );  
7  
8 -- Setting session time zone for TSLTZ demonstration  
9 ALTER SESSION SET TIME_ZONE = '-05:00'; -- e.g., EST  
10  
11 INSERT INTO AppLogs (eventOccurred, serverReportedTime, userLocalEntryTime)  
12 VALUES (  
13     TIMESTAMP '2024-07-15 10:30:45.123',  
14     TIMESTAMP '2024-07-15 15:30:45.987654 UTC',  
15     SYSTIMESTAMP -- Will be stored based on DB timezone, displayed based on session  
16 );  
17  
18 SELECT logId, eventOccurred, serverReportedTime, userLocalEntryTime FROM AppLogs;  
19 -- If you then: ALTER SESSION SET TIME_ZONE = 'Europe/London';  
20 -- And re-run the SELECT, userLocalEntryTime will display differently.
```

Listing 4: Defining and using TIMESTAMP columns

Why Use Them? (Advantages in Oracle) `TIMESTAMP` offers higher precision than `DATE`. `TIMESTAMP WITH TIME ZONE` is vital for applications dealing with events across multiple explicit time zones where preserving the original zone is key. `TIMESTAMP WITH LOCAL TIME ZONE` is excellent for global applications where users in different time zones need to see times relative to their local context, with data stored consistently.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Complexity:** More complex than `DATE`. Choose the simplest type that meets requirements.
- **`TIMESTAMP WITH LOCAL TIME ZONE (TSLTZ) Behavior:`** Understanding depends on `SESSIONTIMEZONE` and `DBTIMEZONE`. Data is normalized to `DBTIMEZONE` on insert and converted to `SESSIONTIMEZONE` on select. Ensure these are managed correctly.
- **Storage:** `TIMESTAMP WITH TIME ZONE` takes more space (13 bytes) than `TIMESTAMP` (11 bytes) or `DATE` (7 bytes). `TSLTZ` is usually 11 bytes (like `TIMESTAMP`).

2.4. Large Object Types: CLOB and BLOB (Brief Introduction)

What Are They? (Meanings & Values in Oracle) `CLOB` (Character Large Object)⁶ stores large blocks of character data (e.g., XML documents, long text descriptions). Uses database character set. `NCLOB` is the national character set version (e.g., for large documents in Japanese: ' 大量の日本語テキスト'). `BLOB` (Binary Large Object) stores large blocks of unstructured binary data (e.g., images, audio, video files).

PostgreSQL Parallel

PostgreSQL's `TEXT` type can handle very large character strings, similar to `CLOB`. For binary data, PostgreSQL has `BYTEA`. PostgreSQL also has a "Large Object" facility (using `OIDs`) which is a different mechanism for storing very large data, often in separate storage. Oracle's `LOBs` are more integrated as column types.

⁶Oracle SQL Language Reference details `CLOBs` and `BLOBs`. They are stored out-of-line or in-line depending on size and table setup.

Relations: How They Play with Others (in Oracle) LOBs are often manipulated using the DBMS_LOB package (covered in a later PL/SQL module). They are typically used for specific columns designed to hold large amounts of data. Standard string functions have limited direct applicability.

How to Use Them: Structures & Syntax (in Oracle)

```
1 CREATE TABLE DocumentArchive (  
2     docId NUMBER PRIMARY KEY,  
3     docTitle VARCHAR2(255),  
4     docContent CLOB,  
5     docScan BLOB,  
6     docNotes NCLOB -- For notes in national character set, e.g. large Chinese texts N' 全球配件 (Global  
7 );
```

Listing 5: Defining LOB columns

Further manipulation (inserting large data, reading parts) typically involves PL/SQL and DBMS_LOB. This is a foundational introduction; deeper LOB handling comes later.

Why Use Them? (Advantages in Oracle) Designed to efficiently store and manage very large data objects that don't fit well into standard VARCHAR2 or RAW types. Oracle provides robust mechanisms (like DBMS_LOB) for manipulating them.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Performance:** Improper use or very frequent access to large LOBs can impact performance.
- **Manipulation Complexity:** Requires specific LOB APIs (DBMS_LOB) for many operations beyond simple insert/select of small LOBs.
- **Storage Considerations:** LOBs can be stored in-row (if small enough) or out-of-row. This affects performance and storage management.

3. DUAL Table: Oracle's Little Helper

What Is It? (Meanings & Values in Oracle) DUAL⁷ is a special single-row, single-column table in Oracle. The column is named DUMMY and its data type is VARCHAR2(1). Its single row contains the value 'X'. Its primary purpose is to provide a FROM clause for SELECT statements where you want to compute an expression, call a function, or get a pseudo-column like SYSDATE without querying an actual user table.

PostgreSQL Parallel

PostgreSQL (and standard SQL) does not require a FROM clause for such operations. You can simply write `SELECT 1+1;` or `SELECT CURRENT_DATE;`. Oracle's SQL parser strictly requires a FROM clause for every SELECT statement, and DUAL serves this purpose. This is a common point of adjustment for PG users.

Relations: How They Play with Others (in Oracle) DUAL is often used with:

- Functions that don't require table data (e.g., SYSDATE, sequence NEXTVAL, USER).
- Simple calculations or literal selections.

⁷Oracle Database Concepts guide often explains DUAL. It's a public table owned by SYS but accessible to all users.

- In PL/SQL, sometimes for `SELECT INTO` a variable if the value comes from a function/expression.

It relates to data types because the expressions selected from `DUAL` will have a data type (e.g., `SELECT SYSDATE FROM DUAL;` returns a `DATE`).

How to Use It: Structures & Syntax (in Oracle)

```

1  -- Get current date and time
2  SELECT SYSDATE FROM DUAL;
3
4  -- Perform a calculation
5  SELECT (10 * 5) + 3 AS result FROM DUAL;
6
7  -- Get the current user
8  SELECT USER FROM DUAL;
9
10 -- (Assuming a sequence mySequence exists)
11 -- Get the next value from a sequence
12 -- CREATE SEQUENCE mySequence START WITH 1 INCREMENT BY 1;
13 -- SELECT mySequence.NEXTVAL FROM DUAL;
14
15 -- Select a literal value
16 SELECT 'Hello Oracle' AS greeting FROM DUAL;
17 SELECT N' こんにちは世界' AS japanese_greeting FROM DUAL; -- NVARCHAR2 literal

```

Listing 6: Common uses of `DUAL` table

Why Use It? (Advantages in Oracle) It's the Oracle-idiomatic way to satisfy the syntactical requirement of a `FROM` clause when you're not actually querying a persistent table. "'`DUAL`' is a table, tiny and neat, for 'SELECT's without tables, can't be beat. PG just 'SELECTS', no 'FROM' in sight, but Oracle's '`DUAL`' makes syntax right!" It's always available and guaranteed to have one row, making it reliable.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Overuse in PL/SQL:** In PL/SQL, you often don't need `SELECT ... FROM DUAL` to assign expression results to variables. Direct assignment is usually more efficient:

```

1  -- Less efficient in PL/SQL
2  DECLARE
3      vToday DATE;
4  BEGIN
5      SELECT SYSDATE INTO vToday FROM DUAL;
6  END;
7  /
8
9  -- More efficient in PL/SQL
10 DECLARE
11     vToday DATE;
12 BEGIN
13     vToday := SYSDATE; -- Preferred
14 END;
15 /
16

```

- **Misunderstanding its purpose:** It's purely a syntactic construct. The actual work (e.g., calculation, function call) is done by Oracle; `DUAL` is just the vehicle.

4. NULL Handling: Embracing the Void (Oracle Style)

`NULL` represents an unknown, missing, or inapplicable value. Handling `NULL`s correctly is vital. Oracle provides standard SQL functions and some of its own.

4.1. NVL(expr1, expr2)

What Is It? (Meanings & Values in Oracle) NVL⁸ is an Oracle-specific function. If expr1 is NULL, NVL returns expr2. If expr1 is not NULL, NVL returns expr1. Both expressions must be of the same data type, or expr2 must be implicitly convertible to the data type of expr1.

PostgreSQL Parallel

This is directly analogous to PostgreSQL's COALESCE(expr1, expr2) when COALESCE is used with only two arguments. Many PG users will naturally reach for COALESCE, which also works in Oracle. NVL is just Oracle's traditional, shorter version for this specific two-argument case.

Relations: How They Play with Others (in Oracle) NVL is used with any data type that can be NULL (VARCHAR2, NUMBER, DATE, etc.). It's commonly used in SELECT lists to provide default values for display, or in WHERE clauses or calculations to avoid issues with NULLs.

How to Use It: Structures & Syntax (in Oracle)

```
1 CREATE TABLE ProductDiscounts (  
2     productName VARCHAR2(100),  
3     discountPercentage NUMBER(4,2) -- Can be NULL  
4 );  
5 INSERT INTO ProductDiscounts VALUES ('Laptop', 0.10);  
6 INSERT INTO ProductDiscounts VALUES ('Mouse', NULL);  
7  
8 -- Display 0 if discountPercentage is NULL  
9 SELECT  
10     productName,  
11     NVL(discountPercentage, 0) AS effectiveDiscount  
12 FROM ProductDiscounts;  
13 -- Output for Mouse will show 0 for effectiveDiscount
```

Listing 7: Using NVL

Why Use It? (Advantages in Oracle) It's a concise, Oracle-specific way to substitute a value for NULL. "If a NULL appears, don't despair or yell, NVL is there to save you well." It's widely used in existing Oracle code.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Data Type Coercion:** If expr1 and expr2 are different data types, Oracle attempts implicit conversion. If expr1 is, say, a NUMBER and expr2 is a string like 'N/A', NVL(numberCol, 'N/A') will try to convert 'N/A' to a NUMBER if numberCol is NULL, leading to an error. Ensure types are compatible or use explicit conversion (e.g., NVL(TO_CHAR(numberCol), 'N/A')).
- **Only Two Arguments:** Unlike COALESCE, NVL only accepts two arguments.
- **Short-Circuiting (or lack thereof):** Both arguments to NVL are evaluated before the function determines which to return. If expr2 is a complex or costly function call, it gets executed even if expr1 is not NULL. COALESCE and CASE short-circuit (evaluate expressions only as needed).

⁸Check Oracle SQL Language Reference for NVL. It's Oracle-specific.

4.2. NVL2(expr1, expr2, expr3)

What Is It? (Meanings & Values in Oracle) NVL2⁹ is another Oracle-specific function. If `expr1` is NOT NULL, NVL2 returns `expr2`. If `expr1` IS NULL, NVL2 returns `expr3`. The data type of the return value is typically determined by `expr2` if it's not NULL, and `expr3` must be convertible to that type (or vice-versa).

PostgreSQL Parallel

There's no direct single function in PostgreSQL like NVL2. The equivalent logic is achieved using a CASE expression: `CASE WHEN expr1 IS NOT NULL THEN expr2 ELSE expr3 END`.

Relations: How They Play with Others (in Oracle) Similar to NVL, NVL2 can be used with various data types and in different parts of a SQL statement. It's essentially a compact IF-THEN-ELSE based on nullity.

How to Use It: Structures & Syntax (in Oracle)

```
1 SELECT
2     employeeName, -- Assuming employeeName column from a table like EmployeeRoster
3     commissionRate, -- Assuming commissionRate column
4     NVL2(commissionRate, 'Eligible for Commission', 'Salary Only') AS commissionStatus
5 FROM EmployeeRoster; -- Assume EmployeeRoster table from exercises
6
7 -- If commissionRate is not NULL, status is 'Eligible for Commission'
8 -- If commissionRate is NULL, status is 'Salary Only'
```

Listing 8: Using NVL2

Why Use It? (Advantages in Oracle) Provides a very concise way to return one of two values based on whether an expression is NULL or not. "NVL2 is clever, with choices three." Can sometimes make queries shorter than an equivalent CASE expression.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Data Type Compatibility:** `expr2` and `expr3` must be of compatible data types. Oracle will try to implicitly convert, which can lead to errors if types are mismatched (e.g., returning a NUMBER vs. a VARCHAR2). Use explicit `TO_CHAR`, `TO_NUMBER`, etc., if needed.
- **Readability:** For complex logic, a CASE expression might be more readable than a nested NVL2.
- **Short-Circuiting:** Like NVL, all three expressions might be evaluated. If `expr2` or `expr3` are costly, this can be inefficient. Standard CASE offers better short-circuiting.

4.3. COALESCE(expr1, expr2, ..., exprN)

What Is It? (Meanings & Values in Oracle) COALESCE¹⁰ returns the first non-NULL expression in a list of expressions. It evaluates expressions from left to right and stops as soon as it finds one that is not NULL. If all expressions are NULL, it returns NULL. All expressions must be of the same data type or implicitly convertible to a common data type.

⁹Also Oracle-specific; see SQL Language Reference for NVL2.

¹⁰COALESCE is an ANSI SQL standard function, well-documented in Oracle's SQL Language Reference.

PostgreSQL Parallel

This function is identical in name, syntax, and behavior to PostgreSQL's COALESCE. You already know this one well!

Relations: How They Play with Others (in Oracle) COALESCE is extremely useful for providing fallback values from a series of potentially NULL columns or expressions. It works with all data types.

How to Use It: Structures & Syntax (in Oracle)

```
1 CREATE TABLE ContactInfo (  
2     personId NUMBER,  
3     homePhone VARCHAR2(20),  
4     mobilePhone VARCHAR2(20),  
5     workPhone VARCHAR2(20),  
6     email VARCHAR2(100)  
7 );  
8 INSERT INTO ContactInfo VALUES (1, NULL, '555-1234', NULL, 'user@example.com');  
9 INSERT INTO ContactInfo VALUES (2, '555-8888', NULL, '555-9999', NULL);  
10 INSERT INTO ContactInfo VALUES (3, NULL, NULL, NULL, 'another@example.com');  
11 INSERT INTO ContactInfo VALUES (4, NULL, NULL, NULL, NULL);  
12  
13  
14 -- Get the first available phone number, or email if no phone  
15 SELECT  
16     personId,  
17     COALESCE(homePhone, mobilePhone, workPhone, email, 'No Contact Info') AS preferredContact  
18 FROM ContactInfo;
```

Listing 9: Using COALESCE

Why Use It? (Advantages in Oracle) It's the ANSI SQL standard way to find the first non-NULL value, making queries more portable. It handles multiple fallback options elegantly. It also short-circuits: expressions are evaluated only until a non-NULL one is found, which can be a performance benefit if later expressions are costly. "But COALESCE is standard, for all to agree."

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Data Type Promotion:** The data type of the result is determined by the data types of the arguments, following Oracle's type precedence rules. Ensure all expressions are compatible to avoid unexpected conversions or errors. For example, COALESCE(numberColumn, dateColumn) would likely cause an error if numberColumn is NULL and dateColumn is not.

5. Conditional Expressions: Making Choices in SQL

Conditional expressions allow you to implement if-then-else logic directly within SQL statements.

5.1. DECODE(expr, search1, result1, [search2, result2, ...], [default])

What Is It? (Meanings & Values in Oracle) DECODE¹¹ compares expr to each search value in sequence. If expr equals a search value, DECODE returns the corresponding result. If no match is found, it returns default. If default is omitted and no match is

¹¹DECODE is an Oracle-specific function. See SQL Language Reference.

found, it returns NULL. A unique feature: DECODE treats two NULLs as equal. So, DECODE(col, NULL, 'Is Null', 'Not Null') will return 'Is Null' if 'col' is NULL.

PostgreSQL Parallel

PostgreSQL does not have DECODE. The equivalent functionality is achieved using a "simple CASE" expression: CASE expr WHEN search1 THEN result1 [WHEN search2 THEN result2 ...] [ELSE default] END. The NULL-equals-NULL behavior of DECODE is different from standard SQL comparisons where NULL = NULL is UNKNOWN.

Relations: How They Play with Others (in Oracle) DECODE can be used with various data types for expr, search, and result (with appropriate type compatibility). It's often found in older Oracle code. It can be used in SELECT lists, WHERE clauses, and ORDER BY clauses.

How to Use It: Structures & Syntax (in Oracle)

```
1 CREATE TABLE OrderStatus (  
2     orderId NUMBER,  
3     statusCode VARCHAR2(1) -- 'P': Pending, 'S': Shipped, 'D': Delivered  
4 );  
5 INSERT INTO OrderStatus VALUES (101, 'P');  
6 INSERT INTO OrderStatus VALUES (102, 'S');  
7 INSERT INTO OrderStatus VALUES (103, 'X'); -- Unknown status  
8  
9 SELECT  
10     orderId,  
11     statusCode,  
12     DECODE(statusCode,  
13         'P', 'Pending Shipment',  
14         'S', 'Shipped to Customer',  
15         'D', 'Delivered Successfully',  
16         'Unknown Status or Error') AS fullStatus  
17 FROM OrderStatus;
```

Listing 10: Using DECODE

Why Use It? (Advantages in Oracle) It's compact for simple equality checks. The NULL-equals-NULL comparison can be convenient in specific scenarios, avoiding explicit IS NULL checks. "Old DECODE compares, with a quirky view, where NULL equals NULL, it's strange but true!"

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Readability:** For many conditions or complex logic, DECODE becomes hard to read compared to CASE.
- **Flexibility:** Only performs equality checks. Cannot do range checks (e.g., 'value > 10') or use 'OR'/'AND' within a search condition directly. CASE is far more flexible.
- **Type Conversion:** All search expressions should be convertible to the type of the first expr, and all result expressions should be convertible to the type of the first result.
- **NULL Handling Surprise:** While sometimes useful, its NULL-equals-NULL behavior can be a pitfall if you expect standard SQL NULL logic where NULL is not equal to anything.
- **Oracle Specific:** Not portable to other databases. CASE is the SQL standard.

5.2. CASE Expression

What Is It? (Meanings & Values in Oracle) The CASE¹² expression is the SQL standard way to implement conditional logic. Oracle supports two forms:

1. **Simple CASE:** CASE expr WHEN comparison_expr1 THEN return_expr1 [WHEN comparison_expr2 THEN return_expr2 ...] [ELSE else_expr] END This compares expr for equality against each comparison_expr.
2. **Searched CASE:** CASE WHEN condition1 THEN return_expr1 [WHEN condition2 THEN return_expr2 ...] [ELSE else_expr] END This evaluates each boolean condition. The first one that is true determines the return_expr.

If no condition is met and no ELSE clause is specified, CASE returns NULL.

PostgreSQL Parallel

CASE expressions are identical in syntax and behavior to those in PostgreSQL. This is familiar territory!

Relations: How They Play with Others (in Oracle) CASE is extremely versatile. It can be used with any data types (ensuring type compatibility for return expressions) and in SELECT lists, WHERE clauses, ORDER BY clauses, GROUP BY clauses, and within PL/SQL.

How to Use It: Structures & Syntax (in Oracle)

```
1 CREATE TABLE EmployeePerformance (
2     employeeId NUMBER,
3     rating NUMBER(2,1), -- e.g., 1.0 to 5.0
4     salesAmount NUMBER(10,2)
5 );
6 INSERT INTO EmployeePerformance VALUES (1, 4.5, 50000);
7 INSERT INTO EmployeePerformance VALUES (2, 2.5, 15000);
8 INSERT INTO EmployeePerformance VALUES (3, 3.8, NULL); -- No sales data
9
10 -- Simple CASE (less common for ranges)
11 SELECT employeeId,
12        CASE rating
13            WHEN 5.0 THEN 'Outstanding'
14            WHEN 4.0 THEN 'Exceeds Expectations'
15            ELSE 'Meets or Needs Improvement'
16        END AS ratingTextSimple
17 FROM EmployeePerformance;
18
19 -- Searched CASE (more flexible)
20 SELECT
21     employeeId,
22     rating,
23     salesAmount,
24     CASE
25         WHEN rating >= 4.5 THEN 'Top Performer'
26         WHEN rating >= 3.5 AND rating < 4.5 THEN 'High Value'
27         WHEN rating >= 2.5 AND rating < 3.5 THEN 'Solid Contributor'
28         ELSE 'Needs Improvement'
29     END AS performanceCategory,
30     CASE
31         WHEN salesAmount > 40000 THEN 'High Sales Bonus'
32         WHEN salesAmount BETWEEN 20000 AND 40000 THEN 'Standard Sales Bonus'
33         WHEN salesAmount < 20000 THEN 'Low Sales Warning'
34         WHEN salesAmount IS NULL THEN 'Sales Data Missing' -- Explicit NULL check
35         ELSE 'No Bonus'
36     END AS salesBonusStatus
37 FROM EmployeePerformance;
```

Listing 11: Using CASE expressions

¹²CASE expressions are ANSI SQL standard and detailed in Oracle's SQL Language Reference.

Why Use It? (Advantages in Oracle) CASE is the SQL standard, making queries more portable. It's highly flexible, especially the searched CASE, allowing complex conditions. It's generally more readable than complex DECODE statements. "CASE is the modern, clearer way, to make decisions in SQL each day." It also short-circuits, evaluating conditions and returning as soon as one is true.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Data Type Compatibility of Results:** All THEN expressions (and the ELSE expression) must be of the same data type or implicitly convertible to a common type. If you mix, say, NUMBER and VARCHAR2 results, you'll get an error.
- **Order of WHEN clauses matters** in searched CASE, as the first true condition's result is returned. This can be a pitfall if logic isn't structured carefully for overlapping conditions.
- **Forgetting ELSE:** If no conditions are met and no ELSE is provided, the result is NULL. This might be intended, or it might be an oversight leading to unexpected NULLs.

6. ROWNUM Pseudo-column: Numbering Your Rows

What Is It? (Meanings & Values in Oracle) ROWNUM¹³ is a pseudo-column (not a real column stored in the table) that assigns a sequential number (1, 2, 3, ...) to each row returned by a query. Crucially, ROWNUM values are assigned to rows *after* they pass the WHERE clause of that query block but *before* any ORDER BY or aggregations in that same query block are processed.

PostgreSQL Parallel

PostgreSQL uses LIMIT and OFFSET clauses for restricting the number of rows returned, which are applied *after* ORDER BY. For row numbering similar to analytical functions, PostgreSQL uses ROW_NUMBER() OVER (...). Oracle's ROWNUM is very different in its evaluation timing and usage for Top-N queries. This is a major conceptual shift.

Relations: How They Play with Others (in Oracle) ROWNUM is often used with subqueries to correctly implement Top-N queries (e.g., get the top 10 highest paid employees). It interacts with the WHERE clause directly. Its value is a NUMBER.

How to Use It: Structures & Syntax (in Oracle)

```
1 CREATE TABLE ProductSales (  
2     saleId NUMBER PRIMARY KEY,  
3     productSold VARCHAR2(50),  
4     saleAmount NUMBER(10,2),  
5     saleDate DATE  
6 );  
7 INSERT INTO ProductSales VALUES (1, 'Product A', 100.50, TO_DATE('2023-01-15', 'YYYY-MM-DD'));  
8 INSERT INTO ProductSales VALUES (2, 'Product B', 250.00, TO_DATE('2023-01-10', 'YYYY-MM-DD'));  
9 INSERT INTO ProductSales VALUES (3, 'Product C', 75.25, TO_DATE('2023-02-01', 'YYYY-MM-DD'));  
10 INSERT INTO ProductSales VALUES (4, 'Product A', 100.50, TO_DATE('2023-02-05', 'YYYY-MM-DD'));  
11 INSERT INTO ProductSales VALUES (5, 'Product D', 500.00, TO_DATE('2023-02-10', 'YYYY-MM-DD'));  
12  
13 -- Incorrect Top-N (ROWNUM applied before ORDER BY)  
14 -- This gets *any* 3 rows then orders them. Not what's usually wanted.
```

¹³ROWNUM is an Oracle pseudo-column. Its behavior is critical to understand; see Oracle SQL Language Reference.

```

15 SELECT productSold, saleAmount, ROWNUM
16 FROM ProductSales
17 WHERE ROWNUM <= 3
18 ORDER BY saleAmount DESC;
19
20 -- Correct Top-N: Get 3 highest sales
21 SELECT productSold, saleAmount
22 FROM (
23     SELECT productSold, saleAmount
24     FROM ProductSales
25     ORDER BY saleAmount DESC
26 )
27 WHERE ROWNUM <= 3; -- ROWNUM applied to already sorted result set
28
29 -- Pagination: Get rows 3 and 4 (e.g., 2nd page of 2 rows)
30 SELECT productSold, saleAmount, rn
31 FROM (
32     SELECT productSold, saleAmount, ROWNUM AS rn
33     FROM (
34         SELECT productSold, saleAmount
35         FROM ProductSales
36         ORDER BY saleAmount DESC
37     ) orderedSales
38 )
39 WHERE rn BETWEEN 3 AND 4;

```

Listing 12: Using ROWNUM for Top-N and pagination

The key is that for ROWNUM to work with ordering, the ORDER BY must be in a subquery, and ROWNUM applied in an outer query.

Why Use It? (Advantages in Oracle) It's Oracle's traditional method for limiting rows and implementing pagination. When understood correctly, it's effective. "ROWNUM counts rows as they fly on by, before the 'ORDER' can make them lie. For Top-N true, a subquery's the art, or 'ROWNUM' will play tricks on your heart."

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Assignment Timing:** This is the biggest pitfall. ROWNUM is assigned *before* ORDER BY in the same query block. So, 'SELECT ..., ROWNUM FROM myTable WHERE ROWNUM <= 5 ORDER BY someColumn;' gives you 5 arbitrary rows, then sorts them. Not the top 5 based on 'someColumn'.
- **Cannot Use ROWNUM > N Directly:** A condition like WHERE ROWNUM > 1 will never return rows. Why? For a row to get ROWNUM = 2, it must first have been assigned ROWNUM = 1. But if the WHERE clause filters out ROWNUM = 1, no row can ever become ROWNUM = 2. This is why pagination requires nesting, aliasing ROWNUM in a subquery, then filtering on that alias.
- **Non-Deterministic Without ORDER BY:** If used without an ORDER BY in a subquery, the rows selected by ROWNUM <= N can be arbitrary and may change between executions if the underlying data access path changes.
- **Analytic Functions are Often Better:** For complex ranking or row numbering within groups, Oracle's analytic functions (like ROW_NUMBER() OVER (ORDER BY ...), covered later) are far more powerful and often clearer than complex ROWNUM tricks.

7. Comments: Annotating Your Code

What Are They? (Meanings & Values in Oracle) Comments are non-executable text within your SQL code used for documentation, explaining logic, or temporarily disabling parts of a statement. Oracle supports two types:

1. **Single-line comments:** Start with two hyphens (--) and extend to the end of the line.
2. **Multi-line (or block) comments:** Start with /* and end with */. Can span multiple lines or be used inline.

PostgreSQL Parallel

These comment styles are identical to those used in PostgreSQL and many other programming languages. This is one area with perfect harmony!

Relations: How They Play with Others (in Oracle) Comments can be placed almost anywhere in SQL or PL/SQL code where whitespace is allowed, without affecting execution. They are purely for human readability and maintainability.

How to Use It: Structures & Syntax (in Oracle)

```
1 -- This is a single-line comment: Selecting employee details
2 SELECT
3     employeeId,
4     firstName, /* This is an inline multi-line comment */
5     lastName,
6     salary -- Another single-line comment for salary
7 FROM EmployeeRoster -- Assuming EmployeeRoster table
8 WHERE departmentName = 'IT' -- Filter for IT department
9 /*
10  This is a multi-line comment block.
11  It can span several lines and is useful for
12  more detailed explanations or for commenting out
13  larger sections of code temporarily.
14  AND hireDate > TO_DATE('2020-01-01', 'YYYY-MM-DD')
15  */
16 ORDER BY lastName;
```

Listing 13: Using comments in Oracle SQL

Why Use Them? (Advantages in Oracle) Essential for code readability, Maintainability, and collaboration. Good comments explain the “why” behind complex logic, not just the “what.” They help future you (and others!) understand the code.

Watch Out! (Disadvantages & Pitfalls in Oracle)

- **Outdated Comments:** The biggest pitfall is comments that are not updated when the code changes, leading to misleading documentation. Strive to keep comments accurate.
- **Over-commenting Obvious Code:** Commenting ‘i := i + 1; – Increment i’ is usually not helpful. Focus on complex logic or business rules.
- **Nesting Multi-line Comments (Not Supported):** Most SQL dialects, including Oracle’s, do not support nesting /* ... */ comments directly. Some tools might appear to handle it, but it’s not standard. E.g., /* outer /* inner */ still outer? */ – the first */ ends the entire comment.

8. Bridging from PostgreSQL: Core Syntax & Foundational Concepts Summary

This section provides a consolidated view of the key differences encountered in this module for a PostgreSQL user transitioning to Oracle.

Key Transitional Points

• Data Types:

- VARCHAR2 vs. PG VARCHAR/TEXT: Oracle's empty string (' ') is NULL. PG distinguishes them. Oracle has BYTE/CHAR length semantics for VARCHAR2.
- NVARCHAR2: Oracle's explicit type for national character sets (often Unicode), crucial for text like 'Günther', 'Résumé', ' こんにちは', or ' 全球配件' if the main DB character set wouldn't support them or if a specific national character set (like AL16UTF16) is desired. PG typically handles this via database encoding (e.g., UTF-8) in standard VARCHAR/TEXT.
- NUMBER vs. PG's specialized numerics (INTEGER, NUMERIC, etc.): Oracle's NUMBER(p,s) is a versatile single type.
- DATE in Oracle includes time (hours, minutes, seconds). PG's DATE is date-only. Oracle DATE \approx PG TIMESTAMP(0) WITHOUT TIME ZONE.
- TIMESTAMP (Oracle) \approx PG TIMESTAMP WITHOUT TIME ZONE.
- TIMESTAMP WITH TIME ZONE (Oracle) stores the specified time zone. PG's TIMESTAMPTZ normalizes to UTC for storage. Oracle's TIMESTAMP WITH LOCAL TIME ZONE is closer to PG's TIMESTAMPTZ behavior (normalizes to DB timezone, displays in session timezone).
- CLOB/BLOB vs. PG TEXT/BYTEA: Oracle's LOBs are distinct types often with dedicated storage characteristics and manipulation via DBMS_LOB. NCLOB for large character data in national character sets (e.g., storing a book in こんにちは).

- **DUAL Table:** Oracle requires FROM DUAL for SELECT statements not querying a user table (e.g., SELECT SYSDATE FROM DUAL;). PostgreSQL allows SELECT SYSDATE; directly.

• NULL Handling:

- NVL(a,b) is Oracle-specific for COALESCE(a,b).
- NVL2(a,b,c) is Oracle-specific for CASE WHEN a IS NOT NULL THEN b ELSE c END.
- COALESCE is standard and works identically in both.

• Conditional Expressions:

- DECODE is Oracle-specific, older. It treats NULL as equal to NULL. PG uses CASE.
- CASE is standard and works identically.

- **ROWNUM Pseudo-column:** Oracle-specific for row limiting. Assigned *before* ORDER BY in the same query block. PostgreSQL uses LIMIT/OFFSET (applied *after* ORDER BY) or ROW_NUMBER() OVER (...). This difference is fundamental for Top-N queries.

- **Comments (--, /* */):** Identical in both systems.

This module has laid the groundwork. As we progress, we'll build upon these core

concepts, exploring more Oracle-specific functions and features. Remember that Oracle Live SQL or an Oracle XE (Express Edition) setup with SQL Developer are great platforms to practice these examples. Keep your PostgreSQL wisdom as a reference, but embrace the Oracle way!