

Contents

Preface

Audience	liii
Documentation Accessibility	liii
Related Documents	liv
Conventions	liv

Part I Basic Database Administration

1 Getting Started with Database Administration

1.1	Changes on Oracle Database Release 23ai for Oracle Database Administrator's Guide	1-1
1.1.1	New Features in 23ai	1-2
1.1.2	Deprecated Features	1-2
1.1.3	Desupported Features	1-2
1.2	Types of Oracle Database Users	1-3
1.2.1	Database Administrators	1-3
1.2.2	Security Officers	1-4
1.2.3	Network Administrators	1-4
1.2.4	Application Developers	1-4
1.2.5	Application Administrators	1-5
1.2.6	Database Users	1-5
1.3	Tasks of a Database Administrator	1-5
1.3.1	Task 1: Evaluate the Database Server Hardware	1-6
1.3.2	Task 2: Install the Oracle Database Software	1-6
1.3.3	Task 3: Plan the Database	1-7
1.3.4	Task 4: Create and Open the Database	1-7
1.3.5	Task 5: Back Up the Database	1-8
1.3.6	Task 6: Enroll System Users	1-8
1.3.7	Task 7: Implement the Database Design	1-8
1.3.8	Task 8: Back Up the Fully Functional Database	1-8
1.3.9	Task 9: Tune Database Performance	1-8
1.3.10	Task 10: Download and Install Release Updates and Release Update Revisions	1-9

1.3.11	Task 11: Roll Out to Additional Hosts	1-9
1.4	SQL Statements	1-10
1.4.1	Submitting Commands and SQL to the Database	1-10
1.4.2	About SQL*Plus	1-11
1.4.3	Connecting to the Database with SQL*Plus	1-11
1.4.3.1	About Connecting to the Database with SQL*Plus	1-12
1.4.3.2	Step 1: Open a Command Window	1-12
1.4.3.3	Step 2: Set Operating System Environment Variables	1-12
1.4.3.4	Step 3: Start SQL*Plus	1-13
1.4.3.5	Step 4: Submit the SQL*Plus CONNECT Command	1-13
1.5	Identifying Your Oracle Database Software Release	1-18
1.5.1	About Oracle Database Release Numbers	1-18
1.5.2	Checking Your Current Release Number	1-20
1.6	About Database Administrator Security and Privileges	1-21
1.6.1	The Database Administrator's Operating System Account	1-21
1.6.2	Administrative User Accounts	1-21
1.6.2.1	About Administrative User Accounts	1-22
1.6.2.2	SYS	1-23
1.6.2.3	SYSTEM	1-23
1.6.2.4	SYSBACKUP, SYSDG, SYSKM, and SYSRAC	1-23
1.6.2.5	The DBA Role	1-24
1.7	Database Administrator Authentication	1-25
1.7.1	Administrative Privileges	1-25
1.7.2	Operations Authorized by Administrative Privileges	1-26
1.7.3	Authentication Methods for Database Administrators	1-29
1.7.3.1	About Authentication Methods for Database Administrators	1-29
1.7.3.2	Nonsecure Remote Connections	1-30
1.7.3.3	Local Connections and Secure Remote Connections	1-31
1.7.4	Using Operating System Authentication	1-31
1.7.4.1	Operating System Groups	1-31
1.7.4.2	Preparing to Use Operating System Authentication	1-33
1.7.4.3	Connecting Using Operating System Authentication	1-33
1.7.5	Using Password File Authentication	1-34
1.7.5.1	Preparing to Use Password File Authentication	1-34
1.7.5.2	Connecting Using Password File Authentication	1-36
1.8	Creating and Maintaining a Database Password File	1-36
1.8.1	ORAPWD Syntax and Command Line Argument Descriptions	1-37
1.8.2	Creating a Database Password File with ORAPWD	1-43
1.8.3	Sharing and Disabling the Database Password File	1-45
1.8.4	Keeping Administrator Passwords Synchronized with the Data Dictionary	1-45
1.8.5	Adding Users to a Database Password File	1-47
1.8.6	Granting and Revoking Administrative Privileges	1-47

1.8.7	Viewing Database Password File Members	1-48
1.8.8	Removing a Database Password File	1-49
1.9	Data Utilities	1-49

2 Configuring Automatic Restart of an Oracle Database

2.1	About Oracle Restart	2-1
2.1.1	Oracle Restart Overview	2-2
2.1.2	About Startup Dependencies	2-3
2.1.3	About Starting and Stopping Components with Oracle Restart	2-3
2.1.4	About Starting and Stopping Oracle Restart	2-4
2.1.5	Oracle Restart Configuration	2-4
2.1.6	Oracle Restart Integration with Oracle Data Guard	2-6
2.1.7	Fast Application Notification with Oracle Restart	2-7
2.1.7.1	Overview of Fast Application Notification	2-7
2.1.7.2	Application High Availability with Services and FAN	2-8
2.2	Configuring Oracle Restart	2-11
2.2.1	About Configuring Oracle Restart	2-13
2.2.2	Preparing to Run SRVCTL	2-13
2.2.3	Obtaining Help for SRVCTL	2-14
2.2.4	Adding Components to the Oracle Restart Configuration	2-15
2.2.5	Removing Components from the Oracle Restart Configuration	2-16
2.2.6	Disabling and Enabling Oracle Restart Management for a Component	2-17
2.2.7	Viewing Component Status	2-18
2.2.8	Viewing the Oracle Restart Configuration for a Component	2-18
2.2.9	Modifying the Oracle Restart Configuration for a Component	2-19
2.2.10	Managing Environment Variables in the Oracle Restart Configuration	2-20
2.2.10.1	About Environment Variables in the Oracle Restart Configuration	2-20
2.2.10.2	Setting and Unsetting Environment Variables	2-21
2.2.10.3	Viewing Environment Variables	2-21
2.2.11	Creating and Deleting Database Services with SRVCTL	2-22
2.2.12	Enabling FAN Events in an Oracle Restart Environment	2-23
2.2.13	Automating the Failover of Connections Between Primary and Standby Databases	2-24
2.2.14	Enabling Clients for Fast Connection Failover	2-25
2.2.14.1	About Enabling Clients for Fast Connection Failover	2-25
2.2.14.2	Enabling Fast Connection Failover for JDBC Clients	2-26
2.2.14.3	Enabling Fast Connection Failover for Oracle Call Interface Clients	2-27
2.2.14.4	Enabling Fast Connection Failover for ODP.NET Clients	2-28
2.3	Starting and Stopping Components Managed by Oracle Restart	2-29
2.4	Stopping and Restarting Oracle Restart for Maintenance Operations	2-31
2.5	SRVCTL Command Reference for Oracle Restart	2-33

2.5.1	add	2-35
2.5.1.1	srvctl add asm	2-36
2.5.1.2	srvctl add database	2-37
2.5.1.3	srvctl add listener	2-39
2.5.1.4	srvctl add ons	2-40
2.5.1.5	srvctl add service	2-41
2.5.2	config	2-45
2.5.2.1	srvctl config asm	2-45
2.5.2.2	srvctl config database	2-46
2.5.2.3	srvctl config listener	2-47
2.5.2.4	srvctl config ons	2-47
2.5.2.5	srvctl config service	2-47
2.5.3	disable	2-48
2.5.3.1	srvctl disable asm	2-49
2.5.3.2	srvctl disable database	2-49
2.5.3.3	srvctl disable diskgroup	2-50
2.5.3.4	srvctl disable listener	2-50
2.5.3.5	srvctl disable ons	2-51
2.5.3.6	srvctl disable service	2-51
2.5.4	downgrade	2-52
2.5.4.1	srvctl downgrade database	2-52
2.5.5	enable	2-53
2.5.5.1	srvctl enable asm	2-53
2.5.5.2	srvctl enable database	2-54
2.5.5.3	srvctl enable diskgroup	2-54
2.5.5.4	srvctl enable listener	2-55
2.5.5.5	srvctl enable ons	2-55
2.5.5.6	srvctl enable service	2-55
2.5.6	getenv	2-56
2.5.6.1	srvctl getenv asm	2-57
2.5.6.2	srvctl getenv database	2-57
2.5.6.3	srvctl getenv listener	2-58
2.5.7	modify	2-58
2.5.7.1	srvctl modify asm	2-59
2.5.7.2	srvctl modify database	2-59
2.5.7.3	srvctl modify listener	2-60
2.5.7.4	srvctl modify ons	2-61
2.5.7.5	srvctl modify service	2-62
2.5.8	remove	2-66
2.5.8.1	srvctl remove asm	2-66
2.5.8.2	srvctl remove database	2-67
2.5.8.3	srvctl remove diskgroup	2-68

2.5.8.4	srvctl remove listener	2-68
2.5.8.5	srvctl remove ons	2-69
2.5.8.6	srvctl remove service	2-69
2.5.9	setenv	2-70
2.5.9.1	srvctl setenv asm	2-71
2.5.9.2	srvctl setenv database	2-71
2.5.9.3	srvctl setenv listener	2-72
2.5.10	start	2-73
2.5.10.1	srvctl start asm	2-73
2.5.10.2	srvctl start database	2-74
2.5.10.3	srvctl start diskgroup	2-75
2.5.10.4	srvctl start home	2-75
2.5.10.5	srvctl start listener	2-76
2.5.10.6	srvctl start ons	2-76
2.5.10.7	srvctl start service	2-77
2.5.11	status	2-78
2.5.11.1	srvctl status asm	2-78
2.5.11.2	srvctl status database	2-79
2.5.11.3	srvctl status diskgroup	2-79
2.5.11.4	srvctl status home	2-80
2.5.11.5	srvctl status listener	2-80
2.5.11.6	srvctl status ons	2-81
2.5.11.7	srvctl status service	2-81
2.5.12	stop	2-82
2.5.12.1	srvctl stop asm	2-83
2.5.12.2	srvctl stop database	2-83
2.5.12.3	srvctl stop diskgroup	2-84
2.5.12.4	srvctl stop home	2-85
2.5.12.5	srvctl stop listener	2-86
2.5.12.6	srvctl stop ons	2-86
2.5.12.7	srvctl stop service	2-87
2.5.13	unsetenv	2-88
2.5.13.1	srvctl unsetenv asm	2-89
2.5.13.2	srvctl unsetenv database	2-89
2.5.13.3	srvctl unsetenv listener	2-90
2.5.14	update	2-90
2.5.14.1	srvctl update database	2-91
2.5.15	upgrade	2-91
2.5.15.1	srvctl upgrade database	2-91
2.6	CRSCTL Command Reference	2-92
2.6.1	check	2-92
2.6.2	config	2-93

2.6.3	disable	2-93
2.6.4	enable	2-93
2.6.5	start	2-93
2.6.6	stop	2-93

3 Managing Processes

3.1	About Dedicated and Shared Server Processes	3-2
3.1.1	Dedicated Server Processes	3-2
3.1.2	Shared Server Processes	3-3
3.2	About Database Resident Connection Pooling	3-5
3.2.1	Comparing DRCP to Dedicated Server and Shared Server	3-7
3.3	About Proxy Resident Connection Pooling	3-8
3.4	Configuring Oracle Database for Shared Server	3-9
3.4.1	Initialization Parameters for Shared Server	3-10
3.4.2	Memory Management for Shared Server	3-10
3.4.3	Enabling Shared Server	3-10
3.4.3.1	About Determining a Value for SHARED_SERVERS	3-12
3.4.3.2	Decreasing the Number of Shared Server Processes	3-12
3.4.3.3	Limiting the Number of Shared Server Processes	3-12
3.4.3.4	Limiting the Number of Shared Server Sessions	3-13
3.4.3.5	Protecting Shared Memory	3-13
3.4.4	Configuring Dispatchers	3-14
3.4.4.1	DISPATCHERS Initialization Parameter Attributes	3-15
3.4.4.2	Determining the Number of Dispatchers	3-16
3.4.4.3	Setting the Initial Number of Dispatchers	3-16
3.4.4.4	Altering the Number of Dispatchers	3-17
3.4.4.5	Shutting Down Specific Dispatcher Processes	3-19
3.4.5	Disabling Shared Server	3-19
3.4.6	Shared Server Data Dictionary Views	3-19
3.5	Configuring Database Resident Connection Pooling	3-20
3.5.1	Database Resident Connection Pooling Initialization Parameters	3-21
3.5.2	Enabling Database Resident Connection Pooling	3-22
3.5.3	Configuring the Connection Pool for Database Resident Connection Pooling	3-23
3.5.3.1	Configuration Parameters for Database Resident Connection Pooling	3-24
3.5.4	Using Multi-Pool Database Resident Connection Pooling	3-26
3.5.5	Data Dictionary Views for Database Resident Connection Pooling	3-27
3.5.6	Determining the States of Connections in the Connection Pool	3-28
3.6	About Oracle Database Background Processes	3-29
3.7	Managing Prespawned Processes	3-30
3.7.1	About Managing Prespawned Processes	3-30
3.7.2	Managing Pools for Prespawned Processes	3-31

3.8	Managing Processes for Parallel SQL Execution	3-32
3.8.1	About Parallel Execution Servers	3-33
3.8.2	Altering Parallel Execution for a Session	3-33
3.8.2.1	Disabling Parallel SQL Execution	3-34
3.8.2.2	Enabling Parallel SQL Execution	3-34
3.8.2.3	Forcing Parallel SQL Execution	3-34
3.9	Managing Processes for External Procedures	3-35
3.9.1	About External Procedures	3-35
3.9.2	DBA Tasks to Enable External Procedure Calls	3-36
3.10	Terminating Sessions	3-37
3.10.1	About Terminating Sessions	3-37
3.10.2	Identifying Which Session to Terminate	3-38
3.10.3	Terminating an Active Session	3-38
3.10.4	Terminating an Inactive Session	3-39
3.10.5	Cancelling a SQL Statement in a Session	3-40
3.11	Process and Session Data Dictionary Views	3-41

4 Managing Memory

4.1	About Memory Management	4-2
4.2	Memory Architecture Overview	4-3
4.3	Using Unified Memory Management	4-5
4.4	Using Automatic Memory Management	4-5
4.4.1	About Automatic Memory Management	4-6
4.4.2	Enabling Automatic Memory Management	4-6
4.4.3	Monitoring and Tuning Automatic Memory Management	4-9
4.5	Configuring Memory Manually	4-10
4.5.1	About Manual Memory Management	4-10
4.5.2	Using Automatic Shared Memory Management	4-11
4.5.2.1	About Automatic Shared Memory Management	4-12
4.5.2.2	Components and Granules in the SGA	4-12
4.5.2.3	Setting Maximum SGA Size	4-13
4.5.2.4	Setting SGA Target Size	4-13
4.5.2.5	Enabling Automatic Shared Memory Management	4-15
4.5.2.6	Setting Minimums for Automatically Sized SGA Components	4-17
4.5.2.7	Dynamic Modification of SGA_TARGET	4-17
4.5.2.8	Modifying Parameters for Automatically Sized Components	4-18
4.5.2.9	Modifying Parameters for Manually Sized Components	4-19
4.5.3	Using Manual Shared Memory Management	4-19
4.5.3.1	About Manual Shared Memory Management	4-20
4.5.3.2	Enabling Manual Shared Memory Management	4-20
4.5.3.3	Setting the Buffer Cache Initialization Parameters	4-20

4.5.3.4	Specifying the Shared Pool Size	4-23
4.5.3.5	Specifying the Large Pool Size	4-24
4.5.3.6	Specifying the Java Pool Size	4-24
4.5.3.7	Specifying the Streams Pool Size	4-24
4.5.3.8	Specifying the Vector Pool Size	4-24
4.5.3.9	Specifying Miscellaneous SGA Initialization Parameters	4-24
4.5.4	Using Automatic PGA Memory Management	4-25
4.5.5	Using Manual PGA Memory Management	4-27
4.6	Using Force Full Database Caching Mode	4-27
4.6.1	About Force Full Database Caching Mode	4-28
4.6.2	Before Enabling Force Full Database Caching Mode	4-29
4.6.3	Enabling Force Full Database Caching Mode	4-29
4.6.4	Disabling Force Full Database Caching Mode	4-30
4.7	Configuring Database Smart Flash Cache	4-30
4.7.1	When to Configure Database Smart Flash Cache	4-31
4.7.2	Sizing Database Smart Flash Cache	4-31
4.7.3	Tuning Memory for Database Smart Flash Cache	4-31
4.7.4	Database Smart Flash Cache Initialization Parameters	4-32
4.7.5	Database Smart Flash Cache in an Oracle Real Applications Clusters Environment	4-33
4.8	Improving Query Response Time with the Server Result Cache	4-34
4.8.1	About the Server Result Cache	4-34
4.8.2	Using the Server Result Cache	4-34
4.8.3	Specifying the Result Cache Maximum Size	4-35
4.8.4	Specifying the Use of Temporary Segments for Query Results	4-36
4.9	Improving Query Performance with Oracle Database In-Memory	4-37
4.10	Enabling High Performance Data Streaming with the Memoptimized Rowstore	4-38
4.11	Memory Management Reference	4-39
4.11.1	Platforms That Support Automatic Memory Management	4-39
4.11.2	Memory Management Data Dictionary Views	4-39
4.12	Configuring and Using True Cache	4-40

5 Managing Users and Securing the Database

5.1	The Importance of Establishing a Security Policy for Your Database	5-1
5.2	Managing Users and Resources	5-1
5.3	User Privileges and Roles	5-2
5.4	Auditing Database Activity	5-2
5.5	Predefined User Accounts	5-3

6 Monitoring the Database

6.1	Monitoring Errors and Alerts	6-1
6.1.1	Monitoring Errors with Trace Files and the Alert Log	6-2
6.1.1.1	About Monitoring Errors with Trace Files and the Alert Log	6-2
6.1.1.2	Controlling the Size of an Alert Log	6-3
6.1.1.3	Controlling the Size of Trace Files	6-4
6.1.1.4	Controlling When Oracle Database Writes to Trace Files	6-5
6.1.1.5	Reading the Trace File for Shared Server Sessions	6-6
6.1.2	Monitoring a Database with Server-Generated Alerts	6-6
6.1.2.1	About Monitoring a Database with Server-Generated Alerts	6-6
6.1.2.2	Setting and Retrieving Thresholds for Server-Generated Alerts	6-7
6.1.2.3	Viewing Server-Generated Alerts	6-9
6.1.2.4	Server-Generated Alerts Data Dictionary Views	6-10
6.2	Monitoring Performance	6-10
6.2.1	Monitoring Locks	6-10
6.2.2	About Monitoring Wait Events	6-11
6.2.3	Performance Monitoring Data Dictionary Views	6-11
6.3	Monitoring Quarantined Objects	6-12
6.3.1	About Object Quarantine	6-12
6.3.2	Viewing Quarantined Objects	6-13
6.4	Automatically Monitoring Schema Objects	6-13

7 Diagnosing and Resolving Problems

7.1	About the Oracle Database Fault Diagnosability Infrastructure	7-1
7.1.1	Fault Diagnosability Infrastructure Overview	7-2
7.1.2	Incidents and Problems	7-3
7.1.2.1	About Incidents and Problems	7-3
7.1.2.2	Incident Flood Control	7-4
7.1.2.3	Related Problems Across the Topology	7-5
7.1.3	Fault Diagnosability Infrastructure Components	7-5
7.1.3.1	Automatic Diagnostic Repository (ADR)	7-6
7.1.3.2	Alert Log	7-7
7.1.3.3	Attention Log	7-7
7.1.3.4	Trace Files, Dumps, and Core Files	7-8
7.1.3.5	DDL Log	7-9
7.1.3.6	Debug Log	7-10
7.1.3.7	Other ADR Contents	7-10
7.1.3.8	Enterprise Manager Support Workbench	7-10
7.1.3.9	ADRCI Command-Line Utility	7-11
7.1.4	Structure, Contents, and Location of the Automatic Diagnostic Repository	7-11

7.2	About Investigating, Reporting, and Resolving a Problem	7-15
7.2.1	Roadmap — Investigating, Reporting, and Resolving a Problem	7-16
7.2.2	Task 1: View Critical Error Alerts in Cloud Control	7-18
7.2.3	Task 2: View Problem Details	7-19
7.2.4	Task 3: (Optional) Gather Additional Diagnostic Information	7-19
7.2.5	Task 4: (Optional) Create a Service Request	7-20
7.2.6	Task 5: Package and Upload Diagnostic Data to Oracle Support	7-20
7.2.7	Task 6: Track the Service Request and Implement Any Repairs	7-21
7.3	Diagnosing Problems	7-22
7.3.1	Identifying Problems Reactively	7-23
7.3.1.1	Viewing Problems with the Support Workbench	7-23
7.3.1.2	Adding Problems Manually to the Automatic Diagnostic Repository	7-24
7.3.1.3	Creating Incidents Manually	7-26
7.3.1.4	Using DBMS_HCHECK to Identify Data Dictionary Inconsistencies	7-26
7.3.2	Identifying Problems Proactively with Health Monitor	7-31
7.3.2.1	About Health Monitor	7-32
7.3.2.2	Running Health Checks Manually	7-33
7.3.2.3	Viewing Checker Reports	7-35
7.3.2.4	Health Monitor Views	7-38
7.3.2.5	Health Check Parameters Reference	7-38
7.3.3	Gathering Additional Diagnostic Data	7-39
7.3.3.1	Viewing the Alert Log	7-40
7.3.3.2	Finding Trace Files	7-40
7.3.4	Creating Test Cases with SQL Test Case Builder	7-41
7.3.4.1	Purpose of SQL Test Case Builder	7-42
7.3.4.2	Concepts for SQL Test Case Builder	7-42
7.3.4.3	User Interfaces for SQL Test Case Builder	7-46
7.3.4.4	Running SQL Test Case Builder	7-48
7.4	Reporting Problems	7-51
7.4.1	Incident Packages	7-52
7.4.1.1	About Incident Packages	7-52
7.4.1.2	About Correlated Diagnostic Data in Incident Packages	7-53
7.4.1.3	About Quick Packaging and Custom Packaging	7-54
7.4.1.4	About Correlated Packages	7-55
7.4.2	Packaging and Uploading Problems with Custom Packaging	7-55
7.4.3	Viewing and Modifying Incident Packages	7-59
7.4.3.1	Viewing Package Details	7-60
7.4.3.2	Accessing the Customize Package Page	7-60
7.4.3.3	Editing Incident Package Files (Copying Out and In)	7-60
7.4.3.4	Adding an External File to an Incident Package	7-61
7.4.3.5	Removing Incident Package Files	7-62
7.4.3.6	Viewing and Updating the Incident Package Activity Log	7-63

7.4.4	Creating, Editing, and Uploading Correlated Packages	7-63
7.4.5	Deleting Correlated Packages	7-64
7.4.6	Setting Incident Packaging Preferences	7-65
7.5	Resolving Problems	7-65
7.5.1	About Automatic Error Mitigation	7-66
7.5.2	Repairing SQL Failures with the SQL Repair Advisor	7-70
7.5.2.1	About the SQL Repair Advisor	7-71
7.5.2.2	Running the SQL Repair Advisor Using Cloud Control	7-71
7.5.2.3	Running the SQL Repair Advisor Using the DBMS_SQLDIAG Package Subprograms	7-72
7.5.2.4	Viewing, Disabling, or Removing a SQL Patch Using Cloud Control	7-74
7.5.2.5	Disabling or Removing a SQL Patch Using DBMS_SQLDIAG Package Subprograms	7-75
7.5.2.6	Exporting and Importing a Patch Using DBMS_SQLDIAG Package Subprograms	7-76
7.5.3	Repairing Data Corruptions with the Data Recovery Advisor	7-77
7.5.4	Quarantine for Execution Plans for SQL Statements Consuming Excessive System Resources	7-78
7.5.4.1	About Quarantine for Execution Plans for SQL Statements	7-79
7.5.4.2	Creating a Quarantine Configuration for an Execution Plan of a SQL Statement	7-81
7.5.4.3	Specifying Quarantine Thresholds in a Quarantine Configuration	7-82
7.5.4.4	Enabling and Disabling a Quarantine Configuration	7-84
7.5.4.5	Viewing the Details of a Quarantine Configuration	7-84
7.5.4.6	Deleting a Quarantine Configuration	7-85
7.5.4.7	Viewing the Details of Quarantined Execution Plans of SQL Statements	7-86
7.5.4.8	Transferring Quarantine Configurations from One Database to Another Database	7-86
7.5.4.9	Example: Quarantine for an Execution Plan of a SQL Statement Consuming Excessive System Resources	7-88
7.5.5	Viewing Attention Log Information	7-91
7.6	Diagnosis and Tracing in a PDB Using Package DBMS_USERDIAG	7-91
7.6.1	About DBMS_USERDIAG	7-92
7.6.2	Examples of Using DBMS_USERDIAG	7-92

Part II Oracle Database Structure and Storage

8 Managing Control Files

8.1	What Is a Control File?	8-1
8.2	Guidelines for Control Files	8-2
8.2.1	Provide File Names for the Control Files	8-2
8.2.2	Multiplex Control Files on Different Disks	8-3

8.2.3	Back Up Control Files	8-3
8.2.4	Manage the Size of Control Files	8-3
8.3	Creating Control Files	8-4
8.3.1	Creating Initial Control Files	8-4
8.3.2	Creating Additional Copies, Renaming, and Relocating Control Files	8-5
8.3.3	Creating New Control Files	8-5
8.3.3.1	When to Create New Control Files	8-5
8.3.3.2	The CREATE CONTROLFILE Statement	8-6
8.3.3.3	Creating New Control Files	8-7
8.4	Troubleshooting After Creating Control Files	8-8
8.4.1	Checking for Missing or Extra Files	8-8
8.4.2	Handling Errors During CREATE CONTROLFILE	8-9
8.5	Backing Up Control Files	8-9
8.6	Recovering a Control File Using a Current Copy	8-9
8.6.1	Recovering from Control File Corruption Using a Control File Copy	8-9
8.6.2	Recovering from Permanent Media Failure Using a Control File Copy	8-10
8.7	Dropping Control Files	8-10
8.8	Control Files Data Dictionary Views	8-11

9 Managing the Redo Log

9.1	What Is the Redo Log?	9-2
9.1.1	Redo Threads	9-2
9.1.2	Redo Log Contents	9-2
9.1.3	How Oracle Database Writes to the Redo Log	9-3
9.1.3.1	Active (Current) and Inactive Redo Log Files	9-4
9.1.3.2	Log Switches and Log Sequence Numbers	9-4
9.2	Planning the Redo Log	9-5
9.2.1	Multiplexing Redo Log Files	9-6
9.2.1.1	Responding to Redo Log Failure	9-7
9.2.1.2	Valid and Invalid Configurations	9-8
9.2.2	Placing Redo Log Members on Different Disks	9-8
9.2.3	Planning the Size of Redo Log Files	9-9
9.2.4	Planning the Block Size of Redo Log Files	9-9
9.2.5	Choosing the Number of Redo Log Files	9-10
9.2.6	Controlling Archive Lag	9-11
9.2.6.1	Setting the ARCHIVE_LAG_TARGET Initialization Parameter	9-11
9.2.6.2	Factors Affecting the Setting of ARCHIVE_LAG_TARGET	9-12
9.3	Creating Redo Log Groups and Members	9-12
9.3.1	Creating Redo Log Groups	9-13
9.3.2	Creating Redo Log Members	9-13
9.4	Relocating and Renaming Redo Log Members	9-14

9.5	Dropping Redo Log Groups and Members	9-15
9.5.1	Dropping Log Groups	9-16
9.5.2	Dropping Redo Log Members	9-16
9.6	Forcing Log Switches	9-17
9.7	Verifying Blocks in Redo Log Files	9-17
9.8	Clearing a Redo Log File	9-18
9.9	Reduction of Redo Generation for Direct Path Operations	9-19
9.10	Redo Log Data Dictionary Views	9-20

10 Managing Archived Redo Log Files

10.1	What Is the Archived Redo Log?	10-2
10.2	Choosing Between NOARCHIVELOG and ARCHIVELOG Mode	10-2
10.2.1	Running a Database in NOARCHIVELOG Mode	10-3
10.2.2	Running a Database in ARCHIVELOG Mode	10-3
10.3	Controlling Archiving	10-4
10.3.1	Setting the Initial Database Archiving Mode	10-5
10.3.2	Changing the Database Archiving Mode	10-5
10.3.3	Performing Manual Archiving	10-6
10.3.4	Adjusting the Number of Archiver Processes	10-7
10.4	Specifying Archive Destinations	10-7
10.4.1	Setting Initialization Parameters for Archive Destinations	10-7
10.4.1.1	Method 1: Using the LOG_ARCHIVE_DEST_n Parameter	10-8
10.4.1.2	Method 2: Using LOG_ARCHIVE_DEST and LOG_ARCHIVE_DUPLEX_DEST	10-10
10.4.2	Expanding Alternate Destinations with Log Archive Destination Groups	10-11
10.4.2.1	About Log Archive Destination Groups	10-11
10.4.2.2	Specifying Log Archive Destination Groups	10-12
10.4.3	Understanding Archive Destination Status	10-13
10.4.4	Specifying Alternate Destinations	10-14
10.5	About Log Transmission Modes	10-14
10.5.1	Normal Transmission Mode	10-15
10.5.2	Standby Transmission Mode	10-15
10.6	Managing Archive Destination Failure	10-15
10.6.1	Specifying the Minimum Number of Successful Destinations	10-16
10.6.1.1	Specifying Mandatory and Optional Destinations	10-16
10.6.1.2	Specifying the Number of Successful Destinations: Scenarios	10-17
10.6.2	Rearchiving to a Failed Destination	10-18
10.7	Controlling Trace Output Generated by the Archivelog Process	10-19
10.8	Viewing Information About the Archived Redo Log	10-19
10.8.1	Archived Redo Log Files Views	10-20

11 Managing Tablespaces

11.1	Guidelines for Managing Tablespaces	11-2
11.1.1	Use Multiple Tablespaces	11-2
11.1.2	Assign Tablespace Quotas to Users	11-3
11.2	Creating Tablespaces	11-3
11.2.1	About Creating Tablespaces	11-4
11.2.2	Locally Managed Tablespaces	11-5
11.2.2.1	About Locally Managed Tablespaces	11-5
11.2.2.2	Creating a Locally Managed Tablespace	11-6
11.2.2.3	Specifying Segment Space Management in Locally Managed Tablespaces	11-7
11.2.3	Bigfile Tablespaces	11-8
11.2.3.1	About Bigfile Tablespaces	11-8
11.2.3.2	Creating a Bigfile Tablespace	11-9
11.2.3.3	Identifying a Bigfile Tablespace	11-10
11.2.4	Tablespaces with Default Compression Attributes	11-10
11.2.4.1	About Tablespaces with Default Compression Attributes	11-10
11.2.4.2	Creating Tablespaces with Default Compression Attributes	11-10
11.2.5	Encrypted Tablespaces	11-11
11.2.5.1	About Encrypted Tablespaces	11-11
11.2.5.2	Creating Encrypted Tablespaces	11-13
11.2.5.3	Viewing Information About Encrypted Tablespaces	11-14
11.2.6	Temporary Tablespaces	11-14
11.2.6.1	About Temporary Tablespaces	11-15
11.2.6.2	Creating a Locally Managed Temporary Tablespace	11-17
11.2.6.3	Creating a Bigfile Temporary Tablespace	11-18
11.2.6.4	Viewing Space Usage for Temporary Tablespaces	11-18
11.2.7	Temporary Tablespace Groups	11-18
11.2.7.1	Multiple Temporary Tablespaces: Using Tablespace Groups	11-18
11.2.7.2	Creating a Tablespace Group	11-19
11.2.7.3	Changing Members of a Tablespace Group	11-19
11.2.7.4	Assigning a Tablespace Group as the Default Temporary Tablespace	11-20
11.3	Consider Storing Tablespaces in the In-Memory Column Store	11-20
11.4	Specifying Nonstandard Block Sizes for Tablespaces	11-21
11.5	Controlling the Writing of Redo Records	11-21
11.6	Altering Tablespace Availability	11-22
11.6.1	Taking Tablespaces Offline	11-22
11.6.2	Bringing Tablespaces Online	11-24
11.7	Using Read-Only Tablespaces	11-24

11.7.1	About Read-Only Tablespaces	11-25
11.7.2	Making a Tablespace Read-Only	11-25
11.7.3	Making a Read-Only Tablespace Writable	11-27
11.7.4	Creating a Read-Only Tablespace on a WORM Device	11-28
11.7.5	Delaying the Opening of Data Files in Read-Only Tablespaces	11-28
11.7.6	Using Read-Only Tablespaces on Object Storage	11-29
11.7.6.1	Enabling a Database for Using Object Storage	11-29
11.7.6.2	Accessing Data in Object Storage	11-32
11.7.6.3	Dropping a Read-Only Tablespace and Its Data Files in Object Storage	11-35
11.8	Altering and Maintaining Tablespaces	11-35
11.8.1	Increasing the Size of a Tablespace	11-36
11.8.2	Altering a Locally Managed Tablespace	11-36
11.8.3	Altering a Bigfile Tablespace	11-37
11.8.4	Shrinking a Tablespace	11-37
11.8.5	Altering a Locally Managed Temporary Tablespace	11-41
11.8.6	Shrinking a Locally Managed Temporary Tablespace	11-42
11.9	Renaming Tablespaces	11-42
11.10	Dropping Tablespaces	11-43
11.11	Managing Lost Write Protection with Shadow Tablespaces	11-44
11.11.1	About Shadow Lost Write Protection	11-45
11.11.2	Creating Shadow Tablespaces for Shadow Lost Write Protection	11-47
11.11.3	Enabling Shadow Lost Write Protection for a Database	11-47
11.11.4	Enabling Shadow Lost Write Protection for Tablespaces and Data Files	11-49
11.11.5	Disabling Shadow Lost Write Protection for a Database	11-50
11.11.6	Removing or Suspending Shadow Lost Write Protection	11-51
11.11.7	Dropping a Shadow Tablespace	11-52
11.12	Managing the SYSAUX Tablespace	11-52
11.12.1	Monitoring Occupants of the SYSAUX Tablespace	11-53
11.12.2	Moving Occupants Out Of or Into the SYSAUX Tablespace	11-53
11.12.3	Controlling the Size of the SYSAUX Tablespace	11-54
11.13	Correcting Problems with Locally Managed Tablespaces	11-54
11.13.1	Diagnosing and Repairing Locally Managed Tablespace Problems	11-55
11.13.2	Scenario 1: Fixing Bitmap When Allocated Blocks are Marked Free (No Overlap)	11-56
11.13.3	Scenario 2: Dropping a Corrupted Segment	11-57
11.13.4	Scenario 3: Fixing Bitmap Where Overlap is Reported	11-57
11.13.5	Scenario 4: Correcting Media Corruption of Bitmap Blocks	11-57
11.13.6	Scenario 5: Migrating from a Dictionary-Managed to a Locally Managed Tablespace	11-57
11.14	Migrating the SYSTEM Tablespace to a Locally Managed Tablespace	11-58
11.15	Viewing Information About Tablespaces	11-59
11.15.1	Tablespace Data Dictionary Views	11-59

11.15.2	Example 1: Listing Tablespaces and Default Storage Parameters	11-60
11.15.3	Example 2: Listing the Data Files and Associated Tablespaces of a Database	11-60
11.15.4	Example 3: Displaying Statistics for Free Space (Extents) of Each Tablespace	11-61

12 Managing Data Files and Temp Files

12.1	Guidelines for Managing Data Files	12-2
12.1.1	About Data Files	12-2
12.1.2	Determine the Number of Data Files	12-3
12.1.2.1	About Determining the Number of Data Files	12-3
12.1.2.2	Determine a Value for the DB_FILES Initialization Parameter	12-4
12.1.2.3	Consider Possible Limitations When Adding Data Files to a Tablespace	12-4
12.1.2.4	Consider the Performance Impact of the Number of Data Files	12-4
12.1.3	Determine the Size of Data Files	12-5
12.1.4	Place Data Files Appropriately	12-5
12.1.5	Store Data Files Separate from Redo Log Files	12-5
12.2	Creating Data Files and Adding Data Files to a Tablespace	12-5
12.3	Changing Data File Size	12-6
12.3.1	Enabling and Disabling Automatic Extension for a Data File	12-6
12.3.2	Manually Resizing a Data File	12-7
12.4	Altering Data File Availability	12-8
12.4.1	About Altering Data File Availability	12-8
12.4.2	Bringing Data Files Online or Taking Offline in ARCHIVELOG Mode	12-9
12.4.3	Taking Data Files Offline in NOARCHIVELOG Mode	12-9
12.4.4	Altering the Availability of All Data Files or Temp Files in a Tablespace	12-10
12.5	Renaming and Relocating Data Files	12-10
12.5.1	Renaming and Relocating Online Data Files	12-11
12.5.2	Renaming and Relocating Offline Data Files	12-13
12.5.2.1	Procedures for Renaming and Relocating Offline Data Files in a Single Tablespace	12-14
12.5.2.2	Renaming and Relocating Offline Data Files in Multiple Tablespaces	12-16
12.6	Dropping Data Files	12-17
12.7	Verifying Data Blocks in Data Files	12-18
12.8	Copying Files Using the Database Server	12-18
12.8.1	About Copying Files Using the Database Server	12-19
12.8.2	Copying a File on a Local File System	12-19
12.8.3	Third-Party File Transfer	12-20
12.8.4	Advanced File Transfer Mechanisms	12-21
12.8.5	File Transfer and the DBMS_SCHEDULER Package	12-21
12.9	Mapping Files to Physical Devices	12-22
12.9.1	Overview of Oracle Database File Mapping Interface	12-23
12.9.2	How the Oracle Database File Mapping Interface Works	12-23

12.9.2.1	Components of File Mapping	12-23
12.9.2.2	Mapping Structures	12-25
12.9.2.3	Example of Mapping Structures	12-26
12.9.2.4	Configuration ID	12-27
12.9.3	Using the Oracle Database File Mapping Interface	12-27
12.9.3.1	Enabling File Mapping	12-28
12.9.3.2	Using the DBMS_STORAGE_MAP Package	12-29
12.9.3.3	Obtaining Information from the File Mapping Views	12-29
12.9.4	File Mapping Examples	12-31
12.9.4.1	Example 1: Map All Database Files that Span a Device	12-31
12.9.4.2	Example 2: Map a File Into Its Corresponding Devices	12-32
12.9.4.3	Example 3: Map a Database Object	12-32
12.10	Data Files Data Dictionary Views	12-33

13 Transporting Data

13.1	About Transporting Data	13-1
13.1.1	Purpose of Transporting Data	13-2
13.1.2	Transporting Data: Scenarios	13-2
13.1.2.1	Scenarios for Full Transportable Export/import	13-2
13.1.2.2	Scenarios for Transportable Tablespaces or Transportable Tables	13-4
13.1.3	Transporting Data Across Platforms	13-8
13.1.4	General Limitations on Transporting Data	13-10
13.1.5	Compatibility Considerations for Transporting Data	13-12
13.2	Transporting Databases	13-13
13.2.1	Introduction to Full Transportable Export/Import	13-13
13.2.2	Limitations on Full Transportable Export/import	13-14
13.2.3	Transporting a Database Using an Export Dump File	13-15
13.2.4	Transporting a Database Over the Network	13-21
13.3	Transporting Tablespaces Between Databases	13-26
13.3.1	Introduction to Transportable Tablespaces	13-27
13.3.2	Limitations on Transportable Tablespaces	13-27
13.3.3	Transporting Tablespaces Between Databases	13-28
13.3.3.1	Task 1: Pick a Self-Contained Set of Tablespaces	13-30
13.3.3.2	Task 2: Generate a Transportable Tablespace Set	13-31
13.3.3.3	Task 3: Transport the Export Dump File	13-33
13.3.3.4	Task 4: Transport the Tablespace Set	13-33
13.3.3.5	Task 5: (Optional) Restore Tablespaces to Read/Write Mode	13-34
13.3.3.6	Task 6: Import the Tablespace Set	13-34
13.4	Transporting Tables, Partitions, or Subpartitions Between Databases	13-36
13.4.1	Introduction to Transportable Tables	13-37
13.4.2	Limitations on Transportable Tables	13-37

13.4.3	Transporting Tables, Partitions, or Subpartitions Using an Export Dump File	13-38
13.4.4	Transporting Tables, Partitions, or Subpartitions Over the Network	13-43
13.5	Converting Data Between Platforms	13-48
13.5.1	Converting Data Between Platforms Using the DBMS_FILE_TRANSFER Package	13-48
13.5.2	Converting Data Between Platforms Using RMAN	13-50
13.5.2.1	Converting Tablespaces on the Source System After Export	13-51
13.5.2.2	Converting Data Files on the Target System Before Import	13-52
13.6	Guidelines for Transferring Data Files	13-53

14 Managing Undo

14.1	What Is Undo?	14-2
14.2	Introduction to Automatic Undo Management	14-2
14.2.1	Overview of Automatic Undo Management	14-2
14.2.2	The Undo Retention Period	14-4
14.2.2.1	About the Undo Retention Period	14-4
14.2.2.2	Automatic Tuning of Undo Retention	14-5
14.2.2.3	Retention Guarantee	14-6
14.2.2.4	Undo Retention Tuning and Alert Thresholds	14-7
14.2.2.5	Tracking the Tuned Undo Retention Period	14-7
14.3	Setting the Minimum Undo Retention Period	14-8
14.4	Sizing a Fixed-Size Undo Tablespace	14-8
14.4.1	Activating the Undo Advisor PL/SQL Interface	14-9
14.5	Managing Undo Tablespaces	14-10
14.5.1	Creating an Undo Tablespace	14-10
14.5.1.1	About Creating an Undo Tablespace	14-11
14.5.1.2	Using CREATE DATABASE to Create an Undo Tablespace	14-11
14.5.1.3	Using the CREATE UNDO TABLESPACE Statement	14-12
14.5.2	Altering an Undo Tablespace	14-12
14.5.3	Dropping an Undo Tablespace	14-13
14.5.4	Switching Undo Tablespaces	14-13
14.5.5	Establishing User Quotas for Undo Space	14-14
14.5.6	Managing Space Threshold Alerts for the Undo Tablespace	14-14
14.6	Migrating to Automatic Undo Management	14-15
14.7	Managing Temporary Undo	14-15
14.7.1	About Managing Temporary Undo	14-15
14.7.2	Enabling and Disabling Temporary Undo	14-16
14.8	Undo Space Data Dictionary Views	14-17

15 Using Oracle Managed Files

15.1	About Oracle Managed Files	15-1
15.1.1	What Is Oracle Managed Files?	15-2
15.1.2	Who Can Use Oracle Managed Files?	15-2
15.1.3	What Is a Logical Volume Manager?	15-3
15.1.4	What Is a File System?	15-3
15.1.5	Benefits of Using Oracle Managed Files	15-3
15.1.6	Oracle Managed Files and Existing Functionality	15-4
15.2	Enabling the Creation and Use of Oracle Managed Files	15-4
15.2.1	Initialization Parameters That Enable Oracle Managed Files	15-5
15.2.2	Setting the DB_CREATE_FILE_DEST Initialization Parameter	15-6
15.2.3	Setting the DB_RECOVERY_FILE_DEST Parameter	15-6
15.2.4	Setting the DB_CREATE_ONLINE_LOG_DEST_n Initialization Parameters	15-6
15.3	Creating Oracle Managed Files	15-7
15.3.1	When Oracle Database Creates Oracle Managed Files	15-7
15.3.2	How Oracle Managed Files Are Named	15-8
15.3.3	Creating Oracle Managed Files at Database Creation	15-9
15.3.3.1	Specifying Control Files at Database Creation	15-10
15.3.3.2	Specifying Redo Log Files at Database Creation	15-11
15.3.3.3	Specifying the SYSTEM and SYSAUX Tablespace Data Files at Database Creation	15-12
15.3.3.4	Specifying the Undo Tablespace Data File at Database Creation	15-12
15.3.3.5	Specifying the Default Temporary Tablespace Temp File at Database Creation	15-13
15.3.3.6	CREATE DATABASE Statement Using Oracle Managed Files: Examples	15-13
15.3.4	Creating Data Files for Tablespaces Using Oracle Managed Files	15-15
15.3.4.1	About Creating Data Files for Tablespaces Using Oracle Managed Files	15-15
15.3.4.2	CREATE TABLESPACE: Examples	15-16
15.3.4.3	CREATE UNDO TABLESPACE: Example	15-17
15.3.4.4	ALTER TABLESPACE: Example	15-18
15.3.5	Creating Temp Files for Temporary Tablespaces Using Oracle Managed Files	15-18
15.3.5.1	About Creating Temp Files for Temporary Tablespaces Using Oracle Managed Files	15-18
15.3.5.2	CREATE TEMPORARY TABLESPACE: Example	15-19
15.3.5.3	ALTER TABLESPACE... ADD TEMPFILE: Example	15-19
15.3.6	Creating Control Files Using Oracle Managed Files	15-20
15.3.6.1	About Creating Control Files Using Oracle Managed Files	15-20
15.3.6.2	CREATE CONTROLFILE Using NORESETLOGS Keyword: Example	15-21
15.3.6.3	CREATE CONTROLFILE Using RESETLOGS Keyword: Example	15-21
15.3.7	Creating Redo Log Files Using Oracle Managed Files	15-22
15.3.7.1	Using the ALTER DATABASE ADD LOGFILE Statement	15-22
15.3.7.2	Using the ALTER DATABASE OPEN RESETLOGS Statement	15-23

15.3.8	Creating Archived Logs Using Oracle Managed Files	15-23
15.4	Operation of Oracle Managed Files	15-23
15.4.1	Dropping Data Files and Temp Files	15-24
15.4.2	Dropping Redo Log Files	15-24
15.4.3	Renaming Files	15-24
15.4.4	Managing Standby Databases	15-25
15.5	Scenarios for Using Oracle Managed Files	15-25
15.5.1	Scenario 1: Create and Manage a Database with Multiplexed Redo Logs	15-25
15.5.2	Scenario 2: Create and Manage a Database with Database and Fast Recovery Areas	15-29
15.5.3	Scenario 3: Adding Oracle Managed Files to an Existing Database	15-30

16 Using Persistent Memory Database

16.1	About Persistent Memory Database	16-1
16.1.1	What Is Persistent Memory Database?	16-1
16.1.2	What Is Oracle Persistent Memory Filestore?	16-1
16.1.3	What Is Directly Mapped Buffer Cache?	16-2
16.1.4	Benefits of Using Persistent Memory Database	16-2
16.2	Setting Initialization Parameters for Persistent Memory Database	16-2
16.2.1	Persistent Memory Database Initialization Parameters	16-2
16.3	Creating a PMEM Filestore for an Oracle Database	16-3
16.3.1	Creating a PMEM Filestore Before Creating the Database	16-3
16.3.2	Creating a Database on PMEM Storage Using Oracle DBCA	16-4
16.3.3	Creating an Oracle Database in the PMEM Filestore	16-5
16.3.4	Migrating an Oracle Database to a PMEM Filestore	16-5
16.4	Managing a PMEM Filestore	16-5
16.4.1	Viewing Information About a PMEM Filestore	16-5
16.4.2	Mounting a PMEM Filestore	16-6
16.4.3	Dismounting a PMEM Filestore	16-6
16.4.4	Changing the Attributes of a PMEM Filestore	16-6
16.4.5	Dropping a PMEM Filestore	16-6

Part III Schema Objects

17 Managing Schema Objects

17.1	About Common and Local Objects	17-2
17.2	About the Container for Schema Objects	17-2
17.3	Creating Multiple Tables and Views in a Single Operation	17-2
17.4	Analyzing Tables, Indexes, and Clusters	17-3
17.4.1	About Analyzing Tables, Indexes, and Clusters	17-4

17.4.2	Using DBMS_STATS to Collect Table and Index Statistics	17-4
17.4.3	Validating Tables, Indexes, Clusters, and Materialized Views	17-5
17.4.4	Cross Validation of a Table and an Index with a Query	17-6
17.4.5	Listing Chained Rows of Tables and Clusters	17-6
17.4.5.1	Creating a CHAINED_ROWS Table	17-7
17.4.5.2	Eliminating Migrated or Chained Rows in a Table	17-7
17.5	Truncating Tables and Clusters	17-8
17.5.1	Using DELETE to Truncate a Table	17-9
17.5.2	Using DROP and CREATE to Truncate a Table	17-9
17.5.3	Using TRUNCATE	17-9
17.6	Enabling and Disabling Triggers	17-10
17.6.1	About Enabling and Disabling Triggers	17-11
17.6.2	Enabling Triggers	17-12
17.6.3	Disabling Triggers	17-12
17.7	Managing Integrity Constraints	17-13
17.7.1	Integrity Constraint States	17-13
17.7.1.1	About Integrity Constraint States	17-14
17.7.1.2	About Disabling Constraints	17-14
17.7.1.3	About Enabling Constraints	17-15
17.7.1.4	About the Enable Novalidate Constraint State	17-15
17.7.1.5	Efficient Use of Integrity Constraints: A Procedure	17-15
17.7.2	Setting Integrity Constraints Upon Definition	17-16
17.7.2.1	Disabling Constraints Upon Definition	17-17
17.7.2.2	Enabling Constraints Upon Definition	17-17
17.7.3	Modifying, Renaming, or Dropping Existing Integrity Constraints	17-18
17.7.3.1	Disabling and Enabling Constraints	17-18
17.7.3.2	Renaming Constraints	17-19
17.7.3.3	Dropping Constraints	17-19
17.7.4	Deferring Constraint Checks	17-19
17.7.4.1	Set All Constraints Deferred	17-20
17.7.4.2	Check the Commit (Optional)	17-20
17.7.5	Reporting Constraint Exceptions	17-20
17.7.6	Viewing Constraint Information	17-22
17.8	Renaming Schema Objects	17-22
17.9	Managing Object Dependencies	17-23
17.9.1	About Object Dependencies and Object Invalidation	17-24
17.9.2	Manually Recompiling Invalid Objects with DDL	17-25
17.9.3	Manually Recompiling Invalid Objects with PL/SQL Package Procedures	17-25
17.10	Managing Object Name Resolution	17-26
17.11	Switching to a Different Schema	17-28
17.12	Managing Editions	17-28
17.12.1	About Editions and Edition-Based Redefinition	17-29

17.12.2	DBA Tasks for Edition-Based Redefinition	17-29
17.12.3	Setting the Database Default Edition	17-30
17.12.4	Querying the Database Default Edition	17-30
17.12.5	Setting the Edition Attribute of a Database Service	17-30
17.12.5.1	About Setting the Edition Attribute of a Database Service	17-31
17.12.5.2	Setting the Edition Attribute During Database Service Creation	17-31
17.12.5.3	Setting the Edition Attribute of an Existing Database Service	17-31
17.12.6	Using an Edition	17-32
17.12.7	Editions Data Dictionary Views	17-32
17.13	Displaying Information About Schema Objects	17-33
17.13.1	Using a PL/SQL Package to Display Information About Schema Objects	17-33
17.13.2	Schema Objects Data Dictionary Views	17-34
17.13.2.1	Example 1: Displaying Schema Objects By Type	17-34
17.13.2.2	Example 2: Displaying Dependencies of Views and Synonyms	17-35

18 Managing Space for Schema Objects

18.1	Managing Tablespace Alerts	18-1
18.1.1	About Managing Tablespace Alerts	18-2
18.1.2	Setting Alert Thresholds	18-3
18.1.3	Viewing Alerts	18-5
18.1.4	Limitations	18-5
18.2	Managing Resumable Space Allocation	18-6
18.2.1	Resumable Space Allocation Overview	18-6
18.2.1.1	How Resumable Space Allocation Works	18-7
18.2.1.2	What Operations are Resumable?	18-8
18.2.1.3	What Errors are Correctable?	18-8
18.2.1.4	Resumable Space Allocation and Distributed Operations	18-9
18.2.1.5	Parallel Execution and Resumable Space Allocation	18-9
18.2.2	Enabling and Disabling Resumable Space Allocation	18-9
18.2.2.1	About Enabling and Disabling Resumable Space Allocation	18-10
18.2.2.2	Setting the RESUMABLE_TIMEOUT Initialization Parameter	18-10
18.2.2.3	Using ALTER SESSION to Enable and Disable Resumable Space Allocation	18-11
18.2.3	Using a LOGON Trigger to Set Default Resumable Mode	18-12
18.2.4	Detecting Suspended Statements	18-12
18.2.4.1	Notifying Users: The AFTER SUSPEND System Event and Trigger	18-13
18.2.4.2	Using Views to Obtain Information About Suspended Statements	18-13
18.2.4.3	Using the DBMS_RESUMABLE Package	18-14
18.2.5	Operation-Suspended Alert	18-14
18.2.6	Resumable Space Allocation Example: Registering an AFTER SUSPEND Trigger	18-15
18.3	Reclaiming Unused Space	18-16

18.3.1	About Reclaimable Unused Space	18-17
18.3.2	The Segment Advisor	18-17
18.3.2.1	About the Segment Advisor	18-18
18.3.2.2	Using the Segment Advisor	18-18
18.3.2.3	Automatic Segment Advisor	18-18
18.3.2.4	Running the Segment Advisor Manually	18-19
18.3.2.5	Viewing Segment Advisor Results	18-24
18.3.2.6	Configuring the Automatic Segment Advisor	18-29
18.3.2.7	Viewing Automatic Segment Advisor Information	18-30
18.3.3	Shrinking Database Segments Online	18-31
18.3.4	Deallocating Unused Space	18-33
18.4	Dropping Unused Object Storage	18-34
18.5	Understanding Space Usage of Data Types	18-35
18.6	Displaying Information About Space Usage for Schema Objects	18-35
18.6.1	Using PL/SQL Packages to Display Information About Schema Object Space Usage	18-35
18.6.2	Schema Objects Space Usage Data Dictionary Views	18-36
18.6.2.1	Example 1: Displaying Segment Information	18-37
18.6.2.2	Example 2: Displaying Extent Information	18-37
18.6.2.3	Example 3: Displaying the Free Space (Extents) in a Tablespace	18-38
18.7	Capacity Planning for Database Objects	18-38
18.7.1	Estimating the Space Use of a Table	18-39
18.7.2	Estimating the Space Use of an Index	18-40
18.7.3	Obtaining Object Growth Trends	18-40

19 Managing Tables

19.1	About Tables	19-2
19.2	Guidelines for Managing Tables	19-3
19.2.1	Design Tables Before Creating Them	19-5
19.2.2	Specify the Type of Table to Create	19-5
19.2.3	Specify the Location of Each Table	19-6
19.2.4	Consider Parallelizing Table Creation	19-6
19.2.5	Consider Using NOLOGGING When Creating Tables	19-7
19.2.6	Consider Using Table Compression	19-7
19.2.6.1	About Table Compression	19-8
19.2.6.2	Examples Related to Table Compression	19-11
19.2.6.3	Compression and Partitioned Tables	19-13
19.2.6.4	Determining If a Table Is Compressed	19-13
19.2.6.5	Determining Which Rows Are Compressed	19-13
19.2.6.6	Changing the Compression Level	19-14
19.2.6.7	Adding and Dropping Columns in Compressed Tables	19-15

19.2.6.8	Exporting and Importing Hybrid Columnar Compression Tables	19-15
19.2.6.9	Restoring a Hybrid Columnar Compression Table	19-16
19.2.6.10	Notes and Restrictions for Compressed Tables	19-17
19.2.6.11	Packing Compressed Tables	19-17
19.2.7	Managing Table Compression Using Enterprise Manager Cloud Control	19-18
19.2.7.1	Table Compression and Enterprise Manager Cloud Control	19-18
19.2.7.2	Viewing the Compression Summary at the Database Level	19-19
19.2.7.3	Viewing the Compression Summary at the Tablespace Level	19-19
19.2.7.4	Estimating the Compression Ratio	19-20
19.2.7.5	Compressing an Object	19-20
19.2.7.6	Viewing Compression Advice	19-21
19.2.7.7	Initiating Automatic Data Optimization on an Object	19-21
19.2.8	Consider Using Segment-Level and Row-Level Compression Tiering	19-21
19.2.9	Consider Using Attribute-Clustered Tables	19-23
19.2.10	Consider Using Zone Maps	19-24
19.2.11	Consider Storing Tables in the In-Memory Column Store	19-25
19.2.12	Consider Using Invisible Columns	19-25
19.2.12.1	Understand Invisible Columns	19-25
19.2.12.2	Invisible Columns and Column Ordering	19-26
19.2.13	Consider Encrypting Columns That Contain Sensitive Data	19-28
19.2.14	Understand Deferred Segment Creation	19-29
19.2.15	Materializing Segments	19-32
19.2.16	Estimate Table Size and Plan Accordingly	19-32
19.2.17	Restrictions to Consider When Creating Tables	19-32
19.3	Creating Tables	19-33
19.3.1	Example: Creating a Table	19-34
19.3.2	Creating a Temporary Table	19-35
19.3.2.1	Overview of Temporary Tables	19-35
19.3.2.2	Considerations When Creating Temporary Tables	19-36
19.3.2.3	Creating Global Temporary Tables	19-36
19.3.2.4	Creating Private Temporary Tables	19-38
19.3.3	Parallelizing Table Creation	19-40
19.4	Loading Tables	19-41
19.4.1	Methods for Loading Tables	19-41
19.4.2	Improving INSERT Performance with Direct-Path INSERT	19-43
19.4.2.1	About Direct-Path INSERT	19-43
19.4.2.2	How Direct-Path INSERT Works	19-44
19.4.2.3	Loading Data with Direct-Path INSERT	19-45
19.4.2.4	Logging Modes for Direct-Path INSERT	19-46
19.4.2.5	Additional Considerations for Direct-Path INSERT	19-47
19.4.3	Using Conventional Inserts to Load Tables	19-49
19.4.4	Avoiding Bulk INSERT Failures with DML Error Logging	19-49

19.4.4.1	Inserting Data with DML Error Logging	19-50
19.4.4.2	Error Logging Table Format	19-51
19.4.4.3	Creating an Error Logging Table	19-52
19.4.4.4	Error Logging Restrictions and Caveats	19-53
19.5	Optimizing the Performance of Bulk Updates	19-54
19.6	Automatically Collecting Statistics on Tables	19-55
19.7	Altering Tables	19-56
19.7.1	Reasons for Using the ALTER TABLE Statement	19-57
19.7.2	Altering Physical Attributes of a Table	19-58
19.7.3	Moving a Table to a New Segment or Tablespace	19-58
19.7.3.1	About Moving a Table to a New Segment or Tablespace	19-58
19.7.3.2	Moving a Table	19-59
19.7.3.3	Moving a Table Partition or Subpartition Online	19-60
19.7.4	Manually Allocating Storage for a Table	19-61
19.7.5	Modifying an Existing Column Definition	19-61
19.7.6	Adding Table Columns	19-62
19.7.7	Renaming Table Columns	19-63
19.7.8	Dropping Table Columns	19-63
19.7.8.1	Removing Columns from Tables	19-64
19.7.8.2	Marking Columns Unused	19-65
19.7.8.3	Removing Unused Columns	19-65
19.7.8.4	Dropping Columns in Compressed Tables	19-66
19.7.9	Placing a Table in Read-Only Mode	19-66
19.8	Redefining Tables Online	19-67
19.8.1	About Redefining Tables Online	19-68
19.8.2	Features of Online Table Redefinition	19-69
19.8.3	Privileges Required for the DBMS_REDEFINITION Package	19-71
19.8.4	Restrictions for Online Redefinition of Tables	19-71
19.8.5	Performing Online Redefinition with the REDEF_TABLE Procedure	19-73
19.8.6	Redefining Tables Online with Multiple Procedures in DBMS_REDEFINITION	19-74
19.8.6.1	Performing Online Redefinition with Multiple Procedures in DBMS_REDEFINITION	19-74
19.8.6.2	Constructing a Column Mapping String	19-77
19.8.6.3	Handling Virtual Private Database (VPD) Policies During Online Redefinition	19-78
19.8.6.4	Creating Dependent Objects Automatically	19-79
19.8.6.5	Creating Dependent Objects Manually	19-79
19.8.7	Results of the Redefinition Process	19-80
19.8.8	Performing Intermediate Synchronization	19-81
19.8.9	Refreshing Dependent Materialized Views During Online Table Redefinition	19-81
19.8.10	Monitoring Online Table Redefinition Progress	19-85
19.8.11	Restarting Online Table Redefinition After a Failure	19-88

19.8.12	Rolling Back Online Table Redefinition	19-91
19.8.12.1	About Online Table Redefinition Rollback	19-92
19.8.12.2	Performing Online Table Redefinition Rollback	19-92
19.8.13	Terminating Online Table Redefinition and Cleaning Up After Errors	19-95
19.8.14	Online Redefinition of One or More Partitions	19-96
19.8.14.1	Rules for Online Redefinition of a Single Partition	19-97
19.8.15	Online Table Redefinition Examples	19-98
19.9	Researching and Reversing Erroneous Table Changes	19-119
19.10	Recovering Tables Using Oracle Flashback Table	19-119
19.11	Dropping Tables	19-120
19.12	Using Flashback Drop and Managing the Recycle Bin	19-121
19.12.1	What Is the Recycle Bin?	19-122
19.12.2	Enabling and Disabling the Recycle Bin	19-123
19.12.3	Viewing and Querying Objects in the Recycle Bin	19-124
19.12.4	Purging Objects in the Recycle Bin	19-124
19.12.5	Restoring Tables from the Recycle Bin	19-125
19.13	Managing Index-Organized Tables	19-126
19.13.1	What Are Index-Organized Tables?	19-127
19.13.2	Creating Index-Organized Tables	19-128
19.13.2.1	About Creating Index-Organized Tables	19-128
19.13.2.2	Example: Creating an Index-Organized Table	19-129
19.13.2.3	Restrictions for Index-Organized Tables	19-129
19.13.2.4	Creating Index-Organized Tables That Contain Object Types	19-130
19.13.2.5	Choosing and Monitoring a Threshold Value	19-131
19.13.2.6	Using the INCLUDING Clause	19-131
19.13.2.7	Parallelizing Index-Organized Table Creation	19-132
19.13.2.8	Using Prefix Compression	19-133
19.13.3	Maintaining Index-Organized Tables	19-134
19.13.3.1	Altering Index-Organized Tables	19-134
19.13.3.2	Moving (Rebuilding) Index-Organized Tables	19-134
19.13.4	Creating Secondary Indexes on Index-Organized Tables	19-135
19.13.4.1	About Secondary Indexes on Index-Organized Tables	19-136
19.13.4.2	Creating a Secondary Index on an Index-Organized Table	19-136
19.13.4.3	Maintaining Physical Guesses in Logical Rowids	19-136
19.13.4.4	Specifying Bitmap Indexes on Index-Organized Tables	19-137
19.13.5	Analyzing Index-Organized Tables	19-137
19.13.5.1	Collecting Optimizer Statistics for Index-Organized Tables	19-137
19.13.5.2	Validating the Structure of Index-Organized Tables	19-138
19.13.6	Using the ORDER BY Clause with Index-Organized Tables	19-138
19.13.7	Converting Index-Organized Tables to Regular Tables	19-139
19.14	Managing Partitioned Tables	19-139
19.15	Managing External Tables	19-139

19.15.1	About External Tables	19-140
19.15.2	Creating External Tables	19-142
19.15.3	Altering External Tables	19-146
19.15.4	Preprocessing External Tables	19-147
19.15.5	Overriding Parameters for External Tables in a Query	19-149
19.15.6	Using Inline External Tables	19-149
19.15.7	Partitioning External Tables	19-150
19.15.7.1	About Partitioning External Tables	19-150
19.15.7.2	Restrictions for Partitioned External Tables	19-152
19.15.7.3	Creating a Partitioned External Table	19-153
19.15.7.4	Altering a Partitioned External Table	19-157
19.15.8	Dropping External Tables	19-157
19.15.9	System and Object Privileges for External Tables	19-157
19.15.10	Using SQL*Loader for External Tables with Partition Values in File Paths	19-158
19.16	Managing Hybrid Partitioned Tables	19-158
19.17	Managing Immutable Tables	19-159
19.17.1	About Immutable Tables	19-159
19.17.2	Guidelines for Managing Immutable Tables	19-160
19.17.2.1	Specify the Retention Period for the Immutable Table	19-160
19.17.2.2	Specify the Retention Period for Rows in the Immutable Table	19-161
19.17.2.3	Restrictions for Immutable Tables	19-161
19.17.3	Creating Immutable Tables	19-162
19.17.4	Altering Immutable Tables	19-163
19.17.5	Adding and Dropping User Columns in Immutable Tables	19-163
19.17.6	Creating Row Versions in Immutable Tables	19-164
19.17.7	Deleting Rows from Immutable Tables	19-164
19.17.8	Dropping Immutable Tables	19-165
19.17.9	Immutable Tables Data Dictionary Views	19-165
19.18	Managing Blockchain Tables	19-166
19.18.1	About Blockchain Tables	19-168
19.18.1.1	Benefits of Using Blockchain Tables	19-168
19.18.1.2	Chaining Rows in Blockchain Tables	19-170
19.18.1.3	Hidden Columns in Blockchain Tables	19-171
19.18.2	Guidelines for Managing Blockchain Tables	19-174
19.18.2.1	Specify the Retention Period for the Blockchain Table	19-174
19.18.2.2	Specify the Retention Period for Rows in the Blockchain Table	19-175
19.18.2.3	Exporting and Importing Blockchain Tables with Oracle Data Pump	19-175
19.18.2.4	Restrictions for Blockchain Tables	19-176
19.18.3	Creating Blockchain Tables	19-177
19.18.4	Adding and Dropping User Columns in Blockchain Tables	19-179
19.18.5	Creating Row Versions in Blockchain Tables	19-179
19.18.6	Creating User Chains in Blockchain Tables	19-180

19.18.7	Altering Blockchain Tables	19-181
19.18.8	Adding Certificates Used to Sign Blockchain Table Rows	19-181
19.18.9	Adding the Certificate of a Certificate Authority to the Database	19-182
19.18.10	Deleting Certificates in Blockchain Tables	19-183
19.18.11	Adding a User Signature to Blockchain Table Rows	19-183
19.18.12	Allowing a Delegate to Sign Blockchain Table Rows	19-185
19.18.13	Countersigning Blockchain Table Rows	19-186
19.18.14	Validating Data in Blockchain Tables	19-187
19.18.15	Verifying the Integrity of Blockchain Tables	19-188
19.18.15.1	Generating a Signed Digest for Blockchain Tables	19-189
19.18.15.2	Verifying Blockchain Table Rows Created in a Specified Time Period	19-190
19.18.16	Deleting Rows from Blockchain Tables	19-191
19.18.17	Dropping Blockchain Tables	19-192
19.18.18	Setting the Table Retention Threshold	19-193
19.18.19	Determining the Data Format for Row Content to Compute Row Hash	19-193
19.18.20	Determining the Data Format to Compute Row Signature	19-195
19.18.21	Displaying the Byte Values of Data in Blockchain Tables	19-195
19.18.22	Creating a Regular Table with Blockchain History Log	19-197
19.18.23	Blockchain Tables Data Dictionary Views	19-197
19.19	Tables Data Dictionary Views	19-198

20 Managing Indexes

20.1	About Indexes	20-1
20.2	Guidelines for Managing Indexes	20-2
20.2.1	Create Indexes After Inserting Table Data	20-3
20.2.2	Index the Correct Tables and Columns	20-4
20.2.3	Order Index Columns for Performance	20-5
20.2.4	Limit the Number of Indexes for Each Table	20-5
20.2.5	Drop Indexes That Are No Longer Required	20-5
20.2.6	Indexes and Deferred Segment Creation	20-6
20.2.7	Estimate Index Size and Set Storage Parameters	20-6
20.2.8	Specify the Tablespace for Each Index	20-7
20.2.9	Consider Parallelizing Index Creation	20-7
20.2.10	Consider Creating Indexes with NOLOGGING	20-7
20.2.11	Understand When to Use Unusable or Invisible Indexes	20-8
20.2.12	Understand When to Create Multiple Indexes on the Same Set of Columns	20-9
20.2.13	Consider Costs and Benefits of Coalescing or Rebuilding Indexes	20-10
20.2.14	Consider Cost Before Disabling or Dropping Constraints	20-11
20.2.15	Consider Using the In-Memory Column Store to Reduce the Number of Indexes	20-11
20.3	Creating Indexes	20-12

20.3.1	Prerequisites for Creating Indexes	20-13
20.3.2	Creating an Index Explicitly	20-13
20.3.3	Creating a Unique Index Explicitly	20-14
20.3.4	Creating an Index Associated with a Constraint	20-14
20.3.4.1	About Creating an Index Associated with a Constraint	20-15
20.3.4.2	Specifying Storage Options for an Index Associated with a Constraint	20-15
20.3.4.3	Specifying the Index Associated with a Constraint	20-15
20.3.5	Creating a Large Index	20-16
20.3.6	Creating an Index Online	20-17
20.3.7	Creating a Function-Based Index	20-17
20.3.8	Creating a Compressed Index	20-18
20.3.8.1	Creating an Index Using Prefix Compression	20-19
20.3.8.2	Creating an Index Using Advanced Index Compression	20-20
20.3.9	Creating an Unusable Index	20-21
20.3.10	Creating an Invisible Index	20-22
20.3.11	Creating Multiple Indexes on the Same Set of Columns	20-23
20.3.12	Creating a Vector Index	20-24
20.4	Altering Indexes	20-24
20.4.1	About Altering Indexes	20-24
20.4.2	Altering Storage Characteristics of an Index	20-25
20.4.3	Rebuilding an Existing Index	20-26
20.4.4	Making an Index Unusable	20-26
20.4.5	Making an Index Invisible or Visible	20-28
20.4.6	Renaming an Index	20-29
20.4.7	Monitoring Index Usage	20-29
20.5	Monitoring Space Use of Indexes	20-30
20.6	Dropping Indexes	20-31
20.7	Managing Automatic Indexes	20-31
20.7.1	About Automatic Indexing	20-32
20.7.2	How Automatic Indexing Works	20-33
20.7.3	Configuring Automatic Indexing in an Oracle Database	20-34
20.7.4	Generating Automatic Indexing Reports	20-38
20.7.5	Views Containing the Automatic Indexing Information	20-40
20.8	Indexes Data Dictionary Views	20-41

21 Managing Clusters

21.1	About Clusters	21-1
21.2	Guidelines for Managing Clusters	21-3
21.2.1	Choose Appropriate Tables for the Cluster	21-3
21.2.2	Choose Appropriate Columns for the Cluster Key	21-4

21.2.3	Specify the Space Required by an Average Cluster Key and Its Associated Rows	21-4
21.2.4	Specify the Location of Each Cluster and Cluster Index Rows	21-4
21.2.5	Estimate Cluster Size and Set Storage Parameters	21-5
21.3	Creating Clusters and Objects That Use Them	21-5
21.3.1	Creating Clusters	21-5
21.3.2	Creating Clustered Tables	21-6
21.3.3	Creating Cluster Indexes	21-7
21.4	Altering Clusters and Objects That Use Them	21-7
21.4.1	Altering Clusters	21-7
21.4.2	Altering Clustered Tables	21-8
21.4.3	Altering Cluster Indexes	21-8
21.5	Dropping Clusters and Objects That Use Them	21-9
21.5.1	Dropping Clusters	21-9
21.5.2	Dropping Clustered Tables	21-10
21.5.3	Dropping Cluster Indexes	21-10
21.6	Clusters Data Dictionary Views	21-11

22 Managing Hash Clusters

22.1	About Hash Clusters	22-1
22.2	When to Use Hash Clusters	22-2
22.2.1	Situations Where Hashing Is Useful	22-2
22.2.2	Situations Where Hashing Is Not Advantageous	22-3
22.3	Creating Different Types of Hash Clusters	22-3
22.3.1	Creating Hash Clusters	22-4
22.3.2	Creating a Sorted Hash Cluster	22-4
22.3.3	Creating Single-Table Hash Clusters	22-7
22.3.4	Controlling Space Use Within a Hash Cluster	22-7
22.3.4.1	Choosing the Key	22-8
22.3.4.2	Setting HASH IS	22-8
22.3.4.3	Setting SIZE	22-8
22.3.4.4	Setting HASHKEYS	22-9
22.3.4.5	Controlling Space in Hash Clusters	22-9
22.3.5	Estimating Size Required by Hash Clusters	22-10
22.4	Altering Hash Clusters	22-11
22.5	Dropping Hash Clusters	22-11
22.6	Hash Clusters Data Dictionary Views	22-11

23 Managing Views, Sequences, and Synonyms

23.1	Managing Views	23-1
------	----------------	------

23.1.1	About Views	23-2
23.1.2	Creating Views and Join Views	23-2
23.1.2.1	Creating Views	23-3
23.1.2.2	Creating Join Views	23-4
23.1.2.3	Expansion of Defining Queries at View Creation Time	23-4
23.1.2.4	Creating Views with Errors	23-5
23.1.3	Replacing Views	23-5
23.1.4	Using Views in Queries	23-6
23.1.5	DML Statements and Join Views	23-7
23.1.5.1	Updating a Join View	23-7
23.1.5.2	Key-Preserved Tables	23-9
23.1.5.3	Rules for DML Statements and Join Views	23-10
23.1.5.4	Updating Views That Involve Outer Joins	23-13
23.1.5.5	Using the UPDATABLE_ COLUMNS Views	23-14
23.1.6	Altering Views	23-15
23.1.7	Dropping Views	23-15
23.2	Managing Sequences	23-16
23.2.1	About Sequences	23-16
23.2.2	Creating Sequences	23-17
23.2.3	Altering Sequences	23-17
23.2.4	Using Sequences	23-18
23.2.4.1	Referencing a Sequence	23-19
23.2.4.2	Caching Sequence Numbers	23-20
23.2.4.3	Making a Sequence Scalable	23-23
23.2.5	Dropping Sequences	23-25
23.3	Managing Synonyms	23-25
23.3.1	About Synonyms	23-26
23.3.2	Creating Synonyms	23-26
23.3.3	Using Synonyms in DML Statements	23-27
23.3.4	Dropping Synonyms	23-27
23.4	Views, Synonyms, and Sequences Data Dictionary Views	23-28

24 Repairing Corrupted Data

24.1	Options for Repairing Data Block Corruption	24-1
24.2	About the DBMS_REPAIR Package	24-2
24.2.1	DBMS_REPAIR Procedures	24-2
24.2.2	Limitations and Restrictions for DBMS_REPAIR Procedures	24-3
24.3	Using the DBMS_REPAIR Package	24-3
24.3.1	Task 1: Detect and Report Corruptions	24-3
24.3.1.1	About Detecting and Reporting Corruptions	24-4

24.3.1.2	DBMS_REPAIR: Using the CHECK_OBJECT and ADMIN_TABLES Procedures	24-4
24.3.1.3	DB_VERIFY: Performing an Offline Database Check	24-4
24.3.1.4	ANALYZE: Reporting Corruption	24-5
24.3.1.5	DB_BLOCK_CHECKING Initialization Parameter	24-5
24.3.2	Task 2: Evaluate the Costs and Benefits of Using DBMS_REPAIR	24-6
24.3.3	Task 3: Make Objects Usable	24-6
24.3.3.1	Corruption Repair: Using the FIX_CORRUPT_BLOCKS and SKIP_CORRUPT_BLOCKS Procedures	24-7
24.3.3.2	Implications When Skipping Corrupt Blocks	24-7
24.3.4	Task 4: Repair Corruptions and Rebuild Lost Data	24-7
24.3.4.1	Recover Data Using the DUMP_ORPHAN_KEYS Procedures	24-7
24.3.4.2	Fix Segment Bitmaps Using the SEGMENT_FIX_STATUS Procedure	24-7
24.4	DBMS_REPAIR Examples	24-8
24.4.1	Examples: Building a Repair Table or Orphan Key Table	24-8
24.4.1.1	About Repair Tables or Orphan Key Tables	24-8
24.4.1.2	Example: Creating a Repair Table	24-9
24.4.1.3	Example: Creating an Orphan Key Table	24-9
24.4.2	Example: Detecting Corruption	24-10
24.4.3	Example: Fixing Corrupt Blocks	24-11
24.4.4	Example: Finding Index Entries Pointing to Corrupt Data Blocks	24-11
24.4.5	Example: Skipping Corrupt Blocks	24-12

Part IV Database Resource Management and Task Scheduling

25 Managing Automated Database Maintenance Tasks

25.1	About Automated Maintenance Tasks	25-2
25.2	About Maintenance Windows	25-3
25.3	Configuring Automated Maintenance Tasks	25-4
25.3.1	Enabling and Disabling Maintenance Tasks for all Maintenance Windows	25-4
25.3.2	Enabling and Disabling Maintenance Tasks for Specific Maintenance Windows	25-5
25.4	Configuring Maintenance Windows	25-5
25.4.1	Modifying a Maintenance Window	25-5
25.4.2	Creating a New Maintenance Window	25-6
25.4.3	Removing a Maintenance Window	25-7
25.5	Configuring Resource Allocations for Automated Maintenance Tasks	25-7
25.5.1	About Resource Allocations for Automated Maintenance Tasks	25-8
25.5.2	Changing Resource Allocations for Automated Maintenance Tasks	25-9
25.6	Automated Maintenance Tasks Reference	25-9
25.6.1	Predefined Maintenance Windows	25-9

26 Managing Resources with Oracle Database Resource Manager

26.1	About Oracle Database Resource Manager	26-2
26.1.1	CDB and PDB Resource Management	26-3
26.1.2	Purpose of Resource Management	26-4
26.1.2.1	Purpose of Resource Management for a CDB	26-4
26.1.2.2	Purpose of Resource Management for PDBs	26-6
26.1.3	Consumer Groups, Plans, and Plan Directives	26-7
26.1.3.1	About the Elements of Resource Manager	26-7
26.1.3.2	About Resource Consumer Groups	26-8
26.1.3.3	About Resource Plan Directives	26-9
26.1.3.4	About Resource Plans	26-20
26.1.3.5	About Subplans	26-29
26.1.4	User Interface for PDB Resource Management	26-30
26.1.4.1	About Resource Manager Administration Privileges	26-31
26.1.4.2	DBMS_RESOURCE_MANAGER for CDBs and PDBs	26-32
26.1.4.3	Initialization Parameters for PDB-Level Resources	26-32
26.2	Enabling Oracle Database Resource Manager and Switching Plans	26-39
26.3	Assigning Sessions to Resource Consumer Groups	26-42
26.3.1	Overview of Assigning Sessions to Resource Consumer Groups	26-42
26.3.2	Assigning an Initial Resource Consumer Group	26-43
26.3.3	Specifying Session-to-Consumer Group Mapping Rules	26-43
26.3.3.1	About Session-to-Consumer Group Mapping Rules	26-43
26.3.3.2	Creating Consumer Group Mapping Rules	26-44
26.3.3.3	Modifying and Deleting Consumer Group Mapping Rules	26-46
26.3.3.4	Creating Mapping Rule Priorities	26-46
26.3.4	Switching Resource Consumer Groups	26-47
26.3.4.1	Manually Switching Resource Consumer Groups	26-48
26.3.4.2	Enabling Users or Applications to Manually Switch Consumer Groups	26-49
26.3.5	Specifying Automatic Consumer Group Switching	26-50
26.3.5.1	Specifying Automatic Switching with Mapping Rules	26-50
26.3.5.2	Specifying Automatic Switching by Setting Resource Limits	26-51
26.3.6	Granting and Revoking the Switch Privilege	26-54
26.3.6.1	About Granting and Revoking the Switch Privilege	26-54
26.3.6.2	Granting the Switch Privilege	26-55
26.3.6.3	Revoking Switch Privileges	26-55
26.4	Managing Resource Plans	26-56
26.4.1	Managing CDB Resource Plans	26-56
26.4.1.1	Creating a CDB Resource Plan for Managing PDBs	26-57
26.4.1.2	Creating a CDB Resource Plan for Managing PDBs: Scenario	26-57

26.4.1.3	Creating a CDB Resource Plan with PDB Performance Profiles	26-60
26.4.1.4	Creating a CDB Resource Plan for PDB Performance Profiles: Scenario	26-61
26.4.1.5	Enabling a CDB Resource Plan	26-64
26.4.1.6	Modifying a CDB Resource Plan	26-64
26.4.1.7	Disabling a CDB Resource Plan	26-76
26.4.1.8	Viewing Information About Plans and Directives in a CDB	26-77
26.4.2	Managing PDB Resource Plans	26-79
26.4.2.1	Creating a PDB Resource Plan	26-79
26.4.2.2	Enabling a PDB Resource Plan	26-80
26.4.2.3	Modifying a PDB Resource Plan	26-81
26.4.2.4	Disabling a PDB Resource Plan	26-82
26.4.3	Creating a Simple Resource Plan	26-82
26.4.4	Creating a Complex Resource Plan	26-84
26.4.4.1	About the Pending Area	26-85
26.4.4.2	Creating a Pending Area	26-86
26.4.4.3	Creating Resource Consumer Groups	26-86
26.4.4.4	Mapping Sessions to Consumer Groups	26-87
26.4.4.5	Creating a Resource Plan	26-88
26.4.4.6	Creating Resource Plan Directives	26-89
26.4.4.7	Validating the Pending Area	26-94
26.4.4.8	Submitting the Pending Area	26-95
26.4.4.9	Clearing the Pending Area	26-96
26.5	Putting It All Together: Oracle Database Resource Manager Examples	26-96
26.5.1	Multilevel Plan Example	26-97
26.5.2	Examples of Using the Utilization Limit Attribute	26-99
26.5.3	Example of Using Several Resource Allocation Methods	26-104
26.5.4	Example of Managing Parallel Statements Using Directive Attributes	26-104
26.5.5	An Oracle-Supplied Mixed Workload Plan	26-107
26.6	Managing Multiple Database Instances on a Single Server	26-108
26.6.1	About Instance Caging	26-108
26.6.2	Enabling Instance Caging	26-109
26.7	Maintaining Consumer Groups, Plans, and Directives	26-109
26.7.1	Updating a Consumer Group	26-110
26.7.2	Deleting a Consumer Group	26-110
26.7.3	Updating a Plan	26-111
26.7.4	Deleting a Plan	26-111
26.7.5	Updating a Resource Plan Directive	26-112
26.7.6	Deleting a Resource Plan Directive	26-113
26.8	Viewing Database Resource Manager Configuration and Status	26-113
26.8.1	About Resource Manager Views	26-114
26.8.2	Viewing Consumer Groups Granted to Users or Roles	26-120
26.8.3	Viewing Plan Information	26-120

26.8.4	Viewing Current Consumer Groups for Sessions	26-121
26.8.5	Viewing the Currently Active Plans	26-121
26.8.6	Monitoring PDBs Managed by Oracle Database Resource Manager	26-121
26.8.6.1	About Resource Manager Views for PDBs	26-122
26.8.6.2	Monitoring CPU Usage for PDBs	26-123
26.8.6.3	Monitoring Parallel Execution for PDBs	26-124
26.8.6.4	Monitoring the I/O Generated by PDBs	26-125
26.8.6.5	Monitoring Memory Usage for PDBs	26-126
26.9	Interacting with Operating-System Resource Control	26-126
26.9.1	Guidelines for Using Operating-System Resource Control	26-126
26.10	Oracle Database Resource Manager Reference	26-127
26.10.1	Predefined Resource Plans and Consumer Groups	26-127
26.10.2	Predefined Consumer Group Mapping Rules	26-129
26.10.3	Resource Manager Data Dictionary Views	26-130
26.11	Operating System CPU Resource Management	26-131

27 Oracle Scheduler Concepts

27.1	Overview of Oracle Scheduler	27-1
27.2	Jobs and Supporting Scheduler Objects	27-3
27.2.1	About Jobs and Supporting Scheduler Objects	27-4
27.2.2	Programs	27-4
27.2.3	Schedules	27-5
27.2.4	Jobs	27-5
27.2.4.1	About Jobs	27-6
27.2.4.2	Specifying a Job Action	27-6
27.2.4.3	Specifying a Job Schedule	27-6
27.2.4.4	Specifying a Job Destination	27-7
27.2.4.5	Specifying a Job Credential	27-7
27.2.5	Destinations	27-8
27.2.5.1	About Destinations	27-8
27.2.5.2	About Destinations and Scheduler Agents	27-9
27.2.6	File Watchers	27-10
27.2.7	Credentials	27-10
27.2.8	Chains	27-11
27.2.9	Job Classes	27-12
27.2.10	Windows	27-14
27.2.10.1	About Windows	27-14
27.2.10.2	Overlapping Windows	27-15
27.2.11	Groups	27-18
27.2.11.1	About Groups	27-18
27.2.11.2	Destination Groups	27-18

27.2.11.3	Window Groups	27-18
27.2.12	Incompatibilities	27-19
27.3	More About Jobs	27-19
27.3.1	Job Categories	27-20
27.3.1.1	Database Jobs	27-21
27.3.1.2	External Jobs	27-22
27.3.1.3	Multiple-Destination Jobs	27-24
27.3.1.4	Chain Jobs	27-26
27.3.1.5	Detached Jobs	27-26
27.3.1.6	Lightweight Jobs	27-27
27.3.1.7	In-Memory Jobs	27-28
27.3.1.8	Script Jobs	27-28
27.3.2	Job Instances	27-29
27.3.3	Job Arguments	27-30
27.3.4	How Programs, Jobs, and Schedules are Related	27-30
27.4	Scheduler Architecture	27-31
27.4.1	Scheduler Components	27-32
27.4.2	The Job Table	27-32
27.4.3	The Job Coordinator	27-32
27.4.3.1	About The Job Coordinator	27-32
27.4.3.2	Job Coordinator Actions	27-33
27.4.3.3	Maximum Number of Scheduler Job Processes	27-33
27.4.4	How Jobs Execute	27-34
27.4.5	After Jobs Complete	27-35
27.4.6	Using the Scheduler in Real Application Clusters Environments	27-35
27.4.6.1	The Scheduler and Real Application Clusters	27-35
27.4.6.2	Service Affinity when Using the Scheduler	27-36
27.5	Processes to Close a PDB	27-37
27.6	Scheduler Support for Oracle Data Guard	27-37

28 Scheduling Jobs with Oracle Scheduler

28.1	About Scheduler Objects and Their Naming	28-2
28.2	Creating, Running, and Managing Jobs	28-2
28.2.1	Job Tasks and Their Procedures	28-3
28.2.2	Creating Jobs	28-4
28.2.2.1	Overview of Creating Jobs	28-4
28.2.2.2	Specifying Job Actions, Schedules, Programs, and Styles	28-6
28.2.2.3	Specifying Scheduler Job Credentials	28-9
28.2.2.4	Specifying Destinations	28-10
28.2.2.5	Creating Multiple-Destination Jobs	28-13
28.2.2.6	Setting Job Arguments	28-14

28.2.2.7	Setting Additional Job Attributes	28-15
28.2.2.8	Creating Detached Jobs	28-15
28.2.2.9	Creating Multiple Jobs in a Single Transaction	28-16
28.2.2.10	Techniques for External Jobs	28-17
28.2.3	Altering Jobs	28-19
28.2.4	Running Jobs	28-20
28.2.5	Stopping Jobs	28-20
28.2.6	Stopping External Jobs	28-21
28.2.7	Stopping a Chain Job	28-22
28.2.8	Dropping Jobs	28-22
28.2.9	Dropping Running Jobs	28-22
28.2.10	Dropping Multiple Jobs	28-22
28.2.11	Disabling Jobs	28-23
28.2.12	Enabling Jobs	28-24
28.2.13	Copying Jobs	28-25
28.3	Creating and Managing Programs to Define Jobs	28-25
28.3.1	Program Tasks and Their Procedures	28-25
28.3.2	Creating Programs with Scheduler	28-26
28.3.2.1	Creating Programs	28-26
28.3.2.2	Defining Program Arguments	28-26
28.3.3	Altering Programs	28-28
28.3.4	Dropping Programs	28-28
28.3.5	Disabling Programs	28-28
28.3.6	Enabling Programs	28-29
28.4	Creating and Managing Schedules to Define Jobs	28-29
28.4.1	Schedule Tasks and Their Procedures	28-30
28.4.2	Creating Schedules	28-30
28.4.3	Altering Schedules	28-31
28.4.4	Dropping Schedules	28-31
28.4.5	Setting the Repeat Interval	28-31
28.4.5.1	About Setting the Repeat Interval	28-32
28.4.5.2	Using the Scheduler Calendaring Syntax	28-32
28.4.5.3	Using a PL/SQL Expression	28-35
28.4.5.4	Differences Between PL/SQL Expression and Calendaring Syntax Behavior	28-35
28.4.5.5	Repeat Intervals and Daylight Savings	28-36
28.5	Using Events to Start Jobs	28-36
28.5.1	About Events	28-37
28.5.2	Starting Jobs with Events Raised by Your Application	28-37
28.5.2.1	About Events Raised by Your Application	28-38
28.5.2.2	Creating an Event-Based Job	28-39
28.5.2.3	Altering an Event-Based Job	28-40

28.5.2.4	Creating an Event Schedule	28-40
28.5.2.5	Altering an Event Schedule	28-41
28.5.2.6	Passing Event Messages into an Event-Based Job	28-41
28.5.3	Starting a Job When a File Arrives on a System	28-42
28.5.3.1	About File Watchers	28-43
28.5.3.2	Enabling File Arrival Events from Remote Systems	28-44
28.5.3.3	Creating File Watchers and File Watcher Jobs	28-44
28.5.3.4	File Arrival Example	28-47
28.5.3.5	Managing File Watchers	28-48
28.5.3.6	Viewing File Watcher Information	28-49
28.6	Creating and Managing Job Chains	28-50
28.6.1	About Creating and Managing Job Chains	28-51
28.6.2	Chain Tasks and Their Procedures	28-52
28.6.3	Creating Chains	28-53
28.6.4	Defining Chain Steps	28-53
28.6.5	Adding Rules to a Chain	28-55
28.6.6	Setting an Evaluation Interval for Chain Rules	28-58
28.6.7	Enabling Chains	28-58
28.6.8	Creating Jobs for Chains	28-59
28.6.9	Dropping Chains	28-60
28.6.10	Running Chains	28-60
28.6.11	Dropping Chain Rules	28-61
28.6.12	Disabling Chains	28-61
28.6.13	Dropping Chain Steps	28-61
28.6.14	Stopping Chains	28-62
28.6.15	Stopping Individual Chain Steps	28-62
28.6.16	Pausing Chains	28-62
28.6.17	Skipping Chain Steps	28-63
28.6.18	Running Part of a Chain	28-64
28.6.19	Monitoring Running Chains	28-64
28.6.20	Handling Stalled Chains	28-65
28.7	Using Incompatibility Definitions	28-65
28.7.1	Creating a Job or Program Incompatibility	28-66
28.7.2	Adding a Job or Program to an Incompatibility	28-66
28.7.3	Removing a Job or Program from an Incompatibility	28-67
28.7.4	Dropping an Incompatibility	28-67
28.8	Managing Job Resources	28-68
28.8.1	Creating or Dropping a Resource	28-68
28.8.2	Altering a Resource	28-69
28.8.3	Setting a Resource Constraint for a Job	28-69
28.9	Prioritizing Jobs	28-70
28.9.1	Managing Job Priorities with Job Classes	28-71

28.9.1.1	Job Class Tasks and Their Procedures	28-71
28.9.1.2	Creating Job Classes	28-72
28.9.1.3	Altering Job Classes	28-72
28.9.1.4	Dropping Job Classes	28-72
28.9.2	Setting Relative Job Priorities Within a Job Class	28-73
28.9.3	Managing Job Scheduling and Job Priorities with Windows	28-73
28.9.3.1	About Job Scheduling and Job Priorities with Windows	28-74
28.9.3.2	Window Tasks and Their Procedures	28-75
28.9.3.3	Creating Windows	28-75
28.9.3.4	Altering Windows	28-76
28.9.3.5	Opening Windows	28-76
28.9.3.6	Closing Windows	28-77
28.9.3.7	Dropping Windows	28-78
28.9.3.8	Disabling Windows	28-78
28.9.3.9	Enabling Windows	28-79
28.9.4	Managing Job Scheduling and Job Priorities with Window Groups	28-79
28.9.4.1	Window Group Tasks and Their Procedures	28-80
28.9.4.2	Creating Window Groups	28-80
28.9.4.3	Dropping Window Groups	28-81
28.9.4.4	Adding a Member to a Window Group	28-81
28.9.4.5	Removing a Member from a Window Group	28-81
28.9.4.6	Enabling a Window Group	28-82
28.9.4.7	Disabling a Window Group	28-82
28.9.5	Allocating Resources Among Jobs Using Resource Manager	28-82
28.9.6	Example of Resource Allocation for Jobs	28-83
28.10	Monitoring Jobs	28-84
28.10.1	About Monitoring Jobs	28-84
28.10.2	The Job Log	28-84
28.10.2.1	Viewing the Job Log	28-85
28.10.2.2	Run Details	28-86
28.10.2.3	Precedence of Logging Levels in Jobs and Job Classes	28-86
28.10.3	Monitoring Multiple Destination Jobs	28-87
28.10.4	Monitoring Job State with Events Raised by the Scheduler	28-88
28.10.4.1	About Job State Events	28-88
28.10.4.2	Altering a Job to Raise Job State Events	28-90
28.10.4.3	Consuming Job State Events with your Application	28-90
28.10.5	Monitoring Job State with E-mail Notifications	28-91
28.10.5.1	About E-mail Notifications	28-91
28.10.5.2	Adding E-mail Notifications for a Job	28-92
28.10.5.3	Removing E-mail Notifications for a Job	28-93
28.10.5.4	Viewing Information About E-mail Notifications	28-94

29 Administering Oracle Scheduler

29.1	Configuring Oracle Scheduler	29-2
29.1.1	Setting Oracle Scheduler Privileges	29-2
29.1.2	Setting Scheduler Preferences	29-3
29.1.3	Using the Oracle Scheduler Agent to Run Remote Jobs	29-5
29.1.3.1	Enabling and Disabling Databases for Remote Jobs	29-6
29.1.3.2	Installing and Configuring the Scheduler Agent on a Remote Host	29-9
29.1.3.3	Performing Tasks with the Scheduler Agent	29-10
29.2	Monitoring and Managing the Scheduler	29-13
29.2.1	Viewing the Currently Active Window and Resource Plan	29-13
29.2.2	Finding Information About Currently Running Jobs	29-13
29.2.3	Monitoring and Managing Window and Job Logs	29-15
29.2.3.1	Job Log	29-15
29.2.3.2	Window Log	29-16
29.2.3.3	Purging Logs	29-17
29.2.4	DBMS_SCHEDULER In-Memory Trace	29-18
29.2.5	Managing Scheduler Security	29-20
29.3	Import/Export and the Scheduler	29-20
29.4	Troubleshooting the Scheduler	29-21
29.4.1	A Job Does Not Run	29-21
29.4.1.1	About Job States	29-21
29.4.1.2	Viewing the Job Log	29-22
29.4.1.3	Troubleshooting Remote Jobs	29-22
29.4.1.4	About Job Recovery After a Failure	29-23
29.4.2	A Program Becomes Disabled	29-24
29.4.3	A Window Fails to Take Effect	29-24
29.5	Examples of Using the Scheduler	29-24
29.5.1	Examples of Creating Job Classes	29-24
29.5.2	Examples of Setting Attributes	29-25
29.5.3	Examples of Creating Chains	29-27
29.5.4	Examples of Creating Jobs and Schedules Based on Events	29-28
29.5.5	Example of Creating a Job In an Oracle Data Guard Environment	29-29
29.6	Scheduler Reference	29-30
29.6.1	Scheduler Privileges	29-30
29.6.2	Scheduler Data Dictionary Views	29-32

30 Managing Transactions

30.1	Priority Transactions	30-1
30.1.1	Using Priority Transactions	30-2
30.1.1.1	Setting Transaction Priority	30-2

30.1.1.2	Setting System-Level Wait Targets	30-3
30.1.1.3	Acknowledging the Automatic Rollback	30-4
30.1.1.4	Setting Priority Transaction Mode	30-4
30.1.1.5	Using Priority Transaction Mode to Determine System-Level Wait Targets	30-5
30.1.2	Monitoring Priority Transactions	30-6
30.1.2.1	Statistics Incremented in ROLLBACK Mode	30-6
30.1.2.2	Statistics Incremented in TRACK Mode	30-7
30.1.3	Priority Transaction Behavior	30-7
30.1.3.1	Behavior of Priority Transactions for Distributed Transactions	30-7
30.1.3.2	Behavior for XA Transactions	30-8
30.1.4	Priority Transaction Restrictions	30-8
30.2	Automatic Transaction Quarantine	30-8
30.2.1	Monitoring Quarantined Transactions	30-11
30.2.2	Resolving Quarantined Transactions	30-12
30.2.3	Dropping Quarantined Transactions	30-13
30.2.4	Transaction Quarantine Escalation	30-13

Part V Distributed Database Management

31 Distributed Database Concepts

31.1	Distributed Database Architecture	31-1
31.1.1	Homogenous Distributed Database Systems	31-2
31.1.1.1	About Homogenous Distributed Database Systems	31-2
31.1.1.2	Distributed Databases Versus Distributed Processing	31-3
31.1.1.3	Distributed Databases Versus Replicated Databases	31-4
31.1.2	Heterogeneous Distributed Database Systems	31-4
31.1.2.1	About Heterogeneous Distributed Database Systems	31-5
31.1.2.2	Heterogeneous Services	31-5
31.1.2.3	Transparent Gateway Agents	31-5
31.1.2.4	Generic Connectivity	31-6
31.1.3	Client/Server Database Architecture	31-6
31.2	Database Links	31-7
31.2.1	What Are Database Links?	31-8
31.2.2	What Are Shared Database Links?	31-9
31.2.3	Why Use Database Links?	31-10
31.2.4	Global Database Names in Database Links	31-10
31.2.5	Global Name as a Loopback Database Link	31-12
31.2.6	Names for Database Links	31-12
31.2.7	Types of Database Links	31-13
31.2.8	Users of Database Links	31-14

31.2.8.1	Overview of Database Link Users	31-14
31.2.8.2	Connected User Database Links	31-15
31.2.8.3	Fixed User Database Links	31-16
31.2.8.4	Current User Database Links	31-16
31.2.9	Creation of Database Links: Examples	31-17
31.2.10	Schema Objects and Database Links	31-18
31.2.10.1	Naming of Schema Objects Using Database Links	31-19
31.2.10.2	Authorization for Accessing Remote Schema Objects	31-19
31.2.10.3	Synonyms for Schema Objects	31-19
31.2.10.4	Schema Object Name Resolution	31-20
31.2.11	Database Link Restrictions	31-20
31.3	Distributed Database Administration	31-21
31.3.1	Site Autonomy	31-22
31.3.2	Distributed Database Security	31-22
31.3.2.1	Authentication Through Database Links	31-23
31.3.2.2	Authentication Without Passwords	31-24
31.3.2.3	Supporting User Accounts and Roles	31-24
31.3.2.4	Centralized User and Privilege Management	31-25
31.3.2.5	Data Encryption	31-28
31.3.3	Auditing Database Links	31-29
31.3.4	Administration Tools	31-30
31.3.4.1	Cloud Control and Distributed Databases	31-30
31.3.4.2	Third-Party Administration Tools	31-30
31.3.4.3	SNMP Support	31-30
31.4	Transaction Processing in a Distributed System	31-31
31.4.1	Remote SQL Statements	31-32
31.4.2	Distributed SQL Statements	31-32
31.4.3	Shared SQL for Remote and Distributed Statements	31-33
31.4.4	Remote Transactions	31-33
31.4.5	Distributed Transactions	31-33
31.4.6	Two-Phase Commit Mechanism	31-34
31.4.7	Database Link Name Resolution	31-34
31.4.7.1	About Database Link Name Resolution	31-35
31.4.7.2	Name Resolution When the Global Database Name Is Complete	31-35
31.4.7.3	Name Resolution When the Global Database Name Is Partial	31-35
31.4.7.4	Name Resolution When No Global Database Name Is Specified	31-36
31.4.7.5	Terminating the Search for Name Resolution	31-36
31.4.8	Schema Object Name Resolution	31-37
31.4.8.1	About Schema Object Name Resolution	31-37
31.4.8.2	Example of Global Object Name Resolution: Complete Object Name	31-37
31.4.8.3	Example of Global Object Name Resolution: Partial Object Name	31-38
31.4.9	Global Name Resolution in Views, Synonyms, and Procedures	31-39

31.4.9.1	About Global Name Resolution in Views, Synonyms, and Procedures	31-39
31.4.9.2	What Happens When Global Names Change	31-40
31.4.9.3	Scenarios for Global Name Changes	31-40
31.5	Distributed Database Application Development	31-41
31.5.1	Transparency in a Distributed Database System	31-42
31.5.1.1	Location Transparency	31-42
31.5.1.2	SQL and COMMIT Transparency	31-43
31.5.2	PL/SQL and Remote Procedure Calls (RPCs)	31-43
31.5.3	Distributed Query Optimization	31-44
31.6	Character Set Support for Distributed Environments	31-44
31.6.1	About Character Set Support for Distributed Environments	31-45
31.6.2	Client/Server Environment	31-45
31.6.3	Homogeneous Distributed Environment	31-46
31.6.4	Heterogeneous Distributed Environment	31-46

32 Managing a Distributed Database

32.1	Managing Global Names in a Distributed System	32-1
32.1.1	Understanding How Global Database Names Are Formed	32-2
32.1.2	Determining Whether Global Naming Is Enforced	32-2
32.1.3	Viewing a Global Database Name	32-3
32.1.4	Changing the Domain in a Global Database Name	32-3
32.1.5	Changing a Global Database Name: Scenario	32-4
32.2	Creating Database Links	32-6
32.2.1	Obtaining Privileges Necessary for Creating Database Links	32-7
32.2.2	Specifying Link Types	32-7
32.2.2.1	Creating Private Database Links	32-8
32.2.2.2	Creating Public Database Links	32-8
32.2.2.3	Creating Global Database Links	32-9
32.2.3	Specifying Link Users	32-9
32.2.3.1	Creating Fixed User Database Links	32-10
32.2.3.2	Creating Connected User and Current User Database Links	32-10
32.2.4	Using Connection Qualifiers to Specify Service Names Within Link Names	32-11
32.3	Using Shared Database Links	32-12
32.3.1	Determining Whether to Use Shared Database Links	32-13
32.3.2	Creating Shared Database Links	32-14
32.3.3	Configuring Shared Database Links	32-14
32.3.3.1	Creating Shared Links to Dedicated Servers	32-14
32.3.3.2	Creating Shared Links to Shared Servers	32-15
32.4	Managing Database Links	32-16
32.4.1	Closing Database Links	32-17
32.4.2	Dropping Database Links	32-17

32.4.2.1	Dropping a Private Database Link	32-17
32.4.2.2	Dropping a Public Database Link	32-18
32.4.3	Limiting the Number of Active Database Link Connections	32-18
32.5	Viewing Information About Database Links	32-19
32.5.1	Determining Which Links Are in the Database	32-19
32.5.2	Determining Which Link Connections Are Open	32-20
32.5.3	Determining the Host of Outgoing Database Links	32-21
32.5.4	Determining Information About Incoming Database Links	32-22
32.5.5	Determining the Source of High SCN Activity for Incoming Database Links	32-23
32.6	Creating Location Transparency	32-24
32.6.1	Using Views to Create Location Transparency	32-24
32.6.2	Using Synonyms to Create Location Transparency	32-25
32.6.2.1	Creating Synonyms	32-26
32.6.2.2	Managing Privileges and Synonyms	32-27
32.6.3	Using Procedures to Create Location Transparency	32-27
32.6.3.1	Using Local Procedures to Reference Remote Data	32-28
32.6.3.2	Using Local Procedures to Call Remote Procedures	32-28
32.6.3.3	Using Local Synonyms to Reference Remote Procedures	32-29
32.6.3.4	Managing Procedures and Privileges	32-29
32.7	Managing Statement Transparency	32-29
32.8	Managing a Distributed Database: Examples	32-31
32.8.1	Example 1: Creating a Public Fixed User Database Link	32-31
32.8.2	Example 2: Creating a Public Fixed User Shared Database Link	32-31
32.8.3	Example 3: Creating a Public Connected User Database Link	32-32
32.8.4	Example 4: Creating a Public Connected User Shared Database Link	32-32
32.8.5	Example 5: Creating a Public Current User Database Link	32-33

33 Developing Applications for a Distributed Database System

33.1	Managing the Distribution of Application Data	33-1
33.2	Controlling Connections Established by Database Links	33-2
33.3	Maintaining Referential Integrity in a Distributed System	33-2
33.4	Tuning Distributed Queries	33-3
33.4.1	Using Collocated Inline Views	33-4
33.4.2	Using Cost-Based Optimization	33-4
33.4.2.1	How Does Cost-Based Optimization Work?	33-5
33.4.2.2	Rewriting Queries for Cost-Based Optimization	33-5
33.4.2.3	Setting Up Cost-Based Optimization	33-6
33.4.3	Using Hints	33-7
33.4.3.1	About Using Hints	33-7
33.4.3.2	Using the NO_MERGE Hint	33-8
33.4.3.3	Using the DRIVING_SITE Hint	33-8

33.4.4	Analyzing the Execution Plan	33-8
33.4.4.1	Generating the Execution Plan	33-9
33.4.4.2	Viewing the Execution Plan	33-9
33.5	Handling Errors in Remote Procedures	33-10

34 Distributed Transactions Concepts

34.1	What Are Distributed Transactions?	34-1
34.1.1	DML and DDL Transactions	34-2
34.1.2	Transaction Control Statements	34-3
34.2	Session Trees for Distributed Transactions	34-3
34.2.1	About Session Trees for Distributed Transactions	34-3
34.2.2	Clients	34-4
34.2.3	Database Servers	34-5
34.2.4	Local Coordinators	34-5
34.2.5	Global Coordinator	34-5
34.2.6	Commit Point Site	34-5
34.2.6.1	About the Commit Point Site	34-6
34.2.6.2	How a Distributed Transaction Commits	34-7
34.2.6.3	Commit Point Strength	34-7
34.3	Two-Phase Commit Mechanism	34-9
34.3.1	About the Two-Phase Commit Mechanism	34-9
34.3.2	Prepare Phase	34-10
34.3.2.1	About Prepare Phase	34-10
34.3.2.2	Types of Responses in the Prepare Phase	34-10
34.3.2.3	Steps in the Prepare Phase	34-12
34.3.3	Commit Phase	34-12
34.3.3.1	Steps in the Commit Phase	34-13
34.3.3.2	Guaranteeing Global Database Consistency	34-13
34.3.4	Forget Phase	34-14
34.4	In-Doubt Transactions	34-14
34.4.1	About In-Doubt Transactions	34-14
34.4.2	Automatic Resolution of In-Doubt Transactions	34-15
34.4.2.1	Failure During the Prepare Phase	34-15
34.4.2.2	Failure During the Commit Phase	34-16
34.4.3	Manual Resolution of In-Doubt Transactions	34-17
34.4.4	Relevance of System Change Numbers for In-Doubt Transactions	34-17
34.5	Distributed Transaction Processing: Case Study	34-18
34.5.1	About the Distributed Transaction Processing Case Study	34-18
34.5.2	Stage 1: Client Application Issues DML Statements	34-18
34.5.3	Stage 2: Oracle Database Determines Commit Point Site	34-19
34.5.4	Stage 3: Global Coordinator Sends Prepare Response	34-20

34.5.5	Stage 4: Commit Point Site Commits	34-21
34.5.6	Stage 5: Commit Point Site Informs Global Coordinator of Commit	34-21
34.5.7	Stage 6: Global and Local Coordinators Tell All Nodes to Commit	34-21
34.5.8	Stage 7: Global Coordinator and Commit Point Site Complete the Commit	34-22

35 Managing Distributed Transactions

35.1	Specifying the Commit Point Strength of a Node	35-2
35.2	Naming Transactions	35-2
35.3	Viewing Information About Distributed Transactions	35-3
35.3.1	Determining the ID Number and Status of Prepared Transactions	35-3
35.3.2	Tracing the Session Tree of In-Doubt Transactions	35-5
35.4	Deciding How to Handle In-Doubt Transactions	35-6
35.4.1	Discovering Problems with a Two-Phase Commit	35-6
35.4.2	Determining Whether to Perform a Manual Override	35-7
35.4.3	Analyzing the Transaction Data	35-7
35.4.3.1	Find a Node that Committed or Rolled Back	35-8
35.4.3.2	Look for Transaction Comments	35-8
35.4.3.3	Look for Transaction Advice	35-8
35.5	Manually Overriding In-Doubt Transactions	35-9
35.5.1	Manually Committing an In-Doubt Transaction	35-9
35.5.1.1	Privileges Required to Commit an In-Doubt Transaction	35-9
35.5.1.2	Committing Using Only the Transaction ID	35-9
35.5.1.3	Committing Using an SCN	35-10
35.5.2	Manually Rolling Back an In-Doubt Transaction	35-10
35.6	Purging Pending Rows from the Data Dictionary	35-11
35.6.1	About Purging Pending Rows from the Data Dictionary	35-11
35.6.2	Executing the PURGE_LOST_DB_ENTRY Procedure	35-12
35.6.3	Determining When to Use DBMS_TRANSACTION	35-12
35.7	Manually Committing an In-Doubt Transaction: Example	35-13
35.7.1	Step 1: Record User Feedback	35-14
35.7.2	Step 2: Query DBA_2PC_PENDING	35-14
35.7.2.1	Determining the Global Transaction ID	35-15
35.7.2.2	Determining the State of the Transaction	35-15
35.7.2.3	Looking for Comments or Advice	35-15
35.7.3	Step 3: Query DBA_2PC_NEIGHBORS on Local Node	35-16
35.7.3.1	Obtaining Database Role and Database Link Information	35-16
35.7.3.2	Determining the Commit Point Site	35-17
35.7.4	Step 4: Querying Data Dictionary Views on All Nodes	35-17
35.7.4.1	Checking the Status of Pending Transactions at sales	35-17
35.7.4.2	Determining the Coordinators and Commit Point Site at sales	35-18
35.7.4.3	Checking the Status of Pending Transactions at HQ	35-19

35.7.5	Step 5: Commit the In-Doubt Transaction	35-19
35.7.6	Step 6: Check for Mixed Outcome Using DBA_2PC_PENDING	35-20
35.8	Data Access Failures Due to Locks	35-20
35.8.1	Transaction Timeouts	35-20
35.8.2	Locks from In-Doubt Transactions	35-21
35.9	Simulating Distributed Transaction Failure	35-21
35.9.1	Forcing a Distributed Transaction to Fail	35-21
35.9.2	Disabling and Enabling RECO	35-22
35.10	Managing Read Consistency	35-23
35.11	Enhancing Distributed Transaction Security	35-23

Part VI Managing Read-Only Materialized Views

36 Read-Only Materialized View Concepts

36.1	Replication Databases	36-1
36.2	Read-Only Materialized Views	36-2
36.3	The Uses of Materialized Views	36-3
36.3.1	Ease Network Loads	36-3
36.3.2	Enable Data Subsetting	36-3
36.3.3	Enable Disconnected Computing	36-3
36.4	Available Materialized Views	36-4
36.4.1	About the Available Materialized Views	36-4
36.4.2	Primary Key Materialized Views	36-4
36.4.3	Object Materialized Views	36-5
36.4.4	ROWID Materialized Views	36-6
36.4.5	Complex Materialized Views	36-6
36.4.5.1	About Complex Materialized Views	36-6
36.4.5.2	A Comparison of Simple and Complex Materialized Views	36-8
36.5	Users and Privileges Related to Materialized Views	36-9
36.5.1	Required Privileges for Materialized View Operations	36-10
36.5.2	Creator Is Owner	36-11
36.5.3	Creator Is Not Owner	36-11
36.5.4	Refresher Is Owner	36-11
36.5.5	Refresher Is Not Owner	36-12
36.6	Data Subsetting with Materialized Views	36-12
36.6.1	About Data Subsetting with Materialized Views	36-12
36.6.2	Materialized Views with Subqueries	36-13
36.6.2.1	Many to One Subqueries	36-13
36.6.2.2	One to Many Subqueries	36-14
36.6.2.3	Many to Many Subqueries	36-15

36.6.2.4	Materialized Views with Subqueries and Unions	36-16
36.6.3	Restrictions for Materialized Views with Subqueries	36-19
36.6.4	Restrictions for Materialized Views with Unions Containing Subqueries	36-20
36.6.4.1	Examples of Materialized Views with Unions Containing Subqueries	36-20
36.7	Materialized View Refresh	36-21
36.8	Refresh Groups	36-21
36.9	Materialized View Log	36-21
36.10	Materialized Views and User-Defined Data Types	36-22
36.10.1	How Materialized Views Work with Object Types and Collections	36-22
36.10.2	Type Agreement at Replication Databases	36-23
36.10.3	Column Subsetting of Masters with Column Objects	36-24
36.10.4	Materialized Views Based on Object Tables	36-25
36.10.4.1	About Materialized Views Based on Object Tables	36-25
36.10.4.2	Materialized Views Based on Object Tables Created Without Using the OF type Clause	36-25
36.10.4.3	OID Preservation in Object Materialized Views	36-26
36.10.5	Materialized Views with Collection Columns	36-26
36.10.5.1	Restrictions for Materialized Views with Collection Columns	36-28
36.10.6	Materialized Views with REF Columns	36-28
36.10.6.1	About Materialized Views with REF Columns	36-28
36.10.6.2	Scoped REF Columns	36-29
36.10.6.3	Unscoped REF Columns	36-30
36.10.6.4	Logging REF Columns in the Materialized View Log	36-30
36.10.6.5	REFs Created Using the WITH ROWID Clause	36-30
36.11	Materialized View Registration at a Master Database	36-30
36.11.1	Viewing Information about Registered Materialized Views	36-31
36.11.2	Internal Mechanisms	36-31
36.11.3	Manual Materialized View Registration	36-31

37 Read-Only Materialized View Architecture

37.1	Master Database Mechanisms	37-1
37.1.1	Master Database Objects	37-1
37.1.2	Master Table	37-2
37.1.3	Internal Trigger for the Materialized View Log	37-2
37.1.4	Materialized View Logs	37-2
37.1.4.1	About Materialized View Logs	37-3
37.1.4.2	Columns Logged in the Materialized View Log	37-4
37.1.4.3	Restriction on Import of Materialized View Logs to a Different Schema	37-5
37.2	Materialized View Database Mechanisms	37-5
37.2.1	Indexes for Materialized Views	37-6
37.3	Organizational Mechanisms	37-6

37.3.1	Refresh Groups	37-6
37.3.2	Refresh Group Size	37-6
37.4	Refresh Process	37-7
37.4.1	About the Refresh Process	37-7
37.4.2	Refresh Types	37-7
37.4.2.1	Complete Refresh	37-8
37.4.2.2	Fast Refresh	37-8
37.4.2.3	Force Refresh	37-9
37.4.3	Initiating a Refresh	37-9
37.4.3.1	Scheduled Refresh	37-9
37.4.3.2	On-Demand Refresh	37-10
37.4.4	Constraints and Refresh	37-11

38 Planning for Read-Only Materialized Views

38.1	Considerations for Master Tables	38-1
38.1.1	Primary Keys and Master Tables	38-1
38.1.2	Foreign Keys and Master Tables	38-1
38.1.3	Data Type Considerations for Master Tables	38-1
38.1.4	Unsupported Table Types	38-3
38.2	Planning for Master Databases and Materialized View Databases	38-3
38.2.1	Characteristics of Master Databases and Materialized View Databases	38-4
38.2.2	Advantages of Master Databases	38-4
38.2.3	Advantages of Materialized View Databases	38-4
38.2.4	Preparing for Materialized Views	38-4
38.2.4.1	Required Schemas at Materialized View Database	38-5
38.2.4.2	Required Database Links for Materialized Views	38-5
38.2.4.3	Required Privileges	38-7
38.2.4.4	Sufficient Job Processes	38-7
38.2.5	Creating Materialized View Logs	38-8
38.2.6	Logging Columns in a Materialized View Log	38-9

39 Creating and Managing Read-Only Materialized Views

39.1	Creating Read-Only Materialized Views	39-1
39.2	Creating Refresh Groups	39-3
39.3	Refreshing Materialized Views	39-4
39.4	Determining the Fast Refresh Capabilities of a Materialized View	39-5
39.5	Adding a New Materialized View Database	39-6
39.6	Monitoring Materialized View Logs	39-7
39.6.1	Listing Information About the Materialized View Logs at a Master Database	39-7
39.6.2	Listing the Materialized Views that Use a Materialized View Log	39-8

39.7	Monitoring Materialized Views	39-9
39.7.1	Listing Information About Materialized Views	39-9
39.7.1.1	Listing Master Database Information For Materialized Views	39-9
39.7.1.2	Listing the Properties of Materialized Views	39-10
39.7.2	Listing Information About the Refresh Groups at a Materialized View Database	39-11
39.7.3	Determining the Job ID for Each Refresh Job at a Materialized View Database	39-12
39.7.4	Determining Which Materialized Views Are Currently Refreshing	39-12

40 Troubleshooting Problems with Read-Only Materialized Views

40.1	Diagnosing Problems with Database Links	40-1
40.2	Problems Creating Materialized Views	40-2
40.3	Refresh Problems	40-2
40.3.1	Common Refresh Problems	40-2
40.3.2	Automatic Refresh Retries	40-3
40.3.3	Fast Refresh Errors at New Materialized View Databases	40-3
40.3.4	Materialized Views Continually Refreshing	40-4
40.3.5	Materialized View Logs Growing Too Large	40-4
40.4	Advanced Troubleshooting of Refresh Problems	40-4

Part VII Appendixes

A Support for DBMS_JOB

A.1	Oracle Scheduler Replaces DBMS_JOB	A-1
A.1.1	Configuring DBMS_JOB	A-2
A.1.2	Using Both DBMS_JOB and Oracle Scheduler	A-2
A.2	Moving from DBMS_JOB to Oracle Scheduler	A-2
A.2.1	Creating a Job	A-3
A.2.2	Altering a Job	A-3
A.2.3	Removing a Job from the Job Queue	A-4

B Blockchain Tables Reference

B.1	Blockchain Tables Column Content	B-1
B.2	Blockchain Tables Row Content	B-3
B.3	Format of the Signed Digest in Blockchain Tables	B-4

Index
