Managing Views, Sequences, and Synonyms

You can create and manage views, sequences, and synonyms with Oracle Database.

Managing Views

You can perform tasks such as creating views, replacing views, altering views, and dropping views.

Managing Sequences

You can perform tasks such as creating sequences, altering sequences, using sequences, and dropping sequences.

Managing Synonyms

You can perform tasks such as creating synonyms, using synonyms, and dropping synonyms.

Views, Synonyms, and Sequences Data Dictionary Views
 You can query data dictionary views for information about views, synonyms, and sequences.

23.1 Managing Views

You can perform tasks such as creating views, replacing views, altering views, and dropping views.

Live SQL:

To view and run examples related to managing views on Oracle Live SQL, go to *Oracle Live SQL: Creating, Replacing, and Dropping a View.*

About Views

A **view** is a logical representation of a table or combination of tables. In essence, a view is a stored query.

Creating Views and Join Views

You can create views using the CREATE VIEW statement. Each view is defined by a query that references tables, materialized views, or other views. You can also create join views that specify multiple base tables or views in the FROM clause.

Replacing Views

You can replace a view by dropping it and re-creating it or by issuing a CREATE VIEW statement that contains the OR REPLACE clause.

Using Views in Queries

You can query a view. You can also perform data manipulation language (DML) operations on views, with some restrictions.

DML Statements and Join Views

Restrictions apply when issuing DML statements on join views.

Altering Views

You use the ALTER VIEW statement only to explicitly recompile a view that is invalid.

Dropping Views

You can drop a view with the DROP VIEW statement.

23.1.1 About Views

A **view** is a logical representation of a table or combination of tables. In essence, a view is a stored query.

A view derives its data from the tables on which it is based. These tables are called **base tables**. Base tables might in turn be actual tables or might be views themselves. All operations performed on a view actually affect the base table of the view. You can use views in almost the same way as tables. You can query, update, insert into, and delete from views, just as you can standard tables.

Views can provide a different representation (such as subsets or supersets) of the data that resides within other tables and views. Views are very powerful because they allow you to tailor the presentation of data to different types of users.

Note:

One special type of view is the editioning view, which is used only to support online upgrade of applications using edition-based redefinition. The remainder of this section on managing views describes all views except editioning views. See *Oracle Database Development Guide* for a discussion of editioning views and edition-based redefinition.

See Also:

Oracle Database Concepts for an overview of views

23.1.2 Creating Views and Join Views

You can create views using the CREATE VIEW statement. Each view is defined by a query that references tables, materialized views, or other views. You can also create join views that specify multiple base tables or views in the FROM clause.

Creating Views

You can create a view with the CREATE VIEW statement.

Creating Join Views

You can also create views that specify multiple base tables or views in the FROM clause of a CREATE VIEW statement. These are called **join views**.

Expansion of Defining Queries at View Creation Time

When a view is created, Oracle Database expands any wildcard (*) in a top-level view query into a column list. The resulting query is stored in the data dictionary; any subqueries are left intact.

Creating Views with Errors

If there are no syntax errors in a CREATE VIEW statement, then the database can create the view even if the defining query of the view cannot be executed. In this case, the view is considered "created with errors."

23.1.2.1 Creating Views

You can create a view with the CREATE VIEW statement.

To create a view, you must meet the following requirements:

- To create a view in your schema, you must have the CREATE VIEW privilege. To create a view in another user's schema, you must have the CREATE ANY VIEW system privilege. You can acquire these privileges explicitly or through a role.
- The owner of the view (whether it is you or another user) must have been explicitly granted privileges to access all objects referenced in the view definition. The owner *cannot* have obtained these privileges through roles. Also, the functionality of the view depends on the privileges of the view owner. For example, if the owner of the view has only the INSERT privilege for Scott's emp table, then the view can be used only to insert new rows into the emp table, not to SELECT, UPDATE, or DELETE rows.
- If the owner of the view intends to grant access to the view to other users, the owner must have received the object privileges to the base objects with the GRANT OPTION or the system privileges with the ADMIN OPTION.

You can create views using the CREATE VIEW statement. Each view is defined by a query that references tables, materialized views, or other views. As with all subqueries, the query that defines a view cannot contain the FOR UPDATE clause.

The following statement creates a view on a subset of data in the hr.departments table:

```
CREATE VIEW departments_hq AS

SELECT department_id, department_name, location_id

FROM hr.departments

WHERE location_id = 1700

WITH CHECK OPTION CONSTRAINT departments_hq_cnst;
```

The query that defines the <code>departments_hq</code> view references only rows in location 1700. Furthermore, the <code>CHECK OPTION</code> creates the view with the constraint (named <code>departments_hq_cnst</code>) so that <code>INSERT</code> and <code>UPDATE</code> statements issued against the view cannot result in rows that the view cannot select. For example, the following <code>INSERT</code> statement successfully inserts a row into the <code>departments</code> table with the <code>departments_hq</code> view, which contains all rows with location 1700:

```
INSERT INTO departments hq VALUES (300, 'NETWORKING', 1700);
```

However, the following INSERT statement returns an error because it attempts to insert a row for location 2700, which cannot be selected using the departments hq view:

```
INSERT INTO departments_hq VALUES (301, 'TRANSPORTATION', 2700);
```

The view could have been constructed specifying the WITH READ ONLY clause, which prevents any updates, inserts, or deletes from being done to the base table through the view. If no WITH clause is specified, the view, with some restrictions, is inherently updatable.

You can also create views with invisible columns. For example, the following statements creates the departments has man view and makes the manager id column invisible:

See Also:

- Oracle Database SQL Language Reference for syntax and semantics of the CREATE VIEW statement
- "Understand Invisible Columns"

23.1.2.2 Creating Join Views

You can also create views that specify multiple base tables or views in the FROM clause of a CREATE VIEW statement. These are called **join views**.

The following statement creates the division1_staff view that joins data from the emp and dept tables:

```
CREATE VIEW division1_staff AS

SELECT ename, empno, job, dname

FROM emp, dept

WHERE emp.deptno IN (10, 30)

AND emp.deptno = dept.deptno;
```

An **updatable join view** is a join view where <code>UPDATE</code>, <code>INSERT</code>, and <code>DELETE</code> operations are allowed. See "Updating a Join View" for further discussion.

23.1.2.3 Expansion of Defining Queries at View Creation Time

When a view is created, Oracle Database expands any wildcard (*) in a top-level view query into a column list. The resulting query is stored in the data dictionary; any subqueries are left intact.

The column names in an expanded column list are enclosed in quotation marks to account for the possibility that the columns of the base object were originally entered with quotes and require them for the query to be syntactically correct.

As an example, assume that the dept view is created as follows:

```
CREATE VIEW dept AS SELECT * FROM scott.dept;
```

The database stores the defining query of the dept view as:

```
SELECT "DEPTNO", "DNAME", "LOC" FROM scott.dept;
```

Views created with errors do not have wildcards expanded. However, if the view is eventually compiled without errors, wildcards in the defining query are expanded.



23.1.2.4 Creating Views with Errors

If there are no syntax errors in a CREATE VIEW statement, then the database can create the view even if the defining query of the view cannot be executed. In this case, the view is considered "created with errors."

For example, when a view is created that refers to a nonexistent table or an invalid column of an existing table, or when the view owner does not have the required privileges, the view can be created anyway and entered into the data dictionary. However, the view is not yet usable.

To create a view with errors, you must include the FORCE clause of the CREATE VIEW statement.

```
CREATE FORCE VIEW AS ...;
```

By default, views with errors are created as INVALID. When you try to create such a view, the database returns a message indicating the view was created with errors. If conditions later change so that the query of an invalid view can be executed, the view can be recompiled and be made valid (usable). For information changing conditions and their impact on views, see "Managing Object Dependencies".

23.1.3 Replacing Views

You can replace a view by dropping it and re-creating it or by issuing a CREATE VIEW statement that contains the OR REPLACE clause.

To replace a view, you must have all of the privileges required to drop and create a view. If the definition of a view must change, the view must be replaced; you cannot use an ALTER VIEW statement to change the definition of a view. You can replace views in the following ways:

You can drop and re-create the view.



When a view is dropped, all grants of corresponding object privileges are revoked from roles and users. After the view is re-created, privileges must be regranted.

You can redefine the view with a CREATE VIEW statement that contains the OR REPLACE clause. The OR REPLACE clause replaces the current definition of a view and preserves the current security authorizations. For example, assume that you created the sales_staff view as shown earlier, and, in addition, you granted several object privileges to roles and other users. However, now you must redefine the sales_staff view to change the department number specified in the WHERE clause. You can replace the current version of the sales staff view with the following statement:

```
CREATE OR REPLACE VIEW sales_staff AS

SELECT empno, ename, deptno

FROM emp

WHERE deptno = 30

WITH CHECK OPTION CONSTRAINT sales_staff_cnst;
```

Before replacing a view, consider the following effects:

 Replacing a view replaces the view definition in the data dictionary. All underlying objects referenced by the view are not affected.

- If a constraint in the CHECK OPTION was previously defined but not included in the new view definition, the constraint is dropped.
- All views dependent on a replaced view become invalid (not usable). In addition, dependent PL/SQL program units may become invalid, depending on what was changed in the new version of the view. For example, if only the WHERE clause of the view changes, dependent PL/SQL program units remain valid. However, if any changes are made to the number of view columns or to the view column names or data types, dependent PL/SQL program units are invalidated. See "Managing Object Dependencies" for more information on how the database manages such dependencies.

23.1.4 Using Views in Queries

You can query a view. You can also perform data manipulation language (DML) operations on views, with some restrictions.

To issue a query or an INSERT, UPDATE, or DELETE statement against a view, you must have the SELECT, READ, INSERT, UPDATE, or DELETE object privilege for the view, respectively, either explicitly or through a role.

Views can be queried in the same manner as tables. For example, to query the <code>Division1 staff view</code>, enter a valid <code>SELECT statement</code> that references the view:

ENAME	EMPNO	JOB	DNAME
CLARK KING MILLER ALLEN WARD JAMES TURNER MARTIN	7782 7839 7934 7499 7521 7900 7844 7654	MANAGER PRESIDENT CLERK SALESMAN SALESMAN CLERK SALESMAN	ACCOUNTING ACCOUNTING ACCOUNTING SALES SALES SALES SALES SALES SALES SALES
BLAKE	7698	MANAGER	SALES

With some restrictions, rows can be inserted into, updated in, or deleted from a base table using a view. The following statement inserts a new row into the emp table using the sales staff view:

```
INSERT INTO sales_staff
    VALUES (7954, 'OSTER', 30);
```

Restrictions on DML operations for views use the following criteria in the order listed:

- 1. If a view is defined by a query that contains SET or DISTINCT operators, a GROUP BY clause, or a group function, then rows cannot be inserted into, updated in, or deleted from the base tables using the view.
- 2. If a view is defined with WITH CHECK OPTION, a row cannot be inserted into, or updated in, the base table (using the view), if the view cannot select the row from the base table.
- 3. If a NOT NULL column that does not have a DEFAULT clause is omitted from the view, then a row cannot be inserted into the base table using the view.
- 4. If the view was created by using an expression, such as DECODE (deptno, 10, "SALES", ...), then rows cannot be inserted into or updated in the base table using the view.



The constraint created by WITH CHECK OPTION of the sales_staff view only allows rows that have a department number of 30 to be inserted into, or updated in, the emp table. Alternatively, assume that the sales_staff view is defined by the following statement (that is, excluding the deptno column):

```
CREATE VIEW sales_staff AS
    SELECT empno, ename
    FROM emp
    WHERE deptno = 10
    WITH CHECK OPTION CONSTRAINT sales staff cnst;
```

Considering this view definition, you can update the <code>empno</code> or <code>ename</code> fields of existing records, but you cannot insert rows into the <code>emp</code> table through the <code>sales_staff</code> view because the view does not let you alter the <code>deptno</code> field. If you had defined a <code>DEFAULT</code> value of 10 on the <code>deptno</code> field, then you could perform inserts.

When a user attempts to reference an invalid view, the database returns an error message to the user:

```
ORA-04063: view 'view_name' has errors
```

This error message is returned when a view exists but is unusable due to errors in its query (whether it had errors when originally created or it was created successfully but became unusable later because underlying objects were altered or dropped).

23.1.5 DML Statements and Join Views

Restrictions apply when issuing DML statements on join views.

Updating a Join View

An updatable join view (also referred to as a **modifiable join view**) is a view that contains multiple tables in the top-level FROM clause of the SELECT statement, and is not restricted by the WITH READ ONLY clause.

Key-Preserved Tables

A table is key-preserved if every key of the table can also be a key of the result of the join that is based on the table. So, a key-preserved table has its keys preserved through a join.

Rules for DML Statements and Join Views

The general rule is that any UPDATE, DELETE, or INSERT statement on a join view can modify only one underlying base table.

- Updating Views That Involve Outer Joins
 Views that involve outer joins are modifiable in some cases.
- Using the UPDATABLE_ COLUMNS Views
 A set of views can assist you in identifying inherently updatable join views.

23.1.5.1 Updating a Join View

An updatable join view (also referred to as a **modifiable join view**) is a view that contains multiple tables in the top-level FROM clause of the SELECT statement, and is not restricted by the WITH READ ONLY clause.

The rules for updatable join views are shown in the following table. Views that meet these criteria are said to be inherently updatable.

Rule	Description
General Rule	Any INSERT, UPDATE, or DELETE operation on a join view can modify only one underlying base table at a time.
UPDATE Rule	Rows from a join view can be updated if the join column keys in the base tables are unique. That is, the WHERE clause in the UPDATE statement must be deterministic. If the base tables are not key-preserved, you must ensure that the join column keys are unique. If the view is defined with the WITH CHECK OPTION clause, then all join columns and all columns of repeated tables are not updatable.
DELETE Rule	Rows from a join view can be deleted as long as there is exactly one key- preserved table in the join. The key preserved table can be repeated in the FROM clause. If the view is defined with the WITH CHECK OPTION clause and the key preserved table is repeated, then the rows cannot be deleted from the view.
INSERT Rule	An INSERT statement must not explicitly or implicitly refer to the columns of a non-key-preserved table. If the join view is defined with the WITH CHECK OPTION clause, INSERT statements are not permitted.

There are data dictionary views that indicate whether the columns in a join view are inherently updatable. See "Using the UPDATABLE COLUMNS Views" for descriptions of these views.



There are some additional restrictions and conditions that can affect whether a join view is inherently updatable. Specifics are listed in the description of the CREATE VIEW statement in the *Oracle Database SQL Language Reference*.

If a view is not inherently updatable, it can be made updatable by creating an INSTEAD OF trigger on it. See *Oracle Database PL/SQL Language Reference* for information about triggers.

Additionally, if a view is a join on other nested views, then the other nested views must be mergeable into the top level view. For a discussion of mergeable and unmergeable views, and more generally, how the optimizer optimizes statements that reference views, see the *Oracle Database SQL Tuning Guide*.

Examples illustrating the rules for inherently updatable join views, and a discussion of key-preserved tables, are presented in following sections. The examples in these sections work only if you explicitly define the primary and foreign keys in the tables, or define unique indexes. The following statements create the appropriately constrained table definitions for emp and dept.

```
CREATE TABLE dept (
     deptno NUMBER(4) PRIMARY KEY,
     dname
                  VARCHAR2 (14),
     loc
                  VARCHAR2(13));
CREATE TABLE emp (
                 NUMBER (4) PRIMARY KEY,
     empno
     ename
                  VARCHAR2(10),
                  VARCHAR2(9),
     job
                  NUMBER (4),
     mgr
                  NUMBER (7,2),
     sal
```



```
comm NUMBER(7,2),
deptno NUMBER(2),
FOREIGN KEY (DEPTNO) REFERENCES DEPT(DEPTNO));
```

You could also omit the primary and foreign key constraints listed in the preceding example, and create a UNIQUE INDEX on dept (deptno) to make the following examples work.

The following statement created the emp dept join view which is referenced in the examples:

```
CREATE VIEW emp_dept AS

SELECT emp.empno, emp.ename, emp.deptno, emp.sal, dept.dname, dept.loc

FROM emp, dept

WHERE emp.deptno = dept.deptno

AND dept.loc IN ('DALLAS', 'NEW YORK', 'BOSTON');
```

23.1.5.2 Key-Preserved Tables

A table is key-preserved if every key of the table can also be a key of the result of the join that is based on the table. So, a key-preserved table has its keys preserved through a join.



It is not necessary that the key or keys of a table be selected for it to be key preserved. It is sufficient that if the key or keys were selected, then they would also be keys of the result of the join.

The concept of a key-preserved table is fundamental to understanding the restrictions on modifying join views. Each row in a key-preserved table appears at most only once in a join view based on this table. If a table ${\tt T}$ is joined to a table ${\tt S}$ using the condition ${\tt T.col1} = {\tt S.col3}$, then ${\tt T}$ is key-preserved if the join key is ${\tt S.col3}$ are unique. The data in a table is not relevant when determining whether a table is key-preserved. Instead, the constraints on the table determine if a table is key-preserved. Key-preserved tables are joined with a source table using the primary key or unique key of the source table.

For example, in the <code>emp_dept</code> view, because <code>emp</code> is joined with the primary key of <code>dept</code>, <code>emp</code> is a key-preserved table. If, in the <code>emp</code> table, there was at most one employee in each department, then <code>deptno</code> would be unique in the result of a join of <code>emp</code> and <code>dept</code>, but <code>dept</code> would still not be a key-preserved table.

If you select all rows from the emp dept view, the results are:

EMPNO	ENAME	DEPTNO	DNAME	LOC
7782	CLARK	10	ACCOUNTING	NEW YORK
7839	KING	10	ACCOUNTING	NEW YORK
7934	MILLER	10	ACCOUNTING	NEW YORK
7369	SMITH	20	RESEARCH	DALLAS
7876	ADAMS	20	RESEARCH	DALLAS
7902	FORD	20	RESEARCH	DALLAS
7788	SCOTT	20	RESEARCH	DALLAS
7566	JONES	20	RESEARCH	DALLAS
8 rows sele	ected.			

In this view, emp is a key-preserved table, because empno is a key of the emp table, and also a key of the result of the join. dept is *not* a key-preserved table, because although deptno is a key of the dept table, it is not a key of the join.



23.1.5.3 Rules for DML Statements and Join Views

The general rule is that any UPDATE, DELETE, or INSERT statement on a join view can modify only one underlying base table.

- UPDATE Statements and Join Views
 Examples illustrate UPDATE statements that can modify join views.
- DELETE Statements and Join Views

For most join views, a delete is successful only if there is *one* and *only one* key-preserved table in the join. The key-preserved table can be repeated in the FROM clause.

INSERT Statements and Join Views
 Examples illustrate INSERT statements that can modify join views.

23.1.5.3.1 UPDATE Statements and Join Views

Examples illustrate **UPDATE** statements that can modify join views.

Starting with Oracle Database Release 21c, it is not mandatory for all updatable columns in a join view to map to columns of a key-preserved table. The columns in a non-key-preserved table can be updated if the <code>UPDATE</code> operation only updates columns from a single table and the update is deterministic, meaning that it updates each row only once.

The following example shows an UPDATE statement that successfully modifies the <code>emp_dept</code> view:

```
UPDATE emp_dept
    SET    sal = sal * 1.10
    WHERE deptno = 10;
```

The following UPDATE statement successfully modifies the LOC column in the DEPT table (the non-key-preserved table) because it updates only one row in the EMP table:

```
UPDATE emp_dept
   SET loc = 'BOSTON'
WHERE ename = 'SMITH';
```

The following UPDATE statement results in an ORA-30926 error because the update operation is non-deterministic:

```
UPDATE emp_dept
    SET loc = 'BOSTON'
WHERE ename = 'S%';
```

A row from the dept table is joined to multiple rows from emp table (with ename = 'SCOTT' and ename = 'SMITH'). Therefore, an attempt is made to modify the same row multiple times. To make the UPDATE is deterministic, ensure that a row from the dept table is only joined to one row from emp table.

In general, all updatable columns of a join view must map to columns of a key-preserved table. If the view is defined using the WITH CHECK OPTION clause, then all join columns and all columns taken from tables that are referenced more than once in the view are not modifiable.

So, for example, if the <code>emp_dept</code> view were defined using <code>WITH CHECK OPTION</code>, the following <code>UPDATE</code> statement would fail:

```
UPDATE emp_dept
    SET    deptno = 10
    WHERE    ename = 'SMITH';
```

The statement fails because it is trying to update a join column.



Oracle Database SQL Language Reference for syntax and additional information about the UPDATE statement

23.1.5.3.2 DELETE Statements and Join Views

For most join views, a delete is successful only if there is *one and only one* key-preserved table in the join. The key-preserved table can be repeated in the FROM clause.

The following DELETE statement works on the emp dept view:

```
DELETE FROM emp_dept
    WHERE ename = 'SMITH';
```

This DELETE statement on the <code>emp_dept</code> view is valid because it can be translated to a <code>DELETE</code> operation on the base <code>emp</code> table, and because the <code>emp</code> table is the only key-preserved table in the join.

In the following view, a DELETE operation is permitted, because although there are two key-preserved tables, they are the same table. That is, the key-preserved table is repeated. In this case, the delete statement operates on the first table in the FROM clause (e1, in this example):

```
CREATE VIEW emp_emp AS

SELECT el.ename, e2.empno, e2.deptno
FROM emp e1, emp e2

WHERE el.empno = e2.empno;
```

If a view is defined using the WITH CHECK OPTION clause and the key-preserved table is repeated, rows cannot be deleted from such a view.

```
CREATE VIEW emp_mgr AS

SELECT el.ename, e2.ename mname

FROM emp e1, emp e2

WHERE e1.mgr = e2.empno

WITH CHECK OPTION;
```



Note:

- If the DELETE statement uses the same column in its WHERE clause that was used to create the view as a join condition, then the delete operation can be successful when there are different key-preserved tables in the join. In this case, the DELETE statement operates on the first table in the FROM clause, and the tables in the FROM clause can be different from the tables in the WHERE clause.
- The DELETE statement is successful, even if it does not use the WHERE clause.
- The DELETE statement is successful, even if it uses a different column in its WHERE clause than the one that was used to create the view as a join condition.
- The DELETE statement operates on the second table in the FROM clause in all the cases, because no primary key is defined on the second table.
- If a primary key is defined on the second table, then the DELETE statement operates on the first table in the FROM clause.

See Also:

Oracle Database SQL Language Reference for syntax and additional information about the DELETE statement

23.1.5.3.3 INSERT Statements and Join Views

Examples illustrate INSERT statements that can modify join views.

The following INSERT statement on the emp dept view succeeds:

```
INSERT INTO emp_dept (ename, empno, deptno)
VALUES ('KURODA', 9010, 40);
```

This statement works because only one key-preserved base table is being modified (emp), and 40 is a valid deptno in the dept table (thus satisfying the FOREIGN KEY integrity constraint on the emp table).

An INSERT statement, such as the following, would fail for the same reason that such an UPDATE on the base emp table would fail: the FOREIGN KEY integrity constraint on the emp table is violated (because there is no deptno 77).

```
INSERT INTO emp_dept (ename, empno, deptno)
VALUES ('KURODA', 9010, 77);
```

The following INSERT statement would fail with an error (ORA-01776 cannot modify more than one base table through a join view):

```
INSERT INTO emp_dept (empno, ename, loc)
   VALUES (9010, 'KURODA', 'BOSTON');
```

An INSERT cannot implicitly or explicitly refer to columns of a non-key-preserved table. If the join view is defined using the WITH CHECK OPTION clause, then you cannot perform an INSERT to it.



Oracle Database SQL Language Reference for syntax and additional information about the INSERT statement

23.1.5.4 Updating Views That Involve Outer Joins

Views that involve outer joins are modifiable in some cases.

For example:

```
CREATE VIEW emp_dept_oj1 AS
    SELECT empno, ename, e.deptno, dname, loc
    FROM emp e, dept d
    WHERE e.deptno = d.deptno (+);
```

The statement:

```
SELECT * FROM emp_dept_oj1;
```

Results in:

EMPNO	ENAME	DEPTNO	DNAME	LOC
7369	SMITH	40	OPERATIONS	BOSTON
7499	ALLEN	30	SALES	CHICAGO
7566	JONES	20	RESEARCH	DALLAS
7654	MARTIN	30	SALES	CHICAGO
7698	BLAKE	30	SALES	CHICAGO
7782	CLARK	10	ACCOUNTING	NEW YORK
7788	SCOTT	20	RESEARCH	DALLAS
7839	KING	10	ACCOUNTING	NEW YORK
7844	TURNER	30	SALES	CHICAGO
7876	ADAMS	20	RESEARCH	DALLAS
7900	JAMES	30	SALES	CHICAGO
7902	FORD	20	RESEARCH	DALLAS
7934	MILLER	10	ACCOUNTING	NEW YORK
7521	WARD	30	SALES	CHICAGO
14 rows	selected.			

Columns in the base <code>emp</code> table of <code>emp_dept_oj1</code> are modifiable through the view, because <code>emp</code> is a key-preserved table in the join.

The following view also contains an outer join:

```
CREATE VIEW emp_dept_oj2 AS
SELECT e.empno, e.ename, e.deptno, d.dname, d.loc
FROM emp e, dept d
WHERE e.deptno (+) = d.deptno;
```

The following statement:

```
SELECT * FROM emp_dept_oj2;
```

Results in:

EMPNO	ENAME	DEPTNO	DNAME	LOC



7782	CLARK	10	ACCOUNTING	NEW YORK
7839	KING	10	ACCOUNTING	NEW YORK
7934	MILLER	10	ACCOUNTING	NEW YORK
7369	SMITH	20	RESEARCH	DALLAS
7876	ADAMS	20	RESEARCH	DALLAS
7902	FORD	20	RESEARCH	DALLAS
7788	SCOTT	20	RESEARCH	DALLAS
7566	JONES	20	RESEARCH	DALLAS
7499	ALLEN	30	SALES	CHICAGO
7698	BLAKE	30	SALES	CHICAGO
7654	MARTIN	30	SALES	CHICAGO
7900	JAMES	30	SALES	CHICAGO
7844	TURNER	30	SALES	CHICAGO
7521	WARD	30	SALES	CHICAGO
			OPERATIONS	BOSTON

15 rows selected.

In this view, emp is no longer a key-preserved table, because the empno column in the result of the join can have nulls (the last row in the preceding SELECT statement). So, UPDATE, DELETE, and INSERT operations cannot be performed on this view.

In the case of views containing an outer join on other nested views, a table is key preserved if the view or views containing the table are merged into their outer views, all the way to the top. A view which is being outer-joined is currently merged only if it is "simple." For example:

```
SELECT col1, col2, ... FROM T;
```

The select list of the view has no expressions.

If you are in doubt whether a view is modifiable, then you can select from the USER UPDATABLE COLUMNS view to see if it is. For example:

```
SELECT owner, table_name, column_name, updatable FROM USER_UPDATABLE_COLUMNS
    WHERE TABLE_NAME = 'EMP_DEPT_VIEW';
```

This returns output similar to the following:

OWNER	TABLE_NAME	COLUMN_NAM	UPD
SCOTT	EMP_DEPT_V	EMPNO	NO
SCOTT	EMP_DEPT_V	ENAME	NO
SCOTT	EMP_DEPT_V	DEPTNO	NO
SCOTT	EMP_DEPT_V	DNAME	NO
SCOTT	EMP_DEPT_V	LOC	NO
5 rows sele	cted.		

23.1.5.5 Using the UPDATABLE_ COLUMNS Views

A set of views can assist you in identifying inherently updatable join views.

View	Description
DBA_UPDATABLE_COLUMNS	Shows all columns in all tables and views that are modifiable.
ALL_UPDATABLE_COLUMNS	Shows all columns in all tables and views accessible to the user that are modifiable.
USER_UPDATABLE_COLUMNS	Shows all columns in all tables and views in the user's schema that are modifiable.

The updatable columns in view emp_dept are shown below.



SELECT COLUMN_NAME, UPDATABLE
FROM USER_UPDATABLE_COLUMNS
WHERE TABLE NAME = 'EMP DEPT';

COLUMN_NAME	UPD
EMPNO	YES
ENAME	YES
DEPTNO	YES
SAL	YES
DNAME	NO
LOC	NO

6 rows selected.



Oracle Database Reference for complete descriptions of the updatable column views

23.1.6 Altering Views

You use the ALTER VIEW statement only to explicitly recompile a view that is invalid.

To change the definition of a view, see "Replacing Views".

The ALTER VIEW statement lets you locate recompilation errors before run time. To ensure that the alteration does not affect the view or other objects that depend on it, you can explicitly recompile a view after altering one of its base tables.

To use the ALTER VIEW statement, the view must be in your schema, or you must have the ALTER ANY TABLE system privilege.



Oracle Database SQL Language Reference for syntax and additional information about the ALTER VIEW statement

23.1.7 Dropping Views

You can drop a view with the DROP VIEW statement.

You can drop any view contained in your schema. To drop a view in another user's schema, you must have the DROP ANY VIEW system privilege. Drop a view using the DROP VIEW statement. For example, the following statement drops the emp dept view:

DROP VIEW emp_dept;





Oracle Database SQL Language Reference for syntax and additional information about the DROP VIEW statement

23.2 Managing Sequences

You can perform tasks such as creating sequences, altering sequences, using sequences, and dropping sequences.

About Sequences

Sequences are database objects from which multiple users can generate unique integers. The sequence generator generates sequential numbers, which can be used to generate unique primary keys automatically, and to coordinate keys across multiple rows or tables.

Creating Sequences

Create a sequence using the CREATE SEQUENCE statement.

Altering Sequences

Alter a sequence using the ALTER SEQUENCE statement.

Using Sequences

A sequence can be accessed and incremented by multiple users.

Dropping Sequences

If a sequence is no longer required, you can drop the sequence using the DROP SEQUENCE statement.

23.2.1 About Sequences

Sequences are database objects from which multiple users can generate unique integers. The sequence generator generates sequential numbers, which can be used to generate unique primary keys automatically, and to coordinate keys across multiple rows or tables.

Without sequences, sequential values can only be produced programmatically. A new primary key value can be obtained by selecting the most recently produced value and incrementing it. This method requires a lock during the transaction and causes multiple users to wait for the next value of the primary key; this waiting is known as **serialization**. If developers have such constructs in applications, then you should encourage the developers to replace them with access to sequences. Sequences eliminate serialization and improve the concurrency of an application.



Oracle Database Concepts for an overview of sequences

23.2.2 Creating Sequences

Create a sequence using the CREATE SEQUENCE statement.

To create a sequence in your schema, you must have the CREATE SEQUENCE system privilege. To create a sequence in another user's schema, you must have the CREATE ANY SEQUENCE privilege.

For example, the following statement creates a sequence used to generate employee numbers for the <code>empno</code> column of the <code>emp</code> table:

```
CREATE SEQUENCE emp_sequence
INCREMENT BY 1
START WITH 1
NOMAXVALUE
NOCYCLE
CACHE 10;
```

Notice that several parameters can be specified to control the function of sequences. You can use these parameters to indicate whether the sequence is ascending or descending, the starting point of the sequence, the minimum and maximum values, and the interval between sequence values. The NOCYCLE option indicates that the sequence cannot generate more values after reaching its maximum or minimum value.

The CACHE clause preallocates a set of sequence numbers and keeps them in memory so that sequence numbers can be accessed faster. When the last of the sequence numbers in the cache has been used, the database reads another set of numbers into the cache.

The database might skip sequence numbers if you choose to cache a set of sequence numbers. For example, when an instance abnormally shuts down (for example, when an instance failure occurs or a Shutdown Abort statement is issued), sequence numbers that have been cached but not used are lost. Also, sequence numbers that have been used but not saved are lost as well. The database might also skip cached sequence numbers after an export and import. See *Oracle Database Utilities* for details.

See Also:

- Oracle Database SQL Language Reference for the CREATE SEQUENCE statement syntax
- Oracle Real Application Clusters Administration and Deployment Guide for information about using sequences in an Oracle Real Application Clusters environment

23.2.3 Altering Sequences

Alter a sequence using the ALTER SEQUENCE statement.

To alter a sequence, your schema must contain the sequence, you must have the ALTER object privilege on the sequence, or you must have the ALTER ANY SEQUENCE system privilege. You can alter a sequence to change any of the parameters that define how it generates sequence numbers. To change the starting point of a sequence, you can either drop the sequence and then re-create it, or use the RESTART clause to restart the sequence. For an ascending

sequence, the RESTART clause resets NEXTVAL to MINVALUE. For a descending sequence, NEXTVAL is reset to MAXVALUE.

The following example alters the emp sequence sequence:

```
ALTER SEQUENCE emp_sequence
INCREMENT BY 10
MAXVALUE 10000
CYCLE
CACHE 20;
```



Oracle Database SQL Language Reference for syntax and additional information about the ALTER SEQUENCE statement

23.2.4 Using Sequences

A sequence can be accessed and incremented by multiple users.

To use a sequence, your schema must contain the sequence or you must have been granted the <code>SELECT</code> object privilege for another user's sequence. Once a sequence is defined, it can be accessed and incremented by multiple users (who have <code>SELECT</code> object privilege for the sequence containing the sequence) with no waiting. The database does not wait for a transaction that has incremented a sequence to complete before that sequence can be incremented again.

The examples outlined in the following sections show how sequences can be used in parent/child table relationships. Assume an order entry system is partially comprised of two tables, orders_tab (parent table) and line_items_tab (child table), that hold information about customer orders. A sequence named order seq is defined by the following statement:

```
CREATE SEQUENCE Order_seq
START WITH 1
INCREMENT BY 1
NOMAXVALUE
NOCYCLE
CACHE 20;
```

Referencing a Sequence

A sequence is referenced in SQL statements with the NEXTVAL and CURRVAL pseudocolumns; each new sequence number is generated by a reference to the sequence pseudocolumn NEXTVAL, while the current sequence number can be repeatedly referenced using the pseudo-column CURRVAL.

Caching Sequence Numbers

Caching sequence numbers can improve access time.

Making a Sequence Scalable

A sequence can be made scalable by specifying the scale clause in the create sequence or alter sequence statement.

23.2.4.1 Referencing a Sequence

A sequence is referenced in SQL statements with the NEXTVAL and CURRVAL pseudocolumns; each new sequence number is generated by a reference to the sequence pseudocolumn NEXTVAL, while the current sequence number can be repeatedly referenced using the pseudocolumn CURRVAL.

NEXTVAL and CURRVAL are not reserved words or keywords and can be used as pseudocolumn names in SQL statements such as SELECT, INSERT, or UPDATE.

- Generating Sequence Numbers with NEXTVAL
 To generate and use a sequence number, reference seq_name.NEXTVAL in a SQL statement.
- Using Sequence Numbers with CURRVAL
 To use or refer to the current sequence value of your session, reference seq_name.CURRVAL in a SQL statement.
- Uses and Restrictions of NEXTVAL and CURRVAL
 CURRVAL and NEXTVAL can be used in specific places, and restrictions apply to their use.

23.2.4.1.1 Generating Sequence Numbers with NEXTVAL

To generate and use a sequence number, reference seq_name.NEXTVAL in a SQL statement.

For example, assume a customer places an order. The sequence number can be referenced in a values list. For example:

```
INSERT INTO Orders_tab (Orderno, Custno)
    VALUES (Order seq.NEXTVAL, 1032);
```

Or, the sequence number can be referenced in the SET clause of an UPDATE statement. For example:

```
UPDATE Orders_tab
    SET Orderno = Order_seq.NEXTVAL
    WHERE Orderno = 10112;
```

The sequence number can also be referenced outermost SELECT of a query or subquery. For example:

```
SELECT Order_seq.NEXTVAL FROM dual;
```

As defined, the first reference to order_seq.NEXTVAL returns the value 1. Each subsequent statement that references order_seq.NEXTVAL generates the next sequence number (2, 3, 4,...). The pseudo-column NEXTVAL can be used to generate as many new sequence numbers as necessary. However, only a single sequence number can be generated for each row. In other words, if NEXTVAL is referenced more than once in a single statement, then the first reference generates the next number, and all subsequent references in the statement return the same number.

Once a sequence number is generated, the sequence number is available only to the session that generated the number. Independent of transactions committing or rolling back, other users referencing <code>order_seq.Nextval</code> obtain unique values. If two users are accessing the same sequence concurrently, then the sequence numbers each user receives might have gaps because sequence numbers are also being generated by the other user.

23.2.4.1.2 Using Sequence Numbers with CURRVAL

To use or refer to the current sequence value of your session, reference seq_name .CURRVAL in a SQL statement.

CURRVAL can only be used if *seq_name*.NEXTVAL has been referenced in the current user session (in the current or a previous transaction). CURRVAL can be referenced as many times as necessary, including multiple times within the same statement. The next sequence number is not generated until NEXTVAL is referenced. Continuing with the previous example, you would finish placing the customer's order by inserting the line items for the order:

```
INSERT INTO Line_items_tab (Orderno, Partno, Quantity)
    VALUES (Order_seq.CURRVAL, 20321, 3);
INSERT INTO Line_items_tab (Orderno, Partno, Quantity)
    VALUES (Order seq.CURRVAL, 29374, 1);
```

Assuming the INSERT statement given in the previous section generated a new sequence number of 347, both rows inserted by the statements in this section insert rows with order numbers of 347.

23.2.4.1.3 Uses and Restrictions of NEXTVAL and CURRVAL

CURRVAL and NEXTVAL can be used in specific places, and restrictions apply to their use.

CURRVAL and NEXTVAL can be used in the following places:

- VALUES clause of INSERT statements
- The SELECT list of a SELECT statement
- A view guery or materialized view guery

However, the use of CURRVAL and NEXTVAL in a materialized view query makes the materialized view complex. Therefore, it cannot be fast refreshed.

The SET clause of an UPDATE statement

CURRVAL and NEXTVAL cannot be used in these places:

- A subquery
- A SELECT statement with the DISTINCT operator
- A SELECT statement with a GROUP BY or ORDER BY clause
- A SELECT statement that is combined with another SELECT statement with the UNION, INTERSECT, or MINUS set operator
- The where clause of a select statement
- The condition of a CHECK constraint

23.2.4.2 Caching Sequence Numbers

Caching sequence numbers can improve access time.

About Caching Sequence Numbers

Sequence numbers can be kept in the sequence cache in the System Global Area (SGA). Sequence numbers can be accessed more quickly in the sequence cache than they can be read from disk.

About Automatic Sizing of the Sequence Cache

Automatic resizing of the sequence cache improves performance significantly for fast insert workloads that use sequences.

• The Number of Entries in the Sequence Cache

When an application accesses a sequence in the sequence cache, the sequence numbers are read quickly. However, if an application accesses a sequence that is not in the cache, then the sequence must be read from disk to the cache before the sequence numbers are used.

The Number of Values in Each Sequence Cache Entry

When a sequence is read into the sequence cache, sequence values are generated and stored in a cache entry. These values can then be accessed quickly.

23.2.4.2.1 About Caching Sequence Numbers

Sequence numbers can be kept in the sequence cache in the System Global Area (SGA). Sequence numbers can be accessed more quickly in the sequence cache than they can be read from disk.

The sequence cache consists of entries. Each entry can hold many sequence numbers for a single sequence.

Follow these guidelines for fast access to all sequence numbers:

- Be sure the sequence cache can hold all the sequences used concurrently by your applications.
- Increase the number of values for each sequence held in the sequence cache.

23.2.4.2.2 About Automatic Sizing of the Sequence Cache

Automatic resizing of the sequence cache improves performance significantly for fast insert workloads that use sequences.

The automatic sequence cache size on each instance is dynamically computed based on the rate of usage of sequence numbers. Each instance caches the maximum of the manually configured sequence cache size and the projected cache size requirement for the next 10 seconds. Based on the sequence usage, the sequence cache size can shrink or grow. The minimum size to which the cache can shrink is the manually configured cache size. To prevent the sequence cache size from growing indefinitely, the cache size and each increment in the cache size is capped.

For cycle sequences, the upper bound for the automatic sequence cache size is the size of one cycle.

23.2.4.2.3 The Number of Entries in the Sequence Cache

When an application accesses a sequence in the sequence cache, the sequence numbers are read quickly. However, if an application accesses a sequence that is not in the cache, then the sequence must be read from disk to the cache before the sequence numbers are used.

If your applications use many sequences concurrently, then your sequence cache might not be large enough to hold all the sequences. In this case, access to sequence numbers might often require disk reads. For fast access to all sequences, be sure your cache has enough entries to hold all the sequences used concurrently by your applications.



Automatic Sizing of Sequence Cache

Starting with Oracle Database 21c, the size of the sequence cache on each instance is dynamically computed. The automatic sequence cache size is based on the rate of usage of sequence numbers. Each instance caches the maximum of the manually configured sequence cache size and the projected cache size requirement for the next 10 seconds. Based on the sequence usage, the sequence cache size can shrink or grow. To prevent the sequence cache size from growing indefinitely, the cache size and each increment in the cache size is capped.

Restrictions on automatic sequence cache size include the following:

- For cycle sequences, the upper bound for the automatic sequence cache size is the size of one cycle.
- Automatic sequence caching is not available for ordered sequences on Oracle Real Application Clusters (Oracle RAC).



Automatic tuning of sequence cache sizes is available with Oracle Autonomous Database only.

23.2.4.2.4 The Number of Values in Each Sequence Cache Entry

When a sequence is read into the sequence cache, sequence values are generated and stored in a cache entry. These values can then be accessed quickly.

The number of sequence values stored in the cache is determined by the CACHE parameter in the CREATE SEQUENCE statement. The default value for this parameter is 20.

This CREATE SEQUENCE statement creates the seq2 sequence so that 50 values of the sequence are stored in the SEQUENCE cache:

```
CREATE SEQUENCE seq2
CACHE 50;
```

The first 50 values of seq2 can then be read from the cache. When the 51st value is accessed, the next 50 values will be read from disk.

Choosing a high value for CACHE lets you access more successive sequence numbers with fewer reads from disk to the sequence cache. However, if there is an instance failure, then all sequence values in the cache are lost. Cached sequence numbers also could be skipped after an export and import if transactions continue to access the sequence numbers while the export is running.

If you use the NOCACHE option in the CREATE SEQUENCE statement, then the values of the sequence are not stored in the sequence cache. In this case, every access to the sequence requires a disk read. Such disk reads slow access to the sequence. This CREATE SEQUENCE statement creates the SEQ3 sequence so that its values are never stored in the cache:

```
CREATE SEQUENCE seq3
NOCACHE;
```



23.2.4.3 Making a Sequence Scalable

A sequence can be made *scalable* by specifying the SCALE clause in the CREATE SEQUENCE or ALTER SEQUENCE statement.

A scalable sequence is particularly efficient when used to generate unordered primary or unique keys for data ingestion workloads having high level of concurrency. Single Oracle database instances as well as Oracle RAC databases benefit from this feature. Scalable sequences significantly reduce the sequence and index block contention and provide better data load scalability compared to the solution of configuring a very large sequence cache using the CACHE clause of CREATE SEQUENCE or ALTER SEQUENCE statement.



In addition to using a scalable sequence, you can also partition the data to increase the performance of a data load operation.

The following is the syntax for defining a scalable sequence:

```
CREATE | ALTER SEQUENCE sequence_name
...
SCALE [EXTEND | NOEXTEND] | NOSCALE
```

When the SCALE clause is specified, a 6 digit numeric scalable sequence offset number is prefixed to the digits of the sequence:

```
scalable sequence number = 6 digit scalable sequence offset number ||
normal sequence number
```

where,

- | is the concatenation operator.
- 6 digit scalable sequence offset number = 3 digit *instance* offset number | | 3 digit session offset number.

The 3 digit *instance* offset number is generated as [(instance id % 100) + 100]. The 3 digit *session* offset number is generated as [session id % 1000].

Additionally, you can also specify EXTEND or NOEXTEND option for the SCALE clause:

EXTEND option

When the EXTEND option is specified for the SCALE clause, the scalable sequence values are of the length [X digits + Y digits], where X is the number of digits in the scalable sequence offset number (default is 6 digits), and Y is the number of digits specified in the MAXVALUE clause.



For example, for an ascending scalable sequence with MINVALUE of 1, MAXVALUE of 100 (3 digits), and EXTEND option specified, the scalable sequence values will be of 9 digits (6 digit scalable sequence offset number + 3 digit MAXVALUE) and will be of the form:

```
6 digit scalable sequence offset number || 001
6 digit scalable sequence offset number || 002
6 digit scalable sequence offset number || 003
...
6 digit scalable sequence offset number || 100
```

NOEXTEND option

When the NOEXTEND option is specified for the SCALE clause, which is the default option, the number of scalable sequence digits cannot exceed the number of digits specified in the MAXVALUE clause.

For example, for an ascending scalable sequence with MINVALUE of 1, MAXVALUE of 1000000 (7 digits), and NOEXTEND option specified, the scalable sequence values will be of 7 digits, because MAXVALUE of 1000000 contains 7 digits, and will be of the form:

```
6 digit scalable sequence offset number || 1
6 digit scalable sequence offset number || 2
6 digit scalable sequence offset number || 3
...
6 digit scalable sequence offset number || 9
```

Note that the NEXTVAL operation on this scalable sequence after the sequence value of [6 digit scalable sequence offset number $\mid\mid$ 9] will report the following error message, because the next scalable sequence value is [6 digit scalable sequence offset number $\mid\mid$ 10], which contains 8 digits and is greater than MAXVALUE of 1000000 that contains 7 digits:

ORA-64603: NEXTVAL cannot be instantiated for SQ. Widen the sequence by 1 digits or alter sequence with SCALE EXTEND.

Note:

The NOEXTEND option is useful for integration with the existing applications where sequences are used to populate fixed width columns.

To convert an existing scalable sequence to a non-scalable sequence, use the NOSCALE clause in the ALTER SEQUENCE statement.

Note:

Oracle recommends that you should not specify ordering for a scalable sequence, because scalable sequence numbers are globally unordered.

To know whether a sequence is scalable or whether a scalable sequence is extendable, check the values of the following columns of the <code>DBA_SEQUENCES</code>, <code>USER_SEQUENCES</code>, and <code>ALL_SEQUENCES</code> views.

Table 23-1 Columns Related to Scalable Sequences in the DBA_SEQUENCES, USER_SEQUENCES, and ALL_SEQUENCES Views

Column Name	Description
SCALE_FLAG	Indicates whether the sequence is a scalable sequence:
	• Y
	• N
EXTEND_FLAG	Indicates whether the scalable sequence is <i>extendable</i> , that is, whether the EXTEND option is applied for the scalable sequence, so that sequence values can extend beyond the value specified for MAXVALUE:
	• Y
	• N

23.2.5 Dropping Sequences

If a sequence is no longer required, you can drop the sequence using the ${\tt DROP}\ {\tt SEQUENCE}$ statement.

You can drop any sequence in your schema. To drop a sequence in another schema, you must have the DROP ANY SEQUENCE system privilege. For example, the following statement drops the order seq sequence:

DROP SEQUENCE order_seq;

When a sequence is dropped, its definition is removed from the data dictionary. Any synonyms for the sequence remain, but return an error when referenced.



Oracle Database SQL Language Reference for syntax and additional information about the DROP SEQUENCE statement

23.3 Managing Synonyms

You can perform tasks such as creating synonyms, using synonyms, and dropping synonyms.

- About Synonyms
 A synonym is an alias for a schema object.
- Creating Synonyms
 Create a synonym using the CREATE SYNONYM statement.
- Using Synonyms in DML Statements
 A synonym can be referenced in a DML statement the same way that the underlying object of the synonym can be referenced.

Dropping Synonyms

Drop a synonym that is no longer required using DROP SYNONYM statement. To drop a private synonym, omit the PUBLIC keyword. To drop a public synonym, include the PUBLIC keyword.

23.3.1 About Synonyms

A synonym is an alias for a schema object.

Synonyms can provide a level of security by masking the name and owner of an object and by providing location transparency for remote objects of a distributed database. Also, they are convenient to use and reduce the complexity of SQL statements for database users.

Synonyms allow underlying objects to be renamed or moved, where only the synonym must be redefined and applications based on the synonym continue to function without modification.

You can create both public and private synonyms. A **public** synonym is owned by the special user group named PUBLIC and is accessible to every user in a database. A **private** synonym is contained in the schema of a specific user and available only to the user and to grantees for the underlying object.

Synonyms themselves are not securable. When you grant object privileges on a synonym, you are really granting privileges on the underlying object, and the synonym is acting only as an alias for the object in the GRANT statement.



Oracle Database Concepts for a more complete description of synonyms

23.3.2 Creating Synonyms

Create a synonym using the CREATE SYNONYM statement.

To create a private synonym in your own schema, you must have the CREATE SYNONYM privilege. To create a private synonym in another user's schema, you must have the CREATE ANY SYNONYM privilege. To create a public synonym, you must have the CREATE PUBLIC SYNONYM system privilege.

When you create a synonym, the underlying schema object need not exist, nor do you need privileges to access the object for the CREATE SYNONYM statement to succeed. The following statement creates a public synonym named public_emp on the emp table contained in the schema of jward:

CREATE PUBLIC SYNONYM public_emp FOR jward.emp

When you create a synonym for a remote procedure or function, you must qualify the remote object with its schema name. Alternatively, you can create a local public synonym on the database where the remote object resides, in which case the database link must be included in all subsequent calls to the procedure or function.





Oracle Database SQL Language Reference for syntax and additional information about the CREATE SYNONYM statement

23.3.3 Using Synonyms in DML Statements

A synonym can be referenced in a DML statement the same way that the underlying object of the synonym can be referenced.

You can successfully use any private synonym contained in your schema or any public synonym, assuming that you have the necessary privileges to access the underlying object, either explicitly, from an enabled role, or from PUBLIC. You can also reference any private synonym contained in another schema if you have been granted the necessary object privileges for the underlying object.

You can reference another user's synonym using only the object privileges that you have been granted. For example, if you have only the SELECT privilege on the jward.emp table, and the synonym jward.employee is created for jward.emp, you can query the jward.employee synonym, but you cannot insert rows using the jward.employee synonym.

For example, if a synonym named employee refers to a table or view, then the following statement is valid:

```
INSERT INTO employee (empno, ename, job)
   VALUES (emp sequence.NEXTVAL, 'SMITH', 'CLERK');
```

If the synonym named $fire_emp$ refers to a standalone procedure or package procedure, then you could execute it with the command

```
EXECUTE Fire_emp(7344);
```

23.3.4 Dropping Synonyms

Drop a synonym that is no longer required using DROP SYNONYM statement. To drop a private synonym, omit the PUBLIC keyword. To drop a public synonym, include the PUBLIC keyword.

You can drop any private synonym in your own schema. To drop a private synonym in another user's schema, you must have the DROP ANY SYNONYM system privilege. To drop a public synonym, you must have the DROP PUBLIC SYNONYM system privilege.

For example, the following statement drops the private synonym named emp:

```
DROP SYNONYM emp;
```

The following statement drops the public synonym named public emp:

```
DROP PUBLIC SYNONYM public_emp;
```

When you drop a synonym, its definition is removed from the data dictionary. All objects that reference a dropped synonym remain. However, they become invalid (not usable). For more information about how dropping synonyms can affect other schema objects, see "Managing Object Dependencies".





Oracle Database SQL Language Reference for syntax and additional information about the DROP SYNONYM statement

23.4 Views, Synonyms, and Sequences Data Dictionary Views

You can query data dictionary views for information about views, synonyms, and sequences.

The following views display information about views, synonyms, and sequences:

View	Description
DBA_VIEWS	DBA view describes all views in the database. ALL view is
ALL_VIEWS	restricted to views accessible to the current user. USER view is
USER_VIEWS	restricted to views owned by the current user.
DBA_SYNONYMS	These views describe synonyms.
ALL_SYNONYMS	
USER_SYNONYMS	
DBA_SEQUENCES	These views describe sequences.
ALL_SEQUENCES	
USER_SEQUENCES	
DBA_UPDATABLE_COLUMNS	These views describe all columns in join views that are
ALL_UPDATABLE_COLUMNS	updatable.
USER_UPDATABLE_COLUMNS	

