9

Using SQL Data Types in Database Applications

This chapter explains how to choose the correct SQL data types for database columns that you create for your database applications.

Topics:

- Using the Correct and Most Specific Data Type
- · Representing Character Data
- · Representing Numeric Data
- Representing Date and Time Data
- · Representing Specialized Data
- Identifying Rows by Address
- Displaying Metadata for SQL Operators and Functions

✓ Note:

Oracle precompilers recognize, in embedded SQL programs, data types other than SQL and PL/SQL data types. These **external data types** are associated with host variables.

See Also:

- Oracle Database SQL Language Reference for information about data type conversion
- PL/SQL Data Types
- Data Types
- Overview of Precompilers

9.1 Using the Correct and Most Specific Data Type

Using the correct and most specific data type for each database column that you create for your database application increases data integrity, decreases storage requirements, and improves performance.

Topics:

How the Correct Data Type Increases Data Integrity

- How the Most Specific Data Type Decreases Storage Requirements
- How the Correct Data Type Improves Performance

9.1.1 How the Correct Data Type Increases Data Integrity

The correct data type increases data integrity by acting as a constraint. For example, if you use a datetime data type for a column of dates, then only dates can be stored in that column. However, if you use a character or numeric data type for the column, then eventually someone will store a character or numeric value that does not represent a date. You could write code to prevent this problem, but it is more efficient to use the correct data type. Therefore, store characters in character data types, numbers in numeric data types, and dates and times in datetime data types.

See Also:

Maintaining Data Integrity in Database Applications, for information about data integrity and constraints

9.1.2 How the Most Specific Data Type Decreases Storage Requirements

In addition to using the correct data type, use the most specific length or precision; for example:

- When creating a VARCHAR2 column intended for strings of at most *n* characters, specify VARCHAR2 (*n*).
- When creating a column intended for integers, use the data type NUMBER (38) rather than NUMBER.

Besides acting as constraints and thereby increasing data integrity, length and precision affect storage requirements.

If you give every column the maximum length or precision for its data type, then your application needlessly allocates many megabytes of RAM. For example, suppose that a query selects 10 VARCHAR2 (4000) columns and a bulk fetch operation returns 100 rows. The RAM that your application must allocate is $10 \times 4,000 \times 100$ —almost 4 MB. In contrast, if the column length is 80, the RAM that your application must allocate is $10 \times 80 \times 100$ —about 78 KB. This difference is significant for a single query, and your application will process many queries concurrently. Therefore, your application must allocate the 4 MB or 78 KB of RAM *for each connection*.

Therefore, do not give a column the maximum length or precision for its data type only because you might need to increase that property later. If you must change a column after creating it, then use the ALTER TABLE statement. For example, to increase the length of a column, use:

ALTER TABLE table_name MODIFY column_name VARCHAR2(larger_number)



Note:

The maximum length of the VARCHAR2, NVARCHAR2, and RAW data types is 32,767 bytes if the MAX_STRING_SIZE initialization parameter is EXTENDED.

See Also:

- Oracle Database SQL Language Reference for more information about ALTERTABLE
- Oracle Database SQL Language Reference for more information about extended data types

9.1.3 How the Correct Data Type Improves Performance

The correct data type improves performance because the incorrect data type can result in the incorrect execution plan.

Example 9-1 performs the same conceptual operation—selecting rows whose dates are between December 31, 2000 and January 1, 2001—for three columns with different data types and shows the execution plan for each query. In the three execution plans, compare Rows (cardinality), Cost, and Operation.

Example 9-1 Performance Comparison of Three Data Types

Create a table that stores the same dates in three columns: str_date, with data type VARCHAR2; date date, with data type DATE, and number date, with data type NUMBER:

Create an index on each column:

```
CREATE INDEX t_str_date_idx ON t(str_date);
CREATE INDEX t_date_date_idx ON t(date_date);
CREATE INDEX t_number_date_idx ON t(number_date);
```

Gather statistics for the table:

```
BEGIN
   DBMS_STATS.GATHER_TABLE_STATS (
    'HR',
    'T',
    method_opt => 'for all indexed columns size 254',
    cascade => TRUE
   );
```



```
END;
```

Show the execution plans of subsequent SQL statements (SQL*Plus command):

```
SET AUTOTRACE ON EXPLAIN
```

Select the rows for which the dates in str_date are between December 31, 2000 and January 1, 2001:

```
SELECT * FROM t WHERE str_date BETWEEN '20001231' AND '20010101' ORDER BY str date;
```

Result and execution plan:

2 rows selected.

Execution Plan

Plan hash value: 948745535

Id	Operation	Name		Rows		Bytes		Cost	(%CPU)	Time	
1 1 1	SELECT STATEMENT SORT ORDER BY TABLE ACCESS FUL	İ	İ	236	İ	11092	İ	216	(8)	00:00:01	İ

Predicate Information (identified by operation id):

```
2 - filter("STR DATE"<='20010101' AND "STR DATE">='20001231')
```

Select the rows for which the dates in number_date are between December 31, 2000 and January 1, 2001:

```
SELECT * FROM t WHERE number_date BETWEEN 20001231 AND 20010101; ORDER BY str_date;
```

Result and execution plan:

Plan hash value: 948745535

2 rows selected.

Execution Plan



```
|* 2 | TABLE ACCESS FULL| T | 234 | 10998 | 218 (9)| 00:00:01 |
Predicate Information (identified by operation id):
_____
  2 - filter("NUMBER DATE"<=20010101 AND "NUMBER DATE">=20001231)
Select the rows for which the dates in date date are between December 31, 2000 and
January 1, 2001:
SELECT * FROM t WHERE date date
 BETWEEN TO DATE ('20001231', 'yyyymmdd')
 AND TO DATE('20010101','yyyymmdd');
 ORDER BY str date;
Result and execution plan (reformatted to fit the page):
STR DATE DATE DATE NUMBER DATE DATA
20001231 31-DEC-00 20001231 *********************************
2 rows selected.
Execution Plan
Plan hash value: 2411593187

      Id
      | Operation
      | Name
      | Rows
      | Bytes
      |

      1 | SORT ORDER BY
      | 1 | 47 |

      2 | TABLE ACCESS BY INDEX ROWID BATCHED | T
      | 1 | 47 |

| Id | Operation
Cost (%CPU) | Time |
    4 (25) | 00:00:01 |
    4 (25) | 00:00:01 |
    3 (0) | 00:00:01 |
    2 (0) | 00:00:01 |
Predicate Information (identified by operation id):
  3 - access("DATE_DATE">=TO_DATE(' 2000-12-31 00:00:00',
  'syyyy-mm-dd hh24:mi:ss') AND
  "DATE DATE"<=TO DATE(' 2001-01-01 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
```

Performance improved for the final query because, for the DATE data type, the optimizer could determine that there was only one day between December 31, 2000 and January 1, 2001. Therefore, it performed an index range scan, which is faster than a full table scan.

See Also:

- EXPLAIN PLAN Statement
- Oracle Database SQL Tuning Guide for more information about Full Table Scans
- Oracle Database SQL Tuning Guide for more information about Index Range Scans

9.2 Representing Character Data

Table 9-1 summarizes the SQL data types that store character data.

Table 9-1 SQL Character Data Types

Data Types	Values Stored
CHAR	Fixed-length character literals
VARCHAR2	Variable-length character literals
NCHAR	Fixed-length Unicode character literals
NVARCHAR2	Variable-length Unicode character literals
CLOB	Single-byte and multibyte character strings of up to (4 gigabytes - 1) * (the value obtained from DBMS_LOB.GETCHUNKSIZE)
NCLOB	Single-byte and multibyte Unicode character strings of up to (4 gigabytes - 1) * (the value obtained from DBMS_LOB.GETCHUNKSIZE)
LONG	Variable-length character data of up to 2 gigabytes - 1. Provided only for backward compatibility.



Do not use the VARCHAR data type. Use the VARCHAR2 data type instead. Although the VARCHAR data type is currently synonymous with VARCHAR2, the VARCHAR data type is scheduled to be redefined as a separate data type used for variable-length character strings compared with different comparison semantics.

When choosing between CHAR and VARCHAR2, consider:

Space usage

Oracle Database blank-pads values stored in CHAR columns but not values stored in VARCHAR2 columns. Therefore, VARCHAR2 columns use space more efficiently than CHAR columns.

Performance

Because of the blank-padding difference, a full table scan on a large table containing VARCHAR2 columns might read fewer data blocks than a full table scan on a table containing the same data stored in CHAR columns. If your application often performs full table scans on large tables containing character data, then you might be able to improve performance by storing data in VARCHAR2 columns rather than in CHAR columns.



Comparison semantics

When you need ANSI compatibility in comparison semantics, use the CHAR data type. When trailing blanks are important in string comparisons, use the VARCHAR2 data type.

For a client/server application, if the character set on the client side differs from the character set on the server side, then Oracle Database converts CHAR, VARCHAR2, and LONG data from the database character set (determined by the NLS_LANGUAGE parameter) to the character set defined for the user session.

See Also:

- Oracle Database SQL Language Reference for more information about comparison semantics for these data types
- Large Objects (LOBs) for more information about CLOB and NCLOB data types
- LONG and LONG RAW Data Types for more information about LONG data type

9.3 Representing Numeric Data

The SQL data types that store numeric data are NUMBER, BINARY FLOAT, and BINARY DOUBLE.

The NUMBER data type stores real numbers in either a fixed-point or floating-point format. NUMBER offers up to 38 decimal digits of precision. In a NUMBER column, you can store positive and negative numbers of magnitude 1×10^{-130} through 9.99 $\times 10^{125}$, and 0. All Oracle Database platforms support NUMBER values.

The BINARY_FLOAT and BINARY_DOUBLE data types store floating-point numbers in the single-precision (32-bit) IEEE 754 format and the double-precision (64-bit) IEEE 754 format, respectively. High-precision values use less space when stored as BINARY_FLOAT and BINARY_DOUBLE than when stored as NUMBER. Arithmetic operations on floating-point numbers are usually faster for BINARY_FLOAT and BINARY_DOUBLE values than for NUMBER values.

In client interfaces that Oracle Database supports, arithmetic operations on BINARY_FLOAT and BINARY_DOUBLE values are performed by the native instruction set that the hardware vendor supplies. The term native floating-point data type includes BINARY_FLOAT and BINARY_DOUBLE data types and all implementations of these types in supported client interfaces.

Native floating-point data types conform substantially with the Institute of Electrical and Electronics Engineers (IEEE) Standard for Binary Floating-Point Arithmetic, IEEE Standard 754-1985 (IEEE754).

Note:

Oracle recommends using $BINARY_FLOAT$ and $BINARY_DOUBLE$ instead of FLOAT, a subtype of NUMBER.

Topics:

Floating-Point Number Components

- Floating-Point Number Formats
- Representing Special Values with Native Floating-Point Data Types
- Comparing Native Floating-Point Values
- Arithmetic Operations with Native Floating-Point Data Types
- Conversion Functions for Native Floating-Point Data Types
- Client Interfaces for Native Floating-Point Data Types



Oracle Database SQL Language Reference for more information about data types

9.3.1 Floating-Point Number Components

The formula for a floating-point value is:

(-1) sign. significand.base exponent

For example, the floating-point value 4.31 is represented:

 $(-1)^{0}.431.10^{-2}$

The components of the preceding representation are:

Component Name	Component Value
Sign	0
Significand	431
Base	10
Exponent	-2

9.3.2 Floating-Point Number Formats

A floating-point number format specifies how the components of a floating-point number are represented, thereby determining the range and precision of the values that the format can represent. The **range** is the interval bounded by the smallest and largest values and the **precision** is the number of significant digits. Both range and precision are finite. If a floating-point number is too precise for a given format, then the number is rounded.

How the number is rounded depends on the base of its format, which can be either decimal or binary. A number stored in decimal format is rounded to the nearest decimal place (for example, 1000, 10, or 0.01). A number stored in binary format is rounded to the nearest binary place (for example, 1024, 512, or 1/64).

NUMBER values are stored in decimal format. For calculations that need decimal rounding, use the NUMBER data type.

Native floating-point values are stored in binary format.

Table 9-2 shows the range and precision of the IEEE 754 single- and double-precision formats and Oracle Database NUMBER. Range limits are expressed as positive numbers, but they also



apply to absolute values of negative numbers. (The notation "number e exponent" means $number * 10^{exponent}$.)

Table 9-2 Range and Precision of Floating-Point Data Types

Range and Precision	Single-precision 32-bit ¹	Double-precision 64-bit ¹	Oracle Database NUMBER Data Type
Maximum positive normal number	3.40282347e+38	1.7976931348623157e+308	< 1.0e126
Minimum positive normal number	1.17549435e-38	2.2250738585072014e-308	1.0e-130
Maximum positive subnormal number	1.17549421e-38	2.2250738585072009e-308	not applicable
Minimum positive subnormal number	1.40129846e-45	4.9406564584124654e-324	not applicable
Precision (decimal digits)	6 - 9	15 - 17	38 - 40

¹ These numbers are from the *IEEE Numerical Computation Guide*.

9.3.2.1 Binary Floating-Point Formats

This formula determines the value of a floating-point number that uses a binary format:

$$(-1)^{\text{sign}} 2^{\mathbb{E}}$$
 (bit₀ bit₁ bit₂ ... bit_{p-1})

Table 9-3 describes the components of the preceding formula.

Table 9-3 Binary Floating-Point Format Components

Component	Component Value
sign	0 or 1
E (exponent)	For single-precision (32-bit) data type, an integer from -126 through 127. For double-precision (64-bit) data type, an integer from -1022 through 1023.
bit _i	0 or 1. (The bit sequence represents a number in base 2.)
p (precision)	For single-precision data type, 24. For double-precision data type, 53.

The leading bit of the significand, b_0 , must be set (1), except for subnormal numbers (explained later). Therefore, the leading bit is not stored, and a binary format provides n bits of precision while storing only n-1 bits. The IEEE 754 standard defines the in-memory formats for single-precision and double-precision data types, as Table 9-4 shows.

Table 9-4 Summary of Binary Format Storage Parameters

Data Type	Sign Bit	Exponent Bits	Significand Bits	Total Bits
Single-precision	1	8	24 (23 stored)	32
Double-precision	1	11	53 (52 stored)	64



Note:

Oracle Database does not support the extended single- and double-precision formats that the IEEE 754 standard defines.

A significand whose leading bit is set is called **normalized**. The IEEE 754 standard defines **subnormal numbers** (also called **denormal numbers**) that are too small to represent with normalized significands. If the significand of a subnormal number were normalized, then its exponent would be too large. Subnormal numbers preserve this property: If x-y==0.0 (using floating-point subtraction), then x==y.

9.3.3 Representing Special Values with Native Floating-Point Data Types

The IEEE 754 standard supports the special values shown in Table 9-5.

Table 9-5 Special Values for Native Floating-Point Formats

Value	Meaning
+INF	Positive infinity
-INF	Negative infinity
+0	Positive zero
-0	Negative zero
NaN	Not a number

Each value in Table 9-5 is represented by a specific bit pattern, except NaN. NaN, the result of any undefined operation, is represented by many bit patterns. Some of these bits patterns have the sign bit set and some do not, but the sign bit has no meaning.

The IEEE 754 standard distinguishes between quiet NaNs (which do not raise additional exceptions as they propagate through most operations) and signaling NaNs (which do). The IEEE 754 standard specifies action for when exceptions are enabled and action for when they are disabled.

In Oracle Database, exceptions cannot be enabled. Oracle Database acts as the IEEE 754 standard specifies for when exceptions are disabled. In particular, Oracle Database does not distinguish between quiet and signaling NaNs. You can use Oracle Call Interface (OCI) to retrieve NaN values from Oracle Database, but whether a retrieved NaN value is signaling or quiet depends on the client platform and is beyond the control of Oracle Database.

The IEEE 754 standard defines these classes of special values:

- Zero
- Subnormal
- Normal
- Infinity
- NaN

The values in each class in the preceding list are larger than the values in the classes that precede it in the list (ignoring signs), except NaN. NaN is unordered with other classes of special values and with itself.

In Oracle Database:

- All NaNs are quiet.
- Any non-NaN value < NaN
- Any NaN == any other NaN
- All Nans are converted to the same bit pattern.
- -0 is converted to +0.
- IEEE 754 exceptions are not raised.

See Also:

Oracle Database SQL Language Reference for information about floating-point conditions, which let you determine whether an expression is infinite or is the undefined result of an operation (is not a number or NaN).

9.3.4 Comparing Native Floating-Point Values

When comparing numeric expressions, Oracle Database uses numeric precedence to determine whether the condition compares NUMBER, BINARY FLOAT, or BINARY DOUBLE values.

Comparisons ignore the sign of zero (-0 equals +0).

See Also:

- Oracle Database SQL Language Reference for more information about Numeric Precedence
- Oracle Database SQL Language Reference for more information about Comparison Conditions

9.3.5 Arithmetic Operations with Native Floating-Point Data Types

IEEE 754 does not require floating-point arithmetic to be exactly reproducible. Therefore, results of operations can be delivered to a destination that uses a range greater than the range that the operands of the operation use.

You can compute the result of a double-precision multiplication at an extended double-precision destination, but the result must be rounded as if the destination were single-precision or double-precision. The range of the result (that is, the number of bits used for the exponent) can use the range supported by the wider (extended double-precision) destination; however, this might cause a double-rounding error in which the least significant bit of the result is incorrect.

This situation can occur only for double-precision multiplication and division on hardware that implements the IA-32 and IA-64 instruction set architecture. Therefore, except for this case, arithmetic for these data types is reproducible across platforms. When the result of a computation is NaN, all platforms produce a value for which IS NAN is true. However, all platforms do not have to use the same bit pattern.



See Also:

Oracle Database SQL Language Reference for general information about arithmetic operations

9.3.6 Conversion Functions for Native Floating-Point Data Types

Oracle Database defines functions that convert between floating-point and other data types, including string formats that use decimal precision (but precision might be lost during the conversion). For example:

- TO BINARY DOUBLE, described in Oracle Database SQL Language Reference
- TO BINARY FLOAT, described in Oracle Database SQL Language Reference
- TO CHAR, described in Oracle Database SQL Language Reference
- TO NUMBER, described in Oracle Database SQL Language Reference

Oracle Database can raise exceptions during conversion. The IEEE 754 standard defines these exceptions:

- Invalid
- Inexact
- Divide by zero
- Underflow
- Overflow

However, Oracle Database does not raise these exceptions for native floating-point data types. Generally, operations that raise exceptions produce the values described in Table 9-6.

Table 9-6 Values Resulting from Exceptions

Exception	Value
Underflow	0
Overflow	-INF, +INF
Invalid Operation	NaN
Divide by Zero	-INF, +INF, NaN
Inexact	Any value – rounding was performed

9.3.7 Client Interfaces for Native Floating-Point Data Types

Oracle Database supports native floating-point data types in these client interfaces:

SQL and PL/SQL

Support for BINARY_FLOAT and BINARY_DOUBLE includes their use as attributes of Abstract Data Types (ADTs), which you create with the SQL statement CREATE TYPE (fully described in *Oracle Database PL/SQL Language Reference*).

Oracle Call Interface (OCI)

For information about using BINARY_FLOAT and BINARY_DOUBLE with OCI, see Oracle Call Interface Programmer's Guide.

Oracle C++ Call Interface (OCCI)

For information about using BINARY_FLOAT with OCCI, see *Oracle C++ Call Interface Programmer's Guide*.

For information about using BINARY_DOUBLE with OCCI, see *Oracle C++ Call Interface Programmer's Guide*.

Pro*C/C++ precompiler

To use BINARY_FLOAT and BINARY_DOUBLE, set the Pro*C/C++ precompiler command line option NATIVE_TYPES to YES when you compile your application. For information about the NATIVE TYPES option, see Pro*C/C++ Programmer's Guide.

Oracle JDBC

For information about using BINARY_FLOAT and BINARY_DOUBLE with Oracle JDBC, see Oracle Database JDBC Developer's Guide.

9.4 Representing Date and Time Data

Oracle Database stores DATE and TIMESTAMP (datetime) data in a binary format that represents the century, year, month, day, hour, minute, second, and optionally, fractional seconds and timezones.

Table 9-7 summarizes the SQL datetime data types.

Table 9-7 SQL Datetime Data Types

Date Type	Usage
DATE	For storing datetime values in a table—for example, dates of jobs.
TIMESTAMP	For storing datetime values that are precise to fractional seconds—for example, times of events that must be compared to determine the order in which they occurred.
TIMESTAMP WITH TIME ZONE	For storing datetime values that must be gathered or coordinated across geographic regions.
TIMESTAMP WITH LOCAL TIME ZONE	For storing datetime values when the time zone is insignificant—for example, in an application that schedules teleconferences, where participants see the start and end times for their own time zone.
	Appropriate for two-tier applications in which you want to display dates and times that use the time zone of the client system. Usually inappropriate for three-tier applications, because data displayed in a web browser is formatted according to the time zone of the web server, not the time zone of the browser. The web server is the database client, so its local time is used.
INTERVAL YEAR TO MONTH	For storing the difference between two datetime values, where only the year and month are significant—for example, to set a reminder for a date 18 months in the future, or check whether 6 months have elapsed since a particular date.



Table 9-7 (Cont.) SQL Datetime Data Types

Date Type	Usage
INTERVAL DAY TO SECOND	For storing the precise difference between two datetime values—for example, to set a reminder for a time 36 hours in the future or to record the time between the start and end of a race. To represent long spans of time with high precision, use a large number of days.

✓ See Also:

- Oracle Call Interface Programmer's Guide for more information about Oracle Database internal date types
- Oracle Database SQL Language Reference for more information about date and time data types

Topics:

- Displaying Current Date and Time
- · Inserting and Displaying Dates
- Inserting and Displaying Times
- Arithmetic Operations with Datetime Data Types
- Conversion Functions for Datetime Data Types
- Importing_ Exporting_ and Comparing Datetime Types

9.4.1 Displaying Current Date and Time

The simplest way to display the current date and time is:

SELECT TO CHAR(SYSDATE, format model) FROM DUAL

The default format model depends on the initialization parameter NLS DATE FORMAT.

The standard Oracle Database default date format is DD-MON-RR. The RR datetime format element lets you store 20th century dates in the 21st century by specifying only the last two digits of the year. For example, in the datetime format DD-MON-YY, 13-NOV-54 refers to the year 1954 in a query issued between 1950 and 2049, but to the year 2054 in a query issued between 2050 and 2149.



For program correctness and to avoid problems with SQL injection and dynamic SQL, Oracle recommends specifying a format model for every datetime value.

The simplest way to display the current date and time using a format model is:

SELECT TO CHAR(SYSDATE, format model) FROM DUAL

Example 9-2 uses TO_CHAR with a format model to display SYSDATE in a format with the qualifier BC or AD. (By default, SYSDATE is displayed without this qualifier.)

Example 9-2 Displaying Current Date and Time

```
SELECT TO CHAR (SYSDATE, 'DD-MON-YYYY BC') NOW FROM DUAL;
```

Result:

NOW ------18-MAR-2009 AD

1 row selected.



Tip:

When testing code that uses SYSDATE, it can be helpful to set SYSDATE to a constant. Do this with the initialization parameter FIXED_DATE.

See Also:

- Oracle Database SQL Language Reference for more information about SYSDATE
- Oracle Database Globalization Support Guide for information about NLS DATE FORMAT
- Oracle Database SQL Language Reference for more information about TO_CHAR
- Oracle Database SQL Language Reference for information about datetime format models
- Oracle Database SQL Language Reference for more information about the RR datetime format element
- Oracle Database Referencefor more information about FIXED DATE

9.4.2 Inserting and Displaying Dates

When you display and insert dates, Oracle recommends using the ${\tt TO_CHAR}$ and ${\tt TO_DATE}$ functions, respectively, with datetime format models.

Example 9-3 creates a table with a DATE column and inserts a date into it, specifying a format model. Then the example displays the date with and without specifying a format model.

Example 9-3 Inserting and Displaying Dates

Create table:

DROP TABLE dates; CREATE TABLE dates (d DATE);

Insert date specified into table, specifying a format model:

A

1998-OCT-27

1 row selected.

Caution:

Be careful when using the YY datetime format element, which indicates the year in the current century. For example, in the 21st century, the format DD-MON-YY, 31-DEC-92 is December 31, 2092 (not December 31, 1992, as you might expect). To store 20th century dates in the 21st century by specifying only the last two digits of the year, use the RR datetime format element (the default).

See Also:

- Oracle Database Globalization Support Guide for information about NLS_DATE_FORMAT
- Oracle Database SQL Language Reference for more information about TO CHAR
- Oracle Database SQL Language Reference for more information about TO DATE
- Oracle Database SQL Language Reference for information about datetime format models
- Oracle Database SQL Language Reference for more information about the RR datetime format element

9.4.3 Inserting and Displaying Times

When you display and insert times, Oracle recommends using the ${\tt TO_CHAR}$ and ${\tt TO_DATE}$ functions, respectively, with datetime format models.

In a date column:

The default time is 12:00:00 A.M. (midnight).

The default time applies to any value in the column that has no time portion, either because none was specified or because the value was truncated.

The default day is the first day of the current month.

The default date applies to any value in the column that has no date portion, because none was specified.

Example 9-4 creates a table with a DATE column and inserts three dates into it, specifying a different format model for each date. The first format model has both date and time portions, the second has no time portion, and the third has no date portion. Then the example displays the three dates, specifying a format model that includes both date and time portions.

Example 9-4 Inserting and Displaying Dates and Times

Create table:

```
DROP TABLE birthdays; CREATE TABLE birthdays (name VARCHAR2(20), day DATE);
```

Insert three dates, specifying a different format model for each date:

Display both date and time portions of stored datetime values:

```
SELECT name,
    TO_CHAR(day, 'Mon DD, RRRR') DAY,
    TO_CHAR(day, 'HH:MI A.M.') TIME
FROM birthdays;
```

Result:

NAME	DAY	TIME
Annie	Nov 13, 1992	10:56 A.M.
Bobby	Apr 05, 2002	12:00 A.M.
Cindy	Nov 01, 2010	08:25 P.M.

3 rows selected.

9.4.4 Arithmetic Operations with Datetime Data Types

The results of arithmetic operations on datetime values are determined by the rules in Oracle Database SQL Language Reference.

SQL has many datetime functions that you can use in datetime expressions. For example, the function ${\tt ADD_MONTHS}$ returns the date that is a specified number of months from a specified date. .

See Also:

- Oracle Database SQL Language Reference for the complete list of datetime functions
- Oracle Database SQL Language Reference

9.4.5 Conversion Functions for Datetime Data Types

Table 9-8 summarizes the SQL functions that convert to or from datetime data types. For more information about these functions, see *Oracle Database SQL Language Reference*.

Table 9-8 SQL Conversion Functions for Datetime Data Types

Function	Converts	То
NUMTODSINTERVAL	NUMBER	INTERVAL DAY TO SECOND
NUMTOYMINTERVAL	NUMBER	INTERVAL DAY TO MONTH
TO_CHAR	DATE	VARCHAR2
	TIMESTAMP	
	TIMESTAMP WITH TIME ZONE	
	TIMESTAMP WITH LOCAL TIME ZONE	
	INTERVAL DAY TO SECOND	
	INTERVAL YEAR TO MONTH	
TO_DATE	CHAR	DATE
	VARCHAR2	
	NCHAR	
	NVARCHAR2	
TO_DSINTERVAL	CHAR	INTERVAL DAY TO SECOND
	VARCHAR2	
	NCHAR	
	NVARCHAR2	
TO_TIMESTAMP	CHAR	TIMESTAMP
	VARCHAR2	
	NCHAR	
	NVARCHAR2	



Table 9-8 (Cont.) SQL Conversion Functions for Datetime Data Types

Function	Converts	То
TO_TIMESTAMP_TZ	CHAR	TIMESTAMP WITH TIME ZONE
	VARCHAR2	
	NCHAR	
	NVARCHAR2	
TO_YMINTERVAL	CHAR	INTERVAL DAY TO MONTH
	VARCHAR2	
	NCHAR	
	NVARCHAR2	

9.4.6 Importing, Exporting, and Comparing Datetime Types

You can import, export, and compare TIMESTAMP WITH TIME ZONE and TIMESTAMP WITH LOCAL TIME ZONE values without worrying about time zone offsets, because the database stores these values in normalized format.

When importing, exporting, and comparing DATE and TIMESTAMP values, you must adjust them to account for any time zone differences between source and target databases, because the database does not store their time zones.

9.5 Representing Specialized Data

Topics:

- · Representing Spatial Data
- Representing Large Amounts of Data
- Representing Searchable Text
- Representing XML Data
- Representing Dynamically Typed Data
- Representing ANSI_ DB2_ and SQL/DS Data

9.5.1 Representing Spatial Data

Spatial data is used by location-enabled applications, geographic information system (GIS) applications, and geoimaging applications.



Oracle Database SQL Language Reference for information about representing spatial data in Oracle Database

9.5.2 Representing Large Amounts of Data

For representing large amounts of data, Oracle Database provides:

- Large Objects (LOBs)
- LONG and LONG RAW Data Types (for backward compatibility)

9.5.2.1 Large Objects (LOBs)

Large Objects (LOBs) are data types that are designed to store large amounts of data in a way that lets your application access and manipulate it efficiently.

Table 9-9 summarizes the LOBs.

Table 9-9 Large Objects (LOBs)

Data Type	Description
BLOB	Binary large object
	Stores any kind of data in binary format.
	Typically used for multimedia data such as images, audio, and video.
CLOB	Character large object
	Stores string data in the database character set format.
	Used for large strings or documents that use the database character set exclusively.
NCLOB	National character large object
	Stores string data in National Character Set format.
	Used for large strings or documents in the National Character Set.
BFILE	External large object
	Stores a binary file outside the database in the host operating system file system. Applications have read-only access to BFILEs.
	Used for static data that applications do not manipulate, such as image data.
	Any kind of data (that is, any operating system file) can be stored in a ${\tt BFILE}$. For example,
	you can store character data in a BFILE and then load the BFILE data into a CLOB, specifying the character set when loading.

An instance of type BLOB, CLOB, or NCLOB can be either **temporary** (declared in the scope of your application) or **persistent** (created and stored in the database).

See Also:

- Oracle Database SecureFiles and Large Objects Developer's Guide for information about using LOBs in application development
- Oracle Database SQL Language Reference for more information about LOB functions



9.5.2.2 LONG and LONG RAW Data Types



All forms of LONG data types (LONG, LONG RAW, LONG VARCHAR, LONG VARRAW) were deprecated in Oracle8i Release 8.1.6. For succeeding releases, the LONG data type was provided for backward compatibility with existing applications. In new applications developed with later releases, Oracle strongly recommends that you use CLOB and NCLOB data types for large amounts of character data.

For more information, see:

Migrating Columns from LONGs to LOBs

LONG columns store variable-length character strings containing up to 2 gigabytes - 1 bytes. .

The LONG RAW (and RAW) data types store data that is not to be explicitly converted by Oracle Database when moving data between different systems. These data types are intended for binary data or byte strings.

See Also:

Oracle Database SQL Language Reference for more information about data types

9.5.3 Representing JSON Data

A new JSON data type to natively store JSON data in Oracle Database.

Oracle Database stores JSON data type to natively store JSON data.

If you have information stored in JSON format, then you can use the Oracle-supplied type $\tt JSON$ type to natively store in the database. Oracle provides indexing, a rich set of packages and operators that can operate on JSON data.

With JSON values, you can use:

- JSON type standard functions
- IS JSON/IS NOT JSON constraints

See Also:

 Oracle Database JSON Developer's Guide for information about Oracle JSON data type and how you can use it to store, generate, manipulate, manage, and query JSON data in the database



9.5.4 Representing Searchable Text

Rather than writing low-level code to do full-text searches, you can use Oracle Text. Oracle Text provides indexing, word and theme searching, and viewing capabilities for text in query applications and document classification applications. You can also use Oracle Text to search XML data.

See Also:

Oracle Text Application Developer's Guide for more information about Oracle Text

9.5.5 Representing XML Data

If you have information stored as files in XML format, or want to store an ADT in XML format, then you can use the Oracle-supplied type XMLType.

With XMLType values, you can use:

- XMLType member functions (see Oracle XML DB Developer's Guide).
- SQL XML functions (see Oracle Database SQL Language Reference)
- PL/SQL DBMS_XML packages (see Oracle Database PL/SQL Packages and Types Reference)

See Also:

- Oracle XML DB Developer's Guide for information about Oracle XML DB and how you can use it to store, generate, manipulate, manage, and query XML data in the database
- Oracle XML Developer's Kit Programmer's Guide
- Oracle Database SQL Language Reference for more information about XMLType

9.5.6 Representing Dynamically Typed Data

Some languages allow data types to change at runtime, and some let a program check the type of a variable. For example, C has the union keyword and the void * pointer, and Java has the typeof operator and wrapper types such as Number.

In Oracle Database, you can create variables and columns that can hold data of any type and test their values to determine their underlying representation. For example, a single table column can have a numeric value in one row, a string value in another row, and an object in another row.

You can use the Oracle-supplied ADT SYS.ANYDATA to represent values of any scalar type or ADT. SYS.ANYDATA has methods that accept scalar values of any type, and turn them back into scalars or objects. Similarly, you can use the Oracle-supplied ADT SYS.ANYDATASET to represent values of any collection type.

To check and manipulate type information, use the DBMS TYPES package, as in Example 9-5.

With OCI, use the OCIAnyData and OCIAnyDataSet interfaces.

Example 9-5 Accessing Information in a SYS.ANYDATA Column

```
CREATE OR REPLACE TYPE employee type AS
 OBJECT (empno NUMBER, ename VARCHAR2(10));
DROP TABLE mytab;
CREATE TABLE mytab (id NUMBER, data SYS.ANYDATA);
INSERT INTO mytab (id, data)
VALUES (1, SYS.ANYDATA.ConvertNumber(5));
INSERT INTO mytab (id, data)
VALUES (2, SYS.ANYDATA.ConvertObject(Employee_type(5555, 'john')));
CREATE OR REPLACE PROCEDURE p IS
 CURSOR cur IS SELECT id, data FROM mytab;
 v id
                             mytab.id%TYPE;
 v data
                              mytab.data%TYPE;
 v_type
                              SYS.ANYTYPE;
 v_typecode
                             PLS INTEGER;
 v_typename
                              VARCHAR2 (60);
                              PLS INTEGER;
 v dummy
                              NUMBER;
 v n
 v employee
                              employee type;
 non_null_anytype_for_NUMBER exception;
 unknown typename
                             exception;
BEGIN
  FOR x IN cur LOOP
   FETCH cur INTO v id, v data;
   EXIT WHEN cur%NOTFOUND;
    /* typecode signifies type represented by v data.
       GetType also produces a value of type SYS.ANYTYPE with methods you
       can call to find precision and scale of a number, length of a
       string, and so on. */
       v typecode := v data.GetType (v type /* OUT */);
    /* Compare typecode to DBMS TYPES constants to determine type of data
       and decide how to display it. */
    CASE v typecode
      WHEN DBMS TYPES.TYPECODE NUMBER THEN
        IF v type IS NOT NULL THEN -- This condition should never happen.
          RAISE non_null_anytype_for_NUMBER;
        END IF;
      -- For each type, there is a Get method.
      v dummy := v data.GetNUMBER (v n /* OUT */);
      DBMS OUTPUT.PUT LINE
        (TO_CHAR(v_id) || ': NUMBER = ' || TO_CHAR(v_n) );
       WHEN DBMS TYPES.TYPECODE OBJECT THEN
         v typename := v data.GetTypeName();
         IF v typename NOT IN ('HR.EMPLOYEE TYPE') THEN
          RAISE unknown typename;
         END IF;
         v dummy := v data.GetObject (v employee /* OUT */);
```

```
DBMS OUTPUT.PUT LINE
           (TO CHAR(v id) || ': user-defined type = ' || v typename ||
            '(' || v_employee.empno || ', ' || v_employee.ename || ')');
   END CASE:
 END LOOP;
EXCEPTION
 WHEN non null anytype for NUMBER THEN
   RAISE Application Error (-20000,
      'Paradox: the return AnyType instance FROM GetType ' ||
      'should be NULL for all but user-defined types');
 WHEN unknown typename THEN
   RAISE_Application_Error( -20000, 'Unknown user-defined type ' ||
     v_typename || ' - program written to handle only HR.EMPLOYEE_TYPE');
END;
SELECT t.data.gettypename() AS "Type Name" FROM mytab t;
Result:
Type Name
SYS.NUMBER
HR.EMPLOYEE TYPE
2 rows selected.
```

See Also:

- Oracle Database Object-Relational Developer's Guide for more information about these ADTs
- Oracle Call Interface Programmer's Guide for more information about OCIAnyData and OCIAnyDataSet interfaces
- Oracle Database PL/SQL Packages and Types Reference for more information about DBMS_TYPES package

9.5.7 Representing ANSI, DB2, and SQL/DS Data

SQL statements that create tables and clusters can use ANSI data types and data types from the IBM products SQL/DS and DB2 (except those noted after this paragraph). Oracle Database converts the ANSI or IBM data type to the equivalent Oracle data type, records the Oracle data type as the name of the column data type, and stores the column data in the Oracle data type.

Note:

SQL statements cannot use the SQL/DS and DB2 data types TIME, GRAPHIC, VARGRAPHIC, and LONG VARGRAPHIC, because they have no equivalent Oracle data types.

See Also:

Oracle Database SQL Language Reference for conversion details

9.6 Identifying Rows by Address

The fastest way to access the row of a database table is by its address, or **rowid**. If a row is larger than one data block, then its rowid identifies its initial row piece.

To see rowids, query the ROWID pseudocolumn. Each value in the ROWID pseudocolumn is a string that represents the address of a row. The data type of the string is either ROWID or UROWID.

Note:

The rowid for a row may change for a number of reasons, which may be user initiated or internally by the database engine. You cannot depend on the rowid to be pointing to the same row or a valid row at all after any of these operations has occurred.

See Also:

- Oracle Database Concepts for an overview of the ROWID pseudocolumn
- Oracle Database Concepts for an overview of rowid data types
- Oracle Database SQL Language Reference for more information about the ROWID pseudocolumn
- Oracle Database SQL Language Reference for more information about the ROWID data type
- Oracle Database SQL Language Reference for more information about the UROWID data type
- Oracle Call Interface Programmer's Guide for information about using the ROWID data type in C
- Pro*C/C++ Programmer's Guide for information about using the ROWID data type with the Pro*C/C++ precompiler
- Oracle Database JDBC Developer's Guide for information about using the ROWID data type in Java
- Oracle Database Concepts for more information about HCC

9.7 Displaying Metadata for SQL Operators and Functions

The dynamic performance view V\$SQLFN_METADATA displays metadata about SQL operators and functions. For every function that V\$SQLFN_METADATA displays, the dynamic performance view V\$SQLFN_ARG_METADATA has one row of metadata about each function argument. If a

These views let third-party tools leverage SQL functions without maintaining their metadata in the application layer.

Topics:

- ARGn Data Type
- DISP TYPE Data Type
- SQL Data Type Families

✓ See Also:

- Oracle Database Reference for more information about V\$SQLFN METADATA
- Oracle Database Reference for more information about V\$SQLFN ARG METADATA

9.7.1 ARGn Data Type

In the view V\$SQLFN_METADATA, the column DATATYPE is the data type of the function (that is, the data type that the function returns). This data type can be an Oracle data type, data type family, or ARGn. ARGn is the data type of the nth argument of the function. For example:

- The MAX function returns a value that has the data type of its first argument, so the MAX function has return data type ARG1.
- The DECODE function returns a value that has the data type of its third argument, so the DECODE function has data type ARG3.

See Also:

- SQL Data Type Families
- Oracle Database SQL Language Reference for more information about MAX function
- Oracle Database SQL Language Reference for more information about DECODE function

9.7.2 DISP_TYPE Data Type

In the view V\$SQLFN_METADATA, the column DISP_TYPE is the data type of an argument that can be any expression. An expression is either a single value or a combination of values and SQL functions that has a single value.



Table 9-10 Display Types of SQL Functions

Display Type	Description	Example
NORMAL	FUNC (A, B,)	LEAST(A,B,C)
ARITHMETIC	A FUNC B)	A+B
PARENTHESIS	FUNC ()	SYS_GUID()
RELOP	A FUNC B	A IN B
CASE_LIKE	CASE statement or DECODE decode	
NOPAREN	FUNC	SYSDATE

9.7.3 SQL Data Type Families

Often, a SQL function argument can have any data type in a data type family. Table 9-11 shows the SQL data type families and their member data types.

Table 9-11 SQL Data Type Families

Family	Data Types	
STRING	• CHARACTER	
	• VARCHAR2	
	• CLOB	
	• NCHAR	
	• NVARCHAR2	
	• NCLOB	
	• LONG	
NUMERIC	• NUMBER	
	• BINARY_FLOAT	
	• BINARY_DOUBLE	
DATETYPE	• DATE	
	• TIMESTAMP	
	TIMESTAMP WITH TIME ZONE	
	TIMESTAMP WITH LOCAL TIME ZONE	
	INTERVAL YEAR TO MONTH	
	INTERVAL DAY TO SECOND	
BINARY	• BLOB	
	• RAW	
	• LONG RAW	

