Using the PL/SQL Hierarchical Profiler

You can use the PL/SQL hierarchical profiler to identify bottlenecks and performance-tuning opportunities in PL/SQL applications.

The profiler reports the dynamic execution profile of a PL/SQL program organized by function calls, and accounts for SQL and PL/SQL execution times separately. No special source or compile-time preparation is required; any PL/SQL program can be profiled.

This chapter describes the PL/SQL hierarchical profiler and explains how to use it to collect and analyze profile data for a PL/SQL program.

Topics:

- Overview of PL/SQL Hierarchical Profiler
- Collecting Profile Data
- Understanding Raw Profiler Output
- Analyzing Profile Data
- plshprof Utility

16.1 Overview of PL/SQL Hierarchical Profiler

Nonhierarchical (**flat**) profilers record the time that a program spends within each subprogram —the **function time** or **self time** of each subprogram. Function time is helpful, but often inadequate. For example, it is helpful to know that a program spends 40% of its time in the subprogram <code>INSERT_ORDER</code>, but it is more helpful to know which subprograms call <code>INSERT_ORDER</code> often and the total time the program spends under <code>INSERT_ORDER</code> (including its descendant subprograms). Hierarchical profilers provide such information.

The PL/SQL hierarchical profiler:

- Reports the dynamic execution profile of your PL/SQL program, organized by subprogram calls
- Accounts for SQL and PL/SQL execution times separately
- Requires no special source or compile-time preparation
- Stores results in database tables (hierarchical profiler tables) for custom report generation by integrated development environment (IDE) tools (such as SQL Developer and third-party tools)
- Provides subprogram-level execution summary information, such as:
 - Number of calls to the subprogram
 - Time spent in the subprogram itself (function time or self time)
 - Time spent in the subprogram itself and in its descendent subprograms (subtree time)
 - Detailed parent-children information, for example:
 - * All callers of a given subprogram (parents)



- * All subprograms that a given subprogram called (children)
- * How much time was spent in subprogram x when called from y
- * How many calls to subprogram x were from y

The PL/SQL hierarchical profiler is implemented by the <code>DBMS_HPROF</code> package and has two components:

Data collection

The data collection component is an intrinsic part of the PL/SQL Virtual Machine. The DBMS_HPROF package provides APIs to turn hierarchical profiling on and off and write the **raw profiler output** to a file or raw profiler data table.

Analyzer

The analyzer component processes the raw profiler output and produce analyzed results. The analyzer component analyzes:

 Raw profiler data located in the raw profiler data file and raw profiler data table into HTML CLOB report, analyzed report file, and hierarchical profiler analysis tables.



To generate simple HTML reports directly from raw profiler output, without using the Analyzer, you can use the plshprof command-line utility.

16.2 Collecting Profile Data

To collect profile data from your PL/SQL program for the PL/SQL hierarchical profiler, follow these steps:

- Ensure that you have these privileges:
 - EXECUTE privilege on the DBMS HPROF package
 - WRITE privilege on the directory that you specify when you call DBMS HPROF.START PROFILING
- 2. Use the DBMS_HPROF.START_PROFILING PL/SQL API to start hierarchical profiler data collection in a session.
- 3. Run your PL/SQL program long enough to get adequate code coverage.

To get the most accurate measurements of elapsed time, avoid unrelated activity on the system on which your PL/SQL program is running.

 Use the DBMS_HPROF.STOP_PROFILING PL/SQL API to stop hierarchical profiler data collection.

Example 16-1 Profiling a PL/SQL Procedure

```
BEGIN
   /* Start profiling.
     Write raw profiler output to file test.trc in a directory
     that is mapped to directory object PLSHPROF_DIR
        (see note after example). */

   DBMS_HPROF.START_PROFILING('PLSHPROF_DIR', 'test.trc');
END;
```



```
/
-- Run procedure to be profiled
BEGIN
test;
END;
/
BEGIN
-- Stop profiling
DBMS_HPROF.STOP_PROFILING;
END;
/
```

Example 16-2 Profiling with Raw Profiler Data Table

```
analyze runid number;
trace id number;
-- create raw profiler data and analysis tables
-- call create tables with force it =>FALSE ( default) when
-- raw profiler data and analysis tables do not exist already
DBMS HPROF . CREATE TABLES ;
-- Start profiling .
-- Write raw profiler data in raw profiler data table
trace id := DBMS HPROF . START PROFILING ;
-- Run procedure to be profiled
test ;
-- Stop profiling
DBMS HPROF . STOP PROFILING ;
-- analyzes trace_id entry in raw profiler data table and writes
-- hierarchical profiler information in hprof 's analysis tables.
analyze runid := DBMS HPROF . ANALYZE(trace id );
END;
```

Consider this PL/SQL procedure, test:

```
CREATE OR REPLACE PROCEDURE test AUTHID DEFINER IS n NUMBER;

PROCEDURE foo IS
BEGIN
SELECT COUNT(*) INTO n FROM EMPLOYEES;
END foo;

BEGIN -- test
FOR i IN 1..3 LOOP
foo;
END LOOP;
END test;
/
```

Consider the PL/SQL procedure that analyzes and generates HTML CLOB report from raw profiler data table

```
declare
reportclob clob ;
trace_id number;
begin
-- create raw profiler data and analysis tables
-- force_it =>TRUE will dropped the tables if table existed
DBMS_HPROF . CREATE_TABLES (force_it =>TRUE );
-- Start profiling .
```



```
-- Write raw profiler data in raw profiler data table
trace_id := DBMS_HPROF . START_PROFILING;
-- Run procedure to be profiled
test;
-- Stop profiling
DBMS_HPROF . STOP_PROFILING;
-- analyzes trace_id entry in raw profiler data table and produce
-- analyzed HTML report in reportclob .
DBMS_HPROF .ANALYZE (trace_id , reportclob );
end;
//
```

The SQL script in Example 16-1 profiles the execution of the PL/SQL procedure test.

Note:

A directory object is an alias for a file system path name. For example, if you are connected to the database AS SYSDBA, this CREATE DIRECTORY statement creates the directory object PLSHPROF_DIR and maps it to the file system directory /private/plshprof/results:

```
CREATE DIRECTORY PLSHPROF_DIR as '/private/plshprof/results';
```

To run the SQL script in Example 16-1, you must have READ and WRITE privileges on both PLSHPROF_DIR and the directory to which it is mapped. if you are connected to the database AS SYSDBA, this GRANT statement grants READ and WRITE privileges on PLSHPROF_DIR to HR:

GRANT READ, WRITE ON DIRECTORY PLSHPROF DIR TO HR;

See Also:

- Oracle Database SecureFiles and Large Objects Developer's Guide for more information about using directory objects
- Oracle Database PL/SQL Packages and Types Reference for more information about DBMS_HPROF.START_PROFILING and DBMS_HPROF.STOP_PROFILING

16.3 Understanding Raw Profiler Output

Raw profiler output is intended to be processed by the analyzer component of the PL/SQL hierarchical profiler. However, even without such processing, it provides a simple function-level trace of the program. This topic explains how to understand raw profiler output.

Note:

The raw profiler format shown in this chapter is intended only to illustrate conceptual features of raw profiler output. Format specifics are subject to change at each Oracle Database release.

The SQL script in Example 16-1 wrote this raw profiler output to the file test.trc:

```
P#V PLSHPROF Internal Version 1.0
P#! PL/SQL Timer Started
P#C PLSQL."".""."_plsql_vm"
P#C PLSQL."".""." anonymous block"
P#X 77
P#C PLSQL."HR"."TEST"::7."TEST"#980980e97e42f8ec #1
P#X 4
P#C PLSQL."HR"."TEST"::7."TEST.FOO"#980980e97e42f8ec #4
P#C SQL."HR"."TEST"::7."__static_sql_exec_line6" #6."3r6qf2qhr3cm1"
P#! SELECT COUNT(*) FROM EMPLOYEES
P#X 279
P#R
P#X 58
P#R
P#X 3
P#C PLSQL."HR"."TEST"::7."TEST.FOO"#980980e97e42f8ec #4
P#C SQL."HR"."TEST"::7." static sql exec line6" #6."3r6qf2qhr3cm1"
P#! SELECT COUNT(*) FROM EMPLOYEES
P#X 121
P#R
P#X 5
P#R
P#X 2
P#C PLSQL."HR"."TEST"::7."TEST.FOO"#980980e97e42f8ec #4
P#C SQL."HR"."TEST"::7."__static_sql_exec_line6" #6."3r6qf2qhr3cm1"
P#! SELECT COUNT(*) FROM EMPLOYEES
P#X 117
P#R
P#X 4
P#R
P#X 2
P#R
P#X 2
P#R
P#X 3
P#R
P#C PLSQL."".""."_plsql_vm"
P#C PLSQL.""."". anonymous_block"
P#C PLSQL."SYS"."DBMS HPROF"::11."STOP PROFILING"#980980e97e42f8ec #453
P#R
P#R
P#R
P#! PL/SQL Timer Stopped
```

Table 16-1 Raw Profiler Output File Indicators

| Indicator | Meaning |
|-----------|---|
| P#V | PLSHPROF banner with version number |
| P#C | Call to a subprogram (call event) |
| P#R | Return from a subprogram (return event) |

Table 16-1 (Cont.) Raw Profiler Output File Indicators

| Indicator | Meaning |
|-----------|---|
| P#X | Elapsed time between preceding and following events |
| P#! | Comment |

Call events (P#C) and return events (P#R) are properly nested (like matched parentheses). If an unhandled exception causes a called subprogram to exit, the profiler still reports a matching return event.

Each call event (P#C) entry in the raw profiler output includes this information:

- Namespace to which the called subprogram belongs
- Name of the PL/SQL module in which the called subprogram is defined
- Type of the PL/SQL module in which the called subprogram is defined
- Name of the called subprogram

This name might be one of the special function names in Special Function Names.

 Hexadecimal value that represents the hash algorithm of the signature of the called subprogram

The DBMS_HPROF.analyze PL/SQL API stores the hash value in a hierarchical profiler table, which allows both you and DBMS_HPROF.analyze to distinguish between **overloaded subprograms** (subprograms with the same name).

Line number at which the called subprogram is defined in the PL/SQL module

PL/SQL hierarchical profiler does not measure time spent at individual lines within modules, but you can use line numbers to identify the source locations of subprograms in the module (as IDE/Tools do) and to distinguish between overloaded subprograms.

For example, consider this entry in the preceding example of raw profiler output:

```
P#C PLSQL."HR"."TEST"::7."TEST.FOO"#980980e97e42f8ec #4
```

The components of the preceding entry have these meanings:

| Component | Meaning |
|-------------------|--|
| PLSQL | PLSQL is the namespace to which the called subprogram belongs. |
| "HR"."TEST" | HR.TEST is the name of the PL/SQL module in which the called subprogram is defined. |
| 7 | 7 is the internal enumerator for the module type of HR.TEST. Examples of module types are procedure, package, and package body. |
| "TEST.FOO" | TEST.FOO is the name of the called subprogram. |
| #980980e97e42f8ec | #980980e97e42f8ec is a hexadecimal value that represents the hash algorithm of the signature of TEST.FOO. |
| #4 | 4 is the line number in the PL/SQL module ${\tt HR.TEST}$ at which ${\tt TEST.FOO}$ is defined. |



Note:

When a subprogram is inlined, it is not reported in the profiler output.

When a call to a DETERMINISTIC function is "optimized away," it is not reported in the profiler output.

See Also:

- Namespaces of Tracked Subprograms
- Analyzing Profile Data for more information about analyze PL/SQL API
- Oracle Database PL/SQL Language Reference for information about subprogram inlining
- Oracle Database PL/SQL Language Reference for information about DETERMINISTIC functions

16.3.1 Namespaces of Tracked Subprograms

The subprograms that the PL/SQL hierarchical profiler tracks are classified into the namespaces PLSQL and SQL, as follows:

- Namespace PLSQL includes:
 - PL/SQL subprogram calls
 - PL/SQL triggers
 - PL/SQL anonymous blocks
 - Remote subprogram calls
 - Package initialization blocks
- Namespace SQL includes SQL statements executed from PL/SQL, such as queries, data manipulation language (DML) statements, data definition language (DDL) statements, and native dynamic SQL statements

16.3.2 Special Function Names

PL/SQL hierarchical profiler tracks certain operations as if they were functions with the names and namespaces shown in Table 16-2.

Table 16-2 Function Names of Operations that the PL/SQL Hierarchical Profiler Tracks

| Function Name | Namespace |
|-----------------------------------|---|
| plsql_vm | PL/SQL |
| anonymous_block | PL/SQL |
| pkg_init | PL/SQL |
| static_sql_exec_line <i>line#</i> | SQL |
| | plsql_vm anonymous_block pkg_init |



Table 16-2 (Cont.) Function Names of Operations that the PL/SQL Hierarchical Profiler Tracks

| ction Name I | |
|--------------|------------------|
| | Namespace SQL |
| | SQL SQL |
| | 1 |

16.4 Analyzing Profile Data

The analyzer component of the PL/SQL hierarchical profiler, DBMS_HPROF.analyze, processes the raw profiler output and stores the results in the hierarchical database tables described in Table 16-3.

Table 16-3 PL/SQL Hierarchical Profiler Database Tables

| Table | Description |
|--------------------------|---|
| DBMSHP_RUNS | Top-level information for this run of DBMS_HPROF.analyze. For column descriptions, see Table 16-4. |
| DBMSHP_FUNCTION_INFO | Information for each subprogram profiled in this run of DBMS_HPROF.analyze. For column descriptions, see Table 16-5. |
| DBMSHP_PARENT_CHILD_INFO | Parent-child information for each subprogram profiled in this run of DBMS_HPROF.analyze. For column descriptions, see Table 16-6. |

Topics:

- Creating Hierarchical Profiler Tables
- Understanding Hierarchical Profiler Tables



To generate simple HTML reports directly from raw profiler output, without using the Analyzer, you can use the plshprof command-line utility. For details, see plshprof Utility.

16.4.1 Creating Hierarchical Profiler Tables

The following steps explain how to create hierarchical profiler tables in Table 16-3 and the other data structures required for persistently storing profile data, privileges required to run the DBMS HPROF package, and generate custom reports:

- 1. Hierarchical profiler tables in Table 16-3 and other data structures required for persistently storing profile data can be created in the following ways.
 - a. Call the DBMS_HPROF.CREATE_TABLES procedure.
 - b. Run the script dbmshptab.sql (located in the directory rdbms/admin).



Running the script ${\tt abmshptab.sql}$ drops any previously created hierarchical profiler tables.

Note:

The <code>dbmshptab.sql</code> (located in the directory <code>rdbms/admin</code>) has been deprecated. This script contains the statements to drop the tables and sequences along with the deprecation notes.

- 2. Ensure that you have these privileges:
 - EXECUTE privilege on the DBMS HPROF package
 - READ privilege on the directory that DBMS HPROF.analyze specifies
- 3. Use the PL/SQL API DBMS_HPROF.analyze to analyze a single raw profiler output file and store the results in hierarchical profiler tables.

(For an example of a raw profiler output file, see test.trc in Understanding Raw Profiler Output.)

For more information about DBMS_HPROF.analyze, see Oracle Database PL/SQL Packages and Types Reference.

4. Use the hierarchical profiler tables to generate custom reports.

Example 16-3 Invoking DBMS_HPROF.analyze

The anonymous block in Example 16-3:

- Invokes the function DBMS HPROF.analyze function, which:
 - Analyzes the profile data in the raw profiler output file test.trc (from Understanding Raw Profiler Output), which is in the directory that is mapped to the directory object PLSHPROF_DIR, and stores the data in the hierarchical profiler tables in Table 16-3.
 - Returns a unique identifier that you can use to query the hierarchical profiler tables in Table 16-3. (By querying these hierarchical profiler tables, you can produce customized reports.)
- Stores the unique identifier in the variable runid, which it prints.

16.4.2 Understanding Hierarchical Profiler Tables

These topics explain how to use the hierarchical profiler tables in Table 16-3:

- Hierarchical Profiler Database Table Columns
- Distinguishing Between Overloaded Subprograms
- Hierarchical Profiler Tables for Sample PL/SQL Procedure
- Examples of Calls to DBMS_HPROF.analyze with Options

16.4.2.1 Hierarchical Profiler Database Table Columns

Table 16-4 describes the columns of the hierarchical profiler table <code>DBMSHP_RUNS</code>, which contains one row of top-level information for each run of <code>DBMS HPROF.analyze</code>.

The primary key for the hierarchical profiler table DBMSHP RUNS is RUNID.

Table 16-4 DBMSHP_RUNS Table Columns

| Column Name | Column Data Type | Column Contents |
|--------------------|------------------|---|
| RUNID | NUMBER | Unique identifier for this run of DBMS_HPROF.analyze, generated from DBMSHP_RUNNUMBER sequence. |
| RUN_TIMESTAMP | TIMESTAMP(6) | Time stamp for this run of DBMS_HPROF.analyze. |
| RUN_COMMENT | VARCHAR2 (2047) | Comment that you provided for this run of DBMS_HPROF.analyze. |
| TOTAL_ELAPSED_TIME | NUMBER(38) | Total elapsed time for this run of DBMS_HPROF.analyze. |

Table 16-5 describes the columns of the hierarchical profiler table <code>DBMSHP_FUNCTION_INFO</code>, which contains one row of information for each subprogram profiled in this run of <code>DBMS_HPROF.analyze</code>. If a subprogram is overloaded, Table 16-5 has a row for each variation of that subprogram. Each variation has its own <code>LINE#</code> and <code>HASH</code> (see <code>Distinguishing Between Overloaded Subprograms.)</code>

The primary key for the hierarchical profiler table DBMSHP FUNCTION INFO is RUNID, SYMBOLID.

Table 16-5 DBMSHP_FUNCTION_INFO Table Columns

| Column Name | Column Data Type | Column Contents |
|-------------|------------------|---|
| RUNID | NUMBER | References RUNID column of DBMSHP_RUNS table. For a description of that column, see Table 16-4. |
| SYMBOLID | NUMBER | Symbol identifier for subprogram (unique for this run of DBMS_HPROF.analyze). |
| OWNER | VARCHAR2(128) | Owner of module in which each subprogram is defined (for example, SYS or HR). |
| MODULE | VARCHAR2 (128) | Module in which subprogram is defined (for example, DBMS_LOB, UTL_HTTP, or MY_PACKAGE). |
| TYPE | VARCHAR2(32) | Type of module in which subprogram is defined (for example, PACKAGE, PACKAGE_BODY, or PROCEDURE). |



Table 16-5 (Cont.) DBMSHP_FUNCTION_INFO Table Columns

| Column Name | Column Data Type | Column Contents |
|-----------------------|------------------|---|
| FUNCTION | VARCHAR2(4000) | Name of subprogram (not necessarily a function) (for example, INSERT_ORDER, PROCESS_ITEMS, or TEST). |
| | | This name might be one of the special function names in Special Function Names. |
| | | For subprogram B defined within subprogram A, this name is A.B. |
| | | A recursive call to function X is denoted $X@n$, where n is the recursion depth. For example, $X@1$ is the first recursive call to X . |
| LINE# | NUMBER | Line number in OWNER.MODULE at which FUNCTION is defined. |
| HASH | RAW(32) | Hash code for signature of subprogram (unique for this run of DBMS_HPROF.analyze). |
| NAMESPACE | VARCHAR2(32) | Namespace of subprogram. For details, see Namespaces of Tracked Subprograms. |
| SUBTREE_ELAPSED_TIME | NUMBER(38) | Elapsed time, in microseconds, for subprogram, including time spent in descendant subprograms. |
| FUNCTION_ELAPSED_TIME | NUMBER(38) | Elapsed time, in microseconds, for subprogram, excluding time spent in descendant subprograms. |
| CALLS | NUMBER (38) | Number of calls to subprogram. |
| SQL_ID | VARCHAR2(13) | SQL Identifier of the SQL statement. |
| SQL_TEXT | VARCHAR2(4000) | First 50 characters of the SQL statement. |

Table 16-6 describes the columns of the hierarchical profiler table <code>DBMSHP_PARENT_CHILD_INFO</code>, which contains one row of parent-child information for each unique parent-child subprogram combination profiled in this run of <code>DBMS_HPROF.analyze</code>.

Table 16-6 DBMSHP_PARENT_CHILD_INFO Table Columns

| Column Name | Column Data Type | Column Contents |
|-------------|------------------|--|
| RUNID | NUMBER | References RUNID column of DBMSHP_FUNCTION_INFO table. For a description of that column, see Table 16-5. |
| PARENTSYMID | NUMBER | Parent symbol ID. |
| | | RUNID, PARENTSYMID references DBMSHP_FUNCTION_INFO(RUNID, SYMBOLID). |
| CHILDSYMID | NUMBER | Child symbol ID. RUNID, CHILDSYMID references DBMSHP_FUNCTION_INFO(RUNID, SYMBOLID). |



Table 16-6 (Cont.) DBMSHP_PARENT_CHILD_INFO Table Columns

| Column Name | Column Data Type | Column Contents |
|-----------------------|------------------|---|
| SUBTREE_ELAPSED_TIME | NUMBER(38) | Elapsed time, in microseconds, for RUNID, CHILDSYMID when called from RUNID, PARENTSYMID, including time spent in descendant subprograms. |
| FUNCTION_ELAPSED_TIME | NUMBER(38) | Elapsed time, in microseconds, for RUNID, CHILDSYMID when called from RUNID, PARENTSYMID, excluding time spent in descendant subprograms. |
| CALLS | NUMBER (38) | Number of calls to RUNID, CHILDSYMID from RUNID, PARENTSYMID. |

16.4.2.2 Distinguishing Between Overloaded Subprograms

Overloaded subprograms are different subprograms with the same name.

Suppose that a program declares three subprograms named <code>compute</code>—the first at line 50, the second at line 76, and the third at line 100. In the <code>DBMSHP_FUNCTION_INFO</code> table, each of these subprograms has <code>compute</code> in the <code>FUNCTION</code> column. To distinguish between the three subprograms, use either the <code>LINE#</code> column (which has 50 for the first subprogram, 76 for the second, and 100 for the third) or the <code>HASH</code> column (which has a unique value for each subprogram).

In the profile data for two different runs, the LINE# and HASH values for a function might differ. The LINE# value of a function changes if you insert or delete lines before the function definition. The HASH value changes only if the signature of the function changes; for example, if you change the parameter list.



 ${\it Oracle\ Database\ PL/SQL\ Language\ Reference\ for\ more\ information\ about\ overloaded\ subprograms}$

16.4.2.3 Hierarchical Profiler Tables for Sample PL/SQL Procedure

The hierarchical profiler tables for the PL/SQL procedure test in Collecting Profile Data are shown in Example 16-4 through Example 16-6.

Consider the third row of the table <code>DBMSHP_PARENT_CHILD_INFO</code> (Example 16-6). The <code>RUNID</code> column shows that this row corresponds to the third run. The columns <code>PARENTSYMID</code> and <code>CHILDSYMID</code> show that the symbol IDs of the parent (caller) and child (called subprogram) are 2 and 1, respectively. The table <code>DBMSHP_FUNCTION_INFO</code> (Example 16-5) shows that for the third run, the symbol IDs 2 and 1 correspond to procedures <code>__plsql_vm</code> and <code>__anonymous_block</code>, respectively. Therefore, the information in this row is about calls from the procedure <code>__plsql_vm</code> to the <code>__anonymous_block</code> (defined within <code>__plsql_vm</code>) in the module <code>HR.TEST</code>. This row shows that, when called from the procedure <code>__plsql_vm</code>, the function time for the procedure <code>__anonymous_block</code> is 44 microseconds, and the time spent in the anonymous <code>block</code> subtree (including descendants) is 121 microseconds.

Example 16-4 DBMSHP_RUNS Table for Sample PL/SQL Procedure

| RUNID | RUN_TIMESTAMP | TOTAL_ELAPSED_TIME RUN_COMMENT |
|-------|------------------------------|--------------------------------|
| | | |
| 1 | 20-DEC-12 11.37.26.688381 AM | 7 |
| 2 | 20-DEC-12 11.37.26.700523 AM | 9 |
| 3 | 20-DEC-12 11.37.26.706824 AM | 123 |

Example 16-5 DBMSHP_FUNCTION_INFO Table for Sample PL/SQL Procedure

| RI | UNID | SYMBOLID | OWNER | MODULE | | TYPE | FUNCTION | |
|---------|-------|-----------|----------|--------|-------|-----------------|-----------------|-----------|
| | 1 | 1 | HR | PKG | | PACKAGE BODY | MYPROC | |
| | 2 | 1 | HR | PKG | | PACKAGE BODY | MYFUNC | |
| | 2 | 2 | HR | PKG | | PACKAGE BODY | MYPROC | |
| | 3 | 1 | | | | | anonymous_block | c |
| | 3 | 2 | | | | | plsql_vm | |
| | 3 | 3 | HR | PKG | | PACKAGE BODY | MYFUNC | |
| | 3 | | HR | | | PACKAGE BODY | MYPROC | |
| | 3 | 5 | HR | TEST2 | | PROCEDURE | TEST2 | |
| LINE# | CALLS | HASH | | | | | | NAMESPACE |
| 2 | 1 | 9689BA467 | A19CD19 | | | | | PLSQL |
| 7 | 1 | 28DC3402B | AEB2B0D | | | | | PLSQL |
| 2 | 1 | 9689BA467 | A19CD19 | | | | | PLSQL |
| 0 | 1 | | | | | | | PLSQL |
| 0 | 1 | | | | | | | PLSQL |
| 7 | 1 | 28DC3402B | AEB2B0D | | | | | PLSQL |
| 2 | 2 | 9689BA467 | A19CD19 | | | | | PLSQL |
| 1 | 1 | 980980E97 | E42F8EC | | | | | PLSQL |
| NAMESPA | ACE | SUBTREE | _ELAPSED | _TIME | FUNC' | rion_elapsed_ti | ME | |
| PLSQL | | | | 7 | | | 7 | |
| PLSQL | | | | 9 | | | 8 | |
| PLSQL | | | | 1 | | | 1 | |
| PLSQL | | | | 121 | | | 44 | |
| PLSQL | | | | 123 | | | 2 | |
| PLSQL | | | | 9 | | | 8 | |
| PLSQL | | | | 8 | | | 8 | |
| PLSQL | | | | 77 | | | 61 | |

Example 16-6 DBMSHP_PARENT_CHILD_INFO Table for Sample PL/SQL Procedure

| RUNID | PARENTSYMID | CHILDSYMID | SUBTREE_ELAPSED_TIME | FUNCTION_ELAPSED_TIME | CALLS |
|-------|-------------|------------|----------------------|-----------------------|-------|
| 2 | 1 | 2 | 1 | 1 | 1 |
| 3 | 1 | 5 | 77 | 61 | 1 |
| 3 | 2 | 1 | 121 | 44 | 1 |
| 3 | 3 | 4 | 1 | 1 | 1 |
| 3 | 5 | 3 | 9 | 8 | 1 |
| 3 | 5 | 4 | 7 | 7 | 1 |

16.4.2.4 Examples of Calls to DBMS_HPROF.analyze with Options

For an example of a call to DBMS HPROF.analyze without options, see Example 16-3.

Example 16-7 creates a package, creates a procedure that invokes subprograms in the package, profiles the procedure, and uses <code>DBMS_HRPROF.analyze</code> to analyze the raw profiler output file. The raw profiler output file is in the directory corresponding to the <code>PLSHPROF_DIR</code> directory object.

Example 16-7 Invoking DBMS_HPROF.analyze with Options

```
-- Create package
CREATE OR REPLACE PACKAGE pkg AUTHID DEFINER IS
  PROCEDURE myproc (n IN out NUMBER);
  FUNCTION myfunc (v VARCHAR2) RETURN VARCHAR2;
  FUNCTION myfunc (n PLS_INTEGER) RETURN PLS_INTEGER;
END pkg;
CREATE OR REPLACE PACKAGE BODY pkg IS
  PROCEDURE myproc (n IN OUT NUMBER) IS
  BEGIN
   n := n + 5;
  END;
  FUNCTION myfunc (v VARCHAR2) RETURN VARCHAR2 IS
   n NUMBER;
  BEGIN
   n := LENGTH(v);
   myproc(n);
   IF n > 20 THEN
     RETURN SUBSTR(v, 1, 20);
     RETURN v || '...';
    END IF;
  END;
  FUNCTION myfunc (n PLS_INTEGER) RETURN PLS_INTEGER IS
    i PLS INTEGER;
   PROCEDURE myproc (n IN out PLS_INTEGER) IS
   BEGIN
     n := n + 1;
   END;
  BEGIN
   i := n;
   myproc(i);
   RETURN i;
  END;
END pkg;
-- Create procedure that invokes package subprograms
CREATE OR REPLACE PROCEDURE test2 AUTHID DEFINER IS
 x NUMBER := 5;
  y VARCHAR2 (32767);
BEGIN
  pkg.myproc(x);
  y := pkg.myfunc('hello');
END;
-- Profile test2
BEGIN
  DBMS_HPROF.START_PROFILING('PLSHPROF_DIR', 'test2.trc');
END;
BEGIN
 test2;
END;
```



```
BEGIN
  DBMS HPROF.STOP PROFILING;
END:
-- If not done, create hierarchical profiler tables (see Creating Hierarchical Profiler Tables).
-- Call DBMS HPROF.analyze with options
DECLARE
  runid NUMBER;
BEGIN
  -- Analyze only subtrees rooted at trace entry "HR"."PKG"."MYPROC"
  runid := DBMS HPROF.analyze('PLSHPROF DIR', 'test2.trc',
                               trace => '"HR"."PKG"."MYPROC"');
  -- Analyze up to 20 calls to subtrees rooted at trace entry
  -- "HR". "PKG". "MYFUNC". Because "HR". "PKG". "MYFUNC" is overloaded,
  -- both overloads are considered for analysis.
  runid := DBMS HPROF.analyze('PLSHPROF DIR', 'test2.trc',
                               collect => 20,
                               trace => '"HR"."PKG"."MYFUNC"');
  -- Analyze second call to PL/SQL virtual machine
  runid := DBMS HPROF.analyze('PLSHPROF DIR', 'test2.trc',
                               skip \Rightarrow 1, collect \Rightarrow 1,
                               trace => '"".""."" plsql vm"');
END;
```

16.5 plshprof Utility

The plshprof command-line utility (located in the directory <code>\$ORACLE_HOME/bin/</code>) generates simple HTML reports from either one or two raw profiler output files. (For an example of a raw profiler output file, see <code>test.trc</code> in Collecting Profile Data.)

You can browse the generated HTML reports in any browser. The browser's navigational capabilities, combined with well chosen links, provide a powerful way to analyze performance of large applications, improve application performance, and lower development costs.

Topics:

- plshprof Options
- HTML Report from a Single Raw Profiler Output File
- HTML Difference Report from Two Raw Profiler Output Files

16.5.1 plshprof Options

The command to run the plshprof utility is:

```
plshprof [option...] profiler output filename 1 profiler output filename 2
```

Each option is one of these:

| Option | Description | Default |
|------------------|--|---|
| -skip count | Skips first count calls. Use only with - trace symbol. | 0 |
| -collect count | Collects information for count calls. Use only with -trace symbol. | 1 |
| -output filename | Specifies name of output file | <pre>symbol.html or tracefile1.html</pre> |
| -summary | Prints only elapsed time | None |
| -trace symbol | Specifies function name of tree root | Not applicable |

Suppose that your raw profiler output file, test.trc, is in the current directory. You want to analyze and generate HTML reports, and you want the root file of the HTML report to be named report.html. Use this command (% is the prompt):

16.5.2 HTML Report from a Single Raw Profiler Output File

To generate a PL/SQL hierarchical profiler HTML report from a single raw profiler output file, use these commands:

```
% cd target_directory
% plshprof -output html root filename profiler output filename
```

target directory is the directory in which you want the HTML files to be created.

html root filename is the name of the root HTML file to be created.

profiler output filename is the name of a raw profiler output file.

The preceding plshprof command generates a set of HTML files. Start browsing them from html root filename.html.

Topics:

- First Page of Report
- Function-Level Reports
- Understanding PL/SQL Hierarchical Profiler SQL-Level Reports
- Module-Level Reports
- Namespace-Level Reports
- Parents and Children Report for a Function

16.5.2.1 First Page of Report

The first page of an HTML report from a single raw profiler output file includes summary information and hyperlinks to other pages of the report.

Sample First Page

PL/SQL Elapsed Time (microsecs) Analysis

824 microsecs (elapsed time) & 12 function calls



[%] plshprof -output report test.trc

The PL/SQL Hierarchical Profiler produces a collection of reports that present information derived from the profiler output log in a variety of formats. These reports have been found to be the most generally useful as starting points for browsing:

- Function Elapsed Time (microsecs) Data sorted by Total Subtree Elapsed Time (microsecs)
- Function Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs)
- SQL ID Elapsed Time (microsecs) Data sorted by SQL ID

In addition, the following reports are also available:

- Function Elapsed Time (microsecs) Data sorted by Function Name
- Function Elapsed Time (microsecs) Data sorted by Total Descendants Elapsed Time (microsecs)
- Function Elapsed Time (microsecs) Data sorted by Total Function Call Count
- Function Elapsed Time (microsecs) Data sorted by Mean Subtree Elapsed Time (microsecs)
- Function Elapsed Time (microsecs) Data sorted by Mean Function Elapsed Time (microsecs)
- Function Elapsed Time (microsecs) Data sorted by Mean Descendants Elapsed Time (microsecs)
- Module Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs)
- Module Elapsed Time (microsecs) Data sorted by Module Name
- Module Elapsed Time (microsecs) Data sorted by Total Function Call Count
- Namespace Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs)
- Namespace Elapsed Time (microsecs) Data sorted by Namespace
- Namespace Elapsed Time (microsecs) Data sorted by Total Function Call Count
- Parents and Children Elapsed Time (microsecs) Data

16.5.2.2 Function-Level Reports

The function-level reports provide a flat view of the profile information. Each function-level report includes this information for each function:

- Function time (time spent in the function itself, also called "self time")
- Descendants time (time spent in the descendants of the function)
- Subtree time (time spent in the subtree of the function—function time plus descendants time)
- Number of calls to the function
- Function name

The function name is hyperlinked to the Parents and Children Report for the function.

- SQL ID
- SQL Text (First 50 characters of the SQL text).



Each function-level report is sorted on a particular attribute; for example, function time or subtree time.

This sample report is sorted in descending order of the total subtree elapsed time for the function, which is why information in the Subtree and Ind% columns is in **bold type**:

Sample Report

Function Elapsed Time (microsecs) Data sorted by Total Subtree Elapsed Time (microsecs)

824 microsecs (elapsed time) & 12 function calls

| Subtr | Ind % | Functi on | Descend ant | Ind % | Cal Is | Ind % | Function Name | SQL ID | SQL TEXT |
|-------|-----------|--------------|----------------|-----------|-----------|-----------|---|-------------------|---|
| 824 | 100 % | 10 | 814 | 98.8 % | 2 | 16.7 % | plsq_vm | - | |
| 814 | 98.8 % | 165 | 649 | 78.8 % | 2 | 16.7 % | anonymous_block | | |
| 649 | 78.8 % | 11 | 638 | 77.4 % | 1 | 8.3 % | HR.TEST.TEST (Line 1) | | |
| 638 | 77.4 % | 121 | 517 | 62.7 % | 3 | 25.0 % | HR.TEST.TEST.FOO (Line 4) | | |
| 517 | 62.7 % | 517 | 0 | 0.0 % | 3 | 25.0 % | HR.TESTstatic_sql_ exec_line5 (Line 6) | 3r6qf2qhr3c m1 | SELECT COUNT(*) FROM EMPLOYEE S |
| 0 | 0.0 % | 0 | 0 | 0.0 % | 1 | 8.3 % | SYS.DBMS_HPROF.S TOP_PROFILING (Line 453) | | |

16.5.2.3 Module-Level Reports

Each module-level report includes this information for each module (for example, package or type):

- Module time (time spent in the module—sum of the function times of all functions in the module)
- · Number of calls to functions in the module

Each module-level report is sorted on a particular attribute; for example, module time or module name.

This sample report is sorted by module time, which is why information in the Module, Ind%, and Cum% columns is in **bold type**:

Sample Report

Module Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs)

166878 microsecs (elapsed time) & 1099 function calls



| Module | Ind% | Cum% | Calls | Ind% | Module Name |
|--------|-------|-------|-------|-------|----------------|
| 84932 | 50.9% | 50.9% | 6 | 0.5% | HR.P |
| 67749 | 40.6% | 91.5% | 216 | 19.7% | SYS.DBMS_LOB |
| 13340 | 8.0% | 99.5% | 660 | 60.1% | SYS.UTL_FILE |
| 839 | 0.5% | 100% | 214 | 19.5% | SYS.UTL_RAW |
| 18 | 0.0% | 100% | 2 | 0.2% | HR.UTILS |
| 0 | 0.0% | 100% | 1 | 0.1% | SYS.DBMS_HPROF |

16.5.2.4 Namespace-Level Reports

Each namespace-level report includes this information for each namespace:

- Namespace time (time spent in the namespace—sum of the function times of all functions in the namespace)
- Number of calls to functions in the namespace

Each namespace-level report is sorted on a particular attribute; for example, namespace time or number of calls to functions in the namespace.

This sample report is sorted by function time, which is why information in the Function, Ind%, and Cum% columns is in **bold type**:

Sample Report

Namespace Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs)

166878 microsecs (elapsed time) & 1099 function calls

| Function | Ind% | Cum% | Calls | Ind% | Namespace |
|----------|-------|-------|-------|-------|-----------|
| 93659 | 56.1% | 56.1% | 1095 | 99.6% | PLSQL |
| 73219 | 43.9% | 100% | 4 | 0.4% | SQL |

16.5.2.5 Parents and Children Report for a Function

For each function tracked by the profiler, the Parents and Children Report provides information about parents (functions that call it) and children (functions that it calls). For each parent, the report gives the function's execution profile (subtree time, function time, descendants time, and number of calls). For each child, the report gives the execution profile for the child when called from this function (but not when called from other functions).

The execution profile for a function includes the same information for that function as a function-level report includes for each function.

This Sample Report is a fragment of a Parents and Children Report that corresponds to a function named HR.P.UPLOAD. The first row has this summary information:

- There are two calls to the function HR.P.UPLOAD.
- The total subtree time for the function is 166,860 microseconds—11,713 microseconds (7.0%) in the function itself and 155,147 microseconds (93.0%) in its descendants.



After the row "Parents" are the execution profile rows for the two parents of HR.P.UPLOAD, which are HR.UTILS.COPY IMAGE and HR.UTILS.COPY FILE.

The first parent execution profile row, for HR.UTILS.COPY IMAGE, shows:

- HR.UTILS.COPY_IMAGE calls HR.P.UPLOAD once, which is 50% of the number of calls to HR.P.UPLOAD.
- The subtree time for HR.P.UPLOAD when called from HR.UTILS.COPY_IMAGE is 106,325 microseconds, which is 63.7% of the total subtree time for HR.P.UPLOAD.
- The function time for HR.P.UPLOAD when called from HR.UTILS.COPY_IMAGE is 6,434 microseconds, which is 54.9% of the total function time for HR.P.UPLOAD.

After the row "Children" are the execution profile rows for the children of HR.P.UPLOAD when called from HR.P.UPLOAD.

The third child execution profile row, for ${\tt SYS.UTL_FILE.GET_RAW}$, shows:

- HR.P.UPLOAD calls SYS.UTL FILE.GET RAW 216 times.
- The subtree time, function time and descendants time for SYS.UTL_FILE.GET_RAW when called from HR.P.UPLOAD are 12,487 microseconds, 3,969 microseconds, and 8,518 microseconds, respectively.
- Of the total descendants time for HR.P.UPLOAD (155,147 microseconds), the child SYS.UTL FILE.GET RAW is responsible for 12,487 microsecs (8.0%).

Sample Report

HR.P.UPLOAD (Line 3)

| Subtree | Ind% | Function | Ind% | Descendant | Ind% | Calls | Ind% | Function Name |
|-----------|-------|----------|-------|------------|-------|-------|-------|--|
| 166860 | 100% | 11713 | 7.0% | 155147 | 93.0% | 2 | 0.2% | HR.P.UPLOAD (Line 3) |
| Parents: | | | | | | | | |
| 106325 | 63.7% | 6434 | 54.9% | 99891 | 64.4% | 1 | 50.0% | HR.UTILS.COPY_IMA GE (Line 3) |
| 60535 | 36.3% | 5279 | 45.1% | 55256 | 35.6% | 1 | 50.0% | HR.UTILS.COPY_FILE (Line 8)) |
| Children: | | | | | | | | |
| 71818 | 46.3% | 71818 | 100% | 0 | N/A | 2 | 100% | HR.Pstatic_sql_exec _line38 (Line 38) |
| 67649 | 43.6% | 67649 | 100% | 0 | N/A | 214 | 100% | SYS.DBMS_LOB.WRIT EAPPEND (Line 926) |
| 12487 | 8.0% | 3969 | 100% | 8518 | 100% | 216 | 100% | SYS.UTL_FILE.GET_R AW (Line 1089) |
| 1401 | 0.9% | 1401 | 100% | 0 | N/A | 2 | 100% | HR.Pstatic_sql_exec _line39 (Line 39) |
| 839 | 0.5% | 839 | 100% | 0 | N/A | 214 | 100% | SYS.UTL_FILE.GET_R AW (Line 246) |
| 740 | 0.5% | 73 | 100% | 667 | 100% | 2 | 100% | SYS.UTL_FILE.FOPEN (Line 422) |
| 113 | 0.1% | 11 | 100% | 102 | 100% | 2 | 100% | SYS.UTL_FILE.FCLOS E (Line 585) |



| Subtree | Ind% | Function | Ind% | Descendant | Ind% | Calls | Ind% | Function Name |
|---------|------|----------|------|------------|------|-------|------|--|
| 100 | 0.1% | 100 | 100% | 0 | N/A | 2 | 100% | SYS.DBMS_LOB.CRE ATETEMPORARY (Line 536) |

See Also:

Function-Level Reports

16.5.2.6 Understanding PL/SQL Hierarchical Profiler SQL-Level Reports

Understanding DBMS_HPROF.ANALYZE SQL-level reports.

The PL/SQL Hierarchical Profiler SQL-level report provides the list of all the SQLs collected during profiling. along with the abbreviated SQL text, and the elapsed time (microsecs) sorted by SQL ID. The SQL ID is useful if other SQL statistics must be retrieved in other tables, for example, for SQL tuning purpose. The first 50 characters of the SQL text is included in the report. You can use the Function-Level reports to get the details surrounding where the SQL is called, and its location in the source code if needed.

Sample Report

SQL ID Elapsed Time (microsecs) Data sorted by SQL ID

824 microsecs (elapsed time) & 12 function calls

| SQL ID | SQL TEXT | Function | Ind% | Calls | Ind% |
|---------------|---|----------|-------|-------|-------|
| 3r6qf2qhr3cm1 | SELECT COUNT(*) FROM EMPLOYEES | 679 | 82.4% | 3 | 25.0% |

16.5.3 HTML Difference Report from Two Raw Profiler Output Files

To generate a PL/SQL hierarchical profiler HTML difference report from two raw profiler output files, use these commands:

```
% cd target directory
```

% plshprof -output html root filename profiler output filename 1 profiler output filename 2

target directory is the directory in which you want the HTML files to be created.

html root filename is the name of the root HTML file to be created.

profiler_output_filename_1 and profiler_output_filename_2 are the names of raw profiler output files.

The preceding plshprof command generates a set of HTML files. Start browsing them from html root filename.html.



Topics:

- Difference Report Conventions
- First Page of Difference Report
- Function-Level Difference Reports
- Module-Level Difference Reports
- Namespace-Level Difference Reports
- Parents and Children Difference Report for a Function

16.5.3.1 Difference Report Conventions

Difference reports use these conventions:

- In a report title, Delta means difference, or change.
- A positive value indicates that the number increased (regressed) from the first run to the second run.
- A negative value for a difference indicates that the number decreased (improved) from the first run to the second run.
- The symbol # after a function name means that the function was called in only one run.

16.5.3.2 First Page of Difference Report

The first page of an HTML difference report from two raw profiler output files includes summary information and hyperlinks to other pages of the report.

Sample First Page

PL/SQL Elapsed Time (microsecs) Analysis - Summary Page

This analysis finds a net **regression** of **2709589** microsecs (elapsed time) or **80%** (**3393719** versus **6103308**). Here is a summary of the 7 most important individual function regressions and improvements:

Regressions: 3399382 microsecs (elapsed time)

| Function | Rel% | Ind% | Calls | Rel% | Function Name |
|----------|--------|-------|-------|--------|------------------|
| 2075627 | +941% | 61.1% | 0 | | HR.P.G (Line 35) |
| 1101384 | +54.6% | 32.4% | 5 | +55.6% | HR.P.H (Line 18) |
| 222371 | | 6.5% | 1 | | HR.P.J (Line 10) |

Improvements: 689793 microsecs (elapsed time)

| Function | Rel% | Ind% | Calls | Rel% | Function Name |
|----------|--------|-------|-------|--------|---------------------|
| -467051 | -50.0% | 67.7% | -2 | -50.0% | HR.P.F (Line 25) |
| -222737 | | 32.3% | -1 | | HR.P.I (Line 2)# |
| -5 | -21.7% | 0.0% | 0 | | HR.P.TEST (Line 46) |



The PL/SQL Timing Analyzer produces a collection of reports that present information derived from the profiler's output logs in a variety of formats. The following reports have been found to be the most generally useful as starting points for browsing:

- Function Elapsed Time (microsecs) Data for Performance Regressions
- Function Elapsed Time (microsecs) Data for Performance Improvements

Also, the following reports are also available:

- Function Elapsed Time (microsecs) Data sorted by Function Name
- Function Elapsed Time (microsecs) Data sorted by Total Subtree Elapsed Time (microsecs) Delta
- Function Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs) Delta
- Function Elapsed Time (microsecs) Data sorted by Total Descendants Elapsed Time (microsecs) Delta
- Function Elapsed Time (microsecs) Data sorted by Total Function Call Count Delta
- Module Elapsed Time (microsecs) Data sorted by Module Name
- Module Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs)
 Delta
- Module Elapsed Time (microsecs) Data sorted by Total Function Call Count Delta
- Namespace Elapsed Time (microsecs) Data sorted by Namespace
- Namespace Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs)
- Namespace Elapsed Time (microsecs) Data sorted by Total Function Call Count
- File Elapsed Time (microsecs) Data Comparison with Parents and Children

16.5.3.3 Function-Level Difference Reports

Each function-level difference report includes, for each function, the change in these values from the first run to the second run:

- Function time (time spent in the function itself, also called "self time")
- Descendants time (time spent in the descendants of the function)
- Subtree time (time spent in the subtree of the function—function time plus descendants time)
- Number of calls to the function
- Mean function time

The mean function time is the function time divided by number of calls to the function.

Function name

The function name is hyperlinked to the Parents and Children Difference Report for the function.

The report in Sample Report 1 shows the difference information for all functions that performed better in the first run than they did in the second run. Note that:

• For HR.P.G, the function time increased by 2,075,627 microseconds (941%), which accounts for 61.1% of all regressions.



- For HR.P.H, the function time and number of calls increased by 1,101,384 microseconds (54.6%) and 5 (55.6%), respectively, but the mean function time improved by 1,346 microseconds (-0.6%).
- HR.P.J was called in only one run.

Sample Report 1

Function Elapsed Time (microsecs) Data for Performance Regressions

| Subtre e | Functio n | Rel% | Ind% | Cum% | Descendan t | Call s | Rel% | Mean Function | Rel% | Function Name |
|-------------|--------------|------------|-----------|-------|----------------|-----------|------------|------------------|-----------|-------------------|
| 407578 7 | 2075627 | +941% | 61.1 % | 61.1% | 2000160 | 0 | | 2075627 | +941 % | HR.P.G (Line 35) |
| 110138 4 | 1101384 | +54.6 % | 32.4 % | 93.5% | 0 | 5 | +55.6 % | -1346 | -0.6% | HR.P.H (Line 18) |
| 222371 | 222371 | | 6.5% | 100% | 0 | 1 | | | | HR.P.J (Line 10)# |

The report in Sample Report 2 shows the difference information for all functions that performed better in the second run than they did in the first run.

Sample Report 2

Function Elapsed Time (microsecs) Data for Performance Improvements

| Subtree | Functio n | Rel% | Ind% | Cum% | Descendan t | Call s | Rel% | Mean Function | Rel% | Function Name |
|--------------|--------------|--------|-----------|-------|----------------|-----------|--------|------------------|-------|----------------------|
| -136582 7 | -467051 | -50.0% | 67.7 % | 67.7% | -898776 | -2 | -50.0% | -32 | 0.0% | HR.P.F (Line 25) |
| -222737 | -222737 | | 32.3 % | 100% | 0 | -1 | | | | HR.P.I (Line 2) |
| 2709589 | -5 | -21.7% | 0.0% | 100% | 2709594 | 0 | | -5 | -20.8 | HR.P.TEST (Line 46)# |

The report in Sample Report 3 summarizes the difference information for all functions.

Sample Report 3

Function Elapsed Time (microsecs) Data sorted by Total Function Call Count Delta

| Subtree | Functio n | Rel% | Ind% | Descendan t | Call s | Rel% | Mean Function | Rel% | Function Name |
|--------------|--------------|------------|-----------|----------------|-----------|------------|------------------|------------|---------------------|
| 1101384 | 1101384 | +54.6 % | 32.4 % | 0 | 5 | +55.6 % | -1346 | -0.6% | HR.P.H (Line 18) |
| -136582 7 | -467051 | +50.0 % | 67.7 % | -898776 | -2 | -50.0% | -32 | -0.0% | HR.P.F (Line 25) |
| -222377 | -222377 | | 32.3 % | 0 | -1 | | | | HR.P.I (Line 2)# |
| 222371 | 222371 | | 6.5% | 0 | 1 | | | | HR.P.J(Line 10)# |
| 4075787 | 2075627 | +941% | 61.1 % | 2000160 | 0 | | 2075627 | +941 % | HR.P.G (Line 35) |
| 2709589 | -5 | -21.7% | 0.0% | 2709594 | 0 | | -5 | -20.8 % | HR.P.TEST (Line 46) |
| | | | | | | | | | |



| Subtree | Functio n | Rel% | Ind% | Descendan t | Call s | Rel% | Mean Function | Rel% | Function Name |
|---------|--------------|------|------|----------------|-----------|------|------------------|------|---|
| 0 | 0 | | | 0 | 0 | | | | SYS.DBMS_HPROF.STOP_P ROFILING (Line 53) |

16.5.3.4 Module-Level Difference Reports

Each module-level report includes, for each module, the change in these values from the first run to the second run:

- Module time (time spent in the module—sum of the function times of all functions in the module)
- Number of calls to functions in the module

Sample Report

Module Elapsed Time (microsecs) Data sorted by Total Function Elapsed Time (microsecs) Delta

| Module | Calls | Module Name |
|---------|-------|----------------|
| 2709589 | 3 | HR.P |
| 0 | 0 | SYS.DBMS_HPROF |

16.5.3.5 Namespace-Level Difference Reports

Each namespace-level report includes, for each namespace, the change in these values from the first run to the second run:

- Namespace time (time spent in the namespace—sum of the function times of all functions in the namespace)
- Number of calls to functions in the namespace

Sample Report

Namespace Elapsed Time (microsecs) Data sorted by Namespace

| Function | Calls | Namespace |
|----------|-------|-----------|
| 2709589 | 3 | PLSQL |

16.5.3.6 Parents and Children Difference Report for a Function

The Parents and Children Difference Report for a function shows changes in the execution profiles of these from the first run to the second run:

- Parents (functions that call the function)
- Children (functions that the function calls)

Execution profiles for children include only information from when this function calls them, not for when other functions call them.

The execution profile for a function includes this information:



- Function time (time spent in the function itself, also called "self time")
- Descendants time (time spent in the descendants of the function)
- Subtree time (time spent in the subtree of the function—function time plus descendants time)
- Number of calls to the function
- Function name

The sample report is a fragment of a Parents and Children Difference Report that corresponds to a function named ${\tt HR.P.X.}$

The first row, a summary of the difference between the first and second runs, shows regression: function time increased by 1,094,099 microseconds (probably because the function was called five more times).

The "Parents" rows show that HR.P.G called HR.P.X nine more times in the second run than it did in the first run, while HR.P.F called it four fewer times.

The "Children" rows show that ${\tt HR.P.X}$ called each child five more times in the second run than it did in the first run.

Sample Report

HR.P.X (Line 11)

| Subtree | Function | Descendant | Calls | Function Name |
|-----------|----------|------------|-------|------------------|
| 3322196 | 1094099 | 2228097 | 5 | HR.P.X (Line 11) |
| Parents: | | | | |
| 6037490 | 1993169 | 4044321 | 9 | HR.P.G (Line 38) |
| -2715294 | -899070 | -1816224 | -4 | HR.P.F (Line 28) |
| Children: | | | | |
| 1125489 | 1125489 | 0 | 5 | HR.P.J (Line 10) |
| 1102608 | 1102608 | 0 | 5 | HR.P.I (Line 2) |

The Parents and Children Difference Report for a function is accompanied by a Function Comparison Report, which shows the execution profile of the function for the first and second runs and the difference between them. This example is the Function Comparison Report for the function ${\tt HR.P.X:}$

Sample Report

Elapsed Time (microsecs) for HR.P.X (Line 11) (20.1% of total regression)

| HR.P.X (Line 11) | First Trace | Ind% | Second Trace | Ind% | Diff | Diff% |
|--|----------------|-------|-----------------|-------|---------|--------|
| Function Elapsed Time (microsecs) | 1999509 | 26.9% | 3093608 | 24.9% | 1094099 | +54.7% |
| Descendants Elapsed Time (microsecs) | 4095943 | 55.1% | 6324040 | 50.9% | 2228097 | +54.4% |
| Subtree Elapsed Time (microsecs) | 6095452 | 81.9% | 9417648 | 75.7% | 3322196 | +54.5% |
| Function Calls | 9 | 25.0% | 14 | 28.6% | 5 | +55.6% |
| Mean Function Elapsed Time (microsecs) | 222167.7 | | 220972.0 | | -1195.7 | -0.5% |



| HR.P.X (Line 11) | First Trace | Ind% | Second Trace | Ind% | Diff | Diff% |
|---|----------------|------|-----------------|------|---------|-------|
| Mean Descendants Elapsed Time (microsecs) | 455104.8 | | 451717.1 | | -3387.6 | -0.7% |
| Mean Subtree Elapsed Time (microsecs) | 677272.4 | | 672689.1 | | -4583.3 | -0.7% |

