DBMS_DATA_MINING

The DBMS_DATA_MINING package is the application programming interface for creating, evaluating, and querying Oracle Machine Learning for SQL models.

In Oracle Database Release 21c, Oracle Data Mining has been rebranded to Oracle Machine Learning for SQL (Oracle Machine Learning for SQL). The PL/SQL package name, however, has not changed and remains DBMS DATA MINING.

This chapter contains the following topics:

- Overview
- Security Model
- Mining Functions
- Model Settings
- Algorithm Specific Settings
- Solver Settings
- Datatypes
- Summary of DBMS DATA MINING Subprograms

✓ See Also:

- Oracle Machine Learning for SQL Concepts
- Oracle Machine Learning for SQL User's Guide
- DBMS_DATA_MINING_TRANSFORM
- DBMS_PREDICTIVE_ANALYTICS

DBMS_DATA_MINING Overview

Oracle Machine Learning for SQL supports both supervised and unsupervised machine learning. Supervised machine learning predicts a target value based on historical data. Unsupervised machine learning discovers natural groupings and does not use a target. You can use Oracle Machine Learning for SQL procedures on structured data and unstructured text.

Supervised machine learning techniques include:

- Classification
- Regression
- Feature Selection (Attribute Importance)
- Time Series

Unsupervised machine learning techniques include:

- Clustering
- Association
- Feature Extraction
- · Anomaly Detection

The steps you use to build and apply a machine learning model depend on the machine learning technique and the algorithm being used. The algorithms supported by Oracle Machine Learning for SQL are listed in the following table.

Table 62-1 Oracle Machine Learning for SQL Algorithms

Algorithm	Abbreviation	Function
Apriori	AR	Association
CUR Matrix Decomposition	CUR	Attribute importance
Decision Tree	DT	Classification
Expectation Maximization	EM	Clustering
Explicit Semantic Analysis	ESA	Feature extraction, classification
Exponential Smoothing	ESM	Time series
Generalized Linear Models	GLM	Classification, regression
k-Means	KM	Clustering
Minimum Descriptor Length	MDL	Attribute importance
Multivariate State Estimation Technique - Sequential Probability Ratio Test	MSET-SPRT	Anomaly detection, classification
Naive Bayes	NB	Classification
Neural Network	NN	Classification, regression
Non-Negative Matrix Factorization	NMF	Feature extraction
Orthogonal Partitioning Clustering	O-Cluster	Clustering
Random Forest	RF	Classification
Singular Value Decomposition and Principal Component Analysis	SVD and PCA	Feature extraction
Support Vector Machine	SVM	Classification, regression, anomaly detection
XGBoost	XGBoost	Classification, regression

Oracle Machine Learning for SQL supports more than one algorithm for the classification, regression, clustering, and feature extraction machine learning techniques. Each of these machine learning techniques has a default algorithm, as shown in the following table.

Table 62-2 Oracle Machine Learning for SQL Default Algorithms

Mining Function	Default Algorithm
Classification	Naive Bayes
Clustering	k-Means
Feature Extraction	Non-Negative Matrix Factorization
Feature Selection	Minimum Descriptor Length



Table 62-2 (Cont.) Oracle Machine Learning for SQL Default Algorithms

Mining Function	Default Algorithm
Regression	Support Vector Machine
Time Series	Exponential Smoothing

DBMS_DATA_MINING Security Model

The DBMS_DATA_MINING package is owned by user SYS and is installed as part of database installation. Execution privilege on the package is granted to public. The routines in the package are run with invokers' rights (run with the privileges of the current user).

The DBMS_DATA_MINING package exposes APIs that are leveraged by the Oracle Machine Learning for SQL. Users who wish to create machine learning models in their own schema require the CREATE MINING MODEL system privilege. Users who wish to create machine learning models in other schemas require the CREATE ANY MINING MODEL system privilege.

Users have full control over managing models that exist within their own schema. Additional system privileges necessary for managing machine learning models in other schemas include ALTER ANY MINING MODEL, DROP ANY MINING MODEL, SELECT ANY MINING MODEL, COMMENT ANY MINING MODEL, and AUDIT ANY.

Individual object privileges on machine learning models, ALTER MINING MODEL and SELECT MINING MODEL, can be used to selectively grant privileges on a model to a different user.



Oracle Data Mining User's Guide for more information about the security features of Oracle Machine Learning for SQL $\,$

DBMS_DATA_MINING — Machine Learning Functions

A machine learning **function** refers to the methods for solving a given class of machine learning problems.

The machine learning function must be specified when a model is created. You specify a machine learning function with the mining_function parameter of the CREATE_MODEL Procedure or the CREATE_MODEL2 Procedure.

Table 62-3 Machine Learning Functions

Value	Description
ASSOCIATION	Association is a descriptive machine learning function. An association model identifies relationships and the probability of their occurrence within a data set.
	Association models use the Apriori algorithm.



Table 62-3 (Cont.) Machine Learning Functions

Value	Description
ATTRIBUTE_IMPORTANCE	Attribute importance is a predictive machine learning function, also known as feature selection. An attribute importance model identifies the relative importance of an attribute in predicting a given outcome.
	Attribute importance models can use Minimum Description Length (MDL) or CUR Matrix Decomposition. MDL is the default.
CLASSIFICATION	Classification is a predictive machine learning function. A classification model uses historical data to predict a categorical target.
	Classification models can use: Decision Tree, logistic regression, Multivariate State Estimation Technique - Sequential Probability Ratio Test, Naive Bayes, Support Vector Machine (SVM), or XGBoost. The default is Naive Bayes.
	The classification function can also be used for anomaly detection . For anomaly detection, you can use the Multivariate State Estimation Technique - Sequential Probability Ratio Test algorithm or the SVM algorithm with a null target (One-Class SVM), or the EM algorithm with a null target (EM Anomaly).
CLUSTERING	Clustering is a descriptive machine learning function. A clustering model identifies natural groupings within a data set.
	Clustering models can use k -Means, O-Cluster, or Expectation Maximization. The default is k -Means.
FEATURE_EXTRACTION	Feature extraction is a descriptive machine learning function. A feature extraction model creates an optimized data set on which to base a model.
	Feature extraction models can use Explicit Semantic Analysis, Non-Negative Matrix Factorization, Singular Value Decomposition, or Principal Component Analysis. Non-Negative Matrix Factorization is the default.
REGRESSION	Regression is a predictive machine learning function. A regression model uses historical data to predict a numerical target.
	Regression models can use linear regression, Support Vector Machine, or XGBoost. The default is Support Vector Machine.
TIME_SERIES	Time series is a predictive machine learning function. A time series model forecasts the future values of a time-ordered series of historical numeric data over a user-specified time window. Time series models use the Exponential Smoothing algorithm.

Oracle Machine Learning for SQL Concepts for more information about mining functions



DBMS_DATA_MINING — Model Settings

Oracle Machine Learning for SQL uses settings to specify the algorithm and other characteristics of a model. Some settings are general, some are specific to a machine learning function, and some are specific to an algorithm.

All settings have default values. If you want to override one or more of the settings for a model, then you must create a settings table. The settings table must have the column names and data types shown in the following table.

Table 62-4 Required Columns in the Model Settings Table

O. I N	D.U. T
Column Name	Data Type
SETTING_NAME	VARCHAR2(30)
SETTING_VALUE	VARCHAR2 (4000)

The information you provide in the settings table is used by the model at build time. The name of the settings table is an optional argument to the CREATE_MODEL Procedure. You can also provide these settings through the CREATE MODEL2 Procedure.

The settings used by a model can be found by querying the data dictionary view <code>ALL_MINING_MODEL_SETTINGS</code>. This view displays the model settings used by the machine learning models to which you have access. All of the default and user-specified setting values are included in the view.

See Also:

- ALL MINING MODEL SETTINGS in Oracle Database Reference
- Oracle Machine Learning for SQL User's Guide for information about specifying model settings

DBMS_DATA_MINING — Algorithm Names

The ${\tt ALGO}\,$ NAME setting specifies the model algorithm.

The values for the ALGO NAME setting are listed in the following table.

Table 62-5 Algorithm Names

ALGO_NAME Value	Description	Machine Learning Function
ALGO_AI_MDL	Minimum Description Length	Attribute importance
ALGO_APRIORI_ASSOCIATION_RULES	Apriori	Association rules
ALGO_CUR_DECOMPOSITION	CUR Matrix Decomposition	Attribute importance
ALGO_DECISION_TREE	Decision Tree	Classification
ALGO_EXPECTATION_MAXIMIZATION	Expectation Maximization	Clustering, Classification



Table 62-5 (Cont.) Algorithm Names

ALGO_NAME Value	Description	Machine Learning Function
ALGO_EXPLICIT_SEMANTIC_ANALYS	Explicit Semantic Analysis	Feature extraction Classification
ALGO_EXPONENTIAL_SMOOTHING	Exponential Smoothing	Time series
ALGO_EXTENSIBLE_LANG	Language used for extensible algorithm	All mining functions supported
ALGO_GENERALIZED_LINEAR_MODEL	Generalized Linear Model	Classification, regression; also feature selection and generation
ALGO_KMEANS	Enhanced k-Means	Clustering
ALGO_MSET_SPRT	Multivariate State Estimation Technique - Sequential Probability Ratio Test	Classification
ALGO_NAIVE_BAYES	Naive Bayes	Classification
ALGO_NEURAL_NETWORK	Neural Network	Classification
ALGO_NONNEGATIVE_MATRIX_FACTOR	Non-Negative Matrix Factorization	Feature extraction
ALGO_O_CLUSTER	O-Cluster	Clustering
ALGO_RANDOM_FOREST	Random Forest	Classification
ALGO_SINGULAR_VALUE_DECOMP	Singular Value Decomposition	Feature extraction
ALGO_SUPPORT_VECTOR_MACHINES	Support Vector Machine	Classification and regression
ALGO_XGBOOST	XGBoost	Classification and regression

Oracle Machine Learning for SQL Concepts for information about algorithms

DBMS_DATA_MINING — Automatic Data Preparation

Oracle Machine Learning for SQL supports fully Automatic Data Preparation (ADP), user-directed general data preparation, and user-specified embedded data preparation. The PREP_* settings enable the user to request fully automated or user-directed general data preparation. By default, fully Automatic Data Preparation (PREP AUTO ON) is enabled.

When you enable ADP, the model uses heuristics to transform the build data according to the requirements of the algorithm. Instead of fully ADP, the user can request that the data be shifted and/or scaled with the PREP_SCALE* and PREP_SHIFT* settings. The transformation instructions are stored with the model and reused whenever the model is applied. The model settings can be viewed in USER MINING MODEL SETTINGS.

You can choose to supplement Automatic Data Preparations by specifying additional transformations in the xform_list parameter when you build the model. See "CREATE_MODEL Procedure" and "CREATE_MODEL2 Procedure".

If you do not use ADP and do not specify transformations in the xform_list parameter to CREATE MODEL, you must implement your own transformations separately in the build, test, and

scoring data. You must take special care to implement the exact same transformations in each data set.

If you do not use ADP, but you do specify transformations in the $xform_list$ parameter to CREATE_MODEL, OML4SQL embeds the transformation definitions in the model and prepares the test and scoring data to match the build data.

The values for the PREP * setting are described in the following table.

Table 62-6 PREP_* Setting

Setting Name	Setting Value	Description
PREP_AUTO	PREP_AUTO_ON PREP_AUTO_OFF	This setting enables fully automated data preparation. The default is PREP_AUTO_ON.
PREP_SCALE_2DNUM	• PREP_SCALE_STDDEV • PREP_SCALE_RANGE	This setting enables scaling data preparation for two-dimensional numeric columns. PREP_AUTO must be OFF for this setting to take effect. The following are the possible values:
		 PREP_SCALE_STDDEV: A request to divide the column values by the standard deviation of the column and is often provided together with PREP_SHIFT_MEAN to yield z-score normalization.
		 PREP_SCALE_RANGE: A request to divide the column values by the range of values and is often provided together with PREP_SHIFT_MIN to yield a range of [0,1].
PREP_SCALE_NNUM	PREP_SCALE_MAXABS	This setting enables scaling data preparation for nested numeric columns. PREP_AUTO must be OFF for this setting to take effect. If specified, then the valid value for this setting is PREP_SCALE_MAXABS, which yields data in the range of [-1,1].
PREP_SHIFT_2DNUM	• PREP_SHIFT_MEAN • PREP_SHIFT_MIN	This setting enables centering data preparation for two-dimensional numeric columns. PREP_AUTO must be OFF for this setting to take effect. The following are the possible values: PREP_SHIFT_MEAN: Results in subtracting the average of the column from each value. PREP_SHIFT_MIN: Results in subtracting the minimum of the column from each value.

See Also:

Oracle® Machine Learning for SQL for information about data transformations

DBMS_DATA_MINING — Machine Learning Function Settings

The settings described in this table apply to a machine learning function.

Table 62-7 Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
Association	ASSO_MAX_RULE_LENGTH	An integer between 2 to 20, inclusive, represented as	Sets the upper limit on the length of association rules. Useful in managing the compute time. Shorter rules (fewer antecedents) take less memory and compute time.
		a character	Expression:
		string	TO_CHAR(3)
			Default is 4.
Association	ASSO_MIN_CONFIDENCE	A floating point number between 0 and 1, inclusive, expressed as a character string	Defines the minimum confidence threshold for association rules. This parameter reduces the number of rules generated, focusing only on those that meet the minimum confidence specification. It reduces model size. Expression:
			TO_CHAR(0.4)
			Default is 0.1.
Association	ASSO_MIN_SUPPORT	A floating point number	Specifies the minimum support threshold for association rules.
		between 0 and 1, inclusive, expressed as a	Expression:
			TO_CHAR(0.2)
		character string	Default is 0.1.
Association	ASSO_MIN_SUPPORT_INT	a positive integer	Determines the minimum absolute support each rule must meet, expressed as an integer. Sets a concrete baseline for the number of occurrences required for an itemset to be considered, ensuring rules are based on sufficiently frequent patterns. The default is 1.
Association	ASSO_MIN_REV_CONFIDENCE	A floating point number between 0 and 1, inclusive, expressed as a character string.	Establishes the minimum reverse confidence for each association rule. This setting filters rules to ensure that they are significant not just in the context of the antecedent but also regarding the consequent, enhancing the relevance of the rule.
			The Reverse Confidence of a rule is defined as the number of transactions in which the rule occurs divided by the number of transactions in which the consequent occurs.
			The value is real number between 0 and 1.
			Expression:
			TO_CHAR(0.45)
		The default is 0.	



Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
Association	ASSO_IN_RULES	NULL	Specifies a list of items that must be included in each association rule. It accepts a comma-separated string of items; at least one of these must appear in every reported rule, either as an antecedent or a consequent. This parameter is useful for focusing the analysis on rules that involve specific items of interest, thereby tailoring the association rules to specific analytical needs or hypotheses.
			If not set, the filtering is not applied by default. For example,
			<pre>INSERT INTO sett_tab (setting_name, setting_value) VALUES</pre>
Association	ASSO_EX_RULES	NULL	Defines a list of items to be excluded from each association rule. Accepts a comma-separated string listing items that should not appear in any reported association rules. Essential for omitting specific items from rule generation, which can be particularly useful when certain items are known to be irrelevant or misleading in the context of the analysis.
			The default is NULL.
			For example,
			<pre>INSERT INTO sett_tab (setting_name, setting_value) VALUES (dbms_data_mining.asso_ex_rules, '''a'',''b''');</pre>
Association	ASSO_ANT_IN_RULES	NULL	Determines inclusion criteria for the antecedent part of each association rule. Accepts a comma-separated list of items where at least one must appear in the antecedent of each rule. Targets specific items to always be considered as potential causes in the rules, refining the focus of the analysis. The default is NULL. For example,
			<pre>INSERT INTO sett_tab (setting_name, setting_value) VALUES</pre>
			<pre>(dbms_data_mining.asso_ant_in_rules, '''a'',''b''');</pre>

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description	
Association	ASSO_ANT_EX_RULES	NULL	Sets exclusion criteria for the antecedent in association rules. Comma-separated string items to be excluded from the antecedent of rule. Prevents specified items from being co as causes in the rules, avoiding irrelevant or redundant combinations.	
			The default is NULL.	
			The following example illustrates how you can define the parameter when you are using the CREATE_MODEL procedure. You must create a table with setting name (a constant) and setting value and then use the CREATE_MODEL procedure.	
			<pre>INSERT INTO sett_tab (setting_name, setting_value) VALUES</pre>	,
			<pre>(dbms_data_mining.asso_ant_ex_rules, '''a'',''b''');</pre>	
			The following example illustrates how you can the parameter when you are using the CREATE_MODEL2 procedure using string value	
			<pre>%script BEGIN DBMS_DATA_MINING.DROP_MODEL('AR_SH_ LE');</pre>	_SAMP
			EXCEPTION WHEN OTHERS THEN NULL; EN	ND;
			DECLARE v_set1st DBMS_DATA_MINING.SETTING_LIST; BEGIN	
			<pre>v_setlst('ALGO_NAME') 'ALGO_APRIORI_ASSOCIATION_RULES';</pre>	:=
			<pre>v_setlst('PREP_AUTO') 'ON';</pre>	:=
			<pre>v_setlst('ASSO_MIN_SUPPORT') '0.04';</pre>	:=
			<pre>v_setlst('ASSO_MIN_CONFIDENCE') '0.1';</pre>	:=
			<pre>v_setlst('ASSO_MAX_RULE_LENGTH') '2';</pre>	:=

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
			<pre>v_set1st('ASSO_ANT_EX_RULES') := '''a'',''b''';</pre>
			<pre>v_setlst('ODMS_ITEM_ID_COLUMN_NAME'):= 'PROD_NAME';</pre>
			<pre>v_set1st('ASSO_AGGREGATES') := 'AMOUNT_SOLD';</pre>
			DBMS_DATA_MINING.CREATE_MODEL2(
			MODEL_NAME =>
			'AR_SH_SAMPLE',
			MINING_FUNCTION =>
			'ASSOCIATION',
			DATA_QUERY => 'select *
			from SALES_TRANS_CUST',
			SET_LIST => v_set1st, CASE ID COLUMN NAME =>
			'CUST ID');

END;



Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
Association	ASSO_CONS_IN_RULES	NULL	Establishes inclusion criteria for the consequent pa of each association rule. Specifies a list of items where at least one must appear in the consequent each rule. Ensures that certain items are always considered as potential outcomes in the rules, focusing on specific effects of interest.
			The default is NULL.
			The following example illustrates how you can define the parameter when you are using the CREATE_MODEL procedure. You must create a table with setting name (a constant) and setting value and then use the CREATE_MODEL procedure.
			<pre>INSERT INTO sett_tab (setting_name, setting_value) VALUES</pre>
			<pre>(dbms_data_mining.asso_cons_in_rules, '''a'',''b''');</pre>
			The following example illustrates how you can define the parameter when you are using the CREATE_MODEL2 procedure using string values.
			%script
			BEGIN
			<pre>DBMS_DATA_MINING.DROP_MODEL('AR_SH_SAME LE');</pre>
			EXCEPTION WHEN OTHERS THEN NULL; END; / DECLARE
			v setlst
			DBMS_DATA_MINING.SETTING_LIST; BEGIN
			<pre>v_setlst('ALGO_NAME') := 'ALGO_APRIORI_ASSOCIATION_RULES';</pre>
			<pre>v_set1st('PREP_AUTO') := 'ON';</pre>
			<pre>v_set1st('ASSO_MIN_SUPPORT') := '0.04';</pre>
			<pre>v_setlst('ASSO_MIN_CONFIDENCE') := '0.1';</pre>
			<pre>v_setlst('ASSO_MAX_RULE_LENGTH') := '2';</pre>

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description	
			v_setlst('ASSO_CONS_IN_RU	ULES') :=
			<pre>v_setlst('ODMS_ITEM_ID_CC 'PROD_NAME';</pre>	OLUMN_NAME'):=
			<pre>v_set1st('ASSO_AGGREGATES 'AMOUNT_SOLD';</pre>	s') :=
			DBMS DATA MINING.CREA	ATE MODEL2 (
			MODEL NAME	=>
			'AR SH SAMPLE',	
			MINING FUNCTION	=>
			'ASSOCIATION',	
			DATA_QUERY	=> 'select *
			from SALES_TRANS_CUST',	
			SET_LIST	=> v_set1st,
			CASE_ID_COLUMN_NA	AME =>
			'CUST_ID');	

END;

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description	
Association	ASSO_CONS_EX_RULES	NULL	Sets exclusion criteria for the consequent in association rules. Comma-separated string indicat items to be excluded from the consequent of each rule. Omits specific items from being outcomes in rules, removing non-relevant or misleading results.	the
			The excluding rule can be used to reduce the rules that must be stored, but the user may be required build an extra model for running different including excluding rules.	to
			The default is NULL.	
			The following example illustrates how you can defit the parameter when you are using the CREATE_MODEL procedure. You must create a table with setting name (a constant) and setting value are then use the CREATE_MODEL procedure.	le
			<pre>INSERT INTO sett_tab (setting_name, setting_value) VALUES</pre>	
			<pre>(dbms_data_mining.asso_cons_ex_rules, '''a'',''b''');</pre>	
			The following example illustrates how you can define the parameter when you are using the CREATE_MODEL2 procedure using string values.	ine
			%script BEGIN	
			<pre>DBMS_DATA_MINING.DROP_MODEL('AR_SH_SAM LE');</pre>	1P
			EXCEPTION WHEN OTHERS THEN NULL; END;	
			DECLARE v_setlst DBMS_DATA_MINING.SETTING_LIST; BEGIN	
			<pre>v_setlst('ALGO_NAME') := 'ALGO_APRIORI_ASSOCIATION_RULES';</pre>	=
			<pre>v_set1st('PREP_AUTO') := 'ON';</pre>	=
			<pre>v_setlst('ASSO_MIN_SUPPORT') := '0.04';</pre>	=
			<pre>v_set1st('ASSO_MIN_CONFIDENCE') := '0.1';</pre>	=

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
			<pre>v_set1st('ASSO_MAX_RULE_LENGTH') := '2';</pre>
			<pre>v_set1st('ASSO_CONS_EX_RULES') := '''a'',''b''';</pre>
			<pre>v_set1st('ODMS_ITEM_ID_COLUMN_NAME'):= 'PROD_NAME';</pre>
			<pre>v_set1st('ASSO_AGGREGATES') := 'AMOUNT_SOLD';</pre>
			DBMS_DATA_MINING.CREATE_MODEL2(
			MODEL_NAME =>
			'AR_SH_SAMPLE', MINING FUNCTION =>
			'ASSOCIATION',
			DATA QUERY => 'select *
			from SALES_TRANS_CUST',
			SET_LIST => v_set1st,
			CASE_ID_COLUMN_NAME =>
			'CUST_ID');
			END;



Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
Association	ASSO_AGGREGATES	NULL	Defines columns for aggregation in association rules. Comma-separated list of column names for aggregation, limited to 10 columns. Aggregates additional data alongside items, providing more context to the association rules, but may increase memory and computational requirements.
			You can specify ASSO_AGGREGATES if ODMS_ITEM_ID_COLUMN_NAME is set indicating transactional input data. See DBMS_DATA_MINING Global Settings. The data table must have valid column names such as ITEM_ID and CASE_ID which are derived from ODMS_ITEM_ID_COLUMN_NAME and case_id_column_name respectively. Numeric value are supported. ITEM_VALUE is not a mandatory value. The default is NULL.
			For example, if the following is your Transactional Data table:
			CREATE OR REPLACE VIEW SALES_TRANS_CUST AS
			SELECT DISTINCT CUST_ID, PROD_NAME, PROD_CATEGORY
			FROM (SELECT A.CUST_ID, B.PROD_NAME, B.PROD_CATEGORY
			FROM SH.SALES A, SH.PRODUCTS B
			WHERE A.PROD_ID = B.PROD_ID AND
			A.CUST_ID BETWEEN 100001 AND 104500);
			Then, you can create your model similar to:
			%script BEGIN DBMS DATA MINING.DROP MODEL('AR SH SAMP
			LE'); EXCEPTION WHEN OTHERS THEN NULL; END;
			DECLARE v_set1st DBMS DATA MINING.SETTING LIST;

v_set1st('ALGO_NAME')

:=

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description	
			'ALGO_APRIORI_ASSOCIATION_RULES';	
			<pre>v_set1st('PREP_AUTO') 'ON';</pre>	=
			<pre>v_set1st('ASSO_MIN_SUPPORT') '0.04';</pre>	=
			<pre>v_set1st('ASSO_MIN_CONFIDENCE') '0.1';</pre>	=
			<pre>v_set1st('ASSO_MAX_RULE_LENGTH') : '2';</pre>	=
			<pre>v_set1st('ODMS_ITEM_ID_COLUMN_NAME'): 'PROD_NAME';</pre>	=
			<pre>v_set1st('ASSO_AGGREGATES') : 'AMOUNT_SOLD';</pre>	=
			<pre>DBMS_DATA_MINING.CREATE_MODEL2(</pre>	
			The default is NULL.	
			For each item, the user may supply several colunto aggregate. It requires more memory to buffer textra data. Also, the performance impact can be because of the larger input data set and more operation.	he
Association	ASSO_ABS_ERROR	0 <asso_abs_e RRORMAX(ASSO _MIN_SUPPORT , ASSO_MIN_CON FIDENCE).</asso_abs_e 	sampling. Balances accuracy and computational efficiency rule sampling; smaller values lead to larger samp	in oles on DR, ize.

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
Association	ASSO_CONF_LEVEL	0 ASSO_CONF_LE VEL 1	Sets the confidence level for an association rules sample. A higher confidence level increases sample size. Any value between 0.9 and 1 is suitable. The default value is 0.95.
Classification	CLAS_COST_TABLE_NAME	table_name	(Decision tree only) Names a user-created cost matrix table for model building. The cost matrix specifies misclassification costs. This parameter tailors decision tree models to prioritize certain types of misclassifications, enhancing model effectiveness in specific scenarios.
			Only decision tree models can use a cost matrix at build time. All classification algorithms can use a cost matrix at apply time.
			See "ADD_COST_MATRIX Procedure" for the column requirements.
			See Oracle Machine Learning for SQL Concepts for information about costs.
Classification	CLAS_PRIORS_TABLE_NAME	table_name	(Naive Bayes) Names a user-created table for storing prior probabilities. Adjusts for distribution differences between build and scoring data. Aligns model training more closely with real-world data distributions, improving prediction accuracy for Naive Bayes models.
			See Oracle Machine Learning for SQL User's Guide for the column requirements.
			See Oracle Machine Learning for SQL Concepts for additional information about priors.
Classification	CLAS_WEIGHTS_TABLE_NAME	table_name	(GLM and SVM only) Names a user-created table for storing weights for target values. Weights bias the model towards higher weighted classes. This parameter adjusts GLM and SVM models to focus on or balance between different classes, enhancing model performance for specific targets.
			See Oracle Machine Learning for SQL User's Guide for the column requirements.
			See Oracle Machine Learning for SQL Concepts for additional information about class weights.
Classification	CLAS_WEIGHTS_BALANCED	ON OFF	Indicates balancing of target distribution in the model. Relevant for rare targets; can improve average accuracy (average of per-class accuracy instead of overall accuracy which favors the dominant class). Particularly useful in data sets with imbalanced classes, ensuring rare events are adequately captured in the model. The default value is OFF.



Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
Classification	CLAS_MAX_SUP_BINS	For Decision Tree: Specify an integer	Specifies the maximum number of bins for each attribute. Manages the granularity of data binning, affecting model complexity and potentially influencing model accuracy and computation time.
		between 2 and 2147483647, inclusive represented as a character string. For Random Forest: Specify an integer between 2 and 254, inclusive represented as a character	The default value is 32. Expression: For Decision Tree: 2 <= a number <=2147483647 For Random Forest: 2 <= a number <=254 See, DBMS_DATA_MINING — Automatic Data Preparation
Clustering	CLUS_NUM_CLUSTERS	string. An integer greater than or equal to 1 expressed as a character string	The maximum number of leaf clusters generated by a clustering algorithm. The algorithm may return fewer clusters, depending on the data. Expression: TO_CHAR(9) Enhanced k-Means usually produces the exact number of clusters specified by CLUS_NUM_CLUSTERS, unless there are fewer distinct data points. When Expectation maximization (EM) is used for clustering, it may return fewer clusters than the number specified by CLUS_NUM_CLUSTERS depending on the data. The number of clusters returned by EM cannot be greater than the number of components, which is governed by algorithm-specific settings. (See Expectation Maximization Settings for Learning table) Depending on these settings, there may be fewer clusters than components. If component clustering is disabled, the number of clusters equals the number of components. The setting can be used only for EM Clustering algorithm. For EM Clustering algorithm, the default value of CLUS_NUM_CLUSTERS is system-determined. For k-Means and O-Cluster, the default is 10.

Table 62-7 (Cont.) Machine Learning Function Settings

Machine Learning Function	Setting Name	Setting Value	Description
Feature extraction	FEAT_NUM_FEATURES	An integer greater than or	The number of features to be extracted by a feature extraction model.
		equal to 1	Expression:
		expressed as a character string	TO_CHAR(8)
The algo	The default is estimated from the data by the algorithm. If the matrix rank is smaller than this number, fewer features will be returned.		
			For CUR Matrix Decomposition, the FEAT_NUM_FEATURES value is the same as the CURS_SVD_RANK value.



Oracle Machine Learning for SQL Concepts for information about machine learning functions

${\tt DBMS_DATA_MINING--Global\ Settings}$

The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.

Table 62-8 Global Settings

Setting Name	Setting Value	Description
ODMS_BOXCOX	ODMS_BOXCOX_ENABLE	This setting enables the Box-Cox variance-stabilization
	ODMS_BOXCOX_DISABLE	transformation. It is useful when the variance increases the target value increases. It reduces variance and transforms a multiplicative relationship with the target, was simpler additive relationship. This setting is applicable only to the Exponential Smoothing algorithm. When a value for EXSM_MODEL setting is not specified, the default value is ODMS_BOXCOX_ENABLE and when a value for the EXSM_MODEL setting is provided, the default value is ODMS_BOXCOX_DISABLE.
ODMS_EXPLOSION_MIN_SUPP	A positive integer	It is the minimum required support for categorical values that must be included in the explosion mapping. It removes categorical values with insufficient row instances to have a statistically significant effect on the model, however, they could potentially degrade performance. The default is system determined depending on the number of rows in the dataset. A value of 1 results into mapping all categorical values.

Table 62-8 (Cont.) Global Settings

Setting Name	Setting Value	Description
ODMS_ITEM_ID_COLUMN_NAME	column_name	(Association rules only) Name of a column that contains the items in a transaction. When this setting is specified, the algorithm expects the data to be presented in a native transactional format, consisting of two columns:

- Case ID, either categorical or numeric
- Item ID, either categorical or numeric



Oracle Machine Learning does not support BOOLEAN values for this setting.

A typical example of transactional data is market basket data, wherein a case represents a basket that may contain many items. Each item is stored in a separate row, and many rows may be needed to represent a case. The case ID values do not uniquely identify each row. Transactional data is also called multi-record case data.

Association rules function is normally used with transactional data, but it can also be applied to single-record case data (similar to other algorithms).

For more information about single-record and multi-record case data, see *Oracle SQL Developer Data Modeler User's Guide*.



Table 62-8 (Cont.) Global Settings

Setting Value Description **Setting Name** (Association rules only) Name of a column that contains a ODMS ITEM VALUE COLUMN NA column_name value associated with each item in a transaction. This setting is only used when a value has been specified for ODMS ITEM ID COLUMN NAME indicating that the data is presented in native transactional format. If ASSO AGGREGATES is used, then the build data must include the following three columns and the columns specified in the AGGREGATES setting. Case ID, either categorical or numeric Item ID, either categorical or numeric, specified by ODMS ITEM ID COLUMN NAME Item value, either categorical or numeric, specified by ODMS ITEM VALUE COLUMN NAME Note: Oracle Machine Learning does not support BOOLEAN values for this setting. If ASSO AGGREGATES, Case ID, and Item ID column are present, then the Item Value column may or may not appear. The Item Value column may specify information such as the number of items (for example, three apples) or the type of the item (for example, macintosh apples). For details on ASSO AGGREGATES, see DBMS_DATA_MINING - Mining Function Settings. ODMS MISSING VALUE TREATM ODMS MISSING VALUE MEA Indicates how to treat missing values in the training data. This setting does not affect the scoring data. The default ENT N MODE value is ODMS MISSING VALUE AUTO. ODMS MISSING VALUE DEL ODMS MISSING VALUE MEAN MODE replaces missing ETE ROW values with the mean (numeric attributes) or the mode ODMS_MISSING_VALUE_AUT (categorical attributes) both at build time and apply time where appropriate. ODMS MISSING VALUE AUTO performs different strategies for different algorithms. When ODMS MISSING VALUE TREATMENT is set to ODMS MISSING VALUE DELETE ROW, the rows in the training data that contain missing values are deleted. However, if you want to replicate this missing value treatment in the scoring data, then you must perform the

transformation explicitly.

all algorithms.

The value <code>ODMS_MISSING_VALUE_DELETE_ROW</code> applies to



Table 62-8 (Cont.) Global Settings

Setting Name	Setting Value	Description
ODMS_ROW_WEIGHT_COLUMN_NA ME	column_name	(GLM only) Name of a column in the training data that contains a weighting factor for the rows. The column data type must be numeric. Oracle Machine Learning does not support BOOLEAN values for this setting.
		Row weights can be used as a compact representation of repeated rows, as in the design of experiments where a specific configuration is repeated several times. Row weights can also be used to emphasize certain rows during model construction. For example, to bias the model towards rows that are more recent and away from potentially obsolete data.
ODMS_TEXT_POLICY_NAME	The name of an Oracle Text POLICY created using	Affects how individual tokens are extracted from unstructured text.
	CTX_DDL.CREATE_POLICY.	For details about CTX_DDL.CREATE_POLICY, see Oracle Text Reference.
ODMS_TEXT_MAX_FEATURES	1 <= value	The maximum number of distinct features, across all text attributes, to use from a document set passed to CREATE_MODEL. The default is 3000. ESA has the default value of 300000.
ODMS_TEXT_MIN_DOCUMENTS	Non-negative value	This is a text processing setting the controls how in how many documents a token needs to appear to be used as a feature.
		The default is 1. ESA has a default of 3.
ODMS_PARTITION_COLUMNS	Comma separated list of machine learning attributes	This setting indicates a request to build a partitioned model. The setting value is a comma-separated list of the machine learning attributes used to determine the in-list partition key values. Oracle Machine Learning supports numeric and categorical values including BOOLEAN for this setting. These machine learning attributes are taken from the input columns unless an XFORM_LIST parameter is passed to CREATE_MODEL or CREATE_MODEL2. If the XFORM_LIST parameter is passed to during model building, then the machine learning attributes are taken from the attributes produced by these transformations.
ODMS_MAX_PARTITIONS	1< value <= 1000000	This setting indicates the maximum number of partitions allowed for the model. The default is 1000 .
ODMS_SAMPLING	ODMS_SAMPLING_ENABLE ODMS_SAMPLING_DISABLE	This setting allows the user to request a sampling of the build data. The default is <code>ODMS_SAMPLING_DISABLE</code> .
ODMS_SAMPLE_SIZE	0 < Value	This setting determines how many rows will be sampled (approximately). It can be set only if <code>ODMS_SAMPLING</code> is enabled. The default value is the system determined.



Table 62-8 (Cont.) Global Settings

Setting Name	Setting Value	Description
ODMS_PARTITION_BUILD_TYPE	ODMS_PARTITION_BUILD_I NTRA	This setting controls the parallel build of partitioned models.
	ODMS_PARTITION_BUILD_I	ODMS_PARTITION_BUILD_INTRA — Each partition is built in parallel using all replicas.
	ODMS_PARTITION_BUILD_H YBRID	ODMS_PARTITION_BUILD_INTER — Each partition is built entirely in a single slave, but multiple partitions may be built at the same time since multiple replicas are active.
		ODMS_PARTITION_BUILD_HYBRID — It is a combination of the other two types and is recommended for most situations to adapt to dynamic environments.
		The default mode is <code>ODMS_PARTITION_BUILD_HYBRID</code>
ODMS_TABLESPACE_NAME	tablespace_name	This setting controls the storage specifications.
		If you explicitly sets this to the name of a tablespace (for which you have sufficient quota), then the specified tablespace storage creates the resulting model content. If you do not provide this setting, then the default tablespace of the user creates the resulting model content.
ODMS_RANDOM_SEED	The value must be a non- negative integer	The hash function with a random number seed generates a random number with uniform distribution. Users can control the random number seed by this setting. The default is 0.
		This setting is used by Random Forest, Neural Network, and CUR Matrix Decomposition.
ODMS_DETAILS	• ODMS_ENABLE • ODMS_DISABLE	This setting reduces the space that is used while creating a model, especially a partitioned model. The default value is <code>ODMS_ENABLE</code> .
		When the setting is <code>ODMS_ENABLE</code> , it creates model tables and views when the model is created. You can query the model with SQL. When the setting is <code>ODMS_DISABLE</code> , model views are not created and tables relevant to model details are not created either.
		The reduction in space depends on the model. Reduction on the order of 10x can be achieved.

Oracle Machine Learning for SQL Concepts for information about GLM

Oracle Machine Learning for SQL Concepts for information about association rules

Oracle Machine Learning for SQL User's Guide for information about machine learning unstructured text

DBMS_DATA_MINING — Algorithm Specific Model Settings

Oracle Machine Learning for SQL uses algorithm specific settings to define the characteristics of a model.

All settings have default values. If you want to override one or more of the settings for a model, then you must specify those settings.

The information you provide in the settings table is used by the model at build time. The name of the settings table is an optional argument to the CREATE_MODEL Procedure. You can also provide these settings through the CREATE_MODEL2 Procedure.

The settings used by a model can be found by querying the data dictionary view <code>ALL_MINING_MODEL_SETTINGS</code>. This view displays the model settings used by the machine learning models to which you have access. All of the default and user-specified setting values are included in the view.

See Also:

- ALL MINING MODEL SETTINGS in Oracle Database Reference
- Oracle Machine Learning for SQL User's Guide for information about specifying model settings

DBMS_DATA_MINING — Algorithm Settings: ALGO_EXTENSIBLE_LANG

The settings listed in the following table configure the behavior of the machine learning model with an extensible algorithm. The model is built in the R language.

The RALG_*_FUNCTION specifies the R script that is used to build, score, and view an R model and must be registered in the Oracle Machine Learning for R script repository. The R scripts are registered through Oracle Machine Learning for R with special privileges. When ALGO_EXTENSIBLE_LANG is set to R in the MINING_MODEL_SETTING table, the machine learning model is built in the R language. After the R model is built, the names of the R scripts are recorded in the MINING_MODEL_SETTING table in the SYS schema. The scripts must exist in the script repository for the R model to function. The amount of R memory used to build, score, and view the R model through these R scripts can be controlled by Oracle Machine Learning for R.

All algorithm-independent DBMS_DATA_MINING subprograms can operate on an R model for machine learning functions such as association, attribute importance, classification, clustering, feature extraction, and regression.

The supported DBMS DATA MINING subprograms include, but are not limited, to the following:

- ADD COST MATRIX Procedure
- COMPUTE_CONFUSION_MATRIX Procedure
- COMPUTE_LIFT Procedure
- COMPUTE ROC Procedure
- CREATE_MODEL Procedure
- DROP_MODEL Procedure



- EXPORT_MODEL Procedure
- GET_MODEL_COST_MATRIX Function
- IMPORT_MODEL Procedure
- REMOVE_COST_MATRIX Procedure
- RENAME_MODEL Procedure

Table 62-9 ALGO_EXTENSIBLE_LANG Settings

Setting Name	Setting Value	Description
RALG_BUILD_FUNCTION	R_BUILD_FUNCTION_SCRIPT_NA ME	Specifies the name of an existing registered R script for the R algorithm machine learning model build function. The R script defines an R function for the first input argument for training data and returns an R model object. For clustering and feature extraction machine learning function model build, the R attributes dm\$nclus and dm\$nfeat must be set on the R model to indicate the number of clusters and features respectively. The RALG_BUILD_FUNCTION must be set along with ALGO_EXTENSIBLE_LANG in the model_setting_table.
RALG_BUILD_PARAMETER	SELECT <i>value</i> param_name,FROM DUAL	Specifies a list of numeric and string scalar for optional input parameters of the model build function.
RALG_SCORE_FUNCTION	R_SCORE_FUNCTION_SCRIPT_NA ME	Specifies the name of an existing registered R script to score data. The script returns a data.frame containing the corresponding prediction results. The setting is used to score data for machine learning functions such as regression, classification, clustering, and feature extraction. This setting does not apply to the association and the attribute importance functions
RALG_WEIGHT_FUNCTION	R_WEIGHT_FUNCTION_SCRIPT_N AME	Specifies the name of an existing registered R script for the R algorithm that computes the weight (contribution) for each attribute in scoring. The script returns a data.frame containing the contributing weight for each attribute in a row. This function setting is needed for the PREDICTION_DETAILS SQL function.
RALG_DETAILS_FUNCTION	R_DETAILS_FUNCTION_SCRIPT_ NAME	Specifies the name of an existing registered R script for the R algorithm that produces the model information. This setting is required to generate a model view.
RALG_DETAILS_FORMAT	SELECT <i>type_value</i> column_name, FROM DUAL	Specifies the SELECT query for the list of numeric and string scalars for the output column type and the column name of the generated mode view. This setting is required to generate a model view.

Oracle Machine Learning for SQL User's Guide



DBMS_DATA_MINING — Algorithm Settings: CUR Matrix Decomposition

The following settings affects the behavior of the CUR Matrix Decomposition algorithm.

The Constant Value column specifies constants using the prefix DBMS_DATA_MINING. For example, DBMS_DATA_MINING.CURS_ROW_IMP_DISABLE. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the DBMS_DATA_MINING prefix, in single quotes. For example, 'CURS_ROW_IMP_DISABLE'.



The distinction between **Constant Value** and **String Value Equivalent** for this algorithm is applicable to Oracle Database 19c and Oracle Database 21c.

Table 62-10 CUR Matrix Decomposition Settings

Setting Name	Constant Value	String Value Equivalent	Description
CURS_APPROX_ATTR_N UM	The value must be a positive integer	The value must be a positive integer	Defines the approximate number of attributes to be selected.
			The default value is the number of attributes.
CURS_ROW_IMPORTANC E	CURS_ROW_IMP_ENABL E	CURS_ROW_IMP_ENABL E	Defines the flag indicating whether or not to perform row selection.
			Enables row selection.
			The default value is CURS_ROW_IMP_DISABLE.
	CURS_ROW_IMP_DISAB	CURS_ROW_IMP_DISAB	Disables row selection.
CURS_APPROX_ROW_NU	The value must be a positive integer	The value must be a positive integer	Defines the approximate number of rows to be selected. This parameter is only used when users decide to perform row selection (CURS_ROW_IMP_ENABLE).
			The default value is the total number of rows.
CURS_SVD_RANK	The value must be a positive integer	The value must be a positive integer	Defines the rank parameter used in the column/row leverage score calculation.
			If users do not provide an input value, the value is determined by the system.

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts

DBMS_DATA_MINING — Algorithm Settings: Decision Tree

These settings configure the behavior of the Decision Tree algorithm. Note that the Decision Tree settings are also used to configure the behavior of Random Forest as it constructs each individual decision tree.

The Constant Value column specifies constants using the prefix DBMS_DATA_MINING. For example, DBMS_DATA_MINING.TREE_IMPURITY_ENTROPY. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the DBMS_DATA_MINING prefix, in single quotes. For example, 'TREE IMPURITY ENTROPY'.



The distinction between **Constant Value** and **String Value Equivalent** for this algorithm is applicable to Oracle Database 19c and Oracle Database 21c.

Table 62-11 Decision Tree Settings

Setting Name	Constant Value	String Value Equivalent	Description
TREE_IMPURITY_METRIC	TREE_IMPURITY_ENTROP Y	TREE_IMPURIT Y_ENTROPY	Tree impurity metric for Decision Tree. Tree algorithms seek the best test question for splitting data at each node. The best splitter and split values are those that result in the largest increase in target value homogeneity (purity) for the entities in the node. Purity is by a metric. By default, the algorithm uses TREE_IMPURITY_GINI.
	TREE_IMPURITY_GINI	TREE_IMPURIT Y_GINI	Decision trees can use either Gini (TREE_IMPURITY_GINI) or entropy (TREE_IMPURITY_ENTROPY) as the purity metric.
TREE_TERM_MAX_DEPTH	For Decision Tree: 2<= a number <=20 For Random Forest: 2<= a number <=100	For Decision Tree: 2<= a number <=20 For Random	Criteria for splits: maximum tree depth (the maximum number of nodes between the root and any leaf node, including the leaf node). For Decision Tree, the default is 7. For Random Forest, the default is 16.
		Forest: 2<= a number <=100	
TREE_TERM_MINPCT_NOD E	0<= a number<=10	0<= a number<=10	The minimum number of training rows in a node expressed as a percentage of the rows in the training data. Default is 0.05, indicating 0.05%.



Table 62-11 (Cont.) Decision Tree Settings

Setting Name	Constant Value	String Value Equivalent	Description
TREE_TERM_MINPCT_SPL	0 < a number <=20	0 < a number <=20	The minimum number of rows required to consider splitting a node expressed as a percentage of the training rows. Default is 0.1, indicating 0.1%.
TREE_TERM_MINREC_NOD E	a number>=0	a number>=0	The minimum number of rows in a node. Default is 10.
TREE_TERM_MINREC_SPL	a number > 1	a number > 1	Criteria for splits: minimum number of records in a parent node expressed as a value. No split is attempted if the number of records is below this value. Default is 20.

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts for information about Decision Tree

DBMS_DATA_MINING — Algorithm Settings: Expectation Maximization

These algorithm settings configure the behavior of the Expectation Maximization algorithm.



Oracle Data Mining Concepts for information about Expectation Maximization

Table 62-12 Expectation Maximization Settings for Data Preparation and Analysis

Setting Name	Constant Value	String Value	Description
		Equivalent	
EMCS_ATTRIBUTE_FILTER	EMCS_ATTR_FILTER_EN ABLE	EMCS_ATTR_FI LTER_ENABLE	Whether or not to include uncorrelated attributes in the model. When EMCS_ATTRIBUTE_FILTER is enabled, uncorrelated attributes are not included.
			Note: This setting applies only to attributes that are not nested.
			For Clustering, the default is system-determined.
			For anomaly detection, the default is EMCS_ATTR_FILTER_DISABLE.
	EMCS_ATTR_FILTER_DI SABLE	EMCS_ATTR_FI LTER_DISABLE	Includes uncorrelated attributes in the model.
EMCS_MAX_NUM_ATTR_2D	An integer greater than or equal to 1, represented as a	An integer greater than or equal to 1,	Maximum number of correlated attributes to include in the model.
	character string.	represented as a character	Note: This setting applies only to attributes that are not nested (2D).
		string.	Default is 50. Expression:
			TO CHAR(40)
EMCS_NUM_DISTRIBUTION	EMCS_NUM_DISTR_BERN OULLI	EMCS_NUM_DIS TR_BERNOULLI	The distribution for modeling numeric attributes. Applies to the input table or view as a whole and does not allow per-attribute specifications.
			The options include Bernoulli, Gaussian, or system-determined distribution. When Bernoulli or Gaussian distribution is chosen, all numeric attributes are modeled using the same type of distribution.
			Default is EMCS_NUM_DISTR_SYSTEM.
	EMCS_NUM_DISTR_GAUS SIAN	EMCS_NUM_DIS TR_GAUSSIAN	Models all numeric attributes using Gaussian distribution.
	EMCS_NUM_DISTR_SYST EM	EMCS_NUM_DIS TR_SYSTEM	When the distribution is system-determined, individual attributes may use different distributions (either Bernoulli or Gaussian), depending on the data.
EMCS_NUM_EQUIWIDTH_BI	An integer between 1 to 255, inclusive, represented as a character string.	An integer between 1 to 255, inclusive, represented as a character string.	Number of equi-width bins that will be used for gathering cluster statistics for numeric columns. Default is 11. Expression: TO_CHAR(20)



Table 62-12 (Cont.) Expectation Maximization Settings for Data Preparation and Analysis

Setting Name	Constant Value	String Value Equivalent	Description
EMCS_NUM_PROJECTIONS	An integer greater than or equal to 1, represented as a character string.	n An integer greater than or equal to 1, represented as a character	Specifies the number of projections that will be used for each nested column. If a column has fewer distinct attributes than the specified number of projections, the data will not be projected. The setting applies to all nested columns.
		string.	Default is 50.
			Expression:
			TO_CHAR(40)
EMCS_NUM_QUANTILE_BIN S	An integer between 1 to 255, inclusive, represented as a	55, inclusive, between 1 to epresented as a 255, inclusive, represented as	Specifies the number of quantile bins that will be used for modeling numeric columns with multivalued Bernoulli distributions.
	character string.		Default is system-determined.
		a character string.	Expression:
		Stillig.	TO_CHAR(20)
EMCS_NUM_TOPN_BINS	An integer between 1 to 255, inclusive, represented as a	between 1 to	Specifies the number of top-N bins that will be used for modeling categorical columns with multivalued Bernoulli distributions.
	represented as	Default is system-determined.	
		a character string.	Expression:
		oung.	TO_CHAR(10)

Table 62-13 Expectation Maximization Settings for Learning

Setting Name	Constant Value	String Value Equivalent	Description
EMCS_CONVERGENCE_C RITERION	EMCS_CONV_CRIT_HEL DASIDE	EMCS_CONV_CRIT_HEL DASIDE	The convergence criterion for EM. The convergence criterion may be based on a held-aside data set, or it may be Bayesian Information Criterion.
			EMCS_CONV_CRIT_HELDASIDE: Uses a heldaside data set for convergence criterion.
			Default is system determined.
	EMCS_CONV_CRIT_BIC	EMCS_CONV_CRIT_BIC	Uses the Bayesian Information Criterion (BIC) for convergence.
EMCS_LOGLIKE_IMPRO VEMENT	A floating point number between 0 and 1 expressed as a character string	A floating point number between 0 and 1 expressed as a character string	When the convergence criterion is based on a held-aside data set (EMCS_CONVERGENCE_CRITERION = EMCS_CONV_CRIT_HELDASIDE), this setting specifies the percentage improvement in the value of the log likelihood function that is required for adding a new component to the model. Default value is 0.001.
			Expression:
			TO_CHAR(0.003)



Table 62-13 (Cont.) Expectation Maximization Settings for Learning

Setting Name	Constant Value	String Value Equivalent	Description
EMCS_NUM_COMPONENT S	An integer greater than or equal to 1, represented as a character string	An integer greater than or equal to 1, represented as a character string	Maximum number of components in the model. If model search is enabled, the algorithm automatically determines the number of components based on improvements in the likelihood function or based on regularization, up to the specified maximum.
			For EM Clustering, the number of components must be greater than or equal to the number of clusters.
			Default is 20 for both EM Clustering and EM Anomaly.
			Expression:
			TO_CHAR(20)
EMCS_NUM_ITERATION S	An integer greater than or equal to 1,	An integer greater than or equal to 1,	Specifies the maximum number of iterations in the EM algorithm.
	represented as a	represented as a	Default is 100.
	character string	character string	Expression:
			TO_CHAR(50)
EMCS_MODEL_SEARCH	EMCS_MODEL_SEARCH_ ENABLE	EMCS_MODEL_SEARCH_ ENABLE	This setting enables model search in EM where different model sizes are explored and a best size is selected.
			The default is EMCS_MODEL_SEARCH_DISABLE.
	EMCS_MODEL_SEARCH_DISABLE (default).	EMCS_MODEL_SEARCH_DISABLE (default).	The model search in EM is disabled.
EMCS_REMOVE_COMPON ENTS	EMCS_REMOVE_COMPS_ ENABLE (default)	EMCS_REMOVE_COMPS_ ENABLE (default)	This setting allows the EM algorithm to remove a small component from the solution. The default is EMCS REMOVE COMPS ENABLE.
	EMCS_REMOVE_COMPS_ DISABLE	EMCS_REMOVE_COMPS_ DISABLE	
EMCS_RANDOM_SEED	Non-negative integer	Non-negative integer	This setting controls the seed of the random generator used in EM. The default is 0.

Table 62-14 Expectation Maximization Settings for Component Clustering

Setting Name	Constant Value	String Value Equivalent	Description
EMCS_CLUSTER_COMPO NENTS	EMCS_CLUSTER_COMP_ ENABLE	EMCS_CLUSTER_COMP_ ENABLE	Enables or disables the grouping of EM components into high-level clusters. When disabled, the components themselves are treated as clusters. The setting can be used only for EM Clustering.
			When component clustering is enabled, model scoring through the SQL CLUSTER function will produce assignments to the higher level clusters.
			Default is EMCS_CLUSTER_COMP_ENABLE.



Table 62-14 (Cont.) Expectation Maximization Settings for Component Clustering

Setting Name	Constant Value	String Value Equivalent	Description
	EMCS_CLUSTER_COMP_ DISABLE	EMCS_CLUSTER_COMP_ DISABLE	When clustering is disabled, the CLUSTER function will produce assignments to the original components.
EMCS_CLUSTER_THRES H	Specify an integer greater than or equal to 1, represented as a character string	Specify an integer greater than or equal to 1, represented as a character string	Dissimilarity threshold that controls the clustering of EM components. When the dissimilarity measure is less than the threshold, the components are combined into a single cluster. The setting can be used only for EM Clustering.
			A lower threshold may produce more clusters that are more compact. A higher threshold may produce fewer clusters that are more spread out. Default is 2.
			Expression:
			TO_CHAR(3)
EMCS_LINKAGE_FUNCT	EMCS_LINKAGE_SINGL	EMCS_LINKAGE_SINGL	Allows the specification of a linkage function for the agglomerative clustering step.
			EMCS_LINKAGE_SINGLE uses the nearest distance within the branch. The clusters tend to be larger and have arbitrary shapes.
			Default is EMCS_LINKAGE_SINGLE.
	EMCS_LINKAGE_AVERA GE	EMCS_LINKAGE_AVERA GE	EMCS_LINKAGE_AVERAGE uses the average distance within the branch. There is less chaining effect and the clusters are more compact.
	EMCS_LINKAGE_COMPL ETE	EMCS_LINKAGE_COMPL ETE	EMCS_LINKAGE_COMPLETE uses the maximum distance within the branch. The clusters are smaller and require strong component overlap.

Table 62-15 Expectation Maximization Settings for Cluster Statistics

Setting Name	Constant Value	Description
EMCS_CLUSTER_STATISTICS	EMCS_CLUS_STATS_ENAB	Enables or disables the gathering of descriptive statistics for clusters (centroids, histograms, and rules). When
	EMCS_CLUS_STATS_DISA BLE	statistics are disabled, model size is reduced, and
		Default is EMCS_CLUS_STATS_ENABLE.
EMCS_MIN_PCT_ATTR_SUPPORT	A floating point number between 0 and 1 expressed as a character string	Minimum support required for including an attribute in the cluster rule. The support is the percentage of the data rows assigned to a cluster that must have non-null values for the attribute. The setting can be used only for EM Clustering. Default is 0.1.
		Expression:
		TO_CHAR(0.9)



Table 62-16 Expectation Maximization Settings for Anomaly Detection

Setting Name	Constant Value	String Value Equivalent	Description
EMCS_OUTLIER_RATE	OUTLIER_RATE A floating point number between 0 and 1 between 0 and 1 expressed as a character string A floating point number between 0 and 1 expressed as a character string	.	The desired rate of outliers in the training data. The setting can be used only for EM Anomaly.
			Default is 0.05.
		character string	Expression:
			TO_CHAR(0.07)

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.

DBMS_DATA_MINING — Algorithm Settings: Explicit Semantic Analysis

Explicit Semantic Analysis (ESA) is a useful technique for extracting meaningful and interpretable features.

The settings listed in the following table configure the ESA values.

Table 62-17 Explicit Semantic Analysis Settings

Setting Name	Setting Value	String Value Equivalent	Description
ESAS_EMBEDDINGS	ESAS_EMBEDDINGS_ENAB LE	ESAS_EMBEDDINGS_ENABLE	This setting applies to feature extraction models. The default value is ESAS_EMBEDDINGS_DISABLE. When you set ESAS_EMBEDDINGS_ENABLE: ESA generates embeddings during scoring The FEATURE_ID of the generated embeddings is of the data type NUMBER The CASE_ID_COLUMN_NAME argument of the DBMS_DATA_MINING.CREATE MODEL and DBMS_DATA_MINING.CREATE MODEL2 function is optional.
	ESAS_EMBEDDINGS_DISA BLE	ESAS_EMBEDDINGS_DISABLE	Disables the use of embeddings for ESA. This setting is useful when embeddings are not required or desired for the analysis



Table 62-17	(Cont.)	Explici	t Semantic Anal	ysis Settings
--------------------	---------	---------	-----------------	---------------

Setting Name	Setting Value	String Value Equivalent	Description
ESAS_EMBEDDING_SIZE	A positive integer less than or equal to 4096	A positive integer less than or equal to 4096	This setting applies to feature extraction models. This setting specifies the size of the vectors representing embeddings. You can set this parameter only if you have enabled ESAS_EMBEDDINGS. The default size is 1024. If this value is less than the number of distinct features in the training set, then the actual number of explicit features is used as the size of embedding vectors instead.
ESAS_MIN_ITEMS	Text input 100	Text input 100	This setting determines the
	Non-text input is 0	Non-text input is 0	minimum number of non-zero entries that need to be present in an input row. The default is 100 for text input and 0 for non-text input.
ESAS_TOPN_FEATURES	A positive integer	A positive integer	This setting controls the maximum number of features per attribute. The default is 1000.
ESAS_VALUE_THRESHOLD	Non-negative number	Non-negative number	This setting thresholds a small value for attribute weights in the transformed build data. The default is 1e-8.

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts for information about ESA.

DBMS_DATA_MINING — Algorithm Settings: Exponential Smoothing

These settings configure the behavior of the Exponential Smoothing (ESM) algorithm.

The settings listed in the following table specify the setting names and possible values for Exponential Smoothing. You can specify the Setting Value using the prefix <code>DBMS_DATA_MINING</code>. For example, <code>DBMS_DATA_MINING.EXSM_SIMPLE</code>. Alternatively, you can specify the Setting Value without the <code>DBMS_DATA_MINING</code> prefix, in single quotes. For example, <code>'EXSM_SIMPLE'</code>.

For Global settings, see DBMS_DATA_MINING — Global Settings.

Table 62-18 Exponential Smoothing Settings

Setting Name	Setting Value	Description
EXSM_MODEL	EXSM_SIMPLE	This setting specifies the model.
	EXSM_SIMPLE_MULT_ERR	EXSM_SIMPLE: Forecasts data as a weighted
	EXSM_HOLT	moving average, with the influence of past observations declining exponentially with the
	EXSM_HOLT_DAMPED	length of time since the observation occurred.
	EXSM_MULT_TREND	Errors in estimation are assumed to be norma
	EXSM_MULT_TREND_DAMPED	distributed, with constant mean and variance. is appropriate for data with no clear trend or
	EXSM_SEASON_ADD	seasonal pattern.
	EXSM_SEASON_MUL	EXSM_SIMPLE_MULT_ERR: Forecasts data as
	EXSM_WINTERS	weighted moving average, with the influence of
	EXSM_WINTERS_DAMPED	past observations declining exponentially with the length of time since the observation
	EXSM_ADDWINTERS	occurred. Errors in estimation are assumed to
	EXSM_ADDWINTERS_DAMPED	be proportional to the level of the prior estimate
	EXSM_WINTERS_MUL_TREND	EXSM_HOLT: Applies Holt's linear exponential smoothing method, designed to forecast data
	EXSM_WINTERS_MUL_TREND_DMP	with an underlying linear trend.
		EXSM HOLT DAMPED: Applies Holt's linear
		exponential smoothing with a damping factor to progressively reduce the strength of the trend over time.
		EXSM_MULT_TREND: Applies an exponential smoothing framework with a multiplicative tren component, effectively capturing data where trends are not linear but grow or decay over time.
		EXSM_MULT_TREND_DAMPED: Applies an exponential smoothing algorithm with a multiplicative trend that diminishes over time, providing a conservative approach to trend estimation.
		EXSM_SEASON_ADD: Applies an exponential
		smoothing with an additive seasonal component, isolating and accounting for seasonal variations without incorporating a trend.
		EXSM_SEASON_MUL: Executes exponential smoothing with a multiplicative seasonal component, capturing seasonal effects that increase or decrease in proportion to the leve the series.
		EXSM_WINTERS: Applies the Holt-Winters method with additive trends and multiplicative seasonality, offering a robust model for data whoth linear trend and proportional seasonal variation.
		EXSM_WINTERS_DAMPED: Applies the Holt-Winters method with a damped trend and multiplicative seasonality, moderating the linestrend over time while still capturing proportion seasonal changes.



Table 62-18 (Cont.) Exponential Smoothing Settings

Setting Name	Setting Value	Description
		EXSM_ADDWINTERS: Applies the Holt-Winters additive model to simultaneously smooth data with linear trends and additive seasonal effects.
		EXSM_ADDWINTERS_DAMPED: Applies the Holt-Winters additive approach with a damping mechanism, reducing the impact of the trend and seasonal components over time.
		EXSM_WINTERS_MULT_TREND: Applies the Holt-Winters model with both trend and seasonality components being multiplicative, suited for series where the seasonal variations and trends are both increasing or decreasing proportional to level.
		EXSM_WINTERS_MUL_TREND_DMP: Applies the Holt-Winters model with a damped multiplicative trend, effectively moderating the exponential increase or decrease of both trend and seasonal components over time.
		The default value is EXSM_SIMPLE.
EXSM_SEASONALITY	positive integer > 1	This setting specifies a positive integer value as the length of seasonal cycle. The value it takes must be larger than 1. For example, setting value 4 means that every group of four observations forms a seasonal cycle.
		This setting is only applicable and must be provided for models with seasonality, otherwise the model throws an error.
		When EXSM_INTERVAL is not set, this setting applies to the original input time series. When EXSM_INTERVAL is set, this setting applies to the accumulated time series.



Table 62-18 (Cont.) Exponential Smoothing Settings

Setting Name	Setting Value	Description
EXSM_INTERVAL_ EXSM_INTERVAL_	EXSM_INTERVAL_YEAR EXSM_INTERVAL_QTR EXSM_INTERVAL_MONTH EXSM_INTERVAL_WEEK	This setting only applies and must be provided when the time column (case_id column) has datetime type. It specifies the spacing interval of the accumulated equally spaced time series. The model throws an error if the time column of
	EXSM_INTERVAL_DAY EXSM_INTERVAL_HOUR	input table is of datetime type and setting EXSM_INTERVAL is not provided.
	EXSM_INTERVAL_MINUTE EXSM_INTERVAL_SECOND	The model throws an error if the time column of input table is of oracle number type and setting EXSM_INTERVAL is provided.
		EXSM_INTERVAL_YEAR: This option sets the spacing interval of the accumulated time series to one year. When selected, the data is aggregated or summarized on a yearly basis.
		EXSM_INTERVAL_QTR: This option sets the spacing interval to a quarter, aggregating the data for every three months.
		EXSM_INTERVAL_MONTH: This option adjusts the spacing interval to one month. The accumulated time series represent aggregated or summarized data for each month.
		EXSM_INTERVAL_WEEK: With this option data is aggregated or summarized on a weekly basis, setting the spacing interval to one week.
		EXSM_INTERVAL_DAY: This option adjusts the spacing interval to one day. It's suitable for scenarios where daily aggregated insights are required.
		EXSM_INTERVAL_HOUR: For more granular insights, this option sets the spacing interval to one hour. It's especially useful when analyzing data that changes significantly within a day.
		EXSM_INTERVAL_MINUTE: With this option the spacing is set to one minute. This provides a very detailed view of data, suitable for applications like high-frequency trading or real-time monitoring systems.
		EXSM_INTERVAL_SECOND: For most granular details, this options sets the spacing interval to one second. It's tailored for scenarios requiring real-time or near-real-time analysis.



Table 62-18 (Cont.) Exponential Smoothing Settings

Setting Name	Setting Value	Description
EXSM_INITVL_OPTIMIZE	EXSM_INITVL_OPTIMIZE_ENABLE	The setting EXSM_INITVL_OPTIMIZE
	EXSM_INITVL_OPTIMIZE_DISABLE	determines whether initial values are optimized during model build. The default value is
		EXSM_INITVL_OPTIMIZE_ENABLE.



EXSM_INITVL_OP TIMIZE can only be set to EXSM INITVL OP TIMIZE DISABLE if the user has set EXSM MODEL to EXSM HW or EXSM HW ADDSEA. If EXSM MODEL is set to another model type or is not specified, error 40213 (conflicting settings) is thrown and the model is not built.



Table 62-18 (Cont.) Exponential Smoothing Settings

Setting Name	Setting Value	Description
EXSM_ACCUMULATE	EXSM_ACCU_TOTAL	This setting only applies and must be provided
	EXSM_ACCU_STD	when the time column has datetime type. It specifies how to generate the value of the
	EXSM_ACCU_MAX	accumulated time series from the input time
	EXSM_ACCU_MIN	series.
	EXSM_ACCU_AVG	${\tt EXSM_ACCU_TOTAL:} \ \textbf{This option calculates the}$
	EXSM_ACCU_MEDIAN	total sum of the time series values within a specified interval. When selected, it will
	EXSM_ACCU_COUNT	aggregate the data by summing up all the individual values in the datetime range.
		EXSM_ACCU_STD: This option computes the standard deviation of the time series values within a specified interval. It helps you understand the amount of variation or dispersio in your data.
		EXSM_ACCU_MAX: By selecting this option, the maximum value of the time series within a specified interval will be determined. It helps in identifying the peak value in the given range.
		EXSM_ACCU_MIN: This option focuses on determining the minimum value of the time series within a specified interval. It is useful for identifying the lowest value in the time series fo the given datetime range.
	EXSM_ACCU_AVG: This specifies the average value of your time series within a specified interval. It calculates the mean value of all data points in the specified range.	
	EXSM_ACCU_MEDIAN: This option provides the median of the time series values within the give interval. The median gives a central value, which can be especially useful if your data contains outliers.	
		EXSM_ACCU_COUNT: This option counts the number of time series values within the specific interval. It is helpful if you want to know how many data points are present in a certain datetime range.
		The default value is EXSM_ACCU_TOTAL.



Table 62-18 (Cont.) Exponential Smoothing Settings

Setting Name	Setting Value	Description
EXSM_SETMISSING	Specify an option: EXSM_MISS_MIN EXSM_MISS_MAX EXSM_MISS_AVG EXSM_MISS_MEDIAN EXSM_MISS_LAST EXSM_MISS_FIRST EXSM_MISS_PREV	This setting specifies how to handle missing values, which may come from input data and/or the accumulation process of time series. You can specify either a number or an option. If a number is specified, all the missing values are set to that number. EXSM_MISS_MIN: Replaces missing value with minimum of the accumulated time series. EXSM_MISS_MAX: Replaces missing value with maximum of the accumulated time series.
	EXSM_MISS_NEXT EXSM_MISS_AUTO	EXSM_MISS_AVG: Replaces missing value with average of the accumulated time series. EXSM_MISS_MEDIAN: Replaces missing value with median of the accumulated time series. EXSM_MISS_LAST: Replaces missing value with last non-missing value of the accumulated time series.
		EXSM_MISS_FIRST: Replaces missing value with first non-missing value of the accumulated time series.
		EXSM_MISS_PREV: Replaces missing value with the previous non-missing value of the accumulated time series.
		EXSM_MISS_NEXT: Replaces missing value with the next non-missing value of the accumulated time series.
		EXSM_MISS_AUTO: EXSM model treats the inpudata as an irregular (non-uniformly spaced) time series.
		If this setting is not provided, EXSM_MISS_AUTO is the default value. In such a case, the model treats the input time series as irregular time series, viewing missing values as gaps.
EXSM_PREDICTION_STEP	It must be set to a number between 1-30.	This setting specifies how many steps ahead the predictions are to be made.
		If it is not set, the default value is 1: the model gives one-step-ahead prediction. A value greater than 30 results in an error.
EXSM_CONFIDENCE_LEVEL	It must be a number between 0 and 1, exclusive.	This setting specifies the desired confidence level for prediction.
		The lower and upper bounds of the specified confidence interval is reported. If this setting is not specified, the default confidence level is 95%



Table 62-18 (Cont.) Exponential Smoothing Settings

Setting Name	Setting Value	Description
EXSM_OPT_CRITERION	EXSM_OPT_CRIT_LIK EXSM_OPT_CRIT_MSE EXSM_OPT_CRIT_AMSE	This setting specifies the desired optimization criterion. The optimization criterion is useful as a diagnostic for comparing models' fit to the same data.
	EXSM_OPT_CRIT_SIG EXSM_OPT_CRIT_MAE	EXSM_OPT_CRIT_LIK: This represents the negative double of the logarithm of the likelihood associated with a given model.
		EXSM_OPT_CRIT_MSE: This provides the mean squared error pertaining to the model.
		EXSM_OPT_CRIT_AMSE: This denotes the average of the mean squared error over a time window as specified by the user.
		EXSM_OPT_CRIT_SIG: This metric captures the standard deviation of the residuals of the model.
		EXSM_OPT_CRIT_MAE: This metric conveys the average absolute error associated with the model. It measures the size of the error.
		The default value is EXSM_OPT_CRIT_LIK.
EXSM_NMSE	positive integer	This setting specifies the length of the window used in computing the error metric average mean square error (AMSE).
EXSM_SERIES_LIST	Comma delimited list of time series columns	This setting allows you to forecast up to twenty predictor series in addition to the target series. The column names in EXSM_SERIES_LIST are enclosed in single quotes.



The list is enclosed in single quotes, not the individual column names.

For example:

INSERT INTO <settings_table_name
VALUES(dbms_data_mining.exsm_series
_list,
'<column1>,<column2>,<column3>,<col
umn4>');

The prefix DM\$ must be added to the build and scoring data sets. The column names must be less than 125 characters long. See Model Detail Views for Exponential Smoothing.



Table 62-18 (Cont.) Exponential Smoothing Settings

Setting Name	Setting Value	Description
EXSM_BACKCAST_OUTPUT	EXSM_BACKCAST_OUTPUT_ENABLE EXSM_BACKCAST_OUTPUT_DISABLE	This setting enables the user to optionally suppress the output of backcast values. Backcasts are the model estimates for historical data. See Backcasts in Time Series for information on backcasts. Suppressing the output of backcast values can provide a potentially large reduction in the memory and storage requirements for a partitioned ESM model with a huge number of partitions. The default value is EXSM_BACKCAST_OUTPUT_ENABLE.

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.

See Also:

Oracle Machine Learning for SQL Concepts for information about ESM.

https://github.com/oracle-samples/oracle-db-examples/tree/main/machine-learning/sql browse to the release folder and click the oml4sql-time-series-exponential-smoothing.sql example.

DBMS_DATA_MINING — Algorithm Settings: Generalized Linear Model

The settings listed in the following table configure the behavior of the Generalized Linear Model algorithm.

The settings listed in the following table specify the setting names and possible values for Generalized Linear Model. The Constant Value column specifies constants using the prefix <code>DBMS_DATA_MINING</code>. Alternatively, you can specify the corresponding string value from the String Value Equivalent column.

For Global settings, see DBMS DATA MINING — Global Settings.

For generic machine learning function settings, see DBMS_DATA_MINING — Machine Learning Functions.



Table 62-19 DBMS_DATA_MINING GLM Settings

Setting Name	Constant Value	String Value Equivalent	Description
GLMS_CONF_LEVEL	A floating point number between 0 and 1 expressed as a character string	A floating point number between 0 and 1 expressed as a character string	The confidence level for coefficient confidence intervals.
			The default confidence level is 0.95.
			Expression:
			TO_CHAR(0.98)
GLMS_FTR_GEN_METHOD	GLMS_FTR_GEN_QUADRATIC GLMS FTR GEN CUBIC	GLMS_FTR_GEN_QUADRATIC GLMS_FTR_GEN_CUBIC	Whether feature generation is quadratic or cubic.
		<u></u>	When feature generation is enabled, the algorithm automatically chooses the most appropriate feature generation method based on the data.
GLMS_FTR_GENERATION	GLMS_FTR_GENERATION_EN ABLE	ABLE	Whether or not feature generation is enabled for GLM. By default, feature generation is not enabled.
	GLMS_FTR_GENERATION_DI SABLE	GLMS_FTR_GENERATION_DI SABLE	Note: Feature generation can only be enabled when feature selection is also enabled.
GLMS_FTR_SEL_CRIT	GLMS_FTR_SEL_AIC	GLMS_FTR_SEL_AIC	Feature selection penalty criterion for adding a feature to the model.
	GLMS_FTR_SEL_SBIC	GLMS_FTR_SEL_SBIC	When feature selection is enabled,
	GLMS_FTR_SEL_RIC GLMS_FTR_SEL_ALPHA_INV	GLMS_FTR_SEL_RIC GLMS_FTR_SEL_ALPHA_INV	the algorithm automatically chooses the penalty criterion based on the data.
GLMS_FTR_SELECTION	GLMS_FTR_SELECTION_ENA BLE	GLMS_FTR_SELECTION_ENA BLE	Whether or not feature selection is enabled for GLM.
	GLMS_FTR_SELECTION_DIS ABLE	GLMS_FTR_SELECTION_DIS ABLE	By default, feature selection is not enabled.
GLMS_MAX_FEATURES	An integer greater than 0 and less than or equal to 2000, represented as a character string	An integer greater than 0 and less than or equal to 2000, represented as a character string	When feature selection is enabled, this setting specifies the maximum number of features that can be selected for the final model.
	J		By default, the algorithm limits the number of features to ensure sufficient memory.
			Expression:
			TO_CHAR(200)
GLMS_PRUNE_MODEL	GLMS_PRUNE_MODEL_ENABL E GLMS_PRUNE_MODEL_DISAB LE	E	Prune enable or disable for features in the final model. Pruning is based on T-Test statistics for linear regression, or Wald Test statistics for logistic regression. Features are pruned in a loop until all features are statistically significant with respect to the full data. When feature selection is enabled,
			When feature selection is enabled the algorithm automatically prunes based on the data.



Table 62-19 (Cont.) DBMS_DATA_MINING GLM Settings

Setting Name	Constant Value	String Value Equivalent	Description
GLMS_REFERENCE_CLAS S_NAME	target_value	target_value	The target value used as the reference class in a binary logistic regression model. Probabilities are produced for the other class.
			By default, the algorithm chooses the value with the highest prevalence (the most cases) for the reference class.
	GLMS_RIDGE_REG_ENABLE	GLMS_RIDGE_REG_ENABLE	Enable or disable ridge regression.
ON	GLMS_RIDGE_REG_DISABLE	GLMS_RIDGE_REG_DISABLE	Ridge applies to both regression and classification machine learning functions.
			When ridge is enabled, prediction bounds are not produced by the PREDICTION_BOUNDS SQL function.
			Note : Ridge may only be enabled when feature selection is not specified, or has been explicitly disabled. If ridge regression and feature selection are both explicitly enabled, then an exception is raised.
GLMS_RIDGE_VALUE	An integer greater than 0 represented as a character string	An integer greater than 0 represented as a character string	The value of the ridge parameter. This setting is only used when the algorithm is configured to use ridge regression.
			If ridge regression is enabled internally by the algorithm, then the ridge parameter is determined by the algorithm.
			Expression: TO_CHAR(5)
	GLMS_ROW_DIAG_ENABLE	GLMS_ROW_DIAG_ENABLE	Enable or disable row diagnostics.
S	<pre>GLMS_ROW_DIAG_DISABLE (default).</pre>	<pre>GLMS_ROW_DIAG_DISABLE (default).</pre>	
GLMS_CONV_TOLERANCE	The range is (0, 1) non-inclusive.	The range is $(0, 1)$ noninclusive.	Convergence Tolerance setting of the GLM algorithm
			The default value is system-determined.
GLMS_NUM_ITERATIONS	A positive integer	A positive integer	Maximum number of iterations for the GLM algorithm. The default value is system-determined.
GLMS_BATCH_ROWS	0 or a positive integer	0 or a positive integer	Number of rows in a batch used by the SGD solver. The value of this parameter sets the size of the batch for the SGD solver. An input of 0 triggers a data driven batch size estimate.



Table 62-19 (Cont.) DBMS_DATA_MINING GLM Settings

Setting Name	Constant Value	String Value Equivalent	Description
GLMS_SOLVER	GLMS_SOLVER_SGD (StochasticGradient Descent) GLMS_SOLVER_CHOL (Cholesky) GLMS_SOLVER_QR	GLMS_SOLVER_SGD (StochasticGradient Descent) GLMS_SOLVER_CHOL (Cholesky) GLMS_SOLVER_QR GLMS_SOLVER_LBFGS_ADMM	This setting allows the user to choose the GLM solver. The solver cannot be selected if GLMS_FTR_SELECTION setting is enabled. GLMS_SOLVER_SGD: Optimizes generalized linear models by iteratively updating parameters using a subset of the data to minimize errors. GLMS_SOLVER_CHOL: Solves generalized linear models using the Cholesky decomposition method, which provides a stable and efficient solution by transforming the right-hand of the equation into a lower triangular matrix and its conjugate transpose. GLMS_SOLVER_QR: Utilizes the QR decomposition technique to solve generalized linear models, ensuring numerical stability and accuracy by decomposing the problem into an orthonormal matrix Q and upper triangular matrix R. GLMS_SOLVER_LBFGS_ADMM: Combines L-BFGS, an approximation of the Broyden-Fletcher-Goldfarb-Shanno optimization algorithm, with ADMM for solving large-scale generalized linear model problems efficiently. The default value is system determined.
GLMS_SPARSE_SOLVER	GLMS_SPARSE_SOLVER_ENA BLE GLMS_SPARSE_SOLVER_DIS ABLE (default).	GLMS_SPARSE_SOLVER_ENA BLE GLMS_SPARSE_SOLVER_DIS ABLE (default).	This setting allows the user to use sparse solver if it is available. The default value is GLMS_SPARSE_SOLVER_DISABLE.



Table 62-19 (Cont.) DBMS_DATA_MINING GLM Settings

Setting Name	Constant Value	String Value Equivalent	Description
GLMS_LINK_FUNCTION	GLMS_IDENTITY_LINK	GLMS_IDENTITY_LINK	This setting allows the user to
	GLMS_LOGIT_LINK	GLMS_LOGIT_LINK	specify the link function for building
	GLMS_PROBIT_LINK	GLMS_PROBIT_LINK	a GLM model. The link functions are specific to the mining function. For
	GLMS_CLOGLOG_LINK	GLMS_CLOGLOG_LINK	classification, the following are
	GLMS_CAUCHIT_LINK	GLMS_CAUCHIT_LINK	<pre>applicable: • GLMS_LOGIT_LINK (default)</pre>
			• GLMS_PROBIT_LINK
			 GLMS_CLOGLOG_LINK
			 GLMS_CAUCHIT_LINK
			For regression, the following is applicable: GLMS_IDENTITY_LINK (default)

- DBMS_DATA_MINING Machine Learning Functions
 - A machine learning **function** refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 - The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.
- DBMS_DATA_MINING Algorithm Settings: Neural Network
 The settings listed in the following table configure the behavior of the Neural Network algorithm.
- DBMS_DATA_MINING Solver Settings: LBFGS
 The settings listed in the following table configure the behavior of L-BFGS. Neural Network and Generalized Linear Model (GLM) use these settings.
- DBMS_DATA_MINING Solver Settings: ADMM
 The settings listed in the following table configure the behavior of Alternating Direction Method of Multipliers (ADMM). The Generalized Linear Model (GLM) algorithm uses these settings.
- Oracle Machine Learning for SQL Concepts



Oracle Machine Learning for SQL Concepts for information about GLM.

DBMS_DATA_MINING — Algorithm Settings: *k*-Means

The settings listed in the following table configure the behavior of the k-Means algorithm.



Table 62-20 k-Means Settings

Setting Name	Setting Value	Description
KMNS_CONV_TOLERANCE	TO_CHAR(0 <numeric_expr<1)< td=""><td>Minimum Convergence Tolerance for <i>k</i>-Means. The algorithm iterates until the minimum Convergence Tolerance is satisfied or until the maximum number of iterations, specified in KMNS_ITERATIONS, is reached.</td></numeric_expr<1)<>	Minimum Convergence Tolerance for <i>k</i> -Means. The algorithm iterates until the minimum Convergence Tolerance is satisfied or until the maximum number of iterations, specified in KMNS_ITERATIONS, is reached.
		Decreasing the Convergence Tolerance produces a more accurate solution but may result in longer run times.
		The default Convergence Tolerance is 0.001.
KMNS_DISTANCE	KMNS_COSINE	Distance function for k-Means.
	KMNS_EUCLIDEAN	The default distance function is KMNS_EUCLIDEAN.
KMNS_ITERATIONS	TO_CHAR(positive_numeric_expr)	Maximum number of iterations for <i>k</i> -Means. The algorithm iterates until either the maximum number of iterations is reached or the minimum Convergence Tolerance, specified in KMNS_CONV_TOLERANCE, is satisfied.
		The default number of iterations is 20.
KMNS_MIN_PCT_ATTR_SUPPORT	TO_CHAR(0<=numeric_expr<=1)	Minimum percentage of attribute values that must be non-null in order for the attribute to be included in the rule description for the cluster.
		If the data is sparse or includes many missing values, a minimum support that is too high can cause very short rules or even empty rules.
		The default minimum support is 0.1.
KMNS_NUM_BINS	TO_CHAR(numeric_expr>0)	Number of bins in the attribute histogram produced by <i>k</i> -means. The bin boundaries for each attribute are computed globally on the entire training data set. The binning method is equi-width. All attributes have the same number of bins with the exception of attributes with a single value that have only one bin.
		The default number of histogram bins is 11.
KMNS_SPLIT_CRITERION	KMNS_SIZE KMNS_VARIANCE	Split criterion for <i>k</i> -means. The split criterion controls the initialization of new <i>k</i> -Means clusters. The algorithm builds a binary tree and adds one new cluster at a time.
		When the split criterion is based on size, the new cluster is placed in the area where the largest current cluster is located. When the split criterion is based on the variance, the new cluster is placed in the area of the most spread-out cluster.
		The default split criterion is the KMNS_VARIANCE.
KMNS_RANDOM_SEED	Non-negative integer	This setting controls the seed of the random generator used during the <i>k</i> -Means initialization. It must be a nonnegative integer value. The default is 0.



An exception is raised if Winsorize is enabled and other distance functions

are set.

Table 62-20 (Cont.) k-Means Settings

Setting Name	Setting Value	Description
KMNS_DETAILS	KMNS_DETAILS_NONE KMNS DETAILS HIERARCHY	This setting determines the level of cluster detail that are computed during the build.
	KMNS_DETAILS_ALL	KMNS_DETAILS_NONE: No cluster details are computed. Only the scoring information is persisted.
		KMNS_DETAILS_HIERARCHY: Cluster hierarchy and cluster record counts are computed. This is the default value.
		KMNS_DETAILS_ALL: Cluster hierarchy, record counts, descriptive statistics (means, variances, modes, histograms, and rules) are computed.
KMNS_WINSORIZE	KMNS_WINSORIZE_ENABLE KMNS_WINSORIZE_DISABLE	To winorize data, enable or disable this parameter. Data is restricted in a window size of six standard deviations around the mean value when winsorize is enabled. This functionality can be used with AUTO_DATA_PREP turned ON and OFF. The values outside the range are replaced with the ends of the interval. Winsorize is not enabled by default.
		Winsorize is only available when the KMNS_EUCLIDEAN distance function is used.

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.

See Also:

- For generic machine learning function settings related to Clustering, see DBMS_DATA_MINING — Machine Learning Functions.
- Oracle Machine Learning for SQL Concepts for information about k-Means



DBMS_DATA_MINING - Algorithm Settings: Multivariate State Estimation Technique - Sequential Probability Ratio Test

Settings that configure the training calibration behavior of the Multivariate State Estimation Technique - Sequential Probability Ratio Test algorithm.

The Constant Value column specifies constants using the prefix <code>DBMS_DATA_MINING</code>. For example, <code>DBMS_DATA_MINING.MSET_ADB_HEIGHT</code>. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the <code>DBMS_DATA_MINING</code> prefix, in single quotes. For example, <code>'MSET_ADB_HEIGHT'</code>.



Table 62-21 MSET-SPRT Settings

Setting Name	Setting Value	String Value Equivalent	Description
MSET_ADB_HEIGHT	A positive double	A positive double	Estimates the band within which signal values normally oscillate.
			The default value is 0.05.
MSET_ALERT_COUNT	A positive integer	A positive integer	The number of the last <i>n</i> signals (the alert window) that should have passed the threshold to raise an alert. The alert count should be lower or equal to the alert window.
			The default value is 5.
MSET_ALERT_WINDOW	A positive integer greater than or equal to		The number of signals to consider in the SPRT hypothesis consolidation logic.
	MSET_ALERT_COUNT	MSET_ALERT_COUNT	The default value is 5.
MSET_ALPHA_PROB	A positive double	A positive double	False Alarm Probability FAP (false positive).
	between 0 and 1	between 0 and 1	The default is 0.01.
MSET_BETA_PROB	A positive double	A positive double	Missed Alarm Probability MAP (false negative).
	between 0 and 1	between 0 and 1	The default is 0.10.
MSET_HELDASIDE	A positive integer	A positive integer	The approximate number of data rows used for MSET model calibration.
			You can use <code>ODMS_RANDOM_SEED</code> to change the held-aside sample.
			The default value is 10000.
MSET_MEMORY_VECTOR S	A positive integer	A positive integer	The default value is data driven.



Table 62-21 (Cont.) MSET-SPRT Settings

Setting Name	Setting Value	String Value Equivalent	Description
MSET_PROJECTION_TH RESHOLD	A positive integer >0, <=10000	A positive integer >0, <=10000	Specifies whether to use random projections. When the number of sensors exceeds the setting value, random projections are used. To turn off random projections, set the threshold to a value that is equal to or greater than the number of sensors.
			The default value is 500.
MSET_STD_TOLERANCE	A positive integer	A positive integer	The tolerance in standard deviations used in the SPRT calculation.
			The default value is 3.

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.

DBMS_DATA_MINING — Algorithm Settings: Naive Bayes

The settings listed in the following table configure the behavior of the Naive Bayes algorithm.

The Constant Value column specifies constants using the prefix <code>DBMS_DATA_MINING</code>. For example, <code>DBMS_DATA_MINING.NABS_PAIRWISE_THRESHOLD</code>. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the <code>DBMS_DATA_MINING</code> prefix, in single quotes. For example, 'NABS_PAIRWISE_THRESHOLD'.



Table 62-22 Naive Bayes Settings

Setting Name	Setting Value	String Value Equivalent	Description
NABS_PAIRWISE_THRE SHOLD	A floating point number between 0 and 1, inclusive, expressed as	between 0 and 1,	Value of pairwise threshold for NB algorithm Default is 0.
	a character string	a character string	Expression: TO_CHAR(0.5)



Table 62-22 (Cont.) Naive Bayes Settings

Setting Name	Setting Value	String Value Equivalent	Description
	.	A floating point number	Value of singleton threshold for NB algorithm
ESHOLD	between 0 and 1, inclusive, expressed as	between 0 and 1, inclusive, expressed as	Default value is 0.
	a character string a character string	Expression:	
		g	TO_CHAR(0.5)

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts for information about Naive Bayes

DBMS_DATA_MINING — Algorithm Settings: Neural Network

The settings listed in the following table configure the behavior of the Neural Network algorithm.

The Constant Value column specifies constants using the prefix DBMS_DATA_MINING. For example, DBMS_DATA_MINING.NNET_SOLVER_ADAM. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the DBMS_DATA_MINING prefix, in single quotes. For example, 'NNET_SOLVER_ADAM'.



Table 62-23 DBMS_DATA_MINING Neural Network Settings

Setting Name	Constant Value	String Value Equivalents	Description
NNET_SOLVER	One of the following	NNET_SOLVER_ADAM	Specifies the method of optimization.
	strings:		The default value is system determined.
	NNET_SOLVER_ADAM		NNET_SOLVER_ADAM: Uses the Adam optimization method.



Table 62-23 (Cont.) DBMS_DATA_MINING Neural Network Settings

Setting Name	Constant Value	String Value Equivalents	Description
	NNET_SOLVER_LBFGS	NNET_SOLVER_LBFGS	Uses the Limited-memory Broyden-Fletcher-Goldfarb-Shanno (LBFGS) optimization method.
NNET_ACTIVATIONS	One or more of the following strings: NNET_ACTIVATIONS_A RCTAN	NNET_ACTIVATIONS_A RCTAN	Specifies the activation functions for the hidden layers. You can specify a single activation function, which is then applied to each hidden layer, or you can specify an activation function for each layer individually. Different layers can have different activation functions.
			To apply a different activation function to one or more of the layers, you must specify an activation function for each layer. The number of activation functions you specify must be consistent with the NNET_HIDDEN_LAYERS and NNET_NODES_PER_LAYER values.
			For example, if you have three hidden layers, you could specify the use of the same activation function for all three layers with the following settings value:
			('NNET_ACTIVATIONS', 'NNET_ACTIVATIONS_TANH')
			The following settings value specifies a different activation function for each layer:
			('NNET_ACTIVATIONS', '''NNET_ACTIVATIONS_TANH'', ''NNET_ACTIVATIONS_LOG_SIG'', ''NNET_ACTIVATIONS_ARCTAN''')

Note:

You specify the different activation functions as strings within a single string. All quotes are single and two single quotes are used to escape a single quote in SQL statements and PL/SQL blocks.

 ${\tt NNET_ACTIVATIONS_ARCTAN:}$ Uses the arctangent activation function.

The default value is NNET_ACTIVATIONS_LOG_SIG.

Table 62-23 (Cont.) DBMS_DATA_MINING Neural Network Settings

Setting Name	Constant Value	String Value Equivalents	Description
	NNET_ACTIVATIONS_B IPOLAR_SIG	NNET_ACTIVATIONS_B IPOLAR_SIG	Uses the bipolar sigmoid activation function.
	NNET_ACTIVATIONS_L INEAR	NNET_ACTIVATIONS_L INEAR	Uses the linear activation function.
	NNET_ACTIVATIONS_L OG_SIG	NNET_ACTIVATIONS_L OG_SIG	Uses the logistic sigmoid activation function.
	NNET_ACTIVATIONS_R ELU	NNET_ACTIVATIONS_R ELU	Uses the rectified linear unit activation function.
	NNET_ACTIVATIONS_T ANH	NNET_ACTIVATIONS_T ANH	Uses the hyperbolic tangent activation function.
NNET_HELDASIDE_MAX _FAIL	A positive integer	A positive integer	With NNET_REGULARIZER_HELDASIDE, the training process is stopped early if the network performance on the validation data fails to improve or remains the same for NNET_HELDASIDE_MAX_FAIL epochs in a row.
			The default value is 6.
NNET_HELDASIDE_RAT IO	An integer greater than 0 and less than or equal to 1, represented as a character string	0 and less than or	Define the held ratio for the held-aside method. The default value is 0.25. Expression: TO CHAR (0.45)
NNET_HIDDEN_LAYERS	A positive integer	A positive integer	Defines the topology by the number of hidden layers.
			The default value is 1.
NNET_ITERATIONS	A positive integer	A positive integer	Specifies the maximum number of iterations in the Neural Network algorithm.
			For the <code>DMSSET_NN_SOLVER_LBFGS</code> solver, the default value is $\overline{200}$.
			For the <code>DMSSET_NN_SOLVER_ADAM</code> solver, the default value is 10000 .
NNET_NODES_PER_LAY ER	A positive integer or a list of positive integers	A positive integer or a list of positive integers	Defines the topology by the number of nodes per layer. Different layers can have different numbers of nodes.
			To specify the same number of nodes for each layer, you can provide a single value, which is then applied to each layer.
			To specify a different number of nodes for one or more layers, provide a list of comma-separated positive integers, one for each layer. For example, '10, 20, 5' for three layers. The setting values must be consistent with the NNET_HIDDEN_LAYERS value.
			The default number of nodes per layer is the number of attributes or 50 (if the number of attributes > 50).



Table 62-23 (Cont.) DBMS_DATA_MINING Neural Network Settings

Setting Name	Constant Value	String Value Equivalents	Description
NNET_REG_LAMBDA	An integer greater than or equal to 0 represented as a	An integer greater than or equal to 0 represented as a	Defines the L2 regularization parameter lambda. This can not be set together with NNET REGULARIZER HELDASIDE.
	character string	character string	The default value is 1.
			Expression:
			TO_CHAR(2)
NNET_REGULARIZER	One of the following strings:	NNET_REGULARIZER_H ELDASIDE	Regularization setting for Neural Network algorithm.
	NNET_REGULARIZER_H ELDASIDE		NNET_REGULARIZER_HELDASIDE: Uses a held-aside method for regularization. If the total number of training rows is greater than 50000, the default is NNET_REGULARIZER_HELDASIDE.
	NNET_REGULARIZER_L 2	NNET_REGULARIZER_L 2	Applies L2 regularization, which penalizes the sum of squared weights.
	NNET_REGULARIZER_N ONE	NNET_REGULARIZER_N ONE	Disables regularization. If the total number of training rows is less than or equal to 50000, the default is NNET_REGULARIZER_NONE.
NNET_TOLERANCE	A floating point number between 0 and 1 expressed as a character string	A floating point number between 0 and 1 expressed as a character string	Defines the convergence tolerance setting of the Neural Network algorithm.
			The default value is 0.000001.
			Expression:
			TO_CHAR(0.00004)
NNET_WEIGHT_LOWER_ BOUND	A real number	A real number	The setting specifies the lower bound of the region where weights are randomly initialized. NNET_WEIGHT_LOWER_BOUND and NNET_WEIGHT_UPPER_BOUND must be set together. Setting one and not setting the other raises an error. NNET_WEIGHT_LOWER_BOUND must not be greater than NNET_WEIGHT_UPPER_BOUND. The default value is -sqrt(6/(1_nodes+r_nodes)). The value of 1_nodes for: input layer dense attributes is (1+number of dense attributes) input layer sparse attributes is number of sparse attributes each hidden layer is (1+number of nodes in that hidden layer) The value of r_nodes is the number of nodes in the layer that the weight is connecting to.



Table 62-23	(Cont.) DBM	S_DATA	_MINING Neural	Network Settings
--------------------	-------------	----------	----------------	-------------------------

Setting Name	Constant Value	String Value Equivalents	Description
NNET_WEIGHT_UPPER_ BOUND	A real number	A real number	This setting specifies the upper bound of the region where weights are initialized. It should be set in pairs with NNET_WEIGHT_LOWER_BOUND and its value must not be smaller than the value of NNET_WEIGHT_LOWER_BOUND. If not specified, the values of NNET_WEIGHT_LOWER_BOUND and NNET_WEIGHT_UPPER_BOUND are system determined.
			The default value is sqrt(6/(1_nodes+r_nodes)). See
			NNET_WEIGHT_LOWER_BOUND.

- DBMS_DATA_MINING Machine Learning Functions
 - A machine learning **function** refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 - The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.
- DBMS_DATA_MINING Solver Settings: LBFGS
 The settings listed in the following table configure the behavior of L-BFGS. Neural Network and Generalized Linear Model (GLM) use these settings.



Oracle Machine Learning for SQL Concepts for information about Neural Network.

DBMS_DATA_MINING — Algorithm Settings: Non-Negative Matrix Factorization

The settings listed in the following table configure the behavior of the Non-negative Matrix Factorization algorithm.

The Constant Value column specifies constants using the prefix DBMS_DATA_MINING. For example, DBMS_DATA_MINING.NMFS_NONNEG_SCORING_ENABLE. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the DBMS_DATA_MINING prefix, in single quotes. For example, 'NMFS_NONNEG_SCORING_ENABLE'.



You can query the data dictionary view *_MINING_MODEL_SETTINGS (using the ALL, USER, or DBA prefix) to find the setting values for a model. See *Oracle Database Reference* for information about *_MINING_MODEL_SETTINGS.

Table 62-24 NMF Settings

Setting Name	Constant Value	String Value Equivalent	Description
NMFS_CONV_TOLERANC E	A floating point number between 0 and 0.5 expressed as a character string	A floating point number between 0 and 0.5 expressed as a character string	Convergence tolerance for NMF algorithm Default is 0.05 Expression: TO_CHAR(0.02)
NMFS_NONNEGATIVE_S CORING	NMFS_NONNEG_SCORIN G_ENABLE	NMFS_NONNEG_SCORIN G_ENABLE	Whether negative numbers should be allowed in scoring results. When set to NMFS_NONNEG_SCORING_ENABLE, negative feature values will be replaced with zeros. Default is NMFS_NONNEG_SCORING_ENABLE
	NMFS_NONNEG_SCORIN G_DISABLE	NMFS_NONNEG_SCORIN G_DISABLE	When set to NMFS_NONNEG_SCORING_DISABLE, negative feature values will be allowed.
NMFS_NUM_ITERATION S	An integer between 1 to 500, inclusive, represented as a character string	An integer between 1 to 500, inclusive, represented as a character string	Number of iterations for NMF algorithm Default is 50 Expression: TO_CHAR(80)
NMFS_RANDOM_SEED	An integer represented as a character string	An integer represented as a character string	Random seed for NMF algorithm. Default is -1. Expression: TO_CHAR(2)

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts for information about NMF

DBMS_DATA_MINING — Algorithm Settings: O-Cluster

The settings in the table configure the behavior of the O-Cluster algorithm.

The Constant Value column specifies constants using the prefix DBMS_DATA_MINING. For example, DBMS DATA MINING.OCLT SENSITIVITY. Alternatively, you can specify the

corresponding string value from the **String Value Equivalent** column without the DBMS DATA MINING prefix, in single quotes. For example, 'OCLT SENSITIVITY'.



The distinction between **Constant Value** and **String Value Equivalent** for this algorithm is applicable to Oracle Database 19c and Oracle Database 21c.

Table 62-25 O-CLuster Settings

Setting Name	Constant Value	String Value Equivalent	Description
OCLT_SENSITIVITY	A floating point number between 0 and 1 expressed as a character string	A floating point number between 0 and 1 expressed as a character string	A fraction that specifies the peak density required for separating a new cluster. The fraction is related to the global uniform density. Default is 0.5. Example: TO_CHAR(0.9)

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts for information about O-Cluster

DBMS_DATA_MINING — Algorithm Settings: Random Forest

These settings configure the behavior of the Random Forest algorithm. Random Forest makes use of the Decision Tree settings to configure the construction of individual trees.

The Constant Value column specifies constants using the prefix DBMS_DATA_MINING. For example, DBMS_DATA_MINING.RFOR_MTRY. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the DBMS_DATA_MINING prefix, in single quotes. For example, 'RFOR MTRY'.



Table 62-26 Random Forest Settings

Setting Name	Constant Value	String Value Equivalent	Description
RFOR_MTRY	a number >= 0	a number >= 0	Size of the random subset of columns to be considered when choosing a split at a node. For each node, the size of the pool remains the same, but the specific candidate columns change. The default is half of the columns in the model signature. The special value 0 indicates that the candidate pool includes all columns.
RFOR_NUM_TREES	1<= a number <=65535	1<= a number <=65535	Number of trees in the forest Default is 20.
RFOR_SAMPLING_RATI	<pre>0< a fraction<=1</pre>	0< a fraction<=1	Fraction of the training data to be randomly sampled for use in the construction of an individual tree. The default is half of the number of rows in the training data.

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.
- DBMS_DATA_MINING Algorithm Settings: Decision Tree
 These settings configure the behavior of the Decision Tree algorithm. Note that the
 Decision Tree settings are also used to configure the behavior of Random Forest as it constructs each individual decision tree.



Oracle Machine Learning for SQL Concepts for information about Random Forest

DBMS_DATA_MINING — Algorithm Constants and Settings: Singular Value Decomposition

The following settings configure the behavior of the Singular Value Decomposition algorithm.



Table 62-27 Singular Value Decomposition Settings

Setting Name	Constant Value	String Value Equivalent	Description
SVDS_U_MATRIX_OUTP UT	SVDS_U_MATRIX_ENAB	SVDS_U_MATRIX_ENAB	Indicates whether or not to persist the U Matrix produced by SVD.
			The U matrix in SVD has as many rows as the number of rows in the build data. To avoid creating a large model, the U matrix is persisted only when SVDS_U_MATRIX_OUTPUT is enabled.
			When SVDS_U_MATRIX_OUTPUT is enabled, the build data must include a case ID. If no case ID is present and the U matrix is requested, then an exception is raised.
			Default is SVDS_U_MATRIX_DISABLE.
	SVDS_U_MATRIX_DISA BLE	SVDS_U_MATRIX_DISA BLE	Does not persist the U Matrix.
SVDS_SCORING_MODE	SVDS_SCORING_SVD	SVDS_SCORING_SVD	Whether to use SVD or PCA scoring for the model.
			When the build data is scored with SVD, the projections will be the same as the U matrix.
			Default is SVDS_SCORING_SVD.
	SVDS_SCORING_PCA	SVDS_SCORING_PCA	When the build data is scored with PCA, the projections will be the product of the U and S matrices.



Table 62-27 (Cont.) Singular Value Decomposition Settings

Setting Name	Constant Value	String Value Equivalent	Description	
SVDS_SOLVER	SVDS_SOLVER_TSSVD	SVDS_SOLVER_TSSVD	This setting indicates the solver to be used for computing SVD of the data. In the case of PCA the solver setting indicates the type of SVD solver used to compute the PCA for the data. When this setting is not specified the solver type selection is data driven. If the number of attributes is greater than 3240, then the default wide solver is used. Otherwise, the default narrow solver is selected.	
			The following are the group of solvers:	
			 Narrow data solvers: for matrices with up to 11500 attributes (TSEIGEN) or up to 8100 attributes (TSSVD). 	
			 Wide data solvers: for matrices up to 1 million attributes. 	
			For narrow data solvers:	
			 Tall-Skinny SVD uses QR computation TSVD (SVDS_SOLVER_TSSVD) 	
			 Tall-Skinny SVD uses eigenvalue computation, TSEIGEN (SVDS SOLVER TSEIGEN), is the default 	
			solver for narrow data.	
			For wide data solvers:	
			 Stochastic SVD uses QR computation SSVD (SVDS_SOLVER_SSVD), is the default 	
			solver for wide data solvers.	
			 Stochastic SVD uses eigenvalue computations, STEIGEN (SVDS_SOLVER_STEIGEN). 	
	SVDS_SOLVER_TSEIGE N	SVDS_SOLVER_TSEIGE N	Tall-Skinny SVD using eigenvalue computation for matrices with up to 11500 attributes. This is the default solver for narrow data.	
	SVDS_SOLVER_SSVD	SVDS_SOLVER_SSVD	Stochastic SVD using QR computation for matrices with up to 1 million attributes. This is the default solver for wide data.	
	SVDS_SOLVER_STEIGE	SVDS_SOLVER_STEIGE	Stochastic SVD using eigenvalue computations for matrices with up to 1 million attributes.	
SVDS_TOLERANCE	Range [0, 1]	Range [0, 1]	This setting is used to prune features. Define the minimum value the eigenvalue of a feature as a share of the first eigenvalue to not to prune. Default value is data driven.	
SVDS_RANDOM_SEED	Range [0 - 4,294,967,296]	Range [0 - 4,294,967,296]	The random seed value is used for initializing the sampling matrix used by the Stochastic SVD solver. The default is 0. The SVD Solver must be set to SSVD or STEIGEN.	



Table 62-27 (Cont.) Singular Value Decomposition Settings	Table 62-27	(Cont.)	Singular	Value Decom	position	Settings
---	--------------------	---------	----------	-------------	----------	----------

Setting Name	Constant Value	String Value Equivalent	Description
SVDS_OVER_SAMPLING	Range [1, 5000].	Range [1, 5000].	This setting is configures the number of columns in the sampling matrix used by the Stochastic SVD solver. The number of columns in this matrix is equal to the requested number of features plus the oversampling setting. The SVD Solver must be set to SSVD or STEIGEN.
SVDS_POWER_ITERATI ONS	Range [0, 20].	Range [0, 20].	The power iteration setting improves the accuracy of the SSVD solver. The default is 2. The SVD Solver must be set to SSVD or STEIGEN.

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts

DBMS_DATA_MINING — Algorithm Settings: Support Vector Machine

The settings listed in the following table configure the behavior of the Support Vector Machine algorithm.

The Constant Value column specifies constants using the prefix DBMS_DATA_MINING. For example, DBMS_DATA_MINING.SVMS_GAUSSIAN. Alternatively, you can specify the corresponding string value from the String Value Equivalent column without the DBMS_DATA_MINING prefix, in single quotes. For example, 'SVMS_GAUSSIAN'.





Table 62-28 SVM Settings

Setting Name	Constant Value	String Value Equivalent	Description
SVMS_COMPLEXITY_FA CTOR	An integer greater than 0 represented as a character string	An integer greater than 0 represented as a character string	Regularization setting that balances the complexity of the model against model robustness to achieve good generalization on new data. SVM uses a data-driven approach to finding the complexity factor.
			Value of complexity factor for SVM algorithm (both classification and regression).
			Default value estimated from the data by the algorithm.
			Expression:
			TO_CHAR(20)
SVMS_CONV_TOLERANC	An integer greater than	An integer greater than	Convergence tolerance for SVM algorithm.
E	0 represented as a character string	0 represented as a character string	Default is 0.0001.
	character string	Character String	Expression:
			TO_CHAR(0.005)
SVMS_EPSILON	An integer greater than 0 represented as a character string	An integer greater than 0 represented as a character string	Regularization setting for regression, similar to complexity factor. Epsilon specifies the allowable residuals, or noise, in the data. Value of epsilon factor for SVM regression.
			Default is 0.1.
			Expression: TO_CHAR(0.5)
SVMS_KERNEL_FUNCTI	SVMS_GAUSSIAN	SVMS_GAUSSIAN	Kernel for Support Vector Machine. Linear or Gaussian.
			SVMS_GAUSSIAN: Uses the Gaussian kernel for SVM.
			The default value is SVMS_LINEAR.
	SVMS_LINEAR	SVMS_LINEAR	Uses the Linear kernel for SVM. This is the default option.
SVMS_OUTLIER_RATE	A floating point number between 0 and 1 expressed as a	A floating point number between 0 and 1 expressed as a	The desired rate of outliers in the training data. Valid for One-Class SVM models only (anomaly detection).
	character string	character string	Default is 0.01.
			Expression:
			TO_CHAR(0.04)
SVMS_STD_DEV	An integer greater than 0 represented as a character string	An integer greater than 0 represented as a character string	Controls the spread of the Gaussian kernel function. SVM uses a data-driven approach to find a standard deviation value that is on the same scale as distances between typical cases.
			Value of standard deviation for SVM algorithm.
			This is applicable only for Gaussian kernel.
			Default value estimated from the data by the algorithm.
			Expression:
			TO_CHAR(6)



Table 62-28 (Cont.) SVM Settings

Setting Name	Constant Value	String Value Equivalent	Description
SVMS_NUM_ITERATION S	A positive integer	A positive integer	This setting sets an upper limit on the number of SVM iterations. The default is system determined because it depends on the SVM solver.
SVMS_NUM_PIVOTS	Range [1; 10000]	Range [1; 10000]	This setting sets an upper limit on the number of pivots used in the Incomplete Cholesky decomposition. It can be set only for non-linear kernels. The default value is 200.
SVMS_BATCH_ROWS	A positive integer	A positive integer	This setting applies to SVM models with linear kernel. This setting sets the size of the batch for the SGD solver. An input of 0 triggers a data driven batch size estimate. The default is 20000.
SVMS_REGULARIZER	SVMS_REGULARIZER_L 1	SVMS_REGULARIZER_L 1	This setting controls the type of regularization that the SGD SVM solver uses. The setting can be used only for linear SVM models. The default is system determined because it depends on the potential model size. SVMS REGULARIZER L1: Uses L1
			regularization.
	SVMS_REGULARIZER_L 2	SVMS_REGULARIZER_L 2	Uses L2 regularization.
SVMS_SOLVER	SVMS_SOLVER_SGD (Sub-Gradient Descend)	SVMS_SOLVER_SGD (Sub-Gradient Descend)	Enables to choose the SVM solver. The SGD solver cannot be selected if the kernel is nonlinear. The default value is system determined. SVMS_SOLVER_SGD: Uses Sub-Gradient Descent solver.
	SVMS_SOLVER_IPM (Interior Point Method)	SVMS_SOLVER_IPM (Interior Point Method)	Uses Interior Point Method solver.

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



Oracle Machine Learning for SQL Concepts for information about SVM



DBMS_DATA_MINING — Algorithm Settings: XGBoost

Settings that configure the behavior of the XGBoost gradient boosting algorithm.

The Constant Name column specifies constants using the prefix <code>DBMS_DATA_MINING</code>. For example, <code>DBMS_DATA_MINING</code>. <code>xgboost_booster</code>. Alternatively, you can specify the corresponding string value from the String Name Equivalent column without the <code>DBMS_DATA_MINING</code> prefix, in single quotes. For example, 'booster'.



The distinction between **Constant Value** and **String Value Equivalent** for this algorithm is applicable to Oracle Database 19c and Oracle Database 21c.

The XGBoost settings are case sensitive. Enter the settings as they appear in the settings table. These settings match the XGBoost settings available in open source. OML4SQL XGBoost is based on the 1.7.4 version of XGBoost.

For Global settings, see DBMS_DATA_MINING — Global Settings.

For generic machine learning technique settings, see DBMS_DATA_MINING — Machine Learning Functions.

Table 62-29 General Settings

Constant Name	String Name Equivalent	Setting Value	Description	
xgboost_booster	booster	A string that is one of the following: dart gblinear gbtree	The booster to use: dart gblinear gbtree The dart and gbtree boosters use tree-based models whereas gblinear uses linear functions. The default value is gbtree.	
xgboost_num_round	num_round	A non-negative integer.	The number of rounds for boosting. The default value is 10.	

Table 62-30 Settings for Tree Boosting

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_alpha	alpha	A non-negative number	L1 regularization term on weights. Increasing this value makes the model more conservative. The default value is 0.
xgboost_colsample_ bylevel	colsample_bylevel	A number in the range (0, 1]	Subsample ratio of columns for each split, in each level. Subsampling occurs each time a new split is made. This parameter has no effect when tree_method is set to hist. The default value is 1.



Table 62-30 (Cont.) Settings for Tree Boosting

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_colsample_ bynode	colsample_bynode	A number in the range (0, 1]	The subsample ratio of columns for each node (split). Subsampling occurs once every time a new split is evaluated. Columns are subsampled from the set of columns chosen for the current level.
			The default value is 1.
<pre>xgboost_colsample_ bytree</pre>	colsample_bytree	A number in the range (0, 1]	Subsample ratio of columns when constructing each tree. Subsampling occurs once in every boosting iteration.
			The default value is 1.
xgboost_eta	eta	A number in the range [0, 1]	Step-size shrinkage used in the update step to prevent overfitting. After each boosting step, eta shrinks the feature weights to make the boosting process more conservative.
			The default value is 0.3.
xgboost_gamma	gamma	A number in the range $[0, \infty]$	Minimum loss reduction required to make a further partition on a leaf node of the tree. The larger gamma value is, the more conservative the algorithm is.
			The default value is 0.
xgboost_grow_polic Y	grow_policy	A string; one of the following:	Controls the way new nodes are added to the tree:
		depthwiselossguide	 depthwise splits at nodes closest to the root
			• lossguide splits at nodes with the highest loss change Valid only if tree method is set to hist.
			The default value is depthwise.
xgboost_interactio n_constraints	interaction_constr aints	[[x0,x1,x2], [x0,x4],[x5,x6]] where xn are feature names or columns	This setting specifies permitted interactions in the model. Specify the constrains in the form of a nested list where each inner list is a group of features (column names) that are allowed to interact with each other. If a single column is passed in the interactions then, the input is ignored.
			Here, features x0, x1, and x2 are allowed to interact with each other but with no other feature Similarly, x0 and x4 are allowed to interact with each other but with no other feature and so on. This setting is applicable to 2-Dimensional features. An error occurs if you pass columns of non-supported type and non-existing feature names.
xgboost_lambda	lambda	A non-negative number	L2 regularization term on weights.
_		-	The default value is 1.



Table 62-30 (Cont.) Settings for Tree Boosting

Constant Name	String Name Equivalent	Setting Value	Description	
xgboost_max_bin	max_bin	A non-negative integer	Maximum number of discrete bins to bucket continuous features. Increasing this number improves the optimality of splits at the cost of higher computation time.	
			This parameter is v is set to hist.	valid only when tree_method
			The default value is	s 256.
xgboost_max_delta_ step	max_delta_step	A number in the range $[0, \infty]$	Maximum delta ste output.	p allowed for each leaf
			Setting this to a positive value can help make the update step more conservative. Usually parameter is not needed, but it might help in logistic regression when the class is extreme imbalanced. Setting it to value from 1 to 10 might help control the update. The default value is 0, which means there is constraint.	
xgboost_max_depth	max_depth	An integer in the range $[0, \infty]$	nge Maximum depth of a tree. Increasing makes the model more complex and to overfit.	
			Setting this to 0 ind	licates no limit.
				You must set a max_depth limit when the grow_policy setting is depthwise.
			The default value is 6.	
xgboost_max_leaves	max_leaves	A non-negative number	r Maximum number of nodes to add. Use this setting only when grow_policy is set to lossguide.	
1		A mount and the	The default value is	
<pre>xgboost_min_child_ weight</pre>	min_child_weight	A number in the range [0, ∞]	Minimum sum of instance weight (hessian) needed in a child. If the tree partition step results in a leaf node with a sum of instance weight less than min_child_weight, then the building process stops partitioning. In a linear regression task, this corresponds to the minimum number of instances needed in each node. The larger min_child_weight is, the more conservative the algorithm is. The default value is 1.	



Table 62-30 (Cont.) Settings for Tree Boosting

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_decrease_c onstraints	num_parallel_tree	[x4,x5]	This setting specifies the features (column names) that must obey decreasing constraint. The feature names are separated by a comma. For example, setting value 'x4,x5' sets decreasing constraint on features x4 and x5. This setting applies to numeric columns and 2-Dimensional features. An error occurs if you pass columns of non-supported type and non-existing feature names.
<pre>xgboost_increase_c onstraints</pre>	<pre>increase_constrain ts</pre>	[x0,x3]	This setting specifies the features (column names) that must obey increasing constraint. The feature names are separated by a comma. For example, setting value 'x0,x3' sets increasing constraint on features x0 and x3. This setting is applicable to 2-Dimensional features. An error occurs if you pass columns of nonsupported type and non-existing feature names.
<pre>xgboost_num_parall el_tree</pre>	num_parallel_tree	A non-negative integer	Number of parallel trees constructed during each iteration. Use this option to support a boosted random forest. The default value is 1.
xgboost_scale_pos_ weight	scale_pos_weight	A non-negative number	Controls the balance of positive and negative weights, which is useful for unbalanced classes. A typical value to consider: sum (negative cases) / sum (positive cases). The default value is 1.
xgboost_sketch_eps	sketch_eps	A number in the range (0, 1)	Increases enumeration accuracy. Valid only for the approximate greedy tree method.
			Compared to directly selecting the number of bins, this setting comes with a theoretical guarantee with sketch accuracy. You usually do not need to change this setting, but you might consider setting a lower number for more accurate enumeration.
			The default value is 0.03.
xgboost_subsample	subsample	A number in the range (0, 1]	Subsample ratio of the training instances. A setting of 0.5 means that XGBoost randomly samples half of the training data prior to growing trees, which prevents overfitting. Subsampling occurs once in every boosting iteration. The default value is 1.



Table 62-30 (Cont.) Settings for Tree Boosting

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_tree_metho d	tree_method	A string that is one of the following: approx auto exact hist	 Tree construction algorithm used in XGBoost: approx: Approximate greedy algorithm using sketching and histogram. auto: Use a heuristic to choose the faster algorithm:
xgboost_updater	updater	A comma-separated string; one or more of the following: grow_colmaker grow_histmaker grow_skmaker grow_quantile_histmaker prune sync	Defines the sequence of tree updaters to run, which provides a modular way to construct and to modify the trees. This is an advanced parameter that is usually set automatically, depending on some other parameters. However, you can also explicitly specify a settting. The setting values are: grow_colmaker: Non-distributed column-based construction of trees. grow_histmaker: Distributed tree construction with row-based data splitting based on a global proposal of histogram counting. grow_skmaker: Uses the approximate sketching algorithm. grow_quantile_histmaker: Grow tree using quantized histogram. prune: Prunes the splits where loss < min_split_loss (or gamma). sync: Synchronizes trees in all distributed nodes.

Table 62-31 Settings for the Dart Booster

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_one_drop	one_drop	A number that is 0 or 1	When set to 1, at least one tree is always dropped during the dropout. When set to 0, at least one tree is not always dropped during the dropout. The default value is 0.



Table 62-31 (Cont.) Settings for the Dart Booster

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_normalize_ type	normalize_type	A string; either: • forest • tree	 Type of normalization algorithm: forest: New trees have the same weight as the sum of the dropped trees (forest): The weight of new trees is 1 / (1 + learning_rate) Dropped trees are scaled by a factor of 1 / (1 + learning_rate) tree: New trees have the same weight as dropped trees: The weight of new trees is 1 / (k + learning_rate) Dropped trees are scaled by a factor of k / (k + learning_rate) The default value is tree.
xgboost_rate_drop	rate_drop	A number in the range [0.0, 1.0]	Dropout rate (a fraction of the previous trees to drop during the dropout). The default value is 0.0.
xgboost_sample_typ e	sample_type	A string; either: uniform weighted	Type of sampling algorithm: uniform: Dropped trees are selected uniformly weighted: Dropped trees are selected in proportion to weight The default value is uniform.
xgboost_skip_drop	skip_drop	A number in the range [0.0, 1.0]	Probability of skipping the dropout procedure during a boosting iteration. If a dropout is skipped, new trees are added in the same manner as gbtree. A non-zero skip_drop has higher priority than rate_drop or one_drop. The default value is 0.0.

Table 62-32 Settings for the Linear Booster

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_alpha	alpha	A non-negative number	L1 regularization term on weights, normalized to the number of training examples. Increasing this value makes the model more conservative. The default value is 0.



Table 62-32 (Cont.) Settings for the Linear Booster

Constant Name	String Name Equivalent	Setting Value	Description
xgboost_feature_se lector	feature_selector	A string that is one of the following: cyclic greedy random shuffle thrifty	 Feature selection and ordering method: cyclic: Deterministic selection by cycling through the features one at a time. greedy: Selects the coordinate with the greatest gradient magnitude. This method: Has 0 (num_feature^2) complexity Is fully deterministic Allows restricting the selection to the top_k features per group with the largest magnitude of univariate weight change, by setting the top_k parameter; doing so reduces the complexity to 0 (num_feature*top_k). random: A random (with replacement) coordinate selector. shuffle: Similar to cyclic but with random feature shuffling prior to each update. thrifty: Thrifty, approximately-greedy feature selector. Prior to cyclic updates, reorders features in descending magnitude of their univariate weight changes. This operation is multithreaded and is a linear complexity approximation of the quadratic greedy selection. Restricts the selection per group to the top_k features with the largest magnitude of univariate weight change. The default value is cyclic.
xgboost_lambda	lambda	A non-negative number	L2 regularization term on weights, normalized to the number of training examples. Increasing this value makes the model more conservative. The default value is 0.
xgboost_top_k	top_k	A non-negative integer	Number of top features to select for the <code>greedy</code> or <code>thrifty</code> feature selector. The value of 0 uses all of the features. The default value is 0.
xgboost_updater	updater	A string that is one of the following: coord_descent shotgun	Algorithm to fit the linear model: coord_descent: Ordinary coordinate descent algorithm; multithreaded but still produces a deterministic solution. shotgun: Parallel coordinate descent algorithm based on the shotgun algorithm; uses "hogwild" parallelism and therefore produces a nondeterministic solution on each run. The default value is shotgun.



Table 62-33 Settings for Tweedie Regression

Constant Name	String Name Equivalent	Setting Value	Description
<pre>xgboost_tweedie_va riance_power</pre>	<pre>tweedie_variance_p ower</pre>	A number in the range (1, 2)	Controls the variance of the Tweedie distribution var(y) ~ E(y)^tweedie_variance_power.
			A setting closer to 1 shifts towards a Poisson distribution.
			A setting closer to 2 shifts towards a gamma distribution.
			The default value is 1.5.

Some XGBoost objectives apply only to classification function models and other objectives apply only to regression function models. If you specify an incompatible objective value, an error is raised. In the DBMS_DATA_MINING.CREATE_MODEL procedure, if you specify DBMS_DATA_MINING.CLASSIFICATION as the function, then the only objective values that you can use are the binary and multi values. The one exception is binary: logitraw, which produces a continuous value and applies only to a regression model. If you specify DBMS_DATA_MINING.REGRESSION as the function, then you can specify binary: logitraw or any of the count, rank, reg, and survival values as the objective.



Table 62-34 Settings for Learning Tasks

Setting Name	String Name Equivalent	Setting Value	Description
xgboost_objective	objective	For a classification model, a string that is one of the following: binary:hinge binary:logistice multi:softmax multi:softmax multi:softprob For a regression model, a string that is one of the following: binary:logitraw count:poisson rank:map rank:ndcg rank:pairwise reg:gamma reg:logistic reg:tweedie survival:aft survival:cox reg:squarederror reg:squaredlogerror	 Settings for a Classification model: binary:hinge: Hinge loss for binary classification. This setting makes predictions of 0 or 1, rather than producing probabilities. binary:logistic: Logistic regression for binary classification. The output is the probability. multi:softmax: Performs multiclass classification using the softmax objective; you must also set num_class (number_of_classes). multi:softprob::Same as softmax, except the output is a vector of ndata * nclass, which can be further reshaped to an ndata * nclass matrix. The result contains the predicted probability of each data point belonging to each class. The default objective value for classification i multi:softprob. Settings for a Regression model: binary:logitraw: Logistic regression for binary classification; the output is the score before logistic transformation. count:poisson: Poisson regression for count data; the output is the mean of the Poisson distribution. The max_delta_ster value is set to 0.7 by default in Poisson regression to safeguard optimization. rank:map: Using LambdaMART, performs list-wise ranking in which the Mean Average Precision (MAP) is maximized. rank:ndcg: Using LambdaMART, performs list-wise ranking in which the Normalized Discounted Cumulative Gain (NDCG) is maximized. rank:pairwise: Performs ranking by minimizing the pairwise loss. reg:gamma: Gamma regression with log-link; the output is the mean of the gamma distribution. This setting might be useful for any outcome that might be gamma-distributed, such as modeling insurance claims severity. reg:logistic: Logistic regression. reg:tweedie: Tweedie regression with log-link. This setting might be useful for any outcome that might be Tweedie-distributed, such as modeling total loss in insurance. survival:aft: Applies the Accelerated Failure Time (AFT) model for censore



Table 62-34 (Cont.) Settings for Learning Tasks

Setting Name	String Name Equivalent	Setting Value	Description
			 option, eval_metric uses aft-nloglik as the default value. survival:cox: Cox regression for right-censored survival time data (negative values are considered right-censored). Predictions are returned on the hazard ratio scale (that is, as HR = exp(marginal_prediction) in the proportional hazard function h(t) = h0(t) * HR). reg:squarederror: Regression with squared loss. reg:squaredlogerror: Regression with squared log loss. All input labels must be greater than -1. The default objective value for regression is reg:squarederror.
xgboost_aft_loss_d istribution	aft_loss_distribut ion	[normal, logistic, extreme]	Specifies the distribution of the Z term in the AFT model. It specifies the Probabilty Density Function used by survival:aft objective and aft-nloglik evaluation metric. The default value is normal.
xgboost_aft_loss_d istribution_scale	aft_loss_distribut ion_scale	A positive number	Specifies the scaling factor σ , which scales the size of Z term in the AFT model. The default value is 1.
xgboost_aft_right_ bound_column_name	aft_right_bound_co lumn_name	column_name	Specifies the column containing the right bounds of the labels for an AFT model. You cannot select this parameter for a non-AFT model. Note: Oracle Machine Learning does not support BOOLEAN values for this setting.
xgboost_base_score	base_score	A number	Initial prediction score of all instances, global bias. For a sufficient number of iterations, changing this value does not have much effect. The default value is 0.5.



Table 62-34 (Cont.) Settings for Learning Tasks

Setting Name String Name Equivalent	Setting Value	Description
	A comma-separated string; one or more of the following: aft-nloglik auc aucpr cox-nloglik error error@t gamma-deviance gamma-nloglik logloss mae map map@n merror mlogloss ndcg ndcg@n poisson- nloglik rmse tweedie- nloglik@rho ndcg- map- rmsle	Evaluation metrics for validation data. You can specify one or more of these evaluation metrics aft-nloglik: Sets the eval_metric to negative log likelihood of AFT model. auc: Area under the curve. aucpr: Area under the PR curve. cox-nloglik: Negative partial log-likelihood for Cox proportional hazards regression. error: Binary classification error rate, calculated as the number of wrong cases divided by the number of all cases. For the predictions, the evaluation regards the instances with a prediction value larger tha 0.5 as positive instances, and the others as negative instances. error@t: You can specify a binary classification threshold value other than 0.5 by specifying a numerical value t; for example, error@0.8. gamma-deviance: Residual deviance for gamma regression. logloss: Negative log-likelihood. mae: Mean absolute error. map@m: Assigns the integer n as the cut-off value for the top positions in the lists for evaluation. merror: Multiclass classification error rate calculated as the number of wrong cases divided by the number of all cases; the objective must be multi:softprob or multi:softmax. mlogloss: Multiclass logloss; the objective must be multi:softprob or multi:softmax. ndcg: Normalized Discounted Cumulative Gain. ndcg@n: Assigns the integer n as the cut-ovalue for the top positions in the lists for evaluation. poisson-nloglik: Negative log-likelihood for Poisson regression rmse: Root Mean Square Error.



Table 62-34 (Cont.) Settings for Learning Tasks

Setting Name	String Name Equivalent	Setting Value	Description
			rho must be a number in the range (1, 2); for example, tweedie-nloglik@1.8. ndcg- and map-: In XGBoost, NDCG and MAP will evaluate the score of a list without any positive samples as 1. By adding "-" in the evaluation metric XGBoost will evaluate these score as 0 to be consistent under some conditions. rmsle: It is root mean square log error. This is the default metric of reg:squaredlogerror objective. This metric reduces errors generated by outliers in dataset. But because log function is employed, rmsle might output nan when prediction value is less than -1. A default metric is assigned according to the objective: error for classification mean average precision for ranking rmse for regression
xgboost_seed	seed	A non-negative integer	Random number seed. The default value is 0.

Related Topics

- DBMS_DATA_MINING Machine Learning Functions
 A machine learning function refers to the methods for solving a given class of machine learning problems.
- DBMS_DATA_MINING Global Settings
 The configuration settings in this table are applicable to any type of model, but are currently only implemented for specific algorithms.



https://github.com/oracle/oracle-db-examples/tree/master/machine-learning/sql/, select the release, and browse for an example of XGBoost.

DBMS_DATA_MINING — Solver Settings

Oracle Machine Learning for SQL algorithms can use different solvers. Solver settings can be provided at build time in the settings table.

Related Topics

DBMS_DATA_MINING - Solver Settings: Adam
 These settings configure the behavior of the Adaptive Moment Estimation (Adam) solver.

DBMS_DATA_MINING — Solver Settings: ADMM

The settings listed in the following table configure the behavior of Alternating Direction Method of Multipliers (ADMM). The Generalized Linear Model (GLM) algorithm uses these settings.

DBMS_DATA_MINING — Solver Settings: LBFGS
 The settings listed in the following table configure the behavior of L-BFGS. Neural Network and Generalized Linear Model (GLM) use these settings.

DBMS DATA MINING - Solver Settings: Adam

These settings configure the behavior of the Adaptive Moment Estimation (Adam) solver.

Neural Network models use these settings.

Table 62-35 DBMS_DATA_MINING Adam Settings

Setting Name	Setting Value	Description
ADAM_ALPHA	A non-negative double precision floating point number in the interval (0; 1]	The learning rate for Adam. The default value is 0.001.
ADAM_BATCH_ROWS	A positive integer	The number of rows per batch. The default value is 10000.
ADAM_BETA1	A positive double precision floating point number in the interval [0; 1)	The exponential decay rate for the 1st moment estimates. The default value is 0.9.
ADAM_BETA2	A positive double precision floating point number in the interval [0; 1)	The exponential decay rate for the 2nd moment estimates. The default value is 0.99.
ADAM_GRADIENT_TOLERANCE	A positive double precision floating point number	The gradient infinity norm tolerance for Adam. The default value is 1E-9.

Related Topics

DBMS_DATA_MINING — Algorithm Settings: Neural Network
 The settings listed in the following table configure the behavior of the Neural Network algorithm.

DBMS DATA MINING — Solver Settings: ADMM

The settings listed in the following table configure the behavior of Alternating Direction Method of Multipliers (ADMM). The Generalized Linear Model (GLM) algorithm uses these settings.

Table 62-36 DBMS_DATA_MINING ADMM Settings

Settings Name	Setting Value	Description
ADMM_CONSENSUS	A positive integer	It is a ADMM's consensus parameter. The value must be a positive number. The default value is 0.1.



Table 62-36 (Cont.) DBMS_DATA_MINING ADMM Settings

Settings Name	Setting Value	Description
ADMM_ITERATIONS	A positive integer	The number of ADMM iterations. The value must be a positive integer. The default value is 50.
ADMM_TOLERANCE	A positive integer	It is a tolerance parameter. The value must be a positive number. The default value is 0.0001

Related Topics

- DBMS_DATA_MINING Algorithm Settings: Generalized Linear Model
 The settings listed in the following table configure the behavior of the Generalized Linear Model algorithm.
- Oracle Machine Learning for SQL Concepts



Oracle Machine Learning for SQL Concepts for information about neural network

DBMS_DATA_MINING — Solver Settings: LBFGS

The settings listed in the following table configure the behavior of L-BFGS. Neural Network and Generalized Linear Model (GLM) use these settings.

Table 62-37 DBMS_DATA_MINING L-BFGS Settings

Setting Name	Setting Value	Description
LBFGS_GRADIENT_TOLERANCE	An integer greater than 0 represented as a character string	Defines gradient infinity norm tolerance for L-BFGS. Default value is 1E-9.
		Expression:
		TO_CHAR (0.00000002)
LBFGS_HISTORY_DEPTH	A positive integer.	Defines the number of historical copies kept in L-BFGS solver.
		The default value is 20.
LBFGS_SCALE_HESSIAN	LBFGS_SCALE_HESSIAN_ENABLE	Defines whether to scale Hessian in L-
	LBFGS SCALE HESSIAN DISABLE	BFGS or not.
		Default value is LBFGS SCALE HESSIAN ENABLE.

Related Topics

DBMS_DATA_MINING — Algorithm Settings: Neural Network
 The settings listed in the following table configure the behavior of the Neural Network algorithm.

DBMS_DATA_MINING — Algorithm Settings: Generalized Linear Model
 The settings listed in the following table configure the behavior of the Generalized Linear Model algorithm.

See Also:

Oracle Machine Learning for SQL Concepts for information about neural network

DBMS_DATA_MINING Datatypes

The DBMS_DATA_MINING package defines object data types for processing transactional data. The package also defines a type for user-specified transformations. These types are called DM NESTED n, where n identifies the Oracle data type of the nested attributes.

The Oracle Machine Learning for SQL object data types are described in the following table:

Table 62-38 DBMS_DATA_MINING Summary of Data Types

Datatype	Description
DM_NESTED_BINARY_DOUBLE	The name and value of a numerical attribute of type BINARY_DOUBLE.
DM_NESTED_BINARY_DOUBLES	A collection of DM_NESTED_BINARY_DOUBLE.
DM_NESTED_BINARY_FLOAT	The name and value of a numerical attribute of type BINARY_FLOAT.
DM_NESTED_BINARY_FLOATS	A collection of DM_NESTED_BINARY_FLOAT.
DM_NESTED_CATEGORICAL	The name and value of a categorical attribute of type CHAR, VARCHAR, or VARCHAR2.
DM_NESTED_CATEGORICALS	A collection of DM_NESTED_CATEGORICAL.
DM_NESTED_NUMERICAL	The name and value of a numerical attribute of type NUMBER or FLOAT.
DM_NESTED_NUMERICALS	A collection of DM_NESTED_NUMERICAL.
ORA_MINING_VARCHAR2_NT	A table of VARCHAR2 (4000).
TRANSFORM_LIST	A list of user-specified transformations for a model. Accepted as a parameter by the CREATE_MODEL Procedure.
	This collection type is defined in the DBMS_DATA_MINING_TRANSFORM package.

For more information about processing nested data, see *Oracle Machine Learning for SQL User's Guide*.

Note:

Starting from Oracle Database 12c Release 2, *GET_MODEL_DETAILS are deprecated and are replaced with *Model Detail Views*. See *Oracle Machine Learning for SQL User's Guide*.

Deprecated Types

This topic contains tables listing deprecated types.

The DBMS_DATA_MINING package defines object datatypes for storing information about model attributes. Most of these types are returned by the table functions GET_n , where n identifies the type of information to return. These functions take a model name as input and return the requested information as a collection of rows.

For a list of the GET functions, see "Summary of DBMS_DATA_MINING Subprograms".

All the table functions use pipelining, which causes each row of output to be materialized as it is read from model storage, without waiting for the generation of the complete table object. For more information on pipelined, parallel table functions, consult the *Oracle Database PL/SQL Language Reference*.

Table 62-39 DBMS_DATA_MINING Summary of Deprecated Datatypes

Datatype	Description
DM_CENTROID	The centroid of a cluster.
DM_CENTROIDS	A collection of DM_CENTROID. A member of DM_CLUSTER.
DM_CHILD	A child node of a cluster.
DM_CHILDREN	A collection of DM_CHILD. A member of DM_CLUSTER.
DM_CLUSTER	A cluster. A cluster includes DM_PREDICATES, DM_CHILDREN, DM_CENTROIDS, and DM_HISTOGRAMS. It also includes a DM_RULE.
	See also, DM_CLUSTER Fields.
DM_CLUSTERS	A collection of DM_CLUSTER. Returned by GET_MODEL_DETAILS_KM Function, GET_MODEL_DETAILS_OC Function, and GET_MODEL_DETAILS_EM Function.
	See also, DM_CLUSTER Fields.
DM_CONDITIONAL	The conditional probability of an attribute in a Naive Bayes model
DM_CONDITIONALS	A collection of DM_CONDITIONAL. Returned by GET_MODEL_DETAILS_NB Function.
DM_COST_ELEMENT	The actual and predicted values in a cost matrix.
DM_COST_MATRIX	A collection of DM_COST_ELEMENT. Returned by GET_MODEL_COST_MATRIX Function.
DM_EM_COMPONENT	A component of an Expectation Maximization model.
DM_EM_COMPONENT_SET	A collection of DM_EM_COMPONENT. Returned by GET_MODEL_DETAILS_EM_COMP Function.
DM_EM_PROJECTION	A projection of an Expectation Maximization model.
DM_EM_PROJECTION_SET	A collection of DM_EM_PROJECTION. Returned by GET_MODEL_DETAILS_EM_PROJ Function.
DM_GLM_COEFF	The coefficient and associated statistics of an attribute in a Generalized Linear Model.
DM_GLM_COEFF_SET	A collection of DM_GLM_COEFF. Returned by GET_MODEL_DETAILS_GLM Function.
DM_HISTOGRAM_BIN	A histogram associated with a cluster.



Table 62-39 (Cont.) DBMS_DATA_MINING Summary of Deprecated Datatypes

DM_HISTOGRAMS	A collection of DM_HISTOGRAM_BIN. A member of DM_CLUSTER.
DV 185V	
DI TEEL	See also, DM_CLUSTER Fields.
DM_ITEM	An item in an association rule.
DM_ITEMS	A collection of DM_ITEM.
DM_ITEMSET	A collection of DM_ITEMS.
DM_ITEMSETS	A collection of DM_ITEMSET. Returned by GET_FREQUENT_ITEMSETS Function.
DM_MODEL_GLOBAL_DETAIL	High-level statistics about a model.
DM_MODEL_GLOBAL_DETAILS	A collection of DM_MODEL_GLOBAL_DETAIL. Returned by GET_MODEL_DETAILS_GLOBAL Function.
DM_NB_DETAIL	Information about an attribute in a Naive Bayes model.
DM_NB_DETAILS	A collection of DM_DB_DETAIL. Returned by GET_MODEL_DETAILS_NB Function.
DM_NMF_ATTRIBUTE	An attribute in a feature of a Non-Negative Matrix Factorization model.
DM_NMF_ATTRIBUTE_SET	A collection of DM_NMF_ATTRIBUTE. A member of DM_NMF_FEATURE.
DM_NMF_FEATURE	A feature in a Non-Negative Matrix Factorization model.
DM_NMF_FEATURE_SET	A collection of DM_NMF_FEATURE. Returned by GET_MODEL_DETAILS_NMF Function.
DM_PREDICATE	Antecedent and consequent in a rule.
DM_PREDICATES	A collection of DM_PREDICATE. A member of DM_RULE and DM_CLUSTER. Predicates are returned by GET_ASSOCIATION_RULES Function, GET_MODEL_DETAILS_EM Function, GET_MODEL_DETAILS_KM Function, and GET_MODEL_DETAILS_OC Function. See also, DM_CLUSTER Fields.
DM_RANKED_ATTRIBUTE	An attribute ranked by its importance in an Attribute Importance model.
DM_RANKED_ATTRIBUTES	A collection of DM_RANKED_ATTRIBUTE. Returned by GET_MODEL_DETAILS_AI Function.
DM RULE	A rule that defines a conditional relationship.
_	The rule can be one of the association rules returned by GET_ASSOCIATION_RULES Function, or it can be a rule associated with a cluster in the collection of clusters returned by GET_MODEL_DETAILS_KM Function and GET_MODEL_DETAILS_OC Function. See also, DM_CLUSTER Fields.
DM_RULES	A collection of DM_RULE. Returned by GET_ASSOCIATION_RULES Function.
	See also, DM_CLUSTER Fields.
DM_SVD_MATRIX	A factorized matrix S, V, or U returned by a Singular Value Decomposition model.

Table 62-39 (Cont.) DBMS_DATA_MINING Summary of Deprecated Datatypes

Datatype	Description
DM_SVD_MATRIX_SET	A collection of DM_SVD_MATRIX. Returned by GET_MODEL_DETAILS_SVD Function.
DM_SVM_ATTRIBUTE	The name, value, and coefficient of an attribute in a Support Vector Machine model.
DM_SVM_ATTRIBUTE_SET	A collection of DM_SVM_ATTRIBUTE. Returned by GET_MODEL_DETAILS_SVM Function. Also a member of DM_SVM_LINEAR_COEFF.
DM_SVM_LINEAR_COEFF	The linear coefficient of each attribute in a Support Vector Machine model.
DM_SVM_LINEAR_COEFF_SET	A collection of DM_SVM_LINEAR_COEFF. Returned by GET_MODEL_DETAILS_SVM Function for an SVM model built using the linear kernel.
DM_TRANSFORM	The transformation and reverse transformation expressions for an attribute.
DM_TRANSFORMS	A collection of DM_TRANSFORM. Returned by GET_MODEL_TRANSFORMATIONS Function.

Return Values for Clustering Algorithms

The table contains description of $\mathtt{DM}_\mathtt{CLUSTER}$ return value columns, nested table columns, and rows.

Table 62-40 DM_CLUSTER Return Values for Clustering Algorithms

Return Value	Description		
DM_CLUSTERS	A set of rows of type DM_CLUSTER. The rows have the following columns:		
	(id cluster_id record_count parent tree_level dispersion split_predicate child centroid histogram rule	NUMBER, NUMBER, NUMBER, NUMBER, DM_PREDICATES, DM_CHILDREN, DM_CENTROIDS,	
DM_PREDICATE		_	imns each return nested tables of type PREDICATE, have the following columns:
	attrik condit attrik attrik attrik		VARCHAR2(4000), NUMBER,

DM_CLUSTER Fields

The following table describes ${\tt DM_CLUSTER}$ fields.

Table 62-41 DM_CLUSTER Fields

Column Name	Description
id	Cluster identifier
cluster_id	The ID of a cluster in the model
record_count	Specifies the number of records
parent	Parent ID
tree_level	Specifies the number of splits from the root
dispersion	A measure used to quantify whether a set of observed occurrences are dispersed compared to a standard statistical model.
split_predicate	The split_predicate column of DM_CLUSTER returns a nested table of type DM_PREDICATES. Each row, of type DM_PREDICATE, has the following columns:
	<pre>(attribute_name</pre>
	Note: The Expectation Maximization algorithm uses all the fields except dispersion and split predicate.
child	The child column of DM_CLUSTER returns a nested table of type DM_CHILDREN. The rows, of type DM_CHILD, have a single column of type NUMBER, which contains the identifiers of each child.
centroid	The centroid column of DM_CLUSTER returns a nested table of type DM_CENTROIDS. The rows, of type DM_CENTROID, have the following columns:
	(attribute_name VARCHAR2(4000), attribute_subname VARCHAR2(4000), mean NUMBER, mode_value VARCHAR2(4000), variance NUMBER)



Table 62-41 (Cont.) DM_CLUSTER Fields

Column Name	Description
histogram	The histogram column of DM_CLUSTER returns a nested table of type DM_HISTOGRAMS. The rows, of type DM_HISTOGRAM_BIN, have the following columns:
	(attribute_name VARCHAR2(4000), attribute_subname VARCHAR2(4000), bin_id NUMBER, lower_bound NUMBER, upper_bound NUMBER, label VARCHAR2(4000), count NUMBER)
rule	The rule column of DM_CLUSTER returns a single row of type DM_RULE. The columns are:
	(rule_id INTEGER, antecedent DM_PREDICATES, consequent DM_PREDICATES, rule_support NUMBER, rule_confidence NUMBER, rule_lift NUMBER, antecedent_support NUMBER, consequent_support NUMBER, number_of_items INTEGER)

Usage Notes

- The table function pipes out rows of type DM_CLUSTER. For information on Oracle Machine Learning for SQL data types and piped output from table functions, see "Data Types".
- For descriptions of predicates (DM_PREDICATE) and rules (DM_RULE), see GET ASSOCIATION RULES Function.

Summary of DBMS_DATA_MINING Subprograms

This table summarizes the subprograms included in the DBMS DATA MINING package.

The GET_* interfaces are replaced by model views. Oracle recommends that users leverage model detail views instead. For more information, refer to Model Detail Views in *Oracle Machine Learning for SQL User's Guide* and Static Data Dictionary Views: ALL_ALL_TABLES to ALL_OUTLINES in *Oracle Database Reference*.

Table 62-42 DBMS_DATA_MINING Package Subprograms

Subprogram	Purpose
ADD_COST_MATRIX Procedure	Adds a cost matrix to a classification model
ADD_PARTITION Procedure	Adds single or multiple partitions in an existing partition model
ALTER_REVERSE_EXPRESSION Procedure	Changes the reverse transformation expression to an expression that you specify
APPLY Procedure	Applies a model to a data set (scores the data)

Table 62-42 (Cont.) DBMS_DATA_MINING Package Subprograms

Subprogram	Purpose
COMPUTE_CONFUSION_MATRIX Procedure	Computes the confusion matrix for a classification model
COMPUTE_CONFUSION_MATRIX_PART Procedure	Computes the evaluation matrix for partitioned models
COMPUTE_LIFT Procedure	Computes lift for a classification model
COMPUTE_LIFT_PART Procedure	Computers lift for partitioned models
COMPUTE_ROC Procedure	Computes Receiver Operating Characteristic (ROC) for a classification model
COMPUTE_ROC_PART Procedure	Computes Receiver Operating Characteristic (ROC) for a partitioned model
CREATE_MODEL Procedure	Creates a model
CREATE_MODEL2 Procedure	Creates a model without extra persistent stages
Create Model Using Registration Information	Fetches setting information from JSON object
DROP_ALGORITHM Procedure	Drops the registered algorithm information.
DROP_PARTITION Procedure	Drops a single partition
DROP_MODEL Procedure	Drops a model
EXPORT_MODEL Procedure	Exports a model to a dump file
EXPORT_SERMODEL Procedure	Exports a model in a serialized format
FETCH_JSON_SCHEMA Procedure	Fetches and reads JSON schema from all_mining_algorithms view
GET_MODEL_COST_MATRIX Function	Returns the cost matrix for a model
IMPORT_MODEL Procedure	Imports a model into a user schema
IMPORT_ONNX_MODEL Procedure	Imports an ONNX model into the Database
IMPORT_SERMODEL Procedure	Imports a serialized model back into the database
JSON Schema for R Extensible Algorithm	Displays flexibility in creating JSON schema for R Extensible
REGISTER_ALGORITHM Procedure	Registers a new algorithm
RANK_APPLY Procedure	Ranks the predictions from the APPLY results for a classification model
REMOVE_COST_MATRIX Procedure	Removes a cost matrix from a model
RENAME_MODEL Procedure	Renames a model

Deprecated GET_MODEL_DETAILS

Starting from Oracle Database 12c Release 2, the following $\texttt{GET}_{\texttt{MODEL}}$ DETAILS are deprecated:

Table 62-43 Deprecated GET_MODEL_DETAILS Functions

Subprogram	Purpose
GET_ASSOCIATION_RULES Function	Returns the rules from an association model



Table 62-43 (Cont.) Deprecated GET_MODEL_DETAILS Functions

Subprogram	Purpose
GET_FREQUENT_ITEMSETS Function	Returns the frequent itemsets for an association model
GET_MODEL_DETAILS_AI Function	Returns details about an attribute importance model
GET_MODEL_DETAILS_EM Function	Returns details about an Expectation Maximization model
GET_MODEL_DETAILS_EM_COMP Function	Returns details about the parameters of an Expectation Maximization model
GET_MODEL_DETAILS_EM_PROJ Function	Returns details about the projects of an Expectation Maximization model
GET_MODEL_DETAILS_GLM Function	Returns details about a Generalized Linear Model model
GET_MODEL_DETAILS_GLOBAL Function	Returns high-level statistics about a model
GET_MODEL_DETAILS_KM Function	Returns details about a k-Means model
GET_MODEL_DETAILS_NB Function	Returns details about a Naive Bayes model
GET_MODEL_DETAILS_NMF Function	Returns details about a Non-Negative Matrix Factorization model
GET_MODEL_DETAILS_OC Function	Returns details about an O-Cluster model
GET_MODEL_SETTINGS Function	Returns the settings used to build the given model This function is replaced with USER/ALL/ DBA MINING MODEL SETTINGS
GET_MODEL_SIGNATURE Function	Returns the list of columns from the build input table
	This function is replaced with USER/ALL/ DBA_MINING_MODEL_ATTRIBUTES
GET_MODEL_DETAILS_SVD Function	Returns details about a Singular Value Decomposition model
GET_MODEL_DETAILS_SVM Function	Returns details about a Support Vector Machine model with a linear kernel
GET_MODEL_TRANSFORMATIONS Function	Returns the transformations embedded in a model This function is replaced with USER/ALL/ DBA_MINING_MODEL_XFORMS
GET_MODEL_DETAILS_XML Function	Returns details about a Decision Tree model
GET_TRANSFORM_LIST Procedure	Converts between two different transformation specification formats

Related Topics

- Oracle Machine Learning for SQL User's Guide
- Oracle Database Reference

ADD_COST_MATRIX Procedure

The ADD_COST_MATRIX procedure associates a cost matrix table with a classification model. The cost matrix biases the model by assigning costs or benefits to specific model outcomes.

The cost matrix is stored with the model and taken into account when the model is scored.

You can also specify a cost matrix inline when you invoke an Oracle Machine Learning for SQL function for scoring. To view the scoring matrix for a model, query the DM\$VC prefixed model view. Refer to Model Detail View for Classification Algorithm.

To obtain the default scoring matrix for a model, query the DM\$VC prefixed model view. To remove the default scoring matrix from a model, use the REMOVE_COST_MATRIX procedure. See REMOVE_COST_MATRIX Procedure.

See Also:

- "Biasing a Classification Model" in Oracle Machine Learning for SQL Concepts for more information about costs
- Oracle Database SQL Language Reference for syntax of inline cost matrix
- Specifying Costs in Oracle Machine Learning for SQL User's Guide

Syntax

Parameters

Table 62-44 ADD_COST_MATRIX Procedure Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is assumed.
cost_matrix_table_name	Name of the cost matrix table (described in Table 62-45).
cost_matrix_schema_name	Schema of the cost matrix table. If no schema is specified, then the current schema is used.
partition_name	Name of the partition in a partitioned model

Usage Notes

- 1. If the model is not in your schema, then ADD_COST_MATRIX requires the ALTER ANY MINING MODEL system privilege or the ALTER object privilege for the machine learning model.
- 2. The cost matrix table must have the columns shown in Table 62-45.



Table 62-45 Required Columns in a Cost Matrix Table

Column Name	Data Type
ACTUAL_TARGET_VALUE	Valid target data type
PREDICTED_TARGET_VALUE	Valid target data type
COST	NUMBER, FLOAT, BINARY_DOUBLE, or BINARY_FLOAT



Oracle Machine Learning for SQL User's Guide for valid target data types

3. The types of the actual and predicted target values must be the same as the type of the model target. For example, if the target of the model is <code>BINARY_DOUBLE</code>, then the actual and predicted values must be <code>BINARY_DOUBLE</code>. If the actual and predicted values are <code>CHAR</code> or <code>VARCHAR</code>, then <code>ADD_COST_MATRIX</code> treats them as <code>VARCHAR2</code> internally.

If the types do not match, or if the actual or predicted value is not a valid target value, then the $\mathtt{ADD_COST_MATRIX}$ procedure raises an error.

Note:

If a reverse transformation is associated with the target, then the actual and predicted values must be consistent with the target after the reverse transformation has been applied.

See "Reverse Transformations and Model Transparency" under the "About Transformation Lists" section in DBMS_DATA_MINING_TRANSFORM Operational Notes for more information.

- 4. Since a benefit can be viewed as a negative cost, you can specify a benefit for a given outcome by providing a negative number in the costs column of the cost matrix table.
- 5. All classification algorithms can use a cost matrix for scoring. The Decision Tree algorithm can also use a cost matrix at build time. If you want to build a Decision Tree model with a cost matrix, specify the cost matrix table name in the CLAS_COST_TABLE_NAME setting in the settings table for the model. See Table 62-7.
 - The cost matrix used to create a Decision Tree model becomes the default scoring matrix for the model. If you want to specify different costs for scoring, use the REMOVE_COST_MATRIX procedure to remove the cost matrix and the ADD_COST_MATRIX procedure to add a new one.
- 6. Scoring on a partitioned model is partition-specific. Scoring cost matrices can be added to or removed from an individual partition in a partitioned model. If PARTITION_NAME is NOT NULL, then the model must be a partitioned model. The COST_MATRIX is added to that partition of the partitioned model.

If the PARTITION_NAME is NULL, but the model is a partitioned model, then the COST_MATRIX table is added to every partition in the model.



Example

This example creates a cost matrix table called <code>COSTS_NB</code> and adds it to a Naive Bayes model called <code>NB_SH_CLAS_SAMPLE</code>. The model has a binary target: 1 means that the customer responds to a promotion; 0 means that the customer does not respond. The cost matrix assigns a cost of .25 to misclassifications of customers who do not respond and a cost of .75 to misclassifications of customers who do respond. This means that it is three times more costly to misclassify responders than it is to misclassify non-responders.

```
CREATE TABLE costs nb (
 actual_target_value
                             NUMBER,
                        NUMBER,
 predicted_target_value
                             NUMBER);
INSERT INTO costs nb values (0, 0, 0);
INSERT INTO costs nb values (0, 1, .25);
INSERT INTO costs nb values (1, 0, .75);
INSERT INTO costs nb values (1, 1, 0);
COMMIT;
EXEC dbms data mining.add cost matrix('nb sh clas sample', 'costs nb');
SELECT cust gender, COUNT(*) AS cnt, ROUND(AVG(age)) AS avg age
  FROM mining data apply v
  WHERE PREDICTION (nb sh clas sample COST MODEL
     USING cust_marital_status, education, household_size) = 1
  GROUP BY cust gender
  ORDER BY cust_gender;
      CNT AVG AGE
 -----
       72 39
555 44
```

ADD PARTITION Procedure

ADD_PARTITION procedure supports a single or multiple partition addition to an existing partitioned model.

The ADD_PARTITION procedure derives build settings and user-defined expressions from the existing model. The target column must exist in the input data query when adding partitions to a supervised model.

Syntax

Parameters

Table 62-46 ADD_PARTITION Procedure Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.



Table 62-46 (Cont.) ADD_PARTITION Procedure Parameters

Description	
An arbitrary SQL statement that provides data to the model build. The user must have privilege to evaluate this query.	
Allows users to control the conditional behavior of ADD for cases where rows in the input dataset conflict with existing partitions in the model. The following are the possible values:	
REPLACE: Replaces the existing partition for which the conflicting keys are found. The state of the	
 ERROR: Terminates the ADD operation without adding any partitions. IGNORE: Eliminates the rows having the conflicting keys. 	



For better performance, Oracle recommends using DROP_PARTITION followed by the ADD_PARTITION instead of using the REPLACE option.

ALTER_REVERSE_EXPRESSION Procedure

This procedure replaces a reverse transformation expression with an expression that you specify. If the attribute does not have a reverse expression, the procedure creates one from the specified expression.

You can also use this procedure to customize the output of clustering, feature extraction, and anomaly detection models.

Syntax

```
DBMS_DATA_MINING.ALTER_REVERSE_EXPRESSION (

model_name VARCHAR2,
expression CLOB,
attribute_name VARCHAR2 DEFAULT NULL,
attribute_subname VARCHAR2 DEFAULT NULL);
```

Parameters

Table 62-47 ALTER_REVERSE_EXPRESSION Procedure Parameters

D	Provide the second seco
Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, your own schema is used.
expression	An expression to replace the reverse transformation associated with the attribute.
attribute_name	Name of the attribute. Specify NULL if you wish to apply expression to a cluster, feature, or One-Class SVM prediction.
attribute_subname	Name of the nested attribute if attribute_name is a nested column, otherwise NULL.

Usage Notes

 For purposes of model transparency, Oracle Machine Learning for SQL provides reverse transformations for transformations that are embedded in a model. Reverse transformations are applied to the attributes returned in model detail views and to the scored target of predictive models.

See Also:

- "About Transformation Lists" under DBMS_DATA_MINING_TRANSFORM Operational Notes
- Model Detail Views in Oracle Machine Learning for SQL User's Guide
- 2. If you alter the reverse transformation for the target of a model that has a cost matrix, you must specify a transformation expression that has the same type as the actual and predicted values in the cost matrix. Also, the reverse transformation that you specify must result in values that are present in the cost matrix.



"ADD_COST_MATRIX Procedure" and Oracle Machine Learning for SQL Concepts for information about cost matrixes.

- 3. To prevent reverse transformation of an attribute, you can specify NULL for expression.
- 4. The reverse transformation expression can contain a reference to a PL/SQL function that returns a valid Oracle data type. For example, you could define a function like the following for a categorical attribute named blood_pressure that has values 'Low', 'Medium' and 'High'.

```
CREATE OR REPLACE FUNCTION numx(c char) RETURN NUMBER IS
BEGIN

CASE c WHEN ''Low'' THEN RETURN 1;

WHEN ''Medium'' THEN RETURN 2;

WHEN ''High'' THEN RETURN 3;

ELSE RETURN null;

END CASE;

END numx;
```

Then you could invoke ALTER REVERSE EXPRESION for blood pressure as follows.

5. You can use ALTER_REVERSE_EXPRESSION to label clusters produced by clustering models and features produced by feature extraction.

You can use <code>ALTER_REVERSE_EXPRESSION</code> to replace the zeros and ones returned by anomaly-detection models. By default, anomaly-detection models label anomalous records with 0 and all other records with 1.





Oracle Machine Learning for SQL Concepts for information about anomaly detection

Examples

 In this example, the target (affinity_card) of the model CLASS_MODEL is manipulated internally as yes or no instead of 1 or 0 but returned as 1s and 0s when scored. The ALTER_REVERSE_EXPRESSION procedure causes the target values to be returned as TRUE or FALSE.

```
DECLARE
        v_xlst dbms_data_mining_transform.TRANSFORM_LIST;
  BEGIN
    dbms data mining transform.SET TRANSFORM(v xlst,
          'affinity card', NULL,
          'decode(affinity_card, 1, ''yes'', ''no'')',
          'decode(affinity card, ''yes'', 1, 0)');
    dbms data mining.CREATE MODEL(
     model_name => 'CLASS_MODEL',
mining_function => dbms_data_mining.classification,
data_table_name => 'mining_data_build',
case_id_column_name => 'cust_id',
target_column_name => 'affinity_card',
settings_table_name => 'NULL'
      settings_table_name => NULL,
      data_schema_name => 'oml user',
      settings_schema_name => NULL,
      xform list => v xlst );
  END;
SELECT cust income level, occupation,
          PREDICTION (CLASS MODEL USING *) predict response
      FROM mining data test WHERE age = 60 AND cust gender IN 'M'
      ORDER BY cust income level;
                                                         PREDICT RESPONSE
CUST INCOME LEVEL
                              OCCUPATION
A: Below 30,000
                              Transp.
                           Transp.
Transp.
Sales
Handler
Crafts
E: 90,000 - 109,999
E: 90,000 - 109,999
                                                                           1
G: 130,000 - 149,999
G: 130,000 - 149,999
                              Prof.
H: 150,000 - 169,999
J: 190,000 - 249,999
                              Prof.
                                                                           1
J: 190,000 - 249,999
                               Sales
BEGIN
  dbms data mining.ALTER REVERSE EXPRESSION (
     model_name => 'CLASS MODEL',
     expression => 'decode(affinity card, ''yes'', ''TRUE'', ''FALSE'')',
     attribute name => 'affinity card');
END;
column predict response on
column predict response format a20
SELECT cust income level, occupation,
             PREDICTION(CLASS MODEL USING *) predict_response
      FROM mining data test WHERE age = 60 AND cust gender IN 'M'
```

ORDER BY cust income level;

CUST_INCOME_LEVEL	OCCUPATION	PREDICT_RESPONSE
A: Below 30,000	Transp.	TRUE
E: 90,000 - 109,999	Transp.	TRUE
E: 90,000 - 109,999	Sales	TRUE
G: 130,000 - 149,999	Handler	FALSE
G: 130,000 - 149,999	Crafts	FALSE
H: 150,000 - 169,999	Prof.	TRUE
J: 190,000 - 249,999	Prof.	TRUE
J: 190,000 - 249,999	Sales	TRUE

2. This example specifies labels for the clusters that result from the sh_clus model. The labels consist of the word "Cluster" and the internal numeric identifier for the cluster.

```
BEGIN
  dbms data mining.ALTER REVERSE EXPRESSION( 'sh clus', '''Cluster ''||value');
END;
SELECT cust id, cluster id(sh clus using *) cluster id
  FROM sh aprep num
      WHERE cust id < 100011
      ORDER by cust id;
CUST ID CLUSTER ID
_____
 100001 Cluster 18
 100002 Cluster 14
 100003 Cluster 14
 100004 Cluster 18
 100005 Cluster 19
 100006 Cluster 7
 100007 Cluster 18
 100008 Cluster 14
 100009 Cluster 8
100010 Cluster 8
```

APPLY Procedure

The APPLY procedure applies a machine learning model to the data of interest, and generates the results in a table. The APPLY procedure is also referred to as **scoring**.

For predictive machine learning functions, the APPLY procedure generates predictions in a target column. For descriptive machine learning functions such as Clustering, the APPLY process assigns each case to a cluster with a probability.

In Oracle Machine Learning for SQL, the APPLY procedure is not applicable to Association models and Attribute Importance models.

Note:

Scoring can also be performed directly in SQL using the OML4SQL functions. See

- Oracle Machine Learning for SQL Functions in Oracle Database SQL Language Reference
- Scoring and Deployment in Oracle Machine Learning for SQL User's Guide

Syntax

Parameters

Table 62-48 APPLY Procedure Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
data_table_name	Name of table or view containing the data to be scored
case_id_column_name	Name of the case identifier column
result_table_name	Name of the table in which to store apply results
data_schema_name	Name of the schema containing the data to be scored

Usage Notes

- 1. The data provided for APPLY must undergo the same preprocessing as the data used to create and test the model. When you use Automatic Data Preparation, the preprocessing required by the algorithm is handled for you by the model: both at build time and apply time. (See "Automatic Data Preparation".)
- 2. APPLY creates a table in the user's schema to hold the results. The columns are algorithm-specific.

The columns in the results table are listed in Table 62-49 through Table 62-53. The case ID column name in the results table will match the case ID column name provided by you. The type of the incoming case ID column is also preserved in APPLY output.



Make sure that the case ID column does not have the same name as one of the columns that will be created by APPLY. For example, when applying a Classification model, the case ID in the scoring data must not be PREDICTION or PROBABILITY (See Table 62-49).

- 3. The data type for the PREDICTION, CLUSTER_ID, and FEATURE_ID output columns is influenced by any reverse expression that is embedded in the model by the user. If the user does not provide a reverse expression that alters the scored value type, then the types will conform to the descriptions in the following tables. See "ALTER_REVERSE_EXPRESSION Procedure".
- 4. If the model is partitioned, the result_table_name can contain results from different partitions depending on the data from the input data table. An additional column called PARTITION_NAME is added to the result table indicating the partition name that is associated with each row.

For a non-partitioned model, the behavior does not change.

Classification

The results table for Classification has the columns described in Table 62-49. If the target of the model is categorical, the PREDICTION column will have a VARCHAR2 data type. If the target has a binary type, the PREDICTION column will have the binary type of the target.

Table 62-49 APPLY Results Table for Classification

Column Name	Data type
Case ID column name	Type of the case ID
PREDICTION	Type of the target
PROBABILITY	BINARY_DOUBLE

Anomaly Detection

The results table for Anomaly Detection has the columns described in Table 62-50.

Table 62-50 APPLY Results Table for Anomaly Detection

Column Name	Data Type
Case ID column name	Type of the case ID
PREDICTION	NUMBER
PROBABILITY	BINARY_DOUBLE

Regression

The results table for Regression has the columns described in APPLY Procedure.

Table 62-51 APPLY Results Table for Regression

Column Name	Data Type
Case ID column name	Type of the case ID
PREDICTION	Type of the target

Clustering

Clustering is an unsupervised machine learning function, and hence there are no targets. The results of an APPLY procedure contain simply the cluster identifier corresponding to a case, and the associated probability. The results table has the columns described in Table 62-52.

Table 62-52 APPLY Results Table for Clustering

Column Name	Data Type
Case ID column name	Type of the case ID
CLUSTER_ID	NUMBER
PROBABILITY	BINARY_DOUBLE

Feature Extraction



Feature Extraction is also an unsupervised machine learning function, hence there are no targets. The results of an APPLY procedure will contain simply the feature identifier corresponding to a case, and the associated match quality. The results table has the columns described in Table 62-53.

Table 62-53 APPLY Results Table for Feature Extraction

Column Name	Data Type	
Case ID column name	Type of the case ID	
FEATURE_ID	NUMBER	
MATCH_QUALITY	BINARY_DOUBLE	

Examples

This example applies the GLM Regression model <code>GLMR_SH_REGR_SAMPLE</code> to the data in the <code>MINING_DATA_APPLY_V</code> view. The <code>APPLY</code> results are output of the table <code>REGRESSION_APPLY_RESULT</code>.

```
SQL> BEGIN
      DBMS DATA MINING.APPLY (
      model name => 'glmr sh regr sample',
      data table name => 'mining_data_apply_v',
      case id column name => 'cust id',
      result table name => 'regression apply result');
   END;
SQL> SELECT * FROM regression apply result WHERE cust id > 101485;
  CUST ID PREDICTION
_____
   101486 22.8048824
   101487 25.0261101
   101488 48.6146619
   101489 51.82595
   101490 22.6220714
   101491 61.3856816
   101492 24.1400748
   101493 58.034631
   101494 45.7253149
   101495 26.9763318
   101496 48.1433425
   101497 32.0573434
   101498 49.8965531
   101499 56.270656
   101500 21.1153047
```

COMPUTE_CONFUSION_MATRIX Procedure

This procedure computes a confusion matrix, stores it in a table in the user's schema, and returns the model accuracy.

A confusion matrix is a test metric for classification models. It compares the predictions generated by the model with the actual target values in a set of test data. The confusion matrix lists the number of times each class was correctly predicted and the number of times it was predicted to be one of the other classes.

COMPUTE CONFUSION MATRIX accepts three input streams:

- The predictions generated on the test data. The information is passed in three columns:
 - Case ID column
 - Prediction column
 - Scoring criterion column containing either probabilities or costs
- The known target values in the test data. The information is passed in two columns:
 - Case ID column
 - Target column containing the known target values
- (Optional) A cost matrix table with predefined columns. See the Usage Notes for the column requirements.

See Also:

Oracle Machine Learning for SQL Concepts for more details about confusion matrixes and other test metrics for classification

```
"COMPUTE_LIFT Procedure"
```

"COMPUTE_ROC Procedure"

Syntax

Parameters

Table 62-54 COMPUTE_CONFUSION_MATRIX Procedure Parameters

Parameter	Description
accuracy	Output parameter containing the overall percentage accuracy of the predictions.
apply_result_table_name	Table containing the predictions.
target_table_name	Table containing the known target values from the test data.
case_id_column_name	Case ID column in the apply results table. Must match the case identifier in the targets table.

Table 62-54 (Cont.) COMPUTE_CONFUSION_MATRIX Procedure Parameters

Parameter	Description
target_column_name	Target column in the targets table. Contains the known target values from the test data.
confusion_matrix_table_name	Table containing the confusion matrix. The table will be created by the procedure in the user's schema.
	The columns in the confusion matrix table are described in the Usage Notes.
score_column_name	Column containing the predictions in the apply results table.
	The default column name is PREDICTION, which is the default name created by the APPLY procedure (See "APPLY Procedure").
score_criterion_column_name	Column containing the scoring criterion in the apply results table. Contains either the probabilities or the costs that determine the predictions.
	By default, scoring is based on probability; the class with the highest probability is predicted for each case. If scoring is based on cost, the class with the lowest cost is predicted.
	The score_criterion_type parameter indicates whether probabilities or costs will be used for scoring.
	The default column name is 'PROBABILITY', which is the default name created by the APPLY procedure (See "APPLY Procedure").
	See the Usage Notes for additional information.
cost_matrix_table_name	(Optional) Table that defines the costs associated with misclassifications. If a cost matrix table is provided and the score_criterion_type parameter is set to 'COSTS', the costs in this table will be used as the scoring criteria. The columns in a cost matrix table are described in the
	Usage Notes.
apply_result_schema_name	Schema of the apply results table.
	If null, the user's schema is assumed.
target_schema_name	Schema of the table containing the known targets.
	If null, the user's schema is assumed.
cost_matrix_schema_name	Schema of the cost matrix table, if one is provided.
	If null, the user's schema is assumed.
score_criterion_type	Whether to use probabilities or costs as the scoring criterion. Probabilities or costs are passed in the column identified in the score_criterion_column_name parameter.
	The default value of score_criterion_type is 'PROBABILITY'. To use costs as the scoring criterion, specify 'COST'.
	If score_criterion_type is set to 'COST' but no cost matrix is provided and if there is a scoring cost matrix associated with the model, then the associated costs are used for scoring.
	See the Usage Notes and the Examples.



Usage Notes

- The predictive information you pass to <code>COMPUTE_CONFUSION_MATRIX</code> may be generated using SQL <code>PREDICTION</code> functions, the <code>DBMS_DATA_MINING.APPLY</code> procedure, or some other mechanism. As long as you pass the appropriate data, the procedure can compute the confusion matrix.
- Instead of passing a cost matrix to COMPUTE_CONFUSION_MATRIX, you can use a scoring cost
 matrix associated with the model. A scoring cost matrix can be embedded in the model or
 it can be defined dynamically when the model is applied. To use a scoring cost matrix,
 invoke the SQL PREDICTION COST function to populate the score criterion column.
- The predictions that you pass to COMPUTE_CONFUSION_MATRIX are in a table or view specified in apply result table name.

```
CREATE TABLE apply_result_table_name AS (

case_id_column_name VARCHAR2,

score_column_name VARCHAR2,

score_criterion_column_name VARCHAR2);
```

A cost matrix must have the columns described in Table 62-55.

Table 62-55 Columns in a Cost Matrix

Column Name	Data Type
actual_target_value	Type of the target column in the build data
<pre>predicted_target_valu e</pre>	Type of the predicted target in the test data. The type of the predicted target must be the same as the type of the actual target unless the predicted target has an associated reverse transformation.
cost	BINARY_DOUBLE



Oracle Machine Learning for SQL User's Guide for valid target data types

Oracle Machine Learning for SQL Concepts for more information about cost matrixes

• The confusion matrix created by COMPUTE_CONFUSION_MATRIX has the columns described in Table 62-56.

Table 62-56 Columns in a Confusion Matrix

Column Name	Data Type
actual_target_value	Type of the target column in the build data
<pre>predicted_target_valu e</pre>	Type of the predicted target in the test data. The type of the predicted target is the same as the type of the actual target unless the predicted target has an associated reverse transformation.
value	BINARY_DOUBLE





Oracle Machine Learning for SQL Concepts for more information about confusion matrixes

Examples

These examples use the Naive Bayes model nb sh clas sample.

Compute a Confusion Matrix Based on Probabilities

The following statement applies the model to the test data and stores the predictions and probabilities in a table.

```
CREATE TABLE nb_apply_results AS

SELECT cust_id,

PREDICTION(nb_sh_clas_sample USING *) prediction,

PREDICTION_PROBABILITY(nb_sh_clas_sample USING *) probability

FROM mining data test v;
```

Using probabilities as the scoring criterion, you can compute the confusion matrix as follows.

```
DECLARE

v_accuracy NUMBER;

BEGIN

DBMS_DATA_MINING.COMPUTE_CONFUSION_MATRIX (

accuracy => v_accuracy,

apply_result_table_name => 'nb_apply_results',

target_table_name => 'mining_data_test_v',

case_id_column_name => 'cust_id',

target_column_name => 'affinity_card',

confusion_matrix_table_name => 'nb_confusion_matrix',

score_column_name => 'PREDICTION',

score_criterion_column_name => 'PROBABILITY'

cost_matrix_table_name => null,

apply_result_schema_name => null,

cost_matrix_schema_name => null,

score_criterion_type => 'PROBABILITY');

DBMS_OUTPUT.PUT_LINE('**** MODEL_ACCURACY ****: ' || ROUND(v_accuracy,4));

END;
```

The confusion matrix and model accuracy are shown as follows.

```
**** MODEL ACCURACY ****: .7847

SQL>SELECT * from nb_confusion_matrix;
ACTUAL_TARGET_VALUE PREDICTED_TARGET_VALUE VALUE

1 0 60
0 0 891
1 1 286
0 1 263
```

Compute a Confusion Matrix Based on a Cost Matrix Table

The confusion matrix in the previous example shows a high rate of false positives. For 263 cases, the model predicted 1 when the actual value was 0. You could use a cost matrix to minimize this type of error.

The cost matrix table nb_cost_matrix specifies that a false positive is 3 times more costly than a false negative.

```
      SQL> SELECT * from nb_cost_matrix;

      ACTUAL_TARGET_VALUE
      PREDICTED_TARGET_VALUE
      COST

      0
      0
      0

      0
      1
      .75

      1
      0
      .25

      1
      1
      0
```

This statement shows how to generate the predictions using APPLY.

This statement computes the confusion matrix using the cost matrix table. The score criterion column is named 'PROBABILITY', which is the name generated by APPLY.

The resulting confusion matrix shows a decrease in false positives (212 instead of 263).

**** MODEL ACCURACY ****: .798

Compute a Confusion Matrix Based on Embedded Costs

You can use the ADD_COST_MATRIX procedure to embed a cost matrix in a model. The embedded costs can be used instead of probabilities for scoring. This statement adds the previously-defined cost matrix to the model.

```
BEGIN DBMS_DATA_MINING.ADD_COST_MATRIX ('nb_sh_clas_sample', 'nb_cost_matrix');END;/
```

The following statement applies the model to the test data using the embedded costs and stores the results in a table.

```
CREATE TABLE nb_apply_results AS

SELECT cust_id,

PREDICTION(nb_sh_clas_sample COST MODEL USING *) prediction,

PREDICTION_COST(nb_sh_clas_sample COST MODEL USING *) cost

FROM mining data test v;
```

You can compute the confusion matrix using the embedded costs.

```
DECLARE
   v accuracy
                      NUMBER;
   BEGIN
       DBMS DATA MINING.COMPUTE CONFUSION MATRIX (
             accuracy
                                            => v accuracy,
             confusion_matrix_table_name => 'nb_confusion_matrix',
             score column name => 'PREDICTION',
             score_criterion_column_name => 'COST',
             cost_matrix_table_name => null,
apply_result_schema_name => null,
target_schema_name => null,
cost_matrix_schema_name => null,
score_criterion_type => 'COST');
   END;
The results are:
**** MODEL ACCURACY ****: .798
SQL> SELECT * FROM nb confusion matrix;
ACTUAL TARGET VALUE PREDICTED TARGET VALUE VALUE

    1
    0
    91

    0
    0
    942

    1
    1
    255
```

COMPUTE_CONFUSION_MATRIX_PART Procedure

The <code>COMPUTE_CONFUSION_MATRIX_PART</code> procedure computes a confusion matrix, stores it in a table in the user's schema, and returns the model accuracy.

COMPUTE_CONFUSION_MATRIX_PART provides support to computation of evaluation metrics perpartition for partitioned models. For non-partitioned models, refer to COMPUTE_CONFUSION_MATRIX Procedure.

A confusion matrix is a test metric for classification models. It compares the predictions generated by the model with the actual target values in a set of test data. The confusion matrix

lists the number of times each class was correctly predicted and the number of times it was predicted to be one of the other classes.

COMPUTE CONFUSION MATRIX PART accepts three input streams:

- The predictions generated on the test data. The information is passed in three columns:
 - Case ID column
 - Prediction column
 - Scoring criterion column containing either probabilities or costs
- The known target values in the test data. The information is passed in two columns:
 - Case ID column
 - Target column containing the known target values
- (Optional) A cost matrix table with predefined columns. See the Usage Notes for the column requirements.

See Also:

Oracle Machine Learning for SQL Concepts for more details about confusion matrixes and other test metrics for classification

```
"COMPUTE_LIFT_PART Procedure"
"COMPUTE_ROC_PART Procedure"
```

Syntax

Parameters

Table 62-57 COMPUTE_CONFUSION_MATRIX_PART Procedure Parameters

Parameter	Description
accuracy	Output parameter containing the overall percentage accuracy of the predictions
	The output argument is changed from NUMBER to
	DM_NESTED_NUMERICALS



Table 62-57 (Cont.) COMPUTE_CONFUSION_MATRIX_PART Procedure Parameters

Parameter	Description
apply_result_table_name	Table containing the predictions
target_table_name	Table containing the known target values from the test data
case_id_column_name	Case ID column in the apply results table. Must match the case identifier in the targets table.
target_column_name	Target column in the targets table. Contains the known target values from the test data.
confusion_matrix_table_name	Table containing the confusion matrix. The table will be created by the procedure in the user's schema.
	The columns in the confusion matrix table are described in the Usage Notes.
score_column_name	Column containing the predictions in the apply results table.
	The default column name is PREDICTION, which is the default name created by the APPLY procedure (See "APPLY Procedure").
score_criterion_column_name	Column containing the scoring criterion in the apply results table. Contains either the probabilities or the costs that determine the predictions.
	By default, scoring is based on probability; the class with the highest probability is predicted for each case. If scoring is based on cost, then the class with the lowest cost is predicted.
	The score_criterion_type parameter indicates whether probabilities or costs will be used for scoring.
	The default column name is PROBABILITY, which is the default name created by the APPLY procedure (See "APPLY Procedure").
	See the Usage Notes for additional information.
score_partition_column_name	(Optional) Parameter indicating the column which contains the name of the partition. This column slices the input test results such that each partition has independent evaluation matrices computed.
cost_matrix_table_name	(Optional) Table that defines the costs associated with misclassifications. If a cost matrix table is provided and the score_criterion_type parameter is set to COSTS, the costs in this table will be used as the scoring criteria. The columns in a cost matrix table are described in the
	Usage Notes.
apply_result_schema_name	Schema of the apply results table. If null, then the user's schema is assumed.
target schema name	Schema of the table containing the known targets.
carget_screma_name	If null, then the user's schema is assumed.
cost matrix schema name	Schema of the cost matrix table, if one is provided.
	If null, then the user's schema is assumed.



Table 62-57 (Cont.) COMPUTE_CONFUSION_MATRIX_PART Procedure Parameters

Parameter	Description
score_criterion_type	Whether to use probabilities or costs as the scoring criterion. Probabilities or costs are passed in the column identified in the score_criterion_column_name parameter.
	The default value of score_criterion_type is PROBABILITY. To use costs as the scoring criterion, specify COST.
	If score_criterion_type is set to COST but no cost matrix is provided and if there is a scoring cost matrix associated with the model, then the associated costs are used for scoring.
	See the Usage Notes and the Examples.

Usage Notes

- The predictive information you pass to <code>COMPUTE_CONFUSION_MATRIX_PART</code> may be generated using SQL <code>PREDICTION</code> functions, the <code>DBMS_DATA_MINING.APPLY</code> procedure, or some other mechanism. As long as you pass the appropriate data, the procedure can compute the confusion matrix.
- Instead of passing a cost matrix to <code>COMPUTE_CONFUSION_MATRIX_PART</code>, you can use a scoring cost matrix associated with the model. A scoring cost matrix can be embedded in the model or it can be defined dynamically when the model is applied. To use a scoring cost matrix, invoke the SQL <code>PREDICTION_COST</code> function to populate the score criterion column.
- The predictions that you pass to COMPUTE_CONFUSION_MATRIX_PART are in a table or view specified in apply result table name.

```
CREATE TABLE apply_result_table_name AS (

case_id_column_name VARCHAR2,

score_column_name VARCHAR2,

score_criterion_column_name VARCHAR2);
```

A cost matrix must have the columns described in Table 62-55.

Table 62-58 Columns in a Cost Matrix

Column Name	Data Type
actual_target_value	Type of the target column in the test data
<pre>predicted_target_valu e</pre>	Type of the predicted target in the test data. The type of the predicted target must be the same as the type of the actual target unless the predicted target has an associated reverse transformation.
cost	BINARY_DOUBLE





Oracle Machine Learning for SQL User's Guide for valid target data types

Oracle Machine Learning for SQL Concepts for more information about cost matrixes

The confusion matrix created by COMPUTE_CONFUSION_MATRIX_PART has the columns described in Table 62-56.

Table 62-59 Columns in a Confusion Matrix Part

Column Name	Data Type
actual_target_value	Type of the target column in the test data
<pre>predicted_target_valu e</pre>	Type of the predicted target in the test data. The type of the predicted target is the same as the type of the actual target unless the predicted target has an associated reverse transformation.
value	BINARY_DOUBLE



Oracle Machine Learning for SQL Concepts for more information about confusion matrixes

Examples

These examples use the Naive Bayes model nb sh clas sample.

Compute a Confusion Matrix Based on Probabilities

The following statement applies the model to the test data and stores the predictions and probabilities in a table.

```
CREATE TABLE nb_apply_results AS

SELECT cust_id,

PREDICTION(nb_sh_clas_sample USING *) prediction,

PREDICTION_PROBABILITY(nb_sh_clas_sample USING *) probability

FROM mining_data_test_v;
```

Using probabilities as the scoring criterion, you can compute the confusion matrix as follows.



```
cost_matrix_table_name => null,
apply_result_schema_name => null,
target_schema_name => null,
cost_matrix_schema_name => null,
score_criterion_type => 'PROBABILITY');
DBMS_OUTPUT.PUT_LINE('**** MODEL ACCURACY ****: ' || ROUND(v_accuracy,4));
END;
//
```

The confusion matrix and model accuracy are shown as follows.

```
**** MODEL ACCURACY ****: .7847

SELECT * FROM NB_CONFUSION_MATRIX;

ACTUAL_TARGET_VALUE PREDICTED_TARGET_VALUE VALUE

1 0 60
0 0 891
1 1 286
0 1 263
```

Compute a Confusion Matrix Based on a Cost Matrix Table

The confusion matrix in the previous example shows a high rate of false positives. For 263 cases, the model predicted 1 when the actual value was 0. You could use a cost matrix to minimize this type of error.

The cost matrix table <code>nb_cost_matrix</code> specifies that a false positive is 3 times more costly than a false negative.

```
SELECT * from NB_COST_MATRIX;

ACTUAL_TARGET_VALUE PREDICTED_TARGET_VALUE COST

0 0 0 0
0 1 .75
1 0 .25
1 1 0
```

This statement shows how to generate the predictions using APPLY.

This statement computes the confusion matrix using the cost matrix table. The score criterion column is named 'PROBABILITY', which is the name generated by APPLY.

```
score_criterion_column_name => 'PROBABILITY',
score_partition_column_name => 'PARTITION_NAME'
cost_matrix_table_name => 'nb_cost_matrix',
apply_result_schema_name => null,
target_schema_name => null,
cost_matrix_schema_name => null,
score_criterion_type => 'COST');
DBMS_OUTPUT_PUT_LINE('**** MODEL ACCURACY ****: ' || ROUND(v_accuracy,4));
END;
//
```

The resulting confusion matrix shows a decrease in false positives (212 instead of 263).

```
**** MODEL ACCURACY ****: .798

SELECT * FROM NB_CONFUSION_MATRIX;

ACTUAL_TARGET_VALUE PREDICTED_TARGET_VALUE VALUE

1 0 91
0 0 942
1 1 255
0 1 212
```

Compute a Confusion Matrix Based on Embedded Costs

You can use the ADD_COST_MATRIX procedure to embed a cost matrix in a model. The embedded costs can be used instead of probabilities for scoring. This statement adds the previously-defined cost matrix to the model.

```
BEGIN
DBMS_DATA_MINING.ADD_COST_MATRIX ('nb_sh_clas_sample', 'nb_cost_matrix');
END;/
```

The following statement applies the model to the test data using the embedded costs and stores the results in a table.

```
CREATE TABLE nb_apply_results AS

SELECT cust_id,

PREDICTION(nb_sh_clas_sample COST MODEL USING *) prediction,

PREDICTION_COST(nb_sh_clas_sample COST MODEL USING *) cost

FROM mining_data_test_v;
```

You can compute the confusion matrix using the embedded costs.

```
END;
```

The results are:

```
**** MODEL ACCURACY ****: .798

SELECT * FROM NB_CONFUSION_MATRIX;

ACTUAL_TARGET_VALUE PREDICTED_TARGET_VALUE VALUE

1 0 91
0 0 942
1 1 255
0 1 212
```

COMPUTE_LIFT Procedure

This procedure computes lift and stores the results in a table in the user's schema.

Lift is a test metric for binary classification models. To compute lift, one of the target values must be designated as the positive class. COMPUTE_LIFT compares the predictions generated by the model with the actual target values in a set of test data. Lift measures the degree to which the model's predictions of the positive class are an improvement over random chance.

Lift is computed on scoring results that have been ranked by probability (or cost) and divided into quantiles. Each quantile includes the scores for the same number of cases.

COMPUTE_LIFT calculates quantile-based and cumulative statistics. The number of quantiles and the positive class are user-specified. Additionally, COMPUTE_LIFT accepts three input streams:

- The predictions generated on the test data. The information is passed in three columns:
 - Case ID column
 - Prediction column
 - Scoring criterion column containing either probabilities or costs associated with the predictions
- The known target values in the test data. The information is passed in two columns:
 - Case ID column
 - Target column containing the known target values
- (Optional) A cost matrix table with predefined columns. See the Usage Notes for the column requirements.

See Also:

Oracle Machine Learning for SQL Concepts for more details about lift and test metrics for classification

"COMPUTE CONFUSION MATRIX Procedure"

"COMPUTE_ROC Procedure"



Syntax

Table 62-60 COMPUTE_LIFT Procedure Parameters

Parameter	Description
apply_result_table_name	Table containing the predictions.
target_table_name	Table containing the known target values from the test data.
case_id_column_name	Case ID column in the apply results table. Must match the case identifier in the targets table.
target_column_name	Target column in the targets table. Contains the known target values from the test data.
lift_table_name	Table containing the lift statistics. The table will be created by the procedure in the user's schema.
	The columns in the lift table are described in the Usage Notes.
positive_target_value	The positive class. This should be the class of interest, for which you want to calculate lift.
	If the target column is a NUMBER, you can use the TO_CHAR() operator to provide the value as a string.
score_column_name	Column containing the predictions in the apply results table.
	The default column name is 'PREDICTION', which is the default name created by the APPLY procedure (See "APPLY Procedure").
score_criterion_column_name	Column containing the scoring criterion in the apply results table. Contains either the probabilities or the costs that determine the predictions.
	By default, scoring is based on probability; the class with the highest probability is predicted for each case. If scoring is based on cost, the class with the lowest cost is predicted.
	The score_criterion_type parameter indicates whether probabilities or costs will be used for scoring.
	The default column name is 'PROBABILITY', which is the default name created by the APPLY procedure (See "APPLY Procedure").
	See the Usage Notes for additional information.

Table 62-60 (Cont.) COMPUTE_LIFT Procedure Parameters

Parameter	Description
num_quantiles	Number of quantiles to be used in calculating lift. The default is 10.
cost_matrix_table_name	(Optional) Table that defines the costs associated with misclassifications. If a cost matrix table is provided and the score_criterion_type parameter is set to 'COST', the costs will be used as the scoring criteria.
	The columns in a cost matrix table are described in the Usage Notes.
apply_result_schema_name	Schema of the apply results table. If null, the user's schema is assumed.
target_schema_name	Schema of the table containing the known targets. If null, the user's schema is assumed.
cost_matrix_schema_name	Schema of the cost matrix table, if one is provided. If null, the user's schema is assumed.
score_criterion_type	Whether to use probabilities or costs as the scoring criterion. Probabilities or costs are passed in the column identified in the score_criterion_column_name parameter.
	The default value of score_criterion_type is 'PROBABILITY'. To use costs as the scoring criterion, specify 'COST'.
	If score_criterion_type is set to 'COST' but no cost matrix is provided and if there is a scoring cost matrix associated with the model, then the associated costs are used for scoring.
	See the Usage Notes and the Examples.

Usage Notes

- The predictive information you pass to COMPUTE_LIFT may be generated using SQL PREDICTION functions, the DBMS_DATA_MINING.APPLY procedure, or some other mechanism. As long as you pass the appropriate data, the procedure can compute the lift.
- Instead of passing a cost matrix to <code>COMPUTE_LIFT</code>, you can use a scoring cost matrix associated with the model. A scoring cost matrix can be embedded in the model or it can be defined dynamically when the model is applied. To use a scoring cost matrix, invoke the <code>SQL PREDICTION COST</code> function to populate the score criterion column.
- The predictions that you pass to COMPUTE_LIFT are in a table or view specified in apply results table name.

```
CREATE TABLE apply_result_table_name AS (

case_id_column_name VARCHAR2,

score_column_name VARCHAR2,

score_criterion_column_name VARCHAR2);
```

A cost matrix must have the columns described in Table 62-61.

Table 62-61 Columns in a Cost Matrix

Column Name	Data Type
actual_target_value	Type of the target column in the build data
<pre>predicted_target_valu e</pre>	Type of the predicted target in the test data. The type of the predicted target must be the same as the type of the actual target unless the predicted target has an associated reverse transformation.
cost	NUMBER



Oracle Machine Learning for SQL Concepts for more information about cost matrixes

The table created by COMPUTE_LIFT has the columns described in Table 62-62

Table 62-62 Columns in a Lift Table

Column Name	Data Type
quantile_number	NUMBER
probability_threshold	NUMBER
gain_cumulative	NUMBER
quantile_total_count	NUMBER
quantile_target_count	NUMBER
percent_records_cumulative	NUMBER
lift_cumulative	NUMBER
target_density_cumulative	NUMBER
targets_cumulative	NUMBER
non_targets_cumulative	NUMBER
lift_quantile	NUMBER
target_density	NUMBER

See Also:

Oracle Machine Learning for SQL Concepts for details about the information in the lift table

• When a cost matrix is passed to COMPUTE_LIFT, the cost threshold is returned in the probability_threshold column of the lift table.

Examples

This example uses the Naive Bayes model nb sh clas sample.

The example illustrates lift based on probabilities. For examples that show computation based on costs, see "COMPUTE_CONFUSION_MATRIX Procedure".

The following statement applies the model to the test data and stores the predictions and probabilities in a table.

```
CREATE TABLE nb_apply_results AS
    SELECT cust_id, t.prediction, t.probability
    FROM mining data test v, TABLE(PREDICTION SET(nb sh clas sample USING *)) t;
```

Using probabilities as the scoring criterion, you can compute lift as follows.

This query displays some of the statistics from the resulting lift table.

COMPUTE_LIFT PART Procedure

The COMPUTE_LIFT_PART procedure computes lift and stores the results in a table in the user's schema. This procedure provides support to the computation of evaluation metrics per-partition for partitioned models.

Lift is a test metric for binary classification models. To compute lift, one of the target values must be designated as the positive class. COMPUTE_LIFT_PART compares the predictions generated by the model with the actual target values in a set of test data. Lift measures the degree to which the model's predictions of the positive class are an improvement over random chance.

Lift is computed on scoring results that have been ranked by probability (or cost) and divided into quantiles. Each quantile includes the scores for the same number of cases.

COMPUTE_LIFT_PART calculates quantile-based and cumulative statistics. The number of quantiles and the positive class are user-specified. Additionally, COMPUTE_LIFT_PART accepts three input streams:

- The predictions generated on the test data. The information is passed in three columns:
 - Case ID column
 - Prediction column
 - Scoring criterion column containing either probabilities or costs associated with the predictions
- The known target values in the test data. The information is passed in two columns:
 - Case ID column
 - Target column containing the known target values
- (Optional) A cost matrix table with predefined columns. See the Usage Notes for the column requirements.

See Also:

Oracle Machine Learning for SQL Concepts for more details about Lift and test metrics for classification

```
"COMPUTE_LIFT Procedure"
```

"COMPUTE_CONFUSION_MATRIX Procedure"

"COMPUTE_CONFUSION_MATRIX_PART Procedure"

"COMPUTE_ROC Procedure"

"COMPUTE_ROC_PART Procedure"

Syntax



Table 62-63 COMPUTE_LIFT_PART Procedure Parameters

Parameter	Description
apply_result_table_name	Table containing the predictions
target_table_name	Table containing the known target values from the test data
case_id_column_name	Case ID column in the apply results table. Must match the case identifier in the targets table.
target_column_name	Target column in the targets table. Contains the known target values from the test data.
lift_table_name	Table containing the Lift statistics. The table will be created by the procedure in the user's schema. The columns in the Lift table are described in the Usage Notes.
positive_target_value	The positive class. This should be the class of interest, for which you want to calculate Lift.
	If the target column is a NUMBER, then you can use the TO_CHAR() operator to provide the value as a string.
score_column_name	Column containing the predictions in the apply results table.
	The default column name is PREDICTION, which is the default name created by the APPLY procedure (See "APPLY Procedure").
score_criterion_column_name	Column containing the scoring criterion in the apply results table. Contains either the probabilities or the costs that determine the predictions.
	By default, scoring is based on probability; the class with the highest probability is predicted for each case. If scoring is based on cost, then the class with the lowest cost is predicted.
	The score_criterion_type parameter indicates whether probabilities or costs will be used for scoring.
	The default column name is PROBABILITY, which is the default name created by the APPLY procedure (See "APPLY Procedure").
	See the Usage Notes for additional information.
score_partition_column_name	Optional parameter indicating the column containing the name of the partition. This column slices the input test results such that each partition has independent evaluation matrices computed.
num_quantiles	Number of quantiles to be used in calculating Lift. The default is 10.
cost_matrix_table_name	(Optional) Table that defines the costs associated with misclassifications. If a cost matrix table is provided and the score_criterion_type parameter is set to COST, then the costs will be used as the scoring criteria.
	The columns in a cost matrix table are described in the Usage Notes.
apply_result_schema_name	Schema of the apply results table If null, then the user's schema is assumed.



Table 62-63 (Cont.) COMPUTE_LIFT_PART Procedure Parameters

Parameter	Description
target_schema_name	Schema of the table containing the known targets
	If null, then the user's schema is assumed.
cost_matrix_schema_name	Schema of the cost matrix table, if one is provided
	If null, then the user's schema is assumed.
score_criterion_type	Whether to use probabilities or costs as the scoring criterion. Probabilities or costs are passed in the column identified in the score_criterion_column_name parameter.
	The default value of score_criterion_type is PROBABILITY. To use costs as the scoring criterion, specify COST.
	If score_criterion_type is set to COST but no cost matrix is provided and if there is a scoring cost matrix associated with the model, then the associated costs are used for scoring.
	See the Usage Notes and the Examples.

Usage Notes

- The predictive information you pass to COMPUTE_LIFT_PART may be generated using SQL PREDICTION functions, the DBMS_DATA_MINING.APPLY procedure, or some other mechanism. As long as you pass the appropriate data, the procedure can compute the Lift.
- Instead of passing a cost matrix to COMPUTE_LIFT_PART, you can use a scoring cost matrix associated with the model. A scoring cost matrix can be embedded in the model or it can be defined dynamically when the model is applied. To use a scoring cost matrix, invoke the SQL PREDICTION COST function to populate the score criterion column.
- The predictions that you pass to COMPUTE_LIFT_PART are in a table or view specified in apply_results_table_name.

A cost matrix must have the columns described in Table 62-61.

Table 62-64 Columns in a Cost Matrix

Column Name	Data Type
actual_target_value	Type of the target column in the test data
<pre>predicted_target_valu e</pre>	Type of the predicted target in the test data. The type of the predicted target must be the same as the type of the actual target unless the predicted target has an associated reverse transformation.
cost	NUMBER



See Also:

Oracle Machine Learning for SQL Concepts for more information about cost matrixes

The table created by COMPUTE LIFT PART has the columns described in Table 62-62

Table 62-65 Columns in a COMPUTE_LIFT_PART Table

Column Name	Data Type
quantile_number	NUMBER
probability_threshold	NUMBER
gain_cumulative	NUMBER
quantile_total_count	NUMBER
quantile_target_count	NUMBER
percent_records_cumulative	NUMBER
lift_cumulative	NUMBER
target_density_cumulative	NUMBER
targets_cumulative	NUMBER
non_targets_cumulative	NUMBER
lift_quantile	NUMBER
target_density	NUMBER

See Also:

Oracle Machine Learning for SQL Concepts for details about the information in the Lift table

• When a cost matrix is passed to COMPUTE_LIFT_PART, the cost threshold is returned in the probability_threshold column of the Lift table.

Examples

This example uses the Naive Bayes model nb sh clas sample.

The example illustrates Lift based on probabilities. For examples that show computation based on costs, see "COMPUTE_CONFUSION_MATRIX Procedure".

For a partitioned model example, see "COMPUTE_CONFUSION_MATRIX_PART Procedure".

The following statement applies the model to the test data and stores the predictions and probabilities in a table.

```
CREATE TABLE nb_apply_results AS

SELECT cust_id, t.prediction, t.probability

FROM mining_data_test_v, TABLE(PREDICTION_SET(nb_sh_clas_sample USING *)) t;
```

Using probabilities as the scoring criterion, you can compute Lift as follows.

This query displays some of the statistics from the resulting Lift table.

QUANTILE_NUMBER	PROBABILITY_THRESHOLD	GAIN_CUMULATIVE	QUANTILE_TOTAL_COUNT
1	.989335775	.15034965	55
2	.980534911	.26048951	55
3	.968506098	.374125874	55
4	.958975196	.493006993	55
5	.946705997	.587412587	55
6	.927454174	.66958042	55
7	.904403627	.748251748	55
8	.836482525	.839160839	55
10	.500184953	1	54

COMPUTE ROC Procedure

This procedure computes the receiver operating characteristic (ROC), stores the results in a table in the user's schema, and returns a measure of the model accuracy.

ROC is a test metric for binary classification models. To compute ROC, one of the target values must be designated as the positive class. <code>COMPUTE_ROC</code> compares the predictions generated by the model with the actual target values in a set of test data.

ROC measures the impact of changes in the probability threshold. The probability threshold is the decision point used by the model for predictions. In binary classification, the default probability threshold is 0.5. The value predicted for each case is the one with a probability greater than 50%.

ROC can be plotted as a curve on an X-Y axis. The false positive rate is placed on the X axis. The true positive rate is placed on the Y axis. A false positive is a positive prediction for a case

that is negative in the test data. A true positive is a positive prediction for a case that is positive in the test data.

COMPUTE ROC accepts two input streams:

- The predictions generated on the test data. The information is passed in three columns:
 - Case ID column
 - Prediction column
 - Scoring criterion column containing probabilities
- The known target values in the test data. The information is passed in two columns:
 - Case ID column
 - Target column containing the known target values

See Also:

Oracle Machine Learning for SQL Concepts for more details about ROC and test metrics for classification

"COMPUTE CONFUSION MATRIX Procedure"

"COMPUTE LIFT Procedure"

Syntax

Table 62-66 COMPUTE_ROC Procedure Parameters

Parameter	Description
roc_area_under_the_curve	Output parameter containing the area under the ROC curve (AUC). The AUC measures the likelihood that an actual positive will be predicted as positive.
	The greater the AUC, the greater the flexibility of the model in accommodating trade-offs between positive and negative class predictions. AUC can be especially important when one target class is rarer or more important to identify than another.
apply_result_table_name	Table containing the predictions.



Table 62-66 (Cont.) COMPUTE_ROC Procedure Parameters

Parameter	Description
target_table_name	Table containing the known target values from the test data.
case_id_column_name	Case ID column in the apply results table. Must match the case identifier in the targets table.
target_column_name	Target column in the targets table. Contains the known target values from the test data.
roc_table_name	Table containing the ROC output. The table will be created by the procedure in the user's schema.
	The columns in the ROC table are described in the Usage Notes.
positive_target_value	The positive class. This should be the class of interest, for which you want to calculate ROC.
	If the target column is a NUMBER, you can use the ${\tt TO_CHAR}$ () operator to provide the value as a string.
score_column_name	Column containing the predictions in the apply results table.
	The default column name is 'PREDICTION', which is the default name created by the APPLY procedure (See "APPLY Procedure").
score_criterion_column_name	Column containing the scoring criterion in the apply results table. Contains the probabilities that determine the predictions.
	The default column name is 'PROBABILITY', which is the default name created by the APPLY procedure (See "APPLY Procedure").
apply_result_schema_name	Schema of the apply results table.
	If null, the user's schema is assumed.
target_schema_name	Schema of the table containing the known targets.
	If null, the user's schema is assumed.

Usage Notes

- The predictive information you pass to <code>COMPUTE_ROC</code> may be generated using SQL <code>PREDICTION</code> functions, the <code>DBMS_DATA_MINING.APPLY</code> procedure, or some other mechanism. As long as you pass the appropriate data, the procedure can compute the receiver operating characteristic.
- The predictions that you pass to COMPUTE_ROC are in a table or view specified in apply_results_table_name.

The table created by COMPUTE ROC has the columns shown in Table 62-67.

Table 62-67 COMPUTE_ROC Output

Column	Datatype
probability	BINARY_DOUBLE
true_positives	NUMBER
false_negatives	NUMBER
false_positives	NUMBER
true_negatives	NUMBER
true_positive_fraction	NUMBER
false_positive_fraction	NUMBER



Oracle Machine Learning for SQL Concepts for details about the output of ${\tt COMPUTE\ ROC}$

ROC is typically used to determine the most desirable probability threshold. This can be
done by examining the true positive fraction and the false positive fraction. The true
positive fraction is the percentage of all positive cases in the test data that were correctly
predicted as positive. The false positive fraction is the percentage of all negative cases in
the test data that were incorrectly predicted as positive.

Given a probability threshold, the following statement returns the positive predictions in an apply result table ordered by probability.

```
SELECT case_id_column_name
    FROM apply_result_table_name
    WHERE probability > probability_threshold
    ORDER BY probability DESC;
```

There are two approaches to identifying the most desirable probability threshold. Which
approach you use depends on whether or not you know the relative cost of positive versus
negative class prediction errors.

If the costs are known, you can apply the relative costs to the ROC table to compute the minimum cost probability threshold. Suppose the relative cost ratio is: Positive Class Error Cost / Negative Class Error Cost = 20. Then execute a query like this.

```
WITH cost AS (
   SELECT probability_threshold, 20 * false_negatives + false_positives cost
   FROM ROC_table
GROUP BY probability_threshold),
   minCost AS (
        SELECT min(cost) minCost
        FROM cost)
   SELECT max(probability_threshold)probability_threshold
        FROM cost, minCost
   WHERE cost = minCost;
```

If relative costs are not well known, you can simply scan the values in the ROC table (in sorted order) and make a determination about which of the displayed trade-offs (misclassified positives versus misclassified negatives) is most desirable.

Examples

This example uses the Naive Bayes model nb sh clas sample.

The following statement applies the model to the test data and stores the predictions and probabilities in a table.

```
CREATE TABLE nb_apply_results AS
    SELECT cust_id, t.prediction, t.probability
    FROM mining_data_test_v, TABLE(PREDICTION_SET(nb_sh_clas_sample USING *)) t;
```

Using the predictions and the target values from the test data, you can compute ROC as follows.

The resulting AUC and a selection of columns from the ROC table are shown as follows.



COMPUTE_ROC_PART Procedure

The <code>COMPUTE_ROC_PART</code> procedure computes Receiver Operating Characteristic (ROC), stores the results in a table in the user's schema, and returns a measure of the model accuracy. This procedure provides support to computation of evaluation metrics per-partition for partitioned models.

ROC is a test metric for binary classification models. To compute ROC, one of the target values must be designated as the positive class. <code>COMPUTE_ROC_PART</code> compares the predictions generated by the model with the actual target values in a set of test data.

ROC measures the impact of changes in the probability threshold. The probability threshold is the decision point used by the model for predictions. In binary classification, the default probability threshold is 0.5. The value predicted for each case is the one with a probability greater than 50%.

ROC can be plotted as a curve on an x-y axis. The false positive rate is placed on the x-axis. The true positive rate is placed on the y-axis. A false positive is a positive prediction for a case that is negative in the test data. A true positive is a positive prediction for a case that is positive in the test data.

COMPUTE ROC PART accepts two input streams:

- The predictions generated on the test data. The information is passed in three columns:
 - Case ID column
 - Prediction column
 - Scoring criterion column containing probabilities
- The known target values in the test data. The information is passed in two columns:
 - Case ID column
 - Target column containing the known target values

See Also:

Oracle Machine Learning for SQL Concepts for more details about ROC and test metrics for Classification

```
"COMPUTE_ROC Procedure"

"COMPUTE_CONFUSION_MATRIX Procedure"

"COMPUTE_LIFT_PART Procedure"

"COMPUTE_LIFT Procedure"
```

Syntax

```
DBMS_DATA_MINING.compute_roc_part(
roc_area_under_curve OUT DM_NESTED_NUMERICALS,
apply_result_table_name IN VARCHAR2,
target_table_name IN VARCHAR2,
case_id_column_name IN VARCHAR2,
target_column_name IN VARCHAR2,
```



```
roc_table_name IN VARCHAR2,
positive_target_value IN VARCHAR2,
score_column_name IN VARCHAR2 DEFAULT 'PREDICTION',
score_criterion_column_name IN VARCHAR2 DEFAULT 'PROBABILITY',
score_partition_column_name IN VARCHAR2 DEFAULT 'PARTITION_NAME',
apply_result_schema_name IN VARCHAR2 DEFAULT NULL,
target_schema_name IN VARCHAR2 DEFAULT NULL);
```

Table 62-68 COMPUTE_ROC_PART Procedure Parameters

Parameter	Description
roc_area_under_the_curve	Output parameter containing the area under the ROC curve (AUC). The AUC measures the likelihood that an actual positive will be predicted as positive.
	The greater the AUC, the greater the flexibility of the model in accommodating trade-offs between positive and negative class predictions. AUC can be especially important when one target class is rarer or more important to identify than another.
	The output argument is changed from NUMBER to DM_NESTED_NUMERICALS.
apply_result_table_name	Table containing the predictions.
target_table_name	Table containing the known target values from the test data.
case_id_column_name	Case ID column in the apply results table. Must match the case identifier in the targets table.
target_column_name	Target column in the targets table. Contains the known target values from the test data.
roc_table_name	Table containing the ROC output. The table will be created by the procedure in the user's schema.
	The columns in the ROC table are described in the Usage Notes.
positive_target_value	The positive class. This should be the class of interest, for which you want to calculate ROC.
	If the target column is a NUMBER, then you can use the TO_CHAR() operator to provide the value as a string.
score_column_name	Column containing the predictions in the apply results table. The default column name is PREDICTION, which is the default name created by the APPLY procedure (See "APPLY Procedure").
score_criterion_column_name	Column containing the scoring criterion in the apply results table. Contains the probabilities that determine the predictions.
	The default column name is PROBABILITY, which is the default name created by the APPLY procedure (See "APPLY Procedure").
score_partition_column_name	Optional parameter indicating the column which contains the name of the partition. This column slices the input test results such that each partition has independent evaluation matrices computed.
apply_result_schema_name	Schema of the apply results table. If null, then the user's schema is assumed.

Table 62-68 (Cont.) COMPUTE_ROC_PART Procedure Parameters

Parameter	Description
target_schema_name	Schema of the table containing the known targets.
	If null, then the user's schema is assumed.

Usage Notes

- The predictive information you pass to COMPUTE_ROC_PART may be generated using SQL
 PREDICTION functions, the DBMS_DATA_MINING.APPLY procedure, or some other mechanism.
 As long as you pass the appropriate data, the procedure can compute the receiver operating characteristic.
- The predictions that you pass to COMPUTE_ROC_PART are in a table or view specified in apply results table name.

• The COMPUTE_ROC_PART table has the following columns:

Table 62-69 COMPUTE_ROC_PART Output

Column	Data Type
probability	BINARY_DOUBLE
true_positives	NUMBER
false_negatives	NUMBER
false_positives	NUMBER
true_negatives	NUMBER
true_positive_fraction	NUMBER
false_positive_fraction	NUMBER



Oracle Machine Learning for SQL Concepts for details about the output of ${\tt COMPUTE\ ROC\ PART}$

ROC is typically used to determine the most desirable probability threshold. This can be
done by examining the true positive fraction and the false positive fraction. The true
positive fraction is the percentage of all positive cases in the test data that were correctly
predicted as positive. The false positive fraction is the percentage of all negative cases in
the test data that were incorrectly predicted as positive.

Given a probability threshold, the following statement returns the positive predictions in an apply result table ordered by probability.

```
SELECT case_id_column_name
FROM apply_result_table_name
```



```
WHERE probability > probability_threshold ORDER BY probability DESC;
```

 There are two approaches to identify the most desirable probability threshold. The approach you use depends on whether you know the relative cost of positive versus negative class prediction errors.

If the costs are known, then you can apply the relative costs to the ROC table to compute the minimum cost probability threshold. Suppose the relative cost ratio is: Positive Class Error Cost / Negative Class Error Cost = 20. Then execute a query as follows:

```
WITH cost AS (
   SELECT probability_threshold, 20 * false_negatives + false_positives cost
   FROM ROC_table
GROUP BY probability_threshold),
   minCost AS (
        SELECT min(cost) minCost
        FROM cost)
        SELECT max(probability_threshold)probability_threshold
        FROM cost, minCost
        WHERE cost = minCost;
```

If relative costs are not well known, then you can simply scan the values in the ROC table (in sorted order) and make a determination about which of the displayed trade-offs (misclassified positives versus misclassified negatives) is most desirable.

```
SELECT * FROM ROC_table
ORDER BY probability threshold;
```

Examples

This example uses the Naive Bayes model nb sh clas sample.

The following statement applies the model to the test data and stores the predictions and probabilities in a table.

```
CREATE TABLE nb_apply_results AS

SELECT cust_id, t.prediction, t.probability

FROM mining data test v, TABLE(PREDICTION SET(nb sh clas sample USING *)) t;
```

Using the predictions and the target values from the test data, you can compute ROC as follows.

```
DECLARE

v_area_under_curve NUMBER;

BEGIN

DBMS_DATA_MINING.COMPUTE_ROC_PART (

roc_area_under_curve => v_area_under_curve,
apply_result_table_name => 'nb_apply_results',
target_table_name => 'mining_data_test_v',
case_id_column_name => 'cust_id',
target_column_name => 'affinity_card',
roc_table_name => 'nb_roc',
positive_target_value => '1',
score_column_name => 'PREDICTION',
score_criterion_column_name => 'PROBABILITY');
score_partition_column_name => 'PARTITION_NAME'
DBMS_OUTPUT.PUT_LINE('**** AREA_UNDER_ROC_CURVE_****: ' ||
ROUND(v_area_under_curve, 4));
```



```
END;
```

The resulting AUC and a selection of columns from the ROC table are shown as follows.

CREATE_MODEL Procedure

This procedure creates an Oracle Machine Learning for SQL model with a given machine learning function.

Syntax

```
DBMS_DATA_MINING.CREATE_MODEL (

model_name IN VARCHAR2,
mining_function IN VARCHAR2,
data_table_name IN VARCHAR2,
case_id_column_name IN VARCHAR2,
target_column_name IN VARCHAR2 DEFAULT NULL,
settings_table_name IN VARCHAR2 DEFAULT NULL,
data_schema_name IN VARCHAR2 DEFAULT NULL,
settings_schema_name IN VARCHAR2 DEFAULT NULL,
xform_list IN TRANSFORM_LIST DEFAULT NULL);
```

Table 62-70 CREATE_MODEL Procedure Parameters

Danamatan	Paradiation
Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
	See the Usage Notes for model naming restrictions.
mining_function	The machine learning function. Values are listed in Table 62-3.
data_table_name	Table or view containing the build data
case_id_column_name	Case identifier column in the build data.
target_column_name	For supervised models, the target column in the build data. \mathtt{NULL} for unsupervised models.

Table 62-70 (Cont.) CREATE_MODEL Procedure Parameters

Parameter	Description	
settings_table_name	Table containing build settings for the model. NULL if there is no settings table (only default settings are used).	
data_schema_name	Schema hosting the build data. If $\mathtt{NULL},$ then the user's schema is assumed.	
settings_schema_name	Schema hosting the settings table. If ${\tt NULL}$ then the user's schema is assumed.	
xform_list	A list of transformations to be used in addition to or instead of automatic transformations, depending on the value of the PREP_AUTO setting. (See "Automatic Data Preparation".)	
	The datatype of xform_list is TRANSFORM_LIST, which consists of records of type TRANSFORM_REC. Each TRANSFORM_REC specifies the transformation information for a single attribute.	
	TYPE TRANFORM_REC IS RECORD (attribute_name VARCHAR2(4000), attribute_subname VARCHAR2(4000), expression EXPRESSION_REC, reverse_expression EXPRESSION_REC, attribute_spec VARCHAR2(4000));	
	The expression field stores a SQL expression for transforming the attribute. The reverse_expression field stores a SQL expression for reversing the transformation in model details and, if the attribute is a target, in the results of scoring. The SQL expressions are manipulated by routines in the DBMS_DATA_MINING_TRANSFORM package:	
	SET_EXPRESSION Procedure	
	GET_EXPRESSION FunctionSET_TRANSFORM Procedure	
	The attribute_spec field identifies individualized treatment for the attribute. See the Usage Notes for details.	
	See Table 63-1 for details about the TRANSFORM_REC type.	

Usage Notes

- 1. You can use the attribute_spec field of the xform_list argument to identify an attribute as unstructured text or to disable Automatic Data Preparation for the attribute. The attribute spec can have the following values:
 - TEXT: Indicates that the attribute contains unstructured text. The TEXT value may
 optionally be followed by POLICY_NAME, TOKEN_TYPE, MAX_FEATURES, and
 MIN_DOCUMENTS parameters.

TOKEN_TYPE has the following possible values: NORMAL, STEM, THEME, SYNONYM, BIGRAM, STEM_BIGRAM. SYNONYM may be optionally followed by a thesaurus name in square brackets.

MAX FEATURES specifies the maximum number of tokens extracted from the text.

MIN_DOCUMENTS specifies the minimal number of documents in which every selected token shall occur. (For information about creating a text policy, see CTX_DDL.CREATE_POLICY in *Oracle Text Reference*).

Oracle Machine Learning for SQL can process columns of VARCHAR2/CHAR, CLOB, BLOB, and BFILE as text. If the column is VARCHAR2 or CHAR and you do not specify TEXT, then OML4SQL processes the column as categorical data. If the column is CLOB, then OML4SQL processes it as text by default (You do not need to specify it as TEXT. However, you do need to provide an Oracle Text Policy in the settings). If the column is BLOB or BFILE, then you must specify it as TEXT, otherwise CREATE_MODEL returns an error.

If you specify TEXT for a nested column or for an attribute in a nested column, then CREATE MODEL returns an error.

NOPREP: Disables ADP for the attribute. When ADP is OFF, the NOPREP value is ignored.

You can specify NOPREP for a nested column, but not for an attribute in a nested column. If you specify NOPREP for an attribute in a nested column when ADP is on, then CREATE MODEL will return an error.

You can obtain information about a model by querying the Data Dictionary views.

```
ALL/USER/DBA_MINING_MODELS
ALL/USER/DBA_MINING_MODEL_ATTRIBUTES
ALL/USER/DBA_MINING_MODEL_SETTINGS
ALL/USER/DBA_MINING_MODEL_VIEWS
ALL/USER/DBA_MINING_MODEL_PARTITIONS
ALL/USER/DBA_MINING_MODEL_XFORMS
```

You can obtain information about model attributes by querying the model details through model views. Refer to *Oracle Machine Learning for SQL User's Guide*.

- 3. The naming rules for models are more restrictive than the naming rules for most database schema objects. A model name must satisfy the following additional requirements:
 - It must be 123 or fewer characters long.
 - It must be a nonquoted identifier. Oracle requires that nonquoted identifiers contain only alphanumeric characters, the underscore (_), dollar sign (\$), and pound sign (#); the initial character must be alphabetic. Oracle strongly discourages the use of the dollar sign and pound sign in nonquoted literals.

Naming requirements for schema objects are fully documented in *Oracle Database SQL Language Reference*.

4. To build a partitioned model, you must provide additional settings.

The setting for partitioning columns are as follows:

```
INSERT INTO settings_table VALUES ('ODMS_PARTITION_COLUMNS', 'GENDER,
AGE');
```

To set user-defined partition number for a model, the setting is as follows:

```
INSERT INTO settings table VALUES ('ODMS MAX PARTITIONS', '10');
```

The default value for maximum number of partitions is 1000.

5. By passing an xform_list to CREATE_MODEL, you can specify a list of transformations to be performed on the input data. If the PREP_AUTO setting is ON, the transformations are used in addition to the automatic transformations. If the PREP_AUTO setting is OFF, the specified transformations are the only ones implemented by the model. In both cases, transformation definitions are embedded in the model and run automatically whenever the



model is applied. See "Automatic Data Preparation". Other transforms that can be specified with xform_list include FORCE_IN. Refer to Oracle Machine Learning for SQL User's Guide.

Examples

The first example builds a classification model using the Support Vector Machine algorithm.

```
-- Create the settings table
CREATE TABLE svm model settings (
 setting name VARCHAR2(30),
 setting value VARCHAR2(30));
-- Populate the settings table
-- Specify SVM. By default, Naive Bayes is used for classification.
-- Specify ADP. By default, ADP is not used.
BEGIN
 INSERT INTO svm model_settings (setting_name, setting_value) VALUES
     (dbms data mining.algo name, dbms data mining.algo support vector machines);
 INSERT INTO svm model settings (setting name, setting value) VALUES
     (dbms data mining.prep auto, dbms data mining.prep auto on);
END;
-- Create the model using the specified settings
BEGIN
 DBMS DATA MINING.CREATE MODEL (
   model_name => 'svm_model',
   case id column name => 'cust id',
   target column name => 'affinity card',
   settings table name => 'svm model settings');
END;
```

You can display the model settings with the following query:

```
SELECT * FROM user_mining_model_settings
WHERE model name IN 'SVM MODEL';
```

MODEL_NAME	SETTING_NAME	SETTING_VALUE	SETTING
SVM MODEL	ALGO NAME	ALGO SUPPORT VECTOR MACHINES	INPUT
0111_110555			1111 01
SVM_MODEL	SVMS_STD_DEV	3.004524	DEFAULT
SVM_MODEL	PREP_AUTO	ON	INPUT
SVM_MODEL	SVMS_COMPLEXITY_FACTOR	1.887389	DEFAULT
SVM_MODEL	SVMS_KERNEL_FUNCTION	SVMS_LINEAR	DEFAULT
SVM_MODEL	SVMS_CONV_TOLERANCE	.001	DEFAULT

The following is an example of querying a model view instead of the older ${\tt GEL\ MODEL\ DETAILS\ SVM\ routine}.$

```
SELECT target_value, attribute_name, attribute_value, coefficient FROM
DM$VLSVM MODEL;
```

The second example creates an anomaly detection model. Anomaly detection uses SVM classification without a target. This example uses the same settings table created for the SVM classification model in the first example.

This query shows that the models created in these examples are the only ones in your schema.

```
SELECT model_name, mining_function, algorithm FROM user_mining_models;
```

MODEL_NAME	MINING_FUNCTION	ALGORITHM
SVM_MODEL	CLASSIFICATION	SUPPORT_VECTOR_MACHINES
ANOMALY_DETECT_MODEL	CLASSIFICATION	SUPPORT_VECTOR_MACHINES

This query shows that only the SVM classification model has a target.

CREATE_MODEL2 Procedure

The CREATE_MODEL2 procedure is an alternate procedure to the CREATE_MODEL procedure, which enables creating a model without extra persistence stages. In the CREATE_MODEL procedure, the input is a table or a view and if such an object is not already present, the user must create it. By using the CREATE_MODEL2 procedure, the user does not need to create such transient database objects.

Syntax



Parameters

Table 62-71 CREATE_MODEL2 Procedure Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then the current schema is used.
	See the Usage Notes, CREATE_MODEL Procedure for model naming restrictions.
mining_function	The machine learning function. Values are listed in DBMS_DATA_MINING — Machine Learning Function Settings.
data_query	A query which provides training data for building the model.
set_list	Specifies the SETTING_LIST
	SETTING_LIST is a table of CLOB index by VARCHAR2 (30); Where the index is the setting name and the CLOB is the setting value for that name.
case_id_column_name	Case identifier column in the build data.
target_column_name	For supervised models, the target column in the build data. ${\tt NULL}$ for unsupervised models.
xform_list	Refer to CREATE_MODEL Procedure.

Usage Notes

Refer to CREATE_MODEL Procedure for Usage Notes.

Examples

The following example uses the Support Vector Machine algorithm.

Create Model Using Registration Information

Create model function fetches the setting information from JSON object.

Usage Notes

If an algorithm is registered, user can create model using the registered algorithm name. Since all R scripts and default setting values are already registered, providing the value through the setting table is not necessary. This makes the use of this algorithm easier.

Examples

The first example builds a Classification model using the GLM algorithm.

```
CREATE TABLE GLM RDEMO SETTINGS CL (
   setting name VARCHAR2(30),
   setting value VARCHAR2(4000));
          INSERT INTO GLM RDEMO SETTINGS CL VALUES
           ('ALGO EXTENSIBLE LANG', 'R');
          INSERT INTO GLM RDEMO SETTINGS CL VALUES
           (dbms data mining.ralg registration algo name, 't1');
          INSERT INTO GLM RDEMO SETTINGS CL VALUES
          (dbms data mining.odms formula,
          'AGE + EDUCATION + HOUSEHOLD SIZE + OCCUPATION');
          INSERT INTO GLM RDEMO SETTINGS CL VALUES
            ('RALG PARAMETER FAMILY', 'binomial(logit)');
    END;
      BEGIN
            DBMS_DATA_MINING.CREATE_MODEL(
            model_name => 'GLM_RDEMO_CLASSIFICATION',
mining_function => dbms_data_mining.classification,
data_table_name => 'mining_data_build_v',
case_id_column_name => 'CUST_ID',
target_column_name => 'AFFINITY_CARD',
settings_table_name => 'GLM_RDEMO_SETTINGS_CL');
       END:
```

DROP ALGORITHM Procedure

This function is used to drop the registered algorithm information.

Syntax

```
DBMS_DATA_MINING.DROP_ALGORITHM (algorithm_name IN VARCHAR2(30), cascade IN BOOLEAN default FALSE)
```

Table 62-72 DROP_ALGORITHM Procedure Parameters



Table 62-72 (Cont.) DROP_ALGORITHM Procedure Parameters

Parameter	Description
cascade	If the cascade option is TRUE, all the models with this algorithms are forced to drop. There after, the algorithm is dropped. The default value is FALSE.

Usage Note

- To drop a machine learning model, you must be the owner or you must have the RQADMIN privilege. See Oracle Machine Learning for SQL User's Guide for information about privileges for machine learning.
- Make sure a model is not built on the algorithm, then drop the algorithm from the system table.
- If you try to drop an algorithm with a model built on it, then an error is displayed.

DROP_PARTITION Procedure

Syntax

Parameters

Table 62-73 DROP PARTITION Procedure Parameters

Parameters	Description
model_name	Name of the machine learning model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Name of the partition that must be dropped.

DROP_MODEL Procedure

This procedure deletes the specified machine learning model.

Syntax

```
DBMS_DATA_MINING.DROP_MODEL (model_name IN VARCHAR2, force IN BOOLEAN DEFAULT FALSE);
```

Parameters

Table 62-74 DROP_MODEL Procedure Parameters

Parameter	Description
model_name	Name of the machine learning model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.

Table 62-74 (Cont.) DROP_MODEL Procedure Parameters

Parameter	Description
force	Forces the machine learning model to be dropped even if it is invalid. A machine learning model may be invalid if a serious system error interrupted the model build process.

Usage Note

To drop a machine learning model, you must be the owner or you must have the DROP ANY MINING MODEL privilege. See *Oracle Data Mining User's Guide* for information about privileges for Oracle Machine Learning for SQL.

Example

You can use the following command to delete a valid machine learning model named nb sh clas sample that exists in your schema.

```
BEGIN
   DBMS_DATA_MINING.DROP_MODEL(model_name => 'nb_sh_clas_sample');
END;
//
```

EXPORT_MODEL Procedure

This procedure exports the specified machine learning models to a dump file set.

To import the models from the dump file set, use the IMPORT_MODEL Procedure. EXPORT MODEL and IMPORT MODEL use Oracle Data Pump technology.

When Oracle Data Pump is used to export/import an entire schema or database, the machine learning models in the schema or database are included. However, <code>EXPORT_MODEL</code> and <code>IMPORT_MODEL</code> are the only utilities that support the export/import of individual models.



Oracle Database Utilities for information about Oracle Data Pump

Oracle Machine Learning for SQL User's Guide for more information about exporting and importing machine learning models

Syntax

```
DBMS_DATA_MINING.EXPORT_MODEL (
filename IN VARCHAR2,
directory IN VARCHAR2,
model_filter IN VARCHAR2 DEFAULT NULL,
filesize IN VARCHAR2 DEFAULT NULL,
operation IN VARCHAR2 DEFAULT NULL,
remote_link IN VARCHAR2 DEFAULT NULL,
jobname IN VARCHAR2 DEFAULT NULL);
```



Table 62-75 EXPORT_MODEL Procedure Parameters

Parameter	Description
filename	Name of the dump file set to which the models should be exported. The name must be unique within the schema.
	The dump file set can contain one or more files. The number of files in a dump file set is determined by the size of the models being exported (both metadata and data) and a specified or estimated maximum file size. You can specify the file size in the filesize parameter, or you can use the operation parameter to cause Oracle Data Pump to estimate the file size. If the size of the models to export is greater than the maximum file size, one or more additional files are created.
	When the export operation completes successfully, the name of the dump file set is automatically expanded to $filename01.dmp$, even if there is only one file in the dump set. If there are additional files, they are named sequentially as $filename02.dmp$, $filename03.dmp$, and so forth.
directory	Name of a pre-defined directory object that specifies where the dump file set should be created.
	The exporting user must have read/write privileges on the directory object and on the file system directory that it identifies.
	See Oracle Database SQL Language Reference for information about directory objects.
model_filter	Optional parameter that specifies which model or models to export. If you do not specify a value for model_filter, all models in the schema are exported. You can also specify NULL (the default) or 'ALL' to export all models.
	You can export individual models by name and groups of models based on machine learning function or algorithm. For instance, you could export all regression models or all Naive Bayes models. Examples are provided in Table 62-76.
filesize	Optional parameter that specifies the maximum size of a file in the dump file set. The size may be specified in bytes, kilobytes (K), megabytes (M), or gigabytes (G). The default size is 50 MB.
	If the size of the models to export is larger than filesize, one or more additional files are created within the dump set. See the description of the filename parameter for more information.
operation	Optional parameter that specifies whether or not to estimate the size of the files in the dump set. By default the size is not estimated and the value of the filesize parameter determines the size of the files.
	You can specify either of the following values for operation:
	 'EXPORT' — Export all or the specified models. (Default)
	 'ESTIMATE' — Estimate the size of the exporting models.
remote_link	Optional parameter that specifies the name of a database link to a remote system. The default value is NULL. A database link is a schema object in a local database that enables access to objects in a remote database. When you specify a value for remote_link, you can export the models in the remote database. The EXP_FULL_DATABASE role is required for exporting the remote models. The EXP_FULL_DATABASE privilege, the CREATE DATABASE LINK privilege, and other
	privileges may also be required.



Table 62-75 (Cont.) EXPORT_MODEL Procedure Parameters

Parameter	Description
jobname	Optional parameter that specifies the name of the export job. By default, the name has the form <code>username_exp_nnnn</code> , where <code>nnnn</code> is a number. For example, a job name in the <code>SCOTT schema might be SCOTT_exp_134</code> .
	If you specify a job name, it must be unique within the schema. The maximum length of the job name is 30 characters.
	A log file for the export job, named $jobname.log$, is created in the same directory as the dump file set.

Usage Notes

The <code>model_filter</code> parameter specifies which models to export. You can list the models by name, or you can specify all models that have the same machine learning function or algorithm. You can query the <code>USER MINING MODELS</code> view to list the models in your schema.

SQL> describe user_mining_models Name	Null?	Туре
MODEL_NAME	NOT NULL	VARCHAR2(30)
MINING_FUNCTION		VARCHAR2(30)
ALGORITHM		VARCHAR2(30)
CREATION_DATE	NOT NULL	DATE
BUILD_DURATION		NUMBER
MODEL_SIZE		NUMBER
COMMENTS		VARCHAR2 (4000)

Examples of model filters are provided in Table 62-76.

Table 62-76 Sample Values for the Model Filter Parameter

Sample Value	Meaning
'mymodel'	Export the model named mymodel
'name= ''mymodel'''	Export the model named mymodel
<pre>'name IN (''mymodel2'',''mymodel3'')'</pre>	Export the models named mymodel2 and mymodel3
'ALGORITHM_NAME = ''NAIVE_BAYES'''	Export all Naive Bayes models. See Table 62-5 for a list of algorithm names.
'FUNCTION_NAME =''CLASSIFICATION'''	Export all classification models. See Table 62-3 for a list of machine learning functions.

Examples

1. The following statement exports all the models in the oml_user3 schema to a dump file set called models_out in the directory \$ORACLE_HOME/rdbms/log. This directory is mapped to a directory object called DATA_PUMP_DIR. The oml_user3 user has read/write access to the directory and to the directory object.

```
SQL>execute dbms_data_mining.export_model ('models_out', 'DATA_PUMP_DIR');
```

You can exit SQL*Plus and list the resulting dump file and log file.

```
SQL>EXIT
>cd $ORACLE_HOME/rdbms/log
>ls
>oml user3 exp 1027.log models out01.dmp
```

2. The following example uses the same directory object and is run by the same user. This example exports the models called NMF_SH_SAMPLE and SVMR_SH_REGR_SAMPLE to a different dump file set in the same directory.

3. The following examples show how to export models with specific algorithm and machine learning function names.

EXPORT_SERMODEL Procedure

This procedure exports the model in a serialized format so that they can be moved to another platform for scoring.

When exporting a model in serialized format, the user must pass in an empty BLOB locator and specify the model name to be exported. If the model is partitioned, the user can optionally select an individual partition to export, otherwise all partitions are exported. The returned BLOB contains the content that can be deployed.

Syntax

Parameters

Table 62-77 EXPORT_SERMODEL Procedure Parameters

Parameter	Description
model data	Provides serialized model data.
model_name	Name of the machine learning model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Name of the partition that must be exported.



Examples

The following statement exports all of the models in a serialized format.

```
DECLARE
  v_blob blob;
BEGIN
  dbms_lob.createtemporary(v_blob, FALSE);
  dbms_data_mining.export_sermodel(v_blob, 'MY_MODEL');
-- save v_blob somewhere (e.g., bfile, etc.)
  dbms_lob.freetemporary(v_blob);
END;
//
```

See Also:

Oracle Machine Learning for SQL User's Guide for more information about exporting and importing machine learning models

FETCH_JSON_SCHEMA Procedure

User can fetch and read JSON schema from the <code>ALL_MINING_ALGORITHMS</code> view. This function returns the pre-registered JSON schema for R extensible algorithms.

Syntax

DBMS_DATA_MINING.FETCH_JSON_SCHEMA RETURN CLOB;

Parameters

Table 62-78 FETCH_JSON_SCHEMA Procedure Parameters

Parameter	Description	
RETURN	This function returns the pre-registered JSON schema for R extensibility.	
	The default value is CLOB.	

Usage Note

If a user wants to register a new algorithm using the algorithm registration function, they must fetch and follow the pre-registered JSON schema using this function, when they create the required JSON object metadata, and then pass it to the registration function.

GET_ASSOCIATION_RULES Function

The GET_ASSOCIATION_RULES function returns the rules produced by an association model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

You can specify filtering criteria to <code>GET_ASSOCIATION_RULES</code> to return a subset of the rules. Filtering criteria can improve the performance of the table function. If the number of rules is large, then the greatest performance improvement will result from specifying the <code>topn</code> parameter.

Syntax

Table 62-79 GET_ASSOCIATION_RULES Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
	This is the only required parameter of <code>GET_ASSOCIATION_RULES</code> . All other parameters specify optional filters on the rules to return.
topn	Returns the <i>n</i> top rules ordered by confidence and then support, both descending. If you specify a sort order, then the top <i>n</i> rules are derived after the sort is performed.
	If topn is specified and no maximum or minimum rule length is specified, then the only columns allowed in the sort order are RULE_CONFIDENCE and RULE_SUPPORT. If topn is specified and a maximum or minimum rule length is specified, then RULE_CONFIDENCE, RULE_SUPPORT, and NUMBER_OF_ITEMS are allowed in the sort order.
rule_id	Identifier of the rule to return. If you specify a value for rule_id, do not specify values for the other filtering parameters.
min_confidence	Returns the rules with confidence greater than or equal to this number.
min_support	Returns the rules with support greater than or equal to this number.
max_rule_length	Returns the rules with a length less than or equal to this number.
	Rule length refers to the number of items in the rule (See NUMBER_OF_ITEMS in Table 62-80). For example, in the rule A=>B (if A, then B), the number of items is 2.
	If max_rule_length is specified, then the NUMBER_OF_ITEMS column is permitted in the sort order.
min_rule_length	Returns the rules with a length greater than or equal to this number. See max_rule_length for a description of rule length.
	If min_rule_length is specified, then the NUMBER_OF_ITEMS column is permitted in the sort order.

Table 62-79 (Cont.) GET_ASSOCIATION_RULES Function Parameters

Parameter	Description
sort_order	Sorts the rules by the values in one or more of the returned columns. Specify one or more column names, each followed by ASC for ascending order or DESC for descending order. (See Table 62-80 for the column names.)
	For example, to sort the result set in descending order first by the NUMBER_OF_ITEMS column, then by the RULE_CONFIDENCE column, you must specify:
	ORA_MINING_VARCHAR2_NT('NUMBER_OF_ITEMS DESC', 'RULE_CONFIDENCE DESC')
	If you specify topn, the results will vary depending on the sort order.
	By default, the results are sorted by Confidence in descending order, then by Support in descending order.
antecedent_items	Returns the rules with these items in the antecedent.
consequent_items	Returns the rules with this item in the consequent.
min_lift	Returns the rules with lift greater than or equal to this number.
partition_name	Specifies a partition in a partitioned model.

Return Values

The object type returned by <code>GET_ASSOCIATION_RULES</code> is described in Table 62-80. For descriptions of each field, see the Usage Notes.

Table 62-80 GET_ASSOCIATION RULES Function Return Values

Return Value	Description	
DM_RULES	A set of rows of type <code>DM_RULE</code> . The rows have the following columns:	
	_	REDICATES, REDICATES, ER, ER, ER, ER, ER, ER, ER,
DM_PREDICATES	DM_PREDICATES. The rows, of (attribute_name attribute_subname conditional_operator attribute_num_value	NUMBER, VARCHAR2(4000), NUMBER,



Usage Notes

- 1. This table function pipes out rows of type <code>DM_RULES</code>. For information on machine learning data types and piped output from table functions, see "Datatypes".
- 2. The columns returned by GET_ASSOCIATION_RULES are described as follows:

Column in DM_RULES	Description
rule_id	Unique identifier of the rule
antecedent	The independent condition in the rule. When this condition exists, the dependent condition in the consequent also exists.
	The condition is a combination of attribute values called a predicate (DM_PREDICATE). The predicate specifies a condition for each attribute. The condition may specify equality (=), inequality (<>), greater than (>), less than (<), greater than or equal to (>=), or less than or equal to (<=) a given value.
	Support and Confidence for each attribute condition in the antecedent is returned in the predicate. Support is the number of transactions that satisfy the antecedent. Confidence is the likelihood that a transaction will satisfy the antecedent.
	Note: The occurrence of the attribute as a DM_PREDICATE indicates the presence of the item in the transaction. The actual value for attribute_num_value or attribute_str_value is meaningless. For example, the following predicate indicates that 'Mouse Pad' is present in the transaction <i>even though</i> the attribute value is NULL.
	<pre>DM_PREDICATE('PROD_NAME',</pre>
consequent	The dependent condition in the rule. This condition exists when the antecedent exists.
	The consequent, like the antecedent, is a predicate (DM_PREDICATE).
	Support and confidence for each attribute condition in the consequent is returned in the predicate. Support is the number of transactions that satisfy the consequent. Confidence is the likelihood that a transaction will satisfy the consequent.
rule_support	The number of transactions that satisfy the rule.
rule_confidence	The likelihood of a transaction satisfying the rule.
rule_lift	The degree of improvement in the prediction over random chance when the rule is satisfied.
antecedent_support	The ratio of the number of transactions that satisfy the antecedent to the total number of transactions.
consequent_support	The ratio of the number of transactions that satisfy the consequent to the total number of transactions.
number_of_items	The total number of attributes referenced in the antecedent and consequent of the rule.

Examples

The following example demonstrates an association model build followed by several invocations of the GET ASSOCIATION RULES table function:

```
-- prepare a settings table to override default settings CREATE TABLE market_settings AS SELECT \star
```

```
FROM TABLE (DBMS DATA MINING.GET DEFAULT SETTINGS)
WHERE setting_name LIKE 'ASSO_%';
-- update the value of the minimum confidence
UPDATE market settings
   SET setting value = TO CHAR(0.081)
 WHERE setting name = DBMS DATA MINING.asso min confidence;
-- build an AR model
DBMS DATA MINING. CREATE MODEL (
 model name => 'market model',
  function => DBMS DATA MINING.ASSOCIATION,
 data table name => 'market build',
 case id column name => 'item id',
 target column name => NULL,
  settings table name => 'market settings');
END:
-- View the (unformatted) rules
SELECT rule id, antecedent, consequent, rule support,
       rule confidence
  FROM TABLE (DBMS DATA MINING.GET ASSOCIATION RULES ('market model'));
```

In the previous example, you view all rules. To view just the top 20 rules, use the following statement.

The following query uses the association model AR SH SAMPLE.

The query returns three rules, shown as follows:

```
13 DM PREDICATES (
      DM_PREDICATE('CUSTPRODS', 'Mouse Pad', '= ', 1, NULL, NULL, NULL),
      DM_PREDICATE('CUSTPRODS', 'Standard Mouse', '= ', 1, NULL, NULL, NULL))
   DM PREDICATES (
      DM PREDICATE('CUSTPRODS', 'Extension Cable', '= ', 1, NULL, NULL, NULL))
   .15532
           .84393 2.7075
                               .18404 .3117
11 DM PREDICATES (
     DM PREDICATE ('CUSTPRODS', 'Standard Mouse', '= ', 1, NULL, NULL, NULL))
   DM PREDICATES (
      DM PREDICATE('CUSTPRODS', 'Extension Cable', '= ', 1, NULL, NULL, NULL))
   .18085
              .56291 1.8059
                                 .32128 .3117 1
   DM PREDICATES (
      DM PREDICATE('CUSTPRODS', 'Mouse Pad', '= ', 1, NULL, NULL, NULL))
   DM PREDICATES (
      DM PREDICATE('CUSTPRODS', 'Extension Cable', '= ', 1, NULL, NULL, NULL))
                               .32234 .3117
     .17766 .55116 1.7682
```



Table 62-80 for the DM RULE column data types.

GET_FREQUENT_ITEMSETS Function

The GET_FREQUENT_ITEMSETS function returns a set of rows that represent the frequent itemsets from an association model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead..

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

For a detailed description of frequent itemsets, consult *Oracle Machine Learning for SQL Concepts*.

Syntax

Parameters

Table 62-81 GET_FREQUENT_ITEMSETS Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
topn	When not \mathtt{NULL} , return the top n rows ordered by support in descending order
max_itemset_length	Maximum length of an item set.
partition_name	Specifies a partition in a partitioned model.



The ${\tt partition_name}$ columns applies only when the model is partitioned.



Return Values

Table 62-82 GET_FREQUENT_ITEMSETS Function Return Values

Return Value Description DM_ITEMSETS A set of rows of type DM_ITEMSET. The rows have the following columns: (partition_name VARCHAR2(128) itemsets_id NUMBER, items DM_ITEMS, support NUMBER, number_of_items NUMBER)



The partition_name columns applies only when the model is partitioned.

The items column returns a nested table of type ${\tt DM_ITEMS}$. The rows have type ${\tt DM}$ ITEM:

```
(attribute_name VARCHAR2(4000), attribute_subname VARCHAR2(4000), attribute_num_value NUMBER, attribute_str_value VARCHAR2(4000))
```

Usage Notes

This table function pipes out rows of type DM_ITEMSETS. For information on machine learning data types and piped output from table functions, see "Data Types".

Examples

The following example demonstrates an association model build followed by an invocation of GET_FREQUENT_ITEMSETS table function from Oracle SQL.



```
settings_table_name => 'market_settings');
END;
/-- View the (unformatted) Itemsets from SQL*Plus
SELECT itemset_id, items, support, number_of_items
FROM TABLE (DBMS_DATA_MINING.GET_FREQUENT_ITEMSETS('market_model'));
```

In the example above, you view all itemsets. To view just the top 20 itemsets, use the following statement:

```
-- View the top 20 (unformatted) Itemsets from SQL*Plus
SELECT itemset_id, items, support, number_of_items
FROM TABLE(DBMS DATA MINING.GET FREQUENT ITEMSETS('market model', 20));
```

GET_MODEL_COST_MATRIX Function

The \mathtt{GET}_{-}^* interfaces are replaced by model views, and Oracle recommends that users leverage the views instead.

The GET_MODEL_COST_MATRIX function is replaced by the DM\$VC prefixed view, Scoring Cost Matrix. The cost matrix used when building a Decision Tree is made available by the DM\$VM prefixed view, Decision Tree build cost matrix.

Refer to Model Detail View for Classification Algorithm.

The <code>GET_MODEL_COST_MATRIX</code> function returns the rows of a cost matrix associated with the specified model.

By default, this function returns the scoring cost matrix that was added to the model with the ADD_COST_MATRIX procedure. If you wish to obtain the cost matrix used to create a model, specify cost matrix type create as the matrix type. See Table 62-83.

See also ADD_COST_MATRIX Procedure.

Syntax

Parameters

Table 62-83 GET MODEL COST MATRIX Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
matrix_type	The type of cost matrix.
	COST_MATRIX_TYPE_SCORE — cost matrix used for scoring. (Default.)
	COST_MATRIX_TYPE_CREATE — cost matrix used to create the model (Decision Tree only).
partition_name	Name of the partition in a partitioned model

Return Values

Table 62-84 GET MODEL COST MATRIX Function Return Values

Return Value	Description	
DM_COST_MATRIX	A set of rows of type <code>DM_COST_ELEMENT</code> . The rows have the following columns:	
	actual VARCHAR2(4000), NUMBER, predicted VARCHAR2(4000), cost NUMBER)	

Usage Notes

Only Decision Tree models can be built with a cost matrix. If you want to build a Decision Tree model with a cost matrix, specify the cost matrix table name in the <code>CLAS_COST_TABLE_NAME</code> setting in the settings table for the model. See Table 62-7.

The cost matrix used to create a Decision Tree model becomes the default scoring matrix for the model. If you want to specify different costs for scoring, you can use the REMOVE_COST_MATRIX procedure to remove the cost matrix and the ADD_COST_MATRIX procedure to add a new one.

The GET_MODEL_COST_MATRIX may return either the build or scoring cost matrix defined for a model or model partition.

If you do not specify a partitioned model name, then an error is displayed.

Example

This example returns the scoring cost matrix associated with the Naive Bayes model NB_SH_CLAS_SAMPLE.

```
column actual format a10
column predicted format a10
SELECT *
    FROM TABLE(dbms_data_mining.get_model_cost_matrix('nb_sh_clas_sample'))
    ORDER BY predicted, actual;
```

ACTUAL	PREDICTED	COST
0	0	.00
1	0	.75
0	1	.25
1	1	.00

GET_MODEL_DETAILS_AI Function

The GET_MODEL_DETAILS_AI function returns a set of rows that provide the details of an attribute importance model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

Syntax

partition_name IN VARCHAR2 DEFAULT NULL)
RETURN dm_ranked_attributes pipelined;

Parameters

Table 62-85 GET_MODEL_DETAILS_AI Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model.

Return Values

Table 62-86 GET_MODEL_DETAILS_AI Function Return Values

Return Value	Description	
DM_RANKED_ATTRIBUTES	A set of rows of type <code>DM_RANKED_ATTRIBUTE</code> . The rows have the following columns:	
	<pre>(attribute_name attribute_subname importance_value rank</pre>	VARCHAR2 (4000, VARCHAR2 (4000), NUMBER, NUMBER (38))

Examples

The following example returns model details for the attribute importance model AI_SH_sample, which was created by the sample program dmaidemo.sql.

```
SELECT attribute_name, importance_value, rank
   FROM TABLE(DBMS_DATA_MINING.GET_MODEL_DETAILS_AI('AI_SH_sample'))
   ORDER BY RANK;
```

ATTRIBUTE_NAME	IMPORTANCE_VALUE	RANK
HOUSEHOLD SIZE	.151685183	1
CUST_MARITAL_STATUS	.145294546	2
YRS RESIDENCE	.07838928	3
AGE	.075027496	4
Y_BOX_GAMES	.063039952	5
EDUCATION	.059605314	6
HOME_THEATER_PACKAGE	.056458722	7
OCCUPATION	.054652937	8
CUST_GENDER	.035264741	9
BOOKKEEPING_APPLICATION	.019204751	10
PRINTER_SUPPLIES	0	11
OS_DOC_SET_KANJI	00050013	12
FLAT_PANEL_MONITOR	00509564	13
BULK_PACK_DISKETTES	00540822	14
COUNTRY_NAME	01201116	15
CUST_INCOME_LEVEL	03951311	16



GET_MODEL_DETAILS_EM Function

The <code>GET_MODEL_DETAILS_EM</code> function returns a set of rows that provide statistics about the clusters produced by an expectation maximization model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

By default, the EM algorithm groups components into high-level clusters, and <code>GET_MODEL_DETAILS_EM</code> returns only the high-level clusters with their hierarchies. Alternatively, you can configure EM model to disable the grouping of components into high-level clusters. In this case, <code>GET_MODEL_DETAILS_EM</code> returns the components themselves as clusters with their hierarchies. See Table 62-12.

Syntax

Parameters

Table 62-87 GET MODEL DETAILS EM Function Parameters

Parameter	Description	
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.	
cluster_id	The ID of a cluster in the model. When a valid cluster ID is specified, only the details of this cluster are returned. Otherwise, the details for all clusters are returned.	
attribute	The name of an attribute. When a valid attribute name is specified, only the details of this attribute are returned. Otherwise, the details for all attributes are returned	
centroid	This parameter accepts the following values:	
	 1: Details about centroids are returned (default) 	
	 0: Details about centroids are not returned 	
histogram	This parameter accepts the following values:	
	 1: Details about histograms are returned (default) 	
	 0: Details about histograms are not returned 	
rules	This parameter accepts the following values:	
	 2: Details about rules are returned (default) 	
	 1: Rule summaries are returned 	
	 0: No information about rules is returned 	

Parameter	Description
attribute_subname	The name of a nested attribute. The full name of a nested attribute has the form:
	attribute_name.attribute_subname
	where attribute_name is the name of the column and attribute_subname is the name of the nested attribute in that column. If the attribute is not nested, then attribute_subname is null.
topn_attributes	Restricts the number of attributes returned in the centroid, histogram, and rules objects. Only the n attributes with the highest confidence values in the rules are returned.
	If the number of attributes included in the rules is less than $topn$, then, up to n additional attributes in alphabetical order are returned.
	If both the attribute and topn_attributes parameters are specified, then topn_attributes is ignored.
partition_name	Specifies a partition in a partitioned model.

Usage Notes

- 1. For information on Oracle Machine Learning for SQL data types and return values for Clustering algorithms piped output from table functions, see "Data Types".
- 2. GET_MODEL_DETAILS functions preserve model transparency by automatically reversing the transformations applied during the build process. Thus the attributes returned in the model details are the original attributes (or a close approximation of the original attributes) used to build the model.
- 3. When cluster statistics are disabled (EMCS_CLUSTER_STATISTICS is set to EMCS_CLUS_STATS_DISABLE), GET_MODEL_DETAILS_EM does not return centroids, histograms, or rules. Only taxonomy (hierarchy) and cluster counts are returned.
- **4.** When the partition_name is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

GET MODEL DETAILS EM COMP Function

he <code>GET_MODEL_DETAILS_EM_COMP</code> table function returns a set of rows that provide details about the parameters of an expectation maximization model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

Syntax



Parameters

Table 62-88 GET_MODEL_DETAILS_EM_COMP Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model to retrieve details for.

Return Values

Table 62-89 GET_MODEL_DETAILS_EM_COMP Function Return Values

Return Value	Description	
DM_EM_COMPONENT_SET	A set of rows of type <code>DM_EM_COMPONENT</code> . The rows have the following columns:	
	<pre>(info_type component_id cluster_id</pre>	VARCHAR2(30), NUMBER, NUMBER,
	<pre>attribute_name covariate_name attribute_value</pre>	VARCHAR2 (4000), VARCHAR2 (4000), VARCHAR2 (4000),
	value	NUMBER)

Usage Notes

1. This table function pipes out rows of type <code>DM_EM_COMPONENT</code>. For information on Oracle Machine Learning for SQL data types and piped output from table functions, see "Data Types".

The columns in each row returned by $\texttt{GET}_\texttt{MODEL}_\texttt{DETAILS}_\texttt{EM}_\texttt{COMP}$ are described as follows:

Column in DM_EM_COMPONENT	Description	
info_type	The type of information in the row. The following information types are supported:	
	• cluster	
	• prior	
	• mean	
	• covariance	
	 frequency 	
component_id	Unique identifier of a component	
cluster_id	Unique identifier of the high-level leaf cluster for each component	
attribute_name	Name of an original attribute or a derived feature ID. The derived feature ID is used in models built on data with nested columns. The derived feature definitions can be obtained from the GET_MODEL_DETAILS_EM_PROJ Function.	



Column in DM_EM_COMPONENT	Description	
covariate_name	Name of an original attribute or a derived feature ID used in variance/covariance definition	
attribute_value	Categorical value or bin interval for binned numerical attributes	
value	Encodes different information depending on the value of info_type, as follows:	
	 cluster — The value field is NULL 	
	 prior — The value field returns the component prior 	
	 mean — The value field returns the mean of the attribute specified in attribute name 	
	 covariance — The value field returns the covariance of the attributes specified in attribute name and covariate name. Using the same attribute in attribute name and covariate name, returns the variance. 	
	 frequency— The value field returns the multivalued Bernoulli frequency parameter for the attribute/value combination specified by attribute_name and attribute_value 	
	See Usage Note 2 for details.	

The following table shows which fields are used for each info_type. The blank cells represent NULLS.

info_type	component_i d	cluster_id	attribute_ name	covariate_n ame	attribute_val ue	value
cluster	Х	Χ				
prior	X	Χ				Χ
mean	X	Χ	Χ			Χ
covariance	X	Χ	Χ	Χ		Χ
frequency	X	Χ	Χ		X	X

- 3. GET_MODEL_DETAILS functions preserve model transparency by automatically reversing the transformations applied during the build process. Thus the attributes returned in the model details are the original attributes (or a close approximation of the original attributes) used to build the model.
- **4.** When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

GET_MODEL_DETAILS_EM_PROJ Function

The GET_MODEL_DETAILS_EM_PROJ function returns a set of rows that provide statistics about the projections produced by an expectation maximization model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

Syntax

```
DBMS_DATA_MINING.get_model_details_em_proj(
          model name IN VARCHAR2,
```

partition_name IN VARCHAR2 DEFAULT NULL)
RETURN DM EM PROJECTION SET PIPELINED;

Parameters

Table 62-90 GET_MODEL_DETAILS_EM_PROJ Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model

Return Values

Table 62-91 GET_MODEL_DETAILS_EM_PROJ Function Return Values

Return Value	Description	
DM_EM_PROJECTION_SET	A set of rows of type DM columns:	_EM_PROJECTION. The rows have the following
	(feature_name	VARCHAR2(4000),
	attribute_name	VARCHAR2(4000),
	attribute_subname	VARCHAR2(4000),
	attribute_value	VARCHAR2(4000),
	coefficient	NUMBER)
	See Usage Notes for de	tails.

Usage Notes

1. This table function pipes out rows of type DM_EM_PROJECTION. For information on machine learning data types and piped output from table functions, see "Datatypes".

The columns in each row returned by $\texttt{GET}_{\texttt{MODEL}}$ $\texttt{DETAILS}_{\texttt{EM}}$ PROJ are described as follows:

Column in DM_EM_PROJECTION	Description
feature_name	Name of a derived feature. The feature maps to the attribute_name returned by the GET_MODEL_DETAILS_EM Function.
attribute_name	Name of a column in the build data
attribute_subname	Subname in a nested column
attribute_value	Categorical value
coefficient	Projection coefficient. The representation is sparse; only the non-zero coefficients are returned.

GET_MODEL_DETAILS functions preserve model transparency by automatically reversing the
transformations applied during the build process. Thus the attributes returned in the model
details are the original attributes (or a close approximation of the original attributes) used to
build the model.

The coefficients are related to the transformed, not the original, attributes. When returned directly with the model details, the coefficients may not provide meaningful information.

3. When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

Related Topics

Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_GLM Function

The GET_MODEL_DETAILS_GLM function returns the coefficient statistics for a generalized linear model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

The same set of statistics is returned for both linear and logistic regression, but statistics that do not apply to the machine learning function are returned as NULL. For more details, see the Usage Notes.

Syntax

Parameters

Table 62-92 GET_MODEL_DETAILS_GLM Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model



Return Values

Table 62-93 GET_MODEL_DETAILS_GLM Return Values

Return Value	Description	
DM_GLM_COEFF_SET	A set of rows of type DM_GL columns:	M_COEFF. The rows have the following
	<pre>(class attribute_name attribute_subname attribute_value feature_expression coefficient std_error test_statistic p_value VIF std_coefficient lower_coeff_limit upper_coeff_limit exp_lower_coeff_limit exp_upper_coeff_limit</pre>	VARCHAR2 (4000), VARCHAR2 (4000), NUMBER, NUMBER, NUMBER, NUMBER, NUMBER, NUMBER, NUMBER, NUMBER, NUMBER, BINARY_DOUBLE, BINARY_DOUBLE,

GET_MODEL_DETAILS_GLM returns a row of statistics for each attribute and one extra row for the intercept, which is identified by a null value in the attribute name. Each row has the DM GLM COEFF data type. The statistics are described in Table 62-94.

Table 62-94 DM_GLM_COEFF Data Type Description

Column	Description
class	The non-reference target class for logistic regression. The model is built to predict the probability of this class.
	The other class (the reference class) is specified in the model setting GLMS_REFERENCE_CLASS_NAME. See Table 62-19.
	For Linear Regression, class is null.
attribute_name	The attribute name when there is no subname, or first part of the attribute name when there is a subname. The value of attribute_name is also the name of the column in the case table that is the source for this attribute.
	For the intercept, attribute_name is null. Intercepts are equivalent to the bias term in SVM models.
attribute_subname	The name of an attribute in a nested table. The full name of a nested attribute has the form:
	attribute_name.attribute_subname
	where <code>attribute_name</code> is the name of the nested column in the case table that is the source for this attribute.
	If the attribute is not nested, then attribute_subname is null. If the attribute is an intercept, then both the attribute_name and the attribute_subname are null.



Table 62-94 (Cont.) DM_GLM_COEFF Data Type Description

Column	Description
attribute_value	The value of the attribute (categorical attribute only).
	For numeric attributes, attribute_value is null.
feature_expression	The feature name constructed by the algorithm when feature generation is enabled and higher-order features are found. If feature selection is not enabled, then the feature name is simply the fully-qualified attribute name (attribute_name.attribute_subname if the attribute is in a nested column).
	For categorical attributes, the algorithm constructs a feature name that has the following form:
	fully-qualified_attribute_name.attribute_value
	For numeric attributes, the algorithm constructs a name for the higher- order feature by taking the product of the resulting values:
	(attrib1)*(attrib2))*
	where attrib1 and attrib2 are fully-qualified attribute names.
coefficient	The linear coefficient estimate.
std_error	Standard error of the coefficient estimate.
test_statistic	For linear regression, the t-value of the coefficient estimate. For logistic regression, the Wald chi-square value of the coefficient estimate.
p-value	Probability of the test_statistic. Used to analyze the significance of specific attributes in the model.
VIF	Variance Inflation Factor. The value is zero for the intercept. For logistic regression, \mathbb{VIF} is null. VIF is not computed if the solver is Cholesky.
std_coefficient	Standardized estimate of the coefficient.
lower_coeff_limit	Lower confidence bound of the coefficient.
upper_coeff_limit	Upper confidence bound of the coefficient.
exp_coefficient	Exponentiated coefficient for logistic regression. For linear regression, exp_coefficient is null.
exp_lower_coeff_limit	Exponentiated coefficient for lower confidence bound of the coefficient for logistic regression. For linear regression, exp_lower_coeff_limit is null.
exp_upper_coeff_limit	Exponentiated coefficient for upper confidence bound of the coefficient for logistic regression. For linear regression, exp_lower_coeff_limit is null.

Usage Notes

Not all statistics are necessarily returned for each coefficient. Statistics will be null if:

- They do not apply to the machine learning function. For example, exp_coefficient does not apply to linear regression.
- They cannot be computed from a theoretical standpoint. For information on ridge regression, see Table 62-19.
- They cannot be computed because of limitations in system resources.
- Their values would be infinity.

 When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

Examples

The following example returns some of the model details for the GLM regression model GLMR SH Regr sample.

```
SET line 120
SET pages 99
column attribute_name format a30
column attribute_subname format a20
column attribute_value format a20
col coefficient format 990.9999
col std_error format 990.9999
SQL> SELECT * FROM
(SELECT attribute_name, attribute_value, coefficient, std_error
    FROM DM$VDGLMR_SH_REGR_SAMPLE order by 1,2)
WHERE rownum < 11;
```

ATTRIBUTE_NAME	ATTRIBUTE_VALUE	COEFFICIENT	STD_ERROR
AFFINITY_CARD		-0.5797	0.5283
BOOKKEEPING_APPLICATION		-0.4689	3.8872
BULK_PACK_DISKETTES		-0.9819	2.5430
COUNTRY_NAME	Argentina	-1.2020	1.1876
COUNTRY_NAME	Australia	-0.0071	5.1146
COUNTRY NAME	Brazil	5.2931	1.9233
COUNTRY NAME	Canada	4.0191	2.4108
COUNTRY_NAME	China	0.8706	3.5889
COUNTRY_NAME	Denmark	-2.9822	3.1803
COUNTRY_NAME	France	-1.1044	7.1811

Related Topics

Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_GLOBAL Function

The <code>GET_MODEL_DETAILS_GLOBAL</code> function returns statistics about the model as a whole. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

Global details are available for Generalized Linear Models, Association Rules, Singular Value Decomposition, and Expectation Maximization. There are new Global model views which show global information for all algorithms. Oracle recommends that users leverage the views instead. Refer to Model Details View Global.

Syntax



Parameters

Table 62-95 GET_MODEL_DETAILS_GLOBAL Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model.

Return Values

Table 62-96 GET_MODEL_DETAILS_GLOBAL Function Return Values

Return Value	Description	
DM_MODEL_GLOBAL_DETAILS	A collection of rows of type <code>DM_MODEL_GLOBAL_DETAIL</code> . Th rows have the following columns:	
	<pre>(global_detail_name VARCHAR2(30), global_detail_value NUMBER)</pre>	

Examples

The following example returns the global model details for the GLM regression model ${\tt GLMR}$ SH Regr sample.

```
SELECT *
 FROM TABLE (dbms data mining.get model details global(
            'GLMR SH Regr sample'))
ORDER BY global detail name;
GLOBAL DETAIL NAME GLOBAL DETAIL VALUE
______
                                    .731412557
ADJUSTED R SQUARE
AIC
                                     5931.814
                                   18.1711243
COEFF VAR
CORRECTED_TOTAL_DF
                                        1499
                                  278740.504
CORRECTED TOT SS
DEPENDENT MEAN
                                      38.892
ERROR DF
                                         1433
ERROR_MEAN_SQUARE
                                  49.9440956
ERROR SUM SQUARES
                                    71569.8891
F VALUE
                                    62.8492452
GMSEP
                                     52.280819
HOCKING_SP
                                    .034877162
                                    52.1749319
JΡ
MODEL CONVERGED
                                            1
MODEL DF
                                            66
MODEL F P VALUE
MODEL MEAN SQUARE
                                    3138.94871
MODEL SUM SQUARES
                                    207170.615
NUM PARAMS
                                            67
NUM ROWS
                                          1500
ROOT MEAN SQ
                                    7.06711367
R SQ
                                    .743238288
SBIC
                                    6287.79977
VALID COVARIANCE MATRIX
```



Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_KM Function

The GET_MODEL_DETAILS_KM function returns a set of rows that provide the details of a k-means clustering model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

You can provide input to <code>GET_MODEL_DETAILS_KM</code> to request specific information about the model, thus improving the performance of the query. If you do not specify filtering parameters, then <code>GET_MODEL_DETAILS_KM</code> returns all the information about the model.

Syntax

```
DBMS_DATA_MINING.get_model_details_km(
    model_name VARCHAR2,
    cluster_id NUMBER    DEFAULT NULL,
    attribute   VARCHAR2    DEFAULT NULL,
    centroid   NUMBER    DEFAULT 1,
    histogram   NUMBER    DEFAULT 1,
    rules         NUMBER    DEFAULT 2,
    attribute_subname    VARCHAR2    DEFAULT NULL,
    topn_attributes    NUMBER    DEFAULT NULL,
    partition_name    VARCHAR2    DEFAULT NULL)
RETURN dm clusters PIPELINED;
```

Parameters

Table 62-97 GET MODEL DETAILS KM Function Parameters

Parameter	Description	
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.	
cluster_id	The ID of a cluster in the model. When a valid cluster ID is specified, only the details of this cluster are returned. Otherwise the details for all clusters are returned.	
attribute	The name of an attribute. When a valid attribute name is specified, only the details of this attribute are returned. Otherwise, the details for all attributes are returned	
centroid	 This parameter accepts the following values: 1: Details about centroids are returned (default) 0: Details about centroids are not returned 	
histogram	 This parameter accepts the following values: 1: Details about histograms are returned (default) 0: Details about histograms are not returned 	
rules	This parameter accepts the following values: 2: Details about rules are returned (default) 1: Rule summaries are returned 0: No information about rules is returned	



Table 62-97 (Cont.) GET_MODEL_DETAILS_KM Function Parameters

Parameter	Description
attribute_subname	The name of a nested attribute. The full name of a nested attribute has the form:
	attribute_name.attribute_subname
	where <code>attribute_name</code> is the name of the column and <code>attribute_subname</code> is the name of the nested attribute in that column.
	If the attribute is not nested, attribute_subname is null.
topn_attributes	Restricts the number of attributes returned in the centroid, histogram, and rules objects. Only the n attributes with the highest confidence values in the rules are returned.
	If the number of attributes included in the rules is less than $topn$, then up to n additional attributes in alphabetical order are returned.
	If both the attribute and topn_attributes parameters are specified, then topn_attributes is ignored.
partition_name	Specifies a partition in a partitioned model.

Usage Notes

- 1. The table function pipes out rows of type DM_CLUSTERS. For information on machine learning data types and Return Value for Clustering Algorithms piped output from table functions, see "Data Types".
- 2. When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

Examples

The following example returns model details for the k-means clustering model KM SH Clus sample.

```
SELECT T.id clu_id,
    T.record_count rec_cnt,
    T.parent parent,
    T.tree_level tree_level,
    T.dispersion dispersion

FROM (SELECT *
    FROM TABLE (DBMS_DATA_MINING.GET_MODEL_DETAILS_KM( 'KM_SH_Clus_sample'))

ORDER BY id) T

WHERE ROWNUM < 6;

CLU_ID REC_CNT PARENT TREE_LEVEL DISPERSION

1 1500 1 5.9152211
2 638 1 2 3.98458982
3 862 1 2 5.83732097
4 376 3 3 5.05192137
5 486 3 3 5.42901522
```

Related Topics

Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_NB Function

The <code>GET_MODEL_DETAILS_NB</code> function returns a set of rows that provide the details of a naive Bayes model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

Syntax

```
DBMS_DATA_MINING.get_model_details_nb(
          model_name IN VARCHAR2,
          partition_name IN VARCHAR2 DEFAULT NULL)
RETURN DM NB Details PIPELINED;
```

Parameters

Table 62-98 GET_MODEL_DETAILS_NB Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model

Return Values

Table 62-99 GET_MODEL_DETAILS_NB Function Return Values

Return Value	Description		
DM_NB_DETAILS	A set of rows of type <code>DM_NB_DETAIL</code> . The rows have the following columns:		
		VARCHAR2(4000),	
	<pre>(attribute_name attribute_subname attribute_str_value attribute_num_value conditional_probability</pre>	NUMBER,	

Usage Notes

- The table function pipes out rows of type DM_NB_DETAILS. For information on machine learning data types and piped output from table functions, see "Data Types".
- When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

Examples

The following query is from the sample program dmnbdemo.sql. It returns model details about the model $NB_SH_Clas_sample$. For information about the sample programs, see *Oracle Machine Learning for SQL User's Guide*.

The query creates labels from the bin boundary tables that were used to bin the training data. It replaces the attribute values with the labels. For numeric bins, the labels are (lower_boundary,upper_boundary]; for categorical bins, the label matches the value it represents. (This method of categorical label representation will only work for cases where one value corresponds to one bin.) The target was not binned.

```
אידיא
      bin label view AS (
      SELECT col, bin, (DECODE(bin, '1', '[', '(') || lv || ',' || val || ']') label
         FROM (SELECT col,
                          bin,
                          LAST VALUE (val) OVER (
                          PARTITION BY col ORDER BY val
                          ROWS BETWEEN UNBOUNDED PRECEDING AND 1 PRECEDING) lv,
                  FROM nb sh sample num)
     UNION ALL
     SELECT col, bin, val label
       FROM nb sh sample cat
    model details AS (
     SELECT T.target attribute name
              NVL(TO_CHAR(T.target_attribute_num_value,T.target_attribute_str_value)) tval,
              C.attribute name
              NVL(L.label, NVL(C.attribute_str_value, C.attribute_num_value)) pval,
              T.prior probability
              C.conditional probability
                                                                                                   condp
       FROM TABLE (DBMS DATA MINING.GET MODEL DETAILS NB ('NB SH Clas sample')) T,
              TABLE (T.conditionals) C,
              bin label view L
      WHERE C.attribute name = L.col (+) AND
              (NVL(C.attribute str value, C.attribute num value) = L.bin(+))
     ORDER BY 1,2,3,4,5,6
     SELECT tname, tval, pname, pval, priorp, condp
       FROM model details
      WHERE ROWNUM < 11;
          TVAL PNAME
                                                              PVAL PRIORP CONDP
 TNAME
AFFINITY_CARD 0 AGE (24,30] .6500 .1714
AFFINITY_CARD 0 AGE (30,35] .6500 .1509
AFFINITY_CARD 0 AGE (35,40] .6500 .1125
AFFINITY_CARD 0 AGE (40,46] .6500 .1134
AFFINITY_CARD 0 AGE (46,53] .6500 .1071
AFFINITY_CARD 0 AGE (53,90] .6500 .1312
AFFINITY_CARD 0 AGE [17,24] .6500 .2134
AFFINITY_CARD 0 BOOKKEEPING_APPLICATION 0 .6500 .1500
AFFINITY_CARD 0 BOOKKEEPING_APPLICATION 1 .6500 .8500
AFFINITY_CARD 0 BULK_PACK_DISKETTES 0 .6500 .3670
```

Related Topics

Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_NMF Function

The <code>GET_MODEL_DETAILS_NMF</code> function returns a set of rows that provide the details of a non-negative matrix factorization model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

Syntax

```
DBMS_DATA_MINING.get_model_details_nmf(
          model_name IN VARCHAR2,
          partition_name VARCHAR2 DEFAULT NULL)
RETURN DM NMF Feature Set PIPELINED;
```

Parameters

Table 62-100 GET_MODEL_DETAILS_NMF Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model

Return Values

Table 62-101 GET_MODEL_DETAILS_NMF Function Return Values

```
Return Value
                      Description
                     A set of rows of DM NMF FEATURE. The rows have the following columns:
DM NMF FEATURE SET
                      (feature id
                                           NUMBER,
                      mapped_feature_id VARCHAR2(4000),
                       attribute set
                                        DM NMF ATTRIBUTE SET)
                      The attribute set column of DM NMF FEATURE returns a nested table of
                      type DM NMF ATTRIBUTE SET. The rows, of type DM NMF ATTRIBUTE, have
                      the following columns:
                           (attribute_name VARCHAR2(4000),
                            attribute_subname VARCHAR2(4000),
                            attribute_value VARCHAR2(4000),
                            coefficient
                                               NUMBER)
```

Usage Notes

- The table function pipes out rows of type DM_NMF_FEATURE_SET. For information on machine learning data types and piped output from table functions, see "Data Types".
- When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

Examples

The following example returns model details for the feature extraction model NMF SH Sample.

Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_OC Function

The <code>GET_MODEL_DETAILS_OC</code> function returns a set of rows that provide the details of an O-cluster clustering model. The rows are an enumeration of the clustering patterns generated during the creation of the model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

You can provide input to <code>GET_MODEL_DETAILS_OC</code> to request specific information about the model, thus improving the performance of the query. If you do not specify filtering parameters, then <code>GET_MODEL_DETAILS_OC</code> returns all the information about the model.

Syntax

```
DBMS_DATA_MINING.get_model_details_oc(
    model_name VARCHAR2,
    cluster_id NUMBER DEFAULT NULL,
    attribute VARCHAR2 DEFAULT NULL,
    centroid NUMBER DEFAULT 1,
    histogram NUMBER DEFAULT 1,
    rules NUMBER DEFAULT 2,
    topn_attributes NUMBER DEFAULT NULL,
    partition_name VARCHAR2 DEFAULT NULL)
RETURN dm clusters PIPELINED;
```

Parameters

Table 62-102 GET_MODEL_DETAILS_OC Function Parameters

Parameter	Description	
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.	
cluster_id	The ID of a cluster in the model. When a valid cluster ID is specified, only the details of this cluster are returned. Otherwise the details for all clusters are returned.	
attribute	The name of an attribute. When a valid attribute name is specified, only the details of this attribute are returned. Otherwise, the details for all attributes are returned	
centroid	This parameter accepts the following values:	
	1: Details about centroids are returned (default)0: Details about centroids are not returned	
histogram	This parameter accepts the following values:	
	 1: Details about histograms are returned (default) 	
	 0: Details about histograms are not returned 	
rules	This parameter accepts the following values:	
	 2: Details about rules are returned (default) 	
	 1: Rule summaries are returned 	
	 0: No information about rules is returned 	
topn_attributes	Restricts the number of attributes returned in the centroid, histogram, and rules objects. Only the n attributes with the highest confidence values in the rules are returned.	
	If the number of attributes included in the rules is less than $topn$, then up to n additional attributes in alphabetical order are returned.	
	If both the attribute and topn_attributes parameters are specified, then topn_attributes is ignored.	
partition_name	Specifies a partition in a partitioned model.	

Usage Notes

- For information about machine learning data types and return values for clustering algorithms piped output from table functions, see "Data Types".
- 2. When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

Examples

The following example returns model details for the clustering model OC SH Clus sample.

For each cluster in this example, the split predicate indicates the attribute and the condition used to assign records to the cluster's children during model build. It provides an important piece of information on how the population within a cluster can be divided up into two smaller clusters.

```
TABLE(a.split_predicate) sp
  ORDER BY a.id, op, s_value)
WHERE ROWNUM < 11;</pre>
```

CLU_ID	ATTRIBUTE_NAME	OP	S_VALUE
1	OCCUPATION	IN	?
1	OCCUPATION	IN	Armed-F
1	OCCUPATION	IN	Cleric.
1	OCCUPATION	IN	Crafts
2	OCCUPATION	IN	?
2	OCCUPATION	IN	Armed-F
2	OCCUPATION	IN	Cleric.
3	OCCUPATION	IN	Exec.
3	OCCUPATION	IN	Farming
3	OCCUPATION	IN	Handler

Oracle Machine Learning for SQL User's Guide

GET_MODEL_SETTINGS Function

The GET_MODEL_SETTINGS function returns the settings used to build the given model. Starting from Oracle Database 12c Release 2, this function is deprecated. See "Static Data Dictionary Views: ALL ALL TABLES to ALL OUTLINES" in *Oracle Database Reference*.

Syntax

FUNCTION get_model_settings (model_name IN VARCHAR2)
 RETURN DM_Model_Settings PIPELINED;

Parameters

Table 62-103 GET_MODEL_SETTINGS Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.

Return Values

Table 62-104 GET_MODEL_SETTINGS Function Return Values

Return Value	Description	
DM_MODEL_SETTINGS	A set of rows of type DM_MODE following columns:	L_SETTINGS. The rows have the
	DM_MODEL_SETTINGS TABLE OF Name	F SYS.DM_MODEL_SETTING Type
	SETTING_NAME SETTING_VALUE	VARCHAR2 (30) VARCHAR2 (4000)

Usage Notes

- This table function pipes out rows of type DM_MODEL_SETTINGS. For information on machine learning data types and piped output from table functions, see "DBMS_DATA_MINING Datatypes".
- 2. The setting names/values include both those specified by the user and any defaults assigned by the build process.

Examples

The following example returns model model settings for an example naive Bayes model.

SETTING_NAME	SETTING_VALUE
ALGO_NAME	ALGO_NAIVE_BAYES
PREP_AUTO	ON
ODMS_MAX_PARTITIONS	1000
NABS_SINGLETON_THRESHOLD	0
CLAS_WEIGHTS_BALANCED	OFF
NABS_PAIRWISE_THRESHOLD	0
ODMS_PARTITION_COLUMNS	GENDER, Y_BOX_GAMES
ODMS MISSING VALUE TREATMENT	ODMS MISSING VALUE AUTO
ODMS_SAMPLING	ODMS_SAMPLING_DISABLE
_	

9 rows selected.

Related Topics

Oracle Database Reference

GET_MODEL_SIGNATURE Function

The GET_MODEL_SIGNATURE function returns the list of columns from the build input table that were used by the build process to train the model. Starting from Oracle Database 12c Release 2, this function is deprecated. See "Static Data Dictionary Views: ALL_ALL_TABLES to ALL OUTLINES" in Oracle Database Reference.

Syntax

```
FUNCTION get_model_signature (model_name IN VARCHAR2) RETURN DM Model Signature PIPELINED;
```

Parameters

Table 62-105 GET_MODEL_SIGNATURE Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.



Return Values

Table 62-106 GET_MODEL_SIGNATURE Function Return Values

Return Value	Description	
DM_MODEL_SIGNATURE	A set of rows of type DM_MODE following columns:	EL_SIGNATURE. The rows have the
	DM_MODEL_SIGNATURE TABLE OF SYS.DM MODEL SIGNATURE ATTRIBUTE	
	Name	Туре
	ATTRIBUTE_NAME	VARCHAR2 (130)
	ATTRIBUTE_TYPE	VARCHAR2 (106)

Usage Notes

- This table function pipes out rows of type DM_MODEL_SIGNATURE. For information on machine learning data types and piped output from table functions, see "DBMS_DATA_MINING Datatypes".
- 2. The signature names or types include only those attributes used by the build process.

Examples

The following example returns model settings for an example naive Bayes model.

ATTRIBUTE_NAME	ATTRIBUTE_TYPE
AGE	NUMBER
ANNUAL INCOME	NUMBER
AVERAGE ITEMS PURCHASED	NUMBER
	NUMBER
BULK PACK DISKETTES	NUMBER
BULK PURCH AVE AMT	NUMBER
DISABLE COOKIES	NUMBER
EDUCATION	VARCHAR2
FLAT PANEL MONITOR	NUMBER
GENDER	VARCHAR2
HOME_THEATER_PACKAGE	NUMBER
HOUSEHOLD_SIZE	VARCHAR2
MAILING_LIST	NUMBER
MARITAL_STATUS	VARCHAR2
NO_DIFFERENT_KIND_ITEMS	NUMBER
OCCUPATION	VARCHAR2
OS_DOC_SET_KANJI	NUMBER
PETS	NUMBER
PRINTER_SUPPLIES	NUMBER
PROMO_RESPOND	NUMBER
SHIPPING_ADDRESS_COUNTRY	VARCHAR2
SR_CITIZEN	NUMBER
TOP_REASON_FOR_SHOPPING	VARCHAR2
WKS_SINCE_LAST_PURCH	NUMBER
WORKCLASS	VARCHAR2
YRS_RESIDENCE	NUMBER
Y_BOX_GAMES	NUMBER
27 rows selected.	



Oracle Database Reference

GET_MODEL_DETAILS_SVD Function

The GET_MODEL_DETAILS_SVD function returns a set of rows that provide the details of a singular value decomposition model. Oracle recommends to use model details view settings. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

Refer to Model Details View for Singular Value Decomposition.

Syntax

Parameters

Table 62-107 GET_MODEL_DETAILS_SVD Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
matrix_type	Specifies which of the three SVD matrix types to return. Values are: U , S , V , and $NULL$. When matrix_type is null (default), all matrices are returned.
	The U matrix is only computed when the SVDS_U_MATRIX_OUTPUT setting is enabled. It is not computed by default. If the model does not contain U matrices and you set matrix_type to U, an empty set of rows is returned. See Table 62-27.
partition_name	A partition in a partitioned model.

Return Values

Table 62-108 GET_MODEL_DETAILS_SVD Function Return Values

Return Value	Description	
DM_SVD_MATRIX_SET	A set of rows of type Dicolumns:	M_SVD_MATRIX. The rows have the following
	<pre>(matrix_type feature_id mapped_feature_id attribute_name attribute_subname case_id value variance pct_cum_variance</pre> See Usage Notes for d	VARCHAR2 (4000), VARCHAR2 (4000), VARCHAR2 (4000), NUMBER, NUMBER, NUMBER)

Usage Notes

1. This table function pipes out rows of type DM_SVD_MATRIX. For information on machine learning data types and piped output from table functions, see "Data Types".

The columns in each row returned by <code>GET_MODEL_DETAILS_SVD</code> are described as follows:

Column in DM_SVD_MATRIX_SET	Description
matrix_type	The type of matrix. Possible values are S , V , and U . This field is never null.
feature_id	The feature that the matrix entry refers to.
mapped_feature_id	A descriptive name for the feature.
attribute_name	Column name in the V matrix component bases. This field is null for the S and U matrices.
attribute_subname	Subname in the V matrix component bases. This is relevant only in the case of a nested column. This field is null for the S and U matrices.
case_id	Unique identifier of the row in the build data described by the U matrix projection. This field is null for the S and V matrices.
value	The matrix entry value.
variance	The variance explained by a component. It is non-null only for S matrix entries. This column is non-null only for S matrix entries and for SVD models with setting dbms_data_mining.svds_scoring_mode set to dbms_data_mining.svds_scoring_pca and the build data is centered, either manually or because the setting dbms_data_mining.prep_auto is set to dbms_data_mining.prep_auto_on.
pct_cum_variance	The percent cumulative variance explained by the components thus far. The components are ranked by the explained variance in descending order.
	This column is non-null only for S matrix entries and for SVD models with setting dbms_data_mining.svds_scoring_mode set to dbms_data_mining.svds_scoring_pca and the build data is centered, either manually or because the setting dbms_data_mining.prep_auto is set to dbms_data_mining.prep_auto_on.

- 2. The output of <code>GET_MODEL_DETAILS</code> is in sparse format. Zero values are not returned. Only the diagonal elements of the <code>S</code> matrix, the non-zero coefficients in the <code>V</code> matrix bases, and the non-zero <code>U</code> matrix projections are returned.
 - There is one exception: If the data row does not produce non-zero **U** Matrix projections, the case ID for that row is returned with <code>NULL</code> for the <code>feature_id</code> and <code>value</code>. This is done to avoid losing any records from the original data.
- 3. GET_MODEL_DETAILS functions preserve model transparency by automatically reversing the transformations applied during the build process. Thus the attributes returned in the model details are the original attributes (or a close approximation of the original attributes) used to build the model.
- **4.** When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the preferred partition name.

Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_SVM Function

The GET_MODEL_DETAILS_SVM function returns a set of rows that provide the details of a linear support vector machines (SVM) model. If invoked for nonlinear SVM, it returns ORA-40215. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views in Oracle Machine Learning for SQL User's Guide.

In linear SVM models, only nonzero coefficients are stored. This reduces storage and speeds up model loading. As a result, if an attribute is missing in the coefficient list returned by GET_MODEL_DETAILS_SVM, then the coefficient of this attribute should be interpreted as zero.

Syntax

Parameters

Table 62-109 GET_MODEL_DETAILS_SVM Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
reverse_coef	Whether or not <code>GET_MODEL_DETAILS_SVM</code> should transform the attribute coefficients using the original attribute transformations.
	When reverse_coef is set to 0 (default), GET_MODEL_DETAILS_SVM returns the coefficients directly from the model without applying transformations.
	When reverse_coef is set to 1, GET_MODEL_DETAILS_SVM transforms the coefficients and bias by applying the normalization shifts and scales that were generated using automatic data preparation.
	See Usage Note 4.
partition_name	Specifies a partition in a partitioned model.

Return Values

Table 62-110 GET MODEL DETAILS SVM Function Return Values

Return Value	Description
DM_SVM_LINEAR_COEFF_ SET	A set of rows of type <code>DM_SVM_LINEAR_COEFF</code> . The rows have the following columns:
	(class VARCHAR2 (4000), attribute_set DM_SVM_ATTRIBUTE_SET) The attribute_set column returns a nested table of type DM_SVM_ATTRIBUTE_SET. The rows, of type DM_SVM_ATTRIBUTE, have the following columns:
	(attribute_name VARCHAR2(4000), attribute_subname VARCHAR2(4000), attribute_value VARCHAR2(4000), coefficient NUMBER)
	See Usage Notes.

Usage Notes

- This table function pipes out rows of type DM_SVM_LINEAR_COEFF. For information on machine learning data types and piped output from table functions, see "Data Types".
- 2. The class column of DM_SVM_LINEAR_COEFF contains classification target values. For SVM Regression models, class is null. For each classification target value, a set of coefficients is returned. For binary classification, one-class classification, and regression models, only a single set of coefficients is returned.
- 3. The attribute value column in DM SVM ATTRIBUTE SET is used for categorical attributes.
- 4. GET_MODEL_DETAILS functions preserve model transparency by automatically reversing the transformations applied during the build process. Thus the attributes returned in the model details are the original attributes (or a close approximation of the original attributes) used to build the model.
 - The coefficients are related to the transformed, not the original, attributes. When returned directly with the model details, the coefficients may not provide meaningful information. If you want <code>GET_MODEL_DETAILS_SVM</code> to transform the coefficients such that they relate to the original attributes, set the <code>reverse coef parameter</code> to 1.
- 5. When the value is NULL for a partitioned model, an exception is thrown. When the value is not null, it must contain the desired partition name.

Examples

The following example returns model details for the SVM classification model SVMC_SH_Clas_sample, which was created by the sample program dmsvcdem.sql. For information about the sample programs, see *Oracle Machine Learning for SQL User's Guide*.

```
WITH
  mod_dtls AS (
  SELECT *
    FROM TABLE(DBMS_DATA_MINING.GET_MODEL_DETAILS_SVM('SVMC_SH_Clas_sample'))
),
  model details AS (
```



CLASS	ANAME	AVAL	COEFF
1			-2.85
1	BOOKKEEPING_APPLICATION		1.11
1	OCCUPATION	Other	94
1	HOUSEHOLD_SIZE	4-5	.88
1	CUST_MARITAL_STATUS	Married	.82
1	YRS_RESIDENCE		.76
1	HOUSEHOLD_SIZE	6-8	74
1	OCCUPATION	Exec.	.71
1	EDUCATION	11th	71
1	EDUCATION	Masters	.63

Oracle Machine Learning for SQL User's Guide

GET_MODEL_DETAILS_XML Function

This function returns an XML object that provides the details of a decision tree model. Starting from Oracle Database 12c Release 2, this function is deprecated. Use model detail views instead.

See Model Detail Views for Decision Tree in Oracle Machine Learning for SQL User's Guide.

Syntax

```
DBMS_DATA_MINING.get_model_details_xml(
          model_name IN VARCHAR2,
          partition_name IN VARCHAR2 DEFAULT NULL)
RETURN XMLType;
```

Parameters

Table 62-111 GET_MODEL_DETAILS_XML Function Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model.

Return Values

Table 62-112 GET MODEL DETAILS XML Function Return Value

Return Value	Description
XMLTYPE	The XML definition for the decision tree model. See "XMLTYPE" for details.
	The XML definition conforms to the Data Mining Group Predictive Model Markup Language (PMML) version 2.1 specification. The specification is available at https://dmg.org.
	If a nested attribute is used as a splitter, the attribute will appear in the XML document as field=" <column_name>'.<subname>", as opposed to the non-nested attributes which appear in the document as field="<column_name>".</column_name></subname></column_name>
	Note: The column names are surrounded by single quotes and a period separates the column_name from the subname.
	The rest of the document style remains unchanged.

Usage Notes

Special characters that cannot be displayed by Oracle XML are converted to '#'.

Examples

The following statements in SQL*Plus return the details of the decision tree model dt sh clas sample.

Note: The """ characters you will see in the XML output are a result of SQL*Plus behavior. To display the XML in proper format, cut and past it into a file and open the file in a browser.

```
column dt details format a320
SELECT
dbms_data_mining.get_model_details_xml('dt_sh_clas_sample')
AS DT DETAILS
FROM dual;
DT DETAILS
______
<PMML version="2.1">
 <Header copyright="Copyright (c) 2004, Oracle Corporation. All rights</pre>
     reserved."/>
 <DataDictionary numberOfFields="9">
   <DataField name="AFFINITY CARD" optype="categorical"/>
   <DataField name="AGE" optype="continuous"/>
   <DataField name="BOOKKEEPING APPLICATION" optype="continuous"/>
   <DataField name="CUST MARITAL STATUS" optype="categorical"/>
   <DataField name="EDUCATION" optype="categorical"/>
   <DataField name="HOUSEHOLD SIZE" optype="categorical"/>
   <DataField name="OCCUPATION" optype="categorical"/>
   <DataField name="YRS RESIDENCE" optype="continuous"/>
   <DataField name="Y BOX_GAMES" optype="continuous"/>
 </DataDictionary>
```

```
<TreeModel modelName="DT SH CLAS SAMPLE" functionName="classification"</pre>
     splitCharacteristic="binarySplit">
    <Extension name="buildSettings">
      <Setting name="TREE IMPURITY METRIC" value="TREE IMPURITY GINI"/>
      <Setting name="TREE TERM MAX DEPTH" value="7"/>
      <Setting name="TREE TERM MINPCT NODE" value=".05"/>
      <Setting name="TREE TERM MINPCT SPLIT" value=".1"/>
      <Setting name="TREE TERM MINREC NODE" value="10"/>
      <Setting name="TREE TERM MINREC SPLIT" value="20"/>
      <costMatrix>
        <costElement>
          <actualValue>0</actualValue>
          <predictedValue>0</predictedValue>
          <cost>0</cost>
        </costElement>
        <costElement>
          <actualValue>0</actualValue>
          cpredictedValue>1</predictedValue>
          <cost>1</cost>
        </costElement>
        <costElement>
          <actualValue>1</actualValue>
          <predictedValue>0</predictedValue>
          <cost>8</cost>
        </costElement>
        <costElement>
          <actualValue>1</actualValue>
          cpredictedValue>1</predictedValue>
          <cost>0</cost>
        </costElement>
      </costMatrix>
    </Extension>
    <MiningSchema>
      </Node>
    </Node>
  </TreeModel>
</PMML>
```

GET MODEL TRANSFORMATIONS Function

This function returns the transformation expressions embedded in the specified model. Starting from Oracle Database 12c Release 2, this function is deprecated. See "Static Data Dictionary Views: ALL ALL TABLES to ALL OUTLINES" in *Oracle Database Reference*.

All GET_* interfaces are replaced by model views, and Oracle recommends that users reference the model views to retrieve the relevant information. The GET MODEL TRANSFORMATIONS function is replaced by the following:

- USER(/DBA/ALL)_MINING_MODEL_XFORMS: provides the user-embedded transformations
- DM\$VX prefixed model view: provides text feature extraction information
- D\$VN prefixed mode view: provides normalization and missing value information
- DM\$VB: provides binning information

See Also:

"About Transformation Lists" in DBMS_DATA_MINING_TRANSFORM Operational Notes

"GET_TRANSFORM_LIST Procedure"

"CREATE_MODEL Procedure"

"ALL_MINING_MODEL_XFORMS" in Oracle Database Reference

"DBA_MINING_MODEL_XFORMS" in Oracle Database Reference

"USER_MINING_MODEL_XFORMS" in Oracle Database Reference

Model Details View for Binning

Normalization and Missing Value Handling

Data Preparation for Text Features

Syntax

```
DBMS_DATA_MINING.get_model_transformations(
          model_name IN VARCHAR2,
          partition_name IN VARCHAR2 DEFAULT NULL)
RETURN DM_Transforms PIPELINED;
```

Parameters

Table 62-113 GET MODEL TRANSFORMATIONS Function Parameters

Parameter	Description
model_name	Indicates the name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
partition_name	Specifies a partition in a partitioned model

Return Values

Table 62-114 GET_MODEL_TRANSFORMATIONS Function Return Value

Return Value	Description		
DM_TRANSFORMS	The transformation expressions embedded in model_name.		
	The DM_TRANSFORMS type is a table of DM_TRANSFORM objects. Each DM_TRANSFORM has these fields:		
	attribute_name attribute_subname expression reverse_expression	VARCHAR2 (4000) VARCHAR2 (4000) CLOB CLOB	

Usage Notes

When Automatic Data Preparation (ADP) is enabled, both automatic and user-defined transformations may be associated with an attribute. In this case, the user-defined transformations are evaluated before the automatic transformations.

When invoked for a partitioned model, the partition name parameter must be specified.

Examples

In this example, several columns in the SH.CUSTOMERS table are used to create a naive Bayes model. A transformation expression is specified for one of the columns. The model does not use ADP.

```
CREATE OR REPLACE VIEW mining data AS
  SELECT cust id, cust year of birth, cust income level, cust credit limit
  FROM sh.customers;
describe mining data
                                   Null? Type
CUST ID
                                     NOT NULL NUMBER
                                    NOT NULL NUMBER (4)
CUST YEAR OF BIRTH
CUST INCOME LEVEL
                                              VARCHAR2 (30)
CUST CREDIT LIMIT
                                               NUMBER
CREATE TABLE settings nb(
     setting name VARCHAR2(30),
     setting value VARCHAR2(30));
BEGIN
    INSERT INTO settings nb (setting name, setting value) VALUES
         (dbms data mining.algo name, dbms data mining.algo naive bayes);
    INSERT INTO settings nb (setting name, setting value) VALUES
         (dbms data mining.prep auto, dbms data mining.prep auto off);
    COMMIT:
END;
DECLARE
   mining data xforms dbms data mining transform.TRANSFORM LIST;
   dbms_data_mining_transform.SET TRANSFORM (
        expression => 'cust_year_of_birth + 10',
        reverse expression => 'cust year of birth - 10');
   dbms_data_mining.CREATE_MODEL (
      model_name => 'new_model',
mining_function => dbms_data_mining.classification,
data_table_name => 'mining_data',
       case id column name => 'cust id',
       target column name => 'cust income level',
       settings_table_name => 'settings nb',
       data schema name => nulL,
       settings_schema_name => null,
       xform_list => mining_data_xforms );
 END:
SELECT attribute name, TO CHAR(expression), TO CHAR(reverse expression)
     FROM TABLE (dbms data mining.GET MODEL TRANSFORMATIONS('new model'));
```

```
ATTRIBUTE_NAME TO_CHAR(EXPRESSION) TO_CHAR(REVERSE_EXPRESSION)

CUST_YEAR_OF_BIRTH cust_year_of_birth + 10 cust_year_of_birth - 10
```

Oracle Database Reference

GET_TRANSFORM_LIST Procedure

This procedure converts transformation expressions specified as <code>DM_TRANSFORMS</code> to a transformation list (<code>TRANSFORM_LIST</code>) that can be used in creating a model. <code>DM_TRANSFORMS</code> is returned by the <code>GET_MODEL_TRANSFORMATIONS</code> function.

You can also use routines in the <code>DBMS_DATA_MINING_TRANSFORM</code> package to construct a transformation list.

```
See Also:
"About Transformation Lists" in DBMS_DATA_MINING_TRANSFORM
"GET_MODEL_TRANSFORMATIONS Function"
"CREATE_MODEL Procedure"
```

Syntax

Parameters

Table 62-115 GET_TRANSFORM_LIST Procedure Parameters

Parameter	Description		
xform_list	A list of transformation specifications that can be embedded in a model. Accepted as a parameter to the CREATE_MODEL Procedure. The TRANSFORM_LIST type is a table of TRANSFORM_REC objects. Each TRANSFORM_REC has these fields:		
	attribute_name		
	<u> </u>		

Table 62-115 (Cont.) GET_TRANSFORM_LIST Procedure Parameters

Parameter	Description	
model_xforms	A list of embedded transformation expressions returned by the GET_MODEL_TRANSFORMATIONS Function for a specific model.	
	The DM_TRANSFORMS type is a table of DM_TRANSFORM objects. Each DM_TRANSFORM has these fields:	
	attribute_name attribute_subname expression reverse_expression	VARCHAR2 (4000) VARCHAR2 (4000) CLOB CLOB

Examples

In this example, a model mod1 is trained using several columns in the SH.CUSTOMERS table. The model uses ADP, which automatically bins one of the columns.

A second model mod2 is trained on the same data without ADP, but it uses a transformation list that was obtained from mod1. As a result, both mod1 and mod2 have the same embedded transformation expression.

```
CREATE OR REPLACE VIEW mining data AS
    SELECT cust id, cust year of birth, cust income level, cust credit limit
     FROM sh.customers;
describe mining data
                                       Null? Type
 Name
 CUST ID
                                         NOT NULL NUMBER
 CUST YEAR OF BIRTH
                                         NOT NULL NUMBER (4)
 CUST INCOME LEVEL
                                                  VARCHAR2 (30)
CUST CREDIT LIMIT
                                                  NUMBER
CREATE TABLE setmod1(setting name VARCHAR2(30), setting value VARCHAR2(30));
  INSERT INTO setmod1 VALUES (dbms data mining.algo name, dbms data mining.algo naive bayes);
  INSERT INTO setmod1 VALUES (dbms data mining.prep auto,dbms data mining.prep auto on);
   dbms data mining.CREATE MODEL (
             case_id_column_name => 'cust_id',
target_column_name => 'cust_income_level',
              settings_table_name => 'setmod1');
    COMMIT:
END;
CREATE TABLE setmod2(setting name VARCHAR2(30), setting value VARCHAR2(30));
BEGIN
  INSERT INTO setmod2
     VALUES (dbms_data_mining.algo_name, dbms_data_mining.algo_naive_bayes);
  COMMIT:
END;
DECLARE
  v xform list
                  dbms_data_mining_transform.TRANSFORM_LIST;
                    DM TRANSFORMS;
```

```
BEGIN
  EXECUTE IMMEDIATE
   'SELECT dm_transform(attribute_name, attribute_subname,expression, reverse_expression)
    FROM TABLE(dbms data mining.GET MODEL TRANSFORMATIONS (''mod1''))'
    BULK COLLECT INTO dmxf;
  dbms data mining.GET TRANSFORM LIST (
       xform_list
model xforms
                            => v xform list,
                           => dmxf);
       model xforms
  dbms data mining.CREATE MODEL (
                           => 'mod2',
       model name
        case_id_column_name => 'cust_id',
        target column name => 'cust income level',
        settings_table_name => 'setmod2',
        xform list
                     => v xform list);
END;
-- Transformation expression embedded in mod1
SELECT TO CHAR (expression) FROM TABLE (dbms data mining.GET MODEL TRANSFORMATIONS ('mod1'));
TO CHAR (EXPRESSION)
CASE WHEN "CUST YEAR OF BIRTH"<1915 THEN 0 WHEN "CUST YEAR OF BIRTH"<=1915 THEN 0
WHEN "CUST YEAR OF BIRTH"<=1920.5 THEN 1 WHEN "CUST YEAR OF BIRTH"<=1924.5 THEN 2
.5 THEN 29 WHEN "CUST YEAR OF BIRTH" IS NOT NULL THEN 30 END
-- Transformation expression embedded in mod2
SELECT TO CHAR(expression) FROM TABLE (dbms data mining.GET MODEL TRANSFORMATIONS('mod2'));
TO CHAR (EXPRESSION)
______
CASE WHEN "CUST YEAR OF BIRTH"<1915 THEN 0 WHEN "CUST YEAR OF BIRTH"<=1915 THEN 0
WHEN "CUST YEAR OF BIRTH"<=1920.5 THEN 1 WHEN "CUST YEAR OF BIRTH"<=1924.5 THEN 2
.5 THEN 29 WHEN "CUST YEAR OF BIRTH" IS NOT NULL THEN 30 END
-- Reverse transformation expression embedded in mod1
SELECT TO CHAR (reverse expression) FROM TABLE (dbms_data_mining.GET_MODEL_TRANSFORMATIONS('mod1'));
TO CHAR (REVERSE EXPRESSION)
DECODE("CUST YEAR OF BIRTH",0,'(; 1915), [1915; 1915]',1,'(1915; 1920.5]',2,'(1
920.5; 1924.5]',3,'(1924.5; 1928.5]',4,'(1928.5; 1932.5]',5,'(1932.5; 1936.5]',6
8,'(1987.5; 1988.5]',29,'(1988.5; 1989.5]',30,'(1989.5; )',NULL,'NULL')
-- Reverse transformation expression embedded in mod2
SELECT TO CHAR(reverse expression) FROM TABLE (dbms data mining.GET MODEL TRANSFORMATIONS('mod2'));
TO CHAR (REVERSE EXPRESSION)
DECODE ("CUST YEAR OF BIRTH",0,'(; 1915), [1915; 1915]',1,'(1915; 1920.5]',2,'(1
920.5; 1924.5]',3,'(1924.5; 1928.5]',4,'(1928.5; 1932.5]',5,'(1932.5; 1936.5]',6
```

```
.
.
8,'(1987.5; 1988.5]',29,'(1988.5; 1989.5]',30,'(1989.5; )',NULL,'NULL')
```

IMPORT_MODEL Procedure

This procedure imports one or more machine learning models. The procedure is overloaded. You can call it to import machine learning models from a dump file set, or you can call it to import a single machine learning model from a PMML document.

Import from a dump file set

You can import machine learning models from a dump file set that was created by the EXPORT_MODEL Procedure. IMPORT_MODEL and EXPORT_MODEL use Oracle Data Pump technology to export to and import from a dump file set.

When Oracle Data Pump is used directly to export/import an entire schema or database, the machine learning models in the schema or database are included. EXPORT_MODEL and IMPORT MODEL export/import machine learning models only.

Import from PMML

You can import a machine learning model represented in Predictive Model Markup Language (PMML). The model must be of type RegressionModel, either linear regression or binary logistic regression.

PMML is an XML-based standard specified by the Data Mining Group (https://dmg.org). Applications that are PMML-compliant can deploy PMML-compliant models that were created by any vendor. Oracle Machine Learning for SQL supports the core features of PMML 3.1 for regression models.

See Also:

Oracle Machine Learning for SQL User's Guide for more information about exporting and importing machine learning models

Oracle Database Utilities for information about Oracle Data Pump

https://dmg.org/dmg-faq.html for more information about PMML

Syntax

Imports a machine learning model from a dump file set:

```
DBMS_DATA_MINING.IMPORT_MODEL (
filename IN VARCHAR2,
directory IN VARCHAR2,
model_filter IN VARCHAR2 DEFAULT NULL,
operation IN VARCHAR2 DEFAULT NULL,
remote_link IN VARCHAR2 DEFAULT NULL,
jobname IN VARCHAR2 DEFAULT NULL,
schema_remap IN VARCHAR2 DEFAULT NULL,
tablespace remap IN VARCHAR2 DEFAULT NULL);
```

Imports a machine learning model from a PMML document:

Parameters

Table 62-116 IMPORT_MODEL Procedure Parameters

Parameter	Description
filename	Name of the dump file set from which the models should be imported. The dump file set must have been created by the <code>EXPORT_MODEL</code> procedure or the <code>expdp</code> export utility of Oracle Data Pump.
	The dump file set can contain one or more files. (Refer to "EXPORT_MODEL Procedure" for details.) If the dump file set contains multiple files, you can specify 'filename%U' instead of listing them. For example, if your dump file set contains 3 files, archive01.dmp, archive02.dmp, and archive03.dmp, you can import them by specifying 'archive%U'.
directory	Name of a pre-defined directory object that specifies where the dump file set is located. Both the exporting and the importing user must have read/write access to the directory object and to the file system directory that it identifies.
	Note: The target database must have also have read/write access to the file system directory.
model_filter	Optional parameter that specifies one or more models to import. If you do not specify a value for model_filter, all models in the dump file set are imported. You can also specify NULL (the default) or 'ALL' to import all models.
	The value of ${\tt model_filter}$ can be one or more model names. The following are valid filters.
	<pre>'mymodel1' 'name IN (''mymodel2'',''mymodel3'')'</pre>
	The first causes IMPORT_MODEL to import a single model named mymodel1. The second causes IMPORT_MODEL to import two models, mymodel2 and mymodel3.
operation	Optional parameter that specifies whether to import the models or the SQL statements that create the models. By default, the models are imported.
	You can specify either of the following values for operation:
	 'IMPORT' — Import the models (Default) 'SQL FILE' — Write the SQL DDL for creating the models to a text file. The
	text file is named job_name.sql and is located in the dump set directory.
remote_link	Optional parameter that specifies the name of a database link to a remote system. The default value is <code>NULL</code> . A database link is a schema object in a local database that enables access to objects in a remote database. When you specify a value for <code>remote_link</code> , you can import models into the local database from the remote database. The import is fileless; no dump file is involved. The <code>IMP_FULL_DATABASE</code> role is required for importing the remote models. The <code>EXP_FULL_DATABASE</code> privilege, the <code>CREATE DATABASE LINK</code> privilege, and other privileges may also be required. (See Example 2.)
jobname	Optional parameter that specifies the name of the import job. By default, the name has the form <code>username_imp_nnnn</code> , where <code>nnnn</code> is a number. For example, a job name in the <code>SCOTT</code> schema might be <code>SCOTT</code> imp 134.
	If you specify a job name, it must be unique within the schema. The maximum length of the job name is 30 characters.
	A log file for the import job, named <code>jobname.log</code> , is created in the same directory as the dump file set.

Table 62-116 (Cont.) IMPORT_MODEL Procedure Parameters

Parameter	Description
schema_remap	Optional parameter for importing into a different schema. By default, models are exported and imported within the same schema.
	If the dump file set belongs to a different schema, you must specify a schema mapping in the form <code>export_user:import_user</code> . For example, you would specify 'SCOTT:MARY' to import a model exported by SCOTT into the MARY schema.
	Note: In some cases, you may need to have the <code>IMP_FULL_DATABASE</code> privilege or the <code>SYS</code> role to import a model from a different schema.
tablespace_remap	Optional parameter for importing into a different tablespace. By default, models are exported and imported within the same tablespace.
	If the dump file set belongs to a different tablespace, you must specify a tablespace mapping in the form <code>export_tablespace:import_tablespace</code> . For example, you would specify 'TBLSPC01:TBLSPC02' to import a model that was exported from tablespace TBLSPC01 into tablespace TBLSPC02.
	Note: In some cases, you may need to have the <code>IMP_FULL_DATABASE</code> privilege or the <code>SYS</code> role to import a model from a different tablespace.
model_name	Name for the new model that will be created in the database as a result of an import from PMML The name must be unique within the user's schema.
pmmldoc	The PMML document representing the model to be imported. The PMML document has an XMLTYPE object type. See "XMLTYPE" for details.
strict_check	Whether or not an error occurs when the PMML document contains sections that are not part of core PMML (for example, Output or Targets). OML4SQL supports only core PMML; any non-core features may affect the scoring representation.
	If the PMML does not strictly conform to core PMML and <code>strict_check</code> is set to <code>TRUE</code> , then <code>IMPORT_MODEL</code> returns an error. If strict_check is <code>FALSE</code> (the default), then the error is suppressed. The model may be imported and scored.

Examples

This example shows a model being exported and imported within the schema oml_user2. Then the same model is imported into the oml_user3 schema. The oml_user3 user has the IMP_FULL_DATABASE privilege. The oml_user2 user has been assigned the USER2 tablespace; oml_user3 has been assigned the USER3 tablespace.

```
SQL>EXECUTE DBMS DATA MINING.IMPORT MODEL (
            filename => 'NMF SH SAMPLE out01.dmp',
            directory => 'DATA_PUMP_DIR',
            model filter => 'name = ''NMF SH SAMPLE''');
-- connect as different user
-- import same model into that schema
SQL> connect oml user3
Enter password: oml user3 password
Connected.
SQL>EXECUTE DBMS DATA MINING.IMPORT MODEL (
            filename => 'NMF SH SAMPLE out01.dmp',
            directory => 'DATA PUMP DIR',
            model filter => 'name = ''NMF SH SAMPLE''',
            operation => 'IMPORT',
            remote link => NULL,
            jobname => 'nmf imp job',
            schema remap => 'oml user2:oml user3',
            tablespace_remap => 'USER2:USER3');
```

The following example shows user MARY importing all models from a dump file, <code>model_exp_001.dmp</code>, which was created by user <code>SCOTT</code>. User MARY has been assigned a tablespace named <code>USER2</code>; user <code>SCOTT</code> was assigned the tablespace <code>USERS</code> when the models were exported into the dump file <code>model_exp_001.dmp</code>. The dump file is located in the file system directory mapped to a directory object called <code>DM_DUMP</code>. If user <code>MARY</code> does not have <code>IMP_FULL_DATABASE</code> privileges, <code>IMPORT_MODEL</code> will raise an error.

2. This example shows how the user xuser could import the model oml_user.rlmod from a remote database. The SQL*Net connection alias for the remote database is R1DB. The user xuser is assigned the SYSAUX tablespace; the user oml_user is assigned the TBS_1 tablespace.

R1MOD

3. This example shows how a PMML document called SamplePMML1.xml could be imported from a location referenced by directory object PMMLDIR into the schema of the current user. The imported model will be called PMMLMODEL1.

IMPORT_ONNX_MODEL Procedure

This procedure enables you to import an ONNX model into the Database.

Syntax

```
DBMS_DATA_MINING.IMPORT_ONNX_MODEL(
model_name IN VARCHAR2,
model_data IN BLOB,
metadata IN JSON);
```

Parameters

Table 62-117 IMPORT_ONNX_MODEL Procedure Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.
model_data	It is a BLOB holding the ONNX representation of the model. The BLOB contains the identical byte sequence as the one stored in an ONNX file.
metadata	A JSON description of the metadata describing the model. The metadata at minimum must describe the machine learning function supported by the model. The model's metadata parameters are described in JSON Metadata Parameters for ONNX Models.

Example

The following example illustrates a code snippet of using the

DBMS_DATA_MINING.IMPORT_ONNX_MODEL procedure. The complete step-by-step example is illustrated in Import ONNX Models and Generate Embeddings and Alternate Method to Import ONNX Models.



For a complete example to illustrate how you can define a BLOB variable and use it in the IMPORT ONNX MODEL procedure, you can have the following:

Usage Notes

The name of the model follows the same restrictions as those used for other machine learning models, namely:

- The schema name, if provided, is limited to 128 characters.
- The model name is limited to 123 characters and must follow the rules of unquoted identifiers: they contain only alphanumeric characters, the underscore (_), dollar sign (\$), and pound sign (#). The initial character must be alphabetic.
- The model size is limited to 1 gigabyte.
- The model must not depend on external initializers. To know more about initializers and other ONNX concepts, see https://onnx.ai/onnx/intro/concepts.html.

IMPORT_SERMODEL Procedure

This procedure imports the serialized format of the model back into a database.

The import routine takes the serialized content in the BLOB and the name of the model to be created with the content. This import does not create model views or tables that are needed for querying model details. The import procedure only provides the ability to score the model.

Syntax

```
DBMS_DATA_MINING.IMPORT_SERMODEL (

model_data IN BLOB,

model name IN VARCHAR2,);
```

Parameters

Table 62-118 IMPORT_SERMODEL Procedure Parameters

Parameter	Description
model_data	Provides model data in BLOB format.
model_name	Name of the machine learning model in the form [schema_name.]model_name. If you do not specify a schema, then your own schema is used.



Examples

The following statement imports the serialized format of the models.

```
declare
  v_blob blob;
BEGIN
  dbms_lob.createtemporary(v_blob, FALSE);
-- fill in v_blob from somewhere (e.g., bfile, etc.)
  dbms_data_mining.import_sermodel(v_blob, 'MY_MODEL');
  dbms_lob.freetemporary(v_blob);
END;
//
```

Related Topics

EXPORT SERMODEL Procedure

This procedure exports the model in a serialized format so that they can be moved to another platform for scoring.

See Also:

Oracle Machine Learning for SQL User's Guide for more information about exporting and importing machine learning models

JSON Schema for R Extensible Algorithm

Provides some flexibility when creating a new JSON object following the JSON schema.

Usage Note

Some flexibility when creating a new JSON object is as follows:

- Partial registration is allowed. For example, the detail function can be missing.
- Different orders are allowed. For example, the detail function can be written before the build function or after it.

Example 62-1 JSON Schema

JSON schema 1.1 for R extensible algorithm:

```
"name" : { "type" : "string"}}
                                              },
        "function_language": {"type": "string" },
        "mining function": {
                 "type" : "array",
                 "items" : [
                     { "type" : "object",
                         "properties" : {
                            "mining_function_name" : { "type" : "string"},
                            "build function": {
                                    "type": "object",
                                    "properties": {
                                         "function body": { "type": "CLOB" }
                                                         }
                                     },
        "detail function": {
                 "type" : "array",
                  "items" : [
                       {"type": "object",
                        "properties": {
                              "function body": { "type": "CLOB" },
                              "view columns": { "type" : "array",
                                                                     "items" : {
"type" : "object",
"properties" : {
"name" : { "type" : "string"},
 "type" : { "type" : "string",
                 "enum" : ["VARCHAR2",
                                   "NUMBER",
                                   "DATE",
                                   "BOOLEAN"]
               }
                                                                               }
                                                            }
                                             }
                     ]
        },
       "score_function": {
                 "type": "object",
                 "properties": {
                       "function_body": { "type": "CLOB" }
                 },
```

```
"weight function": {
                         "type": "object",
                         "properties": {
                             "function body": { "type": "CLOB" },
                 }
                                }
           } ]
        },
       "algo setting": {
                "type" : "array",
                "items" : [
                     { "type" : "object",
                        "properties" : {
                           "name"
                                               : { "type" : "string"},
                           "name display": { "type" : "object",
                                                           "properties" : {
                                                           "language" :
{ "type" : "string",
   "enum" : ["English", "Spanish", "French"],
   "default" : "English"},
                                                           "name" : { "type" :
"string"}}
                           "type" : { "type" : "string",
                                           "enum" : ["string", "integer",
"number", "boolean"]},
                           "optional": {"type" : "BOOLEAN",
                                                 "default" : "FALSE"},
                           "value" : { "type" : "string"},
                           "min value" : { "type": "object",
                                                        "properties": {
                                                              "min_value":
{"type": "number"},
                                                               "inclusive":
{ "type": "boolean",
     "default" : TRUE },
                                                    },
                            "max value" : {"type": "object",
                                                       "properties": {
                                                            "max value":
{"type": "number"},
                                                            "inclusive":
{ "type": "boolean",
   "default" : TRUE},
                                                              }
                                                      },
```

```
"categorical choices" : { "type": "array",
                                                                     "items": {
                                                                         "type":
"string"
                                                                 },
                           "description_display": { "type" : "object",
"properties" : {
"language" : { "type" : "string",
            "enum" : ["English", "Spanish", "French"],
            "default" : "English"},
                                                                     "name" :
{ "type" : "string"}}
                    }
                 ]
          }
    }
}
```

Example 62-2 JSON object example

The following is an JSON object example that must be passed to the registration procedure:

```
{ "algo name display"
                               {"English", "t1"},
                         "function language"
                                                        "R",
                         "mining function" : {
  "mining_function_name" : "CLASSIFICATION",
                         "build_function" : {"function_body": "function(dat,
formula, family)
{
                                                          set.seed(1234);
                                           mod <- glm(formula = formula,</pre>
data=dat,
                                                       family=
eval(parse(text=family))); mod}"},
           "score function" : { "function_body": "function(mod, dat) {
                                              res <- predict (mod, newdata =
dat,
type=''response
                                          '');
                                              res2=data.frame(1-res, res);
res2}"}}
                           "algo setting" :
                                            [{"name"
"dbms data mining.odms m
                                                   issing_value_treatment",
```

```
"name display"
                                              : {"English",
"dbms data mining.odms missing value
_treatment"},
                            "type"
                                                     : "string",
                                                   : "TRUE",
                             "optional"
                             "value"
"dbms data mining.odms missing value mean mode",
                            "categorical choices"
     "dbms_data_mining.odms_missing_value_mean_mode",
"dbms data mining.odms missing value auto",
"dbms data mining.odms missing value delete row"],
                            "description"
                                                           : {"English",
                                                                       "how to
treat missing values"}
                         },
{"name"
                       : "RALG PARAMETER FAMILY",
                            "name display" : {"English",
"RALG PARAMETER FAMILY" },
                            "type"
                                                    : "string",
                            "optional"
"value"
                                                  : "TRUE",
                                                    : "",
                            "description" : {"English", "R family
parameter in build function"}
],
                        }
```

REGISTER_ALGORITHM Procedure

Use this function to register a new algorithm by providing the algorithm name, machine learning function, and all other algorithm metadata.

Syntax

Parameters

Table 62-119 REGISTER ALGORITHM Procedure Parameters

Parameter	Description
algorithm_name	Name of the algorithm.
algorithm_metadata	Metadata of the algorithm.
algorithm_description	Description of the algorithm.

Usage Notes

The registration procedure performs the following:

- Checks whether algorithm metadata has correct JSON syntax.
- Checks whether the input JSON object follows the predefined JSON schema.
- Checks whether current user has RQADMIN privilege.
- Checks duplicate algorithms so that the same algorithm is not registered twice.
- Checks for missing entries. For example, algorithm name, algorithm type, metadata, and build function.

Register Algorithms After the JSON Object Is Created

SQL users can register new algorithms by creating a JSON object following the JSON schema and passing it to the REGISTER ALGORITHM procedure.

```
BEGIN
  DBMS DATA MINING.register algorithm(
    algorithm_name
                                       't1',
    algorithm metadata
                                 =>
    '{"function language" : "R",
      "mining function" :
        { "mining function name" : "CLASSIFICATION",
           "build function" : {"function body": "function(dat, formula,
family) { set.seed(1234);
                                         mod <- glm(formula = formula,</pre>
data=dat,
family=eval(parse(text=family)));
mod}"},
           "score function" : {"function body": "function(mod, dat) {
                                            res <- predict (mod, newdata =
dat, type=''response'');
                                            res2=data.frame(1-res, res);
res2}"}}
   }',
    algorithm description => 't1');
END;
```

RANK_APPLY Procedure

This procedure ranks the results of an APPLY operation based on a top-N specification for predictive and descriptive model results.

For classification models, you can provide a cost matrix as input, and obtain the ranked results with costs applied to the predictions.

Syntax



apply_result_schema_name	IN VARCHAR2	DEFAULT NULL,
cost matrix schema name	IN VARCHAR2	DEFAULT NULL);

Parameters

Table 62-120 RANK_APPLY Procedure Parameters

Parameter	Description
apply_result_table_name	Name of the table or view containing the results of an APPLY operation on the test data set (see Usage Notes)
case_id_column_name	Name of the case identifier column. This must be the same as the one used for generating APPLY results.
score_column_name	Name of the prediction column in the apply results table
<pre>score_criterion_column_n ame</pre>	Name of the probability column in the apply results table
<pre>ranked_apply_result_tab_ name</pre>	Name of the table containing the ranked apply results
top_N	Top N predictions to be considered from the \mathtt{APPLY} results for precision recall computation
cost_matrix_table_name	Name of the cost matrix table
apply_result_schema_name	Name of the schema hosting the APPLY results table
cost_matrix_schema_name	Name of the schema hosting the cost matrix table

Usage Notes

You can use RANK_APPLY to generate ranked apply results, based on a top-N filter and also with application of cost for predictions, if the model was built with costs.

The behavior of RANK_APPLY is similar to that of APPLY with respect to other DDL-like operations such as CREATE_MODEL, DROP_MODEL, and RENAME_MODEL. The procedure does not depend on the model; the only input of relevance is the apply results generated in a fixed schema table from APPLY.

The main intended use of RANK_APPLY is for the generation of the final APPLY results against the scoring data in a production setting. You can apply the model against test data using APPLY, compute various test metrics against various cost matrix tables, and use the candidate cost matrix for RANK APPLY.

The schema for the apply results from each of the supported algorithms is listed in subsequent sections. The <code>case_id</code> column will be the same case identifier column as that of the apply results.

Classification Models — NB and SVM

For numerical targets, the ranked results table will have the definition as shown:

(case_id VARCHAR2/NUMBER,
prediction NUMBER,
probability NUMBER,
cost NUMBER,
rank INTEGER)

For categorical targets, the ranked results table will have the following definition:



```
(case_id VARCHAR2/NUMBER,
prediction VARCHAR2,
probability NUMBER,
cost NUMBER,
rank INTEGER)
```

Clustering Using k-Means or O-Cluster

Clustering is an unsupervised machine learning function, and hence there are no targets. The results of an APPLY operation contains simply the cluster identifier corresponding to a case, and the associated probability. Cost matrix is not considered here. The ranked results table will have the definition as shown, and contains the cluster ids ranked by top-N.

```
(case_id VARCHAR2/NUMBER,
cluster_id NUMBER,
probability NUMBER,
rank INTEGER)
```

Feature Extraction using NMF

Feature extraction is also an unsupervised machine learning function, and hence there are no targets. The results of an APPLY operation contains simply the feature identifier corresponding to a case, and the associated match quality. Cost matrix is not considered here. The ranked results table will have the definition as shown, and contains the feature ids ranked by top-N.

```
(case_id VARCHAR2/NUMBER,
feature_id NUMBER,
match_quality NUMBER,
rank INTEGER)
```

Examples

```
BEGIN
/\star build a model with name census_model.
* (See example under CREATE MODEL)
/* if training data was pre-processed in any manner,
* perform the same pre-processing steps on apply
 * data also.
 * (See examples in the section on DBMS DATA MINING TRANSFORM)
/* apply the model to data to be scored */
DBMS DATA MINING.RANK APPLY(
 apply_result_table_name => 'census_apply_result',
 case_id_column_name => 'person_id',
score_column_name => 'prediction',
 score criterion column name => 'probability
 ranked apply result tab name => 'census ranked apply result',
                              => 3,
  cost_matrix_table_name => 'census_cost_matrix');
END;
-- View Ranked Apply Results
SELECT *
  FROM census ranked apply result;
```



REMOVE_COST_MATRIX Procedure

The REMOVE_COST_MATRIX procedure removes the default scoring matrix from a classification model.

See Also:

- "ADD_COST_MATRIX Procedure"
- "REMOVE_COST_MATRIX Procedure"

Syntax

```
DBMS_DATA_MINING.REMOVE_COST_MATRIX (
          model name IN VARCHAR2);
```

Parameters

Table 62-121 Remove_Cost_Matrix Procedure Parameters

Parameter	Description
model_name	Name of the model in the form [schema_name.]model_name. If you do not specify a schema, your own schema is used.

Usage Notes

If the model is not in your schema, then REMOVE_COST_MATRIX requires the ALTER ANY MINING MODEL system privilege or the ALTER object privilege for the machine learning model.

Example

The naive Bayes model NB_SH_CLAS_SAMPLE has an associated cost matrix that can be used for scoring the model.

```
SQL>SELECT *
    FROM TABLE(dbms_data_mining.get_model_cost_matrix('nb_sh_clas_sample'))
    ORDER BY predicted, actual;
```

ACTUAL	PREDICTED	COST
0	0	0
1	0	.75
0	1	.25
1	1	0

You can remove the cost matrix with REMOVE COST MATRIX.

```
SQL>EXECUTE dbms_data_mining.remove_cost_matrix('nb_sh_clas_sample');

SQL>SELECT *
     FROM TABLE(dbms_data_mining.get_model_cost_matrix('nb_sh_clas_sample'))
     ORDER BY predicted, actual;
no rows selected
```



RENAME_MODEL Procedure

This procedure changes the name of the machine learning model indicated by *model_name* to the name that you specify as *new_model_name*.

If a model with new_model_name already exists, then the procedure optionally renames new_model_name to versioned_model_name before renaming model_name to new_model_name.

The model name is in the form [schema_name.]model_name. If you do not specify a schema, your own schema is used. For machine learning model naming restrictions, see the Usage Notes for "CREATE_MODEL Procedure".

Syntax

Parameters

Table 62-122 RENAME_MODEL Procedure Parameters

Parameter	Description
model_name	Model to be renamed.
new_model_name	New name for the model model_name.
versioned_model_name	New name for the model <code>new_model_name</code> if it already exists.

Usage Notes

If you attempt to rename a model while it is being applied, then the model will be renamed but the apply operation will return indeterminate results.

Examples

1. This example changes the name of model census model to census model 2012.

```
BEGIN
   DBMS_DATA_MINING.RENAME_MODEL(
    model_name => 'census_model',
    new_model_name => 'census_model_2012');
END;
//
```

2. In this example, there are two classification models in the user's schema: clas_mod, the working model, and clas_mod_tst, a test model. The RENAME_MODEL procedure preserves clas mod as clas mod old and makes the test model the new working model.



