# **Getting Started**

This chapter discusses the compatibility of Oracle Java Database Connectivity (JDBC) driver versions, database versions, and Java Development Kit (JDK) versions.

It also describes the basics of testing a client installation and configuration and running a simple application. This chapter contains the following sections:

- Version Compatibility for Oracle JDBC Drivers
- Verifying a JDBC Client Installation
- · Basic Steps in JDBC
- Sample: Connecting\_ Querying\_ and Processing the Results
- Support for JSON-Relational Duality Views
- · Support for Fixed Character Semantic
- Support for Java Virtual Threads
- Support for Annotations
- Support for Oracle True Cache
- Support for the Bequeath Protocol
- Support for Invisible Columns
- Support for Verifying JSON Data
- Support for Implicit Results
- Support for Lightweight Connection Validation
- Support for Deprioritization of Database Nodes
- Support for Oracle Connection Manager in Traffic Director Mode
- Stored Procedure Calls in JDBC Programs
- About Processing SQL Exceptions

# 2.1 RDBMS and JDK Version Compatibility for Oracle JDBC Drivers

Oracle Database Release 23ai JDBC drivers are certified with all the supported Oracle Database releases (23ai, 21c, and 19c).

The following table describes the JDBC and Oracle Database interoperability matrix or the certification matrix:

JDBC Driver Version	Database 23.x	Database 21.x	Database 19.x
JDBC 23	Yes	Yes	Yes
JDBC 21.x	Yes	Yes	Yes
JDBC 19.x	Yes	Yes	Yes



Oracle JDBC Drivers are always compliant to the latest JDK version for every new release. For some versions, JDBC drivers support multiple JDK versions. The following table describes the release-specific JDBC JAR files and supported JDK versions for various Oracle Database versions:



ojdbc8.jar support with JDK 11, JDK 17, and JDK 19 is limited only to the JDBC 4.2 APIs because ojdbc8.jar does not support JDBC 4.3 APIs.

Oracle JDBC Version	Release-Specific JDBC JAR File with Supported JDK Versions
23.x	ojdbc17.jar JDK 17, JDK 19, and JDK 21
	ojdbc11.jar with JDK 11
	ojdbc8.jar with JDK 8 and JDK 11
21.x	ojdbc11.jar with JDK 11, JDK 17, and JDK 19 ojdbc8.jar with JDK 8 and JDK 11
19.x	ojdbc10.jar with JDK 11 and JDK 17 ojdbc8.jar with JDK 8, JDK 11, JDK 17, and JDK 19

### **Related Topics**

- Oracle Universal Connection Pool Developer's Guide
- Oracle JDBC FAQ

## 2.2 Verifying a JDBC Client Installation

This section describes the steps that you must perform to verify a JDBC client installation.

- Checking the Environment Variables
- Ensuring that the Java Code Can Be Compiled and Run
- Determining the Version of the JDBC Driver
- Testing the JDBC and Database Connection

### Note:

- If you use the JDBC Thin driver, then there is no additional installation on the client computer. If you use the JDBC Oracle Call Interface (OCI) driver, then you must also install the Oracle client software. This includes Oracle Net and the OCI libraries.
- The JDBC Thin driver requires a TCP/IP listener to be running on the computer, where the database is installed.

This section describes the steps for verifying an Oracle client installation of the JDBC drivers, assuming that you have already installed the driver of your choice. Installation of an Oracle

JDBC driver is platform-specific. You must follow the installation instructions for the driver you want to install in your platform-specific documentation.

### 2.2.1 Checking the Environment Variables

This section describes the environment variables that you must set for the JDBC OCI driver and the JDBC Thin driver, focusing on Solaris, Linux, and Microsoft Windows platforms.

#### **JDBC Thin Driver**

You must set the CLASSPATH environment variable for using the JDBC Thin driver, in a way similar to the following:

```
jdbc/lib/ojdbc11.jar
jlib/orai18n.jar
```



If you use the JTA features and the JNDI features, then you must specify jta.jar and jndi.jar in your CLASSPATH environment variable.

#### **JDBC OCI Driver**

You must set the CLASSPATH environment variable for using the JDBC OCI driver, in a way similar to the following:

```
ORACLE_HOME/jdbc/lib/ojdbc11.jar
ORACLE HOME/jlib/orai18n.jar
```



If you use the JTA features and the JNDI features, then you must specify jta.jar and jndi.jar in your CLASSPATH environment variable.

To use the JDBC OCI driver, you must also set the value for the library path environment variable in a way similar to the following:

On Solaris or Linux, set the LD LIBRARY PATH environment variable as follows:

```
ORACLE HOME/lib
```

This directory contains the libocijdbc11.so shared object library.

On Microsoft Windows, set the PATH environment variable as follows:

```
ORACLE_HOME\bin
```

This directory contains the ocijdbc11.dl1 dynamic link library.

All of the JDBC OCI demonstration programs can be run in the Instant Client mode by including the JDBC OCI Instant Client data shared library on the library path environment variable.



#### Setting Permission for the Server-Side Thin Driver

To use the JDBC server-side Thin driver, you must set permissions because the JDBC server-side Thin driver opens a socket for its connection to the database. So, Oracle Database enforces the Java security model and performs a check for a <code>SocketPermission</code> object.

The following is an example of how the permission can be granted for the user HR:

```
CREATE ROLE jdbcthin;
CALL dbms_java.grant_permission('JDBCTHIN', 'java.net.SocketPermission', '*', 'connect');
GRANT jdbcthin TO HR;
```

Note that JDBCTHIN in the grant\_permission call must be in uppercase. The asterisk (\*) is a pattern. You can restrict the user by granting permission to connect to only specific computers or ports.

#### **Related Topics**

- Features Specific to JDBC OCI Driver
- Oracle Database Java Developer's Guide

### 2.2.2 Ensuring that the Java Code Can Be Compiled and Run

To further ensure that Java is set up properly on your client system, go to the samples directory under the <code>ORACLE\_HOME/jdbc/demo</code> directory. Now, type the following commands on the command line, one after the other, to see if the Java compiler and the Java interpreter run without error:

```
javac
java
```

Each of the preceding commands should display a list of options and parameters and then exit. Ideally, verify that you can compile and run a simple test program, such as jdbc/demo/samples/generic/SelectExample.

### 2.2.3 Determining the Version of the JDBC Driver

It is very important to use the correct version of the JDBC driver in your applications.

Use the following commands to determine the version of the JDBC driver:

```
java -jar ojdbc8.jarjava -jar ojdbc11.jarjava -jar ojdbc17.jar
```

You can also call the getDriverVersion method of the OracleDatabaseMetaData class, as shown in the following sample code:



```
OracleDataSource ods = new OracleDataSource();
  ods.setURL("jdbc:oracle:thin:HR/<password>@<host>:<service>");
  Connection conn = ods.getConnection();

// Create Oracle DatabaseMetaData object
  DatabaseMetaData meta = conn.getMetaData();

// gets driver info:
  System.out.println("JDBC driver version is " + meta.getDriverVersion());
}
```

### 2.2.4 Testing the JDBC and Database Connection

The samples directory contains sample programs for a particular Oracle JDBC driver. One of the programs, <code>JdbcCheckup.java</code>, is designed to test JDBC and the database connection. The program queries for the user name, password, and the name of the database to which you want to connect. The program connects to the database, queries for the string <code>"Hello World"</code>, and prints it to the screen.

Go to the samples directory, and compile and run the JdbcCheckup.java program. If the results of the query print without error, then your Java and JDBC installations are correct.

Although JdbcCheckup.java is a simple program, it demonstrates several important functions by performing the following:

- Imports the necessary Java classes, including JDBC classes
- Creates a DataSource instance
- Connects to the database
- Runs a simple query
- Prints the guery results to your screen

The JdbcCheckup.java program, which uses the JDBC OCI driver, is as follows:

```
\mbox{\ensuremath{^{\star}}} This sample can be used to check the JDBC installation.
 * Just run it and provide the connect information. It will select
 * "Hello World" from the database.
// You need to import the java.sql and JDBC packages to use JDBC
import java.sql.*;
import oracle.jdbc.*;
import oracle.jdbc.pool.OracleDataSource;
// We import java.io to be able to read from the command line
import java.io.*;
class JdbcCheckup
 public static void main(String args[]) throws SQLException, IOException
    // Prompt the user for connect information
    System.out.println("Please enter information to test connection to
                           the database");
    String user;
    String password;
    String database;
```

```
user = readEntry("user: ");
  int slash_index = user.indexOf('/');
  if (slash index != -1)
    password = user.substring(slash index + 1);
    user = user.substring(0, slash index);
  else
    password = readEntry("password: ");
  database = readEntry("database(a TNSNAME entry): ");
  System.out.print("Connecting to the database...");
  System.out.flush();
  System.out.println("Connecting...");
  // Open an OracleDataSource and get a connection
  OracleDataSource ods = new OracleDataSource();
  ods.setURL("jdbc:oracle:oci:@" + database);
  ods.setUser(user);
  ods.setPassword(password);
  Connection conn = ods.getConnection();
  System.out.println("connected.");
  // Create a statement
  Statement stmt = conn.createStatement();
  // Do the SQL "Hello World" thing
  ResultSet rset = stmt.executeQuery("select 'Hello World' from dual");
  while (rset.next())
    System.out.println(rset.getString(1));
  // close the result set, the statement and the connection
  rset.close();
  stmt.close();
  conn.close();
  System.out.println("Your JDBC installation is correct.");
// Utility function to read a line from standard input
static String readEntry(String prompt)
  try
    StringBuffer buffer = new StringBuffer();
    System.out.print(prompt);
    System.out.flush();
    int c = System.in.read();
    while (c != '\n' \&\& c != -1)
      buffer.append((char)c);
      c = System.in.read();
    return buffer.toString().trim();
  catch (IOException e)
    return "";
}
```

# 2.3 Basic Steps in JDBC

After verifying the JDBC client installation, you can start creating your JDBC applications. When using Oracle JDBC drivers, you must include certain driver-specific information in your programs. This section describes, in the form of a tutorial, where and how to add the information. The tutorial guides you through the steps to create code that connects to and queries a database from the client.

You must write code to perform the following tasks:

- 1. Importing Packages
- 2. Opening a Connection to a Database
- 3. Creating a Statement Object
- 4. Running a Query and Retrieving a Result Set Object
- 5. Processing the Result Set Object
- 6. Closing the Result Set and Statement Objects
- 7. Making Changes to the Database
- 8. About Committing Changes
- 9. Closing the Connection



You must supply Oracle driver-specific information for the first three tasks that enable your program to use the JDBC application programming interface (API) to access a database. For the other tasks, you can use standard JDBC Java code, as you would for any Java application.

### 2.3.1 Importing Packages

Regardless of which Oracle JDBC driver you use, include the import statements shown in Table 2-1 at the beginning of your program using the following syntax:

import <package\_name>;

Table 2-1 Import Statements for JDBC Driver

Import statement	Provides
<pre>import java.sql.*;</pre>	Standard JDBC packages.
<pre>import java.math.*;</pre>	The BigDecimal and BigInteger classes. You can omit this package if you are not going to use these classes in your application.
<pre>import oracle.jdbc.*;</pre>	Oracle extensions to JDBC. This is optional.
<pre>import oracle.jdbc.pool.*;</pre>	OracleDataSource.
<pre>import oracle.sql.*;</pre>	Oracle type extensions. This is optional.



The Oracle packages listed as optional provide access to the extended functionality provided by Oracle JDBC drivers, but are not required for the example presented in this section.



It is better to import only the classes your application needs, rather than using the wildcard asterisk (\*). This guide uses the asterisk (\*) for simplicity, but this is not the recommended way of importing classes and interfaces.

### 2.3.2 Opening a Connection to a Database

First, you must create an <code>OracleDataSource</code> instance. Then, open a connection to the database using the <code>OracleDataSource.getConnection</code> method. The properties of the retrieved connection are derived from the <code>OracleDataSource</code> instance. If you set the URL connection property, then all other properties, including <code>TNSEntryName</code>, <code>DatabaseName</code>, <code>ServiceName</code>, <code>ServerName</code>, <code>PortNumber</code>, <code>Network</code> <code>Protocol</code>, and driver type are ignored.

#### Specifying a Database URL, User Name, and Password

The following code sets the URL, user name, and password for a data source:

```
OracleDataSource ods = new OracleDataSource();
ods.setURL(url);
ods.setUser(user);
ods.setPassword(password);
```

The following example connects user HR with password hr to a database with service orcl through port 5221 of the host myhost, using the JDBC Thin driver:

```
OracleDataSource ods = new OracleDataSource();
String url = "jdbc:oracle:thin:@myhost:5221/orcl";
ods.setURL(url);
ods.setUser("HR");
ods.setPassword("hr");
Connection conn = ods.getConnection();
```

### Note:

The user name and password specified in the arguments override any user name and password specified in the URL.

#### Specifying a Database URL that Includes User Name and Password

The following example connects user HR with password hr to a database host whose Transparent Network Substrate (TNS) entry is myTNSEntry, using the JDBC Oracle Call Interface (OCI) driver. In this case, the URL includes the user name and password and is the only input parameter.

```
String url = "jdbc:oracle:oci:HR/<password>@myTNSEntry");
ods.setURL(url);
Connection conn = ods.getConnection();
```



If you want to connect using the Thin driver, then you must specify the port number. For example, if you want to connect to the database on the host myhost that has a TCP/IP listener on port 5221 and the service identifier is orcl, then provide the following code:

```
String URL = "jdbc:oracle:thin:HR/<password>@myhost:5221/orcl");
ods.setURL(URL);
Connection conn = ods.getConnection();
```

#### **Related Topics**

- Data Sources and URLs
- Data Sources and URLs

### 2.3.3 Creating a Statement Object

Once you connect to the database and, in the process, create a <code>Connection</code> object, the next step is to create a <code>Statement</code> object. The <code>createStatement</code> method of the JDBC <code>Connection</code> object returns an object of the JDBC <code>Statement</code> type. To continue the example from the previous section, where the <code>Connection</code> object <code>conn</code> was created, here is an example of how to create the <code>Statement</code> object:

Statement stmt = conn.createStatement();

### 2.3.4 Running a Query and Retrieving a Result Set Object

To query the database, use the executeQuery method of the Statement object. This method takes a SQL statement as input and returns a JDBC ResultSet object.

### Note:

- The method used to execute a Statement object depends on the type of SQL statement being executed. If the Statement object represents a SQL query returning a ResultSet object, the executeQuery method should be used. If the SQL is known to be a DDL statement or a DML statement returning an update count, the executeUpdate method should be used. If the type of the SQL statement is not known, the execute method should be used.
- In case of a standard JDBC driver, if the SQL string being executed does not return a ResultSet object, then the executeQuery method throws a SQLException exception. In case of an Oracle JDBC driver, the executeQuery method does not throw a SQLException exception even if the SQL string being executed does not return a ResultSet object.

To continue the example, once you create the Statement object stmt, the next step is to run a query that returns a ResultSet object with the contents of the first\_name column of a table of employees named EMPLOYEES:

ResultSet rset = stmt.executeQuery ("SELECT first name FROM employees");



### 2.3.5 Processing the Result Set Object

Once you run your query, use the next() method of the ResultSet object to iterate through the results. This method steps through the result set row by row, detecting the end of the result set when it is reached.

To pull data out of the result set as you iterate through it, use the appropriate <code>getXXX</code> methods of the <code>ResultSet</code> object, where <code>XXX</code> corresponds to a Java data type.

For example, the following code will iterate through the ResultSet object, rset, from the previous section and will retrieve and print each employee name:

```
while (rset.next())
    System.out.println (rset.getString(1));
```

The next() method returns false when it reaches the end of the result set. The employee names are materialized as Java String values.

### 2.3.6 Closing the Result Set and Statement Objects

You must explicitly close the <code>ResultSet</code> and <code>Statement</code> objects after you finish using them. This applies to all <code>ResultSet</code> and <code>Statement</code> objects you create when using Oracle JDBC drivers. The drivers do not have finalizer methods. The cleanup routines are performed by the <code>close</code> method of the <code>ResultSet</code> and <code>Statement</code> classes. If you do not explicitly close the <code>ResultSet</code> and <code>Statement</code> objects, serious memory leaks could occur. You could also run out of cursors in the database. Closing both the result set and the statement releases the corresponding cursor in the database. If you close only the result set, then the cursor is not released.

For example, if your ResultSet object is rset and your Statement object is stmt, then close the result set and statement with the following lines of code:

```
rset.close();
stmt.close();
```

When you close a Statement object that a given Connection object creates, the connection itself remains open.



Typically, you should put close statements in a finally clause.

## 2.3.7 Making Changes to the Database

#### **DML Operations**

To perform DML (Data Manipulation Language) operations, such as INSERT or UPDATE operations, you can create either a Statement object or a PreparedStatement object. PreparedStatement objects enable you to run a statement with varying sets of input parameters. The prepareStatement method of the JDBC Connection object lets you define a statement that takes variable bind parameters and returns a JDBC PreparedStatement object with your statement definition.

Use the setXXX methods on the PreparedStatement object to bind data to the prepared statement to be sent to the database.

The following example shows how to use a prepared statement to run INSERT operations that add two rows to the EMPLOYEES table.

```
// Prepare to insert new names in the EMPLOYEES table
PreparedStatement pstmt = null;
   pstmt = conn.prepareStatement ("insert into EMPLOYEES (EMPLOYEE_ID, FIRST_NAME)
values (?, ?)");
   // Add LESLIE as employee number 1500
   pstmt.setInt (1, 1500); // The first ? is for EMPLOYEE ID
   pstmt.setString (2, "LESLIE"); // The second ? is for FIRST_NAME
   // Do the insertion
   pstmt.execute();
   // Add MARSHA as employee number 507
   pstmt.setString (2, "MARSHA"); // The second ? is for FIRST NAME
   // Do the insertion
   pstmt.execute();
finally{
      if(pstmt!=null)
   // Close the statement
   pstmt.close();
}
```

### **DDL Operations**

To perform data definition language (DDL) operations, you must create a Statement object. The following example shows how to create a table in the database:

### Note:

You can also use a PreparedStatement object to perform DDL operations. However, you should not use a PreparedStatement object because the useful part of such an object is that it can have parameters and a DDL operation does not have any parameters.

Also, due to a Database limitation, if you use a PreparedStatement object for a DDL operation, then it only works for the first time it is executed. So, you should use only Statement objects for DDL operations.

The following example shows how to prepare your DDL statements before any reexecution:

```
//
Statement stmt = null;
PreparedStatement pstmt = null;
   pstmt = conn.prepareStatement ("insert into EMPLOYEES (EMPLOYEE ID, FIRST NAME)
values (?, ?)");
   stmt = conn.createStatement("truncate table EMPLOYEES");
   // Add LESLIE as employee number 1500
   pstmt.setInt (1, 1500); // The first ? is for EMPLOYEE ID
   pstmt.setString (2, "LESLIE"); // The second ? is for FIRST NAME
   pstmt.execute();
   stmt.executeUpdate();
   // Add MARSHA as employee number 507
   pstmt.setInt (1, 507); // The first ? is for EMPLOYEE ID
   pstmt.setString (2, "MARSHA"); // The second ? is for FIRST NAME
   pstmt.execute();
   stmt.executeUpdate();
finally{
if (pstmt!=null)
    // Close the statement
    pstmt.close();
```

### **Related Topics**

- The setObject and setOracleObject Methods
- Other setXXX Methods

### 2.3.8 About Committing Changes

By default, data manipulation language (DML) operations are committed automatically as soon as they are run. This is known as the auto-commit mode. If auto-commit mode is on and you perform a COMMIT or ROLLBACK operation using the commit or rollback method on a connection object, then you get the following error messages:

Table 2-2 Error Messages for Operations Performed When Auto-Commit Mode is ON

Operation	Error Messages	
COMMIT	Could not commit with auto-commit set on	
ROLLBACK	Could not rollback with auto-commit set on	

If a SQLException is raised during a COMMIT or ROLLBACK operation with the error messages as mentioned in the preceding table, then check the auto-commit status of the connection because you get an exception when these operations are performed on a connection that has auto-commit value set to true.

This exception is raised for any one of the following cases:

- When auto-commit status is set to true and commit or rollback method is called
- When the default status of auto-commit is not changed and commit or rollback method is called
- When the value of the COMMIT\_ON\_ACCEPT\_CHANGES property is true and commit or rollback method is called after calling the acceptChanges method on a rowset

However, you can disable auto-commit mode with the following method call on the Connection object:

```
conn.setAutoCommit(false);
```

If you disable the auto-commit mode, then you must manually commit or roll back changes with the appropriate method call on the Connection object:

```
conn.commit();

Or:
conn.rollback();
```

A COMMIT or ROLLBACK operation affects all DML statements run since the last COMMIT or ROLLBACK.

### Note:

- If the auto-commit mode is disabled and you close the connection without explicitly committing or rolling back your last changes, then an implicit COMMIT operation is run.
- Any data definition language (DDL) operation always causes an implicit COMMIT.
   If the auto-commit mode is disabled, then this implicit COMMIT will commit any pending DML operations that had not yet been explicitly committed or rolled back.

### **Related Topics**

Disabling Auto-Commit Mode

### 2.3.8.1 Changing Commit Behavior

When a transaction updates the database, it generates a redo entry corresponding to this update. Oracle Database buffers this redo in memory until the completion of the transaction. When you commit the transaction, the Log Writer (LGWR) process writes the redo entry for the commit to disk, along with the accumulated redo entries of all changes in the transaction. By default, Oracle Database writes the redo to disk before the call returns to the client. This behavior introduces latency in the commit because the application must wait for the redo entry to be persisted on disk.

If your application requires very high transaction throughput and you are willing to trade commit durability for lower commit latency, then you can change the behavior of the default COMMIT operation, depending on the needs of your application. You can change the behavior of the COMMIT operation with the following options:

- WAIT
- NOWAIT
- WRITEBATCH
- WRITEIMMED

These options let you control two different aspects of the commit phase:

- Whether the COMMIT call should wait for the server to process it or not. This is achieved by using the WAIT or NOWAIT option.
- Whether the Log Writer should batch the call or not. This is achieved by using the WRITEIMMED or WRITEBATCH option.

You can also combine different options together. For example, if you want the COMMIT call to return without waiting for the server to process it and also the log writer to process the commits in batch, then you can use the NOWAIT and WRITEBATCH options together. For example:

```
((OracleConnection)conn).commit(
    EnumSet.of(
        OracleConnection.CommitOption.WRITEBATCH,
        OracleConnection.CommitOption.NOWAIT));
```

### Note:

you cannot use the WAIT and NOWAIT options together because they have opposite meanings. If you do so, then the JDBC driver will throw an exception. The same applies to the WRITEIMMED and WRITEBATCH options.

### 2.3.9 Closing the Connection

You must close the connection to the database after you have performed all the required operations and no longer require the connection. You can close the connection by using the close method of the Connection object, as follows:

```
conn.close();
```





Typically, you should put close statements in a finally clause.

# 2.4 Sample: Connecting, Querying, and Processing the Results

The steps in the preceding sections are illustrated in the following example, which uses the Oracle JDBC Thin driver to create a data source, connects to the database, creates a Statement object, runs a query, and processes the result set.

Note that the code for creating the Statement object, running the query, returning and processing the ResultSet object, and closing the statement and connection uses the standard JDBC API.

```
import java.sql.Connection;
import java.sql.ResultSet;
import java.sql.Statement;
import java.sql.SQLException;
import oracle.jdbc.pool.OracleDataSource;
class JdbcTest
  public static void main (String args []) throws SQLException
OracleDataSource ods = null;
Connection conn = null;
Statement stmt = null;
ResultSet rset = null;
      // Create DataSource and connect to the local database
      ods = new OracleDataSource();
      ods.setURL("jdbc:oracle:thin:@localhost:5221/orcl");
      ods.setUser("HR");
      ods.setPassword("hr");
      conn = ods.getConnection();
try
      // Query the employee names
      stmt = conn.createStatement ();
      rset = stmt.executeQuery ("SELECT first name FROM employees");
      // Print the name out
      while (rset.next ())
         System.out.println (rset.getString (1));
      //Close the result set, statement, and the connection
finally{
      if(rset!=null) rset.close();
      if(stmt!=null) stmt.close();
      if(conn!=null) conn.close();
}
```

If you want to adapt the code for the OCI driver, then replace the call to the OracleDataSource.setURL method with the following:

```
ods.setURL("jdbc:oracle:oci:@MyHostString");
```

where, MyHostString is an entry in the TNSNAMES.ORA file.

# 2.5 Support for JSON-Relational Duality Views

Duality views combine the advantages of using JSON documents with those of the relational model, while avoiding the limitations of each. JSON-relational duality underpins collections of documents with relational storage: active, updatable, hierarchical documents are based on a foundation of normalized relations.



Oracle Database JSON-Relational Duality Developer's Guide

You can retrieve information about JSON-relational duality views using the following JDBC methods:

- isDualityView()
- getJsonSchema()

### See Also:

Oracle Database JDBC Java API Reference

The following code snippet shows how to use these methods:

### See Also:

oracle-db-examples folder on GitHub for more examples

# 2.6 Support for Fixed Character Semantic

Starting from Oracle Database Release 23ai, you can change the default parameter mapping of a Java String to its Oracle Type, using the new connection property oracle.jdbc.mapStringParameterToCHAR.

The default value of this property is false. If you set this property to true, then the JDBC driver uses fixed character semantic when the setString method is called. Calls to the setObject method, with a parameter of type String, also uses a fixed character semantic. This is essential because when you run a query using collation or Extended Binary Coded Decimal Interchange Code (EBCDIC), you must typecast a Java String to CHAR instead VARCHAR2 because VARCHAR2 is incompatible with DB2 VARCHAR.



Oracle Database JDBC Java API Reference for more information about this property

# 2.7 Support for Java Virtual Threads

Starting from Oracle Database Release 21c, JDBC drivers support virtual threads. A virtual thread is an instance of <code>java.lang.Thread</code>, which is not tied to a specific operating system (OS) thread.

A virtual thread too runs the code on an OS thread; however, when the code running in a virtual thread calls a blocking I/O operation, then the Java run time suspends the virtual thread until it can be resumed. So, you can use virtual threads for long-running tasks that are blocked most of the time, often waiting for I/O operations to complete.



Java SE Core Libraries

# 2.8 Support for Annotations

Annotations are a lightweight declarative facility for developers to centrally register usage properties for database schema objects.

You can add annotations to schema objects when you create new objects (using the CREATE statement) or modify existing objects (using the ALTER statement). An individual annotation has a name and an optional value, both of which are free-form text fields. A schema object can have multiple annotations.



Application Usage Annotations



Starting from Oracle Database Release 23ai, the JDBC drivers support annotations. You can use the following methods to work with annotations:

```
oracle.jdbc.AdditionalDatabaseMetaData.getAnnotations
(java.lang.String objectName, java.lang.String domainName,
java.lang.String domainOwner) throws java.sql.SQLException

oracle.jdbc.OracleResultSetMetaData.getAnnotations
(java.lang.String objectName, java.lang.String columnName,
java.lang.String domainName, java.lang.String domainOwner) throws
java.sql.SQLException
```

These methods return the annotations associated with the specified table or view. If there is no annotation available for the specified object, then it returns null.

# 2.9 Support for Oracle True Cache

Oracle True Cache is an in-memory, consistent, and automatically managed cache for Oracle Database. It satisfies queries by using data only from its buffer cache.

True Cache is a fully functional, read-only replication of the primary database, which is mostly diskless. It exploits the fact that the applications rarely need the most current data, and can use the cached data instead. Queries that use the cached data, can be issued to a True Cache instance that is in the middle tier. So, the applications need to maintain two data sources, one to the primary database and the other to the True Cache instance.



Oracle Database Oracle True Cache User's Guide

When configured, the JDBC driver can execute queries on both the True Cache instance and the primary database. Applications maintain only one logical connection, which is aware of both the primary database and the True Cache instance. A query is executed on the True Cache instance if the JDBC driver logical connection is in a read-only mode, otherwise, the query is executed on the primary database. This improves the scalability and the application-response-time because the number of queries sent to the primary database is reduced.

For enabling the True Cache functionality, you must set the value of the new oracle.jdbc.useTrueCacheDriverConnection property to true. Once you enable the True Cache functionality, the JDBC driver uses the standard java.sql.Connection.isReadOnly(boolean) and java.sql.Connection.isReadOnly() methods to mark a connection as read-only. By default, the read-only mode is false for a connection.

The following code snippet shows how to use the True Cache functionality:

```
...
OracleDataSource ods = new oracle.jdbc.pool.OracleDataSource();
ods.setURL(DB_URL);
ods.setUser(DB_USER);
ods.setPassword(DB_PASSWORD);
Properties props = new Properties();
```



```
props.setProperty("oracle.jdbc.useTrueCacheDriverConnection", "true");
ods.setConnectionProperties(props);
// this is a True Cache driver connection and it can be used to execute
// queries on both the primary database and a True Cache instance
Connection conn = ods.getConnection();
Statement stmt = conn.createStatement();
// Default value of connection read-only flag is false which means
// the SQL_QUERY1 is executed on the primary database
ResultSet rs = stmt.executeQuery(SQL_QUERY1);
// set the read-only flag to true, in order to execute SQL_QUERY2
// on a True Cache instance
conn.setReadOnly(true);
ResultSet rs1 = stmt.executeQuery(SQL_QUERY2);
...
```

### See Also:

- setReadOnly(boolean) method
- isReadOnly method

## 2.10 Support for the Bequeath Protocol

Starting from Oracle Database Release 23ai, the JDBC thin driver supports the Bequeath protocol (BEQ) for applications running on Linux platforms.

To connect to the Database using the Bequeath Protocol, you must set the value of the ORACLE\_HOME variable, so that the driver can locate the Oracle server process executable. Typically, the ORACLE\_HOME variable points to the database installation location, that is, /var/lib/oracle/dbhome. You can reset the location in the following two ways:

- In the connection URL
- In the environment of the current application

The second mandatory variable, which you must enable, is the <code>ORACLE\_SID</code>. Similar to setting the <code>ORACLE\_HOME</code> variable, you can set the <code>ORACLE\_SID</code> in the connection URL or in the current application environment. To establish a bequeath connection, the BEQ protocol must be enabled, which is the default setting in the authentication services property.

The following example shows how you can set the <code>ORACLE\_HOME</code> variable and the <code>ORACLE\_SID</code> in the connection URL:

```
jdbc:oracle:thin:@(DESCRIPTION=(LOCAL=YES) (ADDRESS=(PROTOCOL=beq))
(ENVS=ORACLE_HOME=/var/lib/oracle/dbhome,ORACLE_SID=oraclesid))
```



Oracle JDBC Java API Reference

# 2.11 Support for Invisible Columns

Starting from this release, Oracle Database supports invisible columns. Using this feature, you can add a column to the table in hidden mode and make it visible later. JDBC provides APIs to retrieve information about invisible columns. To get information about whether a column is invisible or not, you can use the <code>isColumnInvisible</code> method available in the <code>oracle.jdbc.OracleResultSetMetaData</code> interface in the following way:

### **Example**

```
Connection conn = DriverManager.getConnection(jdbcURL, user, password);
Statement stmt = conn.createStatement ();
stmt.executeQuery ("create table hiddenColsTable (a varchar(20), b int invisible)");
stmt.executeUpdate("insert into hiddenColsTable (a,b) values('somedata',1)");
stmt.executeUpdate("insert into hiddenColsTable (a,b) values('newdata',2)");
System.out.println ("Invisible columns information");
try
     ResultSet rset = stmt.executeQuery("SELECT a, b FROM hiddenColsTable");
     OracleResultSetMetaData rsmd = (OracleResultSetMetaData)rset.getMetaData();
     while (rset.next())
       System.out.println("column1 value:" + rset.getString(1));
       System.out.println("Visibility:" + rsmd.isColumnInvisible(1));
       System.out.println("column2 value:" + rset.getInt(2));
        System.out.println("Visibility:" + rsmd.isColumnInvisible(2));
catch (Exception ex)
       System.out.println("Exception :" + ex);
       ex.printStackTrace();
```

Alternatively, you can also use the <code>getColumns</code> method available in the <code>oracle.jdbc.OracleDatabaseMetaData</code> class to retrieve information about invisible columns.

### **Example**

```
Connection conn = DriverManager.getConnection(jdbcURL, user, password);
Statement stmt = conn.createStatement ();
stmt.executeQuery ("create table hiddenColsTable (a varchar(20), b int invisible)");
stmt.executeUpdate("insert into hiddenColsTable (a,b) values('somedata',1)");
stmt.executeUpdate("insert into hiddenColsTable (a,b) values('newdata',2)");

System.out.println ("getColumns for table with invisible columns");
try
{
    DatabaseMetaData dbmd = conn.getMetaData();
    ResultSet rs = dbmd.getColumns(null, "HR", "hiddenColsTable", null);
    OracleResultSetMetaData rsmd = (OracleResultSetMetaData)rs.getMetaData();
    int colCount = rsmd.getColumnCount();
    System.out.println("colCount: " + colCount);
    String[] columnNames = new String [colCount];

for (int i = 0; i < colCount; ++i)</pre>
```

```
{
    columnNames[i] = rsmd.getColumnName (i + 1);
}

while (rs.next())
{
    for (int i = 0; i < colCount; ++i)
        System.out.println(columnNames[i] +":" +rs.getString (columnNames[i]));
}

catch (Exception ex)
{
    System.out.println("Exception: " + ex);
    ex.printStackTrace();
}</pre>
```

### Note:

The server-side internal driver, kprb does not support fetching information about invisible columns.

# 2.12 Support for Verifying JSON Data

Starting from Oracle Database Release 18c, JDBC drivers can verify whether a column returned in the ResultSet is a JSON column or not. To get information about whether a column is JSON or not, you can use the isColumnJSON method available in the oracle.jdbc.OracleResultSetMetaData interface in the following way:

### Example 2-1 Example

```
public void test (Connection conn)
     throws Exception{
    try {
     show ("tkpjb26776242 - start");
     createTable(conn);
     String sql = "SELECT col1, col2, col3, col4, col5, col6, col7, col8
FROM tkpjb26776242 tab";
     Statement stmt = conn.createStatement();
     ResultSet rs = stmt.executeQuery(sql);
     ResultSetMetaData rsmd = rs.getMetaData();
     OracleResultSetMetaData orsmd = (OracleResultSetMetaData)rsmd;
     int colCnt = orsmd.getColumnCount();
     show("Table has " + colCnt + " columns.");
     for (int i = 1; i <= colCnt; i++) {
        String columnName = orsmd.getColumnName(i);
        String typeName = orsmd.getColumnTypeName(i);
        boolean invisible = orsmd.isColumnInvisible(i);
```

```
boolean json = orsmd.isColumnJSON(i);
       show(columnName + " " + typeName + (invisible?" INVISIBLE":"") +
(json?"
         JSON":""));
     rs.close();
     stmt.close();
     show ("tkpjb26776242 - end");
   finally {
     dropTable(conn);
 }
 private void createTable(Connection conn) throws Exception{
   String sql = " create table tkpjb26776242 tab ( "
                + " col1 clob, "
                + " col2 clob , "
                + " col3 clob INVISIBLE, "
                 + " col4 clob INVISIBLE, "
                + " col5 varchar2(200), "
                + " col6 varchar2(200), "
                + " col7 varchar2(200) INVISIBLE, "
                + " col8 varchar2(200) INVISIBLE, "
                + " check (col2 IS JSON), "
                + " check (col4 IS JSON), "
                 + " check (col6 IS JSON), "
                 + " check (col8 IS JSON))";
   Util.doSQL(conn, sql);
 private void dropTable(Connection conn) throws Exception{
   String sql = " drop table tkpjb26776242 tab";
   Util.trySQL(conn, sql);
    . . .
```

# 2.13 Support for Implicit Results

Oracle Database supports results of SQL statements executed in a stored procedure to be returned implicitly to the client applications without the need to explicitly use a REF CURSOR.

You can use the following methods to retrieve and process the implicit results returned by PL/SQL procedures or blocks:

Method	Description	
getMoreResults	Checks if there are more results available in the result set	

Method	Description	
getMoreResults(int)	Checks if there are more results available in the result set, like the overloaded method. This method accepts an int parameter that can have one of the following values:	
	KEEP CURRENT RESULT	
	CLOSE ALL RESULTS	
	• CLOSE_CURRENT_RESULT	
getResultSet	Iteratively retrieves each implicit result from an executed PL/SQL statement	

### Note:

- The server-side internal driver, kprb does not support fetching information about implicit results.
- Only SELECT queries can be returned implicitly.
- Applications retrieve each result set sequentially, but can fetch rows from any result set independent of the sequence.

Suppose you have a procedure called foo as the following:

```
create procedure foo as
  c1 sys_refcursor;
  c2 sys_refcursor;
begin
  open c1 for select * from hr.employees;
  dbms_sql.return_result(c1); --return to client
  -- open 1 more cursor
  open c2 for select * from hr.departments;
  dbms_sql.return_result (c2); --return to client
end;
```

The following code snippet demonstrates how to retrieve the implicit results returned by PL/SQL procedures using the <code>getMoreResults</code> methods:

#### **Example 1**

```
}
```

#### Suppose you have another procedure called foo as the following:

```
create or replace procedure foo asc1 sys_refcursor; c2 sys_refcursor; c3 sys_refcursor;
begin    open c1 for 'select * from hr.employees';
dbms_sql.return_result (c1);-- cursor 2open c2 for 'select * from hr.departments';
dbms_sql.return_result (c2);-- cursor 3open c3 for 'select first_name from hr.employees';
dbms_sql.return_result (c3); end;
```

The following code snippet demonstrates how to retrieve the implicit results returned by PL/SQL procedures using the <code>qetMoreResults(int)</code> methods:

#### **Example 2**

```
String sql = "begin foo; end;";
Connection conn = DriverManager.getConnection(jdbcURL, user, password);
try {
       Statement stmt = conn.createStatement ();
       stmt.executeQuery (sql);
       ResultSet rs = null;
       boolean retval = stmt.getMoreResults(Statement.KEEP CURRENT RESULT))
       if (retval)
        {
            rs = stmt.getResultSet();
            System.out.println("ResultSet");
            while (rs.next())
                /* get results */
            }
        }
       /* closes open results */
       retval = stmt.getMoreResults(Statement.CLOSE_ALL_RESULTS);
       if (retval)
        {
            System.out.println("More ResultSet available");
            rs = stmt.getResultSet();
            System.out.println("ResultSet");
            while (rs.next())
                /* get results */
            }
        }
        /* close current result set */
       retval = stmt.getMoreResults(Statement.CLOSE CURRENT RESULT);
       if (retval)
            System.out.println("More ResultSet available");
            rs = stmt.getResultSet();
            while (rs.next())
                /* get Results */
            }
```

}

# 2.14 Support for Lightweight Connection Validation

Starting from Oracle Database Release 18c, JDBC Thin driver supports lightweight connection validation. Lightweight connection validation enables JDBC applications to verify connection validity by sending a zero length NS data packet that does not require a round-trip to the database.

For the releases of Oracle Database earlier than 18c, when you call the isValid(timeout) method to test the validity of a connection, Oracle JDBC driver uses a ping-pong protocol, which is an expensive operation as it makes a full round-trip to the database. Since Oracle Database Release 18c, the isValid(timeout) method instead sends an empty packet to the database and does not wait to receive it back. So, connection validation is faster, which results in better application performance.

Lightweight connection validation is disabled by default. To enable this feature, you must set the <code>oracle.jdbc.defaultConnectionValidation</code> connection property value to <code>SOCKET</code>. If this property is set, then the JDBC driver performs lightweight connection validation, when you call the <code>isValid(timeout)</code> method.

### Note:

- Lightweight connection validation checks only the underlying socket health. When the <code>isValid(timeout)</code> method returns <code>true</code>, that is, if a connection is termed as valid, this validation only guarantees that the server is not unreachable (dead socket). It does not provide any status about the server processes, like whether they are running or not. However, by default, that is, when lightweight connection validation is not enabled, the <code>isValid(timeout)</code> method does check whether the network between the client and the server is intact or not.
- Only the JDBC Thin driver supports this feature.

#### **New APIs for Lightweight Connection Validation**

oracle.jdbc.defaultConnectionValidation

This connection property specifies the level of connection validation. The possible values for this property are: NONE, LOCAL, SOCKET, NETWORK, SERVER, and COMPLETE. These values are case-sensitive, and setting any value other than these values throws an exception. The default value is NETWORK.

 public boolean isValid(ConnectionValidation validation\_level, int timeout) throws SQLException

The new variation of the existing <code>isValid(timeout)</code> method accepts two parameters: level of validation (validation\_level) and timeout. The first parameter specifies the level of connection validation.



#### **Example 2-2** Example of Lightweight Connection Validation

The following code snippet demonstrates how to implement lightweight connection mechanism:

```
OracleDataSource ods = new OracleDataSource();
ods.setURL(url);
ods.setUser(user);
ods.setPassword(password);
Connection conn = ods.getConnection();
try{
    boolean isValid = ((OracleConnection)conn).

isValid(ConnectionValidation.SOCKET,timeout);
    System.out.println("Connection isValid = "+isValid);
}
catch (Exception ex)
{
    System.out.println("Exception :" + ex);
    ex.printStackTrace();
}
```

# 2.15 Support for Deprioritization of Database Nodes

Starting from Oracle Database 12c Release 2 (12.2.0.1), JDBC drivers support deprioritization of database nodes. When a node fails, JDBC deprioritizes it for the next 10 minutes, which is the default expiry time. For example, if there are three nodes A, B, C, and node A is down, then connections are allocated first from nodes B and C, and then from node A. After the default expiry time, node A is no longer deprioritized, that is, connections are allocated from all the three nodes on availability basis. Also, during the default expiry time, if a connection attempt to node A succeeds, then node A is no longer considered to be a deprioritized node. You can specify the default expiry time for deprioritization using the oracle.net.DOWN HOSTS TIMEOUT system property.

For example, in the following URL, <code>scan\_listener0</code> has <code>ip1</code>, <code>ip2</code>, and <code>ip3</code> IP addresses configured, after retrieving its IP addresses. Now, if <code>ip1</code> is deprioritized, then the order of trying IP addresses will be <code>ip2</code>, <code>ip3</code>, and then <code>ip1</code>. If all IP addresses are unavailable, then the whole host is tried last, after trying <code>node 1</code> and <code>node 2</code>.

# 2.16 Support for Oracle Connection Manager in Traffic Director Mode

Starting from Oracle Database Release 18c, JDBC Drivers support Oracle Connection Manager in Traffic Director Mode, which is a proxy server that resides between the Database clients and the Database instances.

A JDBC client connects to the Oracle Connection Manager in Traffic Director Mode, which in turn connects to the target Oracle Database. The client sends requests in the form of Two-Task Common (TTC) messages that Oracle Connection Manager in Traffic Director Mode intercepts, parses, and then relays to the appropriate target database. Once the responses arrive from the database, Oracle Connection Manager in Traffic Director Mode transfers the responses back to the clients through TTC messages.

The following image illustrates the architecture of Oracle Connection Manager in Traffic Director Mode:

Figure 2-1 Architecture of Oracle Connection Manager in Traffic Director Mode



### See Also:

- Oracle Database Net Services Administrator's Guide for more information about configuring the cman.ora file to set up Oracle Connection Manager in Traffic Director Mode
- Oracle Database Net Services Reference for more information about Oracle Connection Manager in Traffic Director Mode parameters

This chapter describes the following concepts:

- Modes of Running Oracle Connection Manager in Traffic Director Mode
- Benefits of Oracle Connection Manager in Traffic Director Mode

# 2.16.1 Modes of Running Oracle Connection Manager in Traffic Director Mode

You can run Oracle Connection Manager in Traffic Director Mode in the following connection modes:

#### Pooled connection mode

The pooled connection mode uses a new feature called Proxy Resident Connection Pooling, which is a proxy-enabled mode of Database Resident Connection Pooling (DRCP). The Proxy Resident Connection Pooling reduces the connection load on the database as it multiplexes a large number of client connections over a fewer number of database connections. Any application using Oracle Database 12c Release 1 (12.1) JDBC drivers and later can use this connection mode.



The pooled connection mode yields best results when you use it with clients using DRCP-aware connection pools.

### Nonpooled or dedicated connection mode

You can use the nonpooled or dedicated connection mode with applications using any supported Oracle Database JDBC driver. However, some capabilities, such as connection multiplexing, are not available in this mode.

### **Related Topics**

Overview of Database Resident Connection Pooling

### See Also:

- Database Admin Guide
- Universal Connection Pool Developer's Guide

### 2.16.2 Benefits of Oracle Connection Manager in Traffic Director Mode

This section describes how Oracle Connection Manager in Traffic Director Mode provides benefits to your applications.

- Transparent performance enhancements: Oracle Connection Manager in Traffic
  Director Mode auto enables statement caching, row prefetching, and result set caching
  during both pooled and nonpooled modes.
- Connection multiplexing: Using Proxy Resident Connection Pooling (PRCP), Oracle Connection Manager in Traffic Director Mode (for the pooled mode) provides transparent



connection-time load balancing and run-time load balancing with the database. If you use multiple instances of Oracle Connection Manager in Traffic Director Mode, then you can increase the scalability of your application through the implementation of client-side connection-time load balancing or a load balancer like BIG-IP or NGINX.

- Zero application downtime: Oracle Connection Manager in Traffic Director Mode
  provides zero application downtime during planned database maintenance as well as
  unplanned database outages. For Unplanned database outages, it offers zero application
  downtime for read-mostly workloads. For planned database maintenance or pluggable
  database (PDB) relocation, it uses different techniques for the pooled mode and the
  nonpooled mode, as described in this section.
  - Pooled mode:

For planned outages, Oracle Connection Manager in Traffic Director Mode responds to the Oracle Notification Service (ONS) events. It uses database connections from the proxy resident connection pool and redirects the requests to the appropriate databases. When the requests complete, it drains the connections from the pool.

For PDB relocations, Oracle Connection Manager in Traffic Director Mode uses an inband client notification mechanism, which works even when ONS is not configured. This feature is available only for Oracle Database release 18c and later.

Nonpooled or dedicated mode

In this mode, Oracle Connection Manager in Traffic Director Mode leverages either of the following features of Oracle Database:

- \* Continuous application availability to stop the service at the request boundary
- \* Transparent Application Failover (TAF) to reconnect and restore simple states
- High Availability: Oracle Connection Manager in Traffic Director Mode implements the following techniques to avoid single point of failure, and in turn, assures high availability:
  - Multiple instances of Oracle Connection Manager in Traffic Director Mode use a load balancer or client-side load balancing in the connection string.
  - Oracle Connection Manager in Traffic Director Mode instances support rolling upgrade.
  - For planned outages, Oracle Connection Manager in Traffic Director Mode implements graceful close of existing connections from client.
  - Oracle Connection Manager in Traffic Director Mode sends in-band notifications to Oracle Database release 18c and later clients, and ONS notifications to all supported clients prior to Oracle Database release 18c.
- Security: Oracle Connection Manager in Traffic Director Mode provides security to your applications in the following ways:
  - It supports the Transmission Control Protocol Secure (TCPS) protocol.
  - It creates a firewall based on the IP address, service name, and Transport Layer Security (TLS) wallets.
  - It provides protection against denial-of-service and fuzzing attacks.
  - It provides secure tunneling of database traffic across Oracle Database on-premises and Oracle Cloud.
- **Tenant isolation:** Oracle Connection Manager in Traffic Director Mode provides tenant isolation for increased memory and enhanced processing power.



# 2.17 Stored Procedure Calls in JDBC Programs

This section describes how Oracle JDBC drivers support the following kinds of stored procedures:

- PL/SQL Stored Procedures
- Java Stored Procedures

### 2.17.1 PL/SQL Stored Procedures

JDBC supports the invocation of PL/SQL procedures/functions and anonymous blocks, using either JDBC escape syntax or PL/SQL block syntax. The following PL/SQL calls would work with any Oracle JDBC driver:

As an example of using the Oracle syntax, here is a PL/SQL code snippet that creates a stored function. The PL/SQL function gets a character sequence and concatenates a suffix to it:

```
create or replace function foo (val1 char)
return char as
begin
   return val1 || 'suffix';
end:
```

The function invocation in your JDBC program should look like the following:

```
OracleDataSource ods = new OracleDataSource();
ods.setURL("jdbc:oracle:oci:@<hoststring>");
ods.setUser("HR");
ods.setPassword("hr");
Connection conn = ods.getConnection();

CallableStatement cs = conn.prepareCall ("begin ? := foo(?); end;");
cs.registerOutParameter(1, Types.CHAR);
cs.setString(2, "aa");
cs.execute();
String result = cs.getString(1);
```

### 2.17.2 Java Stored Procedures

You can use JDBC to call Java stored procedures through the SQL interface. The syntax for calling Java stored procedures is the same as the syntax for calling PL/SQL stored procedures, presuming they have been properly published. That is, you have written call specifications to publish them to the Oracle data dictionary. Applications can call Java stored procedures using the Native Java Interface for direct invocation of static Java methods.

# 2.18 About Processing SQL Exceptions

To handle error conditions, Oracle JDBC drivers throw SQL exceptions, producing instances of the <code>java.sql.SQLException</code> class or its subclass. Errors can originate either in the JDBC driver or in the database itself. Resulting messages describe the error and identify the method that threw the error. Additional run-time information can also be appended.

JDBC 3.0 defines only a single exception, SQLException. However, there are large categories of errors and it is useful to distinguish them. Therefore, in JDBC 4.0, a set of subclasses of the SQLException exception is introduced to identify the different categories of errors.

Basic exception handling can include retrieving the error message, retrieving the error code, retrieving the SQL state, and printing the stack trace. The SQLException class includes functionality to retrieve all of this information, when available.

### **Retrieving Error Information**

You can retrieve basic error information with the following methods of the SQLException class:

- getMessage class includes functionality to retrieve all of this information, when available.
- getErrorCode class includes functionality to retrieve all of this information, when available.
- getSQLState class includes functionality to retrieve all of this information, when available.

The following example prints output from a getMessage method call:

```
catch(SQLException e)
{
    System.out.println("exception: " + e.getMessage());
}
```

This would print the output, such as the following, for an error originating in the JDBC driver:

```
exception: Invalid column type
```



Error message text is available in alternative languages and character sets supported by Oracle.

### **Printing the Stack Trace**

The SQLException class provides the printStackTrace() method for printing a stack trace. This method prints the stack trace of the Throwable object to the standard error stream. You can also specify a java.io.PrintStream object or java.io.PrintWriter object for output.

The following code fragment illustrates how you can catch SQL exceptions and print the stack trace.

```
try { <some code> }
catch(SQLException e) { e.printStackTrace (); }
```

To illustrate how the JDBC drivers handle errors, assume the following code uses an incorrect column index:

```
// Iterate through the result and print the employee names
// of the code

try {
  while (rset.next ())
      System.out.println (rset.getString (5)); // incorrect column index
}
catch(SQLException e) { e.printStackTrace (); }
```

Assuming the column index is incorrect, running the program would produce the following error text:

```
java.sql.SQLException: Invalid column index
at oracle.jdbc.OracleDriver.OracleResultSetImpl.getDate(OracleResultSetImpl.java:1556)
at Employee.main(Employee.java:41)
```

#### **Related Topics**

- JDBC Error Messages
  - All the JDBC error messages are now listed in the Database Error Portal.
- Oracle Database Error Messages Reference