

# 6

## Document-Identifier Field for Duality Views

A document supported by a duality view always includes, at its top level, a **document-identifier** field, `_id`, which corresponds to the *identifying columns* (primary-key columns, identity columns, or columns with a unique constraint or unique index) of the *root* table underlying the view. The field value can take different forms.

(An **identity column** is one whose numeric value is generated automatically and uniquely for each table row. You declare it using keywords `GENERATED BY DEFAULT ON NULL AS IDENTITY`.)

Often there is only one such identifying column and it is often a primary-key column. If there is more than one primary-key column then we sometimes speak of the primary key being **composite**.

- If there is only *one identifying column* then you use that as the value of field `_id` when you define the duality view.
- Alternatively, you can use an object as the value of field `_id`. The members of the object specify fields whose values are the identifying columns. An error is raised if there is not a field for each of the identifying columns.

If there is only one identifying column, you can nevertheless use an object value for `_id`; doing so lets you provide a meaningful field name.

### Note:

A duality view must have an `_id` field at its top level, to uniquely identify a given row of its root table, and thus the corresponding document.

In order to *replicate* a duality view, you also need to ensure that each document *subobject*<sup>1</sup> has a top-level field whose value is the identifying columns for that table. That is, the columns corresponding to such a table **row-identifier** field need to uniquely identify a row of the table that underlies that subobject.

Just as for a document-identifier field, the columns corresponding to a row-identifier field can be primary-key columns, identity columns, or columns with a unique constraint or unique index, for their table.

A document-identifier field must be named `_id`. A row-identifier field can have any name, but if its name is `_id` then it's up to *you to ensure* that the corresponding columns uniquely identify a table row.

### Example 6-1 Document Identifier Field `_id` With Primary-Key Column Value

For duality view `race_dv`, the value of a single primary-key column, `race_id`, is used as the value of field `_id`. A document supported by the view would look like this: `{"_id" : 1,...}`.

<sup>1</sup> This of course doesn't apply to document subobjects that are explicitly present a JSON-type column that's embedded in the document.

**GraphQL:**

```
CREATE JSON RELATIONAL DUALITY VIEW race_dv AS
  race {_id      : race_id
        name     : name
        laps     : laps @NOUPDATE
        date     : race_date
        podium   : podium @NOCHECK,
        result   : ...};
```

**SQL:**

```
CREATE JSON RELATIONAL DUALITY VIEW race_dv AS
  SELECT JSON {'_id'      : r.race_id,
              'name'     : r.name,
              'laps'     : r.laps WITH NOUPDATE,
              'date'     : r.race_date,
              'podium'   : r.podium WITH NOCHECK,
              'result'   : ...}
  FROM race r;
```

**Example 6-2 Document Identifier Field \_id With Object Value**

For duality view `race_dv`, the value of field `_id` is an object with a single member, which maps the single primary-key column, `race_id`, to a meaningful field name, `raceId`. A document supported by the view would look like this: `{"_id" : {"raceId" : 1}, ...}`.

**GraphQL:**

```
CREATE JSON RELATIONAL DUALITY VIEW race_dv AS
  race {_id @nest {race_id}
        name     : name
        laps     : laps @NOUPDATE
        date     : race_date
        podium   : podium @NOCHECK,
        result   : ...};
```

**SQL:**

```
CREATE JSON RELATIONAL DUALITY VIEW race_dv AS
  SELECT JSON {'_id'      : {'raceId' : r.race_id},
              'name'     : r.name,
              'laps'     : r.laps WITH NOUPDATE,
              'date'     : r.race_date,
              'podium'   : r.podium WITH NOCHECK,
              'result'   : ...}
  FROM race r;
```

An *alternative* car-racing design might instead use a `race` table that has *multiple* identifying columns, `race_id` and `date`, which together identify a row. In that case, a document supported by the view would look like this: `{"_id" : {"raceId" : 1, "date" : "2022-03-20T00:00:00"}, ...}`.

**GraphQL:**

```
CREATE JSON RELATIONAL DUALITY VIEW race_dv AS
  race {_id @nest {raceId: race_id, date: race_date}
        name      : name
        laps       : laps @NOUPDATE
        podium     : podium @NOCHECK,
        result     : ...};
```

**SQL:**

```
CREATE JSON RELATIONAL DUALITY VIEW race_dv AS
  SELECT JSON {'_id'      : {'raceId' : r.race_id, 'date' : r.race_date},
              'name'     : r.name,
              'laps'     : r.laps WITH NOUPDATE,
              'podium'   : r.podium WITH NOCHECK,
              'result'   : ...}
  FROM race r;
```

**Related Topics**

- [Car-Racing Example, JSON Documents](#)  
The car-racing example has three kinds of documents: a team document, a driver document, and a race document.

**See Also:**

Mongo DB API Collections Supported by JSON-Relational Duality Views  
in *Oracle Database API for MongoDB*