7

PL/SQL Static SQL

Static SQL is a PL/SQL feature that allows SQL syntax directly in a PL/SQL statement.

This chapter describes static SQL and explains how to use it.

Topics

- Description of Static SQL
- Cursors Overview
- · Processing Query Result Sets
- Cursor Variables
- CURSOR Expressions
- Transaction Processing and Control
- Autonomous Transactions



"Resolution of Names in Static SQL Statements"

Description of Static SQL

Static SQL has the same syntax as SQL, except as noted.

Topics

- Statements
- Pseudocolumns

Statements

These are the PL/SQL static SQL statements, which have the same syntax as the corresponding SQL statements, except as noted:

SELECT (this statement is also called a query)

For the PL/SQL syntax, see "SELECT INTO Statement".

- Data manipulation language (DML) statements:
 - INSERT

For the PL/SQL syntax, see "INSERT Statement Extension".

UPDATE

For the PL/SQL syntax, see "UPDATE Statement Extensions".

DELETE

For the PL/SQL syntax, see "DELETE Statement Extension".

MERGE (for syntax, see Oracle Database SQL Language Reference)



Oracle Database SQL Language Reference defines DML differently.

- Transaction control language (TCL) statements:
 - COMMIT (for syntax, see Oracle Database SQL Language Reference)
 - ROLLBACK (for syntax, see Oracle Database SQL Language Reference)
 - SAVEPOINT (for syntax, see Oracle Database SQL Language Reference)
 - SET TRANSACTION (for syntax, see Oracle Database SQL Language Reference)
- LOCK TABLE (for syntax, see Oracle Database SQL Language Reference)

A PL/SQL static SQL statement can have a PL/SQL identifier wherever its SQL counterpart can have a placeholder for a bind variable. The PL/SQL identifier must identify either a variable or a formal parameter.

To use PL/SQL identifiers for table names, column names, and so on, use the EXECUTE IMMEDIATE statement, explained in "Native Dynamic SQL"

Note:

After PL/SQL code runs a DML statement, the values of some variables are undefined. For example:

- After a FETCH or SELECT statement raises an exception, the values of the define variables after that statement are undefined.
- After a DML statement that affects zero rows, the values of the OUT bind variables are undefined, unless the DML statement is a BULK or multiple-row operation.

Example 7-1 Static SQL Statements

In this example, a PL/SQL anonymous block declares three PL/SQL variables and uses them in the static SQL statements INSERT, UPDATE, DELETE. The block also uses the static SQL statement COMMIT.



```
UPDATE employees_temp
SET first_name = 'Robert'
WHERE employee_id = emp_id;

DELETE FROM employees_temp
WHERE employee_id = emp_id
RETURNING first_name, last_name
INTO emp_first_name, emp_last_name;

COMMIT;
DBMS_OUTPUT.PUT_LINE (emp_first_name || ' ' || emp_last_name);
END;
//

Result:
Robert Henry
```

Pseudocolumns

A pseudocolumn behaves like a table column, but it is not stored in the table.

For general information about pseudocolumns, including restrictions, see *Oracle Database SQL Language Reference*.

Static SQL includes these SQL pseudocolumns:

- CURRVAL and NEXTVAL, described in "CURRVAL and NEXTVAL in PL/SQL".
- LEVEL, described in Oracle Database SQL Language Reference
- OBJECT VALUE, described in Oracle Database SQL Language Reference



ROWID, described in Oracle Database SQL Language Reference

```
See Also:
"Simulating CURRENT OF Clause with ROWID Pseudocolumn"
```

ROWNUM, described in Oracle Database SQL Language Reference

CURRVAL and NEXTVAL in PL/SQL

After a sequence is created, you can access its values in SQL statements with the CURRVAL pseudocolumn, which returns the current value of the sequence, or the NEXTVAL pseudocolumn, which increments the sequence and returns the new value.

To reference these pseudocolumns, use dot notation—for example, sequence name.CURRVAL.

Note:

Each time you reference <code>sequence_name.NEXTVAL</code>, the sequence is incremented immediately and permanently, whether you commit or roll back the transaction.

You can use <code>sequence_name.CURRVAL</code> and <code>sequence_name.NEXTVAL</code> in a PL/SQL expression wherever you can use a <code>NUMBER</code> expression. However:

- Using sequence_name.CURRVAL or sequence_name.NEXTVAL to provide a default value for an ADT method parameter causes a compilation error.
- PL/SQL evaluates every occurrence of sequence_name.CURRVAL and sequence_name.NEXTVAL (unlike SQL, which evaluates a sequence expression for every row in which it appears).

See Also:

- Oracle Database SQL Language Reference for general information about sequences
- Oracle Database SQL Language Reference for CURRVAL and NEXTVAL complete syntax

Example 7-2 CURRVAL and NEXTVAL Pseudocolumns

This example generates a sequence number for the sequence ${\tt HR.EMPLOYEES_SEQ}$ and refers to that number in multiple statements.

```
DROP TABLE employees temp;
CREATE TABLE employees temp AS
 SELECT employee id, first name, last name
 FROM employees;
DROP TABLE employees temp2;
CREATE TABLE employees temp2 AS
  SELECT employee_id, first_name, last_name
 FROM employees;
DECLARE
  seq value NUMBER;
BEGIN
  -- Generate initial sequence number
  seq value := employees seq.NEXTVAL;
  -- Print initial sequence number:
  DBMS OUTPUT.PUT LINE (
    'Initial sequence value: ' || TO CHAR(seq value)
  -- Use NEXTVAL to create unique number when inserting data:
     INSERT INTO employees temp (employee id, first name, last name)
     VALUES (employees seq.NEXTVAL, 'Lynette', 'Smith');
```

```
-- Use CURRVAL to store same value somewhere else:
     INSERT INTO employees temp2 VALUES (employees seq.CURRVAL,
                                         'Morgan', 'Smith');
  /* Because NEXTVAL values might be referenced
    by different users and applications,
     and some NEXTVAL values might not be stored in database,
     there might be gaps in sequence. */
  -- Use CURRVAL to specify record to delete:
     seq value := employees_seq.CURRVAL;
     DELETE FROM employees temp2
    WHERE employee_id = seq_value;
  -- Update employee id with NEXTVAL for specified record:
    UPDATE employees temp
    SET employee id = employees seq.NEXTVAL
    WHERE first name = 'Lynette'
    AND last name = 'Smith';
  -- Display final value of CURRVAL:
     seq value := employees seq.CURRVAL;
     DBMS OUTPUT.PUT LINE (
       'Ending sequence value: ' || TO CHAR(seq value)
    );
END;
```

Cursors Overview

A **cursor** is a pointer to a private SQL area that stores information about processing a specific SELECT or DML statement.

Note:

The cursors that this topic explains are session cursors. A **session cursor** lives in session memory until the session ends, when it ceases to exist.

A cursor that is constructed and managed by PL/SQL is an **implicit cursor**. A cursor that you construct and manage is an **explicit cursor**.

You can get information about any session cursor from its attributes (which you can reference in procedural statements, but not in SQL statements).

To list the session cursors that each user session currently has opened and parsed, query the dynamic performance view VSOPEN CURSOR.

The number of cursors that a session can have open simultaneously is determined by:

The amount of memory available to the session

The value of the initialization parameter OPEN CURSORS

Note:

Generally, PL/SQL parses an explicit cursor only the first time the session opens it and parses a SQL statement (creating an implicit cursor) only the first time the statement runs.

All parsed SQL statements are cached. A SQL statement is reparsed only if it is aged out of the cache by a new SQL statement. Although you must close an explicit cursor before you can reopen it, PL/SQL need not reparse the associated query. If you close and immediately reopen an explicit cursor, PL/SQL does not reparse the associated query.

Topics

- Implicit Cursors
- Explicit Cursors

See Also:

- Oracle Database Reference for information about the dynamic performance view V\$OPEN CURSOR
- Oracle Database Reference for information about the initialization parameter OPEN_CURSORS

Implicit Cursors

An **implicit cursor** is a session cursor that is constructed and managed by PL/SQL. PL/SQL opens an implicit cursor every time you run a SELECT or DML statement. You cannot control an implicit cursor, but you can get information from its attributes.

The syntax of an implicit cursor attribute value is SQLattribute (therefore, an implicit cursor is also called a **SQL** cursor). SQLattribute always refers to the most recently run SELECT or DML statement. If no such statement has run, the value of SQLattribute is NULL.

An implicit cursor closes after its associated statement runs; however, its attribute values remain available until another SELECT or DML statement runs.

The most recently run SELECT or DML statement might be in a different scope. To save an attribute value for later use, assign it to a local variable immediately. Otherwise, other operations, such as subprogram invocations, might change the value of the attribute before you can test it.

The implicit cursor attributes are:

- SQL%ISOPEN Attribute: Is the Cursor Open?
- SQL%FOUND Attribute: Were Any Rows Affected?
- SQL%NOTFOUND Attribute: Were No Rows Affected?



- SQL%ROWCOUNT Attribute: How Many Rows Were Affected?
- SQL%BULK ROWCOUNT (see "Getting Number of Rows Affected by FORALL Statement"
- SQL%BULK_EXCEPTIONS (see "Handling FORALL Exceptions After FORALL Statement Completes"



"Implicit Cursor Attribute" for complete syntax and semantics

SQL%ISOPEN Attribute: Is the Cursor Open?

SQL%ISOPEN always returns FALSE, because an implicit cursor always closes after its associated statement runs.

SQL%FOUND Attribute: Were Any Rows Affected?

SQL%FOUND returns:

- NULL if no SELECT or DML statement has run
- TRUE if a SELECT statement returned one or more rows or a DML statement affected one or more rows
- FALSE otherwise

Example 7-3 uses SQL%FOUND to determine if a DELETE statement affected any rows.

Example 7-3 SQL%FOUND Implicit Cursor Attribute

```
DROP TABLE dept temp;
CREATE TABLE dept temp AS
  SELECT * FROM departments;
CREATE OR REPLACE PROCEDURE p (
  dept no NUMBER
) AUTHID CURRENT USER AS
BEGIN
 DELETE FROM dept temp
 WHERE department id = dept no;
 IF SQL%FOUND THEN
    DBMS OUTPUT.PUT LINE (
      'Delete succeeded for department number ' || dept no
 ELSE
   DBMS OUTPUT.PUT LINE ('No department number ' || dept no);
 END IF;
END;
BEGIN
 p(270);
 p(400);
END;
```

Result:

Delete succeeded for department number 270 No department number 400

SQL%NOTFOUND Attribute: Were No Rows Affected?

SQL%NOTFOUND (the logical opposite of SQL%FOUND) returns:

- NULL if no SELECT or DML statement has run
- FALSE if a SELECT statement returned one or more rows or a DML statement affected one or more rows
- TRUE otherwise

The SQL%NOTFOUND attribute is not useful with the PL/SQL SELECT INTO statement, because:

- If the SELECT INTO statement returns no rows, PL/SQL raises the predefined exception NO DATA FOUND immediately, before you can check SQL%NOTFOUND.
- A SELECT INTO statement that invokes a SQL aggregate function always returns a value (possibly NULL). After such a statement, the SQL%NOTFOUND attribute is always FALSE, so checking it is unnecessary.

SQL%ROWCOUNT Attribute: How Many Rows Were Affected?

SOL%ROWCOUNT returns:

- NULL if no SELECT or DML statement has run
- Otherwise, the number of rows returned by a SELECT statement or affected by a DML statement (an INTEGER)



If a server is Oracle Database 12c or later and its client is Oracle Database 11g release 2 or earlier (or the reverse), then the maximum number that SQL%ROWCOUNT returns is 4,294,967,295.

Example 7-4 uses SQL%ROWCOUNT to determine the number of rows that were deleted.

If a SELECT INTO statement without a BULK COLLECT clause returns multiple rows, PL/SQL raises the predefined exception TOO_MANY_ROWS and SQL%ROWCOUNT returns 1, not the actual number of rows that satisfy the guery.

The value of SOL%ROWCOUNT attribute is unrelated to the state of a transaction. Therefore:

- When a transaction rolls back to a savepoint, the value of SQL%ROWCOUNT is not restored to the value it had before the savepoint.
- When an autonomous transaction ends, SQL%ROWCOUNT is not restored to the original value in the parent transaction.

Example 7-4 SQL%ROWCOUNT Implicit Cursor Attribute

DROP TABLE employees_temp;
CREATE TABLE employees_temp AS
 SELECT * FROM employees;



```
DECLARE
  mgr_no NUMBER(6) := 122;
BEGIN
  DELETE FROM employees_temp WHERE manager_id = mgr_no;
  DBMS_OUTPUT.PUT_LINE
    ('Number of employees deleted: ' || TO_CHAR(SQL%ROWCOUNT));
END;
//
Result:
Number of employees deleted: 8
```

Explicit Cursors

An **explicit cursor** is a session cursor that you construct and manage. You must declare and define an explicit cursor, giving it a name and associating it with a query (typically, the query returns multiple rows). Then you can process the query result set in either of these ways:

- Open the explicit cursor (with the OPEN statement), fetch rows from the result set (with the FETCH statement), and close the explicit cursor (with the CLOSE statement).
- Use the explicit cursor in a cursor FOR LOOP statement (see "Processing Query Result Sets
 With Cursor FOR LOOP Statements".

You cannot assign a value to an explicit cursor, use it in an expression, or use it as a formal subprogram parameter or host variable. You *can* do those things with a cursor variable (see "Cursor Variables").

Unlike an implicit cursor, you can reference an explicit cursor or cursor variable by its name. Therefore, an explicit cursor or cursor variable is called a **named cursor**.

Topics

- Declaring and Defining Explicit Cursors
- Opening and Closing Explicit Cursors
- Fetching Data with Explicit Cursors
- Variables in Explicit Cursor Queries
- When Explicit Cursor Queries Need Column Aliases
- Explicit Cursors that Accept Parameters
- Explicit Cursor Attributes

Declaring and Defining Explicit Cursors

You can either declare an explicit cursor first and then define it later in the same block, subprogram, or package, or declare and define it at the same time.

An **explicit cursor declaration**, which only declares a cursor, has this syntax:

```
CURSOR cursor_name [ parameter_list ] RETURN return_type;
```

An explicit cursor definition has this syntax:

```
CURSOR cursor_name [ parameter_list ] [ RETURN return_type ]
IS select statement;
```



If you declared the cursor earlier, then the explicit cursor definition defines it; otherwise, it both declares and defines it.

Example 7-5 declares and defines three explicit cursors.

See Also:

- "Explicit Cursor Declaration and Definition" for the complete syntax and semantics of explicit cursor declaration and definition
- "Explicit Cursors that Accept Parameters"

Example 7-5 Explicit Cursor Declaration and Definition

```
CURSOR c1 RETURN departments%ROWTYPE;
                                          -- Declare c1
 CURSOR c2 IS
                                          -- Declare and define c2
   SELECT employee id, job id, salary FROM employees
   WHERE salary > 2000;
 CURSOR c1 RETURN departments%ROWTYPE IS -- Define c1,
   SELECT * FROM departments
                                          -- repeating return type
   WHERE department id = 110;
 CURSOR c3 RETURN locations%ROWTYPE;
                                          -- Declare c3
 CURSOR c3 IS
                                          -- Define c3,
   SELECT * FROM locations
                                          -- omitting return type
   WHERE country_id = 'JP';
BEGIN
 NULL;
END;
```

Opening and Closing Explicit Cursors

After declaring and defining an explicit cursor, you can open it with the OPEN statement, which does the following:

- 1. Allocates database resources to process the query
- 2. Processes the query; that is:
 - a. Identifies the result set

If the query references variables or cursor parameters, their values affect the result set. For details, see "Variables in Explicit Cursor Queries" and "Explicit Cursors that Accept Parameters".

- b. If the query has a FOR UPDATE clause, locks the rows of the result set For details, see "SELECT FOR UPDATE and FOR UPDATE Cursors".
- 3. Positions the cursor before the first row of the result set

You close an open explicit cursor with the CLOSE statement, thereby allowing its resources to be reused. After closing a cursor, you cannot fetch records from its result set or reference its attributes. If you try, PL/SQL raises the predefined exception INVALID CURSOR.

You can reopen a closed cursor. You must close an explicit cursor before you try to reopen it. Otherwise, PL/SQL raises the predefined exception CURSOR ALREADY OPEN.

See Also:

- "OPEN Statement" for its syntax and semantics
- "CLOSE Statement" for its syntax and semantics

Fetching Data with Explicit Cursors

After opening an explicit cursor, you can fetch the rows of the query result set with the FETCH statement. The basic syntax of a FETCH statement that returns one row is:

```
FETCH cursor name INTO into clause
```

The <code>into_clause</code> is either a list of variables or a single record variable. For each column that the query returns, the variable list or record must have a corresponding type-compatible variable or field. The <code>%TYPE</code> and <code>%ROWTYPE</code> attributes are useful for declaring variables and records for use in <code>FETCH</code> statements.

The FETCH statement retrieves the current row of the result set, stores the column values of that row into the variables or record, and advances the cursor to the next row.

Typically, you use the FETCH statement inside a LOOP statement, which you exit when the FETCH statement runs out of rows. To detect this exit condition, use the cursor attribute %NOTFOUND (described in "%NOTFOUND Attribute: Has No Row Been Fetched?"). PL/SQL does not raise an exception when a FETCH statement returns no rows.

Example 7-6 fetches the result sets of two explicit cursors one row at a time, using FETCH and %NOTFOUND inside LOOP statements. The first FETCH statement retrieves column values into variables. The second FETCH statement retrieves column values into a record. The variables and record are declared with %TYPE and %ROWTYPE, respectively.

Example 7-7 fetches the first five rows of a result set into five records, using five FETCH statements, each of which fetches into a different record variable. The record variables are declared with %ROWTYPE.

See Also:

- "FETCH Statement" for its complete syntax and semantics
- "FETCH Statement with BULK COLLECT Clause" for information about FETCH statements that return more than one row at a time

Example 7-6 FETCH Statements Inside LOOP Statements

```
DECLARE

CURSOR c1 IS

SELECT last_name, job_id FROM employees

WHERE REGEXP_LIKE (job_id, 'S[HT]_CLERK')

ORDER BY last name;
```



```
v lastname employees.last name%TYPE; -- variable for last name
 v_jobid employees.job_id%TYPE; -- variable for job_id
  CURSOR c2 IS
   SELECT * FROM employees
   WHERE REGEXP LIKE (job id, '[ACADFIMKSA] M[ANGR]')
   ORDER BY job id;
 v employees employees%ROWTYPE; -- record variable for row of table
BEGIN
 OPEN c1;
 LOOP -- Fetches 2 columns into variables
   FETCH c1 INTO v lastname, v jobid;
   EXIT WHEN c1%NOTFOUND;
   DBMS OUTPUT.PUT LINE ( RPAD (v lastname, 25, ' ') || v jobid );
 END LOOP;
 CLOSE c1;
 DBMS OUTPUT.PUT LINE( '-----');
 OPEN c2;
 LOOP -- Fetches entire row into the v_employees record
   FETCH c2 INTO v employees;
   EXIT WHEN c2%NOTFOUND;
   DBMS OUTPUT.PUT_LINE( RPAD(v_employees.last_name, 25, ' ') ||
                             v employees.job id );
 END LOOP;
 CLOSE c2;
END;
Result:
Atkinson
                        ST CLERK
Bell
                       SH CLERK
                        ST CLERK
Bissot
. . .
Walsh
                       SH CLERK
Higgins
                       AC MGR
Gruenberg
                       FI MGR
Martinez
                       MK MAN
Errazuriz
                        SA MAN
```

Example 7-7 Fetching Same Explicit Cursor into Different Variables

```
DECLARE
   CURSOR c IS
     SELECT e.job_id, j.job_title
     FROM employees e, jobs j
   WHERE e.job_id = j.job_id AND e.manager_id = 100
     ORDER BY last_name;
-- Record variables for rows of cursor result set:
   job1 c%ROWTYPE;
   job2 c%ROWTYPE;
   job3 c%ROWTYPE;
   job4 c%ROWTYPE;
   job5 c%ROWTYPE;
   job5 c%ROWTYPE;
```

```
OPEN c:
 FETCH c INTO job1; -- fetches first row
 FETCH c INTO job2; -- fetches second row
 FETCH c INTO job3; -- fetches third row
 FETCH c INTO job4; -- fetches fourth row
 FETCH c INTO job5; -- fetches fifth row
 CLOSE c;
 DBMS_OUTPUT.PUT_LINE(job1.job_title || ' (' || job1.job_id || ')');
 DBMS_OUTPUT.PUT_LINE(job2.job_title || ' (' || job2.job_id || ')');
 DBMS_OUTPUT.PUT_LINE(job3.job_title || ' (' || job3.job_id || ')');
 DBMS_OUTPUT.PUT_LINE(job4.job_title || ' (' || job4.job_id || ')');
 DBMS_OUTPUT.PUT_LINE(job5.job_title || ' (' || job5.job_id || ')');
Result:
Sales Manager (SA MAN)
Sales Manager (SA MAN)
Stock Manager (ST MAN)
Administration Vice President (AD VP)
Stock Manager (ST MAN)
PL/SQL procedure successfully completed.
```

Variables in Explicit Cursor Queries

An explicit cursor query can reference any variable in its scope. When you open an explicit cursor, PL/SQL evaluates any variables in the query and uses those values when identifying the result set. Changing the values of the variables later does not change the result set.

In Example 7-8, the explicit cursor query references the variable factor. When the cursor opens, factor has the value 2. Therefore, sal_multiple is always 2 times sal, despite that factor is incremented after every fetch.

To change the result set, you must close the cursor, change the value of the variable, and then open the cursor again, as in Example 7-9.

Example 7-8 Variable in Explicit Cursor Query—No Result Set Change

```
DECLARE
 sal
                employees.salary%TYPE;
 sal multiple employees.salary%TYPE;
               INTEGER := 2;
 factor
 CURSOR c1 IS
    SELECT salary, salary*factor FROM employees
   WHERE job id LIKE 'AD %';
BEGIN
 OPEN c1; -- PL/SQL evaluates factor
 T<sub>1</sub>OOP
   FETCH c1 INTO sal, sal_multiple;
   EXIT WHEN c1%NOTFOUND;
    DBMS OUTPUT.PUT LINE('factor = ' || factor);
    DBMS OUTPUT.PUT LINE('sal = ' || sal);
    DBMS OUTPUT.PUT LINE('sal multiple = ' || sal multiple);
    factor := factor + 1; -- Does not affect sal_multiple
```

```
END LOOP;
 CLOSE c1;
END;
Result:
factor = 2
          = 4400
sal
sal multiple = 8800
factor = 3
sal
           = 24000
sal multiple = 48000
factor = 4
sal
           = 17000
sal multiple = 34000
factor = 5
          = 17000
sal
sal multiple = 34000
```

Example 7-9 Variable in Explicit Cursor Query—Result Set Change

```
DECLARE
 sal
              employees.salary%TYPE;
 sal_multiple employees.salary%TYPE;
 factor
             INTEGER := 2;
 CURSOR c1 IS
   SELECT salary, salary*factor FROM employees
   WHERE job id LIKE 'AD %';
 DBMS OUTPUT.PUT LINE('factor = ' | factor);
 OPEN c1; -- PL/SQL evaluates factor
   FETCH c1 INTO sal, sal multiple;
   EXIT WHEN c1%NOTFOUND;
   DBMS OUTPUT.PUT LINE('sal = ' | sal);
   DBMS OUTPUT.PUT LINE('sal multiple = ' || sal multiple);
 END LOOP;
 CLOSE c1;
 factor := factor + 1;
 DBMS OUTPUT.PUT LINE('factor = ' || factor);
 OPEN c1; -- PL/SQL evaluates factor
 LOOP
   FETCH c1 INTO sal, sal_multiple;
   EXIT WHEN c1%NOTFOUND;
                               = ' || sal);
   DBMS OUTPUT.PUT LINE('sal
   DBMS OUTPUT.PUT LINE('sal multiple = ' || sal multiple);
 END LOOP;
 CLOSE c1;
END;
Result:
factor = 2
sal = 4400
sal multiple = 8800
          = 24000
```

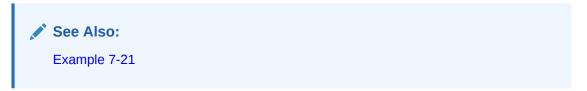
```
sal_multiple = 48000
sal = 17000
sal_multiple = 34000
sal = 17000
sal_multiple = 34000
factor = 3
sal = 4400
sal_multiple = 13200
sal_multiple = 72000
sal_multiple = 72000
sal_multiple = 51000
sal_multiple = 51000
sal_multiple = 51000
```

When Explicit Cursor Queries Need Column Aliases

When an explicit cursor query includes a virtual column (an expression), that column must have an alias if either of the following is true:

- You use the cursor to fetch into a record that was declared with %ROWTYPE.
- You want to reference the virtual column in your program.

In Example 7-10, the virtual column in the explicit cursor needs an alias for both of the preceding reasons.



Example 7-10 Explicit Cursor with Virtual Column that Needs Alias

```
DECLARE
 CURSOR c1 IS
   SELECT employee id,
          (salary * .05) raise
   FROM employees
   WHERE job id LIKE '% MAN'
   ORDER BY employee id;
 emp rec c1%ROWTYPE;
BEGIN
 OPEN c1;
 LOOP
   FETCH c1 INTO emp rec;
   EXIT WHEN c1%NOTFOUND;
   DBMS OUTPUT.PUT LINE (
      'Raise for employee #' || emp_rec.employee_id ||
      ' is $' || emp_rec.raise
   );
 END LOOP;
 CLOSE c1;
END:
Result:
Raise for employee #114 is $550
Raise for employee #120 is $400
```



```
Raise for employee #121 is $410
Raise for employee #122 is $395
Raise for employee #123 is $325
Raise for employee #124 is $368.445
Raise for employee #145 is $700
Raise for employee #146 is $675
Raise for employee #147 is $600
Raise for employee #148 is $550
Raise for employee #149 is $525
Raise for employee #201 is $650
```

Explicit Cursors that Accept Parameters

You can create an explicit cursor that has formal parameters, and then pass different actual parameters to the cursor each time you open it. In the cursor query, you can use a formal cursor parameter anywhere that you can use a constant. Outside the cursor query, you cannot reference formal cursor parameters.



Tip:

To avoid confusion, use different names for formal and actual cursor parameters.

Example 7-11 creates an explicit cursor whose two formal parameters represent a job and its maximum salary. When opened with a specified job and maximum salary, the cursor query selects the employees with that job who are overpaid (for each such employee, the query selects the first and last name and amount overpaid). Next, the example creates a procedure that prints the cursor query result set (for information about procedures, see PL/SQL Subprograms). Finally, the example opens the cursor with one set of actual parameters, prints the result set, closes the cursor, opens the cursor with different actual parameters, prints the result set, and closes the cursor.

Topics

- Formal Cursor Parameters with Default Values
- Adding Formal Cursor Parameters with Default Values

See Also:

- "Explicit Cursor Declaration and Definition" for more information about formal cursor parameters
- "OPEN Statement" for more information about actual cursor parameters

Example 7-11 Explicit Cursor that Accepts Parameters

```
DECLARE

CURSOR c (job VARCHAR2, max_sal NUMBER) IS

SELECT last_name, first_name, (salary - max_sal) overpayment

FROM employees

WHERE job_id = job

AND salary > max_sal

ORDER BY salary;
```



```
PROCEDURE print overpaid IS
   last_name_ employees.last_name%TYPE;
   first name employees.first_name%TYPE;
   overpayment
                 employees.salary%TYPE;
 BEGIN
   LOOP
     FETCH c INTO last name , first name , overpayment ;
     EXIT WHEN c%NOTFOUND;
     DBMS_OUTPUT.PUT_LINE(last_name_ || ', ' || first_name_ ||
      ' (by ' || overpayment_ || ')');
   END LOOP;
 END print overpaid;
BEGIN
 DBMS OUTPUT.PUT LINE('----');
 DBMS OUTPUT.PUT LINE('Overpaid Stock Clerks:');
 DBMS OUTPUT.PUT LINE('----');
 OPEN c('ST_CLERK', 5000);
 print overpaid;
 CLOSE c;
 DBMS OUTPUT.PUT LINE('-----');
 DBMS OUTPUT.PUT LINE ('Overpaid Sales Representatives:');
 DBMS OUTPUT.PUT LINE('----');
 OPEN c('SA_REP', 10000);
 print overpaid;
 CLOSE c;
END;
Result:
Overpaid Stock Clerks:
_____
Overpaid Sales Representatives:
_____
Vishney, Clara (by 500)
Abel, Ellen (by 1000)
Ozer, Lisa (by 1500)
PL/SQL procedure successfully completed.
```

Formal Cursor Parameters with Default Values

When you create an explicit cursor with formal parameters, you can specify default values for them. When a formal parameter has a default value, its corresponding actual parameter is optional. If you open the cursor without specifying the actual parameter, then the formal parameter has its default value.

Example 7-12 creates an explicit cursor whose formal parameter represents a location ID. The default value of the parameter is the location ID of company headquarters.

Example 7-12 Cursor Parameters with Default Values

```
DECLARE

CURSOR c (location NUMBER DEFAULT 1700) IS

SELECT d.department_name,

e.last_name manager,

l.city

FROM departments d, employees e, locations 1
```

```
WHERE 1.location id = location
     AND l.location id = d.location id
     AND d.department_id = e.department_id
   ORDER BY d.department id;
 PROCEDURE print depts IS
   dept name departments.department name%TYPE;
   mgr_name employees.last_name%TYPE;
   city name locations.city%TYPE;
 BEGIN
   LOOP
     FETCH c INTO dept_name, mgr_name, city_name;
     EXIT WHEN c%NOTFOUND;
     DBMS OUTPUT.PUT LINE(dept name || ' (Manager: ' || mgr name || ')');
   END LOOP;
 END print depts;
BEGIN
 DBMS OUTPUT.PUT LINE ('DEPARTMENTS AT HEADQUARTERS:');
 DBMS OUTPUT.PUT LINE('----');
 OPEN c;
 print depts;
 DBMS OUTPUT.PUT LINE('----');
 CLOSE c;
 DBMS OUTPUT.PUT LINE('DEPARTMENTS IN CANADA:');
 DBMS_OUTPUT_LINE('----');
 OPEN c(1800); -- Toronto
 print depts;
 CLOSE c;
 OPEN c(1900); -- Whitehorse
 print depts;
 CLOSE c;
END;
```

Result is similar to:

```
DEPARTMENTS AT HEADOUARTERS:
Administration (Manager: Whalen)
Purchasing (Manager: Himuro)
Purchasing (Manager: Tobias)
Purchasing (Manager: Baida)
Purchasing (Manager: Li)
Purchasing (Manager: Colmenares)
Purchasing (Manager: Khoo)
Executive (Manager: Yang)
Executive (Manager: Garcia)
Executive (Manager: King)
Finance (Manager: Urman)
Finance (Manager: Sciarra)
Finance (Manager: Chen)
Finance (Manager: Faviet)
Finance (Manager: Gruenberg)
Finance (Manager: Popp)
Accounting (Manager: Higgins)
Accounting (Manager: Gietz)
_____
DEPARTMENTS IN CANADA:
```

```
Marketing (Manager: Davis)
Marketing (Manager: Martinez)

PL/SQL procedure successfully completed.
```

Adding Formal Cursor Parameters with Default Values

If you add formal parameters to a cursor, and you specify default values for the added parameters, then you need not change existing references to the cursor. Compare Example 7-13 to Example 7-11.

Example 7-13 Adding Formal Parameter to Existing Cursor

```
DECLARE
 CURSOR c (job VARCHAR2, max sal NUMBER,
          hired DATE DEFAULT TO DATE('31-DEC-1999', 'DD-MON-YYYY')) IS
   SELECT last name, first name, (salary - max sal) overpayment
   FROM employees
   WHERE job id = job
   AND salary > max sal
   AND hire_date > hired
   ORDER BY salary;
  PROCEDURE print overpaid IS
   last_name_ employees.last_name%TYPE;
   first name employees.first name%TYPE;
   overpayment_
                 employees.salary%TYPE;
 BEGIN
   LOOP
     FETCH c INTO last_name_, first_name_, overpayment_;
     EXIT WHEN c%NOTFOUND;
     DBMS_OUTPUT.PUT_LINE(last_name_ || ', ' || first_name_ ||
       ' (by ' || overpayment || ')');
   END LOOP;
 END print overpaid;
BEGIN
  DBMS OUTPUT.PUT LINE('-----');
  DBMS OUTPUT.PUT LINE ('Overpaid Sales Representatives:');
  DBMS OUTPUT.PUT LINE('----');
  OPEN c('SA REP', 10000); -- existing reference
 print overpaid;
  CLOSE c;
  DBMS OUTPUT.PUT LINE('-----');
  DBMS OUTPUT.PUT LINE('Overpaid Sales Representatives Hired After 2014:');
  DBMS OUTPUT.PUT LINE('----');
 OPEN c('SA REP', 10000, To_DATE('31-DEC-2014', 'DD-MON-YYYY'));
                       -- new reference
 print overpaid;
  CLOSE c;
END;
Result:
Overpaid Sales Representatives:
Vishney, Clara (by 500)
Abel, Ellen (by 1000)
Ozer, Lisa (by 1500)
```

```
Overpaid Sales Representatives Hired After 2014:

Vishney, Clara (by 500)
Ozer, Lisa (by 1500)
```

PL/SQL procedure successfully completed.

Explicit Cursor Attributes

The syntax for the value of an explicit cursor attribute is <code>cursor_name</code> immediately followed by <code>attribute</code> (for example, <code>c1%ISOPEN</code>).



Explicit cursors and cursor variables (named cursors) have the same attributes. This topic applies to all named cursors except where noted.

The explicit cursor attributes are:

- %ISOPEN Attribute: Is the Cursor Open?
- %FOUND Attribute: Has a Row Been Fetched?
- %NOTFOUND Attribute: Has No Row Been Fetched?
- %ROWCOUNT Attribute: How Many Rows Were Fetched?

If an explicit cursor is not open, referencing any attribute except %ISOPEN raises the predefined exception INVALID_CURSOR.



"Named Cursor Attribute" for complete syntax and semantics of named cursor (explicit cursor and cursor variable) attributes

%ISOPEN Attribute: Is the Cursor Open?

%ISOPEN returns TRUE if its explicit cursor is open; FALSE otherwise.

%ISOPEN is useful for:

Checking that an explicit cursor is not already open before you try to open it.

If you try to open an explicit cursor that is already open, PL/SQL raises the predefined exception ${\tt CURSOR_ALREADY_OPEN}$. You must close an explicit cursor before you can reopen it.



The preceding paragraph does not apply to cursor variables.

Checking that an explicit cursor is open before you try to close it.

Example 7-14 opens the explicit cursor c1 only if it is not open and closes it only if it is open.

Example 7-14 %ISOPEN Explicit Cursor Attribute

```
DECLARE

CURSOR c1 IS

SELECT last_name, salary FROM employees

WHERE ROWNUM < 11;

the_name employees.last_name%TYPE;
the_salary employees.salary%TYPE;
BEGIN

IF NOT c1%ISOPEN THEN

OPEN c1;
END IF;

FETCH c1 INTO the_name, the_salary;

IF c1%ISOPEN THEN

CLOSE c1;
END IF;

END;
//
```

%FOUND Attribute: Has a Row Been Fetched?

%FOUND returns:

- NULL after the explicit cursor is opened but before the first fetch
- TRUE if the most recent fetch from the explicit cursor returned a row
- FALSE otherwise

%FOUND is useful for determining whether there is a fetched row to process.

Example 7-15 loops through a result set, printing each fetched row and exiting when there are no more rows to fetch.

Example 7-15 %FOUND Explicit Cursor Attribute

```
DECLARE
 CURSOR c1 IS
    SELECT last name, salary FROM employees
   WHERE ROWNUM < 11
   ORDER BY last name;
 my_ename employees.last_name%TYPE;
 my salary employees.salary%TYPE;
BEGIN
 OPEN c1;
 LOOP
   FETCH c1 INTO my ename, my salary;
   IF c1%FOUND THEN -- fetch succeeded
     DBMS_OUTPUT.PUT_LINE('Name = ' || my_ename || ', salary = ' || my_salary);
   ELSE -- fetch failed
     EXIT;
   END IF;
 END LOOP;
END;
```

Result:

```
Name = Faviet, salary = 9000
Name = Garcia, salary = 17000
Name = Gruenberg, salary = 12008
Name = Jackson, salary = 4800
Name = James, salary = 9000
Name = King, salary = 24000
Name = Miller, salary = 6000
Name = Nguyen, salary = 4200
Name = Williams, salary = 4800
Name = Yang, salary = 17000
```

%NOTFOUND Attribute: Has No Row Been Fetched?

%NOTFOUND (the logical opposite of %FOUND) returns:

- NULL after the explicit cursor is opened but before the first fetch
- FALSE if the most recent fetch from the explicit cursor returned a row
- TRUE otherwise

*NOTFOUND is useful for exiting a loop when FETCH fails to return a row, as in Example 7-16.

Example 7-16 %NOTFOUND Explicit Cursor Attribute

```
DECLARE
  CURSOR cl IS
    SELECT last name, salary FROM employees
    WHERE ROWNUM < 11
    ORDER BY last name;
   my ename employees.last name%TYPE;
   my salary employees.salary%TYPE;
BEGIN
  OPEN c1;
  LOOP
    FETCH c1 INTO my ename, my salary;
    IF c1%NOTFOUND THEN -- fetch failed
      EXIT;
    ELSE -- fetch succeeded
      DBMS OUTPUT.PUT LINE
        ('Name = ' || my ename || ', salary = ' || my salary);
    END IF;
  END LOOP;
END;
Result:
Name = Faviet, salary = 9000
Name = Garcia, salary = 17000
Name = Gruenberg, salary = 12008
Name = Jackson, salary = 4800
Name = James, salary = 9000
Name = King, salary = 24000
Name = Miller, salary = 6000
```



```
Name = Nguyen, salary = 4200
Name = Williams, salary = 4800
Name = Yang, salary = 17000
```

%ROWCOUNT Attribute: How Many Rows Were Fetched?

%ROWCOUNT returns:

- Zero after the explicit cursor is opened but before the first fetch
- Otherwise, the number of rows fetched (an INTEGER)



If a server is Oracle Database 12c or later and its client is Oracle Database 11g2 or earlier (or the reverse), then the maximum number that SQL%ROWCOUNT returns is 4,294,967,295.

Example 7-17 numbers and prints the rows that it fetches and prints a message after fetching the fifth row.

Example 7-17 %ROWCOUNT Explicit Cursor Attribute

```
DECLARE
  CURSOR c1 IS
    SELECT last name FROM employees
    WHERE ROWNUM < 11
    ORDER BY last name;
  name employees.last_name%TYPE;
BEGIN
  OPEN c1;
  LOOP
   FETCH c1 INTO name;
   EXIT WHEN c1%NOTFOUND OR c1%NOTFOUND IS NULL;
    DBMS OUTPUT.PUT LINE(c1%ROWCOUNT || '. ' || name);
    IF c1%ROWCOUNT = 5 THEN
       DBMS OUTPUT.PUT LINE('--- Fetched 5th row ---');
   END IF;
  END LOOP;
  CLOSE c1;
END;
```

Result:

```
    Abel
    Ande
    Atkinson
    Baida
    Banda
    Fetched 5th row ---
    Bates
    Bell
```



- 8. Bernstein
- 9. Bissot
- 10. Bloom

Processing Query Result Sets

In PL/SQL, as in traditional database programming, you use cursors to process query result sets. However, in PL/SQL, you can use either implicit or explicit cursors.

The former need less code, but the latter are more flexible. For example, explicit cursors can accept parameters.

The following PL/SQL statements use implicit cursors that PL/SQL defines and manages for you:

- SELECT INTO
- Implicit cursor FOR LOOP

The following PL/SQL statements use explicit cursors:

Explicit cursor FOR LOOP

You define the explicit cursor, but PL/SQL manages it while the statement runs.

OPEN, FETCH, and CLOSE

You define and manage the explicit cursor.



If a query returns no rows, PL/SQL raises the exception NO DATA FOUND.

Topics

- Processing Query Result Sets With SELECT INTO Statements
- Processing Query Result Sets With Cursor FOR LOOP Statements
- Processing Query Result Sets With Explicit Cursors, OPEN, FETCH, and CLOSE
- Processing Query Result Sets with Subqueries

See Also:

- "Explicit Cursors that Accept Parameters"
- Oracle Database Development Guide for information about returning result sets to clients
- "Exception Handler" for information about handling exceptions



Using an implicit cursor, the SELECT INTO statement retrieves values from one or more database tables (as the SQL SELECT statement does) and stores them in variables (which the SQL SELECT statement does not do).

Topics

- Handling Single-Row Result Sets
- Handling Large Multiple-Row Result Sets



"SELECT INTO Statement" for its complete syntax and semantics

Handling Single-Row Result Sets

If you expect the query to return only one row, then use the SELECT INTO statement to store values from that row in either one or more scalar variables, or one record variable.

If the query might return multiple rows, but you care about only the nth row, then restrict the result set to that row with the clause WHERE ROWNUM=n.

See Also:

- "Assigning Values to Variables with the SELECT INTO Statement"
- "Using SELECT INTO to Assign a Row to a Record Variable"
- Oracle Database SQL Language Reference for more information about the ROWNUM pseudocolumn

Handling Large Multiple-Row Result Sets

If you must assign a large quantity of table data to variables, Oracle recommends using the SELECT INTO statement with the BULK COLLECT clause.

This statement retrieves an entire result set into one or more collection variables.

For more information, see "SELECT INTO Statement with BULK COLLECT Clause".

Processing Query Result Sets With Cursor FOR LOOP Statements

The cursor FOR LOOP statement lets you run a SELECT statement and then immediately loop through the rows of the result set.

This statement can use either an implicit or explicit cursor (but not a cursor variable).

If you use the Select statement only in the cursor for loop statement, then specify the Select statement inside the cursor for loop statement. This form of the cursor for loop statement



uses an implicit cursor, and is called an **implicit cursor for loop statement**. Because the implicit cursor is internal to the statement, you cannot reference it with the name SQL.

If you use the SELECT statement multiple times in the same PL/SQL unit, then define an explicit cursor for it and specify that cursor in the cursor FOR LOOP statement. This form of the cursor FOR LOOP statement is called an explicit cursor FOR LOOP statement. You can use the same explicit cursor elsewhere in the same PL/SQL unit.

The cursor FOR LOOP statement implicitly declares its loop index as a %ROWTYPE record variable of the type that its cursor returns. This record is local to the loop and exists only during loop execution. Statements inside the loop can reference the record and its fields. They can reference virtual columns only by aliases.

After declaring the loop index record variable, the FOR LOOP statement opens the specified cursor. With each iteration of the loop, the FOR LOOP statement fetches a row from the result set and stores it in the record. When there are no more rows to fetch, the cursor FOR LOOP statement closes the cursor. The cursor also closes if a statement inside the loop transfers control outside the loop or if PL/SQL raises an exception.



"Cursor FOR LOOP Statement" for its complete syntax and semantics

Note:

When an exception is raised inside a cursor FOR LOOP statement, the cursor closes before the exception handler runs. Therefore, the values of explicit cursor attributes are not available in the handler.

Example 7-18 Implicit Cursor FOR LOOP Statement

In this example, an implicit cursor FOR LOOP statement prints the last name and job ID of every clerk whose manager has an ID greater than 120.

```
BEGIN
 FOR item IN (
   SELECT last name, job id
   FROM employees
   WHERE job id LIKE '%CLERK%'
   AND manager id > 120
   ORDER BY last name
 )
 LOOP
    DBMS OUTPUT.PUT LINE
      ('Name = ' || item.last name || ', Job = ' || item.job id);
 END LOOP;
END;
Result:
```

```
Name = Atkinson, Job = ST CLERK
Name = Bell, Job = SH CLERK
Name = Bissot, Job = ST CLERK
```

```
Name = Walsh, Job = SH CLERK
```

Example 7-19 Explicit Cursor FOR LOOP Statement

This example is like Example 7-18, except that it uses an explicit cursor FOR LOOP statement.

```
DECLARE
  CURSOR c1 IS
    SELECT last name, job id FROM employees
    WHERE job id LIKE '%CLERK%' AND manager id > 120
    ORDER BY last_name;
BEGIN
  FOR item IN c1
  LOOP
    DBMS OUTPUT.PUT LINE
      ('Name = ' || item.last name || ', Job = ' || item.job id);
  END LOOP;
END;
Result:
Name = Atkinson, Job = ST CLERK
Name = Bell, Job = SH CLERK
Name = Bissot, Job = ST CLERK
Name = Walsh, Job = SH CLERK
```

Example 7-20 Passing Parameters to Explicit Cursor FOR LOOP Statement

This example declares and defines an explicit cursor that accepts two parameters, and then uses it in an explicit cursor FOR LOOP statement to display the wages paid to employees who earn more than a specified wage in a specified department.

```
DECLARE
 CURSOR c1 (job VARCHAR2, max wage NUMBER) IS
    SELECT * FROM employees
   WHERE job id = job
   AND salary > max_wage;
BEGIN
 FOR person IN c1('ST CLERK', 3000)
 LOOP
    -- process data record
    DBMS OUTPUT.PUT LINE (
      'Name = ' || person.last_name || ', salary = ' ||
     person.salary || ', Job Id = ' || person.job id
   );
 END LOOP;
END;
Result:
Name = Nayer, salary = 3200, Job Id = ST CLERK
Name = Bissot, salary = 3300, Job Id = ST CLERK
Name = Mallin, salary = 3300, Job Id = ST_CLERK
Name = Ladwig, salary = 3600, Job Id = ST_CLERK
Name = Stiles, salary = 3200, Job Id = ST_CLERK
Name = Rajs, salary = 3500, Job Id = ST CLERK
Name = Davies, salary = 3100, Job Id = ST CLERK
```

Example 7-21 Cursor FOR Loop References Virtual Columns

In this example, the implicit cursor FOR LOOP references virtual columns by their aliases, full name and dream salary.

Result:

```
Stephen King dreams of making 240000
Lex Garcia dreams of making 170000
Neena Yang dreams of making 170000
Alexander James dreams of making 90000
Bruce Miller dreams of making 60000
```

Processing Query Result Sets With Explicit Cursors, OPEN, FETCH, and CLOSE

For full control over query result set processing, declare explicit cursors and manage them with the statements OPEN, FETCH, and CLOSE.

This result set processing technique is more complicated than the others, but it is also more flexible. For example, you can:

- Process multiple result sets in parallel, using multiple cursors.
- Process multiple rows in a single loop iteration, skip rows, or split the processing into multiple loops.
- Specify the query in one PL/SQL unit but retrieve the rows in another.

For instructions and examples, see "Explicit Cursors".

Processing Query Result Sets with Subqueries

If you process a query result set by looping through it and running another query for each row, then you can improve performance by removing the second query from inside the loop and making it a subquery of the first query.

While an ordinary subquery is evaluated for each table, a **correlated subquery** is evaluated for each row.

For more information about subqueries, see Oracle Database SQL Language Reference.

Example 7-22 Subquery in FROM Clause of Parent Query

This example defines explicit cursor c1 with a query whose FROM clause contains a subquery.

```
DECLARE
 CURSOR c1 IS
   SELECT t1.department_id, department_name, staff
    FROM departments t1,
         ( SELECT department_id, COUNT(*) AS staff
           FROM employees
           GROUP BY department id
    WHERE (t1.department id = t2.department id) AND staff >= 5
    ORDER BY staff;
BEGIN
  FOR dept IN c1
  LOOP
     DBMS OUTPUT.PUT LINE ('Department = '
      || dept.department name || ', staff = ' || dept.staff);
  END LOOP;
END;
Result:
Department = IT, staff = 5
Department = Finance, staff = 6
Department = Purchasing, staff = 6
Department = Sales, staff = 34
Department = Shipping, staff = 45
```

Example 7-23 Correlated Subquery

This example returns the name and salary of each employee whose salary exceeds the departmental average. For each row in the table, the correlated subquery computes the average salary for the corresponding department.

```
DECLARE
  CURSOR c1 IS
    SELECT department id, last name, salary
    FROM employees t
    WHERE salary > ( SELECT AVG(salary)
                     FROM employees
                     WHERE t.department id = department id
    ORDER BY department_id, last_name;
BEGIN
  FOR person IN c1
  LOOP
    DBMS OUTPUT.PUT LINE('Making above-average salary = ' || person.last name);
  END LOOP;
END;
Result:
Making above-average salary = Martinez
Making above-average salary = Li
Making above-average salary = Bell
Making above-average salary = Higgins
```

Cursor Variables

A **cursor variable** is like an explicit cursor, except that:

It is not limited to one query.

You can open a cursor variable for a query, process the result set, and then use the cursor variable for another query.

- · You can assign a value to it.
- You can use it in an expression.
- It can be a subprogram parameter.

You can use cursor variables to pass query result sets between subprograms.

It can be a host variable.

You can use cursor variables to pass query result sets between PL/SQL stored subprograms and their clients.

It cannot accept parameters.

You cannot pass parameters to a cursor variable, but you can pass whole queries to it. The queries can include variables.

A cursor variable has this flexibility because it is a pointer; that is, its value is the address of an item, not the item itself.

Before you can reference a cursor variable, you must make it point to a SQL work area, either by opening it or by assigning it the value of an open PL/SQL cursor variable or open host cursor variable.



Cursor variables and explicit cursors are not interchangeable—you cannot use one where the other is expected.

Topics

- Creating Cursor Variables
- Opening and Closing Cursor Variables
- Fetching Data with Cursor Variables
- Assigning Values to Cursor Variables
- Variables in Cursor Variable Queries
- Querying a Collection
- Cursor Variable Attributes
- Cursor Variables as Subprogram Parameters
- Cursor Variables as Host Variables



See Also:

- "Explicit Cursors" for more information about explicit cursors
- "Restrictions on Cursor Variables"
- Oracle Database Development Guide for advantages of cursor variables
- Oracle Database Development Guide for disadvantages of cursor variables

Creating Cursor Variables

To create a cursor variable, either declare a variable of the predefined type SYS_REFCURSOR or define a REF CURSOR type and then declare a variable of that type.



Informally, a cursor variable is sometimes called a REF CURSOR).

The basic syntax of a REF CURSOR type definition is:

```
TYPE type name IS REF CURSOR [ RETURN return type ]
```

For the complete syntax and semantics, see "Cursor Variable Declaration".

If you specify $return_type$, then the REF CURSOR type and cursor variables of that type are **strong**; if not, they are **weak**. SYS_REFCURSOR and cursor variables of that type are weak.

With a strong cursor variable, you can associate only queries that return the specified type. With a weak cursor variable, you can associate any query.

Weak cursor variables are more error-prone than strong ones, but they are also more flexible. Weak REF CURSOR types are interchangeable with each other and with the predefined type SYS_REFCURSOR. You can assign the value of a weak cursor variable to any other weak cursor variable.

You can assign the value of a strong cursor variable to another strong cursor variable only if both cursor variables have the same type (not merely the same return type).



You can partition weak cursor variable arguments to table functions only with the PARTITION BY ANY clause, not with PARTITION BY RANGE or PARTITION BY HASH.

For syntax and semantics, see "PARALLEL_ENABLE Clause".

Example 7-24 Cursor Variable Declarations

This example defines strong and weak REF CURSOR types, variables of those types, and a variable of the predefined type SYS REFCURSOR.



```
DECLARE

TYPE empcurtyp IS REF CURSOR RETURN employees%ROWTYPE; -- strong type

TYPE genericcurtyp IS REF CURSOR; -- weak type

cursor1 empcurtyp; -- strong cursor variable

cursor2 genericcurtyp; -- weak cursor variable

my_cursor SYS_REFCURSOR; -- weak cursor variable

TYPE deptcurtyp IS REF CURSOR RETURN departments%ROWTYPE; -- strong type

dept_cv deptcurtyp; -- strong cursor variable

BEGIN

NULL;

END;

/
```

Example 7-25 Cursor Variable with User-Defined Return Type

In this example, *EmpRecTyp* is a user-defined RECORD type.

```
DECLARE
   TYPE EmpRecTyp IS RECORD (
    employee_id NUMBER,
    last_name VARCHAR2(25),
    salary NUMBER(8,2));

   TYPE EmpCurTyp IS REF CURSOR RETURN EmpRecTyp;
   emp_cv EmpCurTyp;
BEGIN
   NULL;
END;
//
```

Opening and Closing Cursor Variables

After declaring a cursor variable, you can open it with the OPEN FOR statement, which does the following:

- Associates the cursor variable with a query (typically, the query returns multiple rows)
 The query can include placeholders for bind variables, whose values you specify in the USING clause of the OPEN FOR statement.
- 2. Allocates database resources to process the query
- 3. Processes the guery; that is:
 - a. Identifies the result set

If the query references variables, their values affect the result set. For details, see "Variables in Cursor Variable Queries".

- b. If the query has a FOR UPDATE clause, locks the rows of the result set For details, see "SELECT FOR UPDATE and FOR UPDATE Cursors".
- 4. Positions the cursor before the first row of the result set

You need not close a cursor variable before reopening it (that is, using it in another OPEN FOR statement). After you reopen a cursor variable, the query previously associated with it is lost.

When you no longer need a cursor variable, close it with the CLOSE statement, thereby allowing its resources to be reused. After closing a cursor variable, you cannot fetch records from its result set or reference its attributes. If you try, PL/SQL raises the predefined exception INVALID_CURSOR.

You can reopen a closed cursor variable.

See Also:

- "OPEN FOR Statement" for its syntax and semantics
- "CLOSE Statement" for its syntax and semantics

Fetching Data with Cursor Variables

After opening a cursor variable, you can fetch the rows of the query result set with the FETCH statement.

The return type of the cursor variable must be compatible with the <code>into_clause</code> of the <code>FETCH</code> statement. If the cursor variable is strong, PL/SQL catches incompatibility at compile time. If the cursor variable is weak, PL/SQL catches incompatibility at run time, raising the predefined exception <code>ROWTYPE_MISMATCH</code> before the first fetch.

See Also:

- "Fetching Data with Explicit Cursors"
- "FETCH Statement" for its complete syntax and semantics
- "FETCH Statement with BULK COLLECT Clause" for information about FETCH statements that return more than one row at a time

Example 7-26 Fetching Data with Cursor Variables

This example uses one cursor variable to do what Example 7-6 does with two explicit cursors. The first OPEN FOR statement includes the query itself. The second OPEN FOR statement references a variable whose value is a query.

```
DECLARE
    cv SYS_REFCURSOR; -- cursor variable

v_lastname employees.last_name%TYPE; -- variable for last_name
    v_jobid employees.job_id%TYPE; -- variable for job_id

query_2 VARCHAR2(200) :=
    'SELECT * FROM employees
    WHERE REGEXP_LIKE (job_id, ''[ACADFIMKSA]_M[ANGR]'')
    ORDER BY job_id';

v_employees employees%ROWTYPE; -- record variable row of table

BEGIN

OPEN cv FOR
    SELECT last_name, job_id FROM employees
    WHERE REGEXP_LIKE (job_id, 'S[HT]_CLERK')
    ORDER BY last name;
```



```
LOOP -- Fetches 2 columns into variables
   FETCH cv INTO v lastname, v jobid;
   EXIT WHEN cv%NOTFOUND;
   DBMS_OUTPUT.PUT_LINE( RPAD(v_lastname, 25, ' ') || v_jobid );
  END LOOP;
  DBMS OUTPUT.PUT LINE( '----');
  OPEN cv FOR query_2;
  LOOP -- Fetches entire row into the v employees record
   FETCH cv INTO v employees;
   EXIT WHEN cv % NOTFOUND;
   DBMS OUTPUT.PUT_LINE( RPAD(v_employees.last_name, 25, ' ') ||
                            v employees.job id );
  END LOOP;
  CLOSE cv;
END;
Result:
                      ST CLERK
Atkinson
Bell
                       SH CLERK
Bissot
                       ST CLERK
                       SH CLERK
Walsh
                      AC MGR
Gruenberg
                      FI MGR
Martinez
                      MK MAN
Errazuriz
                       SA MAN
```

Example 7-27 Fetching from Cursor Variable into Collections

This example fetches from a cursor variable into two collections (nested tables), using the BULK COLLECT clause of the FETCH statement.

```
DECLARE

TYPE empcurtyp IS REF CURSOR;

TYPE namelist IS TABLE OF employees.last_name%TYPE;

TYPE sallist IS TABLE OF employees.salary%TYPE;

emp_cv empcurtyp;

names namelist;

sals sallist;

BEGIN

OPEN emp_cv FOR

SELECT last_name, salary FROM employees

WHERE job_id = 'SA_REP'

ORDER BY salary DESC;

FETCH emp_cv BULK COLLECT INTO names, sals;

CLOSE emp_cv;

-- loop through the names and sals collections
```



```
FOR i IN names.FIRST .. names.LAST
LOOP

DBMS_OUTPUT.PUT_LINE

('Name = ' || names(i) || ', salary = ' || sals(i));
END LOOP;
END;

/

Result:

Name = Ozer, salary = 11500
Name = Abel, salary = 11000
Name = Vishney, salary = 10500
...
Name = Kumar, salary = 6100
```

Assigning Values to Cursor Variables

You can assign to a PL/SQL cursor variable the value of another PL/SQL cursor variable or host cursor variable.

The syntax is:

```
target_cursor_variable := source_cursor_variable;
```

If <code>source_cursor_variable</code> is open, then after the assignment, <code>target_cursor_variable</code> is also open. The two cursor variables point to the same SQL work area.

If $source_cursor_variable$ is not open, opening $target_cursor_variable$ after the assignment does not open $source_cursor_variable$.

Variables in Cursor Variable Queries

The query associated with a cursor variable can reference any variable in its scope.

When you open a cursor variable with the OPEN FOR statement, PL/SQL evaluates any variables in the query and uses those values when identifying the result set. Changing the values of the variables later does not change the result set.

To change the result set, you must change the value of the variable and then open the cursor variable again for the same query, as in Example 7-29.

Example 7-28 Variable in Cursor Variable Query—No Result Set Change

This example opens a cursor variable for a query that references the variable factor, which has the value 2. Therefore, sal_multiple is always 2 times sal, despite that factor is incremented after every fetch.

```
DECLARE

sal employees.salary%TYPE;

sal_multiple employees.salary%TYPE;

factor INTEGER := 2;

cv SYS_REFCURSOR;

BEGIN

OPEN cv FOR

SELECT salary, salary*factor

FROM employees

WHERE job id LIKE 'AD %'; -- PL/SQL evaluates factor
```



```
LOOP
    FETCH cv INTO sal, sal multiple;
    EXIT WHEN cv%NOTFOUND;
    DBMS_OUTPUT.PUT_LINE('factor = ' || factor);
DBMS_OUTPUT.PUT_LINE('sal = ' || sal);
    DBMS OUTPUT.PUT LINE('sal multiple = ' || sal multiple);
    factor := factor + 1; -- Does not affect sal multiple
  END LOOP;
  CLOSE cv;
END;
Result:
factor = 2
            = 4400
sal
sal_multiple = 8800
factor = 3
sal
            = 24000
sal_multiple = 48000
factor = 4
sal = 17000
sal multiple = 34000
factor = 5
            = 17000
sal multiple = 34000
```

Example 7-29 Variable in Cursor Variable Query—Result Set Change

```
DECLARE
 sal
               employees.salary%TYPE;
  sal multiple employees.salary%TYPE;
             INTEGER := 2;
  cv SYS REFCURSOR;
  DBMS OUTPUT.PUT LINE('factor = ' || factor);
  OPEN CV FOR
    SELECT salary, salary*factor
    FROM employees
   WHERE job_id LIKE 'AD_%'; -- PL/SQL evaluates factor
   FETCH cv INTO sal, sal_multiple;
   EXIT WHEN cv%NOTFOUND;
                                = ' || sal);
   DBMS_OUTPUT.PUT_LINE('sal
   DBMS OUTPUT.PUT LINE('sal multiple = ' || sal multiple);
  END LOOP;
  factor := factor + 1;
  DBMS OUTPUT.PUT LINE('factor = ' || factor);
  OPEN CV FOR
   SELECT salary, salary*factor
   FROM employees
    WHERE job_id LIKE 'AD_%'; -- PL/SQL evaluates factor
   FETCH cv INTO sal, sal multiple;
```

```
EXIT WHEN cv%NOTFOUND;
   DBMS OUTPUT.PUT LINE('sal
                            = ' || sal);
   DBMS_OUTPUT.PUT_LINE('sal_multiple = ' || sal_multiple);
 END LOOP;
 CLOSE cv;
END;
Result:
factor = 2
sal
           = 4400
sal multiple = 8800
sal = 24000
sal multiple = 48000
sal = 17000
sal multiple = 34000
sal = 17000
sal multiple = 34000
factor = 3
sal
           = 4400
sal_multiple = 13200
= 24000
sal_multiple = 72000
sal = 17000
sal_multiple = 51000
sal = 17000
sal multiple = 51000
```

Querying a Collection

You can query a collection if all of the following are true:

- The data type of the collection was either created at schema level or declared in a package specification.
- The data type of the collection element is either a scalar data type, a user-defined type, or a record type.

In the query FROM clause, the collection appears in table_collection_expression as the argument of the TABLE operator.

Note:

In SQL contexts, you cannot use a function whose return type was declared in a package specification.



See Also:

- Oracle Database SQL Language Reference for information about the table collection expression
- "CREATE PACKAGE Statement" for information about the CREATE PACKAGE statement
- "PL/SQL Collections and Records" for information about collection types and collection variables
- Example 8-9, "Querying a Collection with Native Dynamic SQL"

Example 7-30 Querying a Collection with Static SQL

In this example, the cursor variable is associated with a query on an associative array of records. The nested table type, mytab, is declared in a package specification.

```
CREATE OR REPLACE PACKAGE pkg AUTHID DEFINER AS
 TYPE rec IS RECORD(f1 NUMBER, f2 VARCHAR2(30));
 TYPE mytab IS TABLE OF rec INDEX BY pls integer;
END;
DECLARE
 v1 pkg.mytab; -- collection of records
 v2 pkg.rec;
 c1 SYS REFCURSOR;
BEGIN
 v1(1).f1 := 1;
 v1(1).f2 := 'one';
 OPEN c1 FOR SELECT * FROM TABLE(v1);
 FETCH c1 INTO v2;
 CLOSE c1;
 DBMS OUTPUT.PUT LINE('Values in record are ' || v2.f1 || ' and ' || v2.f2);
END;
Result:
```

Cursor Variable Attributes

A cursor variable has the same attributes as an explicit cursor (see Explicit Cursor Attributes.). The syntax for the value of a cursor variable attribute is <code>cursor_variable_name</code> immediately followed by <code>attribute</code> (for example, <code>cv%ISOPEN</code>). If a cursor variable is not open, referencing any attribute except <code>%ISOPEN</code> raises the predefined exception <code>INVALID</code> CURSOR.

Cursor Variables as Subprogram Parameters

Values in record are 1 and one

You can use a cursor variable as a subprogram parameter, which makes it useful for passing query results between subprograms.

For example:

- You can open a cursor variable in one subprogram and process it in a different subprogram.
- In a multilanguage application, a PL/SQL subprogram can use a cursor variable to return a result set to a subprogram written in a different language.

Note:

The invoking and invoked subprograms must be in the same database instance. You cannot pass or return cursor variables to subprograms invoked through database links.

A

Caution:

Because cursor variables are pointers, using them as subprogram parameters increases the likelihood of subprogram parameter aliasing, which can have unintended results. For more information, see "Subprogram Parameter Aliasing with Cursor Variable Parameters".

When declaring a cursor variable as the formal parameter of a subprogram:

- If the subprogram opens or assigns a value to the cursor variable, then the parameter mode must be IN OUT.
- If the subprogram only fetches from, or closes, the cursor variable, then the parameter mode can be either IN OT IN OUT.

Corresponding formal and actual cursor variable parameters must have compatible return types. Otherwise, PL/SQL raises the predefined exception ROWTYPE MISMATCH.

To pass a cursor variable parameter between subprograms in different PL/SQL units, define the REF CURSOR type of the parameter in a package. When the type is in a package, multiple subprograms can use it. One subprogram can declare a formal parameter of that type, and other subprograms can declare variables of that type and pass them to the first subprogram.

See Also:

- •
- "Subprogram Parameters" for more information about subprogram parameters
- "CURSOR Expressions" for information about CURSOR expressions, which can be actual parameters for formal cursor variable parameters
- PL/SQL Packages, for more information about packages

Example 7-31 Procedure to Open Cursor Variable for One Query

This example defines, in a package, a REF CURSOR type and a procedure that opens a cursor variable parameter of that type.



```
CREATE OR REPLACE PACKAGE emp_data AUTHID DEFINER AS

TYPE empcurtyp IS REF CURSOR RETURN employees%ROWTYPE;

PROCEDURE open_emp_cv (emp_cv IN OUT empcurtyp);

END emp_data;

/

CREATE OR REPLACE PACKAGE BODY emp_data AS

PROCEDURE open_emp_cv (emp_cv IN OUT EmpCurTyp) IS

BEGIN

OPEN emp_cv FOR SELECT * FROM employees;

END open_emp_cv;

END emp_data;

/
```

Example 7-32 Opening Cursor Variable for Chosen Query (Same Return Type)

In this example ,the stored procedure opens its cursor variable parameter for a chosen query. The queries have the same return type.

```
CREATE OR REPLACE PACKAGE emp data AUTHID DEFINER AS
 TYPE empcurtyp IS REF CURSOR RETURN employees%ROWTYPE;
  PROCEDURE open_emp_cv (emp_cv IN OUT empcurtyp, choice INT);
END emp_data;
CREATE OR REPLACE PACKAGE BODY emp data AS
 PROCEDURE open emp cv (emp cv IN OUT empcurtyp, choice INT) IS
   IF choice = 1 THEN
     OPEN emp cv FOR SELECT *
     FROM employees
     WHERE commission pct IS NOT NULL;
    ELSIF choice = 2 THEN
     OPEN emp cv FOR SELECT *
     FROM employees
     WHERE salary > 2500;
   ELSIF choice = 3 THEN
     OPEN emp cv FOR SELECT *
     FROM employees
     WHERE department id = 100;
   END IF;
 END;
END emp data;
```

Example 7-33 Opening Cursor Variable for Chosen Query (Different Return Types)

In this example, the stored procedure opens its cursor variable parameter for a chosen query. The queries have the different return types.

```
CREATE OR REPLACE PACKAGE admin_data AUTHID DEFINER AS

TYPE gencurtyp IS REF CURSOR;

PROCEDURE open_cv (generic_cv IN OUT gencurtyp, choice INT);

END admin_data;

/

CREATE OR REPLACE PACKAGE BODY admin_data AS

PROCEDURE open_cv (generic_cv IN OUT gencurtyp, choice INT) IS

BEGIN

IF choice = 1 THEN

OPEN generic_cv FOR SELECT * FROM employees;

ELSIF choice = 2 THEN

OPEN generic_cv FOR SELECT * FROM departments;

ELSIF choice = 3 THEN

OPEN generic_cv FOR SELECT * FROM jobs;

END IF;
```



```
END;
END admin_data;
/
```

Cursor Variables as Host Variables

You can use a cursor variable as a host variable, which makes it useful for passing query results between PL/SQL stored subprograms and their clients.

When a cursor variable is a host variable, PL/SQL and the client (the host environment) share a pointer to the SQL work area that stores the result set.

To use a cursor variable as a host variable, declare the cursor variable in the host environment and then pass it as an input host variable (bind variable) to PL/SQL. Host cursor variables are compatible with any query return type (like weak PL/SQL cursor variables).

A SQL work area remains accessible while any cursor variable points to it, even if you pass the value of a cursor variable from one scope to another. For example, in Example 7-34, the Pro*C program passes a host cursor variable to an embedded PL/SQL anonymous block. After the block runs, the cursor variable still points to the SQL work area.

If you have a PL/SQL engine on the client side, calls from client to server impose no restrictions. For example, you can declare a cursor variable on the client side, open and fetch from it on the server side, and continue to fetch from it on the client side. You can also reduce network traffic with a PL/SQL anonymous block that opens or closes several host cursor variables in a single round trip. For example:

```
/* PL/SQL anonymous block in host environment */
BEGIN
   OPEN :emp_cv FOR SELECT * FROM employees;
   OPEN :dept_cv FOR SELECT * FROM departments;
   OPEN :loc_cv FOR SELECT * FROM locations;
END;
//
```

Because the cursor variables still point to the SQL work areas after the PL/SQL anonymous block runs, the client program can use them. When the client program no longer needs the cursors, it can use a PL/SQL anonymous block to close them. For example:

```
/* PL/SQL anonymous block in host environment */
BEGIN
   CLOSE :emp_cv;
   CLOSE :dept_cv;
   CLOSE :loc_cv;
END;
//
```

This technique is useful for populating a multiblock form, as in Oracle Forms. For example, you can open several SQL work areas in a single round trip, like this:

```
/* PL/SQL anonymous block in host environment */
BEGIN
   OPEN :c1 FOR SELECT 1 FROM DUAL;
   OPEN :c2 FOR SELECT 1 FROM DUAL;
   OPEN :c3 FOR SELECT 1 FROM DUAL;
END;
/
```





If you bind a host cursor variable into PL/SQL from an Oracle Call Interface (OCI) client, then you cannot fetch from it on the server side unless you also open it there on the same server call.

Example 7-34 Cursor Variable as Host Variable in Pro*C Client Program

In this example, a Pro*C client program declares a cursor variable and a selector and passes them as host variables to a PL/SQL anonymous block, which opens the cursor variable for the selected query.

```
EXEC SQL BEGIN DECLARE SECTION;
  SQL CURSOR generic cv; -- Declare host cursor variable.
 int choice; -- Declare selector.
EXEC SQL END DECLARE SECTION;
EXEC SQL ALLOCATE :generic cv; -- Initialize host cursor variable.
-- Pass host cursor variable and selector to PL/SQL block.
EXEC SQL EXECUTE
BEGIN
 IF : choice = 1 THEN
   OPEN :generic cv FOR SELECT * FROM employees;
 ELSIF :choice = 2 THEN
   OPEN : generic cv FOR SELECT * FROM departments;
 ELSIF : choice = 3 THEN
   OPEN :generic_cv FOR SELECT * FROM jobs;
 END IF;
END:
END-EXEC;
```

CURSOR Expressions

A CURSOR expression returns a nested cursor.

It has this syntax:

```
CURSOR ( subquery )
```

You can use a CURSOR expression in a SELECT statement that is not a subquery (as in Example 7-35) or pass it to a function that accepts a cursor variable parameter (see "Passing CURSOR Expressions to Pipelined Table Functions"). You cannot use a cursor expression with an implicit cursor.

See Also:

Oracle Database SQL Language Reference for more information about CURSOR expressions, including restrictions

Example 7-35 CURSOR Expression

This example declares and defines an explicit cursor for a query that includes a cursor expression. For each department in the departments table, the nested cursor returns the last name of each employee in that department (which it retrieves from the employees table).

```
DECLARE
 TYPE emp cur typ IS REF CURSOR;
  emp_cur emp_cur_typ;
  dept name departments.department name%TYPE;
  emp name employees.last name%TYPE;
  CURSOR c1 IS
    SELECT department name,
     CURSOR ( SELECT e.last name
                FROM employees e
                WHERE e.department id = d.department id
                ORDER BY e.last name
              ) employees
    FROM departments d
    WHERE department name LIKE 'A%'
    ORDER BY department name;
BEGIN
 LOOP -- Process each row of query result set
    FETCH c1 INTO dept name, emp cur;
    EXIT WHEN c1%NOTFOUND;
    DBMS OUTPUT.PUT LINE('Department: ' || dept name);
    LOOP -- Process each row of subquery result set
     FETCH emp cur INTO emp name;
     EXIT WHEN emp cur%NOTFOUND;
     DBMS OUTPUT.PUT_LINE('-- Employee: ' || emp_name);
    END LOOP;
 END LOOP;
 CLOSE c1;
END;
Result:
Department: Accounting
-- Employee: Gietz
-- Employee: Higgins
Department: Administration
-- Employee: Whalen
```

Transaction Processing and Control

Transaction processing is an Oracle Database feature that lets multiple users work on the database concurrently, and ensures that each user sees a consistent version of data and that all changes are applied in the right order.

A **transaction** is a sequence of one or more SQL statements that Oracle Database treats as a unit: either all of the statements are performed, or none of them are.

Different users can write to the same data structures without harming each other's data or coordinating with each other, because Oracle Database locks data structures automatically. To maximize data availability, Oracle Database locks the minimum amount of data for the minimum amount of time.

You rarely must write extra code to prevent problems with multiple users accessing data concurrently. However, if you do need this level of control, you can manually override the Oracle Database default locking mechanisms.

Topics

- COMMIT Statement
- ROLLBACK Statement
- SAVEPOINT Statement
- Implicit Rollbacks
- SET TRANSACTION Statement
- Overriding Default Locking

See Also:

- Oracle Database Concepts for more information about transactions
- Oracle Database Concepts for more information about transaction processing
- Oracle Database Concepts for more information about the Oracle Database locking mechanism
- Oracle Database Concepts for more information about manual data locks

COMMIT Statement

The COMMIT statement ends the current transaction, making its changes permanent and visible to other users.



A transaction can span multiple blocks, and a block can contain multiple transactions.

The WRITE clause of the COMMIT statement specifies the priority with which Oracle Database writes to the redo log the information that the commit operation generates.



The default PL/SQL commit behavior for nondistributed transactions is BATCH NOWAIT if the COMMIT_LOGGING and COMMIT_WAIT database initialization parameters have not been set.



See Also:

- Oracle Database Concepts for more information about committing transactions
- Oracle Database Concepts for information about distributed transactions
- Oracle Database SQL Language Reference for information about the COMMIT statement
- Oracle Data Guard Concepts and Administration for information about ensuring no loss of data during a failover to a standby database

Example 7-36 COMMIT Statement with COMMENT and WRITE Clauses

In this example, a transaction transfers money from one bank account to another. It is important that the money both leaves one account and enters the other, hence the COMMIT WRITE IMMEDIATE NOWAIT Statement.

```
DROP TABLE accounts;
CREATE TABLE accounts (
  account id NUMBER(6),
 balance NUMBER (10,2)
INSERT INTO accounts (account_id, balance)
VALUES (7715, 6350.00);
INSERT INTO accounts (account_id, balance)
VALUES (7720, 5100.50);
CREATE OR REPLACE PROCEDURE transfer (
  from acct NUMBER,
  to_acct NUMBER,
 amount NUMBER
) AUTHID CURRENT USER AS
BEGIN
  UPDATE accounts
  SET balance = balance - amount
  WHERE account id = from acct;
  UPDATE accounts
  SET balance = balance + amount
  WHERE account id = to acct;
  COMMIT WRITE IMMEDIATE NOWAIT;
END;
```

Query before transfer:

SELECT * FROM accounts;

Result:

BEGIN

BALANCE	ACCOUNT_ID
6350	7715
5100.5	7720



```
transfer(7715, 7720, 250);
END;
/
```

Query after transfer:

```
SELECT * FROM accounts;
```

Result:

ACCOUNT_ID	BALANCE
7715	6100
7720	5350.5

ROLLBACK Statement

The ROLLBACK statement ends the current transaction and undoes any changes made during that transaction.

If you make a mistake, such as deleting the wrong row from a table, a rollback restores the original data. If you cannot finish a transaction because a SQL statement fails or PL/SQL raises an exception, a rollback lets you take corrective action and perhaps start over.



Oracle Database SQL Language Reference for more information about the ROLLBACK statement

Example 7-37 ROLLBACK Statement

This example inserts information about an employee into three different tables. If an INSERT statement tries to store a duplicate employee number, PL/SQL raises the predefined exception DUP_VAL_ON_INDEX. To ensure that changes to all three tables are undone, the exception handler runs a ROLLBACK.

```
DROP TABLE emp_name;
CREATE TABLE emp_name AS
    SELECT employee_id, last_name
    FROM employees;

CREATE UNIQUE INDEX empname_ix
ON emp_name (employee_id);

DROP TABLE emp_sal;
CREATE TABLE emp_sal AS
    SELECT employee_id, salary
    FROM employees;

CREATE UNIQUE INDEX empsal_ix
ON emp_sal (employee_id);

DROP TABLE emp_job;
CREATE TABLE emp_job AS
    SELECT employee_id, job_id
```



```
FROM employees;
CREATE UNIQUE INDEX empjobid ix
ON emp job (employee id);
DECLARE
 emp id NUMBER(6);
 emp_lastname VARCHAR2(25);
 emp_salary NUMBER(8,2);
emp_jobid VARCHAR2(10);
BEGIN
 SELECT employee_id, last_name, salary, job_id
 INTO emp id, emp lastname, emp salary, emp jobid
 FROM employees
 WHERE employee id = 120;
 INSERT INTO emp name (employee id, last name)
 VALUES (emp id, emp lastname);
  INSERT INTO emp sal (employee id, salary)
 VALUES (emp id, emp salary);
  INSERT INTO emp job (employee id, job id)
 VALUES (emp id, emp jobid);
EXCEPTION
 WHEN DUP VAL ON INDEX THEN
   ROLLBACK;
    DBMS OUTPUT.PUT LINE('Inserts were rolled back');
END;
```

SAVEPOINT Statement

The SAVEPOINT statement names and marks the current point in the processing of a transaction.

Savepoints let you roll back part of a transaction instead of the whole transaction. The number of active savepoints for each session is unlimited.

When you roll back to a savepoint, any savepoints marked after that savepoint are erased. The savepoint to which you roll back is not erased. A simple rollback or commit erases all savepoints.

If you mark a savepoint in a recursive subprogram, new instances of the SAVEPOINT statement run at each level in the recursive descent, but you can only roll back to the most recently marked savepoint.

Savepoint names are undeclared identifiers. Reusing a savepoint name in a transaction moves the savepoint from its old position to the current point in the transaction, which means that a rollback to the savepoint affects only the current part of the transaction.

See Also:

Oracle Database SQL Language Reference for more information about the ${\tt SET}$ TRANSACTION SQL statement

Example 7-38 SAVEPOINT and ROLLBACK Statements

This example marks a savepoint before doing an insert. If the <code>INSERT</code> statement tries to store a duplicate value in the <code>employee_id</code> column, PL/SQL raises the predefined exception <code>DUP_VAL_ON_INDEX</code> and the transaction rolls back to the savepoint, undoing only the <code>INSERT</code> statement.

```
DROP TABLE emp name;
CREATE TABLE emp name AS
 SELECT employee id, last name, salary
 FROM employees;
CREATE UNIQUE INDEX empname ix
ON emp name (employee id);
DECLARE
 emp id
           employees.employee id%TYPE;
 emp_lastname employees.last_name%TYPE;
 emp salary employees.salary%TYPE;
BEGIN
  SELECT employee_id, last_name, salary
 INTO emp_id, emp_lastname, emp_salary
 FROM employees
 WHERE employee id = 120;
 UPDATE emp name
 SET salary = salary * 1.1
 WHERE employee id = emp id;
 DELETE FROM emp name
 WHERE employee id = 130;
 SAVEPOINT do insert;
 INSERT INTO emp name (employee id, last name, salary)
 VALUES (emp id, emp lastname, emp salary);
EXCEPTION
 WHEN DUP VAL ON INDEX THEN
   ROLLBACK TO do insert;
 DBMS OUTPUT.PUT LINE('Insert was rolled back');
END;
```

Example 7-39 Reusing SAVEPOINT with ROLLBACK

```
DROP TABLE emp_name;

CREATE TABLE emp_name AS

SELECT employee_id, last_name, salary

FROM employees;

CREATE UNIQUE INDEX empname_ix

ON emp_name (employee_id);

DECLARE

emp_id employees.employee_id%TYPE;

emp_lastname employees.last_name%TYPE;

emp_salary employees.salary%TYPE;

BEGIN

SELECT employee id, last name, salary
```



```
INTO emp id, emp_lastname, emp_salary
 FROM employees
 WHERE employee_id = 120;
 SAVEPOINT my_savepoint;
 UPDATE emp name
 SET salary = salary * 1.1
 WHERE employee id = emp id;
 DELETE FROM emp name
 WHERE employee id = 130;
 SAVEPOINT my_savepoint;
 INSERT INTO emp name (employee id, last name, salary)
 VALUES (emp id, emp lastname, emp salary);
EXCEPTION
 WHEN DUP VAL ON INDEX THEN
   ROLLBACK TO my savepoint;
   DBMS OUTPUT.PUT LINE('Transaction rolled back.');
END;
```

Implicit Rollbacks

Before running an INSERT, UPDATE, DELETE, or MERGE statement, the database marks an implicit savepoint (unavailable to you). If the statement fails, the database rolls back to the savepoint.

Usually, just the failed SQL statement is rolled back, not the whole transaction. If the statement raises an unhandled exception, the host environment determines what is rolled back.

The database can also roll back single SQL statements to break deadlocks. The database signals an error to a participating transaction and rolls back the current statement in that transaction.

Before running a SQL statement, the database must parse it, that is, examine it to ensure it follows syntax rules and refers to valid schema objects. Errors detected while running a SQL statement cause a rollback, but errors detected while parsing the statement do not.

If you exit a stored subprogram with an unhandled exception, PL/SQL does not assign values to OUT parameters, and does not do any rollback.

For information about handling exceptions, see PL/SQL Error Handling

SET TRANSACTION Statement

You use the SET TRANSACTION statement to begin a read-only or read-write transaction, establish an isolation level, or assign your current transaction to a specified rollback segment.

Read-only transactions are useful for running multiple queries while other users update the same tables.

During a read-only transaction, all queries refer to the same snapshot of the database, providing a multi-table, multi-query, read-consistent view. Other users can continue to query or update data as usual. A commit or rollback ends the transaction.

The SET TRANSACTION statement must be the first SQL statement in a read-only transaction and can appear only once in a transaction. If you set a transaction to READ ONLY, subsequent

queries see only changes committed before the transaction began. The use of READ ONLY does not affect other users or transactions.

Only the SELECT, OPEN, FETCH, CLOSE, LOCK TABLE, COMMIT, and ROLLBACK statements are allowed in a read-only transaction. Queries cannot be FOR UPDATE.



Oracle Database SQL Language Reference for more information about the SQL statement SET TRANSACTION

Example 7-40 SET TRANSACTION Statement in Read-Only Transaction

In this example, a read-only transaction gather order totals for the day, the past week, and the past month. The totals are unaffected by other users updating the database during the transaction. The orders table is in the sample schema OE.

```
DECLARE
 daily order total NUMBER(12,2);
 weekly_order_total NUMBER(12,2);
 monthly order total NUMBER(12,2);
  COMMIT; -- end previous transaction
  SET TRANSACTION READ ONLY NAME 'Calculate Order Totals';
  SELECT SUM (order total)
  INTO daily order total
  FROM orders
  WHERE order_date = SYSDATE;
  SELECT SUM (order total)
  INTO weekly_order_total
  FROM orders
  WHERE order date = SYSDATE - 7;
  SELECT SUM (order total)
  INTO monthly_order_total
  FROM orders
  WHERE order date = SYSDATE - 30;
  COMMIT; -- ends read-only transaction
END:
```

Overriding Default Locking

By default, Oracle Database locks data structures automatically, which lets different applications write to the same data structures without harming each other's data or coordinating with each other.

If you must have exclusive access to data during a transaction, you can override default locking with these SQL statements:

- LOCK TABLE, which explicitly locks entire tables.
- SELECT with the FOR UPDATE clause (SELECT FOR UPDATE), which explicitly locks specific rows
 of a table.

Topics

- LOCK TABLE Statement
- SELECT FOR UPDATE and FOR UPDATE Cursors
- Simulating CURRENT OF Clause with ROWID Pseudocolumn

LOCK TABLE Statement

The LOCK TABLE statement explicitly locks one or more tables in a specified lock mode so that you can share or deny access to them.

The lock mode determines what other locks can be placed on the table. For example, many users can acquire row share locks on a table at the same time, but only one user at a time can acquire an exclusive lock. While one user has an exclusive lock on a table, no other users can insert, delete, or update rows in that table.

A table lock never prevents other users from querying a table, and a query never acquires a table lock. Only if two different transactions try to modify the same row does one transaction wait for the other to complete. The LOCK TABLE statement lets you specify how long to wait for another transaction to complete.

Table locks are released when the transaction that acquired them is either committed or rolled back.

See Also:

- Oracle Database Development Guide for more information about locking tables explicitly
- Oracle Database SQL Language Reference for more information about the LOCK TABLE statement

SELECT FOR UPDATE and FOR UPDATE Cursors

The SELECT statement with the FOR UPDATE clause (SELECT FOR UPDATE statement) selects the rows of the result set and locks them. SELECT FOR UPDATE lets you base an update on the existing values in the rows, because it ensures that no other user can change those values before you update them. You can also use SELECT FOR UPDATE to lock rows that you do not want to update, as in Example 10-6.



In tables compressed with Hybrid Columnar Compression (HCC), DML statements lock compression units rather than rows. HCC, a feature of certain Oracle storage systems, is described in *Oracle Database Concepts*.

By default, the SELECT FOR UPDATE statement waits until the requested row lock is acquired. To change this behavior, use the NOWAIT, WAIT, or SKIP LOCKED clause of the SELECT FOR UPDATE

statement. For information about these clauses, see *Oracle Database SQL Language Reference*.

When SELECT FOR UPDATE is associated with an explicit cursor, the cursor is called a **FOR UPDATE** cursor. Only a FOR UPDATE cursor can appear in the CURRENT OF clause of an UPDATE or DELETE statement. (The CURRENT OF clause, a PL/SQL extension to the WHERE clause of the SQL statements UPDATE and DELETE, restricts the statement to the current row of the cursor.)

When SELECT FOR UPDATE queries multiple tables, it locks only rows whose columns appear in the FOR UPDATE clause.

Simulating CURRENT OF Clause with ROWID Pseudocolumn

The rows of the result set are locked when you open a FOR UPDATE cursor, not as they are fetched. The rows are unlocked when you commit or roll back the transaction. After the rows are unlocked, you cannot fetch from the FOR UPDATE cursor, as Example 7-41 shows (the result is the same if you substitute ROLLBACK for COMMIT).

The workaround is to simulate the CURRENT OF clause with the ROWID pseudocolumn (described in *Oracle Database SQL Language Reference*). Select the rowid of each row into a UROWID variable and use the rowid to identify the current row during subsequent updates and deletes, as in Example 7-42. (To print the value of a UROWID variable, convert it to VARCHAR2, using the ROWIDTOCHAR function described in *Oracle Database SQL Language Reference*.)



When you update a row in a table compressed with Hybrid Columnar Compression (HCC), the ROWID of the row changes. HCC, a feature of certain Oracle storage systems, is described in *Oracle Database Concepts*.

A

Caution:

Because no FOR UPDATE clause locks the fetched rows, other users might unintentionally overwrite your changes.



The extra space needed for read consistency is not released until the cursor is closed, which can slow down processing for large updates.

Example 7-41 FETCH with FOR UPDATE Cursor After COMMIT Statement

```
DROP TABLE emp;
CREATE TABLE emp AS SELECT * FROM employees;
DECLARE
CURSOR c1 IS
SELECT * FROM emp
FOR UPDATE OF salary
```

```
ORDER BY employee id;
  emp_rec emp%ROWTYPE;
BEGIN
  OPEN c1;
  LOOP
    FETCH c1 INTO emp_rec; -- fails on second iteration
    EXIT WHEN c1%NOTFOUND;
    DBMS OUTPUT.PUT LINE (
      'emp_rec.employee_id = ' ||
      TO_CHAR(emp_rec.employee_id)
    UPDATE emp
    SET salary = salary * 1.05
    WHERE employee id = 105;
    COMMIT; -- releases locks
  END LOOP;
END;
Result:
emp_rec.employee_id = 100
DECLARE
ERROR at line 1:
ORA-01002: fetch out of sequence
ORA-06512: at line 11
```

Example 7-42 Simulating CURRENT OF Clause with ROWID Pseudocolumn

```
DROP TABLE emp;
CREATE TABLE emp AS SELECT * FROM employees;
DECLARE
  CURSOR c1 IS
   SELECT last_name, job_id, rowid
   FROM emp; -- no FOR UPDATE clause
  my lastname employees.last name%TYPE;
  my_jobid
               employees.job_id%TYPE;
  my_rowid
                UROWID;
BEGIN
  OPEN c1;
  LOOP
    FETCH c1 INTO my_lastname, my_jobid, my_rowid;
    EXIT WHEN c1%NOTFOUND;
    UPDATE emp
    SET salary = salary * 1.02
    WHERE rowid = my_rowid; -- simulates WHERE CURRENT OF c1
   COMMIT;
  END LOOP;
  CLOSE c1;
END;
```

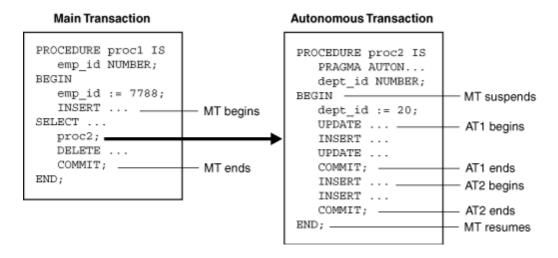
Autonomous Transactions

An **autonomous transaction** is an independent transaction started by another transaction, the main transaction.

Autonomous transactions do SQL operations and commit or roll back, without committing or rolling back the main transaction.

Figure 7-1 shows how control flows from the main transaction (MT) to an autonomous routine (proc2) and back again. The autonomous routine commits two autonomous transactions (AT1 and AT2).

Figure 7-1 Transaction Control Flow



Note:

Although an autonomous transaction is started by another transaction, it is not a nested transaction, because:

- It does not share transactional resources (such as locks) with the main transaction.
- It does not depend on the main transaction.
 - For example, if the main transaction rolls back, nested transactions roll back, but autonomous transactions do not.
- Its committed changes are visible to other transactions immediately.
 - A nested transaction's committed changes are not visible to other transactions until the main transaction commits.
- Exceptions raised in an autonomous transaction cause a transaction-level rollback, not a statement-level rollback.

Topics

Advantages of Autonomous Transactions



- Transaction Context
- Transaction Visibility
- Declaring Autonomous Routines
- Controlling Autonomous Transactions
- Autonomous Triggers
- Invoking Autonomous Functions from SQL



Oracle Database Development Guide for more information about autonomous transactions

Advantages of Autonomous Transactions

After starting, an autonomous transaction is fully independent. It shares no locks, resources, or commit-dependencies with the main transaction. You can log events, increment retry counters, and so on, even if the main transaction rolls back.

Autonomous transactions help you build modular, reusable software components. You can encapsulate autonomous transactions in stored subprograms. An invoking application needs not know whether operations done by that stored subprogram succeeded or failed.

Transaction Context

The main transaction shares its context with nested routines, but not with autonomous transactions. When one autonomous routine invokes another (or itself, recursively), the routines share no transaction context. When an autonomous routine invokes a nonautonomous routine, the routines share the same transaction context.

Transaction Visibility

Changes made by an autonomous transaction become visible to other transactions when the autonomous transaction commits. These changes become visible to the main transaction when it resumes, if its isolation level is set to READ COMMITTED (the default).

If you set the isolation level of the main transaction to SERIALIZABLE, changes made by its autonomous transactions are *not* visible to the main transaction when it resumes:

SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;



- Transaction properties apply only to the transaction in which they are set.
- Cursor attributes are not affected by autonomous transactions.



Declaring Autonomous Routines

To declare an autonomous routine, use the AUTONOMOUS TRANSACTION pragma.

For information about this pragma, see "AUTONOMOUS TRANSACTION Pragma".



Tip:

For readability, put the AUTONOMOUS_TRANSACTION pragma at the top of the declarative section. (The pragma is allowed anywhere in the declarative section.)

You cannot apply the AUTONOMOUS_TRANSACTION pragma to an entire package or ADT, but you can apply it to each subprogram in a package or each method of an ADT.

Example 7-43 Declaring Autonomous Function in Package

This example marks a package function as autonomous.

```
CREATE OR REPLACE PACKAGE emp actions AUTHID DEFINER AS -- package specification
 FUNCTION raise salary (emp id NUMBER, sal raise NUMBER)
 RETURN NUMBER;
END emp_actions;
CREATE OR REPLACE PACKAGE BODY emp actions AS -- package body
 -- code for function raise salary
 FUNCTION raise_salary (emp_id NUMBER, sal_raise NUMBER)
 RETURN NUMBER IS
   PRAGMA AUTONOMOUS TRANSACTION;
   new sal NUMBER(8,2);
 BEGIN
   UPDATE employees SET salary =
     salary + sal raise WHERE employee id = emp id;
   COMMIT;
   SELECT salary INTO new sal FROM employees
     WHERE employee id = emp id;
   RETURN new sal;
 END raise salary;
END emp actions;
```

Example 7-44 Declaring Autonomous Standalone Procedure

This example marks a standalone subprogram as autonomous.

```
CREATE OR REPLACE PROCEDURE lower_salary
  (emp_id NUMBER, amount NUMBER)
AUTHID DEFINER AS
  PRAGMA AUTONOMOUS_TRANSACTION;
BEGIN
  UPDATE employees
  SET salary = salary - amount
  WHERE employee_id = emp_id;
  COMMIT;
END lower_salary;
//
```



Example 7-45 Declaring Autonomous PL/SQL Block

This example marks a schema-level PL/SQL block as autonomous. (A nested PL/SQL block cannot be autonomous.)

```
DROP TABLE emp;
CREATE TABLE emp AS SELECT * FROM employees;

DECLARE

PRAGMA AUTONOMOUS_TRANSACTION;
emp_id NUMBER(6) := 200;
amount NUMBER(6,2) := 200;

BEGIN

UPDATE employees
SET salary = salary - amount
WHERE employee_id = emp_id;

COMMIT;
END;
/
```

Controlling Autonomous Transactions

The first SQL statement in an autonomous routine begins a transaction. When one transaction ends, the next SQL statement begins another transaction. All SQL statements run since the last commit or rollback comprise the current transaction. To control autonomous transactions, use these statements, which apply only to the current (active) transaction:

- COMMIT
- ROLLBACK [TO savepoint_name]
- SAVEPOINT savepoint name
- SET TRANSACTION

Topics

- Entering and Exiting Autonomous Routines
- Committing and Rolling Back Autonomous Transactions
- Savepoints
- Avoiding Errors with Autonomous Transactions

Entering and Exiting Autonomous Routines

When you enter the executable section of an autonomous routine, the main transaction suspends. When you exit the routine, the main transaction resumes.

If you try to exit an active autonomous transaction without committing or rolling back, the database raises an exception. If the exception is unhandled, or if the transaction ends because of some other unhandled exception, then the transaction rolls back.

To exit normally, the routine must explicitly commit or roll back all autonomous transactions. If the routine (or any routine invoked by it) has pending transactions, then PL/SQL raises an exception and the pending transactions roll back.

Committing and Rolling Back Autonomous Transactions

COMMIT and ROLLBACK end the active autonomous transaction but do not exit the autonomous routine. When one transaction ends, the next SQL statement begins another transaction. A single autonomous routine can contain several autonomous transactions, if it issues several COMMIT statements.

Savepoints

The scope of a savepoint is the transaction in which it is defined. Savepoints defined in the main transaction are unrelated to savepoints defined in its autonomous transactions. In fact, the main transaction and an autonomous transaction can use the same savepoint names.

You can roll back only to savepoints marked in the current transaction. In an autonomous transaction, you cannot roll back to a savepoint marked in the main transaction. To do so, you must resume the main transaction by exiting the autonomous routine.

When in the main transaction, rolling back to a savepoint marked before you started an autonomous transaction does *not* roll back the autonomous transaction. Remember, autonomous transactions are fully independent of the main transaction.

Avoiding Errors with Autonomous Transactions

To avoid some common errors, remember:

- If an autonomous transaction tries to access a resource held by the main transaction, a
 deadlock can occur. The database raises an exception in the autonomous transaction,
 which rolls back if the exception is unhandled.
- The database initialization parameter TRANSACTIONS specifies the maximum number of concurrent transactions. That number might be exceeded because an autonomous transaction runs concurrently with the main transaction.
- If you try to exit an active autonomous transaction without committing or rolling back, the database raises an exception. If the exception is unhandled, the transaction rolls back.
- You cannot run a PIPE ROW statement in an autonomous routine while an autonomous transaction is open. You must close the autonomous transaction before running the PIPE ROW statement. This is normally accomplished by committing or rolling back the autonomous transaction before running the PIPE ROW statement.

Autonomous Triggers

A trigger must be autonomous to run TCL or DDL statements.

To run DDL statements, the trigger must use native dynamic SQL.



See Also:

- PL/SQL Triggers, for general information about triggers
- "Description of Static SQL" for general information about TCL statements
- Oracle Database SQL Language Reference for information about DDL statements
- "Native Dynamic SQL" for information about native dynamic SQL

One use of triggers is to log events transparently—for example, to log all inserts into a table, even those that roll back.

Example 7-46 Autonomous Trigger Logs INSERT Statements

In this example, whenever a row is inserted into the EMPLOYEES table, a trigger inserts the same row into a log table. Because the trigger is autonomous, it can commit changes to the log table regardless of whether they are committed to the main table.

```
DROP TABLE emp;
CREATE TABLE emp AS SELECT * FROM employees;
-- Log table:
DROP TABLE log;
CREATE TABLE log (
 log_id NUMBER(6),
 up date DATE,
 new sal NUMBER(8,2),
  old sal NUMBER (8,2)
);
-- Autonomous trigger on emp table:
CREATE OR REPLACE TRIGGER log sal
  BEFORE UPDATE OF salary ON emp FOR EACH ROW
  PRAGMA AUTONOMOUS TRANSACTION;
BEGIN
  INSERT INTO log (
   log id,
   up date,
   new sal,
   old sal
  VALUES (
    :old.employee id,
    SYSDATE,
    :new.salary,
    :old.salary
  );
  COMMIT;
END;
UPDATE emp
SET salary = salary * 1.05
WHERE employee id = 115;
```



COMMIT;