# PL/SQL Optimization and Tuning

This chapter explains how the PL/SQL compiler optimizes your code and how to write efficient PL/SQL code and improve existing PL/SQL code.

### **Topics**

- PL/SQL Optimizer
- Candidates for Tuning
- Minimizing CPU Overhead
- Bulk SQL and Bulk Binding
- Chaining Pipelined Table Functions for Multiple Transformations
- Overview of Polymorphic Table Functions
- Updating Large Tables in Parallel
- Collecting Data About User-Defined Identifiers
- Profiling and Tracing PL/SQL Programs
- Compiling PL/SQL Units for Native Execution



Oracle Database Development Guide for disadvantages of cursor variables

# PL/SQL Optimizer

Prior to Oracle Database 10g release 1, the PL/SQL compiler translated your source text to system code without applying many changes to improve performance. Now, PL/SQL uses an optimizer that can rearrange code for better performance.

The optimizer is enabled by default. In rare cases, if the overhead of the optimizer makes compilation of very large applications too slow, you can lower the optimization by setting the compilation parameter PLSQL\_OPTIMIZE\_LEVEL=1 instead of its default value 2. In even rarer cases, PL/SQL might raise an exception earlier than expected or not at all. Setting PLSQL OPTIMIZE LEVEL=1 prevents the code from being rearranged.

# See Also:

- Oracle Database Reference for information about the PLSQL\_OPTIMIZE\_LEVEL compilation parameter
- Oracle Database Development Guide for examples of changing the PLSQL OPTIMIZE LEVEL compilation parameter
- Oracle Database Reference for information about the static dictionary view ALL PLSQL OBJECT SETTINGS

# Subprogram Inlining

One optimization that the compiler can perform is subprogram inlining.

Subprogram inlining replaces a subprogram invocation with a copy of the invoked subprogram (if the invoked and invoking subprograms are in the same program unit). To allow subprogram inlining, either accept the default value of the PLSQL\_OPTIMIZE\_LEVEL compilation parameter (which is 2) or set it to 3.

With PLSQL\_OPTIMIZE\_LEVEL=2, you must specify each subprogram to be inlined with the INLINE pragma:

```
PRAGMA INLINE (subprogram, 'YES')
```

If *subprogram* is overloaded, then the preceding pragma applies to every subprogram with that name.

With PLSQL\_OPTIMIZE\_LEVEL=3, the PL/SQL compiler seeks opportunities to inline subprograms. You need not specify subprograms to be inlined. However, you can use the INLINE pragma (with the preceding syntax) to give a subprogram a high priority for inlining, and then the compiler inlines it unless other considerations or limits make the inlining undesirable.

If a particular subprogram is inlined, performance almost always improves. However, because the compiler inlines subprograms early in the optimization process, it is possible for subprogram inlining to preclude later, more powerful optimizations.

If subprogram inlining slows the performance of a particular PL/SQL program, then use the PL/SQL hierarchical profiler to identify subprograms for which you want to turn off inlining. To turn off inlining for a subprogram, use the INLINE pragma:

```
PRAGMA INLINE (subprogram, 'NO')
```

The INLINE pragma affects only the immediately following declaration or statement, and only some kinds of statements.

When the INLINE pragma immediately precedes a declaration, it affects:

- Every invocation of the specified subprogram in that declaration
- Every initialization value in that declaration except the default initialization values of records

When the INLINE pragma immediately precedes one of these statements, the pragma affects every invocation of the specified subprogram in that statement:

Assignment



- CALL
- Conditional
- CASE
- CONTINUE WHEN
- EXECUTE IMMEDIATE
- EXIT WHEN
- LOOP
- RETURN

The INLINE pragma does not affect statements that are not in the preceding list.

Multiple pragmas can affect the same declaration or statement. Each pragma applies its own effect to the statement. If PRAGMA INLINE (subprogram, 'YES') and PRAGMA INLINE (identifier, 'NO') have the same subprogram, then 'NO' overrides 'YES'. One PRAGMA

INLINE (*subprogram*, 'NO') overrides any number of occurrences of PRAGMA INLINE (*subprogram*, 'YES'), and the order of these pragmas is not important.

# See Also:

- Oracle Database Development Guide for more information about PL/SQL hierarchical profiler
- Oracle Database Reference for information about the PLSQL\_OPTIMIZE\_LEVEL compilation parameter
- Oracle Database Reference for information about the static dictionary view ALL PLSQL OBJECT SETTINGS

### Example 13-1 Specifying that Subprogram Is To Be Inlined

In this example, if  $PLSQL\_OPTIMIZE\_LEVEL=2$ , the INLINE pragma affects the procedure invocations p1(1) and p1(2), but not the procedure invocations p1(3) and p1(4).

```
PROCEDURE p1 (x PLS_INTEGER) IS ...

PRAGMA INLINE (p1, 'YES');

x:= p1(1) + p1(2) + 17; -- These 2 invocations to p1 are inlined ...

x:= p1(3) + p1(4) + 17; -- These 2 invocations to p1 are not inlined ...
```

#### Example 13-2 Specifying that Overloaded Subprogram Is To Be Inlined

In this example, if PLSQL\_OPTIMIZE\_LEVEL=2, the INLINE pragma affects both functions named p2.

```
FUNCTION p2 (p boolean) return PLS_INTEGER IS ...
FUNCTION p2 (x PLS_INTEGER) return PLS_INTEGER IS ...
...
PRAGMA INLINE(p2, 'YES');
x := p2(true) + p2(3);
...
```



### Example 13-3 Specifying that Subprogram Is Not To Be Inlined

In this example, the INLINE pragma affects the procedure invocations p1(1) and p1(2), but not the procedure invocations p1(3) and p1(4).

```
PROCEDURE p1 (x PLS_INTEGER) IS ...

PRAGMA INLINE (p1, 'NO');

x:= p1(1) + p1(2) + 17; -- These 2 invocations to p1 are not inlined ...

x:= p1(3) + p1(4) + 17; -- These 2 invocations to p1 might be inlined ...
```

### Example 13-4 PRAGMA INLINE ... 'NO' Overrides PRAGMA INLINE ... 'YES'

In this example, the second INLINE pragma overrides both the first and third INLINE pragmas.

```
PROCEDURE p1 (x PLS_INTEGER) IS ...

PRAGMA INLINE (p1, 'YES');

PRAGMA INLINE (p1, 'NO');

PRAGMA INLINE (p1, 'YES');

x:= p1(1) + p1(2) + 17; -- These 2 invocations to p1 are not inlined
```

# **Candidates for Tuning**

The following kinds of PL/SQL code are very likely to benefit from tuning:

Older code that does not take advantage of new PL/SQL language features.



#### Tip:

Before tuning older code, benchmark the current system and profile the older subprograms that your program invokes (see "Profiling and Tracing PL/SQL Programs"). With the many automatic optimizations of the PL/SQL optimizer (described in "PL/SQL Optimizer"), you might see performance improvements before doing any tuning.

Older dynamic SQL statements written with the DBMS SQL package.

If you know at compile time the number and data types of the input and output variables of a dynamic SQL statement, then you can rewrite the statement in native dynamic SQL, which runs noticeably faster than equivalent code that uses the <code>DBMS\_SQL</code> package (especially when it can be optimized by the compiler). For more information, see PL/SQL Dynamic SQL.

Code that spends much time processing SQL statements.

See "Tune SQL Statements".

Functions invoked in queries, which might run millions of times.

See "Tune Function Invocations in Queries".

Code that spends much time looping through query results.

See "Tune Loops".

Code that does many numeric computations.

See "Tune Computation-Intensive PL/SQL Code".

 Code that spends much time processing PL/SQL statements (as opposed to issuing database definition language (DDL) statements that PL/SQL passes directly to SQL).

See "Compiling PL/SQL Units for Native Execution".

# Minimizing CPU Overhead

### **Topics**

- Tune SQL Statements
- Tune Function Invocations in Queries
- Tune Subprogram Invocations
- Tune Loops
- Tune Computation-Intensive PL/SQL Code
- Use SQL Character Functions
- Put Least Expensive Conditional Tests First

# Tune SQL Statements

The most common cause of slowness in PL/SQL programs is slow SQL statements. To make SQL statements in a PL/SQL program as efficient as possible:

Use appropriate indexes.

For details, see Oracle Database Performance Tuning Guide.

Use query hints to avoid unnecessary full-table scans.

For details, see Oracle Database SQL Language Reference.

Collect current statistics on all tables, using the subprograms in the DBMS STATS package.

For details, see Oracle Database Performance Tuning Guide.

- Analyze the execution plans and performance of the SQL statements, using:
  - EXPLAIN PLAN statement

For details, see Oracle Database Performance Tuning Guide.

SQL Trace facility with TKPROF utility

For details, see Oracle Database Performance Tuning Guide.

 Use bulk SQL, a set of PL/SQL features that minimizes the performance overhead of the communication between PL/SQL and SQL.

For details, see "Bulk SQL and Bulk Binding".

# Tune Function Invocations in Queries

Functions invoked in queries might run millions of times. Do not invoke a function in a query unnecessarily, and make the invocation as efficient as possible.

Create a function-based index on the table in the query. The CREATE INDEX statement might take a while, but the query can run much faster because the function value for each row is cached.



If the query passes a column to a function, then the query cannot use user-created indexes on that column, so the query might invoke the function for every row of the table (which might be very large). To minimize the number of function invocations, use a nested query. Have the inner query filter the result set to a small number of rows, and have the outer query invoke the function for only those rows.

# See Also:

- Oracle Database SQL Language Reference for more information about CREATE INDEX statement syntax
- "PL/SQL Function Result Cache" for information about caching the results of PL/SQL functions

### Example 13-5 Nested Query Improves Performance

In this example, the two queries produce the same result set, but the second query is more efficient than the first. (In the example, the times and time difference are very small, because the EMPLOYEES table is very small. For a very large table, they would be significant.)

```
DECLARE
 starting_time TIMESTAMP WITH TIME ZONE;
 ending_time     TIMESTAMP WITH TIME ZONE;
BEGIN
  -- Invokes SQRT for every row of employees table:
 SELECT SYSTIMESTAMP INTO starting time FROM DUAL;
 FOR item IN (
   SELECT DISTINCT(SQRT(department_id)) col_alias
   FROM employees
   ORDER BY col alias
   DBMS OUTPUT.PUT LINE('Square root of dept. ID = ' || item.col alias);
 END LOOP;
 SELECT SYSTIMESTAMP INTO ending time FROM DUAL;
 DBMS OUTPUT.PUT LINE('Time = ' || TO CHAR(ending time - starting time));
 -- Invokes SQRT for every distinct department id of employees table:
 SELECT SYSTIMESTAMP INTO starting_time FROM DUAL;
 FOR item IN (
   SELECT SQRT(department_id) col alias
   FROM (SELECT DISTINCT department id FROM employees)
   ORDER BY col alias
    IF item.col alias IS NOT NULL THEN
     DBMS_OUTPUT.PUT_LINE('Square root of dept. ID = ' || item.col_alias);
   END IF;
 END LOOP;
 SELECT SYSTIMESTAMP INTO ending time FROM DUAL;
```



```
DBMS OUTPUT.PUT LINE('Time = ' || TO CHAR(ending time - starting time));
END:
```

#### Result is similar to:

```
Square root of dept. ID = 3.16227766016837933199889354443271853372
Square root of dept. ID = 4.47213595499957939281834733746255247088
Square root of dept. ID = 5.47722557505166113456969782800802133953
Square root of dept. ID = 6.32455532033675866399778708886543706744
Square root of dept. ID = 7.07106781186547524400844362104849039285
Square root of dept. ID = 7.74596669241483377035853079956479922167
Square root of dept. ID = 8.36660026534075547978172025785187489393
Square root of dept. ID = 8.94427190999915878563669467492510494176
Square root of dept. ID = 9.48683298050513799599668063329815560116
Square root of dept. ID = 10
Square root of dept. ID = 10.48808848170151546991453513679937598475
Time = +000000000 00:00:00.046000000
Square root of dept. ID = 3.16227766016837933199889354443271853372
Square root of dept. ID = 4.47213595499957939281834733746255247088
Square root of dept. ID = 5.47722557505166113456969782800802133953
Square root of dept. ID = 6.32455532033675866399778708886543706744
Square root of dept. ID = 7.07106781186547524400844362104849039285
Square root of dept. ID = 7.74596669241483377035853079956479922167
Square root of dept. ID = 8.36660026534075547978172025785187489393
Square root of dept. ID = 8.94427190999915878563669467492510494176
Square root of dept. ID = 9.48683298050513799599668063329815560116
Square root of dept. ID = 10
Square root of dept. ID = 10.48808848170151546991453513679937598475
Time = +0000000000 00:00:00.000000000
```

# Tune Subprogram Invocations

If a subprogram has OUT or IN OUT parameters, you can sometimes decrease its invocation overhead by declaring those parameters with the NOCOPY hint.

When OUT or IN OUT parameters represent large data structures such as collections, records, and instances of ADTs, copying them slows execution and increases memory use—especially for an instance of an ADT.

For each invocation of an ADT method, PL/SOL copies every attribute of the ADT. If the method is exited normally, then PL/SQL applies any changes that the method made to the attributes. If the method is exited with an unhandled exception, then PL/SQL does not change the attributes.

If your program does not require that an OUT or IN OUT parameter retain its pre-invocation value if the subprogram ends with an unhandled exception, then include the NOCOPY hint in the parameter declaration. The NOCOPY hint requests (but does not ensure) that the compiler pass the corresponding actual parameter by reference instead of value.



## Caution:

Do not rely on NOCOPY (which the compiler might or might not obey for a particular invocation) to ensure that an actual parameter or ADT attribute retains its preinvocation value if the subprogram is exited with an unhandled exception. Instead, ensure that the subprogram handle all exceptions.

# See Also:

- "NOCOPY" for more information about NOCOPY hint
- Oracle Database Object-Relational Developer's Guide for information about using NOCOPY with member methods of ADTs

### **Example 13-6 NOCOPY Subprogram Parameters**

In this example, if the compiler obeys the NOCOPY hint for the invocation of do\_nothing2, then the invocation of do nothing2 is faster than the invocation of do nothing1.

```
DECLARE
 TYPE EmpTabTyp IS TABLE OF employees%ROWTYPE;
 emp tab EmpTabTyp := EmpTabTyp(NULL); -- initialize
 t1 NUMBER;
 t2 NUMBER;
 t3 NUMBER;
 PROCEDURE get time (t OUT NUMBER) IS
   t := DBMS_UTILITY.get_time;
 END:
 PROCEDURE do_nothing1 (tab IN OUT EmpTabTyp) IS
 BEGIN
   NULL;
 END;
 PROCEDURE do nothing2 (tab IN OUT NOCOPY EmpTabTyp) IS
 BEGIN
   NULL;
 END;
BEGIN
 SELECT * INTO emp tab(1)
 FROM employees
 WHERE employee_id = 100;
 emp tab.EXTEND(49999, 1); -- Copy element 1 into 2..50000
 get time(t1);
 do_nothing1(emp_tab); -- Pass IN OUT parameter
 get time(t2);
 do_nothing2(emp_tab); -- Pass IN OUT NOCOPY parameter
 get time(t3);
 DBMS OUTPUT.PUT LINE ('Call Duration (secs)');
 DBMS OUTPUT.PUT LINE ('----');
 DBMS OUTPUT.PUT LINE ('Just IN OUT: ' || TO CHAR((t2 - t1)/100.0));
 DBMS OUTPUT.PUT LINE ('With NOCOPY: ' || TO CHAR((t3 - t2))/100.0);
END;
```

# Tune Loops

Because PL/SQL applications are often built around loops, it is important to optimize both the loops themselves and the code inside them.

If you must loop through a result set more than once, or issue other queries as you loop through a result set, you might be able to change the original query to give you exactly the results you want. Explore the SQL set operators that let you combine multiple queries, described in *Oracle Database SQL Language Reference*.

You can also use subqueries to do the filtering and sorting in multiple stages—see "Processing Query Result Sets with Subqueries".



"Bulk SQL and Bulk Binding"

# Tune Computation-Intensive PL/SQL Code

These recommendations apply especially (but not only) to computation-intensive PL/SQL code.

## **Topics**

- Use Data Types that Use Hardware Arithmetic
- Avoid Constrained Subtypes in Performance-Critical Code
- Minimize Implicit Data Type Conversion

# Use Data Types that Use Hardware Arithmetic

Avoid using data types in the NUMBER data type family (described in "NUMBER Data Type Family"). These data types are represented internally in a format designed for portability and arbitrary scale and precision, not for performance. Operations on data of these types use library arithmetic, while operations on data of the types PLS\_INTEGER, BINARY\_FLOAT and BINARY DOUBLE use hardware arithmetic.

For local integer variables, use PLS\_INTEGER, described in "PLS\_INTEGER and BINARY\_INTEGER Data Types". For variables used in performance-critical code, that can never have the value NULL, and do not need overflow checking, use SIMPLE\_INTEGER, described in "SIMPLE INTEGER Subtype of PLS INTEGER".

For floating-point variables, use BINARY\_FLOAT or BINARY\_DOUBLE, described in *Oracle Database SQL Language Reference*. For variables used in performance-critical code, that can never have the value NULL, and that do not need overflow checking, use SIMPLE\_FLOAT or SIMPLE\_DOUBLE, explained in "Additional PL/SQL Subtypes of BINARY\_FLOAT and BINARY\_DOUBLE".



BINARY\_FLOAT and BINARY\_DOUBLE and their subtypes are less suitable for financial code where accuracy is critical, because they do not always represent fractional values precisely, and handle rounding differently than the NUMBER types.

Many SQL numeric functions (described in *Oracle Database SQL Language Reference*) are overloaded with versions that accept BINARY\_FLOAT and BINARY\_DOUBLE parameters. You can

speed up computation-intensive code by passing variables of these data types to such functions, and by invoking the conversion functions <code>TO\_BINARY\_FLOAT</code> (described in *Oracle Database SQL Language Reference*) and <code>TO\_BINARY\_DOUBLE</code> (described in *Oracle Database SQL Language Reference*) when passing expressions to such functions.

# Avoid Constrained Subtypes in Performance-Critical Code

In performance-critical code, avoid constrained subtypes (described in "Constrained Subtypes"). Each assignment to a variable or parameter of a constrained subtype requires extra checking at run time to ensure that the value to be assigned does not violate the constraint.



PL/SQL Predefined Data Types includes predefined constrained subtypes

# Minimize Implicit Data Type Conversion

At run time, PL/SQL converts between different data types implicitly (automatically) if necessary. For example, if you assign a PLS\_INTEGER variable to a NUMBER variable, then PL/SQL converts the PLS\_INTEGER value to a NUMBER value (because the internal representations of the values differ).

Whenever possible, minimize implicit conversions. For example:

• If a variable is to be either inserted into a table column or assigned a value from a table column, then give the variable the same data type as the table column.



## Tip:

Declare the variable with the %TYPE attribute, described in "%TYPE Attribute".

- Make each literal the same data type as the variable to which it is assigned or the expression in which it appears.
- Convert values from SQL data types to PL/SQL data types and then use the converted values in expressions.

For example, convert NUMBER values to PLS\_INTEGER values and then use the PLS\_INTEGER values in expressions. PLS\_INTEGER operations use hardware arithmetic, so they are faster than NUMBER operations, which use library arithmetic. For more information about the PLS\_INTEGER data type, see "PLS\_INTEGER and BINARY\_INTEGER Data Types".

- Before assigning a value of one SQL data type to a variable of another SQL data type, explicitly convert the source value to the target data type, using a SQL conversion function (for information about SQL conversion functions, see *Oracle Database SQL Language Reference*).
- Overload your subprograms with versions that accept parameters of different data types and optimize each version for its parameter types. For information about overloaded subprograms, see "Overloaded Subprograms".

# See Also:

- Oracle Database SQL Language Reference for information about implicit conversion of SQL data types (which are also PL/SQL data types)
- "Subtypes with Base Types in Same Data Type Family"

# Use SQL Character Functions

SQL has many highly optimized character functions, which use low-level code that is more efficient than PL/SQL code. Use these functions instead of writing PL/SQL code to do the same things.

# See:

- Oracle Database SQL Language Reference for information about SQL character functions that return character values
- Oracle Database SQL Language Reference for information about SQL character functions that return NLS character values
- Oracle Database SQL Language Reference for information about SQL character functions that return number values
- Example 7-6 for an example of PL/SQL code that uses SQL character function REGEXP LIKE

# Put Least Expensive Conditional Tests First

PL/SQL stops evaluating a logical expression as soon as it can determine the result. Take advantage of this short-circuit evaluation by putting the conditions that are least expensive to evaluate first in logical expressions whenever possible. For example, test the values of PL/SQL variables before testing function return values, so that if the variable tests fail, PL/SQL need not invoke the functions:

IF boolean\_variable OR (number > 10) OR boolean\_function(parameter) THEN ...



"Short-Circuit Evaluation"

# **Bulk SQL and Bulk Binding**

**Bulk SQL** minimizes the performance overhead of the communication between PL/SQL and SQL. The PL/SQL features that comprise bulk SQL are the FORALL statement and the BULK COLLECT clause. Assigning values to PL/SQL variables that appear in SQL statements is called **binding**.



PL/SQL and SQL communicate as follows: To run a SELECT INTO or DML statement, the PL/SQL engine sends the query or DML statement to the SQL engine. The SQL engine runs the query or DML statement and returns the result to the PL/SQL engine.

The FORALL statement sends DML statements from PL/SQL to SQL in batches rather than one at a time. The BULK COLLECT clause returns results from SQL to PL/SQL in batches rather than one at a time. If a query or DML statement affects four or more database rows, then bulk SQL can significantly improve performance.



You cannot perform bulk SQL on remote tables.

PL/SQL binding operations fall into these categories:

Binding Category	When This Binding Occurs
In-bind	When an INSERT, UPDATE, or MERGE statement stores a PL/SQL or host variable in the database
Out-bind	When the RETURNING INTO clause of an INSERT, UPDATE, MERGE, or DELETE statement assigns a database value to a PL/SQL or host variable
DEFINE	When a ${\tt SELECT}$ or ${\tt FETCH}$ statement assigns a database value to a PL/SQL or host variable

For in-binds and out-binds, bulk SQL uses **bulk binding**; that is, it binds an entire collection of values at once. For a collection of n elements, bulk SQL uses a single operation to perform the equivalent of n SELECT INTO or DML statements. A query that uses bulk SQL can return any number of rows, without using a FETCH statement for each one.



Parallel DML is disabled with bulk SQL.

### **Topics**

- FORALL Statement
- BULK COLLECT Clause
- Using FORALL Statement and BULK COLLECT Clause Together
- · Client Bulk-Binding of Host Arrays

# **FORALL Statement**

The FORALL statement, a feature of bulk SQL, sends DML statements from PL/SQL to SQL in batches rather than one at a time.

To understand the FORALL statement, first consider the FOR LOOP statement in Example 13-7. It sends these DML statements from PL/SQL to SQL one at a time:

```
DELETE FROM employees_temp WHERE department_id = 10;
DELETE FROM employees_temp WHERE department_id = 30;
DELETE FROM employees_temp WHERE department_id = 70;
```

Now consider the FORALL statement in Example 13-8. It sends the same three DML statements from PL/SQL to SQL as a batch.

A FORALL statement is usually much faster than an equivalent FOR LOOP statement. However, a FOR LOOP statement can contain multiple DML statements, while a FORALL statement can contain only one. The batch of DML statements that a FORALL statement sends to SQL differ only in their VALUES and WHERE clauses. The values in those clauses must come from existing, populated collections.

## Note:

The DML statement in a FORALL statement can reference multiple collections, but performance benefits apply only to collection references that use the FORALL index variable as an index.

Example 13-9 inserts the same collection elements into two database tables, using a FOR LOOP statement for the first table and a FORALL statement for the second table and showing how long each statement takes. (Times vary from run to run.)

In Example 13-10, the FORALL statement applies to a subset of a collection.

#### **Topics**

- Using FORALL Statements for Sparse Collections
- Unhandled Exceptions in FORALL Statements
- · Handling FORALL Exceptions Immediately
- Handling FORALL Exceptions After FORALL Statement Completes
- Getting Number of Rows Affected by FORALL Statement

# See Also:

- "FORALL Statement" for its complete syntax and semantics, including restrictions
- "Implicit Cursors" for information about implicit cursor attributes in general and other implicit cursor attributes that you can use with the FORALL statement

#### Example 13-7 DELETE Statement in FOR LOOP Statement

```
DROP TABLE employees_temp;

CREATE TABLE employees_temp AS SELECT * FROM employees;

DECLARE

TYPE NumList IS VARRAY(20) OF NUMBER;

depts NumList := NumList(10, 30, 70); -- department numbers

BEGIN

FOR i IN depts.FIRST..depts.LAST LOOP
```



```
DELETE FROM employees_temp
WHERE department_id = depts(i);
END LOOP;
END;
//
```

### Example 13-8 DELETE Statement in FORALL Statement

```
DROP TABLE employees_temp;
CREATE TABLE employees_temp AS SELECT * FROM employees;

DECLARE
    TYPE NumList IS VARRAY(20) OF NUMBER;
    depts NumList := NumList(10, 30, 70); -- department numbers
BEGIN
    FORALL i IN depts.FIRST..depts.LAST
        DELETE FROM employees_temp
    WHERE department_id = depts(i);
END;
//
```

# Example 13-9 Time Difference for INSERT Statement in FOR LOOP and FORALL Statements

```
DROP TABLE parts1;
CREATE TABLE parts1 (
 pnum INTEGER,
 pname VARCHAR2 (15)
);
DROP TABLE parts2;
CREATE TABLE parts2 (
  pnum INTEGER,
  pname VARCHAR2(15)
);
DECLARE
  TYPE NumTab IS TABLE OF parts1.pnum%TYPE INDEX BY PLS INTEGER;
  TYPE NameTab IS TABLE OF parts1.pname%TYPE INDEX BY PLS INTEGER;
  pnums NumTab;
  pnames NameTab;
  iterations CONSTANT PLS INTEGER := 50000;
  t1 INTEGER;
  t2 INTEGER;
  t3 INTEGER;
BEGIN
  FOR j IN 1..iterations LOOP -- populate collections
   pnums(j) := j;
   pnames(j) := 'Part No. ' || TO CHAR(j);
  END LOOP;
  t1 := DBMS_UTILITY.get_time;
  FOR i IN 1..iterations LOOP
    INSERT INTO parts1 (pnum, pname)
    VALUES (pnums(i), pnames(i));
  END LOOP;
  t2 := DBMS UTILITY.get time;
  FORALL i IN 1..iterations
    INSERT INTO parts2 (pnum, pname)
    VALUES (pnums(i), pnames(i));
```

```
t3 := DBMS_UTILITY.get_time;

DBMS_OUTPUT.PUT_LINE('Execution Time (secs)');

DBMS_OUTPUT.PUT_LINE('------');

DBMS_OUTPUT.PUT_LINE('FOR LOOP: ' || TO_CHAR((t2 - t1)/100));

DBMS_OUTPUT.PUT_LINE('FORALL: ' || TO_CHAR((t3 - t2)/100));

COMMIT;

END;
```

### Result is similar to:

```
Execution Time (secs)
-----
FOR LOOP: 5.97
FORALL: .07
```

PL/SQL procedure successfully completed.

### Example 13-10 FORALL Statement for Subset of Collection

```
DROP TABLE employees_temp;
CREATE TABLE employees_temp AS SELECT * FROM employees;

DECLARE
   TYPE NumList IS VARRAY(10) OF NUMBER;
   depts NumList := NumList(5,10,20,30,50,55,57,60,70,75);
BEGIN
   FORALL j IN 4..7
        DELETE FROM employees_temp WHERE department_id = depts(j);
END;
//
```

# Using FORALL Statements for Sparse Collections

If the FORALL statement bounds clause references a sparse collection, then specify only existing index values, using either the INDICES OF or VALUES OF clause.

You can use INDICES OF for any collection except an associative array indexed by string. You can use VALUES OF only for a collection of PLS INTEGER elements indexed by PLS INTEGER.

A collection of PLS\_INTEGER elements indexed by PLS\_INTEGER can be an **index collection**; that is, a collection of pointers to elements of another collection (the **indexed collection**).

Index collections are useful for processing different subsets of the same collection with different FORALL statements. Instead of copying elements of the original collection into new collections that represent the subsets (which can use significant time and memory), represent each subset with an index collection and then use each index collection in the VALUES OF clause of a different FORALL statement.

```
See Also:
```

"Sparse Collections and SQL%BULK\_EXCEPTIONS"



### Example 13-11 FORALL Statements for Sparse Collection and Its Subsets

This example uses a FORALL statement with the INDICES OF clause to populate a table with the elements of a sparse collection. Then it uses two FORALL statements with VALUES OF clauses to populate two tables with subsets of a collection.

```
DROP TABLE valid orders;
CREATE TABLE valid orders (
 cust_name VARCHAR2(32),
 amount NUMBER (10,2)
);
DROP TABLE big orders;
CREATE TABLE big orders AS
 SELECT * FROM valid orders
 WHERE 1 = 0;
DROP TABLE rejected orders;
CREATE TABLE rejected orders AS
 SELECT * FROM valid orders
 WHERE 1 = 0;
DECLARE
  SUBTYPE cust name IS valid orders.cust name%TYPE;
 TYPE cust typ IS TABLE OF cust name;
  cust_tab cust_typ; -- Collection of customer names
  SUBTYPE order amount IS valid orders.amount%TYPE;
  TYPE amount typ IS TABLE OF NUMBER;
  amount_tab amount_typ; -- Collection of order amounts
  TYPE index pointer t IS TABLE OF PLS INTEGER;
  /* Collections for pointers to elements of cust tab collection
     (to represent two subsets of cust_tab): */
 big order tab
                      index pointer t := index pointer t();
  rejected order tab index pointer t := index pointer t();
  PROCEDURE populate data collections IS
  BEGIN
    cust tab := cust typ(
      'Company1', 'Company2', 'Company3', 'Company4', 'Company5'
    amount tab := amount typ(5000.01, 0, 150.25, 4000.00, NULL);
  END:
BEGIN
 populate_data_collections;
  DBMS OUTPUT.PUT LINE ('--- Original order data ---');
  FOR i IN 1..cust tab.LAST LOOP
   DBMS OUTPUT.PUT_LINE (
      'Customer #' || i || ', ' || cust tab(i) || ': $' || amount tab(i)
   );
 END LOOP;
  -- Delete invalid orders:
  FOR i IN 1..cust tab.LAST LOOP
```



```
IF amount tab(i) IS NULL OR amount tab(i) = 0 THEN
     cust tab.delete(i);
     amount_tab.delete(i);
   END IF;
  END LOOP;
  -- cust_tab is now a sparse collection.
  DBMS OUTPUT.PUT LINE ('--- Order data with invalid orders deleted ---');
  FOR i IN 1..cust_tab.LAST LOOP
   IF cust tab.EXISTS(i) THEN
     DBMS_OUTPUT.PUT_LINE (
       'Customer #' || i || ', ' || cust tab(i) || ': $' || amount tab(i)
     );
   END IF;
 END LOOP;
  -- Using sparse collection, populate valid orders table:
  FORALL i IN INDICES OF cust tab
   INSERT INTO valid orders (cust name, amount)
   VALUES (cust tab(i), amount tab(i));
 populate_data_collections; -- Restore original order data
  -- cust tab is a dense collection again.
  /* Populate collections of pointers to elements of cust tab collection
     (which represent two subsets of cust tab): */
  FOR i IN cust_tab.FIRST .. cust_tab.LAST LOOP
    IF amount_tab(i) IS NULL OR amount_tab(i) = 0 THEN
     rejected_order_tab.EXTEND;
      rejected_order_tab(rejected_order_tab.LAST) := i;
    END IF;
    IF amount tab(i) > 2000 THEN
     big order tab.EXTEND;
     big order tab(big order tab.LAST) := i;
  END LOOP;
  /* Using each subset in a different FORALL statement,
     populate rejected_orders and big_orders tables: */
 FORALL i IN VALUES OF rejected order tab
    INSERT INTO rejected orders (cust name, amount)
    VALUES (cust tab(i), amount tab(i));
 FORALL i IN VALUES OF big order tab
    INSERT INTO big orders (cust name, amount)
    VALUES (cust tab(i), amount tab(i));
END;
Result:
--- Original order data ---
Customer #1, Company1: $5000.01
Customer #2, Company2: $0
Customer #3, Company3: $150.25
```

```
Customer #4, Company4: $4000

Customer #5, Company5: $
--- Data with invalid orders deleted ---
Customer #1, Company1: $5000.01

Customer #3, Company3: $150.25

Customer #4, Company4: $4000
```

## Verify that correct order details were stored:

```
SELECT cust_name "Customer", amount "Valid order amount"
FROM valid_orders
ORDER BY cust name;
```

#### Result:

Customer	Valid	order	amount
Company1 Company3 Company4			5000.01 150.25 4000

3 rows selected.

### Query:

```
SELECT cust_name "Customer", amount "Big order amount"
FROM big_orders
ORDER BY cust_name;
```

#### Result:

Customer	Big	order	amount
Company1			5000.01
Company4			4000

2 rows selected.

#### Query:

```
SELECT cust_name "Customer", amount "Rejected order amount"
FROM rejected_orders
ORDER BY cust_name;
```

### Result:

Customer	Rejected	order	amount
Company2			0
Company5			

2 rows selected.

# Unhandled Exceptions in FORALL Statements

In a FORALL statement without the SAVE EXCEPTIONS clause, if one DML statement raises an unhandled exception, then PL/SQL stops the FORALL statement and rolls back all changes made by previous DML statements.

For example, the FORALL statement in Example 13-8 processes these DML statements in this order, unless one of them raises an unhandled exception:

```
DELETE FROM employees_temp WHERE department_id = depts(10);

DELETE FROM employees_temp WHERE department_id = depts(30);

DELETE FROM employees_temp WHERE department_id = depts(70);
```

If the third statement raises an unhandled exception, then PL/SQL rolls back the changes that the first and second statements made. If the second statement raises an unhandled exception, then PL/SQL rolls back the changes that the first statement made and never runs the third statement.

You can handle exceptions raised in a FORALL statement in either of these ways:

- As each exception is raised (see "Handling FORALL Exceptions Immediately")
- After the FORALL statement completes execution, by including the SAVE EXCEPTIONS clause (see "Handling FORALL Exceptions After FORALL Statement Completes")

# Handling FORALL Exceptions Immediately

To handle exceptions raised in a FORALL statement immediately, omit the SAVE EXCEPTIONS clause and write the appropriate exception handlers.

If one DML statement raises a handled exception, then PL/SQL rolls back the changes made by that statement, but does not roll back changes made by previous DML statements.

In Example 13-12, the FORALL statement is designed to run three UPDATE statements. However, the second one raises an exception. An exception handler handles the exception, displaying the error message and committing the change made by the first UPDATE statement. The third UPDATE statement never runs.

For information about exception handlers, see PL/SQL Error Handling.

# Example 13-12 Handling FORALL Exceptions Immediately

```
DROP TABLE emp_temp;
CREATE TABLE emp_temp (
 deptno NUMBER (2),
 job VARCHAR2 (18)
);
CREATE OR REPLACE PROCEDURE p AUTHID DEFINER AS
 TYPE NumList IS TABLE OF NUMBER;
                NumList := NumList(10, 20, 30);
 error_message VARCHAR2(100);
BEGIN
 -- Populate table:
 INSERT INTO emp temp (deptno, job) VALUES (10, 'Clerk');
 INSERT INTO emp_temp (deptno, job) VALUES (20, 'Bookkeeper');
 INSERT INTO emp temp (deptno, job) VALUES (30, 'Analyst');
 COMMIT:
 -- Append 9-character string to each job:
 FORALL j IN depts.FIRST..depts.LAST
   UPDATE emp temp SET job = job || ' (Senior)'
   WHERE deptno = depts(j);
EXCEPTION
 WHEN OTHERS THEN
   error message := SQLERRM;
```

```
DBMS OUTPUT.PUT LINE (error message);
    COMMIT; -- Commit results of successful updates
    RAISE;
END;
Result:
Procedure created.
Invoke procedure:
BEGIN
  p;
END:
Result:
ORA-12899: value too large for column "HR". "EMP TEMP". "JOB" (actual: 19,
maximum: 18)
BEGIN
ERROR at line 1:
ORA-12899: value too large for column "HR". "EMP TEMP". "JOB" (actual: 19,
maximum: 18)
ORA-06512: at "HR.P", line 27
ORA-06512: at line 2
Query:
SELECT * FROM emp temp;
Result:
   DEPTNO JOB
_____
       10 Clerk (Senior)
        20 Bookkeeper
        30 Analyst
3 rows selected.
```

# Handling FORALL Exceptions After FORALL Statement Completes

To allow a FORALL statement to continue even if some of its DML statements fail, include the SAVE EXCEPTIONS clause. When a DML statement fails, PL/SQL does not raise an exception; instead, it saves information about the failure. After the FORALL statement completes, PL/SQL raises a single exception for the FORALL statement (ORA-24381).

In the exception handler for ORA-24381, you can get information about each individual DML statement failure from the implicit cursor attribute SQL%BULK EXCEPTIONS.

SQL%BULK\_EXCEPTIONS is like an associative array of information about the DML statements that failed during the most recently run FORALL statement.

SQL%BULK\_EXCEPTIONS.COUNT is the number of DML statements that failed. If SQL%BULK\_EXCEPTIONS.COUNT is not zero, then for each index value i from 1 through SQL%BULK\_EXCEPTIONS.COUNT:

- SQL%BULK EXCEPTIONS (i).ERROR INDEX is the number of the DML statement that failed.
- SQL%BULK\_EXCEPTIONS (i).ERROR\_CODE is the Oracle Database error code for the failure.

For example, if a FORALL SAVE EXCEPTIONS statement runs 100 DML statements, and the tenth and sixty-fourth ones fail with error codes ORA-12899 and ORA-19278, respectively, then:

- SQL%BULK\_EXCEPTIONS.COUNT = 2
- SQL%BULK EXCEPTIONS (1).ERROR INDEX = 10
- SQL%BULK EXCEPTIONS(1).ERROR CODE = 12899
- SQL%BULK EXCEPTIONS(2).ERROR INDEX = 64
- SQL%BULK EXCEPTIONS(2).ERROR CODE = 19278

# Note:

After a FORALL statement without the SAVE EXCEPTIONS clause raises an exception, SQL%BULK\_EXCEPTIONS.COUNT = 1.

With the error code, you can get the associated error message with the SQLERRM function (described in "SQLERRM Function"):

```
SQLERRM(-(SQL%BULK_EXCEPTIONS(i).ERROR_CODE))
```

However, the error message that SQLERRM returns excludes any substitution arguments (compare the error messages in Example 13-12 and Example 13-13).

Example 13-13 is like Example 13-12 except:

- The FORALL statement includes the SAVE EXCEPTIONS clause.
- The exception-handling part has an exception handler for ORA-24381, the internally
  defined exception that PL/SQL raises implicitly when a bulk operation raises and saves
  exceptions. The example gives ORA-24381 the user-defined name dml errors.
- The exception handler for dml\_errors uses SQL%BULK\_EXCEPTIONS and SQLERRM (and some local variables) to show the error message and which statement, collection item, and string caused the error.

### Example 13-13 Handling FORALL Exceptions After FORALL Statement Completes

```
CREATE OR REPLACE PROCEDURE p AUTHID DEFINER AS

TYPE NumList IS TABLE OF NUMBER;
depts NumList := NumList(10, 20, 30);

error_message VARCHAR2(100);
bad_stmt_no PLS_INTEGER;
bad_deptno emp_temp.deptno%TYPE;
bad_job emp_temp.job%TYPE;

dml_errors EXCEPTION;
PRAGMA EXCEPTION_INIT(dml_errors, -24381);

BEGIN
-- Populate table:

INSERT INTO emp_temp (deptno, job) VALUES (10, 'Clerk');
INSERT INTO emp_temp (deptno, job) VALUES (20, 'Bookkeeper');
INSERT INTO emp_temp (deptno, job) VALUES (30, 'Analyst');
```



```
COMMIT;
  -- Append 9-character string to each job:
  FORALL j IN depts.FIRST..depts.LAST SAVE EXCEPTIONS
    UPDATE emp temp SET job = job || ' (Senior)'
    WHERE deptno = depts(j);
EXCEPTION
  WHEN dml errors THEN
    FOR i IN 1..SQL%BULK EXCEPTIONS.COUNT LOOP
      error message := SQLERRM(-(SQL%BULK EXCEPTIONS(i).ERROR CODE));
      DBMS_OUTPUT.PUT_LINE (error_message);
      bad stmt no := SQL%BULK EXCEPTIONS(i).ERROR INDEX;
      DBMS_OUTPUT.PUT_LINE('Bad statement #: ' || bad stmt no);
      bad deptno := depts(bad stmt no);
      DBMS OUTPUT.PUT LINE('Bad department #: ' || bad deptno);
      SELECT job INTO bad job FROM emp temp WHERE deptno = bad deptno;
      DBMS OUTPUT.PUT LINE('Bad job: ' || bad job);
    END LOOP;
    COMMIT; -- Commit results of successful updates
    WHEN OTHERS THEN
      DBMS OUTPUT.PUT LINE('Unrecognized error.');
      RAISE;
END;
Result:
Procedure created.
Invoke procedure:
BEGIN
  p;
END;
Result:
ORA-12899: value too large for column (actual: , maximum: )
Bad statement #: 2
Bad department #: 20
Bad job: Bookkeeper
PL/SQL procedure successfully completed.
Query:
SELECT * FROM emp temp;
Result:
    DEPTNO JOB
       10 Clerk (Senior)
```

```
20 Bookkeeper
30 Analyst (Senior)
3 rows selected.
```

# Sparse Collections and SQL%BULK\_EXCEPTIONS

If the FORALL statement bounds clause references a sparse collection, then to find the collection element that caused a DML statement to fail, you must step through the elements one by one until you find the element whose index is  $SQL\BULK\_EXCEPTIONS(i).ERROR\_INDEX$ . Then, if the FORALL statement uses the VALUES OF clause to reference a collection of pointers into another collection, you must find the element of the other collection whose index is  $SQL\BULK\_EXCEPTIONS(i).ERROR\_INDEX$ .

# Getting Number of Rows Affected by FORALL Statement

After a FORALL statement completes, you can get the number of rows that each DML statement affected from the implicit cursor attribute  $SQL\BULK$  ROWCOUNT.

To get the total number of rows affected by the FORALL statement, use the implicit cursor attribute SQL%ROWCOUNT, described in "SQL%ROWCOUNT Attribute: How Many Rows Were Affected?".

 $SQL\BULK_ROWCOUNT$  is like an associative array whose *i*th element is the number of rows affected by the *i*th DML statement in the most recently completed FORALL statement. The data type of the element is INTEGER.



If a server is Oracle Database 12c or later and its client is Oracle Database 11g release 2 or earlier (or the reverse), then the maximum number that SQL%BULK ROWCOUNT returns is 4,294,967,295.

Example 13-14 uses SQL%BULK\_ROWCOUNT to show how many rows each DELETE statement in the FORALL statement deleted and SOL%ROWCOUNT to show the total number of rows deleted.

Example 13-15 uses SQL%BULK\_ROWCOUNT to show how many rows each INSERT SELECT construct in the FORALL statement inserted and SQL%ROWCOUNT to show the total number of rows inserted.

### Example 13-14 Showing Number of Rows Affected by Each DELETE in FORALL



```
);
 END LOOP;
 DBMS OUTPUT.PUT LINE('Total rows deleted: ' || SQL%ROWCOUNT);
END;
Result:
Statement #1 deleted 6 rows.
Statement #2 deleted 45 rows.
Statement #3 deleted 5 rows.
Total rows deleted: 56
Example 13-15 Showing Number of Rows Affected by Each INSERT SELECT in
FORALL
DROP TABLE emp by dept;
CREATE TABLE emp by dept AS
  SELECT employee_id, department_id
 FROM employees
 WHERE 1 = 0;
DECLARE
 TYPE dept tab IS TABLE OF departments.department id%TYPE;
  deptnums dept tab;
BEGIN
 SELECT department id BULK COLLECT INTO deptnums FROM departments;
 FORALL i IN 1..deptnums.COUNT
    INSERT INTO emp by dept (employee id, department id)
      SELECT employee_id, department_id
     FROM employees
     WHERE department id = deptnums(i)
     ORDER BY department_id, employee_id;
 FOR i IN 1..deptnums.COUNT LOOP
    -- Count how many rows were inserted for each department; that is,
    -- how many employees are in each department.
    DBMS OUTPUT.PUT LINE (
      'Dept '||deptnums(i)||': inserted '||
     SQL%BULK ROWCOUNT(i)||' records'
   );
 END LOOP;
 DBMS OUTPUT.PUT LINE('Total records inserted: ' || SQL%ROWCOUNT);
END;
Result:
Dept 10: inserted 1 records
Dept 20: inserted 2 records
Dept 30: inserted 6 records
Dept 40: inserted 1 records
Dept 50: inserted 45 records
Dept 60: inserted 5 records
Dept 70: inserted 1 records
Dept 80: inserted 34 records
Dept 90: inserted 3 records
Dept 100: inserted 6 records
Dept 110: inserted 2 records
Dept 120: inserted 0 records
```

Dept 130: inserted 0 records

```
Dept 140: inserted 0 records
Dept 150: inserted 0 records
Dept 160: inserted 0 records
Dept 170: inserted 0 records
Dept 180: inserted 0 records
Dept 190: inserted 0 records
Dept 200: inserted 0 records
Dept 210: inserted 0 records
Dept 220: inserted 0 records
Dept 230: inserted 0 records
Dept 240: inserted 0 records
Dept 250: inserted 0 records
Dept 260: inserted 0 records
Dept 270: inserted 0 records
Dept 280: inserted 0 records
Total records inserted: 106
```

# **BULK COLLECT Clause**

The BULK COLLECT clause, a feature of bulk SQL, returns results from SQL to PL/SQL in batches rather than one at a time.

The BULK COLLECT clause can appear in:

- SELECT INTO statement
- FETCH statement
- RETURNING INTO clause of:
  - DELETE statement
  - INSERT statement
  - UPDATE statement
  - MERGE statement
  - EXECUTE IMMEDIATE statement

With the BULK COLLECT clause, each of the preceding statements retrieves an entire result set and stores it in one or more collection variables in a single operation (which is more efficient than using a loop statement to retrieve one result row at a time).



PL/SQL processes the BULK COLLECT clause similar to the way it processes a FETCH statement inside a LOOP statement. PL/SQL does not raise an exception when a statement with a BULK COLLECT clause returns no rows. You must check the target collections for emptiness, as in Example 13-22.

#### **Topics**

- SELECT INTO Statement with BULK COLLECT Clause
- FETCH Statement with BULK COLLECT Clause
- RETURNING INTO Clause with BULK COLLECT Clause



# SELECT INTO Statement with BULK COLLECT Clause

The SELECT INTO statement with the BULK COLLECT clause (also called the SELECT BULK COLLECT INTO statement) selects an entire result set into one or more collection variables.

For more information, see "SELECT INTO Statement".



### Caution:

The SELECT BULK COLLECT INTO statement is vulnerable to aliasing, which can cause unexpected results. For details, see "SELECT BULK COLLECT INTO Statements and Aliasing".

Example 13-16 uses a SELECT BULK COLLECT INTO statement to select two database columns into two collections (nested tables).

Example 13-17 uses a SELECT BULK COLLECT INTO statement to select a result set into a nested table of records.

### **Topics**

- SELECT BULK COLLECT INTO Statements and Aliasing
- Row Limits for SELECT BULK COLLECT INTO Statements
- **Guidelines for Looping Through Collections**

#### Example 13-16 Bulk-Selecting Two Database Columns into Two Nested Tables

```
DECLARE
 TYPE NumTab IS TABLE OF employees.employee id%TYPE;
 TYPE NameTab IS TABLE OF employees.last name%TYPE;
  enums NumTab;
 names NameTab;
 PROCEDURE print first n (n POSITIVE) IS
 BEGIN
    IF enums.COUNT = 0 THEN
     DBMS OUTPUT.PUT LINE ('Collections are empty.');
     DBMS OUTPUT.PUT LINE ('First ' || n || ' employees:');
     FOR i IN 1 .. n LOOP
        DBMS OUTPUT.PUT LINE (
         ' Employee #' || enums(i) || ': ' || names(i));
     END LOOP;
    END IF;
 END;
BEGIN
  SELECT employee_id, last_name
 BULK COLLECT INTO enums, names
 FROM employees
```

# ORDER BY employee id; print first n(3); print\_first\_n(6); END; Result: First 3 employees: Employee #100: King Employee #101: Yang Employee #102: Garcia First 6 employees: Employee #100: King Employee #101: Yang Employee #102: Garcia Employee #103: James Employee #104: Miller Employee #105: Williams

### Example 13-17 Bulk-Selecting into Nested Table of Records

```
DECLARE
  CURSOR c1 IS
    SELECT first name, last name, hire date
    FROM employees;
  TYPE NameSet IS TABLE OF c1%ROWTYPE;
  stock managers NameSet; -- nested table of records
BEGIN
  -- Assign values to nested table of records:
  SELECT first name, last name, hire date
   BULK COLLECT INTO stock_managers
    FROM employees
   WHERE job id = 'ST MAN'
    ORDER BY hire date;
  -- Print nested table of records:
    FOR i IN stock managers.FIRST .. stock managers.LAST LOOP
      DBMS OUTPUT.PUT LINE (
        stock managers(i).hire date || ' ' ||
        stock managers(i).last name || ', ' ||
        stock managers(i).first name
     );
    END LOOP; END;
```

#### Result:

```
01-MAY-13 Kaufling, Payam
18-JUL-14 Weiss, Matthew
10-APR-15 Fripp, Adam
10-OCT-15 Vollman, Shanta
16-NOV-17 Mourgos, Kevin
```

# SELECT BULK COLLECT INTO Statements and Aliasing

#### In a statement of the form

```
SELECT column BULK COLLECT INTO collection FROM table ...
```

column and collection are analogous to IN NOCOPY and OUT NOCOPY subprogram parameters, respectively, and PL/SQL passes them by reference. As with subprogram parameters that are passed by reference, aliasing can cause unexpected results.



"Subprogram Parameter Aliasing with Parameters Passed by Reference"

In Example 13-18, the intention is to select specific values from a collection, numbers1, and then store them in the same collection. The unexpected result is that all elements of numbers1 are deleted. For workarounds, see Example 13-19 and Example 13-20.

Example 13-19 uses a cursor to achieve the result intended by Example 13-18.

Example 13-20 selects specific values from a collection, numbers1, and then stores them in a different collection, numbers2. Example 13-20 runs faster than Example 13-19.

### Example 13-18 SELECT BULK COLLECT INTO Statement with Unexpected Results

```
CREATE OR REPLACE TYPE numbers type IS
 TABLE OF INTEGER
CREATE OR REPLACE PROCEDURE p (i IN INTEGER) AUTHID DEFINER IS
 numbers1 numbers type := numbers type (1, 2, 3, 4, 5);
BEGIN
  DBMS OUTPUT.PUT LINE('Before SELECT statement');
 DBMS OUTPUT.PUT LINE('numbers1.COUNT() = ' || numbers1.COUNT());
 FOR j IN 1..numbers1.COUNT() LOOP
   DBMS OUTPUT.PUT LINE('numbers1(' || j || ') = ' || numbers1(j));
 END LOOP;
 --Self-selecting BULK COLLECT INTO clause:
 SELECT a.COLUMN VALUE
 BULK COLLECT INTO numbers1
 FROM TABLE (numbers1) a
 WHERE a.COLUMN VALUE > p.i
 ORDER BY a.COLUMN VALUE;
  DBMS OUTPUT.PUT LINE('After SELECT statement');
  DBMS OUTPUT.PUT LINE('numbers1.COUNT() = ' || numbers1.COUNT());
```



```
END p;
Invoke p:
BEGIN
  p(2);
END;
Result:
Before SELECT statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
After SELECT statement
numbers1.COUNT() = 0
PL/SQL procedure successfully completed.
Invoke p:
BEGIN
  p(10);
END;
Result:
Before SELECT statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
After SELECT statement
numbers1.COUNT() = 0
Example 13-19 Cursor Workaround for Example 13-18
CREATE OR REPLACE TYPE numbers type IS
  TABLE OF INTEGER
CREATE OR REPLACE PROCEDURE p (i IN INTEGER) AUTHID DEFINER IS
  numbers1 numbers_type := numbers_type(1,2,3,4,5);
  CURSOR c IS
    SELECT a.COLUMN VALUE
    FROM TABLE (numbers1) a
    WHERE a.COLUMN_VALUE > p.i
    ORDER BY a.COLUMN VALUE;
  BEGIN
    DBMS OUTPUT.PUT LINE('Before FETCH statement');
    DBMS_OUTPUT.PUT_LINE('numbers1.COUNT() = ' || numbers1.COUNT());
    FOR j IN 1..numbers1.COUNT() LOOP
      DBMS_OUTPUT.PUT_LINE('numbers1(' || j || ') = ' || numbers1(j));
    END LOOP;
```

```
OPEN c;
  FETCH c BULK COLLECT INTO numbers1;
  CLOSE c;
  DBMS OUTPUT.PUT LINE('After FETCH statement');
  DBMS OUTPUT.PUT LINE('numbers1.COUNT() = ' || numbers1.COUNT());
  IF numbers1.COUNT() > 0 THEN
    FOR j IN 1..numbers1.COUNT() LOOP
      DBMS_OUTPUT.PUT_LINE('numbers1(' || j || ') = ' || numbers1(j));
    END LOOP;
  END IF;
END p;
Invoke p:
BEGIN
 p(2);
END;
Result:
Before FETCH statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
After FETCH statement
numbers1.COUNT() = 3
numbers1(1) = 3
numbers1(2) = 4
numbers1(3) = 5
Invoke p:
BEGIN
  p(10);
END;
Result:
Before FETCH statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
After FETCH statement
numbers1.COUNT() = 0
Example 13-20 Second Collection Workaround for Example 13-18
CREATE OR REPLACE TYPE numbers type IS
  TABLE OF INTEGER
CREATE OR REPLACE PROCEDURE p (i IN INTEGER) AUTHID DEFINER IS
```

```
numbers1 numbers type := numbers type (1,2,3,4,5);
 numbers2 numbers type := numbers type(0,0,0,0,0);
BEGIN
  DBMS_OUTPUT.PUT_LINE('Before SELECT statement');
  DBMS OUTPUT.PUT LINE('numbers1.COUNT() = ' || numbers1.COUNT());
  FOR j IN 1..numbers1.COUNT() LOOP
   DBMS_OUTPUT.PUT_LINE('numbers1(' || j || ') = ' || numbers1(j));
  END LOOP;
  DBMS_OUTPUT.PUT_LINE('numbers2.COUNT() = ' || numbers2.COUNT());
  FOR j IN 1..numbers2.COUNT() LOOP
   DBMS OUTPUT.PUT LINE('numbers2(' || j || ') = ' || numbers2(j));
 END LOOP;
  SELECT a.COLUMN VALUE
 BULK COLLECT INTO numbers2
                                  -- numbers2 appears here
 FROM TABLE (numbers1) a
                                -- numbers1 appears here
 WHERE a.COLUMN VALUE > p.i
 ORDER BY a.COLUMN_VALUE;
  DBMS OUTPUT.PUT LINE('After SELECT statement');
 DBMS OUTPUT.PUT LINE('numbers1.COUNT() = ' | numbers1.COUNT());
  IF numbers1.COUNT() > 0 THEN
    FOR j IN 1..numbers1.COUNT() LOOP
      DBMS_OUTPUT.PUT_LINE('numbers1(' \mid \mid j \mid \mid \mid ') = ' \mid \mid numbers1(j));
    END LOOP;
  END IF;
 DBMS_OUTPUT.PUT_LINE('numbers2.COUNT() = ' || numbers2.COUNT());
  IF numbers2.COUNT() > 0 THEN
    FOR j IN 1..numbers2.COUNT() LOOP
      DBMS OUTPUT.PUT LINE('numbers2(' || j || ') = ' || numbers2(j));
   END LOOP;
 END IF;
END p;
Invoke p:
BEGIN
 p(2);
END;
Result:
Before SELECT statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
numbers2.COUNT() = 5
numbers2(1) = 0
numbers2(2) = 0
numbers2(3) = 0
```

```
numbers2(4) = 0
numbers2(5) = 0
After SELECT statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
numbers2.COUNT() = 3
numbers2(1) = 3
numbers2(2) = 4
numbers2(3) = 5
PL/SQL procedure successfully completed.
Invoke p:
BEGIN
  p(10);
END;
Result:
Before SELECT statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
numbers2.COUNT() = 5
numbers2(1) = 0
numbers2(2) = 0
numbers2(3) = 0
numbers2(4) = 0
numbers2(5) = 0
After SELECT statement
numbers1.COUNT() = 5
numbers1(1) = 1
numbers1(2) = 2
numbers1(3) = 3
numbers1(4) = 4
numbers1(5) = 5
numbers2.COUNT() = 0
```

# Row Limits for SELECT BULK COLLECT INTO Statements

A SELECT BULK COLLECT INTO statement that returns a large number of rows produces a large collection. To limit the number of rows and the collection size, use one of these:

- ROWNUM pseudocolumn (described in Oracle Database SQL Language Reference)
- SAMPLE clause (described in Oracle Database SQL Language Reference)
- FETCH FIRST clause (described in Oracle Database SQL Language Reference)

**Example 13-21** shows several ways to limit the number of rows that a SELECT BULK COLLECT INTO statement returns.

### Example 13-21 Limiting Bulk Selection with ROWNUM, SAMPLE, and FETCH FIRST

```
DECLARE

TYPE SalList IS TABLE OF employees.salary%TYPE;
sals SalList;

BEGIN

SELECT salary BULK COLLECT INTO sals FROM employees
WHERE ROWNUM <= 50;

SELECT salary BULK COLLECT INTO sals FROM employees
SAMPLE (10);

SELECT salary BULK COLLECT INTO sals FROM employees
FETCH FIRST 50 ROWS ONLY;

END;
```

# **Guidelines for Looping Through Collections**

When a result set is stored in a collection, it is easy to loop through the rows and refer to different columns. This technique can be very fast, but also very memory-intensive. If you use it often:

 To loop once through the result set, use a cursor FOR LOOP (see "Processing Query Result Sets With Cursor FOR LOOP Statements").

This technique avoids the memory overhead of storing a copy of the result set.

- Instead of looping through the result set to search for certain values or filter the results into a smaller set, do the searching or filtering in the query of the SELECT INTO statement.
  - For example, in simple queries, use WHERE clauses; in queries that compare multiple result sets, use set operators such as INTERSECT and MINUS. For information about set operators, see *Oracle Database SQL Language Reference*.
- Instead of looping through the result set and running another query for each result row, use
  a subquery in the query of the SELECT INTO statement (see "Processing Query Result Sets
  with Subqueries").
- Instead of looping through the result set and running another DML statement for each result row, use the FORALL statement (see "FORALL Statement").

# FETCH Statement with BULK COLLECT Clause

The FETCH statement with the BULK COLLECT clause (also called the FETCH BULK COLLECT statement) fetches an entire result set into one or more collection variables.

For more information, see "FETCH Statement".

Example 13-22 uses a FETCH BULK COLLECT statement to fetch an entire result set into two collections (nested tables).

Example 13-23 uses a FETCH BULK COLLECT statement to fetch a result set into a collection (nested table) of records.

### Example 13-22 Bulk-Fetching into Two Nested Tables

```
DECLARE
   TYPE NameList IS TABLE OF employees.last_name%TYPE;
   TYPE SalList IS TABLE OF employees.salary%TYPE;
```

```
CURSOR c1 IS
    SELECT last name, salary
    FROM employees
    WHERE salary > 10000
    ORDER BY last name;
  names NameList;
  sals SalList;
  TYPE RecList IS TABLE OF c1%ROWTYPE;
  recs RecList;
  v limit PLS INTEGER := 10;
  PROCEDURE print results IS
  BEGIN
    -- Check if collections are empty:
    IF names IS NULL OR names.COUNT = 0 THEN
      DBMS OUTPUT.PUT LINE('No results!');
    ELSE
      DBMS OUTPUT.PUT LINE('Result: ');
      FOR i IN names.FIRST .. names.LAST
       DBMS OUTPUT.PUT LINE(' Employee ' || names(i) || ': $' || sals(i));
      END LOOP;
    END IF;
  END;
BEGIN
  DBMS_OUTPUT.PUT_LINE ('--- Processing all results simultaneously ---');
  OPEN c1;
  FETCH c1 BULK COLLECT INTO names, sals;
  CLOSE c1;
  print results();
  DBMS OUTPUT.PUT LINE ('--- Processing ' || v limit || ' rows at a time
---');
  OPEN c1;
  LOOP
   FETCH c1 BULK COLLECT INTO names, sals LIMIT v limit;
   EXIT WHEN names.COUNT = 0;
   print results();
  END LOOP;
  CLOSE c1;
  DBMS OUTPUT.PUT LINE ('--- Fetching records rather than columns ---');
  OPEN c1;
  FETCH c1 BULK COLLECT INTO recs;
  FOR i IN recs.FIRST .. recs.LAST
  LOOP
    -- Now all columns from result set come from one record
    DBMS OUTPUT.PUT LINE (
     ' Employee ' || recs(i).last name || ': $' || recs(i).salary
   );
  END LOOP;
END;
/
```

#### Result:

```
--- Processing all results simultaneously ---
Employee Abel: $11000
Employee Cambrault: $11000
Employee Errazuriz: $12000
Employee Garcia: $17000
Employee Gruenberg: $12008
Employee Higgins: $12008
Employee King: $24000
Employee Li: $11000
Employee Martinez: $13000
Employee Ozer: $11500
Employee Partners: $13500
Employee Singh: $14000
Employee Vishney: $10500
Employee Yang: $17000
Employee Zlotkey: $10500
--- Processing 10 rows at a time ---
Result:
Employee Abel: $11000
Employee Cambrault: $11000
Employee Errazuriz: $12000
Employee Garcia: $17000
Employee Gruenberg: $12008
Employee Higgins: $12008
Employee King: $24000
Employee Li: $11000
Employee Martinez: $13000
Employee Ozer: $11500
Employee Partners: $13500
Employee Singh: $14000
Employee Vishney: $10500
Employee Yang: $17000
Employee Zlotkey: $10500
--- Fetching records rather than columns ---
Employee Abel: $11000
Employee Cambrault: $11000
Employee Errazuriz: $12000
Employee Garcia: $17000
Employee Gruenberg: $12008
Employee Higgins: $12008
Employee King: $24000
Employee Li: $11000
Employee Martinez: $13000
Employee Ozer: $11500
Employee Partners: $13500
Employee Singh: $14000
Employee Vishney: $10500
Employee Yang: $17000
Employee Zlotkey: $10500
```

### Example 13-23 Bulk-Fetching into Nested Table of Records

```
DECLARE
  CURSOR c1 IS
    SELECT first name, last name, hire date
    FROM employees;
  TYPE NameSet IS TABLE OF c1%ROWTYPE;
  stock managers NameSet; -- nested table of records
  TYPE cursor var type is REF CURSOR;
  cv cursor var type;
BEGIN
  -- Assign values to nested table of records:
  OPEN CV FOR
    SELECT first name, last_name, hire_date
    FROM employees
   WHERE job id = 'ST MAN'
    ORDER BY hire date;
  FETCH cv BULK COLLECT INTO stock_managers;
  CLOSE cv;
  -- Print nested table of records:
    FOR i IN stock managers.FIRST .. stock managers.LAST LOOP
      DBMS OUTPUT.PUT LINE (
        stock managers(i).hire date || ' ' ||
        stock managers(i).last name || ', ' ||
        stock managers(i).first name
      );
    END LOOP; END;
Result:
01-MAY-13 Kaufling, Payam
18-JUL-14 Weiss, Matthew
10-APR-15 Fripp, Adam
10-OCT-15 Vollman, Shanta
16-NOV-17 Mourgos, Kevin
```

## Row Limits for FETCH BULK COLLECT Statements

A FETCH BULK COLLECT statement that returns a large number of rows produces a large collection. To limit the number of rows and the collection size, use the LIMIT clause.

In Example 13-24, with each iteration of the LOOP statement, the FETCH statement fetches ten rows (or fewer) into associative array empids (overwriting the previous values). Note the exit condition for the LOOP statement.

### Example 13-24 Limiting Bulk FETCH with LIMIT

```
DECLARE
  TYPE numtab IS TABLE OF NUMBER INDEX BY PLS INTEGER;
  CURSOR c1 IS
   SELECT employee id
   FROM employees
    WHERE department id = 80
    ORDER BY employee id;
  empids numtab;
BEGIN
  OPEN c1;
  LOOP -- Fetch 10 rows or fewer in each iteration
   FETCH c1 BULK COLLECT INTO empids LIMIT 10;
    DBMS OUTPUT.PUT LINE ('----- Results from One Bulk Fetch -----');
    FOR i IN 1..empids.COUNT LOOP
     DBMS OUTPUT.PUT LINE ('Employee Id: ' || empids(i));
   END LOOP;
   EXIT WHEN c1%NOTFOUND;
  END LOOP;
  CLOSE c1;
END;
Result:
----- Results from One Bulk Fetch -----
Employee Id: 145
Employee Id: 146
Employee Id: 147
Employee Id: 148
Employee Id: 149
Employee Id: 150
Employee Id: 151
Employee Id: 152
Employee Id: 153
Employee Id: 154
----- Results from One Bulk Fetch -----
Employee Id: 155
Employee Id: 156
Employee Id: 157
Employee Id: 158
Employee Id: 159
Employee Id: 160
Employee Id: 161
Employee Id: 162
Employee Id: 163
Employee Id: 164
----- Results from One Bulk Fetch -----
Employee Id: 165
Employee Id: 166
Employee Id: 167
Employee Id: 168
Employee Id: 169
Employee Id: 170
Employee Id: 171
Employee Id: 172
Employee Id: 173
Employee Id: 174
----- Results from One Bulk Fetch -----
Employee Id: 175
```

```
Employee Id: 176
Employee Id: 177
Employee Id: 179
```

## RETURNING INTO Clause with BULK COLLECT Clause

The RETURNING INTO clause with the BULK COLLECT clause (also called the RETURNING BULK COLLECT INTO clause) can appear in an INSERT, UPDATE, MERGE, DELETE, or EXECUTE IMMEDIATE statement. With the RETURNING BULK COLLECT INTO clause, the statement stores its result set in one or more collections.

For more information, see "RETURNING INTO Clause".

Example 13-25 uses a DELETE statement with the RETURNING BULK COLLECT INTO clause to delete rows from a table and return them in two collections (nested tables).

Example 13-26 uses the keywords OLD and NEW to return the values of employee salaries before and after an UPDATE statement with the RETURNING BULK COLLECT INTO clause.

### Example 13-25 Returning Deleted Rows in Two Nested Tables

```
DROP TABLE emp temp;
CREATE TABLE emp_temp AS
SELECT * FROM employees
ORDER BY employee id;
DECLARE
  TYPE NumList IS TABLE OF employees.employee id%TYPE;
  enums NumList;
  TYPE NameList IS TABLE OF employees.last name%TYPE;
  names NameList;
BEGIN
  DELETE FROM emp temp
  WHERE department id = 30
  RETURNING employee id, last name
  BULK COLLECT INTO enums, names;
  DBMS OUTPUT.PUT LINE ('Deleted ' || SQL%ROWCOUNT || ' rows:');
  FOR i IN enums.FIRST .. enums.LAST
  LOOP
    DBMS OUTPUT.PUT LINE ('Employee #' || enums(i) || ': ' || names(i));
  END LOOP;
END;
Result:
Deleted 6 rows:
Employee #114: Li
Employee #115: Khoo
Employee #116: Baida
Employee #117: Tobias
Employee #118: Himuro
Employee #119: Colmenares
```



### Example 13-26 Returning NEW and OLD Values of Updated Rows

```
DROP TABLE emp temp;
CREATE TABLE emp temp AS
SELECT * FROM employees
ORDER BY employee id;
DECLARE
  TYPE SalList IS TABLE OF employees.salary%TYPE;
  old sals SalList;
  new sals SalList;
  TYPE NameList IS TABLE OF employees.last name%TYPE;
  names NameList;
BEGIN
  UPDATE emp temp SET salary = salary * 1.15
  WHERE salary < 2500
  RETURNING OLD salary, NEW salary, last name
  BULK COLLECT INTO old sals, new sals, names;
  DBMS OUTPUT.PUT LINE('Updated ' || SQL%ROWCOUNT || ' rows: ');
  FOR i IN old sals.FIRST .. old sals.LAST
    DBMS OUTPUT.PUT LINE(names(i) || ': Old Salary $' || old sals(i) ||
            ', New Salary $' || new sals(i));
  END LOOP;
END;
Result:
Landry: Old Salary $2400, New Salary $2760
Markle: Old Salary $2200, New Salary $2530
Olson: Old Salary $2100, New Salary $2415
Gee: Old Salary $2400, New Salary $2760
Philtanker: Old Salary $2200, New Salary $2530
```

# Using FORALL Statement and BULK COLLECT Clause Together

In a FORALL statement, the DML statement can have a RETURNING BULK COLLECT INTO clause. For each iteration of the FORALL statement, the DML statement stores the specified values in the specified collections—without overwriting the previous values, as the same DML statement would do in a FOR LOOP statement.

In Example 13-27, the FORALL statement runs a DELETE statement that has a RETURNING BULK COLLECT INTO clause. For each iteration of the FORALL statement, the DELETE statement stores the employee\_id and department\_id values of the deleted row in the collections e\_ids and d ids, respectively.

Example 13-28 is like Example 13-27 except that it uses a FOR LOOP statement instead of a FORALL statement.

#### Example 13-27 DELETE with RETURN BULK COLLECT INTO in FORALL Statement

```
DROP TABLE emp_temp;
CREATE TABLE emp_temp AS
```

```
SELECT * FROM employees
ORDER BY employee id, department id;
DECLARE
  TYPE NumList IS TABLE OF NUMBER;
  depts NumList := NumList(10,20,30);
  TYPE enum t IS TABLE OF employees.employee id%TYPE;
  e ids enum t;
  TYPE dept t IS TABLE OF employees.department id%TYPE;
  d ids dept t;
BEGIN
  FORALL j IN depts.FIRST..depts.LAST
   DELETE FROM emp temp
    WHERE department id = depts(j)
    RETURNING employee id, department id
    BULK COLLECT INTO e ids, d ids;
  DBMS OUTPUT.PUT LINE ('Deleted ' | SQL%ROWCOUNT | ' rows:');
  FOR i IN e ids.FIRST .. e ids.LAST
  LOOP
    DBMS OUTPUT.PUT LINE (
      'Employee #' || e_ids(i) || ' from dept #' || d_ids(i)
  END LOOP;
END;
Result:
Deleted 9 rows:
Employee #200 from dept #10
Employee #201 from dept #20
Employee #202 from dept #20
Employee #114 from dept #30
Employee #115 from dept #30
Employee #116 from dept #30
Employee #117 from dept #30
Employee #118 from dept #30
Employee #119 from dept #30
```

### Example 13-28 DELETE with RETURN BULK COLLECT INTO in FOR LOOP Statement

```
DROP TABLE emp_temp;
CREATE TABLE emp_temp AS
SELECT * FROM employees
ORDER BY employee_id, department_id;

DECLARE
   TYPE NumList IS TABLE OF NUMBER;
   depts NumList := NumList(10,20,30);
```

```
TYPE enum t IS TABLE OF employees.employee id%TYPE;
  e ids enum t;
  TYPE dept t IS TABLE OF employees.department id%TYPE;
  d ids dept t;
BEGIN
  FOR j IN depts.FIRST..depts.LAST LOOP
   DELETE FROM emp temp
    WHERE department id = depts(j)
   RETURNING employee id, department id
    BULK COLLECT INTO e_ids, d ids;
  END LOOP;
  DBMS OUTPUT.PUT LINE ('Deleted ' || SQL%ROWCOUNT || ' rows:');
  FOR i IN e ids.FIRST .. e ids.LAST
  LOOP
    DBMS OUTPUT.PUT LINE (
      'Employee #' || e_ids(i) || ' from dept #' || d_ids(i)
  END LOOP;
END;
Result:
Deleted 6 rows:
Employee #114 from dept #30
Employee #115 from dept #30
Employee #116 from dept #30
Employee #117 from dept #30
Employee #118 from dept #30
Employee #119 from dept #30
```

# Client Bulk-Binding of Host Arrays

Client programs (such as OCI and Pro\*C programs) can use PL/SQL anonymous blocks to bulk-bind input and output host arrays. This is the most efficient way to pass collections to and from the database server.

In the client program, declare and assign values to the host variables to be referenced in the anonymous block. In the anonymous block, prefix each host variable name with a colon (:) to distinguish it from a PL/SQL collection variable name. When the client program runs, the database server runs the PL/SQL anonymous block.

In Example 13-29, the anonymous block uses a FORALL statement to bulk-bind a host input array. In the FORALL statement, the DELETE statement refers to four host variables: scalars lower, upper, and emp id and array depts.

#### Example 13-29 Anonymous Block Bulk-Binds Input Host Array

```
BEGIN
  FORALL i IN :lower..:upper
   DELETE FROM employees
  WHERE department id = :depts(i);
```

END;

# Chaining Pipelined Table Functions for Multiple Transformations

Chaining pipelined table functions is an efficient way to perform multiple transformations on data.



You cannot run a pipelined table function over a database link. The reason is that the return type of a pipelined table function is a SQL user-defined type, which can be used only in a single database (as explained in *Oracle Database Object-Relational Developer's Guide*). Although the return type of a pipelined table function might appear to be a PL/SQL type, the database actually converts that PL/SQL type to a corresponding SQL user-defined type.

### **Topics**

- Overview of Table Functions
- Creating Pipelined Table Functions
- Pipelined Table Functions as Transformation Functions
- Chaining Pipelined Table Functions
- Fetching from Results of Pipelined Table Functions
- Passing CURSOR Expressions to Pipelined Table Functions
- DML Statements on Pipelined Table Function Results
- NO\_DATA\_NEEDED Exception

# Overview of Table Functions

A **table function** is a user-defined PL/SQL function that returns a collection of rows (an associative array, nested table or varray).

You can select from this collection as if it were a database table by invoking the table function inside the TABLE clause in a SELECT statement. The TABLE operator is optional.

#### For example:

```
SELECT * FROM TABLE(table_function_name(parameter_list))
```

Alternatively, the same query can be written without the TABLE operator as follow:

```
SELECT * FROM table_function_name(parameter_list)
```

A table function can take a collection of rows as input (that is, it can have an input parameter that is a nested table, varray, or cursor variable). Therefore, output from table function tf1 can be input to table function tf2, and output from tf2 can be input to table function tf3, and so on.

To improve the performance of a table function, you can:



• Enable the function for parallel execution, with the PARALLEL ENABLE option.

Functions enabled for parallel execution can run concurrently.

Stream the function results directly to the next process, with Oracle Streams.

Streaming eliminates intermediate staging between processes.

Pipeline the function results, with the PIPELINED option.

A **pipelined table function** returns a row to its invoker immediately after processing that row and continues to process rows. Response time improves because the entire collection need not be constructed and returned to the server before the query can return a single result row. (Also, the function needs less memory, because the object cache need not materialize the entire collection.)



#### **Caution:**

A pipelined table function always references the current state of the data. If the data in the collection changes after the cursor opens for the collection, then the cursor reflects the changes. PL/SQL variables are private to a session and are not transactional. Therefore, read consistency, well known for its applicability to table data, does not apply to PL/SQL collection variables.

# See Also:

- Chaining Pipelined Table Functions
- Oracle Database SQL Language Reference for more information about the TABLE clause of the SELECT statement
- Oracle Database Data Cartridge Developer's Guide for information about using pipelined and parallel table functions

# **Creating Pipelined Table Functions**

A pipelined table function must be either a standalone function or a package function.

### **PIPELINED Option (Required)**

For a standalone function, specify the PIPELINED option in the CREATE FUNCTION Statement (for syntax, see "CREATE FUNCTION Statement"). For a package function, specify the PIPELINED option in both the function declaration and function definition (for syntax, see "Function Declaration and Definition").

## PARALLEL\_ENABLE Option (Recommended)

To improve its performance, enable the pipelined table function for parallel execution by specifying the PARALLEL ENABLE option.

### **AUTONOMOUS\_TRANSACTION Pragma**

If the pipelined table function runs DML statements, then make it autonomous, with the AUTONOMOUS TRANSACTION pragma (described in "AUTONOMOUS TRANSACTION Pragma").

Then, during parallel execution, each instance of the function creates an independent transaction.

#### **DETERMINISTIC Option (Recommended)**

Multiple invocations of a pipelined table function, in either the same query or separate queries, cause multiple executions of the underlying implementation. If the function is deterministic, specify the DETERMINISTIC option, described in "DETERMINISTIC Clause".

#### **Parameters**

Typically, a pipelined table function has one or more cursor variable parameters. For information about cursor variables as function parameters, see "Cursor Variables as Subprogram Parameters".

# See Also:

- "Cursor Variables" for general information about cursor variables
- "Subprogram Parameters" for general information about subprogram parameters

### **RETURN Data Type**

The data type of the value that a pipelined table function returns must be a collection type defined either at schema level or inside a package (therefore, it cannot be an associative array type). The elements of the collection type must be SQL data types, not data types supported only by PL/SQL (such as PLS\_INTEGER). For information about collection types, see "Collection Types". For information about SQL data types, see *Oracle Database SQL Language Reference*.

You can use SQL data types ANYTYPE, ANYDATA, and ANYDATASET to dynamically encapsulate and access type descriptions, data instances, and sets of data instances of any other SQL type, including object and collection types. You can also use these types to create unnamed types, including anonymous collection types. For information about these types, see *Oracle Database PL/SQL Packages and Types Reference*.

#### **PIPE ROW Statement**

Inside a pipelined table function, use the PIPE ROW statement to return a collection element to the invoker without returning control to the invoker. See "PIPE ROW Statement" for its syntax and semantics.

### **RETURN Statement**

As in every function, every execution path in a pipelined table function must lead to a RETURN statement, which returns control to the invoker. However, in a pipelined table function, a RETURN statement need not return a value to the invoker. See "RETURN Statement" for its syntax and semantics.



## **Example**

### Example 13-30 Creating and Invoking Pipelined Table Function

This example creates a package that includes a pipelined table function, f1, and then selects from the collection of rows that f1 returns.

```
CREATE OR REPLACE PACKAGE pkg1 AUTHID DEFINER AS
  TYPE numset t IS TABLE OF NUMBER;
  FUNCTION f1(x NUMBER) RETURN numset_t PIPELINED;
END pkg1;
Create a pipelined table function f1 that returns a collection of elements (1,2,3,... x).
CREATE OR REPLACE PACKAGE BODY pkg1 AS
  FUNCTION f1(x NUMBER) RETURN numset_t PIPELINED IS
  BEGIN
    FOR i IN 1..x LOOP
      PIPE ROW(i);
    END LOOP;
    RETURN;
 END f1;
END pkg1;
SELECT * FROM TABLE(pkg1.f1(5));
Result:
COLUMN_VALUE
           1
           2
           3
            4
            5
5 rows selected.
SELECT * FROM pkg1.f1(2);
Result:
COLUMN VALUE
           1
            2
```



# Pipelined Table Functions as Transformation Functions

A pipelined table function with a cursor variable parameter can serve as a transformation function. Using the cursor variable, the function fetches an input row. Using the PIPE ROW statement, the function pipes the transformed row or rows to the invoker. If the FETCH and PIPE ROW statements are inside a LOOP statement, the function can transform multiple input rows.

In Example 13-31, the pipelined table function transforms each selected row of the <code>employees</code> table to two nested table rows, which it pipes to the <code>SELECT</code> statement that invokes it. The actual parameter that corresponds to the formal cursor variable parameter is a <code>CURSOR</code> expression; for information about these, see "Passing CURSOR Expressions to Pipelined Table Functions".

### **Example 13-31 Pipelined Table Function Transforms Each Row to Two Rows**

```
CREATE OR REPLACE PACKAGE refcur pkg AUTHID DEFINER IS
 TYPE refcur t IS REF CURSOR RETURN employees%ROWTYPE;
 TYPE outrec typ IS RECORD (
   var num NUMBER(6),
   var char1 VARCHAR2(30),
   var char2 VARCHAR2(30)
  TYPE outrecset IS TABLE OF outrec typ;
  FUNCTION f trans (p refcur t) RETURN outrecset PIPELINED;
END refcur_pkg;
CREATE OR REPLACE PACKAGE BODY refcur pkg IS
 FUNCTION f trans (p refcur t) RETURN outrecset PIPELINED IS
    out rec outrec typ;
    in rec p%ROWTYPE;
 BEGIN
    LOOP
     FETCH p INTO in rec; -- input row
     EXIT WHEN p%NOTFOUND;
     out rec.var num := in rec.employee id;
     out rec.var char1 := in rec.first name;
     out rec.var char2 := in rec.last name;
     PIPE ROW(out rec); -- first transformed output row
     out rec.var char1 := in rec.email;
     out rec.var char2 := in rec.phone number;
     PIPE ROW(out rec); -- second transformed output row
    END LOOP;
    CLOSE p;
   RETURN;
 END f trans;
END refcur pkg;
SELECT * FROM TABLE (
  refcur pkg.f trans (
    CURSOR (SELECT * FROM employees WHERE department id = 60)
```



```
);
```

#### Result:

VAR_NUM	VAR_CHAR1	VAR_CHAR2
102	Alexander	James
103	AJAMES	1.590.555.0103
104	Bruce	Miller
104	BMILLER	1.590.555.0104
105	David	Williams
105	DWILLIAMS	1.590.555.0105
106	Valli	Jackson
106	VJACKSON	1.590.555.0106
107	Diana	Nguyen
107	DNGUYEN	1.590.555.0107

10 rows selected.

# **Chaining Pipelined Table Functions**

To **chain** pipelined table functions tf1 and tf2 is to make the output of tf1 the input of tf2. For example:

```
SELECT * FROM TABLE(tf2(CURSOR(SELECT * FROM TABLE(tf1()))));
```

The rows that tf1 pipes out must be compatible actual parameters for the formal input parameters of tf2.

If chained pipelined table functions are enabled for parallel execution, then each function runs in a different process (or set of processes).



"Passing CURSOR Expressions to Pipelined Table Functions"

# Fetching from Results of Pipelined Table Functions

You can associate a named cursor with a query that invokes a pipelined table function. Such a cursor has no special fetch semantics, and such a cursor variable has no special assignment semantics.

However, the SQL optimizer does not optimize across PL/SQL statements. Therefore, in Example 13-32, the first PL/SQL statement is slower than the second—despite the overhead of running two SQL statements in the second PL/SQL statement, and even if function results are piped between the two SQL statements in the first PL/SQL statement.

In Example 13-32, assume that f and g are pipelined table functions, and that each function accepts a cursor variable parameter. The first PL/SQL statement associates cursor variable g with a query that invokes g, and then passes g to g. The second PL/SQL statement passes CURSOR expressions to both g and g.



See Also:

"Cursor Variables as Subprogram Parameters"

### Example 13-32 Fetching from Results of Pipelined Table Functions

```
DECLARE
    r SYS_REFCURSOR;
...
    -- First PL/SQL statement (slower):
BEGIN
    OPEN r FOR SELECT * FROM TABLE(f(CURSOR(SELECT * FROM tab)));
    SELECT * BULK COLLECT INTO rec_tab FROM TABLE(g(r));
    -- NOTE: When g completes, it closes r.
END;
-- Second PL/SQL statement (faster):
SELECT * FROM TABLE(g(CURSOR(SELECT * FROM TABLE(f(CURSOR(SELECT * FROM tab))))));
//
TABLE(f(CURSOR(SELECT * FROM tab))))));
```

# Passing CURSOR Expressions to Pipelined Table Functions

As Example 13-32 shows, the actual parameter for the cursor variable parameter of a pipelined table function can be either a cursor variable or a CURSOR expression, and the latter is more efficient.

Note:

When a SQL SELECT statement passes a CURSOR expression to a function, the referenced cursor opens when the function begins to run and closes when the function completes.

See Also:

"CURSOR Expressions" for general information about CURSOR expressions

Example 13-33 creates a package that includes a pipelined table function with two cursor variable parameters and then invokes the function in a SELECT statement, using CURSOR expressions for actual parameters.

Example 13-34 uses a pipelined table function as an aggregate function, which takes a set of input rows and returns a single result. The SELECT statement selects the function result. (For information about the pseudocolumn COLUMN\_VALUE, see *Oracle Database SQL Language Reference*.)



### Example 13-33 Pipelined Table Function with Two Cursor Variable Parameters

```
CREATE OR REPLACE PACKAGE refcur pkg AUTHID DEFINER IS
 TYPE refcur t1 IS REF CURSOR RETURN employees%ROWTYPE;
 TYPE refcur t2 IS REF CURSOR RETURN departments%ROWTYPE;
 TYPE outrec typ IS RECORD (
   var num NUMBER(6),
   var char1 VARCHAR2(30),
   var char2 VARCHAR2(30)
 TYPE outrecset IS TABLE OF outrec typ;
 FUNCTION g_trans (p1 refcur_t1, p2 refcur_t2) RETURN outrecset PIPELINED;
END refcur pkg;
CREATE OR REPLACE PACKAGE BODY refcur pkg IS
 FUNCTION g trans (
   pl refcur tl,
   p2 refcur t2
 ) RETURN outrecset PIPELINED
   out rec outrec typ;
   in rec1 p1%ROWTYPE;
   in rec2 p2%ROWTYPE;
 BEGIN
   LOOP
     FETCH p2 INTO in rec2;
     EXIT WHEN p2%NOTFOUND;
   END LOOP;
   CLOSE p2;
   LOOP
     FETCH pl INTO in recl;
     EXIT WHEN p1%NOTFOUND;
     -- first row
     out rec.var num := in rec1.employee id;
     out rec.var char1 := in rec1.first name;
     out rec.var char2 := in rec1.last name;
     PIPE ROW(out rec);
     -- second row
     out rec.var num := in rec2.department id;
     out rec.var char1 := in rec2.department name;
     out rec.var char2 := TO CHAR(in rec2.location id);
     PIPE ROW(out rec);
   END LOOP;
   CLOSE p1;
   RETURN;
 END g trans;
END refcur pkg;
SELECT * FROM TABLE (
 refcur pkg.g trans (
```

```
CURSOR (SELECT * FROM employees WHERE department_id = 60),
   CURSOR (SELECT * FROM departments WHERE department_id = 60)
);
```

#### Result:

VAR_NUM	VAR_CHAR1	VAR_CHAR2
103	Alexander	James
60	IT	1400
104	Bruce	Miller
60	IT	1400
105	David	Williams
60	IT	1400
106	Valli	Jackson
60	IT	1400
107	Diana	Nguyen
60	IT	1400
10 rows selected.		

### Example 13-34 Pipelined Table Function as Aggregate Function

```
DROP TABLE gradereport;
CREATE TABLE gradereport (
  student VARCHAR2(30),
  subject VARCHAR2(30),
 weight NUMBER,
  grade NUMBER
INSERT INTO gradereport (student, subject, weight, grade)
VALUES ('Mark', 'Physics', 4, 4);
INSERT INTO gradereport (student, subject, weight, grade)
VALUES ('Mark', 'Chemistry', 4, 3);
INSERT INTO gradereport (student, subject, weight, grade)
VALUES ('Mark', 'Maths', 3, 3);
INSERT INTO gradereport (student, subject, weight, grade)
VALUES ('Mark', 'Economics', 3, 4);
CREATE OR REPLACE PACKAGE pkg gpa AUTHID DEFINER IS
  TYPE gpa IS TABLE OF NUMBER;
  FUNCTION weighted average (input values SYS REFCURSOR)
   RETURN gpa PIPELINED;
END pkg gpa;
CREATE OR REPLACE PACKAGE BODY pkg gpa IS
  FUNCTION weighted_average (input_values SYS_REFCURSOR)
```

```
RETURN gpa PIPELINED
   grade
               NUMBER;
   total NUMBER := 0;
   total weight NUMBER := 0;
   weight NUMBER := 0;
 BEGIN
   LOOP
     FETCH input values INTO weight, grade;
     EXIT WHEN input values%NOTFOUND;
     total weight := total weight + weight; -- Accumulate weighted average
     total := total + grade*weight;
   END LOOP;
   PIPE ROW (total / total weight);
   RETURN; -- returns single result
 END weighted average;
END pkg gpa;
```

This query shows how the table function can be invoked without the optional TABLE operator.

# DML Statements on Pipelined Table Function Results

The "table" that a pipelined table function returns cannot be the target table of a DELETE, INSERT, UPDATE, or MERGE statement. However, you can create a view of such a table and create INSTEAD OF triggers on the view. For information about INSTEAD OF triggers, see "INSTEAD OF DML Triggers".



1 row selected.

Oracle Database SQL Language Reference for information about the CREATE VIEW statement

# NO DATA NEEDED Exception

You must understand the predefined exception  ${\tt NO\_DATA\_NEEDED}$  in two cases:

You include an OTHERS exception handler in a block that includes a PIPE ROW statement

Your code that feeds a PIPE ROW statement must be followed by a clean-up procedure
 Typically, the clean-up procedure releases resources that the code no longer needs.

When the invoker of a pipelined table function needs no more rows from the function, the PIPE ROW statement raises NO\_DATA\_NEEDED. If the pipelined table function does not handle NO\_DATA\_NEEDED, as in Example 13-35, then the function invocation terminates but the invoking statement does not terminate. If the pipelined table function handles NO\_DATA\_NEEDED, its exception handler can release the resources that it no longer needs, as in Example 13-36.

In Example 13-35, the pipelined table function <code>pipe\_rows</code> does not handle the <code>NO\_DATA\_NEEDED</code> exception. The <code>SELECT</code> statement that invokes <code>pipe\_rows</code> needs only four rows. Therefore, during the fifth invocation of <code>pipe\_rows</code>, the <code>PIPE ROW</code> statement raises the exception <code>NO\_DATA\_NEEDED</code>. The fifth invocation of <code>pipe\_rows</code> terminates, but the <code>SELECT</code> statement does not terminate.

If the exception-handling part of a block that includes a PIPE ROW statement includes an OTHERS exception handler to handle unexpected exceptions, then it must also include an exception handler for the expected NO\_DATA\_NEEDED exception. Otherwise, the OTHERS exception handler handles the NO\_DATA\_NEEDED exception, treating it as an unexpected error. The following exception handler reraises the NO\_DATA\_NEEDED exception, instead of treating it as a irrecoverable error:

```
EXCEPTION

WHEN NO_DATA_NEEDED THEN

RAISE;

WHEN OTHERS THEN

-- (Put error-logging code here)

RAISE_APPLICATION_ERROR(-20000, 'Irrecoverable error.');

END;
```

In Example 13-36, assume that the package External Source contains these public items:

- Procedure Init, which allocates and initializes the resources that Next Row needs
- Function Next\_Row, which returns some data from a specific external source and raises the
  user-defined exception Done (which is also a public item in the package) when the external
  source has no more data
- Procedure Clean Up, which releases the resources that Init allocated

The pipelined table function <code>get\_external\_source\_data</code> pipes rows from the external source by invoking <code>External\_Source.Next\_Row</code> until either:

The external source has no more rows.

In this case, the  ${\tt External\_Source.Next\_Row}$  function raises the user-defined exception  ${\tt External\_Source.Done}$ .

• get external source data needs no more rows.

In this case, the PIPE ROW statement in  $get_external_source_data$  raises the NO DATA NEEDED exception.

In either case, an exception handler in block b in get\_external\_source\_data invokes External Source.Clean Up, which releases the resources that Next Row was using.

### Example 13-35 Pipelined Table Function Does Not Handle NO DATA NEEDED

```
CREATE TYPE t IS TABLE OF NUMBER /
CREATE OR REPLACE FUNCTION pipe rows RETURN t PIPELINED AUTHID DEFINER IS
```



```
n NUMBER := 0;
BEGIN
  LOOP
   n := n + 1;
   PIPE ROW (n);
  END LOOP;
END pipe rows;
SELECT COLUMN VALUE
  FROM TABLE(pipe rows())
  WHERE ROWNUM < 5
Result:
COLUMN VALUE
-----
           2
           3
4 rows selected.
```

### Example 13-36 Pipelined Table Function Handles NO\_DATA\_NEEDED

```
CREATE OR REPLACE FUNCTION get_external_source_data
  RETURN t PIPELINED AUTHID DEFINER IS
BEGIN
                                  -- Initialize.
  External Source.Init();
  <<b>> BEGIN
   LOOP
                                  -- Pipe rows from external source.
     PIPE ROW (External Source.Next Row());
  EXCEPTION
    WHEN External Source.Done THEN -- When no more rows are available,
     External_Source.Clean_Up(); -- clean up.
    WHEN NO DATA NEEDED THEN -- When no more rows are needed,
     External_Source.Clean_Up(); -- clean up.
                                   -- Optional, equivalent to RETURN.
     RAISE NO DATA NEEDED;
  END b;
END get external source data;
```

# Overview of Polymorphic Table Functions

Polymorphic table functions (PTF) are table functions whose operands can have more than one type. The return type is determined by the PTF invocation arguments list. The actual arguments to the table type usually determines the row output shape, but not always.

### **Introduction to Polymorphic Table Functions**

Polymorphic Table Functions (PTF) are user-defined functions that can be invoked in the FROM clause of a SQL query block. They are capable of processing tables whose row type is not declared at definition time and producing a result table whose row type may or may not be declared at definition time. Polymorphic table functions leverage dynamic SQL capabilities to create powerful and complex custom functions. This is useful for applications demanding an interface with generic extensions which work for arbitrary input tables or queries.

A PTF author creates an interface to a procedural mechanism that defines a table. The PTF author defines, documents, and implements the PTF.

The query author can only describe the published interface and invoke the PTF function in queries.

The database is the PTF conductor. It manages the compilation and execution states of the PTF. The database and the PTF author can see a family of related SQL invoked procedures, called the PTF component procedures, and possibly additional private data (such as variables and cursors).

### **Types of Polymorphic Table Functions**

The polymorphic table function type is specified based on their formal arguments list semantics:

- If an input TABLE argument has Row Semantics, the input is a single row.
- If an input TABLE argument has Table Semantics, the input is a set of rows. When a Table Semantics PTF is called from a query, the table argument can optionally be extended with either a PARTITION BY clause or an ORDER BY clause or both.

# Polymorphic Table Function Definition

The PTF author defines, documents, and implements the Polymorphic Table Function (PTF).

A PTF has two parts:

- 1. The PL/SQL package which contains the client interface for the PTF implementation.
- 2. The standalone or package function naming the PTF and its associated implementation package.

# Polymorphic Table Function Implementation

The Polymorphic Table Function (PTF) implementation client interface is a set of subprograms with fixed names that every PTF must provide.

### Steps to Implement a Polymorphic Table Function

- 1. Create the implementation package containing the DESCRIBE function (required) and the OPEN, FETCH ROWS, and CLOSE procedures (optional).
- 2. Create the function specification naming the PTF. The function can be created at the top-level after the package has been created, or as a package function in the implementation package (the package created in the first step). Polymorphic table functions do not have a function definition (a FUNCTION BODY), the definition is encapsulated in the associated implementation package.

The function definition specifies:

- The Polymorphic Table Function (PTF) name
- Exactly one formal argument of type TABLE and any number of non TABLE arguments
- The return type of the PTF as TABLE
- The type of PTF function (row or table semantics)
- The PTF implementation package name



# See Also:

- Oracle Database PL/SQL Packages and Types Reference for information about a DESCRIBE Only polymorphic table function
- Oracle Database PL/SQL Packages and Types Reference for more information about how to specify the PTF implementation package and use the DBMS\_TF utilities
- PIPELINED Clause for the standalone or package polymorphic table function creation syntax and semantic

# Polymorphic Table Function Invocation

A polymorphic table function is invoked by specifying its name followed by the argument list in the FROM clause of a SQL query block.

The PTF arguments can be the standard scalar arguments that can be passed to a regular table function, but PTF's can additionally take a table argument. A table argument is either a WITH clause query or a schema-level object that is allowed in a FROM clause (such as tables, views, or table functions).

### **Syntax**

```
table_argument ::= table [ PARTITION BY column_list ] [ORDER BY order_column_list]

column_list ::= identifier | ( identifier[, identifier...])

order_column_list ::= order_column_name | (order_column_name [, order_column_name...])

order_column_name ::= identifier [ ASC | DESC ][ NULLS FIRST | NULLS LAST ]
```

#### **Semantics**

Each identifier is a column in the corresponding table.

```
The PTF has Table Semantics.
```

Query can optionally partition and order Table Semantics PTF input. This is disallowed for Row Semantics PTF input.

A polymorphic table function (PTF) cannot be the target of a DML statement. Any table argument of a PTF is passed in by name.

For example, the noop PTF can be used in a query such as :

```
SELECT *
FROM noop(emp);

Or

WITH e AS
  (SELECT * FROM emp NATURAL JOIN dept)
SELECT t.* FROM noop(e) t;
```

The input table argument must be a basic table name.

The name resolution rules of the table identifier are (in priority order) as follows:

- Identifier is resolved as a column name (such as a correlated column from an outer query block).
- 2. Identifier is resolved as a Common Table Expression (CTE) name in the current or some outer query-block. CTE is commonly known as the WITH clause.
- 3. Identifier is resolved as a schema-level table, view, or table-function (regular or polymorphic, and defined either at the schema-level or inside a package).

Many types of table expressions otherwise allowed in the FROM clause cannot be directly used as a table argument for a PTF (such as ANSI Joins, bind-variables, in-line views, CURSOR operators, TABLE operators). To use such table expressions as a PTF argument, these table expressions must be passed indirectly into a PTF by wrapping them in a CTE and then passing the CTE name into the PTF.

A PTF can be used as a table reference in the FROM clause and thus can be part of the ANSI Join and LATERAL syntax. Additionally, a PTF can be the source table for PIVOT/UNPIVOT and MATCH\_RECOGNIZE. Some table modification clauses that are meant for tables and views (such as SAMPLING, PARTITION, CONTAINERS) are disallowed for PTF.

Direct function composition of PTF is allowed (such as nested PTF cursor expression invocation or PTF(TF()) nesting). However, nested PTF is disallowed (such as PTF(PTF()) nesting).

The scalar arguments of a PTF can be any SQL scalar expression. While the constant scalar values are passed as-is to the DESCRIBE function, all other values are passed as NULLs. This is usually not a problem for the PTF implementation if these values are not row shape determining, but otherwise the DESCRIBE function can raise an error; typically the documentation accompanying the PTF will state which scalar parameters, if any, are shape defining and thus must have constant non-null values. Note, that during query execution (during OPEN, FETCH\_ROWS, CLOSE) the expressions are evaluated and their actual values are passed to these PTF execution procedures. The return type is determined by the PTF invocation arguments list.

Query arguments are passed to PTF using a WITH clause.

The TABLE operator is optional when the table function arguments list or empty list () appears.

# Variadic Pseudo-Operators

A variadic pseudo-operator operates with a variable number of operands.

Starting with Oracle Database Release 18c, we introduce the concept of variadic pseudo-operator into the SQL expression language to support Polymorphic Table Functions (PTF). A pseudo-operator can be used to pass list of identifiers (such as column name) to a PTF. A pseudo-operator can only appear as arguments to PTFs, and are parsed by the SQL compiler like other SQL operators or PL/SQL function invocation. A pseudo-operator has a variable number of arguments but must have at least one. The pseudo-operator does not have any execution function associated with it, and they are completely removed from the SQL cursor after the PTF compilation is finished. During SQL compilation, the pseudo-operators are converted to corresponding DBMS\_TF types and then passed to the DESCRIBE method. There is no output type associated with these operators. It is not possible to embed a pseudo-operator inside a general SQL expression.



# **COLUMNS Pseudo-Operator**

You can use the COLUMNS pseudo-operator to specify arguments to a Polymorphic Table Function (PTF) invocation in the FROM clause of a SQL query block.

The COLUMNS pseudo-operator arguments specify the list of column names, or the list of column names with associated types.

### **Syntax**

```
column_operator ::= COLUMNS ( column_list )

column_list ::= column_name_list | column_type_list

column_name_list ::= identifier [, identifier ... ]

column_type_list::= identifier column_type [, identifier column_type...]
```

#### **Semantics**

The COLUMNS pseudo-operator can only appear as an argument to a PTF. It cannot appear in any other SQL expression than the PTF expression itself.

The column\_type must be a scalar type.

# Polymorphic Table Function Compilation and Execution

The database fulfills the Polymorphic Table Functions (PTF) conductor role. As such, it is responsible for the PTF compilation, execution and its related states.

The database manages:

- The compilation state : This is the immutable state that is generated by DESCRIBE which is needed before execution.
- The execution state: This is the state used by the execution procedures of a Table semantics PTF.

See Also:

 Oracle Database PL/SQL Packages and Types Reference for more information about how the database manages the compilation and execution states of the PTFs

# Polymorphic Table Function Optimization

A polymorphic table function (PTF) provides an efficient and scalable mechanism to extend the analytical capabilities of the database.

The key benefits are:

Minimal data-movement: Only columns of interest are passed to PTF

- Predicates/Projections/Partitioning are/is pushed into underlying table/query (where semantically possible)
- Bulk data transfer into and out of PTF
- Parallelism is based on type of PTF and query specified partitioning (if any)

# Skip\_col Polymorphic Table Function Example

This PTF example demonstrates Row Semantics, Describe Only, package table function, and overloading features.



Oracle Database PL/SQL Packages and Types Reference for more Polymorphic Table Function (PTF) examples

### Example 13-37 Skip\_col Polymorphic Table Function Example

The skip\_col Polymorphic Table Function (PTF) returns all the columns in a table except the columns specified in the PTF input argument. The skip\_col PTF skips columns based on column names (overload 1) or columns data type (overload 2).



You can view and run this example on Oracle Live SQL at 18c Skip\_col Polymorphic Table Function

Create the implementation package named skip\_col\_pkg containing the DESCRIBE function for the skip\_col polymorphic table function (PTF). The DESCRIBE function is invoked to determine the row shape produced by the PTF. It returns a DBMS\_TF.DESCRIBE\_T table. It is overloaded. The FETCH\_ROWS procedure is not required because it does need to produce associated new column values for a given subset of rows.

Create the implementation package body which contains the polymorphic table function definition.

```
CREATE PACKAGE BODY skip col pkg AS
/* OVERLOAD 1: Skip by name
* Package PTF name: skip col pkg.skip col
 * Standalone PTF name: skip col by name
* PARAMETERS:
* tab - The input table
^{\star} col - The name of the columns to drop from the output
 * DESCRIPTION:
   This PTF removes all the input columns listed in col from the output
   of the PTF.
FUNCTION describe(tab IN OUT DBMS TF.TABLE T,
                  col DBMS TF.COLUMNS T)
           RETURN DBMS TF.DESCRIBE T
   new cols DBMS TF.COLUMNS NEW T;
   col id PLS INTEGER := 1;
 BEGIN
   FOR i IN 1 .. tab.column.count() LOOP
     FOR j IN 1 .. col.count() LOOP
       tab.column(i).PASS THROUGH := tab.column(i).DESCRIPTION.NAME !=
       EXIT WHEN NOT tab.column(i).PASS THROUGH;
     END LOOP;
   END LOOP;
   RETURN NULL;
 END;
/* OVERLOAD 2: Skip by type
 * Package PTF name: skip col pkg.skip col
 * Standalone PTF name: skip col by type
 * PARAMETERS:
           - Input table
    type name - A string representing the type of columns to skip
    flip - 'False' [default] => Match columns with given type name
                otherwise => Ignore columns with given type name
* DESCRIPTION:
   This PTF removes the given type of columns from the given table.
```

```
FUNCTION describe(tab
                      IN OUT DBMS TF. TABLE T,
                  RETURN DBMS TF.DESCRIBE T
 AS
   typ CONSTANT VARCHAR2(1024) := UPPER(TRIM(type name));
 BEGIN
   FOR i IN 1 .. tab.column.count() LOOP
      tab.column(i).PASS THROUGH :=
        CASE UPPER(SUBSTR(flip, 1, 1))
          WHEN 'F' THEN DBMS TF.column type name(tab.column(i).DESCRIPTION)!
=typ
                       DBMS TF.column type name(tab.column(i).DESCRIPTION)
         ELSE
=typ
        END /* case */;
   END LOOP;
   RETURN NULL;
 END;
END skip_col_pkg;
```

Create a standalone polymorphic table function named skip\_col\_by\_name for overload 1. Specify exactly one formal argument of type TABLE, specify the return type of the PTF as TABLE, specify a Row Semantics PTF type, and indicate the PTF implementation package to use is skip\_col\_pkg.

Create a standalone polymorphic table function named skip\_col\_by\_type for overload 2. Specify exactly one formal argument of type TABLE, specify the return type of the PTF as TABLE, specify a Row Semantics PTF type, and indicate the PTF implementation package to use is skip\_col\_pkg.

```
CREATE FUNCTION skip_col_by_type(tab TABLE,

type_name VARCHAR2,

flip VARCHAR2 DEFAULT 'False')

RETURN TABLE PIPELINED ROW POLYMORPHIC USING skip col pkg;
```

Invoke the package skip\_col PTF (overload 1) to report from the SCOTT.DEPT table only columns whose type is not NUMBER.



The same result can be achieved by invoking the standalone skip\_col\_by\_type PTF to report from the SCOTT.DEPT table only columns whose type is not NUMBER.

Invoke the package skip\_col PTF (overload 2) to report from the SCOTT. DEPT table only columns whose type is NUMBER.

```
SELECT * FROM skip_col_pkg.skip_col(scott.dept, 'number', flip => 'True');

DEPTNO
-----
10
20
30
40
```

The same result can be achieved by invoking the standalone skip\_col\_by\_type PTF to report from the SCOTT.DEPT table only columns whose type is NUMBER.

```
SELECT * FROM skip_col_by_type(scott.dept, 'number', flip => 'True');

DEPTNO
-----
10
20
30
40
```

Invoke the package skip\_col PTF to report all employees in department 20 from the SCOTT.EMP table all columns except COMM, HIREDATE and MGR.

```
SELECT *
FROM skip_col_pkg.skip_col(scott.emp, COLUMNS(comm, hiredate, mgr))
WHERE deptno = 20;
```

EMPNO 1	ENAME	JOB	SAL	DEPTNO
7369	SMITH	CLERK	800	20
7566	JONES	MANAGER	2975	20
7788	SCOTT	ANALYST	3000	20
7876	ADAMS	CLERK	1100	20
7902	FORD	ANALYST	3000	20

# To\_doc Polymorphic Table Function Example

The to doc PTF example combines a list of specified columns into a single document column.

### Example 13-38 To\_doc Polymorphic Table Function Example

The to\_doc PTF combines a list of columns into a document column constructed like a JSON object.



You can view and run this example on Oracle Live SQL at 18c To\_doc Polymorphic Table Function

Create the implementation package to\_doc\_p containing the DESCRIBE function and FETCH ROWS procedure for the to\_doc polymorphic table function (PTF).

#### The PTF parameters are:

- tab: The input table (The tab parameter is of type DBMS\_TF.TABLE\_T, a table descriptor record type)
- cols (optional): The list of columns to convert to document. (The cols parameter is type DBMS\_TF.COLUMNS\_T, a column descriptor record type)

```
CREATE PACKAGE to_doc_p AS

FUNCTION describe(tab IN OUT DBMS_TF.TABLE_T,

cols IN DBMS_TF.COLUMNS_T DEFAULT NULL)

RETURN DBMS_TF.DESCRIBE_T;

PROCEDURE fetch_rows;
END to_doc_p;
```

Create the package containing the DESCRIBE function and FETCH\_ROWS procedure. The FETCH\_ROWS procedure is required to produce a new column named DOCUMENT in the output rowset. The DESCRIBE function indicates the read columns by annotating them in the input table descriptor, TABLE\_T. Only the indicated read columns will be fetched and thus available for processing during FETCH\_ROWS. The PTF invocation in a query can use the COLUMNS pseudo-operator to indicate which columns the query wants the PTF to read, and this information is passed to the DESCRIBE function which then in turn sets the COLUMN\_T.FOR\_READ boolean flag. Only scalar SQL data types are allowed for the read columns. The COLUMN\_T.PASS\_THROUGH boolean flag indicates columns that are passed from the input table of the PTF to the output, without any modifications.

```
FOR j IN 1 .. cols.count LOOP
        IF (tab.column(i).DESCRIPTION.NAME = cols(j)) THEN
            tab.column(i).FOR READ := TRUE;
            tab.column(i).PASS THROUGH := FALSE;
        END IF;
     END LOOP;
 END LOOP;
 RETURN DBMS TF.describe t(new columns => DBMS TF.COLUMNS NEW T(1 =>
                              DBMS TF.COLUMN_METADATA_T(name
=> 'DOCUMENT')));
END;
PROCEDURE fetch rows AS
     rst DBMS TF.ROW SET T;
     col DBMS TF. TAB VARCHAR2 T;
     rct PLS INTEGER;
BEGIN
     DBMS TF.GET_ROW_SET(rst, row_count => rct);
     FOR rid IN 1 .. rct LOOP
          col(rid) := DBMS TF.ROW TO CHAR(rst, rid);
     END LOOP;
     DBMS TF.PUT COL(1, col);
END;
END to doc p;
```

Create the standalone to\_doc PTF. Specify exactly one formal argument of type TABLE, specify the return type of the PTF as TABLE, specify a Row Semantics PTF type, and indicate the PTF implementation package to use is to\_doc\_p.

Invoke the to\_doc PTF to display all columns of table SCOTT.DEPT as one combined DOCUMENT column.

```
SELECT * FROM to_doc(scott.dept);

DOCUMENT

{"DEPTNO":10, "DNAME":"ACCOUNTING", "LOC":"NEW YORK"}
{"DEPTNO":20, "DNAME":"RESEARCH", "LOC":"DALLAS"}
{"DEPTNO":30, "DNAME":"SALES", "LOC":"CHICAGO"}
{"DEPTNO":40, "DNAME":"OPERATIONS", "LOC":"BOSTON"}
```

For all employees in departments 10 and 30, display the DEPTNO, ENAME and DOCUMENT columns ordered by DEPTNO and ENAME. Invoke the to\_doc PTF with the COLUMNS pseudo-operator to

select columns EMPNO, JOB, MGR, HIREDATE, SAL and COMM of table SCOTT.EMP. The PTF combines these columns into the DOCUMENT column.

```
SELECT deptno, ename, document
     to doc(scott.emp, COLUMNS(empno,job,mgr,hiredate,sal,comm))
WHERE deptno IN (10, 30)
ORDER BY 1, 2;
DEPTNO ENAME
               DOCUMENT
                {"EMPNO":7782, "JOB":"MANAGER", "MGR":7839, "HIREDATE":"09-JUN-81",
   10 CLARK
"SAL":2450}
   10 KING
               {"EMPNO":7839, "JOB":"PRESIDENT", "HIREDATE":"17-NOV-81", "SAL":5000}
   10 MILLER
                {"EMPNO":7934, "JOB":"CLERK", "MGR":7782, "HIREDATE":"23-JAN-82",
"SAL":1300}
   30 ALLEN
               {"EMPNO":7499, "JOB":"SALESMAN", "MGR":7698, "HIREDATE":"20-FEB-81",
"SAL":1600, "COMM":300}
   30 BLAKE {"EMPNO":7698, "JOB":"MANAGER", "MGR":7839, "HIREDATE":"01-MAY-81",
"SAL":2850}
   30 JAMES
               {"EMPNO":7900, "JOB":"CLERK", "MGR":7698, "HIREDATE":"03-DEC-81",
"SAL":950}
   30 MARTIN
               {"EMPNO":7654, "JOB":"SALESMAN", "MGR":7698, "HIREDATE":"28-SEP-81",
"SAL":1250, "COMM":1400}
   30 TURNER {"EMPNO":7844, "JOB":"SALESMAN", "MGR":7698, "HIREDATE":"08-SEP-81",
"SAL":1500, "COMM":0}
   30 WARD {"EMPNO":7521, "JOB":"SALESMAN", "MGR":7698, "HIREDATE":"22-FEB-81",
"SAL":1250, "COMM":500}
```

With the subquery named E, display the DOC\_ID and DOCUMENT columns. Report all clerk employees, their salary, department and department location. Use the to\_doc PTF to combine the NAME, SAL, DEPTNO and LOC columns into the DOCUMENT column.

```
WITH e AS (

SELECT ename name, sal, deptno, loc

FROM scott.emp NATURAL JOIN scott.dept

WHERE job = 'CLERK')

SELECT ROWNUM doc_id, t.*

FROM to_doc(e) t;

DOC_ID DOCUMENT

1 {"NAME":"MILLER", "SAL":1300, "DEPTNO":10, "LOC":"NEW YORK"}

2 {"NAME":"SMITH", "SAL":800, "DEPTNO":20, "LOC":"DALLAS"}

3 {"NAME":"ADAMS", "SAL":1100, "DEPTNO":20, "LOC":"DALLAS"}

4 {"NAME":"JAMES", "SAL":950, "DEPTNO":30, "LOC":"CHICAGO"}
```

Use a subquery block to display c1, c2, c3 column values converted into the DOCUMENT column.

```
WITH t(c1,c2,c3) AS (

SELECT NULL, NULL, NULL FROM dual

UNION ALL

SELECT 1, NULL, NULL FROM dual

UNION ALL

SELECT NULL, 2, NULL FROM dual

UNION ALL

SELECT 0, NULL, 3 FROM dual)
```

For all employees in department 30, display the values of the member with property names ENAME and COMM. The PTF invocation reporting from the SCOTT.EMP table produces the DOCUMENT column which can be used as input to the JSON\_VALUE function. This function selects a scalar value from some JSON data.

# Implicit\_echo Polymorphic Table Function Example

The implicit\_echo PTF example demonstrates that the USING clause is optional when the Polymorphic Table Function and the DESCRIBE function are defined in the same package.

### Example 13-39 Implicit echo Polymorphic Table Function Example

The implicit\_echo PTF, takes in a table and a column and produces a new column with the same value.

This PTF returns the column in the input table tab, and adds to it the column listed in cols but with the column names prefixed with "ECHO\_".

Create the implementation package implicit\_echo\_package containing the DESCRIBE function, implicit\_echo polymorphic table function (PTF) and FETCH\_ROWS procedure.

```
END implicit echo package;
```

Create the package containing the DESCRIBE function containing the input table parameter and the column parameter to be read. This function is invoked to determine the type of rows produced by the Polymorphic Table Function. The function returns a table DBMS\_TF.DESCRIBE\_T. The FETCH\_ROWS procedure is required to produce the indicated read column along with a new column prefixed with "ECHO\_" in the output rowset. The implicit\_echo is the PTF function and contains two arguments, tab and cols, whose values are obtained from the query and this information is passed to the DESCRIBE function. The Row semantics specifies a PTF type but without the USING clause. This function is invoked from the SQL query.

Create the implementation package body <code>implicit\_echo\_package</code> which contains the PTF definition.

```
CREATE PACKAGE BODY implicit echo package AS
FUNCTION DESCRIBE (tab IN OUT DBMS TF. TABLE T,
                 cols IN DBMS TF.COLUMNS T)
         RETURN DBMS TF.DESCRIBE T
AS
  new cols DBMS TF.COLUMNS NEW T;
  col id PLS INTEGER := 1;
BEGIN
 FOR i in 1 .. tab.column.COUNT LOOP
   FOR j in 1 .. cols.COUNT LOOP
     IF (tab.column(i).description.name = cols(j)) THEN
       IF (NOT DBMS TF.SUPPORTED TYPE(tab.column(i).description.type)) THEN
            RAISE APPLICATION ERROR(-20102, 'Unsupported column type['||
                                    tab.column(i).description.type||']');
      END IF;
       tab.column(i).for read := TRUE;
       new cols(col id) := tab.column(i).description;
       new cols(col id).name := prefix||
REGEXP REPLACE (tab.column(i).description.name,
'^"|"$');
                            := col id + 1;
      col id
      EXIT;
    END IF;
    END LOOP;
 END LOOP;
/* VERIFY ALL COLUMNS WERE FOUND */
 IF (col id - 1 != cols.COUNT) then
```

Invoke the PTF to display ENAME column of table SCOTT.EMP and display it along with another column ECHO ENAME having the same value.

```
SELECT ENAME, ECHO ENAME
FROM implicit echo package.implicit echo(SCOTT.EMP, COLUMNS(SCOTT.ENAME));
ENAME
       ECHO ENAME
______
     SMITH
ALLEN
SMITH
ALLEN
WARD
        WARD
       JONES
JONES
MARTIN MARTIN
BLAKE BLAKE
        CLARK
CLARK
SCOTT
       SCOTT
KING
       KING
TURNER TURNER
ADAMS ADAMS
JAMES
       JAMES
FORD
       FORD
MILLER MILLER
```

# **Updating Large Tables in Parallel**

The DBMS\_PARALLEL\_EXECUTE package lets you incrementally update the data in a large table in parallel, in two high-level steps:

- 1. Group sets of rows in the table into smaller chunks.
- 2. Apply the desired UPDATE statement to the chunks in parallel, committing each time you have finished processing a chunk.

This technique is recommended whenever you are updating a lot of data. Its advantages are:

- You lock only one set of rows at a time, for a relatively short time, instead of locking the entire table.
- You do not lose work that has been done if something fails before the entire operation finishes.

- You reduce rollback space consumption.
- You improve performance.

## See Also

Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS\_PARALLEL\_EXECUTE package

# Collecting Data About User-Defined Identifiers

PL/Scope extracts, organizes, and stores data about PL/SQL and SQL identifiers and SQL statements from PL/SQL source text. You can retrieve the identifiers and statements data with the static data dictionary views \* IDENTIFIERS and \*\_STATEMENTS.

# See Also:

- PL/SQL Units and Compilation Parameters for more information about PLSQL\_SETTINGS parameter
- Oracle Database Development Guide for more information about using PL/Scope

# Profiling and Tracing PL/SQL Programs

To help you isolate performance problems in large PL/SQL programs, PL/SQL provides these tools, implemented as PL/SQL packages.

Table 13-1 Profiling and Tracing Tools Summary

Tool	Package	Description
Profiler interface	DBMS_PROFILER	Computes the time that your PL/SQL program spends at each line and in each subprogram.
		You must have CREATE privileges on the units to be profiled.
		Saves runtime statistics in database tables, which you can query.
Trace	DBMS_TRACE	Traces the order in which subprograms run.
interface		You can specify the subprograms to trace and the tracing level.
		Saves runtime statistics in database tables, which you can query.



Table 13-1 (Cont.) Profiling and Tracing Tools Summary

Tool	Package	Description
PL/SQL hierarchical profiler	DBMS_HPROF	Reports the dynamic execution program profile of your PL/SQL program, organized by subprogram invocations. Accounts for SQL and PL/SQL execution times separately.
		Requires no special source or compile-time preparation.
		Generates reports in HTML. Provides the option of storing profiler data and results in relational format in database tables for custom report generation (such as third-party tools offer).
SQL trace	DBMS_APPLICATION_INFO	Uses the DBMS_APPLICATION_INFO package with Oracle Trace and the SQL trace facility to record names of executing modules or transactions in the database for later use when tracking the performance of various modules and debugging.
PL/SQL Basic Block Coverage	DBMS_PLSQL_CODE_COVERAGE	Collects and analyzes basic block coverage data.
Call Stack Utilities	UTL_CALL_STACK	Provides information about currently executing subprograms (such as subprogram names, unit names, owner names, edition names, and error stack information) that you can use to create more revealing error logs and application execution traces.

### **Related Topics**

- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS APPLICATION INFO package
- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS HPROF package
- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS PLSQL CODE COVERAGE package
- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS PROFILER package
- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS TRACE package

Oracle Database PL/SQL Packages and Types Reference for more information about the UTL\_CALL\_STACK package

- COVERAGE Pragma for the syntax and semantics of COVERAGE PRAGMA
- Oracle Database Development Guide for more information about using PL/SQL basic block coverage
- Oracle Database Development Guide for a detailed description of PL/SQL hierarchical profiler
- Oracle Database Development Guide for more information about analyzing and debugging stored subprograms

# Compiling PL/SQL Units for Native Execution

You can usually speed up PL/SQL units by compiling them into native code (processor-dependent system code), which is stored in the SYSTEM tablespace.

You can natively compile any PL/SQL unit of any type, including those that Oracle Database supplies.

Natively compiled program units work in all server environments, including shared server configuration (formerly called "multithreaded server") and Oracle Real Application Clusters (Oracle RAC).

On most platforms, PL/SQL native compilation requires no special set-up or maintenance. On some platforms, the DBA might want to do some optional configuration.

## See Also:

- Oracle Database Administrator's Guide for information about configuring a database
- Platform-specific configuration documentation for your platform

You can test to see how much performance gain you can get by enabling PL/SQL native compilation.

If you have determined that PL/SQL native compilation will provide significant performance gains in database operations, Oracle recommends compiling the entire database for native mode, which requires DBA privileges. This speeds up both your own code and calls to the PL/SQL packages that Oracle Database supplies.

### **Topics**

- Determining Whether to Use PL/SQL Native Compilation
- How PL/SQL Native Compilation Works
- · Dependencies, Invalidation, and Revalidation
- Setting Up a New Database for PL/SQL Native Compilation\*
- Compiling the Entire Database for PL/SQL Native or Interpreted Compilation\*

# Determining Whether to Use PL/SQL Native Compilation

Whether to compile a PL/SQL unit for native or interpreted mode depends on where you are in the development cycle and on what the program unit does.

While you are debugging program units and recompiling them frequently, interpreted mode has these advantages:

- You can use PL/SQL debugging tools on program units compiled for interpreted mode (but not for those compiled for native mode).
- Compiling for interpreted mode is faster than compiling for native mode.



<sup>\*</sup> Requires DBA privileges.

After the debugging phase of development, in determining whether to compile a PL/SQL unit for native mode, consider:

- PL/SQL native compilation provides the greatest performance gains for computationintensive procedural operations. Examples are data warehouse applications and applications with extensive server-side transformations of data for display.
- PL/SQL native compilation provides the least performance gains for PL/SQL subprograms that spend most of their time running SQL.
- When many program units (typically over 15,000) are compiled for native execution, and are simultaneously active, the large amount of shared memory required might affect system performance.

# How PL/SQL Native Compilation Works

Without native compilation, the PL/SQL statements in a PL/SQL unit are compiled into an intermediate form, system code, which is stored in the catalog and interpreted at run time.

With PL/SQL native compilation, the PL/SQL statements in a PL/SQL unit are compiled into native code and stored in the catalog. The native code need not be interpreted at run time, so it runs faster.

Because native compilation applies only to PL/SQL statements, a PL/SQL unit that uses only SQL statements might not run faster when natively compiled, but it does run at least as fast as the corresponding interpreted code. The compiled code and the interpreted code make the same library calls, so their action is the same.

The first time a natively compiled PL/SQL unit runs, it is fetched from the SYSTEM tablespace into shared memory. Regardless of how many sessions invoke the program unit, shared memory has only one copy it. If a program unit is not being used, the shared memory it is using might be freed, to reduce memory load.

Natively compiled subprograms and interpreted subprograms can invoke each other.

PL/SQL native compilation works transparently in an Oracle Real Application Clusters (Oracle RAC) environment.

The PLSQL\_CODE\_TYPE compilation parameter determines whether PL/SQL code is natively compiled or interpreted. For information about this compilation parameters, see "PL/SQL Units and Compilation Parameters".

# Dependencies, Invalidation, and Revalidation

Recompilation is automatic with invalidated PL/SQL modules. For example, if an object on which a natively compiled PL/SQL subprogram depends changes, the subprogram is invalidated. The next time the same subprogram is called, the database recompiles the subprogram automatically. Because the PLSQL\_CODE\_TYPE setting is stored inside the library unit for each subprogram, the automatic recompilation uses this stored setting for code type.

Explicit recompilation does not necessarily use the stored PLSQL\_CODE\_TYPE setting. For the conditions under which explicit recompilation uses stored settings, see "PL/SQL Units and Compilation Parameters".

# Setting Up a New Database for PL/SQL Native Compilation

If you have DBA privileges, you can set up a new database for PL/SQL native compilation by setting the compilation parameter PLSQL CODE TYPE to NATIVE. The performance benefits apply

to the PL/SQL packages that Oracle Database supplies, which are used for many database operations.



If you compile the whole database as NATIVE, Oracle recommends that you set PLSQL CODE TYPE at the system level.

# Compiling the Entire Database for PL/SQL Native or Interpreted Compilation

If you have DBA privileges, you can recompile all PL/SQL modules in an existing database to NATIVE or INTERPRETED, using the dbmsupgnv.sql and dbmsupgin.sql scripts respectively during the process explained in this section. Before making the conversion, review "Determining Whether to Use PL/SQL Native Compilation".

## Note:

- If you compile the whole database as NATIVE, Oracle recommends that you set PLSQL CODE TYPE at the system level.
- If Database Vault is enabled, then you can run dbmsupgnv.sql only if the Database Vault administrator has granted you the DV PATCH ADMIN role.
- The conversion process described here affects only the current container's units.
   Units in other containers are not affected.

During the conversion to native compilation, TYPE specifications are not recompiled by dbmsupgnv.sql to NATIVE because these specifications do not contain executable code.

Package specifications seldom contain executable code so the runtime benefits of compiling to NATIVE are not measurable. You can use the TRUE command-line parameter with the dbmsupgnv.sql script to exclude package specs from recompilation to NATIVE, saving time in the conversion process.

When converting to interpreted compilation, the <code>dbmsupgin.sql</code> script does not accept any parameters and does not exclude any PL/SQL units.

### Note:

The following procedure describes the conversion to native compilation. If you must recompile all PL/SQL modules to interpreted compilation, make these changes in the steps.

- Skip the first step.
- Set the PLSQL\_CODE\_TYPE compilation parameter to INTERPRETED rather than NATIVE.
- Substitute dbmsupgin.sql for the dbmsupgnv.sql script.

1. Ensure that a test PL/SQL unit can be compiled. For example:

```
ALTER PROCEDURE my_proc COMPILE PLSQL_CODE_TYPE=NATIVE REUSE SETTINGS;
```

- 2. Shut down application services, the listener, and the database.
  - Shut down all of the Application services including the Forms Processes, Web Servers, Reports Servers, and Concurrent Manager Servers. After shutting down all of the Application services, ensure that all of the connections to the database were terminated.
  - Shut down the TNS listener of the database to ensure that no new connections are made.
  - Shut down the database in normal or immediate mode as the user SYS. See *Oracle Database Administrator's Guide*.
- 3. Set PLSQL\_CODE\_TYPE to NATIVE in the compilation parameter file. If the database is using a server parameter file, then set this after the database has started.

The value of PLSQL\_CODE\_TYPE does not affect the conversion of the PL/SQL units in these steps. However, it does affect all subsequently compiled units, so explicitly set it to the desired compilation type.

- **4.** Start up the database in upgrade mode, using the <code>UPGRADE</code> option. For information about <code>SQL\*Plus</code> <code>STARTUP</code>, see <code>SQL\*Plus</code> <code>User's</code> Guide and Reference.
- 5. Run this code to list the invalid PL/SQL units. You can save the output of the query for future reference with the SQL SPOOL statement:

```
-- To save the output of the query to a file:

SPOOL pre_update_invalid.log

SELECT o.OWNER, o.OBJECT_NAME, o.OBJECT_TYPE

FROM DBA_OBJECTS o, DBA_PLSQL_OBJECT_SETTINGS s

WHERE o.OBJECT_NAME = s.NAME AND o.STATUS='INVALID';

-- To stop spooling the output: SPOOL OFF
```

If any Oracle supplied units are invalid, try to validate them by recompiling them. For example:

```
ALTER PACKAGE SYS.DBMS OUTPUT COMPILE BODY REUSE SETTINGS;
```

If the units cannot be validated, save the spooled log for future resolution and continue.

6. Run this query to determine how many objects are compiled NATIVE and INTERPRETED (to save the output, use the SQL SPOOL statement):

```
SELECT TYPE, PLSQL_CODE_TYPE, COUNT(*)

FROM DBA_PLSQL_OBJECT_SETTINGS

WHERE PLSQL_CODE_TYPE IS NOT NULL AND ORIGIN_CON_ID=SYS_CONTEXT('USERENV', 'CON_ID')

GROUP BY TYPE, PLSQL_CODE_TYPE

ORDER BY TYPE, PLSQL_CODE_TYPE;
```

Any objects with a NULL plsql code type are special internal objects and can be ignored.

7. Run the \$ORACLE\_HOME/rdbms/admin/dbmsupgnv.sql script as the user SYS to update the plsql\_code\_type setting to NATIVE in the dictionary tables for all PL/SQL units. This process also invalidates the units. Use TRUE with the script to exclude package specifications; FALSE to include the package specifications.

- This update must be done when the database is in <code>UPGRADE</code> mode. The script is guaranteed to complete successfully or rollback all the changes.
- 8. Shut down the database and restart in NORMAL mode.
- 9. Before you run the utlrp.sql script, Oracle recommends that no other sessions are connected to avoid possible problems. You can ensure this with this statement:
  - ALTER SYSTEM ENABLE RESTRICTED SESSION;
- 10. Run the <code>\$ORACLE\_HOME/rdbms/admin/utlrp.sql</code> script as the user <code>SYS</code>. This script recompiles all the PL/SQL modules using a default degree of parallelism. See the comments in the script for information about setting the degree explicitly.
  - If for any reason the script is terminated atypically, rerun the utlrp.sql script to recompile any remaining invalid PL/SQL modules.
- 11. After the compilation completes successfully, verify that there are no invalid PL/SQL units using the query in step 5. You can spool the output of the query to the post\_upgrade\_invalid.log file and compare the contents with the pre\_upgrade\_invalid.log file, if it was created previously.
- 12. Re-run the query in step 6. If recompiling with <code>dbmsupgnv.sql</code>, confirm that all PL/SQL units, except <code>TYPE</code> specifications and package specifications if excluded, are <code>NATIVE</code>. If recompiling with <code>dbmsupgin.sql</code>, confirm that all PL/SQL units are <code>INTERPRETED</code>.
- **13.** Disable the restricted session mode for the database, then start the services that you previously shut down. To disable restricted session mode, use this statement:

ALTER SYSTEM DISABLE RESTRICTED SESSION;

