The ORACLE_LOADER Access Driver

Learn how to control the way external tables are accessed by using the ORACLE_LOADER access driver parameters to modify the default behavior of the access driver.

About the ORACLE LOADER Access Driver

The <code>ORACLE_LOADER</code> access driver provides a set of access parameters unique to external tables of the type <code>ORACLE_LOADER</code>.

access parameters Clause

The access_parameters clause contains comments, record formatting, and field formatting information.

record format info Clause

Learn how to parse, label and manage record information with the <code>record_format_info</code> clause and its subclauses.

field definitions Clause

Learn how to name the fields in the data file and specify how to find them in records using the field definitions clause.

· column transforms Clause

The optional <code>ORACLE_LOADER</code> access drive <code>COLUMN TRANSFORMS</code> clause provides transforms that you can use to describe how to load columns in the external table that do not map directly to columns in the data file.

- Parallel Loading Considerations for the ORACLE_LOADER Access Driver
 The ORACLE_LOADER access driver attempts to divide large data files into chunks that can be processed separately.
- Performance Hints When Using the ORACLE_LOADER Access Driver
 This topic describes some performance hints when using the ORACLE LOADER access driver.
- Restrictions When Using the ORACLE_LOADER Access Driver
 This section lists restrictions to be aware of when you use the ORACLE_LOADER access driver.
- Reserved Words for the ORACLE_LOADER Access Driver
 When identifiers (for example, column or table names) are specified in the external table
 access parameters, certain values are considered to be reserved words by the access
 parameter parser.

15.1 About the ORACLE_LOADER Access Driver

The <code>ORACLE_LOADER</code> access driver provides a set of access parameters unique to external tables of the type <code>ORACLE_LOADER</code>.

You can use the access parameters to modify the default behavior of the access driver. The information you provide through the access driver ensures that data from the data source is processed so that it matches the definition of the external table.

To use the external table management features that the ORACLE_LOADER access parameters provide, you must have some knowledge of the file format and record format (including character sets and field data types) of the data files on your platform. You must also

know enough about SQL to be able to create an external table, and to perform queries against it.

You can find it helpful to use the EXTERNAL_TABLE=GENERATE_ONLY parameter in SQL*Loader to obtain the proper access parameters for a given SQL*Loader control file. When you specify GENERATE_ONLY, all the SQL statements needed to do the load using external tables, as described in the control file, are placed in the SQL*Loader log file. You can edit and customize these SQL statements. You can perform the actual load later without the use of SQL*Loader by executing these statements in SQL*Plus.

Note:

- It is sometimes difficult to understand <code>ORACLE_LOADER</code> access driver parameter syntax without reference to other <code>ORACLE_LOADER</code> access driver parameters. If you have difficulty understanding the syntax of a particular parameter, then refer to it in context with other referenced parameters.
- Be aware that in ORACLE_LOADER access driver parameter examples that show a CREATE TABLE...ORGANIZATION EXTERNAL statement, followed by an example of contents of the data file for the external table, the contents of the data file in the example are not part of the CREATE TABLE statement. They are present in the example only to help complete the example.
- When identifiers (for example, column or table names) are specified in the
 external table access parameters, certain values are considered to be reserved
 words by the access parameter parser. If a reserved word is used as an identifier,
 then it must be enclosed in double quotation marks.

Related Topics

- EXTERNAL TABLE
 - The EXTERNAL_TABLE parameter instructs SQL*Loader whether to load data using the external tables option.
- Reserved Words for the ORACLE_LOADER Access Driver
 When identifiers (for example, column or table names) are specified in the external table
 access parameters, certain values are considered to be reserved words by the access
 parameter parser.
- Oracle Database Administrator's Guide

15.2 access parameters Clause

The access_parameters clause contains comments, record formatting, and field formatting information.

Default

None.

Syntax

The syntax for the access parameters clause is as follows:





Purpose

The description of the data in the data source is separate from the definition of the external table. This means that:

- The source file can contain more or fewer fields than there are columns in the external table
- The data types for fields in the data source can be different from the columns in the external table

The access driver ensures that data from the data source is processed so that it matches the definition of the external table.



These access parameters are collectively referred to as the <code>opaque_format_spec</code> in the SQL CREATE TABLE...ORGANIZATION EXTERNAL statement.

See Also:

Oracle Database SQL Language Reference for information about specifying opaque_format_spec when using the SQL CREATE TABLE...ORGANIZATION EXTERNAL statement

comments

Comments are lines that begin with two hyphens followed by text. Comments must be placed *before* any access parameters, for example:

```
--This is a comment.
--This is another comment.
RECORDS DELIMITED BY NEWLINE
```

All text to the right of the double hyphen is ignored, until the end of the line.

record_format_info

The record_format_info clause is an optional clause that contains information about the record, such as its format, the character set of the data, and what rules are used to exclude records from being loaded. For a full description of the syntax, see record format info Clause.

field_definitions

The field_definitions clause is used to describe the fields in the data file. If a data file field has the same name as a column in the external table, then the data from the field is used for that column. For a full description of the syntax, see field_definitions Clause.



column_transforms

The column_transforms clause is an optional clause used to describe how to load columns in the external table that do not map directly to columns in the data file. This is done using the following transforms: NULL, CONSTANT, CONCAT, and LOBFILE. For a full description of the syntax, see column_transforms Clause.

15.3 record format info Clause

Learn how to parse, label and manage record information with the <code>record_format_info</code> clause and its subclauses.

· Overview of record format info Clause

The record_format_info clause contains information about the record, such as its format, the character set of the data, and what rules are used to exclude records from being loaded.

FIXED Length

Use the record_format_info FIXED clause to identify the records in external tables as all having a fixed size of length bytes.

VARIABLE size

Use the record_format_info VARIABLE clause to indicate that the records have a variable length

DELIMITED BY

Use the record format info DELIMITED BY clause to delimit the end-of-record character.

XMLTAG

Use the <code>record_format_info</code> XMLTAG clause to specify XML tags that are used to load subdocuments from an XML document.

CHARACTERSET

Use the record_format_info CHARACTERSET clause to specify the character set of the data file.

PREPROCESSOR

To specify your own preprocessor program that you want to run for every data file, use the record_format_info PREPROCESSOR clause.

PREPROCESSOR TIMEOUT

To extend the timeout period for preprocessor programs, use the <code>record_format_info</code> <code>PREPROCESSOR_TIMEOUT clause</code>.

EXTERNAL VARIABLE DATA

To load dump files into the Oracle SQL Connector for HDFS that are generated with the ORACLE DATAPUMP access driver, use the EXTERNAL VARIABLE DATA clause.

LANGUAGE

The Language clause allows you to specify a language name (for example, FRENCH), from which locale-sensitive information about the data can be derived.

TERRITORY

The TERRITORY clause allows you to specify a territory name to further determine input data characteristics.

DATA IS...ENDIAN

The DATA IS...ENDIAN clause indicates the endianness of data whose byte order may vary, depending on the platform that generated the data file.

BYTEORDERMARK [CHECK | NOCHECK]

Use the record_format_info BYTEORDERMARK clause to specify whether the data file should be checked for the presence of a byte-order mark (BOM).

STRING SIZES ARE IN

Use the record_format_info STRING SIZES ARE IN clause to indicate whether the lengths specified for character strings are in bytes or characters.

LOAD WHEN

Use the <code>record_format_info LOAD WHEN</code> clause to identify the records that should be passed to the database.

BADFILE | NOBADFILE

Use the <code>record_format_info</code> <code>BADFILE</code> clause to name the file to which records are written when they cannot be loaded because of errors.

• DISCARDFILE | NODISCARDFILE

Use the record_format_info DISCARDFILE clause to name the file to which records are written that fail the condition in the LOAD WHEN clause.

LOGFILE | NOLOGFILE

Use the record_format_info LOGFILE clause to name the file that contains messages generated by the external tables utility while it was accessing data in the data file.

SKIP

Use the ${\tt record_format_info}$ SKIP clause to skip the specified number of records in the data file before loading.

FIELD NAMES

Use the record format info FIELD NAMES clause to specify field order in data files.

READSIZE

The READSIZE parameter specifies the size of the read buffer used to process records.

DATE_CACHE

string

A string is a quoted series of characters or hexadecimal digits.

condition spec

The <code>condition_spec</code> specifies one or more conditions that are joined by Boolean operators.

• [directory object name:] [filename]

The [directory object name:] [filename] clause is used to specify the name of an output file (BADFILE, DISCARDFILE, or LOGFILE).

condition

To compare a range of bytes or a field from the record against a constant string, you can use the <code>ORACLE LOADER condition clause</code>

IO OPTIONS clause

To specify whether the operating system uses direct input/output to read data files from disk, or uses a cache for reading the data files, use the <code>ORACLE_LOADER</code> records clause <code>IO OPTIONS</code>.

DNFS DISABLE | DNFS ENABLE

To disable and enable use of the Direct NFS Client on input data files during an external tables operation, use <code>DNFS_DISABLE</code> or <code>DNFS_ENABLE</code>.

DNFS READBUFFERS

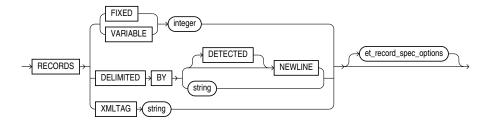
The DNFS_READBUFFERS parameter of the record_format_info clause is used to control the number of read buffers used by the Direct NFS Client.

15.3.1 Overview of record_format_info Clause

The record_format_info clause contains information about the record, such as its format, the character set of the data, and what rules are used to exclude records from being loaded.

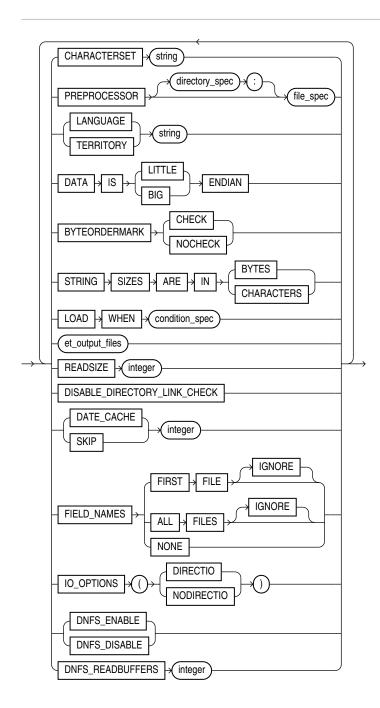
The PREPROCESSOR clause allows you to optionally specify the name of a user-supplied program that will run and modify the contents of a data file so that the <code>ORACLE_LOADER</code> access driver can parse it.

The record_format_info clause is optional. The syntax for the record_format_info clause is as follows:

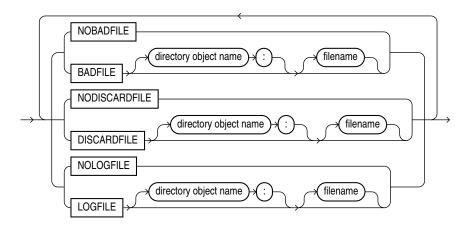


The et_record_spec_options clause allows you to optionally specify additional formatting information. You can specify as many of the formatting options as you want, in any order. The syntax of the options is as follows:





The following <code>et_output_files</code> diagram shows the options for specifying the bad, discard, and log files. For each of these clauses, you must supply either a directory object name or a file name, or both.



15.3.2 FIXED Length

Use the $record_format_info$ FIXED clause to identify the records in external tables as all having a fixed size of length bytes.

Default

None.

Purpose

Enables you to identify the records in external tables as all having a fixed size of length bytes.

Usage Notes

The size specified for FIXED records must include any record termination characters, such as newlines. Compared to other record types, fixed-length fields in fixed-length records are the easiest field and record formats for the access driver to process.

Example

The following is an example of using FIXED records. In this example, we assume that there is a 1-byte newline character at the end of each record in the data file. After the create table command using FIXED, you see an example of the data file that you can load with it.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth CHAR(4))

ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (RECORDS FIXED 20 FIELDS (first_name CHAR(7),
last_name CHAR(8),
year_of_birth CHAR(4)))

LOCATION ('info.dat'));

Alvin Tolliver1976

KennethBaer 1963

Mary Dube 1973
```



15.3.3 VARIABLE size

Use the record_format_info VARIABLE clause to indicate that the records have a variable length

Default

None.

Purpose

Use the VARIABLE clause to indicate that the records have a variable length, and that each record is preceded by a character string containing a number with the count of bytes for the record. The length of the character string containing the count field is the size argument that follows the VARIABLE parameter. Note that size indicates a count of bytes, not characters. The count at the beginning of the record must include any record termination characters, but it does not include the size of the count field itself. The number of bytes in the record termination characters can vary depending on how the file is created and on what platform it is created.

Example

In the following example of using VARIABLE records, there is a 1-byte newline character at the end of each record in the data file. After the SQL example, you see an example of a data file that can be used to load it.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth CHAR(4))

ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (RECORDS VARIABLE 2 FIELDS TERMINATED BY ','

(first_name CHAR(7),
 last_name CHAR(8),
 year_of_birth CHAR(4)))

LOCATION ('info.dat'));

21Alvin, Tolliver, 1976,
19Kenneth, Baer, 1963,
16Mary, Dube, 1973,
```

15.3.4 DELIMITED BY

Use the record format info DELIMITED BY clause to delimit the end-of-record character.

Default

None

Purpose

The DELIMITED BY clause is used to indicate the character that identifies the end of a record.

If you specify <code>DELIMITED BY NEWLINE</code> then the actual value used is platform-specific. On Unix or Linux operating systems, <code>NEWLINE</code> is assumed to be ' \n' . On Microsoft Windows operating systems, <code>NEWLINE</code> is assumed to be ' \n' .

If you are unsure what record delimiter was used when a data file was created, then running an external table query with <code>DELIMITED BY NEWLINE</code> can result in files that are incorrectly loaded.

The query can be run without identifying what record delimiter was used when the data file was created. For example, you can work on a Unix or Linux operating system and use a file that was created in Windows format. If you specify RECORDS DELIMITED BY NEWLINE on the UNIX or Linux operating system, the delimiter is automatically assumed to be '\n'. However, because the file was created in Windows format, in which the records are delimited by '\r\n', the file is incorrectly uploaded to the UNIX or Linux operating system.

To resolve problems of different record delimiters, use this syntax:

RECORDS DELIMITED BY DETECTED NEWLINE

With this syntax, the <code>ORACLE_LOADER</code> access driver scans the data looking first for a Windows delimiter (' \r \n'). If a Windows delimiter is not found, then the access driver looks for a Unix or Linux delimiter (' \n '). The first delimiter found is the one used as the record delimiter.

After a record delimiter is found, the access driver identifies that delimiter as the end of the record. For this reason, if the data contains an embedded delimiter character in a field before the end of the record, then you cannot use the DETECTED keyword. This is because the ORACLE_LOADER access driver incorrectly assumes that the delimiter in the field denotes the end of the record. As a result, the current and all subsequent records in the file cannot parse correctly.

You cannot mix newline delimiters in the same file. When the <code>ORACLE_LOADER</code> access driver finds the first delimiter, then that is the delimiter that it identifies for the records in the file. The access driver then processes all subsequent records in the file by using the same newline character as the delimiter..

If you specify <code>DELIMITED BY string</code>, then <code>string</code> can be either text or a series of hexadecimal digits enclosed within quotation marks and prefixed by <code>OX</code> or <code>X</code>. If the string is text, then the text is converted to the character set of the data file, and the result is used for identifying record boundaries.

If the following conditions are true, then you must use hexadecimal digits to identify the delimiter:

- The character set of the access parameters is different from the character set of the data file.
- Some characters in the delimiter string cannot be translated into the character set of the data file.

The hexadecimal digits are converted into bytes, and there is no character set translation performed on the hexadecimal string.

If the end of the file is found before the record terminator, then the access driver proceeds as if a terminator was found, and all unprocessed data up to the end of the file is considered part of the record.



Do not include any binary data, including binary counts for VARCHAR and VARRAW, in a record that has delimiters. Doing so could cause errors or corruption, because the binary data will be interpreted as characters during the search for the delimiter.



Example

The following is an example of using DELIMITED BY records.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth CHAR(4))

ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (RECORDS DELIMITED BY '|' FIELDS TERMINATED BY ','

(first_name CHAR(7),

last_name CHAR(8),

year_of_birth CHAR(4)))

LOCATION ('info.dat'));
```

Alvin, Tolliver, 1976 | Kenneth, Baer, 1963 | Mary, Dube, 1973

Related Topics

string

A string is a quoted series of characters or hexadecimal digits.

15.3.5 XMLTAG

Use the <code>record_format_info</code> XMLTAG clause to specify XML tags that are used to load subdocuments from an XML document.

Default

None

Purpose

You can use the XMLTAG clause of the ORACLE_LOADER access driver to specify XML tags that are used to load subdocuments from an XML document. The access driver searches the data file for documents enclosed by the tags you identify with the clause, and loads those documents as separate rows in the external table.

The XMLTAG clause accepts a list of one or more strings. The strings are used to build tags that ORACLE_LOADER uses to search for subdocuments in the data file. The tags specified in the access parameters do not include the "<" and ">" delimiters.

The <code>ORACLE_LOADER</code> access driver starts at the beginning of the file, and looks for the first occurrence of any of the tags listed in the <code>XMLTAG</code> clause. When it finds a match, it then searches for the corresponding closing tag. For example, if the tag is <code>"ORDER_ITEM"</code>, then <code>ORACLE_LOADER</code> looks for the text string "<code><ORDER_ITEM>"</code>, starting at the beginning of the file. When it finds an occurrence of "<code><ORDER_ITEM>"</code> it then looks for "<code></ORDER_ITEM>"</code>. Everything found between the <code><ORDER_ITEM></code> and <code></ORDER_ITEM></code> tags is part of the document loaded for the row. <code>ORACLE_LOADER</code> then searches for the next occurrence of any of the tags, starting from the first character after the closing tag.

The <code>ORACLE_LOADER</code> access driver is not parsing the XML document to the elements that match the tag names; it is only doing a string search through a text file. If the external table is being accessed in parallel, then <code>ORACLE_LOADER</code> splits large files up so that different sections are read independently. When it starts reading a section of the data file, it starts looking for one of the tags specified by <code>XMLTAG</code>. If it reaches the end of a section and is still looking for a matching end tag, then <code>ORACLE_LOADER</code> continues reading into the next section until the matching end tag is found.

Restrictions When Using XMLTAG

 The XMLTAG clause cannot be used to load data files that have elements nested inside of documents of the same element. For example, if a data file being loaded with XMLTAG ('FOO') contains the following data:

```
<F00><BAR><F00></F00></BAR></F00>
```

then <code>ORACLE_LOADER</code> extracts everything between the first <code><FOO></code> and the first <code></FOO></code> as a document, which does not constitute a valid document.

Similarly, if XMLTAG ("FOO", "BAR") is specified and the data file contains the following:

```
<FOO><BAR></BAR></FOO>
```

then <BAR> and </BAR> are loaded, but as the document for "FOO".

The limit on how large an extracted sub-document can be is determined by the READSIZE
access parameter. If the ORACLE_LOADER access driver sees a subdocument larger than
READSIZE, then it returns an error.

Example Use of the XMLTAG Clause

Suppose you create an external table ${\mathbb T}$ XT as follows:

```
CREATE TABLE "T XT"
  "C0" VARCHAR2 (2000)
ORGANIZATION external
 TYPE oracle loader
 DEFAULT DIRECTORY DMPDIR
 ACCESS PARAMETERS
   RECORDS
   XMLTAG ("home address", "work address", " home phone ")
    READSIZE 1024
    SKIP 0
   FIELDS NOTRIM
   MISSING FIELD VALUES ARE NULL
      "CO" (1:2000) CHAR(2000)
    )
 location
    't.dat'
) REJECT LIMIT UNLIMITED
exit;
```



Assume the contents of the data file are as follows:

<first name>Lionel</first name><home address>23 Oak St, Tripoli, CT</home
address><last name>Rice</last name>

You could then perform the following SQL query:

15.3.6 CHARACTERSET

Use the record format info CHARACTERSET clause to specify the character set of the data file.

Default

None.

Purpose

The CHARACTERSET string clause identifies the character set of the data file. If a character set is not specified, then the data is assumed to be in the default character set for the database.



The settings of NLS environment variables on the client have no effect on the character set used for the database.

Related Topics

- string
 - A string is a quoted series of characters or hexadecimal digits.
- Oracle Database Globalization Support Guide

15.3.7 PREPROCESSOR

To specify your own preprocessor program that you want to run for every data file, use the record format info PREPROCESSOR clause.

Default

None.

Purpose



Caution:

There are security implications to consider when using the PREPROCESSOR clause.

If the file you want to load contains data records that are not in a format supported by the <code>ORACLE_LOADER</code> access driver, then use the <code>PREPROCESSOR</code> clause to specify a user-supplied preprocessor program that will execute for every data file. Note that the program specification must be enclosed in a shell script if it uses arguments (see the description of <code>file spec</code>).

The preprocessor program converts the data to a record format supported by the access driver and then writes the converted record data to standard output (stdout), which the access driver reads as input.

Syntax

The syntax of the PREPROCESSOR clause is as follows:



directory_spec

Specifies the directory object containing the name of the preprocessor program to execute for every data file. The user accessing the external table must have the EXECUTE privilege for the directory object that is used. If directory_spec is omitted, then the default directory specified for the external table is used.



Caution:

For security reasons, to store preprocessor programs, Oracle strongly recommends that you use a separate directory. Do not use the default directory. Do not store any other files in the directory in which preprocessor programs are stored.

To maintain security, the preprocessor program must reside in a directory object, so that access to it can be controlled . Your operating system administrator must create a directory corresponding to the directory object, and and must verify that the operating system Oracle user for the database has access to that directory. Database administrators then must ensure that only approved users are granted permissions to the directory object associated with the directory path. Although multiple database users can have access to a directory object, only those with the EXECUTE privilege can run a preprocessor in that directory. No existing database user with read-write privileges to a directory object will be able to use the preprocessing feature. As a DBA, you can prevent preprocessors from ever being used by never granting the EXECUTE privilege to anyone for a directory object. Refer to *Oracle Database SQL Language Reference* for information about how to grant the EXECUTE privilege.



file_spec

The name of the preprocessor program. It is appended to the path name associated with the directory object that is being used (either the directory_spec or the default directory for the external table). The file spec cannot contain an absolute or relative directory path.

If the preprocessor program requires any arguments (for example, <code>gunzip -c</code>), then you must specify the program name and its arguments in an executable shell script (or on Microsoft Windows operating systems, in a batch (.bat) file). Shell scripts and batch files have certain requirements, as discussed in the following sections.

It is important to verify that the correct version of the preprocessor program is in the operating system directory.

The following is an example of specifying the PREPROCESSOR clause without using a shell or batch file:

```
SQL> CREATE TABLE xtab (recno varchar2(2000))
2   ORGANIZATION EXTERNAL (
3   TYPE ORACLE_LOADER
4   DEFAULT DIRECTORY data_dir
5   ACCESS PARAMETERS (
6   RECORDS DELIMITED BY NEWLINE
7   PREPROCESSOR execdir:'zcat'
8   FIELDS (recno char(2000)))
9   LOCATION ('foo.dat.gz'))
10   REJECT LIMIT UNLIMITED;
Table created.
```

Using Shell Scripts With the PREPROCESSOR Clause on Linux Operating Systems

To use shell scripts on Linux, the following conditions must be true:

- The shell script must reside in directory spec.
- The full path name must be specified for system commands such as gunzip.
- The preprocessor shell script must have EXECUTE permissions.
- The data file listed in the external table LOCATION clause should be referred to by \$1.

The following example shows how to specify a shell script on the PREPROCESSOR clause when creating an external table.

```
SQL> CREATE TABLE xtab (recno varchar2(2000))

2  ORGANIZATION EXTERNAL (
3  TYPE ORACLE_LOADER
4  DEFAULT DIRECTORY data_dir
5  ACCESS PARAMETERS (
6  RECORDS DELIMITED BY NEWLINE
7  PREPROCESSOR execdir:'uncompress.sh'
8  FIELDS (recno char(2000)))
9  LOCATION ('foo.dat.gz'))
10  REJECT LIMIT UNLIMITED;
Table created.
```



Using Batch Files With The PREPROCESSOR Clause on Windows Operating Systems

To use shell scripts on Microsoft Windows, the following conditions must be true:

- The batch file must reside in directory spec.
- The full path name must be specified for system commands such as gunzip.
- The preprocessor batch file must have EXECUTE permissions.
- The first line of the batch file should contain @echo off. The reason for this requirement is
 that when the batch file is run, the default is to display the commands being executed,
 which has the unintended side-effect of the echoed commands being treated as input to
 the external table access driver.
- To represent the input from the location clause, \$1 should be used. (Note that this differs from Unix and Linux-style shell scripts where the location clause is referenced by \$1.)
- A full path should be specified to any executables in the batch file (sed.exe in the following example). Note also that the MKS Toolkit may not exist on all Microsoft Windows installations, so commands such as sed.exe may not be available.

The batch file used on Microsoft Windows must have either a .bat or .cmd extension. Failure to do so (for example, trying to specify the preprocessor script as sed.sh) results in the following error:

```
SQL> select * from foo;
select * from foo
*
ERROR at line 1:

ORA-29913: error in executing ODCIEXTTABLEFETCH callout
ORA-29400: data cartridge error
KUP-04095: preprocessor command
C:/Temp\sed.sh encountered error
"CreateProcess Failure for Preprocessor:
C:/Temp\sed.sh, errorcode: 193
```

The following is a simple example of using a batch file with the external table PREPROCESSOR option on Windows. In this example a batch file uses the stream editor (sed.exe) utility to perform a simple transformation of the input data.

```
SQL> create table deptXT (deptno char(2),
 2 dname char(14),
 3 loc char (13)
 4)
 5 organization external
 6
 7
    type ORACLE LOADER
 8 default directory def dir1
 9 access parameters
10 (
11 records delimited by newline
12 badfile 'deptXT.bad'
13 logfile 'deptXT.log'
14 preprocessor exec dir: 'sed.bat'
15 fields terminated by ','
16 missing field values are null
```

```
17 )
18 location ('deptXT.dat')
19 )
20 reject limit unlimited;
Table created.

select * from deptxt;
Where deptxt.dat contains:
20,RESEARCH,DALLAS
30,SALES,CHICAGO
40,OPERATIONS,BOSTON
51,OPERATIONS,BOSTON
```

The preprocessor program sed.bat has the following content:

```
@echo off
c:/mksnt/mksnt/sed.exe -e 's/BOSTON/CALIFORNIA/' %1
```

The PREPROCESSOR option passes the input data (deptxt.dat) to sed.bat. If you then select from the deptxt table, the results show that the LOC column in the last two rows, which used to be BOSTON. is now CALIFORNIA.

```
SQL> select * from deptxt;

DE DNAME LOC

20 RESEARCH, DALLAS
30 SALES CHICAGO
40 OPERATIONS CALIFORNIA
51 OPERATIONS CALIFORNIA
4 rows selected.
```

Usage Notes for Parallel Processing with the PREPROCESSOR Clause

External tables treat each data file specified on the LOCATION clause as a single granule. To make the best use of parallel processing with the PREPROCESSOR clause, Oracle recommends that the data that you want to load is split into multiple files (granules). Note that external tables limits the degree of parallelism to the number of data files present. For example, if you specify a degree of parallelism of 16, but have only 10 data files, then in effect the degree of parallelism is 10; this is because 10 child processes are busy, and 6 are idle. To process data more efficiently, avoid idle child processes. If you do specify a degree of parallelism, then try to ensure that the degree of parallelism you specify is no larger than the number of data files, so that all child processes are kept busy. Refer to *Oracle Database VLDB and Partitioning Guide* for more information about granules of parallelism.

Also note that you cannot use the same preprocessor script that you use for file system files to process object store data. If you want to use the preprocessor for object store data, then you must write a preprocessor script that can access the object store data, and modify the data.

For example, on Linux or Unix systems, in this case, \$1 represents a source such as https://www.yoururl.example.com/yourdata:

```
@echo off
#!/bin/sh/your_script_or_plsql_function_to_display_objectstore_contents($1) |
sed -e 's/BOSTON/CALIFORNIA/'
```

With this syntax, the preprocessor obtains your data, and sends it to stdout, and pipes it for the access driver to read.

Restrictions When Using the PREPROCESSOR Clause

- The PREPROCESSOR clause is not available on databases that use the Oracle Database Vault feature.
- The PREPROCESSOR clause does not work in conjunction with the COLUMN TRANSFORMS clause.

Related Topics

- Guidelines for Securing the ORACLE_LOADER Access Driver
- Oracle Database SQL Language Reference GRANT

15.3.8 PREPROCESSOR_TIMEOUT

To extend the timeout period for preprocessor programs, use the <code>record_format_infopereprocessor timeout clause</code>.

Default

None.

Purpose

If you encounter a timeout when running your preprocessor, and you think that the preprocessor requires additional time to run, than you can specify a value (in seconds) for PREPROCESSOR_TIMEOUT to wait for your preprocessor to begin producing output to the access driver.

Syntax

The syntax of the PREPROCESSOR_TIMEOUT clause is as follows, where <code>seconds</code> is a numeric value indicating the number of seconds before a timeout is triggered:

```
PREPROCESSOR TIMEOUT seconds
```

Example

The following is a scenario of how you can use the PREPROCESSOR clause with the PREPROCESSOR_TIMEOUT clause to extend the timeout limit for a preprocessor:

Suppose you have a preprocessor whose purpose is to convert data from lowercase to uppercase:

```
#!/bin/sh
    /bin/cat $1 | /bin/tr '[:lower:]''[:upper:]'
```

Next, suppose you have a department data file with the following content:

```
10,accounting,new yorK
20,research,dallas
30,sales,chicago
40,operations,boston
```

Then you create this data file as an external table:

```
SQL> create table deptXT (deptno char(2),
 2 dname char(14),
 3 loc char(13)
 4)
 5 organization external
 6 (
 7 type ORACLE LOADER
 8 default directory def dir1
 9 access parameters
 10 (
 11 records delimited by newline
 12 badfile 'deptXT.bad'
 13 logfile 'deptXT.log'
 14 preprocessor exec_dir:'tr.sh'
 15 fields terminated by ','
16 missing field values are null
17 )
18 location ('deptxt.dat')
19 )
 20 reject limit unlimited;
Table created.
SOL>
SQL> set echo on
SQL> set feedback on
SQL> select * from deptXT ;
DE DNAME
               LOC
__ _____
10 ACCOUNTING NEW YORK
20 RESEARCH DALLAS
30 SALES CHICAGO
40 OPERATIONS BOSTON
4 rows selected.
```

Note that the department name (DNAME) and location (LOC) data are changed from lowercase to uppercase.

Suppose that as you add data to the department table, the script takes longer to process, and you encounter timeout errors. To resolve this issue, you can add PREPROCESSOR_TIMEOUT to the CREATE TABLE statement. In the following example, PREPROCESSOR_TIMEOUT (in bold font) is set to 300 seconds:



```
records delimited by newline
PREPROCESSOR DEF_DIR1:'tr.sh'
PREPROCESSOR_TIMEOUT 300
fields terminated by ','
missing field values are null
)
LOCATION
   (
    'deptxt.dat'
   )
PARALLEL REJECT LIMIT UNLIMITED;
```

15.3.9 EXTERNAL VARIABLE DATA

To load dump files into the Oracle SQL Connector for HDFS that are generated with the ORACLE DATAPUMP access driver, use the EXTERNAL VARIABLE DATA clause.

Default

None.

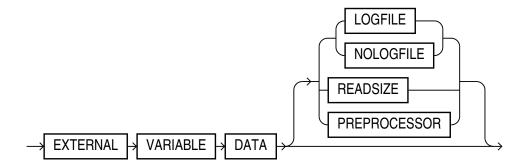
Purpose

When you specify the EXTERNAL VARIABLE DATA clause, the ORACLE_LOADER access driver is used to load dump files that were generated with the ORACLE DATAPUMP access driver.

Note:

The EXTERNAL VARIABLE DATA clause is valid only for use with the Oracle SQL Connector for Hadoop Distributed File System (HDFS). See *Oracle Big Data Connectors User's Guide* for more information about the Oracle SQL Connector for HDFS.

Syntax and Description



You can only use the following access parameters with the EXTERNAL VARIABLE DATA clause:

- LOGFILE | NOLOGFILE
- READSIZE
- PREPROCESSOR



The parameter DISABLE_DIRECTORY_LINK_CHECK is desupported.

Example

In the following example of using the EXTERNAL VARIABLE DATA clause, the following scenario is true:

- The deptxt1.dmp dump file was previously generated by the ORACLE_DATAPUMP access
 driver.
- The thexat program specified by the PREPROCESSOR parameter is a user-supplied program used to manipulate the input data.

```
CREATE TABLE deptxt1
(
    deptno number(2),
    dname varchar2(14),
    loc varchar2(13)
)
ORGANIZATION EXTERNAL
(
    TYPE ORACLE_LOADER
    DEFAULT DIRECTORY dpump_dir
    ACCESS PARAMETERS
    (
        EXTERNAL VARIABLE DATA
        LOGFILE 'deptxt1.log'
        READSIZE=10000
        PREPROCESSOR tkexcat
)
LOCATION ('deptxt1.dmp')
)
REJECT LIMIT UNLIMITED
;
```

Related Topics

LOGFILE | NOLOGFILE

Use the record_format_info LOGFILE clause to name the file that contains messages generated by the external tables utility while it was accessing data in the data file.

READSIZE

The READSIZE parameter specifies the size of the read buffer used to process records.

PREPROCESSOR

To specify your own preprocessor program that you want to run for every data file, use the record format info PREPROCESSOR clause.

15.3.10 LANGUAGE

The LANGUAGE clause allows you to specify a language name (for example, FRENCH), from which locale-sensitive information about the data can be derived.

The following are some examples of the type of information that can be derived from the language name:

- Day and month names and their abbreviations
- Symbols for equivalent expressions for A.M., P.M., A.D., and B.C.
- Default sorting sequence for character data when the ORDER BY SQL clause is specified
- Writing direction (right to left or left to right)
- Affirmative and negative response strings (for example, YES and NO)



Oracle Database Globalization Support Guide for a listing of Oracle-supported languages

15.3.11 TERRITORY

The TERRITORY clause allows you to specify a territory name to further determine input data characteristics.

For example, in some countries a decimal point is used in numbers rather than a comma (for example, 531.298 instead of 531,298).



Oracle Database Globalization Support Guide for a listing of Oracle-supported territories

15.3.12 DATA IS...ENDIAN

The DATA IS...ENDIAN clause indicates the endianness of data whose byte order may vary, depending on the platform that generated the data file.

Purpose

Indicates the endianness of data whose byte order may vary depending on the platform that generated the data file.

Usage Notes

Fields of the following types are affected by this clause:

- INTEGER
- UNSIGNED INTEGER



- FLOAT
- BINARY FLOAT
- DOUBLE
- BINARY DOUBLE
- VARCHAR (numeric count only)
- VARRAW (numeric count only)
- Any character data type in the UTF16 character set
- Any string specified by RECORDS DELIMITED BY string, and in the UTF16 character set

Microsoft Windows-based platforms generate little-endian data. Big-endian platforms include Oracle Solaris and IBM zSeries Based Linux. If the DATA IS...ENDIAN clause is not specified, then the data is assumed to have the same endianness as the platform where the access driver is running. UTF-16 data files can have a mark at the beginning of the file indicating the endianness of the data. If present, then this mark overrides the DATA IS...ENDIAN clause.

15.3.13 BYTEORDERMARK [CHECK | NOCHECK]

Use the record_format_info BYTEORDERMARK clause to specify whether the data file should be checked for the presence of a byte-order mark (BOM).

Default

CHECK

Syntax

BYTEORDERMARK [CHECK | NOCHECK]

Purpose

The BYTEORDERMARK clause is used to specify whether the data file should be checked for the presence of a byte-order mark (BOM). This clause is meaningful only when the character set is Unicode.

BYTEORDERMARK NOCHECK indicates that the data file should not be checked for a BOM and that all the data in the data file should be read as data.

BYTEORDERMARK CHECK indicates that the data file should be checked for a BOM. This is the default behavior for a data file in a Unicode character set.

Usage Notes

The following are examples of some possible scenarios:

If the data is specified as being little or big-endian, and CHECK is specified, and it is
determined that the specified endianness does not match the data file, then an error is
returned. For example, suppose you specify the following:

```
DATA IS LITTLE ENDIAN BYTEORDERMARK CHECK
```

If the BOM is checked in the Unicode data file, and the data is actually big-endian, then an error is returned because you specified little-endian.



- If a BOM is not found, and no endianness is specified with the DATA IS...ENDIAN parameter, then the endianness of the platform is used.
- If BYTE ORDER MARK NOCHECK is specified, and the DATA IS...ENDIAN parameter specified an endianness, then that endian value is used. Otherwise, the endianness of the platform is used.

Related Topics

Understanding how SQL*Loader Manages Byte Ordering
 SQL*Loader can load data from a data file that was created on a system whose byte
 ordering is different from the byte ordering on the system where SQL*Loader is running,
 even if the data file contains certain nonportable data types.

15.3.14 STRING SIZES ARE IN

Use the record_format_info STRING SIZES ARE IN clause to indicate whether the lengths specified for character strings are in bytes or characters.

Default

None.

Syntax

STRING SIZES ARE IN [BYTES | CHARACTERS]

Purpose

The STRING SIZES ARE IN clause is used to indicate whether the lengths specified for character strings are in bytes or characters. If this clause is not specified, then the access driver uses the mode that the database uses. Character types with embedded lengths (such as VARCHAR) are also affected by this clause. If this clause is specified, then the embedded lengths are a character count, not a byte count. Specifying STRING SIZES ARE IN CHARACTERS is needed only when loading multibyte character sets, such as UTF16.

15.3.15 LOAD WHEN

Use the record_format_info LOAD WHEN clause to identify the records that should be passed to the database.

Default

Syntax

The syntax of the LOAD WHEN clause is as follows, where <code>condition_spec</code> are condition specifications:

LOAD WHEN condition spec

Purpose

The LOAD WHEN condition_spec clause is used to identify the records that should be passed to the database. The evaluation method varies:

If the condition_spec references a field in the record, then the clause is evaluated only
after all fields have been parsed from the record, but before any NULLIF or DEFAULTIF
clauses have been evaluated.

If the condition specification references only ranges (and no field names), then the clause
is evaluated before the fields are parsed. This use case is helpful where the records in the
file that you do not want to be loaded cannot be parsed into the current record definition
without errors.

Example

The following is an examples of using LOAD WHEN:

```
LOAD WHEN (empid != BLANKS)
LOAD WHEN ((dept id = "SPORTING GOODS" OR dept id = "SHOES") AND total sales != 0)
```

Related Topics

condition_spec

The <code>condition_spec</code> specifies one or more conditions that are joined by Boolean operators.

15.3.16 BADFILE | NOBADFILE

Use the <code>record_format_info</code> <code>BADFILE</code> clause to name the file to which records are written when they cannot be loaded because of errors.

Default

Create a bad file with default name. See Purpose for details.

Syntax

```
BADFILE name | NOBADFILE
```

Purpose

The BADFILE clause names the file to which records are written when they cannot be loaded because of errors. For example, a record would be written to the bad file if a field in the data file could not be converted to the data type of a column in the external table. The purpose of the bad file is to have one file where all rejected data can be examined and fixed so that it can be loaded. If you do not intend to fix the data, then you can use the NOBADFILE option to prevent creation of a bad file, even if there are bad records.

If you specify the BADFILE clause, then you must supply either a directory object name or file name, or both. See [directory object name:] [filename].

If you specify NOBADFILE, then a bad file is not created.

If neither BADFILE nor NOBADFILE is specified, then the default is to create a bad file if at least one record is rejected. The name of the file is the table name followed by $_{\$p}$, where $_{\$p}$ is replaced with the PID of the process creating the file. The file is given an extension of .bad. If the table name contains any characters that could be interpreted as directory navigation (for example, $_{\%}$, $_{/}$, or $_{/}$), then those characters are not included in the output file name.

Records that fail the LOAD WHEN clause are not written to the bad file, but instead are written to the discard file. Also, any errors in using a record from an external table (such as a constraint violation when using INSERT INTO...AS SELECT... from an external table) will not cause the record to be written to the bad file.



Related Topics

[directory object name:] [filename]

The [directory object name:] [filename] clause is used to specify the name of an output file (BADFILE, DISCARDFILE, or LOGFILE).

15.3.17 DISCARDFILE | NODISCARDFILE

Use the <code>record_format_info</code> <code>DISCARDFILE</code> clause to name the file to which records are written that fail the condition in the <code>LOAD WHEN</code> clause.

Default

Create a discard file with default name. See Purpose for details.

Syntax

DISCARDFILE name | NODISCARDFILE

Purpose

The DISCARDFILE clause names the file to which records are written that fail the condition in the LOAD WHEN clause. The discard file is created when the first record for discard is encountered. If the same external table is accessed multiple times, then the discard file is rewritten each time. If there is no need to save the discarded records in a separate file, then use NODISCARDFILE.

If you specify DISCARDFILE, then you must supply either a directory object name or file name, or both. See [directory object name:] [filename].

If you specify NODISCARDFILE, then a discard file is not created.

If neither DISCARDFILE nor NODISCARDFILE is specified, then the default is to create a discard file if at least one record fails the LOAD WHEN clause. The name of the file is the table name followed by $_{\$p}$, where $_{\$p}$ is replaced with the PID of the process creating the file. The file is given an extension of .dcs. If the table name contains any characters that could be interpreted as directory navigation (for example, $_{\$}$, /, or *), then those characters are not included in the file name.

Related Topics

• [directory object name:] [filename]

The [directory object name:] [filename] clause is used to specify the name of an output file (BADFILE, DISCARDFILE, or LOGFILE).

15.3.18 LOGFILE | NOLOGFILE

Use the record_format_info LOGFILE clause to name the file that contains messages generated by the external tables utility while it was accessing data in the data file.

Default

Use an existing file, or create a log file with default name. See Purpose for details.

Syntax

LOGFILE name | NOLOGFILE



Purpose

The LOGFILE clause names the file that contains messages generated by the external tables utility while it was accessing data in the data file. If a log file already exists by the same name, then the access driver reopens that log file and appends new log information to the end. This is different from bad files and discard files, which overwrite any existing file. The NOLOGFILE clause is used to prevent creation of a log file.

If you specify LOGFILE, then you must supply either a directory object name or file name, or both. See [directory object name:] [filename].

If you specify NOLOGFILE, then a log file is not created.

If neither LOGFILE nor NOLOGFILE is specified, then the default is to create a log file. The name of the file is the table name followed by $_{\$p}$, where $_{\$p}$ is replaced with the PID of the process creating the file. The file is given an extension of .log. If the table name contains any characters that could be interpreted as directory navigation (for example, $_{\$}$, /, or *), then those characters are not included in the file name.

Related Topics

• [directory object name:] [filename]
The [directory object name:] [filename] clause is used to specify the name of an output file (BADFILE, DISCARDFILE, or LOGFILE).

15.3.19 SKIP

Use the record_format_info SKIP clause to skip the specified number of records in the data file before loading.

Default

None (0)

Syntax

The syntax is as follows, where num is the number of records to skip (Default 0).

SKIP = num

Purpose

The SKIP clause skips the specified number of records in the data file before loading. You can specify this clause only when nonparallel access is being made to the data. If there is more than one data file in the same location for the same table, then the SKIP clause causes the ORACLE LOADER driver to skip the specified number of records in the first data file only.

15.3.20 FIELD NAMES

Use the record format info FIELD NAMES clause to specify field order in data files.

Default

NONE

Syntax

FIELD NAMES {FIRST FILE | FIRST IGNORE | ALL FILES | ALL IGNORE | NONE }



The FIELD NAMES options are:

- FIRST FILE Indicates that the first data file contains a list of field names for the data in
 the first record. This list uses the same delimiter as the data in the data file. This record is
 read and used to set up the mapping between the fields in the data file and the columns in
 the target table. This record is skipped when the data is processed. This option can be
 useful if the order of the fields in the data file is different from the order of the columns in
 the table.
- FIRST IGNORE Indicates that the first data file contains a list of field names for the data in the first record, but that the information should be ignored. This record is skipped when the data is processed, but is not used for setting up the fields.
- ALL FILES Indicates that all data files contain a list of field names for the data in the first record. The ordering of the fields in the datafiles can be in any order. The order is specified by the first row in each file, which specifies to the access driver that the fields are in a different order than the columns in the external table.
- ALL IGNORE Indicates that all data files contain a list of field names for the data in the first record, but that the information should be ignored. This record is skipped when the data is processed in every data file, but it is not used for setting up the fields.
- NONE Indicates that the data file contains normal data in the first record. This is the
 default option.

Purpose

Use the FIELD NAMES clause to specify the field order of data files for the first row of the data file using one of the options. For example, if FIELD NAMES FIRST FILE is specified, then only the first data file has the row header. If FIELD NAMES ALL FILES is specified, then all data files will have the row header.

Restrictions

- The FIELD NAMES clause does not trim whitespace between field names in data files.
 - For example, if a data file has field names <code>deptno</code>, <code>dname</code>, <code>loc</code> (with whitespace between field names) then specifying <code>FIELD NAMES</code> can fail with "KUP-04117: Field name LOC was not found in the access parameter field list or table."
- Field names in data files cannot use quotations. For example, the following column field names are not supported:

```
deptno, "dname", loc
```

Embedded delimiters are not supported in the first column header row.

Example

Typically fields in a data file where you want to generate a table with columns (COL1, COL2, COL3) are in the same order in the data file as they will be in the table. However, in the following example, the ordering of data file fields is diffferent in deptxt1.dat and deptxt2.dat. Specifying FIELD NAMES ALL FILES enables data fields in differing field name order in one or more datafiles to be queried correctly:

[admin@example]\$ cat /tmp/deptxt1.dat deptno,dname,loc
10,ACCOUNTING,NEW YORK
20,RESEARCH,DALLAS
30,SALES,CHICAGO



```
40, OPERATIONS, BOSTON
[admin@example]$ cat /tmp/deptxt2.dat
dNamE, 10c, DEPTNO
ACCOUNTING, NEW YORK, 11
RESEARCH, DALLAS, 21
SALES, CHICAGO, 31
OPERATIONS, BOSTON, 41
[admin@example] $ sql @xt
Connected.
Directory created.
SQL> create table deptXT
  2 (
  3
        deptno number(2),
        dname varchar2(14),
        loc varchar2(13)
  6 )
  7 organization external
  8 (
  9
     type ORACLE LOADER
 10 DEFAULT DIRECTORY DATA DIR
 11 access parameters
 12
 13
      records delimited by newline
 14
        field names all files
        logfile 'deptxt.log'
 15
 16
        badfile 'deptxt.bad'
 17
        fields terminated by ','
        missing field values are null
 19
 20
     location ('deptxt?.dat')
 21 )
 22 reject limit unlimited
 23 ;
Table created.
SQL> Rem returns all 8 rows
SQL> select deptno, dname, loc from deptxt order by deptno;
DEPTNO
        DNAME
                      LOC
-----
    10
        ACCOUNTING NEW YORK
    11 ACCOUNTING NEW YORK
       RESEARCH
                     DALLAS
    20
    21 RESEARCH
                      DALLAS
    30 SALES
                      CHICAGO
    31 SALES
                      CHICAGO
    40
         OPERATIONS
                      BOSTON
    41 OPERATIONS
                      BOSTON
```

15.3.21 READSIZE

The READSIZE parameter specifies the size of the read buffer used to process records.

The size of the read buffer must be at least as big as the largest input record the access driver will encounter. The size is specified with an integer indicating the number of bytes. The default value is 512 KB (524288 bytes). You must specify a larger value if any of the records in the data file are larger than 512 KB. There is no limit on how large READSIZE can be, but practically, it is limited by the largest amount of memory that can be allocated by the access driver.

The amount of memory available for allocation is another limit because additional buffers might be allocated. The additional buffer is used to correctly complete the processing of any records that may have been split (either in the data; at the delimiter; or if multi character/byte delimiters are used, in the delimiter itself).

15.3.22 DATE_CACHE

By default, the date cache feature is enabled (for 1000 elements). To completely disable the date cache feature, set it to 0.

DATE_CACHE specifies the date cache size (in entries). For example, DATE_CACHE=5000 specifies that each date cache created can contain a maximum of 5000 unique date entries. Every table has its own date cache, if one is needed. A date cache is created only if at least one date or timestamp value is loaded that requires data type conversion in order to be stored in the table.

The date cache feature is enabled by default. The default date cache size is 1000 elements. If the default size is used and the number of unique input values loaded exceeds 1000, then the date cache feature is automatically disabled for that table. However, if you override the default and specify a nonzero date cache size and that size is exceeded, then the cache is not disabled.

You can use the date cache statistics (entries, hits, and misses) contained in the log file to tune the size of the cache for future similar loads.



Specifying a Value for the Date Cache

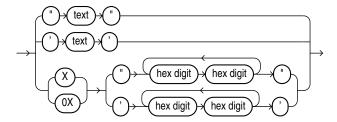
15.3.23 string

A string is a quoted series of characters or hexadecimal digits.

Syntax

The syntax for a string is as follows:





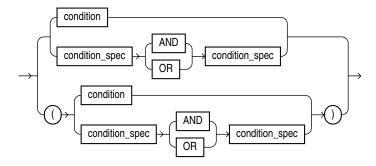
Purpose

If it is a series of characters, then those characters will be converted into the character set of the data file. If it is a series of hexadecimal digits, then there must be an even number of hexadecimal digits. The hexadecimal digits are converted into their binary translation, and the translation is treated as a character string in the character set of the data file. This means that once the hexadecimal digits have been converted into their binary translation, there is no other character set translation that occurs.

15.3.24 condition_spec

The condition spec specifies one or more conditions that are joined by Boolean operators.

This clause is an expression that evaluates to either true or false. The conditions and Boolean operators are evaluated from left to right. (Boolean operators are applied after the conditions are evaluated.) To override the default order of evaluation of Boolean operators, you can use parentheses. The evaluation of condition_spec clauses slows record processing, so these clauses should be used sparingly. The syntax for condition_spec is as follows:



Note that if the condition specification contains any conditions that reference field names, then the condition specifications are evaluated only after all fields have been found in the record, and after blank trimming has been done. It is not useful to compare a field to <code>BLANKS</code> if blanks have been trimmed from the field.

The following are some examples of using <code>condition_spec:</code>

```
empid = BLANKS OR last_name = BLANKS
(dept id = SPORTING GOODS OR dept id = SHOES) AND total sales != 0
```





15.3.25 [directory object name:] [filename]

The [directory object name:] [filename] clause is used to specify the name of an output file (BADFILE, DISCARDFILE, or LOGFILE).

Syntax

[directory object name:] [filename]

- directory object name: The alias for the operating system directory on the database server for reading and writing files.
- filename: The name of the file that you want to create in the directory object.

To help make file names unique in parallel loads, the access driver does some symbol substitution. The symbol substitutions supported for the Linux, Unix, and Microsoft Windows operating systems are as follows (other platforms can have different symbols):

- %p is replaced by the process ID of the current process.
 - For example, if the process ID of the access driver is 12345, then a filename specified as exttab %p.log becomes exttab 12345.log.
- %a is replaced by the agent number of the current process. The agent number is the unique number assigned to each parallel process accessing the external table. This number is padded to the left with zeros to fill three characters.
 - For example, if the third parallel agent is creating a file and you specify bad_data_%a.bad as the file name, then the agent creates a file named bad data 003.bad.
- % is replaced by %. If there is a need to have a percent sign in the file name, then this symbol substitution is used.

If the % character is encountered followed by anything other than one of the preceding characters, then an error is returned.

Purpose

Specifies the name of an output file (BADFILE, DISCARDFILE, or LOGFILE).

Usage Notes

To use this clause, you must supply either a directory object name or file name, or both. The directory object name is the name of a directory object where the user accessing the external table has privileges to write. If the directory object name is omitted, then the value specified for the DEFAULT DIRECTORY clause in the CREATE TABLE...ORGANIZATION EXTERNAL statement is used.

If p or a is not used to create unique file names for output files, and an external table is being accessed in parallel, then it is possible that output files can be corrupted, or that agents may be unable to write to the files.



If you do not specify BADFILE (or DISCARDFILE or LOGFILE), then the SQL_LOADER access driver uses the name of the table, followed by _%p as the name of the file. If no extension is supplied for the file, then a default extension is used. For bad files, the default extension is .bad; for discard files, the default is .dsc; and for log files, the default is .log.

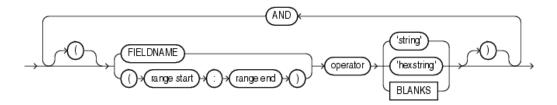
15.3.26 condition

To compare a range of bytes or a field from the record against a constant string, you can use the <code>ORACLE LOADER</code> condition clause

Purpose

Compares a range of bytes or a field from the record against a constant string. The source of the comparison can be either a field in the record, or a byte range in the record. The comparison is done on a byte-by-byte basis. If a string is specified as the target of the comparison, then it is translated into the character set of the data file. If the field has a noncharacter data type, then no data type conversion is performed on either the field value, or the string.

Syntax



range start : range end

The (range start: range end) clause of condition describes a range of bytes or characters in the record, which you want to use for a condition.

15.3.26.1 range start : range end

The (range start: range end) clause of condition describes a range of bytes or characters in the record, which you want to use for a condition.

Purpose

Describes a range of bytes or characters in the record that you want to want to use to create a condition.

Syntax

(range start:range end)

- range start: The starting byte or character offsets into the record.
- range end: The ending byte or character offsets into the record.

Usage Notes

The value that you enter for the STRING SIZES ARE clause determines whether the range refers to bytes, or refers to characters.

The value that you provide for <code>range start</code> must be less than or equal to the value for <code>range end</code>. Finding ranges of characters is faster for data in fixed-width character sets than it is for data in varying-width character sets. If the range refers to parts of the record that do not exist, then the record is rejected when an attempt is made to reference the range. The <code>range start:range end</code> clause must be enclosed in parentheses. For example: (10:13).



In your data file, Oracle recommends that you do not mix binary data (including data types with binary counts, such as VARCHAR) and character data that is in a varying-width character set, or more than one byte wide. When binary and character data with these characteristics are mixed, it is possible that the access driver may not find the correct start for the field, because it treats the binary data as character data when trying to find the start.

The following is an example of using condition with a range clause:

```
LOAD WHEN empid != BLANKS
LOAD WHEN (10:13) = 0x'00000830'
LOAD WHEN PRODUCT COUNT = "MISSING"
```

15.3.27 IO_OPTIONS clause

To specify whether the operating system uses direct input/output to read data files from disk, or uses a cache for reading the data files, use the <code>ORACLE LOADER records clause IO OPTIONS</code>.

Default

If not otherwise specified, then the default IO OPTIONS setting is DIRECTIO.

Purpose

Enables you to specify the input and output (I/O) options that the operating system uses for reading the data files, either by reading files directly from storage, or by reading data files from cache. The only options available for specification are DIRECTIO (the default), and NODIRECTIO.

Syntax

io options (directio|nodirectio)

Usage Notes

When set to DIRECTIO, an attempt is made to open the data file and read it directly from storage. If successful, then the operating system and NFS server (if the file is on an NFS server) do not cache the data read from the file. Accessing data without cacheing it can improve the read performance for the data file, especially if the file is large. If direct I/O is not supported for the data file being read, then the file is opened and read, but the DIRECTIO option is ignored.

If the IO_OPTIONS clause is specified with the NODIRECTIO option, then direct I/O is not used to read the data files, and instead Oracle Database reads files from the operating system cache.

If the IO OPTIONS clause is not specified at all, then the default DIRECTIO option is used.

The following is an example of specifying that the operating system should use direct input/output writes to storage:

```
(
records delimited by newline io_options (directio)
logfile
.
.
.
```

Related Topics

When to Separate Files

15.3.28 DNFS_DISABLE | DNFS_ENABLE

To disable and enable use of the Direct NFS Client on input data files during an external tables operation, use DNFS DISABLE or DNFS ENABLE.

Purpose

Use these parameters to enable and disable use of the Direct NFS Client on input data files during an external tables operation.

Usage Notes

The Direct NFS Client is an API that can be implemented by file servers to enable improved performance when Oracle Database accesses files on those servers.

By default, external tables use the Direct NFS Client interfaces when they read data files over 1 gigabyte in size. For smaller files, the operating system I/O interfaces are used. To use the Direct NFS Client on all input data files, specify DNFS ENABLE.

To disable use of the Direct NFS Client for all data files, specify DNFS DISABLE.

15.3.29 DNFS_READBUFFERS

The DNFS_READBUFFERS parameter of the record_format_info clause is used to control the number of read buffers used by the Direct NFS Client.

Default

The default value for DNFS READBUFFERS is 4.

Purpose

Controls the number of read buffers used by the Direct NFS Client.

The Direct NFS Client is an API that can be implemented by file servers to allow improved performance when Oracle accesses files on those servers.

Usage Notes

It is possible that using larger values for <code>DNFS_READBUFFERS</code> can compensate for inconsistent input and output from the Direct NFS Client file server. However, using larger values can result in increased memory usage.



15.4 field_definitions Clause

Learn how to name the fields in the data file and specify how to find them in records using the field_definitions clause.

Overview of field definitions Clause

In the field_definitions clause, you use the FIELDS parameter to name the fields in the data file, and specify how to find fields in records.

delim_spec

The delim_spec clause is used to find the end (and if ENCLOSED BY is specified, the start) of a field.

trim_spec

The trim_spec clause is used to specify that spaces should be trimmed from the beginning of a text field, the end of a text field, or both.

MISSING FIELD VALUES ARE NULL

The effect of MISSING FIELD VALUES ARE NULL depends on whether POSITION is used to explicitly state field positions.

field list

The field_definitions field_list clause identifies the fields in the data file and their data types.

pos_spec Clause

The ORACLE LOADER pos spec clause indicates the position of the column within the record.

datatype spec Clause

The <code>ORACLE_LOADER</code> datatype_spec clause describes the data type of a field in the data file if the data type is different than the default.

init_spec Clause

The <code>init_spec</code> clause for external tables is used to specify when a field should be set to <code>NULL</code>, or when it should be set to a default value.

LLS Clause

If a field in a data file is a LOB location Specifier (LLS) field, then you can indicate this by using the LLS clause.

15.4.1 Overview of field_definitions Clause

In the field_definitions clause, you use the FIELDS parameter to name the fields in the data file, and specify how to find fields in records.

Default

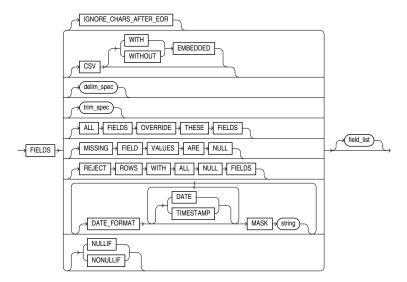
If the field definitions clause is omitted, then the following is assumed:

- The fields are delimited by ','
- The fields are of data type CHAR
- The maximum length of the field is 255
- The order of the fields in the data file is the order in which the fields were defined in the external table
- No blanks are trimmed from the field



Syntax

The syntax for the field definitions clause is as follows:



Example 15-1 External Table Created Without Access Parameters (Default)

In this example, an external table is created without any access parameters. It is followed by a sample data file, info.dat, that can be used to load the table.

CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth CHAR(4)) ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir LOCATION ('info.dat'));

Alvin, Tolliver, 1976 Kenneth, Baer, 1963

Parameters to Specify Fields with field_definition

The sections that follow provide an overview of the field definitions that you can specify with the field definition clause, and some examples of how to use these clauses.

IGNORE_CHARS_AFTER_EOR

This optional parameter specifies that if extraneous characters are found after the **last** end-of-record, **but before the end of the file** that do not satisfy the record definition, then they are ignored.

Error messages are written to the external tables log file if all four of the following conditions apply:

- The IGNORE_CHARS_AFTER_EOR parameter is set, or the field allows free formatting. (Free
 formatting means either that the field is variable length, or the field is specified by a
 delimiter or enclosure characters, and is also variable length).
- Characters remain after the last end-of-record in the file.
- The access parameter MISSING FIELD VALUES ARE NULL is not set.



The field does not have absolute positioning.

The error messages that are written to the external tables log file are as follows:

```
KUP-04021: field formatting error for field Coll
KUP-04023: field start is after end of record
KUP-04101: record 2 rejected in file /home/oracle/datafiles/example.dat
```

CSV

To direct external tables to access the data files as comma-separated-values format files, use the <code>FIELDS CSV</code> clause. To use this clause, the file should be a stream record format file with the normal carriage return string (for example, \n on Unix or Linux operating systems, and either \n or \n on Microsoft Windows operating systems). Record terminators can be included (embedded) in data values. The syntax for the <code>FIELDS CSV</code> clause is as follows:

```
FIELDS CSV [WITH EMBEDDED | WITHOUT EMBEDDED] [TERMINATED BY ','] [OPTIONALLY ENCLOSED BY '"']
```

When using the FIELDS CSV clause, note the following:

- The default is to not use the FIELDS CSV clause.
- The with embedded and without embedded options specify whether record terminators are included (embedded) in the data. The with embedded option is the default.
- If WITH EMBEDDED is used, then embedded record terminators must be enclosed, and intradatafile parallelism is disabled for external table loads.
- The TERMINATED BY ',' and OPTIONALLY ENCLOSED BY '"' options are the defaults. They do not have to be specified. You can override them with different termination and enclosure characters.
- When the CSV clause is used, a delimiter specification is not allowed at the field level and only delimitable data types are allowed. Delimitable data types include CHAR, datetime, interval, and numeric EXTERNAL.
- The TERMINATED BY and ENCLOSED BY clauses cannot be used at the field level when the CSV clause is specified.
- When the CSV clause is specified, the default trimming behavior is LDRTRIM. You can override this default by specifying one of the other external table trim options (NOTRIM, LRTRIM, LTRIM, OT RTRIM).
- The CSV clause must be specified after the IGNORE_CHARS_AFTER_EOR clause, and before the delim spec clause.

delim spec Clause

The delim_spec clause is used to identify how all fields are terminated in the record. The delim_spec specified for all fields can be overridden for a particular field as part of the field_list clause. For a full description of the syntax, refer to the delim_spec clause description.

trim_spec Clause

The trim_spec clause specifies the type of whitespace trimming to be performed by default on all character fields. The trim spec clause specified for all fields can be overridden for



individual fields by specifying a $trim_spec$ clause for those fields. For a full description of the syntax, refer to the $trim_spec$ clause description.

ALL FIELDS OVERRIDE

The ALL FIELDS OVERRIDE clause specifies to the access driver that all fields are present, and that they are in the same order as the columns in the external table. You only need to specify fields that have a special definition. This clause must be specified after the optional trim_spec clause, and before the optional MISSING FIELD VALUES ARE NULL clause.

The following is a sample use of the ALL FIELDS OVERRIDE clause. The only field in this example that requires specification is HIREDATE, which requires data format mask. All the other fields take default values.

```
FIELDS TERMINATED BY "," OPTIONALLY ENCLOSED BY '"' LDRTRIM ALL FIELDS OVERRIDE
REJECT ROWS WITH ALL NULL FIELDS
(
HIREDATE CHAR(20) DATE_FORMAT DATE MASK "DD-Month-YYYY"
)
```

MISSING FIELD VALUES ARE NULL

MISSING FIELD VALUES ARE NULL sets to null any fields for which position is not explicitly stated and there is not enough data to fill them. For a full description the description for MISSING FIELD VALUES ARE NULL.

REJECT ROWS WITH ALL NULL FIELDS

REJECT ROWS WITH ALL NULL FIELDS indicates that a row will not be loaded into the external table if all referenced fields in the row are null. If this parameter is not specified, then the default value is to accept rows with all null fields. The setting of this parameter is written to the log file either as "reject rows with all null fields" or as "rows with all null fields are accepted."

DATE FORMAT

The DATE_FORMAT clause enables you to specify a datetime format mask once at the fields level, and then have that format apply to all fields of that type that do not have their own mask specified. The datetime format mask must be specified after the optional REJECT ROWS WITH ALL NULL FIELDS clause, and before the fields list clause.

The DATE FORMAT can be specified for the following datetime types:

- DATE
- TIME
- TIME
- WITH TIME ZONE
- TIMESTAMP
- TIMESTAMP WITH TIME ZONE



The following example shows a sample use of the DATE_FORMAT clause that applies a date mask of DD-Month-YYYY to any DATE type fields:

```
FIELDS TERMINATED BY "," OPTIONALLY ENCLOSED BY '"' LDRTRIM
REJECT ROWS WITH ALL NULL FIELDS

DATE_FORMAT DATE MASK "DD-Month-YYYY"

(

EMPNO,
ENAME,
JOB,
MGR,
HIREDATE CHAR(20),
SAL,
COMM,
DEPTNO,
PROJNO,
ENTRYDATE CHAR(20)
)
```

NULLIF | NO NULLIF

The NULLIF clause applies to all character fields (for example, CHAR, VARCHAR, VARCHARC, external NUMBER, and datetime).

The syntax is as follows:

```
NULLIF {=|!=}{"char_string"|x'hex_string'|BLANKS}
```

If there is a match using the equal or not equal specification for a field, then the field is set to NULL for that row.

The char string and hex string must be enclosed in single- or double-quotation marks.

If a NULLIF specification is specified at the field level, then it overrides this NULLIF clause.

If there is a field to which you do not want the \mathtt{NULLIF} clause to apply, then you can specify \mathtt{NO} \mathtt{NULLIF} at the field level.

The NULLIF clause must be specified after the optional REJECT ROWS WITH ALL NULL FIELDS clause and before the fields list clause.

The following is an example of using the <code>NULLIF</code> clause in which you specify a field to which you do not want the <code>NULLIF</code> clause to apply. The <code>MGR</code> field is set to <code>NO NULLIF</code>, which means that the <code>NULLIF="NONE"</code> clause does not apply to that field.

```
FIELDS TERMINATED BY "," OPTIONALLY ENCLOSED BY '"' LDRTRIM REJECT ROWS WITH ALL NULL FIELDS
NULLIF = "NONE"
(

EMPNO,
ENAME,
JOB,
MGR
```



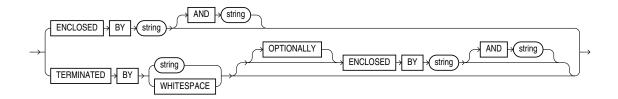
field list Clause

The field_list clause identifies the fields in the data file and their data types. For a full description of the syntax, see the description of the field list clause.

15.4.2 delim_spec

The delim_spec clause is used to find the end (and if ENCLOSED BY is specified, the start) of a field.

Syntax



Usage Notes

If you specify ENCLOSED BY, then the ORACLE_LOADER access driver starts at the current position in the record, and skips over all whitespace looking for the first delimiter. All whitespace between the current position and the first delimiter is ignored. Next, the access driver looks for the second enclosure delimiter (or looks for the first one again if a second one is not specified). Everything between those two delimiters is considered part of the field.

If TERMINATED BY *string* is specified with the ENCLOSED BY clause, then the terminator string must immediately follow the second enclosure delimiter. Any whitespace between the second enclosure delimiter and the terminating delimiter is skipped. If anything other than whitespace is found between the two delimiters, then the row is rejected for being incorrectly formatted.

If TERMINATED BY is specified without the ENCLOSED BY clause, then everything between the current position in the record and the next occurrence of the termination string is considered part of the field.

If OPTIONALLY is specified, then TERMINATED BY must also be specified. The OPTIONALLY parameter means the ENCLOSED BY delimiters can either both be present or both be absent. The terminating delimiter must be present, regardless of whether the ENCLOSED BY delimiters are present. If OPTIONALLY is specified, then the access driver skips over all whitespace, looking for the first non-blank character. After the first non-blank character is found, the access driver checks to see if the current position contains the first enclosure delimiter. If it does, then the access driver finds the second enclosure string. Everything between the first and second enclosure delimiters is considered part of the field. The terminating delimiter must immediately follow the second enclosure delimiter (with optional whitespace allowed between the second enclosure delimiter and the terminating delimiter). If the first enclosure string is not found at the first non-blank character, then the access driver looks for the terminating delimiter. In this case, leading blanks are trimmed.

After the delimiters have been found, the current position in the record is set to the spot after the last delimiter for the field. If TERMINATED BY WHITESPACE was specified, then the current position in the record is set to after all whitespace following the field.

To find out more about the access driver's default trimming behavior, refer to "Trimming Whitespace." You can override this behavior by using with LTRIM and RTRIM.

A missing terminator for the last field in the record is not an error. The access driver proceeds as if the terminator was found. It is an error if the second enclosure delimiter is missing.

The string used for the second enclosure can be included in the data field by including the second enclosure twice. For example, if a field is enclosed by single quotation marks, then it could contain a single quotation mark by specifying two single quotation marks in a row, as shown in the word don't in the following example:

```
'I don''t like green eggs and ham'
```

There is no way to quote a terminator string in the field data without using enclosing delimiters. Because the field parser does not look for the terminating delimiter until after it has found the enclosing delimiters, the field can contain the terminating delimiter.

In general, specifying single characters for the strings is faster than multiple characters. Also, searching data in fixed-width character sets is usually faster than searching data in varying-width character sets.



The use of the backslash character (\) within strings is not supported in external tables

- Example: External Table with Terminating Delimiters
 See how to create an external table that uses terminating delimiters, and a data file with terminating delimiters.
- Example: External Table with Enclosure and Terminator Delimiters

 See how to create an external table that uses both enclosure and terminator delimiters.
- Example: External Table with Optional Enclosure Delimiters
 See how to create an external table that uses optional enclosure delimiters.

Related Topics

Trimming Whitespace
 Blanks, tabs, and other nonprinting characters (such as carriage returns and line feeds)
 constitute whitespace.

15.4.2.1 Example: External Table with Terminating Delimiters

See how to create an external table that uses terminating delimiters, and a data file with terminating delimiters.

This table is created to use terminating delimiters. It is followed by an example of a data file that can be used to load the table.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth CHAR(4))
ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (FIELDS TERMINATED BY WHITESPACE)
LOCATION ('info.dat'));
```

Alvin Tolliver 1976 Kenneth Baer 1963 Mary Dube 1973



15.4.2.2 Example: External Table with Enclosure and Terminator Delimiters

See how to create an external table that uses both enclosure and terminator delimiters.

The following is an example of an external table that uses both enclosure and terminator delimiters. Remember that all whitespace between a terminating string and the first enclosure string is ignored, as is all whitespace between a second enclosing delimiter and the terminator. The example is followed by a sample of the data file that can be used to load it.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth CHAR(4))
ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (FIELDS TERMINATED BY "," ENCLOSED BY "(" AND ")")
LOCATION ('info.dat'));

(Alvin) , (Tolliver), (1976)
(Kenneth), (Baer) , (1963)
(Mary), (Dube) , (1973)
```

15.4.2.3 Example: External Table with Optional Enclosure Delimiters

See how to create an external table that uses optional enclosure delimiters.

This table is an external table that is created to use optional enclosure delimiters. Note that LRTRIM is used to trim leading and trailing blanks from fields. The example is followed by an example of a data file that can be used to load the table.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth CHAR(4))

ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (FIELDS TERMINATED BY ','

OPTIONALLY ENCLOSED BY '(' and ')'

LRTRIM)

LOCATION ('info.dat'));

Alvin , Tolliver , 1976
(Kenneth), (Baer), (1963)
(Mary ), Dube , (1973)
```

15.4.3 trim_spec

The trim_spec clause is used to specify that spaces should be trimmed from the beginning of a text field, the end of a text field, or both.

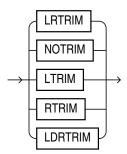
Description

Directs the <code>ORACLE_LOADER</code> access driver to trim spaces from the beginning of a text field, the end of a text field, or both. Spaces include blanks and other non-printing characters, such as tabs, line feeds, and carriage returns.

Default

The default is LDRTRIM. Specifying NOTRIM yields the fastest performance.

Syntax



Options

- NOTRIM Indicates that you want no characters trimmed from the field.
- LRTRIM Indicates that you want both leading and trailing spaces trimmed.
- LTRIM Indicates that you want leading spaces trimmed.
- RTRIM Indicates that you want trailing spaces trimmed.
- LDRTRIM Provides compatibility with SQL*Loader trim features. It is the same as NOTRIM except in the following cases:
 - If the field is not a delimited field, then spaces will be trimmed from the right.
 - If the field is a delimited field with OPTIONALLY ENCLOSED BY specified, and the optional
 enclosures are missing for a particular instance, then spaces are trimmed from the left.

Usage Notes

The trim_spec clause can be specified before the field list to set the default trimming for all fields. If trim_spec is omitted before the field list, then LDRTRIM is the default trim setting. The default trimming can be overridden for an individual field as part of the datatype spec.

If trimming is specified for a field that is all spaces, then the field will be set to NULL.

In the following example, all data is fixed-length; however, the character data will not be loaded with leading spaces. The example is followed by a sample of the data file that can be used to load it.



15.4.4 MISSING FIELD VALUES ARE NULL

The effect of MISSING FIELD VALUES ARE NULL depends on whether POSITION is used to explicitly state field positions.

For example:

- The default behavior is that if field position is not explicitly stated and there is not enough
 data in a record for all fields, then the record is rejected. You can override this behavior by
 using MISSING FIELD VALUES ARE NULL to define as NULL any fields for which there is no
 data available.
- If field position is explicitly stated, then fields for which there are no values are always defined as NULL, regardless of whether MISSING FIELD VALUES ARE NULL is used.

In the following example, the second record is stored with a <code>NULL</code> set for the <code>year_of_birth</code> column, even though the data for the <code>year</code> of birth is missing from the data file. If the <code>MISSINGFIELD VALUES ARE NULL</code> clause were omitted from the access parameters, then the second row would be rejected because it did not have a value for the <code>year_of_birth</code> column. The example is followed by a sample of the data file that can be used to load it.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth INT)
ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (FIELDS TERMINATED BY ","

MISSING FIELD VALUES ARE NULL)
LOCATION ('info.dat'));
```

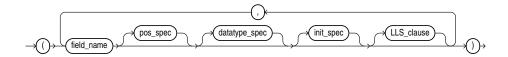
Alvin, Tolliver, 1976 Baer, Kenneth Mary, Dube, 1973

15.4.5 field list

The $field_definitions\ field_list\ clause\ identifies\ the\ fields\ in\ the\ data\ file\ and\ their\ data\ types.$

Syntax

The syntax for the field list clause is as follows:



The field list clauses are as follows:

field_name: A string identifying the name of a field in the data file. If the string is not within quotation marks, then the name is uppercased when matching field names with column names in the external table.

If field_name matches the name of a column in the external table that is referenced in the query, then the field value is used for the value of that external table column. If the name does not match any referenced name in the external table, then the field is not loaded but can be used for clause evaluation (for example WHEN or NULLIF).

- pos_spec: Indicates the position of the column within the record. For a full description of the syntax, see pos_spec Clause.
- datatype_spec: Indicates the data type of the field. If datatype_spec is omitted, then the
 access driver assumes the data type is CHAR (255). For a full description of the syntax, see
 datatype_spec Clause.
- init_spec: Indicates when a field is NULL or has a default value. For a full description of the syntax, see init_spec Clause.
- LLS: When LLS is specified for a field, ORACLE_LOADER does not load the value of the field
 into the corresponding column. Instead, it use the information in the value to determine
 where to find the value of the field. See LLS Clause.

Purpose

The field_list clause identifies the fields in the data file and their data types. Evaluation criteria for the field list clause are as follows:

- If no data type is specified for a field, then it is assumed to be CHAR(1) for a nondelimited field, and CHAR(255) for a delimited field.
- If no field list is specified, then the fields in the data file are assumed to be in the same order as the fields in the external table. The data type for all fields is CHAR (255) unless the column in the database is CHAR or VARCHAR. If the column in the database is CHAR or VARCHAR, then the data type for the field is still CHAR but the length is either 255 or the length of the column, whichever is greater.
- If no field list is specified and no delim_spec clause is specified, then the fields in the data file are assumed to be in the same order as fields in the external table. All fields are assumed to be CHAR(255) and terminated by a comma.

Example

This example shows the definition for an external table with no field_list and a delim_spec. It is followed by a sample of the data file that can be used to load it.

```
CREATE TABLE emp_load (first_name CHAR(15), last_name CHAR(20), year_of_birth INT)
ORGANIZATION EXTERNAL (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_tab_dir

ACCESS PARAMETERS (FIELDS TERMINATED BY "|")
LOCATION ('info.dat'));
```

Alvin|Tolliver|1976 Kenneth|Baer|1963 Mary|Dube|1973

15.4.6 pos spec Clause

The <code>ORACLE_LOADER</code> pos_spec clause indicates the position of the column within the record.

The setting of the STRING SIZES ARE IN clause determines whether pos_spec refers to byte positions or character positions. Using character positions with varying-width character sets takes significantly longer than using character positions with fixed-width character sets. Binary and multibyte character data should not be present in the same data file when pos_spec is used for character positions. If they are, then the results are unpredictable.

pos_spec Clause Syntax
 The syntax for the ORACLE LOADER pos spec clause is as follows.

star

The pos_spec clause start parameter indicates the number of bytes from the beginning of the record to where the field begins.

•

The pos_spec clause * parameter indicates that the field begins at the first byte after the end of the previous field.

increment

The pos_spec clause increment parameter positions the start of the field is a fixed number of bytes from the end of the previous field.

end

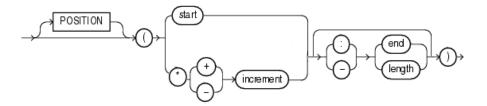
The pos_spec clause end parameter indicates the absolute byte offset into the record for the last byte of the field.

length

The pos_spec clause length parameter value indicates that the end of the field is a fixed number of bytes from the start.

15.4.6.1 pos spec Clause Syntax

The syntax for the ORACLE LOADER pos spec clause is as follows.



15.4.6.2 start

The pos_spec clause start parameter indicates the number of bytes from the beginning of the record to where the field begins.

The start parameter enables you to position the start of the field at an absolute spot in the record, rather than relative to the position of the previous field.

15.4.6.3 *

The pos_spec clause * parameter indicates that the field begins at the first byte after the end of the previous field.

The * parameter is useful with <code>ORACLE_LOADER</code> where you have a varying-length field followed by a fixed-length field. This option cannot be used for the first field in the record.

15.4.6.4 increment

The pos_spec clause increment parameter positions the start of the field is a fixed number of bytes from the end of the previous field.

The increment parameter positions the start of the field at a fixed number of bytes from the end of the previous field. Use *-increment to indicate that the start of the field starts before the current position in the record (this is a costly operation for multibyte character sets). To move the start after the current position, use *+increment

15.4.6.5 end

The pos_spec clause end parameter indicates the absolute byte offset into the record for the last byte of the field.

Use the end parameter to set the absolute byte offset into the record for the last byte of the field. If start is specified along with end, then end cannot be less than start. If * or increment is specified along with end, and the start evaluates to an offset larger than the end for a particular record, then that record will be rejected.

15.4.6.6 length

The pos_spec clause length parameter value indicates that the end of the field is a fixed number of bytes from the start.

Use the <code>length</code> parameter when you want to set fixed-length fields when the start is specified with \star . The following example shows various ways of using <code>pos_spec</code>. It is followed by an example of a data file that you can use to load it.

```
CREATE TABLE emp load (first name CHAR(15),
                      last name CHAR(20),
                      year of birth INT,
                      phone CHAR (12),
                      area code CHAR(3),
                      exchange CHAR(3),
                      extension CHAR(4))
  ORGANIZATION EXTERNAL
  (TYPE ORACLE LOADER
  DEFAULT DIRECTORY ext tab dir
  ACCESS PARAMETERS
     (RECORDS CHARACTERSET we8iso8859p1
     FIELDS RTRIM
            (first name (1:15) CHAR(15),
             last name (*:+20),
             year of birth (36:39),
             phone (40:52),
             area code (*-12: +3),
             exchange (*+1: +3),
             extension (*+1: +4))
   LOCATION ('info.dat'));
Alvin
                                  1976415-922-1982
               Tolliver
Kenneth
             Baer
                                  1963212-341-7912
Mary
               Dube
                                   1973309-672-2341
```

In this example, the declared RECORDS CHARACTERSET, we8iso8859p1, is not a multi-byte character set. It is guaranteed that every character is represented as single byte. The POSITION clause calculations to determine where the data field starts and ends (including the * and + operators) are based on bytes rather than characters (that is, characters must only require 1 byte to represent them, such as the Oracle character set WE8ISO8859P1). If you use a variable length character set (for example, Unicode variants, JIS X 0208-1990, or other multibyte character sets, where the field data contains one or more multibyte characters), then the calculations will be incorrect.

15.4.7 datatype_spec Clause

The <code>ORACLE_LOADER</code> datatype_spec clause describes the data type of a field in the data file if the data type is different than the default.

The data type of the field can be different than the data type of a corresponding column in the external table. The access driver handles the necessary conversions.

datatype spec Clause Syntax

The syntax for the ORACLE LOADER datatype spec clause is as follows:

• [UNSIGNED] INTEGER [EXTERNAL] [(len)]

The datatype_spec clause [UNSIGNED] INTEGER [EXTERNAL] [(len)] defines a field as an integer.

DECIMAL [EXTERNAL] and ZONED [EXTERNAL]

The DECIMAL clause is used to indicate that the field is a packed decimal number. The ZONED clause is used to indicate that the field is a zoned decimal number.

ORACLE DATE

ORACLE DATE is a field containing a date in the Oracle binary date format.

ORACLE NUMBER

ORACLE NUMBER is a field containing a number in the Oracle number format.

Floating-Point Numbers

The following four data types, DOUBLE, FLOAT, BINARY_DOUBLE, and BINARY_FLOAT are floating-point numbers.

DOUBLE

The DOUBLE clause indicates that the field is the same format as the C language DOUBLE data type on the platform where the access driver is executing.

FLOAT [EXTERNAL]

The FLOAT clause indicates that the field is the same format as the C language FLOAT data type on the platform where the access driver is executing.

BINARY DOUBLE

The datatype_spec clause value BINARY_DOUBLE is a 64-bit, double-precision, floating-point number data type.

BINARY FLOAT

The datatype_spec clause value BINARY_FLOAT is a 32-bit, single-precision, floating-point number data type.

RAW

The RAW clause is used to indicate that the source data is binary data.

CHAR

The datatype_spec clause data type CHAR clause is used to indicate that a field is a character data type.

date format spec

The date_format_spec clause is used to indicate that a character string field contains date data, time data, or both, in a specific format.

VARCHAR and VARRAW

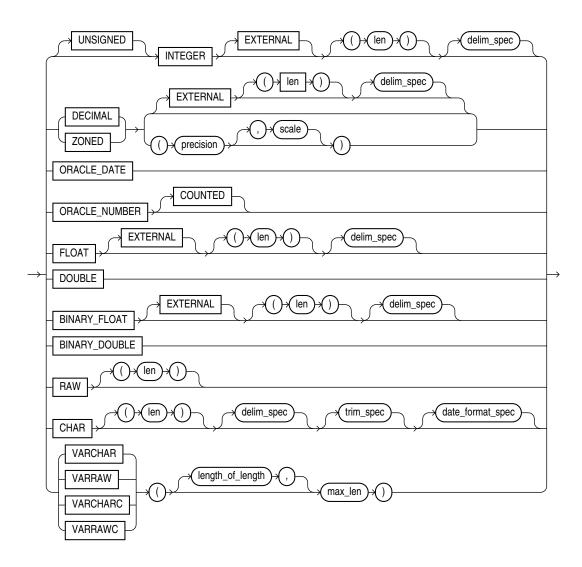
The datatype_spec clause VARCHAR data type defines character data, and the VARRAW data type defines binary data.

VARCHARC and VARRAWC

The datatype_spec clause VARCHARC data type defines character data, and the VARRAWC data type defines binary data.

15.4.7.1 datatype spec Clause Syntax

The syntax for the ORACLE LOADER datatype spec clause is as follows:



If the number of bytes or characters in any field is 0, then the field is assumed to be <code>NULL</code>. The optional <code>DEFAULTIF</code> clause specifies when the field is set to its default value. Also, the optional <code>NULLIF</code> clause specifies other conditions for when the column associated with the field is set to <code>NULL</code>. If the <code>DEFAULTIF</code> or <code>NULLIF</code> clause is <code>TRUE</code>, then the actions of those clauses override whatever values are read from the data file.

Related Topics

init spec Clause

The $init_spec$ clause for external tables is used to specify when a field should be set to NULL, or when it should be set to a default value.

Oracle Database SQL Language Reference



15.4.7.2 [UNSIGNED] INTEGER [EXTERNAL] [(len)]

The datatype_spec clause [UNSIGNED] INTEGER [EXTERNAL] [(len)] defines a field as an integer.

This clause defines a field as an integer. If EXTERNAL is specified, then the number is a character string. If EXTERNAL is not specified, then the number is a binary field. The valid values for len in binary integer fields are 1, 2, 4, and 8. If len is omitted for binary integers, then the default value is whatever the value of sizeof(int) is on the platform where the access driver is running. Use of the DATA IS {BIG|LITTLE} ENDIAN clause may cause the data to be byte-swapped before it is stored.

If EXTERNAL is specified, then the value of *len* is the number of bytes or characters in the number (depending on the setting of the STRING SIZES ARE IN BYTES or CHARACTERS clause). If no length is specified, then the default value is 255.

The default value of the [UNSIGNED] INTEGER [EXTERNAL] [(len)] data type is determined as follows:

- If no length specified, then the default length is 1.
- If no length is specified and the field is delimited with a DELIMITED BY NEWLINE clause, then the default length is 1.
- If no length is specified and the field is delimited with a DELIMITED BY clause, then the default length is 255 (unless the delimiter is NEWLINE, as stated above).

15.4.7.3 DECIMAL [EXTERNAL] and ZONED [EXTERNAL]

The DECIMAL clause is used to indicate that the field is a packed decimal number. The ZONED clause is used to indicate that the field is a zoned decimal number.

The precision field indicates the number of digits in the number. The scale field is used to specify the location of the decimal point in the number. It is the number of digits to the right of the decimal point. If scale is omitted, then a value of 0 is assumed.

Note that there are different encoding formats of zoned decimal numbers depending on whether the character set being used is EBCDIC-based or ASCII-based. If the language of the source data is EBCDIC, then the zoned decimal numbers in that file must match the EBCDIC encoding. If the language is ASCII-based, then the numbers must match the ASCII encoding.

If the ${\tt EXTERNAL}$ parameter is specified, then the data field is a character string whose length matches the precision of the field.

15.4.7.4 ORACLE_DATE

ORACLE DATE is a field containing a date in the Oracle binary date format.

This is the format used by the DTYDAT data type in Oracle Call Interface (OCI) programs. The field is a fixed length of 7.



15.4.7.5 ORACLE NUMBER

ORACLE NUMBER is a field containing a number in the Oracle number format.

The field is a fixed length (the maximum size of an Oracle number field) unless COUNTED is specified, in which case the first byte of the field contains the number of bytes in the rest of the field.

ORACLE_NUMBER is a fixed-length 22-byte field. The length of an ORACLE_NUMBER COUNTED field is one for the count byte, plus the number of bytes specified in the count byte.

15.4.7.6 Floating-Point Numbers

The following four data types, DOUBLE, FLOAT, BINARY_DOUBLE, and BINARY_FLOAT are floating-point numbers.

The following four data types, DOUBLE, FLOAT, BINARY_DOUBLE, and BINARY_FLOAT are floating-point numbers.

DOUBLE and FLOAT are the floating-point formats used natively on the platform in use. They are the same data types used by default for the DOUBLE and FLOAT data types in a C program on that platform. BINARY_FLOAT and BINARY_DOUBLE are floating-point numbers that conform substantially with the Institute for Electrical and Electronics Engineers (IEEE) Standard for Binary Floating-Point Arithmetic, IEEE Standard 754-1985. Because most platforms use the IEEE standard as their native floating-point format, FLOAT and BINARY_FLOAT are the same on those platforms and DOUBLE and BINARY_DOUBLE are also the same.



See Oracle Database SQL Language Reference for more information about floating-point numbers

15.4.7.7 DOUBLE

The DOUBLE clause indicates that the field is the same format as the C language DOUBLE data type on the platform where the access driver is executing.

Use of the DATA IS {BIG | LITTLE} ENDIAN clause may cause the data to be byte-swapped before it is stored. This data type may not be portable between certain platforms.

15.4.7.8 FLOAT [EXTERNAL]

The FLOAT clause indicates that the field is the same format as the C language FLOAT data type on the platform where the access driver is executing.

The FLOAT clause indicates that the field is the same format as the C language FLOAT data type on the platform where the access driver is executing. Use of the DATA IS $\{BIG \mid LITTLE\}$ ENDIAN clause may cause the data to be byte-swapped before it is stored. This data type may not be portable between certain platforms.

If the EXTERNAL parameter is specified, then the field is a character string whose maximum length is 255.

15.4.7.9 BINARY DOUBLE

The datatype_spec clause value BINARY_DOUBLE is a 64-bit, double-precision, floating-point number data type.

Each BINARY_DOUBLE value requires 9 bytes, including a length byte. See the information in the note provided for the FLOAT data type for more details about floating-point numbers.

15.4.7.10 BINARY_FLOAT

The datatype_spec clause value BINARY_FLOAT is a 32-bit, single-precision, floating-point number data type.

Each BINARY_FLOAT value requires 5 bytes, including a length byte. See the information in the note provided for the FLOAT data type for more details about floating-point numbers.

15.4.7.11 RAW

The RAW clause is used to indicate that the source data is binary data.

The *len* for RAW fields is always in number of bytes. When a RAW field is loaded in a character column, the data that is written into the column is the hexadecimal representation of the bytes in the RAW field.

15.4.7.12 CHAR

The datatype_spec clause data type CHAR clause is used to indicate that a field is a character data type.

The length (len) for CHAR fields specifies the largest number of bytes or characters in the field. The len is in bytes or characters, depending on the setting of the STRING SIZES ARE IN clause.

If no length is specified for a field of data type CHAR, then the size of the field is assumed to be 1, unless the field is delimited:

- For a delimited CHAR field, if a length is specified, then that length is used as a maximum.
- For a delimited CHAR field for which no length is specified, the default is 255 bytes.
- For a delimited CHAR field that is greater than 255 bytes, you must specify a maximum length. Otherwise, you receive an error stating that the field in the data file exceeds maximum length.

The following example shows the use of the CHAR clause.

```
SQL> CREATE TABLE emp load
  2
       (employee number
                             CHAR(5),
  3
       employee dob
                             CHAR (20),
       employee last name CHAR(20),
       employee first name CHAR(15),
  6
       employee middle name CHAR(15),
  7
       employee hire date DATE)
  8 ORGANIZATION EXTERNAL
  9
      (TYPE ORACLE LOADER
       DEFAULT DIRECTORY def dir1
10
```

```
11
       ACCESS PARAMETERS
12
         (RECORDS DELIMITED BY NEWLINE
13
          FIELDS (employee number CHAR(2),
14
                  employee dob
                                       CHAR (20),
                  employee last name CHAR(18),
15
                  employee_first name CHAR(11),
16
17
                  employee middle name CHAR(11),
18
                  employee hire date CHAR(10) date format DATE mask "mm/dd/
уууу"
19
20
       LOCATION ('info.dat')
21
22
      );
```

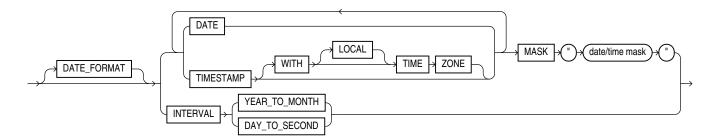
Table created.

15.4.7.13 date format spec

The date_format_spec clause is used to indicate that a character string field contains date data, time data, or both, in a specific format.

This information is used only when a character field is converted to a date or time data type and only when a character string field is mapped into a date column.

The syntax for the date format spec clause is as follows:



For detailed information about the correct way to specify date and time formats, see *Oracle Database SQL Reference*.

DATE

The DATE clause indicates that the string contains a date.

MASK

The MASK clause is used to override the default globalization format mask for the data type.

TIMESTAMP

The TIMESTAMP clause indicates that a field contains a formatted timestamp.

• INTERVAL

The INTERVAL clause indicates that a field contains a formatted interval.

Related Topics

Oracle Database SQL Language Reference

15.4.7.13.1 DATE

The DATE clause indicates that the string contains a date.

15.4.7.13.2 MASK

The MASK clause is used to override the default globalization format mask for the data type.

If a date mask is not specified, then the settings of NLS parameters for the database (not the session settings) for the appropriate globalization parameter for the data type are used. The NLS DATABASE PARAMETERS view shows these settings.

- NLS DATE FORMAT for DATE data types
- NLS_TIMESTAMP_FORMAT for TIMESTAMP data types
- NLS TIMESTAMP TZ FORMAT for TIMESTAMP WITH TIME ZONE data types

Note the following:

- The database setting for the NLS_NUMERIC_CHARACTERS initialization parameter (that is, from the NLS_DATABASE_PARAMETERS view) governs the decimal separator for implicit conversion from character to numeric data types.
- A group separator is not allowed in the default format.

15.4.7.13.3 TIMESTAMP

The TIMESTAMP clause indicates that a field contains a formatted timestamp.

15.4.7.13.4 INTERVAL

The INTERVAL clause indicates that a field contains a formatted interval.

The INTERVAL clause indicates that a field contains a formatted interval. The type of interval can be either YEAR TO MONTH or DAY TO SECOND.

The following example shows a sample use of a complex DATE character string and a TIMESTAMP character string. It is followed by a sample of the data file that can be used to load it.

```
SQL> CREATE TABLE emp load
    (employee number CHAR(5),
 3 employee dob CHAR(20),
     employee last name CHAR(20),
     employee first name CHAR(15),
    employee middle name CHAR(15),
 7 employee hire date DATE,
     rec creation date TIMESTAMP WITH TIME ZONE)
 9 ORGANIZATION EXTERNAL
10 (TYPE ORACLE LOADER
11
     DEFAULT DIRECTORY def dir1
12
     ACCESS PARAMETERS
       (RECORDS DELIMITED BY NEWLINE
        FIELDS (employee number CHAR(2),
                 employee dob
                                  CHAR (20),
16
                 employee last name CHAR(18),
17
                 employee first name CHAR(11),
                 employee middle name CHAR(11),
18
19
                 employee hire date CHAR(22) date format DATE mask "mm/dd/yyyy hh:mi:ss
ΑΜ",
```

```
20
           rec_creation date
                        CHAR(35) date format TIMESTAMP WITH TIME ZONE mask
"DD-MON-RR HH.MI.SSXFF AM TZH:TZM"
21
         )
22
     )
  LOCATION ('infoc.dat')
23
  ) ;
Table created.
SQL> SELECT * FROM emp load;
            EMPLO EMPLOYEE DOB
EMPLOYEE
_____
REC CREATION DATE
_____
56 november, 15, 1980 baker
                            mary alice
01-SEP-04
01-DEC-04 11.22.03.034567 AM -08:00
87 december, 20, 1970 roper
                            lisa
                                      marie
01-JAN-02
01-DEC-02 02.03.00.678573 AM -08:00
```

2 rows selected.

The info.dat file looks like the following. Note that this is 2 long records. There is one space between the data fields (09/01/2004, 01/01/2002) and the time field that follows.

```
56november, 15, 1980 baker mary alice 09/01/2004 08:23:01 AM01-DEC-04 11.22.03.034567 AM -08:00 lisa marie 01/01/2002 02:44:55 PM01-DEC-02 02.03.00.678573 AM -08:00
```

15.4.7.14 VARCHAR and VARRAW

The datatype_spec clause VARCHAR data type defines character data, and the VARRAW data type defines binary data.

The VARCHAR data type has a binary count field followed by character data. The value in the binary count field is either the number of bytes in the field or the number of characters. See STRING SIZES ARE IN for information about how to specify whether the count is interpreted as a count of characters or count of bytes.

The VARRAW data type has a binary count field followed by binary data. The value in the binary count field is the number of bytes of binary data. The data in the VARRAW field is not affected by the DATA IS...ENDIANClause.

The VARIABLE 2 clause in the ACCESS PARAMETERS clause specifies the size of the binary field that contains the length.

The optional <code>length_of_length</code> field in the specification is the number of bytes in the count field. Valid values for <code>length_of_length</code> for <code>VARCHAR</code> are 1, 2, 4, and 8. If <code>length_of_length</code> is not specified, then a value of 2 is used. The count field has the same endianness as specified by the <code>DATA_IS...ENDIAN</code> clause.

The <code>max_len</code> field is used to indicate the largest size of any instance of the field in the data file. For <code>VARRAW</code> fields, <code>max_len</code> is number of bytes. For <code>VARCHAR</code> fields, <code>max_len</code> is either number of characters, or number of bytes, depending on the <code>STRING SIZES ARE IN clause</code>.

The following example shows various uses of VARCHAR and VARRAW. The content of the data file, info.dat, is shown following the example.

```
CREATE TABLE emp_load
             (first name CHAR(15),
             last name CHAR(20),
              resume CHAR(2000),
              picture RAW(2000))
  ORGANIZATION EXTERNAL
  (TYPE ORACLE LOADER
  DEFAULT DIRECTORY ext tab dir
  ACCESS PARAMETERS
     (RECORDS
       VARIABLE 2
        DATA IS BIG ENDIAN
        CHARACTERSET US7ASCII
      FIELDS (first name VARCHAR(2,12),
              last name VARCHAR(2,20),
              resume VARCHAR(4,10000),
              picture VARRAW(4,100000)))
    LOCATION ('info.dat'));
```

Contents of info.dat Data File

The contents of the data file used in the example are as follows:.

```
0005Alvin0008Tolliver0000001DAlvin Tolliver's Resume etc. 0000001013f4690a30bc29d7e40023ab4599ffff
```

It is important to understand that, for the purposes of readable documentation, the binary values for the count bytes and the values for the raw data are shown in the data file in italics, with 2 characters per binary byte. The values in an actual data file would be in binary format, not ASCII. Therefore, if you attempt to use this example by cutting and pasting, then you will receive an error.

Related Topics

STRING SIZES ARE IN

Use the record_format_info STRING SIZES ARE IN clause to indicate whether the lengths specified for character strings are in bytes or characters.

15.4.7.15 VARCHARC and VARRAWC

The datatype_spec clause VARCHARC data type defines character data, and the VARRAWC data type defines binary data.

The VARCHARC data type has a character count field followed by character data. The value in the count field is either the number of bytes in the field or the number of characters. See STRING SIZES ARE IN for information about how to specify whether the count is interpreted as a count of characters, or acount of bytes. The optional <code>length_of_length</code> is either the number of bytes, or the number of characters in the count field for <code>VARCHARC</code>, depending on whether lengths are being interpreted as characters or bytes.

The maximum value for <code>length_of_lengths</code> for <code>VARCHARC</code> is 10 if string sizes are in characters, and 20 if string sizes are in bytes. The default value for <code>length of length</code> is 5.

The VARRAWC data type has a character count field followed by binary data. The value in the count field is the number of bytes of binary data. The <code>length_of_length</code> is the number of bytes in the count field.

The <code>max_len</code> field is used to indicate the largest size of any instance of the field in the data file. For <code>VARRAWC</code> fields, <code>max_len</code> is number of bytes. For <code>VARCHARC</code> fields, <code>max_len</code> is either number of characters or number of bytes depending on the <code>STRING SIZES</code> ARE <code>IN</code> clause.

The following example shows various uses of VARCHARC and VARRAWC. The length of the picture field is 0, which means the field is set to NULL.

```
CREATE TABLE emp_load

(first_name CHAR(15),
last_name CHAR(20),
resume CHAR(2000),
picture RAW (2000))

ORGANIZATION EXTERNAL

(TYPE ORACLE_LOADER
DEFAULT DIRECTORY ext_tab_dir
ACCESS PARAMETERS

(FIELDS (first_name VARCHARC(5,12),
last_name VARCHARC(2,20),
resume VARCHARC(4,10000),
picture VARRAWC(4,100000)))

LOCATION ('info.dat'));
```

00007William05Ricca0035Resume for William Ricca is missing0000

Related Topics

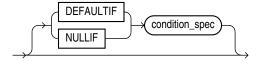
STRING SIZES ARE IN

Use the ${\tt record_format_info}$ STRING SIZES ARE IN clause to indicate whether the lengths specified for character strings are in bytes or characters.

15.4.8 init_spec Clause

The $init_spec$ clause for external tables is used to specify when a field should be set to NULL, or when it should be set to a default value.

The syntax for the init spec clause is as follows:



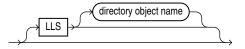
Only one NULLIF clause and only one DEFAULTIF clause can be specified for any field. These clauses behave as follows:

- If NULLIF condition spec is specified and it evaluates to TRUE, then the field is set to NULL.
- If DEFAULTIF condition_spec is specified and it evaluates to TRUE, then the value of the
 field is set to a default value. The default value depends on the data type of the field, as
 follows:
 - For a character data type, the default value is an empty string.
 - For a numeric data type, the default value is a 0.
 - For a date data type, the default value is NULL.
- If a NULLIF clause and a DEFAULTIF clause are both specified for a field, then the NULLIF clause is evaluated first, and the DEFAULTIF clause is evaluated only if the NULLIF clause evaluates to FALSE.

15.4.9 LLS Clause

If a field in a data file is a LOB location Specifier (LLS) field, then you can indicate this by using the LLS clause.

If a field in a data file is a LOB location Specifier (LLS) field, then you can indicate this by using the LLS clause. An LLS field contains the file name, offset, and length of the LOB data in the data file. SQL*Loader uses this information to read data for the LOB column. The LLS clause for ORACLE LOADER has the following syntax:



When the LLS clause is used, <code>ORACLE_LOADER</code> does not load the value of the field into the corresponding column. Instead, it uses the information in the value to determine where to find the value of the field. The LOB can be loaded in part or in whole and it can start from an arbitrary position and for an arbitrary length. <code>ORACLE_LOADER</code> expects the contents of the field to be <code>filename.ext.nnn.mmm/</code> where each element is defined as follows:

- filename.ext is the name of the file that contains the LOB
- nnn is the offset in bytes of the LOB within the file
- mmm is the length of the LOB in bytes A value of -1 means the LOB is NULL. A value of 0 means the lob exists, but is empty.
- The forward slash (/) terminates the field

The LLS clause has an optional DIRECTORY clause which specifies an Oracle directory object:

 If DIRECTORY is specified, then the file must exist there and you must have READ access to that directory object. If DIRECTORY is not specified, then the file must exist in the same directory as the data file.

An error is returned and the row rejected if any of the following are true:

- The file name contains a relative or absolute path specification.
- The file is not found, the offset is invalid, or the length extends beyond the end of the file.
- The contents of the field do not match the expected format.
- The data type for the column associated with an LLS field is not a CLOB, BLOB or NCLOB.

If an LLS field is referenced by a clause for any other field (for example a NULLIF clause), then in the access parameters, the value used for evaluating the clause is the string in the data file, not the data in the file pointed to by that string.

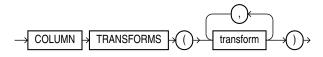
The character set for the data in the file pointed to by the LLS clause is assumed to be the same character set as the data file.

15.5 column transforms Clause

The optional <code>ORACLE_LOADER</code> access drive <code>COLUMN TRANSFORMS</code> clause provides transforms that you can use to describe how to load columns in the external table that do not map directly to columns in the data file.

Syntax

The syntax for the column transforms clause is as follows:





The COLUMN TRANSFORMS clause does not work in conjunction with the PREPROCESSOR clause.

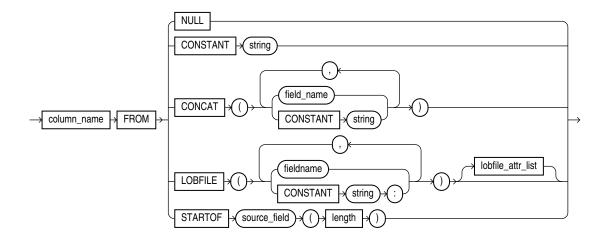
transform

Each transform specified in the transform clause identifies a column in the external table and then a specifies how to calculate the value of the column.

15.5.1 transform

Each transform specified in the transform clause identifies a column in the external table and then a specifies how to calculate the value of the column.

The syntax is as follows:



The NULL transform is used to set the external table column to NULL in every row. The CONSTANT transform is used to set the external table column to the same value in every row. The CONCAT transform is used to set the external table column to the concatenation of constant strings and/or fields in the current record from the data file. The LOBFILE transform is used to load data into a field for a record from another data file. Each of these transforms is explained further in the following sections.

· column name FROM

The column_name uniquely identifies a column in the external table that you want to be loaded.

NULL

When the \mathtt{NULL} transform is specified, every value of the field is set to \mathtt{NULL} for every record.

CONSTANT

The CONSTANT clause transform uses the value of the string specified as the value of the column in the record.

CONCAT

The CONCAT transform concatenates constant strings and fields in the data file together to form one string.

LOBFILE

The LOBFILE transform is used to identify a file whose contents are to be used as the value for a column in the external table.

lobfile_attr_list

The lobfile attr list lists additional attributes of the LOBFILE.

STARTOF source_field (length)

The STARTOF keyword allows you to create an external table in which a column can be a substring of the data in the source field.

15.5.1.1 column_name FROM

The column name uniquely identifies a column in the external table that you want to be loaded.

Note that if the name of a column is mentioned in the transform clause, then that name cannot be specified in the FIELDS clause as a field in the data file.

15.5.1.2 NULL

When the NULL transform is specified, every value of the field is set to NULL for every record.

15.5.1.3 CONSTANT

The CONSTANT clause transform uses the value of the string specified as the value of the column in the record.

If the column in the external table is not a character string type, then the constant string will be converted to the data type of the column. This conversion will be done for every row.

The character set of the string used for data type conversions is the character set of the database.

15.5.1.4 CONCAT

The CONCAT transform concatenates constant strings and fields in the data file together to form one string.

Only fields that are character data types and that are listed in the fields clause can be used as part of the concatenation. Other column transforms cannot be specified as part of the concatenation.

15.5.1.5 LOBFILE

The LOBFILE transform is used to identify a file whose contents are to be used as the value for a column in the external table.

All LOBFILEs are identified by an optional directory object and a file name in the form directory object: filename. The following rules apply to use of the LOBFILE transform:

- Both the directory object and the file name can be either a constant string or the name of a field in the field clause.
- If a constant string is specified, then that string is used to find the LOBFILE for every row in the table.
- If a field name is specified, then the value of that field in the data file is used to find the LOBFILE.
- If a field name is specified for either the directory object or the file name and if the value of that field is NULL, then the column being loaded by the LOBFILE is also set to NULL.
- If the directory object is not specified, then the default directory specified for the external table is used.
- If a field name is specified for the directory object, then the FROM clause also needs to be specified.

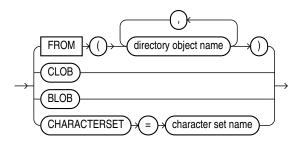
Note that the entire file is used as the value of the LOB column. If the same file is referenced in multiple rows, then that file is reopened and reread in order to populate each column.

15.5.1.6 lobfile_attr_list

The lobfile attr list lists additional attributes of the LOBFILE.

The syntax is as follows:





The FROM clause lists the names of all directory objects that will be used for LOBFILEs. It is used only when a field name is specified for the directory object of the name of the LOBFILE. The purpose of the FROM clause is to determine the type of access allowed to the named directory objects during initialization. If directory object in the value of field is not a directory object in this list, then the row will be rejected.

The CLOB attribute indicates that the data in the LOBFILE is character data (as opposed to RAW data). Character data may need to be translated into the character set used to store the LOB in the database.

The CHARACTERSET attribute contains the name of the character set for the data in the LOBFILES.

The BLOB attribute indicates that the data in the LOBFILE is raw data.

If neither CLOB nor BLOB is specified, then CLOB is assumed. If no character set is specified for character LOBFILEs, then the character set of the data file is assumed.

15.5.1.7 STARTOF source field (length)

The STARTOF keyword allows you to create an external table in which a column can be a substring of the data in the source field.

The length is the length of the substring, beginning with the first byte. It is assumed that length refers to a byte count and that the external table column(s) being transformed use byte length and not character length semantics. (Character length semantics might give unexpected results.)

Only complete character encodings are moved; characters are never split. So if a substring ends in the middle of a multibyte character, then the resulting string will be shortened. For example, if a length of 10 is specified, but the 10th byte is the first byte of a multibyte character, then only the first 9 bytes are returned.

The following example shows how you could use the STARTOF keyword if you only wanted the first 4 bytes of the department name (dname) field:

```
SQL> CREATE TABLE dept (deptno NUMBER(2),
  2
                        dname VARCHAR2(14),
  3
                                VARCHAR2 (13)
                        loc
  4
  5
    ORGANIZATION EXTERNAL
  6
  7
       DEFAULT DIRECTORY def dir1
  8
       ACCESS PARAMETERS
  9
10
         RECORDS DELIMITED BY NEWLINE
 11
         FIELDS TERMINATED BY ','
```



Table created.

If you now perform a SELECT operation from the dept table, only the first four bytes of the dname field are returned:

```
SQL> SELECT * FROM dept;

DEPTNO DNAME LOC

10 ACCO NEW YORK
20 RESE DALLAS
30 SALE CHICAGO
40 OPER BOSTON

4 rows selected.
```

15.6 Parallel Loading Considerations for the ORACLE_LOADER Access Driver

The <code>ORACLE_LOADER</code> access driver attempts to divide large data files into chunks that can be processed separately.

The following file, record, and data characteristics make it impossible for a file to be processed in parallel:

- Sequential data sources (such as a tape drive or pipe)
- Data in any multibyte character set whose character boundaries cannot be determined starting at an arbitrary byte in the middle of a string

This restriction does not apply to any data file with a fixed number of bytes per record.

Records with the VAR format

Specifying a PARALLEL clause is of value only when large amounts of data are involved.

15.7 Performance Hints When Using the ORACLE_LOADER Access Driver

This topic describes some performance hints when using the ORACLE LOADER access driver.

When you monitor performance, the most important measurement is the elapsed time for a load. Other important measurements are CPU usage, memory usage, and I/O rates.

You can alter performance by increasing or decreasing the degree of parallelism. The degree of parallelism indicates the number of access drivers that can be started to process the data files. The degree of parallelism enables you to choose on a scale between slower load with little resource usage and faster load with all resources utilized. The access driver cannot automatically tune itself, because it cannot determine how many resources you want to dedicate to the access driver.

An additional consideration is that the access drivers use large I/O buffers for better performance (you can use the READSIZE clause in the access parameters to specify the size of the buffers). On databases with shared servers, all memory used by the access drivers comes out of the system global area (SGA). For this reason, you should be careful when using external tables on shared servers.

Performance can also sometimes be increased with use of date cache functionality. By using the date cache to specify the number of unique dates anticipated during the load, you can reduce the number of date conversions done when many duplicate date or timestamp values are present in the input data. The date cache functionality provided by external tables is identical to the date cache functionality provided by SQL*Loader. See DATE_CACHE for a detailed description.

In addition to changing the degree of parallelism and using the date cache to improve performance, consider the following information:

- Fixed-length records are processed faster than records terminated by a string.
- Fixed-length fields are processed faster than delimited fields.
- Single-byte character sets are the fastest to process.
- Fixed-width character sets are faster to process than varying-width character sets.
- Byte-length semantics for varying-width character sets are faster to process than character-length semantics.
- Single-character delimiters for record terminators and field delimiters are faster to process than multicharacter delimiters.
- Having the character set in the data file match the character set of the database is faster than a character set conversion.
- Having data types in the data file match the data types in the database is faster than data type conversion.
- Not writing rejected rows to a reject file is faster because of the reduced overhead.
- Condition clauses (including WHEN, NULLIF, and DEFAULTIF) slow down processing.
- The access driver takes advantage of multithreading to streamline the work as much as possible.

15.8 Restrictions When Using the ORACLE_LOADER Access Driver

This section lists restrictions to be aware of when you use the <code>ORACLE_LOADER</code> access driver. Specifically:

Exporting and importing of external tables with encrypted columns is not supported.

- Column processing: By default, the external tables feature fetches all columns defined for an external table. This guarantees a consistent result set for all queries. However, for performance reasons you can decide to process only the referenced columns of an external table, thus minimizing the amount of data conversion and data handling required to execute a query. In this case, a row that is rejected because a column in the row causes a data type conversion error will not get rejected in a different query if the query does not reference that column. You can change this column-processing behavior with the ALTER TABLE command.
- An external table cannot load data into a LONG column.
- SQL strings cannot be specified in access parameters for the ORACLE_LOADER access
 driver. As a workaround, you can use the DECODE clause in the SELECT clause of the
 statement that is reading the external table. Alternatively, you can create a view of the
 external table that uses the DECODE clause and select from that view rather than the
 external table.
- The use of the backslash character (\) within strings is not supported in external tables. See Use of the Backslash Escape Character.
- When identifiers (for example, column or table names) are specified in the external table
 access parameters, certain values are considered to be reserved words by the access
 parameter parser. If a reserved word is used as an identifier, then it must be enclosed in
 double quotation marks.

15.9 Reserved Words for the ORACLE_LOADER Access Driver

When identifiers (for example, column or table names) are specified in the external table access parameters, certain values are considered to be reserved words by the access parameter parser.

If a reserved word is used as an identifier, then it must be enclosed in double quotation marks. The following are the reserved words for the <code>ORACLE LOADER</code> access driver:

- ALL
- AND
- ARE
- ASTERISK
- AT
- ATSIGN
- BADFILE
- BADFILENAME
- BACKSLASH
- BENDIAN
- BIG
- BLANKS
- BY
- BYTES
- BYTESTR
- CHAR



- CHARACTERS
- CHARACTERSET
- CHARSET
- CHARSTR
- CHECK
- CLOB
- COLLENGTH
- COLON
- COLUMN
- COMMA
- CONCAT
- CONSTANT
- COUNTED
- DATA
- DATE
- DATE_CACHE
- DATE_FORMAT
- DATEMASK
- DAY
- DEBUG
- DECIMAL
- DEFAULTIF
- DELIMITBY
- DELIMITED
- DISCARDFILE
- DNFS_ENABLE
- DNFS_DISABLE
- DNFS_READBUFFERS
- DOT
- DOUBLE
- DOUBLETYPE
- DQSTRING
- DQUOTE
- DSCFILENAME
- ENCLOSED
- ENDIAN
- ENDPOS
- EOF



- EQUAL
- EXIT
- EXTENDED_IO_PARAMETERS
- EXTERNAL
- EXTERNALKW
- EXTPARM
- FIELD
- FIELDS
- FILE
- FILEDIR
- FILENAME
- FIXED
- FLOAT
- FLOATTYPE
- FOR
- FROM
- HASH
- HEXPREFIX
- IN
- INTEGER
- INTERVAL
- LANGUAGE
- IS
- LEFTCB
- LEFTTXTDELIM
- LEFTP
- LENDIAN
- LDRTRIM
- LITTLE
- LOAD
- LOBFILE
- LOBPC
- LOBPCCONST
- LOCAL
- LOCALTZONE
- LOGFILE
- LOGFILENAME
- LRTRIM



- LTRIM
- MAKE_REF
- MASK
- MINUSSIGN
- MISSING
- MISSINGFLD
- MONTH
- NEWLINE
- NO
- NOCHECK
- NOT
- NOBADFILE
- NODISCARDFILE
- NOLOGFILE
- NOTEQUAL
- NOTERMBY
- NOTRIM
- NULL
- NULLIF
- OID
- OPTENCLOSE
- OPTIONALLY
- OPTIONS
- OR
- ORACLE_DATE
- ORACLE_NUMBER
- PLUSSIGN
- POSITION
- PROCESSING
- QUOTE
- RAW
- READSIZE
- RECNUM
- RECORDS
- REJECT
- RIGHTCB
- RIGHTTXTDELIM
- RIGHTP



- ROW
- ROWS
- RTRIM
- SCALE
- SECOND
- SEMI
- SETID
- SIGN
- SIZES
- SKIP
- STRING
- TERMBY
- TERMEOF
- TERMINATED
- TERMWS
- TERRITORY
- TIME
- TIMESTAMP
- TIMEZONE
- TO
- TRANSFORMS
- UNDERSCORE
- UINTEGER
- UNSIGNED
- VALUES
- VARCHAR
- VARCHARC
- VARIABLE
- VARRAW
- VARRAWC
- VLENELN
- VMAXLEN
- WHEN
- WHITESPACE
- WITH
- YEAR
- ZONED

