

Oracle® Database

SQL Tuning Guide



23ai
F46737-19
April 2025

ORACLE®

Contributors: Hermann Baer, Bjorn Bolltoft, Ali Cakmak, Sunil Chakkappen, Immanuel Chan, Deba Chatterjee, Chris Chiappa, Dinesh Das, Kurt Engeleiter, Leonidas Galanis, William Endress, Marcus Fallen, Bruce Golbus, Katsumi Inoue, Mark Jefferys, Shantanu Joshi, Adam Kociubes, Praveen Kumar Tupati Jaganath, Keith Laker, Allison Lee, Sue Lee, Cheng Li, Jiakun Li, David McDermid, Colin McGregor, Ajit Mylavarapu, Kantikiran Pasupuleti, Ted Persky, Palash Sharma, Lei Sheng, Ekrem Soylemez, Hong Su, Murali Thiyagarajah, Randy Urbano, Satya Valluri, Sahil Vazirani, Bharath Venkatakrishnan, Hailing Yu, Yuying Zhang, John Zimmerman

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this is software, software documentation, data (as defined in the Federal Acquisition Regulation), or related documentation that is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, then the following notice is applicable:

U.S. GOVERNMENT END USERS: Oracle programs (including any operating system, integrated software, any programs embedded, installed, or activated on delivered hardware, and modifications of such programs) and Oracle computer documentation or other Oracle data delivered to or accessed by U.S. Government end users are "commercial computer software," "commercial computer software documentation," or "limited rights data" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, the use, reproduction, duplication, release, display, disclosure, modification, preparation of derivative works, and/or adaptation of i) Oracle programs (including any operating system, integrated software, any programs embedded, installed, or activated on delivered hardware, and modifications of such programs), ii) Oracle computer documentation and/or iii) other Oracle data, is subject to the rights and limitations specified in the license contained in the applicable contract. The terms governing the U.S. Government's use of Oracle cloud services are defined by the applicable contract for such services. No other rights are granted to the U.S. Government.

This software or hardware is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications that may create a risk of personal injury. If you use this software or hardware in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure its safe use. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software or hardware in dangerous applications.

Oracle®, Java, MySQL, and NetSuite are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Inside are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Epyc, and the AMD logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group.

This software or hardware and documentation may provide access to or information about content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services unless otherwise set forth in an applicable agreement between you and Oracle. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services, except as set forth in an applicable agreement between you and Oracle.

Contents

Preface

Audience	xxv
Documentation Accessibility	xxv
Related Documents	xxv
Conventions	xxvi

Part I SQL Performance Fundamentals

1 Introduction to SQL Tuning

About SQL Tuning	1-1
Purpose of SQL Tuning	1-1
Prerequisites for SQL Tuning	1-1
Tasks and Tools for SQL Tuning	1-2
SQL Tuning Tasks	1-2
SQL Tuning Tools	1-4
Automated SQL Tuning Tools	1-4
Manual SQL Tuning Tools	1-13
User Interfaces to SQL Tuning Tools	1-15
About Automatic Error Mitigation	1-16

2 SQL Performance Methodology

Guidelines for Designing Your Application	2-1
Guideline for Data Modeling	2-1
Guideline for Writing Efficient Applications	2-1
Guidelines for Deploying Your Application	2-3
Guideline for Deploying in a Test Environment	2-3
Guidelines for Application Rollout	2-4

Part II Query Optimizer Fundamentals

3 SQL Processing

About SQL Processing	3-1
SQL Parsing	3-2
Syntax Check	3-2
Semantic Check	3-2
Shared Pool Check	3-3
SQL Optimization	3-5
SQL Row Source Generation	3-5
SQL Execution	3-6
How Oracle Database Processes DML	3-8
How Row Sets Are Fetched	3-8
Read Consistency	3-8
Data Changes	3-9
How Oracle Database Processes DDL	3-9

4 Query Optimizer Concepts

Introduction to the Query Optimizer	4-1
Purpose of the Query Optimizer	4-1
Cost-Based Optimization	4-1
Execution Plans	4-2
Query Blocks	4-2
Query Subplans	4-3
Analogy for the Optimizer	4-3
About Optimizer Components	4-4
Query Transformer	4-4
Estimator	4-5
Selectivity	4-6
Cardinality	4-7
Cost	4-8
Plan Generator	4-8
About Automatic Tuning Optimizer	4-10
About Adaptive Query Optimization	4-11
Adaptive Query Plans	4-11
About Adaptive Query Plans	4-12
Purpose of Adaptive Query Plans	4-12
How Adaptive Query Plans Work	4-13
When Adaptive Query Plans Are Enabled	4-20
Adaptive Statistics	4-20
Dynamic Statistics	4-20
Automatic Reoptimization	4-21

SQL Plan Directives	4-23
When Adaptive Statistics Are Enabled	4-24
About Approximate Query Processing	4-24
Approximate Query Initialization Parameters	4-25
Approximate Query SQL Functions	4-26
About SQL Plan Management	4-28
About Quarantined SQL Plans	4-28
About the Expression Statistics Store (ESS)	4-29

5 Query Transformations

OR Expansion	5-1
View Merging	5-3
Query Blocks in View Merging	5-3
Simple View Merging	5-4
Complex View Merging	5-6
Predicate Pushing	5-9
Subquery Unnesting	5-10
Query Rewrite with Materialized Views	5-11
About Query Rewrite and the Optimizer	5-11
About Initialization Parameters for Query Rewrite	5-12
About the Accuracy of Query Rewrite	5-12
Example of Query Rewrite	5-13
Star Transformation	5-14
About Star Schemas	5-14
Purpose of Star Transformations	5-15
How Star Transformation Works	5-15
Controls for Star Transformation	5-15
Star Transformation: Scenario	5-16
Temporary Table Transformation: Scenario	5-19
In-Memory Aggregation (VECTOR GROUP BY)	5-21
Cursor-Duration Temporary Tables	5-21
Purpose of Cursor-Duration Temporary Tables	5-21
How Cursor-Duration Temporary Tables Work	5-21
Cursor-Duration Temporary Tables: Example	5-22
Table Expansion	5-23
Purpose of Table Expansion	5-23
How Table Expansion Works	5-23
Table Expansion: Scenario	5-24
Table Expansion and Star Transformation: Scenario	5-27
Join Factorization	5-29
Purpose of Join Factorization	5-29

How Join Factorization Works	5-29
Factorization and Join Orders: Scenario	5-30
Factorization of Outer Joins: Scenario	5-31

Part III Query Execution Plans

6 Explaining and Displaying Execution Plans

Introduction to Execution Plans	6-1
Contents of an Execution Plan	6-1
Why Execution Plans Change	6-2
Different Schemas	6-2
Different Costs	6-2
Generating Plan Output Using the EXPLAIN PLAN Statement	6-3
About the EXPLAIN PLAN Statement	6-3
About PLAN_TABLE	6-3
EXPLAIN PLAN Restrictions	6-4
Explaining a SQL Statement: Basic Steps	6-4
Specifying a Statement ID in EXPLAIN PLAN: Example	6-7
Specifying a Different Location for EXPLAIN PLAN Output: Example	6-7
EXPLAIN PLAN Output for a CONTAINERS Query: Example	6-8
Displaying Execution Plans	6-9
About the Display of PLAN_TABLE Output	6-9
DBMS_XPLAN Display Functions	6-10
Plan-Related Views	6-12
Displaying Execution Plans: Basic Steps	6-13
Displaying Adaptive Query Plans: Tutorial	6-15
Display Execution Plans: Examples	6-19
Customizing PLAN_TABLE Output	6-19
Displaying Parallel Execution Plans: Example	6-20
Displaying Bitmap Index Plans: Example	6-22
Displaying Result Cache Plans: Example	6-23
Displaying Plans for Partitioned Objects: Example	6-24
Comparing Execution Plans	6-32
Purpose of Plan Comparison	6-32
User Interface for Plan Comparison	6-32
Comparing Execution Plans: Tutorial	6-36
Comparing Execution Plans: Examples	6-39

7 PLAN_TABLE Reference

PLAN_TABLE Columns	7-1
--------------------	-----

OPERATION and OPTION Columns of PLAN_TABLE	7-4
DISTRIBUTION Column of PLAN_TABLE	7-9

Part IV SQL Operators: Access Paths and Joins

8 Optimizer Access Paths

Introduction to Access Paths	8-1
Table Access Paths	8-2
About Heap-Organized Table Access	8-2
Row Storage in Data Blocks and Segments: A Primer	8-2
Importance of Rowids for Row Access	8-3
Direct Path Reads	8-3
Full Table Scans	8-4
When the Optimizer Considers a Full Table Scan	8-4
How a Full Table Scan Works	8-6
Full Table Scan: Example	8-7
Table Access by Rowid	8-7
When the Optimizer Chooses Table Access by Rowid	8-8
How Table Access by Rowid Works	8-8
Table Access by Rowid: Example	8-8
Sample Table Scans	8-9
When the Optimizer Chooses a Sample Table Scan	8-9
Sample Table Scans: Example	8-9
In-Memory Table Scans	8-10
When the Optimizer Chooses an In-Memory Table Scan	8-10
In-Memory Query Controls	8-10
In-Memory Table Scans: Example	8-11
B-Tree Index Access Paths	8-12
About B-Tree Index Access	8-12
B-Tree Index Structure	8-12
How Index Storage Affects Index Scans	8-13
Unique and Nonunique Indexes	8-14
B-Tree Indexes and Nulls	8-14
Index Unique Scans	8-16
When the Optimizer Considers Index Unique Scans	8-16
How Index Unique Scans Work	8-17
Index Unique Scans: Example	8-18
Index Range Scans	8-19
When the Optimizer Considers Index Range Scans	8-19
How Index Range Scans Work	8-20

Index Range Scan: Example	8-21
Index Range Scan Descending: Example	8-22
Index Full Scans	8-23
When the Optimizer Considers Index Full Scans	8-23
How Index Full Scans Work	8-23
Index Full Scans: Example	8-24
Index Fast Full Scans	8-25
When the Optimizer Considers Index Fast Full Scans	8-25
How Index Fast Full Scans Work	8-25
Index Fast Full Scans: Example	8-25
Index Skip Scans	8-26
When the Optimizer Considers Index Skip Scans	8-26
How Index Skip Scans Work	8-26
Index Skip Scans: Example	8-26
Index Join Scans	8-28
When the Optimizer Considers Index Join Scans	8-28
How Index Join Scans Work	8-29
Index Join Scans: Example	8-29
Bitmap Index Access Paths	8-30
About Bitmap Index Access	8-30
Differences Between Bitmap and B-Tree Indexes	8-30
Purpose of Bitmap Indexes	8-31
Bitmaps and Rowids	8-32
Bitmap Join Indexes	8-33
Bitmap Storage	8-35
Bitmap Conversion to Rowid	8-35
When the Optimizer Chooses Bitmap Conversion to Rowid	8-35
How Bitmap Conversion to Rowid Works	8-35
Bitmap Conversion to Rowid: Example	8-35
Bitmap Index Single Value	8-36
When the Optimizer Considers Bitmap Index Single Value	8-36
How Bitmap Index Single Value Works	8-36
Bitmap Index Single Value: Example	8-36
Bitmap Index Range Scans	8-37
When the Optimizer Considers Bitmap Index Range Scans	8-37
How Bitmap Index Range Scans Work	8-37
Bitmap Index Range Scans: Example	8-38
Bitmap Merge	8-39
When the Optimizer Considers Bitmap Merge	8-39
How Bitmap Merge Works	8-39
Bitmap Merge: Example	8-39
Table Cluster Access Paths	8-40

Cluster Scans	8-40
When the Optimizer Considers Cluster Scans	8-40
How a Cluster Scan Works	8-40
Cluster Scans: Example	8-41
Hash Scans	8-42
When the Optimizer Considers a Hash Scan	8-42
How a Hash Scan Works	8-42
Hash Scans: Example	8-42

9 Joins

About Joins	9-1
Join Trees	9-1
How the Optimizer Executes Join Statements	9-3
How the Optimizer Chooses Execution Plans for Joins	9-4
Join Methods	9-5
Nested Loops Joins	9-5
When the Optimizer Considers Nested Loops Joins	9-6
How Nested Loops Joins Work	9-6
Nested Nested Loops	9-7
Current Implementation for Nested Loops Joins	9-10
Original Implementation for Nested Loops Joins	9-12
Nested Loops Controls	9-14
Hash Joins	9-15
When the Optimizer Considers Hash Joins	9-16
How Hash Joins Work	9-16
How Hash Joins Work When the Hash Table Does Not Fit in the PGA	9-18
Hash Join Controls	9-19
Sort Merge Joins	9-19
When the Optimizer Considers Sort Merge Joins	9-20
How Sort Merge Joins Work	9-20
Sort Merge Join Controls	9-24
Join Types	9-24
Inner Joins	9-24
Equijoins	9-24
Nonequijoins	9-25
Band Joins	9-26
Outer Joins	9-29
Nested Loops Outer Joins	9-30
Hash Join Outer Joins	9-30
Sort Merge Outer Joins	9-32
Full Outer Joins	9-32

Multiple Tables on the Left of an Outer Join	9-33
Semijoins	9-34
When the Optimizer Considers Semijoins	9-34
How Semijoins Work	9-34
Antijoins	9-36
When the Optimizer Considers Antijoins	9-36
How Antijoins Work	9-37
How Antijoins Handle Nulls	9-38
Cartesian Joins	9-41
When the Optimizer Considers Cartesian Joins	9-41
How Cartesian Joins Work	9-42
Cartesian Join Controls	9-42
Join Optimizations	9-43
Bloom Filters	9-43
Purpose of Bloom Filters	9-43
How Bloom Filters Work	9-44
Bloom Filter Controls	9-45
Bloom Filter Metadata	9-45
Bloom Filters: Scenario	9-45
Partition-Wise Joins	9-47
Purpose of Partition-Wise Joins	9-47
How Partition-Wise Joins Work	9-47
In-Memory Join Groups	9-50

Part V Optimizer Statistics

10 Optimizer Statistics Concepts

Introduction to Optimizer Statistics	10-1
About Optimizer Statistics Types	10-2
Table Statistics	10-2
Permanent Table Statistics	10-2
Temporary Table Statistics	10-3
Column Statistics	10-6
Index Statistics	10-7
Types of Index Statistics	10-7
Index Clustering Factor	10-8
Effect of Index Clustering Factor on Cost: Example	10-12
System Statistics	10-12
User-Defined Optimizer Statistics	10-13
How the Database Gathers Optimizer Statistics	10-13

DBMS_STATS Package	10-13
Supplemental Dynamic Statistics	10-14
Online Statistics Gathering	10-15
Online Statistics Gathering for Bulk Loads	10-15
Online Statistics Gathering for Partition Maintenance Operations	10-17
Real-Time Statistics	10-18
When the Database Gathers Optimizer Statistics	10-25
Sources for Optimizer Statistics	10-26
SQL Plan Directives	10-26
When the Database Creates SQL Plan Directives	10-27
How the Database Uses SQL Plan Directives	10-27
SQL Plan Directive Maintenance	10-28
How the Optimizer Uses SQL Plan Directives: Example	10-29
How the Optimizer Uses Extensions and SQL Plan Directives: Example	10-33
When the Database Samples Data	10-37
How the Database Samples Data	10-39

11 Histograms

Purpose of Histograms	11-1
When Oracle Database Creates Histograms	11-1
How Oracle Database Chooses the Histogram Type	11-3
Cardinality Algorithms When Using Histograms	11-4
Endpoint Numbers and Values	11-4
Popular and Nonpopular Values	11-4
Bucket Compression	11-5
Frequency Histograms	11-5
Criteria For Frequency Histograms	11-6
Generating a Frequency Histogram	11-6
Top Frequency Histograms	11-9
Criteria For Top Frequency Histograms	11-9
Generating a Top Frequency Histogram	11-10
Height-Balanced Histograms (Legacy)	11-13
Criteria for Height-Balanced Histograms	11-13
Generating a Height-Balanced Histogram	11-14
Hybrid Histograms	11-17
How Endpoint Repeat Counts Work	11-17
Criteria for Hybrid Histograms	11-19
Generating a Hybrid Histogram	11-20

12 Configuring Options for Optimizer Statistics Gathering

Purpose of Optimizer Statistics Collection	12-1
User Interfaces for Optimizer Statistics Management	12-1
Graphical Interface for Optimizer Statistics Management	12-1
Accessing the Database Home Page in Cloud Control	12-1
Accessing the Optimizer Statistics Console	12-2
Command-Line Interface for Optimizer Statistics Management	12-2
Setting Optimizer Statistics Preferences	12-2
About Optimizer Statistics Preferences	12-2
Purpose of Optimizer Statistics Preferences	12-2
Examples of Statistics Preferences	12-3
DBMS_STATS Procedures for Setting Statistics Preferences	12-4
Statistics Preference Overrides	12-5
Setting Statistics Preferences: Example	12-7
Setting Global Optimizer Statistics Preferences Using Cloud Control	12-9
Setting Object-Level Optimizer Statistics Preferences Using Cloud Control	12-9
Setting Optimizer Statistics Preferences from the Command Line	12-10
Configuring Options for Dynamic Statistics	12-11
About Dynamic Statistics Levels	12-11
Setting Dynamic Statistics Levels Manually	12-13
Disabling Dynamic Statistics	12-15
Dynamic Statistics for PL/SQL Functions	12-16
The Use Case for PL/SQL Dynamic Statistics	12-16
Controls for Dynamic Sampling With PL/SQL Functions	12-17
Setting Global Level PL/SQL Preferences	12-18
Setting Function-Level Preferences	12-19
Setting the PLSQL_FUNCTION_DYNAMIC_STATS Session Level Parameter	12-20
Managing SQL Plan Directives	12-20

13 Gathering Optimizer Statistics

Configuring Automatic Optimizer Statistics Collection	13-1
About Automatic Optimizer Statistics Collection	13-1
Configuring Automatic Optimizer Statistics Collection Using Cloud Control	13-2
Configuring Automatic Optimizer Statistics Collection from the Command Line	13-3
Configuring High-Frequency Automatic Optimizer Statistics Collection	13-5
About High-Frequency Automatic Optimizer Statistics Collection	13-5
Setting Preferences for High-Frequency Automatic Optimizer Statistics Collection	13-6
High-Frequency Automatic Optimizer Statistics Collection: Example	13-7
Gathering Optimizer Statistics Manually	13-10
About Manual Statistics Collection with DBMS_STATS	13-10

Guidelines for Gathering Optimizer Statistics Manually	13-11
Guideline for Setting the Sample Size	13-11
Guideline for Gathering Statistics in Parallel	13-12
Guideline for Partitioned Objects	13-13
Guideline for Frequently Changing Objects	13-13
Guideline for External Tables	13-13
Determining When Optimizer Statistics Are Stale	13-14
Gathering Schema and Table Statistics	13-15
Gathering Statistics for Fixed Objects	13-16
Gathering Statistics for Volatile Tables Using Dynamic Statistics	13-17
Gathering Optimizer Statistics Concurrently	13-18
About Concurrent Statistics Gathering	13-18
Enabling Concurrent Statistics Gathering	13-20
Monitoring Statistics Gathering Operations	13-23
Gathering Incremental Statistics on Partitioned Objects	13-25
Purpose of Incremental Statistics	13-25
How DBMS_STATS Derives Global Statistics for Partitioned tables	13-25
Gathering Statistics for a Partitioned Table: Basic Steps	13-28
Maintaining Incremental Statistics for Partition Maintenance Operations	13-31
Maintaining Incremental Statistics for Tables with Stale or Locked Partition Statistics	13-34
Gathering System Statistics Manually	13-36
About System Statistics	13-36
Guidelines for Gathering System Statistics	13-37
Gathering System Statistics with DBMS_STATS	13-38
About the GATHER_SYSTEM_STATS Procedure	13-38
Gathering Workload Statistics	13-40
Gathering Noworkload Statistics	13-43
Deleting System Statistics	13-44
Running Statistics Gathering Functions in Reporting Mode	13-45

14 Managing Extended Statistics

Managing Column Group Statistics	14-1
About Statistics on Column Groups	14-1
Why Column Group Statistics Are Needed: Example	14-2
Automatic and Manual Column Group Statistics	14-4
User Interface for Column Group Statistics	14-4
Detecting Useful Column Groups for a Specific Workload	14-5
Creating Column Groups Detected During Workload Monitoring	14-8
Creating and Gathering Statistics on Column Groups Manually	14-10
Displaying Column Group Information	14-11
Dropping a Column Group	14-12

Managing Expression Statistics	14-13
About Expression Statistics	14-13
When Expression Statistics Are Useful: Example	14-14
Creating Expression Statistics	14-15
Displaying Expression Statistics	14-16
Dropping Expression Statistics	14-17

15 Controlling the Use of Optimizer Statistics

Locking and Unlocking Optimizer Statistics	15-1
Locking Statistics	15-1
Unlocking Statistics	15-2
Publishing Pending Optimizer Statistics	15-3
About Pending Optimizer Statistics	15-3
User Interfaces for Publishing Optimizer Statistics	15-5
Managing Published and Pending Statistics	15-6
Creating Artificial Optimizer Statistics for Testing	15-9
About Artificial Optimizer Statistics	15-9
Setting Artificial Optimizer Statistics for a Table	15-10
Setting Optimizer Statistics: Example	15-11

16 Managing Historical Optimizer Statistics

Restoring Optimizer Statistics	16-1
About Restore Operations for Optimizer Statistics	16-1
Guidelines for Restoring Optimizer Statistics	16-1
Restrictions for Restoring Optimizer Statistics	16-2
Restoring Optimizer Statistics Using DBMS_STATS	16-2
Managing Optimizer Statistics Retention	16-4
Obtaining Optimizer Statistics History	16-4
Changing the Optimizer Statistics Retention Period	16-5
Purging Optimizer Statistics	16-6
Reporting on Past Statistics Gathering Operations	16-7

17 Importing and Exporting Optimizer Statistics

About Transporting Optimizer Statistics	17-1
Purpose of Transporting Optimizer Statistics	17-1
How Transporting Optimizer Statistics Works	17-1
User Interface for Importing and Exporting Optimizer Statistics	17-2
Transporting Optimizer Statistics to a Test Database: Tutorial	17-3

18 Analyzing Statistics Using Optimizer Statistics Advisor

About Optimizer Statistics Advisor	18-1
Purpose of Optimizer Statistics Advisor	18-2
Problems with a Traditional Script-Based Approach	18-2
Advantages of Optimizer Statistics Advisor	18-3
Optimizer Statistics Advisor Concepts	18-3
Components of Optimizer Statistics Advisor	18-3
Operational Modes for Optimizer Statistics Advisor	18-7
Command-Line Interface to Optimizer Statistics Advisor	18-7
Basic Tasks for Optimizer Statistics Advisor	18-9
Creating an Optimizer Statistics Advisor Task	18-12
Listing Optimizer Statistics Advisor Tasks	18-13
Creating Filters for an Optimizer Advisor Task	18-13
About Filters for Optimizer Statistics Advisor	18-13
Creating an Object Filter for an Optimizer Advisor Task	18-14
Creating a Rule Filter for an Optimizer Advisor Task	18-17
Creating an Operation Filter for an Optimizer Advisor Task	18-19
Executing an Optimizer Statistics Advisor Task	18-23
Generating a Report for an Optimizer Statistics Advisor Task	18-24
Implementing Optimizer Statistics Advisor Recommendations	18-28
Implementing Actions Recommended by Optimizer Statistics Advisor	18-28
Generating a Script Using Optimizer Statistics Advisor	18-30

Part VI Optimizer Controls

19 Influencing the Optimizer

Techniques for Influencing the Optimizer	19-1
Influencing the Optimizer with Initialization Parameters	19-2
About Optimizer Initialization Parameters	19-2
Enabling Optimizer Features	19-8
Choosing an Optimizer Goal	19-9
Controlling Adaptive Optimization	19-10
Influencing the Optimizer with Hints	19-11
About Optimizer Hints	19-12
Purpose of Hints	19-12
Types of Hints	19-13
Scope of Hints	19-14
Guidelines for Join Order Hints	19-14
Reporting on Hints	19-15
Purpose of Hint Usage Reports	19-16

User Interface for Hint Usage Reports	19-16
Reporting on Hint Usage: Tutorial	19-18
Hint Usage Reports: Examples	19-20
SQL Analysis Report	19-29

20 Improving Real-World Performance Through Cursor Sharing

Overview of Cursor Sharing	20-1
About Cursors	20-1
Private and Shared SQL Areas	20-1
Parent and Child Cursors	20-3
About Cursors and Parsing	20-7
About Literals and Bind Variables	20-10
Literals and Cursors	20-10
Bind Variables and Cursors	20-12
Bind Variable Peeking	20-13
About the Life Cycle of Shared Cursors	20-16
Cursor Marked Invalid	20-16
Cursors Marked Rolling Invalid	20-18
CURSOR_SHARING and Bind Variable Substitution	20-20
CURSOR_SHARING Initialization Parameter	20-20
Parsing Behavior When CURSOR_SHARING = FORCE	20-21
Adaptive Cursor Sharing	20-23
Purpose of Adaptive Cursor Sharing	20-23
How Adaptive Cursor Sharing Works: Example	20-23
Bind-Sensitive Cursors	20-25
Bind-Aware Cursors	20-29
Cursor Merging	20-32
Adaptive Cursor Sharing Views	20-33
Real-World Performance Guidelines for Cursor Sharing	20-33
Develop Applications with Bind Variables for Security and Performance	20-33
Do Not Use CURSOR_SHARING = FORCE as a Permanent Fix	20-35
Establish Coding Conventions to Increase Cursor Reuse	20-36
Minimize Session-Level Changes to the Optimizer Environment	20-37

Part VII Monitoring and Tracing SQL

21 Gathering Diagnostic Data with SQL Test Case Builder

Purpose of SQL Test Case Builder	21-1
Concepts for SQL Test Case Builder	21-1
SQL Incidents	21-1

What SQL Test Case Builder Captures	21-2
Output of SQL Test Case Builder	21-3
User Interfaces for SQL Test Case Builder	21-5
Graphical Interface for SQL Test Case Builder	21-5
Accessing the Incident Manager	21-5
Accessing the Support Workbench	21-6
Command-Line Interface to SQL Tuning Advisor	21-6
Running SQL Test Case Builder	21-7

22 Monitoring Database Operations

About Monitoring Database Operations	22-1
About Database Operations	22-1
Purpose of Monitoring Database Operations	22-2
How Database Monitoring Works	22-3
User Interfaces for Database Operations Monitoring	22-4
Monitored SQL Executions Page in Cloud Control	22-4
DBMS_SQL_MONITOR Package	22-5
Attributes of composite Database Operations	22-6
MONITOR and NO_MONITOR Hints	22-7
Views for Monitoring and Reporting on Database Operations	22-8
Basic Tasks in Database Operations Monitoring	22-10
Enabling and Disabling Real-Time Monitoring of Database Operations	22-10
Enabling Monitoring of Database Operations at the System Level	22-10
Enabling and Disabling Monitoring of Database Operations at the Statement Level	22-11
Defining a Composite Database Operation	22-12
Generating and Accessing SQL Monitor Reports	22-15
Monitoring Database Operations: Scenarios	22-19
Reporting on a Simple Database Operation: Scenario	22-19
Reporting on Composite Database Operation: Scenario	22-22
Real-Time SQL Monitoring Across Multiple PDBs	22-24
Enabling or Disabling Real-Time SQL Monitoring on PDBs	22-25
Querying Real-Time SQL Monitor Dictionary Tables and Imported Snapshots	22-25
Views for Examining Real-Time SQL Monitoring	22-26
Reporting Tools for Real-Time SQL Monitoring	22-26
SQL History Monitoring and Reporting	22-27
Enabling and Viewing SQL History Monitoring and Reporting	22-28
DBMS_SQLTUNE APIs for Query History and Reporting	22-28

23 Performing Application Tracing

Overview of End-to-End Application Tracing	23-1
Purpose of End-to-End Application Tracing	23-1
End-to-End Application Tracing for PDBs	23-2
Tools for End-to-End Application Tracing	23-2
Overview of the SQL Trace Facility	23-3
Overview of TKPROF	23-4
Enabling Statistics Gathering for End-to-End Tracing	23-4
Enabling Statistics Gathering for a Client ID	23-4
Enabling Statistics Gathering for Services, Modules, and Actions	23-5
Enabling End-to-End Application Tracing	23-6
Enabling Tracing for a Client Identifier	23-6
Enabling Tracing for a Service, Module, and Action	23-7
Enabling Tracing for a Session	23-8
Enabling Tracing for an Instance or Database	23-9
Generating Output Files Using SQL Trace and TKPROF	23-10
Step 1: Setting Initialization Parameters for Trace File Management	23-10
Step 2: Enabling the SQL Trace Facility	23-11
Step 3: Generating Output Files with TKPROF	23-13
Step 4: Storing SQL Trace Facility Statistics	23-14
Generating the TKPROF Output SQL Script	23-14
Editing the TKPROF Output SQL Script	23-14
Querying the Output Table	23-14
Guidelines for Interpreting TKPROF Output	23-16
Guideline for Interpreting the Resolution of Statistics	23-16
Guideline for Recursive SQL Statements	23-16
Guideline for Deciding Which Statements to Tune	23-17
Guidelines for Avoiding Traps in TKPROF Interpretation	23-18
Guideline for Avoiding the Argument Trap	23-18
Guideline for Avoiding the Read Consistency Trap	23-18
Guideline for Avoiding the Schema Trap	23-18
Guideline for Avoiding the Time Trap	23-19
Application Tracing Utilities	23-20
TRCSESS	23-20
Purpose	23-20
Guidelines	23-20
Syntax	23-21
Options	23-21
Examples	23-21
TKPROF	23-22
Purpose	23-22

Guidelines	23-22
Syntax	23-23
Options	23-23
Output	23-25
Examples	23-28
Views for Application Tracing	23-32
Views Relevant for Trace Statistics	23-33
Views Related to Enabling Tracing	23-33
Controls Over Application Tracing and Access to Trace Events	23-33

Part VIII Automatic SQL Tuning

24 Capturing Workloads in SQL Tuning Sets

About SQL Tuning Sets	24-1
Purpose of SQL Tuning Sets	24-1
Concepts for SQL Tuning Sets	24-1
User Interfaces for SQL Tuning Sets	24-3
Accessing the SQL Tuning Sets Page in Cloud Control	24-3
Command-Line Interface to SQL Tuning Sets	24-3
Basic Tasks for Managing SQL Tuning Sets	24-4
Creating a SQL Tuning Set Using CREATE_SQLSET	24-6
Loading a SQL Tuning Set Using LOAD_SQLSET	24-7
Querying a SQL Tuning Set	24-8
Modifying a SQL Tuning Set Using UPDATE_SQLSET	24-11
Transporting a SQL Tuning Set	24-12
About Transporting SQL Tuning Sets	24-12
Basic Steps for Transporting SQL Tuning Sets	24-12
Basic Steps for Transporting SQL Tuning Sets When the CON_DBID Values Differ	24-13
Transporting SQL Tuning Sets with DBMS_SQLTUNE	24-14
Dropping a SQL Tuning Set Using DROP_SQLSET	24-16

25 Analyzing SQL with SQL Tuning Advisor

About SQL Tuning Advisor	25-1
Purpose of SQL Tuning Advisor	25-1
SQL Tuning Advisor Architecture	25-2
Input to SQL Tuning Advisor	25-3
Output of SQL Tuning Advisor	25-4
Automatic Tuning Optimizer Analyses	25-5
SQL Tuning Advisor Operation	25-14
Automatic and On-Demand SQL Tuning	25-14

Local and Remote SQL Tuning	25-15
Managing the Automatic SQL Tuning Task	25-20
About the Automatic SQL Tuning Task	25-20
Purpose of Automatic SQL Tuning	25-21
Automatic SQL Tuning Concepts	25-21
Command-Line Interface to SQL Tuning Advisor	25-21
Basic Tasks for Automatic SQL Tuning	25-22
Enabling and Disabling the Automatic SQL Tuning Task	25-23
Enabling and Disabling the Automatic SQL Tuning Task Using Cloud Control	25-23
Enabling and Disabling the Automatic SQL Tuning Task from the Command Line	25-24
Configuring the Automatic SQL Tuning Task	25-26
Configuring the Automatic SQL Tuning Task Using Cloud Control	25-26
Configuring the Automatic SQL Tuning Task Using the Command Line	25-27
Viewing Automatic SQL Tuning Reports	25-28
Viewing Automatic SQL Tuning Reports Using the Command Line	25-29
Running SQL Tuning Advisor On Demand	25-31
About On-Demand SQL Tuning	25-32
Purpose of On-Demand SQL Tuning	25-32
User Interfaces for On-Demand SQL Tuning	25-32
Basic Tasks in On-Demand SQL Tuning	25-34
Creating a SQL Tuning Task	25-35
Configuring a SQL Tuning Task	25-37
Executing a SQL Tuning Task	25-39
Monitoring a SQL Tuning Task	25-40
Displaying the Results of an SQL Tuning Task	25-41
The Automatic SQL Tuning Set	25-43

26 Optimizing Access Paths with SQL Access Advisor

About SQL Access Advisor	26-1
Purpose of SQL Access Advisor	26-1
SQL Access Advisor Architecture	26-2
Input to SQL Access Advisor	26-3
Filter Options for SQL Access Advisor	26-3
SQL Access Advisor Recommendations	26-4
SQL Access Advisor Actions	26-5
SQL Access Advisor Repository	26-6
User Interfaces for SQL Access Advisor	26-7
Accessing the SQL Access Advisor: Initial Options Page Using Cloud Control	26-7
Command-Line Interface to SQL Tuning Sets	26-8
Using SQL Access Advisor: Basic Tasks	26-8
Creating a SQL Tuning Set as Input for SQL Access Advisor	26-10

Populating a SQL Tuning Set with a User-Defined Workload	26-11
Creating and Configuring a SQL Access Advisor Task	26-13
Executing a SQL Access Advisor Task	26-15
Viewing SQL Access Advisor Task Results	26-16
Generating and Executing a Task Script	26-20
Performing a SQL Access Advisor Quick Tune	26-21
Using SQL Access Advisor: Advanced Tasks	26-22
Evaluating Existing Access Structures	26-22
Updating SQL Access Advisor Task Attributes	26-23
Creating and Using SQL Access Advisor Task Templates	26-24
Terminating SQL Access Advisor Task Execution	26-26
Interrupting SQL Access Advisor Tasks	26-26
Canceling SQL Access Advisor Tasks	26-27
Deleting SQL Access Advisor Tasks	26-28
Marking SQL Access Advisor Recommendations	26-29
Modifying SQL Access Advisor Recommendations	26-30
SQL Access Advisor Examples	26-31
SQL Access Advisor Reference	26-31
Action Attributes in the DBA_ADVISOR_ACTIONS View	26-31
Categories for SQL Access Advisor Task Parameters	26-33
SQL Access Advisor Constants	26-33

Part IX SQL Management Objects

27 Managing SQL Profiles

About SQL Profiles	27-1
Purpose of SQL Profiles	27-1
Concepts for SQL Profiles	27-2
Statistics in SQL Profiles	27-2
SQL Profiles and Execution Plans	27-2
SQL Profile Recommendations	27-3
SQL Profiles and SQL Plan Baselines	27-5
User Interfaces for SQL Profiles	27-6
Basic Tasks for SQL Profiles	27-6
Implementing a SQL Profile	27-7
About SQL Profile Implementation	27-7
Implementing a SQL Profile	27-8
Listing SQL Profiles	27-9
Altering a SQL Profile	27-10
Dropping a SQL Profile	27-11

28 Overview of SQL Plan Management

About SQL Plan Baselines	28-1
Purpose of SQL Plan Management	28-1
Benefits of SQL Plan Management	28-1
Differences Between SQL Plan Baselines and SQL Profiles	28-2
Plan Capture	28-3
Automatic Initial Plan Capture	28-3
Eligibility for Automatic Initial Plan Capture	28-4
Plan Matching for Automatic Initial Plan Capture	28-5
Manual Plan Capture	28-5
Plan Selection	28-7
Plan Evolution	28-8
Purpose of Plan Evolution	28-8
How Plan Evolution Works	28-8
PL/SQL Subprograms for Plan Evolution	28-9
About Automatic SQL Plan Management	28-10
Use Cases for Automatic SQL Plan Management	28-10
SQL Plan Management (With Background Verification)	28-10
About Real-Time SQL Plan Management	28-10
Storage Architecture for SQL Plan Management	28-11
SQL Management Base	28-11
SQL Statement Log	28-12
SQL Plan History	28-13
Enabled Plans	28-14
Accepted Plans	28-14
Fixed Plans	28-14

29 Managing SQL Plan Baselines

About Managing SQL Plan Baselines	29-1
User Interfaces for SQL Plan Management	29-1
Accessing the SQL Plan Baseline Page in Cloud Control	29-1
DBMS_SPM Package	29-2
Basic Tasks in SQL Plan Management	29-3
Configuring the Capture and Use of SQL Plan Baselines	29-4
Enabling Automatic Initial Plan Capture for SQL Plan Management	29-5
Configuring Filters for Automatic Plan Capture	29-6
Disabling All SQL Plan Baselines	29-8
Configuring SQL Plan Management	29-8

Managing the SPM Evolve Advisor Task	29-9
About the SPM Evolve Advisor Task	29-9
Enabling and Disabling the Automatic SPM Evolve Advisor Task	29-9
Configuring the Automatic SPM Evolve Advisor Task	29-10
Automatic SQL Plan Management	29-13
Configuring Automatic SQL Plan Management	29-15
Enabling Automatic SQL Plan Management: Tutorial	29-15
Monitoring and Reporting in Automatic SQL Plan Management	29-17
Displaying Plans in a SQL Plan Baseline	29-17
Loading SQL Plan Baselines	29-18
About Loading SQL Plan Baselines	29-18
Loading Plans from AWR	29-19
Loading Plans from the Shared SQL Area	29-22
Loading Plans from a SQL Tuning Set	29-24
Loading Plans from a Staging Table	29-26
Evolving SQL Plan Baselines Manually	29-29
About the DBMS_SPM Evolve Functions	29-29
Managing an Evolve Task	29-31
Adding SQL Plans to a Baseline	29-39
Dropping SQL Plan Baselines	29-40
Managing the SQL Management Base	29-42
About Managing the SMB	29-42
Changing the Disk Space Limit for the SMB	29-43
Changing the Plan Retention Policy in the SMB	29-44

30 Migrating Stored Outlines to SQL Plan Baselines

About Stored Outline Migration	30-1
Purpose of Stored Outline Migration	30-1
How Stored Outline Migration Works	30-2
Stages of Stored Outline Migration	30-2
Outline Categories and Baseline Modules	30-3
User Interface for Stored Outline Migration	30-4
Basic Steps in Stored Outline Migration	30-6
Preparing for Stored Outline Migration	30-6
Migrating Outlines to Utilize SQL Plan Management Features	30-7
Migrating Outlines to Preserve Stored Outline Behavior	30-8
Performing Follow-Up Tasks After Stored Outline Migration	30-9

Glossary

Preface

This manual explains how to tune Oracle SQL.

Audience

This document is intended for database administrators and application developers who perform the following tasks:

- Generating and interpreting SQL execution plans
- Managing optimizer statistics
- Influencing the optimizer through initialization parameters or SQL hints
- Controlling cursor sharing for SQL statements
- Monitoring SQL execution
- Performing application tracing
- Managing SQL tuning sets
- Using SQL Tuning Advisor or SQL Access Advisor
- Managing SQL profiles
- Managing SQL baselines

Documentation Accessibility

For information about Oracle's commitment to accessibility, visit the Oracle Accessibility Program website at <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc>.

Access to Oracle Support

Oracle customer access to and use of Oracle support services will be pursuant to the terms and conditions specified in their Oracle order for the applicable services.

Related Documents

This manual assumes that you are familiar with *Oracle Database Concepts*. The following books are frequently referenced:

- *Oracle Database Data Warehousing Guide*
- *Oracle Database VLDB and Partitioning Guide*
- *Oracle Database SQL Language Reference*
- *Oracle Database Reference*

Many examples in this book use the sample schemas, which are installed by default when you select the Basic Installation option with an Oracle Database. See *Oracle Database Sample Schemas* for information on how these schemas were created and how you can use them.

Conventions

The following text conventions are used in this document:

Convention	Meaning
boldface	Boldface type indicates graphical user interface elements associated with an action, or terms defined in text or the glossary.
<i>italic</i>	Italic type indicates book titles, emphasis, or placeholder variables for which you supply particular values.
<code>monospace</code>	Monospace type indicates commands within a paragraph, URLs, code in examples, text that appears on the screen, or text that you enter.