Creating Your Own DBFS Store

You can create your own DBFS Store using DBFS Content Store Provider Interface (DBMS DBFS CONTENT SPI).

· Overview of DBFS Store Creation and Use

In order to customize a DBFS store, you must implement the DBFS Content SPI (DBMS_DBFS_CONTENT_SPI). It is the basis for existing stores such as the DBFS SecureFiles Store and the DBFS Hierarchical Store, as well as any user-defined DBFS stores that you create.

- DBFS Content Store Provider Interface (DBFS Content SPI)
 The DBFS Content SPI (Store Provider Interface) is a specification only and has no package body.
- Creating a Custom Store Provider You can use this example store provider for DBFS, TaBleFileSystem Store Provider ("tbfs"), as a skeleton for custom providers or as a learning tool, to become familiar with the DBFS and its SPI.

23.1 Overview of DBFS Store Creation and Use

In order to customize a DBFS store, you must implement the DBFS Content SPI (DBMS_DBFS_CONTENT_SPI). It is the basis for existing stores such as the DBFS SecureFiles Store and the DBFS Hierarchical Store, as well as any user-defined DBFS stores that you create.

Client-side applications, such the PL/SQL interface, invoke functions and procedures in the DBFS Content API. The DBFS Content API then invokes corresponding subprograms in the DBFS Content SPI to create stores and perform other related functions.

Once you create your DBFS store, you use it much the same way that you would a SecureFiles Store.

See Also:

- DBFS Content API
- DBFS SecureFiles Store

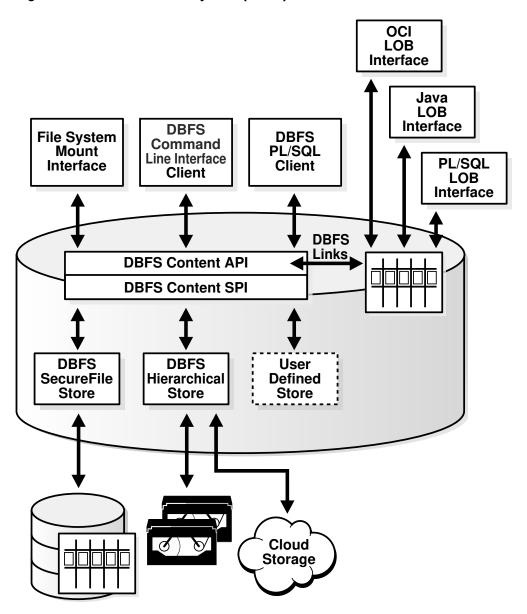


Figure 23-1 Database File System (DBFS)

23.2 DBFS Content Store Provider Interface (DBFS Content SPI)

The DBFS Content SPI (Store Provider Interface) is a specification only and has no package body.

You must implement the package body in order to respond to calls from the DBFS Content API. In other words, DBFS Content SPI is a collection of required program specifications which you must implement using the method signatures and semantics indicated.

You may add additional functions and procedures to the DBFS Content SPI package body as needed. Your implementation may implement other methods and expose other interfaces, but the DBFS Content API will not use these interfaces.

The DBFS Content SPI references various elements such as constants, types, and exceptions defined by the DBFS Content API (package DBMS DBFS CONTENT).

Note that all path name references must be store-qualified, that is, the notion of mount points and full absolute path names has been normalized and converted to store-qualified path names by the DBFS Content API before it invokes any of the Provider SPI methods.

Because the DBFS Content API and SPI implementation is a one-to-many pluggable architecture, the DBFS Content API uses dynamic SQL to invoke methods in the SPI implementation; this may lead to run time errors if your SPI implementation does not follow the specification of SPI implementation given in this document.

There are no explicit initial or final methods to indicate when the DBFS Content API plugs and unplugs a particular SPI implementation. SPI implementations must be able to auto-initialize themselves at any SPI entry point.

See Also:

- Oracle Database PL/SQL Packages and Types Reference for syntax of the DBMS DBFS CONTENT SPI package
- See the file <code>\$ORACLE HOME/rdbms/admin/dbmscapi.sql</code> for more information

23.3 Creating a Custom Store Provider

You can use this example store provider for DBFS, TaBleFileSystem Store Provider ("tbfs"), as a skeleton for custom providers or as a learning tool, to become familiar with the DBFS and its SPI.

This example store provider for DBFS, exposes a relational table containing a BLOB column as a flat, non-hierarchical filesystem, that is, a collection of named files.

To use this example, it is assumed that you have installed the Oracle Database 12c and are familiar with DBFS concepts, and have installed and used <code>dbfs_client</code> and <code>FUSE</code> to mount and access filesystems backed by the standard SFS store provider.

The TaBleFileSystem Store Provider ("tbfs") does not aim to be feature-rich or even complete, it does however provide a sufficient demonstration of what it takes for users of DBFS to write their own custom providers that expose their table(s) through <code>dbfs_client</code> to traditional filesystem programs.

Installation and Setup

You will need certain files for installation and setup of the DBFS TaBleFileSystem Store Provider ("tbfs").

TBFS Use

Once the example store provider for DBFS, TaBleFileSystem Store Provider ("tbfs") is installed, files can be added or removed in several different ways and other changes can be made to the TBFS.

TBFS Internals

The TBFS is simple because its primary purpose is to serve as a teaching and learning example.

Example Scripts

This section describes some example SQL scripts.



23.3.1 Installation and Setup

You will need certain files for installation and setup of the DBFS TaBleFileSystem Store Provider ("tbfs").

The TBFS consists of the following SQL files:

script to create a test user, tablespace, the table backing the filesystem, and son. spec.sql the SPI specification of the tbfs body.sql the SPI implementation of the tbfs capi.sql DBFS register/mount script	tbfs.sql	top-level driver script
body.sql the SPI implementation of the tbfs	-	script to create a test user, tablespace, the table backing the filesystem, and so
	spec.sql	the SPI specification of the tbfs
capi.sql DBFS register/mount script	body.sql	the SPI implementation of the tbfs
	capi.sql	DBFS register/mount script

To install the TBFS, just run tbfs.sql as SYSDBA, in the directory that contains all of the above files. tbfs.sql will load the other SQL files in the proper sequence.

Ignoring any name conflicts, all of the SQL files should load without any compilation errors. All SQL files should also load without any run time errors, depending on the value of the "plsql_warnings" init.ora parameter, you may see various innocuous warnings.

If there are any name conflicts (tablespace name TBFS, datafile name"tbfs.f", user name TBFS, package name TBFS), the appropriate references in the various SQL files must be changed consistently.

23.3.2 TBFS Use

Once the example store provider for DBFS, TaBleFileSystem Store Provider ("tbfs") is installed, files can be added or removed in several different ways and other changes can be made to the TBFS.

A dbfs_client connected as user TBFS will see a simple, non-hierarchical, filesystem backed by an RDBMS table (TBFS.TBFST).

Files can be added or removed from this filesystem through SQL (that is, through DML on the underlying table), through Unix utilities (mediated by $dbfs_client$), or through PL/SQL (using the DBFS APIs).

Changes to the filesystem made through any of the access methods will be visible, in a transactionally consistent manner (that is, at commit/rollback boundaries) to all of the other access methods.

23.3.3 TBFS Internals

The TBFS is simple because its primary purpose is to serve as a teaching and learning example.

However, the implementation shows the path towards a robust, production-quality custom SPI that can plug into the DBFS, and expose existing relational data as Unix filesystems.

The TBFS makes various simplifications in order to remain concise (however, these should not be taken as inviolable limitations of DBFS or the SPI):



- The TBFS SPI package handles only a single table with a hard-coded name (TBFS.TBFST). It is possible to use dynamic SQL and additional configuration information to have a single SPI package support multiple tables, each as a separate filesystem (or even to unify data in multiple tables into a single filesystem).
- The TBFS does not support filesystem hierarchies; it imposes a flat namespace: a
 collection of files, identified by a simple item name, under a virtual "/" root directory.
 Implementing directory hierarchies is significantly more complex because it requires the
 store provider to manage parent/child relationships in a consistent manner.
 - Moreover, existing relational data (the kind of data that TBFS is attempting to expose as a filesystem) does not typically have inter-row relationships that form a natural directory/file hierarchy.
- Because the TBFS supports only a flat namespace, most methods in the SPI are
 unimplemented, and the method bodies raise a
 dbms_dbfs_content.unsupported_operation exception. This exception is also a good
 starting point for you to write your own custom SPI. You can start with a simple SPI
 skeleton cloned from the DBMS_DBFS_CONTENT_SPI package, default all method bodies to
 ones that raise this exception, and subsequently fill in more realistic implementations
 incrementally.
- The table underlying the TBFS is close to being the simplest possible structure (a key/ name column and a LOB column). This means that various properties used or expected by DBFS and dbfs_client must be generated dynamically (the TBFS implementation shows how this is done for the std:quid property).
 - Other properties (such as Unix-style timestamps) are not implemented at all. This still allows a surprisingly functional filesystem to be implemented, but when you write your own custom SPIs, you can easily incorporate support for additional DBFS properties by expanding the structure of their underlying table(s) to include additional columns as needed, or by using existing columns in their existing tables to provide the values for these DBFS properties.
- The TBFS does not implement a rename/move method; adding support for this (a suitable UPDATE statement in the renamePath method) is left as an exercise for the user.
- The TBFS example uses the string "tbfs" in multiple places (tablespace, datafile, user, package, and even filesystem name). All these uses of "tbfs" belong in different namespaces—identifying which namespace corresponds to a specific occurrence of the string. "tbfs" in these examples is also a good learning exercise to make sure that the DBFS concepts are clear in your mind.

23.3.4 Example Scripts

This section describes some example SQL scripts.

- Driver Script
 - The TBFS.SQL script is the top level driver script.
- Creating a Test User, Tablespace and Table to Backup Filesystem
 The TBL.SQL script creates a test user, a tablespace, the table that backs the filesystem and so on.
- Providing SPI Specification
 - The spec.sql script provide the SPI specification of the tbfs.
- SPI Implementation of tbfs
 - The body.sql script provides the SPI implementation of the tbfs.

Registering and Mounting the DBFS

The capi.sql script registers and mounts the DBFS.

23.3.4.1 Driver Script

The TBFS.SQL script is the top level driver script.

The TBFS.SQL script:

```
set echo on;

@tbl
@spec
@body
@capi
quit;
```

23.3.4.2 Creating a Test User, Tablespace and Table to Backup Filesystem

The TBL.SQL script creates a test user, a tablespace, the table that backs the filesystem and so on.

The TBL.SQL script:

```
connect / as sysdba
create tablespace tbfs datafile 'tbfs.f' size 100m
    reuse autoextend on
    extent management local
   segment space management auto;
create user tbfs identified by tbfs;
alter user tbfs default tablespace tbfs;
grant connect, resource, dbfs_role to tbfs;
connect tbfs/tbfs;
drop table tbfst;
purge recyclebin;
create table tbfst(
          varchar2(256)
    key
           primary key
                           (instr(key, '/') = 0),
           check
          blob)
    data
       tablespace tbfs
    lob(data)
        store as securefile
            (tablespace tbfs);
grant select on tbfst to dbfs role;
grant insert on tbfst to dbfs role;
grant delete on tbfst to dbfs role;
grant update on tbfst to dbfs role;
```

23.3.4.3 Providing SPI Specification

The spec.sql script provide the SPI specification of the tbfs.

```
The spec.sql script:
connect / as sysdba;
create or replace package tbfs
  authid current user
    * Lookup store features (see dbms dbfs content.feature XXX). Lookup
    * store id.
    * A store ID identifies a provider-specific store, across
    * registrations and mounts, but independent of changes to the store
    * contents.
    ^{\star} I.e. changes to the store table(s) should be reflected in the
    * store ID, but re-initialization of the same store table(s) should
    * preserve the store ID.
    * Providers should also return a "version" (either specific to a
    * provider package, or to an individual store) based on a standard
     * <a.b.c> naming convention (for <major>, <minor>, and <patch>
     * components).
    */
   function getFeatures(
       store name in varchar2)
          return integer;
   function getStoreId(
       store_name in
    return number;
                                 varchar2)
   function getVersion(
       store_name in varchar2)
          return varchar2;
    * Lookup pathnames by (store name, std guid) or (store mount,
    * std guid) tuples.
    * If the underlying "std guid" is found in the underlying store,
     * this function returns the store-qualified pathname.
     * If the "std guid" is unknown, a "null" value is returned. Clients
     * are expected to handle this as appropriate.
     */
    function getPathByStoreId(
       store_name in varchar2, guid in integer)
          return varchar2;
```

```
* DBFS SPI: space usage.
^{\star} Clients can query filesystem space usage statistics via the
 * "spaceUsage()" method. Providers are expected to support this
 ^{\star} method for their stores (and to make a best effort determination
 * of space usage---esp. if the store consists of multiple
 * tables/indexes/lobs, etc.).
 * "blksize" is the natural tablespace blocksize that holds the
 * store---if multiple tablespaces with different blocksizes are
 * used, any valid blocksize is acceptable.
 ^{\star} "tbytes" is the total size of the store in bytes, and "fbytes" is
 * the free/unused size of the store in bytes. These values are
 * computed over all segments that comprise the store.
* "nfile", "ndir", "nlink", and "nref" count the number of
 * currently available files, directories, links, and references in
 * the store.
 * Since database objects are dynamically growable, it is not easy
 * to estimate the division between "free" space and "used" space.
 * /
procedure spaceUsage(
   store name in
                               varchar2,
                               integer,
   blksize out
   tbytes out fbytes out nfile out
                                integer,
                                integer,
                                integer,
   ndir
              out
                                integer,
   nlink out nref out
                               integer,
                               integer);
/*
 * DBFS SPI: notes on pathnames.
* All pathnames used in the SPI are store-qualified, i.e. a 2-tuple
 * of the form (store name, pathname) (where the pathname is rooted
 * within the store namespace).
 * Stores/providers that support contentID-based access (see
 * "feature content id") also support a form of addressing that is
 * not based on pathnames. Items are identified by an explicit store
 * name, a "null" pathname, and possibly a contentID specified as a
 * parameter or via the "opt content id" property.
 * Not all operations are supported with contentID-based access, and
 * applications should depend only on the simplest create/delete
 * functionality being available.
* DBFS SPI: creation operations
```

```
* The SPI must allow the DBFS API to create directory, file, link,
 * and reference elements (subject to store feature support).
 * All of the creation methods require a valid pathname (see the
 * special exemption for contentID-based access below), and can
 * optionally specify properties to be associated with the pathname
 * as it is created. It is also possible for clients to fetch-back
 * item properties after the creation completes (so that
 * automatically generated properties (e.g. "std_creation_time") are
 * immediately available to clients (the exact set of properties
 * fetched back is controlled by the various "prop xxx" bitmasks in
 * "prop flags").
^{\star} Links and references require an additional pathname to associate
 * with the primary pathname.
* File pathnames can optionally specify a BLOB value to use to
 * initially populate the underlying file content (the provided BLOB
 * may be any valid lob: temporary or permanent). On creation, the
 * underlying lob is returned to the client (if "prop data" is
 * specified in "prop flags").
 ^{\star} Non-directory pathnames require that their parent directory be
 * created first. Directory pathnames themselves can be recursively
 * created (i.e. the pathname hierarchy leading up to a directory
 * can be created in one call).
 * Attempts to create paths that already exist is an error; the one
 * exception is pathnames that are "soft-deleted" (see below for
 * delete operations) --- in these cases, the soft-deleted item is
 * implicitly purged, and the new item creation is attempted.
^{\star} Stores/providers that support contentID-based access accept an
 * explicit store name and a "null" path to create a new element.
 * The contentID generated for this element is available via the
 * "opt content id" property (contentID-based creation automatically
 * implies "prop opt" in "prop flags").
 * The newly created element may also have an internally generated
 * pathname (if "feature lazy path" is not supported) and this path
 * is available via the "std canonical path" property.
 * Only file elements are candidates for contentID-based access.
 */
procedure createFile(
   store name in
                               varchar2,
   path in
                               varchar2,
   properties in out nocopy dbms dbfs content properties t,
   content in out nocopy blob,
   prop flags in
                               integer,
   ctx
              in
                               dbms_dbfs_content_context_t);
procedure createLink(
   store name in
                               varchar2,
   srcPath in
                              varchar2,
```

```
varchar2,
   dst.Pat.h
             in
   properties in out nocopy dbms_dbfs_content_properties_t,
   prop_flags in
                              integer,
                              dbms dbfs content context t);
              in
   ctx
procedure createReference(
   store name in
                              varchar2,
   srcPath in dstPath in
                              varchar2,
                              varchar2,
              in
   properties in out nocopy dbms_dbfs_content_properties_t,
   prop_flags in
                              integer,
   ctx
          in
                              dbms dbfs content context t);
procedure createDirectory(
   store name in
                            varchar2,
   path
          in
                            varchar2,
   properties in out nocopy dbms dbfs content properties t,
   prop flags in
                             integer,
   recurse in
                              integer,
              in
                              dbms dbfs content context t);
 * DBFS SPI: deletion operations
 * The SPI must allow the DBFS API to delete directory, file, link,
 * and reference elements (subject to store feature support).
 * By default, the deletions are "permanent" (get rid of the
 * successfully deleted items on transaction commit), but stores may
 * also support "soft-delete" features. If requested by the client,
 * soft-deleted items are retained by the store (but not typically
 * visible in normal listings or searches).
 * Soft-deleted items can be "restore"d, or explicitly purged.
^{\star} Directory pathnames can be recursively deleted (i.e. the pathname
 * hierarchy below a directory can be deleted in one call).
 * Non-recursive deletions can be performed only on empty
 * directories. Recursive soft-deletions apply the soft-delete to
 * all of the items being deleted.
 * Individual pathnames (or all soft-deleted pathnames under a
 * directory) can be restored or purged via the restore and purge
 * methods.
 * Providers that support filtering can use the provider "filter" to
 * identify subsets of items to delete---this makes most sense for
 * bulk operations (deleteDirectory, restoreAll, purgeAll), but all
 * of the deletion-related operations accept a "filter" argument.
 * Stores/providers that support contentID-based access can also
 * allow file items to be deleted by specifying their contentID.
 */
```

```
procedure deleteFile(
  store name in
                          varchar2,
   path in filter in
                           varchar2,
                           varchar2,
   soft delete in
                           integer,
                          dbms dbfs_content_context_t);
   ctx in
procedure deleteContent(
   store name in
                           varchar2,
   contentID in
                           raw,
   filter in
                           varchar2,
   soft delete in
                           integer,
   ctx in
                           dbms dbfs_content_context_t);
procedure deleteDirectory(
   store name in
                           varchar2,
  path in filter in
                          varchar2,
                          varchar2,
   soft delete in
                          integer,
   recurse in
                          integer,
   ctx
            in
                           dbms dbfs content context t);
procedure restorePath(
   store name in
                          varchar2,
   path in
                          varchar2,
   filter
            in
                           varchar2,
           in
                           dbms dbfs content_context_t);
   ctx
procedure purgePath(
   store name in
                           varchar2,
   path in filter in
                           varchar2,
             in
                           varchar2,
                           dbms dbfs content context t);
   ctx
            in
procedure restoreAll(
   store name in
                          varchar2,
   path in
                           varchar2,
   filter
            in
                          varchar2,
   ctx
            in
                          dbms dbfs content context t);
procedure purgeAll(
   store name in
                          varchar2,
   path in
                          varchar2,
   filter in
                          varchar2,
   ctx in
                          dbms_dbfs_content_context_t);
* DBFS SPI: path get/put operations.
* Existing path items can be accessed (for query or for update) and
* modified via simple get/put methods.
^{\star} All pathnames allow their metadata (i.e. properties) to be
* read/modified. On completion of the call, the client can request
* (via "prop flags") specific properties to be fetched as well.
^{\star} File pathnames allow their data (i.e. content) to be
^{\star} read/modified. On completion of the call, the client can request
* (via the "prop data" bitmaks in "prop flags") a new BLOB locator
* that can be used to continue data access.
```

```
* Files can also be read/written without using BLOB locators, by
 * explicitly specifying logical offsets/buffer-amounts and a
 * suitably sized buffer.
 ^{\star} Update accesses must specify the "forUpdate" flag. Access to link
 * pathnames can be implicitly and internally deferenced by stores
 * (subject to feature support) if the "deref" flag is
 * specified---however, this is dangerous since symbolic links are
 * not always resolvable.
* The read methods (i.e. "getPath" where "forUpdate" is "false"
* also accepts a valid "asof" timestamp parameter that can be used
 * by stores to implement "as of" style flashback queries. Mutating
 * versions of the "getPath" and the "putPath" methods do not
 * support as-of modes of operation.
 * "getPathNowait" implies a "forUpdate", and, if implemented (see
 * "feature nowait"), allows providers to return an exception
 * (ORA-54) rather than wait for row locks.
 */
procedure getPath(
   store name in
                             varchar2,
   path in
                             varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   content out nocopy blob, item_type out integ
                             integer,
   prop_flags in
                             integer,
                            integer,
   forUpdate in
   deref in
                            integer,
             in
                            dbms dbfs content context t);
   ctx
procedure getPathNowait(
   store name in
                            varchar2,
   path in
                            varchar2,
   properties in out nocopy dbms dbfs content properties t,
   content out nocopy blob,
   item type out
                             integer,
   prop flags in
                            integer,
   deref in
                            integer,
            in
                            dbms dbfs content context t);
   ctx
procedure getPath(
   store name in
                            varchar2,
   path in
                             varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   amount in out
                             number,
   offset in numb
buffer out nocopy raw,
                             number,
   prop_flags in
                             integer,
            in
   ctx
                             dbms dbfs content context t);
procedure getPath(
   store name in
                            varchar2,
   path in
                             varchar2,
   properties in out nocopy dbms dbfs content properties t,
   amount in out
                           number,
```

```
offset in number, buffers out nocopy dbms_dbfs_content_raw_t,
   prop_flags in
                            integer,
         in
                             dbms dbfs content context t);
procedure putPath(
   store name in
                            varchar2,
   path
             in
                            varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   content in out nocopy blob,
   item_type out integer,
   prop_flags in
                            integer,
   ctx
        in
                            dbms dbfs content context t);
procedure putPath(
   store name in
                            varchar2,
   path in
                            varchar2,
   properties in out nocopy dbms dbfs content properties t,
   amount in
                           number,
   offset
             in
                           number,
   buffer in
                            raw,
   prop flags in
                            integer,
   ctx in
                            dbms dbfs content context t);
procedure putPath(
   store name in
                            varchar2,
   path in
                            varchar2,
   properties in out nocopy dbms dbfs content properties t,
   written out
                            number,
   offset
              in
                            number,
   buffers in
                           dbms_dbfs_content raw t,
   prop_flags in
                            integer,
                            dbms dbfs content context t);
   ctx in
/*
 * DBFS SPI: rename/move operations.
* Pathnames can be renamed or moved, possibly across directory
 * hierarchies and mount-points, but within the same store.
 * Non-directory pathnames previously accessible via "oldPath" are
 * renamed as a single item subsequently accessible via "newPath";
 * assuming that "newPath" does not already exist.
 * If "newPath" exists and is not a directory, the rename implicitly
 * deletes the existing item before renaming "oldPath". If "newPath"
 * exists and is a directory, "oldPath" is moved into the target
 * directory.
 * Directory pathnames previously accessible via "oldPath" are
 * renamed by moving the directory and all of its children to
 * "newPath" (if it does not already exist) or as children of
 * "newPath" (if it exists and is a directory).
 * Stores/providers that support contentID-based access and lazy
 * pathname binding also support the "setPath" method that
 * associates an existing "contentID" with a new "path".
```

```
procedure renamePath(
   store_name in
                            varchar2,
   oldPath in newPath in
                           varchar2,
              in varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
                            dbms dbfs content context t);
        in
procedure setPath(
   store name in
                             varchar2,
   contentID in
                             raw,
   path
         in
                             varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   ctx in dbms dbfs content context t);
* DBFS SPI: directory navigation and search.
* The DBFS API can list or search the contents of directory
* pathnames, optionally recursing into sub-directories, optionally
 * seeing soft-deleted items, optionally using flashback "as of" a
 * provided timestamp, and optionally filtering items in/out within
 * the store based on list/search predicates.
 */
function list(
                            varchar2,
   store name in
   path in filter in
                            varchar2,
                            varchar2,
             in
   recurse in ctx in
                           integer,
                            dbms dbfs content context t)
      return dbms dbfs content list items t
         pipelined;
function search(
   store name in
                            varchar2,
   path in
                            varchar2,
   filter
             in
                           varchar2,
   recurse in ctx in
                       integer,
dbms_dbfs_content_context_t)
      return dbms dbfs content list items t
          pipelined;
* DBFS SPI: locking operations.
* Clients of the DBFS API can apply user-level locks to any valid
 ^{\star} pathname (subject to store feature support), associate the lock
 * with user-data, and subsequently unlock these pathnames.
 * The status of locked items is available via various optional
 * properties (see "opt lock*" above).
 * It is the responsibility of the store (assuming it supports
```

```
* user-defined lock checking) to ensure that lock/unlock operations
     * are performed in a consistent manner.
     * /
   procedure lockPath(
       store_name in
                                  varchar2,
                                 varchar2,
integer,
varchar2,
       path in
       lock_type in lock_data in
       ctx
               in
                                  dbms_dbfs_content_context_t);
   procedure unlockPath(
       store name in
                                  varchar2,
       path in ctx in
                                 varchar2,
                                 dbms dbfs content context t);
    * DBFS SPI: access checks.
    * Check if a given pathname (store name, path, pathtype) can be
     * manipulated by "operation (see the various
     * "dbms_dbfs_content.op_xxx" opcodes) by "principal".
    ^{\star} This is a convenience function for the DBFS API; a store that
     ^{\star} supports access control still internally performs these checks to
     * guarantee security.
     */
   function checkAccess(
                                  varchar2,
       store_name in
                                 varchar2,
       path in
      pathtype in operation in principal in
                                 integer,
varchar2,
                                 varchar2)
          return integer;
end;
show errors;
create or replace public synonym tbfs
   for sys.tbfs;
grant execute on tbfs
   to dbfs role;
```

23.3.4.4 SPI Implementation of tbfs

/

The body.sql script provides the SPI implementation of the tbfs.

```
The body.sql script:
connect / as sysdba;
create or replace package body tbfs
as
```

```
* Lookup store features (see dbms dbfs content.feature XXX). Lookup
 * store id.
 * A store ID identifies a provider-specific store, across
 * registrations and mounts, but independent of changes to the store
 * contents.
 * I.e. changes to the store table(s) should be reflected in the
 * store ID, but re-initialization of the same store table(s) should
 * preserve the store ID.
 * Providers should also return a "version" (either specific to a
 * provider package, or to an individual store) based on a standard
 * <a.b.c> naming convention (for <major>, <minor>, and <patch>
 * components).
 */
function getFeatures(
   store name in varchar2)
      return integer
is
begin
   return dbms dbfs content.feature locator;
end;
function
         getStoreId(
   store name in varchar2)
      return number
is
begin
  return 1;
end;
function getVersion(
  store name in varchar2)
      return varchar2
is
begin
  return '1.0.0';
 * Lookup pathnames by (store name, std guid) or (store mount,
 * std_guid) tuples.
 * If the underlying "std guid" is found in the underlying store,
 * this function returns the store-qualified pathname.
 * If the "std guid" is unknown, a "null" value is returned. Clients
 * are expected to handle this as appropriate.
 */
function getPathByStoreId(
   store_name in varchar2,
                     in
   guid
                            integer)
     return varchar2
is
```

```
raise dbms dbfs content.unsupported operation;
* DBFS SPI: space usage.
* Clients can query filesystem space usage statistics via the
 * "spaceUsage()" method. Providers are expected to support this
 * method for their stores (and to make a best effort determination
 * of space usage---esp. if the store consists of multiple
 * tables/indexes/lobs, etc.).
 ^{\star} "blksize" is the natural tablespace blocksize that holds the
 * store---if multiple tablespaces with different blocksizes are
* used, any valid blocksize is acceptable.
* "tbytes" is the total size of the store in bytes, and "fbytes" is
 * the free/unused size of the store in bytes. These values are
 * computed over all segments that comprise the store.
 * "nfile", "ndir", "nlink", and "nref" count the number of
 * currently available files, directories, links, and references in
 * the store.
 * Since database objects are dynamically growable, it is not easy
 * to estimate the division between "free" space and "used" space.
 */
procedure spaceUsage(
   store_name in
                               varchar2,
   blksize out
                              integer,
   tbytes out
fbytes out
nfile out
ndir out
nlink out
nref out
                               integer,
                               integer,
                               integer,
                               integer,
                               integer,
   nref
              out
                               integer)
   nblks
               number;
begin
   select count(*) into nfile
       from tbfs.tbfst:
   ndir := 0;
   nlink := 0;
   nref := 0;
   select sum(bytes) into tbytes
      from user segments;
    select sum(blocks) into nblks
      from user segments;
   blksize := tbytes/nblks;
    fbytes := 0;
                                                      /* change as needed */
end;
* DBFS SPI: notes on pathnames.
```

```
* All pathnames used in the SPI are store-qualified, i.e. a 2-tuple
* of the form (store_name, pathname) (where the pathname is rooted
* within the store namespace).
* Stores/providers that support contentID-based access (see
* "feature content id") also support a form of addressing that is
* not based on pathnames. Items are identified by an explicit store
* name, a "null" pathname, and possibly a contentID specified as a
 parameter or via the "opt content id" property.
* Not all operations are supported with contentID-based access, and
* applications should depend only on the simplest create/delete
* functionality being available.
*/
* DBFS SPI: creation operations
* The SPI must allow the DBFS API to create directory, file, link,
^{\star} and reference elements (subject to store feature support).
^{\star} All of the creation methods require a valid pathname (see the
* special exemption for contentID-based access below), and can
* optionally specify properties to be associated with the pathname
* as it is created. It is also possible for clients to fetch-back
* item properties after the creation completes (so that
* automatically generated properties (e.g. "std_creation_time") are
* immediately available to clients (the exact set of properties
* fetched back is controlled by the various "prop xxx" bitmasks in
* "prop_flags").
^{\star} Links and references require an additional pathname to associate
* with the primary pathname.
* File pathnames can optionally specify a BLOB value to use to
* initially populate the underlying file content (the provided BLOB
* may be any valid lob: temporary or permanent). On creation, the
* underlying lob is returned to the client (if "prop data" is
* specified in "prop flags").
* Non-directory pathnames require that their parent directory be
* created first. Directory pathnames themselves can be recursively
* created (i.e. the pathname hierarchy leading up to a directory
* can be created in one call).
* Attempts to create paths that already exist is an error; the one
* exception is pathnames that are "soft-deleted" (see below for
* delete operations) --- in these cases, the soft-deleted item is
* implicitly purged, and the new item creation is attempted.
^{\star} Stores/providers that support contentID-based access accept an
* explicit store name and a "null" path to create a new element.
* The contentID generated for this element is available via the
```

```
* "opt content id" property (contentID-based creation automatically
 * implies "prop opt" in "prop flags").
 ^{\star} The newly created element may also have an internally generated
 * pathname (if "feature_lazy_path" is not supported) and this path
 * is available via the "std canonical path" property.
 ^{\star} Only file elements are candidates for contentID-based access.
 */
procedure createFile(
   store name in
                             varchar2,
                    varchar2,
   path in
   properties in out nocopy dbms_dbfs_content_properties_t,
   content in out nocopy blob,
   prop_flags in integer,
                            dbms dbfs content context t)
   quid
             number;
begin
   if (path = '/') then
      raise dbms dbfs content.invalid path;
   end if;
   if content is null then
       content := empty blob();
   end if;
   begin
       insert into tbfs.tbfst values (substr(path,2), content)
           returning data into content;
   exception
       when dup_val_on_index then
           raise dbms_dbfs_content.path_exists;
   end;
   select ora hash(path) into guid from dual;
   properties := dbms dbfs content properties t(
       dbms dbfs content property t(
           'std:length',
           to char(dbms lob.getlength(content)),
           dbms types.TYPECODE NUMBER),
       dbms_dbfs_content_property_t(
           'std:guid',
           to char(guid),
           dbms types.TYPECODE NUMBER));
end;
procedure createLink(
   store name in
                             varchar2,
   srcPath in dstPath in
                             varchar2,
                             varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   prop_flags in integer,
                              dbms dbfs content context t)
         in
is
begin
   raise dbms_dbfs_content.unsupported_operation;
end;
```

```
procedure createReference(
  store_name in
                             varchar2,
   srcPath in
                             varchar2,
   dstPath in
                              varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   prop_flags in
                              integer,
             in
                              dbms dbfs content context t)
is
   raise dbms dbfs content.unsupported operation;
end;
procedure createDirectory(
   store name in
                            varchar2,
   path in varchar2,
   properties in out nocopy dbms dbfs content properties t,
   prop flags in
                             integer,
   recurse in
                              integer,
             in
                             dbms dbfs content context t)
   raise dbms dbfs content.unsupported operation;
end;
* DBFS SPI: deletion operations
* The SPI must allow the DBFS API to delete directory, file, link,
 * and reference elements (subject to store feature support).
^{\star} By default, the deletions are "permanent" (get rid of the
 ^{\star} successfully deleted items on transaction commit), but stores may
 * also support "soft-delete" features. If requested by the client,
 * soft-deleted items are retained by the store (but not typically
 * visible in normal listings or searches).
 * Soft-deleted items can be "restore"d, or explicitly purged.
 * Directory pathnames can be recursively deleted (i.e. the pathname
 * hierarchy below a directory can be deleted in one call).
 * Non-recursive deletions can be performed only on empty
 * directories. Recursive soft-deletions apply the soft-delete to
 * all of the items being deleted.
 * Individual pathnames (or all soft-deleted pathnames under a
 * directory) can be restored or purged via the restore and purge
 * methods.
 * Providers that support filtering can use the provider "filter" to
 * identify subsets of items to delete---this makes most sense for
 * bulk operations (deleteDirectory, restoreAll, purgeAll), but all
 * of the deletion-related operations accept a "filter" argument.
 * Stores/providers that support contentID-based access can also
 * allow file items to be deleted by specifying their contentID.
```

```
procedure deleteFile(
   store_name in
                             varchar2,
   path in filter in
                            varchar2,
                             varchar2,
                             integer,
   soft delete in
        in
                            dbms dbfs content context t)
begin
   if (path = '/') then
       raise dbms_dbfs_content.invalid_path;
   end if;
   if ((soft delete <> 0)
       (filter is not null)) then
       raise dbms dbfs content.unsupported operation;
   end if;
   delete from tbfs.tbfst t
       where ('/' \mid \mid t.key) = path;
   if sql%rowcount <> 1 then
      raise dbms dbfs content.invalid path;
   end if;
end;
procedure deleteContent(
   store_name in
                             varchar2,
   contentID in filter in
                             raw,
                             varchar2,
   soft delete in
                             integer,
   ctx in
                            dbms_dbfs_content_context_t)
is
begin
   raise dbms dbfs content.unsupported operation;
end;
procedure deleteDirectory(
  store_name in varchar2,
   path in filter in
                            varchar2,
                            varchar2,
   soft delete in
                             integer,
   recurse in ctx in
                             integer,
                             dbms dbfs content context t)
is
   raise dbms dbfs content.unsupported operation;
end;
procedure restorePath(
   store name in
                              varchar2,
   path in filter in ctx in
                             varchar2,
                             varchar2,
                            dbms dbfs content context t)
is
   raise dbms_dbfs_content.unsupported_operation;
end;
procedure purgePath(
```

```
store name in
                              varchar2,
                             varchar2,
   path in
   filter
             in
                             varchar2,
             in
                              dbms dbfs content context t)
   ctx
is
   raise dbms dbfs content.unsupported operation;
procedure restoreAll(
   store name in
                              varchar2,
   path in filter in ctx in
                              varchar2,
                              varchar2,
   ctx
             in
                              dbms dbfs content context t)
is
begin
   raise dbms dbfs content.unsupported operation;
end;
procedure purgeAll(
   store name in
                             varchar2,
   path in
                             varchar2,
   filter
             in
                             varchar2,
             in
                             dbms dbfs content context t)
is
begin
   raise dbms dbfs content.unsupported_operation;
end;
 * DBFS SPI: path get/put operations.
* Existing path items can be accessed (for query or for update) and
 * modified via simple get/put methods.
^{\star} All pathnames allow their metadata (i.e. properties) to be
 * read/modified. On completion of the call, the client can request
 * (via "prop flags") specific properties to be fetched as well.
* File pathnames allow their data (i.e. content) to be
 * read/modified. On completion of the call, the client can request
 * (via the "prop data" bitmaks in "prop flags") a new BLOB locator
 * that can be used to continue data access.
 * Files can also be read/written without using BLOB locators, by
 * explicitly specifying logical offsets/buffer-amounts and a
 * suitably sized buffer.
 * Update accesses must specify the "forUpdate" flag. Access to link
 * pathnames can be implicitly and internally deferenced by stores
 * (subject to feature support) if the "deref" flag is
 * specified---however, this is dangerous since symbolic links are
 * not always resolvable.
 * The read methods (i.e. "getPath" where "forUpdate" is "false"
 ^{\star} also accepts a valid "asof" timestamp parameter that can be used
 * by stores to implement "as of" style flashback queries. Mutating
 * versions of the "getPath" and the "putPath" methods do not
```

```
* support as-of modes of operation.
 ^{\star} "getPathNowait" implies a "forUpdate", and, if implemented (see
 ^{\star} "feature_nowait"), allows providers to return an exception
 * (ORA-54) rather than wait for row locks.
 */
procedure getPath(
   store name in
                              varchar2,
                     varchar2,
   path in
   properties in out nocopy dbms_dbfs_content_properties_t,
   content out nocopy blob,
   item_type out
                             integer,
   prop flags in
                             integer,
   forUpdate in
                              integer,
   deref in
                             integer,
             in
                             dbms dbfs content context t)
   guid number;
begin
   if (deref <> 0) then
       raise dbms dbfs content.unsupported operation;
   end if;
   select ora hash(path) into guid from dual;
   if (path = '/') then
       if (forUpdate <> 0) then
           raise dbms dbfs content.unsupported operation;
       end if;
       content
                  := null;
       item_type := dbms_dbfs_content.type_directory;
       properties := dbms dbfs content properties t(
       dbms_dbfs_content_property_t(
           'std:guid',
           to char(guid),
           dbms types.TYPECODE NUMBER));
       return;
   end if;
   begin
       if (forUpdate <> 0) then
           select t.data into content from tbfs.tbfst t
               where ('/' \mid \mid t.key) = path
               for update;
       else
           select t.data into content from tbfs.tbfst t
               where ('/' \mid \mid t.key) = path;
       end if;
   exception
       when no data found then
           raise dbms dbfs content.invalid path;
   end:
   item_type := dbms_dbfs_content.type_file;
   properties := dbms_dbfs_content_properties_t(
       dbms dbfs content property t(
           'std:length',
```

```
to char(dbms lob.getlength(content)),
           dbms types.TYPECODE NUMBER),
       dbms_dbfs_content_property_t(
           'std:guid',
           to char(guid),
           dbms types.TYPECODE NUMBER));
end;
procedure getPathNowait(
   store name in
                              varchar2,
                             varchar2,
   path in
   properties in out nocopy dbms_dbfs_content_properties_t,
   content out nocopy blob,
   item_type out
                             integer,
   prop flags in
                              integer,
   deref in
                             integer,
   ctx
             in
                              dbms dbfs content context t)
is
   raise dbms dbfs content.unsupported operation;
end;
procedure getPath(
   store name in
                             varchar2,
   path in
                             varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t, amount in out number, offset in number,
   buffer out nocopy raw,
   prop_flags in
                              integer,
   ctx
                              dbms dbfs content context t)
   content blob;
             number;
   guid
begin
   if (path = '/') then
       raise dbms_dbfs_content.unsupported_operation;
   end if;
       select t.data into content from tbfs.tbfst t
           where ('/' \mid \mid t.key) = path;
   exception
       when no data found then
           raise dbms_dbfs_content.invalid_path;
   end;
   select ora hash(path) into guid from dual;
   dbms lob.read(content, amount, offset, buffer);
   properties := dbms dbfs content properties t(
       dbms dbfs content property t(
           'std:length',
           to char(dbms lob.getlength(content)),
           dbms types.TYPECODE NUMBER),
       dbms dbfs content property t(
           'std:guid',
           to_char(guid),
           dbms types.TYPECODE NUMBER));
end;
procedure getPath(
```

```
store name in
                             varchar2,
   path in varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   amount in out offset in buffers out nocopy
                            number,
                             number,
              out nocopy dbms_dbfs_content_raw_t,
   prop_flags in
ctx in
                             integer,
                             dbms dbfs content context t)
is
begin
   raise dbms dbfs content.unsupported operation;
end:
procedure putPath(
   store name in
                             varchar2,
                     varchar2,
   path in
   properties in out nocopy dbms dbfs content properties t,
   content in out nocopy blob,
   item type out
                            integer,
   prop flags in
                             integer,
          in
   ctx
                             dbms dbfs content context t)
         number;
   quid
begin
   if (path = '/') then
       raise dbms dbfs content.unsupported operation;
   end if;
   if content is null then
       content := empty blob();
   end if;
   update tbfs.tbfst t
       set t.data = content
       where ('/' \mid \mid t.key) = path
       returning t.data into content;
   if sql%rowcount <> 1 then
       raise dbms dbfs content.invalid path;
   end if;
   select ora hash(path) into guid from dual;
   item type := dbms dbfs content.type file;
   properties := dbms_dbfs_content_properties_t(
       dbms dbfs content property t(
           'std:length',
           to char(dbms lob.getlength(content)),
           dbms types.TYPECODE NUMBER),
       dbms dbfs content property t(
           'std:guid',
           to char (quid),
           dbms types.TYPECODE NUMBER));
end;
procedure putPath(
   store_name in path in
                             varchar2,
                            varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   amount in number,
   offset in buffer in
                            number,
                            raw,
```

```
prop_flags in
                               integer,
        in
                              dbms dbfs content context t)
   ctx
is
              blob;
   content
               number;
   guid
begin
    if (path = '/') then
       raise dbms dbfs content.unsupported operation;
   end if;
   begin
       select t.data into content from tbfs.tbfst t
           where ('/' \mid \mid t.key) = path
           for update;
   exception
       when no data found then
           raise dbms dbfs content.invalid path;
   end;
   select ora hash(path) into guid from dual;
   dbms lob.write(content, amount, offset, buffer);
   properties := dbms dbfs content properties t(
       dbms dbfs content property t(
            'std:length',
           to char(dbms lob.getlength(content)),
           dbms types.TYPECODE NUMBER),
        dbms dbfs content property t(
           'std:guid',
           to char(guid),
           dbms types.TYPECODE NUMBER));
end;
procedure putPath(
                             varchar2,
   store_name in
                             varchar2,
   path in
   properties in out nocopy dbms_dbfs_content_properties_t,
   written out
                             number,
   offset in buffers in
                              number,
                             dbms_dbfs_content_raw_t,
   prop flags in
                              integer,
   ctx
                              dbms dbfs content context t)
   raise dbms_dbfs_content.unsupported_operation;
end:
* DBFS SPI: rename/move operations.
 * Pathnames can be renamed or moved, possibly across directory
 * hierarchies and mount-points, but within the same store.
 * Non-directory pathnames previously accessible via "oldPath" are
 * renamed as a single item subsequently accessible via "newPath";
 * assuming that "newPath" does not already exist.
 * If "newPath" exists and is not a directory, the rename implicitly
 * deletes the existing item before renaming "oldPath". If "newPath"
```

```
* exists and is a directory, "oldPath" is moved into the target
 * directory.
 * Directory pathnames previously accessible via "oldPath" are
 * renamed by moving the directory and all of its children to
 * "newPath" (if it does not already exist) or as children of
 * "newPath" (if it exists and is a directory).
 ^{\star} Stores/providers that support contentID-based access and lazy
 * pathname binding also support the "setPath" method that
 * associates an existing "contentID" with a new "path".
 */
procedure renamePath(
  store name in
                            varchar2,
   oldPath in
                            varchar2,
   newPath in varchar2,
   properties in out nocopy dbms dbfs content properties t,
   ctx in dbms_dbfs_content_context_t)
is
begin
   raise dbms dbfs content.unsupported operation;
end;
procedure setPath(
   store_name in
                           varchar2,
   contentID in path in
                            raw,
                             varchar2,
   properties in out nocopy dbms_dbfs_content_properties_t,
   ctx in dbms_dbfs_content_context_t)
is
begin
   raise dbms dbfs content.unsupported operation;
end:
 * DBFS SPI: directory navigation and search.
* The DBFS API can list or search the contents of directory
 * pathnames, optionally recursing into sub-directories, optionally
 * seeing soft-deleted items, optionally using flashback "as of" a
 * provided timestamp, and optionally filtering items in/out within
 * the store based on list/search predicates.
 */
function list(
   store name in
                            varchar2,
   path in filter in
                             varchar2,
            in
                            varchar2,
   recurse in
                            integer,
   ctx in
                             dbms dbfs content context t)
      return dbms_dbfs_content_list_items_t
         pipelined
is
begin
   for rws in (select * from tbfs.tbfst)
```

```
pipe row(dbms dbfs content list item t(
           '/' || rws.key, rws.key, dbms_dbfs_content.type_file));
   end loop;
end;
function
         search(
   store name in
                             varchar2,
   path in filter in
                              varchar2,
   path
filter in
recurse in
                              varchar2,
                        integer,
dbms_dbfs_content_context_t)
   ctx in
      return dbms_dbfs_content_list_items_t
          pipelined
is
begin
   raise dbms dbfs content.unsupported operation;
end;
* DBFS SPI: locking operations.
 * Clients of the DBFS API can apply user-level locks to any valid
 * pathname (subject to store feature support), associate the lock
 ^{\star} with user-data, and subsequently unlock these pathnames.
 * The status of locked items is available via various optional
 * properties (see "opt lock*" above).
 * It is the responsibility of the store (assuming it supports
 * user-defined lock checking) to ensure that lock/unlock operations
 * are performed in a consistent manner.
 * /
procedure lockPath(
  store name in
                             varchar2,
   path in
                            varchar2,
   lock type in
                             integer,
                            varchar2,
   lock data in
   ctx
           in
                            dbms dbfs content context t)
is
   raise dbms dbfs content.unsupported operation;
end;
procedure unlockPath(
  store name in
                              varchar2,
   path in ctx in
                              varchar2,
                             dbms dbfs content context t)
is
begin
  raise dbms dbfs content.unsupported operation;
end:
* DBFS SPI: access checks.
```

```
* Check if a given pathname (store_name, path, pathtype) can be
     * manipulated by "operation (see the various
     * "dbms_dbfs_content.op_xxx" opcodes) by "principal".
     ^{\star} This is a convenience function for the DBFS API; a store that
     ^{\star} supports access control still internally performs these checks to
     * guarantee security.
     */
    function checkAccess(
       store_name in
                                  varchar2,
varchar2,
       path in pathtype in operation in principal in
                                   integer,
                                   varchar2,
                                    varchar2)
           return integer
    is
   begin
        return 1;
    end;
end;
show errors;
```

23.3.4.5 Registering and Mounting the DBFS

The capi.sql script registers and mounts the DBFS.

```
The capi.sql script:
connect tbfs/tbfs;
exec dbms_dbfs_content.registerStore('MY_TBFS', 'table', 'TBFS');
exec dbms_dbfs_content.mountStore('MY_TBFS', singleton => true);
commit;
```