F

Extended Examples

The body of the *SQL Language Reference* contains examples for almost every reference topic. This appendix contains lengthy examples that are not appropriate in the context of a single SQL statement. These examples are intended to provide uninterrupted the series of steps that you would use to take advantage of particular Oracle functionality. They do not replace the syntax diagrams and semantics found for each individual SQL statement in the body of the reference. Use the cross-references provided to access additional information, such as privileges required and restrictions, as well as syntax.

This appendix contains the following sections:

- Using Extensible Indexing
- Using XML in SQL Statements

Using Extensible Indexing

This section provides examples of the steps entailed in a simple but realistic extensible indexing scenario.

Suppose you want to rank the salaries in the HR.employees table and then find those that rank between 10 and 20. You could use the DENSE RANK function, as follows:

```
SELECT last_name, salary FROM
  (SELECT last_name, DENSE_RANK() OVER
        (ORDER BY salary DESC) rank_val, salary FROM employees)
WHERE rank val BETWEEN 10 AND 20;
```



This nested query is somewhat complex, and it requires a full scan of the <code>employees</code> table as well as a sort. An alternative would be to use extensible indexing to achieve the same goal. The resulting query will be simpler. The query will require only an index scan and a table access by rowid, and will therefore perform much more efficiently.

The first step is to create the implementation type <code>position_im</code>, including method headers for index definition, maintenance, and creation. Most of the type body uses PL/SQL, which is shown in italics.

The type must created with the AUTHID CURRENT_USER clause because of the EXECUTE IMMEDIATE statement inside the function ODCIINDEXCREATE(). By default that function runs with the definer rights. When the function is called in the subsequent creation of the domain index, the invoker does not have the same rights.

See Also:

- CREATE TYPE and CREATE TYPE BODY
- Oracle Database Data Cartridge Developer's Guide for complete information on the ODCI routines in this statement

```
CREATE OR REPLACE TYPE position im AUTHID CURRENT USER AS OBJECT
  curnum NUMBER,
 howmany NUMBER,
 lower bound NUMBER,
 upper bound NUMBER,
/* lower bound and upper bound are used for the
index-based functional implementation */
  STATIC FUNCTION ODCIGETINTERFACES (ifclist OUT SYS.ODCIOBJECTLIST) RETURN NUMBER,
  STATIC FUNCTION ODCIINDEXCREATE
    (ia SYS.ODCIINDEXINFO, parms VARCHAR2, env SYS.ODCIEnv) RETURN NUMBER,
  STATIC FUNCTION ODCIINDEXTRUNCATE (ia SYS.ODCIINDEXINFO,
                                     env SYS.ODCIEnv) RETURN NUMBER,
  STATIC FUNCTION ODCIINDEXDROP(ia SYS.ODCIINDEXINFO,
                                env SYS.ODCIEnv) RETURN NUMBER,
  STATIC FUNCTION ODCIINDEXINSERT (ia SYS.ODCIINDEXINFO, rid ROWID,
                                  newval NUMBER, env SYS.ODCIEnv) RETURN NUMBER,
  STATIC FUNCTION ODCIINDEXDELETE (ia SYS.ODCIINDEXINFO, rid ROWID, oldval NUMBER,
                                  env SYS.ODCIEnv) RETURN NUMBER,
  STATIC FUNCTION ODCIINDEXUPDATE(ia SYS.ODCIINDEXINFO, rid ROWID, oldval NUMBER,
                                  newval NUMBER, env SYS.ODCIEnv) RETURN NUMBER,
  STATIC FUNCTION ODCIINDEXSTART (SCTX IN OUT position im, ia SYS.ODCIINDEXINFO,
                                 op SYS.ODCIPREDINFO, qi SYS.ODCIQUERYINFO,
                                 strt NUMBER, stop NUMBER, lower_pos NUMBER,
                                 upper pos NUMBER, env SYS.ODCIEnv) RETURN NUMBER,
  MEMBER FUNCTION ODCIINDEXFETCH (SELF IN OUT position im, nrows NUMBER,
                                 rids OUT SYS.ODCIRIDLIST, env SYS.ODCIEnv)
                                 RETURN NUMBER,
 MEMBER FUNCTION ODCIINDEXCLOSE (env SYS.ODCIEnv) RETURN NUMBER
);
CREATE OR REPLACE TYPE BODY position im
   STATIC FUNCTION ODCIGETINTERFACES (ifclist OUT SYS.ODCIOBJECTLIST)
      RETURN NUMBER IS
       ifclist := SYS.ODCIOBJECTLIST(SYS.ODCIOBJECT('SYS','ODCIINDEX2'));
      RETURN ODCICONST.SUCCESS;
   END ODCIGETINTERFACES;
 STATIC FUNCTION ODCIINDEXCREATE (ia SYS.ODCIINDEXINFO, parms VARCHAR2, env SYS.ODCIEnv) RETURN
 NUMBER
 IS
         VARCHAR2 (2000);
  stmt
  BEGIN
/* Construct the SQL statement */
   stmt := 'Create Table ' || ia.INDEXSCHEMA || '.' || ia.INDEXNAME ||
           '_STORAGE_TAB' || '(col_val, base_rowid, constraint pk PRIMARY KEY ' ||
           '(col val, base_rowid)) ORGANIZATION INDEX AS SELECT ' ||
           ia.INDEXCOLS(1).COLNAME || ', ROWID FROM ' ||
           ia.INDEXCOLS(1).TABLESCHEMA || '.' || ia.INDEXCOLS(1).TABLENAME;
   EXECUTE IMMEDIATE stmt;
```



```
RETURN ODCICONST.SUCCESS:
 END:
 STATIC FUNCTION ODCIINDEXDROP(ia SYS.ODCIINDEXINFO, env SYS.ODCIEnv) RETURN NUMBER IS
  stmt VARCHAR2(2000);
 BEGIN
/* Construct the SQL statement */
  stmt := 'DROP TABLE ' || ia.INDEXSCHEMA || '.' || ia.INDEXNAME ||
   ' STORAGE TAB';
/* Execute the statement */
  EXECUTE IMMEDIATE stmt;
  RETURN ODCICONST.SUCCESS;
 END;
 STATIC FUNCTION ODCIINDEXTRUNCATE(ia SYS.ODCIINDEXINFO, env SYS.ODCIEnv) RETURN NUMBER IS
  stmt VARCHAR2(2000);
 BEGIN
/* Construct the SQL statement */
  stmt := 'TRUNCATE TABLE ' || ia.INDEXSCHEMA || '.' || ia.INDEXNAME || ' STORAGE TAB';
  EXECUTE IMMEDIATE stmt;
  RETURN ODCICONST.SUCCESS;
 STATIC FUNCTION ODCIINDEXINSERT (ia SYS. ODCIINDEXINFO, rid ROWID,
                          newval NUMBER, env SYS.ODCIEnv) RETURN NUMBER IS
  stmt VARCHAR2 (2000);
 BEGIN
/* Construct the SQL statement */
   stmt := 'INSERT INTO ' || ia.INDEXSCHEMA || '.' || ia.INDEXNAME ||
          ' STORAGE TAB VALUES (''' || newval || ''' , ''' || rid || ''' )';
/* Execute the SQL statement */
  EXECUTE IMMEDIATE stmt;
  RETURN ODCICONST.SUCCESS;
 END:
 STATIC FUNCTION ODCIINDEXDELETE(ia SYS.ODCIINDEXINFO, rid ROWID, oldval NUMBER,
                                  env SYS.ODCIEnv)
    RETURN NUMBER IS
  stmt VARCHAR2(2000);
 BEGIN
/* Construct the SQL statement */
   stmt := 'DELETE FROM ' || ia.INDEXSCHEMA || '.' || ia.INDEXNAME ||
         'STORAGE TAB WHERE col val = ''' || oldval || ''' AND base rowid = ''' || rid || '''';
/* Execute the statement */
  EXECUTE IMMEDIATE stmt;
  RETURN ODCICONST.SUCCESS;
 END;
 STATIC FUNCTION ODCIINDEXUPDATE (ia SYS.ODCIINDEXINFO, rid ROWID, oldval NUMBER,
                          newval NUMBER, env SYS.ODCIEnv) RETURN NUMBER IS
  stmt VARCHAR2 (2000);
 BEGIN
/* Construct the SQL statement */
  stmt := 'UPDATE ' || ia.INDEXSCHEMA || '.' || ia.INDEXNAME ||
         'STORAGE TAB SET col val = ''' || newval || ''' WHERE f2 = '''|| rid ||'''';
/* Execute the statement */
  EXECUTE IMMEDIATE stmt;
  RETURN ODCICONST.SUCCESS;
 END:
 STATIC FUNCTION ODCIINDEXSTART (SCTX IN OUT position im, ia SYS.ODCIINDEXINFO,
                         op SYS.ODCIPREDINFO, qi SYS.ODCIQUERYINFO,
                         strt NUMBER, stop NUMBER, lower pos NUMBER,
                         upper pos NUMBER, env SYS.ODCIEnv) RETURN NUMBER IS
    rid
                     VARCHAR2 (5072);
    storage tab name VARCHAR2(65);
```

```
lower bound stmt VARCHAR2(2000);
   upper bound stmt VARCHAR2(2000);
   range_query_stmt VARCHAR2(2000);
   lower bound
                   NUMBER;
   upper_bound
                    NUMBER;
   cnum
                    INTEGER;
   nrows
                    INTEGER;
/* Take care of some error cases.
   The only predicates in which position operator can appear are
      op() = 1
                 OR
      op() = 0
                   OR
      op() between 0 and 1
   IF (((strt != 1) AND (strt != 0)) OR
        ((stop != 1) AND (stop != 0)) OR
        ((strt = 1) AND (stop = 0))) THEN
     RAISE APPLICATION ERROR (-20101,
                              'incorrect predicate for position between operator');
   END IF;
   IF (lower pos > upper pos) THEN
     RAISE APPLICATION ERROR (-20101, 'Upper Position must be greater than or
     equal to Lower Position');
   END IF;
   IF (lower pos <= 0) THEN
     RAISE APPLICATION ERROR(-20101, 'Both Positions must be greater than zero');
   storage tab name := ia.INDEXSCHEMA || '.' || ia.INDEXNAME ||
                        ' STORAGE TAB';
   upper bound stmt := 'Select MIN(col val) FROM (Select /*+ INDEX DESC(' ||
                        storage_tab_name || ') */ DISTINCT ' ||
                        'col_val FROM ' || storage_tab_name || ' ORDER BY ' ||
                        'col_val DESC) WHERE rownum <= ' || lower_pos;</pre>
   EXECUTE IMMEDIATE upper_bound_stmt INTO upper_bound;
   IF (lower_pos != upper_pos) THEN
     lower bound stmt := 'Select MIN(col val) FROM (Select /*+ INDEX DESC(' ||
                          storage_tab_name || ') */ DISTINCT ' ||
                          'col val FROM ' || storage tab name ||
                          'WHERE col val < ' || upper bound || 'ORDER BY ' ||
                          'col val DESC) WHERE rownum <= ' ||</pre>
                          (upper pos - lower pos);
     EXECUTE IMMEDIATE lower bound stmt INTO lower bound;
   ELSE
     lower_bound := upper_bound;
   END IF;
   IF (lower bound IS NULL) THEN
     lower bound := upper bound;
   END IF;
   range query stmt := 'Select base rowid FROM ' || storage tab name ||
                        'WHERE col val BETWEEN ' || lower bound || 'AND ' ||
                       upper bound;
   cnum := DBMS SQL.OPEN CURSOR;
   DBMS SQL.PARSE(cnum, range query stmt, DBMS SQL.NATIVE);
/* set context as the cursor number */
   SCTX := position_im(cnum, 0, 0, 0);
/* return success */
   RETURN ODCICONST.SUCCESS;
 MEMBER FUNCTION ODCIINDEXFETCH (SELF IN OUT position_im, nrows NUMBER,
                                 rids OUT SYS.ODCIRIDLIST, env SYS.ODCIEnv)
  RETURN NUMBER IS
```

```
INTEGER;
    cnum
   rid tab DBMS SQL.Varchar2 table;
   rlist SYS.ODCIRIDLIST := SYS.ODCIRIDLIST();
            INTEGER;
   d
            INTEGER;
 BEGIN
    cnum := SELF.curnum;
    IF self.howmany = 0 THEN
     dbms sql.define array(cnum, 1, rid_tab, nrows, 1);
     d := DBMS SQL.EXECUTE(cnum);
   END IF;
   d := DBMS SQL.FETCH ROWS(cnum);
   IF d = nrows THEN
     rlist.extend(d);
     rlist.extend(d+1);
   END IF:
   DBMS SQL.COLUMN VALUE (cnum, 1, rid tab);
   for i in 1..d loop
     rlist(i) := rid tab(i+SELF.howmany);
   end loop;
   SELF.howmany := SELF.howmany + d;
   rids := rlist;
   RETURN ODCICONST.SUCCESS;
 MEMBER FUNCTION ODCIINDEXCLOSE(env SYS.ODCIEnv) RETURN NUMBER IS
   cnum INTEGER;
 BEGIN
   cnum := SELF.curnum;
   DBMS SQL.CLOSE CURSOR (cnum);
   RETURN ODCICONST.SUCCESS;
 END;
END;
```

The next step is to create the functional implementation <code>function_for_position_between</code> for the operator that will be associated with the indextype. (The PL/SQL blocks are shown in parentheses.)

This function is for use with an index-based function evaluation. Therefore, it takes an index context and scan context as parameters.

See Also:

- Oracle Database Data Cartridge Developer's Guide for information on creating index-based functional implementation
- CREATE FUNCTION and Oracle Database PL/SQL Language Reference



```
upper_bound_stmt VARCHAR2(2000);
 col val stmt
                  VARCHAR2 (2000);
                  NUMBER;
 lower_bound
 upper bound
                   NUMBER;
                   NUMBER;
 column_value
BEGIN
 IF (indexctx.IndexInfo IS NOT NULL) THEN
    storage tab name := indexctx.IndexInfo.INDEXSCHEMA || '.' ||
                        indexctx.IndexInfo.INDEXNAME || ' STORAGE TAB';
   IF (scanctx IS NULL) THEN
/\star This is the first call. Open a cursor for future calls.
  First, do some error checking
     IF (lower pos > upper pos) THEN
       RAISE APPLICATION ERROR (-20101,
          'Upper Position must be greater than or equal to Lower Position');
     END IF;
     IF (lower pos <= 0) THEN
       RAISE APPLICATION ERROR (-20101,
          'Both Positions must be greater than zero');
/* Obtain the upper and lower value bounds for the range we're interested in.
      upper_bound_stmt := 'Select MIN(col_val) FROM (Select /*+ INDEX_DESC(' ||
                        storage tab name || ') */ DISTINCT ' ||
                        'col val FROM ' || storage tab name || ' ORDER BY ' ||
                        'col val DESC) WHERE rownum <= ' || lower pos;</pre>
      EXECUTE IMMEDIATE upper bound stmt INTO upper bound;
      IF (lower pos != upper pos) THEN
        lower bound stmt := 'Select MIN(col val) FROM (Select /*+ INDEX DESC(' ||
                            storage_tab_name || ') */ DISTINCT ' ||
                             'col_val FROM ' || storage_tab_name ||
                             'WHERE col_val < ' || upper_bound || 'ORDER BY ' ||
                            'col val DESC) WHERE rownum <= ' ||</pre>
                             (upper_pos - lower_pos);
        EXECUTE IMMEDIATE lower bound stmt INTO lower bound;
     ELSE
       lower bound := upper_bound;
      END IF;
      IF (lower bound IS NULL) THEN
        lower bound := upper bound;
/* Store the lower and upper bounds for future function invocations for
   the positions.
     scanctx := position im(0, 0, lower bound, upper bound);
/* Fetch the column value corresponding to the rowid, and see if it falls
  within the determined range.
   col_val_stmt := 'Select col_val FROM ' || storage tab name ||
                    'WHERE base rowid = ''' || indexctx.Rid || '''';
    EXECUTE IMMEDIATE col val stmt INTO column value;
    IF (column value <= scanctx.upper bound AND</pre>
       column value >= scanctx.lower bound AND
        scanflg = ODCICONST.RegularCall) THEN
     RETURN 1;
   ELSE
     RETURN 0;
   END IF;
    RAISE APPLICATION ERROR(-20101, 'A column that has a domain index of' ||
```

```
END IF;
END;
```

'Position indextype must be the first argument');

Next, create the <code>position_between</code> operator, which uses the <code>function_for_position_between</code> function. The operator takes an indexed <code>NUMBER</code> column as the first argument, followed by a <code>NUMBER</code> lower and upper bound as the second and third arguments.



CREATE OPERATOR

CREATE OR REPLACE OPERATOR position_between
BINDING (NUMBER, NUMBER, NUMBER) RETURN NUMBER
WITH INDEX CONTEXT, SCAN CONTEXT position_im
USING function_for_position_between;

In this CREATE OPERATOR statement, the WITH INDEX CONTEXT, SCAN CONTEXT position_im clause is included so that the index context and scan context are passed in to the functional evaluation, which is index based.

Now create the position indextype indextype for the position operator:



CREATE INDEXTYPE

CREATE INDEXTYPE position_indextype
FOR position_between(NUMBER, NUMBER, NUMBER)
USING position im;

The operator <code>position_between</code> uses an index-based functional implementation. Therefore, a domain index must be defined on the referenced column so that the index information can be passed into the functional evaluation. So the final step is to create the domain index <code>salary_index</code> using the <code>position_indextype</code> indextype:



CREATE INDEX

CREATE INDEX salary_index ON employees(salary)
INDEXTYPE IS position indextype;

Now you can use the position_between operator function to rewrite the original query as follows:



SELECT	last	name,	salary	FROM	empl	oye	es
WHER	RE pos	sition_	betweer	n(sala	ary,	10,	20)=1
ORDE	ER BY	salary	DESC,	last	name	e;	

LAST_NAME	SALARY
Tucker	10000
King	10000
Baer	10000
Bloom	10000
Fox	9600
Bernstein	9500
Sully	9500
Greene	9500
Hunold	9000
Faviet	9000
McEwen	9000
Hall	9000
Hutton	8800
Taylor	8600
Livingston	8400
Gietz	8300
Chen	8200
Fripp	8200
Weiss	8000
Olsen	8000
Smith	8000
Kaufling	7900

Using XML in SQL Statements

This section describes some of the ways you can use XMLType data in the database.

XMLType Tables

The sample schema oe contains a table warehouses, which contains an XMLType column warehouse_spec. Suppose you want to create a separate table with the warehouse_spec information. The following example creates a very simple XMLType table with one CLOB column:

```
CREATE TABLE xwarehouses OF XMLTYPE XMLTYPE STORE AS CLOB;
```

You can insert into such a table using XMLType syntax, as shown in the next statement. (The data inserted in this example corresponds to the data in the warehouse_spec column of the sample table oe.warehouses where warehouse_id = 1.)



See Also:

Oracle XML DB Developer's Guide for information on $\mathtt{XMLType}$ and its member methods

You can query this table with the following statement:

```
SELECT e.getClobVal() FROM xwarehouses e;
```

CLOB columns are subject to all of the restrictions on LOB columns. To avoid these restrictions, create an XMLSchema-based table. The XMLSchema maps the XML elements to their object-relational equivalents. The following example registers an XMLSchema locally. The XMLSchema (xwarhouses.xsd) reflects the same structure as the xwarehouses table. (XMLSchema declarations use PL/SQL and the DBMS_XMLSCHEMA package, so the example is shown in italics.)

See Also:

Oracle XML DB Developer's Guide for information on creating XMLSchemas

```
begin
dbms xmlschema.registerSchema(
 'http://www.example.com/xwarehouses.xsd',
 '<schema xmlns="http://www.w3.org/2001/XMLSchema"
     targetNamespace="http://www.example.com/xwarehouses.xsd"
     xmlns:who="http://www.example.com/xwarehouses.xsd"
     version="1.0">
 <simpleType name="RentalType">
  <restriction base="string">
   <enumeration value="Rented"/>
   <enumeration value="Owned"/>
  </restriction>
 </simpleType>
 <simpleType name="ParkingType">
  <restriction base="string">
   <enumeration value="Street"/>
   <enumeration value="Lot"/>
  </restriction>
 </simpleType>
 <element name = "Warehouse">
   <complexType>
    <sequence>
     <element name = "WarehouseId" type = "positiveInteger"/>
     <element name = "WarehouseName" type = "string"/>
     <element name = "WaterAccess" type = "boolean"/>
     <element name = "VClearance" type = "positiveInteger"/>
```



Now you can create an XMLSchema-based table, as shown in the following example:

```
CREATE TABLE xwarehouses OF XMLTYPE

XMLSCHEMA "http://www.example.com/xwarehouses.xsd"

ELEMENT "Warehouse";
```

By default, Oracle stores this as an object-relational table. Therefore, you can insert into it as shown in the example that follows. (The data inserted in this example corresponds to the data in the warehouse spec column of the sample table oe.warehouses where warehouse id = 1.)

```
INSERT INTO xwarehouses VALUES( xmltype.createxml('<?xml version="1.0"?>
   <who:Warehouse xmlns:who="http://www.example.com/xwarehouses.xsd"</pre>
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.example.com/xwarehouses.xsd
  http://www.example.com/xwarehouses.xsd">
      <WarehouseId>1</WarehouseId>
     <WarehouseName>Southlake, Texas</warehouseName>
     <Building>Owned</Building>
     <Area>25000</Area>
      <Docks>2</Docks>
      <DockType>Rear load
      <WaterAccess>true</WaterAccess>
      <RailAccess>false</RailAccess>
      <Parking>Street</Parking>
      <VClearance>10</VClearance>
      </who:Warehouse>'));
```

You can define constraints on an XMLSchema-based table. To do so, you use the XMLDATA pseudocolumn to refer to the appropriate attribute within the Warehouse XML element:

```
ALTER TABLE xwarehouses ADD (PRIMARY KEY(XMLDATA. "WarehouseId"));
```

Because the data in xwarehouses is stored object relationally, Oracle rewrites queries to this XMLType table to go to the underlying storage when possible. Therefore the following queries would use the index created by the primary key constraint in the preceding example:

```
SELECT * FROM xwarehouses x
   WHERE EXISTSNODE(VALUE(x), '/Warehouse[WarehouseId="1"]',
   'xmlns:who="http://www.example.com/xwarehouses.xsd"') = 1;

SELECT * FROM xwarehouses x
   WHERE EXTRACTVALUE(VALUE(x), '/Warehouse/WarehouseId',
   'xmlns:who="http://www.example.com/xwarehouses.xsd"') = 1;
```

You can also explicitly create indexes on XMLSchema-based tables, which greatly enhance the performance of subsequent queries. You can create object-relational views on XMLType tables, and you can create XMLType views on object-relational tables.

See Also:

- XMLDATA Pseudocolumn for information on the XMLDATA pseudocolumn
- "Creating an XMLType View: Example"
- · Creating an Index on an XMLType Table: Example

XMLType Columns

The sample table <code>oe.warehouses</code> was created with a <code>warehouse_spec</code> column of type <code>XMLType</code>. The examples in this section create a shortened form of the <code>oe.warehouses</code> table, using two different types of storage.

The first example creates a table with an XMLType table stored as a CLOB. This table does not require an XMLSchema, so the content structure is not predetermined:

```
CREATE TABLE xwarehouses (
warehouse_id NUMBER,
warehouse_spec XMLTYPE)

XMLTYPE warehouse_spec STORE AS CLOB
(TABLESPACE example
STORAGE (INITIAL 6144)
CHUNK 4000
NOCACHE LOGGING);
```

The following example creates a similar table, but stores the $\mathtt{XMLType}$ data in an object-relational $\mathtt{XMLType}$ column whose structure is determined by the specified XMLSchema:

```
CREATE TABLE xwarehouses (

warehouse_id NUMBER,

warehouse_spec XMLTYPE)

XMLTYPE warehouse_spec STORE AS OBJECT RELATIONAL

XMLSCHEMA "http://www.example.com/xwarehouses.xsd"

ELEMENT "Warehouse";
```

