7

Functions

Functions are similar to operators in that they manipulate data items and return a result. Functions differ from operators in the format of their arguments. This format enables them to operate on zero, one, two, or more arguments:

```
function(argument, argument, ...)
```

A function without any arguments is similar to a pseudocolumn (refer to Pseudocolumns). However, a pseudocolumn typically returns a different value for each row in the result set, whereas a function without any arguments typically returns the same value for each row.

This chapter contains these sections:

- About SQL Functions
- Single-Row Functions
 - Numeric Functions
 - Character Functions Returning Character Values
 - Character Functions Returning Number Values
 - Character Set Functions
 - Collation Functions
 - Datetime Functions
 - General Comparison Functions
 - Conversion Functions
 - Large Object Functions
 - Collection Functions
 - Hierarchical Functions
 - Oracle Machine Learning for SQL Functions
 - XML Functions
 - JSON Functions
 - Encoding and Decoding Functions
 - NULL-Related Functions
 - Environment and Identifier Functions
- Aggregate Functions
- Analytic Functions
- Object Reference Functions
- Model Functions
- OLAP Functions
- Data Cartridge Functions



About User-Defined Functions

About SQL Functions

SQL functions are built into Oracle Database and are available for use in various appropriate SQL statements. Do not confuse SQL functions with user-defined functions written in PL/SQL.

If you call a SQL function with an argument of a data type other than the data type expected by the SQL function, then Oracle attempts to convert the argument to the expected data type before performing the SQL function.



About User-Defined Functions for information on user functions and Data Conversion for implicit conversion of data types

Nulls in SQL Functions

Most scalar functions return null when given a null argument. You can use the NVL function to return a value when a null occurs. For example, the expression NVL (commission_pct, 0) returns 0 if commission pct is null or the value of commission pct if it is not null.

For information on how aggregate functions handle nulls, see Aggregate Functions .

Syntax for SQL Functions

In the syntax diagrams for SQL functions, arguments are indicated by their data types. When the parameter <code>function</code> appears in SQL syntax, replace it with one of the functions described in this section. Functions are grouped by the data types of their arguments and their return values.

Note:

When you apply SQL functions to LOB columns, Oracle Database creates temporary LOBs during SQL and PL/SQL processing. You should ensure that temporary tablespace quota is sufficient for storing these temporary LOBs for your application.

A SQL function may be collation-sensitive, which means that character value comparison or matching that it performs is controlled by a collation. The particular collation to use by the function is determined from the collations of the function's arguments.

If the result of a SQL function has a character data type, the collation derivation rules define the collation to associate with the result.

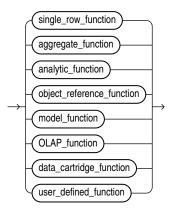
See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation and determination rules for SQL functions

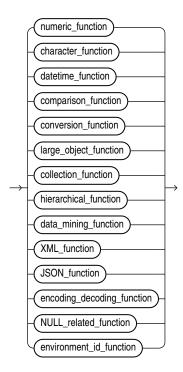


The syntax showing the categories of functions follows:

function::=



single_row_function::=



The sections that follow list the built-in SQL functions in each of the groups illustrated in the preceding diagrams except user-defined functions. All of the built-in SQL functions are then described in alphabetical order.

See Also:

About User-Defined Functions and CREATE FUNCTION



Aggregate Functions

Aggregate functions return a single result row based on groups of rows, rather than on single rows. Aggregate functions can appear in select lists and in ORDER BY and HAVING clauses. They are commonly used with the GROUP BY clause in a SELECT statement, where Oracle Database divides the rows of a queried table or view into groups. In a query containing a GROUP BY clause, the elements of the select list can be aggregate functions, GROUP BY expressions, constants, or expressions involving one of these. Oracle applies the aggregate functions to each group of rows and returns a single result row for each group.

If you omit the GROUP BY clause, then Oracle applies aggregate functions in the select list to all the rows in the queried table or view. You use aggregate functions in the HAVING clause to eliminate groups from the output based on the results of the aggregate functions, rather than on the values of the individual rows of the queried table or view.

See Also:

- Using the GROUP BY Clause: Examples and the HAVING Clause for more information on the GROUP BY clause and HAVING clauses in queries and subqueries
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules for expressions in the ORDER BY clause of an aggregate function

Many (but not all) aggregate functions that take a single argument accept these clauses:

- DISTINCT and UNIQUE, which are synonymous, cause an aggregate function to consider
 only distinct values of the argument expression. The syntax diagrams for aggregate
 functions in this chapter use the keyword DISTINCT for simplicity.
- ALL causes an aggregate function to consider all values, including all duplicates.

For example, the DISTINCT average of 1, 1, 1, and 3 is 2. The ALL average is 1.5. If you specify neither, then the default is ALL.

Some aggregate functions allow the windowing_clause, which is part of the syntax of analytic functions. Refer to windowing_clause for information about this clause.

All aggregate functions except COUNT(*), GROUPING, and GROUPING_ID ignore nulls. You can use the NVL function in the argument to an aggregate function to substitute a value for a null. COUNT and REGR_COUNT never return null, but return either a number or zero. For all the remaining aggregate functions, if the data set contains no rows, or contains only rows with nulls as arguments to the aggregate function, then the function returns null.

The aggregate functions MIN, MAX, SUM, AVG, COUNT, VARIANCE, and STDDEV, when followed by the KEEP keyword, can be used in conjunction with the FIRST or LAST function to operate on a set of values from a set of rows that rank as the FIRST or LAST with respect to a given sorting specification. Refer to FIRST for more information.

You can nest aggregate functions. For example, the following example calculates the average of the maximum salaries of all the departments in the sample schema hr:



This calculation evaluates the inner aggregate (MAX(salary)) for each group defined by the GROUP BY clause (department id), and aggregates the results again.

```
ANY_VALUE
APPROX_COUNT
APPROX_COUNT_DISTINCT
APPROX_COUNT_DISTINCT_AGG
APPROX_COUNT_DISTINCT_DETAIL
APPROX MEDIAN
APPROX PERCENTILE
APPROX PERCENTILE AGG
APPROX_PERCENTILE_DETAIL
APPROX RANK
APPROX_SUM
AVG
BIT_AND_AGG
BIT OR AGG
BIT XOR AGG
BOOLEAN_AND_AGG
BOOLEAN OR AGG
CHECKSUM
COLLECT
CORR
CORR *
COUNT
COVAR POP
COVAR SAMP
CUME DIST
DENSE RANK
EVERY
FIRST
GROUP_ID
GROUPING
GROUPING_ID
JSON ARRAYAGG
JSON OBJECTAGG
KURTOSIS POP
KURTOSIS SAMP
LAST
LISTAGG
MAX
MEDIAN
MIN
PERCENT_RANK
```



PERCENTILE CONT PERCENTILE DISC **RANK REGR** (Linear Regression) Functions SKEWNESS POP SKEWNESS SAMP STATS_BINOMIAL_TEST STATS CROSSTAB STATS_F_TEST STATS KS TEST STATS MODE STATS MW TEST STATS ONE WAY ANOVA STATS_T_TEST_* STATS_WSR_TEST **STDDEV** STDDEV POP STDDEV_SAMP **SUM** SYS OP ZONE ID SYS XMLAGG TO APPROX COUNT DISTINCT TO APPROX PERCENTILE VAR POP VAR SAMP **VARIANCE XMLAGG**

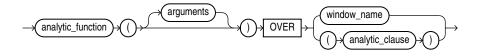
Analytic Functions

Analytic functions compute an aggregate value based on a group of rows. They differ from aggregate functions in that they return multiple rows for each group. The group of rows is called a **window** and is defined by the <code>analytic_clause</code>. For each row, a sliding window of rows is defined. The window determines the range of rows used to perform the calculations for the current row. Window sizes can be based on either a physical number of rows or a logical interval such as time.

Analytic functions are the last set of operations performed in a query except for the final ORDER BY clause. All joins and all WHERE, GROUP BY, and HAVING clauses are completed before the analytic functions are processed. Therefore, analytic functions can appear only in the select list or ORDER BY clause.

Analytic functions are commonly used to compute cumulative, moving, centered, and reporting aggregates.

analytic_function::=

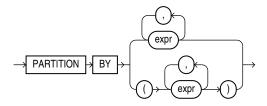




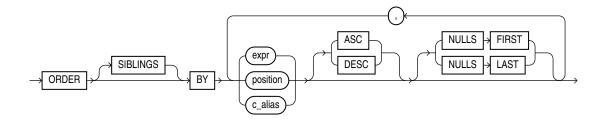
analytic_clause::=



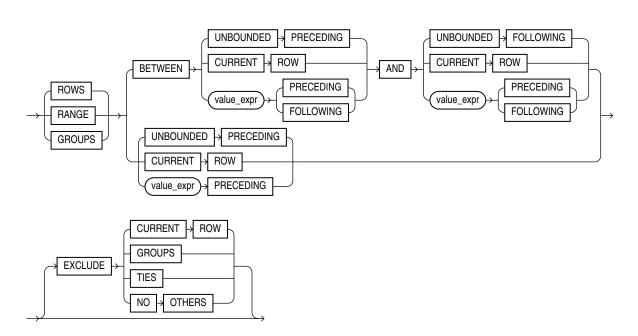
query_partition_clause::=



order_by_clause::=



windowing_clause::=



The semantics of this syntax are discussed in the sections that follow.

analytic_function

Specify the name of an analytic function (see the listing of analytic functions following this discussion of semantics).

arguments

Analytic functions take 0 to 3 arguments. The arguments can be any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the highest numeric precedence and implicitly converts the remaining arguments to that data type. The return type is also that data type, unless otherwise noted for an individual function.



Numeric Precedence for information on numeric precedence and Table 2-9 for more information on implicit conversion

analytic_clause

Use OVER <code>analytic_clause</code> to indicate that the function operates on a query result set. This clause is computed after the <code>FROM</code>, <code>WHERE</code>, <code>GROUP BY</code>, and <code>HAVING</code> clauses. You can specify analytic functions with this clause in the select list or <code>ORDER BY</code> clause. To filter the results of a query based on an analytic function, nest these functions within the parent query, and then filter the results of the nested subquery.

Notes on the analytic clause:

The following notes apply to the analytic clause:

- You cannot nest analytic functions by specifying any analytic function in any part of the <code>analytic_clause</code>. However, you can specify an analytic function in a subquery and compute another analytic function over it.
- You can specify OVER analytic_clause with user-defined analytic functions as well as built-in analytic functions. See CREATE FUNCTION.
- The PARTITION BY and ORDER BY clauses in the analytic clause are collation-sensitive.

✓ See Also:

- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules for the OVER (PARTITION BY ... ORDER BY ...) clause of an analytic function
- window clause in the SELECT statement



query_partition_clause

Use the PARTITION BY clause to partition the query result set into groups based on one or more $value_expr$. If you omit this clause, then the function treats all rows of the query result set as a single group.

To use the <code>query_partition_clause</code> in an analytic function, use the upper branch of the syntax (without parentheses). To use this clause in a model query (in the <code>model_column_clauses</code>) or a partitioned outer join (in the <code>outer_join_clause</code>), use the lower branch of the syntax (with parentheses).

You can specify multiple analytic functions in the same query, each with the same or different PARTITION BY keys.

If the objects being queried have the parallel attribute, and if you specify an analytic function with the <code>query_partition_clause</code>, then the function computations are parallelized as well.

Valid values of *value_expr* are constants, columns, nonanalytic functions, function expressions, or expressions involving any of these.

order_by_clause

Use the <code>order_by_clause</code> to specify how data is ordered within a partition. For all analytic functions you can order the values in a partition on multiple keys, each defined by a <code>value expr</code> and each qualified by an ordering sequence.

Within each function, you can specify multiple ordering expressions. Doing so is especially useful when using functions that rank values, because the second expression can resolve ties between identical values for the first expression.

Whenever the <code>order_by_clause</code> results in identical values for multiple rows, the function behaves as follows:

- CUME_DIST, DENSE_RANK, NTILE, PERCENT_RANK, and RANK return the same result for each of the rows
- ROW_NUMBER assigns each row a distinct value even if there is a tie based on the
 order_by_clause. The value is based on the order in which the row is processed, which
 may be nondeterministic if the ORDER BY does not guarantee a total ordering.
- For all other analytic functions, the result depends on the window specification. If you
 specify a logical window with the RANGE keyword, then the function returns the same result
 for each of the rows. If you specify a physical window with the ROWS keyword, then the
 result is nondeterministic.

Restrictions on the ORDER BY Clause

The following restrictions apply to the ORDER BY clause:

- When used in an analytic function, the <code>order_by_clause</code> must take an expression (<code>expr</code>). The <code>SIBLINGS</code> keyword is not valid (it is relevant only in hierarchical queries). Position (<code>position</code>) and column aliases (<code>c_alias</code>) are also invalid. Otherwise this <code>order_by_clause</code> is the same as that used to order the overall query or subquery.
- An analytic function that uses the RANGE keyword can use multiple sort keys in its ORDER BY clause if it specifies any of the following windows:
 - RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW. The short form of this is RANGE UNBOUNDED PRECEDING.



- RANGE BETWEEN CURRENT ROW AND UNBOUNDED FOLLOWING
- RANGE BETWEEN CURRENT ROW AND CURRENT ROW
- RANGE BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING

Window boundaries other than these four can have only one sort key in the ORDER BY clause of the analytic function. This restriction does not apply to window boundaries specified by the ROW keyword.

ASC | DESC

Specify the ordering sequence (ascending or descending). ASC is the default.

NULLS FIRST | NULLS LAST

Specify whether returned rows containing nulls should appear first or last in the ordering sequence.

NULLS LAST is the default for ascending order, and NULLS FIRST is the default for descending order.

Analytic functions always operate on rows in the order specified in the <code>order_by_clause</code> of the function. However, the <code>order_by_clause</code> of the function does not guarantee the order of the result. Use the <code>order_by_clause</code> of the query to guarantee the final result ordering.



order_by_clause of SELECT for more information on this clause

windowing_clause

Some analytic functions allow the <code>windowing_clause</code>. In the listing of analytic functions at the end of this section, the functions that allow the <code>windowing_clause</code> are followed by an asterisk (*).

ROWS | RANGE | GROUPS

The keywords ROWS, RANGE, and GROUPS are options to define a window frame unit used for calculating the function result. The function is then applied to all the rows in the window. The window moves through the query result set or partition from top to bottom.

- Use ROWS to specify the window frame extent by counting rows forward or backward from the current row. ROWS allows any number of sort keys, of any ordered data types.
- Use RANGE to specify the window frame extent as a logical offset. RANGE allows only one sort key, and its declared data type must allow addition and subtraction operations, for example they must be numeric, datetime, or interval data types.
- Use GROUPS to specify the window frame extent with both ROWS and RANGE characteristics.
 Like ROWS a GROUPS window can have any number of sort keys, or any ordered types. Like RANGE, a GROUPS window does not make cutoffs between adjacent rows with the same values in the sort keys.

You cannot specify this clause unless you have specified the <code>order_by_clause</code>. Some window boundaries defined by the <code>RANGE</code> clause let you specify only one expression in the <code>order_by_clause</code>. Refer to Restrictions on the ORDER BY Clause.



The value returned by an analytic function with a logical offset is always deterministic. However, the value returned by an analytic function with a physical offset may produce nondeterministic results unless the ordering expression results in a unique ordering. You may have to specify multiple columns in the <code>order by clause</code> to achieve this unique ordering.

BETWEEN ... AND

Use the BETWEEN ... AND clause to specify a start point and end point for the window. The first expression (before AND) defines the start point and the second expression (after AND) defines the end point.

If you omit Between and specify only one end point, then Oracle considers it the start point, and the end point defaults to the current row.

UNBOUNDED PRECEDING

Specify UNBOUNDED PRECEDING to indicate that the window starts at the first row of the partition. This is the start point specification and cannot be used as an end point specification.

UNBOUNDED FOLLOWING

Specify UNBOUNDED FOLLOWING to indicate that the window ends at the last row of the partition. This is the end point specification and cannot be used as a start point specification.

CURRENT ROW

As a start point, CURRENT ROW specifies that the window begins at the current row or value (depending on whether you have specified ROW or RANGE, respectively). In this case the end point cannot be <code>value expr PRECEDING</code>.

As an end point, CURRENT ROW specifies that the window ends at the current row or value (depending on whether you have specified ROW or RANGE, respectively). In this case the start point cannot be <code>value exprFOLLOWING</code>.

value_expr PRECEDING or value_expr FOLLOWING

For RANGE or ROW:

- If value_expr FOLLOWING is the start point, then the end point must be value_expr FOLLOWING.
- If value_expr PRECEDING is the end point, then the start point must be value_expr PRECEDING.

If you are defining a logical window defined by an interval of time in numeric format, then you may need to use conversion functions.



NUMTOYMINTERVAL and NUMTODSINTERVAL for information on converting numeric times into intervals

If you specified ROWS:

• value_expr is a physical offset. It must be a constant or expression and must evaluate to a positive numeric value.



• If value_expr is part of the start point, then it must evaluate to a row before the end point.

If you specified RANGE:

- value_expr is a logical offset. It must be a constant or expression that evaluates to a
 positive numeric value or an interval literal. Refer to Literals for information on interval
 literals.
- You can specify only one expression in the order_by_clause.
- If value_expr evaluates to a numeric value, then the ORDER BY expr must be a numeric or DATE data type.
- If value_expr evaluates to an interval value, then the ORDER BY expr must be a DATE data type.

If you omit the windowing_clause entirely, then the default is RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW.

EXCLUDE

You can remove rows, groups, and ties from the window frame with the EXCLUDE options:

- If you specify EXCLUDE CURRENT ROW, and the current row in in the window frame, then the current row is removed from the window frame.
- If you specify EXCLUDE GROUP, then the current row and any peers of the current row are removed from the window frame.
- If you specify EXCLUDE TIES, then the peers of the current row are removed from the window frame. The current row is retained. Note, that if the current row is previously removed from the window frame, it remains removed.
- If you specify EXCLUDE NO OTHERS, then no additional rows are removed from the window frame. This is the default option.

Analytic functions are commonly used in data warehousing environments. In the list of analytic functions that follows, functions followed by an asterisk (*) allow the full syntax, including the windowing clause.

```
AVG *
BIT_AND_AGG*
BIT_OR_AGG*
BIT XOR AGG*
BOOLEAN_AND_AGG*
BOOLEAN OR AGG*
CHECKSUM*
CLUSTER DETAILS
CLUSTER DISTANCE
CLUSTER ID
CLUSTER_PROBABILITY
CLUSTER_SET
CORR *
COUNT *
COVAR POP *
COVAR SAMP*
CUME DIST
DENSE_RANK
EVERY*
```



FEATURE DETAILS FEATURE_ID FEATURE SET FEATURE_VALUE **FIRST** FIRST_VALUE * KURTOSIS_POP* KURTOSIS_SAMP* LAG **LAST** LAST_VALUE * **LEAD LISTAGG** MAX * **MEDIAN** MIN * NTH_VALUE * **NTILE** PERCENT_RANK PERCENTILE_CONT PERCENTILE DISC **PREDICTION** PREDICTION_COST PREDICTION_DETAILS PREDICTION_PROBABILITY PREDICTION_SET **RANK** RATIO_TO_REPORT REGR_ (Linear Regression) Functions * **ROW NUMBER** STDDEV * **SKEWNESS POP*** SKEWNESS_SAMP* STDDEV_POP * STDDEV_SAMP * SUM * VAR POP* VAR_SAMP * **VARIANCE ***

See Also:

Oracle Database Data Warehousing Guide for more information on these functions and for scenarios illustrating their use

Data Cartridge Functions

Data Cartridge functions are useful for Data Cartridge developers. The Data Cartridge functions are:

```
DATAOBJ_TO_MAT_PARTITION DATAOBJ_TO_PARTITION
```

Model Functions

Model functions can be used only in the $model_clause$ of the SELECT statement. The model functions are:

CV ITERATION_NUMBER PRESENTNNV PRESENTV PREVIOUS

Object Reference Functions

Object reference functions manipulate REF values, which are references to objects of specified object types. The object reference functions are:

DEREF MAKE_REF REF REFTOHEX VALUE



Oracle Database Object-Relational Developer's Guide for more information about REF data types

OLAP Functions

OLAP functions returns data from a dimensional object in two-dimension relational format. The OLAP function is:

CUBE_TABLE

Single-Row Functions

Single-row functions return a single result row for every row of a queried table or view. These functions can appear in select lists, WHERE clauses, START WITH and CONNECT BY clauses, and HAVING clauses.

Numeric Functions

Numeric functions accept numeric input and return numeric values. Most numeric functions return NUMBER values that are accurate to 38 decimal digits. The transcendental functions COS, COSH, EXP, LN, LOG, SIN, SINH, SQRT, TAN, and TANH are accurate to 36 decimal digits. The transcendental functions ACOS, ASIN, ATAN, and ATAN2 are accurate to 30 decimal digits. The numeric functions are:

ABS ACOS ASIN ATAN ATAN2 **BITAND** CEIL (number) COS **COSH EXP** FLOOR (number) LN LOG MOD **NANVL POWER REMAINDER ROUND** (number) **SIGN** SIN SINH **SQRT TAN TANH** TRUNC (number) WIDTH_BUCKET

Character Functions Returning Character Values

Character functions that return character values return values of the following data types unless otherwise documented:

- If the input argument is CHAR or VARCHAR2, then the value returned is VARCHAR2.
- If the input argument is NCHAR or NVARCHAR2, then the value returned is NVARCHAR2.

The length of the value returned by the function is limited by the maximum length of the data type returned.

- For functions that return CHAR or VARCHAR2, if the length of the return value exceeds the limit, then Oracle Database truncates it and returns the result without an error message.
- For functions that return CLOB values, if the length of the return values exceeds the limit, then Oracle raises an error and returns no data.

The character functions that return character values are:

```
CHR
CONCAT
INITCAP
LOWER
LPAD
LTRIM
NCHR
NLS INITCAP
NLS_LOWER
NLS UPPER
NLSSORT
REGEXP REPLACE
REGEXP SUBSTR
REPLACE
RPAD
RTRIM
SOUNDEX
SUBSTR
TRANSLATE
TRANSLATE ... USING
TRIM
UPPER
```

Character Functions Returning Number Values

Character functions that return number values can take as their argument any character data type. The character functions that return number values are:

```
ASCII
INSTR
LENGTH
REGEXP_COUNT
REGEXP_INSTR
```

Character Set Functions

The character set functions return information about the character set. The character set functions are:

```
NLS_CHARSET_ID
NLS_CHARSET_ID
NLS_CHARSET_NAME
```

Collation Functions

The collation functions return information about collation settings. The collation functions are:

```
COLLATION
NLS_COLLATION_ID
NLS_COLLATION_NAME
```



Datetime Functions

Datetime functions operate on date (DATE), timestamp (TIMESTAMP, TIMESTAMP WITH TIME ZONE, and TIMESTAMP WITH LOCAL TIME ZONE), and interval (INTERVAL DAY TO SECOND, INTERVAL YEAR TO MONTH) values.

Some of the datetime functions were designed for the Oracle DATE data type (ADD_MONTHS, CURRENT_DATE, LAST_DAY, NEW_TIME, and NEXT_DAY). If you provide a timestamp value as their argument, then Oracle Database internally converts the input type to a DATE value and returns a DATE value. The exceptions are the MONTHS_BETWEEN function, which returns a number, and the ROUND and TRUNC functions, which do not accept timestamp or interval values at all.

The remaining datetime functions were designed to accept any of the three types of data (date, timestamp, and interval) and to return a value of one of these types.

All of the datetime functions that return current system datetime information, such as SYSDATE, SYSTIMESTAMP, CURRENT_TIMESTAMP, and so forth, are evaluated once for each SQL statement, regardless how many times they are referenced in that statement.

The datetime functions are:

ADD MONTHS CEIL (datetime) CURRENT_DATE CURRENT_TIMESTAMP **DBTIMEZONE** EXTRACT (datetime) FLOOR (datetime) FROM TZ LAST DAY **LOCALTIMESTAMP** MONTHS BETWEEN **NEW_TIME NEXT DAY** NUMTODSINTERVAL NUMTOYMINTERVAL ORA_DST_AFFECTED ORA DST CONVERT ORA_DST_ERROR **ROUND** (datetime) **SESSIONTIMEZONE** SYS_EXTRACT_UTC SYSDATE **SYSTIMESTAMP** TO CHAR (datetime) TO_DSINTERVAL TO TIMESTAMP TO_TIMESTAMP_TZ TO YMINTERVAL TRUNC (datetime) TZ_OFFSET



General Comparison Functions

The general comparison functions determine the greatest and or least value from a set of values. The general comparison functions are:

GREATEST LEAST

Conversion Functions

Conversion functions convert a value from one data type to another. Generally, the form of the function names follows the convention *datatype* TO *datatype*. The first data type is the input data type. The second data type is the output data type. The SQL conversion functions are:

ASCIISTR BIN TO NUM **CAST CHARTOROWID COMPOSE CONVERT DECOMPOSE HEXTORAW** NUMTODSINTERVAL **NUMTOYMINTERVAL RAWTOHEX RAWTONHEX** ROWIDTOCHAR **ROWIDTONCHAR** SCN_TO_TIMESTAMP TIMESTAMP_TO_SCN TO BINARY DOUBLE TO_BINARY_FLOAT TO_BLOB (bfile) TO_BLOB (raw) TO_CHAR (bfile|blob) TO_CHAR (character) TO_CHAR (datetime) TO_CHAR (number) TO_CLOB (bfile|blob) TO_CLOB (character) TO DATE TO DSINTERVAL TO LOB TO MULTI_BYTE TO_NCHAR (character) TO_NCHAR (datetime) TO_NCHAR (number) TO_NCLOB TO_NUMBER TO_SINGLE_BYTE

```
TO_TIMESTAMP
TO_TIMESTAMP_TZ
TO_YMINTERVAL
TREAT
UNISTR
VALIDATE_CONVERSION
```

Large Object Functions

The large object functions operate on LOBs. The large object functions are:

```
BFILENAME
EMPTY_BLOB, EMPTY_CLOB
```

Collection Functions

The collection functions operate on nested tables and varrays. The SQL collection functions are:

```
CARDINALITY
COLLECT
POWERMULTISET
POWERMULTISET_BY_CARDINALITY
SET
```

Hierarchical Functions

Hierarchical functions applies hierarchical path information to a result set. The hierarchical function is:

```
SYS CONNECT BY PATH
```

Oracle Machine Learning for SQL Functions

The Oracle Machine Learning for SQL functions use analytics to score data. The functions can apply a mining model schema object to the data, or they can dynamically mine the data by executing an analytic clause. The OML4SQL functions can be applied to models built using the native algorithms of Oracle, as well as those built using R through the extensibility mechanism.

The Oracle Machine Learning for SQL functions are:

```
CLUSTER_DETAILS
CLUSTER_DISTANCE
CLUSTER_ID
CLUSTER_PROBABILITY
CLUSTER_SET
FEATURE_COMPARE
FEATURE_DETAILS
FEATURE_ID
FEATURE_SET
FEATURE_SET
FEATURE_VALUE
ORA_DM_PARTITION_NAME
PREDICTION
```



PREDICTION_BOUNDS
PREDICTION_COST
PREDICTION_DETAILS
PREDICTION_PROBABILITY
PREDICTION_SET
VECTOR_EMBEDDING

See Also:

- Oracle Machine Learning for SQL Concepts to learn about Oracle Machine Learning for SQL
- Oracle Machine Learning for SQL User's Guide for information about scoring

XML Functions

The XML functions operate on or return XML documents or fragments. These functions use arguments that are not defined as part of the ANSI/ISO/IEC SQL Standard but are defined as part of the World Wide Web Consortium (W3C) standards. The processing and operations that the functions perform are defined by the relevant W3C standards. The table below provides a link to the appropriate section of the W3C standard for the rules and guidelines that apply to each of these XML-related arguments. A SQL statement that uses one of these XML functions, where any of the arguments does not conform to the relevant W3C syntax, will result in an error. Of special note is the fact that not every character that is allowed in the value of a database column is considered legal in XML.

Syntax Element	W3C Standard URL	
value_expr	http://www.w3.org/TR/2006/REC-xml-20060816	
Xpath_string	http://www.w3.org/TR/1999/REC-xpath-19991116	
XQuery_string	http://www.w3.org/TR/2007/REC-xquery-semantics-20070123/	
	http://www.w3.org/TR/xquery-update-10/	
namespace_string	http://www.w3.org/TR/2006/REC-xml-names-20060816/	
identifier	http://www.w3.org/TR/2006/REC-xml-20060816/#NT-Nmtoken	

For more information about selecting and querying XML data using these functions, including information on formatting output, refer to *Oracle XML DB Developer's Guide*

The SQL XML functions are:

DEPTH
EXISTSNODE
EXTRACT (XML)
EXTRACTVALUE
PATH
SYS_DBURIGEN
SYS_XMLAGG
SYS_XMLGEN
XMLAGG
XMLAGG
XMLCAST



```
XMLCDATA
XMLCOLATTVAL
XMLCOMMENT
XMLCONCAT
XMLDIFF
XMLELEMENT
XMLEXISTS
XMLFOREST
XMLISVALID
XMLPARSE
XMLPATCH
XMLPI
XMLQUERY
XMLSEQUENCE
XMLSERIALIZE
XMLTABLE
XMLTRANSFORM
```

JSON Functions

JavaScript Object Notation (JSON) functions allow you to guery and generate JSON data.

The following SQL/JSON functions allow you to query JSON data:

```
JSON_QUERY
JSON_TABLE
JSON_VALUE
```

The following SQL/JSON functions allow you to generate JSON data:

```
JSON_ARRAY
JSON_ARRAYAGG
JSON_OBJECT
JSON_OBJECTAGG
JSON Type Constructor
JSON_SCALAR
JSON_SERIALIZE
JSON_TRANSFORM
```

The following Oracle SQL function creates a JSON data guide:

```
JSON_DATAGUIDE
```

Encoding and Decoding Functions

The encoding and decoding functions let you inspect and decode data in the database. The encoding and decoding functions are:

```
DECODE
DUMP
ORA_HASH
STANDARD_HASH
VSIZE
```



NULL-Related Functions

The NULL-related functions facilitate null handling. The NULL-related functions are:

COALESCE LNNVL NANVL NULLIF NVL NVL2

Environment and Identifier Functions

The environment and identifier functions provide information about the instance and session. The environment and identifier functions are:

```
CON_DBID_TO_ID
CON_GUID_TO_ID
CON_NAME_TO_ID
CON_UID_TO_ID
ORA_INVOKING_USER
ORA_INVOKING_USERID
SYS_CONTEXT
SYS_GUID
SYS_TYPEID
UID
USER
USERENV
```

Domain Functions

Purpose

Use the following domain functions to work with usecase domains more efficiently:

- DOMAIN_DISPLAY
- DOMAIN_ORDER
- DOMAIN_NAME
- DOMAIN_CHECK
- DOMAIN_CHECK_TYPE

Vector Functions

Purpose

You can use the following vector functions in Oracle AI Vector Search to create and manipulate vectors:

Vector Distance Functions

VECTOR_DISTANCE

- L1_DISTANCE
- L2_DISTANCE
- COSINE_DISTANCE
- INNER_PRODUCT

Vector Constructors

- TO_VECTOR
- VECTOR

Vector Serializers

- FROM_VECTOR
- VECTOR_SERIALIZE

Other Common Vector Functions

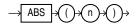
- VECTOR CHUNKS
- VECTOR_DIMS
- VECTOR_DIMENSION_COUNT
- VECTOR_DIMENSION_FORMAT
- VECTOR_EMBEDDING
- VECTOR_NORM



Al Vector Search User's Guide

ABS

Syntax



Purpose

ABS returns the absolute value of n.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.



Table 2-9 for more information on implicit conversion



Examples

The following example returns the absolute value of -15:

ACOS

Syntax



Purpose

ACOS returns the arc cosine of n. The argument n must be in the range of -1 to 1, and the function returns a value in the range of 0 to pi, expressed in radians.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



Table 2-9 for more information on implicit conversion

Examples

The following example returns the arc cosine of .3:

```
SELECT ACOS(.3)"Arc_Cosine"
FROM DUAL;

Arc_Cosine
------
1.26610367
```

ADD_MONTHS

Syntax





Purpose

ADD_MONTHS returns the date *date* plus *integer* months. A month is defined by the session parameter NLS_CALENDAR. The date argument can be a datetime value or any value that can be implicitly converted to DATE. The *integer* argument can be an integer or any value that can be implicitly converted to an integer. The return type is always DATE, regardless of the data type of *date*. If *date* is the last day of the month or if the resulting month has fewer days than the day component of *date*, then the result is the last day of the resulting month. Otherwise, the result has the same day component as *date*.



Table 2-9 for more information on implicit conversion

Examples

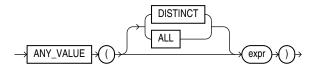
The following example returns the month after the <code>hire_date</code> in the sample table <code>employees:</code>

```
SELECT TO_CHAR(ADD_MONTHS(hire_date, 1), 'DD-MON-YYYY') "Next month"
  FROM employees
  WHERE last_name = 'Baer';

Next Month
-----07-JUL-2002
```

ANY_VALUE

Syntax



Purpose

ANY_VALUE returns a single non-deterministic value of expr. You can use it as an aggregate function.

Use ANY_VALUE to optimize a query that has a GROUP BY clause. ANY_VALUE returns a value of an expression in a group. It is optimized to return the first value.

It ensures that there are no comparisons for any incoming row and also eliminates the necessity to specify every column as part of the <code>GROUP BY</code> clause. Because it does not compare values, <code>ANY_VALUE</code> returns a value more quickly than <code>MIN</code> or <code>MAX</code> in a <code>GROUP BY</code> query.

Semantics

ALL, DISTINCT: These keywords are supported by ANY_VALUE although they have no effect on the result of the query.

 ${\tt expr}$: The expression can be a column, constant, bind variable, or an expression involving them.



NULL values in the expression are ignored.

Supports all of the data types, except for LONG, LOB, FILE, or COLLECTION.

If you use LONG, ORA-00997 is raised.

If you use LOB, FILE, or COLLECTION data types, ORA-00932 is raised.

ANY VALUE follows the same rules as MIN and MAX.

Returns any value within each group based on the GROUP BY specification. Returns NULL if all rows in the group have NULL expression values.

The result of any value is not deterministic.

Restrictions

XMLType and ANYDATA are not supported.

Example 7-1 Using ANY_VALUE As an Aggregate Function

This example uses <code>ANY_VALUE</code> as an aggregate function in a <code>GROUP</code> BY query of the SH schema.

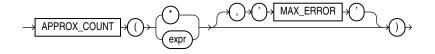
```
SELECT c.cust_id, ANY_VALUE(cust_last_name), SUM(amount_sold)
  FROM customers c, sales s
  WHERE s.cust_id = c.cust_id
  GROUP BY c.cust id;
```

In the following result of the query, only the first eleven rows are shown.

CUST_ID	ANY_VALUE(CUST_LAST_NAME)	SUM (AMOUNT_SOLD)
6950	Sandburg	78
17920	Oliver	3201
66800	Case	2024
37280	Edwards	2256
109850	Lindegreen	757
3910	Oddell	185
84700	Marker	164.4
26380	Remler	118
11600	Орру	158
23030	Rothrock	533
42780	Zanis	182
630 rows	s selected.	

APPROX_COUNT

Syntax





Purpose

APPROX_COUNT returns the approximate count of an expression. If you supply MAX_ERROR as the second argument, then the function returns the maximum error between the actual and approximate count.

You must use this function with a corresponding APPROX_RANK function in the HAVING clause. If a query uses APPROX_COUNT, APPROX_SUM, or APPROX_RANK, then the query must not use any other aggregation functions.

Examples

The following query returns the 10 most common jobs within every department:

APPROX_COUNT_DISTINCT

Syntax



Purpose

APPROX_COUNT_DISTINCT returns the approximate number of rows that contain a distinct value for *expr*.

This function provides an alternative to the COUNT (DISTINCT expr) function, which returns the exact number of rows that contain distinct values of expr. APPROX_COUNT_DISTINCT processes large amounts of data significantly faster than COUNT, with negligible deviation from the exact result.

For expr, you can specify a column of any scalar data type other than BFILE, BLOB, CLOB, LONG, LONG RAW, or NCLOB.

APPROX_COUNT_DISTINCT ignores rows that contain a null value for <code>expr</code>. This function returns a <code>NUMBER</code>.



See Also:

- COUNT for more information on the COUNT (DISTINCT expr) function
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation APPROX_COUNT_DISTINCT uses to compare character values for expr

Examples

The following statement returns the approximate number of rows with distinct values for manager id:

The following statement returns the approximate number of distinct customers for each product:

```
SELECT prod id, APPROX COUNT DISTINCT(cust id) AS "Number of Customers"
 FROM sales
 GROUP BY prod id
 ORDER BY prod id;
  PROD ID Number of Customers
       13
                        2516
       14
                        2030
       15
                        2105
       16
                        2367
       17
                       2093
       18
                       2975
       19
                       2630
                        3791
```

APPROX_COUNT_DISTINCT_AGG

Syntax



Purpose

APPROX_COUNT_DISTINCT_AGG takes as its input a column of details containing information about approximate distinct value counts, and enables you to perform aggregations of those counts.

For detail, specify a column of details created by the APPROX_COUNT_DISTINCT_DETAIL function or the APPROX COUNT DISTINCT AGG function. This column is of data type BLOB.

You can specify this function in a SELECT statement with a GROUP BY clause to aggregate the information contained in the details within each group of rows and return a single detail for each group.

This function returns a BLOB value, called a detail, which contains information about the count aggregations in a special format. You can store details returned by this function in a table or materialized view, and then again use the APPROX_COUNT_DISTINCT_AGG function to further aggregate those details, or use the TO_APPROX_COUNT_DISTINCT function to convert the detail values to human-readable NUMBER values.

See Also:

- APPROX_COUNT_DISTINCT_DETAIL
- TO APPROX COUNT DISTINCT

Examples

Refer to APPROX_COUNT_DISTINCT_AGG: Examples for examples of using the APPROX_COUNT_DISTINCT_AGG function in conjunction with the APPROX_COUNT_DISTINCT_DETAIL and TO APPROX COUNT DISTINCT functions.

APPROX COUNT DISTINCT DETAIL

Syntax



Purpose

APPROX_COUNT_DISTINCT_DETAIL calculates information about the approximate number of rows that contain a distinct value for expr and returns a BLOB value, called a detail, which contains that information in a special format.

For *expr*, you can specify a column of any scalar data type other than BFILE, BLOB, CLOB, LONG, LONG RAW, or NCLOB. This function ignores rows for which the value of *expr* is null.

This function is commonly used with the GROUP BY clause in a SELECT statement. When used in this way, it calculates approximate distinct value count information for *expr* within each group of rows and returns a single detail for each group.

The details returned by APPROX_COUNT_DISTINCT_DETAIL can be used as input to the APPROX_COUNT_DISTINCT_AGG function, which enables you to perform aggregations of the details, or the TO_APPROX_COUNT_DISTINCT function, which converts a detail to a human-readable distinct count value. You can use these three functions together to perform resource-intensive approximate count calculations once, store the resulting details, and then perform efficient aggregations and queries on those details. For example:

Use the APPROX_COUNT_DISTINCT_DETAIL function to calculate approximate distinct value
count information and store the resulting details in a table or materialized view. These
could be highly-granular details, such as city demographic counts or daily sales counts.

- 2. Use the APPROX_COUNT_DISTINCT_AGG function to aggregate the details obtained in the previous step and store the resulting details in a table or materialized view. These could be details of lower granularity, such as state demographic counts or monthly sales counts.
- 3. Use the TO_APPROX_COUNT_DISTINCT function to convert the stored detail values to human-readable NUMBER values. You can use the TO_APPROX_COUNT_DISTINCT function to query detail values created by the APPROX_COUNT_DISTINCT_DETAIL function or the APPROX_COUNT_DISTINCT_AGG function.

See Also:

- APPROX_COUNT_DISTINCT_AGG
- TO_APPROX_COUNT_DISTINCT

Examples

The examples in this section demonstrate how to use the APPROX_COUNT_DISTINCT_DETAIL, APPROX_COUNT_DISTINCT_AGG, and TO_APPROX_COUNT_DISTINCT functions together to perform resource-intensive approximate count calculations once, store the resulting details, and then perform efficient aggregations and queries on those details.

APPROX_COUNT_DISTINCT_DETAIL: Example

The following statement queries the tables <code>sh.times</code> and <code>sh.sales</code> for the approximate number of distinct products sold each day. The <code>APPROX_COUNT_DISTINCT_DETAIL</code> function returns the information in a detail, called <code>daily_detail</code>, for each day that products were sold. The returned details are stored in a materialized view called <code>daily_prod_count_mv</code>.

```
CREATE MATERIALIZED VIEW daily_prod_count_mv AS

SELECT t.calendar_year year,

t.calendar_month_number month,

t.day_number_in_month day,

APPROX_COUNT_DISTINCT_DETAIL(s.prod_id) daily_detail

FROM times t, sales s

WHERE t.time_id = s.time_id

GROUP BY t.calendar year, t.calendar month number, t.day number in month;
```

APPROX_COUNT_DISTINCT_AGG: Examples

The following statement uses the APPROX_COUNT_DISTINCT_AGG function to read the daily details stored in daily_prod_count_mv and create aggregated details that contain the approximate number of distinct products sold each month. These aggregated details are stored in a materialized view called monthly_prod_count_mv.

```
CREATE MATERIALIZED VIEW monthly_prod_count_mv AS

SELECT year,

month,

APPROX_COUNT_DISTINCT_AGG(daily_detail) monthly_detail

FROM daily_prod_count_mv

GROUP BY year, month;
```

The following statement is similar to the previous statement, except it creates aggregated details that contain the approximate number of distinct products sold each year. These aggregated details are stored in a materialized view called annual prod count mv.

```
CREATE MATERIALIZED VIEW annual_prod_count_mv AS SELECT year,

APPROX_COUNT_DISTINCT_AGG(daily_detail) annual_detail FROM daily_prod_count_mv GROUP BY year;
```

TO_APPROX_COUNT_DISTINCT: Examples

The following statement uses the ${\tt TO_APPROX_COUNT_DISTINCT}$ function to query the daily detail information stored in ${\tt daily_prod_count_mv}$ and return the approximate number of distinct products sold each day:

```
SELECT year,

month,

day,

TO_APPROX_COUNT_DISTINCT(daily_detail) "NUM PRODUCTS"

FROM daily_prod_count_mv

ORDER BY year, month, day;
```

YEAR	MONTH	DAY NUM	PRODUCTS
1998	1	1	24
1998	1	2	25
1998	1	3	11
1998	1	4	34
1998	1	5	10
1998	1	6	8
1998	1	7	37
1998	1	8	26
1998	1	9	25
1998	1	10	38

The following statement uses the TO_APPROX_COUNT_DISTINCT function to query the monthly detail information stored in monthly_prod_count_mv and return the approximate number of distinct products sold each month:

YEAR	MONTH	NUM PRODUCTS
1998	1	57
1998	2	56
1998	3	55
1998	4	49
1998	5	49
1998	6	48
1998	7	54
1998	8	56
1998	9	55



1998 10 57

. . .

The following statement uses the <code>TO_APPROX_COUNT_DISTINCT</code> function to query the annual detail information stored in <code>annual_prod_count_mv</code> and return the approximate number of distinct products sold each year:

```
SELECT year,

TO_APPROX_COUNT_DISTINCT(annual_detail) "NUM PRODUCTS"

FROM annual_prod_count_mv

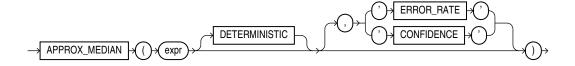
ORDER BY year;

YEAR NUM PRODUCTS

1998 60
1999 72
2000 72
2001 71
```

APPROX_MEDIAN

Syntax



Purpose

APPROX_MEDIAN is an approximate inverse distribution function that assumes a continuous distribution model. It takes a numeric or datetime value and returns an approximate middle value or an approximate interpolated value that would be the middle value once the values are sorted. Nulls are ignored in the calculation.

This function provides an alternative to the MEDIAN function, which returns the exact middle value or interpolated value. APPROX_MEDIAN processes large amounts of data significantly faster than MEDIAN, with negligible deviation from the exact result.

For expr, specify the expression for which the approximate median value is being calculated. The acceptable data types for expr, and the return value data type for this function, depend on the algorithm that you specify with the DETERMINISTIC clause.

DETERMINISTIC

This clause lets you specify the type of algorithm this function uses to calculate the approximate median value.

• If you specify DETERMINISTIC, then this function calculates a deterministic approximate median value. In this case, <code>expr</code> must evaluate to a numeric value, or to a value that can be implicitly converted to a numeric value. The function returns the same data type as the numeric data type of its argument.

• If you omit DETERMINSTIC, then this function calculates a nondeterministic approximate median value. In this case, *expr* must evaluate to a numeric or datetime value, or to a value that can be implicitly converted to a numeric or datetime value. The function returns the same data type as the numeric or datetime data type of its argument.

ERROR_RATE | CONFIDENCE

These clauses let you determine the accuracy of the value calculated by this function. If you specify one of these clauses, then instead of returning the approximate median value for <code>expr</code>, the function returns a decimal value from 0 to 1, inclusive, which represents one of the following values:

- If you specify ERROR_RATE, then the return value represents the error rate for the approximate median value calculation for *expr*.
- If you specify CONFIDENCE, then the return value represents the confidence level for the error rate that is returned when you specify ERROR_RATE.

See Also:

- MEDIAN
- APPROX_PERCENTILE which returns, for a given percentile, the approximate
 value that corresponds to that percentile by way of interpolation. APPROX_MEDIAN
 is the specific case of APPROX_PERCENTILE where the percentile value is 0.5.

Examples

The following query returns the deterministic approximate median salary for each department in the hr.employees table:

```
SELECT department id "Department",
     APPROX MEDIAN(salary DETERMINISTIC) "Median Salary"
 FROM employees
 GROUP BY department id
 ORDER BY department id;
Department Median Salary
 _____
      10
               4400
      20
              6000
      30
               2765
      40
               6500
      50
              3100
      60
               4800
      70
              10000
      80
               9003
              17000
      90
     100
               7739
     110
               8300
                7000
```

The following query returns the error rates for the approximate median salaries that were returned by the previous query:

```
SELECT department_id "Department",
          APPROX_MEDIAN(salary DETERMINISTIC, 'ERROR_RATE') "Error Rate"
FROM employees
```



```
GROUP BY department id
 ORDER BY department id;
Department Error Rate
 _____
       10 .002718282
       20 .021746255
       30 .021746255
       40 .002718282
       50 .019027973
       60 .019027973
       70 .002718282
       80 .021746255
       90 .021746255
      100 .019027973
      110 .019027973
          .002718282
```

The following query returns the confidence levels for the error rates that were returned by the previous query:

```
SELECT department id "Department",
       APPROX MEDIAN (salary DETERMINISTIC, 'CONFIDENCE') "Confidence Level"
 FROM employees
 GROUP BY department id
 ORDER BY department id;
Department Confidence Level
       10 .997281718
20 .999660215
30 .999660215
40 .997281718
                .997281718
                .999611674
        50
        60
                .999611674
                .997281718
        70
                .999660215
        80
                .999660215
        90
       100
                 .999611674
       110
                 .999611674
                 .997281718
```

The following query returns the nondeterministic approximate median hire date for each department in the hr.employees table:

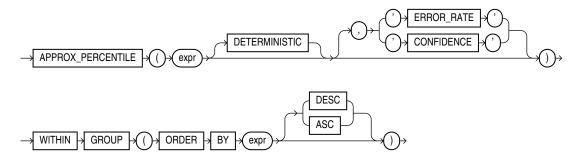
```
SELECT department id "Department",
      APPROX MEDIAN(hire date) "Median Hire Date"
 FROM employees
 GROUP BY department id
 ORDER BY department id;
Department Median Hire Date
            17-SEP-03
17-FEB-04
       10
       20
       30
                24-JUL-05
       40
                 07-JUN-02
                15-MAR-06
       50
       60
                05-FEB-06
       70
                07-JUN-02
                23-MAR-06
       80
       90
                17-JUN-03
```



100 28-SEP-05 110 07-JUN-02 24-MAY-07

APPROX_PERCENTILE

Syntax



Purpose

APPROX_PERCENTILE is an approximate inverse distribution function. It takes a percentile value and a sort specification, and returns the value that would fall into that percentile value with respect to the sort specification. Nulls are ignored in the calculation

This function provides an alternative to the PERCENTILE_CONT and PERCENTILE_DISC functions, which returns the exact results. APPROX_PERCENTILE processes large amounts of data significantly faster than PERCENTILE_CONT and PERCENTILE_DISC, with negligible deviation from the exact result.

The first expr is the percentile value, which must evaluate to a numeric value between 0 and 1.

The second expr, which is part of the ORDER BY clause, is a single expression over which this function calculates the result. The acceptable data types for expr, and the return value data type for this function, depend on the algorithm that you specify with the DETERMINISTIC clause.

DETERMINISTIC

This clause lets you specify the type of algorithm this function uses to calculate the return value.

- If you specify DETERMINISTIC, then this function calculates a deterministic result. In this case, the ORDER BY clause expression must evaluate to a numeric value, or to a value that can be implicitly converted to a numeric value, in the range -2,147,483,648 through 2,147,483,647. The function rounds numeric input to the closest integer. The function returns the same data type as the numeric data type of the ORDER BY clause expression. The return value is not necessarily one of the values of *expr*
- If you omit DETERMINSTIC, then this function calculates a nondeterministic result. In this case, the ORDER BY clause expression must evaluate to a numeric or datetime value, or to a value that can be implicitly converted to a numeric or datetime value. The function returns the same data type as the numeric or datetime data type of the ORDER BY clause expression. The return value is one of the values of *expr*.

ERROR_RATE | CONFIDENCE

These clauses let you determine the accuracy of the result calculated by this function. If you specify one of these clauses, then instead of returning the value that would fall into the



specified percentile value for expr, the function returns a decimal value from 0 to 1, inclusive, which represents one of the following values:

- If you specify ERROR_RATE, then the return value represents the error rate for calculating the value that would fall into the specified percentile value for *expr*.
- If you specify CONFIDENCE, then the return value represents the confidence level for the error rate that is returned when you specify ERROR RATE.

DESC | ASC

Specify the sort specification for the calculating the value that would fall into the specified percentile value. Specify DESC to sort the ORDER BY clause expression values in descending order, or ASC to sort the values in ascending order. ASC is the default.

See Also:

- PERCENTILE CONT and PERCENTILE DISC
- APPROX_MEDIAN, which is the specific case of APPROX_PERCENTILE where the
 percentile value is 0.5

Examples

The following query returns the deterministic approximate 25th percentile, 50th percentile, and 75th percentile salaries for each department in the hr.employees table. The salaries are sorted in ascending order for the interpolation calculation.

```
SELECT department id "Department",
      APPROX PERCENTILE (0.25 DETERMINISTIC)
       WITHIN GROUP (ORDER BY salary ASC) "25th Percentile Salary",
      APPROX PERCENTILE (0.50 DETERMINISTIC)
       WITHIN GROUP (ORDER BY salary ASC) "50th Percentile Salary",
      APPROX PERCENTILE (0.75 DETERMINISTIC)
       WITHIN GROUP (ORDER BY salary ASC) "75th Percentile Salary"
 FROM employees
 GROUP BY department_id
 ORDER BY department id;
Department 25th Percentile Salary 50th Percentile Salary 75th Percentile Salary
 10
                         4400
                                             4400
                                                                  4400
      20
                                            6000
                        6000
                                                                13000
      30
                         2633
                                            2765
                                                                  3100
       40
                         6500
                                             6500
                                                                  6500
      50
                                             3100
                                                                 3599
                         2600
      60
                         4800
                                             4800
                                                                  6000
      70
                        10000
                                            10000
                                                                 10000
      80
                         7400
                                             9003
                                                                 10291
      90
                        17000
                                             17000
                                                                 24000
      100
                         7698
                                             7739
                                                                  8976
      110
                         8300
                                              8300
                                                                 12006
                         7000
                                              7000
                                                                  7000
```

The following query returns the error rates for the approximate 25th percentile salaries that were calculated in the previous query:

```
SELECT department id "Department",
      APPROX PERCENTILE (0.25 DETERMINISTIC, 'ERROR RATE')
        WITHIN GROUP (ORDER BY salary ASC) "Error Rate"
 FROM employees
 GROUP BY department id
 ORDER BY department id;
Department Error Rate
______
       10 .002718282
       20 .021746255
       30 .021746255
       40 .002718282
       50 .019027973
       60 .019027973
       70 .002718282
       80 .021746255
       90 .021746255
      100 .019027973
      110 .019027973
          .002718282
```

The following query returns the confidence levels for the error rates that were calculated in the previous guery:

```
SELECT department id "Department",
       APPROX_PERCENTILE(0.25 DETERMINISTIC, 'CONFIDENCE')
        WITHIN GROUP (ORDER BY salary ASC) "Confidence"
FROM employees
GROUP BY department id
ORDER BY department id;
Department Confidence
       10 .997281718
       20 .999660215
       30 .999660215
       40 .997281718
       50 .999611674
       60 .999611674
       70 .997281718
       80 .999660215
       90 .999660215
       100 .999611674
       110 .999611674
           .997281718
```

The following query returns the nondeterministic approximate 25th percentile, 50th percentile, and 75th percentile salaries for each department in the hr.employees table. The salaries are sorted in ascending order for the interpolation calculation.

```
SELECT department_id "Department",

APPROX_PERCENTILE(0.25)

WITHIN GROUP (ORDER BY salary ASC) "25th Percentile Salary",

APPROX_PERCENTILE(0.50)

WITHIN GROUP (ORDER BY salary ASC) "50th Percentile Salary",

APPROX_PERCENTILE(0.75)

WITHIN GROUP (ORDER BY salary ASC) "75th Percentile Salary"

FROM employees

GROUP BY department_id

ORDER BY department_id;
```

Department	25th	Percentile	Salary	50th	Percentile	Salary	75th	Percentile	Salary
10			4400			4400			4400
20			6000			6000			13000
30			2600			2800			3100
40			6500			6500			6500
50			2600			3100			3600
60			4800			4800			6000
70			10000			10000			10000
80			7300			8800			10000
90			17000			17000			24000
100			7700			7800			9000
110			8300			8300			12008
			7000			7000			7000

APPROX PERCENTILE AGG

Syntax



Purpose

APPROX_PERCENTILE_AGG takes as its input a column of details containing approximate percentile information, and enables you to perform aggregations of that information.

For detail, specify a column of details created by the APPROX_PERCENT_DETAIL function or the APPROX PERCENTILE AGG function. This column is of data type BLOB.

You can specify this function in a SELECT statement with a GROUP BY clause to aggregate the information contained in the details within each group of rows and return a single detail for each group.

This function returns a BLOB value, called a detail, which contains approximate percentile information in a special format. You can store details returned by this function in a table or materialized view, and then again use the APPROX_PERCENTILE_AGG function to further aggregate those details, or use the TO_APPROX_PERCENTILE function to convert the details to specified percentile values.

See Also:

- APPROX_PERCENTILE_DETAIL
- TO_APPROX_PERCENTILE

Examples

Refer to APPROX_PERCENTILE_AGG: Examples for examples of using the APPROX_PERCENTILE_AGG function in conjunction with the APPROX_PERCENTILE_DETAIL and TO APPROX PERCENTILE functions.

APPROX_PERCENTILE_DETAIL

Syntax



Purpose

APPROX_PERCENTILE_DETAIL calculates approximate percentile information for the values of expr and returns a BLOB value, called a detail, which contains that information in a special format.

The acceptable data types for expr depend on the algorithm that you specify with the DETERMINISTIC clause. Refer to the DETERMINISTIC clause for more information.

This function is commonly used with the GROUP BY clause in a SELECT statement. It calculates approximate percentile information for expr within each group of rows and returns a single detail for each group.

The details returned by APPROX_PERCENTILE_DETAIL can be used as input to the APPROX_PERCENTILE_AGG function, which enables you to perform aggregations of the details, or the TO_APPROX_PERCENTILE function, which converts a detail to a specified percentile value. You can use these three functions together to perform resource-intensive approximate percentile calculations once, store the resulting details, and then perform efficient aggregations and queries on those details. For example:

- Use the APPROX_PERCENTILE_DETAIL function to perform approximate percentile
 calculations and store the resulting details in a table or materialized view. These could be
 highly-granular percentile details, such as income percentile information for cities.
- 2. Use the APPROX_PERCENTILE_AGG function to aggregate the details obtained in the previous step and store the resulting details in a table or materialized view. These could be details of lower granularity, such as income percentile information for states.
- 3. Use the TO_APPROX_PERCENTILE function to convert the stored detail values to percentile values. You can use the TO_APPROX_PERCENTILE function to query detail values created by the APPROX_PERCENTILE_DETAIL function or the APPROX_PERCENTILE_AGG function.

DETERMINISTIC

This clause lets you control the type of algorithm used to calculate the approximate percentile values.

- If you specify DETERMINISTIC, then this function calculates deterministic approximate percentile information. In this case, <code>expr</code> must evaluate to a numeric value, or to a value that can be implicitly converted to a numeric value.
- If you omit DETERMINSTIC, then this function calculates nondeterministic approximate percentile information. In this case, <code>expr</code> must evaluate to a numeric or datetime value, or to a value that can be implicitly converted to a numeric or datetime value.



```
See Also:
```

- APPROX_PERCENTILE_AGG
- TO_APPROX_PERCENTILE

Examples

The examples in this section demonstrate how to use the APPROX_PERCENTILE_DETAIL, APPROX_PERCENTILE_AGG, and TO_APPROX_PERCENTILE functions together to perform resource-intensive approximate percentile calculations once, store the resulting details, and then perform efficient aggregations and queries on those details.

APPROX_PERCENTILE_DETAIL: Example

The following statement queries the tables <code>sh.customers</code> and <code>sh.sales</code> for the monetary amounts for products sold to each customer. The <code>APPROX_PERCENTILE_DETAIL</code> function returns the information in a detail, called <code>city_detail</code>, for each city in which customers reside. The returned details are stored in a materialized view called <code>amt_sold</code> by <code>city_mv</code>.

APPROX_PERCENTILE_AGG: Examples

The following statement uses the APPROX_PERCENTILE_AGG function to read the details stored in amt_sold_by_city_mv and create aggregated details that contain the monetary amounts for products sold to customers in each state. These aggregated details are stored in a materialized view called amt_sold_by_state_mv.

The following statement is similar to the previous statement, except it creates aggregated details that contain the approximate monetary amounts for products sold to customers in each country. These aggregated details are stored in a materialized view called

```
amt_sold_by_country_mv.

CREATE MATERIALIZED VIEW amt_sold_by_country_mv AS
    SELECT country,
          APPROX_PERCENTILE_AGG(city_detail) country_detail
    FROM amt_sold_by_city_mv
    GROUP BY country;
```

TO_APPROX_PERCENTILE: Examples



The following statement uses the TO_APPROX_PERCENTILE function to query the details stored in amt_sold_by_city_mv and return approximate 25th percentile, 50th percentile, and 75th percentile values for monetary amounts for products sold to customers in each city:

```
SELECT country,
state,
city,
TO_APPROX_PERCENTILE(city_detail, .25, 'NUMBER') "25th Percentile",
TO_APPROX_PERCENTILE(city_detail, .50, 'NUMBER') "50th Percentile",
TO_APPROX_PERCENTILE(city_detail, .75, 'NUMBER') "75th Percentile"

FROM amt_sold_by_city_mv
ORDER BY country, state, city;

COUNTRY STATE CITY 25th Percentile 50th Percentile 75th Percentile

52769 Kuala Lumpur Kuala Lumpur 19.29 38.1 53.84
52769 Penang Batu Ferringhi 21.51 42.09 57.26
52769 Penang Georgetown 19.15 33.25 56.12
52769 Selangor Klang 18.08 32.06 51.29
52769 Selangor Petaling Jaya 19.29 35.43 60.2
```

The following statement uses the TO_APPROX_PERCENTILE function to query the details stored in amt_sold_by_state_mv and return approximate 25th percentile, 50th percentile, and 75th percentile values for monetary amounts for products sold to customers in each state:

The following statement uses the TO_APPROX_PERCENTILE function to query the details stored in amt_sold_by_country_mv and return approximate 25th percentile, 50th percentile, and 75th percentile values for monetary amounts for products sold to customers in each country:

```
SELECT country,
      TO_APPROX_PERCENTILE(country_detail, .25, 'NUMBER') "25th Percentile",
      TO_APPROX_PERCENTILE(country_detail, .50, 'NUMBER') "50th Percentile",
      TO APPROX PERCENTILE (country detail, .75, 'NUMBER') "75th Percentile"
FROM amt sold by country mv
ORDER BY country;
 COUNTRY 25th Percentile 50th Percentile 75th Percentile
_____ ____
   52769
                              35.43
                 19.1
               19.29 38.99
11.99 44.99
18.08 33.72
15.67 29.61
   52770
   52771
                                           561.47
   52772
                                            54.16
   52773
                                            50.65
```



APPROX_PERCENTILE_AGG takes as its input a column of details containing approximate percentile information, and enables you to perform aggregations of that information. The following statement demonstrates how approximate percentile details can interpreted by APPROX_PERCENTILE_AGG to provide an input to the TO_APPROX_PERCENTILE function. Like the previous example, this query returns approximate 25th percentile values for monetary amounts for products sold to customers in each country. Note that the results are identical to those returned for the 25th percentile in the previous example.

```
SELECT country,

TO_APPROX_PERCENTILE (APPROX_PERCENTILE_AGG(city_detail), .25, 'NUMBER') "25th
Percentile"

FROM amt_sold_by_city_mv

GROUP BY country

ORDER BY country;

COUNTRY 25th Percentile

52769 19.1
52770 19.29
52771 11.99
52772 18.08
52773 15.67
```

Query Rewrite and Materialized Views Based on Approximate Queries: Example

In APPROX_PERCENTILE_DETAIL: Example, the ENABLE QUERY REWRITE clause is specified when creating the materialized view amt_sold_by_city_mv. This enables queries that contain approximation functions, such as APPROX_MEDIAN or APPROX_PERCENTILE, to be rewritten using the materialized view.

For example, ensure that query rewrite is enabled at either the database level or for the current session, and run the following query:

```
SELECT c.country_id country,
          APPROX_MEDIAN(s.amount_sold) amount_median
FROM customers c, sales s
WHERE c.cust_id = s.cust_id
GROUP BY c.country_id;
```

Explain the plan by querying DBMS XPLAN:

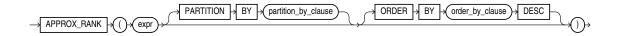
```
SET LINESIZE 300
SET PAGESIZE 0
COLUMN plan_table_output FORMAT A150
SELECT * FROM TABLE(DBMS XPLAN.DISPLAY CURSOR(format=>'BASIC'));
```

As shown in the following plan, the optimizer used the materialized view amt_sold_by_city_mv for the query:



APPROX_RANK

Syntax



Purpose

APPROX RANK returns the approximate value in a group of values.

This function takes an optional PARTITION BY clause followed by a mandatory ORDER BY ... DESC clause. The PARTITION BY key must be a subset of the GROUP BY key. The ORDER BY clause must include either APPROX COUNT or APPROX SUM.

Examples

The query returns the jobs that are among the top 10 total salary per department. For each job, the total salary and ranking is also given.

```
SELECT job_id,

APPROX_SUM(sal),

APPROX_RANK(PARTITION BY department_id ORDER BY APPROX_SUM(salary)

DESC)

FROM employees

GROUP BY department_id, job_id

HAVING

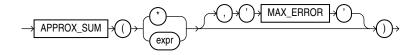
APPROX_RANK(
PARTITION BY department_id

ORDER BY APPROX_SUM (salary)

DESC) <= 10;
```

APPROX SUM

Syntax



Purpose

APPROX_SUM returns the approximate sum of an expression. If you supply MAX_ERROR as the second argument, then the function returns the maximum error between the actual and approximate sum.

You must use this function with a corresponding APPROX_RANK function in the HAVING clause. If a query uses APPROX_COUNT, APPROX_SUM, or APPROX_RANK, then the query must not use any other aggregation functions.

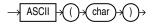
Note that APPROX SUM returns an error when the input is a negative number.

Examples

The following query returns the 10 job types within every department that have the highest aggregate salary:

ASCII

Syntax



Purpose

ASCII returns the decimal representation in the database character set of the first character of char.

char can be of data type CHAR, VARCHAR2, NCHAR, or NVARCHAR2. The value returned is of data type NUMBER. If your database character set is 7-bit ASCII, then this function returns an ASCII value. If your database character set is EBCDIC Code, then this function returns an EBCDIC value. There is no corresponding EBCDIC character function.

This function does not support CLOB data directly. However, CLOBs can be passed in as arguments through implicit data conversion.



Data Type Comparison Rules for more information

Examples

The following example returns employees whose last names begin with the letter L, whose ASCII equivalent is 76:

```
SELECT last_name
  FROM employees
  WHERE ASCII(SUBSTR(last_name, 1, 1)) = 76
```



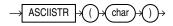
ORDER BY last_name;

LAST_NAME

Ladwig
Landry
Lee
Livingston
Lorentz

ASCIISTR

Syntax



Purpose

ASCIISTR takes as its argument a string, or an expression that resolves to a string, in any character set and returns an ASCII version of the string in the database character set. Non-ASCII characters are converted to the form $\xspace \xspace \xspace \xspace \xspace$ where $\xspace \xspace \xspace \xspace \xspace$ and $\xspace \xspace \xspace \xspace \xspace \xspace \xspace$



- Oracle Database Globalization Support Guide for information on Unicode character sets and character semantics
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of ASCIISTR

Examples

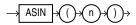
The following example returns the ASCII string equivalent of the text string "ABÄCDE":

```
SELECT ASCIISTR('ABÄCDE')
FROM DUAL;

ASCIISTR('
-----
AB\00C4CDE
```

ASIN

Syntax





Purpose

ASIN returns the arc sine of n. The argument n must be in the range of -1 to 1, and the function returns a value in the range of -pi/2 to pi/2, expressed in radians.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



Table 2-9 for more information on implicit conversion

Examples

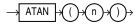
The following example returns the arc sine of .3:

```
SELECT ASIN(.3) "Arc_Sine"
FROM DUAL;

Arc_Sine
-----
.304692654
```

ATAN

Syntax



Purpose

ATAN returns the arc tangent of n. The argument n can be in an unbounded range and returns a value in the range of -pi/2 to pi/2, expressed in radians.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is $\verb|BINARY_FLOAT|$, then the function returns $\verb|BINARY_DOUBLE|$. Otherwise the function returns the same numeric data type as the argument.



ATAN2 for information about the $\mathtt{ATAN2}$ function and Table 2-9 for more information on implicit conversion

Examples

The following example returns the arc tangent of .3:

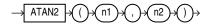


```
SELECT ATAN(.3) "Arc_Tangent"
FROM DUAL;

Arc_Tangent
-----
.291456794
```

ATAN2

Syntax



Purpose

ATAN2 returns the arc tangent of n1 and n2. The argument n1 can be in an unbounded range and returns a value in the range of -pi to pi, depending on the signs of n1 and n2, expressed in radians.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If any argument is <code>BINARY_FLOAT</code> or <code>BINARY_DOUBLE</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns <code>NUMBER</code>.



See Also:

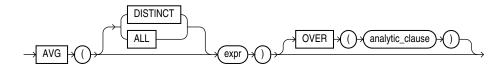
ATAN for information on the ATAN function and Table 2-9 for more information on implicit conversion

Examples

The following example returns the arc tangent of .3 and .2:

AVG

Syntax





See Also:

Analytic Functions for information on syntax, semantics, and restrictions

Purpose

AVG returns average value of expr.

It takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type or an interval data type.

The function returns the same data type as the numeric data type of the argument. If the input is an interval, this returns an interval with the same units as the input.

See Also:

Table 2-9 for more information on implicit conversion

If you specify DISTINCT, then you can specify only the <code>query_partition_clause</code> of the <code>analytic_clause</code>. The <code>order_by_clause</code> and <code>windowing_clause</code> are not allowed.

See Also:

About SQL Expressions for information on valid forms of expr and Aggregate Functions

Vector Aggregate Operations

You can use AVG on vectors to return the average on non-null inputs.

expr must evaluate to VECTOR and must not be BINARY vectors. The returned vector has the same number of dimensions as the input, and the format is always FLOAT64. For flexible number of dimensions, all inputs must have the same number of dimensions within each aggregation group.

NULL vectors are ignored. They are not counted when calculating the average vector. If all inputs within an aggregation group are NULL, the result is NULL for that group. If a certain dimension overflows when applying arithmetic operations, an error is raised.

Rules

- DISTINCT syntax is not allowed.
- Only GROUP BY and GROUP BY ROLLUP are supported.
- Analytic functions are not supported for input arguments of type VECTOR.

See Arithmetic Operatorsof the AI Vector Search User's Guide for examples.



Aggregate Example

The following example calculates the average salary of all employees in the hr.employees table:

Analytic Example

The following example calculates, for each employee in the employees table, the average salary of the employees reporting to the same manager who were hired in the range just before through just after the employee:

```
SELECT manager_id, last_name, hire_date, salary,

AVG(salary) OVER (PARTITION BY manager_id ORDER BY hire_date

ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING) AS c_mavg

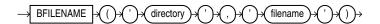
FROM employees

ORDER BY manager id, hire date, salary;
```

MANAGER_ID	LAST_NAME	HIRE_DATE	SALARY	C_MAVG
100	De Haan	13-JAN-01	17000	14000
100	Raphaely	07-DEC-02	11000	11966.6667
100	Kaufling	01-MAY-03	7900	10633.3333
100	Hartstein	17-FEB-04	13000	9633.33333
100	Weiss	18-JUL-04	8000	11666.6667
100	Russell	01-OCT-04	14000	11833.3333
100	Partners	05-JAN-05	13500	13166.6667
100	Errazuriz	10-MAR-05	12000	11233.3333

BFILENAME

Syntax



Purpose

BFILENAME returns a BFILE locator that is associated with a physical LOB binary file on the server file system.

- 'directory' is a database object that serves as an alias for a full path name on the server file system where the files are actually located.
- 'filename' is the name of the file in the server file system.

You must create the directory object and associate a BFILE value with a physical file before you can use them as arguments to BFILENAME in a SQL or PL/SQL statement, DBMS_LOB package, or OCI operation.

You can use this function in two ways:

- In a DML statement to initialize a BFILE column
- In a programmatic interface to access BFILE data by assigning a value to the BFILE locator

The directory argument is case sensitive. You must ensure that you specify the directory object name exactly as it exists in the data dictionary. For example, if an "Admin" directory object was created using mixed case and a quoted identifier in the CREATE DIRECTORY statement, then when using the BFILENAME function you must refer to the directory object as 'Admin'. You must specify the filename argument according to the case and punctuation conventions for your operating system.

See Also:

- Oracle Database SecureFiles and Large Objects Developer's Guide and Oracle Call Interface Developer's Guide for more information on LOBs and for examples of retrieving BFILE data
- CREATE DIRECTORY

Examples

The following example inserts a row into the sample table <code>pm.print_media</code>. The example uses the <code>BFILENAME</code> function to identify a binary file on the server file system in the directory <code>/demo/schema/product_media</code>. The example shows how the directory database object <code>media_dir</code> was created in the <code>pm</code> schema.

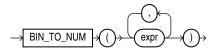
```
CREATE DIRECTORY media_dir AS '/demo/schema/product_media';

INSERT INTO print_media (product_id, ad_id, ad_graphic)

VALUES (3000, 31001, BFILENAME('MEDIA_DIR', 'modem_comp_ad.gif'));
```

BIN_TO_NUM

Syntax



Purpose

BIN_TO_NUM converts a bit vector to its equivalent number. Each argument to this function represents a bit in the bit vector. This function takes as arguments any numeric data type, or any nonnumeric data type that can be implicitly converted to NUMBER. Each expr must evaluate to 0 or 1. This function returns Oracle NUMBER.

BIN_TO_NUM is useful in data warehousing applications for selecting groups of interest from a materialized view using grouping sets.



See Also:

- group_by_clause for information on GROUPING SETS syntax
- Table 2-9 for more information on implicit conversion
- Oracle Database Data Warehousing Guide for information on data aggregation in general

Examples

The following example converts a binary value to a number:

The next example converts three values into a single binary value and uses BIN_TO_NUM to convert that binary into a number. The example uses a PL/SQL declaration to specify the original values. These would normally be derived from actual data sources.

```
SELECT order_status
  FROM orders
  WHERE order_id = 2441;
ORDER STATUS
DECLARE
  warehouse NUMBER := 1;
  ground NUMBER := 1;
  insured NUMBER := 1;
 result NUMBER;
  SELECT BIN TO NUM(warehouse, ground, insured) INTO result FROM DUAL;
  UPDATE orders SET order status = result WHERE order id = 2441;
END;
PL/SQL procedure successfully completed.
SELECT order status
  FROM orders
  WHERE order id = 2441;
ORDER STATUS
-----
```

Refer to the examples for BITAND for information on reversing this process by extracting multiple values from a single column value.

BITAND

Syntax



Purpose

The BITAND function treats its inputs and its output as vectors of bits; the output is the bitwise AND of the inputs.

The types of expr1 and expr2 are NUMBER, and the result is of type NUMBER. If either argument to BITAND is NULL, the result is NULL.

The arguments must be in the range $-(2^{(n-1)})$.. $((2^{(n-1)})-1)$. If an argument is out of this range, the result is undefined.

The result is computed in several steps. First, each argument A is replaced with the value SIGN(A) *FLOOR(ABS(A)). This conversion has the effect of truncating each argument towards zero. Next, each argument A (which must now be an integer value) is converted to an n-bit two's complement binary integer value. The two bit values are combined using a bitwise AND operation. Finally, the resulting n-bit two's complement value is converted back to NUMBER.

Notes on the BITAND Function

- The current implementation of BITAND defines n = 128.
- PL/SQL supports an overload of BITAND for which the types of the inputs and of the result are all BINARY INTEGER and for which n = 32.

Examples

The following example performs an AND operation on the numbers 6 (binary 1,1,0) and 3 (binary 0,1,1):

```
SELECT BITAND(6,3)
FROM DUAL;
BITAND(6,3)
```

This is the same as the following example, which shows the binary values of 6 and 3. The BITAND function operates only on the significant digits of the binary values:

Refer to the example for BIN_TO_NUM for information on encoding multiple values in a single column value.



The following example supposes that the <code>order_status</code> column of the sample table <code>oe.orders</code> encodes several choices as individual bits within a single numeric value. For example, an order still in the warehouse is represented by a binary value 001 (decimal 1). An order being sent by ground transportation is represented by a binary value 010 (decimal 2). An insured package is represented by a binary value 100 (decimal 4). The example uses the <code>DECODE</code> function to provide two values for each of the three bits in the <code>order_status</code> value, one value if the bit is turned on and one if it is turned off.

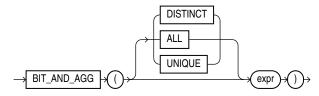
```
SELECT order id, customer id, order status,
   DECODE (BITAND (order status, 1), 1, 'Warehouse', 'PostOffice') "Location",
   DECODE(BITAND(order_status, 2), 2, 'Ground', 'Air') "Method",
   DECODE(BITAND(order_status, 4), 4, 'Insured', 'Certified') "Receipt"
 FROM orders
 WHERE sales rep_id = 160
 ORDER BY order id;
 ORDER ID CUSTOMER ID ORDER STATUS Location Method Receipt
2416
               104
                             6 PostOffice Ground Insured
     2419
               107
                             3 Warehouse Ground Certified
     2420
                108
                             2 PostOffice Ground Certified
     2423
                145
                             3 Warehouse Ground Certified
     2441
                106
                             5 Warehouse Air Insured
     2455
                145
                             7 Warehouse Ground Insured
```

For the Location column, BITAND first compares <code>order_status</code> with 1 (binary 001). Only significant bit values are compared, so any binary value with a 1 in its rightmost bit (any odd number) will evaluate positively and return 1. Even numbers will return 0. The <code>DECODE</code> function compares the value returned by <code>BITAND</code> with 1. If they are both 1, then the location is "Warehouse". If they are different, then the location is "PostOffice".

The Method and Receipt columns are calculated similarly. For Method, BITAND performs the AND operation on order_status and 2 (binary 010). For Receipt, BITAND performs the AND operation on order status and 4 (binary 100).

BIT AND AGG

Syntax



Purpose

 ${\tt BIT_AND_AGG}$ is a bitwise aggregation function that returns the result of a bitwise AND operation.

You can use ${\tt BIT_AND_AGG}$ as part of a <code>GROUP BY</code> query, window function, or as an analytical function. The return type of ${\tt BIT_AND_AGG}$ is always a number.

Semantics

The keywords DISTINCT or UNIQUE ensure that only unique values in expr are used for computation. UNIQUE is an Oracle-specific keyword and not an ANSI standard.

NULL values in the expr column are ignored.

Returns NULL if all rows in the group have NULL expr values.

Floating point values are truncated to the integer prior to aggregation. For instance, the value 4.64 is converted to 4, and the value 4.4 is also converted to 4.

Negative numbers are represented in two's complement form internally prior to performing an aggregate operation. The resultant aggregate could be a negative value.

Range of inputs supported: -2 raised to 127 to (2 raised to 127) -1

Numbers are internally converted to a 128b decimal representation prior to aggregation. The resultant aggregate is converted back into an Oracle Number.

For a given set of values, the result of a bitwise aggregate is always deterministic and independent of ordering.

Example 7-2 Use the BIT_AND_AGG Function

Select two numbers and their bitwise representation:

Perform the bitwise AND operation:

Only the first bit is identical in both rows, thus the result is 001, which is the number 1.

BITMAP BIT POSITION

Syntax



Purpose

Use BITMAP_BIT_POSITION to construct the one-to-one mapping between a number and a bit position.

The argument *expr* is of type NUMBER. It is the absolute bit position in the bitmap.

BITMAP BIT POSITION returns a NUMBER, the relative bit position.

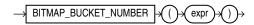
If expr is NULL, the function returns NULL.

If *expr* is not an integer, you will see the following error message:

Invalid value has been passed to a BITMAP COUNT DISTINCT related operator.

BITMAP BUCKET NUMBER

Syntax



Purpose

Use BITMAP_BUCKET_NUMBER to construct a one-to-one mapping between a number and a bit position in a bitmap.

The argument *expr* is of type NUMBER. It represents the absolute bit position in the bitmap.

BITMAP BUCKET NUMBER returns a number. It represents the relative bit position.

If expr is NULL, the function returns NULL.

If *expr* is not an integer, you will see the following error message:

Invalid value has been passed to a BITMAP COUNT DISTINCT related operator.

BITMAP_CONSTRUCT_AGG

Syntax



Purpose

BITMAP_CONSTRUCT_AGG is an aggregation function that operates on bit positions and returns the bitmap representation of the set of all input bit positions. It essentially maintains a bitmap and sets into it all the input bit positions. It returns the representation of the bitmap.

The argument expr is of type NUMBER.

The return type is of type BLOB.

If expr is NULL, the function returns NULL.

Restrictions



• The argument must be of NUMBER type. If the input value cannot be converted to a natural number, error ORA-62575 is raised:

62575, 00000, "Invalid value has been passed to a BITMAP COUNT DISTINCT related operator."

// *Cause: An attempt was made to pass an invalid value to a BITMAP COUNT DISTINCT operator.

// *Action: Pass only natural number values to BITMAP CONSTRUCT AGG.

• If the bitmap exceeds the maximum value of a BLOB, you will see error ORA-62577:

62577, 00000, "The bitmap size exceeds maximum size of its SQL data type." // *Cause: An attempt was made to construct a bitmap larger than its maximum SQL type size.

// *Action: Break the input to BITMAP_CONSTRUCT_AGG into smaller ranges.

BITMAP COUNT

Syntax



Purpose

BITMAP COUNT is a scalar function that returns the 1-bit count for the input bitmap.

The argument expr is of type BLOB.

It returns a NUMBER representing the count of bits set in its input.

If expr is NULL, it returns 0.

Restrictions

The argument must be of type BLOBtype. The argument is expected to be a bitmap produced by BITMAP_CONSTRUCT_AGG or, recursively, by BITMAP_OR_AGG. Any other input results in ORA-62578:

62578, 00000, "The input is not a valid bitmap produced by BITMAP COUNT DISTINCT related operators."

// *Cause: An attempt was made to pass a bitmap that was not produced by one of the BITMAP COUNT DISTINCT operators.

// *Action: Only pass bitmaps constructed via BITMAP_CONSTRUCT_AGG or BITMAP OR AGG to BITMAP COUNT DISTINCT related operators.

BITMAP OR AGG

Syntax





Purpose

 ${\tt BITMAP_OR_AGG}$ is an aggregation function that operates on bitmaps and computes the OR of its inputs.

The argument expr must be of type BLOB.

The return type is of type BLOB. It returns the bitmap representing the OR of all the bitmaps it has aggregated.

The output of BITMAP_OR_AGG is not human-readable. It is meant to be processed by further aggregations via BITMAP_OR_AGG or by the scalar function BITMAP_COUNT.

If expr is NULL, the function returns NULL.

Restrictions

The argument must be of type BLOB. The argument is expected to be a bitmap produced by BITMAP_CONSTRUCT_AGG or, recursively, by BITMAP_OR_AGG. Any other input results in ORA-62578:

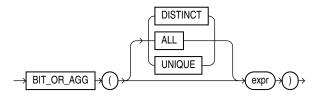
```
62578, 00000, "The input is not a valid bitmap produced by BITMAP COUNT DISTINCT related operators."
```

// *Cause: An attempt was made to pass a bitmap that was not produced by one of the BITMAP COUNT DISTINCT operators.

// *Action: Only pass bitmaps constructed via BITMAP_CONSTRUCT_AGG or BITMAP OR AGG to BITMAP COUNT DISTINCT related operators.

BIT_OR_AGG

Syntax



Purpose

BIT OR AGG is a bitwise aggregation function that returns the result of a bitwise OR operation.

You can use BIT_OR_AGG as part of a GROUP BY query, window function, or as an analytical function. The return type of BIT_OR_AGG is always a number.

Semantics

The keywords DISTINCT or UNIQUE ensure that only unique values in expr are used for computation. UNIQUE is an Oracle-specific keyword and not an ANSI standard.

NULL values in the expr column are ignored.

Returns NULL if all rows in the group have NULL expr values.

Floating point values are truncated to the integer prior to aggregation. For instance, the value 4.64 is converted to 4 and the value 4.4 is also converted to 4.



Negative numbers are represented in two's complement form internally prior to performing an aggregate operation. The resultant aggregate could be a negative value.

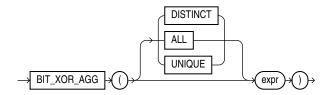
Range of inputs supported: -2 raised to 127 to (2 raised to 127) -1

Numbers are internally converted to a 128b decimal representation prior to aggregation. The resultant aggregate is converted back into an Oracle Number.

For a given set of values, the result of a bitwise aggregate is always deterministic and independent of ordering.

BIT_XOR_AGG

Syntax



Purpose

BIT_XOR_AGG is a bitwise aggregation function that returns the result of a bitwise XOR operation.

You can use $\texttt{BIT_XOR_AGG}$ as part of a <code>GROUP BY</code> query, window function, or as an analytical function. The return type of BIT XOR AGG is always a number.

Semantics

The keywords DISTINCT or UNIQUE ensure that only unique values in expr are used for computation. BIT_XOR_AGG could potentially return a different value when DISTINCT is present. UNIQUE is an Oracle-specific keyword and not an ANSI standard.

NULL values in the expr column are ignored.

Returns NULL if all rows in the group have NULL expr values.

Floating point values are truncated to the integer prior to aggregation. For instance, the value 4.64 is converted to 4 and the value 4.4 is also converted to 4.

Negative numbers are represented in two's complement form internally prior to performing an aggregate operation. The resultant aggregate could be a negative value.

Range of inputs supported: -2 raised to 127 to (2 raised to 127) -1

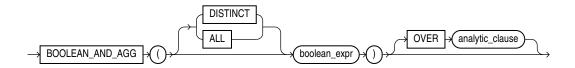
Numbers are internally converted to a 128b decimal representation prior to aggregation. The resultant aggregate is converted back into an Oracle Number.

For a given set of values, the result of a bitwise aggregate is always deterministic and independent of ordering.



BOOLEAN_AND_AGG

Syntax



Purpose

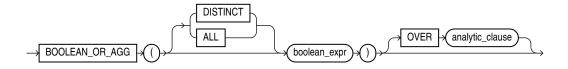
BOOLEAN_AND_AGG returns 'TRUE' if the boolean_expr evaluates to true for every row that qualifies. Otherwise it returns 'FALSE'. You can use it as an aggregate or analytic function.

Examples

```
SELECT BOOLEAN AND AGG(c2)
   FROM t;
SELECT BOOLEAN_AND_AGG(c2)
   FROM t
   WHERE c1 = 0;
SELECT BOOLEAN_AND_AGG(c2)
   FROM t
   WHERE c2 IS FALSE;
SELECT BOOLEAN_AND_AGG(c2)
   FROM t
   WHERE c2 IS FALSE OR c2 IS NULL;
SELECT BOOLEAN_AND_AGG(c2)
   FROM t
   WHERE c2 IS NOT TRUE OR c2 IS NULL;
SELECT BOOLEAN_AND_AGG(c2)
   FROM t
   WHERE c2 IS NOT FALSE OR c2 IS NULL;
SELECT BOOLEAN AND AGG(c2 OR c2 OR (c2))
   WHERE c2 IS NOT FALSE OR c2 IS NULL;
```

BOOLEAN_OR_AGG

Syntax



Purpose

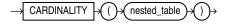
BOOLEAN_OR_AGG returns 'TRUE' if the boolean_expr evaluates to true for at least one row that qualifies. Otherwise it returns 'FALSE'. You can use it as an aggregate or analytic function.

Examples

```
SELECT BOOLEAN OR AGG(c2)
   FROM t;
SELECT BOOLEAN OR AGG(c2)
   FROM t
   WHERE c1 = 0;
SELECT BOOLEAN_OR_AGG(c2)
   FROM t
   WHERE c2 IS TRUE;
SELECT BOOLEAN OR AGG(c2)
   WHERE c2 IS TRUE OR c2 IS NULL;
SELECT BOOLEAN OR AGG(c2)
   FROM t
   WHERE c2 IS NOT FALSE OR c2 IS NULL;
SELECT BOOLEAN OR AGG(c2)
   WHERE c2 IS NOT TRUE OR c2 IS NULL;
SELECT BOOLEAN OR AGG(c2 OR c2)
   WHERE c2 IS NOT TRUE OR c2 IS NULL;
```

CARDINALITY

Syntax



Purpose

CARDINALITY returns the number of elements in a nested table. The return type is NUMBER. If the nested table is empty, or is a null collection, then CARDINALITY returns NULL.

Examples

The following example shows the number of elements in the nested table column ad textdocs ntab of the sample table pm.print media:

```
SELECT product_id, CARDINALITY(ad_textdocs_ntab) cardinality
FROM print_media
ORDER BY product_id;

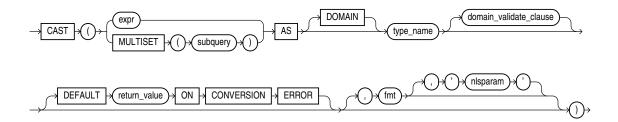
PRODUCT_ID CARDINALITY
```



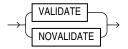
2056	3
2268	3
3060	3
3106	3

CAST

Syntax



domain_validate_clause::=



Purpose

CAST lets you convert built-in data types or collection-typed values of one type into another built-in data type or collection type. You can cast an unnamed operand (such as a date or the result set of a subquery) or a named collection (such as a varray or a nested table) into a type-compatible data type or named collection. The $type_name$ must be the name of a built-in data type, collection type, or domain name and the operand must be a built-in data type or must evaluate to a collection value.

For the operand, <code>expr</code> can be either a built-in data type, a collection type, or an instance of an <code>ANYDATA</code> type. If <code>expr</code> is an instance of an <code>ANYDATA</code> type, then <code>CAST</code> tries to extract the value of the <code>ANYDATA</code> instance and return it if it matches the cast target type, otherwise, null will be returned. <code>MULTISET</code> informs Oracle Database to take the result set of the subquery and return a collection value. Table 7-1 shows which built-in data types can be cast into which other built-in data types. (<code>CAST</code> does not support <code>LONG</code>, <code>LONG</code> <code>RAW</code>, or the Oracle-supplied types.)

CAST does not directly support any of the LOB data types. When you use CAST to convert a CLOB value into a character data type or a BLOB value into the RAW data type, the database implicitly converts the LOB value to character or raw data and then explicitly casts the resulting value into the target data type. If the resulting value is larger than the target type, then the database returns an error.

When you use CAST ... MULTISET to get a collection value, each select list item in the query passed to the CAST function is converted to the corresponding attribute type of the target collection element type.

The cells with an 'X' indicate the possible conversions from source to destination data type using CAST.



Table 7-1 Casting Built-In Data Types

Destination Data Type	from BINARY_F LOAT, BINARY_D OUBLE	from CHAR, VARCHAR 2	from NUMBER/ INTEGER	from DATETIME / INTERVAL (Note 1)	from RAW	from ROWID, UROWID (Note 2)	from NCHAR, NVARCHA R2	from BOOLEAN
to BINARY_FLO AT, BINARY_DO UBLE	X (Note 3)	X (Note 3)	X (Note 3)				X (Note 3)	X(Note 4)
to CHAR, VARCHAR2	X	X	X	X	X	X	X	X (Note 5)
to NUMBER/ INTEGER	X (Note 3)	X (Note 3)	X (Note 3)				X (Note 3)	X(Note 4)
to DATETIME/ INTERVAL		X (Note 3)		X (Note 3)				
to RAW		X			X		X	-
to ROWID, UROWID		X				X		-
to NCHAR, NVARCHAR2	X		X	X	X	X	X	X(Note 5)
to BOOLEAN	X(Note 4)	X(Note 6)	x (Note 4)				X(Note 6)	- X

Note 1: Datetime/interval includes DATE, TIMESTAMP, TIMESTAMP WITH TIMEZONE, TIMESTAMP WITH LOCAL TIME ZONE, INTERVAL DAY TO SECOND, and INTERVAL YEAR TO MONTH.

Note 2: You cannot cast a UROWID to a ROWID if the UROWID contains the value of a ROWID of an index-organized table.

Note 3: You can specify the DEFAULT return_value ON CONVERSION ERROR clause for this type of conversion. You can specify the fmt and nlsparam clauses for this type of conversion with the following exceptions: you cannot specify fmt when converting to INTERVAL DAY TO SECOND, and you cannot specify fmt or nlsparam when converting to INTERVAL YEAR TO MONTH.

Note 4: Casting Between Boolean and Numeric

When casting BOOLEAN to numeric:

- If the boolean value is true, then resulting value is 1.
- If the boolean value is false, then resulting value is 0.

When casting numeric to BOOLEAN:

- If the numeric value is non-zero (e.g., 1, 2, -3, 1.2), then resulting value is true.
- If the numeric value is zero, then resulting value is false.

Note 5: Casting Between Boolean and Char(n), NCHAR(n)

When casting BOOLEAN to VARCHAR (n), NVARCHAR (n)

- If the boolean value is true and n is not less than 4, then resulting value is true.
- If the boolean value is false and n is not less than 5, then resulting value is false.
- Otherwise, a data exception error is raised.



Note 6: Casting Character Strings to Boolean

When casting a character string to boolean, you must trim both leading and trailing spaces of the character string first. If the resulting character string is one of the accepted literals used to determine a valid boolean value, then the result is that valid boolean value.

If you want to cast a named collection type into another named collection type, then the elements of both collections must be of the same type.

See Also:

- Boolean Data Type
- Implicit Data Conversion for information on how Oracle Database implicitly converts collection type data into character data and Security Considerations for Data Conversion
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of CAST when it is a character value

MULTISET

If the result set of <code>subquery</code> can evaluate to multiple rows, then you must specify the <code>MULTISET</code> keyword. The rows resulting from the subquery form the elements of the collection value into which they are cast. Without the <code>MULTISET</code> keyword, the subquery is treated as a scalar subquery.

Restriction on MULTISET

If you specify the MULTISET keyword, then you cannot specify the DEFAULT return_value ON CONVERSION ERROR, fmt, or nlsparam clauses.

DOMAIN

The DOMAIN clause specifies that type_name is a domain. type_name must be the name of a domain that the user has execute privileges on.

domain_validate_clause

This clause is only valid when casting to domain types. It controls whether domain constraints are applied when converting <code>expr</code> to <code>type_name</code>. If <code>type_name</code> is a domain with constraints, and <code>domain_validate_clause</code> is not specified, enabled constraints will be applied to <code>expr</code>. Any disabled constraints are ignored.

VALIDATE

All the domain constraints are applied to *expr*, regardless of their state.

NOVALIDATE

None of the domain constraints are applied to *expr*, regardless of their state.

DEFAULT return_value ON CONVERSION ERROR

This clause allows you to specify the value returned by this function if an error occurs while converting <code>expr</code> to <code>type_name</code>. This clause has no effect if an error occurs while evaluating <code>expr</code>.



This clause is valid if expr evaluates to a character string of type CHAR, VARCHAR2, NCHAR, or NVARCHAR2, and type_name is binary_double, binary_float, date, interval day to second, interval year to month, number, timestamp, timestamp with time zone, or timestamp with local time zone.

The return_value can be a string literal, null, constant expression, or a bind variable, and must evaluate to null or a character string of type CHAR, VARCHAR2, NCHAR, or NVARCHAR2. If return_value cannot be converted to type_name, then the function returns an error.

fmt and nlsparam

The fmt argument lets you specify a format model and the nlsparam argument lets you specify NLS parameters. If you specify these arguments, then they are applied when converting expr and return value, if specified, to type name.

You can specify fmt and nlsparam if type name is one of the following data types:

• BINARY DOUBLE

If you specify BINARY_DOUBLE, then the optional fmt and nlsparam arguments serve the same purpose as for the TO_BINARY_DOUBLE function. Refer to TO_BINARY_DOUBLE for more information.

BINARY FLOAT

If you specify BINARY_FLOAT, then the optional fmt and nlsparam arguments serve the same purpose as for the TO_BINARY_FLOAT function. Refer to TO_BINARY_FLOAT for more information.

DATE

If you specify DATE, then the optional fmt and nlsparam arguments serve the same purpose as for the TO DATE function. Refer to TO_DATE for more information.

• NUMBER

If you specify NUMBER, then the optional *fmt* and *nlsparam* arguments serve the same purpose as for the TO NUMBER function. Refer to TO_NUMBER for more information.

• TIMESTAMP

If you specify <code>TIMESTAMP</code>, then the optional fmt and nlsparam arguments serve the same purpose as for the <code>TO_TIMESTAMP</code> function. If you omit fmt, then expr must be in the default format of the <code>TIMESTAMP</code> data type, which is determined explicitly by the <code>NLS_TIMESTAMP_FORMAT</code> parameter or implicitly by the <code>NLS_TERRITORY</code> parameter. Refer to <code>TO_TIMESTAMP</code> for more information.

• TIMESTAMP WITH TIME ZONE

If you specify TIMESTAMP WITH TIME ZONE, then the optional fmt and nlsparam arguments serve the same purpose as for the TO_TIMESTAMP_TZ function. If you omit fmt, then expr must be in the default format of the TIMESTAMP WITH TIME ZONE data type, which is determined explicitly by the NLS_TIMESTAMP_TZ_FORMAT parameter or implicitly by the NLS_TIMESTAMP_TZ_FORMAT parameter or implicitly by the NLS_TIMESTAMP_TZ for more information.

• TIMESTAMP WITH LOCAL TIME ZONE

If you specify TIMESTAMP WITH LOCAL TIME ZONE then the optional fmt and nlsparam arguments serve the same purpose as for the TO_TIMESTAMP function. If you omit fmt, then expr must be in the default format of the TIMESTAMP data type, , which is determined explicitly by the NLS_TIMESTAMP_FORMAT parameter or implicitly by the NLS_TERRITORY parameter. Refer to TO_TIMESTAMP for more information.



Built-In Data Type Examples

The following examples use the CAST function with scalar data types. The first example converts text to a timestamp value by applying the format model provided in the session parameter NLS_TIMESTAMP_FORMAT. If you want to avoid dependency on this NLS parameter, then you can use the TO_DATE as shown in the second example.

```
SELECT CAST('22-OCT-1997'

AS TIMESTAMP WITH LOCAL TIME ZONE)

FROM DUAL;

SELECT CAST(TO_DATE('22-OCT-1997', 'DD-Mon-YYYY')

AS TIMESTAMP WITH LOCAL TIME ZONE)

FROM DUAL;
```

In the preceding example, TO_DATE converts from text to DATE, and CAST converts from DATE to TIMESTAMP WITH LOCAL TIME ZONE, interpreting the date in the session time zone (SESSIONTIMEZONE).

```
SELECT product_id, CAST(ad_sourcetext AS VARCHAR2(30)) text
FROM print_media
   ORDER BY product id;
```

The following examples return a default value if an error occurs while converting the specified value to the specified data type. In these examples, the conversions occurs without error.

```
SELECT CAST(200

AS NUMBER
DEFAULT 0 ON CONVERSION ERROR)

FROM DUAL;

SELECT CAST('January 15, 1989, 11:00 A.M.'
AS DATE
DEFAULT NULL ON CONVERSION ERROR,
'Month dd, YYYY, HH:MI A.M.')

FROM DUAL;

SELECT CAST('1999-12-01 11:00:00 -8:00'
AS TIMESTAMP WITH TIME ZONE
DEFAULT '2000-01-01 01:00:00 -8:00' ON CONVERSION ERROR,
'YYYY-MM-DD HH:MI:SS TZH:TZM',
'NLS_DATE_LANGUAGE = American')
FROM DUAL;
```

In the following example, an error occurs while converting 'N/A' to a NUMBER value. Therefore, the CAST function returns the default value of 0.

```
SELECT CAST('N/A'
AS NUMBER
DEFAULT '0' ON CONVERSION ERROR)
FROM DUAL;
```

The following example converts data types VARCHAR2, NUMBER as BOOLEAN:

```
SELECT
CAST ( 'yes' AS BOOLEAN ),
CAST ( true AS NUMBER ),
CAST ( false AS VARCHAR2(10) );

CAST('YES'ASBOOLEAN) CAST(TRUEASNUMBER) CAST(FALSE
```



```
TRUE
                   1 FALSE
```

Collection Examples

The CAST examples that follow build on the cust address typ found in the sample order entry schema, oe.

```
CREATE TYPE address_book_t AS TABLE OF cust_address_typ;
CREATE TYPE address array t AS VARRAY(3) OF cust address typ;
CREATE TABLE cust_address (
  custno
         NUMBER,
  street address VARCHAR2(40),
 postal_code VARCHAR2(10), city VARCHAR2(30),
 CREATE TABLE cust short (custno NUMBER, name VARCHAR2(31));
CREATE TABLE states (state_id NUMBER, addresses address_array_t);
This example casts a subquery:
SELECT s.custno, s.name,
      CAST(MULTISET(SELECT ca.street_address,
                           ca.postal code,
                           ca.city,
                           ca.state_province,
                           ca.country id
                      FROM cust address ca
                       WHERE s.custno = ca.custno)
      AS address book t)
  FROM cust short s
  ORDER BY s.custno;
CAST converts a varray type column into a nested table:
SELECT CAST(s.addresses AS address_book_t)
  FROM states s
  WHERE s.state id = 111;
The following objects create the basis of the example that follows:
```

```
CREATE TABLE projects
  (employee id NUMBER, project name VARCHAR2(10));
CREATE TABLE emps short
  (employee id NUMBER, last name VARCHAR2(10));
CREATE TYPE project table typ AS TABLE OF VARCHAR2 (10);
```

The following example of a MULTISET expression uses these objects:

```
SELECT e.last name,
       CAST (MULTISET (SELECT p.project name
                       FROM projects p
```



The following example casts the string 'yes' to a boolean value, the boolean value true to a NUMBER and the boolean value false to VARCHAR2 (10):

```
SELECT

CAST ( 'yes' AS BOOLEAN ),

CAST ( true AS NUMBER ),

CAST ( false AS VARCHAR2(10) );

CAST('YES'ASBOOLEAN) CAST(TRUEASNUMBER) CAST(FALSE

TRUE 1 FALSE
```

Domain Examples

The following example creates the domain <code>DAY_OF_WEEK</code> with a disabled check constraint. The first query omits the <code>domain_validate_clause</code>, so uses the constraint state to determine whether to verify the value. As this is disabled, the database does not check the value.

The second query uses the VALIDATE clause. This applies the constraint to "N/A", even though it's disabled. The value "N/A" is not in the list permitted by the constraints, so CAST raises an exception.

The following example creates the domain DAY_OF_WEEK with an enabled check constraint. The first query omits the $domain_validate_clause$, so uses the constraint state to determine whether to verify the value. As this is enabled, the database applies the constraint to "N/A". This is not in the list of permitted values so CAST raises an error.

The second query uses the NOVALIDATE clause. This ignores the constraint even though it is enabled and the statement completes without error.

```
CREATE DOMAIN day_of_week AS VARCHAR2(3 CHAR)
  CONSTRAINT CHECK (day_of_week IN('MON','TUE','WED','THU','FRI','SAT','SUN'))
  ENABLE;

SELECT CAST ( 'N/A' AS day_of_week ) use_constraint_state;
```



```
ORA-11513: CAST AS DOMAIN has failed due to domain constraints.

SELECT CAST ('N/A' AS DOMAIN day_of_week NOVALIDATE) ignore_constraints;

IGNORE_CONSTRAINTS
-----N/A
```

CEIL (datetime)

Syntax



Purpose

CEIL (datetime) returns the date or the timestamp rounded up to the unit specified by the second argument fmt, the format model. If the input value is already truncated to the specified unit, then the return value is the same as the input. That is, if datetime = TRUNC (datetime, fmt), then CEIL (datetime, fmt) = datetime. For example, CEIL (DATE '2023-02-01', 'MONTH') returns February 1 2023.

This function is not sensitive to the NLS_CALENDAR session parameter. It operates according to the rules of the Gregorian calendar. The value returned is always of data type DATE, even if you specify a different datetime data type for the argument. If you do not specify the second argument, the default format model 'DD' is used.



Refer to CEIL, FLOOR, ROUND, and TRUNC Date Functions for the permitted format models to use in *fmt*.

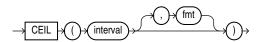
Examples

For these examples NLS DATE FORMAT is set:



CEIL (interval)

Syntax



Purpose

CEIL (interval) returns the interval rounded up to the unit specified by the second argument fmt, the format model. If the first argument is truncated to the units of fmt, the output equals the input. For example, CEIL (INTERVAL '+123-0' YEAR(3) TO MONTH) returns 123 years and no months (+123-00).

The result of <code>CEIL(interval)</code> is never smaller than <code>interval</code>. The result precision for year and day is the input precision for year plus one and day plus one, since <code>CEIL(interval)</code> can have overflow. If an interval already has the maximum precision for year and day, the statement compiles but errors at runtime.

For INTERVAL YEAR TO MONTH, fmt can only be year. The default fmt is year.

For INTERVAL DAY TO SECOND, fmt can be day, hour, and minute. The default fmt is day. Note that fmt does not support second.

CEIL (interval) supports the format models of ROUND and TRUNC.



Refer to CEIL, FLOOR, ROUND, and TRUNC Date Functions for the permitted format models to use in fmt.

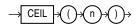
Examples



+05 00:00:00.000000

CEIL (number)

Syntax



Purpose

CEIL returns the smallest integer that is greater than or equal to n. The number n can always be written as the difference of an integer k and a positive fraction f such that $0 \le f \le 1$ and f is f. The value of CEIL is the integer f. Thus, the value of CEIL is f itself if and only if f is precisely an integer.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.



Table 2-9 for more information on implicit conversion and FLOOR (number)

Examples

The following example returns the smallest integer greater than or equal to the order total of a specified order:

CHARTOROWID

Syntax



Purpose

CHARTOROWID converts a value from CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type to ROWID data type.



This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.



Data Type Comparison Rules for more information.

Examples

The following example converts a character rowid representation to a rowid. (The actual rowid is different for each database instance.)

CHECKSUM

Syntax



Purpose

Use CHECKSUM to detect changes in a table. The order of the rows in the table does not affect the result. You can use CHECKSUM with DISTINCT, as part of a GROUP BY query, as a window function, or an analytical function.

Semantics

ALL: Applies the aggregate function to all values. ALL is the default option.

DISTINCT or UNIQUE: Returns the checksum of unique values. UNIQUE is an Oracle-specific keyword and not an ANSI standard.

 ${\tt expr}$: Can be a column, constant, bind variable, or an expression involving them. All data types except ADT and JSON are supported.

The return data type is an Oracle number (converted from an (8-byte) signed long long) regardless of the data type of expr.

NULL values in expr column are ignored.

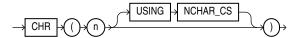
It returns NULL if expr is NULL.

The output of the CHECKSUM function is deterministic and independent of the ordering of the input rows.



CHR

Syntax



Purpose

CHR returns the character having the binary equivalent to n as a VARCHAR2 value in either the database character set or, if you specify USING NCHAR CS, the national character set.

For single-byte character sets, if n > 256, then Oracle Database returns the binary equivalent of $n \mod 256$. For multibyte character sets, n must resolve to one entire code point. Invalid code points are not validated, and the result of specifying invalid code points is indeterminate.

This function takes as an argument a NUMBER value, or any value that can be implicitly converted to NUMBER, and returns a character.



Use of the CHR function (either with or without the optional USING NCHAR_CS clause) results in code that is not portable between ASCII- and EBCDIC-based machine architectures.

See Also:

- NCHR and Table 2-9 for more information on implicit conversion
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of CHR

Examples

The following example is run on an ASCII-based machine with the database character set defined as WE8ISO8859P1:

```
SELECT CHR(67)||CHR(65)||CHR(84) "Dog"
FROM DUAL;

Dog
---
CAT
```

To produce the same results on an EBCDIC-based machine with the WE8EBCDIC1047 character set, the preceding example would have to be modified as follows:

```
SELECT CHR(195)||CHR(193)||CHR(227) "Dog" FROM DUAL;
```



```
Dog
---
CAT
```

For multibyte character sets, this sort of concatenation gives different results. For example, given a multibyte character whose hexadecimal value is ala2 (al representing the first byte and a2 the second byte), you must specify for n the decimal equivalent of 'ala2', or 41378:

```
SELECT CHR(41378)
FROM DUAL;
```

You cannot specify the decimal equivalent of a1 concatenated with the decimal equivalent of a2, as in the following example:

```
SELECT CHR(161) | | CHR(162)
FROM DUAL;
```

However, you can concatenate whole multibyte code points, as in the following example, which concatenates the multibyte characters whose hexadecimal values are ala2 and ala3:

```
SELECT CHR(41378) | | CHR(41379) FROM DUAL;
```

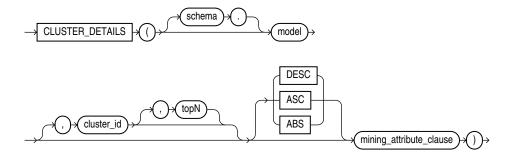
The following example assumes that the national character set is UTF16:

```
SELECT CHR (196 USING NCHAR_CS)
FROM DUAL;
CH
--
A
```

CLUSTER DETAILS

Syntax

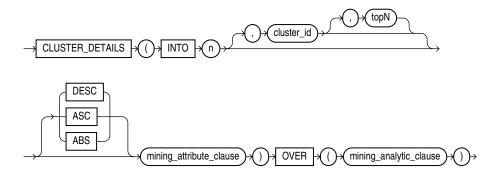
cluster_details::=



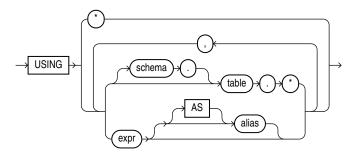


Analytic Syntax

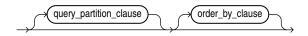
cluster_details_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

Analytic Functions for information on the syntax, semantics, and restrictions of $mining\ analytic\ clause$

Purpose

CLUSTER_DETAILS returns cluster details for each row in the selection. The return value is an XML string that describes the attributes of the highest probability cluster or the specified <code>cluster_id</code>.



topN

If you specify a value for topN, the function returns the N attributes that most influence the cluster assignment (the score). If you do not specify topN, the function returns the 5 most influential attributes.

DESC, ASC, or ABS

The returned attributes are ordered by weight. The weight of an attribute expresses its positive or negative impact on cluster assignment. A positive weight indicates an increased likelihood of assignment. A negative weight indicates a decreased likelihood of assignment.

By default, CLUSTER_DETAILS returns the attributes with the highest positive weights (DESC). If you specify ASC, the attributes with the highest negative weights are returned. If you specify ABS, the attributes with the greatest weights, whether negative or positive, are returned. The results are ordered by absolute value from highest to lowest. Attributes with a zero weight are not included in the output.

Syntax Choice

CLUSTER_DETAILS can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- Syntax Use the first syntax to score the data with a pre-defined model. Supply the name of a clustering model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of clusters to compute, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order by clause. (See analytic_clause::=.)

The syntax of the <code>CLUSTER_DETAILS</code> function can use an optional <code>GROUPING</code> hint when scoring a partitioned model. See <code>GROUPING</code> Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See mining attribute clause::=.)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about clustering.





The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

This example lists the attributes that have the greatest impact (more that 20% probability) on cluster assignment for customer ID 100955. The query invokes the <code>CLUSTER_DETAILS</code> and <code>CLUSTER_SET</code> functions, which apply the clustering model <code>em_sh_clus_sample</code>.

```
SELECT S.cluster id, probability prob,
       CLUSTER DETAILS (em sh clus sample, S.cluster id, 5 USING T.*) det
FROM
  (SELECT v.*, CLUSTER SET(em sh clus sample, NULL, 0.2 USING *) pset
   FROM mining data apply v v
  WHERE cust id = 100955) T,
  TABLE (T.pset) S
ORDER BY 2 DESC;
CLUSTER ID PROB DET
        14 .6761 <Details algorithm="Expectation Maximization" cluster="14">
                 <a href="AGE" actualValue="51" weight=".676" rank="1"/>
                 <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".557" rank="2"/>
                 <Attribute name="FLAT PANEL MONITOR" actualValue="0" weight=".412" rank="3"/>
                 <a href="Attribute name="Y BOX GAMES" actualValue="0" weight=".171" rank="4"/>
                 <Attribute name="BOOKKEEPING APPLICATION" actualValue="1" weight="-.003"rank="5"/>
                 </Details>
         3 .3227 <Details algorithm="Expectation Maximization" cluster="3">
                 <Attribute name="YRS RESIDENCE" actualValue="3" weight=".323" rank="1"/>
                 <Attribute name="BULK PACK DISKETTES" actualValue="1" weight=".265" rank="2"/>
                 <Attribute name="EDUCATION" actualValue="HS-grad" weight=".172" rank="3"/>
                 <Attribute name="AFFINITY CARD" actualValue="0" weight=".125" rank="4"/>
                 <Attribute name="OCCUPATION" actualValue="Crafts" weight=".055" rank="5"/>
                 </Details>
```

Analytic Example

This example divides the customer database into four segments based on common characteristics. The clustering functions compute the clusters and return the score without a predefined clustering model.

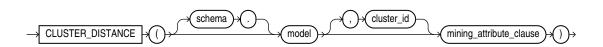


```
rank="3"/>
           <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".268" rank="4"/>
           <Attribute name="Y BOX GAMES" actualValue="0" weight=".179" rank="5"/>
           </Details>
100002
       6 <Details algorithm="K-Means Clustering" cluster="6">
           <Attribute name="CUST GENDER" actualValue="F" weight=".945" rank="1"/>
           <Attribute name="CUST MARITAL STATUS" actualValue="NeverM" weight=".856" rank="2"/>
           <a href="Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".468" rank="3"/>
           <Attribute name="AFFINITY CARD" actualValue="0" weight=".012" rank="4"/>
           <Attribute name="CUST INCOME LEVEL" actualValue="L: 300\,000 and above" weight=".009"</pre>
           rank="5"/>
           </Details>
100003
        7 <Details algorithm="K-Means Clustering" cluster="7">
           <Attribute name="CUST MARITAL STATUS" actualValue="NeverM" weight=".862" rank="1"/>
           <Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".423" rank="2"/>
           <Attribute name="HOME THEATER PACKAGE" actualValue="0" weight=".113" rank="3"/>
           <Attribute name="AFFINITY CARD" actualValue="0" weight=".007" rank="4"/>
           <Attribute name="CUST ID" actualValue="100003" weight=".006" rank="5"/>
           </Details>
```

CLUSTER_DISTANCE

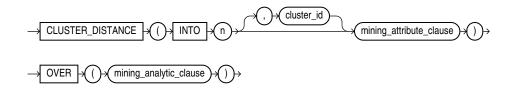
Syntax

cluster distance::=

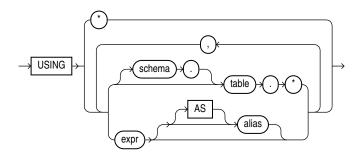


Analytic Syntax

cluster_distance_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

Analytic Functions for information on the syntax, semantics, and restrictions of $mining_analytic_clause$

Purpose

CLUSTER_DISTANCE returns a cluster distance for each row in the selection. The cluster distance is the distance between the row and the centroid of the highest probability cluster or the specified <code>cluster id</code>. The distance is returned as <code>BINARY DOUBLE</code>.

Syntax Choice

CLUSTER_DISTANCE can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- Syntax Use the first syntax to score the data with a pre-defined model. Supply the name of a clustering model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of clusters to compute, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order by clause. (See analytic_clause::=.)

The syntax of the <code>CLUSTER_DISTANCE</code> function can use an optional <code>GROUPING</code> hint when scoring a partitioned model. See <code>GROUPING</code> Hint.

mining attribute clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, this data is also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See mining_attribute_clause::=.)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about clustering.



The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

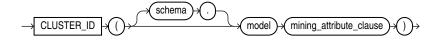
This example finds the 10 rows that are most anomalous as measured by their distance from their nearest cluster centroid.

```
SELECT cust_id
 FROM (
   SELECT cust_id,
           rank() over
             (order by CLUSTER DISTANCE(km sh clus sample USING *) desc) rnk
      FROM mining_data_apply_v)
 WHERE rnk <= 11
 ORDER BY rnk;
  CUST_ID
   100579
   100050
   100329
   100962
   101251
   100179
   100382
   100713
   100629
   100787
   101478
```

CLUSTER_ID

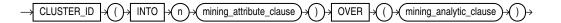
Syntax

cluster_id::=



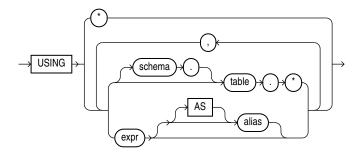
Analytic Syntax

cluster_id_analytic::=

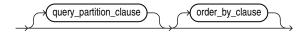




mining_attribute_clause::=



mining_analytic_clause::=



See Also:

Analytic Functions for information on the syntax, semantics, and restrictions of $mining_analytic_clause$

Purpose

CLUSTER_ID returns the identifier of the highest probability cluster for each row in the selection. The cluster identifier is returned as an Oracle NUMBER.

Syntax Choice

CLUSTER_ID can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- **Syntax** Use the first syntax to score the data with a pre-defined model. Supply the name of a clustering model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of clusters to compute, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order by clause. (See analytic_clause::=.)

The syntax of the <code>CLUSTER_ID</code> function can use an optional <code>GROUPING</code> hint when scoring a partitioned model. See <code>GROUPING</code> Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See mining_attribute_clause::=.)



See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about clustering.

Note:

The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

The following example lists the clusters into which the customers in mining_data_apply_v have been grouped.

```
SELECT CLUSTER_ID(km_sh_clus_sample USING *) AS clus, COUNT(*) AS cnt
FROM mining_data_apply_v
GROUP BY CLUSTER_ID(km_sh_clus_sample USING *)
ORDER BY cnt DESC;
```

CLUS	CNT
2	580
10	216
6	186
8	115
19	110
12	101
18	81
16	39
17	38
14	34

Analytic Example

This example divides the customer database into four segments based on common characteristics. The clustering functions compute the clusters and return the score without a predefined clustering model.

```
SELECT * FROM (

SELECT cust_id,

CLUSTER_ID(INTO 4 USING *) OVER () cls,

CLUSTER_DETAILS(INTO 4 USING *) OVER () cls_details

FROM mining_data_apply_v)

WHERE cust_id <= 100003

ORDER BY 1;

CUST_ID CLS CLS_DETAILS

CUST_ID CLS CLS_DETAILS

Attribute name="FLAT_PANEL_MONITOR" cluster="5">

<Attribute name="FLAT_PANEL_MONITOR" actualValue="0" weight=".349" rank="1"/>

<Attribute name="BULK_PACK_DISKETTES" actualValue="0" weight=".33" rank="2"/>

<Attribute name="CUST_INCOME_LEVEL" actualValue="G: 130\,000 - 149\,999"

weight=".291" rank="3"/>
```

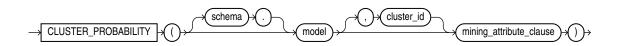


```
<Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".268" rank="4"/>
           <Attribute name="Y BOX GAMES" actualValue="0" weight=".179" rank="5"/>
           </Details>
         6 <Details algorithm="K-Means Clustering" cluster="6">
100002
           <Attribute name="CUST GENDER" actualValue="F" weight=".945" rank="1"/>
           <Attribute name="CUST_MARITAL STATUS" actualValue="NeverM" weight=".856" rank="2"/>
           <Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".468" rank="3"/>
           <Attribute name="AFFINITY CARD" actualValue="0" weight=".012" rank="4"/>
           <a href="Attribute name="CUST INCOME LEVEL" actualValue="L: 300\,000 and above"</a>
              weight=".009" rank="5"/>
           </Details>
100003
        7 <Details algorithm="K-Means Clustering" cluster="7">
           <a href="Attribute name="CUST MARITAL STATUS" actualValue="NeverM" weight=".862" rank="1"/>
           <Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".423" rank="2"/>
           <Attribute name="HOME THEATER PACKAGE" actualValue="0" weight=".113" rank="3"/>
           <Attribute name="AFFINITY CARD" actualValue="0" weight=".007" rank="4"/>
           <Attribute name="CUST ID" actualValue="100003" weight=".006" rank="5"/>
```

CLUSTER PROBABILITY

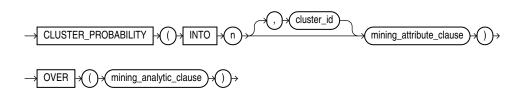
Syntax

cluster probability::=

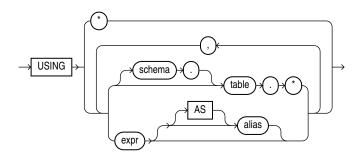


Analytic Syntax

cluster_prob_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

Analytic Functions for information on the syntax, semantics, and restrictions of $mining_analytic_clause$

Purpose

CLUSTER_PROBABILITY returns a probability for each row in the selection. The probability refers to the highest probability cluster or to the specified $cluster_id$. The cluster probability is returned as BINARY DOUBLE.

Syntax Choice

CLUSTER_PROBABILITY can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- Syntax Use the first syntax to score the data with a pre-defined model. Supply the name of a clustering model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of clusters to compute, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order_by_clause. (See analytic_clause::=.)

The syntax of the CLUSTER_PROBABILITY function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See mining_attribute_clause::=.)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about clustering.



The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

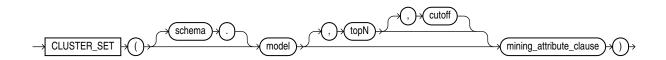
The following example lists the ten most representative customers, based on likelihood, of cluster 2.

```
SELECT cust_id
 FROM (SELECT cust id, rank() OVER (ORDER BY prob DESC, cust id) rnk clus2
   FROM (SELECT cust id, CLUSTER PROBABILITY(km sh clus sample, 2 USING *) prob
          FROM mining_data_apply_v))
WHERE rnk_clus2 <= 10
ORDER BY rnk clus2;
  CUST_ID
   100256
   100988
   100889
   101086
   101215
   100390
   100985
   101026
   100601
   100672
```

CLUSTER_SET

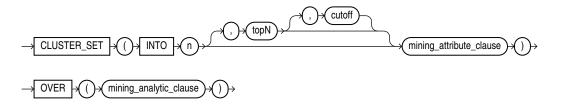
Syntax

cluster_set::=

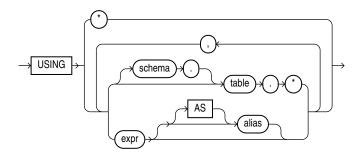


Analytic Syntax

cluster_set_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

Analytic Functions for information on the syntax, semantics, and restrictions of $mining_analytic_clause$

Purpose

CLUSTER_SET returns a set of cluster ID and probability pairs for each row in the selection. The return value is a varray of objects with field names <code>CLUSTER_ID</code> and <code>PROBABILITY</code>. The cluster identifier is an Oracle <code>NUMBER</code>; the probability is <code>BINARY DOUBLE</code>.

topN and cutoff

You can specify topN and cutoff to limit the number of clusters returned by the function. By default, both topN and cutoff are null and all clusters are returned.

- topN is the N most probable clusters. If multiple clusters share the Nth probability, then the function chooses one of them.
- *cutoff* is a probability threshold. Only clusters with probability greater than or equal to *cutoff* are returned. To filter by *cutoff* only, specify NULL for *topN*.

To return up to the N most probable clusters that are greater than or equal to cutoff, specify both topN and cutoff.

Syntax Choice

CLUSTER_SET can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

• **Syntax** — Use the first syntax to score the data with a pre-defined model. Supply the name of a clustering model.

• Analytic Syntax — Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of clusters to compute, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order by clause. (See analytic_clause::=.)

The syntax of the CLUSTER_SET function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See mining_attribute_clause::=.)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about clustering.

Note:

The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

This example lists the attributes that have the greatest impact (more that 20% probability) on cluster assignment for customer ID 100955. The query invokes the <code>CLUSTER_DETAILS</code> and <code>CLUSTER_SET</code> functions, which apply the clustering model <code>em_sh_clus_sample</code>.

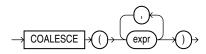
```
SELECT S.cluster id, probability prob,
       CLUSTER DETAILS(em sh clus sample, S.cluster id, 5 USING T.*) det
FROM
  (SELECT v.*, CLUSTER SET(em sh clus sample, NULL, 0.2 USING *) pset
   FROM mining data apply v v
  WHERE cust id = 100955) T,
  TABLE (T.pset) S
ORDER BY 2 DESC;
CLUSTER ID PROB DET
        14 .6761 <Details algorithm="Expectation Maximization" cluster="14">
                 <Attribute name="AGE" actualValue="51" weight=".676" rank="1"/>
                 <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".557" rank="2"/>
                 <Attribute name="FLAT PANEL MONITOR" actualValue="0" weight=".412" rank="3"/>
                 <Attribute name="Y BOX GAMES" actualValue="0" weight=".171" rank="4"/>
                 <Attribute name="BOOKKEEPING_APPLICATION" actualValue="1" weight="-.003"rank="5"/>
                 </Details>
         3 .3227 <Details algorithm="Expectation Maximization" cluster="3">
```



```
<Attribute name="YRS_RESIDENCE" actualValue="3" weight=".323" rank="1"/>
<Attribute name="BULK_PACK_DISKETTES" actualValue="1" weight=".265" rank="2"/>
<Attribute name="EDUCATION" actualValue="HS-grad" weight=".172" rank="3"/>
<Attribute name="AFFINITY_CARD" actualValue="0" weight=".125" rank="4"/>
<Attribute name="OCCUPATION" actualValue="Crafts" weight=".055" rank="5"/>
</Details>
```

COALESCE

Syntax



Purpose

COALESCE returns the first non-null expr in the expression list. You must specify at least two expressions. If all occurrences of expr evaluate to null, then the function returns null.

Oracle Database uses **short-circuit evaluation**. The database evaluates each expr value and determines whether it is NULL, rather than evaluating all of the expr values before determining whether any of them is NULL.

If all occurrences of expr are numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type, then Oracle Database determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.

See Also:

- Table 2-9 for more information on implicit conversion and Numeric Precedence for information on numeric precedence
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of COALESCE when it is a character value

This function is a generalization of the NVL function.

You can also use COALESCE as a variety of the CASE expression. For example,

```
COALESCE(expr1, expr2)
```

is equivalent to:

CASE WHEN expr1 IS NOT NULL THEN expr1 ELSE expr2 END

Similarly,

COALESCE (expr1, expr2, ..., exprn)

where $n \ge 3$, is equivalent to:

CASE WHEN expr1 IS NOT NULL THEN expr1 ELSE COALESCE (expr2, ..., exprn) END



NVL and CASE Expressions

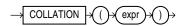
Examples

The following example uses the sample <code>oe.product_information</code> table to organize a clearance sale of products. It gives a 10% discount to all products with a list price. If there is no list price, then the sale price is the minimum price. If there is no minimum price, then the sale price is "5":

Sale	MIN_PRICE	LIST_PRICE	PRODUCT_ID
43.2		48	1769
73	73	40	1770
274.5	247	305	2378
765	731	850	2382
5			3355

COLLATION

Syntax



Purpose

COLLATION returns the name of the derived collation for expr. This function returns named collations and pseudo-collations. If the derived collation is a Unicode Collation Algorithm (UCA) collation, then the function returns the long form of its name. This function is evaluated during compilation of the SQL statement that contains it. If the derived collation is undefined due to a collation conflict while evaluating expr, then the function returns null.

expr must evaluate to a character string of type CHAR, VARCHAR2, LONG, NCHAR, or NVARCHAR2.

This function returns a VARCHAR2 value.



Note:

The <code>COLLATION</code> function returns only the data-bound collation, and not the dynamic collation set by the <code>NLS_SORT</code> parameter. Thus, for a column declared as <code>COLLATE USING_NLS_SORT</code>, the function returns the character value <code>'USING_NLS_SORT'</code>, not the actual value of the session parameter <code>NLS_SORT</code>. You can use the built-in function <code>SYS_CONTEXT('USERENV', 'NLS_SORT')</code> to get the actual value of the session parameter <code>NLS_SORT</code>.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of COLLATION

Examples

The following example returns the derived collation of columns name and id in table id_table:

```
CREATE TABLE id_table
  (name VARCHAR2(64) COLLATE BINARY_AI,
  id VARCHAR2(8) COLLATE BINARY_CI);

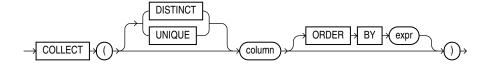
INSERT INTO id_table VALUES('Christopher', 'ABCD1234');

SELECT COLLATION(name), COLLATION(id)
  FROM id_table;

COLLATION COLLATION
------
BINARY AI BINARY CI
```

COLLECT

Syntax



Purpose

COLLECT is an aggregate function that takes as its argument a column of any type and creates a nested table of the input type out of the rows selected. To get accurate results from this function you must use it within a CAST function.

If <code>column</code> is itself a collection, then the output of <code>COLLECT</code> is a nested table of collections. If <code>column</code> is of a user-defined type, then <code>column</code> must have a MAP or <code>ORDER</code> method defined on it in order for you to use the optional <code>DISTINCT</code>, <code>UNIQUE</code>, and <code>ORDER</code> BY clauses.

See Also:

- CAST and Aggregate Functions
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation COLLECT uses to compare character values for the DISTINCT and ORDER BY clauses

Examples

The following example creates a nested table from the varray column of phone numbers in the sample table <code>oe.customers</code>. The nested table includes only the phone numbers of customers with an income level of <code>L: 300,000</code> and above.

The following example creates a nested table from the column of warehouse names in the sample table <code>oe.warehouses</code>. It uses <code>ORDER BY</code> to order the warehouse names.

COMPOSE

Syntax



Purpose

COMPOSE takes as its argument a character value *char* and returns the result of applying the Unicode canonical composition, as described in the Unicode Standard definition D117, to it. If

the character set of the argument is not one of the Unicode character sets, COMPOSE returns its argument unmodified.

COMPOSE does not directly return strings in any of the Unicode normalization forms. To get a string in the NFC form, first call $\tt DECOMPOSE$ with the CANONICAL setting and then $\tt COMPOSE$. To get a string in the NFKC form, first call $\tt DECOMPOSE$ with the $\tt COMPATIBILITY$ setting and then $\tt COMPOSE$.

char can be of any of the data types: CHAR, VARCHAR2, NCHAR, or NVARCHAR2. Other data types are allowed if they can be implicitly converted to VARCHAR2 or NVARCHAR2. The return value of COMPOSE is in the same character set as its argument.

CLOB and NCLOB values are supported through implicit conversion. If char is a character LOB value, then it is converted to a VARCHAR2 value before the COMPOSE operation. The operation will fail if the size of the LOB value exceeds the supported length of the VARCHAR2 in the particular execution environment.

See Also:

- Oracle Database Globalization Support Guide for information on Unicode character sets and character semantics
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of COMPOSE
- DECOMPOSE

Examples

The following example returns the o-umlaut code point:

```
SELECT COMPOSE( 'o' || UNISTR('\0308') )
FROM DUAL;
CO
--
Ö
```

See Also:

UNISTR

CON_DBID_TO_ID

Syntax





Purpose

CON_DBID_TO_ID takes as its argument a container DBID and returns the container ID. For container_dbid, specify a NUMBER value or any value that can be implicitly converted to NUMBER. The function returns a NUMBER value.

This function is useful in a multitenant container database (CDB). If you use this function in a non-CDB, then it returns 0.

Example

The following query displays the ID and DBID for all containers in a CDB. The sample output shown is for the purpose of this example.

The following statement returns the ID for the container with DBID 2256797992:

CON_GUID_TO_ID

Syntax



Purpose

CON_GUID_TO_ID takes as its argument a container GUID (globally unique identifier) and returns the container ID. For container_guid, specify a raw value. The function returns a NUMBER value.

This function is useful in a multitenant container database (CDB). If you use this function in a non-CDB, then it returns 0.

Example

The following query displays the ID and GUID for all containers in a CDB. The GUID is stored as a 16-byte RAW value in the V\$CONTAINERS view. The query returns the 32-character hexadecimal representation of the GUID. The sample output shown is for the purpose of this example.

```
SELECT CON_ID, GUID
FROM V$CONTAINERS;
```



```
CON_ID GUID

1 DB0A9F33DF99567FE04305B4F00A667D
2 D990C280C309591EE04305B4F00A593E
4 D990F4BD938865C1E04305B4F00ACA18
```

The following statement returns the ID for the container whose GUID is represented by the hexadecimal value D990F4BD938865C1E04305B4F00ACA18. The HEXTORAW function converts the GUID's hexadecimal representation to a raw value.

CON_ID_TO_CON_NAME

Syntax

```
CON_ID_TO_CON_NAME () (container_id )
```

Purpose

CON ID TO CON NAME takes as an argument a container CON ID and returns the container NAME.

For CON_ID you must specify a number or an expression that resolves to a number. The function returns a NUMBER value.

This function is useful in a multitentant container database (CDB). If you use this function in a non-CDB, then it returns 0.

Example

```
SELECT CON_ID, NAME FROM V$CONTAINERS;

CON_ID NAME

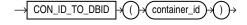
1 CDB$ROOT
2 PDB$SEED
3 CDB1_PDB1
4 SALESPDB
```

The following statement returns the container NAME given the container CON ID 4:

```
SELECT CON_ID_TO_CON_NAME(4) "CON_NAME" FROM DUAL;
CON_NAME
-----
SALESDB
```

CON_ID_TO_DBID

Syntax



Purpose

CON_ID_TO_DBID takes as an argument a container CON_ID and returns the container DBID. For CON_ID you must specify a number or an expression that resolves to a number. The function returns a NUMBER value.

This function is useful in a multitentant container database (CDB). If you use this function in a non-CDB, then it returns 0.

Example

SELECT CON ID, NAME, DBID FROM V\$CONTAINERS;

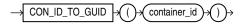
CON_ID	NAME	DBID
1	CDB\$ROOT	2048400776
2	PDB\$SEED	2929762556
3	CDB1_PDB1	3483444080
4	SALESPDB	2221053340

The following statement returns the container DBID given the container CON ID 4:

```
SELECT CON_ID_TO_DBID(4) FROM DUAL;
DBID
-----
2221053340
```

CON_ID_TO_GUID

Syntax



Purpose

CON_ID_TO_GUID takes as an argument a container CON_ID and returns the container's GLOBAL UNIQUE ID (GUID). For CON_ID you must specify a number or an expression that resolves to a number. The function returns a NUMBER value.

This function is useful in a multitentant container database (CDB).

Example

The following query displays the CON ID, NAME and GUID for all containers in a CDB:

SELECT CON_ID, NAME, GUID FROM V\$CONTAINERS;

CON_ID	NAME	GUID
1	CDB\$ROOT	A8C0E03CB11A132FE0532684E80A96B3
2	PDB\$SEED	A8DA5D32F8F5590DE053C4E15A0A6EED
3	CDB1_PDB1	A8DA63CEAD385A5BE053C4E15A0A774A
4	SALESPDB	A8DA9AB18CE85BD0E053C4E15A0AE2C3

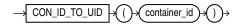


The following statement returns the container \mathtt{GUID} given the container \mathtt{CON} \mathtt{ID} 4:

```
SELECT CON_ID_TO_GUID(4) "CON_GUID" FROM DUAL;
CON_GUID
------
A8DA9AB18CE85BD0E053C4E15A0AE2C3
```

CON_ID_TO_UID

Syntax



Purpose

 $CON_ID_TO_UID$ takes as an argument a container CON_ID and returns the container's UNIQUE ID (UID). For CON_ID you must specify a number or an expression that resolves to a number. The function returns a NUMBER value.

This function is useful in a multitentant container database (CDB).

Example

The following query displays the CON ID, NAME and CON UID for all containers in a CDB:

SELECT CON_ID, NAME, CON_UID FROM V\$CONTAINERS;

CON_ID	NAME	CON_UID
1	CDB\$ROOT	1
2	PDB\$SEED	2929762556
3	CDB1_PDB1	3483444080
4	SALESPDB	2221053340

The following statement returns the container CON UID given the container CON ID 4:

```
SELECT CON_ID_TO_UID(4) "PDB_UID" FROM DUAL;
PDB_UID
------
2221053340
```

CON NAME TO ID

Syntax



Purpose

CON_NAME_TO_ID takes as its argument a container name and returns the container ID. For container_name, specify a string, or an expression that resolves to a string, in any data type. The function returns a NUMBER value.

This function is useful in a multitenant container database (CDB). If you use this function in a non-CDB, then it returns 0.

Example

The following query displays the ID and name for all containers in a CDB. The sample output shown is for the purpose of this example.

```
SELECT CON_ID, NAME
FROM V$CONTAINERS;

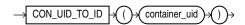
CON_ID NAME

1 CDB$ROOT
2 PDB$SEED
4 SALESPDB
```

The following statement returns the ID for the container named SALESPDB:

CON_UID_TO_ID

Syntax



Purpose

CON_UID_TO_ID takes as its argument a container UID (unique identifier) and returns the container ID. For <code>container_uid</code>, specify a <code>NUMBER</code> value or any value that can be implicitly converted to <code>NUMBER</code>. The function returns a <code>NUMBER</code> value.

This function is useful in a multitenant container database (CDB). If you use this function in a non-CDB, then it returns 0.

Example

The following query displays the ID and UID for all containers in a CDB. The sample output shown is for the purpose of this example.

```
SELECT CON_ID, CON_UID FROM V$CONTAINERS;

CON_ID CON_UID

1 1
2 4054529501
4 2256797992
```

The following query returns the ID for the container with UID 2256797992:

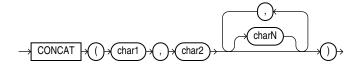
```
SELECT CON_UID_TO_ID(2256797992) "Container ID"
FROM DUAL;
```



Container ID

CONCAT

Syntax



Purpose

CONCAT takes as input two or more arguments and returns the concatenation of all arguments.

The arguments can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. Arguments of other data types are implicitly converted to VARCHAR2 before concatenation.

The string returned is in the same character set as *char1*. Its data type depends on the data types of the arguments.

In concatenations of two or more different data types, Oracle Database returns the data type that results in a lossless conversion. Therefore, if one of the arguments is a LOB, then the returned value is a LOB. If one of the arguments is a national data type, then the returned value is a national data type.

Rules for the Data Types of Return Values

Among all arguments:

- if there is a NCLOB, or if there is a CLOB and a NVARCHAR2 /NCHAR, then the return type is NCLOB.
- otherwise, if there is CLOB, then the return type is CLOB
- otherwise, if there is NVARCHAR2, or if there is a VARCHAR2 and a NCHAR, then the return type is NVARCHAR2
- otherwise, if there is VARCHAR2, then the return type is VARCHAR2
- otherwise, if there is NCHAR, then the return type is NCHAR
- otherwise, the return type is CHAR

Examples of Data Types Returned

CONCAT(CLOB, NCLOB) returns NCLOB

CONCAT(CLOB, NCHAR) returns NCLOB

CONCAT(CLOB, CHAR) returns CLOB

CONCAT (VARCHAR2, NCHAR) returns NVARCHAR2

CONCAT (CHAR, VARCHAR2) returns VARCHAR2

CONCAT (CHAR, VARCHAR2, CLOB) returns CLOB



CONCAT(CHAR, NVARCHAR2, CLOB) returns NCLOB

This function is equivalent to the concatenation operator (||).

See Also:

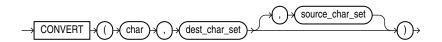
- Concatenation Operator for information on the CONCAT operator
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of CONCAT

Examples

This example concatenates three character strings:

CONVERT

Syntax



Purpose

CONVERT converts a character string from one character set to another.

- The char argument is the value to be converted. It can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB.
- The dest_char set argument is the name of the character set to which char is converted.
- The <code>source_char_set</code> argument is the name of the character set in which <code>char</code> is stored in the database. The default value is the database character set.

The return value for CHAR and VARCHAR2 is VARCHAR2. For NCHAR and NVARCHAR2, it is NVARCHAR2. For CLOB, it is CLOB, and for NCLOB, it is NCLOB.

Both the destination and source character set arguments can be either literals or columns containing the name of the character set.

For complete correspondence in character conversion, it is essential that the destination character set contains a representation of all the characters defined in the source character set. Where a character does not exist in the destination character set, a replacement character appears. Replacement characters can be defined as part of a character set definition.



Note:

Oracle discourages the use of the CONVERT function in the current Oracle Database release. The return value of CONVERT has a character data type, so it should be either in the database character set or in the national character set, depending on the data type. Any <code>dest_char_set</code> that is not one of these two character sets is unsupported. The <code>char</code> argument and the <code>source_char_set</code> have the same requirements. Therefore, the only practical use of the function is to correct data that has been stored in a wrong character set.

Values that are in neither the database nor the national character set should be processed and stored as RAW or BLOB. Procedures in the PL/SQL packages UTL_RAW and UTL_I18N—for example, UTL_RAW.CONVERT—allow limited processing of such values. Procedures accepting a RAW argument in the packages UTL_FILE, UTL_TCP, UTL_HTTP, and UTL_SMTP can be used to output the processed data.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of CONVERT

Examples

The following example illustrates character set conversion by converting a Latin-1 string to ASCII. The result is the same as importing the same string from a WE8ISO8859P1 database to a US7ASCII database.

You can query the V\$NLS VALID VALUES view to get a listing of valid character sets, as follows:

```
SELECT * FROM V$NLS VALID VALUES WHERE parameter = 'CHARACTERSET';
```

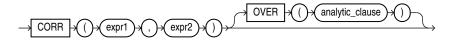
See Also:

Oracle Database Globalization Support Guide for the list of character sets that Oracle Database supports and Oracle Database Reference for information on the V\$NLS VALID VALUES view



CORR

Syntax





Analytic Functions for information on syntax, semantics, and restrictions

Purpose

CORR returns the coefficient of correlation of a set of number pairs. You can use it as an aggregate or analytic function.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.

See Also:

Table 2-9 for more information on implicit conversion and Numeric Precedence for information on numeric precedence

Oracle Database applies the function to the set of (expr1, expr2) after eliminating the pairs for which either expr1 or expr2 is null. Then Oracle makes the following computation:

```
COVAR_POP(expr1, expr2) / (STDDEV_POP(expr1) * STDDEV_POP(expr2))
```

The function returns a value of type NUMBER. If the function is applied to an empty set, then it returns null.

Note:

The CORR function calculates the Pearson's correlation coefficient, which requires numeric expressions as arguments. Oracle also provides the CORR_S (Spearman's rho coefficient) and CORR_K (Kendall's tau-b coefficient) functions to support nonparametric or rank correlation.



See Also:

Aggregate Functions , About SQL Expressions for information on valid forms of expr, and $CORR_*$ for information on the $CORR_S$ and $CORR_K$ functions

Aggregate Example

The following example calculates the coefficient of correlation between the list prices and minimum prices of products by weight class in the sample table <code>oe.product_information</code>:

Analytic Example

The following example shows the correlation between duration at the company and salary by the employee's position. The result set shows the same correlation for each employee in a given job:

CORR_*

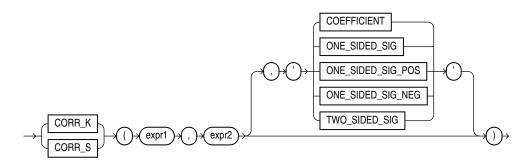
The CORR * functions are:

CORR S

CORR K

Syntax

correlation::=



Purpose

The CORR function (see CORR) calculates the Pearson's correlation coefficient and requires numeric expressions as input. The CORR_* functions support nonparametric or rank correlation. They let you find correlations between expressions that are ordinal scaled (where a ranking of the values is possible). Correlation coefficients take on a value ranging from -1 to 1, where 1 indicates a perfect relationship, -1 a perfect inverse relationship (when one variable increases as the other decreases), and a value close to 0 means no relationship.

These functions take two mandatory arguments, expr1 and expr2, and an optional third argument. The return value of the functions is a NUMBER.

The mandatory arguments are the two variables being analyzed. They can be of any data type that is comparable other than LONG, CLOB, BLOB, BFILE, or VECTOR. If the data type is a user-defined type (UDT), it must have a MAP or ORDER method to be comparable.

The third argument specifies the variant of the result returned by the functions. It is of type VARCHAR2 and must be a constant expression, for example, a character literal. If you omit the third argument, then the default is 'COEFFICIENT'. The allowed argument values and their meaning are shown in Table 7-2 Table 7-2:

See Also:

- Table 2-9 for more information on implicit conversion and Numeric Precedence for information on numeric precedence
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation CORR_K and CORR_S use to compare characters from expr1 with characters from expr2

Table 7-2 CORR_* Return Values

Return Value	Meaning
'COEFFICIENT'	Coefficient of correlation



Table 7-2 (Cont.) CORR_* Return Values

Return Value	Meaning
'ONE_SIDED_SIG'	Positive one-tailed significance of the correlation
'ONE_SIDED_SIG_POS'	Same as ONE_SIDED_SIG
'ONE_SIDED_SIG_NEG'	Negative one-tailed significance of the correlation
'TWO_SIDED_SIG'	Two-tailed significance of the correlation

CORR_S

CORR_S calculates the Spearman's rho correlation coefficient. The input expressions should be a set of (x_i, y_i) pairs of observations. The function first replaces each value with a rank. Each value of x_i is replaced with its rank among all the other x_i s in the sample, and each value of y_i is replaced with its rank among all the other y_i s. Thus, each x_i and y_i take on a value from 1 to n, where n is the total number of pairs of values. Ties are assigned the average of the ranks they would have had if their values had been slightly different. Then the function calculates the linear correlation coefficient of the ranks.

CORR_S Example

Using Spearman's rho correlation coefficient, the following example derives a coefficient of correlation for each of two different comparisons -- salary and commission_pct, and salary and employee id:

CORR K

CORR_K calculates the Kendall's tau-b correlation coefficient. As for $CORR_S$, the input expressions are a set of (x_i, y_i) pairs of observations. To calculate the coefficient, the function counts the number of concordant and discordant pairs. A pair of observations is concordant if the observation with the larger x also has a larger value of y. A pair of observations is discordant if the observation with the larger x has a smaller y.

The significance of tau-b is the probability that the correlation indicated by tau-b was due to chance—a value of 0 to 1. A small value indicates a significant correlation for positive values of tau-b (or anticorrelation for negative values of tau-b).

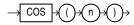
CORR_K Example

Using Kendall's tau-b correlation coefficient, the following example determines whether a correlation exists between an employee's salary and commission percent:

```
COEFFICIENT TWO_SIDED_P_VALUE
------
.603079768 3.4702E-07
```

COS

Syntax



Purpose

cos returns the cosine of n (an angle expressed in radians).

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is $\verb|BINARY_FLOAT|$, then the function returns $\verb|BINARY_DOUBLE|$. Otherwise the function returns the same numeric data type as the argument.



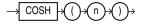
Table 2-9 for more information on implicit conversion

Examples

The following example returns the cosine of 180 degrees:

COSH

Syntax



Purpose

COSH returns the hyperbolic cosine of n.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



See Also:

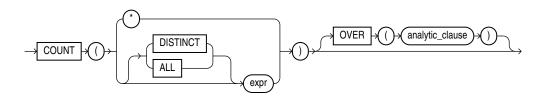
Table 2-9 for more information on implicit conversion

Examples

The following example returns the hyperbolic cosine of zero:

COUNT

Syntax



See Also:

Analytic Functions for information on syntax, semantics, and restrictions

Purpose

COUNT returns the number of rows returned by the query. You can use it as an aggregate or analytic function.

If you specify DISTINCT, then you can specify only the <code>query_partition_clause</code> of the <code>analytic_clause</code>. The <code>order_by_clause</code> and <code>windowing_clause</code> are not allowed.

If you specify expr, then COUNT returns the number of rows where expr is not null. You can count either all rows, or only distinct values of expr.

If you specify the asterisk (*), then this function returns all rows, including duplicates and nulls. COUNT never returns null.



Note:

Before performing a COUNT (DISTINCT expr) operation on a large amount of data, consider using one of the following methods to obtain approximate results more quickly than exact results:

- Set the APPROX_FOR_COUNT_DISTINCT initialization parameter to true before using the COUNT (DISTINCT expr) function. Refer to *Oracle Database Reference* for more information on this parameter.
- Use the APPROX_COUNT_DISTINCT function instead of the COUNT (DISTINCT expr) function. Refer to APPROX_COUNT_DISTINCT.

See Also:

- "About SQL Expressions" for information on valid forms of expr and Aggregate Functions
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation COUNT uses to compare character values for the DISTINCT clause

Aggregate Examples

The following examples use COUNT as an aggregate function:

```
SELECT COUNT(*) "Total"
 FROM employees;
    Total
      107
SELECT COUNT(*) "Allstars"
 FROM employees
 WHERE commission pct > 0;
Allstars
SELECT COUNT(commission pct) "Count"
 FROM employees;
    Count
-----
       35
SELECT COUNT(DISTINCT manager id) "Managers"
 FROM employees;
 Managers
       18
```



Analytic Example

The following example calculates, for each employee in the employees table, the moving count of employees earning salaries in the range 50 less than through 150 greater than the employee's salary.

```
SELECT last_name, salary,

COUNT(*) OVER (ORDER BY salary RANGE BETWEEN 50 PRECEDING AND

150 FOLLOWING) AS mov_count

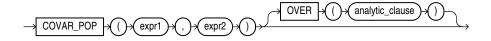
FROM employees

ORDER BY salary, last_name;
```

LAST_NAME	SALARY	MOV_COUNT
Olson	2100	3
Markle	2200	2
Philtanker	2200	2
Gee	2400	8
Landry	2400	8
Colmenares	2500	10
Marlow	2500	10
Patel	2500	10

COVAR POP

Syntax





"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

 $COVAR_POP$ returns the population covariance of a set of number pairs. You can use it as an aggregate or analytic function.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.

See Also:

Table 2-9 for more information on implicit conversion and Numeric Precedence for information on numeric precedence

Oracle Database applies the function to the set of (expr1, expr2) pairs after eliminating all pairs for which either expr1 or expr2 is null. Then Oracle makes the following computation:

```
(SUM(expr1 * expr2) - SUM(expr2) * SUM(expr1) / n) / n
```

where n is the number of (expr1, expr2) pairs where neither expr1 nor expr2 is null.

The function returns a value of type NUMBER. If the function is applied to an empty set, then it returns null.



About SQL Expressions for information on valid forms of expr and Aggregate Functions

Aggregate Example

The following example calculates the population covariance and sample covariance for time employed (SYSDATE - hire_date) and salary using the sample table hr.employees:

Analytic Example

The following example calculates cumulative sample covariance of the list price and minimum price of the products in the sample schema oe:

```
SELECT product id, supplier id,
      COVAR POP(list price, min price)
        OVER (ORDER BY product id, supplier id)
        AS CUM COVP,
       COVAR SAMP(list price, min price)
        OVER (ORDER BY product_id, supplier_id)
        AS CUM COVS
 FROM product_information p
 WHERE category_id = 29
 ORDER BY product id, supplier id;
PRODUCT ID SUPPLIER ID CUM COVP CUM COVS
______
     1774 103088 0
1775 103087 1473.25
1794 103096 1702.77778 1
1825 103093 1926.25
               103087 1473.25 2946.5
               103096 1702.77778 2554.16667
               103093 1926.25 2568.33333
```



```
    2004
    103086
    1591.4
    1989.25

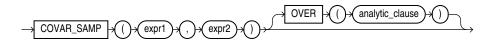
    2005
    103086
    1512.5
    1815

    2416
    103088
    1475.97959
    1721.97619
```

. .

COVAR SAMP

Syntax





Analytic Functions for information on syntax, semantics, and restrictions

Purpose

COVAR_SAMP returns the sample covariance of a set of number pairs. You can use it as an aggregate or analytic function.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.

See Also:

Table 2-9 for more information on implicit conversion and Numeric Precedence for information on numeric precedence

Oracle Database applies the function to the set of (expr1, expr2) pairs after eliminating all pairs for which either expr1 or expr2 is null. Then Oracle makes the following computation:

```
(SUM(expr1 * expr2) - SUM(expr1) * SUM(expr2) / n) / (n-1)
```

where n is the number of (expr1, expr2) pairs where neither expr1 nor expr2 is null.

The function returns a value of type NUMBER. If the function is applied to an empty set, then it returns null.

See Also:

About SQL Expressions for information on valid forms of expr and Aggregate Functions



Aggregate Example

Refer to the aggregate example for COVAR_POP.

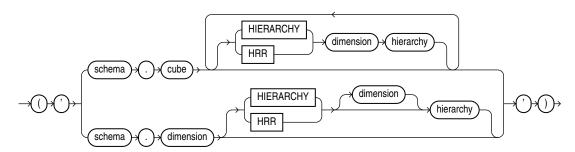
Analytic Example

Refer to the analytic example for COVAR_POP.

CUBE TABLE

Syntax





Purpose

CUBE_TABLE extracts data from a cube or dimension and returns it in the two-dimensional format of a relational table, which can be used by SQL-based applications.

The function takes a single VARCHAR2 argument. The optional hierarchy clause enables you to specify a dimension hierarchy. A cube can have multiple hierarchy clauses, one for each dimension.

You can generate these different types of tables:

- A cube table contains a key column for each dimension and a column for each measure
 and calculated measure in the cube. To create a cube table, you can specify the cube with
 or without a cube hierarchy clause. For a dimension with multiple hierarchies, this clause
 limits the return values to the dimension members and levels in the specified hierarchy.
 Without a hierarchy clause, all dimension members and all levels are included.
- A dimension table contains a key column, and a column for each level and each attribute.
 It also contains a MEMBER_TYPE column, which identifies each member with one of the following codes:
 - L Loaded from a table, view, or synonym
 - A Loaded member and the single root of all hierarchies in the dimension, that is, the "all" aggregate member
 - C Calculated member

All dimension members and all levels are included in the table. To create a dimension table, specify the dimension **without** a dimension hierarchy clause.

• A hierarchy table contains all the columns of a dimension table plus a column for the parent member and a column for each source level. It also contains a MEMBER TYPE

column, as described for dimension tables. Any dimension members and levels that are not part of the named hierarchy are excluded from the table. To create a hierarchy table, specify the dimension **with** a dimension hierarchy clause.

CUBE_TABLE is a table function and is always used in the context of a SELECT statement with this syntax:

```
SELECT ... FROM TABLE(CUBE TABLE('arg'));
```

See Also:

- Oracle OLAP User's Guide for information about dimensional objects and about the tables generated by CUBE TABLE.
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to each character data type column in the table generated by CUBE TABLE

Examples

The following examples require Oracle Database with the OLAP option and the GLOBAL sample schema. Refer to *Oracle OLAP User's Guide* for information on downloading and installing the GLOBAL sample schema.

The following SELECT statement generates a dimension table of CHANNEL in the GLOBAL schema.

The next statement generates a cube table of <code>UNITS_CUBE</code>. It restricts the table to the <code>MARKET</code> and <code>CALENDAR</code> hierarchies.

```
SELECT sales, units, cost, time, customer, product, channel
   FROM TABLE(CUBE_TABLE('global.units_cube HIERARCHY customer market HIERARCHY time calendar'))
   WHERE rownum < 20;</pre>
```

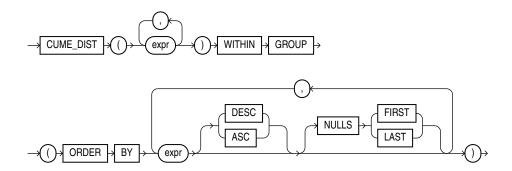
SALES	UNITS	COST	TIME	CUSTOMER	PRODUCT	CHANNEL
24538587.9 24993273.3 25080541.4 26258474 32785170	61320 65265 66122	23147171 23242535.4 24391020.6	CALENDAR_QUARTER_CY1998.Q1 CALENDAR_QUARTER_CY1998.Q2 CALENDAR_QUARTER_CY1998.Q3 CALENDAR_QUARTER_CY1998.Q4 CALENDAR_QUARTER_CY1999.Q1	TOTAL_TOTAL TOTAL_TOTAL TOTAL_TOTAL	TOTAL_TOTAL TOTAL_TOTAL TOTAL_TOTAL	TOTAL_TOTAL TOTAL_TOTAL TOTAL_TOTAL TOTAL_TOTAL TOTAL_TOTAL



CUME_DIST

Aggregate Syntax

cume_dist_aggregate::=



Analytic Syntax

cume_dist_analytic::=





Analytic Functions for information on syntax, semantics, and restrictions

Purpose

CUME_DIST calculates the cumulative distribution of a value in a group of values. The range of values returned by CUME_DIST is >0 to <=1. Tie values always evaluate to the same cumulative distribution value.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle Database determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, makes the calculation, and returns NUMBER.



Table 2-9 for more information on implicit conversion and Numeric Precedence for information on numeric precedence

• As an aggregate function, $CUME_DIST$ calculates, for a hypothetical row r identified by the arguments of the function and a corresponding sort specification, the relative position of

row x among the rows in the aggregation group. Oracle makes this calculation as if the hypothetical row x were inserted into the group of rows to be aggregated over. The arguments of the function identify a single hypothetical row within each aggregate group. Therefore, they must all evaluate to constant expressions within each aggregate group. The constant argument expressions and the expressions in the ORDER BY clause of the aggregate match by position. Therefore, the number of arguments must be the same and their types must be compatible.

• As an analytic function, $CUME_DIST$ computes the relative position of a specified value in a group of values. For a row r, assuming ascending ordering, the $CUME_DIST$ of r is the number of rows with values lower than or equal to the value of r, divided by the number of rows being evaluated (the entire query result set or a partition).

Aggregate Example

The following example calculates the cumulative distribution of a hypothetical employee with a salary of \$15,500 and commission rate of 5% among the employees in the sample table oe.employees:

Analytic Example

The following example calculates the salary percentile for each employee in the purchasing division. For example, 40% of clerks have salaries less than or equal to Himuro.

```
SELECT job_id, last_name, salary, CUME_DIST()

OVER (PARTITION BY job_id ORDER BY salary) AS cume_dist

FROM employees

WHERE job_id LIKE 'PU%'

ORDER BY job_id, last_name, salary, cume_dist;
```

JOB_ID	LAST_NAME	SALARY	CUME_DIST
PU_CLERK	Baida	2900	.8
PU_CLERK	Colmenares	2500	.2
PU_CLERK	Himuro	2600	. 4
PU_CLERK	Khoo	3100	1
PU_CLERK	Tobias	2800	.6
PU_MAN	Raphaely	11000	1

CURRENT_DATE

Syntax



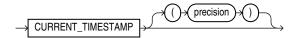
Purpose

CURRENT_DATE returns the current date in the session time zone, in a value in the Gregorian calendar of data type DATE.

The following example illustrates that CURRENT DATE is sensitive to the session time zone:

CURRENT_TIMESTAMP

Syntax



Purpose

CURRENT_TIMESTAMP returns the current date and time in the session time zone, in a value of data type TIMESTAMP WITH TIME ZONE. The time zone offset reflects the current local time of the SQL session. If you omit precision, then the default is 6. The difference between this function and LOCALTIMESTAMP is that CURRENT_TIMESTAMP returns a TIMESTAMP WITH TIME ZONE value while LOCALTIMESTAMP returns a TIMESTAMP value.

In the optional argument, *precision* specifies the fractional second precision of the time value returned.



Examples

The following example illustrates that CURRENT TIMESTAMP is sensitive to the session time zone:



When you use the CURRENT_TIMESTAMP with a format mask, take care that the format mask matches the value returned by the function. For example, consider the following table:

```
CREATE TABLE current test (col1 TIMESTAMP WITH TIME ZONE);
```

The following statement fails because the mask does not include the TIME ZONE portion of the type returned by the function:

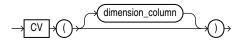
```
INSERT INTO current_test VALUES
  (TO_TIMESTAMP_TZ(CURRENT_TIMESTAMP, 'DD-MON-RR HH.MI.SSXFF PM'));
```

The following statement uses the correct format mask to match the return type of CURRENT TIMESTAMP:

```
INSERT INTO current_test VALUES
  (TO TIMESTAMP TZ(CURRENT TIMESTAMP, 'DD-MON-RR HH.MI.SSXFF PM TZH:TZM'));
```



Syntax



Purpose

The CV function can be used only in the $model_clause$ of a SELECT statement and then only on the right-hand side of a model rule. It returns the current value of a dimension column or a partitioning column carried from the left-hand side to the right-hand side of a rule. This function is used in the $model_clause$ to provide relative indexing with respect to the dimension column. The return type is that of the data type of the dimension column. If you omit the argument, then it defaults to the dimension column associated with the relative position of the function within the cell reference.

The CV function can be used outside a cell reference. In this case, <code>dimension_column</code> is required.

See Also:

- model clause and Model Expressions for the syntax and semantics
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of CV when it is a character value



The following example assigns the sum of the sales of the product represented by the current value of the dimension column (Mouse Pad or Standard Mouse) for years 1999 and 2000 to the sales of that product for year 2001:

```
SELECT country, prod, year, s
FROM sales_view_ref
MODEL

PARTITION BY (country)
DIMENSION BY (prod, year)
MEASURES (sale s)
IGNORE NAV
UNIQUE DIMENSION
RULES UPSERT SEQUENTIAL ORDER
(
    s[FOR prod IN ('Mouse Pad', 'Standard Mouse'), 2001] =
    s[CV(), 1999] + s[CV(), 2000]
)
ORDER BY country, prod, year;
```

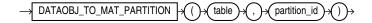
COUNTRY	PROD	YEAR	S
France	Mouse Pad	1998	2509.42
France	Mouse Pad	1999	3678.69
France	Mouse Pad	2000	3000.72
France	Mouse Pad	2001	6679.41
France	Standard Mouse	1998	2390.83
France	Standard Mouse	1999	2280.45
France	Standard Mouse	2000	1274.31
France	Standard Mouse	2001	3554.76
Germany	Mouse Pad	1998	5827.87
Germany	Mouse Pad	1999	8346.44
Germany	Mouse Pad	2000	7375.46
Germany	Mouse Pad	2001	15721.9
Germany	Standard Mouse	1998	7116.11
Germany	Standard Mouse	1999	6263.14
Germany	Standard Mouse	2000	2637.31
Germany	Standard Mouse	2001	8900.45

16 rows selected.

The preceding example requires the view sales_view_ref. Refer to The MODEL clause: Examples to create this view.

DATAOBJ_TO_MAT_PARTITION

Syntax



Purpose

DATAOBJ_TO_MAT_PARTITION is useful only to Data Cartridge developers who are performing data maintenance or query operations on system-partitioned tables that are used to store domain index data. The DML or query operations are triggered by corresponding operations on the base table of the domain index.

This function takes as arguments the name of the base table and the partition ID of the base table partition, both of which are passed to the function by the appropriate ODCIIndex method. The function returns the materialized partition number of the corresponding system-partitioned table, which can be used to perform the operation (DML or query) on that partition of the system-partitioned table.

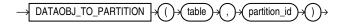
If the base table is interval partitioned, then Oracle recommends that you use this function instead of the <code>DATAOBJ_TO_PARTITION</code> function. The <code>DATAOBJ_TO_PARTITION</code> function determines the absolute partition number, given the physical partition identifier. However, if the base table is interval partitioned, then there might be holes in the partition numbers corresponding to unmaterialized partitions. Because the system partitioned table only has materialized partitions, <code>DATAOBJ_TO_PARTITION</code> numbers can cause a mis-match between the partitions of the base table and the partitions of the underlying system partitioned index storage tables. The <code>DATAOBJ_TO_MAT_PARTITION</code> function returns the materialized partition number (as opposed to the absolute partition number) and helps keep the two tables in sync. Indextypes planning to support local domain indexes on interval partitioned tables should migrate to the use of this function.

See Also:

- DATAOBJ TO PARTITION
- Oracle Database Data Cartridge Developer's Guide for information on the use of the DATAOBJ TO MAT PARTITION function, including examples

DATAOBJ TO PARTITION

Syntax



Purpose

DATAOBJ_TO_PARTITION is useful only to Data Cartridge developers who are performing data maintenance or query operations on system-partitioned tables that are used to store domain index data. The DML or query operations are triggered by corresponding operations on the base table of the domain index.

This function takes as arguments the name of the base table and the partition ID of the base table partition, both of which are passed to the function by the appropriate ODCIIndex method. The function returns the absolute partition number of the corresponding system-partitioned table, which can be used to perform the operation (DML or query) on that partition of the system-partitioned table.

Note:

If the base table is interval partitioned, then Oracle recommends that you instead use the <code>DATAOBJ_TO_MAT_PARTITION</code> function. Refer to <code>DATAOBJ_TO_MAT_PARTITION</code> for more information.



Oracle Database Data Cartridge Developer's Guide for information on the use of the DATAOBJ TO PARTITION function, including examples

DBTIMEZONE

Syntax



Purpose

DBTIMEZONE returns the value of the database time zone. The return type is a time zone offset (a character type in the format '[+|-]TZH:TZM') or a time zone region name, depending on how the user specified the database time zone value in the most recent CREATE DATABASE or ALTER DATABASE statement.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of <code>DBTIMEZONE</code>

Examples

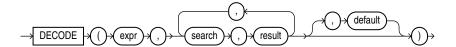
The following example assumes that the database time zone is set to UTC time zone:

```
SELECT DBTIMEZONE FROM DUAL;

DBTIME -----++00:00
```

DECODE

Syntax





Purpose

DECODE compares expr to each search value one by one. If expr is equal to a search, then Oracle Database returns the corresponding result. If no match is found, then Oracle returns default. If default is omitted, then Oracle returns null.

- If expr and search are character data, then Oracle compares them using nonpadded comparison semantics. expr, search, and result can be any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. The string returned is of VARCHAR2 data type and is in the same character set as the first result parameter.
- If the first <code>search-result</code> pair are numeric, then Oracle compares all <code>search-result</code> expressions and the first <code>expr</code> to determine the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.

The *search*, *result*, and *default* values can be derived from expressions. Oracle Database uses **short-circuit evaluation**. The database evaluates each *search* value only before comparing it to *expr*, rather than evaluating all *search* values before comparing any of them with *expr*. Consequently, Oracle never evaluates a *search* if a previous *search* is equal to *expr*.

Oracle automatically converts <code>expr</code> and <code>each</code> <code>search</code> value to the data type of the first <code>search</code> value before comparing. Oracle automatically converts the return value to the same data type as the first <code>result</code>. If the first <code>result</code> has the data type <code>CHAR</code> or if the first <code>result</code> is null, then Oracle converts the return value to the data type <code>VARCHAR2</code>.

In a DECODE function, Oracle considers two nulls to be equivalent. If *expr* is null, then Oracle returns the *result* of the first *search* that is also null.

The maximum number of components in the DECODE function, including expr, searches, results, and default, is 255.

See Also:

- Data Type Comparison Rules for information on comparison semantics
- Data Conversion for information on data type conversion in general
- Floating-Point Numbers for information on floating-point comparison semantics
- Implicit and Explicit Data Conversion for information on the drawbacks of implicit conversion
- COALESCE and CASE Expressions , which provide functionality similar to that of DECODE
- Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules, which define the collation <code>DECODE</code> uses to compare characters from <code>expr</code> with characters from <code>search</code>, and for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value



This example decodes the value warehouse_id. If warehouse_id is 1, then the function returns 'Southlake'; if warehouse_id is 2, then it returns 'San Francisco'; and so forth. If warehouse_id is not 1, 2, 3, or 4, then the function returns 'Non domestic'.

```
SELECT product_id,

DECODE (warehouse_id, 1, 'Southlake',

2, 'San Francisco',

3, 'New Jersey',

4, 'Seattle',

'Non domestic') "Location"

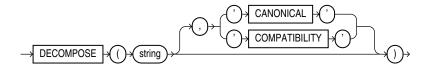
FROM inventories

WHERE product_id < 1775

ORDER BY product id, "Location";
```

DECOMPOSE

Syntax



Purpose

DECOMPOSE takes as its first argument a character value <code>string</code> and returns the result of applying one of the Unicode decompositions to it. The decomposition to apply is determined by the second, optional parameter. If the character set of the first argument is not one of the Unicode character sets, <code>DECOMPOSE</code> returns the argument unmodified.

If the second argument to DECOMPOSE is the string CANONICAL (case-insensitively), DECOMPOSE applies canonical decomposition, as described in the Unicode Standard definition D68, and returns a string in the NFD normalization form. If the second argument is the string COMPATIBILITY, DECOMPOSE applies compatibility decomposition, as described in the Unicode Standard definition D65, and returns a string in the NFKD normalization form. The default behavior is to apply the canonical decomposition.

In a pessimistic case, the return value of <code>DECOMPOSE</code> may be a few times longer than <code>string</code>. If a string to be returned is longer than the maximum length <code>VARCHAR2</code> value in a given runtime environment, the value is silently truncated to the maximum <code>VARCHAR2</code> length.

Both arguments to DECOMPOSE can be of any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. Other data types are allowed if they can be implicitly converted to VARCHAR2 or NVARCHAR2. The return value of DECOMPOSE is in the same character set as its first argument.

CLOB and NCLOB values are supported through implicit conversion. If <code>string</code> is a character LOB value, then it is converted to a <code>VARCHAR2</code> value before the <code>DECOMPOSE</code> operation. The operation will fail if the size of the LOB value exceeds the supported length of the <code>VARCHAR2</code> in the particular execution environment.



See Also:

- Oracle Database Globalization Support Guide for information on Unicode character sets and character semantics
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of DECOMPOSE
- COMPOSE

Examples

The following example decomposes the string "Châteaux" into its component code points:

```
SELECT DECOMPOSE ('Châteaux')
FROM DUAL;

DECOMPOSE
-----
Châteaux
```

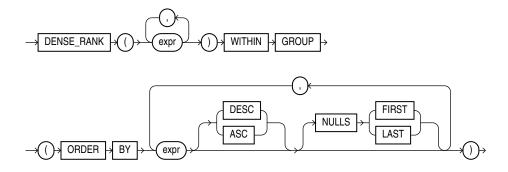


The results of this example can vary depending on the character set of your operating system.

DENSE_RANK

Aggregate Syntax

dense_rank_aggregate::=



Analytic Syntax

dense_rank_analytic::=





Analytic Functions for information on syntax, semantics, and restrictions

Purpose

DENSE_RANK computes the rank of a row in an ordered group of rows and returns the rank as a NUMBER. The ranks are consecutive integers beginning with 1. The largest rank value is the number of unique values returned by the query. Rank values are not skipped in the event of ties. Rows with equal values for the ranking criteria receive the same rank. This function is useful for top-N and bottom-N reporting.

This function accepts as arguments any numeric data type and returns NUMBER.

- As an aggregate function, DENSE_RANK calculates the dense rank of a hypothetical row identified by the arguments of the function with respect to a given sort specification. The arguments of the function must all evaluate to constant expressions within each aggregate group, because they identify a single row within each group. The constant argument expressions and the expressions in the order_by_clause of the aggregate match by position. Therefore, the number of arguments must be the same and types must be compatible.
- As an analytic function, DENSE_RANK computes the rank of each row returned from a query
 with respect to the other rows, based on the values of the value_exprs in the
 order by clause.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules, which define the collation <code>DENSE_RANK</code> uses to compare character values for the <code>ORDER BY</code> clause

Aggregate Example

The following example computes the ranking of a hypothetical employee with the salary \$15,500 and a commission of 5% in the sample table <code>oe.employees</code>:

Analytic Example

The following statement ranks the employees in the sample hr schema in department 60 based on their salaries. Identical salary values receive the same rank. However, no rank values are skipped. Compare this example with the analytic example for RANK .



ORDER BY DENSE RANK, last name;

DEPARTMENT_ID	LAST_NAME	SALARY	DENSE_RANK
60	Lorentz	4200	1
60	Austin	4800	2
60	Pataballa	4800	2
60	Ernst	6000	3
60	Hunold	9000	4

DEPTH

Syntax



Purpose

DEPTH is an ancillary function used only with the <code>UNDER_PATH</code> and <code>EQUALS_PATH</code> conditions. It returns the number of levels in the path specified by the <code>UNDER_PATH</code> condition with the same correlation variable.

The *correlation_integer* can be any NUMBER integer. Use it to correlate this ancillary function with its primary condition if the statement contains multiple primary conditions. Values less than 1 are treated as 1.



EQUALS_PATH Condition, UNDER_PATH Condition, and the related function PATH

Examples

The EQUALS_PATH and UNDER_PATH conditions can take two ancillary functions, DEPTH and PATH. The following example shows the use of both ancillary functions. The example assumes the existence of the XMLSchema warehouses.xsd (created in Using XML in SQL Statements).

```
SELECT PATH(1), DEPTH(2)

FROM RESOURCE_VIEW

WHERE UNDER_PATH(res, '/sys/schemas/OE', 1)=1

AND UNDER_PATH(res, '/sys/schemas/OE', 2)=1;

PATH(1)

DEPTH(2)

...

www.example.com

1

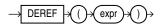
www.example.com/xwarehouses.xsd

2
```



DEREF

Syntax



Purpose

DEREF returns the object reference of argument expr, where expr must return a REF to an object. If you do not use this function in a query, then Oracle Database returns the object ID of the REF instead, as shown in the example that follows.

```
See Also:

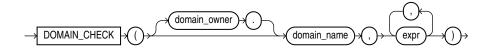
MAKE_REF
```

Examples

The sample schema oe contains an object type <code>cust_address_typ</code>. The REF Constraint <code>Examples</code> create a similar type, <code>cust_address_typ_new</code>, and a table with one column that is a <code>REF</code> to the type. The following example shows how to insert into such a column and how to use <code>DEREF</code> to extract information from the column:

DOMAIN CHECK

Syntax



Purpose

DOMAIN_CHECK first converts the data type of the arguments in <code>expr</code> to the data type of their corresponding domain columns. It then applies the constraint conditions (not null or check constraint) on <code>domain name to expr</code>.

If the domain's constraint is deferred or unvalidated, <code>DOMAIN_CHECK</code> still applies the conditions to <code>expr</code>. If the domain's constraint is disabled, it is not checked as part of <code>DOMAIN CHECK</code>.



Domain Functions

- domain_name must be an identifier and can be specified using domain_owner.domain_name.
 If you specify it without domain_owner, it resolves first to the current user then as a public synonym. If the name cannot be resolved, an error is raised.
- If domain_name refers to a non-existent domain or one that you do not have EXECUTE privileges on, then DOMAIN CHECK will raise an error.
- If the domain column data type is STRICT, then the value is converted to the domain column's data type. For example, if the domain column data type is VARCHAR2 (100) STRICT, then the value is converted to VARCHAR2 (100). Note that the conversion will not automatically trim the input to the maximum length. If the value evaluates to 'abc' for some row and the domain data type is CHAR (2 CHAR), the conversion will fail instead of returning 'ab'.

If the domain column data type is not STRICT, then the value is converted to the most permissive variant of the domain column's data type in terms of length, scale, and precision. For example, if the input value is a VARCHAR2 (30), it is converted to a VARCHAR2 (100) because it is shorter than the domain length. If the input value is a VARCHAR2 (200), it remains a VARCHAR2 (200) because this is larger than the domain length.

- If the data type conversion fails, the error is masked and DOMAIN_CHECK returns FALSE. You can use DOMAIN_CHECK to filter out values that cannot be inserted into a column of the given domain.
 - If the data type conversion succeeds and <code>domain_name</code> does not have any enabled constraint associated with it, <code>DOMAIN CHECK</code> returns <code>TRUE</code>.
- If the data type conversion succeeds and domain_name has enabled constraints that are all satisfied for a given converted value, DOMAIN_CHECK returns TRUE. If any of the domain constraints are not satisfied, it returns FALSE.

MULTI-COLUMN Domains

When calling $DOMAIN_CHECK$ for multicolumn domains, the number if arguments for expr must match the number of columns in the domain. If there is a mismatch, $DOMAIN_CHECK$ raises an error

If domain D has n columns, then you should call DOMAIN_CHECK should be called with D+1 arguments, like DOMAIN_CHECK (D, arg1, ..., argn).

If D does not exist or you have no privilege to access D, then an error is raised. If all the checks return true. TRUE is returned. This means that:



- arg1 is successfully converted to the data type of column 1 in D, arg2 is successfully
 converted to the data type of column 2 in Dand so on to argn is successfully converted to
 the data type of column n in D.
- All of D's enabled constraints are all satisfied with column 1 substituted by arg1 converted to D's column 1 data type, column 2 substituted by arg2 converted to D's column 2 data type, and so on to column n substituted by argn converted to D's column n data type.

The following example creates a domain dgreater with two columns c1 and c2 of type NUMBER and a check constraint that c1 be greater than c2:

```
CREATE DOMAIN dgreater AS (c1 AS NUMBER, c2 AS NUMBER) CHECK (c1 > c2);
```

Then DOMAIN_CHECK (dgreater, 1, 2) returns FALSE because c1 is less than c2 (the check condition fails). DOMAIN_CHECK (dgreater, 2, 1) returns TRUE because because c1 is greater than c2 (the check condition passes).

Flexible Domains

When calling DOMAIN_CHECK for flexible domains, the number of arguments for *expr* must match the number of domain columns plus discriminant columns. If there is a mismatch DOMAIN_CHECK raises an error.

Checking flexible domain constraints is equivalent to checking constraints of the corresponding subdomain.

You must have the EXECUTE privilege on the flexible domain in order to use DOMAIN CHECK.

Operations that require EXECUTE privilege on a flexible domain (such as when associating columns with the flexible domain, or during DOMAIN_CHECK with the first argument the flexible domain name) require EXECUTE privilege on the sub-domains. This is because a flexible domain is translated during its creation to a multi-column domain. Therefore the following rules apply:

- Associating columns to a flex domain is equivalent to associating them to the corresponding multi-column domain.
- Checking flexible domain constraints is equivalent to checking constraints of the corresponding multi-column domain.
- Evaluating flexible domain display and order properties is equivalent to evaluating properties on the corresponding multi-column domain.

Examples

Example 1

The following example creates a strict domain of data type CHAR (3 CHAR):

```
CREATE DOMAIN three chars AS CHAR(3 CHAR) STRICT;
```

Calling DOMAIN_CHECK returns true for strings three characters or shorter. For strings four characters or more long it returns false:

```
SELECT DOMAIN_CHECK (three_chars, 'ab') two_chars,
DOMAIN_CHECK (three_chars, 'abc') three_chars,
DOMAIN CHECK (three chars, 'abcd') four chars;
```



```
TWO_CHARS THREE_CHARS FOUR_CHARS
-----
TRUE TRUE FALSE
```

The following example creates a domain dgreater with two columns c1 and c2 of type NUMBER and a check constraint that c1 be greater than c2:

```
CREATE DOMAIN dgreater AS (
  c1 AS NUMBER, c2 AS NUMBER
)
  CHECK (c1 > c2);
```

The first query passes one expression value. This raises an error because there are two columns in the domain:

```
SELECT DOMAIN_CHECK (dgreater, 1) one_expr;

ORA-11515: incorrect number of columns in domain association list
```

In the second query:

- first lower is FALSE because this fails the domain constraint
- first higher is TRUE because it passes the domain constraint
- letters is FALSE because the values cannot be converted to numbers

```
SELECT DOMAIN_CHECK (dgreater, 1, 2) first_lower,

DOMAIN_CHECK (dgreater, 2, 1) first_higher,

DOMAIN_CHECK (dgreater, 'b', 'a') letters;

FIRST_LOWER FIRST_HIGHER LETTERS

FALSE TRUE FALSE
```

Example 3

The following example creates the domain DAY_OF_WEEK with no domain constraints. All calls to DOMAIN_CHECK return true because all the input values can be converted to CHAR. It is a non-strict domain, so there is no length check.



```
DOMAIN CHECK(day of week, calendar date) nondomain column,
      DOMAIN CHECK(day of week, CAST('MON' AS day of week)) domain value,
      DOMAIN_CHECK(day_of_week, 'mon') nondomain_value
 FROM calendar dates;
DAY DOMAIN COLUMN NONDOMAIN COLUMN DOMAIN VALUE NONDOMAIN VALUE
TRUE
TRUE
               TRUE TRUE
TRUE TRUE
TRUE TRUE
TRUE TRUE
TRUE TRUE
MON TRUE
TUE TRUE
FRI TRUE
                                TRUE
                               TRUE
TRUE
mon TRUE
MON TRUE
                                              TRUE
MON TRUE
                                              TRUE
```

The following example creates the domain DAY_OF_WEEK with a constraint to ensure the values are the uppercase day name abbreviations (MON, TUE, etc.). Validating this constraint is deferred until commit, so you can insert invalid values.

Using DOMAIN_CHECK to test the values for the domain column DAY_OF_WEEK_ABBR returns TRUE for the value that conforms to the constraint (MON) and FALSE for those that do not (tue, fRI):

```
CREATE DOMAIN day of week AS CHAR(3 CHAR)
  CONSTRAINT CHECK(day of week IN ('MON', 'TUE', 'WED', 'THU', 'FRI', 'SAT', 'SUN'))
  INITIALLY DEFERRED;
CREATE TABLE calendar dates (
  calendar date DATE,
  day_of_week_abbr day_of_week
INSERT INTO calendar dates
VALUES (DATE '2023-05-01', 'MON'),
      (DATE'2023-05-02', 'tue'),
      (DATE'2023-05-05', 'fRI');
SELECT day of week abbr,
       DOMAIN CHECK(day_of_week, day_of_week_abbr) domain_column,
       DOMAIN CHECk(day of week, calendar_date) nondomain_column,
       DOMAIN_CHECK(day_of_week, CAST('MON' AS day_of_week)) domain value,
       DOMAIN_CHECK(day_of_week, 'mon') nondomain_value
  FROM calendar dates;
DAY DOMAIN_COLUMN NONDOMAIN_COLUMN DOMAIN_VALUE NONDOMAIN_VALUE
--- ------ -----
MON TRUE FALSE TRUE FALSE tue FALSE FALSE TRUE FALSE FALSE TRUE FALSE
```

Example 5

The following example creates the multicolumn domain currency with two deferred constraints:

```
CREATE DOMAIN currency AS (
amount AS NUMBER(10, 2)
currency_code AS CHAR(3 CHAR)
)

CONSTRAINT supported_currencies_c
CHECK (currency_code IN ('USD', 'GBP', 'EUR', 'JPY'))
DEFERRABLE INITIALLY DEFERRED
```

```
CONSTRAINT non_negative_amounts_c
CHECK ( amount >= 0 )
DEFERRABLE INITIALLY DEFERRED;
```

The columns AMOUNT and CURRENCY_CODE in the table ORDER_ITEMS are associated with domain currency:

```
CREATE TABLE order_items (
  order_id INTEGER,
  product_id INTEGER,
  amount NUMBER(10, 2),
  currency_code CHAR(3 CHAR),
  DOMAIN currency(amount, currency_code)
);
INSERT INTO order_items
VALUES (1, 1, 9.99, 'USD'),
  (2, 2, 1234.56, 'GBP'),
  (3, 3, -999999, 'JPY'),
  (4, 4, 3141592, 'XXX'),
  (5, 5, 2718281, '123');
```

The query makes four calls to DOMAIN CHECK:

```
SELECT order id,
        product id,
         amount,
         currency code,
         DOMAIN CHECk (currency, order id, product id) order product,
         DOMAIN CHECk(currency, amount, currency code) amount currency,
         DOMAIN_CHECk(currency, currency_code, amount) currency_amount,
         DOMAIN CHECk(currency, order id, currency code) order currency
  FROM order items;
  ORDER ID PRODUCT ID AMOUNT CUR ORDER PRODUCT AMOUNT CURRENCY CURRENCY AMOUNT
ORDER CURRENCY
______ ______

      1
      9.99 USD FALSE
      TRUE
      FALSE

      2
      1234.56 GBP FALSE
      TRUE
      FALSE

      3
      -999999 JPY FALSE
      FALSE
      FALSE

      4
      3141592 XXX FALSE
      FALSE
      FALSE

      5
      2718281 123 FALSE
      FALSE
      FALSE

                                                                                                             FALSE
                                                                                       FALSE
                                                                                                             FALSE
```

In the example above:

- ORDER_PRODUCT is FALSE for all rows because the values for PRODUCT_ID do not conform to the supported currencies c constraint.
- AMOUNT_CURRENCY is FALSE for the rows with values that violate the constraints (AMOUNT = -999999, and CURRENCY CODE = "XXX" and "123"). It is TRUE for the valid values.
- CURRENCY_AMOUNT is FALSE for all rows. For the first four rows this is because the values for
 the first argument, CURRENCY_CODE are all letters. These cannot be converted to the type of
 the first column in the domain (NUMBER), leading to a type error. For the fifth row, the
 amount (2718281) does not conform to the supported currencies c constraint.
- ORDER_CURRENCY is FALSE for the row with values that violate the constraints (CURRENCY CODE = "XXX" and "123"). It is TRUE for the valid values.

Example 6

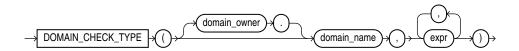


The following statement tries to validate the string "raises an error" against the non-existent domain NOT A DOMAIN. This raises an exception:

SELECT DOMAIN_CHECK(not_a_domain, 'raises an error'); ORA-11504: The domain specified does not exist or the user does not have privileges on the domain for the operation.

DOMAIN_CHECK_TYPE

Syntax



Purpose

Use <code>DOMAIN_CHECK_TYPE</code> to convert the value expression to the data type of the domain column without checking domain constraints. If you want to check constraints, you must use <code>DOMAIN_CHECK</code>.

DOMAIN_CHECK_TYPE takes the same arguments as DOMAIN_CHECK and returns TRUE if the data type of the arguments match the data types of the corresponding domain columns. If the data type match fails, it returns FALSE.



Domain Functions

- domain_name must be an identifier and can be specified using domain_owner.domain_name. If you specify it without domain_owner, it resolves first to the current user then as a public synonym. If the name cannot be resolved, an error is raised.
- If domain_name refers to a non-existent domain or one that you do not have EXECUTE privileges on, then DOMAIN CHECK will raise an error.
- If the domain column data type is STRICT, then the value is converted to the domain column's data type. For example, if the domain column data type is VARCHAR2 (100) STRICT, then the value is converted to VARCHAR2 (100). Note that the conversion will not automatically trim the input to the maximum length. If the value evaluates to 'abc' for some row and the domain data type is CHAR (2 CHAR), the conversion will fail instead of returning 'ab'.

If the domain column data type is not STRICT, then the value is converted to the most permissive variant of the domain column's data type in terms of length, scale and precision. For example, if the input value is a VARCHAR2 (30), it is converted to a VARCHAR2 (100) because it is shorter than the domain length. If the input value is a VARCHAR2 (200), it remains aVARCHAR2 (200) because this is larger than the domain length.

• If the data type conversion fails, the error is masked and <code>DOMAIN_CHECK_TYPE</code> returns <code>FALSE</code>. You can use <code>DOMAIN_CHECK_TYPE</code> to filter out values that cannot be inserted into a column of the given domain..

MULTI-COLUMN Domains

When calling DOMAIN_CHECK_TYPE for multicolumn domains, the number of arguments for expr must match the number of columns in the domain. If there is a mismatch DOMAIN_CHECK_TYPE raises an error.

Flexible Domains

When calling $DOMAIN_CHECK_TYPE$ for flexible domains, the number of arguments for expr must match the number of domain columns plus discriminant columns. If there is a mismatch DOMAIN CHECK TYPE raises an error.

Examples

Example 1

The following example creates a strict domain of data type CHAR (3 CHAR):

```
CREATE DOMAIN three chars AS CHAR(3 CHAR) STRICT;
```

Calling DOMAIN_CHECK_TYPE returns true for strings three characters or shorter. For strings four characters or more long it returns false:

Example 2

The following example creates a domain dgreater with two columns c1 and c2 of type NUMBER and a check constraint that c1 be greater than c2:

```
CREATE DOMAIN dgreater AS (
  c1 AS NUMBER, c2 AS NUMBER
)
  CHECK (c1 > c2);
```

The first query passes one expression value. This raises an error because there are two columns in the domain.

```
SELECT DOMAIN_CHECK_TYPE (dgreater, 1) one_expr;
ORA-11515: incorrect number of columns in domain association list
```

In the second query:

- first_lower and first_higher are both TRUE because the values are numbers. The
 domain constraint is not checked.
- letters is FALSE because the values cannot be converted to numbers.



```
FIRST_LOWER FIRST_HIGHER LETTERS
-----
TRUE TRUE FALSE
```

The following example creates the domain DAY_OF_WEEK with no domain constraints. All calls to DOMAIN_CHECK_TYPE return true because all the input values can be converted to CHAR. It's a non-strict domain, so there is no length check.

```
CREATE DOMAIN day of week AS CHAR(3 CHAR);
CREATE TABLE calendar dates (
  calendar date DATE,
  day of week abbr day of week
INSERT INTO calendar dates
VALUES (DATE '2023-05-01', 'MON'),
      (DATE'2023-05-02', 'tue'),
      (DATE'2023-05-05', 'fRI');
SELECT day of week abbr,
       DOMAIN CHECK TYPE (day of week, day of week abbr) domain column,
       DOMAIN CHECK TYPE (day of week, calendar date) nondomain column,
       DOMAIN CHECK TYPE(day of week, CAST('MON' AS day of week)) domain value,
       DOMAIN CHECK TYPE (day of week, 'mon') nondomain value
  FROM calendar dates;
DAY DOMAIN COLUMN NONDOMAIN COLUMN DOMAIN VALUE NONDOMAIN VALUE
MON TRUE TRUE TRUE TRUE
TUE TRUE TRUE TRUE TRUE
FRI TRUE TRUE TRUE TRUE
mon TRUE TRUE TRUE TRUE
MON TRUE TRUE TRUE TRUE
```

Example 4

The following example creates the domain DAY_OF_WEEK with a constraint to ensure the values are the uppercase day name abbreviations (MON, TUE, etc.).

Validating this constraint is deferred until commit, so you can insert invalid values.

Using DOMAIN CHECK TYPE returns TRUE for all values because they all pass the type check:

The following example creates the multicolumn domain currency with two deferred constraints:

```
CREATE DOMAIN currency AS (
  amount          AS NUMBER(10, 2)
  currency_code AS CHAR(3 CHAR)
)
CONSTRAINT supported_currencies_c
  CHECK ( currency_code IN ( 'USD', 'GBP', 'EUR', 'JPY' ) )
  DEFERRABLE INITIALLY DEFERRED
CONSTRAINT non_negative_amounts_c
  CHECK ( amount >= 0 )
  DEFERRABLE INITIALLY DEFERRED;
```

The columns amount and Currency_Code in the table <code>ORDER_ITEMS</code> are associated with domain <code>currency:</code>

```
CREATE TABLE order_items (
  order_id INTEGER,
  product_id INTEGER,
  amount NUMBER(10, 2),
  currency_code CHAR(3 CHAR),
  DOMAIN currency(amount, currency_code)
);
INSERT INTO order_items
VALUES (1, 1, 9.99, 'USD'),
  (2, 2, 1234.56, 'GBP'),
  (3, 3, -999999, 'JPY'),
  (4, 4, 3141592, 'XXX'),
  (5, 5, 2718281, '123');
```

The query makes four calls to DOMAIN CHECK TYPE:

	_					
1	1	9.99 USD	TRUE	TRUE	FALSE	TRUE
2	2	1234.56 GBP	TRUE	TRUE	FALSE	TRUE
3	3	-999999 JPY	TRUE	TRUE	FALSE	TRUE
4	4	3141592 XXX	TRUE	TRUE	FALSE	TRUE
5	5	2718281 123	TRUE	TRUE	TRUE	TRUE

In the example above:

- ORDER_PRODUCT is TRUE for all rows because the values for ORDER_ID and PRODUCT_ID can be converted to the corresponding column types in the domain (NUMBER and CHAR).
- AMOUNT_CURRENCY is TRUE for all rows because the table columns match the domain columns.
- CURRENCY_AMOUNT is FALSE for the first four rows because the values for the first argument, CURRENCY_CODE are all letters. These cannot be converted to the type of the first column in the domain (NUMBER), leading to a type error. The fifth row is TRUE because the amount (2718281) can be converted to CHAR.
- ORDER_CURRENCY is TRUE for all rows because the types for ORDER_ID and CURRENCY_CODE match the corresponding domain column types (NUMBER and CHAR).

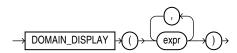
Example 6

The following statement tries to validate the string "raises an error" against the non-existent domain NOT A DOMAIN. This raises an exception:

```
SELECT DOMAIN_CHECK_TYPE(not_a_domain, 'raises an error'); ORA-11504: The domain specified does not exist or the user does not have privileges on the domain for the operation.
```

DOMAIN_DISPLAY

Syntax



Purpose

DOMAIN_DISPLAY returns expr formatted according to the domain's display expression. This returns NULL if the arguments are not associated with a domain or the domain has no display expression.

When calling DOMAIN_DISPLAY for multicolumn domains, all values of expr should be from the same domain. It returns NULL if the number of expr arguments are different from the number of domain columns or they are in a different order in the domain.

To get the display expression for a non-domain value, cast expr to the domain type. This is only possible for single column domains.



See Also:

Domain Functions

Examples

The following example creates the domain <code>DAY_OF_WEEK</code> and associates it with the column <code>CALENDAR_DATES.DAY_OF_WEEK_ABBR</code>. Passing this column to <code>DOMAIN_DISPLAY</code> returns it with the first letter of each word capitalized and all other letters in lowercase. <code>DOMAIN_DISPLAY</code> also returns this format when casting a string to the domain.

All other calls to DOMAIN DISPLAY pass non-domain values, so return NULL.

```
CREATE DOMAIN day of week AS CHAR(3 CHAR)
  DISPLAY INITCAP (day of week);
CREATE TABLE calendar dates (
  calendar date DATE,
  day of week abbr day of week
);
INSERT INTO calendar dates
VALUES (DATE '2023-05-01', 'MON'),
     (DATE'2023-05-02', 'tue'),
      (DATE'2023-05-05', 'fRI');
SELECT day of week abbr,
       DOMAIN DISPLAY(day of week abbr) domain column,
        DOMAIN DISPLAY(calendar date) nondomain column,
       DOMAIN DISPLAY (CAST ('MON' AS day of week)) domain value,
       DOMAIN DISPLAY('MON') nondomain value
  FROM calendar dates;
DAY OF WEEK ABBR DOMAIN COLUMN NONDOMAIN COLUMN DOMAIN VALUE NONDOMAIN VALUE

        Mon
        <null>
        Mon
        <null>

        Tue
        <null>
        Mon
        <null>

        Fri
        <null>
        Mon
        <null>

MON
tue
fRI
```

The following example creates the multicolumn domain CURRENCY with a display expression. The columns AMOUNT and CURRENCY_CODE in ORDER_ITEMS are associated with this domain.

In the query, only the domain_cols expression formats the columns according to the domain expression. All other calls to <code>DOMAIN_DISPLAY</code> have a mismatch between its arguments and the domain columns so return <code>NULL</code>:

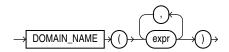
```
CREATE DOMAIN currency AS (
amount AS NUMBER(10, 2)
currency_code AS CHAR(3 CHAR)
)
DISPLAY CASE currency_code
WHEN 'USD' THEN '$'
WHEN 'GBP' THEN '£'
```



```
WHEN 'EUR' THEN '€'
 WHEN 'JPY' THEN '\'
END || TO_CHAR(amount, '999,999,999.00');
CREATE TABLE order items (
 order id INTEGER,
 product_id INTEGER, amount NUMBER(10, 2),
 currency_code CHAR(3 CHAR),
 DOMAIN currency (amount, currency code)
INSERT INTO order items
VALUES (1, 1, \frac{-}{9.99}, 'USD'),
     (2, 2, 1234.56, 'GBP'),
      (3, 3, 4321, 'EUR'),
      (4, 4, 3141592, 'JPY');
SELECT order id,
      product id,
      DOMAIN_DISPLAY(amount, currency_code) domain_cols,
      DOMAIN DISPLAY(currency_code, amount) domain_cols_wrong_order,
      DOMAIN DISPLAY(order id, product id) nondomain cols,
      DOMAIN_DISPLAY(amount) domain_cols_subset
 FROM order items;
 ORDER_ID PRODUCT_ID DOMAIN_COLS DOMAIN_COLS_WRONG_ORDER NONDOMAIN_COLS
DOMAIN COLS SUBSET
1 $ 9.99 <null>
                                                        <null>
<null>
              2 £ 1,234.56 <null>
                                                       <null>
<null>
               3 € 4,321.00 <null>
                                                       <null>
<null>
          4 ¥ 3,141,592.00 <null>
                                                      <null>
                                                                  <null>
```

DOMAIN_NAME

Syntax



Purpose

DOMAIN_NAME returns the fully qualified name of the domain associated with expr. This returns NULL if expr is associated with a domain.

When calling DOMAIN_NAME for multicolumn domains, all values of expr should be from the same domain. It returns NULL if the number of expr arguments are different from the number of domain columns or they are in a different order in the domain.

See Also:

Domain Functions

Examples

The following example creates the domain <code>DAY_OF_WEEK</code> in the schema <code>HR</code> and associates it with the column <code>HR.CALENDAR_DATES.DAY_OF_WEEK_ABBR</code>. Passing this column to <code>DOMAIN_NAME</code> returns the fully qualified name of the domain.

The query casts the string "MON" to DAY OF WEEK to get the domain name.

All other calls to DOMAIN NAME return NULL.

```
CREATE DOMAIN hr.day of week AS CHAR(3 CHAR);
CREATE TABLE hr.calendar dates (
 calendar date DATE,
 day_of_week_abbr hr.day_of_week
INSERT INTO hr.calendar dates
VALUES (DATE '2023-05-01', 'MON');
SELECT day of week abbr,
      DOMAIN NAME(day of week abbr) domain column,
      DOMAIN_NAME(calendar_date) nondomain_column,
      DOMAIN NAME(CAST('MON' AS hr.day of week)) domain value,
      DOMAIN NAME('MON') nondomain value
 FROM hr.calendar dates;
DAY OF WEEK ABBR DOMAIN COLUMN NONDOMAIN COLUMN DOMAIN VALUE NONDOMAIN VALUE
HR.DAY OF WEEK <null> HR.DAY OF WEEK <null>
MON
```

The following example creates the multicolumn domain CURRENCY in the schema CO. The columns AMOUNT and CURRENCY CODE in CO.ORDER ITEMS are associated with this domain.

In the query, the arguments for <code>DOMAIN_NAME</code> only match the domain definition for the domain_cols expression. This returns the fully qualified name of the domain. All other calls to <code>DOMAIN_NAME</code> have a mismatch between its arguments and the domain columns so return <code>NULL</code>:

```
CREATE DOMAIN co.currency AS (
amount AS NUMBER(10, 2)
currency_code AS CHAR(3 CHAR)
);

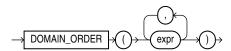
CREATE TABLE co.order_items (
order_id INTEGER,
product_id INTEGER,
amount NUMBER(10, 2),
currency_code CHAR(3 CHAR),
DOMAIN co.currency(amount, currency_code)
);
```



```
INSERT INTO co.order items
VALUES (1, 1, 9.99, 'USD');
SELECT order id,
      product id,
      DOMAIN NAME (amount, currency code) domain cols,
      DOMAIN NAME(currency code, amount) domain cols wrong order,
      DOMAIN_NAME(order_id, product_id) nondomain_cols,
      DOMAIN NAME (amount) domain cols subset
  FROM co.order items
 ORDER BY domain cols;
 ORDER ID PRODUCT_ID DOMAIN_COLS DOMAIN_COLS_WRONG_ORDER NONDOMAIN_COLS
DOMAIN COLS SUBSET
_____
                1 CO.CURRENCY <null>
                                                          <null>
                                                                           <null>
```

DOMAIN_ORDER

Syntax



Purpose

DOMAIN_ORDER returns expr formatted according to the domain's order expression. This returns NULL if the arguments are not associated with a domain or the domain has no order expression.

When calling DOMAIN_ORDER for multicolumn domains, all values of expr should be from the same domain. It returns NULL if the number expr arguments are different from the number of domain columns or they are in a different order in the domain.

To get the display expression for a non-domain value, cast expr to the domain type. This is only possible for single column domains.



Domain Functions

Examples

The following example creates the domain <code>DAY_OF_WEEK</code> and associates it with the column <code>CALENDAR_DATES.DAY_OF_WEEK_ABBR</code>. Passing this column to <code>DOMAIN_ORDER</code> returns the result of the <code>ORDER</code> expression (MON = 0, TUE = 1, etc.). Using this in the <code>ORDER</code> BY returns the rows sorted by their position in the week instead of alphabetically.

The query casts the string "MON" to DAY OF WEEK to get its sort value.

All other calls to DOMAIN ORDER pass non-domain values, so return NULL.

```
CREATE DOMAIN day_of_week AS CHAR(3 CHAR)
  ORDER CASE UPPER(day_of_week)
    WHEN 'MON' THEN 0
    WHEN 'TUE' THEN 1
    WHEN 'WED' THEN 2
    WHEN 'THU' THEN 3
    WHEN 'FRI' THEN 4
    WHEN 'SAT' THEN 5
    WHEN 'SUN' THEN 6
    ELSE 7
 END;
CREATE TABLE calendar dates (
 calendar date DATE,
 day of week abbr day of week
INSERT INTO calendar dates
VALUES (DATE '2023-05-01', 'MON'),
     (DATE'2023-05-02', 'TUE'),
     (DATE'2023-05-05', 'FRI'),
     (DATE'2023-05-08', 'mon');
SELECT day of week abbr,
      DOMAIN_ORDER(day_of_week_abbr) domain_column,
      DOMAIN ORDER(calendar_date) nondomain_column,
      DOMAIN ORDER(CAST('MON' AS day of week)) domain value,
      DOMAIN ORDER ('MON') nondomain value
 FROM calendar dates
 ORDER BY DOMAIN_ORDER(day_of_week_abbr);
DAY OF WEEK ABBR DOMAIN COLUMN NONDOMAIN COLUMN DOMAIN VALUE NONDOMAIN VALUE
0 <null>
                                             0 <null>
MON
                         0 <null>
mon
                                                    0 <null>
TUE
                          1 <null>
                                                    0 <null>
                          4 <null>
                                                     0 <null>
FRI
```

The following example creates the multicolumn domain CURRENCY with an order expression. This sorts the values by currency then value. The columns AMOUNT and CURRENCY_CODE in ORDER ITEMS are associated with this domain.

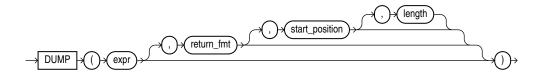
In the query, only the domain_cols expression formats the columns according to the order expression. All other calls to <code>DOMAIN_ORDER</code> have a mismatch between its arguments and the domain columns so return <code>NULL</code>:

```
CREATE DOMAIN currency AS (
  amount        AS NUMBER(10, 2)
  currency_code AS CHAR(3 CHAR)
)
ORDER currency_code || TO_CHAR(amount, '999999999.00');
CREATE TABLE order_items (
  order id        INTEGER,
```

```
product_id
             INTEGER,
 amount NUMBER(10, 2),
 currency_code CHAR(3 CHAR),
 DOMAIN currency (amount, currency code)
);
INSERT INTO order items
VALUES (1, 1, 9.99, 'USD'),
     (2, 2, 1234.56, 'USD'),
      (3, 3, 4321, 'EUR'),
      (4, 4, 3141592, 'JPY'),
      (5, 5, 2718281, 'JPY');
SELECT order id,
     product id,
      DOMAIN ORDER (amount, currency_code) domain_cols,
      DOMAIN ORDER (currency code, amount) domain cols wrong order,
      DOMAIN ORDER (order id, product id) nondomain cols,
     DOMAIN ORDER (amount) domain cols subset
 FROM order items
 ORDER BY domain cols;
 ORDER_ID PRODUCT_ID DOMAIN_COLS DOMAIN_COLS_WRONG_ORDER NONDOMAIN_COLS
DOMAIN COLS SUBSET
 ______
          3 EUR 4321.00 <null>
                                                        <null>
<null>
                 5 JPY 2718281.00 <null>
                                                        <null>
<null>
                 4 JPY 3141592.00 <null>
                                                        <null>
<null>
                 1 USD
                             9.99 <null>
                                                        <null>
<null>
                 2 USD 1234.56 <null>
                                                        <null>
                                                                   <null>
```

DUMP

Syntax



Purpose

DUMP returns a VARCHAR2 value containing the data type code, length in bytes, and internal representation of *expr*. The returned result is always in the database character set. For the data type corresponding to each code, see Table 2-1.

The argument <code>return_fmt</code> specifies the format of the return value and can have any of the following values:

8 returns result in octal notation.

- 10 returns result in decimal notation.
- 16 returns result in hexadecimal notation.
- 17 returns each byte printed as a character if and only if it can be interpreted as a printable character in the character set of the compiler—typically ASCII or EBCDIC. Some ASCII control characters may be printed in the form ^X as well. Otherwise the character is printed in hexadecimal notation. All NLS parameters are ignored. Do not depend on any particular output format for DUMP with return fmt 17.

By default, the return value contains no character set information. To retrieve the character set name of *expr*, add 1000 to any of the preceding format values. For example, a *return_fmt* of 1008 returns the result in octal and provides the character set name of *expr*.

The arguments <code>start_position</code> and <code>length</code> combine to determine which portion of the internal representation to return. The default is to return the entire internal representation in decimal notation.

If expr is null, then this function returns NULL.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.

See Also:

- Data Type Comparison Rules for more information
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of DUMP

Examples

The following examples show how to extract dump information from a string expression and a column:

```
SELECT DUMP ('abc', 1016)
 FROM DUAL;
DUMP('ABC', 1016)
Typ=96 Len=3 CharacterSet=WE8DEC: 61,62,63
SELECT DUMP(last_name, 8, 3, 2) "OCTAL"
 FROM employees
 WHERE last name = 'Hunold'
 ORDER BY employee_id;
OCTAL
______
Typ=1 Len=6: 156,157
SELECT DUMP(last name, 10, 3, 2) "ASCII"
 FROM employees
 WHERE last_name = 'Hunold'
 ORDER BY employee id;
ASCII
```

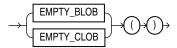


Typ=1 Len=6: 110,111

EMPTY_BLOB, EMPTY_CLOB

Syntax

empty_LOB::=



Purpose

EMPTY_BLOB and EMPTY_CLOB return an empty LOB locator that can be used to initialize a LOB variable or, in an INSERT or UPDATE statement, to initialize a LOB column or attribute to EMPTY. EMPTY means that the LOB is initialized, but not populated with data.



An empty LOB is not the same as a null LOB, and an empty CLOB is not the same as a LOB containing a string of 0 length. For more information, see *Oracle Database SecureFiles and Large Objects Developer's Guide*.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the return value of EMPTY_CLOB

Restriction on LOB Locators

You cannot use the locator returned from this function as a parameter to the <code>DBMS_LOB</code> package or the OCI.

Examples

The following example initializes the <code>ad_photo</code> column of the sample <code>pm.print_media</code> table to <code>EMPTY:</code>

```
UPDATE print_media
  SET ad_photo = EMPTY_BLOB();
```



EVERY

Syntax



Purpose

EVERY returns 'TRUE' if the <code>boolean_expr</code> evaluates to true for every row that qualifies. Otherwise it returns 'FALSE'. You can use it as an aggregate or analytic function. It is the same as the function <code>BOOLEAN AND AGG</code>.

EXISTSNODE



The EXISTSNODE function is deprecated. It is still supported for backward compatibility. However, Oracle recommends that you use the XMLEXISTS function instead. See XMLEXISTS for more information.

Syntax



Purpose

EXISTSNODE determines whether traversal of an XML document using a specified path results in any nodes. It takes as arguments the XMLType instance containing an XML document and a VARCHAR2 XPath string designating a path. The optional <code>namespace_string</code> must resolve to a VARCHAR2 value that specifies a default mapping or namespace mapping for prefixes, which Oracle Database uses when evaluating the XPath expression(s).

The <code>namespace_string</code> argument defaults to the namespace of the root element. If you refer to any subelement in <code>Xpath_string</code>, then you must specify <code>namespace_string</code>, and you must specify the "who" prefix in both of these arguments.



Using XML in SQL Statements for examples that specify <code>namespace_string</code> and use the "who" prefix.

The return value is NUMBER:

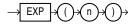
- 0 if no nodes remain after applying the XPath traversal on the document
- 1 if any nodes remain

Examples

The following example tests for the existence of the /Warehouse/Dock node in the XML path of the warehouse spec column of the sample table oe.warehouses:

EXP

Syntax



Purpose

EXP returns e raised to the nth power, where e = 2.71828183... The function returns a value of the same type as the argument.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



Table 2-9 for more information on implicit conversion

Examples

The following example returns e to the 4th power:

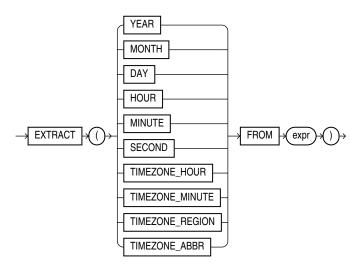
```
SELECT EXP(4) "e to the 4th power"
FROM DUAL;
e to the 4th power
-----
54.59815
```



EXTRACT (datetime)

Syntax

extract datetime::=



Purpose

EXTRACT extracts and returns the value of a specified datetime field from a datetime or interval expression. The expr can be any expression that evaluates to a datetime or interval data type compatible with the requested field:

- If YEAR or MONTH is requested, then expr must evaluate to an expression of data type DATE, TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIME ZONE, Or INTERVAL YEAR TO MONTH.
- If DAY is requested, then *expr* must evaluate to an expression of data type DATE, TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIME ZONE, **or** INTERVAL DAY TO SECOND.
- If HOUR, MINUTE, or SECOND is requested, then expr must evaluate to an expression of data
 type TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIME ZONE, or INTERVAL
 DAY TO SECOND. DATE is not valid here, because Oracle Database treats it as ANSI DATE
 data type, which has no time fields.
- If TIMEZONE_HOUR, TIMEZONE_MINUTE, TIMEZONE_ABBR, TIMEZONE_REGION, or TIMEZONE_OFFSET is requested, then expr must evaluate to an expression of data type TIMESTAMP WITH TIME ZONE or TIMESTAMP WITH LOCAL TIME ZONE.

EXTRACT interprets *expr* as an ANSI datetime data type. For example, EXTRACT treats DATE not as legacy Oracle DATE but as ANSI DATE, without time elements. Therefore, you can extract only YEAR, MONTH, and DAY from a DATE value. Likewise, you can extract TIMEZONE_HOUR and TIMEZONE MINUTE only from the TIMESTAMP WITH TIME ZONE data type.

When you specify <code>TIMEZONE_REGION</code> or <code>TIMEZONE_ABBR</code> (abbreviation), the value returned is a <code>VARCHAR2</code> string containing the appropriate time zone region name or abbreviation. When you specify any of the other datetime fields, the value returned is an integer value of <code>NUMBER</code> data type representing the datetime value in the Gregorian calendar. When extracting from a



datetime with a time zone value, the value returned is in UTC. For a listing of time zone region names and their corresponding abbreviations, query the V\$TIMEZONE_NAMES dynamic performance view.

This function can be very useful for manipulating datetime field values in very large tables, as shown in the first example below.



Time zone region names are needed by the daylight saving feature. These names are stored in two types of time zone files: one large and one small. One of these files is the default file, depending on your environment and the release of Oracle Database you are using. For more information regarding time zone files and names, see *Oracle Database Globalization Support Guide*.

Some combinations of datetime field and datetime or interval value expression result in ambiguity. In these cases, Oracle Database returns UNKNOWN (see the examples that follow for additional information).

See Also:

- Oracle Database Globalization Support Guide for a complete listing of the time zone region names in both files
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of EXTRACT
- Datetime/Interval Arithmetic for a description of datetime_value_expr and interval_value_expr
- Oracle Database Reference for information on the dynamic performance views

Examples

The following example returns from the oe.orders table the number of orders placed in each month:

```
SELECT EXTRACT (month FROM order_date) "Month", COUNT (order_date) "No. of Orders"
 FROM orders
 GROUP BY EXTRACT (month FROM order date)
 ORDER BY "No. of Orders" DESC, "Month";
    Month No. of Orders
        11
         6
                      14
         7
                      14
         3
                      11
         5
                      10
         2
         9
                       9
                       7
         8
        10
```



1 5 12 4

12 rows selected.

The following example returns the year 1998.

The following example selects from the sample table hr.employees all employees who were hired after 2007:

The following example results in ambiguity, so Oracle returns UNKNOWN:

The ambiguity arises because the time zone numerical offset is provided in the expression, and that numerical offset may map to more than one time zone region name.

EXTRACT (XML)



The EXTRACT (XML) function is deprecated. It is still supported for backward compatibility. However, Oracle recommends that you use the XMLQUERY function instead. See XMLQUERY for more information.

Syntax

extract_xml::=



Purpose

EXTRACT (XML) is similar to the EXISTSNODE function. It applies a VARCHAR2 XPath string and returns an XMLType instance containing an XML fragment. You can specify an absolute XPath_string with an initial slash or a relative XPath_string by omitting the initial slash. If you omit the initial slash, then the context of the relative path defaults to the root node. The optional namespace_string is required if the XML you are handling uses a namespace prefix. This argument must resolve to a VARCHAR2 value that specifies a default mapping or namespace mapping for prefixes, which Oracle Database uses when evaluating the XPath expression(s).

Examples

The following example extracts the value of the /Warehouse/Dock node of the XML path of the warehouse spec column in the sample table oe.warehouses:

```
SELECT warehouse_name,

EXTRACT(warehouse_spec, '/Warehouse/Docks') "Number of Docks"

FROM warehouses

WHERE warehouse_spec IS NOT NULL

ORDER BY warehouse_name;

WAREHOUSE_NAME

Number of Docks

New Jersey
San Francisco

Seattle, Washington

Southlake, Texas

Vocks>2</Docks>

Vocks>2</Docks>
Southlake, Texas
```

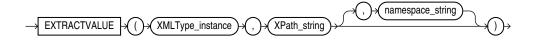
Compare this example with the example for EXTRACTVALUE, which returns the scalar value of the XML fragment.

EXTRACTVALUE



The EXTRACTVALUE function is deprecated. It is still supported for backward compatibility. However, Oracle recommends that you use the XMLTABLE function, or the XMLCAST and XMLQUERY functions instead. See XMLTABLE, XMLCAST, and XMLQUERY for more information.

Syntax



The EXTRACTVALUE function takes as arguments an XMLType instance and an XPath expression and returns a scalar value of the resultant node. The result must be a single node and be either a text node, attribute, or element. If the result is an element, then the element must have a single text node as its child, and it is this value that the function returns. You can specify an absolute XPath_string with an initial slash or a relative XPath_string by omitting the initial slash. If you omit the initial slash, the context of the relative path defaults to the root node.

If the specified XPath points to a node with more than one child, or if the node pointed to has a non-text node child, then Oracle returns an error. The optional <code>namespace_string</code> must resolve to a <code>VARCHAR2</code> value that specifies a default mapping or namespace mapping for prefixes, which Oracle uses when evaluating the XPath expression(s).

For documents based on XML schemas, if Oracle can infer the type of the return value, then a scalar value of the appropriate type is returned. Otherwise, the result is of type VARCHAR2. For documents that are not based on XML schemas, the return type is always VARCHAR2.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of <code>EXTRACTVALUE</code>

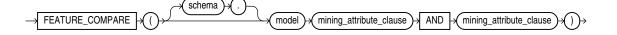
Examples

The following example takes as input the same arguments as the example for EXTRACT (XML). Instead of returning an XML fragment, as does the EXTRACT function, it returns the scalar value of the XML fragment:

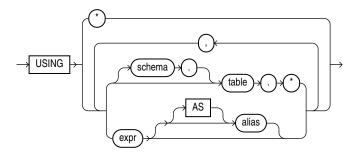
FEATURE_COMPARE

Syntax

feature_compare::=



mining_attribute_clause::=



Purpose

The FEATURE_COMPARE function uses a Feature Extraction model to compare two different documents, including short ones such as keyword phrases or two attribute lists, for similarity or dissimilarity. The FEATURE_COMPARE function can be used with Feature Extraction algorithms such as Singular Value Decomposition (SVD), Principal Component Analysis PCA), Non-Negative Matrix Factorization (NMF), and Explicit Semantic Analysis (ESA). This function is applicable not only to documents, but also to numeric and categorical data.

The input to the <code>FEATURE_COMPARE</code> function is a single feature model built using the Feature Extraction algorithms of Oracle Machine Learning for SQL, such as NMF, SVD, and ESA. The double <code>USING</code> clause provides a mechanism to compare two different documents or constant keyword phrases, or any combination of the two, for similarity or dissimilarity using the extracted features in the model.

The syntax of the FEATURE_COMPARE function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining attribute clause

The <code>mining_attribute_clause</code> identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The <code>mining_attribute_clause</code> behaves as described for the <code>PREDICTION</code> function. See <code>mining_attribute_clause</code>.

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring
- Oracle Machine Learning for SQL Concepts for information about clustering

Note:

The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Examples

An ESA model is built against a 2005 Wiki dataset rendering over 200,000 features. The documents are mined as text and the document titles are considered as the Feature IDs.

The examples show the FEATURE_COMPARE function with the ESA algorithm, which compares a similar set of texts and then a dissimilar set of texts.

Similar texts

SELECT 1-FEATURE_COMPARE(esa_wiki_mod USING 'There are several PGA tour golfers from South Africa' text AND USING 'Nick Price won the 2002 Mastercard Colonial Open' text) similarity FROM DUAL;

SIMILARITY
----.258

The output metric shows the results of a distance calculation. Therefore, a smaller number represents more similar texts. So 1 minus the distance in the queries represents a document similarity metric.

Dissimilar texts

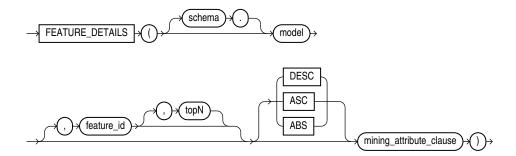
SELECT 1-FEATURE_COMPARE(esa_wiki_mod USING 'There are several PGA tour golfers from South Africa' text AND USING 'John Elway played quarterback for the Denver Broncos' text) similarity FROM DUAL;

SIMILARITY

FEATURE_DETAILS

Syntax

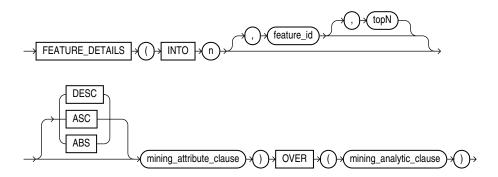
feature_details::=



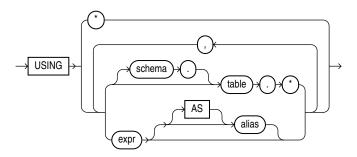


Analytic Syntax

feature_details_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions" for information on the syntax, semantics, and restrictions of $mining\ analytic\ clause$

Purpose

FEATURE_DETAILS returns feature details for each row in the selection. The return value is an XML string that describes the attributes of the highest value feature or the specified feature id.

topN

If you specify a value for topN, the function returns the N attributes that most influence the feature value. If you do not specify topN, the function returns the 5 most influential attributes.



DESC, ASC, or ABS

The returned attributes are ordered by weight. The weight of an attribute expresses its positive or negative impact on the value of the feature. A positive weight indicates a higher feature value. A negative weight indicates a lower feature value.

By default, FEATURE_DETAILS returns the attributes with the highest positive weight (DESC). If you specify ASC, the attributes with the highest negative weight are returned. If you specify ABS, the attributes with the greatest weight, whether negative or positive, are returned. The results are ordered by absolute value from highest to lowest. Attributes with a zero weight are not included in the output.

Syntax Choice

FEATURE_DETAILS can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- Syntax Use the first syntax to score the data with a pre-defined model. Supply the name of a feature extraction model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of features to extract, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order by clause. (See "analytic_clause::=".)

The syntax of the FEATURE_DETAILS function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining attribute clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining_attribute_clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about feature extraction.

Note:

The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.



Example

This example uses the feature extraction model nmf_sh_sample to score the data. The query returns the three features that best represent customer 100002 and the attributes that most affect those features

```
SELECT S.feature id fid, value val,
      FEATURE DETAILS(nmf sh sample, S.feature id, 5 using T.*) det
   FROM
     (SELECT v.*, FEATURE SET(nmf sh sample, 3 USING *) fset
        FROM mining data apply v v
        WHERE cust id = 100002) T,
   TABLE (T.fset) S
ORDER BY 2 DESC;
FID VAL DET
  5 3.492 <Details algorithm="Non-Negative Matrix Factorization" feature="5">
            <attribute name="BULK PACK DISKETTES" actualValue="1" weight=".077" rank="1"/>
            <Attribute name="OCCUPATION" actualValue="Prof." weight=".062" rank="2"/>
            <Attribute name="BOOKKEEPING APPLICATION" actualValue="1" weight=".001" rank="3"/>
            <Attribute name="OS DOC SET KANJI" actualValue="0" weight="0" rank="4"/>
            <Attribute name="YRS_RESIDENCE" actualValue="4" weight="0" rank="5"/>
            </Details>
   3 1.928 <Details algorithm="Non-Negative Matrix Factorization" feature="3">
             <Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".239" rank="1"/>
             <a href="CUST INCOME LEVEL" actualValue="L: 300\,000 and above"</a>
             weight=".051" rank="2"/>
             <Attribute name="FLAT PANEL MONITOR" actualValue="1" weight=".02" rank="3"/>
            <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".006" rank="4"/>
            <a href="AGE" actualValue="41" weight=".004" rank="5"/>
             </Details>
      .816 <Details algorithm="Non-Negative Matrix Factorization" feature="8">
            <Attribute name="EDUCATION" actualValue="Bach." weight=".211" rank="1"/>
            <Attribute name="CUST MARITAL STATUS" actualValue="NeverM" weight=".143" rank="2"/>
            <a tribute name="FLAT PANEL MONITOR" actualValue="1" weight=".137" rank="3"/>
             <Attribute name="CUST GENDER" actualValue="F" weight=".044" rank="4"/>
             <Attribute name="BULK PACK DISKETTES" actualValue="1" weight=".032" rank="5"/>
             </Details>
```

Analytic Example

This example dynamically maps customer attributes into six features and returns the feature mapping for customer 100001.

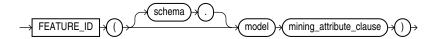
```
SELECT feature_id, value
    SELECT cust id, feature set(INTO 6 USING *) OVER () fset
      FROM mining_data_apply_v),
 TABLE (fset)
 WHERE cust id = 100001
 ORDER BY feature id;
           VALUE
FEATURE ID
           2.670
        1
        2
             .000
        3
           1.792
            .000
        4
        5
             .000
        6 3.379
```



FEATURE_ID

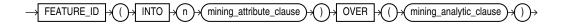
Syntax

feature_id::=

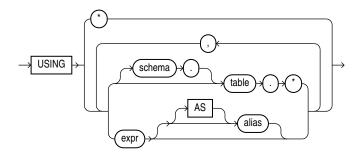


Analytic Syntax

feature_id_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions " for information on the syntax, semantics, and restrictions of mining_analytic_clause

Purpose

FEATURE_ID returns the identifier of the highest value feature for each row in the selection. The feature identifier is returned as an Oracle NUMBER.

Syntax Choice

FEATURE_ID can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- Syntax Use the first syntax to score the data with a pre-defined model. Supply the name of a feature extraction model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of features to extract, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order by clause. (See "analytic_clause::=".)

The syntax of the FEATURE_ID function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining attribute clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining_attribute_clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about feature extraction.

Note:

The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

This example lists the features and corresponding count of customers in a data set.

```
SELECT FEATURE_ID(nmf_sh_sample USING *) AS feat, COUNT(*) AS cnt
FROM nmf_sh_sample_apply_prepared
GROUP BY FEATURE_ID(nmf_sh_sample USING *)
ORDER BY cnt DESC, feat DESC;
```

CNT	FEAT
1443	7
49	2
6	3

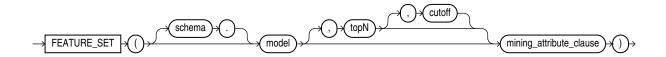


5 1 L 1

FEATURE_SET

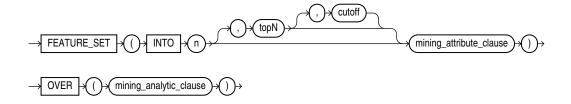
Syntax

feature_set::=

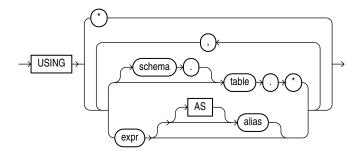


Analytic Syntax

feature_set_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions" for information on the syntax, semantics, and restrictions of $mining_analytic_clause$

Purpose

FEATURE_SET returns a set of feature ID and feature value pairs for each row in the selection. The return value is a varray of objects with field names FEATURE_ID and VALUE. The data type of both fields is NUMBER.

topN and cutoff

You can specify topN and cutoff to limit the number of features returned by the function. By default, both topN and cutoff are null and all features are returned.

- topN is the N highest value features. If multiple features have the Nth value, then the function chooses one of them.
- *cutoff* is a value threshold. Only features that are greater than or equal to *cutoff* are returned. To filter by *cutoff* only, specify NULL for *topN*.

To return up to N features that are greater than or equal to cutoff, specify both topN and cutoff.

Syntax Choice

FEATURE_SET can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- **Syntax** Use the first syntax to score the data with a pre-defined model. Supply the name of a feature extraction model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of features to extract, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order_by_clause. (See "analytic_clause::=".)

The syntax of the FEATURE_SET function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining attribute clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about feature extraction.





The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

This example lists the top features corresponding to a given customer record and determines the top attributes for each feature (based on coefficient > 0.25).

```
WITH
feat tab AS (
SELECT F.feature_id fid,
      A.attribute name attr,
      TO CHAR (A.attribute value) val,
      A.coefficient coeff
  FROM TABLE (DBMS DATA MINING.GET MODEL DETAILS NMF('nmf sh sample')) F,
      TABLE (F.attribute set) A
 WHERE A.coefficient > 0.25
),
feat AS (
SELECT fid,
      CAST(COLLECT(Featattr(attr, val, coeff))
       AS Featattrs) f attrs
 FROM feat tab
GROUP BY fid
),
cust 10 features AS (
SELECT T.cust id, S.feature id, S.value
 FROM (SELECT cust id, FEATURE SET(nmf sh sample, 10 USING *) pset
         FROM nmf sh sample apply prepared
        WHERE cust id = 100002) T,
      TABLE (T.pset) S
SELECT A.value, A.feature id fid,
     B.attr, B.val, B.coeff
 FROM cust 10 features A,
      (SELECT T.fid, F.*
         FROM feat T,
            TABLE (T.f attrs) F) B
 WHERE A.feature id = B.\overline{f}id
ORDER BY A. value DESC, A. feature id ASC, coeff DESC, attr ASC, val ASC;
  VALUE FID ATTR
                                     VAL
                                                               COEFF
 6.8409 7 YRS_RESIDENCE
                                                              1.3879
  6.8409 7 BOOKKEEPING APPLICATION
                                                              .4388
 6.8409 7 CUST_GENDER M
 6.8409 7 COUNTRY NAME
                            United States of America .2848
  6.4975 3 YRS RESIDENCE
 6.4975 3 BOOKKEEPING APPLICATION
 6.4975 3 COUNTRY NAME United States of America .2927
 6.4886 2 YRS RESIDENCE
                                                              1.3285
 6.4886 2 CUST GENDER
                                                               .2819
                                    M
 6.4886 2 PRINTER_SUPPLIES
                                                               .2704
 6.3953 4 YRS RESIDENCE
                                                              1.2931
 5.9640 6 YRS RESIDENCE
                                                              1.1585
  5.9640 6 HOME THEATER_PACKAGE
                                                               .2576
```

5.2424	5 YRS RESIDENCE	1.0067
2.4714	8 YRS_RESIDENCE	.3297
2.3559	1 YRS_RESIDENCE	.2768
2.3559	1 FLAT PANEL MONITOR	.2593

FEATURE_VALUE

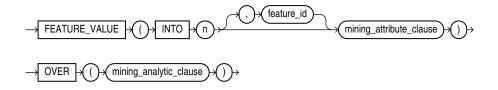
Syntax

feature_value::=

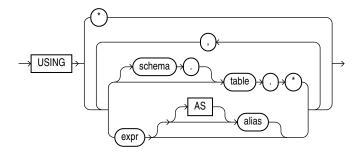


Analytic Syntax

feature_value_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions" for information on the syntax, semantics, and restrictions of $mining_analytic_clause$

Purpose

FEATURE_VALUE returns a feature value for each row in the selection. The value refers to the highest value feature or to the specified <code>feature_id</code>. The feature value is returned as <code>BINARY DOUBLE</code>.

Syntax Choice

FEATURE_VALUE can score the data in one of two ways: It can apply a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- Syntax Use the first syntax to score the data with a pre-defined model. Supply the name of a feature extraction model.
- Analytic Syntax Use the analytic syntax to score the data without a pre-defined model. Include INTO n, where n is the number of features to extract, and mining_analytic_clause, which specifies if the data should be partitioned for multiple model builds. The mining_analytic_clause supports a query_partition_clause and an order by clause. (See "analytic_clause::=".)

The syntax of the FEATURE_VALUE function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, this data is also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining_attribute_clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about feature extraction.

Note:

The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

The following example lists the customers that correspond to feature 3, ordered by match quality.

```
SELECT *
  FROM (SELECT cust_id, FEATURE_VALUE(nmf_sh_sample, 3 USING *) match_quality
        FROM nmf_sh_sample_apply_prepared
        ORDER BY match quality DESC)
```



WHERE ROWNUM < 11;

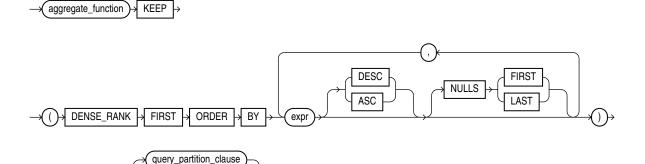
CUST_ID	MATCH_QUALITY
100210	19.4101627
100962	15.2482251
101151	14.5685197
101499	14.4186292
100363	14.4037396
100372	14.3335148
100982	14.1716545
101039	14.1079914
100759	14.0913761
100953	14.0799737

FIRST

OVER

Syntax

first::=





"Analytic Functions" for information on syntax, semantics, and restrictions of the $\mbox{ORDER BY}$ clause and \mbox{OVER} clause

Purpose

FIRST and LAST are very similar functions. Both are aggregate and analytic functions that operate on a set of values from a set of rows that rank as the FIRST or LAST with respect to a given sorting specification. If only one row ranks as FIRST or LAST, then the aggregate operates on the set with only one element.

If you omit the OVER clause, then the FIRST and LAST functions are treated as aggregate functions. You can use these functions as analytic functions by specifying the OVER clause. The $query_partition_clause$ is the only part of the OVER clause valid with these functions. If you include the OVER clause but omit the $query_partition_clause$, then the function is treated as an analytic function, but the window defined for analysis is the entire table.



These functions take as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.

When you need a value from the first or last row of a sorted group, but the needed value is not the sort key, the FIRST and LAST functions eliminate the need for self-joins or views and enable better performance.

- The aggregate_function argument is any one of the MIN, MAX, SUM, AVG, COUNT, VARIANCE, or STDDEV functions. It operates on values from the rows that rank either FIRST or LAST. If only one row ranks as FIRST or LAST, then the aggregate operates on a singleton (nonaggregate) set.
- The KEEP keyword is for semantic clarity. It qualifies aggregate_function, indicating that only the FIRST or LAST values of aggregate function will be returned.
- DENSE_RANK FIRST or DENSE_RANK LAST indicates that Oracle Database will aggregate over
 only those rows with the minimum (FIRST) or the maximum (LAST) dense rank (also called
 olympic rank).



Table 2-9 for more information on implicit conversion and LAST

Aggregate Example

The following example returns, within each department of the sample table hr.employees, the minimum salary among the employees who make the lowest commission and the maximum salary among the employees who make the highest commission:

```
SELECT department_id,

MIN(salary) KEEP (DENSE_RANK FIRST ORDER BY commission_pct) "Worst",

MAX(salary) KEEP (DENSE_RANK LAST ORDER BY commission_pct) "Best"

FROM employees

GROUP BY department_id

ORDER BY department id;
```

Worst	Best
4400	4400
6000	13000
2500	11000
6500	6500
2100	8200
4200	9000
10000	10000
6100	14000
17000	24000
6900	12008
8300	12008
7000	7000
	4400 6000 2500 6500 2100 4200 10000 6100 17000 6900 8300

Analytic Example

The next example makes the same calculation as the previous example but returns the result for each employee within the department:



```
SELECT last_name, department_id, salary,

MIN(salary) KEEP (DENSE_RANK FIRST ORDER BY commission_pct)

OVER (PARTITION BY department_id) "Worst",

MAX(salary) KEEP (DENSE_RANK LAST ORDER BY commission_pct)

OVER (PARTITION BY department_id) "Best"

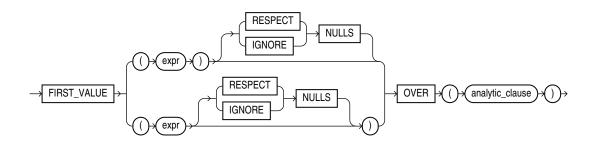
FROM employees

ORDER BY department id, salary, last name;
```

LAST_NAME	DEPARTMENT_ID	SALARY	Worst	Best
Whalen Fay	10 20	4400	4400	4400 13000
Hartstein	20	13000	6000	13000
Gietz	110	8300	8300	12008
Higgins Grant	110	12008 7000	8300 7000	12008 7000

FIRST_VALUE

Syntax



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions, including valid forms of expr

Purpose

FIRST_VALUE is an analytic function. It returns the first value in an ordered set of values. If the first value in the set is null, then the function returns NULL unless you specify IGNORE NULLS. This setting is useful for data densification.

Note:

The two forms of this syntax have the same behavior. The top branch is the ANSI format, which Oracle recommends for ANSI compatibility.

{RESPECT | IGNORE} NULLS determines whether null values of expr are included in or eliminated from the calculation. The default is RESPECT NULLS. If you specify IGNORE NULLS, then

FIRST_VALUE returns the first non-null value in the set, or NULL if all values are null. Refer to "Using Partitioned Outer Joins: Examples" for an example of data densification.

You cannot nest analytic functions by using FIRST_VALUE or any other analytic function for *expr*. However, you can use other built-in function expressions for *expr*. Refer to "About SQL Expressions" for information on valid forms of *expr*.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the return value of ${\tt FIRST_VALUE}$ when it is a character value

Examples

The following example selects, for each employee in Department 90, the name of the employee with the lowest salary.

```
SELECT employee_id, last_name, salary, hire_date,
    FIRST_VALUE(last_name)
    OVER (ORDER BY salary ASC ROWS UNBOUNDED PRECEDING) AS fv
FROM (SELECT * FROM employees
    WHERE department_id = 90
    ORDER BY hire_date);
```

EMPLOYEE_ID	LAST_NAME	SALARY	HIRE_DATE	FV
102	De Haan	17000	13-JAN-01	De Haan
101	Kochhar	17000	21-SEP-05	De Haan
100	King	24000	17-JUN-03	De Haan

The example illustrates the nondeterministic nature of the <code>FIRST_VALUE</code> function. Kochhar and DeHaan have the same salary, so are in adjacent rows. Kochhar appears first because the rows returned by the subquery are ordered by <code>hire_date</code>. However, if the rows returned by the subquery are ordered by <code>hire_date</code> in descending order, as in the next example, then the function returns a different value:

The following two examples show how to make the FIRST_VALUE function deterministic by ordering on a unique key. By ordering within the function by both salary and the unique key employee id, you can ensure the same result regardless of the ordering in the subquery.

```
SELECT employee_id, last_name, salary, hire_date,
    FIRST_VALUE(last_name)
```



```
OVER (ORDER BY salary ASC, employee id ROWS UNBOUNDED PRECEDING) AS fv
 FROM (SELECT * FROM employees
        WHERE department_id = 90
        ORDER BY hire date);
EMPLOYEE ID LAST NAME
                                 SALARY HIRE DATE FV
17000 21-SEP-05 Kochhar
      101 Kochhar
      102 De Haan
                                  17000 13-JAN-01 Kochhar
      100 King
                                   24000 17-JUN-03 Kochhar
SELECT employee_id, last_name, salary, hire_date,
    FIRST VALUE(last name)
      OVER (ORDER BY salary ASC, employee id ROWS UNBOUNDED PRECEDING) AS fv
  FROM (SELECT * FROM employees
        WHERE department id = 90
        ORDER BY hire date DESC);
EMPLOYEE ID LAST NAME
                                SALARY HIRE DATE FV
101 Kochhar
                                17000 21-SEP-05 Kochhar
      102 De Haan
                                  17000 13-JAN-01 Kochhar
      100 King
                                  24000 17-JUN-03 Kochhar
```

The following two examples show that the FIRST_VALUE function is deterministic when you use a logical offset (RANGE instead of ROWS). When duplicates are found for the ORDER BY expression, the FIRST VALUE is the lowest value of expr:

```
SELECT employee_id, last_name, salary, hire_date,
     FIRST VALUE(last name)
       OVER (ORDER BY salary ASC RANGE UNBOUNDED PRECEDING) AS fv
 FROM (SELECT * FROM employees
       WHERE department id = 90
        ORDER BY hire date);
EMPLOYEE ID LAST NAME
                                 SALARY HIRE DATE FV
17000 13-JAN-01 De Haan
      102 De Haan
                                   17000 21-SEP-05 De Haan
      101 Kochhar
      100 King
                                   24000 17-JUN-03 De Haan
SELECT employee id, last name, salary, hire date,
     FIRST_VALUE(last_name)
      OVER (ORDER BY salary ASC RANGE UNBOUNDED PRECEDING) AS fv
 FROM (SELECT * FROM employees
       WHERE department id = 90
       ORDER BY hire date DESC);
EMPLOYEE ID LAST NAME
                                 SALARY HIRE DATE FV
102 De Haan
                                  17000 13-JAN-01 De Haan
      101 Kochhar
                                   17000 21-SEP-05 De Haan
      100 King
                                   24000 17-JUN-03 De Haan
```

FLOOR (datetime)

Syntax



Purpose

FLOOR (datetime) returns the date or the timestamp rounded down to the unit specified by the second argument fmt, the format model. This function is not sensitive to the NLS_CALENDAR session parameter. It operates according to the rules of the Gregorian calendar. The value returned is always of data type DATE, even if you specify a different datetime data type for the first argument. If you do not specify the second argument, the default format model 'DD' is used.

The FLOOR and TRUNC functions are synonymous for dates and timestamps.



Refer to CEIL, FLOOR, ROUND, and TRUNC Date Functions for the permitted format models to use in fmt.

Examples

For these examples NLS DATE FORMAT is set:

FLOOR (interval)

Syntax





Purpose

FLOOR (interval) returns the interval rounded down to the unit specified by the second argument fmt, the format model .

The result of <code>FLOOR(interval)</code> is never larger than <code>interval</code>. The result precision for year and day is the input precision for year plus one and day plus one, since <code>FLOOR(interval)</code> can have overflow. If an interval already has the maximum precision for year and day, the statement compiles but errors at runtime.

For INTERVAL YEAR TO MONTH, fmt can only be year. The default fmt is year.

For INTERVAL DAY TO SECOND, fmt can be day, hour and minute. The default fmt is day. Note that fmt does not support second.

FLOOR (interval) supports the format models of ROUND and TRUNC.

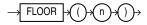


Refer to CEIL, FLOOR, ROUND, and TRUNC Date Functions for the permitted format models to use in fmt.

Examples

FLOOR (number)

Syntax





Purpose

FLOOR returns the largest integer equal to or less than n. The number n can always be written as the sum of an integer k and a positive fraction f such that $0 \le f \le 1$ and n = k + f. The value of FLOOR is the integer k. Thus, the value of FLOOR is n itself if and only if n is precisely an integer.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.



Table 2-9 for more information on implicit conversion and CEIL (number)

Examples

The following example returns the largest integer equal to or less than 15.7:

FROM TZ

Syntax



Purpose

FROM_TZ converts a timestamp value and a time zone to a TIMESTAMP WITH TIME ZONE value. $time_zone_value$ is a character string in the format 'TZH:TZM' or a character expression that returns a string in TZR with optional TZD format.

Examples

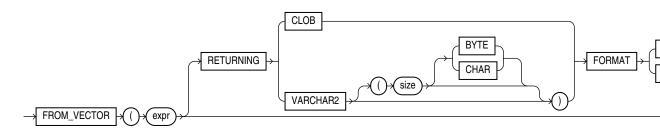
The following example returns a timestamp value to TIMESTAMP WITH TIME ZONE:



FROM_VECTOR

FROM VECTOR takes a vector as input and returns a string of type VARCHAR2 or CLOB as output.

Syntax



Purpose

FROM_VECTOR optionally takes a RETURNING clause to specify the data type of the returned value.

If VARCHAR2 is specified without size, the size of the returned value size is 32767.

You can optionally specify the text format of the output in the FORMAT clause, using the tokens SPARSE or DENSE. Note that the input vector storage format does not need to match the specified output format.

There is no support to convert to CHAR, NCHAR, and NVARCHAR2.

FROM VECTOR is synonymous with VECTOR SERIALIZE.

Parameters

expr must evaluate to a vector. The function returns NULL if expr is NULL.

Examples

```
SELECT FROM_VECTOR(TO_VECTOR('[1, 2, 3]'));

FROM_VECTOR(TO_VECTOR('[1,2,3]'))

[1.0E+000,2.0E+000,3.0E+000]

1 row selected.

SELECT FROM_VECTOR(TO_VECTOR('[1.1, 2.2, 3.3]', 3, FLOAT32));

FROM_VECTOR(TO_VECTOR('[1.1,2.2,3.3]',3,FLOAT32))

[1.10000002E+000,2.20000005E+000,3.29999995E+000]

1 row selected.

SELECT FROM_VECTOR( TO_VECTOR('[1.1, 2.2, 3.3]', 3, FLOAT32) RETURNING VARCHAR2(1000));
```

```
FROM_VECTOR(TO_VECTOR('[1.1,2.2,3.3]',3,FLOAT32)RETURNINGVARCHAR2(1000))
[1.10000002E+000,2.20000005E+000,3.29999995E+000]
1 row selected.
SELECT FROM_VECTOR(TO_VECTOR('[1.1, 2.2, 3.3]', 3, FLOAT32) RETURNING CLOB);
FROM_VECTOR(TO_VECTOR('[1.1,2.2,3.3]',3,FLOAT32)RETURNINGCLOB)
[1.10000002E+000,2.20000005E+000,3.29999995E+000]
1 row selected.
SELECT FROM_VECTOR(TO_VECTOR('[5,[2,4],[1.0,2.0]]', 5, FLOAT64, SPARSE) RETURNING CLOB
FORMAT SPARSE);
FROM_VECTOR(TO_VECTOR('[5,[2,4],[1.0,2.0]]',5,FLOAT64,SPARSE)RETURNINGCLOBFORMAT
[5, [2, 4], [1.0E+000, 2.0E+000]]
1 row selected.
SELECT FROM_VECTOR(TO_VECTOR('[5,[2,4],[1.0,2.0]]', 5, FLOAT64, SPARSE) RETURNING CLOB
FORMAT DENSE);
FROM_VECTOR(TO_VECTOR('[5,[2,4],[1.0,2.0]]',5,FLOAT64,SPARSE)RETURNINGCLOBFORMAT
[0,1.0E+000,0,2.0E+000,0]
1 row selected.
```

Note:

 Applications using Oracle Client 23ai libraries or Thin mode drivers can fetch vector data directly, as shown in the following example:

SELECT dataVec FROM vecTab;

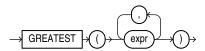
 For applications using Oracle Client 23ai libraries prior to 23ai connected to Oracle Database 23ai, use the FROM_VECTOR to fetch vector data, as shown by the following example:

SELECT FROM VECTOR(dataVec) FROM vecTab;



GREATEST

Syntax



Purpose

GREATEST returns the greatest of a list of one or more expressions. Oracle Database uses the first expr to determine the return type. If the first expr is numeric, then Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type before the comparison, and returns that data type. If the first expr is not numeric, then each expr after the first is implicitly converted to the data type of the first expr before the comparison.

Oracle Database compares each expr using nonpadded comparison semantics. The comparison is binary by default and is linguistic if the NLS_COMP parameter is set to LINGUISTIC and the NLS_SORT parameter has a setting other than BINARY. Character comparison is based on the numerical codes of the characters in the database character set and is performed on whole strings treated as one sequence of bytes, rather than character by character. If the value returned by this function is character data, then its data type is VARCHAR2 if the first expr is a character data type and NVARCHAR2 if the first expr is a national character data type.

See Also:

- "Data Type Comparison Rules" for more information on character comparison
- Table 2-9 for more information on implicit conversion and "Floating-Point Numbers" for information on binary-float comparison semantics
- "LEAST", which returns the least of a list of one or more expressions
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation GREATEST uses to compare character values for expr, and for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value

Examples

The following statement selects the string with the greatest value:

```
SELECT GREATEST('HARRY', 'HARRIOT', 'HAROLD') "Greatest"
FROM DUAL;

Greatest
------
HARRY
```

In the following statement, the first argument is numeric. Oracle Database determines that the argument with the highest numeric precedence is the second argument, converts the

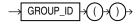
remaining arguments to the data type of the second argument, and returns the greatest value as that data type:

```
SELECT GREATEST (1, '3.925', '2.4') "Greatest"
FROM DUAL;

Greatest
-----
3.925
```

GROUP ID

Syntax



Purpose

GROUP_ID distinguishes duplicate groups resulting from a GROUP BY specification. It is useful in filtering out duplicate groupings from the query result. It returns an Oracle NUMBER to uniquely identify duplicate groups. This function is applicable only in a SELECT statement that contains a GROUP BY clause.

If *n* duplicates exist for a particular grouping, then $GROUP_ID$ returns numbers in the range 0 to n-1.

Examples

The following example assigns the value 1 to the duplicate co.country_region grouping from a query on the sample tables sh.countries and sh.sales:

COUNTRY_REGION	COUNTRY_SUBREGION	Revenue	G
Americas	Northern America	944.6	0
Americas	Northern America	944.6	1
Europe	Western Europe	566.39	0
Europe	Western Europe	566.39	1

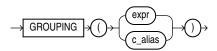
To ensure that only rows with $GROUP_ID < 1$ are returned, add the following HAVING clause to the end of the statement :

```
{\tt HAVING\ GROUP\_ID()} < 1
```



GROUPING

Syntax



Purpose

GROUPING distinguishes superaggregate rows from regular grouped rows. GROUP BY extensions such as ROLLUP and CUBE produce superaggregate rows where the set of all values is represented by null. Using the GROUPING function, you can distinguish a null representing the set of all values in a superaggregate row from a null in a regular row.

The expr in the <code>GROUPING</code> function must match one of the expressions in the <code>GROUP BY</code> clause. The function returns a value of 1 if the value of expr in the row is a null representing the set of all values. Otherwise, it returns zero. The data type of the value returned by the <code>GROUPING</code> function is Oracle <code>NUMBER</code>. Refer to the <code>SELECT</code> $group_by_clause$ for a discussion of these terms.

Examples

In the following example, which uses the sample tables <code>hr.departments</code> and <code>hr.employees</code>, if the <code>GROUPING</code> function returns 1 (indicating a superaggregate row rather than a regular row from the table), then the string "All Jobs" appears in the "JOB" column instead of the null that would otherwise appear:

```
SELECT
   DECODE(GROUPING(department_name), 1, 'ALL DEPARTMENTS', department_name)
   AS department,
   DECODE(GROUPING(job_id), 1, 'All Jobs', job_id) AS job,
   COUNT(*) "Total Empl",
   AVG(salary) * 12 "Average Sal"
   FROM employees e, departments d
   WHERE d.department_id = e.department_id
   GROUP BY ROLLUP (department_name, job_id)
   ORDER BY department, job;
```

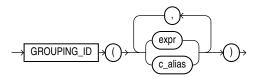
DEPARTMENT	JOB	Total Empl	Average Sal
ALL DEPARTMENTS	All Jobs	106	77481.0566
Accounting	AC_ACCOUNT	1	99600
Accounting	AC_MGR	1	144096
Accounting	All Jobs	2	121848
Administration	AD_ASST	1	52800
Administration	All Jobs	1	52800
Executive	AD_PRES	1	288000
Executive	AD_VP	2	204000
Executive	All Jobs	3	232000
Finance	All Jobs	6	103216
Finance	FI_ACCOUNT	5	95040

. . .



GROUPING_ID

Syntax



Purpose

GROUPING_ID returns a number corresponding to the GROUPING bit vector associated with a row. GROUPING_ID is applicable only in a SELECT statement that contains a GROUP BY extension, such as ROLLUP or CUBE, and a GROUPING function. In queries with many GROUP BY expressions, determining the GROUP BY level of a particular row requires many GROUPING functions, which leads to cumbersome SQL. GROUPING ID is useful in these cases.

GROUPING_ID is functionally equivalent to taking the results of multiple GROUPING functions and concatenating them into a bit vector (a string of ones and zeros). By using GROUPING_ID you can avoid the need for multiple GROUPING functions and make row filtering conditions easier to express. Row filtering is easier with GROUPING_ID because the desired rows can be identified with a single condition of GROUPING_ID = n. The function is especially useful when storing multiple levels of aggregation in a single table.

Examples

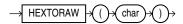
The following example shows how to extract grouping IDs from a query of the sample table sh.sales:

CHANNEL_ID	PROMO_ID	S_SALES	GC	GP	GCP	GPC
2	999	25797563.2	0	0	0	0
2		25797563.2	0	1	1	2
3	999	55336945.1	0	0	0	0
3		55336945.1	0	1	1	2
4	999	13370012.5	0	0	0	0
4		13370012.5	0	1	1	2
	999	94504520.8	1	0	2	1
		94504520.8	1	1	3	3



HEXTORAW

Syntax



Purpose

HEXTORAW converts char containing hexadecimal digits in the CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type to a raw value.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.



"Data Type Comparison Rules " for more information.

Examples

The following example creates a simple table with a raw column, and inserts a hexadecimal value that has been converted to RAW:

```
CREATE TABLE test (raw_col RAW(10));
INSERT INTO test VALUES (HEXTORAW('7D'));
```

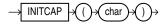
The following example converts hexadecimal digits to a raw value and casts the raw value to VARCHAR2:



"RAW and LONG RAW Data Types" and RAWTOHEX

INITCAP

Syntax





Purpose

INITCAP returns *char*, with the first letter of each word in uppercase, all other letters in lowercase. Words are delimited by white space or characters that are not alphanumeric.

char can be of any of the data types <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, or <code>NVARCHAR2</code>. The return value is the same data type as char. The database sets the case of the initial characters based on the binary mapping defined for the underlying character set. For linguistic-sensitive uppercase and lowercase, refer to <code>NLS_INITCAP</code>.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.

See Also:

- "Data Type Comparison Rules" for more information.
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of INITCAP

Examples

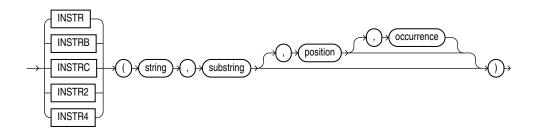
The following example capitalizes each word in the string:

```
SELECT INITCAP('the soap') "Capitals"
FROM DUAL;

Capitals
-----
The Soap
```

INSTR

Syntax



Purpose

The INSTR functions search string for substring. The search operation is defined as comparing the substring argument with substrings of string of the same length for equality until a match is found or there are no more substrings left. Each consecutive compared substring of string begins one character to the right (for forward searches) or one character to the left (for backward searches) from the first character of the previous compared substring. If



a substring that is equal to *substring* is found, then the function returns an integer indicating the position of the first character of this substring. If no such substring is found, then the function returns zero.

- position is an nonzero integer indicating the character of string where Oracle Database begins the search—that is, the position of the first character of the first substring to compare with substring. If position is negative, then Oracle counts backward from the end of string and then searches backward from the resulting position.
- occurrence is an integer indicating which occurrence of substring in string Oracle should search for. The value of occurrence must be positive. If occurrence is greater than 1, then the database does not return on the first match but continues comparing consecutive substrings of string, as described above, until match number occurrence has been found.

INSTR accepts and returns positions in characters as defined by the input character set, with the first character of string having position 1. INSTRB uses bytes instead of characters. INSTRC uses Unicode complete characters. INSTR2 uses UCS2 code points. INSTR4 uses UCS4 code points.

string can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The exceptions are INSTRC, INSTR2, and INSTR4, which do not allow string to be a CLOB or NCLOB.

substring can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB.

The value returned is of NUMBER data type.

Both position and occurrence must be of data type NUMBER, or any data type that can be implicitly converted to NUMBER, and must resolve to an integer. The default values of both position and occurrence are 1, meaning Oracle begins searching at the first character of string for the first occurrence of substring. The return value is relative to the beginning of string, regardless of the value of position.

See Also:

- Oracle Database Globalization Support Guide for more on character length.
- Oracle Database SecureFiles and Large Objects Developer's Guide for more on character length.
- Table 2-9 for more information on implicit conversion
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation the INSTR functions use to compare the substring argument with substrings of string

Examples

The following example searches the string CORPORATE FLOOR, beginning with the third character, for the string "OR". It returns the position in CORPORATE FLOOR at which the second occurrence of "OR" begins:

```
SELECT INSTR('CORPORATE FLOOR','OR', 3, 2) "Instring"
FROM DUAL;
Instring
```



```
-----
14
```

In the next example, Oracle counts backward from the last character to the third character from the end, which is the first \circ in <code>FLOOR</code>. Oracle then searches backward for the second occurrence of <code>OR</code>, and finds that this second occurrence begins with the second character in the search string :

The next example assumes a double-byte database character set.

ITERATION_NUMBER

Syntax

```
\rightarrow ITERATION_NUMBER \rightarrow
```

Purpose

The <code>ITERATION_NUMBER</code> function can be used only in the <code>model_clause</code> of the <code>SELECT</code> statement and then only when <code>ITERATE(number)</code> is specified in the <code>model_rules_clause</code>. It returns an integer representing the completed iteration through the model <code>rules</code>. The <code>ITERATION_NUMBER</code> function returns 0 during the first iteration. For each subsequent iteration, the <code>ITERATION_NUMBER</code> function returns the equivalent of <code>iteration_number</code> plus one.

```
See Also:
```

model_clause and "Model Expressions" for the syntax and semantics

Examples

The following example assigns the sales of the Mouse Pad for the years 1998 and 1999 to the sales of the Mouse Pad for the years 2001 and 2002 respectively:

```
SELECT country, prod, year, s
FROM sales_view_ref
MODEL
PARTITION BY (country)
DIMENSION BY (prod, year)
MEASURES (sale s)
IGNORE NAV
```



```
UNIQUE DIMENSION
RULES UPSERT SEQUENTIAL ORDER ITERATE(2)

(
    s['Mouse Pad', 2001 + ITERATION_NUMBER] =
    s['Mouse Pad', 1998 + ITERATION_NUMBER]
)
ORDER BY country, prod, year;
```

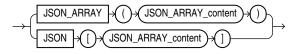
COUNTRY	PROD	YEAR	S
France	Mouse Pad	1998	2509.42
France	Mouse Pad	1999	3678.69
France	Mouse Pad	2000	3000.72
France	Mouse Pad	2001	2509.42
France	Mouse Pad	2002	3678.69
France	Standard Mouse	1998	2390.83
France	Standard Mouse	1999	2280.45
France	Standard Mouse	2000	1274.31
France	Standard Mouse	2001	2164.54
Germany	Mouse Pad	1998	5827.87
Germany	Mouse Pad	1999	8346.44
Germany	Mouse Pad	2000	7375.46
Germany	Mouse Pad	2001	5827.87
Germany	Mouse Pad	2002	8346.44
Germany	Standard Mouse	1998	7116.11
Germany	Standard Mouse	1999	6263.14
Germany	Standard Mouse	2000	2637.31
Germany	Standard Mouse	2001	6456.13

¹⁸ rows selected.

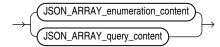
The preceding example requires the view <code>sales_view_ref</code>. Refer to "The MODEL clause: Examples" to create this view.

JSON_ARRAY

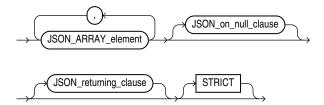
Syntax



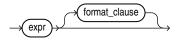
JSON_ARRAY_content



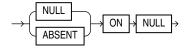
JSON_ARRAY_enumeration_content::=



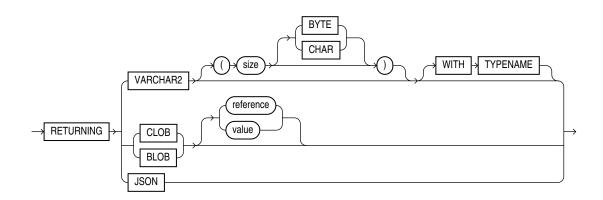
JSON_ARRAY_element



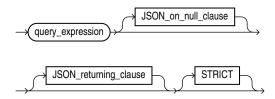
JSON_on_null_clause::=



JSON_returning_clause::=



JSON_ARRAY_query_content::=





Purpose

The SQL/JSON function JSON_ARRAY takes as its input a sequence of SQL scalar expressions or *one* collection type instance, VARRAY or NESTED TABLE.

It converts each expression to a JSON value, and returns a JSON array that contains those JSON values.

If an ADT has a member which is a collection than the type mapping creates a JSON object for the ADT with a nested JSON array for the collection member.

If a collection contains ADT instances then the type mapping will create a JSON array of JSON objects.



Generation of JSON Data Using SQL of the JSON Developer's Guide.

JSON_ARRAY_content

Use this clause to define the input to the JSON ARRAY function.

JSON_ARRAY_element

expr

For *expr*, you can specify any SQL expression that evaluates to a JSON object, a JSON array, a numeric literal, a text literal, date, timestamp, or null. This function converts a numeric literal to a JSON number value, and a text literal to a JSON string value. The date and timestamp data types are printed in the generated JSON object or array as JSON Strings following the ISO 8601 date format.

format clause

You can specify FORMAT JSON to indicate that the input string is JSON, and will therefore not be quoted in the output.

JSON on null clause

Use this clause to specify the behavior of this function when <code>expr</code> evaluates to null.

- NULL ON NULL If you specify this clause, then the function returns the JSON null value.
- ABSENT ON NULL If you specify this clause, then the function omits the value from the JSON array. This is the default.

JSON_returning_clause

Use this clause to specify the type of return value. One of :

- BLOB to return a binary large object of the AL32UTF8 character set.
- CLOB to return a character large object containing single-byte or multi-byte characters.
- VARCHAR2 specifying the size as a number of bytes or characters. The default is bytes. If you omit this clause, or specify the clause without specifying the size value, then JSON_ARRAY returns a character string of type VARCHAR2 (4000). Refer to VARCHAR2 Data Type for more information. Note that when specifying the VARCHAR2 data type elsewhere in



SQL, you are required to specify a size. However, in the <code>JSON_returning_clause</code> you can omit the size.

- BOOLEAN
- JSON
- VECTOR

STRICT

Specify the STRICT clause to verify that the output of the JSON generation function is correct JSON. If the check fails, a syntax error is raised.

Refer to JSON_OBJECT for examples.

Examples

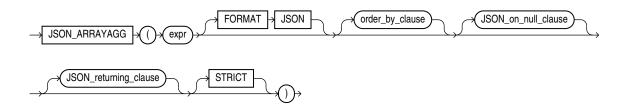
The following example constructs a JSON array from a JSON object, a JSON array, a numeric literal, a text literal, and null:

```
SELECT JSON_ARRAY (
    JSON_OBJECT('percentage' VALUE .50),
    JSON_ARRAY(1,2,3),
    100,
    'California',
    null
    NULL ON NULL
    ) "JSON Array Example"
    FROM DUAL;

JSON Array Example
[{"percentage":0.5},[1,2,3],100,"California",null]
```

JSON ARRAYAGG

Syntax

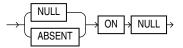


(See order_by_clause::= in the documentation on SELECT for the syntax of this clause)

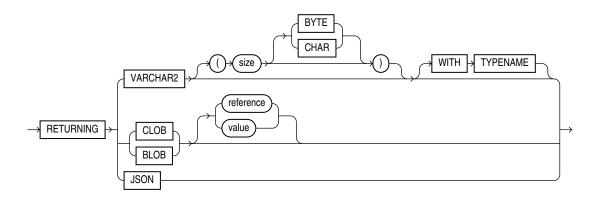


For JSON_ARRAYAGG, the <code>order_by</code> clause must refer to columns. Positional orders are not supported.

JSON on null clause::=



JSON_returning_clause::=



Purpose

The SQL/JSON function <code>JSON_ARRAYAGG</code> is an aggregate function. It takes as its input a column of SQL expressions, converts each expression to a JSON value, and returns a single JSON array that contains those JSON values.

expr

For expr, you can specify any SQL expression that evaluates to a JSON object, a JSON array, a numeric literal, a text literal, or null. This function converts a numeric literal to a JSON number value and a text literal to a JSON string value.

FORMAT JSON

Use this optional clause to indicate that the input string is JSON, and will therefore not be quoted in the output.

order_by_clause

This clause allows you to order the JSON values within the JSON array returned by the statement. Refer to the *order_by_clause* in the documentation on SELECT for the full semantics of this clause.

JSON on null clause

Use this clause to specify the behavior of this function when <code>expr</code> evaluates to null.

- NULL ON NULL If you specify this clause, then the function returns the JSON null value.
- ABSENT ON NULL If you specify this clause, then the function omits the value from the JSON array. This is the default.



JSON_returning_clause

Use this clause to specify the data type of the character string returned by this function. You can specify the following data types:

• VARCHAR2 [(size [BYTE, CHAR])]

When specifying the VARCHAR2 data type elsewhere in SQL, you are required to specify a size. However, in this clause you can omit the size.

- CLOB to return a character large object containing single-byte or multi-byte characters.
- BLOB to return a binary large object of the AL32UTF8 character set.
- JSON to return JSON data.

You must set the database initialization parameter compatible to 20 or greater to use the JSON type.

If you omit this clause, or if you specify VARCHAR2 but omit the size value, then JSON_ARRAYAGG returns a character string of type VARCHAR2 (4000).

Refer to "Data Types" for more information on the preceding data types.

STRICT

Specify the STRICT clause to verify that the output of the JSON generation function is correct JSON. If the check fails, a syntax error is raised.

Refer to JSON_OBJECT for examples.

WITH UNIQUE KEYS

Specify WITH UNIQUE KEYS to guarantee that generated JSON objects have unique keys.

Examples

The following statements creates a table id table, which contains ID numbers:

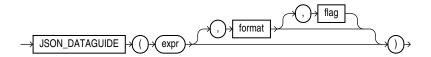
```
CREATE TABLE id_table (id NUMBER);
INSERT INTO id_table VALUES(624);
INSERT INTO id_table VALUES(null);
INSERT INTO id_table VALUES(925);
INSERT INTO id table VALUES(585);
```

The following example constructs a JSON array from the ID numbers in table id table:



JSON_DATAGUIDE

Syntax



Purpose

The aggregate function <code>JSON_DATAGUIDE</code> computes the data guide of a set of <code>JSON</code> data. The data guide is returned as a <code>CLOB</code> which can be in either flat or hierarchical format depending on the passing format parameter.

expr

expr is a SQL expression that evaluates to a JSON object or a JSON array. It can also be a JSON column in a table.

format options

Use the format options to specify the format of the data guide that will be returned. It must be one of the following values:

- dbms_json.format_flat for a flat format.
- dbms json.format hierarchical for a hierarchical format.
- dbms_json.format_schema for a data guide of a JSON schema that you can use to validate JSON documents.

If the parameter is the absent, the default is dbms json.format flat.

See Data-Guide Formats and Ways of Creating a Data Guide of the JSON Developer's Guide.

flag options

flag can have the following values:

- Specify DBMS_JSON.PRETTY to improve readability of the returned data guide with appropriate indentation.
- Specify DBMS_JSON.GEOJSON for the data guide to auto detect the GeoJSON type. The corresponding view column created by the data guide will be of sdo geometry type.
- Specify DBMS_JSON.GATHER_STATS for the data guide to collect statistical information. The data guide report generated with DBMS_JSON.GATHER_STATS has a new field o:sample_size, in addition to all of the other statistical fields that you get with DBMS_JSON.get index dataguide.
- Specify DBMS_JSON.DETECT_DATETIME for the data guide to detect temporal types. The data guide reports a JSON field value that conforms to the ISO 8601 format as a timestamp type, not a string type.
- All values DBMS_JSON.PRETTY, DBMS_JSON.GEOJSON, and DBMS_JSON.GATHER_STATS, and DBMS_JSON.DETECT_DATETIME can be combined with a plus sign. For example,



DBMS_JSON.GEOJSON+DBMS_JSON.PRETTY, **or**DBMS_JSON.GEOJSON+DBMS_JSON.PRETTY+DBMS_JSON.GATHER_STATS.

See Also:

JSON Data Guide

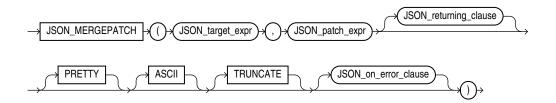
Examples

The following example uses the <code>j_purchaseorder</code> table, which is created in "Creating a Table That Contains a JSON Document: Example". This table contains a column of JSON data called <code>po document</code>. This example returns a flat data guide for each year group.

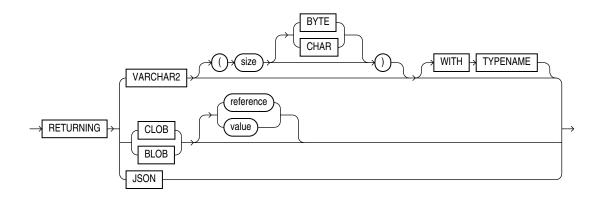
```
SELECT EXTRACT(YEAR FROM date_loaded) YEAR,
      JSON_DATAGUIDE(po_document) "DATA GUIDE"
  FROM j purchaseorder
  GROUP BY extract (YEAR FROM date loaded)
  ORDER BY extract (YEAR FROM date loaded) DESC;
YEAR DATA GUIDE
2016 [
         "o:path" : "$.PO_ID",
         "type" : "number",
         "o:length" : 4
         "o:path" : "$.PO_Ref",
"type" : "string",
         "o:length" : 16
       },
         "o:path" : "$.PO_Items",
         "type" : "array",
         "o:length" : 64
         "o:path" : "$.PO Items.Part No",
         "type" : "number",
         "o:length" : 16
       },
         "o:path" : "$.PO Items.Item Quantity",
         "type" : "number",
         "o:length" : 2
```

JSON_MERGEPATCH

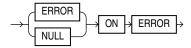
Syntax



JSON_returning_clause::=



json_on_error_clause::=



Purpose

You can use the JSON_MERGEPATCH function to update specific portions of a JSON document. You pass it a JSON Merge Patch document in JSON_patch_expr, which specifies the changes to make to a specified JSON document, the JSON target expr.

JSON_MERGEPATCH evaluates the patch document against the target document to produce the result document. If the target or the patch document is NULL, then the result is also NULL.

You can input any SQL data type that supports JSON data: JSON, VARCHAR2, CLOB, or BLOB. The function returns any of the SQL data types as output.

Data type JSON is available only if database initialization parameter compatible is 20 or greater.

The default return type depends on the input data type. If the input type is JSON, then JSON is also the default return type. Otherwise, VARCHAR2 is the default return type.

The ${\it JSON_returning_clause}$ specifies the return type of the operator. The default return type is ${\it VARCHAR2}$ (4000).

The PRETTY keyword specifies that the result should be formatted for human readability.

The ASCII keyword specifies that non-ASCII characters should be output using JSON escape sequences.

The TRUNCATE keyword specifies that the result document should be truncated to fit in the specified return type.

The <code>JSON_on_error_clause</code> optionally controls the handling of errors that occur during the processing of the target and patch documents.

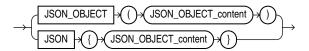
- NULL ON ERROR Returns null when an error occurs. This is the default.
- ERROR ON ERROR Returns the appropriate Oracle error when an error occurs.

See Also:

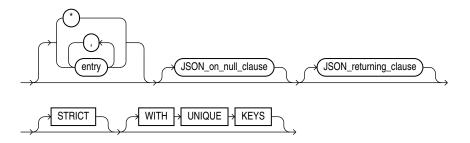
- RFC 7396 JSON Merge Patch
- Updating a JSON Document with JSON Merge Patch

JSON_OBJECT

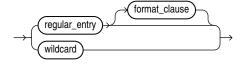
Syntax



json_object_content::=

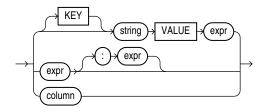


entry::=

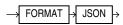


regular_entry::=

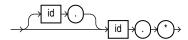




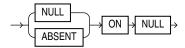
format_clause::=



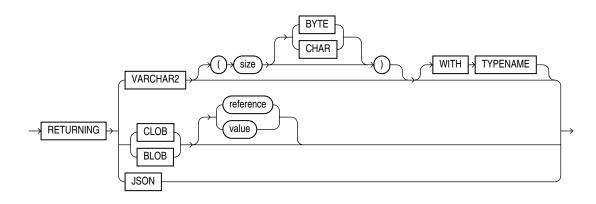
wildcard::=



JSON_on_null_clause::=



JSON_returning_clause::=



Purpose

The SQL/JSON function <code>JSON_OBJECT</code> takes as its input either a sequence of key-value pairs or one object type instance. A collection type cannot be passed to <code>JSON_OBJECT</code>.

It returns a JSON object that contains an object member for each of those key-value pairs.

entry

regular_entry: Use this clause to specify a property key-value pair.

regular_entry

- KEY is optional and is provided for semantic clarity.
- Use the optional expr to specify the property key name as a case-sensitive text literal.
- Use *expr* to specify the property value. For *expr*, you can specify any expression that evaluates to a SQL numeric literal, text literal, date, or timestamp. The date and timestamp data types are printed in the generated JSON object or array as JSON strings following the ISO date format. If *expr* evaluates to a numeric literal, then the resulting property value is a JSON number value; otherwise, the resulting property value is a case-sensitive JSON string value enclosed in double quotation marks.

You can use the colon to separate JSON OBJECT entries.

Example

```
SELECT JSON_OBJECT(
'name' : first_name || ' ' || last_name,
'email' : email,
'phone' : phone_number,
'hire_date' : hire_date
)
FROM employees
WHERE employee_id = 140;
```

format clause

Specify FORMAT JSON after an input expression to declare that the value that results from it represents JSON data, and will therefore not be quoted in the output.

wildcard

Wildcard entries select multiple columns and can take the form of *, table.*, view.*, or t_alias.*. Use wildcard entries to map all the columns from a table, subquery, or view to a JSON object without explicitly naming all of the columns in the query. In this case wildcard entries are used in the same way that they are used directly in a select_list.

Example 1

In the resulting JSON object, the key names are equal to the names of the corresponding columns.

```
SELECT JSON_OBJECT(*)
FROM employees
WHERE employee id = 140;
```

Output 1

```
{"EMPLOYEE_ID":140,"FIRST_NAME":"Joshua","LAST_NAME":"Patel","EMAIL":"JPAT EL","PHONE_NUMBER":"650.121.1834","HIRE_DATE":"2006-04-06T00:00:00","JOB_ID":"ST_CLERK","SALARY":2500,"COMMISSION_PCT":null,"MAN AGER ID":123,"DEPARTMENT ID":50}
```

Example 2



This query selects columns from a specific table in a join query.

```
SELECT JSON_OBJECT('NAME' VALUE first_name, d.*)
FROM employees e, departments d
WHERE e.department_id = d.department_id
AND e.employee id =140
```

Example 3

This query converts the departments table to a single JSON array value.

```
SELECT JSON_ARRAYAGG(JSON_OBJECT(*))
FROM departments
```

JSON_on_null_clause

Use this clause to specify the behavior of this function when expr evaluates to null.

NULL ON NULL - When NULL ON NULL is specified, then a JSON NULL value is used as a
value for the given key.

```
SELECT JSON OBJECT('key1' VALUE NULL) evaluates to {"key1": null}
```

 ABSENT ON NULL - If you specify this clause, then the function omits the property key-value pair from the JSON object.

JSON_returning_clause

Use this clause to specify the type of return value. One of :

- VARCHAR2 specifying the size as a number of bytes or characters. The default is bytes. If you omit this clause, or specify the clause without specifying the size value, then JSON_ARRAY returns a character string of type VARCHAR2 (4000). Refer to VARCHAR2 Data Type for more information. Note that when specifying the VARCHAR2 data type elsewhere in SQL, you are required to specify a size. However, in the JSON_returning_clause you can omit the size.
- CLOB to return a character large object containing single-byte or multi-byte characters.
- BLOB to return a binary large object of the AL32UTF8 character set.
- WITH TYPENAME

STRICT

Specify the STRICT clause to verify that the output of the JSON generation function is correct JSON. If the check fails, a syntax error is raised.

Example 1: Output string appears within quotes, because FORMAT JSON is not used



Example 2: No quotes around output string when FORMAT JSON is used.

Example 3: JSON Syntax error when FORMAT JSON STRICT is used.

```
SELECT JSON_OBJECT ('name' value 'Foo' FORMAT JSON STRICT ) FROM DUAL Output:
ORA-40441: JSON syntax error
```

WITH UNIQUE KEYS

Specify WITH UNIQUE KEYS to guarantee that generated JSON objects have unique keys.

Example

The following example returns JSON objects that each contain two property key-value pairs:

JSON_OBJECT Column Entries

In some cases you might want to have JSON object key names match the names of the table columns to avoid repeating the column name in the key value expression. For example:

```
SELECT JSON_OBJECT(
'first_name' VALUE first_name,
'last_name' VALUE last_name,
'email' VALUE email,
'hire_date' VALUE hire_date
)
FROM employees
WHERE employee_id = 140;
{"first_name":"Joshua","last_name":"Patel","email":"JPATEL","hire_date":"2006-
```



```
04-
06T00:00:00"}
```

In such cases you can use a shortcut, where a single column value may be specified as input and the corresponding object entry key is inferred from the name of the column. For example:

```
SELECT JSON_OBJECT(first_name, last_name, email, hire_date)
FROM employees
WHERE employee_id = 140;

{"first_name":"Joshua","last_name":"Patel","email":"JPATEL","hire_date":"2006-04-
06T00:00:00"}
```

You can use quoted or non-quoted identifiers for column names. If you use non-quoted identifiers, then the case-sensitive value of the identifier, as written in the query, is used to generate the corresponding object key value. However for the purpose of referencing the column value, the identifier is still case-insensitive. For example:

```
SELECT JSON_OBJECT(eMail)
FROM employees
WHERE employee_id = 140
{"eMail":"JPATEL"}
```

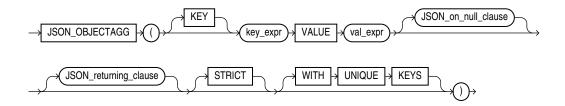
Notice that the capital 'M' as typed in the column name is preserved.



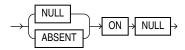
Generation of JSON Data Using SQL

JSON_OBJECTAGG

Syntax

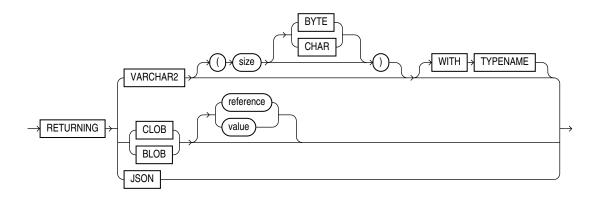


JSON_on_null_clause::=





JSON_returning_clause::=



Purpose

The SQL/JSON function <code>JSON_OBJECTAGG</code> is an aggregate function. It takes as its input a property key-value pair. Typically, the property key, the property value, or both are columns of SQL expressions. This function constructs an object member for each key-value pair and returns a single JSON object that contains those object members.

[KEY] string VALUE expr

Use this clause to specify property key-value pairs.

- KEY is optional and is provided for semantic clarity.
- Use string to specify the property key name as a case-sensitive text literal.
- Use <code>expr</code> to specify the property value. For <code>expr</code>, you can specify any expression that evaluates to a SQL numeric literal, text literal, date, or timestamp. The date and timestamp data types are printed in the generated JSON object or array as JSON Strings following the ISO 8601 date format. If <code>expr</code> evaluates to a numeric literal, then the resulting property value is a JSON number value; otherwise, the resulting property value is a case-sensitive JSON string value enclosed in double quotation marks.

FORMAT JSON

Use this optional clause to indicate that the input string is JSON, and will therefore not be quoted in the output.

JSON on null clause

Use this clause to specify the behavior of this function when expr evaluates to null.

- NULL ON NULL When NULL ON NULL is specified, then a JSON NULL value is used as a
 value for the given key.
- ABSENT ON NULL If you specify this clause, then the function omits the property key-value pair from the JSON object.

JSON returning clause

Use this clause to specify the data type of the character string returned by this function. You can specify the following data types:

VARCHAR2 [(size [BYTE, CHAR])]



When specifying the VARCHAR2 data type elsewhere in SQL, you are required to specify a size. However, in this clause you can omit the size.

- CLOB to return a character large object containing single-byte or multi-byte characters.
- BLOB to return a binary large object of the AL32UTF8 character set.
- JSON to return JSON data.

You must set the database initialization parameter compatible to 20 or greater to use the JSON data type.

If you omit this clause, or if you specify VARCHAR2 but omit the size value, then JSON_OBJECTAGG returns a character string of type VARCHAR2 (4000).

Refer to "Data Types" for more information on the preceding data types.

STRICT

Specify the STRICT clause to verify that the output of the JSON generation function is correct JSON. If the check fails, a syntax error is raised.

Refer to JSON_OBJECT for examples.

WITH UNIQUE KEYS

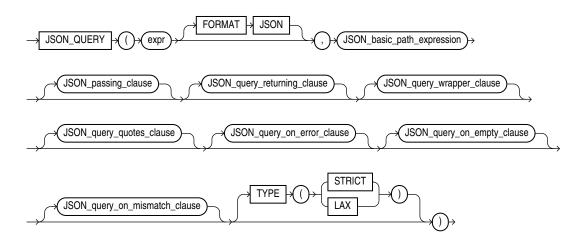
Specify WITH UNIQUE KEYS to guarantee that generated JSON objects have unique keys.

Examples

The following example constructs a JSON object whose members contain department names and department numbers:

JSON QUERY

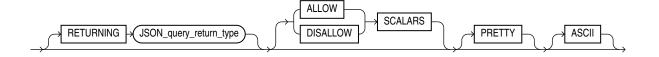
Syntax



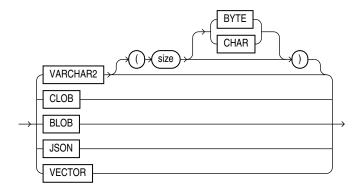


(JSON basic path expression: See Oracle Database JSON Developer's Guide)

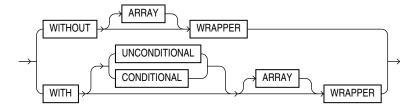
JSON_query_returning_clause::=



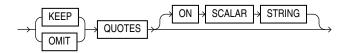
JSON_query_return_type::=



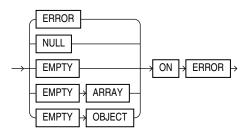
JSON_query_wrapper_clause::=



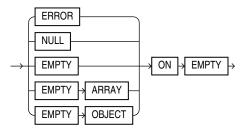
JSON_query_quotes_clause::=



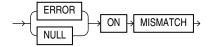
JSON_query_on_error_clause::=



JSON_query_on_empty_clause::=



JSON_query_on_mismatch_clause::=



Purpose

JSON_QUERY selects and returns one or more values from JSON data and returns those values. You can use JSON QUERY to retrieve fragments of a JSON document.

See Also:

- Query JSON Data
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character value returned by JSON QUERY

expr

Use *expr* to specify the JSON data you want to query.

expr is a SQL expression that returns an instance of a SQL data type, one of JSON, VARCHAR2, CLOB, or BLOB. It can be a table or view column value, a PL/SQL variable, or a bind variable with proper casting.

If expr is null, then the function returns null.

If expr is not a text literal of well-formed JSON data using strict or lax syntax, then the function returns null by default. You can use the $JSON_query_on_error_clause$ to override this default behavior. Refer to $JSON_query_on_error_clause$.

FORMAT JSON

You must specify FORMAT JSON if expr is a column of data type BLOB.

JSON_basic_path_expression

Use this clause to specify a SQL/JSON path expression. The function uses the path expression to evaluate expr and find one or more JSON values that match, or satisfy the path



expression. The path expression must be a text literal. See *Oracle Database JSON Developer's Guide* for the full semantics of *JSON basic path expression*.

JSON_query_returning_clause

Use this clause to specify the data type and format of the character string returned by this function.

RETURNING

You can use the RETURNING clause to specify the data type of the returned instance, one of JSON, VARCHAR2, CLOB, or BLOB.

The default return type depends on the input data type. If the input type is JSON, then JSON is also the default return type. Otherwise VARCHAR2 (4000) is the default return type.

When specifying the VARCHAR2 data type elsewhere in SQL, you are required to specify a size. However, in this clause you can omit the size. In this case, <code>JSON_QUERY</code> returns a character string of type <code>VARCHAR2(4000)</code>.

Refer to "VARCHAR2 Data Type" for more information.

If the data type is not large enough to hold the return character string, then <code>JSON_QUERY</code> returns null by default. You can use the <code>JSON_query_on_error_clause</code> to override this default behavior. Refer to the <code>JSON_query_on_error_clause</code>.

PRETTY

Specify PRETTY to pretty-print the return character string by inserting newline characters and indenting.

ASCII

Specify ASCII to automatically escape any non-ASCII Unicode characters in the return character string, using standard ASCII Unicode escape sequences.

JSON_query_wrapper_clause

Use this clause to control whether this function wraps the values matched by the path expression in an array wrapper—that is, encloses the sequence of values in square brackets ([]).

- Specify WITHOUT WRAPPER to omit the array wrapper. You can specify this clause only if the path expression matches a single JSON object or JSON array. This is the default.
- Specify WITH WRAPPER to include the array wrapper. You must specify this clause if the path
 expression matches a single scalar value (a value that is not a JSON object or JSON
 array) or multiple values of any type.
- Specifying the WITH UNCONDITIONAL WRAPPER clause is equivalent to specifying the WITH WRAPPER clause. The UNCONDITIONAL keyword is provided for semantic clarity.
- Specify WITH CONDITIONAL WRAPPER to include the array wrapper only if the path expression matches a single scalar value or multiple values of any type. If the path expression matches a single JSON object or JSON array, then the array wrapper is omitted.

The ARRAY keyword is optional and is provided for semantic clarity.

If the function returns a single scalar value, or multiple values of any type, and you do not specify WITH [UNCONDITIONAL | CONDITIONAL] WRAPPER, then the function returns null by



default. You can use the <code>JSON_query_on_error_clause</code> to override this default behavior. Refer to the <code>JSON_query_on_error_clause</code>.

JSON_query_on_error_clause

Use this clause to specify the value returned by this function when the following errors occur:

- expr is not well-formed JSON data using strict or lax JSON syntax
- No match is found when the JSON data is evaluated using the SQL/JSON path expression. You can override the behavior for this type of error by specifying the JSON_query_on_empty_clause.
- The return value data type is not large enough to hold the return character string
- The function matches a single scalar value or, multiple values of any type, and the WITH [UNCONDITIONAL | CONDITIONAL] WRAPPER clause is not specified

You can specify the following clauses:

- NULL ON ERROR Returns null when an error occurs. This is the default.
- ERROR ON ERROR Returns the appropriate Oracle error when an error occurs.
- EMPTY ON ERROR Specifying this clause is equivalent to specifying EMPTY ARRAY ON ERROR.
- EMPTY ARRAY ON ERROR Returns an empty JSON array ([]) when an error occurs.
- EMPTY OBJECT ON ERROR Returns an empty JSON object ({}) when an error occurs.

JSON_query_on_empty_clause

Use this clause to specify the value returned by this function if no match is found when the JSON data is evaluated using the SQL/JSON path expression. This clause allows you to specify a different outcome for this type of error than the outcome specified with the <code>JSON_query_on_error_clause</code>.

You can specify the following clauses:

- NULL ON EMPTY Returns null when no match is found.
- ERROR ON EMPTY Returns the appropriate Oracle error when no match is found.
- EMPTY ON EMPTY Specifying this clause is equivalent to specifying EMPTY ARRAY ON EMPTY.
- EMPTY ARRAY ON EMPTY Returns an empty JSON array ([]) when no match is found.
- EMPTY OBJECT ON EMPTY Returns an empty JSON object ({}) when no match is found.

If you omit this clause, then the <code>JSON_query_on_error_clause</code> determines the value returned when no match is found.

TYPE Clause

For a full discussion of STRICT and LAX syntax see About Strict and Lax JSON Syntax, and TYPE Clause for SQL Functions and Conditions

Examples

The following query returns the context item, or the specified string of JSON data. The path expression matches a single JSON object, which does not require an array wrapper. Note that the JSON data is converted to strict JSON syntax in the returned value—that is, the object property names are enclosed in double quotation marks.



The following query returns the value of the member with property name a. The path expression matches a scalar value, which must be enclosed in an array wrapper. Therefore, the WITH WRAPPER clause is specified.

```
SELECT JSON_QUERY('{a:100, b:200, c:300}', '$.a' WITH WRAPPER) AS value FROM DUAL;

VALUE

[100]
```

The following query returns the values of all object members. The path expression matches multiple values, which together must be enclosed in an array wrapper. Therefore, the WITH WRAPPER clause is specified.

```
SELECT JSON_QUERY('{a:100, b:200, c:300}', '$.*' WITH WRAPPER) AS value FROM DUAL;

VALUE

[100,200,300]
```

The following query returns the context item, or the specified string of JSON data. The path expression matches a single JSON array, which does not require an array wrapper.

```
SELECT JSON_QUERY('[0,1,2,3,4]', '$') AS value FROM DUAL;

VALUE

[0,1,2,3,4]
```

The following query is similar to the previous query, except the WITH WRAPPER clause is specified. Therefore, the JSON array is wrapped in an array wrapper.

The following query returns all elements in a JSON array. The path expression matches multiple values, which together must be enclosed in an array wrapper. Therefore, the WITH WRAPPER clause is specified.

```
SELECT JSON_QUERY('[0,1,2,3,4]', '$[*]' WITH WRAPPER) AS value
   FROM DUAL;

VALUE
[0,1,2,3,4]
```



The following query returns the elements at indexes 0, 3 through 5, and 7 in a JSON array. The path expression matches multiple values, which together must be enclosed in an array wrapper. Therefore, the WITH WRAPPER clause is specified.

The following query returns the fourth element in a JSON array. The path expression matches a scalar value, which must be enclosed in an array wrapper. Therefore, the WITH WRAPPER clause is specified.

```
SELECT JSON_QUERY('[0,1,2,3,4]', '$[3]' WITH WRAPPER) AS value FROM DUAL;

VALUE

[3]
```

The following query returns the first element in a JSON array. The WITH CONDITIONAL WRAPPER clause is specified and the path expression matches a single JSON object. Therefore, the value returned is not wrapped in an array. Note that the JSON data is converted to strict JSON syntax in the returned value—that is, the object property name is enclosed in double quotation marks.

The following query returns all elements in a JSON array. The WITH CONDITIONAL WRAPPER clause is specified and the path expression matches multiple JSON objects. Therefore, the value returned is wrapped in an array.

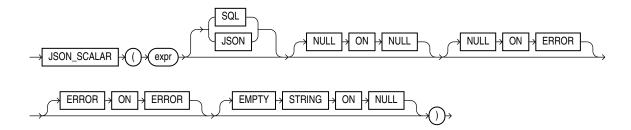
The following query is similar to the previous query, except that the value returned is of data type VARCHAR2 (100).



The following query returns the fourth element in a JSON array. However, the supplied JSON array does not contain a fourth element, which results in an error. The EMPTY ON ERROR clause is specified. Therefore, the query returns an empty JSON array.

JSON_SCALAR

Syntax



Purpose

JSON_SCALAR accepts a SQL scalar value as input and returns a corresponding JSON scalar value as a JSON type instance. The value can be an Oracle-specific JSON-language type, such as a date, which is not part of the JSON standard.

To use $\scalebox{JSON_SCALAR}$ you must set the database initialization parameter compatible is at least 20. Otherwise it raises an error.

The argument to JSON_SCALAR can be an instance of any of these SQL data types: BINARY_DOUBLE, BINARY_FLOAT, BLOB, CLOB, DATE, INTERVAL YEAR TO MONTH, INTERVAL DAY TO SECOND, JSON, NUMBER, RAW, TIMESTAMP, VARCHAR, VARCHAR2, or VECTOR.

The returned JSON type instance is a JSON-language scalar value supported by Oracle.

If the argument to <code>JSON_SCALAR</code> is a SQL NULL value, then you can obtain a return value as follows:

- SQL NULL, the default behavior
- JSON null, using keywords NULL ON NULL
- An empty JSON string, " ", using keywords EMPTY STRING ON NULL

The default behavior of returning SQL $\tt NULL$ is the only exception to the rule that a JSON scalar value is returned.

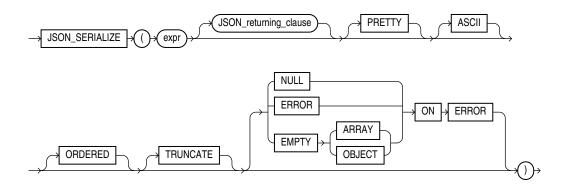


See Oracle SQL Function JSON_SCALAR of the JSON Developer's Guide.

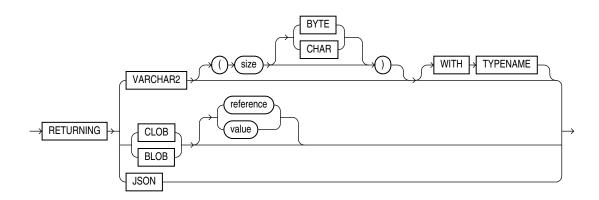
JSON_SERIALIZE

Syntax

json_serialize



json_returning_clause



Purpose

 $json_serialize$ takes JSON data of any SQL data type (BLOB, CLOB, JSON, or VARCHAR2) as input and returns a textual representation of it. You typically use it to transform the result of a query.

You can use <code>json_serialize</code> to convert binary JSON data to textual form (<code>CLOB</code> or <code>VARCHAR2</code>), or to transform textual JSON data by pretty-printing it or escaping non-ASCII Unicode characters in it.

When Oracle SQL function $vector_serialize$ is applied to a JSON type instance, any non-standard Oracle scalar JSON value is returned as a standard JSON scalar value.

When you apply *vector_serialize* to a VECTOR type instance, it returns a textual JSON array of numbers.



Note:

You can serialize a VECTOR instance to a textual JSON array of numbers using SQL function <code>vector_serialize</code>. (Function json_serialize serializes only JSON data.) See <code>VECTOR_SERIALIZE</code>

See Also:

Oracle SQL Function JSON_SERIALIZE of the JSON Developer's Guide.

expr

expr is the input expression. Can be any one of type JSON, VARCHAR2, CLOB, or BLOB.

JSON_returning_clause::=

You can use the <code>JSON_returning_clause</code> to specify the return type of the function. One of <code>BOOLEAN</code>, <code>BLOB</code>, <code>CLOB</code>, <code>JSON</code>, or <code>VARCHAR2</code>.

The default return type is VARCHAR2 (4000).

If the return type is RAW or BLOB, it contains UTF8 encoded JSON text.

PRETTY

Specify PRETTY if you want the result to be formatted for human readability.

ASCII

Specify ASCII if you want non-ASCII characters to be output using JSON escape sequences.

ORDERED

Specify ORDERED if you want to reorder key-value pairs alphabetically in ascending order. You can combine ORDERED with PRETTY and ASCII.

Example

```
SELECT JSON_SERIALIZE('{price:20, currency:" €"}' ASCII PRETTY ORDERED) from dual; {
   "currency" : "\u20AC",
   "price" : 20
}
```

TRUNCATE

Specify TRUNCATE, if you want the textual output in the result document to fit into the buffer of the specified return type.

JSON_on_error_clause::=

Specify JSON on error clause to control the handling of processing errors.

ERROR ON ERROR is the default.

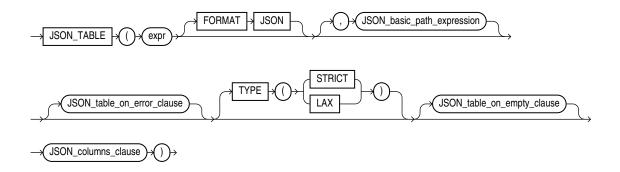
EMPTY ON ERROR is not supported.

If you specify TRUNCATE with JSON_on_error_clause, then a value too large for the return type will be truncated to fit into the buffer instead of raising an error.

Example

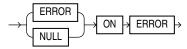
JSON_TABLE

Syntax

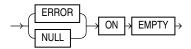


(JSON_basic_path_expression: See Oracle Database JSON Developer's Guide, JSON_table_on_error_clause::=, JSON_columns_clause::=)

JSON_table_on_error_clause::=



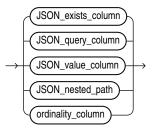
JSON_table_on_empty_clause::=



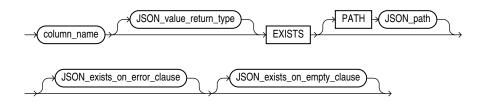
JSON_columns_clause::=



JSON column definition::=

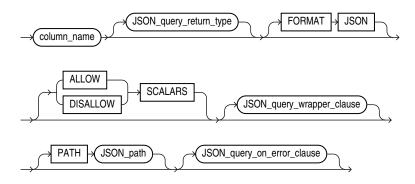


JSON_exists_column::=



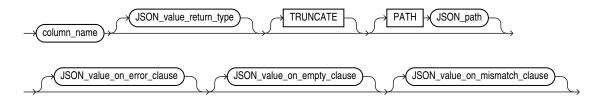
(JSON_value_return_type::=—part of JSON_VALUE, JSON_basic_path_expression: See Oracle Database JSON Developer's Guide, JSON_exists_on_error_clause::=—part of JSON_EXISTS)

JSON_query_column::=



(JSON_query_return_type::=, JSON_query_wrapper_clause::=, and JSON_query_on_error_clause::=—part of JSON_QUERY, JSON_basic_path_expression: See Oracle Database JSON Developer's Guide)

JSON_value_column::=



(JSON_value_return_type::= and JSON_value_on_error_clause::=—part of JSON_VALUE, JSON basic path expression: See Oracle Database JSON Developer's Guide)

JSON_nested_path::=

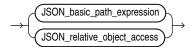


(JSON_basic_path_expression: See Oracle Database JSON Developer's Guide, JSON columns clause::=)

ordinality_column::=



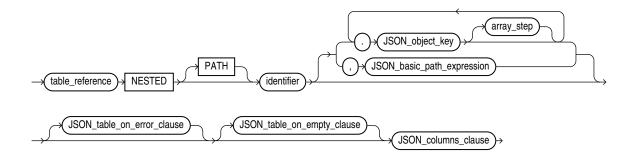
JSON_path ::=



JSON_relative_object_access ::=



nested_clause ::=



Purpose

The SQL/JSON function <code>JSON_TABLE</code> creates a relational view of JSON data. It maps the result of a JSON data evaluation into relational rows and columns. You can query the result returned by the function as a virtual relational table using SQL. The main purpose of <code>JSON_TABLE</code> is to

create a row of relational data for each object inside a JSON array and output JSON values from within that object as individual SQL column values.

You must specify JSON_TABLE only in the FROM clause of a SELECT statement. The function first applies a path expression, called a **SQL/JSON row path expression**, to the supplied JSON data. The JSON value that matches the row path expression is called a **row source** in that it generates a row of relational data. The COLUMNS clause evaluates the row source, finds specific JSON values within the row source, and returns those JSON values as SQL values in individual columns of a row of relational data.

The COLUMNS clause enables you to search for JSON values in different ways by using the following clauses:

- JSON_exists_column Evaluates JSON data in the same manner as the JSON_EXISTS condition, that is, determines if a specified JSON value exists, and returns either a VARCHAR2 column of values 'true' or 'false', or a NUMBER column of values 1 or 0.
- JSON_query_column Evaluates JSON data in the same manner as the JSON_QUERY function, that is, finds one or more specified JSON values, and returns a column of character strings that contain those JSON values.
- JSON_value_column Evaluates JSON data in the same manner as the JSON_VALUE function, that is, finds a specified scalar JSON value, and returns a column of those JSON values as SQL values.
- JSON_nested_path Allows you to flatten JSON values in a nested JSON object or JSON
 array into individual columns in a single row along with JSON values from the parent object
 or array. You can use this clause recursively to project data from multiple layers of nested
 objects or arrays into a single row.
- ordinality column Returns a column of generated row numbers.

The column definition clauses allow you to specify a name for each column of data that they return. You can reference these column names elsewhere in the SELECT statement, such as in the SELECT list and the WHERE clause.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to each character data type column in the table generated by JSON TABLE

expr

Use this clause to specify the JSON data to be evaluated. For expr, specify an expression that evaluates to a text literal. If expr is a column, then the column must be of data type VARCHAR2, CLOB, or BLOB. If expr is null, then the function returns null.

If expr is not a text literal of well-formed JSON data using strict or lax syntax, then the function returns null by default. You can use the <code>JSON_table_on_error_clause</code> to override this default behavior. Refer to <code>JSON table on error clause</code>.

FORMAT JSON

You must specify FORMAT JSON if expr is a column of data type BLOB.



PATH

Use the PATH clause to delineate a portion of the row that you want to use as the column content. The absence of the PATH clause does not change the behavior with a path of '\$.<column-name>', where <column-name> is the column name. The name of the object field that is targeted is taken implicitly as the column name. See *Oracle Database JSON Developer's Guide* for the full semantics of PATH.

JSON_basic_path_expression

The JSON_basic_path_expression is a text literal. See *Oracle Database JSON Developer*'s *Guide* for the full semantics of this clause.

JSON relative object access

Specify this row path expression to enable simple dot notation. The value of <code>JSON_relative_object_access</code> is evaluated as a <code>JSON/Path</code> expression relative to the current row item.

For more information on the <code>JSON_object_key clause</code>, refer to <code>JSON Object Access Expressions</code> .

JSON table on error clause

Use this clause to specify the value returned by the function when errors occur:

- NULL ON ERROR
 - If the input is not well-formed JSON text, no more rows will be returned as soon as the error is detected. Note that since JSON_TABLE supports streaming evaluation, rows may be returned prior to encountering the portion of the input with the error.
 - If no match is found when the row path expression is evaluated, no rows are returned.
 - Sets the default error behavior for all column expressions to NULL ON ERROR
- ERROR ON ERROR
 - If the input is not well-formed JSON text, an error will be raised.
 - If no match is found when the row path expression is evaluated, an error will be raised
 - Sets the default error behavior for all column expressions to ERROR ON ERROR

TYPE Clause

For a full discussion of STRICT and LAX syntax see About Strict and Lax JSON Syntax, and TYPE Clause for SQL Functions and Conditions

JSON_table_on_empty_clause

Use this clause to specify the value returned by this function if no match is found when the JSON data is evaluated using the SQL/JSON path expression. This clause allows you to specify a different outcome for this type of error than the outcome specified with the <code>JSON_table_on_error_clause</code>.

You can specify the following clauses:

- NULL ON EMPTY Returns null when no match is found.
- ERROR ON EMPTY Returns the appropriate Oracle error when no match is found.



• DEFAULT literal ON EMPTY - Returns literal when no match is found. The data type of literal must match the data type of the value returned by this function.

If you omit this clause, then the <code>JSON_table_on_error_clause</code> determines the value returned when no match is found.

JSON columns clause

Use the COLUMNS clause to define the columns in the virtual relational table returned by the JSON TABLE function.

JSON_exists_column

This clause evaluates JSON data in the same manner as the JSON_EXISTS condition, that is, it determines if a specified JSON value exists. It returns either a VARCHAR2 column of values 'true' or 'false', or a NUMBER column of values 1 or 0.

A value of 'true' or 1 indicates that the JSON value exists and a value of 'false' or 0 indicates that the JSON value does not exist.

You can use the <code>JSON_value_return_type</code> clause to control the data type of the returned column. If you omit this clause, then the data type is <code>VARCHAR2(4000)</code>. Use <code>column_name</code> to specify the name of the returned column. The rest of the clauses of <code>JSON_exists_column</code> have the same semantics here as they have for the <code>JSON_EXISTS</code> condition. For full information on these clauses, refer to "<code>JSON_EXISTS</code> Condition". Also see "Using <code>JSON_exists_column</code>: Examples" for an example.

JSON_query_column

This clause evaluates JSON data in the same manner as the <code>JSON_QUERY</code> function, that is, it finds one or more specified JSON values, and returns a column of character strings that contain those JSON values.

Use <code>column_name</code> to specify the name of the returned column. The rest of the clauses of <code>JSON_query_column</code> have the same semantics here as they have for the <code>JSON_QUERY</code> function. For full information on these clauses, refer to <code>JSON_QUERY</code>. Also see "Using <code>JSON_query_column</code>: Examples" for an example.

JSON value column

This clause evaluates JSON data in the same manner as the $\tt JSON_VALUE$ function, that is, it finds a specified scalar JSON value, and returns a column of those JSON values as SQL values.

Use <code>column_name</code> to specify the name of the returned column. The rest of the clauses of <code>JSON_value_column</code> have the same semantics here as they have for the <code>JSON_VALUE</code> function. For full information on these clauses, refer to <code>JSON_VALUE</code>. Also see "Using <code>JSON_value_column</code>: Examples" for an example.

JSON nested path

Use this clause to flatten JSON values in a nested JSON object or JSON array into individual columns in a single row along with JSON values from the parent object or array. You can use this clause recursively to project data from multiple layers of nested objects or arrays into a single row.

Specify the <code>JSON_basic_path_expression</code> clause to match the nested object or array. This path expression is relative to the SQL/JSON row path expression specified in the <code>JSON_TABLE</code> function.



Use the COLUMNS clause to define the columns of the nested object or array to be returned. This clause is recursive—you can specify the <code>JSON_nested_path</code> clause within another <code>JSON_nested_path</code>: Examples" for an example.

ordinality_column

This clause returns a column of generated row numbers of data type NUMBER. You can specify at most one <code>ordinality_column</code>. Also see "Using <code>JSON_value_column</code>: Examples" for an example of using the <code>ordinality_column</code> clause.

nested_clause

Use the <code>nested_clause</code> as a short-hand syntax for mapping JSON values to relational columns. It reuses the syntax of the <code>JSON_TABLE</code> columns clause and is essentially equivalent to a left-outer ANSI join with <code>JSON_TABLE</code>.

Example 1 using the nested_clause is equivalent to Example 2 using the left-outer join with ${\tt JSON}$ TABLE .

Example 1 Nested_Clause

Example 2 Left-Outer Join With JSON TABLE

```
SELECT t.*
FROM j_purchaseOrder LEFT OUTER JOIN
JSON TABLE(po document COLUMNS(PONumber, Reference, Requestor)) t ON 1=1;
```

When using the <code>nested_clause</code>, the JSON column name following the <code>NESTED</code> keyword will not be included in <code>SELECT * expansion</code>. For example:

The result does not include the JSON column name po_document as one of the columns in the result.

When unnesting JSON column data, the recommendation is to use LEFT OUTER JOIN semantics, so that JSON columns that produce no rows will not filter other non-JSON data from the result. For example, a <code>j_purchaseOrder</code> row with a <code>NULL po_document</code> column will not filter the possibly non-null relational <code>columns</code> id and <code>date loaded</code> from the result.



The columns clause supports all the same features defined for ${\tt JSON_TABLE}$ including nested columns. For example:

Examples

Creating a Table That Contains a JSON Document: Example

This example shows how to create and populate table <code>j_purchaseorder</code>, which is used in the rest of the <code>JSON TABLE</code> examples in this section.

The following statement creates table <code>j_purchaseorder</code>. Column <code>po_document</code> is for storing JSON data and, therefore, has an <code>IS JSON</code> check constraint to ensure that only well-formed JSON is stored in the column.

```
CREATE TABLE j_purchaseorder
  (id RAW (16) NOT NULL,
   date_loaded TIMESTAMP(6) WITH TIME ZONE,
   po document CLOB CONSTRAINT ensure json CHECK (po document IS JSON));
```

The following statement inserts one row, or one JSON document, into table j purchaseorder:

```
INSERT INTO j purchaseorder
  VALUES (
   SYS GUID(),
    SYSTIMESTAMP,
                          : 1600,
: "ABULL-20140421",
: "Alexis Bull",
    '{"PONumber"
      "Reference"
       "Requestor"
       "User"
                             : "ABULL",
       "CostCenter" : "A50",
       "ShippingInstructions" : {"name" : "Alexis Bull",
                                  "Address": {"street" : "200 Sporting Green",
                                                "city" : "South San Francisco",
                                                "state" : "CA",
                                                "zipCode" : 99236,
                                                "country" : "United States of America"},
                                  "Phone" : [{"type" : "Office", "number" : "909-555-7307"},
                                              {"type" : "Mobile", "number" : "415-555-1234"}]},
       "Special Instructions" : null,
       "AllowPartialShipment" : true,
       "LineItems" : [{"ItemNumber" : 1,
                        "Part" : {"Description" : "One Magic Christmas",
                                  "UnitPrice" : 19.95,
"UPCCode" : 13131092899},
                        "Quantity" : 9.0},
                       {"ItemNumber" : 2,
                        "Part" : {"Description" : "Lethal Weapon",
                                  "UnitPrice" : 19.95,
"UPCCode" : 85391628927},
                        "Quantity" : 5.0}]}');
```



Using JSON_query_column: Examples

The statement in this example queries JSON data for a specific JSON property using the JSON query column clause, and returns the property value in a column.

The statement first applies a SQL/JSON row path expression to column po_document, which results in a match to the ShippingInstructions property. The COLUMNS clause then uses the JSON guery column clause to return the Phone property value in a VARCHAR2 (100) column.

```
SELECT jt.phones

FROM j_purchaseorder,

JSON_TABLE(po_document, '$.ShippingInstructions'

COLUMNS

(phones VARCHAR2(100) FORMAT JSON PATH '$.Phone')) AS jt;

PHONES

["type":"Office", "number": "909-555-7307"}, {"type": "Mobile", "number": "415-555-1234"}]
```

Using JSON_value_column: Examples

The statement in this example refines the statement in the previous example by querying JSON data for specific JSON values using the <code>JSON_value_column</code> clause, and returns the JSON values as SQL values in relational rows and columns.

The statement first applies a SQL/JSON row path expression to column <code>po_document</code>, which results in a match to the elements in the JSON array <code>Phone</code>. These elements are JSON objects that contain two members named <code>type</code> and <code>number</code>. The statement uses the <code>COLUMNS</code> clause to return the <code>type</code> value for each object in a <code>VARCHAR2(10)</code> column called <code>phone_type</code>, and the <code>number</code> value for each object in a <code>VARCHAR2(20)</code> column called <code>phone_num</code>. The statement also returns an ordinal column named <code>row number</code>.

Using JSON_exists_column: Examples

The statements in this example test whether a JSON value exists in JSON data using the <code>JSON_exists_column</code> clause. The first example returns the result of the test as a 'true' or 'false' value in a column. The second example uses the result of the test in the <code>WHERE</code> clause.

The following statement first applies a SQL/JSON row path expression to column <code>po_document</code>, which results in a match to the entire context item, or JSON document. It then uses the <code>COLUMNS</code> clause to return the requestor's name and a string value of 'true' or 'false' indicating whether the JSON data for that requestor contains a zip code. The <code>COLUMNS</code> clause first uses the <code>JSON_value_column</code> clause to return the <code>Requestor</code> value in a <code>VARCHAR2(32)</code> column called requestor. It then uses the <code>JSON_exists_column</code> clause to determine if the <code>zipCode</code> object exists and returns the result in a <code>VARCHAR2(5)</code> column called has <code>zip</code>.



```
SELECT requestor, has_zip

FROM j_purchaseorder,

JSON_TABLE(po_document, '$'

COLUMNS

(requestor VARCHAR2(32) PATH '$.Requestor',
   has_zip VARCHAR2(5) EXISTS PATH '$.ShippingInstructions.Address.zipCode'));

REQUESTOR

HAS_ZIP

Alexis Bull

true
```

The following statement is similar to the previous statement, except that it uses the value of has zip in the WHERE clause to determine whether to return the Requestor value:

Using JSON_nested_path: Examples

The following two simple statements demonstrate the functionality of the <code>JSON_nested_path</code> clause. They operate on a simple JSON array that contains three elements. The first two elements are numbers. The third element is a nested JSON array that contains two string value elements.

The following statement does not use the <code>JSON_nested_path</code> clause. It returns the three elements in the array in a single row. The nested array is returned in its entirety.

The following statement is different from the previous statement because it uses the \(\mathref{JSON_nested_path} \) clause to return the individual elements of the nested array in individual columns in a single row along with the parent array elements.



The previous example shows how to use <code>JSON_nested_path</code> with a nested JSON array. The following example shows how to use the <code>JSON_nested_path</code> clause with a nested JSON object by returning the individual elements of the nested object in individual columns in a single row along with the parent object elements.

The following statement uses the <code>JSON_nested_path</code> clause when querying the <code>j_purchaseorder</code> table. It first applies a row path expression to column <code>po_document</code>, which results in a match to the entire context item, or <code>JSON</code> document. It then uses the <code>COLUMNS</code> clause to return the <code>Requestor</code> value in a <code>VARCHAR2(32)</code> column called <code>requestor</code>. It then uses the <code>JSON_nested_path</code> clause to return the property values of the individual objects in each member of the nested <code>Phone</code> array. Note that a row is generated for each member of the nested array, and each row contains the corresponding <code>Requestor</code> value.

```
SELECT jt.*
FROM j purchaseorder,
JSON_TABLE(po_document, '$'
COLUMNS
 (requestor VARCHAR2(32) PATH '$.Requestor',
  NESTED PATH '$.ShippingInstructions.Phone[*]'
    COLUMNS (phone type VARCHAR2(32) PATH '$.type',
          phone num VARCHAR2(20) PATH '$.number')))
AS jt;
               PHONE TYPE
                                PHONE NUM
REQUESTOR
------
Alexis Bull Office 909-555-7307
Alexis Bull
               Mobile
                                 415-555-1234
```

The following example shows the use of simple dot-notation in <code>JSON_nested_path</code> and its equivalent without dot notation.

```
SELECT c.*
FROM customer t,
JSON_TABLE(t.json COLUMNS(
id, name, phone, address,
NESTED orders[*] COLUMNS(
updated, status,
NESTED lineitems[*] COLUMNS(
description, quantity NUMBER, price NUMBER
)
)
))
);
```

The above statement in dot notation is equivalent to the following one without dot notation:

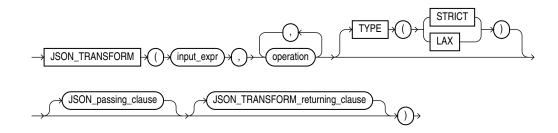
```
SELECT c.*
FROM customer t,
JSON_TABLE(t.json, '$' COLUMNS(
```



```
id PATH '$.id',
name PATH '$.name',
phone PATH '$.name',
address PATH '$.address',
NESTED PATH '$.orders[*]' COLUMNS(
updated PATH '$.updated',
status PATH '$.status',
NESTED PATH '$.lineitems[*]' COLUMNS(
description PATH '$.description',
quantity NUMBER PATH '$.quantity',
price NUMBER PATH '$.price'
)
))
)) c;
```

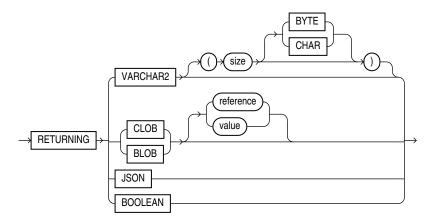
JSON_TRANSFORM

Syntax



(operation::=, JSON_TRANSFORM_returning_clause::=, JSON_passing_clause::=)

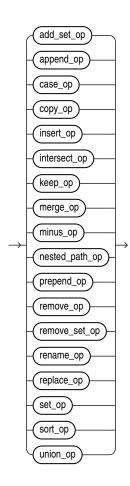
JSON_TRANSFORM_returning_clause::=



JSON_passing_clause::=

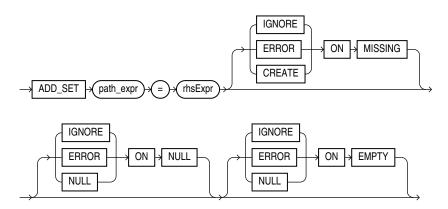
For details on JSON passing clause see JSON_EXISTS Condition.

operation ::=

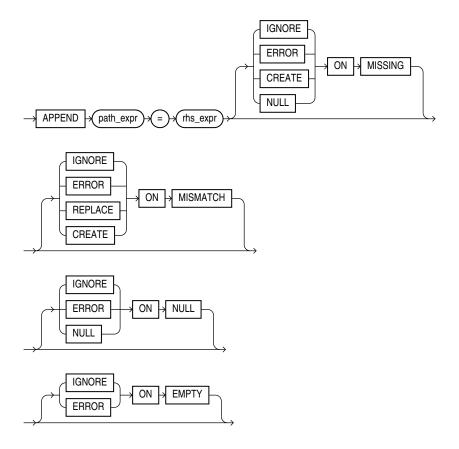


(add_set_op::=,append_op::=, case_op::=,copy_op::=,insert_op::=,intersect_op::=,keep_op::=,merge_op::=,minus_op::=,nest ed_path_op::=,prepend_op::=,remove_op::=,rename_op::=,remove_set_op::=,replace_op::=,s et_op,sort_op,union_op,)

add_set_op::=

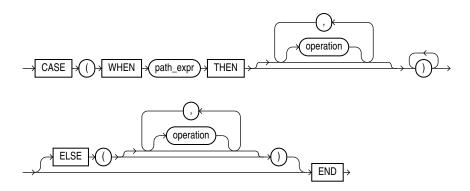


append_op ::=

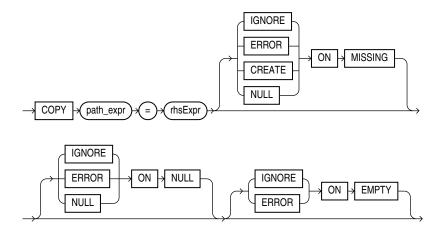


rhs_expr::=

case_op::=

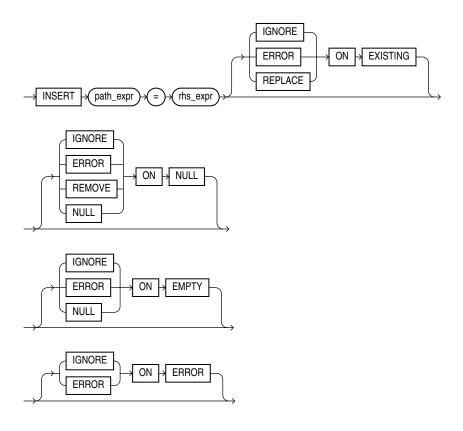


copy_op::=



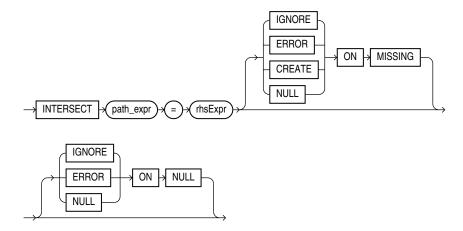
rhs_expr::=

insert_op ::=



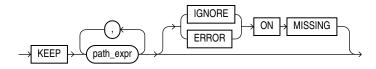
rhs_expr::=

intersect_op::=

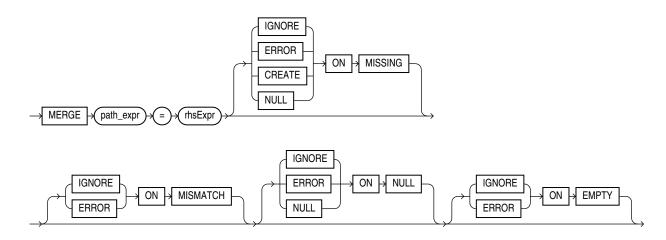


rhs_expr::=

keep_op ::=

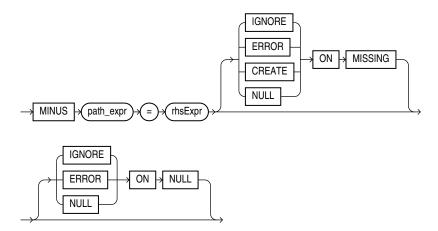


merge_op::=



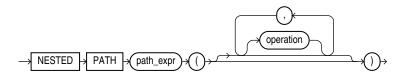
rhs_expr::=

minus_op::=

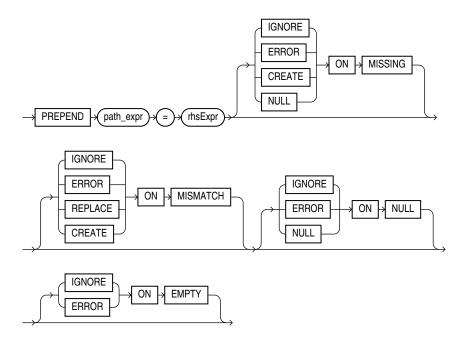


rhs_expr::=

nested_path_op::=

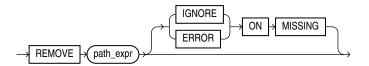


prepend_op::=

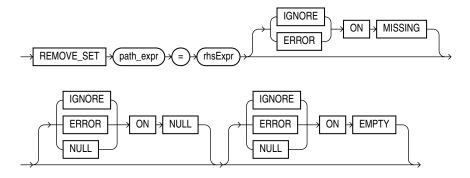


rhs_expr::=

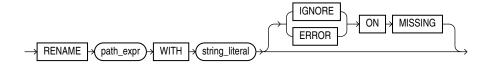
remove_op ::=



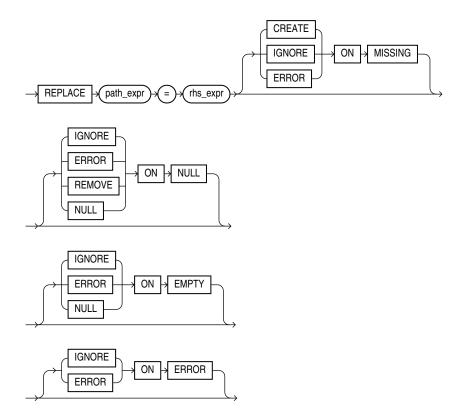
remove_set_op::=



rename_op ::=

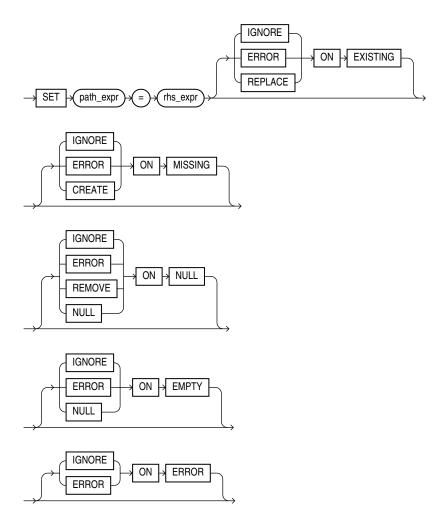


replace_op ::=



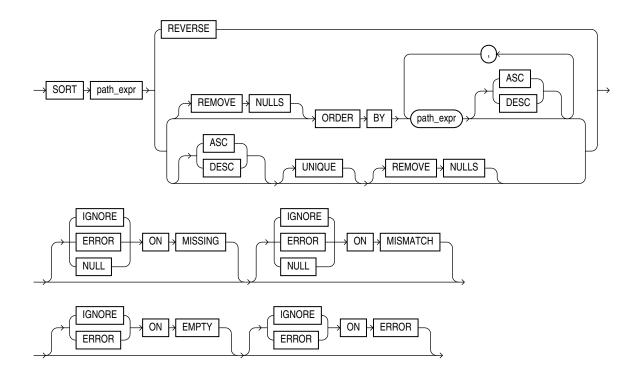
rhs_expr::=

set_op ::=

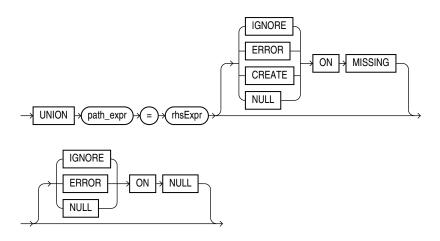


rhs_expr::=

sort_op::=

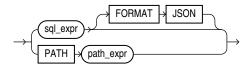


union_op::=



rhs_expr::=

rhs_expr ::=



Purpose

JSON_TRANSFORM modifies JSON documents. You specify operations to perform and SQL/JSON path expressions that target the places to modify. The operations are applied to the input data in the order specified: each operation acts on the data that results from applying all of the preceding operations.

 $\tt JSON_TRANSFORM$ either succeeds completely or not at all. If any of the specified operations raises an error, then none of the operations take effect. $\tt JSON_TRANSFORM$ returns the original data changed according to the operations specified.

You can use the <code>JSON_TRANSFORM</code> within the <code>UPDATE</code> statement to modify documents in a <code>JSON</code> column.

You can use it in a $\tt SELECT$ list, to modify the selected documents. The modified documents can be returned or processed further.

JSON_TRANSFORM can accept as input, and return as output, any SQL data type that supports JSON data: JSON, VARCHAR2, CLOB, or BLOB. Note that data type JSON is available only if database initialization parameter compatible is 20 or greater.

The default return (output) data type is the same as the input data type.



Oracle SQL Function JSON_TRANSFORM of the JSON Developer's Guide for a full discussion with examples.

JSON_TRANSFORM Operations

- Use ADD SET to add missing value to an array, as if adding an element to a set.
- Use APPEND to append the values that are specified by the RHS to the array that is targeted by the LHS path expression.

APPEND has the effect of INSERT for an array position of last+1.

An error is raised if the LHS path expression targets an existing field whose value is not an array.

If the RHS targets an array then the LHS array is updated by appending the elements of the RHS array to it, in order.

• Use Case to set conditions to perform a sequence of JSON TRANSFORM operations.

This is a control operation that conditionally applies other operations, which in turn can modify data.

The syntax is keyword CASE followed by one or more WHEN clauses, followed optionally by an ELSE clause, followed by END.

A WHEN clause is keyword WHEN followed by a path expression, followed by a THEN clause.

The path expression contains a filter condition, which checks for the existence of some data.

A THEN or an ELSE clause is keyword THEN or ELSE, respectively, followed by parentheses (()) containing zero or more JSON TRANSFORM operations.



The operations of a THEN clause are performed if the condition of its WHEN clause is satisfied. The operations of the optional ELSE clause are performed if the condition of no WHEN clause is satisfied.

The syntax of the JSON_TRANSFORM CASE operation is thus essentially the same as an Oracle SQL searched CASE expression, except that it is the predicate that is tested and the resulting effect of each THEN/ELSE branch.

For SQL, the predicate tested is a SQL comparison. For <code>JSON_TRANSFORM</code>, the predicate is a path expression that checks for the existence of some data. (The check is essentially done using <code>JSON_EXISTS</code>.)

For SQL, each <code>THEN/ELSE</code> branch holds a SQL expression to evaluate, and its value is returned as the result of the <code>CASE</code> expression. For json_transform, each <code>THEN/ELSE</code> branch holds a (parenthesized) sequence of <code>JSON_TRANSFORM</code> operations, which are performed in order.

The conditional path expressions of the WHEN clauses are tested in order, until one succeeds (those that follow are not tested). The THEN operations for the successful WHEN test are then performed, in order.

- Use COPY to replace the elements of the array that is targeted by the LHS path expression
 with the values that are specified by the RHS. An error is raised if the LHS path expression
 does not target an array. The operation can accept a sequence of multiple values matched
 by the RHS path expression.
- INSERT Insert the value of the specified SQL expression at the location that's targeted by
 the specified path expression that follows the equal sign (=), which must be either the field
 of an object or an array position (otherwise, an error is raised). By default, an error is
 raised if a targeted object field already exists.

INSERT for an object field has the effect of SET with clause CREATE ON MISSING (default for SET), except that the default behavior for ON EXISTING is ERROR, not REPLACE.)

You can specify an array position past the current end of an array. In that case, the array is lengthened to accommodate insertion of the value at the indicated position, and the intervening positions are filled with JSON null values.

For example, if the input JSON data is {"a":["b"]} then INSERT '\$.a[3]'=42 returns {"a":["b", null, null 42]} as the modified data. The elements at array positions 1 and 2 are null.

- Use INTERSECT to remove all elements of the array that is targeted by the LHS path
 expression that are not equal to any value specified by the RHS. Remove any duplicate
 elements. Note that this is a set operation. The order of all array elements is undefined
 after the operation.
- Use MERGE to add specified fields (name and value) matched by the RHS path expression
 to the object that is targeted by the LHS path expression. Ignore any fields specified by the
 RHS that are already in the targeted LHS object. If the same field is specified more than
 once by the RHS then use only the last one in the sequence of matches.
- Use MINUS to remove all elements of the array that is targeted by the LHS path expression that are equal to a value specified by the RHS. Remove any duplicate elements. Note that this is a set operation. The order of all array elements is undefined after the operation.
- KEEP removes all parts of the input data that are not targeted by at least one of the specified path expressions. A topmost object or array is not removed, it is emptied and becomes an empty object ({}) or array ([]).
- Use NESTED PATH to define a scope (a particular part of your data) in which to apply a sequence of operations.



 Use PREPEND to prepend the values that are specified by the RHS to the array that is targeted by the LHS path expression. The operation can accept a sequence of multiple values matched by the RHS path expression.

An error is raised if the LHS path expression targets an existing field whose value is not an array.

When prepending a single value, PREPEND has the effect of INSERT for an array position of 0

If the RHS targets an array then the LHS array is updated by prepending the elements of the RHS array to it, in order.

- REMOVE Remove the input data that's targeted by the specified path expression. An error is
 raised if you try to remove all of the data, for example you cannot use REMOVE '\$'. By
 default, no error is raised if the targeted data does not exist (IGNORE ON MISSING).
- Use REMOVE_SET to remove all occurrences of a value from an array, as if removing an
 element from a set.
- RENAME renames the field that is targeted by the specified path expression to the value of the SQL expression that follows the equal sign (=). By default, no error is raised if the targeted field does not exist (IGNORE ON MISSING).
- REPLACE replaces the data that's targeted by the specified path expression with the value of
 the specified SQL expression that follows the equal sign (=). By default, no error is raised if
 the targeted data does not exist (IGNORE ON MISSING).

REPLACE has the effect of SET with clause IGNORE ON MISSING.

- SET Set what the LHS specifies to the value specified by what follows the equal sign (=).
 The LHS can be either a SQL/JSON variable or a path expression that targets data. If the
 RHS is a SQL expression then its value is assigned to the LHS variable. When the LHS
 specifies a path expression, the default behavior is to replace existing targeted data with
 the new value, or insert the new value at the targeted location if the path expression
 matches nothing. (See operator INSERT about inserting an array element past the end of
 the array.)
- When the LHS specifies a SQL/JSON variable, the variable is dynamically assigned to
 whatever is specified by the RHS. (The variable is created if it does not yet exist.) The
 variable continues to have that value until it is set to a different value by a subsequent SET
 operation (in the same JSON TRANSFORM invocation).

If the RHS is a path expression then its targeted data is assigned to the variable.

Setting a variable is a control operation; it can affect how subsequent operations modify data, but it does not, itself, directly modify data.

When the LHS specifies a path expression, the default behavior is like that of SQL UPSERT: replace existing targeted data with the new value, or insert the new value at the targeted location if the path expression matches nothing. (See operator INSERT about inserting an array element past the end of the array.)

- SORT sorts the elements of the array targeted by the specified path. The result includes all
 elements of the array (none are dropped); the only possible change is that they are
 reordered.
- Use UNION to add the values specified by the RHS to the array that is targeted by the LHS
 path expression. Remove any duplicate elements. The operation can accept a sequence of
 multiple values matched by the RHS path expression. Note that this is a set operation. The
 order of all array elements is undefined after the operation.



TYPE Clause

For a full discussion of STRICT and LAX syntax see About Strict and Lax JSON Syntax, and TYPE Clause for SQL Functions and Conditions

JSON_passing_clause

You can use <code>JSON_passing_clause</code> to specify SQL bindings of bind variables to SQL/JSON variables similar to the <code>JSON_EXISTS</code> condition and the SQL/JSON query functions.

JSON_TRANSFORM_returning_clause

After you specify the operations you can use <code>JSON_TRANSFORM_returning_clause</code> to specify the return data type.

Examples

Example 1: Update a JSON Column with a Timestamp

UPDATE t SET jcol = JSON TRANSFORM(jcol, SET '\$.lastUpdated' = SYSTIMESTAMP)

Example 2: Remove a Social Security Number before Shipping JSON to a Client

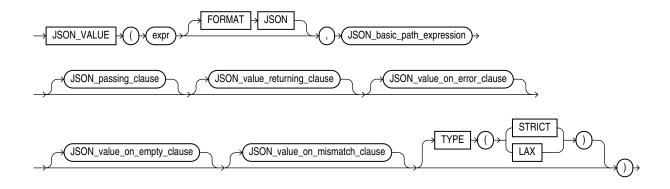
SELECT JSON TRANSFORM (jcol, REMOVE '\$.ssn') FROM t WHERE ...

JSON_TRANSFORM_returning_clause

If the input data is JSON, then the output data type is also JSON. For all other input types, the default output data type is VARCHAR2 (4000).

JSON_VALUE

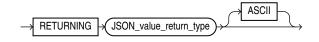
Syntax



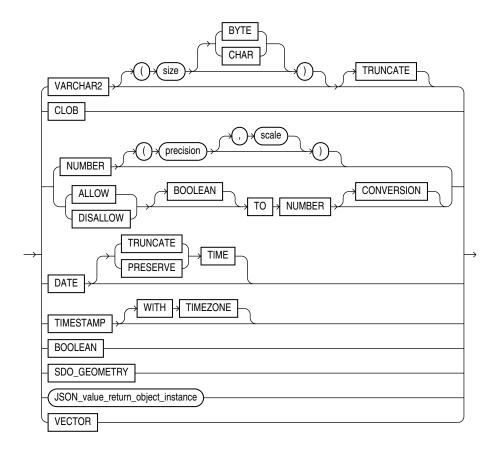
JSON_basic_path_expression::=

(JSON_basic_path_expression: See SQL/JSON Path Expressions)

JSON_value_returning_clause::=



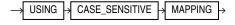
JSON_value_return_type::=



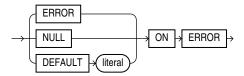
JSON_value_return_object_instance ::=



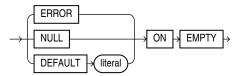
JSON_value_mapper_clause ::=



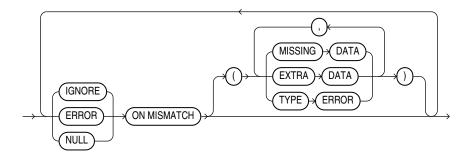
JSON_value_on_error_clause::=



JSON_value_on_empty_clause::=



JSON_value_on_mismatch_clause::=



Purpose

SQL/JSON function JSON_VALUE selects JSON data and returns a SQL scalar or an instance of a user-defined SQL object type or SQL collection type (varray, nested table)



- JSON TRANSFORM of the JSON Developer's Guide.
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the value returned by this function when it is a character value

expr

The first argument to <code>JSON_VALUE expr</code> is a SQL expression that returns an instance of a scalar SQL data type (that is, not an object or collection data type). A scalar value returned from <code>JSON_VALUE</code> can be of any of these data types: <code>BINARY_DOUBLE</code>, <code>BINARY_FLOAT</code>, <code>BOOLEAN</code>, <code>CHAR</code>, <code>CLOB</code>, <code>DATE</code>, <code>INTERVAL</code> DAY TO SECOND, <code>INTERVAL</code> YEAR TO MONTH, <code>NCHAR</code>, <code>NCLOB</code>, <code>NVARCHAR2</code>, <code>NUMBER</code>, <code>RAW1</code>, <code>SDO_GEOMETRY</code>, <code>TIMESTAMP</code>, <code>TIMESTAMP</code> WITH <code>TIME ZONE</code>, <code>VARCHAR2</code>, and <code>VECTOR</code>.

If expr is not a text literal of well-formed JSON data using strict or lax syntax, then the function returns null by default. You can use the $JSON_value_on_error_clause$ to override this default behavior. Refer to the $JSON_value_on_error_clause$.

FORMAT JSON

You must specify FORMAT JSON if expr is a column of data type BLOB.



JSON_basic_path_expression

Use this clause to specify a SQL/JSON path expression. The function uses the path expression to evaluate expr and find a scalar JSON value that matches, or satisfies, the path expression. The path expression must be a text literal. See *Oracle Database JSON Developer's Guide* for the full semantics of JSON basic path expression.

JSON_value_returning_clause

Use this clause to specify the data type and format of the value returned by this function.

RETURNING

Use the RETURNING clause to specify the data type of the return value. If you omit this clause, then JSON VALUE returns a value of type VARCHAR2 (4000).

JSON_value_return_type ::=

You can use JSON value return type to specify the following data types:

VARCHAR2[(size [BYTE, CHAR])]

If you specify this data type, then the scalar value returned by this function can be a character or number value. A number value will be implicitly converted to a VARCHAR2. When specifying the VARCHAR2 data type elsewhere in SQL, you are required to specify a size. However, in this clause you can omit the size. In this case, JSON_VALUE returns a value of type VARCHAR2 (4000).

Specify the optional TRUNCATE clause immediately after VARCHAR2 (N) to truncate the return value to N characters, if the return value is greater than N characters.

Notes on the TRUNCATE clause:

- If the string value is too long, then ORA-40478 is raised.
- If TRUNCATE is present, and the return value is not a character type, then a compile time error is raised.
- If TRUNCATE is present with FORMAT JSON, then the return value may contain data that is not syntactically correct JSON.
- TRUNCATE does not work with EXISTS.
- CLOB

Specify this data type to return a character large object containing single-byte or multi-byte characters.

NUMBER[(precision [, scale])]

If you specify this data type, then the scalar value returned by this function must be a number value. The scalar value returned can also be a JSON Boolean value. Note however, that returning NUMBER for a JSON Boolean value is deprecated.

DATE

If you specify this data type, then the scalar value returned by this function must be a character value that can be implicitly converted to a DATE data type. If the JSON input represents a date with a time component, specify DATE PRESERVE TIME to retain the time component. If you do not want to retain the time component, specify DATE TRUNCATE TIME.



If you specify neither PRESERVE TIME nor TRUNCATE TIME, the time component is not preserved.

TIMESTAMP

If you specify this data type, then the scalar value returned by this function must be a character value that can be implicitly converted to a TIMESTAMP data type.

TIMESTAMP WITH TIME ZONE

If you specify this data type, then the scalar value returned by this function must be a character value that can be implicitly converted to a TIMESTAMP WITH TIME ZONE data type.

BOOLEAN

Specify BOOLEAN to return true, false, or unknown.

SDO GEOMETRY

This data type is used for Oracle Spatial and Graph data. If you specify this data type, then expr must evaluate to a text literal containing GeoJSON data, which is a format for encoding geographic data in JSON. If you specify this data type, then the scalar value returned by this function must be an object of type SDO GEOMETRY.

JSON value return object instance

If JSON_VALUE targets a JSON object, and you specify a user-defined SQL object type as the return type, then JSON_VALUE returns an instance of that object type in object_type_name.

For examples see Using JSON_VALUE To Instantiate a User-Defined Object Type Instance of the JSON Developer's Guide.

See Also:

- SQL/JSON Function JSON_VALUE for a conceptual understanding in the JSON Developer's Guide.
- Refer to "Data Types" for more information on the preceding data types.
- If the data type is not large enough to hold the return value, then this function returns null by default. You can use the <code>JSON_value_on_error_clause</code> to override this default behavior. Refer to the <code>JSON_value_on_error_clause</code>.

ASCII

Specify ASCII to automatically escape any non-ASCII Unicode characters in the return value, using standard ASCII Unicode escape sequences.

JSON_value_on_error_clause

Use this clause to specify the value returned by this function when the following errors occur:

- expr is not well-formed JSON data using strict or lax JSON syntax
- A nonscalar value is found when the JSON data is evaluated using the SQL/JSON path expression
- No match is found when the JSON data is evaluated using the SQL/JSON path expression. You can override the behavior for this type of error by specifying the JSON value on empty clause.



The return value data type is not large enough to hold the return value

You can specify the following clauses:

- NULL ON ERROR Returns null when an error occurs. This is the default.
- ERROR ON ERROR Returns the appropriate Oracle error when an error occurs.
- DEFAULT literal ON ERROR Returns literal when an error occurs. The data type of literal must match the data type of the value returned by this function.

JSON_value_on_empty_clause

Use this clause to specify the value returned by this function if no match is found when the JSON data is evaluated using the SQL/JSON path expression. This clause allows you to specify a different outcome for this type of error than the outcome specified with the JSON value on error clause.

You can specify the following clauses:

- NULL ON EMPTY Returns null when no match is found.
- ERROR ON EMPTY Returns the appropriate Oracle error when no match is found.
- DEFAULT literal ON EMPTY Returns literal when no match is found. The data type of literal must match the data type of the value returned by this function.

If you omit this clause, then the <code>JSON_value_on_error_clause</code> determines the value returned when no match is found.

JSON value on mismatch clause

The <code>JSON_value_on_mismatch_clause</code> applies when a type conversion fails, for example when you try to convert a JSON number to a SQL date.

If the return type of <code>JSON_VALUE</code> is a SQL scalar like <code>NUMBER</code> or <code>DATE</code>, then <code>ON MISMATCH</code> applies for all type conversion errors - no further specification is required. ERROR and <code>NULL</code> are valid options.

Example 1

```
select json_value( '{a:"cat"}','$.a.number()' NULL ON
EMPTY
   ERROR ON MISMATCH DEFAULT -1 ON ERROR ) from dual;
   ORA-01722: invalid number
```

If the return type is an object type, then ON MISMATCH can be further specified with MISSING DATA, EXTRA DATA and TYPE ERROR. You can use it generally to apply to all error cases, or you can use it case by case by specifying different ON MISMATCH clauses for each case.

Example 2

```
IGNORE ON MISMATCH (EXTRA DATA)

ERROR ON MISMATCH (MISSING DATA, TYPE ERROR)
```

The option IGNORE is only valid when the return type is an object type.



TYPE Clause

For a full discussion of STRICT and LAX syntax see About Strict and Lax JSON Syntax, and TYPE Clause for SQL Functions and Conditions

Examples

The following query returns the value of the member with property name a. Because the RETURNING clause is not specified, the value is returned as a VARCHAR2 (4000) data type:

```
SELECT JSON_VALUE('{a:100}', '$.a') AS value
  FROM DUAL;

VALUE
----
100
```

The following query returns the value of the member with property name a. Because the RETURNING NUMBER clause is specified, the value is returned as a NUMBER data type:

The following query returns the value of the member with property name b, which is in the value of the member with property name a:

```
SELECT JSON_VALUE('{a:{b:100}}', '$.a.b') AS value
FROM DUAL;

VALUE
----
100
```

The following query returns the value of the member with property name d in any object:

```
SELECT JSON_VALUE('{a:{b:100}, c:{d:200}, e:{f:300}}', '$.*.d') AS value
FROM DUAL;
VALUE
----
200
```

The following query returns the value of the first element in an array:

```
SELECT JSON_VALUE('[0, 1, 2, 3]', '$[0]') AS value
FROM DUAL;

VALUE
-----
0
```

The following query returns the value of the third element in an array. The array is the value of the member with property name a.

```
SELECT JSON_VALUE('{a:[5, 10, 15, 20]}', '$.a[2]') AS value
FROM DUAL;
```



```
VALUE
```

The following query returns the value of the member with property name a in the second object in an array:

```
SELECT JSON_VALUE('[{a:100}, {a:200}, {a:300}]', '$[1].a') AS value
   FROM DUAL;

VALUE
----
200
```

The following query returns the value of the member with property name $\,c\,$ in any object in an array:

```
SELECT JSON_VALUE('[{a:100}, {b:200}, {c:300}]', '$[*].c') AS value
   FROM DUAL;

VALUE
----
300
```

The following query attempts to return the value of the member that has property name lastname. However, such a member does not exist in the specified JSON data, resulting in no match. Because the ON ERROR clause is not specified, the statement uses the default NULL ON ERROR and returns null.

```
SELECT JSON_VALUE('{firstname:"John"}', '$.lastname') AS "Last Name"
FROM DUAL;
Last Name
```

The following query results in an error because it attempts to return the value of the member with property name lastname, which does not exist in the specified JSON. Because the ON ERROR clause is specified, the statement returns the specified text literal.

```
SELECT JSON_VALUE('{firstname:"John"}', '$.lastname'

DEFAULT 'No last name found' ON ERROR) AS "Last Name"

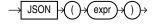
FROM DUAL;

Last Name

-----
No last name found
```

JSON Type Constructor

Syntax





Purpose

The JSON data type constructor, JSON, takes as input a textual JSON value (a scalar, object, or array), parses it, and returns the value as an instance of JSON type. Alternatively, the input can be an instance of SQL type VECTOR, a user-defined PL/SQL type, or a SQL aggregate type.

You can use the JSON data type constructor JSON to parse textual JSON input (a scalar, object, or array), and return it as an instance of type JSON.

Input values must pass the IS JSON test. Input values that fail the IS JSON test are rejected with a syntax error.

To filter out duplicate input values, you must run the IS JSON (WITH UNIQUE KEYS) check on the textual JSON input before using the JSON constructor.

Prerequisites

You can use the constructor JSON only if database initialization parameter compatible is atleast 20.



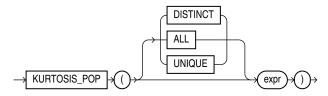
JSON Data Type Constructor of the JSON Developer's Guide.

expr

The input in expr must be a syntactically valid textual representation of type VARCHAR2, CLOB and BLOB. It can also be a literal SQL string. A SQL NULL input value results in a JSON type instance of SQL NULL.

KURTOSIS POP

Syntax



Purpose

The population kurtosis function KURTOSIS_POP is primarily used to determine the characteristics of outliers in a given distribution.

NULL values in expr are ignored.

Returns NULL if all rows in the group have NULL expr values.

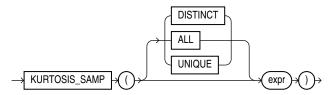
Returns 0 if there are one or two rows in expr.

For a given set of values, the result of population kurtosis (KURTOSIS_POP) and sample kurtosis (KURTOSIS_SAMP) are always deterministic. However, the values of KURTOSIS_POP and

KURTOSIS_SAMP differ. As the number of values in the data set increases, the difference between the computed values of KURTOSIS SAMP and KURTOSIS POP decreases.

KURTOSIS_SAMP

Syntax



Purpose

The sample kurtosis function <code>KURTOSIS_SAMP</code> is primarily used to determine the characteristics of outliers in a given distribution.

NULL values in expr are ignored.

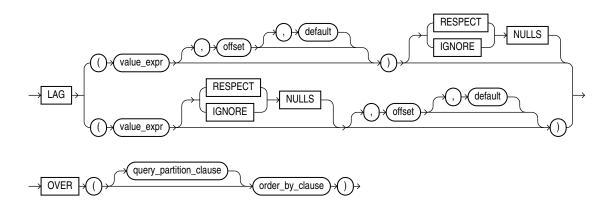
Returns NULL if all rows in the group have NULL expr values.

Returns 0 if there are one or two rows in expr.

For a given set of values, the result of sample kurtosis (KURTOSIS_SAMP) and population kurtosis (KURTOSIS_POP) are always deterministic. However, the values of KURTOSIS_SAMP and KURTOSIS_POP differ. As the number of values in the data set increases, the difference between the computed values of KURTOSIS_SAMP and KURTOSIS_POP decreases.

LAG

Syntax



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions, including valid forms of $value_expr$



Purpose

LAG is an analytic function. It provides access to more than one row of a table at the same time without a self join. Given a series of rows returned from a query and a position of the cursor, LAG provides access to a row at a given physical offset prior to that position.

For the optional <code>offset</code> argument, specify an integer that is greater than zero. If you do not specify <code>offset</code>, then its default is 1. The optional <code>default</code> value is returned if the offset goes beyond the scope of the window. If you do not specify <code>default</code>, then its default is null.

{RESPECT | IGNORE} NULLS determines whether null values of value_expr are included in or eliminated from the calculation. The default is RESPECT NULLS.

You cannot nest analytic functions by using LAG or any other analytic function for $value_expr$. However, you can use other built-in function expressions for $value_expr$.

See Also:

- "About SQL Expressions" for information on valid forms of expr and LEAD
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of LAG when it is a character value

Examples

The following example provides, for each purchasing clerk in the employees table, the salary of the employee hired just before:

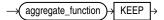
HIRE_DATE	LAST_NAME	SALARY	PREV_SAL
18-MAY-03	Khoo	3100	0
24-JUL-05	Tobias	2800	3100
24-DEC-05	Baida	2900	2800
15-NOV-06	Himuro	2600	2900
10-AUG-07	Colmenares	2500	2600

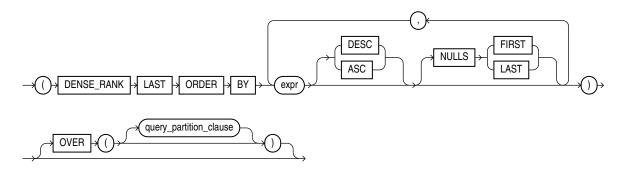


LAST

Syntax

last::=





See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions of the $query_partitioning_clause$

Purpose

FIRST and LAST are very similar functions. Both are aggregate and analytic functions that operate on a set of values from a set of rows that rank as the FIRST or LAST with respect to a given sorting specification. If only one row ranks as FIRST or LAST, then the aggregate operates on the set with only one element.

Refer to FIRST for complete information on this function and for examples of its use.

LAST_DAY

Syntax



Purpose

LAST_DAY returns the date of the last day of the month that contains date. The last day of the month is defined by the session parameter NLS_CALENDAR. The return type is always DATE, regardless of the data type of date.



Examples

The following statement determines how many days are left in the current month.

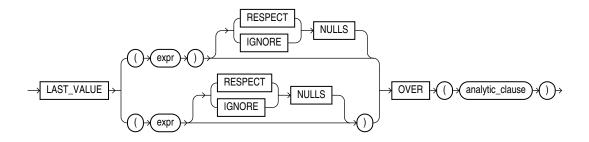
The following example adds 5 months to the hire date of each employee to give an evaluation date:

```
SELECT last_name, hire_date,
        TO_CHAR(ADD_MONTHS(LAST_DAY(hire_date), 5)) "Eval Date"
FROM employees
ORDER BY last_name, hire_date;
```

LAST_NAME	HIRE_DATE	Eval Date
Abel	11-MAY-04	31-OCT-04
Ande	24-MAR-08	31-AUG-08
Atkinson	30-OCT-05	31-MAR-06
Austin	25-JUN-05	30-NOV-05
Baer	07-JUN-02	30-NOV-02
Baida	24-DEC-05	31-MAY-06
Banda	21-APR-08	30-SEP-08
Bates	24-MAR-07	31-AUG-07

LAST_VALUE

Syntax



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions, including valid forms of expr

Purpose

LAST_VALUE is an analytic function that is useful for data densification. It returns the last value in an ordered set of values.



The two forms of this syntax have the same behavior. The top branch is the ANSI format, which Oracle recommends for ANSI compatibility.

{RESPECT | IGNORE} NULLS determines whether null values of <code>expr</code> are included in or eliminated from the calculation. The default is <code>RESPECT NULLS</code>. If the last value in the set is null, then the function returns <code>NULL</code> unless you specify <code>IGNORE NULLS</code>. If you specify <code>IGNORE NULLS</code>, then <code>LAST_VALUE</code> returns the last non-null value in the set, or <code>NULL</code> if all values are null. Refer to "Using Partitioned Outer Joins: Examples" for an example of data densification.

You cannot nest analytic functions by using LAST_VALUE or any other analytic function for expr. However, you can use other built-in function expressions for expr. Refer to "About SQL Expressions" for information on valid forms of expr.

If you omit the <code>windowing_clause</code> of the <code>analytic_clause</code>, it defaults to <code>RANGE BETWEEN</code> UNBOUNDED PRECEDING AND CURRENT ROW. This default sometimes returns an unexpected value, because the last value in the window is at the bottom of the window, which is not fixed. It keeps changing as the current row changes. For expected results, specify the <code>windowing_clause</code> as <code>RANGE BETWEEN UNBOUNDED PRECEDING</code> AND UNBOUNDED FOLLOWING. Alternatively, you can specify the <code>windowing_clause</code> as <code>RANGE BETWEEN CURRENT ROW</code> AND UNBOUNDED FOLLOWING.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value

Examples

The following example returns, for each row, the hire date of the employee earning the lowest salary:

```
SELECT employee_id, last_name, salary, hire_date,
     LAST VALUE(hire date)
       OVER (ORDER BY salary DESC ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED
           FOLLOWING) AS lv
 FROM (SELECT * FROM employees
        WHERE department id = 90
        ORDER BY hire date);
EMPLOYEE ID LAST NAME
                            SALARY HIRE DATE LV
 ------ ----- -----
                                    24000 17-JUN-03 13-JAN-01
      100 King
      101 Kochhar
                                    17000 21-SEP-05 13-JAN-01
      102 De Haan
                                    17000 13-JAN-01 13-JAN-01
```



This example illustrates the nondeterministic nature of the LAST_VALUE function. Kochhar and De Haan have the same salary, so they are in adjacent rows. Kochhar appears first because the rows in the subquery are ordered by hire_date. However, if the rows are ordered by hire_date in descending order, as in the next example, then the function returns a different value:

```
SELECT employee id, last name, salary, hire date,
      LAST VALUE (hire date)
       OVER (ORDER BY salary DESC ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED
            FOLLOWING) AS lv
 FROM (SELECT * FROM employees
        WHERE department id = 90
        ORDER BY hire date DESC);
                                 SALARY HIRE DATE LV
EMPLOYEE ID LAST NAME
100 King
                                   24000 17-JUN-03 21-SEP-05
      102 De Haan
                                  17000 13-JAN-01 21-SEP-05
      101 Kochhar
                                   17000 21-SEP-05 21-SEP-05
```

The following two examples show how to make the LAST_VALUE function deterministic by ordering on a unique key. By ordering within the function by both salary and the unique key employee id, you can ensure the same result regardless of the ordering in the subquery.

```
SELECT employee id, last name, salary, hire date,
     LAST VALUE (hire date)
       OVER (ORDER BY salary DESC, employee id ROWS BETWEEN UNBOUNDED PRECEDING
           AND UNBOUNDED FOLLOWING) AS 1v
 FROM (SELECT * FROM employees
       WHERE department id = 90
        ORDER BY hire date);
EMPLOYEE ID LAST NAME
                          SALARY HIRE DATE LV
     100 King 24000 17-JUN-03 13-JAN-01
      101 Kochhar
                                    17000 21-SEP-05 13-JAN-01
      102 De Haan
                                    17000 13-JAN-01 13-JAN-01
SELECT employee id, last name, salary, hire date,
      LAST VALUE (hire date)
       OVER (ORDER BY salary DESC, employee id ROWS BETWEEN UNBOUNDED PRECEDING
            AND UNBOUNDED FOLLOWING) AS 1v
 FROM (SELECT * FROM employees
        WHERE department id = 90
        ORDER BY hire date DESC);
EMPLOYEE ID LAST NAME
                                  SALARY HIRE DATE LV
100 King
                                    24000 17-JUN-03 13-JAN-01
      101 Kochhar
                                   17000 21-SEP-05 13-JAN-01
      102 De Haan
                                    17000 13-JAN-01 13-JAN-01
```

The following two examples show that the LAST_VALUE function is deterministic when you use a logical offset (RANGE instead of ROWS). When duplicates are found for the ORDER BY expression, the LAST_VALUE is the highest value of expr:

```
SELECT employee_id, last_name, salary, hire_date,

LAST_VALUE(hire_date)

OVER (ORDER BY salary DESC RANGE BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) AS 1v
```



```
FROM (SELECT * FROM employees
          WHERE department_id = 90
          ORDER BY hire_date);
```

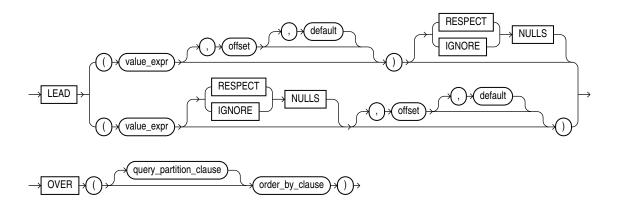
EMPLOYEE_ID	LAST_NAME	SALARY	HIRE_DATE	LV
100	King	24000	17-JUN-03	21-SEP-05
102	De Haan	17000	13-JAN-01	21-SEP-05
101	Kochhar	17000	21-SEP-05	21-SEP-05

WHERE department_id = 90
ORDER BY hire date **DESC**);

EMPLOYEE_ID	LAST_NAME	SALARY	HIRE_DATE	LV
100	King	24000	17-JUN-03	21-SEP-05
102	De Haan	17000	13-JAN-01	21-SEP-05
101	Kochhar	17000	21-SEP-05	21-SEP-05

LEAD

Syntax



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions, including valid forms of $value\ expr$

Purpose

LEAD is an analytic function. It provides access to more than one row of a table at the same time without a self join. Given a series of rows returned from a query and a position of the cursor, LEAD provides access to a row at a given physical offset beyond that position.

If you do not specify <code>offset</code>, then its default is 1. The optional <code>default</code> value is returned if the offset goes beyond the scope of the table. If you do not specify <code>default</code>, then its default value is null.

{RESPECT | IGNORE} NULLS determines whether null values of value_expr are included in or eliminated from the calculation. The default is RESPECT NULLS.

You cannot nest analytic functions by using LEAD or any other analytic function for $value_expr$. However, you can use other built-in function expressions for $value_expr$.

See Also:

- "About SQL Expressions" for information on valid forms of expr and LAG
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of LEAD when it is a character value

Examples

The following example provides, for each employee in Department 30 in the employees table, the hire date of the employee hired just after:

```
SELECT hire_date, last_name,

LEAD(hire_date, 1) OVER (ORDER BY hire_date) AS "NextHired"

FROM employees

WHERE department_id = 30

ORDER BY hire_date;

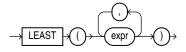
HIRE_DATE_LAST_NAME

Next Hired
```

UIVE DAIE	LASI_NAME	Next Hite
07-DEC-02	Raphaely	18-MAY-03
18-MAY-03	Khoo	24-JUL-05
24-JUL-05	Tobias	24-DEC-05
24-DEC-05	Baida	15-NOV-06
15-NOV-06	Himuro	10-AUG-07
10-AUG-07	Colmenares	

LEAST

Syntax



Purpose

LEAST returns the least of a list of one or more expressions. Oracle Database uses the first expr to determine the return type. If the first expr is numeric, then Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type before the comparison, and returns that data type. If the first expr is not numeric, then each expr after the first is implicitly converted to the data type of the first expr before the comparison.



Oracle Database compares each expr using nonpadded comparison semantics. The comparison is binary by default and is linguistic if the NLS_COMP parameter is set to LINGUISTIC and the NLS_SORT parameter has a setting other than BINARY. Character comparison is based on the numerical codes of the characters in the database character set and is performed on whole strings treated as one sequence of bytes, rather than character by character. If the value returned by this function is character data, then its data type is VARCHAR2 if the first expr is a character data type and NVARCHAR2 if the first expr is a national character data type.

See Also:

- "Data Type Comparison Rules" for more information on character comparison
- Table 2-9 for more information on implicit conversion and "Floating-Point Numbers" for information on binary-float comparison semantics
- "GREATEST", which returns the greatest of a list of one or more expressions
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation LEAST uses to compare character values for expr, and for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value

Examples

The following statement selects the string with the least value:

```
SELECT LEAST('HARRY','HARRIOT','HAROLD') "Least"
FROM DUAL;

Least
-----
HAROLD
```

In the following statement, the first argument is numeric. Oracle Database determines that the argument with the highest numeric precedence is the third argument, converts the remaining arguments to the data type of the third argument, and returns the least value as that data type:

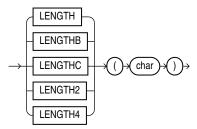
```
SELECT LEAST (1, '2.1', '.000832') "Least"
   FROM DUAL;
Least
-----
.000832
```



LENGTH

Syntax

length::=



Purpose

The LENGTH functions return the length of char. LENGTH calculates length using characters as defined by the input character set. LENGTHB uses bytes instead of characters. LENGTHC uses Unicode complete characters. LENGTH2 uses UCS2 code points. LENGTH4 uses UCS4 code points.

char can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The exceptions are LENGTHC, LENGTH2, and LENGTH4, which do not allow char to be a CLOB or NCLOB. The return value is of data type NUMBER. If char has data type CHAR, then the length includes all trailing blanks. If char is null, then this function returns null.

For more on character length see the following:

- Oracle Database Globalization Support Guide
- Oracle Database SecureFiles and Large Objects Developer's Guide

Restriction on LENGTHB

The LENGTHB function is supported for single-byte LOBs only. It cannot be used with CLOB and NCLOB data in a multibyte character set.

Examples

The following example uses the LENGTH function using a single-byte database character set:

```
SELECT LENGTH('CANDIDE') "Length in characters" FROM DUAL;

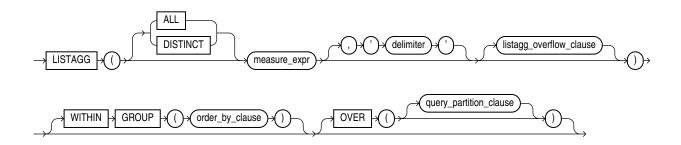
Length in characters
```

The next example assumes a double-byte database character set.



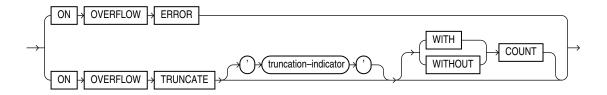
LISTAGG

Syntax



(listagg_overflow_clause::=, order_by_clause::=, query_partition_clause::=)

listagg_overflow_clause::=



See Also:

"Analytic Functions " for information on syntax, semantics, and restrictions of the ORDER BY clause and OVER clause

Purpose

For a specified measure, LISTAGG orders data within each group specified in the ORDER BY clause and then concatenates the values of the measure column.

- As a single-set aggregate function, LISTAGG operates on all rows and returns a single output row.
- As a group-set aggregate, the function operates on and returns an output row for each group defined by the GROUP BY clause.
- As an analytic function, LISTAGG partitions the query result set into groups based on one or more expression in the *query_partition_clause*.

The arguments to the function are subject to the following rules:

- The ALL keyword is optional and is provided for semantic clarity.
- The measure_expr is the measure column and can be any expression. Null values in the measure column are ignored.

• The *delimiter* designates the string that is to separate the measure column values. This clause is optional and defaults to NULL.

If measure_expr is of type RAW, then the delimiter must be of type RAW. You can achieve this by specifying the delimiter as a character string that can be implicitly converted to RAW, or by explicitly converting the delimiter to RAW, for example, using the UTL_RAW.CAST_TO_RAW function.

- The <code>order_by_clause</code> determines the order in which the concatenated values are returned. The function is deterministic only if the <code>ORDER BY</code> column list achieved unique ordering.
- If you specify <code>order_by_clause</code>, you must also specify <code>WITHIN GROUP</code> and vice versa. These two clauses must be specified together or not at all.

The DISTINCT keyword removes duplicate values from the list.

If the measure column is of type RAW, then the return data type is RAW. Otherwise, the return data type is VARCHAR2.

The maximum length of the return data type depends on the value of the MAX_STRING_SIZE initialization parameter. If MAX_STRING_SIZE = EXTENDED, then the maximum length is 32767 bytes for the VARCHAR2 and RAW data types. If MAX_STRING_SIZE = STANDARD, then the maximum length is 4000 bytes for the VARCHAR2 data type and 2000 bytes for the RAW data type. A final delimiter is not included when determining if the return value fits in the return data type.

See Also:

- Extended Data Types for more information on the MAX_STRING_SIZE initialization parameter
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of LISTAGG
- Database Data Warehousing Guide for details.

listagg_overflow_clause

This clause controls how the function behaves when the return value exceeds the maximum length of the return data type.

ON OVERFLOW ERROR If you specify this clause, then the function returns an ORA-01489 error. This is the default.

ON OVERFLOW TRUNCATE If you specify this clause, then the function returns a truncated list of measure values.

• The truncation_indicator designates the string that is to be appended to the truncated list of measure values. If you omit this clause, then the truncation indicator is an ellipsis (...).

If <code>measure_expr</code> is of type <code>RAW</code>, then the truncation indicator must be of type <code>RAW</code>. You can achieve this by specifying the truncation indicator as a character string that can be implicitly converted to <code>RAW</code>, or by explicitly converting the truncation indicator to <code>RAW</code>, for example, using the <code>UTL RAW.CAST TO RAW</code> function.

• If you specify WITH COUNT, then after the truncation indicator, the database appends the number of truncated values, enclosed in parentheses. In this case, the database truncates

- enough measure values to allow space in the return value for a final delimiter, the truncation indicator, and 24 characters for the number value enclosed in parentheses.
- If you specify WITHOUT COUNT, then the database omits the number of truncated values from
 the return value. In this case, the database truncates enough measure values to allow
 space in the return value for a final delimiter and the truncation indicator.

If you do not specify WITH COUNT or WITHOUT COUNT, then the default is WITH COUNT.

Aggregate Examples

The following single-set aggregate example lists all of the employees in Department 30 in the hr.employees table, ordered by hire date and last name:

The following group-set aggregate example lists, for each department ID in the hr.employees table, the employees in that department in order of their hire date:

```
SELECT department id "Dept.",
      LISTAGG(last name, '; ') WITHIN GROUP (ORDER BY hire date) "Employees"
 FROM employees
 GROUP BY department id
 ORDER BY department id;
Dept. Employees
_____
   10 Whalen
   20 Hartstein; Fav
   30 Raphaely; Khoo; Tobias; Baida; Himuro; Colmenares
   40 Mavris
   50 Kaufling; Ladwig; Rajs; Sarchand; Bell; Mallin; Weiss; Davie
      s; Marlow; Bull; Everett; Fripp; Chung; Nayer; Dilly; Bissot
      ; Vollman; Stiles; Atkinson; Taylor; Seo; Fleaur; Matos; Pat
      el; Walsh; Feeney; Dellinger; McCain; Vargas; Gates; Rogers;
      Mikkilineni; Landry; Cabrio; Jones; Olson; OConnell; Sulliv
      an; Mourgos; Gee; Perkins; Grant; Geoni; Philtanker; Markle
   60 Austin; Hunold; Pataballa; Lorentz; Ernst
   70 Baer
```

The following example is identical to the previous example, except it contains the ON OVERFLOW TRUNCATE clause. For the purpose of this example, assume that the maximum length of the return value is an artificially small number of 200 bytes. Because the list of employees for department 50 exceeds 200 bytes, the list is truncated and appended with a final delimiter '; ', the specified truncation indicator '...', and the number of truncated values '(23)'.

```
SELECT department_id "Dept.",

LISTAGG(last_name, '; ' ON OVERFLOW TRUNCATE '...')

WITHIN GROUP (ORDER BY hire_date) "Employees"

FROM employees

GROUP BY department_id

ORDER BY department id;
```



```
Dept. Employees

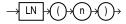
10 Whalen
20 Hartstein; Fay
30 Raphaely; Khoo; Tobias; Baida; Himuro; Colmenares
40 Mavris
50 Kaufling; Ladwig; Rajs; Sarchand; Bell; Mallin; Weiss; Davie
s; Marlow; Bull; Everett; Fripp; Chung; Nayer; Dilly; Bissot
; Vollman; Stiles; Atkinson; Taylor; Seo; Fleaur; ... (23)
70 Baer
...
```

Analytic Example

The following analytic example shows, for each employee hired earlier than September 1, 2003, the employee's department, hire date, and all other employees in that department also hired before September 1, 2003:

LN

Syntax



Purpose

LN returns the natural logarithm of n, where n is greater than 0.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.

Table 2-9 for more information on implicit conversion

Examples

The following example returns the natural logarithm of 95:

LNNVL

Syntax



Purpose

LNNVL provides a concise way to evaluate a condition when one or both operands of the condition may be null. The function can be used in the WHERE clause of a query, or as the WHEN condition in a searched CASE expression. It takes as an argument a condition and returns TRUE if the condition is FALSE or UNKNOWN and FALSE if the condition is TRUE. LNNVL can be used anywhere a scalar expression can appear, even in contexts where the IS [NOT] NULL, AND, or OR conditions are not valid but would otherwise be required to account for potential nulls.

Oracle Database sometimes uses the LNNVL function internally in this way to rewrite NOT IN conditions as NOT EXISTS conditions. In such cases, output from EXPLAIN PLAN shows this operation in the plan table output. The *condition* can evaluate any scalar values but cannot be a compound condition containing AND, OR, or BETWEEN.

The table that follows shows what LNNVL returns given that a = 2 and b is null.

Condition	Truth of Condition	LNNVL Return Value
a = 1	FALSE	TRUE
a = 2	TRUE	FALSE
a IS NULL	FALSE	TRUE
b = 1	UNKNOWN	TRUE
b IS NULL	TRUE	FALSE
a = b	UNKNOWN	TRUE

Examples

Suppose that you want to know the number of employees with commission rates of less than 20%, including employees who do not receive commissions. The following query returns only employees who actually receive a commission of less than 20%:

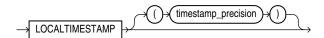


```
SELECT COUNT(*)
  FROM employees
  WHERE commission_pct < .2;
  COUNT(*)
------
11</pre>
```

To include the 72 employees who receive no commission at all, you could rewrite the query using the LNNVL function as follows:

LOCALTIMESTAMP

Syntax



Purpose

LOCALTIMESTAMP returns the current date and time in the session time zone in a value of data type TIMESTAMP. The difference between this function and CURRENT_TIMESTAMP is that LOCALTIMESTAMP returns a TIMESTAMP value while CURRENT_TIMESTAMP returns a TIMESTAMP WITH TIME ZONE value.

The optional argument timestamp_precision specifies the fractional second precision of the time value returned.



CURRENT_TIMESTAMP , "TIMESTAMP Data Type ", and "TIMESTAMP WITH TIME ZONE Data Type " $\,$

Examples

This example illustrates the difference between LOCALTIMESTAMP and CURRENT TIMESTAMP:

```
ALTER SESSION SET TIME_ZONE = '-5:00';

SELECT CURRENT_TIMESTAMP, LOCALTIMESTAMP FROM DUAL;

CURRENT_TIMESTAMP

LOCALTIMESTAMP

04-APR-00 01.27.18.999220 PM -05:00 04-APR-00 01.27.19 PM

ALTER SESSION SET TIME_ZONE = '-8:00';

SELECT CURRENT_TIMESTAMP, LOCALTIMESTAMP FROM DUAL;
```



When you use the LOCALTIMESTAMP with a format mask, take care that the format mask matches the value returned by the function. For example, consider the following table:

```
CREATE TABLE local_test (col1 TIMESTAMP WITH LOCAL TIME ZONE);
```

The following statement fails because the mask does not include the TIME ZONE portion of the return type of the function:

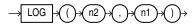
```
INSERT INTO local_test
   VALUES (TO_TIMESTAMP(LOCALTIMESTAMP, 'DD-MON-RR HH.MI.SSXFF'));
```

The following statement uses the correct format mask to match the return type of LOCALTIMESTAMP:

```
INSERT INTO local_test
   VALUES (TO TIMESTAMP(LOCALTIMESTAMP, 'DD-MON-RR HH.MI.SSXFF PM'));
```

LOG

Syntax



Purpose

LOG returns the logarithm, base n2, of n1. The base n2 can be any positive value other than 0 or 1 and n1 can be any positive value.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If any argument is <code>BINARY_FLOAT</code> or <code>BINARY_DOUBLE</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns <code>NUMBER</code>.



Table 2-9 for more information on implicit conversion

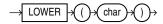
Examples

The following example returns the log of 100:



LOWER

Syntax



Purpose

LOWER returns *char*, with all letters lowercase. *char* can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The return value is the same data type as *char*. The database sets the case of the characters based on the binary mapping defined for the underlying character set. For linguistic-sensitive lowercase, refer to NLS LOWER.



See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of LOWER

Examples

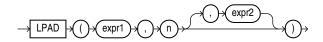
The following example returns a string in lowercase:

```
SELECT LOWER('MR. SCOTT MCMILLAN') "Lowercase"
FROM DUAL;

Lowercase
------
mr. scott mcmillan
```

LPAD

Syntax



Purpose

LPAD returns expr1, left-padded to length n characters with the sequence of characters in expr2. This function is useful for formatting the output of a query.

Both expr1 and expr2 can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The string returned is of VARCHAR2 data type if expr1 is a character data type, NVARCHAR2 if expr1 is a national character data type, and a LOB if expr1 is a LOB data type. The string returned is in the same character set as expr1. The argument n must be a NUMBER integer or a value that can be implicitly converted to a NUMBER integer.



If you do not specify expr2, then the default is a single blank. If expr1 is longer than n, then this function returns the portion of expr1 that fits in n.

The argument n is the total length of the return value as it is displayed on your terminal screen. In most character sets, this is also the number of characters in the return value. However, in some multibyte character sets, the display length of a character string can differ from the number of characters in the string.



See Also:

Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of ${\tt LPAD}$

Examples

The following example left-pads a string with the asterisk (*) and period (.) characters:

```
SELECT LPAD('Page 1',15,'*.') "LPAD example"
FROM DUAL;

LPAD example
-----*
*.*.*.*.*Page 1
```

LTRIM

Syntax



Purpose

LTRIM removes from the left end of *char* all of the characters contained in *set*. If you do not specify *set*, then it defaults to a single blank. Oracle Database begins scanning *char* from its first character and removes all characters that appear in *set* until reaching a character not in *set* and then returns the result.

Both char and set can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The string returned is of VARCHAR2 data type if char is a character data type, NVARCHAR2 if char is a national character data type, and a LOB if char is a LOB data type.



- RTRIM
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation LTRIM uses to compare characters from set with characters from char, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

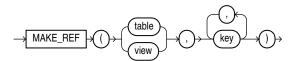
The following example trims all the left-most occurrences of less than sign (<), greater than sign (>), and equal sign (=) from a string:

```
SELECT LTRIM('<====>BROWNING<====>', '<>=') "LTRIM Example"
FROM DUAL;

LTRIM Example
-----BROWNING<====>
```

MAKE_REF

Syntax



Purpose

MAKE_REF creates a REF to a row of an object view or a row in an object table whose object identifier is primary key based. This function is useful, for example, if you are creating an object view



Oracle Database Object-Relational Developer's Guide for more information about object views and DEREF

Examples

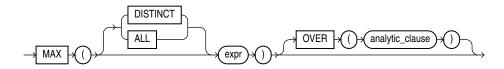
The sample schema oe contains an object view oc_inventories based on inventory_typ. The object identifier is product_id. The following example creates a REF to the row in the oc inventories object view with a product id of 3003:

```
SELECT MAKE_REF (oc_inventories, 3003)
FROM DUAL;
MAKE_REF(OC_INVENTORIES, 3003)
```



MAX

Syntax



See Also:

"Analytic Functions " for information on syntax, semantics, and restrictions

Purpose

MAX returns maximum value of expr. You can use it as an aggregate or analytic function.

See Also:

- "About SQL Expressions " for information on valid forms of expr, "Floating-Point Numbers " for information on binary-float comparison semantics, and "Aggregate Functions"
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation MAX uses to compare character values for expr, and for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value

Aggregate Example

The following example determines the highest salary in the hr.employees table:

Analytic Examples

The following example calculates, for each employee, the highest salary of the employees reporting to the same manager as the employee.

```
SELECT manager_id, last_name, salary,

MAX(salary) OVER (PARTITION BY manager_id) AS mgr_max
```



FROM employees
ORDER BY manager_id, last_name, salary;

MANAGER_ID	LAST_NAME	SALARY	MGR_MAX
100	Cambrault	11000	17000
100	De Haan	17000	17000
100	Errazuriz	12000	17000
100	Fripp	8200	17000
100	Hartstein	13000	17000
100	Kaufling	7900	17000
100	Kochhar	17000	17000

If you enclose this query in the parent query with a predicate, then you can determine the employee who makes the highest salary in each department:

MANAGER_ID	LAST_NAME	SALARY
100	De Haan	17000
100	Kochhar	17000
101	Greenberg	12008
101	Higgins	12008
102	Hunold	9000
103	Ernst	6000
108	Faviet	9000
114	Khoo	3100
120	Nayer	3200
120	Taylor	3200
121	Sarchand	4200
122	Chung	3800
123	Bell	4000
124	Rajs	3500
145	Tucker	10000
146	King	10000
147	Vishney	10500
148	Ozer	11500
149	Abel	11000
201	Fay	6000
205	Gietz	8300
	King	24000

22 rows selected.

MEDIAN

Syntax



"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

MEDIAN is an inverse distribution function that assumes a continuous distribution model. It takes a numeric or datetime value and returns the middle value or an interpolated value that would be the middle value once the values are sorted. Nulls are ignored in the calculation.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If you specify only <code>expr</code>, then the function returns the same data type as the numeric data type of the argument. If you specify the <code>OVER</code> clause, then Oracle Database determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.



Table 2-9 for more information on implicit conversion and "Numeric Precedence" for information on numeric precedence

The result of MEDIAN is computed by first ordering the rows. Using N as the number of rows in the group, Oracle calculates the row number (RN) of interest with the formula RN = (1 + (0.5*(N-1))). The final result of the aggregate function is computed by linear interpolation between the values from rows at row numbers CRN = CEILING(RN) and FRN = FLOOR(RN).

The final result will be:

```
if (CRN = FRN = RN) then
   (value of expression from row at RN)
else
   (CRN - RN) * (value of expression for row at FRN) +
   (RN - FRN) * (value of expression for row at CRN)
```

You can use MEDIAN as an analytic function. You can specify only the <code>query_partition_clause</code> in its <code>OVER</code> clause. It returns, for each row, the value that would fall in the middle among a set of values within each partition.

Compare this function with these functions:

- PERCENTILE_CONT, which returns, for a given percentile, the value that corresponds to that percentile by way of interpolation. MEDIAN is the specific case of PERCENTILE_CONT where the percentile value defaults to 0.5.
- PERCENTILE_DISC, which is useful for finding values for a given percentile without interpolation.

Aggregate Example

The following query returns the median salary for each department in the hr.employees table:

```
SELECT department_id, MEDIAN(salary)
FROM employees
GROUP BY department id
```



ORDER BY department_id;

DEPARTMENT_ID	MEDIAN (SALARY)
10	4400
20	9500
30	2850
40	6500
50	3100
60	4800
70	10000
80	8900
90	17000
100	8000
110	10154
	7000

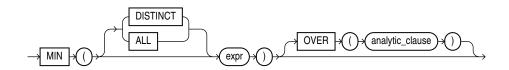
Analytic Example

The following query returns the median salary for each manager in a subset of departments in the hr.employees table:

MANAGER_ID	EMPLOYEE_ID	SALARY	Median by Mgr
100	101	17000	13500
100	102	17000	13500
100	145	14000	13500
100	146	13500	13500
100	147	12000	13500
100	148	11000	13500
100	149	10500	13500
101	108	12008	12008
101	204	10000	12008
101	205	12008	12008
108	109	9000	7800
108	110	8200	7800
108	111	7700	7800
108	112	7800	7800
108	113	6900	7800
145	150	10000	8500
145	151	9500	8500
145	152	9000	8500

MIN

Syntax





"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

MIN returns minimum value of expr. You can use it as an aggregate or analytic function.

See Also:

- "About SQL Expressions" for information on valid forms of expr, "Floating-Point Numbers" for information on binary-float comparison semantics, and "Aggregate Functions"
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation MIN uses to compare character values for expr, and for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value

Aggregate Example

The following statement returns the earliest hire date in the hr.employees table:

```
SELECT MIN(hire_date) "Earliest"
  FROM employees;
Earliest
-----
13-JAN-01
```

Analytic Example

The following example determines, for each employee, the employees who were hired on or before the same date as the employee. It then determines the subset of employees reporting to the same manager as the employee, and returns the lowest salary in that subset.

```
SELECT manager_id, last_name, hire_date, salary,

MIN(salary) OVER(PARTITION BY manager_id ORDER BY hire_date

RANGE UNBOUNDED PRECEDING) AS p_cmin

FROM employees

ORDER BY manager id, last name, hire date, salary;
```

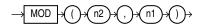
MANAGER_ID	LAST_NAME	HIRE_DATE	SALARY	P_CMIN
100	Cambrault	15-OCT-07	11000	6500
100	De Haan	13-JAN-01	17000	17000
100	Errazuriz	10-MAR-05	12000	7900
100	Fripp	10-APR-05	8200	7900
100	Hartstein	17-FEB-04	13000	7900
100	Kaufling	01-MAY-03	7900	7900
100	Kochhar	21-SEP-05	17000	7900
100	Mourgos	16-NOV-07	5800	5800
100	Partners	05-JAN-05	13500	7900
100	Raphaely	07-DEC-02	11000	11000
100	Russell	01-OCT-04	14000	7900



. . .

MOD

Syntax



Purpose

MOD returns the remainder of n2 divided by n1. Returns n2 if n1 is 0.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.



Table 2-9 for more information on implicit conversion and "Numeric Precedence" for information on numeric precedence

Examples

The following example returns the remainder of 11 divided by 4:

This function behaves differently from the classical mathematical modulus function, if the product of n1 and n2 is negative. The classical modulus can be expressed using the MOD function with this formula:

```
n2 - n1 * FLOOR(n2/n1)
```

The following table illustrates the difference between the MOD function and the classical modulus:

n2	n1	MOD(n2,n1)	Classical Modulus
11	4	3	3
11	-4	3	-1
-11	4	-3	1
-11	-4	-3	-3





FLOOR (number) and REMAINDER , which is similar to MOD, but uses ROUND in its formula instead of FLOOR

MONTHS_BETWEEN

Syntax



Purpose

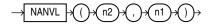
MONTHS_BETWEEN returns number of months between dates <code>date1</code> and <code>date2</code>. The month and the last day of the month are defined by the parameter <code>NLS_CALENDAR</code>. If <code>date1</code> is later than <code>date2</code>, then the result is positive. If <code>date1</code> is earlier than <code>date2</code>, then the result is negative. If <code>date1</code> and <code>date2</code> are either the same days of the month or both last days of months, then the result is always an integer. Otherwise Oracle Database calculates the fractional portion of the result based on a 31-day month and considers the difference in time components <code>date1</code> and <code>date2</code>.

Examples

The following example calculates the months between two dates:

NANVL

Syntax



Purpose

The NANVL function is useful only for floating-point numbers of type BINARY_FLOAT or BINARY_DOUBLE. It instructs Oracle Database to return an alternative value n1 if the input value n2 is NaN (not a number). If n2 is **not** NaN, then Oracle returns n2.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the

highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.



Table 2-9 for more information on implicit conversion, "Floating-Point Numbers" for information on binary-float comparison semantics, and "Numeric Precedence" for information on numeric precedence

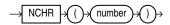
Examples

Using table float_point_demo created for TO_BINARY_DOUBLE, insert a second entry into the table:

The following example returns bin float if it is a number. Otherwise, 0 is returned.

NCHR

Syntax



Purpose

NCHR returns the character having the binary equivalent to number in the national character set. The value returned is always NVARCHAR2. This function is equivalent to using the CHR function with the USING NCHAR_CS clause.

This function takes as an argument a NUMBER value, or any value that can be implicitly converted to NUMBER, and returns a character.



- CHR
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of NCHR

Examples

The following examples return the nchar character 187:

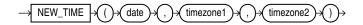
```
SELECT NCHR(187)
  FROM DUAL;

N
-
>
SELECT CHR(187 USING NCHAR_CS)
  FROM DUAL;

C
-
```

NEW TIME

Syntax



Purpose

NEW_TIME returns the date and time in time zone <code>timezone2</code> when date and time in time zone <code>timezone1</code> are <code>date</code>. Before using this function, you must set the <code>NLS_DATE_FORMAT</code> parameter to display 24-hour time. The return type is always <code>DATE</code>, regardless of the data type of <code>date</code>.



This function takes as input only a limited number of time zones. You can have access to a much greater number of time zones by combining the ${\tt FROM_TZ}$ function and the datetime expression. See ${\tt FROM_TZ}$ and the example for "Datetime Expressions".

The arguments timezone1 and timezone2 can be any of these text strings:

- AST, ADT: Atlantic Standard or Daylight Time
- BST, BDT: Bering Standard or Daylight Time

- CST, CDT: Central Standard or Daylight Time
- EST, EDT: Eastern Standard or Daylight Time
- GMT: Greenwich Mean Time
- HST, HDT: Alaska-Hawaii Standard Time or Daylight Time.
- MST, MDT: Mountain Standard or Daylight Time
- NST: Newfoundland Standard Time
- PST, PDT: Pacific Standard or Daylight Time
- YST, YDT: Yukon Standard or Daylight Time

Examples

The following example returns an Atlantic Standard time, given the Pacific Standard time equivalent:

NEXT DAY

Syntax



Purpose

NEXT_DAY returns the date of the first weekday named by <code>char</code> that is later than the date <code>date</code>. The return type is always <code>DATE</code>, regardless of the data type of <code>date</code>. The argument <code>char</code> must be a day of the week in the date language of your session, either the full name or the abbreviation. The minimum number of letters required is the number of letters in the abbreviated version. Any characters immediately following the valid abbreviation are ignored. The return value has the same hours, minutes, and seconds component as the argument <code>date</code>.

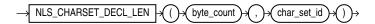
Examples

This example returns the date of the next Tuesday after October 15, 2009:



NLS_CHARSET_DECL_LEN

Syntax



Purpose

NLS_CHARSET_DECL_LEN returns the declaration length (in number of characters) of an NCHAR column. The <code>byte_count</code> argument is the width of the column. The <code>char_set_id</code> argument is the character set ID of the column.

Examples

The following example returns the number of characters that are in a 200-byte column when you are using a multibyte character set:

NLS CHARSET ID

Syntax



Purpose

NLS_CHARSET_ID returns the character set ID number corresponding to character set name string. The string argument is a run-time VARCHAR2 value. The string value 'CHAR_CS' returns the database character set ID number of the server. The string value 'NCHAR_CS' returns the national character set ID number of the server.

Invalid character set names return null.



Oracle Database Globalization Support Guide for a list of character sets

Examples

The following example returns the character set ID of a character set:

```
SELECT NLS_CHARSET_ID('ja16euc')
FROM DUAL;
```



NLS CHARSET NAME

Syntax



Purpose

NLS_CHARSET_NAME returns the name of the character set corresponding to ID number number. The character set name is returned as a VARCHAR2 value in the database character set. If number is not recognized as a valid character set ID, then this function returns null.

This function returns a VARCHAR2 value.



Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of $\tt NLS_CHARSET_NAME$

Examples

The following example returns the character set corresponding to character set ID number 2:

```
SELECT NLS_CHARSET_NAME(2)
FROM DUAL;

NLS_CH
-----
WE8DEC
```

NLS_COLLATION_ID

Syntax



Purpose

NLS_COLLATION_ID takes as its argument a collation name and returns the corresponding collation ID number. Collation IDs are used in the data dictionary tables and in Oracle Call Interface (OCI). Collation names are used in SQL statements and data dictionary views

For expr, specify the collation name as a VARCHAR2 value. You can specify a valid named collation or a pseudo-collation, in any combination of uppercase and lowercase letters.

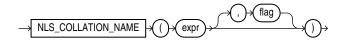
This function returns a NUMBER value. If you specify an invalid collation name, then this function returns null.

Examples

The following example returns the collation ID of collation BINARY CI:

NLS COLLATION NAME

Syntax



Purpose

NLS_COLLATION_NAME takes as its argument a collation ID number and returns the corresponding collation name. Collation IDs are used in the data dictionary tables and in Oracle Call Interface (OCI). Collation names are used in SQL statements and data dictionary views

For expr, specify the collation ID as a NUMBER value.

This function returns a VARCHAR2 value. If you specify an invalid collation ID, then this function returns null.

The optional flag parameter applies only to Unicode Collation Algorithm (UCA) collations. This parameter determines whether the function returns the short form or long form of the collation name. The parameter must be a character expression evaluating to the value 'S', 's', 'L', or 'l', with the following meaning:

- 's' or 's' Returns the short form of the collation name
- 'L' or 'l' Returns the long form of the collation name

If you omit flag, then the default is 'L'.

See Also:

- Oracle Database Globalization Support Guide for more information on UCA collations
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of NLS_COLLATION_NAME



Examples

The following example returns the name of the collation corresponding to collation ID number 81919:

```
SELECT NLS_COLLATION_NAME(81919)
FROM DUAL;
NLS_COLLA
-----
BINARY AI
```

The following example returns the short form of the name of the UCA collation corresponding to collation ID number 208897:

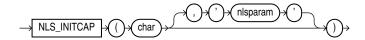
```
SELECT NLS_COLLATION_NAME(208897,'S')
FROM DUAL;

NLS_COLLATION
-----
UCA0610_DUCET
```

The following example returns the long form of the name of the UCA collation corresponding to collation ID number 208897:

NLS INITCAP

Syntax



Purpose

NLS_INITCAP returns *char*, with the first letter of each word in uppercase, all other letters in lowercase. Words are delimited by white space or characters that are not alphanumeric.

Both *char* and *'nlsparam'* can be any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. The string returned is of VARCHAR2 data type and is in the same character set as *char*.

The value of 'nlsparam' can have this form:

```
'NLS_SORT = sort'
```

where *sort* is a named collation. The collation handles special linguistic requirements for case conversions. These requirements can result in a return value of a different length than the *char*. If you omit 'nlsparam', then this function uses the determined collation of the function.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.

- "Data Type Comparison Rules" for more information.
- Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules for NLS_INITCAP, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following examples show how the linguistic sort sequence results in a different return value from the function:

```
SELECT NLS_INITCAP('ijsland') "InitCap"
   FROM DUAL;

InitCap
-----
Ijsland

SELECT NLS_INITCAP('ijsland', 'NLS_SORT = XDutch') "InitCap"
   FROM DUAL;

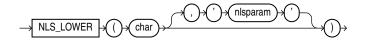
InitCap
-----
IJsland
```

See Also:

Oracle Database Globalization Support Guide for information on collations

NLS_LOWER

Syntax



Purpose

NLS LOWER returns char, with all letters lowercase.

Both char and 'nlsparam' can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The string returned is of VARCHAR2 data type if char is a character data type and a LOB if char is a LOB data type. The return string is in the same character set as char.

The 'nlsparam' can have the same form and serve the same purpose as in the NLS_INITCAP function.



Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules for NLS_LOWER, and for the collation derivation rules, which define the collation assigned to the character return value of this function

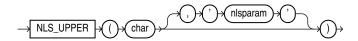
Examples

The following statement returns the lowercase form of the character string 'NOKTASINDA' using the XTurkish linguistic sort sequence. The Turkish uppercase I becoming a small, dotless i.

```
SELECT NLS_LOWER('NOKTASINDA', 'NLS_SORT = XTurkish') "Lowercase"
FROM DUAL;
```

NLS UPPER

Syntax



Purpose

NLS UPPER returns char, with all letters uppercase.

Both char and 'nlsparam' can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The string returned is of VARCHAR2 data type if char is a character data type and a LOB if char is a LOB data type. The return string is in the same character set as char.

The 'nlsparam' can have the same form and serve the same purpose as in the NLS_INITCAP function.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules for NLS_UPPER, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following example returns a string with all the letters converted to uppercase:

```
SELECT NLS_UPPER('große') "Uppercase"
  FROM DUAL;

Upper
----
GROßE

SELECT NLS_UPPER('große', 'NLS_SORT = XGerman') "Uppercase"
  FROM DUAL;
```

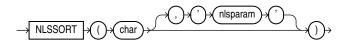


Upperc -----GROSSE

See Also:
NLS_INITCAP

NI SSORT

Syntax



Purpose

NLSSORT returns a collation key for the character value *char* and an explicitly or implicitly specified collation. A collation key is a string of bytes used to sort *char* according to the specified collation. The property of the collation keys is that mutual ordering of two such keys generated for the given collation when compared according to their binary order is the same as mutual ordering of the source character values when compared according to the given collation.

Both char and 'nlsparam' can be any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2.

The value of 'nlsparam' must have the form

'NLS SORT = collation'

where <code>collation</code> is the name of a linguistic collation or <code>BINARY</code>. <code>NLSSORT</code> uses the specified collation to generate the collation key. If you omit <code>'nlsparam'</code>, then this function uses the derived collation of the argument <code>char</code>. If you specify <code>BINARY</code>, then this function returns the <code>char</code> value itself cast to <code>RAW</code> and possibly truncated as described below.

If you specify 'nlsparam', then you can append to the linguistic collation name the suffix _ai to request an accent-insensitive collation or _ci to request a case-insensitive collation. Refer to Oracle Database Globalization Support Guide for more information on accent- and case-insensitive sorting. Using accent-insensitive or case-insensitive collations with the ORDER BY query clause is not recommended as it leads to a nondeterministic sort order.

The returned collation key is of RAW data type. The length of the collation key resulting from a given *char* value for a given collation may exceed the maximum length of the RAW value returned by NLSSORT. In this case, the behavior of NLSSORT depends on the value of the initialization parameter MAX_STRING_SIZE. If MAX_STRING_SIZE = EXTENDED, then the maximum length of the return value is 32767 bytes. If the collation key exceeds this limit, then the function fails with the error "ORA-12742: unable to create the collation key". This error may also be reported for short input strings if they contain a high percentage of Unicode characters with very high decomposition ratios.



Oracle Database Globalization Support Guide for details of when the ORA-12742 error is reported and how to prevent application availability issues that the error could cause

If MAX_STRING_SIZE = STANDARD, then the maximum length of the return value is 2000 bytes. If the value to be returned exceeds the limit, then NLSSORT calculates the collation key for a maximum prefix, or initial substring, of *char* so that the calculated result does not exceed the maximum length. For monolingual collations, for example FRENCH, the prefix length is typically 1000 characters. For multilingual collations, for example GENERIC_M, the prefix is typically 500 characters. For Unicode Collation Algorithm (UCA) collations, for example UCA0610_DUCET, the prefix is typically 285 characters. The exact length may be lower or higher depending on the collation and the characters contained in *char*.

The behavior when MAX_STRING_SIZE = STANDARD implies that two character values whose collation keys (NLSSORT results) are compared to find the linguistic ordering are considered equal if they do not differ in the prefix even though they may differ at some further character position. Because the NLSSORT function is used implicitly to find linguistic ordering for comparison conditions, the BETWEEN condition, the IN condition, ORDER BY, GROUP BY, and COUNT (DISTINCT), those operations may return results that are only approximate for long character values. If you want guarantee that the results of those operations are exact, then migrate your database to use MAX_STRING_SIZE = EXTENDED.

Refer to "Extended Data Types" for more information on the MAX_STRING_SIZE initialization parameter.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.

See Also:

- "Data Type Comparison Rules" for more information.
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules for NLSSORT

Examples

This function can be used to specify sorting and comparison operations based on a linguistic sort sequence rather than on the binary value of a string. The following example creates a test table containing two values and shows how the values returned can be ordered by the $\tt NLSSORT$ function:

```
CREATE TABLE test (name VARCHAR2(15));
INSERT INTO test VALUES ('Gaardiner');
INSERT INTO test VALUES ('Gaberd');
INSERT INTO test VALUES ('Gaasten');

SELECT *
   FROM test
   ORDER BY name;
```



The following example shows how to use the NLSSORT function in comparison operations:

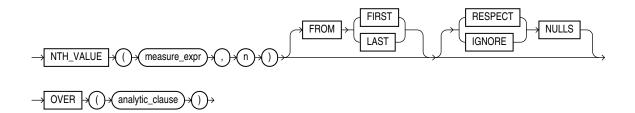
If you frequently use <code>NLSSORT</code> in comparison operations with the same linguistic sort sequence, then consider this more efficient alternative: Set the <code>NLS_COMP</code> parameter (either for the database or for the current session) to <code>LINGUISTIC</code>, and set the <code>NLS_SORT</code> parameter for the session to the desired sort sequence. Oracle Database will use that sort sequence by default for all sorting and comparison operations during the current session:

See Also:

Oracle Database Globalization Support Guide for information on sort sequences

NTH_VALUE

Syntax



See Also:

"Analytic Functions " for information on syntax, semantics, and restrictions of the analytic clause

Purpose

NTH_VALUE returns the measure_expr value of the nth row in the window defined by the analytic clause. The returned value has the data type of the measure expr.

- {RESPECT | IGNORE} NULLS determines whether null values of measure_expr are included in or eliminated from the calculation. The default is RESPECT NULLS.
- n determines the nth row for which the measure value is to be returned. n can be a constant, bind variable, column, or an expression involving them, as long as it resolves to a positive integer. The function returns NULL if the data source window has fewer than n rows. If n is null, then the function returns an error.
- FROM {FIRST | LAST} determines whether the calculation begins at the first or last row of the window. The default is FROM FIRST.

If you omit the <code>windowing_clause</code> of the <code>analytic_clause</code>, it defaults to <code>RANGE BETWEEN</code> UNBOUNDED PRECEDING AND CURRENT ROW. This default sometimes returns an unexpected value for <code>NTH_VALUE</code> ... FROM <code>LAST</code> ... , because the last value in the window is at the bottom of the window, which is not fixed. It keeps changing as the current row changes. For expected results, specify the <code>windowing_clause</code> as <code>RANGE BETWEEN UNBOUNDED PRECEDING</code> AND UNBOUNDED FOLLOWING. Alternatively, you can specify the <code>windowing_clause</code> as <code>RANGE BETWEEN CURRENT ROW</code> AND UNBOUNDED FOLLOWING.

See Also:

- Oracle Database Data Warehousing Guide for more information on the use of this function
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of NTH_VALUE when it is a character value

Examples

The following example shows the minimum <code>amount_sold</code> value for the second <code>channel_id</code> in ascending order for each <code>prod_id</code> between 13 and 16:

```
SELECT prod_id, channel_id, MIN(amount_sold),

NTH_VALUE(MIN(amount_sold), 2) OVER (PARTITION BY prod_id ORDER BY channel_id

ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) nv

FROM sales

WHERE prod_id BETWEEN 13 and 16

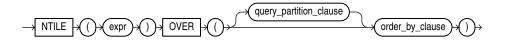
GROUP BY prod_id, channel_id;
```

PROD_ID	CHANNEL_ID	MIN (AMOUNT_SOLD)	NV
13	2	907.34	906.2
13	3	906.2	906.2
13	4	842.21	906.2
14	2	1015.94	1036.72
14	3	1036.72	1036.72
14	4	935.79	1036.72
15	2	871.19	871.19
15	3	871.19	871.19
15	4	871.19	871.19
16	2	266.84	266.84
16	3	266.84	266.84
16	4	266.84	266.84
16	9	11.99	266.84

13 rows selected.

NTILE

Syntax



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions, including valid forms of expr

Purpose

NTILE is an analytic function. It divides an ordered data set into a number of buckets indicated by expr and assigns the appropriate bucket number to each row. The buckets are numbered 1 through expr. The expr value must resolve to a positive constant for each partition. Oracle Database expects an integer, and if expr is a noninteger constant, then Oracle truncates the value to an integer. The return value is NUMBER.

The number of rows in the buckets can differ by at most 1. The remainder values (the remainder of number of rows divided by buckets) are distributed one for each bucket, starting with bucket 1.

If expr is greater than the number of rows, then a number of buckets equal to the number of rows will be filled, and the remaining buckets will be empty.

You cannot nest analytic functions by using NTILE or any other analytic function for expr. However, you can use other built-in function expressions for expr.



"About SQL Expressions" for information on valid forms of expr and Table 2-9 for more information on implicit conversion

Examples

The following example divides into 4 buckets the values in the salary column of the oe.employees table from Department 100. The salary column has 6 values in this department, so the two extra values (the remainder of 6 / 4) are allocated to buckets 1 and 2, which therefore have one more value than buckets 3 or 4.

```
SELECT last_name, salary, NTILE(4) OVER (ORDER BY salary DESC) AS quartile
FROM employees
WHERE department_id = 100
ORDER BY last_name, salary, quartile;
```

LAST_NAME	SALARY	QUARTILE
Chen	8200	2
Faviet	9000	1
Greenberg	12008	1
Popp	6900	4
Sciarra	7700	3
Urman	7800	2

NULLIF

Syntax



Purpose

NULLIF compares expr1 and expr2. If they are equal, then the function returns null. If they are not equal, then the function returns expr1. You cannot specify the literal NULL for expr1.

If both arguments are numeric data types, then Oracle Database determines the argument with the higher numeric precedence, implicitly converts the other argument to that data type, and returns that data type. If the arguments are not numeric, then they must be of the same data type, or Oracle returns an error.

The NULLIF function is logically equivalent to the following CASE expression:

CASE WHEN expr1 = expr2 THEN NULL ELSE expr1 END



- "CASE Expressions"
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation NULLIF uses to compare characters from expr1 with characters from expr2, and for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value

Examples

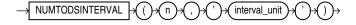
The following example selects those employees from the sample schema hr who have changed jobs since they were hired, as indicated by a job_id in the job_history table different from the current job id in the employees table:

```
SELECT e.last name, NULLIF(j.job id, e.job id) "Old Job ID"
  FROM employees e, job history j
  WHERE e.employee id = j.employee id
  ORDER BY last name, "Old Job ID";
LAST NAME
                        Old Job ID
De Haan
                         IT PROG
                         MK REP
Hartstein
                         ST CLERK
Kaufling
                        AC ACCOUNT
Kochhar
Kochhar
                         AC MGR
Raphaely
                         ST_CLERK
Taylor
                         SA_MAN
Taylor
Whalen
                         AC ACCOUNT
```

NUMTODSINTERVAL

Whalen

Syntax



Purpose

NUMTODSINTERVAL converts n to an INTERVAL DAY TO SECOND literal. The argument n can be any NUMBER value or an expression that can be implicitly converted to a NUMBER value. The argument $interval_unit$ can be of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type. The value for $interval_unit$ specifies the unit of n and must resolve to one of the following string values:

- 'DAY'
- 'HOUR'
- 'MINUTE'
- 'SECOND'



interval_unit is case insensitive. Leading and trailing values within the parentheses are ignored. By default, the precision of the return is 9.



Table 2-9 for more information on implicit conversion

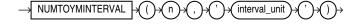
Examples

The following example uses NUMTODSINTERVAL in a COUNT analytic function to calculate, for each employee, the number of employees hired by the same manager within the past 100 days from his or her hire date. Refer to "Analytic Functions" for more information on the syntax of the analytic functions.

MANAGER_ID	LAST_NAME	HIRE_DATE	T_COUNT
149	Abel	11-MAY-04	1
147	Ande	24-MAR-08	3
121	Atkinson	30-OCT-05	2
103	Austin	25-JUN-05	1
124	Walsh	24-APR-06	2
100	Weiss	18-JUL-04	1
101	Whalen	17-SEP-03	1
100	Zlotkey	29-JAN-08	2

NUMTOYMINTERVAL

Syntax



Purpose

NUMTOYMINTERVAL converts number n to an INTERVAL YEAR TO MONTH literal. The argument n can be any NUMBER value or an expression that can be implicitly converted to a NUMBER value. The argument $interval_unit$ can be of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type. The value for $interval_unit$ specifies the unit of n and must resolve to one of the following string values:

- 'YEAR'
- 'MONTH'

interval_unit is case insensitive. Leading and trailing values within the parentheses are ignored. By default, the precision of the return is 9.



Table 2-9 for more information on implicit conversion

Examples

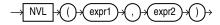
The following example uses NUMTOYMINTERVAL in a SUM analytic function to calculate, for each employee, the total salary of employees hired in the past one year from his or her hire date. Refer to "Analytic Functions" for more information on the syntax of the analytic functions.

```
SELECT last_name, hire_date, salary,
    SUM(salary) OVER (ORDER BY hire_date
    RANGE NUMTOYMINTERVAL(1,'year') PRECEDING) AS t_sal
FROM employees
ORDER BY last_name, hire_date;
```

LAST_NAME	HIRE_DATE	SALARY	T_SAL
Abel	11-MAY-04	11000	90300
Ande	24-MAR-08	6400	112500
Atkinson	30-OCT-05	2800	177000
Austin	25-JUN-05	4800	134700
Walsh	24-APR-06	3100	186200
Weiss	18-JUL-04	8000	70900
Whalen	17-SEP-03	4400	54000
Zlotkey	29-JAN-08	10500	119000

NVL

Syntax



Purpose

NVL lets you replace null (returned as a blank) with a string in the results of a query. If expr1 is null, then NVL returns expr2. If expr1 is not null, then NVL returns expr1.

The arguments expr1 and expr2 can have any data type. If their data types are different, then Oracle Database implicitly converts one to the other. If they cannot be converted implicitly, then the database returns an error. The implicit conversion is implemented as follows:

- If expr1 is character data, then Oracle Database converts expr2 to the data type of expr1 before comparing them and returns VARCHAR2 in the character set of expr1.
- If *expr1* is numeric, then Oracle Database determines which argument has the highest numeric precedence, implicitly converts the other argument to that data type, and returns that data type.



- Table 2-9 for more information on implicit conversion and "Numeric Precedence" for information on numeric precedence
- "COALESCE" and "CASE Expressions", which provide functionality similar to that of NVL
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of NVL when it is a character value

Examples

The following example returns a list of employee names and commissions, substituting "Not Applicable" if the employee receives no commission:

```
SELECT last_name, NVL(TO_CHAR(commission_pct), 'Not Applicable') commission
  FROM employees
  WHERE last_name LIKE 'B%'
  ORDER BY last name;
```

LAST_NAME	COMMISSION
Baer	Not Applicable
Baida	Not Applicable
Banda	.1
Bates	.15
Bell	Not Applicable
Bernstein	.25
Bissot	Not Applicable
Bloom	.2
Bull	Not Applicable

NVL₂

Syntax



Purpose

NVL2 lets you determine the value returned by a query based on whether a specified expression is null or not null. If expr1 is not null, then NVL2 returns expr2. If expr1 is null, then NVL2 returns expr3.

The argument expr1 can have any data type. The arguments expr2 and expr3 can have any data types except LONG.

If the data types of expr2 and expr3 are different, then Oracle Database implicitly converts one to the other. If they cannot be converted implicitly, then the database returns an error. If expr2 is character or numeric data, then the implicit conversion is implemented as follows:

• If expr2 is character data, then Oracle Database converts expr3 to the data type of expr2 before returning a value unless expr3 is a null constant. In that case, a data type

conversion is not necessary, and the database returns VARCHAR2 in the character set of expr2.

• If *expr2* is numeric data, then Oracle Database determines which argument has the highest numeric precedence, implicitly converts the other argument to that data type, and returns that data type.

See Also:

- Table 2-9 for more information on implicit conversion and "Numeric Precedence" for information on numeric precedence
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of NVL2 when it is a character value

Examples

The following example shows whether the income of some employees is made up of salary plus commission, or just salary, depending on whether the <code>commission_pct</code> column of <code>employees</code> is null or not.

```
SELECT last_name, salary,
            NVL2(commission_pct, salary + (salary * commission_pct), salary) income
FROM employees
WHERE last_name like 'B%'
ORDER BY last_name;
```

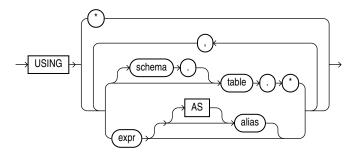
LAST_NAME	SALARY	INCOME
Baer	10000	10000
Baida	2900	2900
Banda	6200	6820
Bates	7300	8395
Bell	4000	4000
Bernstein	9500	11875
Bissot	3300	3300
Bloom	10000	12000
Bull	4100	4100

ORA DM PARTITION NAME

Syntax



mining_attribute_clause::=



Purpose

ORA_DM_PARTITION_NAME is a single row function that works along with other existing functions. This function returns the name of the partition associated with the input row. When ORA DM PARTITION NAME is used on a non-partitioned model, the result is NULL.

The syntax of the <code>ORA_DM_PARTITION_NAME</code> function can use an optional <code>GROUPING</code> hint when scoring a partitioned model. See <code>GROUPING</code> Hint.

mining_attribute_clause

The <code>mining_attribute_clause</code> identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The <code>mining_attribute_clause</code> behaves as described for the <code>PREDICTION</code> function. See <code>mining_attribute_clause</code>.

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring
- Oracle Machine Learning for SQL Concepts for information about clustering

Note:

The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

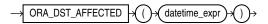
Example

SELECT prediction(mymodel using *) pred, ora_dm_partition_name(mymodel USING
*) pname FROM customers;



ORA_DST_AFFECTED

Syntax



Purpose

ORA_DST_AFFECTED is useful when you are changing the time zone data file for your database. The function takes as an argument a datetime expression that resolves to a TIMESTAMP WITH TIME ZONE value or a VARRAY object that contains TIMESTAMP WITH TIME ZONE values. The function returns 1 if the datetime value is affected by or will result in a "nonexisting time" or "duplicate time" error with the new time zone data. Otherwise, it returns 0.

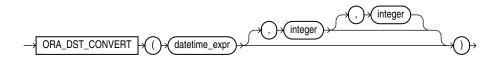
This function can be issued only when changing the time zone data file of the database and upgrading the timestamp with the time zone data, and only between the execution of the DBMS_DST.BEGIN_PREPARE and the DBMS_DST.END_PREPARE procedures or between the execution of the DBMS_DST.BEGIN_UPGRADE and the DBMS_DST.END_UPGRADE procedures.

See Also:

Oracle Database Globalization Support Guide for more information on time zone data files and on how Oracle Database handles daylight saving time, and Oracle Database PL/SQL Packages and Types Reference for information on the DBMS_DST package

ORA DST CONVERT

Syntax



Purpose

ORA_DST_CONVERT is useful when you are changing the time zone data file for your database. The function lets you specify error handling for a specified datetime expression.

- For datetime_expr, specify a datetime expression that resolves to a TIMESTAMP WITH TIME ZONE value or a VARRAY object that contains TIMESTAMP WITH TIME ZONE values.
- The optional second argument specifies handling of "duplicate time" errors. Specify 0 (false) to suppress the error by returning the source datetime value. This is the default. Specify 1 (true) to allow the database to return the duplicate time error.

The optional third argument specifies handling of "nonexisting time" errors. Specify 0
(false) to suppress the error by returning the source datetime value. This is the default.
Specify 1 (true) to allow the database to return the nonexisting time error.

If no error occurs, this function returns a value of the same data type as <code>datetime_expr</code> (a <code>TIMESTAMP WITH TIME ZONE</code> value or a <code>VARRAY</code> object that contains <code>TIMESTAMP WITH TIME ZONE</code> values). The returned datetime value when interpreted with the new time zone file corresponds to <code>datetime expr</code> interpreted with the old time zone file.

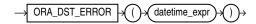
This function can be issued only when changing the time zone data file of the database and upgrading the timestamp with the time zone data, and only between the execution of the DBMS DST.BEGIN UPGRADE and the DBMS DST.END UPGRADE procedures.



Oracle Database Globalization Support Guide for more information on time zone data files and on how Oracle Database handles daylight saving time, and Oracle Database PL/SQL Packages and Types Reference for information on the DBMS_DST package

ORA DST ERROR

Syntax



Purpose

ORA_DST_ERROR is useful when you are changing the time zone data file for your database. The function takes as an argument a datetime expression that resolves to a TIMESTAMP WITH TIME ZONE value or a VARRAY object that contains TIMESTAMP WITH TIME ZONE values, and indicates whether the datetime value will result in an error with the new time zone data. The return values are:

- 0: the datetime value does not result in an error with the new time zone data.
- 1878: the datetime value results in a "nonexisting time" error.
- 1883: the datetime value results in a "duplicate time" error.

This function can be issued only when changing the time zone data file of the database and upgrading the timestamp with the time zone data, and only between the execution of the DBMS_DST.BEGIN_PREPARE and the DBMS_DST.END_PREPARE procedures or between the execution of the DBMS_DST.BEGIN_UPGRADE and the DBMS_DST.END_UPGRADE procedures.





Oracle Database Globalization Support Guide for more information on time zone data files and on how Oracle Database handles daylight saving time, and Oracle Database PL/SQL Packages and Types Reference for information on the DBMS_DST package

ORA_HASH

Syntax



Purpose

ORA_HASH is a function that computes a hash value for a given expression. This function is useful for operations such as analyzing a subset of data and generating a random sample.

- The <code>expr</code> argument determines the data for which you want Oracle Database to compute a hash value. There are no restrictions on the length of data represented by <code>expr</code>, which commonly resolves to a column name. The <code>expr</code> cannot be a <code>LONG</code> or LOB type. It cannot be a user-defined object type unless it is a nested table type. The hash value for nested table types does not depend on the order of elements in the collection. All other data types are supported for <code>expr</code>.
- The optional max_bucket argument determines the maximum bucket value returned by the hash function. You can specify any value between 0 and 4294967295. The default is 4294967295.
- The optional <code>seed_value</code> argument enables Oracle to produce many different results for the same set of data. Oracle applies the hash function to the combination of <code>expr</code> and <code>seed_value</code>. You can specify any value between 0 and 4294967295. The default is 0.

The function returns a NUMBER value.

Examples

The following example creates a hash value for each combination of customer ID and product ID in the sh.sales table, divides the hash values into a maximum of 100 buckets, and returns the sum of the amount_sold values in the first bucket (bucket 0). The third argument (5) provides a seed value for the hash function. You can obtain different hash results for the same query by changing the seed value.



ORA_INVOKING_USER

Syntax

→ ORA_INVOKING_USER

Purpose

ORA_INVOKING_USER returns the name of the database user who invoked the current statement or view. This function takes into account the BEQUEATH property of intervening views referenced in the statement. If this function is invoked from within a definer's rights context, then it returns the name of the owner of the definer's rights object. If the invoking user is a Real Application Security user, then it returns user XS\$NULL.

This function returns a VARCHAR2 value.

See Also:

- BEQUEATH clause of the CREATE VIEW statement
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of ORA INVOKING USER

Examples

The following example returns the name of the database user who invoked the statement:

SELECT ORA_INVOKING_USER FROM DUAL;

ORA_INVOKING_USERID

Syntax

→ ORA_INVOKING_USERID →

Purpose

ORA_INVOKING_USERID returns the identifier of the database user who invoked the current statement or view. This function takes into account the BEQUEATH property of intervening views referenced in the statement.

This function returns a NUMBER value.



See Also:

- ORA_INVOKING_USER to learn how Oracle Database determines the database user who invoked the current statement or view
- BEQUEATH clause of the CREATE VIEW statement

Examples

The following example returns the identifier of the database user who invoked the statement:

SELECT ORA INVOKING USERID FROM DUAL;

PATH

Syntax



Purpose

PATH is an ancillary function used only with the UNDER_PATH and EQUALS_PATH conditions. It returns the relative path that leads to the resource specified in the parent condition.

The <code>correlation_integer</code> can be any <code>NUMBER</code> integer and is used to correlate this ancillary function with its primary condition. Values less than 1 are treated as 1.

See Also:

- EQUALS_PATH Condition and UNDER_PATH Condition
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of PATH

Examples

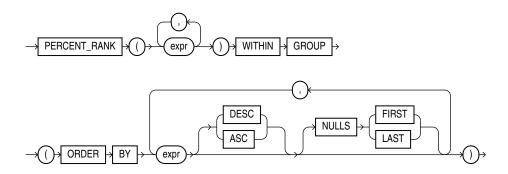
Refer to the related function $\frac{\text{DEPTH}}{\text{DEPTH}}$ for an example using both of these ancillary functions of the EQUALS_PATH and UNDER_PATH conditions.



PERCENT_RANK

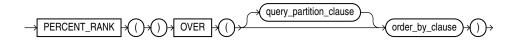
Aggregate Syntax

percent_rank_aggregate::=



Analytic Syntax

percent_rank_analytic::=





"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

PERCENT_RANK is similar to the CUME_DIST (cumulative distribution) function. The range of values returned by PERCENT_RANK is 0 to 1, inclusive. The first row in any set has a PERCENT_RANK of 0. The return value is NUMBER.



Table 2-9 for more information on implicit conversion

As an aggregate function, PERCENT_RANK calculates, for a hypothetical row r identified by the arguments of the function and a corresponding sort specification, the rank of row r minus 1 divided by the number of rows in the aggregate group. This calculation is made as if the hypothetical row r were inserted into the group of rows over which Oracle Database is to aggregate.

The arguments of the function identify a single hypothetical row within each aggregate group. Therefore, they must all evaluate to constant expressions within each aggregate

group. The constant argument expressions and the expressions in the ORDER BY clause of the aggregate match by position. Therefore the number of arguments must be the same and their types must be compatible.

 As an analytic function, for a row r, PERCENT_RANK calculates the rank of r minus 1, divided by 1 less than the number of rows being evaluated (the entire query result set or a partition).



Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules, which define the collation PERCENT_RANK uses to compare character values for the ORDER BY clause

Aggregate Example

The following example calculates the percent rank of a hypothetical employee in the sample table hr.employees with a salary of \$15,500 and a commission of 5%:

Analytic Example

The following example calculates, for each employee, the percent rank of the employee's salary within the department:

```
SELECT department_id, last_name, salary, PERCENT_RANK()

OVER (PARTITION BY department_id ORDER BY salary DESC) AS pr

FROM employees

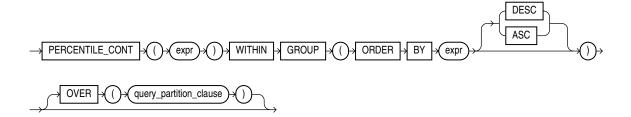
ORDER BY pr, salary, last name;
```

DEPARTMENT_ID LAST_NAME	SALARY	PR
10 Whalen	4400	0
40 Mavris	6500	0
Grant	7000	0
80 Vishney	10500	.181818182
80 Zlotkey	10500	.181818182
30 Khoo	3100	.2
50 Markle	2200	.954545455
50 Philtanker	2200	.954545455
50 Olson	2100	1



PERCENTILE_CONT

Syntax



See Also:

"Analytic Functions " for information on syntax, semantics, and restrictions of the OVER clause

Purpose

PERCENTILE_CONT is an inverse distribution function that assumes a continuous distribution model. It takes a percentile value and a sort specification, and returns an interpolated value that would fall into that percentile value with respect to the sort specification. Nulls are ignored in the calculation.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.

See Also:

Table 2-9 for more information on implicit conversion

The first expr must evaluate to a numeric value between 0 and 1, because it is a percentile value. This expr must be constant within each aggregation group. The ORDER BY clause takes a single expression that must be a numeric or datetime value, as these are the types over which Oracle can perform interpolation.

The result of PERCENTILE_CONT is computed by linear interpolation between values after ordering them. Using the percentile value (P) and the number of rows (N) in the aggregation group, you can compute the row number you are interested in after ordering the rows with respect to the sort specification. This row number (RN) is computed according to the formula RN = (1 + (P*(N-1))). The final result of the aggregate function is computed by linear interpolation between the values from rows at row numbers CRN = CEILING(RN) and FRN = FLOOR(RN).

The final result will be:

```
If (CRN = FRN = RN) then the result is
  (value of expression from row at RN)
```



```
Otherwise the result is

(CRN - RN) * (value of expression for row at FRN) +

(RN - FRN) * (value of expression for row at CRN)
```

You can use the PERCENTILE_CONT function as an analytic function. You can specify only the <code>query_partitioning_clause</code> in its <code>OVER</code> clause. It returns, for each row, the value that would fall into the specified percentile among a set of values within each partition.

The MEDIAN function is a specific case of PERCENTILE_CONT where the percentile value defaults to 0.5. For more information, refer to MEDIAN.

Note:

Before processing a large amount of data with the PERCENTILE_CONT function, consider using one of the following methods to obtain approximate results more quickly than exact results:

- Set the APPROX_FOR_PERCENTILE initialization parameter to PERCENTILE_CONT or ALL before using the PERCENTILE_CONT function. Refer to *Oracle Database Reference* for more information on this parameter.
- Use the APPROX_PERCENTILE function instead of the PERCENTILE_CONT function. Refer to APPROX_PERCENTILE.

Aggregate Example

The following example computes the median salary in each department:

```
SELECT department_id,
     PERCENTILE CONT(0.5) WITHIN GROUP (ORDER BY salary DESC) "Median cont",
     PERCENTILE DISC(0.5) WITHIN GROUP (ORDER BY salary DESC) "Median disc"
 FROM employees
 GROUP BY department id
 ORDER BY department id;
DEPARTMENT ID Median cont Median disc
----- -----
        10
             4400
                        4400
        20
               9500
                       13000
                       2900
        30
              2850
        40
                        6500
               6500
        50
              3100
                        3100
        60
               4800
                        4800
        70
              10000
                       10000
        80
               8900
                        9000
              17000
                        17000
        90
       100
               8000
                         8200
       110
              10154
                        12008
                7000
                         7000
```

PERCENTILE_CONT and PERCENTILE_DISC may return different results. PERCENTILE_CONT returns a computed result after doing linear interpolation. PERCENTILE_DISC simply returns a value from the set of values that are aggregated over. When the percentile value is 0.5, as in this example, PERCENTILE_CONT returns the average of the two middle values for groups with even number of elements, whereas PERCENTILE_DISC returns the value of the first one among the two middle values. For aggregate groups with an odd number of elements, both functions return the value of the middle element.

Analytic Example

In the following example, the median for Department 60 is 4800, which has a corresponding percentile (Percent_Rank) of 0.5. None of the salaries in Department 30 have a percentile of 0.5, so the median value must be interpolated between 2900 (percentile 0.4) and 2800 (percentile 0.6), which evaluates to 2850.

```
SELECT last_name, salary, department_id,

PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY salary DESC)

OVER (PARTITION BY department_id) "Percentile_Cont",

PERCENT_RANK()

OVER (PARTITION BY department_id ORDER BY salary DESC) "Percent_Rank"

FROM employees

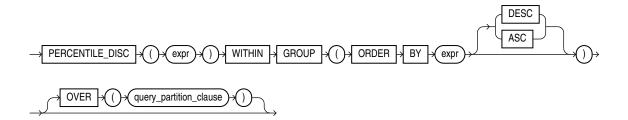
WHERE department_id IN (30, 60)

ORDER BY last_name, salary, department_id;
```

LAST_NAME	SALARY	DEPARTMENT_ID	Percentile_Cont	Percent_Rank
			4000	
Austin	4800	60	4800	.5
Baida	2900	30	2850	. 4
Colmenares	2500	30	2850	1
Ernst	6000	60	4800	.25
Himuro	2600	30	2850	.8
Hunold	9000	60	4800	0
Khoo	3100	30	2850	.2
Lorentz	4200	60	4800	1
Pataballa	4800	60	4800	.5
Raphaely	11000	30	2850	0
Tobias	2800	30	2850	.6

PERCENTILE_DISC

Syntax



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions of the OVER clause

Purpose

PERCENTILE_DISC is an inverse distribution function that assumes a discrete distribution model. It takes a percentile value and a sort specification and returns an element from the set. Nulls are ignored in the calculation.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.

See Also:

Table 2-9 for more information on implicit conversion

The first expr must evaluate to a numeric value between 0 and 1, because it is a percentile value. This expression must be constant within each aggregate group. The ORDER BY clause takes a single expression that can be of any type that can be sorted.

For a given percentile value P, PERCENTILE_DISC sorts the values of the expression in the ORDER BY clause and returns the value with the smallest CUME_DIST value (with respect to the same sort specification) that is greater than or equal to P.



Before processing a large amount of data with the PERCENTILE_DISC function, consider using one of the following methods to obtain approximate results more quickly than exact results:

- Set the APPROX_FOR_PERCENTILE initialization parameter to PERCENTILE_DISC or ALL before using the PERCENTILE_DISC function. Refer to *Oracle Database Reference* for more information on this parameter.
- Use the APPROX_PERCENTILE function instead of the PERCENTILE_DISC function.
 Refer to APPROX_PERCENTILE.

Aggregate Example

See aggregate example for PERCENTILE CONT.

Analytic Example

The following example calculates the median discrete percentile of the salary of each employee in the sample table hr.employees:

```
SELECT last_name, salary, department_id,

PERCENTILE_DISC(0.5) WITHIN GROUP (ORDER BY salary DESC)

OVER (PARTITION BY department_id) "Percentile_Disc",

CUME_DIST() OVER (PARTITION BY department_id

ORDER BY salary DESC) "Cume_Dist"

FROM employees

WHERE department_id in (30, 60)

ORDER BY last name, salary, department id;
```

LAST_NAME	SALARY	DEPARTMENT_ID	Percentile_Disc	Cume_Dist
Austin	4800	60	4800	.8
Baida	2900	30	2900	.5
Colmenares	2500	30	2900	1
Ernst	6000	60	4800	. 4
Himuro	2600	30	2900	.833333333
Baida Colmenares Ernst	2900 2500 6000	30 30 60	2900 2900 4800	.5

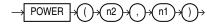


Hunold	9000	60	4800 .2
Khoo	3100	30	2900 .333333333
Lorentz	4200	60	4800 1
Pataballa	4800	60	4800 .8
Raphaely	11000	30	2900 .166666667
Tobias	2800	30	2900 .666666667

The median value for Department 30 is 2900, which is the value whose corresponding percentile (Cume_Dist) is the smallest value greater than or equal to 0.5. The median value for Department 60 is 4800, which is the value whose corresponding percentile is the smallest value greater than or equal to 0.5.

POWER

Syntax



Purpose

POWER returns n2 raised to the n1 power. The base n2 and the exponent n1 can be any numbers, but if n2 is negative, then n1 must be an integer.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If any argument is <code>BINARY_FLOAT</code> or <code>BINARY_DOUBLE</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise, the function returns <code>NUMBER</code>.



Table 2-9 for more information on implicit conversion

Examples

The following example returns 3 squared:

```
SELECT POWER(3,2) "Raised"
FROM DUAL;

Raised
-----
```

POWERMULTISET

Syntax





Purpose

POWERMULTISET takes as input a nested table and returns a nested table of nested tables containing all nonempty subsets (called submultisets) of the input nested table.

- expr can be any expression that evaluates to a nested table.
- If expr resolves to null, then Oracle Database returns NULL.
- If expr resolves to a nested table that is empty, then Oracle returns an error.
- The element types of the nested table must be comparable. Refer to "Comparison Conditions" for information on the comparability of nonscalar types.



This function is not supported in PL/SQL.

Examples

First, create a data type that is a nested table of the cust_address_tab_type data type:

```
CREATE TYPE cust_address_tab_tab_typ
  AS TABLE OF cust_address_tab_typ;
//
```

Now, select the nested table column <code>cust_address_ntab</code> from the <code>customers_demo</code> table using the <code>POWERMULTISET</code> function:

The preceding example requires the customers_demo table and a nested table column containing data. Refer to "Multiset Operators" to create this table and nested table columns.

POWERMULTISET BY CARDINALITY

Syntax



Purpose

POWERMULTISET_BY_CARDINALITY takes as input a nested table and a cardinality and returns a nested table of nested tables containing all nonempty subsets (called submultisets) of the nested table of the specified cardinality.

- expr can be any expression that evaluates to a nested table.
- cardinality can be any positive integer.
- If expr resolves to null, then Oracle Database returns NULL.
- If expr resolves to a nested table that is empty, then Oracle returns an error.
- The element types of the nested table must be comparable. Refer to "Comparison Conditions" for information on the comparability of nonscalar types.



This function is not supported in PL/SQL.

Examples

First, create a data type that is a nested table of the cust address tab type data type:

```
CREATE TYPE cust_address_tab_tab_typ
  AS TABLE OF cust_address_tab_typ;
/
```

Next, duplicate the elements in all the nested table rows to increase the cardinality of the nested table rows to 2:

```
UPDATE customers_demo
SET cust address ntab = cust address ntab MULTISET UNION cust address ntab;
```

Now, select the nested table column <code>cust_address_ntab</code> from the <code>customers_demo</code> table using the <code>powermultiset</code> by <code>Cardinality</code> function:

```
SELECT CAST (POWERMULTISET_BY_CARDINALITY (cust_address_ntab, 2)

AS cust_address_tab_tab_typ)

FROM customers_demo;

CAST (POWERMULTISET_BY_CARDINALITY (CUST_ADDRESS_NTAB, 2) AS CUST_ADDRESS_TAB_TAB_TYP)

(STREET_ADDRESS, POSTAL_CODE, CITY, STATE_PROVINCE, COUNTRY_ID)

CUST_ADDRESS_TAB_TAB_TYP(CUST_ADDRESS_TAB_TYP

(CUST_ADDRESS_TYP('514 W Superior St', '46901', 'Kokomo', 'IN', 'US'),

CUST_ADDRESS_TYP('514 W Superior St', '46901', 'Kokomo', 'IN', 'US')))

CUST_ADDRESS_TAB_TAB_TYP(CUST_ADDRESS_TAB_TYP

(CUST_ADDRESS_TYP('2515 Bloyd Ave', '46218', 'Indianapolis', 'IN', 'US')),

CUST_ADDRESS_TYP('2515 Bloyd Ave', '46218', 'Indianapolis', 'IN', 'US')))

CUST_ADDRESS_TYP('8768 N State Rd 37', '47404', 'Bloomington', 'IN', 'US')))

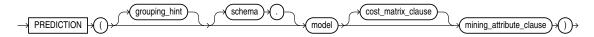
CUST_ADDRESS_TYP('8768 N State Rd 37', '47404', 'Bloomington', 'IN', 'US')))
```

The preceding example requires the <code>customers_demo</code> table and a nested table column containing data. Refer to "Multiset Operators" to create this table and nested table columns.

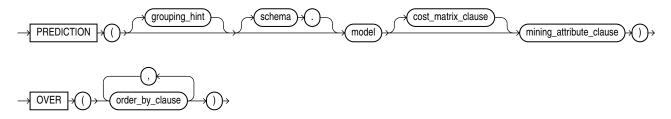
PREDICTION

Syntax

prediction::=

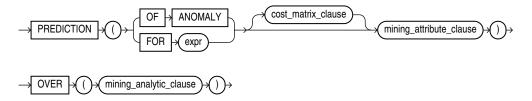


prediction_ordered::=

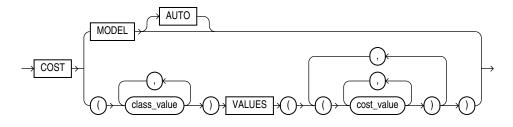


Analytic Syntax

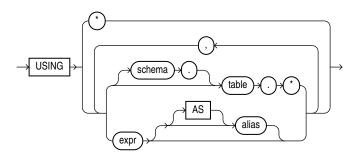
prediction_analytic::=



cost_matrix_clause::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions" for information on the syntax, semantics, and restrictions of $mining\ analytic\ clause$

Purpose

PREDICTION returns a prediction for each row in the selection. The data type of the returned prediction depends on whether the function performs Regression, Classification, or Anomaly Detection.

- **Regression**: Returns the expected target value for each row. The data type of the return value is the data type of the target.
- **Classification**: Returns the most probable target class (or lowest cost target class, if costs are specified) for each row. The data type of the return value is the data type of the target.
- **Anomaly Detection**: Returns 1 or 0 for each row. Typical rows are classified as 1. Rows that differ significantly from the rest of the data are classified as 0.

cost matrix clause

Costs are a biasing factor for minimizing the most harmful kinds of misclassifications. You can specify <code>cost_matrix_clause</code> for Classification or Anomaly Detection. Costs are not relevant for Regression. The <code>cost_matrix_clause</code> behaves as described for "PREDICTION_COST".

Syntax Choice

PREDICTION can score data by applying a mining model object to the data, or it can dynamically score the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

• **Syntax**: Use the *prediction* syntax to score the data with a pre-defined model. Supply the name of a model that performs classification, regression, or anomaly detection.

Use the <code>prediction_ordered</code> syntax for a model that requires ordered data, such as an MSET-SPRT model. The <code>prediction_ordered</code> syntax requires an <code>order_by_clause</code> clause.

Restrictions on the <code>prediction_ordered</code> syntax are that you cannot use it in the <code>WHERE</code> clause of a query. Also, you cannot use a <code>query_partition_clause</code> or a <code>windowing clause</code> with the <code>prediction ordered</code> syntax.

For details about the order by clause, see "Analytic Functions".

Analytic Syntax: Use the analytic syntax to score the data without a pre-defined model. The analytic syntax uses <code>mining_analytic_clause</code>, which specifies if the data should be partitioned for multiple model builds. The <code>mining_analytic_clause</code> supports a query partition clause and an order by clause. (See "analytic_clause::=".)



- For Regression, specify FOR expr, where expr is an expression that identifies a target column that has a numeric data type.
- For Classification, specify FOR expr, where expr is an expression that identifies a target column that has a character data type.
- For Anomaly Detection, specify the keywords OF ANOMALY.

The syntax of the PREDICTION function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining attribute clause identifies the column attributes to use as predictors for scoring.

- If you specify USING *, all the relevant attributes present in the input row are used.
- If you invoke the function with the analytic syntax, the <code>mining_attribute_clause</code> is used both for building the transient models and for scoring.
- It you invoke the function with a pre-defined model, the <code>mining_attribute_clause</code> should include all or some of the attributes that were used to create the model. The following conditions apply:
 - If mining_attribute_clause includes an attribute with the same name but a different data type from the one that was used to create the model, then the data type is converted to the type expected by the model.
 - If you specify more attributes for scoring than were used to create the model, then the extra attributes are silently ignored.
 - If you specify fewer attributes for scoring than were used to create the model, then scoring is performed on a best-effort basis.

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about predictive Oracle Machine Learning for SQL.
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of PREDICTION when it is a character value

Note:

The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

In this example, the model dt_sh_clas_sample predicts the gender and age of customers who are most likely to use an affinity card (target = 1). The PREDICTION function takes into account



the cost matrix associated with the model and uses marital status, education, and household size as predictors.

The cost matrix associated with the model $dt_sh_clas_sample$ is stored in the table $dt_sh_sample_costs$. The cost matrix specifies that the misclassification of 1 is 8 times more costly than the misclassification of 0.

Analytic Example

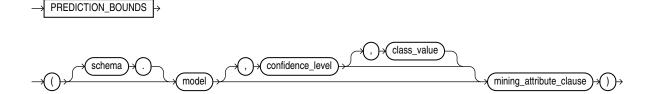
In this example, dynamic regression is used to predict the age of customers who are likely to use an affinity card. The query returns the 3 customers whose predicted age is most different from the actual. The query includes information about the predictors that have the greatest influence on the prediction.

```
SELECT cust id, age, pred age, age-pred age age diff, pred det FROM
   (SELECT cust_id, age, pred_age, pred_det,
        RANK() OVER (ORDER BY ABS(age-pred age) desc) rnk FROM
   (SELECT cust id, age,
          PREDICTION (FOR age USING *) OVER () pred age,
          PREDICTION DETAILS(FOR age ABS USING *) OVER () pred det
   FROM mining data apply v))
 WHERE rnk <= 3;
CUST_ID AGE PRED_AGE AGE_DIFF PRED_DET
______
100910 80 40.67 39.33 <Details algorithm="Support Vector Machines">
                             <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".059"</pre>
                             rank="1"/>
                             <Attribute name="Y BOX GAMES" actualValue="0" weight=".059"</pre>
                             rank="2"/>
                             <Attribute name="AFFINITY CARD" actualValue="0" weight=".059"</pre>
                             rank="3"/>
                             <Attribute name="FLAT PANEL MONITOR" actualValue="1" weight=".059"</pre>
                             rank="4"/>
                             <Attribute name="YRS RESIDENCE" actualValue="4" weight=".059"</pre>
                             rank="5"/>
                             </Details>
101285 79 42.18
                       36.82 <Details algorithm="Support Vector Machines">
                             <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".059"</pre>
                              rank="1"/>
```

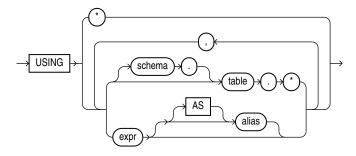
```
<Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".059"</pre>
                                  rank="2"/>
                                 <Attribute name="CUST_MARITAL_STATUS" actualValue="Mabsent"</pre>
                                  weight=".059" rank="3"/>
                                 <a href="Attribute name="Y_BOX_GAMES" actualValue="0" weight=".059"
                                  rank="4"/>
                                  <Attribute name="OCCUPATION" actualValue="Prof." weight=".059"</pre>
                                  rank="5"/>
                                  </Details>
100694
            77 41.04
                          35.96 <Details algorithm="Support Vector Machines">
                                 <Attribute name="HOME_THEATER_PACKAGE" actualValue="1" weight=".059"</pre>
                                  rank="1"/>
                                 <Attribute name="EDUCATION" actualValue="&lt; Bach." weight=".059"</pre>
                                  rank="2"/>
                                 <Attribute name="Y BOX GAMES" actualValue="0" weight=".059"</pre>
                                  rank="3"/>
                                 <Attribute name="CUST ID" actualValue="100694" weight=".059"</pre>
                                  rank="4"/>
                                 <a href="COUNTRY NAME" actualValue="United States of">COUNTRY NAME</a>" actualValue="United States of
                                  America" weight=".059" rank="5"/>
```

PREDICTION BOUNDS

Syntax



mining_attribute_clause::=



Purpose

PREDICTION_BOUNDS applies a Generalized Linear Model (GLM) to predict a class or a value for each row in the selection. The function returns the upper and lower bounds of each prediction in a varray of objects with fields UPPER and LOWER.

GLM can perform either regression or binary classification:

- The bounds for regression refer to the predicted target value. The data type of UPPER and LOWER is the data type of the target.
- The bounds for binary classification refer to the probability of either the predicted target class or the specified <code>class_value</code>. The data type of <code>upper</code> and <code>Lower</code> is <code>Binary_Double</code>.

If the model was built using ridge regression, or if the covariance matrix is found to be singular during the build, then PREDICTION BOUNDS returns NULL for both bounds.

confidence_level is a number in the range (0,1). The default value is 0.95. You can specify class_value while leaving confidence_level at its default by specifying NULL for confidence level.

The syntax of the PREDICTION_BOUNDS function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. This clause behaves as described for the PREDICTION function. (Note that the reference to analytic syntax does not apply.) See "mining_attribute_clause::=".

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring
- Oracle Machine Learning for SQL Concepts for information about Generalized Linear Models

Note:

The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

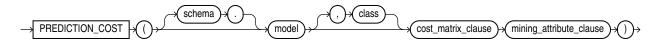
The following example returns the distribution of customers whose ages are predicted with 98% confidence to be greater than 24 and less than 46.



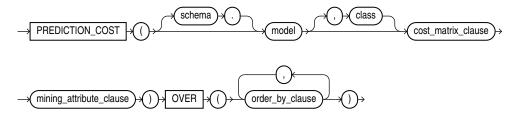
PREDICTION_COST

Syntax

prediction_cost::=

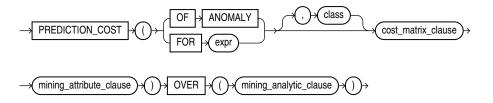


prediction_cost_ordered::=

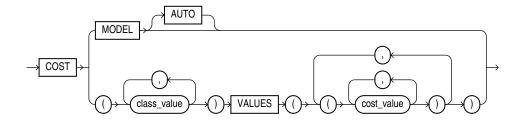


Analytic Syntax

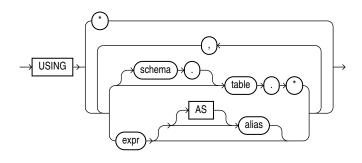
prediction_cost_analytic::=



cost_matrix_clause::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions " for information on the syntax, semantics, and restrictions of mining_analytic_clause

Purpose

PREDICTION_COST returns a cost for each row in the selection. The cost refers to the lowest cost class or to the specified class. The cost is returned as BINARY DOUBLE.

PREDICTION_COST can perform classification or anomaly detection. For classification, the returned cost refers to a predicted target class. For anomaly detection, the returned cost refers to a classification of 1 (for typical rows) or 0 (for anomalous rows).

You can use PREDICTION_COST in conjunction with the PREDICTION function to obtain the prediction and the cost of the prediction.

cost_matrix_clause

Costs are a biasing factor for minimizing the most harmful kinds of misclassifications. For example, false positives might be considered more costly than false negatives. Costs are specified in a cost matrix that can be associated with the model or defined inline in a VALUES clause. All classification algorithms can use costs to influence scoring.

Decision Tree is the only algorithm that can use costs to influence the model build. The cost matrix used to build a Decision Tree model is also the default scoring cost matrix for the model.

The following cost matrix table specifies that the misclassification of 1 is five times more costly than the misclassification of 0.

ACTUAL_TARGET_VALUE	PREDICTED_TARGET_VALUE	COST
0	0	0
0	1	1
1	0	5
1	1	0

In cost matrix clause:

- COST MODEL indicates that scoring should be performed by taking into account the scoring cost matrix associated with the model. If the cost matrix does not exist, then the function returns an error.
- COST MODEL AUTO indicates that the existence of a cost matrix is unknown. If a cost matrix exists, then the function uses it to return the lowest cost prediction. Otherwise the function returns the highest probability prediction.
- The VALUES clause specifies an inline cost matrix for class_value. For example, you could
 specify that the misclassification of 1 is five times more costly than the misclassification of
 0 as follows:

```
PREDICTION (nb model COST (0,1) VALUES ((0, 1), (1, 5)) USING *)
```

If a model that has a scoring cost matrix is invoked with an inline cost matrix, then the inline costs are used.

See Also:

Oracle Machine Learning for SQL User's Guide for more information about costsensitive prediction.

Syntax Choice

PREDICTION_COST can score the data by applying a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

• **Syntax**: Use the *prediction_cost* syntax to score the data with a pre-defined model. Supply the name of a model that performs classification or anomaly detection.

Use the <code>prediction_cost_ordered</code> syntax for a model that requires ordered data, such as an MSET-SPRT model. The <code>prediction_cost_ordered</code> syntax requires an <code>order</code> by <code>clause</code> clause.

Restrictions on the <code>prediction_cost_ordered</code> syntax are that you cannot use it in the <code>WHERE</code> clause of a query. Also, you cannot use a <code>query_partition_clause</code> or a <code>windowing clause</code> with the <code>prediction ordered</code> syntax.

- Analytic Syntax: Use the analytic syntax to score the data without a pre-defined model.
 The analytic syntax uses mining_analytic_clause, which specifies if the data should be
 partitioned for multiple model builds. The mining_analytic_clause supports a
 query_partition_clause and an order_by_clause. (See "analytic_clause::=".)
 - For classification, specify FOR expr, where expr is an expression that identifies a target column that has a character data type.
 - For anomaly detection, specify the keywords OF ANOMALY.

The syntax of the PREDICTION_COST function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining_attribute_clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about classification with costs





The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

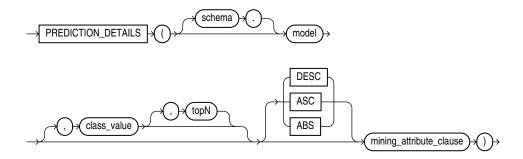
This example predicts the ten customers in Italy who would respond to the least expensive sales campaign (offering an affinity card).

```
SELECT cust_id
FROM (SELECT cust id, rank()
       OVER (ORDER BY PREDICTION_COST(DT_SH_Clas_sample, 1 COST MODEL USING *)
           ASC, cust id) rnk
        FROM mining_data_apply_v
        WHERE country_name = 'Italy')
 WHERE rnk <= 10
 ORDER BY rnk;
  CUST_ID
   100081
   100179
   100185
   100324
   100344
   100554
   100662
   100733
   101250
   101306
```

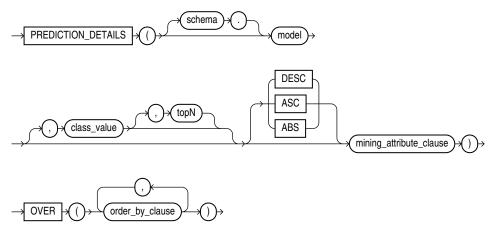
PREDICTION_DETAILS

Syntax

prediction_details::=

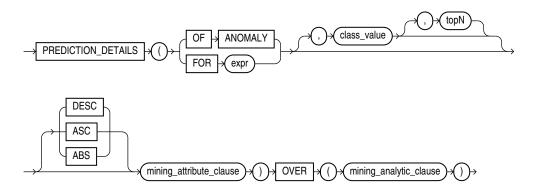


prediction_details_ordered::=

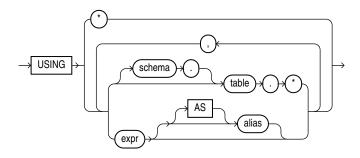


Analytic Syntax

prediction_details_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions" for information on the syntax, semantics, and restrictions of mining_analytic_clause

Purpose

PREDICTION_DETAILS returns prediction details for each row in the selection. The return value is an XML string that describes the attributes of the prediction.

For regression, the returned details refer to the predicted target value. For classification and anomaly detection, the returned details refer to the highest probability class or the specified *class value*.

topN

If you specify a value for topN, the function returns the N attributes that have the most influence on the prediction (the score). If you do not specify topN, the function returns the 5 most influential attributes.

DESC, ASC, or ABS

The returned attributes are ordered by weight. The weight of an attribute expresses its positive or negative impact on the prediction. For regression, a positive weight indicates a higher value prediction; a negative weight indicates a lower value prediction. For classification and anomaly detection, a positive weight indicates a higher probability prediction; a negative weight indicates a lower probability prediction.

By default, PREDICTION_DETAILS returns the attributes with the highest positive weight (DESC). If you specify ASC, the attributes with the highest negative weight are returned. If you specify ABS, the attributes with the greatest weight, whether negative or positive, are returned. The results are ordered by absolute value from highest to lowest. Attributes with a zero weight are not included in the output.

Syntax Choice

PREDICTION_DETAILS can score the data by applying a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- **Syntax**: Use the *prediction_details* syntax to score the data with a pre-defined model. Supply the name of a model that performs classification, regression, or anomaly detection.
 - Use the <code>prediction_details_ordered</code> syntax for a model that requires ordered data, such as an MSET-SPRT model. The <code>prediction_details_ordered</code> syntax requires an <code>order</code> by <code>clause</code> clause.
 - Restrictions on the <code>prediction_details_ordered</code> syntax are that you cannot use it in the <code>WHERE</code> clause of a query. Also, you cannot use a <code>query_partition_clause</code> or a <code>windowing_clause</code> with the <code>prediction_details_ordered</code> syntax.
- Analytic Syntax: Use the analytic syntax to score the data without a pre-defined model.
 The analytic syntax uses mining_analytic_clause, which specifies if the data should be
 partitioned for multiple model builds. The mining_analytic_clause supports a
 query partition clause and an order by clause. (See "analytic_clause::=".)



- For classification, specify FOR expr, where expr is an expression that identifies a target column that has a character data type.
- For regression, specify FOR expr, where expr is an expression that identifies a target column that has a numeric data type.
- For anomaly detection, specify the keywords OF ANOMALY.

The syntax of the PREDICTION_DETAILS function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining_attribute_clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about predictive Oracle Machine Learning for SQL.

Note:

The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

This example uses the model <code>svmr_sh_regr_sample</code> to score the data. The query returns the three attributes that have the greatest influence on predicting a higher value for customer age.

Analytic Syntax

This example dynamically identifies customers whose age is not typical for the data. The query returns the attributes that predict or detract from a typical age.

```
SELECT cust_id, age, pred_age, age-pred_age age_diff, pred_det FROM (SELECT cust_id, age, pred_age, pred_det,
```



```
RANK() OVER (ORDER BY ABS(age-pred age) DESC) rnk
          FROM (SELECT cust id, age,
             PREDICTION(FOR age USING *) OVER () pred_age,
             PREDICTION DETAILS(FOR age ABS USING *) OVER () pred det
             FROM mining_data_apply_v))
    WHERE rnk <= 5;
CUST_ID AGE PRED_AGE AGE_DIFF PRED_DET
_____ ___
100910 80
             40.67
                      39.33 <Details algorithm="Support Vector Machines">
                               <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".059"</pre>
                                rank="1"/>
                               <Attribute name="Y BOX GAMES" actualValue="0" weight=".059"</pre>
                                rank="2"/>
                               <Attribute name="AFFINITY CARD" actualValue="0" weight=".059"</pre>
                                rank="3"/>
                               <Attribute name="FLAT PANEL MONITOR" actualValue="1" weight=".059"</pre>
                                rank="4"/>
                               <Attribute name="YRS RESIDENCE" actualValue="4" weight=".059"</pre>
                                rank="5"/>
                               </Details>
 101285 79 42.18
                      36.82 <Details algorithm="Support Vector Machines">
                                <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".059"</pre>
                                 rank="1"/>
                                <Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".059"</pre>
                                 rank="2"/>
                                <Attribute name="CUST MARITAL STATUS" actualValue="Mabsent"</pre>
                                 weight=".059" rank="3"/>
                                <Attribute name="Y BOX GAMES" actualValue="0" weight=".059"</pre>
                                <Attribute name="OCCUPATION" actualValue="Prof." weight=".059"</pre>
                                 rank="5"/>
                                </Details>
 100694 77 41.04 35.96 <Details algorithm="Support Vector Machines">
                                 <Attribute name="HOME THEATER PACKAGE" actualValue="1"</pre>
                                  weight=".059" rank="1"/>
                                 <Attribute name="EDUCATION" actualValue="&lt; Bach." weight=".059"</pre>
                                  rank="2"/>
                                 <Attribute name="Y BOX GAMES" actualValue="0" weight=".059"</pre>
                                  rank="3"/>
                                 <Attribute name="CUST ID" actualValue="100694" weight=".059"</pre>
                                 <a href="COUNTRY NAME" actualValue="United States of">CAttribute name="COUNTRY NAME" actualValue="United States of</a>
                                  America" weight=".059" rank="5"/>
                                 </Details>
 100308 81 45.33 35.67 <Details algorithm="Support Vector Machines">
                                <Attribute name="HOME THEATER PACKAGE" actualValue="1" weight=".059"</pre>
                                 rank="1"/>
                                <a href="Attribute name="Y BOX GAMES" actualValue="0" weight=".059"
                                 rank="2"/>
                                <a href="Attribute name="HOUSEHOLD SIZE" actualValue="2" weight=".059"</a>
                                 rank="3"/>
                                <Attribute name="FLAT PANEL MONITOR" actualValue="1" weight=".059"</pre>
                                 rank="4"/>
                                <Attribute name="CUST GENDER" actualValue="F" weight=".059"</pre>
                                 rank="5"/>
                                </Details>
101256 90 54.39 35.61 <Details algorithm="Support Vector Machines">
```

```
<Attribute name="YRS_RESIDENCE" actualValue="9" weight=".059"
  rank="1"/>
<Attribute name="HOME_THEATER_PACKAGE" actualValue="1" weight=".059"
  rank="2"/>
<Attribute name="EDUCATION" actualValue="&lt; Bach." weight=".059"
  rank="3"/>
<Attribute name="Y_BOX_GAMES" actualValue="0" weight=".059"
  rank="4"/>
<Attribute name="COUNTRY_NAME" actualValue="United States of
  America" weight=".059" rank="5"/>
</Details>
```

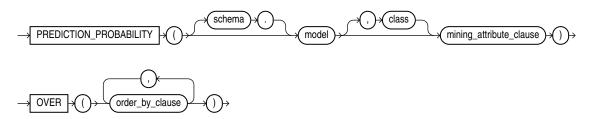
PREDICTION_PROBABILITY

Syntax

prediction_probability::=

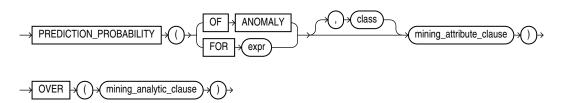


prediction_probability_ordered::=

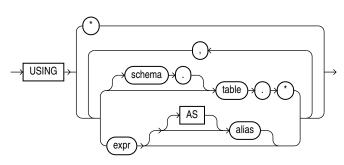


Analytic Syntax

prediction_prob_analytic::=



mining_attribute_clause::=



mining_analytic_clause::=



See Also:

"Analytic Functions" for information on the syntax, semantics, and restrictions of mining analytic clause

Purpose

PREDICTION_PROBABILITY returns a probability for each row in the selection. The probability refers to the highest probability class or to the specified *class*. The data type of the returned probability is BINARY DOUBLE.

PREDICTION_PROBABILITY can perform classification or anomaly detection. For classification, the returned probability refers to a predicted target class. For anomaly detection, the returned probability refers to a classification of 1 (for typical rows) or 0 (for anomalous rows).

You can use PREDICTION_PROBABILITY in conjunction with the PREDICTION function to obtain the prediction and the probability of the prediction.

Syntax Choice

PREDICTION_PROBABILITY can score the data by applying a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

- **Syntax**: Use the *prediction_probability* syntax to score the data with a pre-defined model. Supply the name of a model that performs classification or anomaly detection.
 - Use the <code>prediction_probability_ordered</code> syntax for a model that requires ordered data, such as an MSET-SPRT model. The <code>prediction_probability_ordered</code> syntax requires an <code>order</code> by <code>clause</code> clause.
 - Restrictions on the <code>prediction_probability_ordered</code> syntax are that you cannot use it in the <code>WHERE</code> clause of a query. Also, you cannot use a <code>query_partition_clause</code> or a <code>windowing_clause</code> with the <code>prediction_probability_ordered</code> syntax.
- Analytic Syntax: Use the analytic syntax to score the data without a pre-defined model.
 The analytic syntax uses mining_analytic_clause, which specifies if the data should be
 partitioned for multiple model builds. The mining_analytic_clause supports a
 query_partition_clause and an order_by_clause. (See "analytic_clause::=".)
 - For classification, specify FOR expr, where expr is an expression that identifies a target column that has a character data type.
 - For anomaly detection, specify the keywords OF ANOMALY.

The syntax of the PREDICTION_PROBABILITY function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.



mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining_attribute_clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about predictive Oracle Machine Learning for SQL.

Note:

The following examples are excerpted from the Oracle Machine Learning for SQL sample programs. For information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

The following example returns the 10 customers living in Italy who are most likely to use an affinity card.

```
SELECT cust_id FROM (
  SELECT cust id
  FROM mining_data_apply_v
  WHERE country name = 'Italy'
  ORDER BY PREDICTION PROBABILITY(DT SH Clas sample, 1 USING *)
     DESC, cust_id)
  WHERE rownum < 11;
  CUST ID
   100081
   100179
   100185
   100324
   100344
   100554
   100662
   100733
   101250
   101306
```

Analytic Example

This example identifies rows that are most atypical in the data in mining_data_one_class_v. Each type of marital status is considered separately so that the most anomalous rows per marital status group are returned.

The query returns three attributes that have the most influence on the determination of anomalous rows. The PARTITION BY clause causes separate models to be built and applied for

each marital status. Because there is only one record with status Mabsent, no model is created for that partition (and no details are provided).

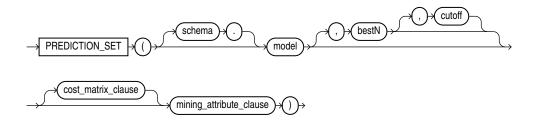
```
SELECT cust_id, cust_marital_status, rank_anom, anom_det FROM
    (SELECT cust id, cust marital status, anom det,
           rank() OVER (PARTITION BY CUST MARITAL STATUS
                        ORDER BY ANOM PROB DESC, cust id) rank anom FROM
     (SELECT cust id, cust marital status,
           PREDICTION PROBABILITY (OF ANOMALY, 0 USING *)
             OVER (PARTITION BY CUST MARITAL STATUS) anom prob,
           PREDICTION DETAILS (OF ANOMALY, 0, 3 USING *)
             OVER (PARTITION BY CUST MARITAL STATUS) anom det
     FROM mining data one class v
   ))
  WHERE rank_anom < 3 order by 2, 3;
CUST ID CUST MARITAL STATUS RANK ANOM ANOM DET
______
102366 Divorc.
                                     <Details algorithm="Support Vector Machines" class="0">
                           1
                                      <Attribute name="COUNTRY NAME" actualValue="United Kingdom"</pre>
                                       weight=".069" rank="1"/>
                                      <Attribute name="AGE" actualValue="28" weight=".013"</pre>
                                       rank="2"/>
                                      <Attribute name="YRS RESIDENCE" actualValue="4"</pre>
                                       weight=".006" rank="3"/>
                                      </Details>
101817 Divorc.
                           2
                                      <Details algorithm="Support Vector Machines" class="0">
                                      <Attribute name="YRS RESIDENCE" actualValue="8"</pre>
                                       weight=".018" rank="1"/>
                                      <Attribute name="EDUCATION" actualValue="PhD" weight=".007"</pre>
                                       rank="2"/>
                                      <Attribute name="CUST INCOME LEVEL" actualValue="K:</pre>
                                       250\,000 - 299\,999" weight=".006" rank="3"/>
                                      </Details>
101713 Mabsent
                                      <Details algorithm="Support Vector Machines" class="0">
101790 Married
                           1
                                      <Attribute name="COUNTRY NAME" actualValue="Canada"</pre>
                                       weight=".063" rank="1"/>
                                      <Attribute name="EDUCATION" actualValue="7th-8th"</pre>
                                       weight=".011" rank="2"/>
                                      <Attribute name="HOUSEHOLD SIZE" actualValue="4-5"</pre>
                                       weight=".011" rank="3"/>
                                      </Details>
```

ORACLE[®]

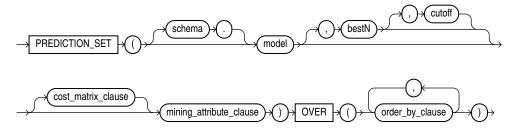
PREDICTION_SET

Syntax

prediction_set::=

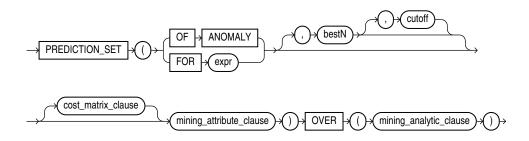


prediction_set_ordered::=

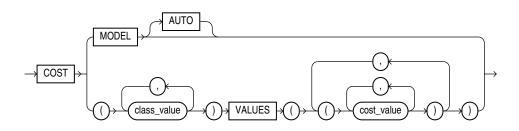


Analytic Syntax

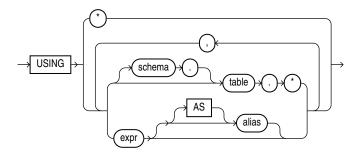
prediction_set_analytic::=



cost_matrix_clause::=



mining_attribute_clause::=



mining_analytic_clause::-



See Also:

"Analytic Functions " for information on the syntax, semantics, and restrictions of mining_analytic_clause

Purpose

PREDICTION_SET returns a set of predictions with either probabilities or costs for each row in the selection. The return value is a varray of objects with field names PREDICTION_ID and PROBABILITY OR COST. The prediction identifier has the data type of the target; the probability and cost fields are BINARY DOUBLE.

PREDICTION_SET can perform classification or anomaly detection. For classification, the return value refers to a predicted target class. For anomaly detection, the return value refers to a classification of 1 (for typical rows) or 0 (for anomalous rows).

bestN and cutoff

You can specify <code>bestN</code> and <code>cutoff</code> to limit the number of predictions returned by the function. By default, both <code>bestN</code> and <code>cutoff</code> are null and all predictions are returned.

- bestN is the N predictions that are either the most probable or the least costly. If multiple predictions share the Nth probability or cost, then the function chooses one of them.
- cutoff is a value threshold. Only predictions with probability greater than or equal to cutoff, or with cost less than or equal to cutoff, are returned. To filter by cutoff only, specify NULL for bestN. If the function uses a cost_matrix_clause with COST MODEL AUTO, then cutoff is ignored.

You can specify bestN with cutoff to return up to the N most probable predictions that are greater than or equal to cutoff. If costs are used, specify bestN with cutoff to return up to the N least costly predictions that are less than or equal to cutoff.



cost matrix clause

You can specify <code>cost_matrix_clause</code> as a biasing factor for minimizing the most harmful kinds of misclassifications. <code>cost_matrix_clause</code> behaves as described for "PREDICTION_COST".

Syntax Choice

PREDICTION_SET can score the data by applying a mining model object to the data, or it can dynamically mine the data by executing an analytic clause that builds and applies one or more transient mining models. Choose **Syntax** or **Analytic Syntax**:

• **Syntax**: Use the *prediction_set* syntax to score the data with a pre-defined model. Supply the name of a model that performs classification or anomaly detection.

Use the <code>prediction_set_ordered</code> syntax for a model that requires ordered data, such as an MSET-SPRT model. The <code>prediction_set_ordered</code> syntax requires an <code>order</code> by <code>clause</code> clause.

Restrictions on the <code>prediction_set_ordered</code> syntax are that you cannot use it in the <code>WHERE</code> clause of a query. Also, you cannot use a <code>query_partition_clause</code> or a <code>windowing clause</code> with the <code>prediction set ordered</code> syntax.

- Analytic Syntax: Use the analytic syntax to score the data without a pre-defined model.
 The analytic syntax uses mining_analytic_clause, which specifies if the data should be
 partitioned for multiple model builds. The mining_analytic_clause supports a
 query partition clause and an order by clause. (See "analytic_clause::=".)
 - For classification, specify FOR expr, where expr is an expression that identifies a target column that has a character data type.
 - For anomaly detection, specify the keywords OF ANOMALY.

The syntax of the PREDICTION_SET function can use an optional GROUPING hint when scoring a partitioned model. See GROUPING Hint.

mining_attribute_clause

mining_attribute_clause identifies the column attributes to use as predictors for scoring. When the function is invoked with the analytic syntax, these predictors are also used for building the transient models. The mining_attribute_clause behaves as described for the PREDICTION function. (See "mining_attribute_clause::=".)

See Also:

- Oracle Machine Learning for SQL User's Guide for information about scoring.
- Oracle Machine Learning for SQL Concepts for information about predictive Oracle Machine Learning for SQL.





The following example is excerpted from the Oracle Machine Learning for SQL sample programs. For more information about the sample programs, see Appendix A in *Oracle Machine Learning for SQL User's Guide*.

Example

This example lists the probability and cost that customers with ID less than 100006 will use an affinity card. This example has a binary target, but such a query is also useful for multiclass classification such as low, medium, and high.

```
SELECT T.cust_id, S.prediction, S.probability, S.cost

FROM (SELECT cust_id,

PREDICTION_SET(dt_sh_clas_sample COST MODEL USING *) pset

FROM mining_data_apply_v

WHERE cust_id < 100006) T,

TABLE(T.pset) S

ORDER BY cust_id, S.prediction;
```

CUST_ID	PREDICTION	PROBABILITY	COST
100001	0	.966183575	.270531401
100001	1	.033816425	.966183575
100002	0	.740384615	2.076923077
100002	1	.259615385	.740384615
100003	0	.909090909	.727272727
100003	1	.090909091	.909090909
100004	0	.909090909	.727272727
100004	1	.090909091	.909090909
100005	0	.272357724	5.821138211
100005	1	.727642276	.272357724

PRESENTNNV

Syntax



Purpose

The PRESENTNNV function can be used only in the $model_clause$ of the SELECT statement and then only on the right-hand side of a model rule. It returns expr1 when $cell_reference$ exists prior to the execution of the $model_clause$ and is not null when PRESENTNNV is evaluated. Otherwise it returns expr2. This function differs from NVL2 in that NVL2 evaluates the data at the time it is executed, rather than evaluating the data as it was prior to the execution of the $model_clause$.



See Also:

- model_clause and "Model Expressions" for the syntax and semantics
- NVL2 for comparison
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of PRESENTINIV when it is a character value

Examples

In the following example, if a row containing sales for the Mouse Pad for the year 2002 exists, and the sales value is not null, then the sales value remains unchanged. If the row exists and the sales value is null, then the sales value is set to 10. If the row does not exist, then the row is created with the sales value set to 10.

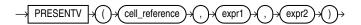
COUNTRY	PROD	YEAR	S
	Marian Park	1000	2500 42
France	Mouse Pad	1998	2509.42
France	Mouse Pad	1999	3678.69
France	Mouse Pad	2000	3000.72
France	Mouse Pad	2001	3269.09
France	Mouse Pad	2002	10
France	Standard Mouse	1998	2390.83
France	Standard Mouse	1999	2280.45
France	Standard Mouse	2000	1274.31
France	Standard Mouse	2001	2164.54
Germany	Mouse Pad	1998	5827.87
Germany	Mouse Pad	1999	8346.44
Germany	Mouse Pad	2000	7375.46
Germany	Mouse Pad	2001	9535.08
Germany	Mouse Pad	2002	10
Germany	Standard Mouse	1998	7116.11
Germany	Standard Mouse	1999	6263.14
Germany	Standard Mouse	2000	2637.31
Germany	Standard Mouse	2001	6456.13

18 rows selected.

The preceding example requires the view <code>sales_view_ref</code>. Refer to "Examples" to create this view.

PRESENTV

Syntax



Purpose

The PRESENTV function can be used only within the <code>model_clause</code> of the <code>SELECT</code> statement and then only on the right-hand side of a model rule. It returns <code>expr1</code> when, prior to the execution of the <code>model_clause</code>, <code>cell_reference</code> exists. Otherwise it returns <code>expr2</code>.

See Also:

- model_clause and "Model Expressions" for the syntax and semantics
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of PRESENTV when it is a character value

Examples

In the following example, if a row containing sales for the Mouse Pad for the year 2000 exists, then the sales value for the Mouse Pad for the year 2001 is set to the sales value for the Mouse Pad for the year 2000. If the row does not exist, then a row is created with the sales value for the Mouse Pad for year 20001 set to 0.

```
SELECT country, prod, year, s
FROM sales_view_ref
MODEL

PARTITION BY (country)
DIMENSION BY (prod, year)
MEASURES (sale s)
IGNORE NAV
UNIQUE DIMENSION
RULES UPSERT SEQUENTIAL ORDER
(
    s['Mouse Pad', 2001] =
        PRESENTV(s['Mouse Pad', 2000], s['Mouse Pad', 2000], 0)
)
ORDER BY country, prod, year;
```

COUNTRY	PROD	YEAR	S
France	Mouse Pad	1998	2509.42
France	Mouse Pad	1999	3678.69
France	Mouse Pad	2000	3000.72
France	Mouse Pad	2001	3000.72
France	Standard Mouse	1998	2390.83
France	Standard Mouse	1999	2280.45
France	Standard Mouse	2000	1274.31
France	Standard Mouse	2001	2164.54
Germany	Mouse Pad	1998	5827.87



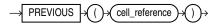
Germany	Mouse Pad	1999	8346.44
Germany	Mouse Pad	2000	7375.46
Germany	Mouse Pad	2001	7375.46
Germany	Standard Mouse	1998	7116.11
Germany	Standard Mouse	1999	6263.14
Germany	Standard Mouse	2000	2637.31
Germany	Standard Mouse	2001	6456.13

16 rows selected.

The preceding example requires the view <code>sales_view_ref</code>. Refer to "The MODEL clause: Examples" to create this view.

PREVIOUS

Syntax



Purpose

The PREVIOUS function can be used only in the <code>model_clause</code> of the <code>SELECT</code> statement and then only in the <code>ITERATE</code> ... [UNTIL] clause of the <code>model_rules_clause</code>. It returns the value of <code>cell reference</code> at the beginning of each iteration.

See Also:

- model_clause and "Model Expressions" for the syntax and semantics
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of PREVIOUS when it is a character value

Examples

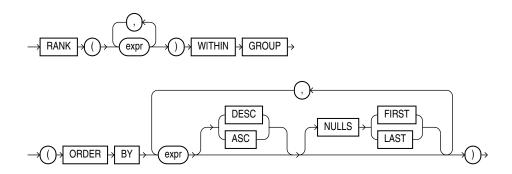
The following example repeats the rules, up to 1000 times, until the difference between the values of cur val at the beginning and at the end of an iteration is less than one:



RANK

Aggregate Syntax

rank_aggregate::=



Analytic Syntax

rank_analytic::=



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

RANK calculates the rank of a value in a group of values. The return type is NUMBER.

See Also:

Table 2-9 for more information on implicit conversion and "Numeric Precedence" for information on numeric precedence

Rows with equal values for the ranking criteria receive the same rank. Oracle Database then adds the number of tied rows to the tied rank to calculate the next rank. Therefore, the ranks may not be consecutive numbers. This function is useful for top-N and bottom-N reporting.

As an aggregate function, RANK calculates the rank of a hypothetical row identified by the
arguments of the function with respect to a given sort specification. The arguments of the
function must all evaluate to constant expressions within each aggregate group, because
they identify a single row within each group. The constant argument expressions and the

- expressions in the ORDER BY clause of the aggregate match by position. Therefore, the number of arguments must be the same and their types must be compatible.
- As an analytic function, RANK computes the rank of each row returned from a query with respect to the other rows returned by the query, based on the values of the <code>value_exprs</code> in the <code>order by clause</code>.



Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules, which define the collation RANK uses to compare character values for the ORDER BY clause

Aggregate Example

The following example calculates the rank of a hypothetical employee in the sample table hr.employees with a salary of \$15,500 and a commission of 5%:

```
SELECT RANK(15500, .05) WITHIN GROUP

(ORDER BY salary, commission_pct) "Rank"

FROM employees;

Rank
------
105
```

Similarly, the following query returns the rank for a \$15,500 salary among the employee salaries:

```
SELECT RANK(15500) WITHIN GROUP
(ORDER BY salary DESC) "Rank of 15500"
FROM employees;

Rank of 15500
```

Analytic Example

The following statement ranks the employees in the sample hr schema in department 60 based on their salaries. Identical salary values receive the same rank and cause nonconsecutive ranks. Compare this example with the analytic example for DENSE RANK.

DEPARTMENT_ID	LAST_NAME	SALARY	RANK
60	Lorentz	4200	1
60	Austin	4800	2
60	Pataballa	4800	2
60	Ernst	6000	4
60	Hunold	9000	5



RATIO_TO_REPORT

Syntax





"Analytic Functions " for information on syntax, semantics, and restrictions, including valid forms of <code>expr</code>

Purpose

RATIO_TO_REPORT is an analytic function. It computes the ratio of a value to the sum of a set of values. If expr evaluates to null, then the ratio-to-report value also evaluates to null.

The set of values is determined by the <code>query_partition_clause</code>. If you omit that clause, then the ratio-to-report is computed over all rows returned by the query.

You cannot nest analytic functions by using RATIO_TO_REPORT or any other analytic function for *expr*. However, you can use other built-in function expressions for *expr*. Refer to "About SQL Expressions" for information on valid forms of *expr*.

Examples

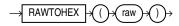
The following example calculates the ratio-to-report value of each purchasing clerk's salary to the total of all purchasing clerks' salaries:

```
SELECT last_name, salary, RATIO_TO_REPORT(salary) OVER () AS rr
  FROM employees
  WHERE job_id = 'PU_CLERK'
  ORDER BY last_name, salary, rr;
```

LAST_NAME	SALARY	RR
Baida	2900	.208633094
Colmenares	2500	.179856115
Himuro	2600	.18705036
Khoo	3100	.223021583
Tobias	2800	.201438849

RAWTOHEX

Syntax



Purpose

RAWTOHEX converts raw to a character value containing its hexadecimal representation.

As a SQL built-in function, RAWTOHEX accepts an argument of any scalar data type other than LONG, LONG RAW, CLOB, NCLOB, BLOB, or BFILE. If the argument is of a data type other than RAW, then this function converts the argument value, which is represented using some number of data bytes, into a RAW value with the same number of data bytes. The data itself is not modified in any way, but the data type is recast to a RAW data type.

This function returns a VARCHAR2 value with the hexadecimal representation of bytes that make up the value of *raw*. Each byte is represented by two hexadecimal digits.



RAWTOHEX functions differently when used as a PL/SQL built-in function. Refer to *Oracle Database Development Guide* for more information.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of RAWTOHEX

Examples

The following hypothetical example returns the hexadecimal equivalent of a RAW column value:

```
SELECT RAWTOHEX(raw_column) "Graphics"
   FROM graphics;

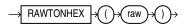
Graphics
-----7D
```



"RAW and LONG RAW Data Types" and HEXTORAW

RAWTONHEX

Syntax





Purpose

RAWTONHEX converts raw to a character value containing its hexadecimal representation. RAWTONHEX (raw) is equivalent to TO_NCHAR (RAWTOHEX (raw)). The value returned is always in the national character set.



RAWTONHEX functions differently when used as a PL/SQL built-in function. Refer to Oracle Database Development Guide for more information.

Examples

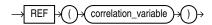
The following hypothetical example returns the hexadecimal equivalent of a RAW column value:

```
SELECT RAWTONHEX(raw_column),
DUMP ( RAWTONHEX (raw_column) ) "DUMP"
FROM graphics;

RAWTONHEX(RA)
DUMP
Typ=1 Len=4: 0,55,0,68
```

REF

Syntax



Purpose

REF takes as its argument a correlation variable (table alias) associated with a row of an object table or an object view. A REF value is returned for the object instance that is bound to the variable or row.

Examples

The sample schema oe contains a type called cust address typ, described as follows:

Attribute	Туре
STREET_ADDRESS	VARCHAR2 (40)
POSTAL_CODE	VARCHAR2(10)
CITY	VARCHAR2(30)
STATE_PROVINCE	VARCHAR2(10)
COUNTRY_ID	CHAR(2)

The following example creates a table based on the sample type $oe.cust_address_typ$, inserts a row into the table, and retrieves a REF value for the object instance of the type in the addresses table:

CREATE TABLE addresses OF cust_address_typ;





Oracle Database Object-Relational Developer's Guide for information on REFs

REFTOHEX

Syntax



Purpose

REFTOHEX converts argument expr to a character value containing its hexadecimal equivalent. expr must return a REF.

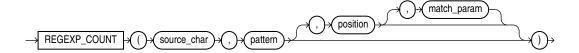
Examples

The sample schema oe contains a warehouse_typ. The following example builds on that type to illustrate how to convert the REF value of a column to a character value containing its hexadecimal equivalent:



REGEXP COUNT

Syntax



Purpose

REGEXP_COUNT complements the functionality of the REGEXP_INSTR function by returning the number of times a pattern occurs in a source string. The function evaluates strings using characters as defined by the input character set. It returns an integer indicating the number of occurrences of pattern. If no match is found, then the function returns 0.

- source_char is a character expression that serves as the search value. It is commonly a character column and can be of any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB.
- pattern is the regular expression. It is usually a text literal and can be of any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. It can contain up to 512 bytes. If the data type of pattern is different from the data type of source_char, then Oracle Database converts pattern to the data type of source char.
 - REGEXP_COUNT ignores subexpression parentheses in *pattern*. For example, the pattern '(123(45))' is equivalent to '12345'. For a listing of the operators you can specify in *pattern*, refer to Oracle Regular Expression Support.
- position is a positive integer indicating the character of source_char where Oracle should begin the search. The default is 1, meaning that Oracle begins the search at the first character of source_char. After finding the first occurrence of pattern, the database searches for a second occurrence beginning with the first character following the first occurrence.
- match_param is a character expression of the data type VARCHAR2 or CHAR that lets you change the default matching behavior of the function.

The value of match param can include one or more of the following characters:

- 'i' specifies case-insensitive matching, even if the determined collation of the condition is case-sensitive.
- 'c' specifies case-sensitive and accent-sensitive matching, even if the determined collation of the condition is case-insensitive or accent-insensitive.
- 'n' allows the period (.), which is the match-any-character character, to match the newline character. If you omit this parameter, then the period does not match the newline character.
- 'm' treats the source string as multiple lines. Oracle interprets the caret (^) and dollar sign (\$) as the start and end, respectively, of any line anywhere in the source string, rather than only at the start or end of the entire source string. If you omit this parameter, then Oracle treats the source string as a single line.
- 'x' ignores whitespace characters. By default, whitespace characters match themselves.



If the value of <code>match_param</code> contains multiple contradictory characters, then Oracle uses the last character. For example, if you specify <code>'ic'</code>, then Oracle uses case-sensitive and accent-sensitive matching. If the value contains a character other than those shown above, then Oracle returns an error.

If you omit match_param, then:

- The default case and accent sensitivity are determined by the determined collation of the REGEXP COUNT function.
- A period (.) does not match the newline character.
- The source string is treated as a single line.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules, which define the collation REGEXP_COUNT uses to compare characters from *source char* with characters from *pattern*

Examples

The following example shows that subexpressions parentheses in pattern are ignored:

In the following example, the function begins to evaluate the source string at the third character, so skips over the first occurrence of pattern:

REGEXP_COUNT simple matching: Examples

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters:

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by a single digit number:

```
select regexp_count('ABC123', '[A-Z][0-9]'), regexp_count('A1B2C3', '[A-Z][0-9]') from
dual;

REGEXP_COUNT('ABC123','[A-Z][0-9]') REGEXP_COUNT('A1B2C3','[A-Z][0-9]')
```



In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by a single digit number only at the beginning of the string:

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by two digits of number only contained within the string:

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by a single digit number within the first two occurrences from the beginning of the string:

Live SQL:

View and run related examples on Oracle Live SQL at *REGEXP_COUNT simple matching*

REGEXP COUNT advanced matching: Examples

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters:

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by a single digit number:

```
select regexp_count('ABC123', '[A-Z][0-9]') Match_string_C1_count,
regexp_count('A1B2C3', '[A-Z][0-9]') Match_strings_A1_B2_C3_count from dual;
```



```
MATCH_STRING_C1_COUNT MATCH_STRINGS_A1_B2_C3_COUNT

1 3
```

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by a single digit number only at the beginning of the string:

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by two digits of number only contained within the string:

In the following example, REGEXP_COUNT validates the supplied string for the given pattern and returns the number of alphabetic letters followed by a single digit number within the first two occurrences from the beginning of the string:

```
select regexp_count('ABC12D3', '([A-Z][0-9]){2}') Char_num_within_2_places,
regexp_count('A1B2C3', '([A-Z][0-9]){2}') Char_num_within_2_places from dual;
CHAR_NUM_WITHIN_2_PLACES CHAR_NUM_WITHIN_2_PLACES
```

Live SQL:

View and run related examples on Oracle Live SQL at REGEXP_COUNT advanced matching

REGEXP_COUNT case-sensitive matching: Examples

The following statements create a table regexp_temp and insert values into it:

```
CREATE TABLE regexp_temp(empName varchar2(20));

INSERT INTO regexp_temp (empName) VALUES ('John Doe');

INSERT INTO regexp temp (empName) VALUES ('Jane Doe');
```

In the following example, the statement queries the employee name column and searches for the lowercase of character 'E':



```
John Doe 1
Jane Doe 2
```

In the following example, the statement queries the employee name column and searches for the lowercase of character 'O':

In the following example, the statement queries the employee name column and searches for the lowercase or uppercase of character 'E':

In the following example, the statement queries the employee name column and searches for the lowercase of string 'DO':

In the following example, the statement queries the employee name column and searches for the lowercase or uppercase of string 'AN':

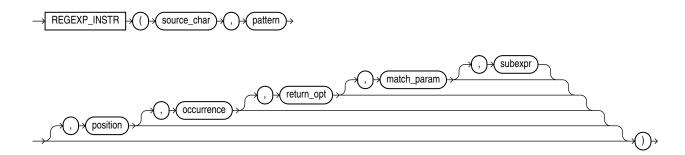
Live SQL:

View and run related examples on Oracle Live SQL at *REGEXP_COUNT* casesensitive matching



REGEXP_INSTR

Syntax



Purpose

REGEXP_INSTR extends the functionality of the INSTR function by letting you search a string for a regular expression pattern. The function evaluates strings using characters as defined by the input character set. It returns an integer indicating the beginning or ending position of the matched substring, depending on the value of the <code>return_option</code> argument. If no match is found, then the function returns 0.

This function complies with the POSIX regular expression standard and the Unicode Regular Expression Guidelines. For more information, refer to Oracle Regular Expression Support.

- source_char is a character expression that serves as the search value. It is commonly a character column and can be of any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB.
- pattern is the regular expression. It is usually a text literal and can be of any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. It can contain up to 512 bytes. If the data type of pattern is different from the data type of source_char, then Oracle Database converts pattern to the data type of source_char. For a listing of the operators you can specify in pattern, refer to Oracle Regular Expression Support.
- position is a positive integer indicating the character of source_char where Oracle should begin the search. The default is 1, meaning that Oracle begins the search at the first character of source_char.
- occurrence is a positive integer indicating which occurrence of pattern in source_char Oracle should search for. The default is 1, meaning that Oracle searches for the first occurrence of pattern. If occurrence is greater than 1, then the database searches for the second occurrence beginning with the first character following the first occurrence of pattern, and so forth. This behavior is different from the INSTR function, which begins its search for the second occurrence at the second character of the first occurrence.
- return option lets you specify what Oracle should return in relation to the occurrence:
 - If you specify 0, then Oracle returns the position of the first character of the occurrence. This is the default.
 - If you specify 1, then Oracle returns the position of the character following the occurrence.
- match_param is a character expression of the data type VARCHAR2 or CHAR that lets you change the default matching behavior of the function. The behavior of this parameter is the



same for this function as for <code>REGEXP_COUNT</code>. Refer to <code>REGEXP_COUNT</code> for detailed information.

• For a pattern with subexpressions, subexpr is an integer from 0 to 9 indicating which subexpression in pattern is the target of the function. The subexpr is a fragment of pattern enclosed in parentheses. Subexpressions can be nested. Subexpressions are numbered in order in which their left parentheses appear in pattern. For example, consider the following expression:

```
0123(((abc)(de)f)ghi)45(678)
```

This expression has five subexpressions in the following order: "abcdefghi" followed by "abcdef", "abc", "de" and "678".

If <code>subexpr</code> is zero, then the position of the entire substring that matches the <code>pattern</code> is returned. If <code>subexpr</code> is greater than zero, then the position of the substring fragment that corresponds to subexpression number <code>subexpr</code> in the matched substring is returned. If <code>pattern</code> does not have at least <code>subexpr</code> subexpressions, the function returns zero. A null <code>subexpr</code> value returns <code>NULL</code>. The default value for <code>subexpr</code> is zero.

See Also:

- INSTR and REGEXP_SUBSTR
- REGEXP_REPLACE and REGEXP_LIKE Condition
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation REGEXP_INSTR uses to compare characters from source char with characters from pattern

Examples

The following example examines the string, looking for occurrences of one or more non-blank characters. Oracle begins searching at the first character in the string and returns the starting position (default) of the sixth occurrence of one or more non-blank characters.

The following example examines the string, looking for occurrences of words beginning with s, r, or p, regardless of case, followed by any six alphabetic characters. Oracle begins searching at the third character in the string and returns the position in the string of the character following the second occurrence of a seven-letter word beginning with s, r, or p, regardless of case.



```
______
28
```

The following examples use the *subexpr* argument to search for a particular subexpression in *pattern*. The first statement returns the position in the source string of the first character in the first subexpression, which is '123':

The next statement returns the position in the source string of the first character in the second subexpression, which is '45678':

The next statement returns the position in the source string of the first character in the fourth subexpression, which is '78':

REGEXP_INSTR pattern matching: Examples

The following statements create a table regexp_temp and insert values into it:

```
CREATE TABLE regexp_temp(empName varchar2(20), emailID varchar2(20));

INSERT INTO regexp_temp (empName, emailID) VALUES ('John Doe', 'johndoe@example.com');

INSERT INTO regexp temp (empName, emailID) VALUES ('Jane Doe', 'janedoe');
```

In the following example, the statement queries the email column and searches for valid email addresses:

In the following example, the statement queries the email column and returns the count of valid email addresses:

```
EMPNAME Valid Email FIELD_WITH_VALID_EMAIL

John Doe johndoe@example.com 1

Jane Doe
```

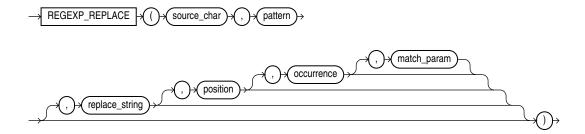


Live SQL:

View and run related examples on Oracle Live SQL at *REGEXP_INSTR* pattern matching

REGEXP REPLACE

Syntax



Purpose

REGEXP_REPLACE extends the functionality of the REPLACE function by letting you search a string for a regular expression pattern. By default, the function returns <code>source_char</code> with every occurrence of the regular expression pattern replaced with <code>replace_string</code>. The string returned is in the same character set as <code>source_char</code>. The function returns <code>VARCHAR2</code> if the first argument is not a LOB and returns <code>CLOB</code> if the first argument is a LOB.

This function complies with the POSIX regular expression standard and the Unicode Regular Expression Guidelines. For more information, refer to Oracle Regular Expression Support.

- source_char is a character expression that serves as the search value. It is commonly a
 character column and can be of any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2,
 CLOB or NCLOB.
- pattern is the regular expression. It is usually a text literal and can be of any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. It can contain up to 512 bytes. If the data type of pattern is different from the data type of source_char, then Oracle Database converts pattern to the data type of source_char. For a listing of the operators you can specify in pattern, refer to Oracle Regular Expression Support.
- replace_string can be of any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. If replace_string is a CLOB or NCLOB, then Oracle truncates replace_string to 32K. The replace_string can contain up to 500 backreferences to subexpressions in the form \n, where n is a number from 1 to 9. If you want to include a backslash (\) in replace_string, then you must precede it with the escape character, which is also a backslash. For example, to replace \2 you would enter \\2. For more information on backreference expressions, refer to the notes to "Oracle Regular Expression Support", Table D-1.
- position is a positive integer indicating the character of source_char where Oracle should begin the search. The default is 1, meaning that Oracle begins the search at the first character of source_char.



- occurrence is a nonnegative integer indicating the occurrence of the replace operation:
 - If you specify 0, then Oracle replaces all occurrences of the match.
 - If you specify a positive integer *n*, then Oracle replaces the *n*th occurrence.

If occurrence is greater than 1, then the database searches for the second occurrence beginning with the first character following the first occurrence of pattern, and so forth. This behavior is different from the INSTR function, which begins its search for the second occurrence at the second character of the first occurrence.

match_param is a character expression of the data type VARCHAR2 or CHAR that lets you change the default matching behavior of the function. The behavior of this parameter is the same for this function as for REGEXP_COUNT. Refer to REGEXP_COUNT for detailed information.

See Also:

- REPLACE
- REGEXP_INSTR, REGEXP_SUBSTR, and REGEXP_LIKE Condition
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation REGEXP_REPLACE uses to compare characters from source_char with characters from pattern, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following example examines phone_number, looking for the pattern xxx.xxx.xxxx. Oracle reformats this pattern with (xxx) xxx-xxxx.

The following example examines country_name. Oracle puts a space after each non-null character in the string.



```
Australia
Belgium
Brazil
Canada
```

The following example examines the string, looking for two or more spaces. Oracle replaces each occurrence of two or more spaces with a single space.

REGEXP_REPLACE pattern matching: Examples

The following statements create a table regexp_temp and insert values into it:

```
CREATE TABLE regexp_temp(empName varchar2(20), emailID varchar2(20));

INSERT INTO regexp_temp (empName, emailID) VALUES ('John Doe', 'johndoe@example.com');

INSERT INTO regexp temp (empName, emailID) VALUES ('Jane Doe', 'janedoe@example.com');
```

The following statement replaces the string 'Jane' with 'John':

```
SELECT empName, REGEXP_REPLACE (empName, 'Jane', 'John') "STRING_REPLACE" FROM regexp_temp;

EMPNAME STRING_REPLACE

John Doe John Doe
Jane Doe John Doe
```

The following statement replaces the string 'John' with 'Jane':

```
Live SQL:
```

View and run a related example on Oracle Live SQL at REGEXP_REPLACE - Pattern Matching

REGEXP_REPLACE: Examples

The following statement replaces all the numbers in a string:

```
WITH strings AS (
SELECT 'abc123' s FROM dual union all
SELECT '123abc' s FROM dual union all
```



```
SELECT 'a1b2c3' s FROM dual
)

SELECT s "STRING", regexp_replace(s, '[0-9]', '') "MODIFIED_STRING"
FROM strings;

STRING MODIFIED_STRING

abc123 abc
abc
alb2c3 abc
```

The following statement replaces the first numeric occurrence in a string:

```
WITH strings AS (
SELECT 'abc123' s from DUAL union all
SELECT '123abc' s from DUAL union all
SELECT 'a1b2c3' s from DUAL
)

SELECT s "STRING", REGEXP_REPLACE(s, '[0-9]', '', 1, 1) "MODIFIED_STRING"
FROM strings;

STRING MODIFIED_STRING

abc123 abc23
123abc 23abc
a1b2c3 ab2c3
```

The following statement replaces the second numeric occurrence in a string:

The following statement replaces multiple spaces in a string with a single space:

```
WITH strings AS (
   SELECT 'Hello World' s FROM dual union all
   SELECT 'Hello World' s FROM dual union all
   SELECT 'Hello, World !' s FROM dual
)
   SELECT s "STRING", regexp_replace(s, ' {2,}', ' ') "MODIFIED_STRING"
   FROM strings;
STRING MODIFIED_STRING
```



```
Hello World Hello World
Hello World Hello World
Hello, World ! Hello, World !
```

The following statement converts camel case strings to a string containing lower case words separated by an underscore:

```
WITH strings as (
 SELECT 'AddressLine1' s FROM dual union all
 SELECT 'ZipCode' s FROM dual union all
 SELECT 'Country' s FROM dual
)
 SELECT s "STRING",
       lower(regexp replace(s, '([A-Z0-9])', ' \1', 2)) "MODIFIED STRING"
 FROM strings;
                  MODIFIED STRING
 STRING
_____
AddressLine1 address line 1
ZipCode
                zip_code
Country
                country
```

The following statement converts the format of a date:

```
WITH date strings AS (
  SELECT '2015-01-01' d from dual union all
  SELECT '2000-12-31' d from dual union all
  SELECT '900-01-01' d from dual
)
  SELECT d "STRING",
        regexp replace(d, '([[:digit:]]+)-([[:digit:]]{2})-([[:digit:]]
{2})', '\3.\2.\1') "MODIFIED STRING"
  FROM date strings;
  STRING
                   MODIFIED STRING
2015-01-01 01.01.2015
2000-12-31
                  31.12.2000
900-01-01
                  01.01.900
```

The following statement replaces all the letters in a string with '1':

```
WITH strings as (
   SELECT 'NEW YORK' s FROM dual union all
   SELECT 'New York' s FROM dual union all
   SELECT 'new york' s FROM dual
)

SELECT s "STRING",
        regexp_replace(s, '[a-z]', '1', 1, 0, 'i') "CASE_INSENSITIVE",
        regexp_replace(s, '[a-z]', '1', 1, 0, 'c') "CASE_SENSITIVE",
        regexp_replace(s, '[a-zA-z]', '1', 1, 0, 'c')
"CASE_SENSITIVE_MATCHING"
FROM strings;
```

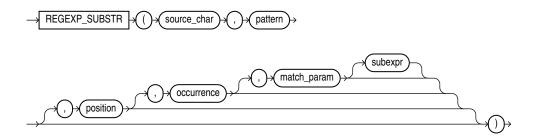
SI	TRING	CZ	ASE_INSE	EN CA	ASE_SENS	SI CA	ASE_SEN	SI
NEW	YORK	111	1111	NEW	YORK	111	1111	
New	York	111	1111	N11	Y111	111	1111	
new	york	111	1111	111	1111	111	1111	



View and run a related example on Oracle Live SQL at REGEXP_REPLACE

REGEXP SUBSTR

Syntax



Purpose

REGEXP_SUBSTR extends the functionality of the SUBSTR function by letting you search a string for a regular expression pattern. It is also similar to REGEXP_INSTR, but instead of returning the position of the substring, it returns the substring itself. This function is useful if you need the contents of a match string but not its position in the source string. The function returns the string as VARCHAR2 or CLOB data in the same character set as <code>source char</code>.

This function complies with the POSIX regular expression standard and the Unicode Regular Expression Guidelines. For more information, refer to Oracle Regular Expression Support.

- source_char is a character expression that serves as the search value. It is commonly a character column and can be of any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB.
- pattern is the regular expression. It is usually a text literal and can be of any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. It can contain up to 512 bytes. If the data type of pattern is different from the data type of source_char, then Oracle Database converts pattern to the data type of source_char. For a listing of the operators you can specify in pattern, refer to Oracle Regular Expression Support.
- position is a positive integer indicating the character of source_char where Oracle should begin the search. The default is 1, meaning that Oracle begins the search at the first character of source_char.
- occurrence is a positive integer indicating which occurrence of pattern in source_char Oracle should search for. The default is 1, meaning that Oracle searches for the first occurrence of pattern.

If occurrence is greater than 1, then the database searches for the second occurrence beginning with the first character following the first occurrence of pattern, and so forth. This behavior is different from the SUBSTR function, which begins its search for the second occurrence at the second character of the first occurrence.

- match_param is a character expression of the data type VARCHAR2 or CHAR that lets you change the default matching behavior of the function. The behavior of this parameter is the same for this function as for REGEXP_COUNT. Refer to REGEXP_COUNT for detailed information.
- For a pattern with subexpressions, subexpr is a nonnegative integer from 0 to 9 indicating
 which subexpression in pattern is to be returned by the function. This parameter has the
 same semantics that it has for the REGEXP_INSTR function. Refer to REGEXP_INSTR for
 more information.

See Also:

- SUBSTR and REGEXP_INSTR
- REGEXP_REPLACE, and REGEXP_LIKE Condition
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation REGEXP_SUBSTR uses to compare characters from source_char with characters from pattern, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following example examines the string, looking for the first substring bounded by commas. Oracle Database searches for a comma followed by one or more occurrences of non-comma characters followed by a comma. Oracle returns the substring, including the leading and trailing commas.

The following example examines the string, looking for http:// followed by a substring of one or more alphanumeric characters and optionally, a period (.). Oracle searches for a minimum of three and a maximum of four occurrences of this substring between http:// and either a slash (/) or the end of the string.



The next two examples use the *subexpr* argument to return a specific subexpression of *pattern*. The first statement returns the first subexpression in *pattern*:

The next statement returns the fourth subexpression in pattern:

REGEXP_SUBSTR pattern matching: Examples

The following statements create a table regexp_temp and insert values into it:

```
CREATE TABLE regexp_temp(empName varchar2(20), emailID varchar2(20));

INSERT INTO regexp_temp (empName, emailID) VALUES ('John Doe', 'johndoe@example.com');
INSERT INTO regexp_temp (empName, emailID) VALUES ('Jane Doe', 'janedoe');
```

In the following example, the statement queries the email column and searches for valid email addresses:

```
SELECT empName, REGEXP_SUBSTR(emailID, '[[:alnum:]]+\@[[:alnum:]]+\.[[:alnum:]]+')
"Valid Email" FROM regexp_temp;

EMPNAME Valid Email
_______
John Doe johndoe@example.com
Jane Doe
```

In the following example, the statement queries the email column and returns the count of valid email addresses:

Live SQL:

View and run related examples on Oracle Live SQL at *REGEXP_SUBSTR* pattern matching

In the following example, numbers and alphabets are extracted from a string:

```
with strings as (
 select 'ABC123' str from dual union all
 select 'A1B2C3' str from dual union all
 select '123ABC' str from dual union all
  select '1A2B3C' str from dual
  select regexp substr(str, '[0-9]') First Occurrence of Number,
        regexp_substr(str, '[0-9].*') Num_Followed_by_String,
        regexp substr(str, '[A-Z][0-9]') Letter Followed by String
  from strings;
FIRST_OCCURRENCE_OF_NUMB NUM_FOLLOWED_BY_STRING LETTER_FOLLOWED_BY_STRIN
        123 C1
1B2C3 A1
1
1
1
           123ABC
1
           1A2B3C
                           A2
```


View and run a related example on Oracle Live SQL at REGEXP_SUBSTR - Extract Numbers and Alphabets

In the following example, passenger names and flight information are extracted from a string:

```
with strings as (
 select 'LHRJFK/010315/JOHNDOE' str from dual union all
  select 'CDGLAX/050515/JANEDOE' str from dual union all
 select 'LAXCDG/220515/JOHNDOE' str from dual union all
  select 'SFOJFK/010615/JANEDOE' str from dual
)
  SELECT regexp_substr(str, '[A-Z]{6}') String_of_6_characters,
        regexp_substr(str, '[0-9]+') First_Matching_Numbers,
        regexp_substr(str, '[A-Z].*$') Letter_by_other_characters,
        regexp substr(str, '/[A-Z].*$') Slash letter and characters
 FROM strings;
STRING OF 6 CHARACTERS FIRST MATCHING NUMBERS LETTER BY OTHER CHARACTERS
SLASH LETTER AND CHARACTERS
______
                                               _____
LHRJFK
                       010315
                                               LHRJFK/010315/JOHNDOE
JOHNDOE
                       050515
CDGLAX
                                              CDGLAX/050515/JANEDOE
JANEDOE
                      220515
                                             LAXCDG/220515/JOHNDOE
LAXCDG
JOHNDOE
SFOJFK
                      010615
                                              SFOJFK/010615/JANEDOE
JANEDOE
```

Live SQL:

View and run a related example on Oracle Live SQL at REGEXP_SUBSTR - Extract Passenger Names and Flight Information

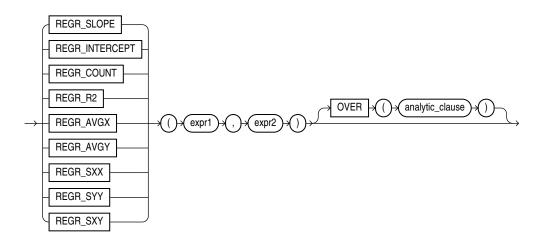
REGR_ (Linear Regression) Functions

The linear regression functions are:

- REGR SLOPE
- REGR INTERCEPT
- REGR COUNT
- REGR R2
- REGR AVGX
- REGR AVGY
- REGR SXX
- REGR SYY
- REGR SXY

Syntax

linear_regr::=



See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

The linear regression functions fit an ordinary-least-squares regression line to a set of number pairs. You can use them as both aggregate and analytic functions.

See Also:

"Aggregate Functions " and "About SQL Expressions " for information on valid forms of expr

These functions take as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.

See Also:

Table 2-9 for more information on implicit conversion and "Numeric Precedence" for information on numeric precedence

Oracle applies the function to the set of (expr1, expr2) pairs after eliminating all pairs for which either expr1 or expr2 is null. Oracle computes all the regression functions simultaneously during a single pass through the data.

expr1 is interpreted as a value of the dependent variable (a y value), and expr2 is interpreted as a value of the independent variable (an x value).

 REGR_SLOPE returns the slope of the line. The return value is a numeric data type and can be null. After the elimination of null (expr1, expr2) pairs, it makes the following computation:

```
COVAR POP(expr1, expr2) / VAR POP(expr2)
```

• REGR_INTERCEPT returns the y-intercept of the regression line. The return value is a numeric data type and can be null. After the elimination of null (expr1, expr2) pairs, it makes the following computation:

```
AVG(expr1) - REGR SLOPE(expr1, expr2) * AVG(expr2)
```

- REGR_COUNT returns an integer that is the number of non-null number pairs used to fit the regression line.
- REGR_R2 returns the coefficient of determination (also called R-squared or goodness of fit)
 for the regression. The return value is a numeric data type and can be null. VAR_POP(expr1)
 and VAR_POP(expr2) are evaluated after the elimination of null pairs. The return values are:

All of the remaining regression functions return a numeric data type and can be null:

• REGR_AVGX evaluates the average of the independent variable (*expr2*) of the regression line. It makes the following computation after the elimination of null (*expr1*, *expr2*) pairs:

```
AVG(expr2)
```



• REGR_AVGY evaluates the average of the dependent variable (expr1) of the regression line. It makes the following computation after the elimination of null (expr1, expr2) pairs:

```
AVG(expr1)
```

REGR_SXY, REGR_SXX, REGR_SYY are auxiliary functions that are used to compute various diagnostic statistics.

- REGR_SXX makes the following computation after the elimination of null (expr1, expr2) pairs:
 REGR_COUNT(expr1, expr2) * VAR POP(expr2)
- REGR_SYY makes the following computation after the elimination of null (expr1, expr2) pairs:
 REGR_COUNT(expr1, expr2) * VAR POP(expr1)
- REGR_SXY makes the following computation after the elimination of null (expr1, expr2) pairs:
 REGR_COUNT(expr1, expr2) * COVAR_POP(expr1, expr2)

The following examples are based on the sample tables sh.sales and sh.products.

General Linear Regression Example

The following example provides a comparison of the various linear regression functions used in their analytic form. The analytic form of these functions can be useful when you want to use regression statistics for calculations such as finding the salary predicted for each employee by the model. The sections that follow on the individual linear regression functions contain examples of the aggregate form of these functions.

```
SELECT job id, employee id ID, salary,
REGR SLOPE (SYSDATE-hire date, salary)
  OVER (PARTITION BY job id) slope,
REGR INTERCEPT (SYSDATE-hire date, salary)
  OVER (PARTITION BY job id) intcpt,
REGR R2 (SYSDATE-hire date, salary)
  OVER (PARTITION BY job id) rsqr,
REGR COUNT (SYSDATE-hire date, salary)
  OVER (PARTITION BY job id) count,
REGR AVGX (SYSDATE-hire date, salary)
  OVER (PARTITION BY job id) avgx,
REGR AVGY (SYSDATE-hire date, salary)
  OVER (PARTITION BY job id) avgy
   FROM employees
   WHERE department_id in (50, 80)
  ORDER BY job id, employee id;
```

JOB_ID	ID	SALARY	SLOPE	INTCPT	RSQR	COUNT	AVGX	AVGY
SA MAN	145	14000	.355	-1707.035	.832	5	12200.000	2626.589
SA MAN	146	13500	.355	-1707.035	.832	5	12200.000	2626.589
SA MAN	147	12000	.355	-1707.035	.832	5	12200.000	2626.589
SA MAN	148	11000	.355	-1707.035	.832	5	12200.000	2626.589
SA_MAN	149	10500	.355	-1707.035	.832	5	12200.000	2626.589
SA REP	150	10000	.257	404.763	.647	29	8396.552	2561.244
SA REP	151	9500	.257	404.763	.647	29	8396.552	2561.244
SA REP	152	9000	.257	404.763	.647	29	8396.552	2561.244
SA REP	153	8000	.257	404.763	.647	29	8396.552	2561.244
SA REP	154	7500	.257	404.763	.647	29	8396.552	2561.244
SA REP	155	7000	.257	404.763	.647	29	8396.552	2561.244
SA_REP	156	10000	.257	404.763	.647	29	8396.552	2561.244



REGR_SLOPE and REGR_INTERCEPT Examples

The following example calculates the slope and regression of the linear regression model for time employed (SYSDATE - hire_date) and salary using the sample table hr.employees. Results are grouped by job id.

REGR_COUNT Examples

The following example calculates the count of by job_id for time employed (SYSDATE - hire date) and salary using the sample table hr.employees. Results are grouped by job id.

REGR_R2 Examples

The following example calculates the coefficient of determination the linear regression of time employed (SYSDATE - hire_date) and salary using the sample table hr.employees:



REGR_AVGY and REGR_AVGX Examples

The following example calculates the average values for time employed (SYSDATE - hire_date) and salary using the sample table hr.employees. Results are grouped by job id:

REGR_SXY, REGR_SXX, and REGR_SYY Examples

The following example calculates three types of diagnostic statistics for the linear regression of time employed (SYSDATE - hire date) and salary using the sample table hr.employees:

REMAINDER

Syntax



Purpose

REMAINDER returns the remainder of *n2* divided by *n1*.

This function takes as arguments any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. Oracle determines the argument with the highest numeric precedence, implicitly converts the remaining arguments to that data type, and returns that data type.

The MOD function is similar to REMAINDER except that it uses FLOOR in its formula, whereas REMAINDER uses ROUND. Refer to MOD.



Table 2-9 for more information on implicit conversion and "Numeric Precedence" for information on numeric precedence

- If n1 = 0 or n2 = infinity, then Oracle returns
 - An error if the arguments are of type NUMBER
 - Nan if the arguments are BINARY FLOAT or BINARY DOUBLE.
- If n1 != 0, then the remainder is n2 (n1*N) where N is the integer nearest n2/n1. If n2/n1 equals x.5, then N is the nearest *even* integer.
- If n2 is a floating-point number, and if the remainder is 0, then the sign of the remainder is the sign of n2. Remainders of 0 are unsigned for NUMBER values.

Examples

Using table float_point_demo, created for the TO_BINARY_DOUBLE "Examples", the following example divides two floating-point numbers and returns the remainder of that operation:

REPLACE

Syntax



Purpose

REPLACE returns char with every occurrence of search_string replaced with replacement_string. If replacement_string is omitted or null, then all occurrences of search_string are removed. If search_string is null, then char is returned.

Both <code>search_string</code> and <code>replacement_string</code>, as well as <code>char</code>, can be any of the data types <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, <code>NVARCHAR2</code>, <code>CLOB</code>, or <code>NCLOB</code>. The string returned is in the same character set as <code>char</code>. The function returns <code>VARCHAR2</code> if the first argument is not a LOB and returns <code>CLOB</code> if the first argument is a LOB.

REPLACE provides functionality related to that provided by the TRANSLATE function. TRANSLATE provides single-character, one-to-one substitution. REPLACE lets you substitute one string for another as well as to remove character strings.

See Also:

- TRANSLATE
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation REPLACE uses to compare characters from char with characters from search_string, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following example replaces occurrences of J with BL:

```
SELECT REPLACE('JACK and JUE','J','BL') "Changes"
FROM DUAL;

Changes
-----
BLACK and BLUE
```

ROUND (datetime)

Syntax

round_datetime::=



Purpose

ROUND returns datetime rounded to the unit specified by the format model fmt.

This function is not sensitive to the NLS_CALENDAR session parameter. It operates according to the rules of the Gregorian calendar. The value returned is always of data type DATE, even if you specify a different datetime data type for date. If you omit fmt, then date is rounded to the nearest day. The date expression must resolve to a DATE value.



"CEIL, FLOOR, ROUND, and TRUNC Date Functions" for the permitted format models to use in fmt

Examples

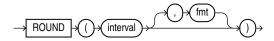
The following example rounds a date to the first day of the following year:

```
SELECT ROUND (TO_DATE ('27-OCT-00'),'YEAR')
"New Year" FROM DUAL;
```

```
New Year
-----01-JAN-01
```

ROUND (interval)

Syntax



Purpose

ROUND (interval) returns the interval rounded up to the unit specified by the second argument fmt, the format model .

For INTERVAL YEAR TO MONTH, fmt can only be year. The default fmt is year.

For INTERVAL DAY TO SECOND, fmt can be day, hour and minute. The default fmt is day. Note that fmt does not support second.

ROUND (interval) rounds up on the mid value of next part of fmt as follows:

- If fmt is year, ROUND (interval) rounds up on the mid value of month which is 6.
- If fmt is day, ROUND (interval) rounds up on the mid value of hour which is 12.
- If fmt is hour, ROUND (interval) rounds up on the mid value of minute which is 30.
- If fmt is minute, ROUND (interval) rounds up on the mid value of second which is 30.

The result precision for year and day is the input precision for year plus one and day plus one respectively, since ROUND (interval) can have overflow. If an interval already has the maximum precision for year and day, the statement compiles but errors at runtime.

See Also:

Refer to CEIL, FLOOR, ROUND, and TRUNC Date Functions for the permitted format models to use in fmt.

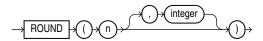
Examples



ROUND (number)

Syntax

round number::=



Purpose

ROUND returns n rounded to integer places to the right of the decimal point. If you omit integer, then n is rounded to zero places. If integer is negative, then n is rounded off to the left of the decimal point.

n can be any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If you omit integer, then the function returns the value ROUND(n, 0) in the same data type as the numeric data type of n. If you include integer, then the function returns NUMBER.

ROUND is implemented using the following rules:

- **1.** If *n* is **0**, then ROUND always returns **0** regardless of *integer*.
- 2. If *n* is negative, then ROUND(n, integer) returns -ROUND(-n, integer).
- 3. If *n* is positive, then

```
ROUND(n, integer) = FLOOR(n * POWER(10, integer) + 0.5) * POWER(10, -integer)
```

ROUND applied to a NUMBER value may give a slightly different result from ROUND applied to the same value expressed in floating-point. The different results arise from differences in internal representations of NUMBER and floating point values. The difference will be 1 in the rounded digit if a difference occurs.

See Also:

- Table 2-9 for more information on implicit conversion
- "Floating-Point Numbers" for more information on how Oracle Database handles BINARY_FLOAT and BINARY_DOUBLE values
- FLOOR (number) and CEIL (number), TRUNC (number) and MOD for information on functions that perform related operations

Examples

The following example rounds a number to one decimal point:

```
Round

Round

15.2
```

The following example rounds a number one digit to the left of the decimal point:

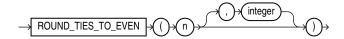
```
SELECT ROUND(15.193,-1) "Round" FROM DUAL;

Round
-----
20
```

ROUND_TIES_TO_EVEN (number)

Syntax

round_ties_to_even ::=



Purpose

ROUND_TIES_TO_EVEN is a rounding function that takes two parameters: n and integer. The function returns n rounded to integer places according to the following rules:

- 1. If integer is positive, n is rounded to integer places to the right of the decimal point.
- 2. If integer is not specified, then n is rounded to 0 places.
- 3. If integer is negative, then n is rounded to integer places to the left of the decimal point.

Restrictions

The function does not support the following types: BINARY FLOAT and BINARY DOUBLE.

Examples

The following example rounds a number to one decimal point to the right:

The following example rounds a number to one decimal point to the left:

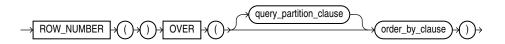
```
SELECT ROUND_TIES_TO_EVEN(45.177,-1) "ROUND_EVEN" FROM DUAL;
```

```
ROUND_TIES_TO_EVEN(45.177,-1)
```

50

ROW NUMBER

Syntax



See Also:

"Analytic Functions " for information on syntax, semantics, and restrictions

Purpose

ROW_NUMBER is an analytic function. It assigns a unique number to each row to which it is applied (either each row in the partition or each row returned by the query), in the ordered sequence of rows specified in the order by clause, beginning with 1.

By nesting a subquery using ROW_NUMBER inside a query that retrieves the ROW_NUMBER values for a specified range, you can find a precise subset of rows from the results of the inner query. This use of the function lets you implement top-N, bottom-N, and inner-N reporting. For consistent results, the query must ensure a deterministic sort order.

Examples

The following example finds the three highest paid employees in each department in the hr.employees table. Fewer than three rows are returned for departments with fewer than three employees.

```
SELECT department_id, first_name, last_name, salary
FROM
(
    SELECT
        department_id, first_name, last_name, salary,
        ROW_NUMBER() OVER (PARTITION BY department_id ORDER BY salary desc) rn
    FROM employees
)
WHERE rn <= 3
ORDER BY department id, salary DESC, last name;</pre>
```

The following example is a join query on the sh.sales table. It finds the sales amounts in 2000 of the five top-selling products in 1999 and compares the difference between 2000 and 1999. The ten top-selling products are calculated within each distribution channel.



```
FROM
    SELECT channel_desc, prod_name, sum(amount_sold) amt,
      ROW NUMBER () OVER (PARTITION BY channel desc
                        ORDER BY SUM(amount sold) DESC) rn
    FROM sales, times, channels, products
    WHERE sales.time id = times.time id
      AND times.calendar year = 1999
      AND channels.channel_id = sales.channel_id
      AND products.prod id = sales.prod id
    GROUP BY channel desc, prod name
  WHERE rn <= 5
  ) top 5 prods 1999 year,
/* The next subquery finds sales per product and per channel in 2000. */
  (SELECT channel desc, prod name, sum(amount sold) amt
    FROM sales, times, channels, products
    WHERE sales.time id = times.time id
      AND times.calendar year = 2000
      AND channels.channel id = sales.channel id
      AND products.prod id = sales.prod id
    GROUP BY channel desc, prod name
  ) sales 2000
WHERE sales_2000.channel_desc = top_5_prods_1999_year.channel_desc
 AND sales_2000.prod_name = top_5_prods_1999_year.prod_name
ORDER BY sales 2000.channel desc, sales 2000.prod name
                                                            AMT 2000 AMT 1999 AMT DIFF
CHANNEL DESC
              PROD NAME
Direct Sales 17" LCD w/built-in HDTV Tuner
                                                             628855.7 1163645.78 -534790.08
Direct Sales
              Envoy 256MB - 40GB
                                                            502938.54 843377.88 -340439.34
                                                           2259566.96 1770349.25 489217.71
Direct Sales Envoy Ambassador
Direct Sales Home Theatre Package with DVD-Audio/Video Play 1235674.15 1260791.44 -25117.29
Direct Sales Mini DV Camcorder with 3.5" Swivel LCD 775851.87 1326302.51 -550450.64
Internet 17" LCD w/built-in HDTV Tuner
                                                           31707.48 160974.7 -129267.22
Internet
                                                           404090.32 155235.25 248855.07
             8.3 Minitower Speaker
                                                            28293.87 154072.02 -125778.15
Internet
             Envoy 256MB - 40GB
             Home Theatre Package with DVD-Audio/Video Play 155405.54 153175.04
Internet
Internet
             Mini DV Camcorder with 3.5" Swivel LCD
                                                           39726.23 189921.97 -150195.74
Partners
             17" LCD w/built-in HDTV Tuner
                                                           269973.97 325504.75 -55530.78
             Envoy Ambassador
                                                          1213063.59 614857.93 598205.66
Partners
Partners
             Home Theatre Package with DVD-Audio/Video Play 700266.58 520166.26 180100.32
             Mini DV Camcorder with 3.5" Swivel LCD 404265.85 520544.11 -116278.26
Partners
                                                           374002.51 340123.02 33879.49
Partners
              Unix/Windows 1-user pack
```

15 rows selected.

ROWIDTOCHAR

Syntax



Purpose

ROWIDTOCHAR converts a rowid value to VARCHAR2 data type. The result of this conversion is always 18 characters long.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of ROWIDTOCHAR

Examples

The following example converts a rowid value in the employees table to a character value. (Results vary for each build of the sample database.)

ROWIDTONCHAR

Syntax



Purpose

ROWIDTONCHAR converts a rowid value to NVARCHAR2 data type. The result of this conversion is always in the national character set and is 18 characters long.



Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of ${\tt ROWIDTONCHAR}$

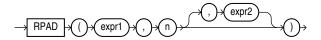
Examples

The following example converts a rowid value to an NVARCHAR2 string:



RPAD

Syntax



Purpose

RPAD returns expr1, right-padded to length n characters with expr2, replicated as many times as necessary. This function is useful for formatting the output of a query.

Both expr1 and expr2 can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The string returned is of VARCHAR2 data type if expr1 is a character data type, NVARCHAR2 if expr1 is a national character data type, and a LOB if expr1 is a LOB data type. The string returned is in the same character set as expr1. The argument n must be a NUMBER integer or a value that can be implicitly converted to a NUMBER integer.

expr1 cannot be null. If you do not specify expr2, then it defaults to a single blank. If expr1 is longer than n, then this function returns the portion of expr1 that fits in n.

The argument n is the total length of the return value as it is displayed on your terminal screen. In most character sets, this is also the number of characters in the return value. However, in some multibyte character sets, the display length of a character string can differ from the number of characters in the string.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of RPAD

Examples

The following example creates a simple chart of salary amounts by padding a single space with asterisks:

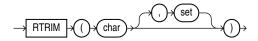
```
SELECT last name, RPAD(' ', salary/1000/1, '*') "Salary"
  FROM employees
  WHERE department id = 80
  ORDER BY last name, "Salary";
LAST NAME
                         Salary
                           *****
Abel
Ande
Banda
Bates
Bernstein
Bloom
Cambrault
Cambrault
Doran
```



Errazuriz	******
Fox	*****
Greene	*****
Hall	*****
Hutton	*****
Johnson	****
King	******

RTRIM

Syntax



Purpose

RTRIM removes from the right end of *char* all of the characters that appear in *set*. This function is useful for formatting the output of a query.

If you do not specify set, then it defaults to a single blank. RTRIM works similarly to LTRIM.

Both char and set can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The string returned is of VARCHAR2 data type if char is a character data type, NVARCHAR2 if char is a national character data type, and a LOB if char is a LOB data type.

See Also:

- LTRIM
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation RTRIM uses to compare characters from set with characters from char, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following example trims all the right-most occurrences of less than sign (<), greater than sign (>), and equal sign (=) from a string:



SCN_TO_TIMESTAMP

Syntax



Purpose

SCN_TO_TIMESTAMP takes as an argument a number that evaluates to a system change number (SCN), and returns the approximate timestamp associated with that SCN. The returned value is of TIMESTAMP data type. This function is useful any time you want to know the timestamp associated with an SCN. For example, it can be used in conjunction with the ORA_ROWSCN pseudocolumn to associate a timestamp with the most recent change to a row.

Notes:

- The usual precision of the result value is 3 seconds.
- The association between an SCN and a timestamp when the SCN is generated is remembered by the database for a limited period of time. This period is the maximum of the auto-tuned undo retention period, if the database runs in the Automatic Undo Management mode, and the retention times of all flashback archives in the database, but no less than 120 hours. The time for the association to become obsolete elapses only when the database is open. An error is returned if the SCN specified for the argument to SCN_TO_TIMESTAMP is too old.

See Also:

ORA_ROWSCN Pseudocolumn and TIMESTAMP_TO_SCN

Examples

The following example uses the <code>ORA_ROWSCN</code> pseudocolumn to determine the system change number of the last update to a row and uses $SCN_TO_TIMESTAMP$ to convert that SCN to a timestamp:

```
SELECT SCN_TO_TIMESTAMP(ORA_ROWSCN) FROM employees
WHERE employee id = 188;
```

You could use such a query to convert a system change number to a timestamp for use in an Oracle Flashback Query:

```
SELECT salary FROM employees WHERE employee_id = 188;
     SALARY
------
3800

UPDATE employees SET salary = salary*10 WHERE employee_id = 188;
```



SESSIONTIMEZONE

Syntax

 \rightarrow SESSIONTIMEZONE \rightarrow

Purpose

SESSIONTIMEZONE returns the time zone of the current session. The return type is a time zone offset (a character type in the format '[+|-]TZH:TZM') or a time zone region name, depending on how the user specified the session time zone value in the most recent ALTER SESSION statement.

Note:

The default client session time zone is an offset even if the client operating system uses a named time zone. If you want the default session time zone to use a named time zone, then set the <code>ORA_SDTZ</code> variable in the client environment to an Oracle time zone region name. Refer to *Oracle Database Globalization Support Guide* for more information on this variable.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of SESSIONTIMEZONE

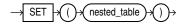
Examples

The following example returns the time zone of the current session:

```
SELECT SESSIONTIMEZONE FROM DUAL;
SESSION
-----
-08:00
```

SET

Syntax



Purpose

SET converts a nested table into a set by eliminating duplicates. The function returns a nested table whose elements are distinct from one another. The returned nested table is of the same type as the input nested table.

The element types of the nested table must be comparable. Refer to "Comparison Conditions" for information on the comparability of nonscalar types.

Examples

The following example selects from the <code>customers_demo</code> table the unique elements of the <code>cust address ntab</code> nested table column:

```
SELECT customer_id, SET(cust_address_ntab) address
FROM customers_demo
ORDER BY customer_id;

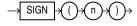
CUSTOMER_ID ADDRESS(STREET_ADDRESS, POSTAL_CODE, CITY, STATE_PROVINCE, COUNTRY_ID)

101 CUST_ADDRESS_TAB_TYP(CUST_ADDRESS_TYP('514 W Superior St', '46901', 'Kokomo', 'IN', 'US'))
102 CUST_ADDRESS_TAB_TYP(CUST_ADDRESS_TYP('2515 Bloyd Ave', '46218', 'Indianapolis', 'IN', 'US'))
103 CUST_ADDRESS_TAB_TYP(CUST_ADDRESS_TYP('8768 N State Rd 37', '47404', 'Bloomington', 'IN', 'US'))
104 CUST_ADDRESS_TAB_TYP(CUST_ADDRESS_TYP('6445 Bay Harbor Ln', '46254', 'Indianapolis', 'IN', 'US'))
105 CUST_ADDRESS_TAB_TYP(CUST_ADDRESS_TYP('4019 W 3Rd St', '47404', 'Bloomington', 'IN', 'US'))
```

The preceding example requires the table <code>customers_demo</code> and a nested table column containing data. Refer to "Multiset Operators" to create this table and nested table column.

SIGN

Syntax



Purpose

SIGN returns the sign of n. This function takes as an argument any numeric data type, or any nonnumeric data type that can be implicitly converted to NUMBER, and returns NUMBER.

For value of NUMBER type, the sign is:

- -1 if n<0
- 0 if n=0
- 1 if n>0

For binary floating-point numbers (BINARY_FLOAT and BINARY_DOUBLE), this function returns the sign bit of the number. The sign bit is:

- -1 if n<0
- +1 if n>=0 or n=NaN

Examples

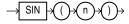
The following example indicates that the argument of the function (-15) is <0:

```
SELECT SIGN(-15) "Sign" FROM DUAL;

Sign
--------------------1
```

SIN

Syntax



Purpose

SIN returns the sine of n (an angle expressed in radians).

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



Table 2-9 for more information on implicit conversion

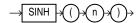
Examples

The following example returns the sine of 30 degrees:



SINH

Syntax



Purpose

SINH returns the hyperbolic sine of *n*.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



Table 2-9 for more information on implicit conversion

Examples

The following example returns the hyperbolic sine of 1:

```
SELECT SINH(1) "Hyperbolic sine of 1" FROM DUAL;

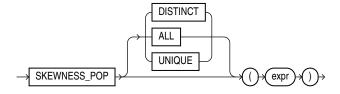
Hyperbolic sine of 1

-----

1.17520119
```

SKEWNESS POP

Syntax



Purpose

SKEWNESS_POP is an aggregate function that is primarily used to determine symmetry in a given distribution.

NULL values in expr are ignored.

Returns NULL if all rows in the group have NULL expr values.

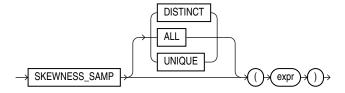
Returns 0 if there are one or two rows in expr.

For a given set of values, the result of population skewness (SKEWNESS_POP) and sample skewness (SKEWNESS_SAMP) are always deterministic. However, the values of SKEWNESS_POP

and <code>SKEWNESS_SAMP</code> differ. As the number of values in the data set increases, the difference between the computed values of <code>SKEWNESS_SAMP</code> and <code>SKEWNESS_POP</code> decreases.

SKEWNESS_SAMP

Syntax



Purpose

SKEWNESS_SAMP is an aggregate function that is primarily used to determine symmetry in a given distribution.

NULL values in expr are ignored.

Returns NULL if all rows in the group have NULL expr values.

Returns 0 if there are one or two rows in expr.

For a given set of values, the result of population skewness (SKEWNESS_POP) and sample skewness (SKEWNESS_SAMP) are always deterministic. However, the values of SKEWNESS_POP and SKEWNESS_SAMP differ. As the number of values in the data set increases, the difference between the computed values of SKEWNESS_SAMP and SKEWNESS_POP decreases.

SOUNDEX

Syntax



Purpose

SOUNDEX returns a character string containing the phonetic representation of *char*. This function lets you compare words that are spelled differently, but sound alike in English.

The phonetic representation is defined in *The Art of Computer Programming*, Volume 3: Sorting and Searching, by Donald E. Knuth, as follows:

- 1. Retain the first letter of the string and remove all other occurrences of the following letters: a, e, h, i, o, u, w, y.
- 2. Assign numbers to the remaining letters (after the first) as follows:

```
b, f, p, v = 1
c, g, j, k, q, s, x, z = 2
d, t = 3
1 = 4
m, n = 5
r = 6
```



- 3. If two or more letters with the same number were adjacent in the original name (before step 1), or adjacent except for any intervening h and w, then retain the first letter and omit rest of all the adjacent letters with same number.
- 4. Return the first four bytes padded with 0.

char can be of any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. The return value is the same data type as char.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.

See Also:

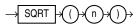
- "Data Type Comparison Rules" for more information.
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of SOUNDEX

Examples

The following example returns the employees whose last names are a phonetic representation of "Smyth":

SQRT

Syntax



Purpose

SQRT returns the square root of *n*.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.



See Also:

Table 2-9 for more information on implicit conversion

- If n resolves to a NUMBER, then the value n cannot be negative. SQRT returns a real number.
- If *n* resolves to a binary floating-point number (BINARY_FLOAT or BINARY_DOUBLE):
 - If $n \ge 0$, then the result is positive.
 - If n = -0, then the result is -0.
 - If n < 0, then the result is NaN.

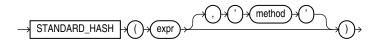
Examples

The following example returns the square root of 26:

```
SELECT SQRT(26) "Square root" FROM DUAL;
Square root
-----
5.09901951
```

STANDARD_HASH

Syntax



Purpose

STANDARD_HASH computes a hash value for a given expression using one of several hash algorithms that are defined and standardized by the National Institute of Standards and Technology. This function is useful for performing authentication and maintaining data integrity in security applications such as digital signatures, checksums, and fingerprinting.

You can use the STANDARD_HASH function to create an index on an extended data type column. Refer to "Creating an Index on an Extended Data Type Column" for more information.

- The *expr* argument determines the data for which you want Oracle Database to compute a hash value. There are no restrictions on the length of data represented by *expr*, which commonly resolves to a column name. The *expr* cannot be a LONG or LOB type. It cannot be a user-defined object type. All other data types are supported for *expr*.
- The optional method argument lets you specify the name of the hash algorithm to be used. Valid algorithms are SHA1, SHA256, SHA384, SHA512 and MD5. If you omit this argument, then SHA1 is used.

The function returns a RAW value.

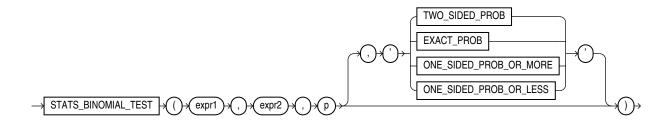




The STANDARD_HASH function is not identical to the one used internally by Oracle Database for hash partitioning.

STATS_BINOMIAL_TEST

Syntax



Purpose

STATS_BINOMIAL_TEST is an exact probability test used for dichotomous variables, where only two possible values exist. It tests the difference between a sample proportion and a given proportion. The sample size in such tests is usually small.

This function takes three required arguments: expr1 is the sample being examined, expr2 contains the values for which the proportion is expected to be, and p is a proportion to test against. The optional fourth argument lets you specify the meaning of the NUMBER value returned by this function, as shown in Table 7-3. For this argument, you can specify a text literal, or a bind variable or expression that evaluates to a constant character value. If you omit the fourth argument, then the default is 'TWO SIDED PROB'.



Appendix C in Oracle Database Globalization Support Guide for the collation determination rules for ${\tt STATS_BINOMIAL_TEST}$

Table 7-3 STATS_BINOMIAL Return Values

Argument	Return Value Meaning
'TWO_SIDED_PROB'	The probability that the given population proportion, p , could result in the observed proportion or a more extreme one.
'EXACT_PROB'	The probability that the given population proportion, p , could result in exactly the observed proportion.
'ONE_SIDED_PROB_OR_MORE'	The probability that the given population proportion, p , could result in the observed proportion or a larger one.
'ONE_SIDED_PROB_OR_LESS'	The probability that the given population proportion, p , could result in the observed proportion or a smaller one.

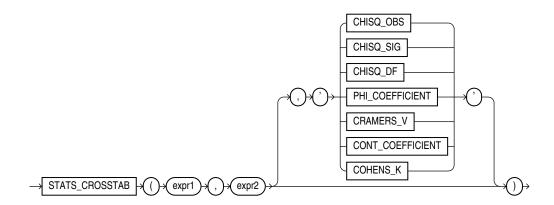
'EXACT_PROB' gives the probability of getting exactly proportion p. In cases where you want to test whether the proportion found in the sample is significantly different from a 50-50 split, p would normally be 0.50. If you want to test only whether the proportion is different, then use the return value 'TWO_SIDED_PROB'. If your test is whether the proportion is more than the value of expr2, then use the return value 'ONE_SIDED_PROB_OR_MORE'. If the test is to determine whether the proportion of expr2 is less, then use the return value 'ONE_SIDED_PROB_OR_MORE'.

STATS_BINOMIAL_TEST Example

The following example determines the probability that reality exactly matches the number of men observed under the assumption that 69% of the population is composed of men:

STATS_CROSSTAB

Syntax



Purpose

Crosstabulation (commonly called crosstab) is a method used to analyze two nominal variables. The <code>STATS_CROSSTAB</code> function takes two required arguments: expr1 and expr2 are the two variables being analyzed. The optional third argument lets you specify the meaning of the <code>NUMBER</code> value returned by this function, as shown in Table 7-4. For this argument, you can specify a text literal, or a bind variable or expression that evaluates to a constant character value. If you omit the third argument, then the default is <code>'CHISQ_SIG'</code>.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules for STATS CROSSTAB

Table 7-4 STATS_CROSSTAB Return Values

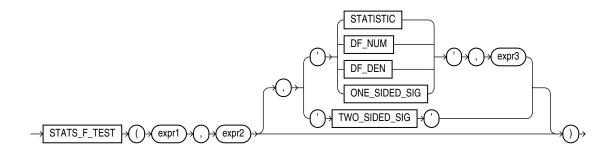
Argument	Return Value Meaning
'CHISQ_OBS'	Observed value of chi-squared
'CHISQ_SIG'	Significance of observed chi-squared
'CHISQ_DF'	Degree of freedom for chi-squared
'PHI_COEFFICIENT'	Phi coefficient
'CRAMERS_V'	Cramer's V statistic
'CONT_COEFFICIENT'	Contingency coefficient
'COHENS_K'	Cohen's kappa

STATS_CROSSTAB Example

The following example determines the strength of the association between gender and income level:

STATS_F_TEST

Syntax



Purpose

STATS_F_TEST tests whether two variances are significantly different. The observed value of f is the ratio of one variance to the other, so values very different from 1 usually indicate significant differences.

This function takes two required arguments: expr1 is the grouping or independent variable and expr2 is the sample of values. The optional third argument lets you specify the meaning of the NUMBER value returned by this function, as shown in Table 7-5. For this argument, you can

specify a text literal, or a bind variable or expression that evaluates to a constant character value. If you omit the third argument, then the default is 'TWO SIDED SIG'.



Appendix C in Oracle Database Globalization Support Guide for the collation determination rules for \mathtt{STATS} F \mathtt{TEST}

Table 7-5 STATS F TEST Return Values

Argument	Return Value Meaning
'STATISTIC'	The observed value of f
'DF_NUM'	Degree of freedom for the numerator
'DF_DEN'	Degree of freedom for the denominator
'ONE_SIDED_SIG'	One-tailed significance of f
'TWO_SIDED_SIG'	Two-tailed significance of f

The one-tailed significance is always in relation to the upper tail. The final argument, expr3, indicates which of the two groups specified by expr1 is the high value or numerator (the value whose rejection region is the upper tail).

The observed value of f is the ratio of the variance of one group to the variance of the second group. The significance of the observed value of f is the probability that the variances are different just by chance—a number between 0 and 1. A small value for the significance indicates that the variances are significantly different. The degree of freedom for each of the variances is the number of observations in the sample minus 1.

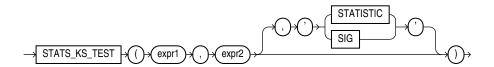
STATS_F_TEST Example

The following example determines whether the variance in credit limit between men and women is significantly different. The results, a p_value not close to zero, and an f_statistic close to 1, indicate that the difference between credit limits for men and women are not significant.



STATS_KS_TEST

Syntax



Purpose

STATS_KS_TEST is a Kolmogorov-Smirnov function that compares two samples to test whether they are from the same population or from populations that have the same distribution. It does not assume that the population from which the samples were taken is normally distributed.

This function takes two required arguments: expx1 classifies the data into the two samples and expx2 contains the values for each of the samples. If expx1 classifies the data into only one sample or into more than two samples, then an error is raised. The optional third argument lets you specify the meaning of the NUMBER value returned by this function, as shown in Table 7-6. For this argument, you can specify a text literal, or a bind variable or expression that evaluates to a constant character value. If you omit the third argument, then the default is 'SIG'.



Appendix C in Oracle Database Globalization Support Guide for the collation determination rules for $\tt STATS KS TEST$

Table 7-6 STATS_KS_TEST Return Values

Argument	Return Value Meaning
'STATISTIC'	Observed value of D
'SIG'	Significance of D

STATS KS TEST Example

Using the Kolmogorov Smirnov test, the following example determines whether the distribution of sales between men and women is due to chance:



STATS_MODE

Syntax



Purpose

STATS_MODE takes as its argument a set of values and returns the value that occurs with the greatest frequency. If more than one mode exists, then Oracle Database chooses one and returns only that one value.

To obtain multiple modes (if multiple modes exist), you must use a combination of other functions, as shown in the hypothetical query:

```
SELECT x FROM (SELECT x, COUNT(x) AS cnt1
  FROM t GROUP BY x)
WHERE cnt1 =
    (SELECT MAX(cnt2) FROM (SELECT COUNT(x) AS cnt2 FROM t GROUP BY x));
```

See Also:

90

100

110

Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules, which define the collation $\mathtt{STATS_MODE}$ uses to compare character values for expr, and for the collation derivation rules, which define the collation assigned to the return value of this function when it is a character value

Examples

The following example returns the mode of salary per department in the hr.employees table:

```
SELECT department id, STATS MODE(salary) FROM employees
  GROUP BY department id
  ORDER BY department id, stats mode(salary);
DEPARTMENT_ID STATS_MODE(SALARY)
_____
         10
         20
                         6000
         30
                         2500
         40
                         6500
         50
                         2500
         60
                         4800
         70
                        10000
         80
                         9500
```

17000

6900

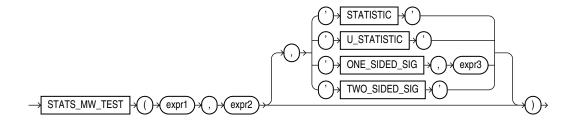
8300 7000

If you need to retrieve all of the modes (in cases with multiple modes), you can do so using a combination of other functions, as shown in the next example:



STATS_MW_TEST

Syntax



Purpose

A Mann Whitney test compares two independent samples to test the null hypothesis that two populations have the same distribution function against the alternative hypothesis that the two distribution functions are different.

The STATS_MW_TEST does not assume that the differences between the samples are normally distributed, as do the STATS_T_TEST_* functions. This function takes two required arguments: expr1 classifies the data into groups and expr2 contains the values for each of the groups. The optional third argument lets you specify the meaning of the NUMBER value returned by this function, as shown in Table 7-7. For this argument, you can specify a text literal, or a bind variable or expression that evaluates to a constant character value. If you omit the third argument, then the default is 'TWO_SIDED_SIG'.



Appendix C in Oracle Database Globalization Support Guide for the collation determination rules for STATS MW TEST

Table 7-7 STATS_MW_TEST Return Values

Argument	Return Value Meaning
'STATISTIC'	The observed value of Z
'U_STATISTIC'	The observed value of <i>U</i>

Table 7-7 (Cont.) STATS_MW_TEST Return Values

Argument	Return Value Meaning
'ONE_SIDED_SIG'	One-tailed significance of Z
'TWO_SIDED_SIG'	Two-tailed significance of Z

The significance of the observed value of Z or U is the probability that the variances are different just by chance—a number between 0 and 1. A small value for the significance indicates that the variances are significantly different. The degree of freedom for each of the variances is the number of observations in the sample minus 1.

The one-tailed significance is always in relation to the upper tail. The final argument, <code>expr3</code>, indicates which of the two groups specified by <code>expr1</code> is the high value (the value whose rejection region is the upper tail).

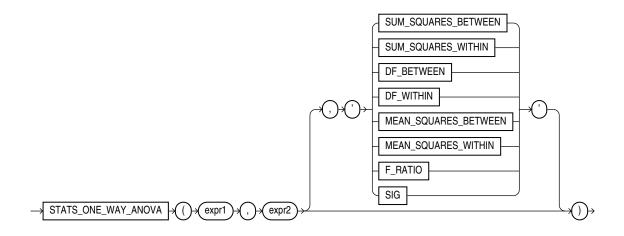
STATS_MW_TEST computes the probability that the samples are from the same distribution by checking the differences in the sums of the ranks of the values. If the samples come from the same distribution, then the sums should be close in value.

STATS_MW_TEST Example

Using the Mann Whitney test, the following example determines whether the distribution of sales between men and women is due to chance:

STATS_ONE_WAY_ANOVA

Syntax



Purpose

The one-way analysis of variance function (STATS_ONE_WAY_ANOVA) tests differences in means (for groups or variables) for statistical significance by comparing two different estimates of variance. One estimate is based on the variances within each group or category. This is known as the **mean squares within** or **mean square error**. The other estimate is based on the variances among the means of the groups. This is known as the **mean squares between**. If the means of the groups are significantly different, then the mean squares between will be larger than expected and will not match the mean squares within. If the mean squares of the groups are consistent, then the two variance estimates will be about the same.

STATS_ONE_WAY_ANOVA takes two required arguments: expr1 is an independent or grouping variable that divides the data into a set of groups and expr2 is a dependent variable (a numeric expression) containing the values corresponding to each member of a group. The optional third argument lets you specify the meaning of the NUMBER value returned by this function, as shown in Table 7-8. For this argument, you can specify a text literal, or a bind variable or expression that evaluates to a constant character value. If you omit the third argument, then the default is 'SIG'.



Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules for STATS ONE WAY_ANOVA

Table 7-8 STATS_ONE_WAY_ANOVA Return Values

Argument	Return Value Meaning
'SUM_SQUARES_BETEEN'	Sum of squares between groups
'SUM_SQUARES_WITHIN'	Sum of squares within groups
'DF_BETWEEN'	Degree of freedom between groups
'DF_WITHIN'	Degree of freedom within groups
'MEAN_SQUARES_BETWEEN'	Mean squares between groups
'MEAN_SQUARES_WITHIN'	Mean squares within groups
'F_RATIO'	Ratio of the mean squares between to the mean squares within (MSB/MSW)
'SIG'	Significance

The significance of one-way analysis of variance is determined by obtaining the one-tailed significance of an f-test on the ratio of the mean squares between and the mean squares within. The f-test should use one-tailed significance, because the mean squares between can be only equal to or larger than the mean squares within. Therefore, the significance returned by STATS_ONE_WAY_ANOVA is the probability that the differences between the groups happened by chance—a number between 0 and 1. The smaller the number, the greater the significance of the difference between the groups. Refer to the STATS_F_TEST for information on performing an f-test.



STATS_ONE_WAY_ANOVA Example

The following example determines the significance of the differences in mean sales within an income level and differences in mean sales between income levels. The results, p_values close to zero, indicate that, for both men and women, the difference in the amount of goods sold across different income levels is significant.

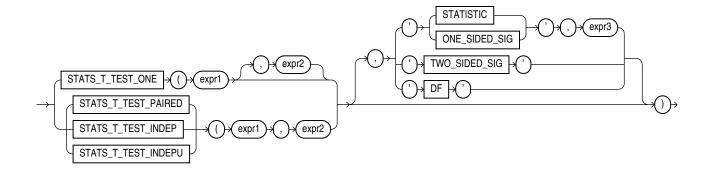
STATS_T_TEST_*

The *t*-test functions are:

- STATS T TEST ONE: A one-sample t-test
- STATS T TEST PAIRED: A two-sample, paired t-test (also known as a crossed t-test)
- STATS_T_TEST_INDEP: A t-test of two independent groups with the same variance (pooled variances)
- STATS_T_TEST_INDEPU: A *t*-test of two independent groups with unequal variance (unpooled variances)

Syntax

stats_t_test::=



Purpose

The t-test measures the significance of a difference of means. You can use it to compare the means of two groups or the means of one group with a constant. Each t-test function takes two expression arguments, although the second expression is optional for the one-sample function (STATS_T_TEST_ONE). Each t-test function takes an optional third argument, which lets you specify the meaning of the NUMBER value returned by the function, as shown in Table 7-9. For

this argument, you can specify a text literal, or a bind variable or expression that evaluates to a constant character value. If you omit the third argument, then the default is 'TWO SIDED SIG'.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules for the $\mathtt{STATS_T_TEST_*}$ functions

Table 7-9 STATS T TEST * Return Values

Argument	Return Value Meaning
'STATISTIC'	The observed value of t
'DF'	Degree of freedom
'ONE_SIDED_SIG'	One-tailed significance of t
'TWO_SIDED_SIG'	Two-tailed significance of t

The two independent STATS_T_TEST_* functions can take a fourth argument (expr3) if the third argument is specified as 'STATISTIC' or 'ONE_SIDED_SIG'. In this case, expr3 indicates which value of expr1 is the high value, or the value whose rejection region is the upper tail.

The significance of the observed value of t is the probability that the value of t would have been obtained by chance—a number between 0 and 1. The smaller the value, the more significant the difference between the means. One-sided significance is always respect to the upper tail. For one-sample and paired t-test, the high value is the first expression. For independent t-test, the high value is the one specified by expr3.

The degree of freedom depends on the type of t-test that resulted in the observed value of t. For example, for a one-sample t-test (STATS_T_TEST_ONE), the degree of freedom is the number of observations in the sample minus 1.

STATS T TEST ONE

In the STATS_T_TEST_ONE function, expr1 is the sample and expr2 is the constant mean against which the sample mean is compared. For this t-test only, expr2 is optional; the constant mean defaults to 0. This function obtains the value of t by dividing the difference between the sample mean and the known mean by the standard error of the mean (rather than the standard error of the difference of the means, as for STATS_T_TEST_PAIRED).

STATS_T_TEST_ONE Example

The following example determines the significance of the difference between the average list price and the constant value 60:



STATS_T_TEST_PAIRED

In the STATS_T_TEST_PAIRED function, expr1 and expr2 are the two samples whose means are being compared. This function obtains the value of t by dividing the difference between the sample means by the standard error of the difference of the means (rather than the standard error of the mean, as for STATS T TEST ONE).

STATS_T_TEST_INDEP and STATS_T_TEST_INDEPU

In the STATS_T_TEST_INDEP and STATS_T_TEST_INDEPU functions, expr1 is the grouping column and expr2 is the sample of values. The pooled variances version (STATS_T_TEST_INDEP) tests whether the means are the same or different for two distributions that have similar variances. The unpooled variances version (STATS_T_TEST_INDEPU) tests whether the means are the same or different even if the two distributions are known to have significantly different variances.

Before using these functions, it is advisable to determine whether the variances of the samples are significantly different. If they are, then the data may come from distributions with different shapes, and the difference of the means may not be very useful. You can perform an f-test to determine the difference of the variances. If they are not significantly different, use STATS_T_TEST_INDEP. If they are significantly different, use STATS_T_TEST_INDEPU. Refer to STATS_F_TEST_for information on performing an f-test.

STATS T TEST INDEP Example

The following example determines the significance of the difference between the average sales to men and women where the distributions are assumed to have similar (pooled) variances:

```
SELECT SUBSTR(cust_income_level, 1, 22) income_level,
    AVG(DECODE(cust_gender, 'M', amount_sold, null)) sold_to_men,
    AVG(DECODE(cust_gender, 'F', amount_sold, null)) sold_to_women,
    STATS_T_TEST_INDEP(cust_gender, amount_sold, 'STATISTIC', 'F') t_observed,
    STATS_T_TEST_INDEP(cust_gender, amount_sold) two_sided_p_value
    FROM sh.customers c, sh.sales s
    WHERE c.cust_id = s.cust_id
    GROUP BY ROLLUP(cust_income_level)
    ORDER BY income level, sold to men, sold to women, t observed;
```

INCOME_LEVEL	SOLD_TO_MEN	SOLD_TO_WOMEN	T_OBSERVED	TWO_SIDED_P_VALUE
A: Below 30,000	105.28349	99.4281447	-1.9880629	.046811482
B: 30,000 - 49,999	102.59651	109.829642	3.04330875	.002341053
C: 50,000 - 69,999	105.627588	110.127931	2.36148671	.018204221
D: 70,000 - 89,999	106.630299	110.47287	2.28496443	.022316997
E: 90,000 - 109,999	103.396741	101.610416	-1.2544577	.209677823
F: 110,000 - 129,999	106.76476	105.981312	60444998	.545545304
G: 130,000 - 149,999	108.877532	107.31377	85298245	.393671218
H: 150,000 - 169,999	110.987258	107.152191	-1.9062363	.056622983
I: 170,000 - 189,999	102.808238	107.43556	2.18477851	.028908566
J: 190,000 - 249,999	108.040564	115.343356	2.58313425	.009794516
K: 250,000 - 299,999	112.377993	108.196097	-1.4107871	.158316973
L: 300,000 and above	120.970235	112.216342	-2.0642868	.039003862
	107.121845	113.80441	.686144393	.492670059
	106.663769	107.276386	1.08013499	.280082357



14 rows selected.

STATS_T_TEST_INDEPU Example

The following example determines the significance of the difference between the average sales to men and women where the distributions are known to have significantly different (unpooled) variances:

```
SELECT SUBSTR(cust_income_level, 1, 22) income_level,

AVG(DECODE(cust_gender, 'M', amount_sold, null)) sold_to_men,

AVG(DECODE(cust_gender, 'F', amount_sold, null)) sold_to_women,

STATS_T_TEST_INDEPU(cust_gender, amount_sold, 'STATISTIC', 'F') t_observed,

STATS_T_TEST_INDEPU(cust_gender, amount_sold) two_sided_p_value

FROM sh.customers c, sh.sales s

WHERE c.cust_id = s.cust_id

GROUP BY ROLLUP(cust_income_level)

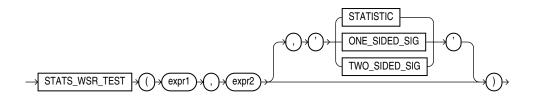
ORDER BY income_level, sold_to_men, sold_to_women, t_observed;
```

INCOME_LEVEL	SOLD_TO_MEN	SOLD_TO_WOMEN	T_OBSERVED	TWO_SIDED_P_VALUE
A: Below 30,000	105.28349	99.4281447	-2.0542592	.039964704
B: 30,000 - 49,999	102.59651	109.829642	2.96922332	.002987742
C: 50,000 - 69,999	105.627588	110.127931	2.3496854	.018792277
D: 70,000 - 89,999	106.630299	110.47287	2.26839281	.023307831
E: 90,000 - 109,999	103.396741	101.610416	-1.2603509	.207545662
F: 110,000 - 129,999	106.76476	105.981312	60580011	.544648553
G: 130,000 - 149,999	108.877532	107.31377	85219781	.394107755
H: 150,000 - 169,999	110.987258	107.152191	-1.9451486	.051762624
I: 170,000 - 189,999	102.808238	107.43556	2.14966921	.031587875
J: 190,000 - 249,999	108.040564	115.343356	2.54749867	.010854966
K: 250,000 - 299,999	112.377993	108.196097	-1.4115514	.158091676
L: 300,000 and above	120.970235	112.216342	-2.0726194	.038225611
	107.121845	113.80441	.689462437	.490595765
	106.663769	107.276386	1.07853782	.280794207

14 rows selected.

STATS_WSR_TEST

Syntax



Purpose

STATS_WSR_TEST is a Wilcoxon Signed Ranks test of paired samples to determine whether the median of the differences between the samples is significantly different from zero. The absolute values of the differences are ordered and assigned ranks. Then the null hypothesis states that the sum of the ranks of the positive differences is equal to the sum of the ranks of the negative differences.

This function takes two required arguments: expr1 and expr2 are the two samples being analyzed. The optional third argument lets you specify the meaning of the NUMBER value returned by this function, as shown in Table 7-10. For this argument, you can specify a text

literal, or a bind variable or expression that evaluates to a constant character value. If you omit the third argument, then the default is 'TWO SIDED SIG'.

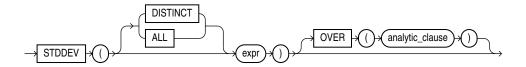
Table 7-10 STATS_WSR_TEST_* Return Values

Argument	Return Value Meaning
'STATISTIC'	The observed value of Z
'ONE_SIDED_SIG'	One-tailed significance of Z
'TWO_SIDED_SIG'	Two-tailed significance of Z

One-sided significance is always with respect to the upper tail. The high value (the value whose rejection region is the upper tail) is <code>expr1</code>.

STDDEV

Syntax





"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

STDDEV returns the sample standard deviation of expr, a set of numbers. You can use it as both an aggregate and analytic function. It differs from STDDEV_SAMP in that STDDEV returns zero when it has only 1 row of input data, whereas STDDEV_SAMP returns null.

Oracle Database calculates the standard deviation as the square root of the variance defined for the VARIANCE aggregate function.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.

See Also:

Table 2-9 for more information on implicit conversion

If you specify DISTINCT, then you can specify only the query_partition_clause of the analytic clause. The order by clause and windowing clause are not allowed.



See Also:

- "Aggregate Functions", VARIANCE, and STDDEV_SAMP
- "About SQL Expressions" for information on valid forms of expr

Aggregate Examples

The following example returns the standard deviation of the salaries in the sample hr.employees table:

```
SELECT STDDEV(salary) "Deviation"
FROM employees;

Deviation
-----
3909.36575
```

Analytic Examples

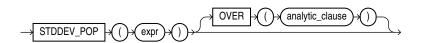
The query in the following example returns the cumulative standard deviation of the salaries in Department 80 in the sample table hr.employees, ordered by hire_date:

```
SELECT last_name, salary,
   STDDEV(salary) OVER (ORDER BY hire_date) "StdDev"
   FROM employees
   WHERE department_id = 30
   ORDER BY last_name, salary, "StdDev";
```

LAST_NAME	SALARY	StdDev
Baida	2900	4035.26125
Colmenares	2500	3362.58829
Himuro	2600	3649.2465
Khoo	3100	5586.14357
Raphaely	11000	0
Tobias	2800	4650.0896

STDDEV_POP

Syntax





"Analytic Functions" for information on syntax, semantics, and restrictions



Purpose

STDDEV_POP computes the population standard deviation and returns the square root of the population variance. You can use it as both an aggregate and analytic function.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.



Table 2-9 for more information on implicit conversion

This function is the same as the square root of the VAR_POP function. When VAR_POP returns null, this function returns null.

See Also:

- "Aggregate Functions" and VAR_POP
- "About SQL Expressions" for information on valid forms of expr

Aggregate Example

The following example returns the population and sample standard deviations of the amount of sales in the sample table sh.sales:

Analytic Example

The following example returns the population standard deviations of salaries in the sample hr.employees table by department:

```
SELECT department id, last name, salary,
  STDDEV POP(salary) OVER (PARTITION BY department id) AS pop std
  FROM employees
  ORDER BY department id, last name, salary, pop std;
DEPARTMENT ID LAST NAME
                                   SALARY POP_STD
------ -----
                                     4400 0
        10 Whalen
                                    6000 3500
13000 3500
        20 Fay
        20 Hartstein
                                     2900 3069.6091
        30 Baida
       100 Urman
                                   7800 1644.18166
       110 Gietz
                                     8300 1850
```



110 Higgins 12000 1850 Grant 7000 0

STDDEV_SAMP

Syntax



See Also:

"Analytic Functions " for information on syntax, semantics, and restrictions

Purpose

STDDEV_SAMP computes the cumulative sample standard deviation and returns the square root of the sample variance. You can use it as both an aggregate and analytic function.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.

See Also:

Table 2-9 for more information on implicit conversion

This function is same as the square root of the VAR_SAMP function. When VAR_SAMP returns null, this function returns null.

See Also:

- "Aggregate Functions" and VAR_SAMP
- "About SQL Expressions" for information on valid forms of expr

Aggregate Example

Refer to the aggregate example for STDDEV_POP.

Analytic Example

The following example returns the sample standard deviation of salaries in the employees table by department:



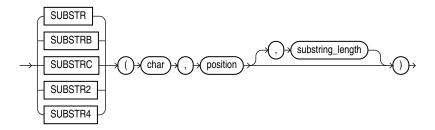
ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS cum_sdev FROM employees
ORDER BY department_id, last_name, hire_date, salary, cum_sdev;

DEPARTMENT_ID	LAST_NAME	HIRE_DATE	SALARY	CUM_SDEV
10	Whalen	17-SEP-03	4400	
20	Fay	17-AUG-05	6000	4949.74747
20	Hartstein	17-FEB-04	13000	
30	Baida	24-DEC-05	2900	4035.26125
30	Colmenares	10-AUG-07	2500	3362.58829
30	Himuro	15-NOV-06	2600	3649.2465
30	Khoo	18-MAY-03	3100	5586.14357
30	Raphaely	07-DEC-02	11000	
100	Greenberg	17-AUG-02	12008	2126.9772
100	Popp	07-DEC-07	6900	1804.13155
100	Sciarra	30-SEP-05	7700	1929.76233
100	Urman	07-MAR-06	7800	1788.92504
110	Gietz	07-JUN-02	8300	2621.95194
110	Higgins	07-JUN-02	12008	
	Grant	24-MAY-07	7000	

SUBSTR

Syntax

substr::=



Purpose

The SUBSTR functions return a portion of *char*, beginning at character *position*, *substring_length* characters long. SUBSTR calculates lengths using characters as defined by the input character set. SUBSTRB uses bytes instead of characters. SUBSTRC uses Unicode complete characters. SUBSTR2 uses UCS2 code points. SUBSTR4 uses UCS4 code points.

- If position is 0, then it is treated as 1.
- If position is positive, then Oracle Database counts from the beginning of char to find the first character.
- If position is negative, then Oracle counts backward from the end of char.
- If <code>substring_length</code> is omitted, then Oracle returns all characters to the end of <code>char</code>. If <code>substring_length</code> is less than 1, then Oracle returns null.

char can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The exceptions are SUBSTRC, SUBSTR2, and SUBSTR4, which do not allow char to be a CLOB or NCLOB. Both position and substring length must be of data type NUMBER, or any data type that can



be implicitly converted to <code>NUMBER</code>, and must resolve to an integer. The return value is the same data type as <code>char</code>, except that for a <code>CHAR</code> argument a <code>VARCHAR2</code> value is returned, and for an <code>NCHAR</code> argument an <code>NVARCHAR2</code> value is returned. Floating-point numbers passed as arguments to <code>SUBSTR</code> are automatically converted to integers.

See Also:

- For a complete description of character length see Oracle Database
 Globalization Support Guide and Oracle Database SecureFiles and Large
 Objects Developer's Guide
- Oracle Database Globalization Support Guide for more information about SUBSTR functions and length semantics in different locales
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of SUBSTR

Examples

The following example returns several specified substrings of "ABCDEFG":

```
SELECT SUBSTR('ABCDEFG',3,4) "Substring"
FROM DUAL;

Substring
------
CDEF

SELECT SUBSTR('ABCDEFG',-5,4) "Substring"
FROM DUAL;

Substring
------
CDEF
```

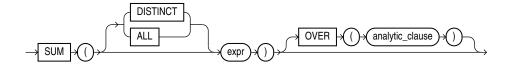
Assume a double-byte database character set:

```
SELECT SUBSTRB('ABCDEFG',5,4.2) "Substring with bytes"
FROM DUAL;

Substring with bytes
-----CD
```

SUM

Syntax





See Also:

"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

SUM returns the sum of values of expr. You can use it as an aggregate or analytic function.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.

Release 23 adds support for INTERVAL interval data types. However interval data types cannot be implicitly converted to a numeric data type. If the input is an INTERVAL, the function returns an INTERVAL with the same units as the input.

See Also:

Table 2-9 for more information on implicit conversion

If you specify DISTINCT, then you can specify only the <code>query_partition_clause</code> of the analytic clause. The order by clause and windowing clause are not allowed.

See Also:

"About SQL Expressions" for information on valid forms of expr and "Aggregate Functions"

Vector Aggregate Operations

You can use SUM to perform vector addition operations on non-null inputs.

expr must evaluate to VECTOR and must not be BINARY vectors. The returned vector has the same number of dimensions as the input, and the format is always FLOAT64. For flexible number of dimensions, all inputs must have the same number of dimensions within each aggregation group.

NULL vectors are ignored. They are not counted when calculating the average vector. If all inputs within an aggregation group are NULL, the result is NULL for that group. If a certain dimension overflows when applying arithmetic operations, an error is raised.

Rules

- DISTINCT syntax is not allowed.
- Only GROUP BY and GROUP BY ROLLUP are supported.
- Analytic functions are not supported for input arguments of type VECTOR.

See Arithmetic Operatorsof the AI Vector Search User's Guide for examples.



Aggregate Example

The following example calculates the sum of all salaries in the sample hr.employees table:

```
SELECT SUM(salary) "Total"
FROM employees;

Total
------
691400
```

Analytic Example

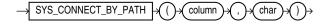
The following example calculates, for each manager in the sample table hr.employees, a cumulative total of salaries of employees who answer to that manager that are equal to or less than the current salary. You can see that Raphaely and Cambrault have the same cumulative total. This is because Raphaely and Cambrault have the identical salaries, so Oracle Database adds together their salary values and applies the same cumulative total to both rows.

```
SELECT manager_id, last_name, salary,
SUM(salary) OVER (PARTITION BY manager_id ORDER BY salary
RANGE UNBOUNDED PRECEDING) l_csum
FROM employees
ORDER BY manager_id, last_name, salary, l_csum;
```

MANAGER_ID	LAST_NAME	SALARY	L_CSUM
100	Cambrault	11000	68900
100	De Haan	17000	155400
100	Errazuriz	12000	80900
100	Fripp	8200	36400
100	Hartstein	13000	93900
100	Kaufling	7900	20200
100	Kochhar	17000	155400
100	Mourgos	5800	5800
100	Partners	13500	107400
100	Raphaely	11000	68900
100	Russell	14000	121400
149	Hutton	8800	39000
149	Johnson	6200	6200
149	Livingston	8400	21600
149	Taylor	8600	30200
201	Fay	6000	6000
205	Gietz	8300	8300
	King	24000	24000

SYS_CONNECT_BY_PATH

Syntax



Purpose

SYS_CONNECT_BY_PATH is valid only in hierarchical queries. It returns the path of a column value from root to node, with column values separated by *char* for each row returned by CONNECT BY condition.

Both column and char can be any of the data types CHAR, VARCHAR2, NCHAR, or NVARCHAR2. The string returned is of VARCHAR2 data type and is in the same character set as column.

See Also:

- "Hierarchical Queries" for more information about hierarchical queries and CONNECT BY conditions
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of SYS CONNECT BY PATH

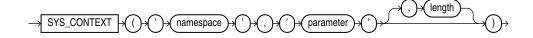
Examples

The following example returns the path of employee names from employee Kochhar to all employees of Kochhar (and their employees):

```
SELECT LPAD(' ', 2*level-1)||SYS CONNECT BY PATH(last name, '/') "Path"
  FROM employees
  START WITH last name = 'Kochhar'
  CONNECT BY PRIOR employee id = manager id;
Path
    /Kochhar/Greenberg/Chen
     /Kochhar/Greenberg/Faviet
     /Kochhar/Greenberg/Popp
     /Kochhar/Greenberg/Sciarra
     /Kochhar/Greenberg/Urman
     /Kochhar/Higgins/Gietz
   /Kochhar/Baer
   /Kochhar/Greenberg
   /Kochhar/Higgins
   /Kochhar/Mavris
   /Kochhar/Whalen
/Kochhar
```

SYS_CONTEXT

Syntax





Purpose

SYS_CONTEXT returns the value of parameter associated with the context namespace at the current instant. You can use this function in both SQL and PL/SQL statements. SYS_CONTEXT must be executed locally.

For namespace and parameter, you can specify either a string or an expression that resolves to a string designating a namespace or an attribute. If you specify literal arguments for namespace and parameter, and you are using SYS_CONTEXT explicitly in a SQL statement—rather than in a PL/SQL function that in turn is in mentioned in a SQL statement—then Oracle Database evaluates SYS_CONTEXT only once per SQL statement execution for each call site that invokes the SYS_CONTEXT function.

The context namespace must already have been created, and the associated parameter and its value must also have been set using the DBMS_SESSION.set_context procedure. The namespace must be a valid identifier. The parameter name can be any string. It is not case sensitive, but it cannot exceed 30 bytes in length.

The data type of the return value is VARCHAR2. The default maximum size of the return value is 256 bytes. You can override this default by specifying the optional <code>length</code> parameter, which must be a <code>NUMBER</code> or a value that can be implicitly converted to <code>NUMBER</code>. The valid range of values is 1 to 4000 bytes. If you specify an invalid value, then Oracle Database ignores it and uses the default.

Oracle provides the following built-in namespaces:

- USERENV Describes the current session. The predefined parameters of namespace
 USERENV are listed in Table 7-11.
- SYS_SESSION_ROLES Indicates whether a specified role is currently enabled for the
 session. Oracle Database evaluates the SYS_SESSION_ROLES context for the current user,
 and assumes the defining user's role when it evaluates SYS_SESSION_ROLES within a
 definer's rights procedure or function. An alternative to using SYS_SESSION_ROLES to find
 the login user's enabled roles in a definer's rights procedure is to use the
 DBMS_SESSION:SESSION_IS_ROLE_ENABLED function. Invoker's rights, procedures or
 functions, and/or code based access control (CBAC) are also alternatives.

See Also:

- Using Code Based Access Control for Definer's Rights and Invoker's Rights
- Oracle Database Security Guide for information on using the application context feature in your application development
- CREATE CONTEXT for information on creating user-defined context namespaces
- Oracle Database PL/SQL Packages and Types Reference for information on the DBMS SESSION.set context procedure
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of SYS CONTEXT



Examples

The following statement returns the name of the user who logged onto the database:

The following example queries the SESSION_ROLES data dictionary view to show that RESOURCE is the only role currently enabled for the session. It then uses the SYS_CONTEXT function to show that the RESOURCE role is currently enabled for the session and the DBA role is not.

```
CONNECT OE
Enter password: password

SELECT role FROM session_roles;

ROLE
______
RESOURCE

SELECT SYS_CONTEXT('SYS_SESSION_ROLES', 'RESOURCE')
FROM DUAL

SYS_CONTEXT('SYS_SESSION_ROLES', 'RESOURCE')
_____
TRUE

SELECT SYS_CONTEXT('SYS_SESSION_ROLES', 'DBA')
FROM DUAL;

SYS_CONTEXT('SYS_SESSION_ROLES', 'DBA')
_____
FALSE
```

Note:

For simplicity in demonstrating this feature, these examples do not perform the password management techniques that a deployed system normally uses. In a production environment, follow the Oracle Database password management guidelines, and disable any sample accounts. See *Oracle Database Security Guide* for password management guidelines and other security recommendations.

The following hypothetical example returns the group number that was set as the value for the attribute $group_no$ in the PL/SQL package that was associated with the context hr_apps when hr_apps was created:

```
SELECT SYS_CONTEXT ('hr_apps', 'group_no') "User Group"
FROM DUAL;
```



Starting with Oracle Database 23ai, users authenticating to the database using the legacy RADIUS API are not granted administrative privileges such as SYSDBA or SYSBACKUP.

In Oracle Database 23ai Oracle introduces a new RADIUS API that uses the latest standards to grant administrative privileges to users.

You must ensure that the database connection to the database uses the new RADIUS API and that you are using the Oracle Database 23ai client to connect to the Oracle Database 23ai server.

Table 7-11 Predefined Parameters of Namespace USERENV

Parameter	Return Value	
ACTION	Identifies the position in the module (application name) and is set through the DBMS_APPLICATION_INFO package or OCI.	
AUDITED_CURSORID	Returns the cursor ID of the SQL that triggered the audit. This parameter is not valid in a fine-grained auditing environment. If you specify it in such an environment, then Oracle Database always returns null.	
AUTHENTICATED_IDENTITY	Returns the identity used in authentication. In the list that follows, the type of user is followed by the value returned:	
	Kerberos-authenticated enterprise user: kerberos principal name	
	 Kerberos-authenticated external user: kerberos principal name; same as the schema name 	
	 SSL-authenticated enterprise user: the DN in the user's PKI certificate 	
	 SSL-authenticated external user: the DN in the user's PKI certificate 	
	 Password-authenticated enterprise user: nickname; same as the login name 	
	 Password-authenticated database user: the database username; same as the schema name 	
	 OS-authenticated external user: the external operating system user name 	
	Radius-authenticated external user: the schema name	
	 Proxy with DN: Oracle Internet Directory DN of the client 	
	 Proxy with certificate: certificate DN of the client 	
	 For single session proxy or dual session proxy without client authentication: database user name if proxy is a local database user; nickname if proxy is an enterprise user. 	
	For dual session proxy with client authentication: database user name if client is a local database user; nickname if client is an enterprise user.	
	SYSDBA/SYSOPER using Password File: login name	
	 SYSDBA/SYSOPER using OS authentication: operating system user name 	
AUTHENTICATION_DATA	Data being used to authenticate the login user. For X.503 certificate authenticated sessions, this field returns the context of the certificate in HEX2 format.	
	Note: You can change the return value of the AUTHENTICATION_DATA attribute using the <code>length</code> parameter of the syntax. Values of up to 4000 are accepted. This is the only attribute of <code>USERENV</code> for which Oracle Database implements such a change.	



Table 7-11 (Cont.) Predefined Parameters of Namespace USERENV

Parameter	Return Value		
AUTHENTICATION_METHOD	Returns the method of authentication. In the list that follows, the type of user is followed by the method returned:		
	 Password-authenticated enterprise user, local database user, or user with the SYSDBA or SYSOPER administrative privilege using a password file; proxy with username using password: PASSWORD 		
	 Password-authenticated enterprise user, local database user, or user with the SYSDBA or SYSOPER administrative privilege using a password file; proxy with username using password: PASSWORD_GLOBAL 		
	 Kerberos-authenticated enterprise user or external user (with no administrative privileges): KERBEROS 		
	 Kerberos-authenticated enterprise user (with administrative privileges): KERBEROS GLOBAL 		
	 Kerberos-authenticated external user (with administrative privileges): KERBEROS EXTERNAL 		
	 SSL-authenticated enterprise or external user (with no administrative privileges): 		
	SSL-authenticated enterprise user (with administrative privileges): SSL GLOBAL		
	SSL-authenticated external user (with administrative privileges): SSL EXTERNAL		
	Radius-authenticated external user: RADIUS		
	 OS-authenticated external user or use with the SYSDBA or SYSOPER administrative privilege: OS 		
	• Proxy authentication: AUTHENTICATION_METHOD used during authentication of PROXY USER with "_PROXY" added at end. For example, if a proxy user uses PASSWORD to connect to the database, then the AUTHENTICATION_METHOD will be PASSWORD PROXY.		
	In the case of dual session proxy without client authentication: PROXYUSER AUTHENTICATED PROXY		
	Background process (job queue slave process): JOB		
	Parallel Query Slave process: PQ SLAVE		
	For non-administrative connections, you can use <code>IDENTIFICATION_TYPE</code> to distinguish between external and enterprise users when the authentication method is <code>PASSWORD</code> , <code>KERBEROS</code> , or <code>SSL</code> . For administrative connections, <code>AUTHENTICATION_METHOD</code> is sufficient for the <code>PASSWORD</code> , <code>SSL_EXTERNAL</code> , and <code>SSL_GLOBAL</code> authentication methods.		
BG_JOB_ID	Job ID of the current session if it was established by an Oracle Database background process. Null if the session was not established by a background process.		
CDB_DOMAIN	CDB_DOMAIN is the DB_DOMAIN of the CDB and is the same for all the PDBs associated with it.		
CDB_NAME	If queried while connected to a multitenant container database (CDB), returns the name of the CDB. Otherwise, returns null.		
CLIENT_IDENTIFIER	Returns an identifier that is set by the application through the DBMS_SESSION.SET_IDENTIFIER procedure, the OCI attribute OCI_ATTR_CLIENT_IDENTIFIER, or Oracle Dynamic Monitoring Service (DMS). This attribute is used by various database components to identify lightweight application users who authenticate as the same database user.		
CLIENT_INFO	Returns up to 64 bytes of user session information that can be stored by an application using the <code>DBMS_APPLICATION_INFO</code> package.		
CLIENT_PROGRAM_NAME	The name of the program used for the database session.		
CLOUD_MIGRATION_MODE	The session parameter to specify ON when CLOUD_MIGRATION_MODE is TRUE, else OFF.		



Table 7-11 (Cont.) Predefined Parameters of Namespace USERENV

Parameter	Return Value	
CLOUD_SERVICE	Only valid for cloud implementations.	
	Returns DWCS on autonomous database management systems (ADW), OLTP on autonomous transaction processing systems (ATP), and JDCS on autonomous JSON database systems.	
CON_ID	Returns the current container ID that the session is connected to.	
CON_NAME	Returns the current container name that the session is connected to.	
CURRENT_BIND	The bind variables for fine-grained auditing. You can specify this attribute only inside the event handler for the fine-grained auditing feature.	
CURRENT_EDITION_ID	The identifier of the current edition.	
CURRENT_EDITION_NAME	The name of the current edition.	
CURRENT_SCHEMA	The name of the currently active default schema. This value may change during the duration of a session through use of an ALTER SESSION SET CURRENT_SCHEMA statement. This may also change during the duration of a session to reflect the owner of any active definer's rights object. When used directly in the body of a view definition, this returns the default schema used when executing the cursor that is using the view; it does not respect views used in the cursor as being definer's rights.	
	Note : Oracle recommends against issuing the SQL statement ALTER SESSION SET CURRENT_SCHEMA from within all types of stored PL/SQL units except logon triggers.	
CURRENT_SCHEMAID	Identifier of the currently active default schema.	
CURRENT_SQL CURRENT_SQLn	CURRENT_SQL returns the first 4K bytes of the current SQL that triggered the fine-grained auditing event. The CURRENT_SQLn attributes return subsequent 4K-byte increments, where n can be an integer from 1 to 7, inclusive. CURRENT_SQL1 returns bytes 4K to 8K; CURRENT_SQL2 returns bytes 8K to 12K, and so forth. You can specify these attributes only inside the event handler for the fine-grained auditing feature.	
CURRENT_SQL_LENGTH	The length of the current SQL statement that triggers fine-grained audit or row-level security (RLS) policy functions or event handlers. You can specify this attribute only inside the event handler for the fine-grained auditing feature.	
CURRENT_USER	The name of the database user whose privileges are currently active. This may change during the duration of a database session as Real Application Security sessions are attached or detached, or to reflect the owner of any active definer's rights object. When no definer's rights object is active, <code>CURRENT_USER</code> returns the same value as <code>SESSION_USER</code> . When used directly in the body of a view definition, this returns the user that is executing the cursor that is using the view; it does not respect views used in the cursor as being definer's rights. For enterprise users, returns schema. If a Real Application Security user is currently active, returns user <code>XS\$NULL</code> .	
CURRENT_USERID	The identifier of the database user whose privileges are currently active.	
DATABASE_ROLE	The database role using the SYS_CONTEXT function with the USERENV namespace.	
	The role is one of the following: PRIMARY, PHYSICAL STANDBY, LOGICAL STANDBY, SNAPSHOT STANDBY, TRUE CACHE.	
DB_DOMAIN	Domain of the database as specified in the DB_DOMAIN initialization parameter.	
DB_NAME	Name of the database as specified in the DB_NAME initialization parameter.	
DB_SUPPLEMENTAL_LOG_LEVEL		



Table 7-11 (Cont.) Predefined Parameters of Namespace USERENV

Parameter	Return Value		
DB_UNIQUE_NAME	Name of the database as specified in the DB_UNIQUE_NAME initialization parameter.		
DBLINK_INFO	Returns the source of a database link session. Specifically, it returns a string of the form		
	SOURCE_GLOBAL_NAME=dblink_src_global_name, DBLINK_NAME=dblink_name, SOURCE_AUDIT_SESSIONID=dblink_src_audit_sessionid		
	For a multitenant database, it returns the string above with an additional field <code>SOURCE_DB_NAME</code> :		
	SOURCE_GLOBAL_NAME=dblink_src_global_name, SOURCE_DB_NAME=source_database_name, DBLINK_NAME=dblink_name, SOURCE_AUDIT_SESSIONID=dblink_src_audit_sessionid		
	where:		
	• dblink_src_global_name is the unique global name of the source database		
	 dblink_name is the name of the database link on the source database 		
	 dblink_src_audit_sessionid is the audit session ID of the session on the source database that initiated the connection to the remote database using dblink name 		
	 SOURCE_DB_NAME is the database identifier of the source database 		
DRAIN_STATUS	Displays the draining status for the current session. Returns DRAINING if the session is a candidate for drain else returns NONE.		
ENTRYID	The current audit entry number. The audit entryid sequence is shared between fine- grained audit records and regular audit records. You cannot use this attribute in distributed SQL statements. The correct auditing entry identifier can be seen only through an audit handler for standard or fine-grained audit.		
ENTERPRISE_IDENTITY	Returns the user's enterprise-wide identity:		
	 For enterprise users: the Oracle Internet Directory DN. 		
	 For external users: the external identity (Kerberos principal name, Radius schema names, OS user name, Certificate DN). 		
	For local users and SYSDBA/SYSOPER logins: NULL.		
	The value of the attribute differs by proxy method:		
	 For a proxy with DN: the Oracle Internet Directory DN of the client For a proxy with certificate: the certificate DN of the client for external users; the Oracle Internet Directory DN for global users 		
	 For a proxy with username: the Oracle Internet Directory DN if the client is an enterprise users; Null if the client is a local database user. 		
FG_JOB_ID	If queried from within a job that was created using the <code>DBMS_JOB</code> package: Returns the job ID of the current session if it was established by a client foreground process. Null if the session was not established by a foreground process.		
	Otherwise: Returns 0.		
GLOBAL_CONTEXT_MEMORY	Returns the number being used in the System Global Area by the globally accessed context.		
GLOBAL_UID	Returns the global user ID (GUID) from Active Directory for Centrally Managed Users (CMU) logins, or from Oracle Internet Directory for Enterprise User Security (EUS) logins. Returns null for all other logins.		
HOST	Name of the host machine from which the client has connected.		



Table 7-11 (Cont.) Predefined Parameters of Namespace USERENV

Parameter	Return Value	
IDENTIFICATION_TYPE	Returns the way the user's schema was created in the database. Specifically, it reflects the IDENTIFIED clause in the CREATE/ALTER USER syntax. In the list that follows, the syntax used during schema creation is followed by the identification type returned:	
	• IDENTIFIED BY password: LOCAL	
	• IDENTIFIED EXTERNALLY: EXTERNAL	
	IDENTIFIED GLOBALLY: GLOBAL SHARED	
	• IDENTIFIED GLOBALLY AS <i>DN</i> : GLOBAL PRIVATE	
	 GLOBAL EXCLUSIVE for exclusive global user mapping. 	
	GLOBAL SHARED for shared user mapping.	
	NONE when the schema is created with no authentication.	
INSTANCE	The instance identification number of the current instance.	
INSTANCE_NAME	The name of the instance.	
IP_ADDRESS	IP address of the machine from which the client is connected. If the client and server are on the same machine and the connection uses IPv6 addressing, then ::1 is returned.	
IS_APPLY_SERVER	Returns TRUE if queried from within a SQL Apply server in a logical standby database. Otherwise, returns FALSE.	
IS_DG_ROLLING_UPGRADE	Returns TRUE if a rolling upgrade of the database software in a Data Guard configuration, initiated by way of the DBMS_ROLLING package, is active. Otherwise, returns FALSE.	
ISDBA	Returns TRUE if the user has been authenticated as having DBA privileges either through the operating system or through a password file.	
LANG	The abbreviated name for the language, a shorter form than the existing 'LANGUAGE' parameter.	
LANGUAGE	The language and territory currently used by your session, along with the database character set, in this form:	
	<pre>language_territory.characterset</pre>	
LDAP_SERVER_TYPE	Returns the configured LDAP server type, one of OID, AD(Active Directory), OID_G, OPENLDAP.	
MODULE	The application name (module) set through the <code>DBMS_APPLICATION_INFO</code> package or OCI.	
NETWORK_PROTOCOL	Network protocol being used for communication, as specified in the 'PROTOCOL=protocol' portion of the connect string.	
NLS_CALENDAR	The current calendar of the current session.	
NLS_CURRENCY	The currency of the current session.	
NLS_DATE_FORMAT	The date format for the session.	
NLS_DATE_LANGUAGE	The language used for expressing dates.	
NLS_SORT	BINARY or the linguistic sort basis.	
NLS_TERRITORY	The territory of the current session.	
ORACLE_HOME	The full path name for the Oracle home directory.	
OS_USER	Operating system user name of the client process that initiated the database session.	
PID	Oracle process ID.	
PLATFORM SLASH	The slash character that is used as the file path delimiter for your platform.	



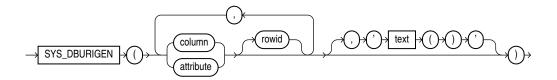
Table 7-11 (Cont.) Predefined Parameters of Namespace USERENV

Parameter	Return Value	
POLICY_INVOKER	The invoker of row-level security (RLS) policy functions.	
PROXY_ENTERPRISE_IDENTITY	Returns the Oracle Internet Directory DN when the proxy user is an enterprise user.	
PROXY_USER	Name of the database user who opened the current session on behalf of SESSION_USER.	
PROXY_USERID	Identifier of the database user who opened the current session on behalf of SESSION_USER.	
RESET_STATE	RESET_STATE can be set using DBMS_APP_CONT_ADMIN.ENABLE_RESET_STATE() procedure call and is related to Application Continuity.	
SCHEDULER_JOB	Returns Y if the current session belongs to a foreground job or background job. Otherwise, returns N.	
SERVER_HOST	The host name of the machine on which the instance is running.	
SERVICE_NAME	The name of the service to which a given session is connected.	
SESSION_DEFAULT_COLLATION	The default collation for the session, which is set by the ALTER SESSION SET DEFAULT_COLLATION statement.	
SESSION_EDITION_ID	The identifier of the session edition.	
SESSION_EDITION_NAME	The name of the session edition.	
SESSION_USER	The name of the session user (the user who logged on). This may change during the duration of a database session as Real Application Security sessions are attached or detached. If a Real Application Security session is currently attached to the database session, returns user XS\$NULL.	
SESSION_USERID	The identifier of the session user (the user who logged on).	
SESSIONID	The auditing session identifier. You cannot use this attribute in distributed SQL statements.	
SID	The session ID.	
STANDBY_MAX_DATA_DELAY	The session parameter to specify allowed time limit to elapse between when changes are committed on primary database and when those changes can be queried on the standby database. If not set, returns null.	
STATEMENTID	The auditing statement identifier. STATEMENTID represents the number of SQL statements audited in a given session. You cannot use this attribute in distributed SQL statements. The correct auditing statement identifier can be seen only through an audit handler for standard or fine-grained audit.	
TERMINAL	The operating system identifier for the client of the current session. In distributed SQL statements, this attribute returns the identifier for your local session. In a distributed environment, this is supported only for remote SELECT statements, not for remote INSERT, UPDATE, or DELETE operations. (The return length of this parameter may vary by operating system.)	
TLS_CIPHERSUITE	Used to retrieve the ciphersuite negotiated during the TLS.	
	Valid ciphersuite values can be found in the TLS chapter of the Database Security Guide.	
TLS_VERSION	Used to retrieve the TLS version negotiated during the TLS session.	
UNIFIED_AUDIT_SESSIONID	If queried while connected to a database that uses unified auditing or mixed mode auditing, returns the unified audit session ID.	
	If queried while connected to a database that uses traditional auditing, returns null.	



SYS_DBURIGEN

Syntax



Purpose

SYS_DBURIGEN takes as its argument one or more columns or attributes, and optionally a rowid, and generates a URL of data type DBURIType to a particular column or row object. You can then use the URL to retrieve an XML document from the database.

All columns or attributes referenced must reside in the same table. They must perform the function of a primary key. They need not actually match the primary key of the table, but they must reference a unique value. If you specify multiple columns, then all but the final column identify the row in the database, and the last column specified identifies the column within the row.

By default the URL points to a formatted XML document. If you want the URL to point only to the text of the document, then specify the optional 'text()'.



In this XML context, the lowercase text is a keyword, not a syntactic placeholder.

If the table or view containing the columns or attributes does not have a schema specified in the context of the query, then Oracle Database interprets the table or view name as a public synonym.

See Also:

Oracle XML DB Developer's Guide for information on the DBURIType data type and XML documents in the database

Examples

The following example uses the SYS_DBURIGen function to generate a URL of data type DBURIType to the email column of the row in the sample table hr.employees where the employee id = 206:

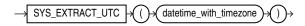
```
SELECT SYS_DBURIGEN(employee_id, email)
   FROM employees
   WHERE employee_id = 206;
SYS_DBURIGEN(EMPLOYEE_ID, EMAIL)(URL, SPARE)
```



DBURITYPE('/PUBLIC/EMPLOYEES/ROW[EMPLOYEE ID=''206'']/EMAIL', NULL)

SYS_EXTRACT_UTC

Syntax



Purpose

SYS_EXTRACT_UTC extracts the UTC (Coordinated Universal Time—formerly Greenwich Mean Time) from a datetime value with time zone offset or time zone region name. If a time zone is not specified, then the datetime is associated with the session time zone.

Examples

The following example extracts the UTC from a specified datetime:

SYS_GUID

Syntax



Purpose

SYS_GUID generates and returns a globally unique identifier (RAW value) made up of 16 bytes. On most platforms, the generated identifier consists of a host identifier, a process or thread identifier of the process or thread invoking the function, and a nonrepeating value (sequence of bytes) for that process or thread.

Examples

The following example adds a column to the sample table hr.locations, inserts unique identifiers into each row, and returns the 32-character hexadecimal representation of the 16-byte RAW value of the global unique identifier:

```
ALTER TABLE locations ADD (uid_col RAW(16));

UPDATE locations SET uid_col = SYS_GUID();

SELECT location_id, uid_col FROM locations
   ORDER BY location_id, uid_col;

LOCATION_ID UID_COL
```



```
1000 09F686761827CF8AE040578CB20B7491

1100 09F686761828CF8AE040578CB20B7491

1200 09F686761829CF8AE040578CB20B7491

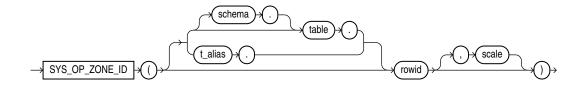
1300 09F68676182ACF8AE040578CB20B7491

1400 09F68676182BCF8AE040578CB20B7491

1500 09F68676182CCF8AE040578CB20B7491
```

SYS_OP_ZONE ID

Syntax



Purpose

SYS_OP_ZONE_ID takes as its argument a rowid and returns a zone ID. The rowid identifies a row in a table. The zone ID identifies the set of contiguous disk blocks, called the zone, that contains the row. The function returns a NUMBER value.

The SYS_OP_ZONE_ID function is used when creating a zone map with the CREATE MATERIALIZED ZONEMAP statement. You must specify SYS_OP_ZONE_ID in the SELECT and GROUP BY clauses of the defining subquery of the zone map.

For rowid, specify the ROWID pseudocolumn of the fact table of the zone map.

Use schema and table to specify the schema and name of the fact table, or t_alias to specify the table alias for the fact table. The specification of these parameters depends on the FROM clause in the defining subquery of the zone map:

- If the FROM clause specifies a table alias for the fact table, then you must also specify the table alias (t_alias) in SYS_OP_ZONE_ID.
- If the FROM clause does not specify a table alias for the fact table, then use <code>table</code> to specify the name of the fact table. You can use the <code>schema</code> qualifier if the fact table is in a schema other than your own. If you omit <code>schema</code>, then the database assumes the fact table is in your own schema. If the <code>FROM</code> clause specifies only one table (the fact table) then you need not specify <code>schema</code> or <code>table</code>.

The optional scale parameter represents the scale of the zone map. It is not necessary to specify this parameter because, by default, $SYS_OP_ZONE_ID$ uses the scale of the zone map being created. If you do specify scale, then it must match the scale of the zone map being created. Refer to the SCALE clause of CREATE MATERIALIZED ZONEMAP for information on specifying the scale of a zone map.



CREATE MATERIALIZED ZONEMAP for more information on creating zone maps

Examples

The following example uses the $SYS_OP_ZONE_ID$ function when creating a basic zone map that tracks the column $time_id$ of the fact table sales. The scale of the zone map is the default value of 10. Therefore, the $SYS_OP_ZONE_ID$ function will default to a scale value of 10.

```
CREATE MATERIALIZED ZONEMAP sales_zmap

AS

SELECT SYS_OP_ZONE_ID(rowid), MIN(time_id), MAX(time_id)

FROM sales

GROUP BY SYS OP ZONE ID(rowid);
```

The following example is similar to the previous example, except that the scale of the zone map being created is specified as 8. Therefore, the SYS_OP_ZONE_ID function will default to a scale value of 8.

```
CREATE MATERIALIZED ZONEMAP sales_zmap
SCALE 8
AS
    SELECT SYS_OP_ZONE_ID(rowid), MIN(time_id), MAX(time_id)
FROM sales
GROUP BY SYS OP ZONE ID(rowid);
```

The following example returns an error because the scale of the zone map being created is specified as 8, which does not match the scale argument of 12 specified in the SYS OP ZONE ID function.

```
CREATE MATERIALIZED ZONEMAP sales_zmap
SCALE 8
AS
SELECT SYS_OP_ZONE_ID(rowid, 12), MIN(time_id), MAX(time_id)
FROM sales
```

The following example creates a join zone map. The fact table is sales and the dimension tables are products and customers. Because the table alias s is specified for the fact table in the FROM clause, the table alias s is also specified in the SYS OP ZONE ID function.

```
CREATE MATERIALIZED ZONEMAP sales_zmap

AS

SELECT SYS_OP_ZONE_ID(s.rowid),

MIN(prod_category), MAX(prod_category),

MIN(country_id), MAX(country_id)

FROM sales s, products p, customers c

WHERE s.prod_id = p.prod_id(+) AND

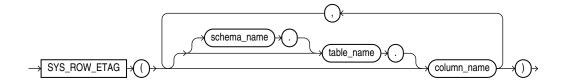
s.cust_id = c.cust_id(+)

GROUP BY SYS OP ZONE ID(s.rowid);
```

GROUP BY SYS_OP_ZONE_ID(rowid, 12);

SYS_ROW_ETAG

Syntax



Purpose

You can use ETAGS with table data, for lock-free row updates using SQL. To do that, use function SYS_ROW_ETAG, to obtain the current state of a given set of columns in a table row as an ETAG hash value. Function SYS_ROW_ETAG calculates an etag (128 bits hash value) for a row using the values of a set of columns in the row that you want the etag to be computed on. You can pass the function the names of the columns in any order.

Function SYS_ROW_ETAG calculates the ETAG value for a row using only the values of those columns in the row: you pass it the names of all columns that you want to be sure no other session tries to update concurrently. This includes the columns that the current session intends to update, but also any other columns on whose value that updating operation logically depends for your application. (The order in which you pass the columns to SYS_ROW_ETAG as arguments is irrelevant.)

Example

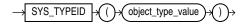
The example below creates table foo with columns c1, c2, and c3 of type NUMBER, and inserts values into the table. It then passes columns c2 and c1 to SYS_ROW_ETAG to get the etag for c2 and c1:

See Also:

Example 4.18 in the JSON-Relational Duality Developer's Guide Using Function SYS_ROW_ETAG To Optimistically Control Concurrent Table Updates

SYS TYPEID

Syntax



Purpose

SYS_TYPEID returns the typeid of the most specific type of the operand. This value is used primarily to identify the type-discriminant column underlying a substitutable column. For



example, you can use the value returned by SYS_TYPEID to build an index on the type-discriminant column.

You can use this function only on object type operands. All final root object types—final types not belonging to a type hierarchy—have a null typeid. Oracle Database assigns to all types belonging to a type hierarchy a unique non-null typeid.



Oracle Database Object-Relational Developer's Guide for more information on typeids

Examples

The following examples use the tables persons and books, which are created in "Substitutable Table and Column Examples". The first query returns the most specific types of the object instances stored in the persons table.

SELECT name, SYS TYPEID(VALUE(p)) "Type id" FROM persons p;

NAME	Type_id
Bob	01
Joe	02
Tim	03

The next query returns the most specific types of authors stored in the table books:

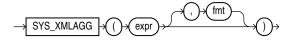
SELECT b.title, b.author.name, SYS_TYPEID(author)
"Type_ID" FROM books b;

TITLE	AUTHOR.NAME	Type_ID
An Autobiography	Bob	01
Business Rules	Joe	02
Mixing School and Work	Tim	03

You can use the SYS_TYPEID function to create an index on the type-discriminant column of a table. For an example, see "Indexing on Substitutable Columns: Examples".

SYS XMLAGG

Syntax



Purpose

SYS_XMLAgg aggregates all of the XML documents or fragments represented by expr and produces a single XML document. It adds a new enclosing element with a default name ROWSET. If you want to format the XML document differently, then specify fmt, which is an instance of the XMLFormat object.



See Also:

SYS_XMLGEN and "XML Format Model" for using the attributes of the XMLFormat type to format SYS XMLAgg results

Examples

The following example uses the SYS_XMLGen function to generate an XML document for each row of the sample table <code>employees</code> where the employee's last name begins with the letter R, and then aggregates all of the rows into a single XML document in the default enclosing element <code>ROWSET</code>:

SYS_XMLGEN

Note:

The SYS_XMLGen function is deprecated. It is still supported for backward compatibility. However, Oracle recommends that you use the SQL/XML generation functions instead. See *Oracle XML DB Developer's Guide* for more information.

Syntax



Purpose

SYS_XMLGen takes an expression that evaluates to a particular row and column of the database, and returns an instance of type XMLType containing an XML document. The expr can be a scalar value, a user-defined type, or an XMLType instance.

- If expr is a scalar value, then the function returns an XML element containing the scalar value
- If expr is a type, then the function maps the user-defined type attributes to XML elements.

 If expr is an XMLType instance, then the function encloses the document in an XML element whose default tag name is ROW.

By default the elements of the XML document match the elements of expr. For example, if expr resolves to a column name, then the enclosing XML element will be the same column name. If you want to format the XML document differently, then specify fmt, which is an instance of the XMLFormat object.



"XML Format Model" for a description of the XMLFormat type and how to use its attributes to format SYS_XMLGen results

Examples

The following example retrieves the employee email ID from the sample table <code>oe.employees</code> where the <code>employee_id</code> value is 205, and generates an instance of an <code>XMLType</code> containing an XML document with an <code>EMAIL</code> element.

SYSDATE

Syntax



Purpose

SYSDATE returns the current date and time set for the operating system on which the database server resides. The data type of the returned value is DATE, and the format returned depends on the value of the NLS_DATE_FORMAT initialization parameter. The function requires no arguments. In distributed SQL statements, this function returns the date and time set for the operating system of your local database. You cannot use this function in the condition of a CHECK constraint.

In a multitenant setup existing PDBs and PDBs created later inherit the timezone of the system.

If you want SYSDATE to return the timezone of the PDB, then you must set the initialization parameter <code>TIME_AT_DBTIMEZONE</code> to <code>TRUE</code> before starting the PDB.

You can change the timezone using ALTER SYSTEM SET TIME_ZONE or ALTER DATABASE db name SET TIME ZONE.



You can set SYSTIMESTAMP to return system time by setting the initialization parameter TIME AT DBTIMEZONE to FALSE and restarting the database.

Note:

- For more see TIME_AT_DBTIMEZONE of the Oracle Database Reference.
- The FIXED_DATE initialization parameter enables you to set a constant date and time that SYSDATE will always return instead of the current date and time. This parameter is useful primarily for testing. Refer to Oracle Database Reference for more information on the FIXED DATE initialization parameter.

Examples

The following example returns the current operating system date and time:

SYSTIMESTAMP

Syntax



Purpose

SYSTIMESTAMP returns the system date, including fractional seconds and time zone, of the system on which the database resides. The return type is TIMESTAMP WITH TIME ZONE.

In a multitenant setup existing PDBs and PDBs created later inherit the timezone of the system.

If you want SYSTIMESTAMP to return the timezone of the PDB, then you must set the initialization parameter TIME_AT_DBTIMEZONE to TRUE before starting the PDB.

You can change the timezone using ALTER SYSTEM SET TIME_ZONE or ALTER DATABASE db name SET TIME ZONE.

You can set SYSTIMESTAMP to return system time by setting the initialization parameter TIME AT DBTIMEZONE to FALSE and restarting the database.



For more see TIME_AT_DBTIMEZONE of the Oracle Database Reference.

Examples

The following example returns the system timestamp:

```
SELECT SYSTIMESTAMP FROM DUAL;

SYSTIMESTAMP

28-MAR-00 12.38.55.538741 PM -08:00
```

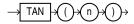
The following example shows how to explicitly specify fractional seconds:

The following example returns the current timestamp in a specified time zone:

The output format in this example depends on the NLS TIMESTAMP TZ FORMAT for the session.

TAN

Syntax



Purpose

TAN returns the tangent of n (an angle expressed in radians).

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



Table 2-9 for more information on implicit conversion

Examples

The following example returns the tangent of 135 degrees:

```
SELECT TAN(135 * 3.14159265359/180)
"Tangent of 135 degrees" FROM DUAL;
```



```
Tangent of 135 degrees
```

TANH

Syntax



Purpose

TANH returns the hyperbolic tangent of n.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If the argument is <code>BINARY_FLOAT</code>, then the function returns <code>BINARY_DOUBLE</code>. Otherwise the function returns the same numeric data type as the argument.



Table 2-9 for more information on implicit conversion

Examples

The following example returns the hyperbolic tangent of .5:

TIMESTAMP_TO_SCN

Syntax



Purpose

TIMESTAMP_TO_SCN takes as an argument a timestamp value and returns the approximate system change number (SCN) associated with that timestamp. The returned value is of data type NUMBER. This function is useful any time you want to know the SCN associated with a particular timestamp.





The association between an SCN and a timestamp when the SCN is generated is remembered by the database for a limited period of time. This period is the maximum of the auto-tuned undo retention period, if the database runs in the Automatic Undo Management mode, and the retention times of all flashback archives in the database, but no less than 120 hours. The time for the association to become obsolete elapses only when the database is open. An error is returned if the timestamp specified for the argument to TIMESTAMP TO SCN is too old.

See Also:

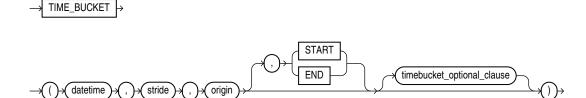
SCN_TO_TIMESTAMP for information on converting SCNs to timestamp

Examples

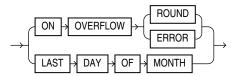
The following example inserts a row into the <code>oe.orders</code> table and then uses <code>TIMESTAMP_TO_SCN</code> to determine the system change number of the insert operation. (The actual SCN returned will differ on each system.)

TIME_BUCKET (datetime)

Syntax



timebucket_optional_clause::=



Purpose

Use TIME BUCKET(datetime) to obtain the datetime over an interval that you specify.

TIME BUCKET has three required arguments, and two optional arguments.

The first argument datetime is the input to the bucket.

The third argument *origin* is an anchor to which all buckets are aligned.

datetime and origin can be DATE, TIMESTAMP, TIMESTAMP WITH TIMEZONE, TIMESTAMP WITH LOCAL TIMEZONE, EPOCH TIME, BINARY_FLOAT, BINARY_DOUBLE, CHAR, expression, or a bind variable.

EPOCH TIME is represented by Oracle type NUMBER, which is the number of seconds that have elapsed since 00:00:00 UTC on 1 January 1970. The supported EPOCH TIME range is from SB8MINVAL(- 9223372036854775808, inclusive) to SB8MAXVAL(9223372036854775807, inclusive).

There are implicit conversions for datetime and origin:

- If it is BINARY_FLOAT or BINARY_DOUBLE, it will be converted to NUMBER implicitly. Note
 that you must account for the loss in precision from implicit conversions.
- If it is CHAR, it will be converted to TIMESTAMP implicitly. Note that CHAR should match the session NLS TIMESTAMP FORMAT. Otherwise an error is raised.

Fractional second is supported only if datetime and origin are EPOCH TIME, BINARY_FLOAT or BINARY_DOUBLE.

The valid range for datetime and origin is from -4712-01-01 00:00:00 inclusive to 9999-12-31 23:59:59:00 inclusive.

• The second argument *stride* is a positive Oracle INTERVAL, ISO 8601 time interval string, expression or bind variable. Fractional second is supported only if *datetime* and *origin* are EPOCH TIME, BINARY FLOAT OR BINARY DOUBLE.

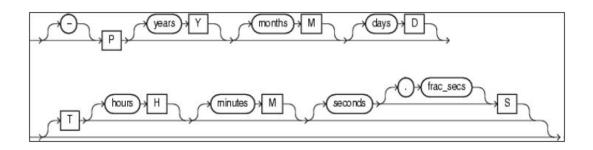
Oracle Interval has two types of valid intervals: Interval Year to Month and Interval DAY to Second. If year or month is specified, all other units are ignored, if specified.

The ISO 8601 time interval string that you specify should match the definition of Oracle INTERVAL. P is required, and no blanks are allowed in the value. If you specify T, then you must specify at least one of hours, minutes, or seconds. hours are based on 24-hour time.

For example, P100DT05H indicates 100 days and 5 hours. P1Y2M'indicates 1 year and 2 months. P1M1DT5H30M30S is equivalent to P1M which indicates 1 month.

The syntax of the ISO 8601 time interval string:





Use only postive values for *stride*. (Although the Oracle INTERVAL and ISO 8601 time interval string can be positive or negative.)

If datetime or origin is EPOCH TIME, BINARY_FLOAT or BINARY_DOUBLE, then stride cannot contain YEAR or MONTH. This is because month is variable and could be one of 28, 29, 30 or 31.

- The fourth argument is optional and specifies whether the start or the end of the time bucket is returned. Specify START to return the start value of the time bucket or END to return the end value. The values are case-insensitive. The default value is START.
- The fifth argument is optional and controls how the buckets (strides) are determined.

ON OVERFLOW ROUND (default): The buckets will be cut on the same day as <code>origin</code> in the corresponding month. For a month that does not have that day, the bucket is rounded to the last day of the month.

ON OVERFLOW ERROR: The buckets will be cut on the same day as <code>origin</code> in the corresponding month. For a month that does not have that day will error out.

LAST DAY OF MONTH: If origin is the last day of the month and stride only contains MONTH and/or YEAR, the buckets will be cut on the corresponding last day of the month.

For example, if origin is '1991-11-30' and stride is 'P1M', then:

For on overflow round, the start of each bucket will be:

```
..., 1991-11-30, 1991-12-30, 1992-01-30, 1992-02-29, 1992-03-30, 1992-04-30,...
```

For on overflow error, the start of each bucket will be:

```
\dots, 1991-11-30, 1991-12-30, 1992-01-30, error (or 1992-02-30), 1992-03-30, 1992-04-30,...
```

For LAST DAY OF MONTH, the start of each bucket will be:

```
..., 1991-11-30, 1991-12-31, 1992-01-31, 1992-02-29, 1992-03-31, 1992-04-30, ...
```

Rules

- The end of each bucket is the same as the beginning of the following bucket. For example, if the bucket is 2 years and the start of the slice is 2000-01-01, then the end of the bucket will be 2002-01- 01, not 2001-12-31. In other words, the bucket contains <code>datetime</code> greater than or equal to the start and less than (but not equal to) the end.
- In general, START of a bucket is always less than END of the bucket. But for the bucket on the two sides of the valid time range, START can be equal to END.



- origin and datetime can be positive or negative as long as it is in the valid range. Errors
 are raised if origin or datetime is outside of the valid range, or if the return value is
 outside of the valid range.
- If the input value is of type TIMESTAMP WITH TIME ZONE OR TIMESTAMP WITH LOCAL TIME ZONE, then a time bucket might cross the daylight saving time boundaries. In this case, the duration of the time bucket is still the same as any other time bucket.
- If origin and datetime are TIMESTAMP WITH TIME ZONE or TIMESTAMP WITH LOCAL TIME ZONE, all arithmetic calculations are based in UTC time.

Examples

The following examples use the NLS_DATE_FORMAT YYYY-MM-DD. Set the date format with ALTER SESSION:

```
ALTER SESSION SET NLS DATE FORMAT='YYYY-MM-DD';
```

Example 1

```
SELECT TIME BUCKET (DATE '2022-06-29', INTERVAL '5' YEAR, DATE '2000-01-01', START);
```

The result is:

```
2020-01-01
```

The 5-year time bucket that contains 2022-06-29 is from 2020-01-01(start) to 2025-01-01(end). The fourth argument START is used, so the start of the time bucket 2020-01-01 is returned.

Example 2

The following two queries are equivalent:

```
SELECT TIME_BUCKET ( DATE \'-2022-06-29', \'P5M', DATE \'-2022-01-01', END );

Or:

SELECT TIME_BUCKET ( DATE \'-2022-06-29', INTERVAL \'5' MONTH, DATE \'-2022-01-01', END);

The result is:

-2022-11-01
```

The 5-month time bucket that contains -2022-06-29 is from 2022-06-01(start) to -2022-11-01(end). The fourth argument END is used, so the end of the time bucket 2022-11-01 is returned.

Example 3

```
SELECT TIME_BUCKET ( DATE '2005-03-10', 'P1Y', DATE '2004-02-29' ON OVERFLOW ERROR );

The result is:

ORA-01839: date not valid for month specified
```



The one-year time bucket that contains '2005-03-10' is from error (or '2005-02-29') (start) to error (or '2006-02-29') (end). Default fourth argument START is used, so the start of the time bucket should be returned which is an error.

Example 4

```
SELECT TIME BUCKET ( DATE '2005-03-10', 'P1Y', DATE '2004-02-29' ON OVERFLOW ROUND );
```

The result is:

2005-02-28

The one-year time bucket that contains '2005-03-10' is from '2005-02-28' (start) to '2006-02-28' (end) since February 29 is rounded to February 28. Default fourth argument START is used, so the start of the time bucket is returned: '2005-02-28'.

Example 5

```
SELECT TIME BUCKET ( DATE '2004-04-02', 'P1Y', DATE '2003-02-28' LAST DAY OF MONTH );
```

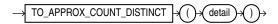
The result is:

2004-02-29

The one-year time bucket that contains \2003-02-28' is from \2004-02-29' (start) to \2005-02-28' (end) since \2004-02-28' is rounded to the last day of that month which is \2004-02-29'. Default fourth argument START is used, so the start of the time bucket is returned: \2004-02-29'.

TO_APPROX_COUNT_DISTINCT

Syntax



Purpose

TO_APPROX_COUNT_DISTINCT takes as its input a detail containing information about an approximate distinct value count, and converts it to a NUMBER value.

For detail, specify a detail of type BLOB, which was created by the APPROX COUNT DISTINCT DETAIL function or the APPROX COUNT DISTINCT AGG function.

See Also:

- APPROX_COUNT_DISTINCT_DETAIL
- TO_APPROX_COUNT_DISTINCT

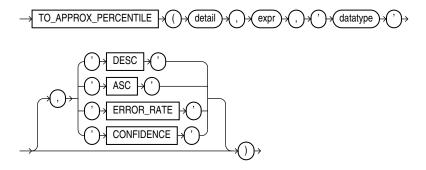


Examples

Refer to TO_APPROX_COUNT_DISTINCT: Examples for examples of using the TO_APPROX_COUNT_DISTINCT function in conjunction with the APPROX_COUNT_DISTINCT_DETAIL and APPROX_COUNT_DISTINCT_AGG functions.

TO APPROX PERCENTILE

Syntax



(datatype::=)

Purpose

TO_APPROX_PERCENTILE takes as its input a detail containing approximate percentile information, a percentile value, and a sort specification, and returns an approximate interpolated value that would fall into that percentile value with respect to the sort specification.

For detail, specify a detail of type BLOB, which was created by the APPROX PERCENTLE DETAIL function or the APPROX PERCENTLE AGG function.

For expr, specify a percentile value, which must evaluate to a numeric value between 0 and 1. If you specify the ERROR_RATE or CONFIDENCE clause, then the percentile value does not apply. In this case, for expr you must specify null or a numeric value between 0 and 1. However, the value will be ignored.

For datatype, specify the data type of the approximate percentile information in the detail. This is the data type of the expression supplied to the <code>APPROX_PERCENTILE_DETAIL</code> function that originated the detail. Valid data types are <code>NUMBER</code>, <code>BINARY_FLOAT</code>, <code>BINARY_DOUBLE</code>, <code>DATE</code>, <code>TIMESTAMP</code>, <code>INTERVAL</code> YEAR TO MONTH, and <code>INTERVAL</code> DAY TO SECOND.

DESC | ASC

Specify the sort specification for the interpolation. Specify DESC for a descending sort order, or ASC for an ascending sort order. ASC is the default.

ERROR_RATE | CONFIDENCE

These clauses let you determine the accuracy of the percentile evaluation of the detail. If you specify one of these clauses, then instead of returning the approximate interpolated value, the function returns a decimal value from 0 to 1, inclusive, which represents one of the following values:



- If you specify ERROR_RATE, then the return value represents the error rate of the percentile evaluation for the detail.
- If you specify CONFIDENCE, then the return value represents the confidence level for the error rate returned when you specify ERROR RATE.

If you specify ERROR RATE or CONFIDENCE, then the percentile value expr is ignored.



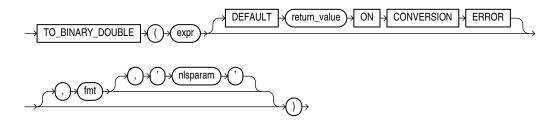
- APPROX_PERCENTILE_DETAIL
- APPROX_PERCENTILE_AGG

Examples

Refer to APPROX_PERCENTILE_AGG: Examples for examples of using the TO_APPROX_PERCENTILE function in conjunction with the APPROX_PERCENTILE_DETAIL and APPROX_PERCENTILE AGG functions.

TO BINARY DOUBLE

Syntax



Purpose

TO BINARY DOUBLE converts *expr* to a double-precision floating-point number.

- expr can be any expression that evaluates to a character string of type CHAR, VARCHAR2, NCHAR, or NVARCHAR2, a numeric value of type NUMBER, BINARY_FLOAT, or BINARY_DOUBLE, BOOLEAN, or null. If expr is BINARY_DOUBLE, then the function returns expr. If expr evaluates to null, then the function returns null. Otherwise, the function converts expr to a BINARY_DOUBLE value.
- The optional DEFAULT return_value on Conversion error clause allows you to specify the value returned by this function if an error occurs while converting expr to BINARY_DOUBLE. This clause has no effect if an error occurs while evaluating expr. The return_value can be an expression or a bind variable, and must evaluate to a character string of type CHAR, VARCHAR2, NCHAR, or NVARCHAR2, a numeric value of type NUMBER, BINARY_FLOAT, or BINARY_DOUBLE, or null. The function converts return_value to BINARY_DOUBLE in the same way it converts expr to BINARY_DOUBLE. If return_value cannot be converted to BINARY_DOUBLE, then the function returns an error.
- The optional 'fmt' and 'nlsparam' arguments serve the same purpose as for the TO_NUMBER function. If you specify these arguments, then expr and return value, if specified, must

each be a character string or null. If either is a character string, then the function uses the fmt and nlsparam arguments to convert the character string to a BINARY_DOUBLE value.

If <code>expr</code> or <code>return_value</code> evaluate to the following character strings, then the function converts them as follows:

- The case-insensitive string 'INF' is converted to positive infinity.
- The case-insensitive string '-INF' is converted to negative identity.
- The case-insensitive string 'NaN' is converted to NaN (not a number).

You cannot use a floating-point number format element (F, f, D, or f) in a character string expr.

Conversions from character strings or NUMBER to BINARY_DOUBLE can be inexact, because the NUMBER and character types use decimal precision to represent the numeric value, and BINARY DOUBLE uses binary precision.

Conversions from BINARY FLOAT to BINARY DOUBLE are exact.

If you specify an expr of type BOOLEAN, then TRUE will be converted to 1 and FALSE will be converted to 0.

```
See Also:
```

TO CHAR (number) and "Floating-Point Numbers"

Examples

The examples that follow are based on a table with three columns, each with a different numeric data type:

The following example converts a value of data type NUMBER to a value of data type BINARY_DOUBLE:

The following example compares extracted dump information from the dec_num and bin double columns:

```
SELECT DUMP(dec_num) "Decimal",
    DUMP(bin double) "Double"
```



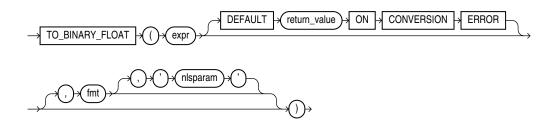
The following example returns the default value of 0 because the specified expression cannot be converted to a BINARY DOUBLE value:

```
SELECT TO_BINARY_DOUBLE('200' DEFAULT 0 ON CONVERSION ERROR) "Value"
FROM DUAL;

Value
-----
0
```

TO_BINARY_FLOAT

Syntax



Purpose

TO BINARY FLOAT converts expr to a single-precision floating-point number.

- expr can be any expression that evaluates to a character string of type CHAR, VARCHAR2,
 NCHAR, or NVARCHAR2, a numeric value of type NUMBER, BINARY_FLOAT, or
 BINARY_DOUBLE,BOOLEAN, or null. If expr is BINARY_FLOAT, then the function returns expr. If
 expr evaluates to null, then the function returns null. Otherwise, the function converts expr
 to a BINARY_FLOAT value.
- The optional DEFAULT return_value on CONVERSION ERROR clause allows you to specify the value returned by this function if an error occurs while converting expr to BINARY_FLOAT. This clause has no effect if an error occurs while evaluating expr. The return_value can be an expression or a bind variable, and must evaluate to a character string of type CHAR, VARCHAR2, NCHAR, or NVARCHAR2, a numeric value of type NUMBER, BINARY_FLOAT, or BINARY_DOUBLE, or null. The function converts return_value to BINARY_FLOAT in the same way it converts expr to BINARY_FLOAT. If return_value cannot be converted to BINARY_FLOAT, then the function returns an error.
- The optional 'fmt' and 'nlsparam' arguments serve the same purpose as for the TO_NUMBER function. If you specify these arguments, then expr and return_value, if specified, must each be a character string or null. If either is a character string, then the function uses the fmt and nlsparam arguments to convert the character string to a BINARY FLOAT value.

If <code>expr</code> or <code>return_value</code> evaluate to the following character strings, then the function converts them as follows:

The case-insensitive string 'INF' is converted to positive infinity.



- The case-insensitive string '-INF' is converted to negative identity.
- The case-insensitive string 'NaN' is converted to NaN (not a number).

You cannot use a floating-point number format element (F, f, D, or d) in a character string expr.

Conversions from character strings or NUMBER to BINARY_FLOAT can be inexact, because the NUMBER and character types use decimal precision to represent the numeric value and BINARY FLOAT uses binary precision.

Conversions from BINARY_DOUBLE to BINARY_FLOAT are inexact if the BINARY_DOUBLE value uses more bits of precision than supported by the BINARY_FLOAT.

If you specify an expr of type BOOLEAN, then TRUE will be converted to 1 and FALSE will be converted to 0.



TO_CHAR (number) and "Floating-Point Numbers "

Examples

Using table $float_point_demo$ created for TO_BINARY_DOUBLE, the following example converts a value of data type NUMBER to a value of data type BINARY_FLOAT:

The following example returns the default value of 0 because the specified expression cannot be converted to a BINARY FLOAT value:

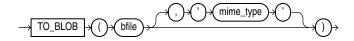
```
SELECT TO_BINARY_FLOAT('200' DEFAULT 0 ON CONVERSION ERROR) "Value"
FROM DUAL;

Value
```

TO_BLOB (bfile)

Syntax

to_blob_bfile::=



Purpose

TO BLOB (bfile) converts a BFILE value to a BLOB value.



For <code>mime_type</code>, specify the MIME type to be set on the <code>BLOB</code> value returned by this function. If you omit <code>mime_type</code>, then a MIME type will not be set on the <code>BLOB</code> value.

Example

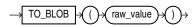
The following hypothetical example returns the BLOB of a BFILE column value $media_col$ in table $media_tab$. It sets the MIME type to JPEG on the resulting BLOB.

```
SELECT TO_BLOB(media_col, 'JPEG') FROM media_tab;
```

TO_BLOB (raw)

Syntax

to blob::=



Purpose



All forms of LONG data types (LONG, LONG RAW, LONG VARCHAR, LONG VARRAW) were deprecated in Oracle8i Release 8.1.6. For succeeding releases, the LONG data type was provided for backward compatibility with existing applications. In new applications developed with later releases, Oracle strongly recommends that you use CLOB and NCLOB data types for large amounts of character data.

TO BLOB (raw) converts LONG RAW and RAW values to BLOB values.

From within a PL/SQL package, you can use TO_BLOB (raw) to convert RAW and BLOB values to BLOB.

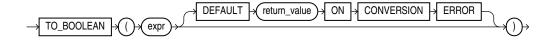
Examples

The following hypothetical example returns the BLOB of a RAW column value:

```
SELECT TO_BLOB(raw_column) blob FROM raw_table;
BLOB
-----
00AADD343CDBBD
```

TO_BOOLEAN

Syntax



Purpose

Use TO_BOOLEAN to explicitly convert character value expressions or numeric value expressions to boolean values.

If *expr* is a string, it must evaluate to the allowed string inputs. See Table 2-6.

expr can take one of the following types, or null:

- A character string of type CHAR, VARCHAR2, NCHAR, NVARCHAR2
- A numeric value of type NUMBER, BINARY FLOAT, or BINARY DOUBLE
- A boolean value of type BOOLEAN.

Examples

```
SELECT TO_BOOLEAN(0), TO_BOOLEAN('true'), TO_BOOLEAN('no');
```

The output is:

The output is:

```
TO_BOOLEAN(
-----
TRUE
```

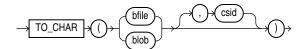
See Also:

- CAST for conversion rules.
- Boolean Data Type for more details on the built-in boolean data type.

TO_CHAR (bfile|blob)

Syntax

to_char_bfile_blob::=





Purpose

TO_CHAR (bfile|blob) converts BFILE or BLOB data to the database character set. The value returned is always VARCHAR2. If the value returned is too large to fit into the VARCHAR2 data type, then the data is truncated.

For csid, specify the character set ID of the BFILE or BLOB data. If the character set of the BFILE or BLOB data is the database character set, then you can specify a value of 0 for csid, or omit csid altogether.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of this function

Example

The following hypothetical example takes as its input a BFILE column media_col in table media_tab, which uses the character set with ID 873. The example returns a VARCHAR2 value that uses the database character set.

SELECT TO CHAR (media col, 873) FROM media tab;

TO_CHAR (boolean)

Syntax



Purpose

Use TO_CHAR (boolean) to explicitly convert a boolean value to a character value of 'TRUE' or 'FALSE'.

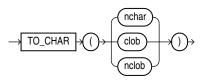
See Also:

- CAST for conversion rules.
- Boolean Data Type for more details on the built-in boolean data type.

TO_CHAR (character)

Syntax

to_char_char::=



Purpose

TO_CHAR (character) converts NCHAR, NVARCHAR2, CLOB, or NCLOB data to the database character set. The value returned is always VARCHAR2.

When you use this function to convert a character LOB into the database character set, if the LOB value to be converted is larger than the target type, then the database returns an error.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following example interprets a simple string as character data:

```
SELECT TO_CHAR('01110') FROM DUAL;

TO_CH
----
01110
```

Compare this example with the first example for TO_CHAR (number) .

The following example converts some CLOB data from the pm.print_media table to the database character set:



TO_CHAR (character) Function: Example

The following statements create a table named <code>empl_temp</code> and populate it with employee details:

```
CREATE TABLE empl temp
     employee id NUMBER(6),
     first_name VARCHAR2(20),
     last_name VARCHAR2(25), email VARCHAR2(25),
     hire date DATE DEFAULT SYSDATE,
     job id VARCHAR2(10),
     clob column CLOB
  );
INSERT INTO empl temp
VALUES (111, 'John', 'Doe', 'example.com', '10-JAN-2015', '1001', 'Experienced
Employee');
INSERT INTO empl temp
VALUES (112, 'John', 'Smith', 'example.com', '12-JAN-2015', '1002', 'Junior
Employee');
INSERT INTO empl temp
VALUES (113, 'Johnnie', 'Smith', 'example.com', '12-JAN-2014', '1002', 'Mid-Career
Employee');
INSERT INTO empl temp
VALUES (115, 'Jane', 'Doe', 'example.com', '15-JAN-2015', '1005', 'Executive
Employee');
```

The following statement converts CLOB data to the database character set:

```
SELECT To_char(clob_column) "CLOB_TO_CHAR"
FROM empl_temp
WHERE employee_id IN ( 111, 112, 115 );

CLOB_TO_CHAR
------
Experienced Employee
Junior Employee
Executive Employee
```

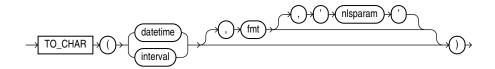
Live SQL:

View and run a related example on Oracle Live SQL at *Using the TO_CHAR Function*

TO_CHAR (datetime)

Syntax

to_char_date::=



Purpose

TO_CHAR (datetime) converts a datetime or interval value of DATE, TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIME ZONE, INTERVAL DAY TO SECOND, or INTERVAL YEAR TO MONTH data type to a value of VARCHAR2 data type in the format specified by the date format fmt. If you omit fmt, then date is converted to a VARCHAR2 value as follows:

- DATE values are converted to values in the default date format.
- TIMESTAMP and TIMESTAMP WITH LOCAL TIME ZONE values are converted to values in the default timestamp format.
- TIMESTAMP WITH TIME ZONE values are converted to values in the default timestamp with time zone format.
- Interval values are converted to the numeric representation of the interval literal.

Refer to "Format Models" for information on datetime formats.

The 'nlsparam' argument specifies the language in which month and day names and abbreviations are returned. This argument can have this form:

```
'NLS DATE LANGUAGE = language'
```

If you omit 'nlsparam', then this function uses the default date language for your session.



"Security Considerations for Data Conversion"

You can use this function in conjunction with any of the XML functions to generate a date in the database format rather than the XML Schema standard format.

See Also:

- Oracle XML DB Developer's Guide for information about formatting of XML dates and timestamps, including examples
- "XML Functions" for a listing of the XML functions
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following example uses this table:

```
CREATE TABLE date_tab (
ts_col TIMESTAMP,
tsltz_col TIMESTAMP WITH LOCAL TIME ZONE,
tstz col TIMESTAMP WITH TIME ZONE);
```

The example shows the results of applying <code>TO_CHAR</code> to different <code>TIMESTAMP</code> data types. The result for a <code>TIMESTAMP</code> WITH <code>LOCAL</code> <code>TIME</code> <code>ZONE</code> column is sensitive to session time zone, whereas the results for the <code>TIMESTAMP</code> and <code>TIMESTAMP</code> WITH <code>TIME</code> <code>ZONE</code> columns are not sensitive to session time zone:

```
ALTER SESSION SET TIME ZONE = '-8:00';
INSERT INTO date tab VALUES (
   TIMESTAMP'1999-12-01 10:00:00',
   TIMESTAMP'1999-12-01 10:00:00',
   TIMESTAMP'1999-12-01 10:00:00');
INSERT INTO date tab VALUES (
  TIMESTAMP'1999-12-02 10:00:00 -8:00',
   TIMESTAMP'1999-12-02 10:00:00 -8:00',
   TIMESTAMP'1999-12-02 10:00:00 -8:00');
SELECT TO CHAR(ts col, 'DD-MON-YYYY HH24:MI:SSxFF') AS ts date,
   TO CHAR(tstz col, 'DD-MON-YYYY HH24:MI:SSxFF TZH:TZM') AS tstz date
   FROM date tab
   ORDER BY ts date, tstz date;
TS DATE
                            TSTZ DATE
01-DEC-1999 10:00:00.000000 01-DEC-1999 10:00:00.000000 -08:00
02-DEC-1999 10:00:00.000000 02-DEC-1999 10:00:00.000000 -08:00
SELECT SESSIONTIMEZONE,
   TO CHAR(tsltz col, 'DD-MON-YYYY HH24:MI:SSxFF') AS tsltz
   FROM date tab
   ORDER BY sessiontimezone, tsltz;
SESSIONTIM TSLTZ
-08:00 01-DEC-1999 10:00:00.000000
-08:00 02-DEC-1999 10:00:00.000000
ALTER SESSION SET TIME ZONE = '-5:00';
SELECT TO CHAR(ts col, 'DD-MON-YYYY HH24:MI:SSxFF') AS ts col,
   TO CHAR(tstz col, 'DD-MON-YYYY HH24:MI:SSxFF TZH:TZM') AS tstz col
   FROM date tab
```



The following example converts an interval literal into a text literal:

```
SELECT TO_CHAR(INTERVAL '123-2' YEAR(3) TO MONTH) FROM DUAL;

TO_CHAR
-----+123-02
```

Using TO_CHAR to Format Dates and Numbers: Example

The following statement converts date values to the format specified in the TO CHAR function:

```
WITH dates AS (

SELECT date'2015-01-01' d FROM dual union

SELECT date'2015-01-10' d FROM dual union

SELECT date'2015-02-01' d FROM dual
)

SELECT d "Original Date",

to_char(d, 'dd-mm-yyyy') "Day-Month-Year",

to_char(d, 'hh24:mi') "Time in 24-hr format",

to_char(d, 'iw-iyyy') "ISO Year and Week of Year"

FROM dates;
```

The following statement converts date and timestamp values to the format specified in the ${\tt TO}$ CHAR function:

```
WITH dates AS (

SELECT date'2015-01-01' d FROM dual union

SELECT date'2015-01-10' d FROM dual union

SELECT date'2015-02-01' d FROM dual union

SELECT timestamp'2015-03-03 23:44:32' d FROM dual union

SELECT timestamp'2015-04-11 12:34:56' d FROM dual
)

SELECT d "Original Date",

to_char(d, 'dd-mm-yyyy') "Day-Month-Year",

to_char(d, 'iw-iyyy') "Time in 24-hr format",

to_char(d, 'iw-iyyy') "ISO Year and Week of Year",

to_char(d, 'Month') "Month Name",

to_char(d, 'Year') "Year"

FROM dates;
```

The following statement extracts the datetime fields specified in the ${\tt EXTRACT}$ function from the input datetime expressions:

```
WITH dates AS (
SELECT date'2015-01-01' d FROM dual union
```

```
SELECT date'2015-01-10' d FROM dual union
SELECT date'2015-02-01' d FROM dual union
SELECT timestamp'2015-03-03 23:44:32' d FROM dual union
SELECT timestamp'2015-04-11 12:34:56' d FROM dual
)
SELECT extract(minute from d) minutes,
extract(hour from d) hours,
extract(day from d) days,
extract(month from d) months,
extract(year from d) years
FROM dates;
```

The following statement displays the input numbers as per the format specified in the TO_CHAR function:

```
WITH nums AS (

SELECT 10 n FROM dual union

SELECT 9.99 n FROM dual union

SELECT 1000000 n FROM dual --one million
)

SELECT n "Input Number N",

to_char(n),

to_char(n, '9,999,999.99') "Number with Commas",

to_char(n, '0,000,000.000') "Zero-padded Number",

to_char(n, '9.9EEEE') "Scientific Notation"

FROM nums;
```

The following statement converts the input numbers as per the format specified in the TO_CHAR function:

```
WITH nums AS (

SELECT 10 n FROM dual union

SELECT 9.99 n FROM dual union

SELECT 1000000 n FROM dual --one million
)

SELECT n "Input Number N",

to_char(n),

to_char(n, '9,999,999.99') "Number with Commas",

to_char(n, '0,000,000.000') "Zero_padded Number",

to_char(n, '9.9EEEE') "Scientific Notation",

to_char(n, '$9,999,990.00') Monetary,

to_char(n, 'X') "Hexadecimal Value"

FROM nums;
```

The following statement converts the input numbers as per the format specified in the ${\tt TO_CHAR}$ function:

```
WITH nums AS (

SELECT 10 n FROM dual union

SELECT 9.99 n FROM dual union

SELECT .99 n FROM dual union

SELECT 1000000 n FROM dual --one million
)

SELECT n "Input Number N",

to_char(n),

to_char(n, '9,999,999.99') "Number with Commas",

to_char(n, '0,000,000.000') "Zero_padded Number",

to_char(n, '9.9EEEE') "Scientific Notation",

to_char(n, '$9,999,990.00') Monetary,

to_char(n, 'XXXXXXX') "Hexadecimal Value"

FROM nums;
```

Live SQL:

View and run a related example on Oracle Live SQL at *Using TO_CHAR to Format Dates and Numbers*

TO_CHAR (datetime) Function: Example

The following statements create a table named <code>empl_temp</code> and populate it with employee details:

```
CREATE TABLE empl temp
     employee id NUMBER(6),
     first name VARCHAR2(20),
     last_name VARCHAR2(25),
     email VARCHAR2(25),
     hire date DATE DEFAULT SYSDATE,
     job id VARCHAR2(10),
     clob column CLOB
  );
INSERT INTO empl temp
VALUES (111, 'John', 'Doe', 'example.com', '10-JAN-2015', '1001', 'Experienced
Employee');
INSERT INTO empl temp
VALUES (112, 'John', 'Smith', 'example.com', '12-JAN-2015', '1002', 'Junior
Employee');
INSERT INTO empl temp
VALUES(113, 'Johnnie', 'Smith', 'example.com', '12-JAN-2014', '1002', 'Mid-Career
Employee');
INSERT INTO empl temp
VALUES(115, 'Jane', 'Doe', 'example.com', '15-JAN-2015', '1005', 'Executive
Employee');
```

The following statement displays dates by using the short and long formats:

```
SELECT hire_date "Default",

TO_CHAR(hire_date,'DS') "Short",

TO_CHAR(hire_date,'DL') "Long"FROM empl_temp

WHERE employee_id IN (111, 112, 115);

Default Short Long

10-JAN-15 1/10/2015 Saturday, January 10, 2015

12-JAN-15 1/12/2015 Monday, January 12, 2015

15-JAN-15 1/15/2015 Thursday, January 15, 2015
```



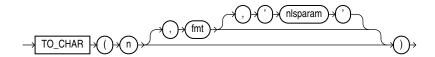


View and run a related example on Oracle Live SQL at *Using the TO_CHAR Function*

TO_CHAR (number)

Syntax

to_char_number::=



Purpose

TO_CHAR (number) converts n to a value of VARCHAR2 data type, using the optional number format fmt. The value n can be of type <code>NUMBER</code>, <code>BINARY_FLOAT</code>, or <code>BINARY_DOUBLE</code>. If you omit fmt, then n is converted to a <code>VARCHAR2</code> value exactly long enough to hold its significant digits.

If n is negative, then the sign is applied after the format is applied. Thus TO_CHAR(-1, '\$9') returns -\$1, rather than \$-1.

Refer to "Format Models" for information on number formats.

The 'nlsparam' argument specifies these characters that are returned by number format elements:

- Decimal character
- Group separator
- Local currency symbol
- International currency symbol

This argument can have this form:

```
'NLS_NUMERIC_CHARACTERS = ''dg''
NLS_CURRENCY = ''text''
NLS ISO CURRENCY = territory '
```

The characters d and g represent the decimal character and group separator, respectively. They must be different single-byte characters. Within the quoted string, you must use two single quotation marks around the parameter values. Ten characters are available for the currency symbol.

If you omit 'nlsparam' or any one of the parameters, then this function uses the default parameter values for your session.

See Also:

- "Security Considerations for Data Conversion"
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following statement uses implicit conversion to combine a string and a number into a number:

```
SELECT TO_CHAR('01110' + 1) FROM DUAL;

TO_C
---
1111
```

Compare this example with the first example for TO_CHAR (character) .

In the next example, the output is blank padded to the left of the currency symbol. In the optional number format fmt, $\[mu]$ designates local currency symbol and $\[mu]$ designates a trailing minus sign. See Table 2-16 for a complete listing of number format elements. The example shows the output in a session in which the session parameter NLS TERRITORY is set to AMERICA.

In the next example, NLS_CURRENCY specifies the string to use as the local currency symbol for the L number format element. NLS_NUMERIC_CHARACTERS specifies comma as the character to use as the decimal separator for the D number format element and period as the character to use as the group separator for the G number format element. These characters are expected in many countries, for example in Germany.

In the next example, NLS_ISO_CURRENCY instructs the database to use the international currency symbol for the territory of POLAND for the C number format element:

```
SELECT TO_CHAR(-10000,'99G999D99C',
   'NLS_NUMERIC_CHARACTERS = '',.''
   NLS_ISO_CURRENCY=POLAND') "Amount"
    FROM DUAL;
Amount
```



```
-10.000,00PLN
```

TO_CHAR (number) Function: Example

The following statements create a table named <code>empl_temp</code> and populate it with employee details:

```
CREATE TABLE empl temp
     employee id NUMBER(6),
     first name VARCHAR2(20),
     last_name VARCHAR2(25),
     email VARCHAR2(25),
     hire_date DATE DEFAULT SYSDATE,
     job id VARCHAR2(10),
     clob column CLOB
  );
INSERT INTO empl temp
VALUES(111, 'John', 'Doe', 'example.com', '10-JAN-2015', '1001', 'Experienced
Employee');
INSERT INTO empl temp
VALUES(112, 'John', 'Smith', 'example.com', '12-JAN-2015', '1002', 'Junior
Employee');
INSERT INTO empl temp
VALUES (113, 'Johnnie', 'Smith', 'example.com', '12-JAN-2014', '1002', 'Mid-Career
Employee');
INSERT INTO empl temp
VALUES (115, 'Jane', 'Doe', 'example.com', '15-JAN-2015', '1005', 'Executive
Employee');
```

The following statement converts numeric data to the database character set:

Live SQL:

View and run a related example on Oracle Live SQL at *Using the TO_CHAR Function*

TO_CLOB (bfile|blob)

Syntax



Purpose

TO_CLOB (bfile|blob) converts BFILE or BLOB data to the database character set and returns the data as a CLOB value.

For csid, specify the character set ID of the BFILE or BLOB data. If the character set of the BFILE or BLOB data is the database character set, then you can specify a value of 0 for csid, or omit csid altogether.

For <code>mime_type</code>, specify the MIME type to be set on the <code>CLOB</code> value returned by this function. If you omit <code>mime_type</code>, then a MIME type will not be set on the <code>CLOB</code> value.

See Also:

Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of this function

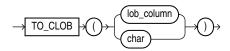
Example

The following hypothetical example returns the CLOB of a BFILE column value docu in table media_tab, which uses the character set with ID 873. It sets the MIME type to text/xml for the resulting CLOB.

SELECT TO CLOB(docu, 873, 'text/xml') FROM media tab;

TO_CLOB (character)

Syntax



Purpose

TO_CLOB (character) converts NCLOB values in a LOB column or other character strings to CLOB values. char can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB.

Oracle Database executes this function by converting the underlying LOB data from the national character set to the database character set.

From within a PL/SQL package, you can use the TO_CLOB (character) function to convert RAW, CHAR, VARCHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB values to CLOB or NCLOB values.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of this function

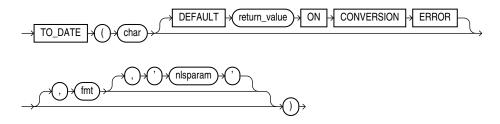
Examples

The following statement converts NCLOB data from the sample pm.print_media table to CLOB and inserts it into a CLOB column, replacing existing data in that column.

```
UPDATE PRINT_MEDIA
SET AD_FINALTEXT = TO_CLOB (AD_FLTEXTN);
```

TO_DATE

Syntax



Purpose

TO_DATE converts char to a value of DATE data type.

For char, you can specify any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type.



This function does not convert data to any of the other datetime data types. For information on other datetime conversions, refer to TO_TIMESTAMP , TO_TIMESTAMP_TZ , TO_DSINTERVAL , and TO_YMINTERVAL .

The optional DEFAULT return_value ON CONVERSION ERROR clause allows you to specify the value this function returns if an error occurs while converting char to DATE. This clause has no effect if an error occurs while evaluating char. The return_value can be an expression or a bind variable, and it must evaluate to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type, or null. The function converts return value to DATE using the same method it uses

to convert char to DATE. If return value cannot be converted to DATE, then the function returns an error.

The fmt is a datetime model format specifying the format of char. If you omit fmt, then char must be in the default date format. The default date format is determined implicitly by the NLS TERRITORY initialization parameter or can be set explicitly by the NLS DATE FORMAT parameter. If fmt is J, for Julian, then char must be an integer.

Caution:

It is good practice always to specify a format mask (fmt) with TO DATE, as shown in the examples in the section that follow, if char is a literal or an expression that evaluates to a known, fixed format, independent of the locale (NLS) configuration of the session. When TO DATE is used without a format mask, the function is valid only if char uses the same format as is determined by the NLS TERRITORY and NLS DATE FORMAT parameters.

However, if *char* corresponds to user input provided by an application, for example, in a bind variable, and the user input is expected to follow the locale (NLS) conventions set for the session provided in the NLS DATE FORMAT parameter, then the format mask should not be specified.

The 'nlsparam' argument specifies the language of the text string that is being converted to a date. This argument can have this form:

```
'NLS DATE LANGUAGE = language'
```

Do not use the TO DATE function with a DATE value for the char argument. The first two digits of the returned DATE value can differ from the original char, depending on fmt or the default date format.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.



See Also:

"Datetime Format Models " and "Data Type Comparison Rules " for more information

Examples

The following example converts a character string into a date:

```
SELECT TO DATE (
  'January 15, 1989, 11:00 A.M.',
  'Month dd, YYYY, HH:MI A.M.',
  'NLS DATE LANGUAGE = American')
   FROM DUAL;
TO DATE ('
15-JAN-89
```



The value returned reflects the default date format if the NLS_TERRITORY parameter is set to 'AMERICA'. Different NLS TERRITORY values result in different default date formats:

```
ALTER SESSION SET NLS_TERRITORY = 'KOREAN';

SELECT TO_DATE(
   'January 15, 1989, 11:00 A.M.',
   'Month dd, YYYY, HH:MI A.M.',
   'NLS_DATE_LANGUAGE = American')
   FROM DUAL;

TO_DATE(
-----89/01/15
```

The following example returns the default value because the specified expression cannot be converted to a DATE value, due to a misspelling of the month:

```
SELECT TO_DATE('Febuary 15, 2016, 11:00 A.M.'

DEFAULT 'January 01, 2016 12:00 A.M.' ON CONVERSION ERROR,

'Month dd, YYYY, HH:MI A.M.') "Value"

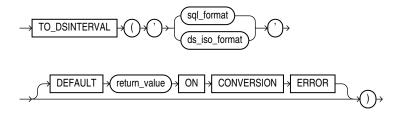
FROM DUAL;

Value

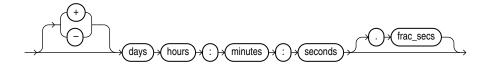
--------
01-JAN-16
```

TO_DSINTERVAL

Syntax

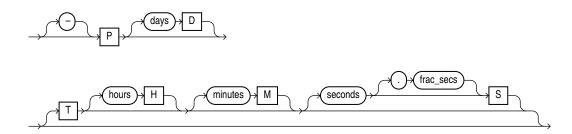


sql_format::=





ds iso format::=



Note:

In earlier releases, the TO_DSINTERVAL function accepted an optional nlsparam clause. This clause is still accepted for backward compatibility, but has no effect.

Purpose

TO DSINTERVAL converts its argument to a value of INTERVAL DAY TO SECOND data type.

For the argument, you can specify any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type.

TO DSINTERVAL accepts argument in one of the two formats:

- SQL interval format compatible with the SQL standard (ISO/IEC 9075)
- ISO duration format compatible with the ISO 8601:2004 standard

In the ISO format, days, hours, minutes and seconds are integers between 0 and 999999999. $frac_secs$ is the fractional part of seconds between .0 and .999999999. No blanks are allowed in the value. If you specify T, then you must specify at least one of the hours, minutes, or seconds values.

The optional DEFAULT $return_value$ ON CONVERSION ERROR clause allows you to specify the value this function returns if an error occurs while converting the argument to an INTERVAL DAY TO SECOND type. This clause has no effect if an error occurs while evaluating the argument. The $return_value$ can be an expression or a bind variable, and it must evaluate to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type. It can be in either the SQL format or ISO format, and need not be in the same format as the function argument. If $return_value$ cannot be converted to an INTERVAL DAY TO SECOND type, then the function returns an error.

Examples

The following example uses the SQL format to select from the hr.employees table the employees who had worked for the company for at least 100 days on November 1, 2002:

SELECT employee_id, last_name FROM employees
WHERE hire date + TO DSINTERVAL('100 00:00:00')

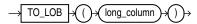


The following example uses the ISO format to display the timestamp 100 days and 5 hours after the beginning of the year 2009:

The following example returns the default value because the specified expression cannot be converted to an INTERVAL DAY TO SECOND value:

TO_LOB

Syntax



Purpose



All forms of LONG data types (LONG, LONG RAW, LONG VARCHAR, LONG VARRAW) were deprecated in Oracle8i Release 8.1.6. For succeeding releases, the LONG data type was provided for backward compatibility with existing applications. In new applications developed with later releases, Oracle strongly recommends that you use CLOB and NCLOB data types for large amounts of character data.

TO_LOB converts LONG or LONG RAW values in the column <code>long_column</code> to LOB values. You can apply this function only to a <code>LONG</code> or <code>LONG</code> RAW column, and only in the select list of a subquery in an <code>INSERT</code> statement.



Before using this function, you must create a LOB column to receive the converted LONG values. To convert LONG values, create a CLOB column. To convert LONG RAW values, create a BLOB column.

You cannot use the TO_LOB function to convert a LONG column to a LOB column in the subquery of a CREATE TABLE ... AS SELECT statement if you are creating an index-organized table. Instead, create the index-organized table without the LONG column, and then use the TO_LOB function in an INSERT ... AS SELECT statement.

You cannot use this function within a PL/SQL package. Instead use the TO_CLOB (character) or TO_BLOB (raw) functions.

See Also:

- the modify_col_properties clause of ALTER TABLE for an alternative method of converting LONG columns to LOB
- INSERT for information on the subquery of an INSERT statement
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

The following syntax shows how to use the ${\tt TO_LOB}$ function on your LONG data in a hypothetical table old table:

```
CREATE TABLE new_table (col1, col2, ... lob_col CLOB);
INSERT INTO new_table (select o.col1, o.col2, ... TO_LOB(o.old_long_col)
    FROM old table o;
```

TO MULTI BYTE

Syntax



Purpose

TO_MULTI_BYTE returns <code>char</code> with all of its single-byte characters converted to their corresponding multibyte characters. <code>char</code> can be of data type <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, or <code>NVARCHAR2</code>. The value returned is in the same data type as <code>char</code>.

Any single-byte characters in *char* that have no multibyte equivalents appear in the output string as single-byte characters. This function is useful only if your database character set contains both single-byte and multibyte characters.

This function does not support CLOB data directly. However, CLOBs can be passed in as arguments through implicit data conversion.

See Also:

- "Data Type Comparison Rules" for more information.
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of TO_MULTI_BYTE

Examples

The following example illustrates converting from a single byte A to a multibyte A in UTF8:

```
SELECT dump(TO_MULTI_BYTE( 'A')) FROM DUAL;

DUMP(TO_MULTI_BYTE('A'))
------
Typ=1 Len=3: 239,188,161
```

TO_NCHAR (boolean)

Syntax



Purpose

Use TO_NCHAR (boolean) to explicitly convert a boolean value to a character value of 'TRUE' or 'FALSE'.

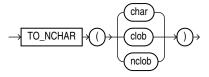
See Also:

- CAST for conversion rules.
- Boolean Data Type for more details on the built-in boolean data type.

TO_NCHAR (character)

Syntax

to_nchar_char::=



Purpose

TO_NCHAR (character) converts a character string, CHAR, VARCHAR2, CLOB, or NCLOB value to the national character set. The value returned is always NVARCHAR2. This function is equivalent to the TRANSLATE ... USING function with a USING clause in the national character set.

See Also:

- "Data Conversion" and TRANSLATE ... USING
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of this function

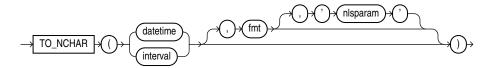
Examples

The following example converts <code>VARCHAR2</code> data from the <code>oe.customers</code> table to the national character set:

TO_NCHAR (datetime)

Syntax

to_nchar_date::=



Purpose

TO_NCHAR (datetime) converts a datetime or interval value of DATE, TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIME ZONE, INTERVAL MONTH TO YEAR, OR INTERVAL DAY TO SECOND data type from the database character set to the national character set.

See Also:

- "Security Considerations for Data Conversion"
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of this function

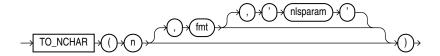
Examples

The following example converts the order_date of all orders whose status is 9 to the national character set:

TO_NCHAR (number)

Syntax

to_nchar_number::=



Purpose

TO_NCHAR (number) converts n to a string in the national character set. The value n can be of type <code>NUMBER</code>, <code>BINARY_FLOAT</code>, or <code>BINARY_DOUBLE</code>. The function returns a value of the same type as the argument. The optional fmt and 'nlsparam' corresponding to n can be of <code>DATE</code>, <code>TIMESTAMP</code>, <code>TIMESTAMP</code> WITH <code>TIME ZONE</code>, <code>TIMESTAMP</code> WITH <code>LOCAL TIME ZONE</code>, <code>INTERVAL MONTH TO YEAR</code>, or <code>INTERVAL DAY TO SECOND data type</code>.

See Also:

- "Security Considerations for Data Conversion"
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

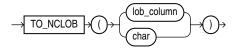
The following example converts the <code>customer_id</code> values from the sample table <code>oe.orders</code> to the national character set:

```
SELECT TO_NCHAR(customer_id) "NCHAR_Customer_ID" FROM orders
WHERE order_status > 9
ORDER BY "NCHAR Customer ID";
```

NCHAR_Customer_ID
102
103
148
148
149

TO_NCLOB

Syntax



Purpose

TO_NCLOB converts CLOB values in a LOB column or other character strings to NCLOB values. *char* can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. Oracle Database implements this function by converting the character set of *char* from the database character set to the national character set.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of this function

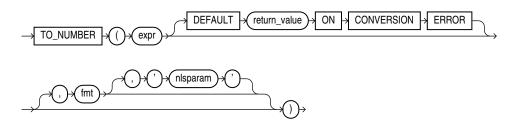
Examples

The following example inserts some character data into an NCLOB column of the pm.print_media table by first converting the data with the TO_NCLOB function:

```
INSERT INTO print_media (product_id, ad_id, ad_fltextn)
    VALUES (3502, 31001,
          TO_NCLOB('Placeholder for new product description'));
```

TO NUMBER

Syntax





Purpose

TO NUMBER converts expr to a value of NUMBER data type.

expr can be any expression that evaluates to a character string of type CHAR, VARCHAR2, NCHAR, or NVARCHAR2, a numeric value of type NUMBER, BINARY_FLOAT, or BINARY_DOUBLE, BOOLEAN, or null. If expr is NUMBER, then the function returns expr. If expr evaluates to null, then the function returns null. Otherwise, the function converts expr to a NUMBER value.

- If you specify an *expr* of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type, then you can optionally specify the format model *fmt*.
- If you specify an *expr* of BINARY_FLOAT or BINARY_DOUBLE data type, then you cannot specify a format model because a float can be interpreted only by its internal representation.
- If you specify an *expr* of type BOOLEAN, then TRUE will be converted to 1 and FALSE will be converted to 0. You cannot specify a format model with inputs of type BOOLEAN.

Refer to "Format Models" for information on number formats.

The 'nlsparam' argument in this function has the same purpose as it does in the TO_CHAR function for number conversions. Refer to TO_CHAR (number) for more information.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.



"Data Type Comparison Rules" for more information.

Examples

The following examples convert character string data into a number:

The following example returns the default value of 0 because the specified expression cannot be converted to a NUMBER value:

```
SELECT TO_NUMBER('2,00' DEFAULT 0 ON CONVERSION ERROR) "Value"
FROM DUAL;

Value
```



TO_SINGLE_BYTE

Syntax



Purpose

TO_SINGLE_BYTE returns char with all of its multibyte characters converted to their corresponding single-byte characters. char can be of data type CHAR, VARCHAR2, NCHAR, or NVARCHAR2. The value returned is in the same data type as char.

Any multibyte characters in *char* that have no single-byte equivalents appear in the output as multibyte characters. This function is useful only if your database character set contains both single-byte and multibyte characters.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.

See Also:

- "Data Type Comparison Rules" for more information.
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of TO_SINGLE_BYTE

Examples

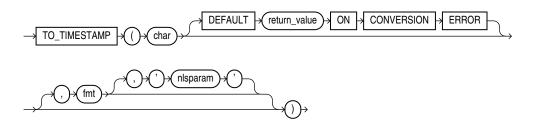
The following example illustrates going from a multibyte A in UTF8 to a single byte ASCII A:

```
SELECT TO_SINGLE_BYTE( CHR(15711393)) FROM DUAL;

T
-
A
```

TO TIMESTAMP

Syntax





Purpose

TO TIMESTAMP converts char to a value of TIMESTAMP data type.

For char, you can specify any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type.

The optional DEFAULT return_value ON CONVERSION ERROR clause allows you to specify the value this function returns if an error occurs while converting char to TIMESTAMP. This clause has no effect if an error occurs while evaluating char. The return_value can be an expression or a bind variable, and it must evaluate to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type, or null. The function converts return_value to TIMESTAMP using the same method it uses to convert char to TIMESTAMP. If return_value cannot be converted to TIMESTAMP, then the function returns an error.

The optional fmt specifies the format of char. If you omit fmt, then char must be in the default format of the TIMESTAMP data type, which is determined by the NLS_TIMESTAMP_FORMAT initialization parameter. The optional 'nlsparam' argument has the same purpose in this function as in the TO CHAR function for date conversion.



Caution:

It is good practice always to specify a format mask (fmt) with TO_TIMESTAMP, as shown in the examples in the section that follow, if char is a literal or an expression that evaluates to a known, fixed format, independent of the locale (NLS) configuration of the session. When TO_TIMESTAMP is used without a format mask, the function is valid only if char uses the same format as is determined by the NLS_TERRITORY and NLS_TIMESTAMP_FORMAT parameters.

However, if char corresponds to user input provided by an application, for example, in a bind variable, and the user input is expected to follow the locale (NLS) conventions set for the session provided in the NLS_TIMESTAMP_FORMAT parameter, then the format mask should not be specified.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.



See Also:

"Data Type Comparison Rules" for more information.

Examples

The following example converts a character string to a timestamp. The character string is not in the default TIMESTAMP format, so the format mask must be specified:

```
SELECT TO_TIMESTAMP ('10-Sep-02 14:10:10.123000', 'DD-Mon-RR HH24:MI:SS.FF')
FROM DUAL;

TO_TIMESTAMP('10-SEP-0214:10:10.123000', 'DD-MON-RRHH24:MI:SS.FF')
```



10-SEP-02 02.10.10.123000000 PM

The following example returns the default value of NULL because the specified expression cannot be converted to a TIMESTAMP value, due to an invalid month specification:

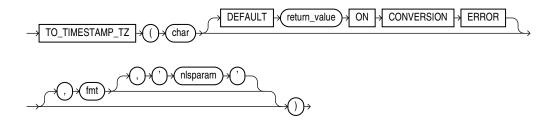
```
SELECT TO_TIMESTAMP ('10-Sept-02 14:10:10.123000'
DEFAULT NULL ON CONVERSION ERROR,
   'DD-Mon-RR HH24:MI:SS.FF',
   'NLS_DATE_LANGUAGE = American') "Value"
   FROM DUAL;
```



NLS_TIMESTAMP_FORMAT initialization parameter for information on the default TIMESTAMP format and "Datetime Format Models" for information on specifying the format mask

TO TIMESTAMP TZ

Syntax



Purpose

TO TIMESTAMP TZ converts char to a value of TIMESTAMP WITH TIME ZONE data type.

For char, you can specify any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type.



This function does not convert character strings to TIMESTAMP WITH LOCAL TIME ZONE. To do this, use a CAST function, as shown in CAST.

The optional DEFAULT return_value ON CONVERSION ERROR clause allows you to specify the value this function returns if an error occurs while converting char to TIMESTAMP WITH TIME ZONE. This clause has no effect if an error occurs while evaluating char. The return_value can be an expression or a bind variable, and it must evaluate to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type, or null. The function converts return_value to TIMESTAMP WITH TIME ZONE using the same method it uses to convert char to TIMESTAMP WITH

TIME ZONE. If return value cannot be converted to TIMESTAMP WITH TIME ZONE, then the function returns an error.

The optional fmt specifies the format of char. If you omit fmt, then char must be in the default format of the TIMESTAMP WITH TIME ZONE data type. The optional 'nlsparam' has the same purpose in this function as in the TO CHAR function for date conversion.



Caution:

It is good practice always to specify a format mask (fmt) with TO TIMESTAMP TZ, as shown in the examples in the section that follow, if char is a literal or an expression that evaluates to a known, fixed format, independent of the locale (NLS) configuration of the session. When to timestamp to is used without a format mask, the function is valid only if char uses the same format as is determined by the NLS TERRITORY and NLS TIMESTAMP TZ FORMAT parameters.

However, if *char* corresponds to user input provided by an application, for example, in a bind variable, and the user input is expected to follow the locale (NLS) conventions set for the session provided in the NLS TIMESTAMP TZ FORMAT parameter, then the format mask should not be specified.

Examples

The following example converts a character string to a value of TIMESTAMP WITH TIME ZONE:

```
SELECT TO TIMESTAMP TZ('1999-12-01 11:00:00 -8:00',
 'YYYY-MM-DD HH:MI:SS TZH:TZM') FROM DUAL;
TO TIMESTAMP TZ('1999-12-0111:00:00-08:00','YYYY-MM-DDHH:MI:SSTZH:TZM')
______
01-DEC-99 11.00.00.00000000 AM -08:00
```

The following example casts a null column in a UNION operation as TIMESTAMP WITH LOCAL TIME ZONE using the sample tables oe.order items and oe.orders:

```
SELECT order id, line item id,
 CAST(NULL AS TIMESTAMP WITH LOCAL TIME ZONE) order date
 FROM order items
UNION
SELECT order id, to number(null), order date
 FROM orders;
 ORDER ID LINE ITEM ID ORDER DATE
     2354
                    1
     2354
                   2
     2354
                    3
     2354
     2354
     2354
     2354
     2354
                    8
     2354
                    9
     2354
                   1.0
     2354
                   11
```



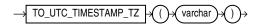
```
2354 12
2354 13
2354 14-JUL-00 05.18.23.234567 PM
2355 1
2355 2
```

The following example returns the default value of NULL because the specified expression cannot be converted to a TIMESTAMP WITH TIME ZONE value, due to an invalid month specification:

```
SELECT TO_TIMESTAMP_TZ('1999-13-01 11:00:00 -8:00' DEFAULT NULL ON CONVERSION ERROR, 'YYYY-MM-DD HH:MI:SS TZH:TZM') "Value" FROM DUAL;
```

TO_UTC_TIMESTAMP TZ

Syntax



Purpose

The SQL function <code>TO_UTC_TIMESTAMP_TZ</code> takes an ISO 8601 date format string as the <code>varchar</code> input and returns an instance of SQL data type <code>TIMESTAMP WITH TIMEZONE</code>. It normalizes the input to UTC time (Coordinated Universal Time, formerly Greenwich Mean Time). Unlike SQL function <code>TO_TIMESTAMP_TZ</code>, the new function assumes that the input string uses the ISO 8601 date format, defaulting the time zone to UTC 0.

A typical use of this function would be to provide its output to SQL function SYS_EXTRACT_UTC, obtaining a UTC time that is then passed as a SQL bind variable to SQL/JSON condition JSON EXISTS, to perform a time-stamp range comparison.

This is the allowed syntax for dates and times:

- Date (only): YYYY-MM-DD
- Date with time: YYYY-MM-DDThh:mm:ss[.s[s[s[s[s]]]]][Z|(+|-)hh:mm]

where:

- YYYY specifies the year, as four decimal digits.
- **MM** specifies the *month*, as two decimal digits, 00 to 12.
- pecifies the day, as two decimal digits, 00 to 31.
- hh specifies the hour, as two decimal digits, 00 to 23.
- mm specifies the minutes, as two decimal digits, 00 to 59.
- ss[.s[s[s[s[s]]]]] specifies the *seconds*, as two decimal digits, 00 to 59, optionally followed by a decimal point and 1 to 6 decimal digits (representing the fractional part of a second).
- z specifies UTC time (time zone 0). (It can also be specified by +00:00, but not by -00:00.)

(+|-) hh:mm specifies the time-zone as difference from UTC. (One of + or - is required.)

For a time value, the time-zone part is optional. If it is absent then UTC time is assumed.

No other ISO 8601 date-time syntax is supported. In particular:

- Negative dates (dates prior to year 1 BCE), which begin with a hyphen (e.g. -2018-10-26T21:32:52), are not supported.
- Hyphen and colon separators are required: so-called "basic" format, YYYYMMDDThhmmss, is not supported.
- Ordinal dates (year plus day of year, calendar week plus day number) are not supported.
- Using more than four digits for the year is not supported.

Supported dates and times include the following:

- 2018-10-26T21:32:52
- 2018-10-26T21:32:52+02:00
- 2018-10-26T19:32:52Z
- 2018-10-26T19:32:52+00:00
- 2018-10-26T21:32:52.12679

Unsupported dates and times include the following:

- 2018-10-26T21:32 (if a time is specified then all of its parts must be present)
- 2018-10-26T25:32:52+02:00 (the hours part, 25, is out of range)
- 18-10-26T21:32 (the year is not specified fully)

Examples



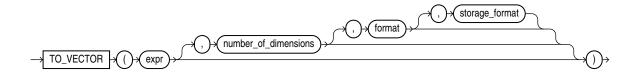
See Also:

- ISO 8601 standard
- ISO 8601 at Wikipedia

TO VECTOR

TO_VECTOR is a constructor that takes a string of type VARCHAR2, CLOB, BLOB, or JSON as input, converts it to a vector, and returns a vector as output. TO_VECTOR also takes another vector as input, adjusts its format, and returns the adjusted vector as output. TO_VECTOR is synonymous with VECTOR.

Syntax



Parameters

- expr must evaluate to one of:
 - A string (of character types or CLOB) that represents a vector.
 - A VECTOR.
 - A BLOB. The BLOB must represent the vector's binary bytes.
 - A JSON array. All elements in the array must be numeric.

If expr is NULL, the result is NULL.

The string representation of the vector must be in the form of an array of non-null numbers enclosed with a bracket and separated by commas, such as [1, 3.4, -05.60, 3e+4]. TO_VECTOR converts a valid string representation of a vector to a vector in the format specified. If no format is specified the default format is used.

- number_of_dimensions must be a numeric value that describes the number of dimensions
 of the vector to construct. The number of dimensions may also be specified as an asterisk
 (*), in which case the dimension is determined by expr.
- format must be one of the following tokens: INT8, FLOAT32, FLOAT64, BINARY, or *. This is the target internal storage format of the vector. If * is used, the format will be FLOAT32.
 - Note that this behavior is different from declaring a vector column. When you declare a column of type VECTOR(3, *), then all inserted vectors will be stored as is without a change in format.
- storage_format must be one of the following tokens: DENSE, SPARSE, or *. If no storage
 format is specified or if * is used, the following will be observed depending on the input
 type:
 - Textual input: the storage format will default to DENSE.



- JSON input: the storage format will default to DENSE.
- VECTOR input: there is no default and the storage format is not changed.
- BLOB input: there is no default and the storage format is not changed.

Examples

```
SELECT TO_VECTOR('[34.6, 77.8]');

TO_VECTOR('[34.6, 77.8]')

[3.45999985E+001, 7.78000031E+001]

SELECT TO_VECTOR('[34.6, 77.8]', 2, FLOAT32);

TO_VECTOR('[34.6, 77.8]', 2, FLOAT32)

[3.45999985E+001, 7.78000031E+001]

SELECT TO_VECTOR('[34.6, 77.8, -89.34]', 3, FLOAT32);

TO_VECTOR('[34.6, 77.8, -89.34]', 3, FLOAT32)

[3.45999985E+001, 7.78000031E+001, -8.93399963E+001]

SELECT TO_VECTOR('[34.6, 77.8, -89.34]', 3, FLOAT32, DENSE);

TO_VECTOR('[34.6, 77.8, -89.34]', 3, FLOAT32, DENSE);

TO_VECTOR('[34.6, 77.8, -89.34]', 3, FLOAT32, DENSE)

[3.45999985E+001, 7.78000031E+001, -8.93399963E+001]
```

Note:

• For applications using Oracle Client libraries prior to 23ai connected to Oracle Database 23ai, use the TO_VECTOR function to insert vector data. For example:

```
INSERT INTO vecTab VALUES(TO_VECTOR('[1.1, 2.9, 3.14]'));
```

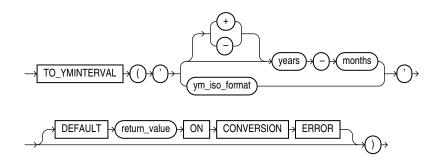
 Applications using Oracle Client 23ai libraries or Thin mode drivers can insert vector data directly as a string or a CLOB. For example:

```
INSERT INTO vecTab VALUES ('[1.1, 2.9, 3.14]');
```

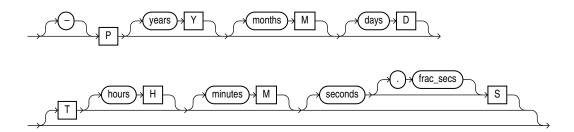


TO_YMINTERVAL

Syntax



ym_iso_format::=



Purpose

TO YMINTERVAL converts its argument to a value of INTERVAL MONTH TO YEAR data type.

For the argument, you can specify any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type.

TO YMINTERVAL accepts argument in one of the two formats:

- SQL interval format compatible with the SQL standard (ISO/IEC 9075)
- ISO duration format compatible with the ISO 8601:2004 standard

In the SQL format, years is an integer between 0 and 99999999, and months is an integer between 0 and 11. Additional blanks are allowed between format elements.

In the ISO format, years and months are integers between 0 and 999999999. Days, hours, minutes, seconds, and frac_secs are non-negative integers, and are ignored, if specified. No blanks are allowed in the value. If you specify T, then you must specify at least one of the hours, minutes, or seconds values.

The optional DEFAULT return_value ON CONVERSION ERROR clause allows you to specify the value this function returns if an error occurs while converting the argument to an INTERVAL MONTH TO YEAR type. This clause has no effect if an error occurs while evaluating the argument. The return_value can be an expression or a bind variable, and it must evaluate to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type. It can be in either the SQL format or ISO format, and need not be in the same format as the function argument. If return_value cannot be converted to an INTERVAL MONTH TO YEAR type, then the function returns an error.



Examples

The following example calculates for each employee in the sample hr.employees table a date one year two months after the hire date:

The following example makes the same calculation using the ISO format:

```
SELECT hire_date, hire_date + TO_YMINTERVAL('P1Y2M') FROM employees;
```

The following example returns the default value because the specified expression cannot be converted to an INTERVAL MONTH TO YEAR value:

TRANSI ATF

Syntax



Purpose

TRANSLATE returns expr with all occurrences of each character in $from_string$ replaced by its corresponding character in to_string . Characters in expr that are not in $from_string$ are not replaced. The argument $from_string$ can contain more characters than to_string . In this case, the extra characters at the end of $from_string$ have no corresponding characters in to_string . If these extra characters appear in expr, then they are removed from the return value.

If a character appears multiple times in <code>from_string</code>, then the <code>to_string</code> mapping corresponding to the first occurrence is used.

You cannot use an empty string for to_string to remove all characters in $from_string$ from the return value. Oracle Database interprets the empty string as null, and if this function has a null argument, then it returns null. To remove all characters in $from_string$, concatenate another character to the beginning of $from_string$ and specify this character as the $to\ string$. For example, TRANSLATE(expr, 'x0123456789', 'x') removes all digits from expr.

TRANSLATE provides functionality related to that provided by the REPLACE function. REPLACE lets you substitute a single string for another single string, as well as remove character strings. TRANSLATE lets you make several single-character, one-to-one substitutions in one operation.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.

See Also:

- "Data Type Comparison Rules" for more information and REPLACE
- Appendix C in Oracle Database Globalization Support Guide for the collation determination rules, which define the collation TRANSLATE uses to compare characters from expr with characters from from_string, and for the collation derivation rules, which define the collation assigned to the character return value of TRANSLATE

Examples

The following statement translates a book title into a string that could be used (for example) as a filename. The $from_string$ contains four characters: a space, asterisk, slash, and apostrophe (with an extra apostrophe as the escape character). The to_string contains only three underscores. This leaves the fourth character in the $from_string$ without a corresponding replacement, so apostrophes are dropped from the returned value.

TRANSLATE ... USING

Syntax



Purpose

TRANSLATE ... USING converts *char* into the character set specified for conversions between the database character set and the national character set.



Note:

The <code>TRANSLATE</code> ... USING function is supported primarily for ANSI compatibility. Oracle recommends that you use the <code>TO_CHAR</code> and <code>TO_NCHAR</code> functions, as appropriate, for converting data to the database or national character set. <code>TO_CHAR</code> and <code>TO_NCHAR</code> can take as arguments a greater variety of data types than <code>TRANSLATE</code> ... <code>USING</code>, which accepts only character data.

The char argument is the expression to be converted.

- Specifying the USING CHAR_CS argument converts char into the database character set. The
 output data type is VARCHAR2.
- Specifying the USING NCHAR_CS argument converts char into the national character set. The output data type is NVARCHAR2.

This function is similar to the Oracle CONVERT function, but must be used instead of CONVERT if either the input or the output data type is being used as NCHAR or NVARCHAR2. If the input contains UCS2 code points or backslash characters (), then use the UNISTR function.

See Also:

- CONVERT and UNISTR
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of TRANSLATE ... USING

Examples

The following statements use data from the sample table <code>oe.product_descriptions</code> to show the use of the <code>TRANSLATE</code> ... <code>USING</code> function:

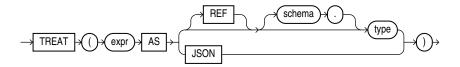
```
CREATE TABLE translate tab (char col VARCHAR2(100),
                       nchar col NVARCHAR2(50));
INSERT INTO translate tab
  SELECT NULL, translated name
    FROM product descriptions
     WHERE product id = 3501;
SELECT * FROM translate tab;
CHAR COL
                NCHAR COL
______
                 C pre SPNIX4.0 - Sys
                 C pro SPNIX4.0 - Sys
                 C til SPNIX4.0 - Sys
                 C voor SPNIX4.0 - Sys
. . .
UPDATE translate tab
  SET char col = TRANSLATE (nchar col USING CHAR CS);
```



```
SELECT * FROM translate tab;
```

TREAT

Syntax



Purpose

You can use the TREAT function to change the declared type of an expression.

Use the keywords AS JSON when you want the expression to return JSON data. This is useful when you want to force some text to be interpreted as JSON data. For example, you can use it to interpret a VARCHAR2 value of {} as an empty JSON object instead of a string.

You must have the EXECUTE object privilege on type to use this function.

- In expr AS JSON, expr is a SQL data type containing JSON, for example CLOB.
- In expr AS type, expr and type must be a user-defined object types, excluding top-level collections.
- *type* must be some supertype or subtype of the declared type of *expr*. If the most specific type of *expr* is *type* (or some subtype of *type*), then TREAT returns *expr*. If the most specific type of *expr* is not *type* (or some subtype of *type*), then TREAT returns NULL.
- You can specify REF only if the declared type of expr is a REF type.
- If the declared type of expr is a REF to a source type of expr, then type must be some subtype or supertype of the source type of expr. If the most specific type of DEREF(expr) is type (or a subtype of type), then TREAT returns expr. If the most specific type of DEREF(expr) is not type (or a subtype of type), then TREAT returns NULL.



"Data Type Comparison Rules " for more information

Examples

The following statement uses the table <code>oe.persons</code>, which is created in "Substitutable Table and Column Examples". The example retrieves the salary attribute of all people in the <code>persons</code> table, the value being null for instances of people that are not employees.



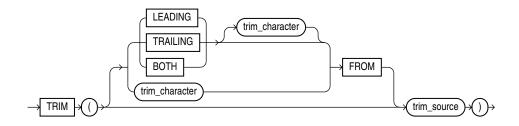
SELECT name, TREAT(VALUE(p) AS employee_t).salary salary
FROM persons p;

NAME	SALARY
Bob	
Joe	100000
Tim	1000

You can use the TREAT function to create an index on the subtype attributes of a substitutable column. For an example, see "Indexing on Substitutable Columns: Examples".

TRIM

Syntax



Purpose

TRIM enables you to trim leading or trailing characters (or both) from a character string. If $trim_character$ or $trim_source$ is a character literal, then you must enclose it in single quotation marks.

- If you specify LEADING, then Oracle Database removes any leading characters equal to trim character.
- If you specify TRAILING, then Oracle removes any trailing characters equal to trim character.
- If you specify BOTH or none of the three, then Oracle removes leading and trailing characters equal to trim character.
- If you do not specify trim character, then the default value is a blank space.
- If you specify only trim source, then Oracle removes leading and trailing blank spaces.
- The function returns a value with data type VARCHAR2. The maximum length of the value is the length of trim source.
- If either trim source or trim character is null, then the TRIM function returns null.

Both trim_character and trim_source can be VARCHAR2 or any data type that can be implicitly converted to VARCHAR2. The string returned is a VARCHAR2 (NVARCHAR2) data type if trim_source is a CHAR or VARCHAR2 (NCHAR or NVARCHAR2) data type, and a CLOB if trim_source is a CLOB data type. The return string is in the same character set as trim source.





Appendix C in *Oracle Database Globalization Support Guide* for the collation determination rules, which define the collation TRIM uses to compare characters from $trim_character$ with characters from $trim_source$, and for the collation derivation rules, which define the collation assigned to the character return value of this function

Examples

This example trims leading zeros from the hire date of the employees in the hr schema:

TRUNC (datetime)

Syntax

trunc datetime::=



Purpose

The TRUNC (datetime) function returns date with the time portion of the day truncated to the unit specified by the format model fmt.

This function is not sensitive to the NLS_CALENDAR session parameter. It operates according to the rules of the Gregorian calendar. The value returned is always of data type DATE, even if you specify a different datetime data type for date. If you do not specify the second argument fmt, then the default format model 'DD' is used and the value returned is date truncated to the day with a time of midnight.

The TRUNC and FLOOR functions are synonymous for dates and timestamps.

Refer to "CEIL, FLOOR, ROUND, and TRUNC Date Functions" for the permitted format models to use in *fmt*.

Examples

The following example truncates a date:



```
SELECT TRUNC(TO_DATE('27-OCT-92','DD-MON-YY'), 'YEAR')
"New Year" FROM DUAL;

New Year
-----01-JAN-92
```

Formatting Dates using TRUNC: Examples

In the following example, the TRUNC function returns the input date with the time portion of the day truncated as specified in the format model:

```
WITH dates AS (

SELECT date'2015-01-01' d FROM dual union

SELECT date'2015-01-10' d FROM dual union

SELECT date'2015-02-01' d FROM dual union

SELECT timestamp'2015-03-03 23:45:00' d FROM dual union

SELECT timestamp'2015-04-11 12:34:56' d FROM dual
)

SELECT d "Original Date",

trunc(d) "Nearest Day, Time Removed",

trunc(d, 'ww') "Nearest Week",

trunc(d, 'iw') "Start of Week",

trunc(d, 'mm') "Start of Month",

trunc(d, 'year') "Start of Year"

FROM dates;
```

In the following example, the input date values are truncated and the ${\tt TO_CHAR}$ function is used to obtain the minute component of the truncated date values:

```
WITH dates AS (

SELECT date'2015-01-01' d FROM dual union

SELECT date'2015-01-10' d FROM dual union

SELECT date'2015-02-01' d FROM dual union

SELECT timestamp'2015-03-03 23:45:00' d FROM dual union

SELECT timestamp'2015-04-11 12:34:56' d FROM dual
)

SELECT d "Original Date",

trunc(d) "Date with Time Removed",

to_char(trunc(d, 'mi'), 'dd-mon-yyyy hh24:mi') "Nearest Minute",

trunc(d, 'iw') "Start of Week",

trunc(d, 'mm') "Start of Year"

FROM dates;
```

The following statement alters the date format for the current session:

```
ALTER SESSION SET nls_date_format = 'dd-mon-yyyy hh24:mi';
```

In the following example, the data is displayed in the new date format:



```
trunc(d, 'year') "Start of Year"
FROM dates;
```

TRUNC (interval)

Syntax



Purpose

TRUNC (interval) returns the interval rounded down to the unit specified by the second argument fmt, the format model .

The absolute value of TRUNC (interval) is never greater than the absolute value of *interval*. The result precision is the same as the input precision, since there is no overflow issue for TRUNC (interval).

For INTERVAL YEAR TO MONTH, fmt can only be year. The default fmt is year.

For INTERVAL DAY TO SECOND, fmt can be day, hour and minute. The default fmt is day. Note that fmt does not support second.



Refer to CEIL, FLOOR, ROUND, and TRUNC Date Functions for the permitted format models to use in fmt.

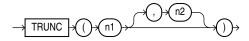
Examples



TRUNC (number)

Syntax

trunc number::=



Purpose

The TRUNC (number) function returns n1 truncated to n2 decimal places. If n2 is omitted, then n1 is truncated to 0 places. n2 can be negative to truncate (make zero) n2 digits left of the decimal point.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. If you omit n2, then the function returns the same data type as the numeric data type of the argument. If you include n2, then the function returns NUMBER.



Table 2-9 for more information on implicit conversion

Examples

The following examples truncate numbers:

```
SELECT TRUNC(15.79,1) "Truncate" FROM DUAL;

Truncate

15.7

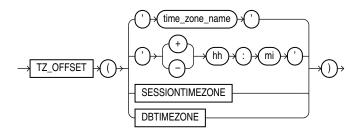
SELECT TRUNC(15.79,-1) "Truncate" FROM DUAL;

Truncate

10
```

TZ_OFFSET

Syntax





Purpose

TZ_OFFSET returns the time zone offset corresponding to the argument based on the date the statement is executed. You can enter a valid time zone region name, a time zone offset from UTC (which simply returns itself), or the keyword SESSIONTIMEZONE or DBTIMEZONE. For a listing of valid values for time_zone_name, query the TZNAME column of the V\$TIMEZONE_NAMES dynamic performance view.



Time zone region names are needed by the daylight saving feature. These names are stored in two types of time zone files: one large and one small. One of these files is the default file, depending on your environment and the release of Oracle Database you are using. For more information regarding time zone files and names, see *Oracle Database Globalization Support Guide*.

See Also:

 Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of TZ_OFFSET

Examples

The following example returns the time zone offset of the US/Eastern time zone from UTC:

```
SELECT TZ_OFFSET('US/Eastern') FROM DUAL;

TZ_OFFS
------
-04:00
```

UID

Syntax



Purpose

UID returns an integer that uniquely identifies the session user (the user who logged on).



USER to learn how Oracle Database determines the session user

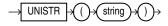
Examples

The following example returns the UID of the session user:

SELECT UID FROM DUAL;

UNISTR

Syntax



Purpose

UNISTR takes as its argument a text literal or an expression that resolves to character data and returns it in the national character set. The national character set of the database can be either AL16UTF16 or UTF8. UNISTR provides support for Unicode string literals by letting you specify the Unicode encoding value of characters in the string. This is useful, for example, for inserting data into NCHAR columns.

The Unicode encoding value has the form '\xxxx' where 'xxxx' is the hexadecimal value of a character in UCS-2 encoding format. Supplementary characters are encoded as two code units, the first from the high-surrogates range (U+D800 to U+DBFF), and the second from the low-surrogates range (U+DC00 to U+DFFF). To include the backslash in the string itself, precede it with another backslash (\\).

For portability and data preservation, Oracle recommends that in the UNISTR string argument you specify only ASCII characters and the Unicode encoding values.

See Also:

- Oracle Database Globalization Support Guide for information on Unicode and national character sets
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of UNISTR

Examples

The following example passes both ASCII characters and Unicode encoding values to the UNISTR function, which returns the string in the national character set:

SELECT UNISTR('abc\00e5\00f1\00f6') FROM DUAL;

UNISTR

abcåñö



UPPER

Syntax



Purpose

UPPER returns *char*, with all letters uppercase. *char* can be any of the data types CHAR, VARCHAR2, NCHAR, NVARCHAR2, CLOB, or NCLOB. The return value is the same data type as *char*. The database sets the case of the characters based on the binary mapping defined for the underlying character set. For linguistic-sensitive uppercase, refer to NLS UPPER.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of UPPER

Examples

The following example returns each employee's last name in uppercase:

```
SELECT UPPER(last_name) "Uppercase"
FROM employees;
```

USER

Syntax



Purpose

USER returns the name of the session user (the user who logged on). This may change during the duration of a database session as Real Application Security sessions are attached or detached. If a Real Application Security session is currently attached to the database session, then it returns user XS\$NULL.

This function returns a VARCHAR2 value.

Oracle Database compares values of this function with blank-padded comparison semantics.

In a distributed SQL statement, the UID and USER functions together identify the user on your local database. You cannot use these functions in the condition of a CHECK constraint.





Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of USER

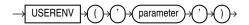
Examples

The following example returns the session user and the user's UID:

SELECT USER, UID FROM DUAL;

USERENV

Syntax



Purpose



USERENV is a legacy function that is retained for backward compatibility. Oracle recommends that you use the SYS_CONTEXT function with the built-in USERENV namespace for current functionality. See SYS_CONTEXT for more information.

USERENV returns information about the current session. This information can be useful for writing an application-specific audit trail table or for determining the language-specific characters currently used by your session. You cannot use USERENV in the condition of a CHECK constraint. Table 7-12 describes the values for the parameter argument.

All calls to userenv return varchar2 data except for calls with the sessionid, sid, and entryid parameters, which return NUMBER.

Table 7-12 Parameters of the USERENV Function

Parameter	Return Value
CLIENT_IN FO	CLIENT_INFO returns up to 64 bytes of user session information that can be stored by an application using the DBMS_APPLICATION_INFO package.
	Caution: Some commercial applications may be using this context value. Refer to the applicable documentation for those applications to determine what restrictions they may impose on use of this context area.
	See Also: Oracle Database Security Guide for more information on application context, CREATE CONTEXT, and SYS_CONTEXT



Table 7-12 (Cont.) Parameters of the USERENV Function

Parameter	Return Value
ENTRYID	The current audit entry number. The audit entryid sequence is shared between fine-grained audit records and regular audit records. You cannot use this attribute in distributed SQL statements.
ISDBA	ISDBA returns 'TRUE' if the user has been authenticated as having DBA privileges either through the operating system or through a password file.
LANG	LANG returns the ISO abbreviation for the language name, a shorter form than the existing 'LANGUAGE' parameter.
LANGUAGE	LANGUAGE returns the language and territory used by the current session along with the database character set in this form:
	language_territory.characterset
SESSIONID	SESSIONID returns the auditing session identifier. You cannot specify this parameter in distributed SQL statements.
SID	SID returns the session ID.
TERMINAL	TERMINAL returns the operating system identifier for the terminal of the current session. In distributed SQL statements, this parameter returns the identifier for your local session. In a distributed environment, this parameter is supported only for remote SELECT statements, not for remote INSERT, UPDATE, or DELETE operations.



Appendix C in *Oracle Database Globalization Support Guide* for the collation derivation rules, which define the collation assigned to the character return value of USERENV

Examples

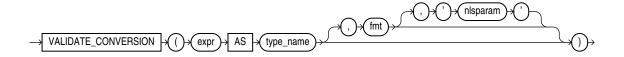
The following example returns the LANGUAGE parameter of the current session:

```
SELECT USERENV('LANGUAGE') "Language" FROM DUAL;

Language
-----
AMERICAN_AMERICA.WE8ISO8859P1
```

VALIDATE_CONVERSION

Syntax



Purpose

VALIDATE_CONVERSION determines whether expr can be converted to the specified data type. If expr can be successfully converted, then this function returns 1; otherwise, this function returns 0. If expr evaluates to null, then this function returns 1. If an error occurs while evaluating expr, then this function returns the error.

For *expr*, specify a SQL expression. The acceptable data types for *expr*, and the purpose of the optional *fmt* and *nlsparam* arguments, depend on the data type you specify for *type name*.

For type_name, specify the data type to which you want to convert expr. You can specify the following data types:

• BINARY DOUBLE

If you specify BINARY_DOUBLE, then <code>expr</code> can be any expression that evaluates to a character string of <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, or <code>NVARCHAR2</code> data type, or a numeric value of type <code>NUMBER</code>, <code>BINARY_FLOAT</code>, or <code>BINARY_DOUBLE</code>. The optional <code>fmt</code> and <code>nlsparam</code> arguments serve the same purpose as for the <code>TO_BINARY_DOUBLE</code> function. Refer to <code>TO_BINARY_DOUBLE</code> for more information.

BINARY FLOAT

If you specify BINARY_FLOAT, then <code>expr</code> can be any expression that evaluates to a character string of <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, or <code>NVARCHAR2</code> data type, or a numeric value of type <code>NUMBER</code>, <code>BINARY_FLOAT</code>, or <code>BINARY_DOUBLE</code>. The optional <code>fmt</code> and <code>nlsparam</code> arguments serve the same purpose as for the <code>TO_BINARY_FLOAT</code> function. Refer to TO <code>BINARY_FLOAT</code> for more information.

DATE

If you specify DATE, then <code>expr</code> can be any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type. The optional <code>fmt</code> and <code>nlsparam</code> arguments serve the same purpose as for the <code>TO_DATE</code> function. Refer to <code>TO_DATE</code> for more information.

INTERVAL DAY TO SECOND

If you specify INTERVAL DAY TO SECOND, then *expr* can be any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type, and must contain a value in either the SQL interval format or the ISO duration format. The optional *fmt* and *nlsparam* arguments do not apply for this data type. Refer to TO_DSINTERVAL for more information on the SQL interval format and the ISO duration format.

• INTERVAL YEAR TO MONTH

If you specify INTERVAL YEAR TO MONTH, then expr can be any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type, and must contain a value in either the SQL interval format or the ISO duration format. The optional fmt and nlsparam arguments do not apply for this data type. Refer to TO_YMINTERVAL for more information on the SQL interval format and the ISO duration format.

NUMBER

If you specify <code>NUMBER</code>, then <code>expr</code> can be any expression that evaluates to a character string of <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, or <code>NVARCHAR2</code> data type, a numeric value of type <code>NUMBER</code>, <code>BINARY_FLOAT</code>, or <code>BINARY_DOUBLE</code> or value of type <code>BOOLEAN</code>. The optional <code>fmt</code> and <code>nlsparam</code> arguments serve the same purpose as for the <code>TO_NUMBER</code> function. Refer to <code>TO_NUMBER</code> for more information.



If expr is a value of type <code>NUMBER</code>, then the <code>VALIDATE_CONVERSION</code> function verifies that expr is a legal numeric value. If expr is not a legal numeric value, then the function returns 0. This enables you to identify corrupt numeric values in your database.

TIMESTAMP

If you specify <code>TIMESTAMP</code>, then <code>expr</code> can be any expression that evaluates to a character string of <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, or <code>NVARCHAR2</code> data type. The optional <code>fmt</code> and <code>nlsparam</code> arguments serve the same purpose as for the <code>TO_TIMESTAMP</code> function. If you omit <code>fmt</code>, then <code>expr</code> must be in the default format of the <code>TIMESTAMP</code> data type, which is determined by the <code>NLS_TIMESTAMP_FORMAT</code> initialization parameter. Refer to <code>TO_TIMESTAMP</code> for more information.

TIMESTAMP WITH TIME ZONE

If you specify TIMESTAMP WITH TIME ZONE, then expr can be any expression that evaluates to a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type. The optional fmt and nlsparam arguments serve the same purpose as for the TO_TIMESTAMP_TZ function. If you omit fmt, then expr must be in the default format of the TIMESTAMP WITH TIME ZONE data type, which is determined by the NLS_TIMESTAMP_TZ_FORMAT initialization parameter. Refer to TO_TIMESTAMP_TZ for more information.

TIMESTAMP WITH LOCAL TIME ZONE

If you specify <code>TIMESTAMP</code>, then <code>expr</code> can be any expression that evaluates to a character string of <code>CHAR</code>, <code>VARCHAR2</code>, <code>NCHAR</code>, or <code>NVARCHAR2</code> data type. The optional <code>fmt</code> and <code>nlsparam</code> arguments serve the same purpose as for the <code>TO_TIMESTAMP</code> function. If you omit <code>fmt</code>, then <code>expr</code> must be in the default format of the <code>TIMESTAMP</code> data type, which is determined by the <code>NLS_TIMESTAMP_FORMAT</code> initialization parameter. Refer to <code>TO_TIMESTAMP</code> for more information.

BOOLEAN

BOOLEAN is supported as a target type. It supports NUMBER type family, VARCHAR type family, and BOOLEAN itself as input.

Examples

In each of the following statements, the specified value can be successfully converted to the specified data type. Therefore, each of these statements returns a value of 1.

```
SELECT VALIDATE_CONVERSION(1000 AS BINARY_DOUBLE)
FROM DUAL;

SELECT VALIDATE_CONVERSION('1234.56' AS BINARY_FLOAT)
FROM DUAL;

SELECT VALIDATE_CONVERSION('July 20, 1969, 20:18' AS DATE,
    'Month dd, YYYY, HH24:MI', 'NLS_DATE_LANGUAGE = American')
FROM DUAL;

SELECT VALIDATE_CONVERSION('200 00:00' AS INTERVAL DAY TO SECOND)
FROM DUAL;

SELECT VALIDATE_CONVERSION('P1Y2M' AS INTERVAL YEAR TO MONTH)
FROM DUAL;

SELECT VALIDATE_CONVERSION('$100,00' AS NUMBER,
    '$999D99', 'NLS_NUMERIC_CHARACTERS = '',.''')
FROM DUAL;
```



The following statement returns 0, because the specified value cannot be converted to BINARY FLOAT:

```
SELECT VALIDATE_CONVERSION('$29.99' AS BINARY_FLOAT)
FROM DUAL;
```

The following statement returns 1, because the specified number format model enables the value to be converted to BINARY FLOAT:

```
SELECT VALIDATE_CONVERSION('$29.99' AS BINARY_FLOAT, '$99D99')
FROM DUAL;
```



Syntax



Purpose

VALUE takes as its argument a correlation variable (table alias) associated with a row of an object table and returns object instances stored in the object table. The type of the object instances is the same type as the object table.

Examples

The following example uses the sample table oe.persons, which is created in "Substitutable Table and Column Examples":



"IS OF *type* Condition " for information on using IS OF type conditions with the VALUE function



VAR_POP

Syntax





"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

VAR_POP returns the population variance of a set of numbers after discarding the nulls in this set. You can use it as both an aggregate and analytic function.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.

See Also:

Table 2-9 for more information on implicit conversion

If the function is applied to an empty set, then it returns null. The function makes the following calculation:

 $SUM((expr - (SUM(expr) / COUNT(expr)))^2) / COUNT(expr)$

See Also:

"About SQL Expressions " for information on valid forms of $\it expr$ and "Aggregate Functions "

Aggregate Example

The following example returns the population variance of the salaries in the employees table:



Analytic Example

The following example calculates the cumulative population and sample variances in the sh.sales table of the monthly sales in 1998:

```
SELECT t.calendar month desc,
  VAR POP(SUM(s.amount sold))
     OVER (ORDER BY t.calendar month desc) "Var Pop",
   VAR SAMP(SUM(s.amount sold))
     OVER (ORDER BY t.calendar month_desc) "Var_Samp"
  FROM sales s, times t
  WHERE s.time_id = t.time_id AND t.calendar_year = 1998
  GROUP BY t.calendar_month_desc
  ORDER BY t.calendar month desc, "Var Pop", "Var Samp";
CALENDAR Var Pop Var Samp
1998-01
                 Ω
1998-02 2269111326 4538222653
1998-03 5.5849E+10 8.3774E+10
1998-04 4.8252E+10 6.4336E+10
1998-05 6.0020E+10 7.5025E+10
1998-06 5.4091E+10 6.4909E+10
1998-07 4.7150E+10 5.5009E+10
1998-08 4.1345E+10 4.7252E+10
1998-09 3.9591E+10 4.4540E+10
1998-10 3.9995E+10 4.4439E+10
1998-11 3.6870E+10 4.0558E+10
1998-12 4.0216E+10 4.3872E+10
```

VAR_SAMP

Syntax





"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

VAR_SAMP returns the sample variance of a set of numbers after discarding the nulls in this set. You can use it as both an aggregate and analytic function.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.



See Also:

Table 2-9 for more information on implicit conversion

If the function is applied to an empty set, then it returns null. The function makes the following calculation:

```
(SUM(expr - (SUM(expr) / COUNT(expr)))^2) / (COUNT(expr) - 1)
```

This function is similar to VARIANCE, except that given an input set of one element, VARIANCE returns 0 and VAR SAMP returns null.



"About SQL Expressions" for information on valid forms of expr and "Aggregate Functions"

Aggregate Example

The following example returns the sample variance of the salaries in the sample employees table.

```
SELECT VAR_SAMP(salary) FROM employees;

VAR_SAMP(SALARY)

______

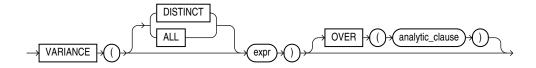
15284813.7
```

Analytic Example

Refer to the analytic example for VAR_POP.

VARIANCE

Syntax





"Analytic Functions" for information on syntax, semantics, and restrictions

Purpose

VARIANCE returns the variance of expr. You can use it as an aggregate or analytic function.



Oracle Database calculates the variance of expr as follows:

- 0 if the number of rows in expr = 1
- VAR SAMP if the number of rows in expr > 1

If you specify <code>DISTINCT</code>, then you can specify only the <code>query_partition_clause</code> of the <code>analytic_clause</code>. The <code>order_by_clause</code> and <code>windowing_clause</code> are not allowed.

This function takes as an argument any numeric data type or any nonnumeric data type that can be implicitly converted to a numeric data type. The function returns the same data type as the numeric data type of the argument.



Table 2-9 for more information on implicit conversion, "About SQL Expressions" for information on valid forms of expr and "Aggregate Functions"

Aggregate Example

The following example calculates the variance of all salaries in the sample employees table:

```
SELECT VARIANCE(salary) "Variance"
FROM employees;

Variance
-----
15283140.5
```

Analytic Example

The following example returns the cumulative variance of salary values in Department 30 ordered by hire date.

```
SELECT last_name, salary, VARIANCE(salary)
        OVER (ORDER BY hire_date) "Variance"
    FROM employees
    WHERE department_id = 30
    ORDER BY last_name, salary, "Variance";
```

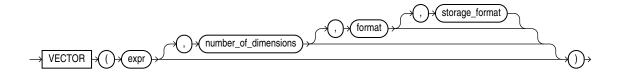
LAST_NAME	SALARY	Variance
Baida	2900	16283333.3
Colmenares	2500	11307000
Himuro	2600	13317000
Khoo	3100	31205000
Raphaely	11000	0
Tobias	2800	21623333.3



VECTOR

VECTOR is synonymous with TO VECTOR.

Syntax



Purpose

See TO_VECTOR for semantics and examples.



Applications using Oracle Client 23ai libraries or Thin mode drivers can insert vector data directly as a string or a CLOB. For example:

```
INSERT INTO vecTab VALUES ('[1.1, 2.9, 3.14]');
```

Examples

```
SELECT VECTOR('[34.6, 77.8]');

VECTOR('[34.6, 77.8]')

[3.45999985E+001, 7.78000031E+001]

SELECT VECTOR('[34.6, 77.8]', 2, FLOAT32);

VECTOR('[34.6, 77.8]', 2, FLOAT32)

[3.45999985E+001, 7.78000031E+001]

SELECT VECTOR('[34.6, 77.8, -89.34]', 3, FLOAT32);

VECTOR('[34.6, 77.8, -89.34]', 3, FLOAT32)

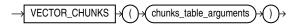
[3.45999985E+001, 7.78000031E+001, -8.93399963E+001]
```



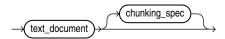
VECTOR_CHUNKS

Use <code>VECTOR_CHUNKS</code> to split plain text into smaller chunks to generate vector embeddings that can be used with vector indexes or hybrid vector indexes.

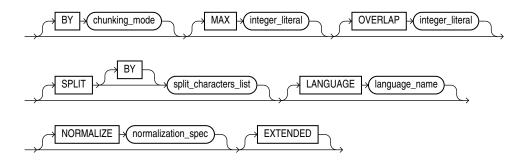
Syntax



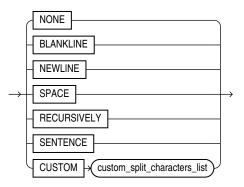
chunks_table_arguments::=



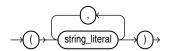
chunking_spec::=



split_characters_list::=

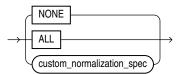


custom_split_characters_list





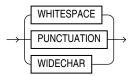
normalization_spec



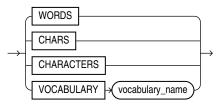
custom_normalization_spec



normalization mode



chunking_mode::=



Purpose

VECTOR_CHUNKS takes a character value as the <code>text_document</code> argument and splits it into chunks using a process controlled by the chunking parameters given in the optional <code>chunking_spec</code>. The chunks are returned as rows of a virtual relational table. Therefore, <code>VECTOR_CHUNKS</code> can only appear in the <code>FROM_clause</code> of a subquery.

The returned virtual table has the following columns:

- CHUNK_OFFSET of data type NUMBER is the position of each chunk in the source document, relative to the start of the document, which has a position of 1.
- CHUNK LENGTH of data type NUMBER is the length of each chunk.
- CHUNK TEXT is a segment of text that has been split off from text document.

The data type of the CHUNK_TEXT column and the length unit used by the values of CHUNK_OFFSET and CHUNK_LENGTH depend on the data type of text_document as listed in the following table:



Table 7-13 Input and Output Data Type Details

Input Data Type	Output Data Type	Offset and Length Unit
VARCHAR2	VARCHAR2	byte
CHAR	VARCHAR2	byte
CLOB	VARCHAR2	character
NVARCHAR2	NVARCHAR2	byte
NCHAR	NVARCHAR2	byte
NCLOB	NVARCHAR2	character

Note:

- For more information about data types, see *Data Types* in the SQL Reference Manual.
- The VARCHAR2 input data type is limited to 4000 bytes unless the MAX_STRING_SIZE parameter is set to EXTENDED, which increases the limit to 32767.

Parameters

All chunking parameters are optional, and the default chunking specifications are automatically applied to your chunk data.

When specifying chunking parameters for this API, ensure that you provide these parameters only in the listed order.



Table 7-14 Chunking Parameters Table

Parameter Description and Acceptable Values

Specifies the mode for splitting your data, that is, to split by counting the number of characters, words, or vocabulary tokens.

Valid values:

CHARACTERS (or CHARS):

Splits by counting the number of characters.

WORDS:

Splits by counting the number of words.

Words are defined as sequences of alphabetic characters, sequences of digits, individual punctuation marks, or symbols. For segmented languages without whitespace word boundaries (such as Chinese, Japanese, or Thai), each native character is considered a word (that is, unigram).

VOCABULARY:

Splits by counting the number of vocabulary tokens.

Vocabulary tokens are words or word pieces, recognized by the vocabulary of the tokenizer that your embedding model uses. You can load your vocabulary file using the <code>VECTOR_CHUNKS</code> helper API <code>DBMS_VECTOR_CHAIN.CREATE_VOCABULARY</code>.

Note: For accurate results, ensure that the chosen model matches the vocabulary file used for chunking. If you are not using a vocabulary file, then ensure that the input length is defined within the token limits of your model.

Default value: WORDS

MAX

ΒY

Specifies a limit on the maximum size of each chunk. This setting splits the input text at a fixed point where the maximum limit occurs in the larger text. The units of MAX correspond to the BY mode, that is, to split data when it reaches the maximum size limit of a certain number of characters, words, numbers, punctuation marks, or vocabulary tokens.

Valid values:

BY CHARACTERS: 50 to 4000 characters

BY WORDS: 10 to 1000 words

BY VOCABULARY: 10 to 1000 tokens

Default value: 100



Table 7-14 (Cont.) Chunking Parameters Table

Parameter Description and Acceptable Values

SPLIT [BY]

Specifies where to split the input text when it reaches the maximum size limit. This helps to keep related data together by defining appropriate boundaries for chunks.

Valid values:

NONE:

Splits at the MAX limit of characters, words, or vocabulary tokens.

NEWLINE, BLANKLINE, and SPACE:

These are single-split character conditions that split at the last split character before the MAX value. Use NEWLINE to split at the end of a line of text. Use BLANKLINE to split at the end of a blank line (sequence of characters, such as two newlines). Use SPACE to split at the end of a blank space.

RECURSIVELY:

This is a multiple-split character condition that breaks the input text using an ordered list of characters (or sequences).

RECURSIVELY is predefined as BLANKLINE, NEWLINE, SPACE, NONE in this order:

- 1. If the input text is more than the MAX value, then split by the first split character.
- 2. If that fails, then split by the second split character.
- 3. And so on.
- 4. If no split characters exist, then split by MAX wherever it appears in the text.
- SENTENCE:

This is an end-of-sentence split condition that breaks the input text at a sentence boundary.

This condition automatically determines sentence boundaries by using knowledge of the input language's sentence punctuation and contextual rules. This language-specific condition relies mostly on end-of-sentence (EOS) punctuations and common abbreviations.

Contextual rules are based on word information, so this condition is only valid when splitting the text by words or vocabulary (not by characters).

Note: This condition obeys the BY WORD and MAX settings, and thus may not determine accurate sentence boundaries in some cases. For example, when a sentence is larger than the MAX value, it splits the sentence at MAX. Similarly, it includes multiple sentences in the text only when they fit within the MAX limit.

CUSTOM:

Splits based on a custom list of characters strings, for example, markup tags. You can provide custom sequences up to a limit of 16 split character strings, with a maximum length of 10 bytes each.

Provide valid text literals as follows:

VECTOR_CHUNKS(c. doc, BY character SPLIT CUSTOM ('<html>' , '</html>')) vc

Default value: RECURSIVELY

OVERLAP

Specifies the amount (as a positive integer literal or zero) of the preceding text that the chunk should contain, if any. This helps in logically splitting up related text (such as a sentence) by including some amount of the preceding chunk text.

The amount of overlap depends on how the maximum size of the chunk is measured (in characters, words, or vocabulary tokens). The overlap begins at the specified SPLIT condition (for example, at NEWLINE).

Valid value: 5% to 20% of MAX

Default value: 0

Table 7-14 (Cont.) Chunking Parameters Table

Parameter Description and Acceptable Values

LANGUAGE

Specifies the language of your input data.

This clause is important, especially when your text contains certain characters (for example, punctuations or abbreviations) that may be interpreted differently in another language.

Valid values:

- NLS-supported language name or its abbreviation, as listed in Oracle Database Globalization Support Guide.
- Custom language name or its abbreviation, as listed in Supported Languages and Data File
 Locations. You use the DBMS_VECTOR_CHAIN.CREATE_LANG_DATA chunker helper API to load
 language-specific data (abbreviation tokens) into the database, for your specified language.

You must use double quotation marks (") for any language name with spaces. For example:

LANGUAGE "simplified chinese"

For one-word language names, quotation marks are not needed. For example:

LANGUAGE american

Default value: NLS LANGUAGE from session

NORMALIZE

Automatically pre-processes or post-processes issues (such as multiple consecutive spaces and smart quotes) that may arise when documents are converted into text. Oracle recommends you to use a normalization mode to extract high-quality chunks.

Valid values:

NONE:

Applies no normalization.

ALL:

Normalizes multi-byte (Unicode) punctuation to standard single-byte.

- Applies all supported normalization modes: PUNCTUATION, WHITESPACE, and WIDECHAR.
 - PUNCTUATION:

Converts quotes, dashes, and other punctuation characters supported in the character set of the text to their common ASCII form. For example:

- * U+2013 (En Dash) maps to U+002D (Hyphen-Minus)
- * U+2018 (Left Single Quotation Mark) maps to U+0027 (Apostrophe)
- * U+2019 (Right Single Quotation Mark) maps to U+0027 (Apostrophe)
- * U+201B (Single High-Reversed-9 Quotation Mark) maps to U+0027 (Apostrophe)
- WHITESPACE:

Minimizes whitespace by eliminating unnecessary characters.

For example, retain blanklines, but remove any extra newlines and interspersed spaces or tabs: " \n " => " \n ""

WIDECHAR:

Normalizes wide, multi-byte digits and (a-z) letters to single-byte.

These are multi-byte equivalents for 0-9 and a-z A-Z, which can show up in Chinese, Japanese, or Korean text.

Default value: NONE

EXTENDED

Increases the output limit of a VARCHAR2 string to 32767 bytes, without requiring you to set the MAX STRING SIZE parameter to EXTENDED.

If EXTENDED is present in <code>chunking_spec</code>, the maximum length of a <code>CHUNK_TEXT</code> column value is 32767 bytes. If it is absent, the maximum length is 4000 bytes if <code>MAX_STRING_SIZE</code> is set to <code>STANDARD</code> and 32767 bytes if <code>MAX_STRING_SIZE</code> is set to <code>EXTENDED</code>.

Examples

VECTOR_CHUNKS can be called for a single character value provided in a character literal or a bind variable as shown in the following example:

```
COLUMN chunk_offset HEADING Offset FORMAT 999

COLUMN chunk_length HEADING Len FORMAT 999

COLUMN chunk_text HEADING Text FORMAT a60

VARIABLE txt VARCHAR2(4000)

EXECUTE :txt := 'An example text value to split with VECTOR_CHUNKS, having over 10 words because the minimum MAX value is 10';

SELECT * FROM VECTOR_CHUNKS(:txt BY WORDS MAX 10);

SELECT * FROM VECTOR_CHUNKS('Another example text value to split with VECTOR_CHUNKS, having over 10 words because the minimum MAX value is 10' BY WORDS MAX 10);
```

To chunk values of a table column, the table needs to be joined with the VECTOR_CHUNKS call using left correlation as shown in the following example:

```
CREATE TABLE documentation_tab (
 id NUMBER,
  text VARCHAR2(2000));
INSERT INTO documentation tab
   VALUES(1, 'sample');
COMMIT;
SET LINESIZE 100;
SET PAGESIZE 20;
COLUMN pos FORMAT 999;
COLUMN siz FORMAT 999;
COLUMN txt FORMAT a60;
PROMPT SQL VECTOR CHUNKS
SELECT D.id id, C.chunk offset pos, C.chunk length siz, C.chunk text txt
FROM documentation tab D, VECTOR CHUNKS (D.text
                                  BY words
                                  MAX 200
                                  OVERLAP 10
                                  SPLIT BY recursively
                                  LANGUAGE american
                                  NORMALIZE all) C;
```

See Also:

- For a complete set of examples on each of the chunking parameters listed in the preceding table, see *Explore Chunking Techniques and Examples of the Al Vector Search User's Guide.*
- To run an end-to-end example scenario using this function, see Convert Text to Chunks With Custom Chunking Specifications of the AI Vector Search User's Guide.

VECTOR_DISTANCE

VECTOR_DISTANCE is the main function that you can use to calculate the distance between two vectors.

Syntax



Purpose

VECTOR_DISTANCE takes two vectors as parameters. You can optionally specify a distance metric to calculate the distance. If you do not specify a distance metric, then the default distance metric is cosine. If the input vectors are BINARY vectors, the default metric is hamming.

You can optionally use the following shorthand vector distance functions:

- L1 DISTANCE
- L2 DISTANCE
- COSINE DISTANCE
- INNER PRODUCT
- HAMMING DISTANCE
- JACCARD DISTANCE

All the vector distance functions take two vectors as input and return the distance between them as a BINARY_DOUBLE.

Note the following caveats:

- If you specify a metric as the third argument, then that metric is used.
- If you do not specify a metric, then the following rules apply:
 - If there is a single column referenced in expr1 and expr2 as in:
 VECTOR_DISTANCE (vec1, :bind), and if there is a vector index defined on vec1, then the metric used when defining the vector index is used.

If no vector index is defined on vec1, then the COSINE metric is used.

- If there are multiple columns referenced in expr1 and expr2 as in:
 VECTOR_DISTANCE(vec1, vec2), or VECTOR_DISTANCE(vec1+vec2, :bind), then for all indexed columns, if their metrics used in the definitions of the indexes are the same, then that metric is used.
 - On the other hand, if the indexed columns do not have a common metric, or none of the columns have an index defined, then the ${\tt COSINE}$ metric is used.
- In a similarity search query, if *expr1* or *expr2* reference an indexed column and you specify a distance metric that conflicts with the metric specified in the vector index, then the vector index is not used and the metric you specified is used to perform an exact search.
- Approximate (index-based) searches can be done if only one column is referenced by either expr1 or expr2, and this column has a vector index defined, and the metric that is



specified in the vector_distance matches the metric used in the definition of the vector index.

Parameters

 expr1 and expr2 must evaluate to vectors and have the same format and number of dimensions.

If you use <code>JACCARD_DISTANCE</code> or the <code>JACCARD</code> metric, then <code>expr1</code> and <code>expr2</code> must evaluate to <code>BINARY</code> vectors.

- This function returns NULL if either expr1 or expr2 is NULL.
- metric must be one of the following tokens:
 - COSINE metric is the default metric. It calculates the cosine distance between two vectors.
 - DOT metric calculates the negated dot product of two vectors. The INNER_PRODUCT function calculates the dot product, as in the negation of this metric.
 - EUCLIDEAN metric, also known as L2 distance, calculates the Euclidean distance between two vectors.
 - EUCLIDEAN_SQUARED metric, also called L2_SQUARED, is the Euclidean distance without taking the square root.
 - HAMMING metric calculates the hamming distance between two vectors by counting the number dimensions that differ between the two vectors.
 - MANHATTAN metric, also known as L1 distance or taxicab distance, calculates the Manhattan distance between two vectors.
 - JACCARD metric calculates the Jaccard distance. The two vectors used in the query must be BINARY vectors.

Shorthand Operators for Distances

Syntax

- expr1 <-> expr2
 - <-> is the Euclidean distance operator: expr1 <-> expr2 is equivalent to
 L2 DISTANCE(expr1, expr2) or VECTOR DISTANCE(expr1, expr2, EUCLIDEAN)
- expr1 <=> expr2
 - <=> is the cosine distance operator: expr1 <=> expr2 is equivalent to
 COSINE DISTANCE(expr1, expr2) or VECTOR DISTANCE(expr1, expr2, COSINE)
- expr1 <#> expr2
 - <#> is the negative dot product operator: expr1 <#> expr2 is equivalent to
 -1*INNER PRODUCT(expr1, expr2) or VECTOR DISTANCE(expr1, expr2, DOT)

Examples Using Shorthand Operators for Distances

```
'[1, 2]' <-> '[0,1]'
v1 <-> '[' || '1,2,3' || ']' is equivalent to v1 <-> '[1, 2, 3]'
v1 <-> '[1,2]' is equivalent to L2_DISTANCE(v1, '[1,2]')
v1 <=> v2 is equivalent to COSINE DISTANCE(v1, v2)
```



```
v1 <#> v2 is equivalent to -1*INNER PRODUCT(v1, v2)
```

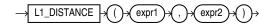
Examples

```
VECTOR DISTANCE with metric EUCLIDEAN is equivalent to L2 DISTANCE:
VECTOR_DISTANCE(expr1, expr2, EUCLIDEAN);
L2 DISTANCE(expr1, expr2);
VECTOR DISTANCE with metric COSINE is equivalent to COSINE DISTANCE:
VECTOR DISTANCE (expr1, expr2, COSINE);
COSINE DISTANCE (expr1, expr2);
VECTOR DISTANCE with metric DOT is equivalent to -1 * INNER PRODUCT:
VECTOR DISTANCE (expr1, expr2, DOT);
-1*INNER PRODUCT(expr1, expr2);
VECTOR DISTANCE with metric Manhattan is equivalent to L1 DISTANCE:
VECTOR DISTANCE (expr1, expr2, MANHATTAN);
L1_DISTANCE(expr1, expr2);
VECTOR DISTANCE with metric HAMMING is equivalent to HAMMING DISTANCE:
VECTOR_DISTANCE(expr1, expr2, HAMMING);
HAMMING DISTANCE (expr1, expr2);
VECTOR DISTANCE with metric JACCARD is equivalent to JACCARD DISTANCE:
VECTOR DISTANCE (expr1, expr2, JACCARD);
JACCARD DISTANCE(expr1, expr2);
```

L1_DISTANCE

L1_DISTANCE is a shorthand version of the VECTOR_DISTANCE function that calculates the distance between two vectors. It takes two vectors as input and returns the distance between them as a BINARY_DOUBLE.

Syntax



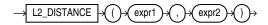
Parameters

- expr1 and expr2 must evaluate to vectors and have the same format and number of dimensions.
- L1 DISTANCE returns NULL, if either expr1 or expr2 is NULL.

L2_DISTANCE

L2_DISTANCE is a shorthand version of the <code>VECTOR_DISTANCE</code> function that calculates the distance between two vectors. It takes two vectors as input and returns the distance between them as a <code>BINARY DOUBLE</code>.

Syntax



Parameters

- expr1 and expr2 must evaluate to vectors that have the same format and number of dimensions.
- L2 DISTANCE returns NULL, if either expr1 or expr2 is NULL.

COSINE_DISTANCE

COSINE_DISTANCE is a shorthand version of the VECTOR_DISTANCE function that calculates the distance between two vectors. It takes two vectors as input and returns the distance between them as a BINARY DOUBLE.

Syntax



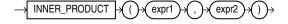
Parameters

- expr1 and expr2 must evaluate to vectors that have the same format and number of dimensions.
- COSINE DISTANCE returns NULL, if either expr1 or expr2 is NULL.

INNER_PRODUCT

INNER_PRODUCT calculates the inner product of two vectors. It takes two vectors as input and returns the inner product as a BINARY_DOUBLE. INNER_PRODUCT (<expr1>, <expr2>) is equivalent to -1 * VECTOR DISTANCE (<math><expr1>, <expr2>, DOT).

Syntax





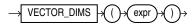
Parameters

- expr1 and expr2 must evaluate to vectors that have the same format and number of dimensions.
- INNER PRODUCT returns NULL, if either expr1 or expr2 is NULL.

VECTOR_DIMS

VECTOR_DIMS returns the number of dimensions of a vector as a NUMBER. VECTOR_DIMS is synonymous with VECTOR_DIMENSION_COUNT.

Syntax



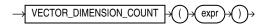
Purpose

Refer to VECTOR_DIMENSION_COUNT for full semantics.

VECTOR_DIMENSION_COUNT

VECTOR DIMENSION COUNT returns the number of dimensions of a vector as a NUMBER.

Syntax



Purpose

VECTOR DIMENSION COUNT is synonymous with VECTOR_DIMS.

Parameters

expr must evaluate to a vector.

If expr is NULL, NULL is returned.

Example



VECTOR_DIMENSION_FORMAT

VECTOR_DIMENSION_FORMAT returns the storage format of the vector. It returns a VARCHAR2, which can be one of the following values: INT8, FLOAT32, FLOAT64, or BINARY.

Syntax



Parameters

expr must evaluate to a vector.

If expr is NULL, NULL is returned.

Examples

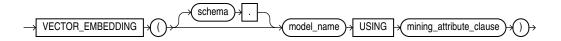
```
SELECT VECTOR DIMENSION FORMAT (TO VECTOR ('[34.6, 77.8]', 2, FLOAT64));
VECTOR_DIMENSION_FORMAT(TO_VECTOR('[34.6,77.8]',2,
FLOAT64
SELECT VECTOR DIMENSION FORMAT(TO VECTOR('[34.6, 77.8, 9]', 3, FLOAT32));
VECTOR_DIMENSION_FORMAT(TO_VECTOR('[34.6,77.8,9]',
FLOAT32
SELECT VECTOR_DIMENSION_FORMAT(TO_VECTOR('[34.6, 77.8, 9.10]', 3, INT8));
VECTOR DIMENSION FORMAT (TO VECTOR ('[34.6,77.8,9.10
INT8
SELECT VECTOR DIMENSION FORMAT (TO VECTOR ('[206, 32]', 16, BINARY));
VECTOR DIMENSION FORMAT(TO VECTOR('[206,32]',16,BI
BINARY
SELECT VECTOR DIMENSION FORMAT(TO VECTOR('[34.6, 77.8, 9, 10]', 3, INT8));
SELECT VECTOR_DIMENSION_FORMAT(TO_VECTOR('[34.6, 77.8, 9, 10]', 3, INT8))
ERROR at line 1:
ORA-51803: Vector dimension count must match the dimension count specified in
the column definition (expected 3 dimensions, specified 4 dimensions).
```



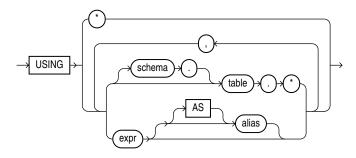
VECTOR_EMBEDDING

Use VECTOR_EMBEDDING to generate a single vector embedding for different data types using embedding or feature extraction machine learning models.

Syntax



mining_attribute_clause::=



Purpose

The function accepts the following types as input:

VARCHAR2 for text embedding models. Oracle automatically converts any other type to VARCHAR2 except for NCLOB, which is automatically converted to NVARCHAR2. Oracle does not expect values whose textual representation exceeds the maximum size of a VARCHAR2, since embedding models support only text that translates to a couple of thousand tokens. An attribute with a type that has no conversion to VARCHAR2 results in a SQL compilation error.

For feature extraction models Oracle Machine Learning for SQL supports standard Oracle data types except DATE, TIMESTAMP, RAW, and LONG. Oracle Machine Learning supports date type (datetime, date, timestamp) for case_id, CLOB/BLOB/FILE that are interpreted as text columns, and the following collection types as well:

- DM NESTED CATEGORICALS
- DM NESTED NUMERICALS
- DM NESTED BINARY DOUBLES
- DM_NESTED_BINARY_FLOATS

The function always returns a VECTOR type, whose dimension is dictated by the model itself. The model stores the dimension information in metadata within the data dictionary.

You can use <code>VECTOR_EMBEDDING</code> in <code>SELECT</code> clauses, in predicates, and as an operand for SQL operations accepting a <code>VECTOR</code> type.

Parameters

model_name refers to the name of the imported embedding model that implements the embedding machine learning function.

mining_attribute_clause

- The mining_attribute_clause argument identifies the column attributes to use as predictors for scoring. This is used as a convenience, as the embedding operator only accepts single input value.
- USING *: all the relevant attributes present in the input (supplied in JSON metadata) are used. This is used as a convenience. For an embedding model, the operator only takes one input value as embedding models have only one column.
- USING expr [AS alias] [, expr [AS alias]]: all the relevant attributes present in the commaseparated list of column expressions are used. This syntax is consistent with the syntax of other machine learning operators. You may specify more than one attribute, however, the embedding model only takes one relevant input. Therefore, you must specify a single mining attribute.

Example

The following example generates vector embeddings with "hello" as the input, utilizing the pretrained ONNX format model my_embedding_model.onnx imported into the Database. For complete example, see Import ONNX Models and Generate Embeddings

See Also:

- Data Requirements for Machine Learning
- Vector Distance Metrics

VECTOR NORM

VECTOR_NORM returns the Euclidean norm of a vector (SQRT(SUM((xi-yi)2))) as a BINARY_DOUBLE. This value is also called magnitude or size and represents the Euclidean distance between the vector and the origin.

Syntax





Parameters

expr must evaluate to a vector.

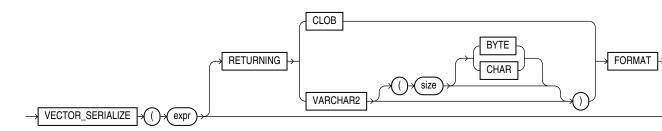
If expr is NULL, NULL is returned.

Example

VECTOR_SERIALIZE

VECTOR SERIALIZE is synonymous with FROM VECTOR.

Syntax



Purpose

See FROM_VECTOR for semantics and examples.

Examples

```
VECTOR_SERIALIZE(VECTOR('[1.1, 2.2, 3.3]',3,FLOAT32)RETURNINGCLOB)

1 row selected.

SELECT VECTOR_SERIALIZE(TO_VECTOR('[5,[2,4],[1.0,2.0]]', 5, FLOAT64, SPARSE) RETURNING CLOB FORMAT SPARSE);

VECTOR_SERIALIZE(TO_VECTOR('[5,[2,4],[1.0,2.0]]',5,FLOAT64,SPARSE)RETURNINGCLOBF

[5,[2,4],[1.0E+000,2.0E+000]]

1 row selected.

SELECT VECTOR_SERIALIZE(TO_VECTOR('[5,[2,4],[1.0,2.0]]',5,FLOAT64,SPARSE)RETURNINGCLOBF

CLOB FORMAT DENSE);

VECTOR_SERIALIZE(TO_VECTOR('[5,[2,4],[1.0,2.0]]',5,FLOAT64,SPARSE)RETURNINGCLOBF

CLOB FORMAT DENSE);

VECTOR_SERIALIZE(TO_VECTOR('[5,[2,4],[1.0,2.0]]',5,FLOAT64,SPARSE)RETURNINGCLOBF

CLOB FORMAT DENSE);

VECTOR_SERIALIZE(TO_VECTOR('[5,[2,4],[1.0,2.0]]',5,FLOAT64,SPARSE)RETURNINGCLOBF

CLOB FORMAT DENSE);
```

VSIZE

Syntax



Purpose

VSIZE returns the number of bytes in the internal representation of expr. If expr is null, then this function returns null.

This function does not support CLOB data directly. However, CLOBS can be passed in as arguments through implicit data conversion.



"Data Type Comparison Rules" for more information

Examples

The following example returns the number of bytes in the <code>last_name</code> column of the employees in department 10:

```
SELECT last_name, VSIZE (last_name) "BYTES"
FROM employees
WHERE department_id = 10
ORDER BY employee_id;
```



LAST_NAME	BYTES
Whalen	6

WIDTH_BUCKET

Syntax



Purpose

WIDTH_BUCKET lets you construct equiwidth histograms, in which the histogram range is divided into intervals that have identical size. (Compare this function with NTILE, which creates equiheight histograms.) Ideally each bucket is a closed-open interval of the real number line. For example, a bucket can be assigned to scores between 10.00 and 19.999 ... to indicate that 10 is included in the interval and 20 is excluded. This is sometimes denoted [10, 20).

For a given expression, <code>WIDTH_BUCKET</code> returns the bucket number into which the value of this expression would fall after being evaluated.

- expr is the expression for which the histogram is being created. This expression must
 evaluate to a numeric or datetime value or to a value that can be implicitly converted to a
 numeric or datetime value. If expr evaluates to null, then the expression returns null.
- min_value and max_value are expressions that resolve to the end points of the acceptable range for expr. Both of these expressions must also evaluate to numeric or datetime values, and neither can evaluate to null.
- num_buckets is an expression that resolves to a constant indicating the number of buckets. This expression must evaluate to a positive integer.



Table 2-9 for more information on implicit conversion

When needed, Oracle Database creates an underflow bucket numbered 0 and an overflow bucket numbered $num_buckets+1$. These buckets handle values less than min_value and more than max_value and are helpful in checking the reasonableness of endpoints.

Examples

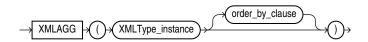
The following example creates a ten-bucket histogram on the <code>credit_limit</code> column for customers in Switzerland in the sample table <code>oe.customers</code> and returns the bucket number ("Credit Group") for each customer. Customers with credit limits greater than or equal to the maximum value are assigned to the overflow bucket, 11:



825	Dreyfuss	500	1
	Barkin	500	1
827	Siegel	500	1
853	Palin	400	1
843	Oates	700	2
844	Julius	700	2
835	Eastwood	1200	3
836	Berenger	1200	3
837	Stanton	1200	3
840	Elliott	1400	3
841	Boyer	1400	3
842	Stern	1400	3
848	Olmos	1800	4
849	Kaurusmdki	1800	4
828	Minnelli	2300	5
829	Hunter	2300	5
850	Finney	2300	5
851	Brown	2300	5
852	Tanner	2300	5
830	Dutt	3500	7
831	Bel Geddes	3500	7
832	Spacek	3500	7
833	Moranis	3500	7
834	Idle	3500	7
838	Nicholson	3500	7
839	Johnson	3500	7
845	Fawcett	5000	11
846	Brando	5000	11
847	Streep	5000	11

XMLAGG

Syntax



Purpose

XMLAgg is an aggregate function. It takes a collection of XML fragments and returns an aggregated XML document. Any arguments that return null are dropped from the result.

XMLAgg is similar to SYS_XMLAgg except that XMLAgg returns a collection of nodes but it does not accept formatting using the XMLFormat object. Also, XMLAgg does not enclose the output in an element tag as does SYS XMLAgg.

Within the <code>order_by_clause</code>, Oracle Database does not interpret number literals as column positions, as it does in other uses of this clause, but simply as number literals.

See Also:

XMLELEMENT and SYS_XMLAGG



Examples

The following example produces a Department element containing Employee elements with employee job ID and last name as the contents of the elements:

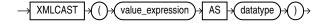
```
SELECT XMLELEMENT ("Department",
  XMLAGG(XMLELEMENT("Employee",
  e.job_id||' '||e.last_name)
  ORDER BY last name))
  as "Dept_list"
  FROM employees e
  WHERE e.department_id = 30;
Dept list
<Department>
 <Employee>PU CLERK Baida
 <Employee>PU CLERK Colmenares
 <Employee>PU CLERK Himuro
 <Employee>PU CLERK Khoo
 <Employee>PU MAN Raphaely
 <Employee>PU CLERK Tobias
</Department>
```

The result is a single row, because XMLAgg aggregates the rows. You can use the GROUP BY clause to group the returned set of rows into multiple groups:

```
SELECT XMLELEMENT ("Department",
    XMLAGG(XMLELEMENT("Employee", e.job_id||' '||e.last_name)))
  AS "Dept list"
  FROM employees e
  GROUP BY e.department id;
Dept_list
           ______
<Department>
 <Employee>AD_ASST Whalen
</Department>
<Department>
 <Employee>MK MAN Hartstein
 <Employee>MK REP Fay
</Department>
<Department>
 <Employee>PU MAN Raphaely
 <Employee>PU CLERK Khoo
 <Employee>PU CLERK Tobias
 <Employee>PU CLERK Baida
 <Employee>PU CLERK Colmenares
 <Employee>PU CLERK Himuro
</Department>
. . .
```

XMLCAST

Syntax



(datatype::=)

Purpose

XMLCast casts value_expression to the scalar SQL data type specified by datatype. The value expression argument is a SQL expression that is evaluated.

datatype

The datatype argument can be of data type NUMBER, VARCHAR2, VARCHAR, CHAR, CLOB, BLOB, REF XMLTYPE, and any of the datetime data types.

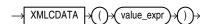
BLOB, or CLOB with options reference or value. The default is reference.

See Also:

- Oracle XML DB Developer's Guide for more information on uses for this function and examples
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the return value of XMLCAST when it is a character value

XMLCDATA

Syntax



Purpose

XMLCData generates a CDATA section by evaluating value_expr. The value_expr must resolve to a string. The value returned by the function takes the following form:

```
<![CDATA[string]]>
```

If the resulting value is not a valid XML CDATA section, then the function returns an error. The following conditions apply to XMLCData:

- The value_expr cannot contain the substring]]>.
- If value expr evaluates to null, then the function returns null.

See Also:

Oracle XML DB Developer's Guide for more information on this function

Examples

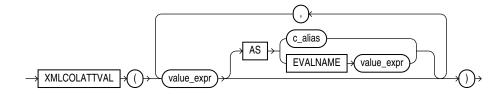
The following statement uses the <code>DUAL</code> table to illustrate the syntax of <code>XMLCData</code>:



```
SELECT XMLELEMENT ("PurchaseOrder",
  XMLAttributes(dummy as "pono"),
  XMLCdata('<!DOCTYPE po_dom_group [</pre>
  <!ELEMENT po dom group(student name) *>
   <!ELEMENT po_purch_name (#PCDATA)>
   <!ATTLIST po name po no ID #REQUIRED>
   <!ATTLIST po name trust 1 IDREF #IMPLIED>
   <!ATTLIST po name trust 2 IDREF #IMPLIED>
   ]>')) "XMLCData" FROM DUAL;
XMLCData
<PurchaseOrder pono="X"><! [CDATA[
<!DOCTYPE po dom group [
  <!ELEMENT po dom group(student name) *>
  <!ELEMENT po purch name (#PCDATA)>
  <!ATTLIST po name po no ID #REQUIRED>
  <!ATTLIST po name trust 1 IDREF #IMPLIED>
  <!ATTLIST po name trust 2 IDREF #IMPLIED>
  ]>
 11>
</PurchaseOrder>
```

XMLCOLATTVAL

Syntax



Purpose

XMLColAttVal creates an XML fragment and then expands the resulting XML so that each XML fragment has the name column with the attribute name.

You can use the AS clause to change the value of the name attribute to something other than the column name. You can do this by specifying c_alias , which is a string literal, or by specifying <code>EVALNAME value_expr</code>. In the latter case, the value expression is evaluated and the result, which must be a string literal, is used as the alias. The alias can be up to 4000 characters if the initialization parameter <code>MAX_STRING_SIZE = STANDARD</code>, and 32767 characters if <code>MAX_STRING_SIZE = EXTENDED</code>. See "Extended Data Types" for more information.

You must specify a value for value expr. If value expr is null, then no element is returned.

Restriction on XMLColAttVal

You cannot specify an object type column for value expr.

Examples

The following example creates an Emp element for a subset of employees, with nested employee_id, last_name, and salary elements as the contents of Emp. Each nested element is named column and has a name attribute with the column name as the attribute value:

Refer to the example for XMLFOREST to compare the output of these two functions.

XMLCOMMENT

Syntax



Purpose

XMLComment generates an XML comment using an evaluated result of value_expr. The value_expr must resolve to a string. It cannot contain two consecutive dashes (hyphens). The value returned by the function takes the following form:

```
<!--string-->
```

If value expr resolves to null, then the function returns null.



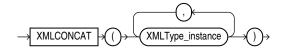
Oracle XML DB Developer's Guide for more information on this function

Examples

The following example uses the DUAL table to illustrate the XMLComment syntax:

XMLCONCAT

Syntax



Purpose

XMLConcat takes as input a series of XMLType instances, concatenates the series of elements for each row, and returns the concatenated series. XMLConcat is the inverse of XMLSequence.

Null expressions are dropped from the result. If all the value expressions are null, then the function returns null.

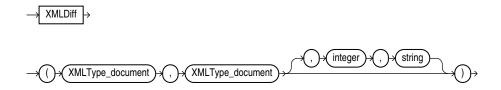


Examples

The following example creates XML elements for the first and last names of a subset of employees, and then concatenates and returns those elements:

XMLDIFF

Syntax



Purpose

The XMLDiff function is the SQL interface for the XmlDiff C API. This function compares two XML documents and captures the differences in XML conforming to an Xdiff schema. The diff document is returned as an XMLType document.

- For the first two arguments, specify the names of two XMLType documents.
- For the *integer*, specify a number representing the hashLevel for a C function XmlDiff. If you do not want hashing, set this argument to 0 or omit it entirely. If you do not want hashing, but you want to specify flags, then you must set this argument to 0.
- For *string*, specify the flags that control the behavior of the function. These flags are specified by one or more names separated by semicolon. The names are the same as the names of constants for XmlDiff function.



Oracle XML Developer's Kit Programmer's Guide for more information on using this function, including examples, and *Oracle Database XML C API Reference* for information on the XML APIs for C

Examples

The following example compares two XML documents and returns the difference as an XMLType document:

```
SELECT XMLDIFF(
 XMLTYPE('<?xml version="1.0"?>
 <br/>
<br/>
bk:book xmlns:bk="http://example.com">
                                    <bk:tr>
                                                                                          <bk:td>
                                                                                                                                                                                     <br/>

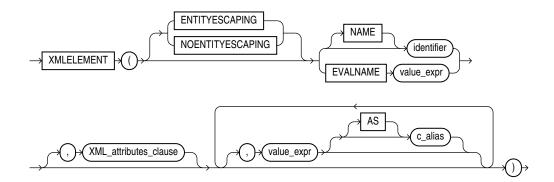
                                                                                                                                                                                                                                                                     Chapter 1.
                                                                                                                                                                                     </bk:chapter>
                                                                                          </bk:td>
                                                                                            <bk:td>
                                                                                                                                                                                                 <br/>

                                                                                                                                                                                                                                                                   Chapter 2.
                                                                                                                                                                                       </bk:chapter>
                                                                                          </bk:td>
                                    </bk:tr>
 </bk:book>'),
 XMLTYPE('<?xml version="1.0"?>
 <br/><bk:book xmlns:bk="http://example.com">
                                    <bk:tr>
                                                                                          <bk:td>
                                                                                                                                                                                       <br/><bk:chapter>
                                                                                                                                                                                                                                                                            Chapter 1.
                                                                                                                                                                                       </bk:chapter>
                                                                                          </bk:td>
                                                                                          <bk:td/>
                                    </bk:tr>
 </bk:book>')
 )
FROM DUAL;
```

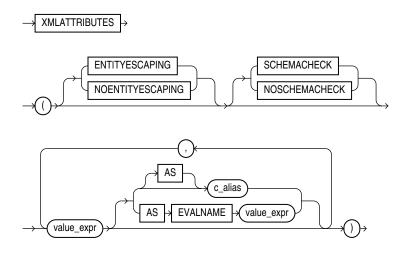


XMLELEMENT

Syntax



XML attributes clause::=



Purpose

XMLElement takes an element name for identifier or evaluates an element name for EVALNAME $value_expr$, an optional collection of attributes for the element, and arguments that make up the content of the element. It returns an instance of type XMLType. XMLElement is similar to SYS_XMLGen except that XMLElement can include attributes in the XML returned, but it does not accept formatting using the XMLFormat object.

The XMLElement function is typically nested to produce an XML document with a nested structure, as in the example in the following section.

For an explanation of the ENTITYESCAPING and NONENTITYESCAPING keywords, refer to *Oracle XML DB Developer's Guide*.

You must specify a value for Oracle Database to use an the enclosing tag. You can do this by specifying <code>identifier</code>, which is a string literal, or by specifying <code>EVALNAME value_expr</code>. In the latter case, the value expression is evaluated and the result, which must be a string literal, is used as the identifier. The identifier does not have to be a column name or column reference. It

cannot be an expression or null. It can be up to 4000 characters if the initialization parameter MAX STRING SIZE = STANDARD, and 32767 characters if MAX STRING SIZE = EXTENDED.

The objects that make up the element content follow the XMLATTRIBUTES keyword. In the $XML_attributes_clause$, if the $value_expr$ is null, then no attribute is created for that value expression. The type of $value_expr$ cannot be an object type or collection. If you specify an alias for $value_expr$ using the AS clause, then the c_alias or the evaluated value expression (EVALNAME $value_expr$) can be up to 4000 characters if the initialization parameter MAX STRING SIZE = STANDARD, and 32767 characters if MAX STRING SIZE = EXTENDED.

```
See Also:

"Extended Data Types" for more information on MAX_STRING_SIZE
```

For the optional value_expr that follows the XML_attributes_clause in the diagram:

- If value_expr is a scalar expression, then you can omit the AS clause, and Oracle uses the column name as the element name.
- If $value_expr$ is an object type or collection, then the AS clause is mandatory, and Oracle uses the specified c alias as the enclosing tag.
- If value expr is null, then no element is created for that value expression.

```
SYS_XMLGEN
```

Examples

The following example produces an Emp element for a series of employees, with nested elements that provide the employee's name and hire date:

```
SELECT XMLELEMENT ("Emp", XMLELEMENT ("Name",
  e.job id||' '||e.last name),
  XMLELEMENT("Hiredate", e.hire_date)) as "Result"
  FROM employees e WHERE employee id > 200;
Result
______
<Emp>
 <Name>MK MAN Hartstein</Name>
 <Hiredate>2004-02-17</Hiredate>
</Emp>
<Emp>
 <Name>MK REP Fay</Name>
 <hiredate>2005-08-17</hiredate>
</Emp>
<Emp>
 <Name>HR REP Mavris</Name>
 <Hiredate>2002-06-07</Hiredate>
</Emp>
```

The following similar example uses the XMLElement function with the XML_attributes_clause to create nested XML elements with attribute values for the top-level element:

Notice that the AS *identifier* clause was not specified for the last_name column. As a result, the XML returned uses the column name last name as the default.

Finally, the next example uses a subquery within the XML_attributes_clause to retrieve information from another table into the attributes of an element:

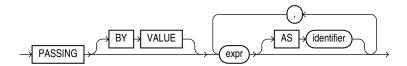


XMLEXISTS

Syntax



XML_passing_clause::=



Purpose

XMLExists checks whether a given XQuery expression returns a nonempty XQuery sequence. If so, the function returns TRUE; otherwise, it returns FALSE. The argument XQuery_string is a literal string, but it can contain XQuery variables that you bind using the XML passing clause.

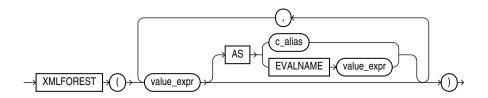
The expr in the $XML_passing_clause$ is an expression returning an XMLType or an instance of a SQL scalar data type that is used as the context for evaluating the XQuery expression. You can specify only one expr in the PASSING clause without an identifier. The result of evaluating each expr is bound to the corresponding identifier in the $XQuery_string$. If any expr that is not followed by an AS clause, then the result of evaluating that expression is used as the context item for evaluating the $XQuery_string$. If expr is a relational column, then its declared collation is ignored by Oracle XML DB.



Oracle XML DB Developer's Guide for more information on uses for this function and examples

XMLFOREST

Syntax





Purpose

XMLForest converts each of its argument parameters to XML, and then returns an XML fragment that is the concatenation of these converted arguments.

- If value_expr is a scalar expression, then you can omit the AS clause, and Oracle
 Database uses the column name as the element name.
- If value_expr is an object type or collection, then the AS clause is mandatory, and Oracle uses the specified expression as the enclosing tag.

You can do this by specifying c_alias , which is a string literal, or by specifying EVALNAME $value_expr$. In the latter case, the value expression is evaluated and the result, which must be a string literal, is used as the identifier. The identifier does not have to be a column name or column reference. It cannot be an expression or null. It can be up to 4000 characters if the initialization parameter MAX_STRING_SIZE = STANDARD, and 32767 characters if MAX_STRING_SIZE = EXTENDED. See "Extended Data Types" for more information.

If value expr is null, then no element is created for that value expr.

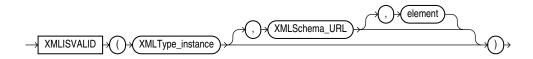
Examples

The following example creates an Emp element for a subset of employees, with nested employee id, last name, and salary elements as the contents of Emp:

Refer to the example for XMLCOLATTVAL to compare the output of these two functions.

XMLISVALID

Syntax



Purpose

XMLISVALID checks whether the input XMLType_instance conforms to the relevant XML schema. It does not change the validation status recorded for XMLType instance.

If the input XML document is determined to be valid, then XMLISVALID returns 1; otherwise, it returns 0. If you provide XMLSchema URL as an argument, then that is used to check

conformance. Otherwise, the XML schema specified by the XML document is used to check conformance.

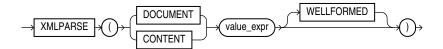
- XMLType_instance is the XMLType instance to be validated.
- XMLSchema URL is the URL of the XML schema against which to check conformance.
- element is the element of the specified schema against which to check conformance. Use this if you have an XML schema that defines more than one top level element, and you want to check conformance against a specific one of those elements.



Oracle XML DB Developer's Guide for information on the use of this function, including examples

XMLPARSE

Syntax



Purpose

XMLParse parses and generates an XML instance from the evaluated result of value_expr. The value_expr must resolve to a string. If value_expr resolves to null, then the function returns null.

- If you specify DOCUMENT, then value expr must resolve to a singly rooted XML document.
- If you specify CONTENT, then value expr must resolve to a valid XML value.
- When you specify WELLFORMED, you are guaranteeing that value_expr resolves to a well-formed XML document, so the database does not perform validity checks to ensure that the input is well formed.

See Also:

Oracle XML DB Developer's Guide for more information on this function

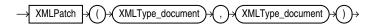
Examples

The following example uses the DUAL table to illustrate the syntax of XMLParse:



XMLPATCH

Syntax



Purpose

The XMLPatch function is the SQL interface for the XmlPatch C API. This function patches an XML document with the changes specified. A patched XMLType document is returned.

- For the first argument, specify the name of the input XMLType document.
- For the second argument, specify the XMLType document containing the changes to be applied to the first document. The changes should conform to the Xdiff XML schema. You can supply the XML output from the Oracle XML Developer's Kit Java method diff().

See Also:

Oracle XML Developer's Kit Programmer's Guide for more information on using this function, including examples, and *Oracle Database XML C API Reference* for information on the XML APIs for C

Examples

The following example patches an $\mathtt{XMLType}$ document with the changes specified in another $\mathtt{XMLType}$ and returns a patched $\mathtt{XMLType}$ document:

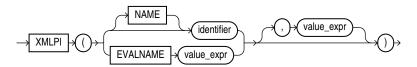
```
SELECT XMLPATCH (
XMLTYPE('<?xml version="1.0"?>
<bk:book xmlns:bk="http://example.com">
   <bk:tr>
        <bk:td>
                <br/><bk:chapter>
                        Chapter 1.
                 </bk:chapter>
        </bk:td>
        <bk:td>
                  <br/><br/>chapter>
                        Chapter 2.
                </bk:chapter>
        </bk:td>
   </bk:tr>
</bk:book>'),
XMLTYPE('<?xml version="1.0"?>
<xd:xdiff xsi:schemaLocation="http://xmlns.oracle.com/xdb/xdiff.xsd</pre>
 http://xmlns.oracle.com/xdb/xdiff.xsd"
 xmlns:xd="http://xmlns.oracle.com/xdb/xdiff.xsd"
```



```
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
xmlns:bk="http://example.com">
    <?oracle-xmldiff operations-in-docorder="true" output-model="snapshot"
    diff-algorithm="global"?>
    <xd:delete-node xd:node-type="element"
    xd:xpath="/bk:book[1]/bk:tr[1]/bk:td[2]/bk:chapter[1]"/>
    </xd:xdiff>')
)
FROM DUAL;
```

XMLPI

Syntax



Purpose

XMLPI generates an XML processing instruction using *identifier* and optionally the evaluated result of *value_expr*. A processing instruction is commonly used to provide to an application information that is associated with all or part of an XML document. The application uses the processing instruction to determine how best to process the XML document.

You must specify a value for Oracle Database to use an the enclosing tag. You can do this by specifying <code>identifier</code>, which is a string literal, or by specifying <code>EVALNAME value_expr</code>. In the latter case, the value expression is evaluated and the result, which must be a string literal, is used as the identifier. The identifier does not have to be a column name or column reference. It cannot be an expression or null. It can be up to 4000 characters if the initialization parameter <code>MAX_STRING_SIZE = STANDARD</code>, and 32767 characters if <code>MAX_STRING_SIZE = EXTENDED</code>. See "Extended Data Types" for more information.

The optional *value_expr* must resolve to a string. If you omit the optional *value_expr*, then a zero-length string is the default. The value returned by the function takes this form:

<?identifier string?>

XMLPI is subject to the following restrictions:

- The *identifier* must be a valid target for a processing instruction.
- You cannot specify xml in any case combination for identifier.
- The identifier cannot contain the consecutive characters ?>.



Oracle XML DB Developer's Guide for more information on this function

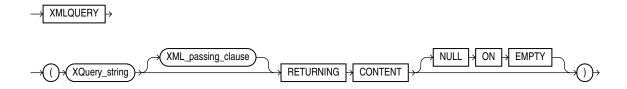
Examples

The following statement uses the DUAL table to illustrate the use of the XMLPI syntax:

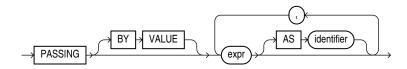


XMLQUERY

Syntax



XML_passing_clause::=



Purpose

XMLQUERY lets you query XML data in SQL statements. It takes an XQuery expression as a string literal, an optional context item, and other bind variables and returns the result of evaluating the XQuery expression using these input values.

- XQuery string is a complete XQuery expression, including prolog.
- The <code>expr</code> in the <code>XML_passing_clause</code> is an expression returning an <code>XMLType</code> or an instance of a SQL scalar data type that is used as the context for evaluating the XQuery expression. You can specify only one <code>expr</code> in the <code>PASSING</code> clause without an identifier. The result of evaluating each <code>expr</code> is bound to the corresponding identifier in the <code>XQuery_string</code>. If any <code>expr</code> that is not followed by an <code>AS</code> clause, then the result of evaluating that expression is used as the context item for evaluating the <code>XQuery_string</code>. If <code>expr</code> is a relational column, then its declared collation is ignored by Oracle XML DB.
- RETURNING CONTENT indicates that the result from the XQuery evaluation is either an XML
 1.0 document or a document fragment conforming to the XML
 1.0 semantics.
- If the result set is empty, then the function returns the SQL NULL value. The NULL ON EMPTY keywords are implemented by default and are shown for semantic clarity.



Oracle XML DB Developer's Guide for more information on this function



Examples

The following statement specifies the <code>warehouse_spec</code> column of the <code>oe.warehouses</code> table in the <code>XML_passing_clause</code> as a context item. The statement returns specific information about the warehouses with area greater than 50K.

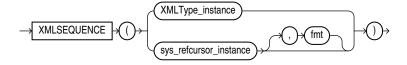
```
SELECT warehouse name,
EXTRACTVALUE (warehouse spec, '/Warehouse/Area'),
XMLQuery(
   'for $i in /Warehouse
  where $i/Area > 50000
  return <Details>
             <Docks num="{$i/Docks}"/>
             <Rail>
              if ($i/RailAccess = "Y") then "true" else "false"
             </Rail>
          </Details>' PASSING warehouse spec RETURNING CONTENT) "Big warehouses"
   FROM warehouses;
WAREHOUSE ID Area
                    Big warehouses
                 25000
          2
                50000
               85700 <Details><Docks></Docks><Rail>false</Rail></Details>
             103000 <Details><Docks num="3"></Docks><Rail>true</Rail></Details>
```

XMLSEQUENCE

Note:

The XMLSEQUENCE function is deprecated. It is still supported for backward compatibility. However, Oracle recommends that you use the XMLTABLE function instead. See XMLTABLE for more information.

Syntax



Purpose

XMLSequence has two forms:

• The first form takes as input an XMLType instance and returns a varray of the top-level nodes in the XMLType. This form is effectively superseded by the SQL/XML standard function XMLTable, which provides for more readable SQL code. Prior to Oracle Database 10g Release 2, XMLSequence was used with SQL function TABLE to do some of what can now be done better with the XMLTable function.

 The second form takes as input a REFCURSOR instance, with an optional instance of the XMLFormat object, and returns as an XMLSequence type an XML document for each row of the cursor.

Because XMLSequence returns a collection of XMLType, you can use this function in a TABLE clause to unnest the collection values into multiple rows, which can in turn be further processed in the SQL query.



Oracle XML DB Developer's Guide for more information on this function, and XMLTABLE

Examples

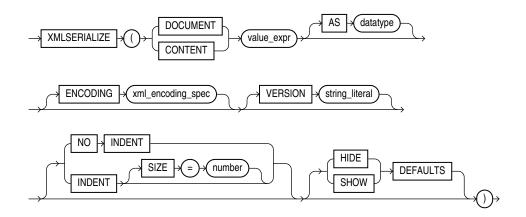
The following example shows how XMLSequence divides up an XML document with multiple elements into VARRAY single-element documents. In this example, the TABLE keyword instructs Oracle Database to consider the collection a table value that can be used in the FROM clause of the subquery:

```
SELECT EXTRACT(warehouse_spec, '/Warehouse') as "Warehouse"
  FROM warehouses WHERE warehouse name = 'San Francisco';
______
<Warehouse>
 <Building>Rented</Building>
 <Area>50000</Area>
 <Docks>1</Docks>
 <DockType>Side load</DockType>
 <WaterAccess>Y</WaterAccess>
 <RailAccess>N</RailAccess>
 <Parking>Lot</Parking>
 <VClearance>12 ft</VClearance>
</Warehouse>
1 row selected.
SELECT VALUE (p)
  FROM warehouses w,
  TABLE (XMLSEQUENCE (EXTRACT (warehouse_spec, '/Warehouse/*'))) p
  WHERE w.warehouse name = 'San Francisco';
VALUE (P)
______
<Building>Rented</Building>
<Area>50000</Area>
<Docks>1</Docks>
<DockType>Side load</DockType>
<WaterAccess>Y</WaterAccess>
<RailAccess>N</RailAccess>
<Parking>Lot</Parking>
<VClearance>12 ft</VClearance>
8 rows selected.
```



XMLSERIALIZE

Syntax



(datatype::=)

Purpose

XMLSerialize creates a string or LOB containing the contents of value expr.

Any lob returned by XMLSERIALIZE will be read-only.

If you specify DOCUMENT, then the value expr must be a valid XML document.

If you specify CONTENT, then the $value_expr$ need not be a singly rooted XML document. However it must be valid XML content.

datatype

The datatype specified can be:

- VARCHAR2 or VARCHAR. but not NVARCHAR2
- BLOB, or CLOB with options reference or value. The default is reference.
- With BLOB, you can specify the ENCODING clause to use the specified encoding in the prolog. The xml encoding spec is an XML encoding declaration (encoding="...").

The default type is CLOB.

Specify the VERSION clause to use the version you provide as $string_literal$ in the XML declaration (<?xml version="..." ...?>).

Specify NO INDENT to strip all insignificant whitespace from the output. Specify INDENT SIZE = N, where N is a whole number, for output that is pretty-printed using a relative indentation of N spaces. If N is 0, then pretty-printing inserts a newline character after each element, placing each element on a line by itself, but omitting all other insignificant whitespace in the output. If INDENT is present without a SIZE specification, then 2-space indenting is used. If you omit this clause, then the behavior (pretty-printing or not) is indeterminate.

HIDE DEFAULTS and SHOW DEFAULTS apply only to XML schema-based data. If you specify SHOW DEFAULTS and the input data is missing any optional elements or attributes for which the XML schema defines default values, then those elements or attributes are included in the output

with their default values. If you specify HIDE DEFAULTS, then no such elements or attributes are included in the output. HIDE DEFAULTS is the default behavior.

See Also:

- Oracle XML DB Developer's Guide for more information on this function
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to the character return value of XMLSERIALIZE

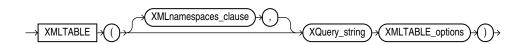
Examples

The following statement uses the DUAL table to illustrate the syntax of XMLSerialize:

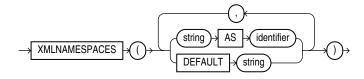
SELECT XMLSERIALIZE(CONTENT XMLTYPE('<Owner>Grandco</Owner>')) AS xmlserialize_doc
FROM DUAL;

XMLTABLE

Syntax



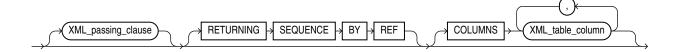
XMLnamespaces_clause::=



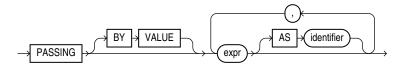
Note:

You can specify at most one DEFAULT string clause.

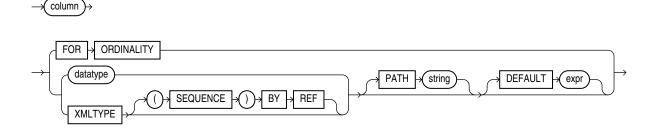
XMLTABLE_options::=



XML_passing_clause::=



XML table column::=



(datatype::=)

Purpose

XMLTable maps the result of an XQuery evaluation into relational rows and columns. You can query the result returned by the function as a virtual relational table using SQL.

- The XMLNAMESPACES clause contains a set of XML namespace declarations. These declarations are referenced by the XQuery expression (the evaluated XQuery_string), which computes the row, and by the XPath expression in the PATH clause of XML_table_column, which computes the columns for the entire XMLTable function. If you want to use qualified names in the PATH expressions of the COLUMNS clause, then you need to specify the XMLNAMESPACES clause.
- *XQuery_string* is a literal string. It is a complete XQuery expression and can include prolog declarations. The value of XQuery_string serves as input to the XMLTable function; it is this XQuery result that is decomposed and stored as relational data.
- The <code>expr</code> in the <code>XML_passing_clause</code> is an expression returning an <code>XMLType</code> or an instance of a SQL scalar data type that is used as the context for evaluating the XQuery expression. You can specify only one <code>expr</code> in the <code>PASSING</code> clause without an identifier. The result of evaluating each <code>expr</code> is bound to the corresponding identifier in the <code>XQuery_string</code>. If any <code>expr</code> that is not followed by an <code>AS</code> clause, then the result of evaluating that expression is used as the context item for evaluating the <code>XQuery_string</code>. This clause supports only passing by value, not passing by reference. Therefore, the <code>BY VALUE</code> keywords are optional and are provided for semantic clarity.
- The optional RETURNING SEQUENCE BY REF clause causes the result of the XQuery evaluation to be returned by reference. This allows you to refer to any part of the source data in the XML_table_column clause.

If you omit this clause, then the result of the XQuery evaluation is returned by value. That is, a copy of the targeted nodes is returned instead of a reference to the actual nodes. In this case, you cannot refer to any data that is not in the returned copy in the

XML_table_column clause. In particular, you cannot refer to data that precedes the targeted nodes in the source data.

- The optional COLUMNS clause defines the columns of the virtual table to be created by XMLTable.
 - If you omit the COLUMNS clause, then XMLTable returns a row with a single XMLType pseudocolumn named COLUMN VALUE.
 - FOR ORDINALITY specifies that column is to be a column of generated row numbers.
 There must be at most one FOR ORDINALITY clause. It is created as a NUMBER column.
 - For each resulting column except the FOR ORDINALITY column, you must specify the column data type, which can be XMLType or any other data type.

If the column data type is XMLType, then specify the XMLTYPE clause. If you specify the optional (SEQUENCE) BY REF clause, then a reference to the source data targeted by the PATH expression is returned as the column content. Otherwise, column contains a copy of that targeted data.

Returning the $\mathtt{XMLType}$ data by reference lets you specify other columns whose paths target nodes in the source data that are outside those targeted by the PATH expression for column.

If the column data type is any other data type, then specify datatype clause.

datatype

The datatype specified can be:

- * BLOB, or CLOB with options reference or value. The default is reference.
- * Any other data type.
- The optional PATH clause specifies that the portion of the XQuery result that is addressed by XQuery expression string is to be used as the column content.

If you omit PATH, then the XQuery expression column is assumed. For example:

```
XMLTable(... COLUMNS xyz)
```

is equivalent to

```
XMLTable(... COLUMNS xyz PATH 'XYZ')
```

You can use different PATH clauses to split the XQuery result into different virtual-table columns.

- The optional DEFAULT clause specifies the value to use when the PATH expression results in an empty sequence. Its expr is an XQuery expression that is evaluated to produce the default value.

See Also:

- Oracle XML DB Developer's Guide for more information on the XMLTable function, including additional examples, and on XQuery in general
- Appendix C in Oracle Database Globalization Support Guide for the collation derivation rules, which define the collation assigned to each character data type column in the table generated by XMLTABLE



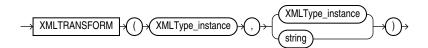
Examples

The following example converts the result of applying the XQuery '/Warehouse' to each value in the warehouse_spec column of the warehouses table into a virtual relational table with columns Water and Rail:

```
SELECT warehouse name warehouse,
  warehouse2. "Water", warehouse2. "Rail"
  FROM warehouses,
  XMLTABLE('/Warehouse'
     PASSING warehouses.warehouse_spec
        "Water" varchar2(6) PATH 'WaterAccess',
       "Rail" varchar2(6) PATH 'RailAccess')
     warehouse2;
WAREHOUSE
                               Water Rail
Southlake, Texas
                               Y
                                     Ν
                               Y
                                     Ν
San Francisco
New Jersey
                               N
                                     Ν
Seattle, Washington
```

XMLTRANSFORM

Syntax



Purpose

XMLTransform takes as arguments an XMLType instance and an XSL style sheet, which is itself a form of XMLType instance. It applies the style sheet to the instance and returns an XMLType.

This function is useful for organizing data according to a style sheet as you are retrieving it from the database.



Oracle XML DB Developer's Guide for more information on this function

Examples

The XMLTransform function requires the existence of an XSL style sheet. Here is an example of a very simple style sheet that alphabetizes elements within a node:

```
CREATE TABLE xsl_tab (col1 XMLTYPE);
INSERT INTO xsl_tab VALUES (
   XMLTYPE.createxml(
  '<?xml version="1.0"?>
   <xsl:stylesheet version="1.0" xmlns:xsl="http://www.w3.org/1999/XSL/Transform" >
```



The next example uses the xsl_tab XSL style sheet to alphabetize the elements in one warehouse spec of the sample table oe.warehouses:

CEIL, FLOOR, ROUND, and TRUNC Date Functions

Table 7-15 lists the format models you can use with the CEIL, FLOOR, ROUND, and TRUNC date functions and the units to which they round and truncate dates. The default model, 'DD', returns the date rounded or truncated to the day with a time of midnight.

Table 7-15 Date Format Models for the CEIL, FLOOR, ROUND, and TRUNC Date Functions

Format Model	Rounding or Truncating Unit
CC SCC	One greater than the first two digits of a four-digit year
SYYYY	Year (rounds up on July 1)
YYYY	
YEAR SYEAR	
YYY	
YY	
Y	

Table 7-15 (Cont.) Date Format Models for the CEIL, FLOOR, ROUND, and TRUNC Date Functions

Format Model	Rounding or Truncating Unit
IYYY IY IY	Year containing the calendar week, as defined by the ISO 8601 standard
Q	Quarter (rounds up on the sixteenth day of the second month of the quarter)
MONTH MON MM RM	Month (rounds up on the sixteenth day)
WW	Same day of the week as the first day of the year
IW	Same day of the week as the first day of the calendar week as defined by the ISO 8601 standard, which is Monday
W	Same day of the week as the first day of the month
DDD DD J	Day
DAY DY D	Starting day of the week
НН НН12 НН24	Hour
MI	Minute

The starting day of the week used by the format models DAY, DY, and D is specified implicitly by the initialization parameter NLS_TERRITORY.



Oracle Database Reference and Oracle Database Globalization Support Guide for information on this parameter

About User-Defined Functions

You can write user-defined functions in PL/SQL, Java, or C to provide functionality that is not available in SQL or SQL built-in functions. User-defined functions can appear in a SQL statement wherever an expression can occur.

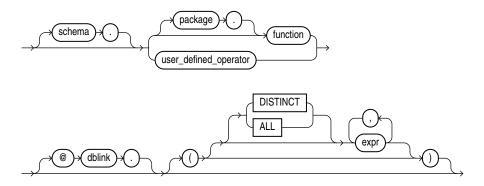
For example, user-defined functions can be used in the following:

- The select list of a SELECT statement
- The condition of a WHERE clause
- CONNECT BY, START WITH, ORDER BY, and GROUP BY clauses
- The VALUES clause of an INSERT statement
- The SET clause of an UPDATE statement



Oracle SQL does not support calling of functions with Boolean parameters or returns. Therefore, if your user-defined functions will be called from SQL statements, you must design them to return numbers (0 or 1) or character strings ('TRUE' or 'FALSE').

user defined function::=



The optional expression list must match attributes of the function, package, or operator.

Restriction on User-defined Functions

The DISTINCT and ALL keywords are valid only with a user-defined aggregate function.

✓ See Also:

- CREATE FUNCTION for information on creating functions, including restrictions on user-defined functions
- Oracle Database Development Guide for a complete discussion of the creation and use of user functions

Prerequisites

User-defined functions must be created as top-level functions or declared with a package specification before they can be named within a SQL statement.

To use a user function in a SQL expression, you must own or have EXECUTE privilege on the user function. To query a view defined with a user function, you must have the READ or SELECT privilege on the view. No separate EXECUTE privileges are needed to select from the view.



CREATE FUNCTION for information on creating top-level functions and CREATE PACKAGE for information on specifying packaged functions

Name Precedence

Within a SQL statement, the names of database columns take precedence over the names of functions with no parameters. For example, if the Human Resources manager creates the following two objects in the hr schema:

```
CREATE TABLE new_emps (new_sal NUMBER, ...);
CREATE FUNCTION new sal RETURN NUMBER IS BEGIN ... END;
```

then in the following two statements, the reference to new sal refers to the column

```
new_emps.new_sal:
SELECT new_sal FROM new_emps;
SELECT new emps.new sal FROM new emps;
```

To access the function new_sal, you would enter:

```
SELECT hr.new sal FROM new emps;
```

Here are some sample calls to user functions that are allowed in SQL expressions:

```
circle_area (radius)
payroll.tax_rate (empno)
hr.employees.tax rate (dependent, empno)@remote
```

Example

To call the tax_rate user function from schema hr, execute it against the ss_no and sal columns in tax table, specify the following:

```
SELECT hr.tax_rate (ss_no, sal)
   INTO income_tax
FROM tax_table WHERE ss_no = tax_id;
```

The INTO clause is PL/SQL that lets you place the results into the variable income tax.

Naming Conventions

If only one of the optional schema or package names is given, then the first identifier can be either a schema name or a package name. For example, to determine whether PAYROLL in the

reference <code>PAYROLL.TAX_RATE</code> is a schema or package name, Oracle Database proceeds as follows:

- 1. Check for the PAYROLL package in the current schema.
- 2. If a PAYROLL package is not found, then look for a schema name PAYROLL that contains a top-level TAX RATE function. If no such function is found, then return an error.
- 3. If the PAYROLL package is found in the current schema, then look for a TAX_RATE function in the PAYROLL package. If no such function is found, then return an error.

You can also refer to a stored top-level function using any synonym that you have defined for it.

