Managing SQL Plan Baselines

This chapter explains the concepts and tasks relating to SQL plan management using the DBMS SPM package.

See Also:

- "Migrating Stored Outlines to SQL Plan Baselines"
- Oracle Database PL/SQL Packages and Types Reference to learn more about DBMS SPM

About Managing SQL Plan Baselines

This topic describes the available interfaces and basic tasks for SQL plan management.

User Interfaces for SQL Plan Management

You can access the DBMS SPM package through Cloud Control or through the command line.

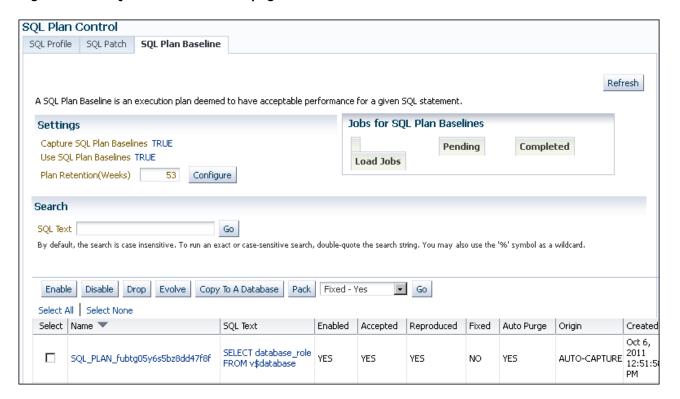
Accessing the SQL Plan Baseline Page in Cloud Control

The SQL Plan Control page in Cloud Control is a GUI that shows information about SQL profiles, SQL patches, and SQL plan baselines.

To access the SQL Plan Baseline page:

- 1. Log in to Cloud Control with the appropriate credentials.
- 2. Under the **Targets** menu, select **Databases**.
- 3. In the list of database targets, select the target for the Oracle Database instance that you want to administer.
- **4.** If prompted for database credentials, then enter the minimum credentials necessary for the tasks you intend to perform.
- 5. From the **Performance** menu, select **SQL**, then **SQL Plan Control**.
 - The SQL Plan Control page appears.
- 6. Click Files to view the SQL Plan Baseline subpage, shown in Figure 29-1.

Figure 29-1 SQL Plan Baseline Subpage



You can perform most SQL plan management tasks in this page or in pages accessed through this page.



- Cloud Control context-sensitive online help to learn about the options on the SQL Plan Baseline subpage
- "Managing the SPM Evolve Advisor Task"

DBMS_SPM Package

On the command line, use the DBMS_SPM and DBMS_XPLAN PL/SQL packages to perform most SQL plan management tasks.

The following table describes the most relevant DBMS_SPM procedures and functions for creating, adding, dropping, and loading SQL plan baselines.

Table 29-1 DBMS_SPM Procedures and Functions

Procedure or Function	Description
CONFIGURE	This procedure changes configuration options for the SMB in name/value format.
CREATE_STGTAB_BASELINE	This procedure creates a staging table that enables you to transport SQL plan baselines from one database to another.



Table 29-1	(Cont.) DBMS	SPM Procedures and Functions
-------------------	--------------	------------------------------

Procedure or Function	Description
ADD_VERIFIED_SQL_PLAN_BASELINE	Finds the plans available for the given SQL_ID in different sources such Cursor Cache, Auto SQL Tuning Set, and Automatic Workload Repository and approves the best plans.
DROP_SQL_PLAN_BASELINE	This function drops some or all plans in a plan baseline.
LOAD_PLANS_FROM_CURSOR_CACHE	This function loads plans in the shared SQL area (also called the cursor cache) into SQL plan baselines.
LOAD_PLANS_FROM_SQLSET	This function loads plans in an STS into SQL plan baselines.
LOAD_PLANS_FROM_AWR	This function loads plans from AWR into SQL plan baselines.
PACK_STGTAB_BASELINE	This function packs SQL plan baselines, which means that it copies them from the SMB into a staging table.
UNPACK_STGTAB_BASELINE	This function unpacks SQL plan baselines, which means that it copies SQL plan baselines from a staging table into the SMB.

Also, you can use <code>DBMS_XPLAN.DISPLAY_SQL_PLAN_BASELINE</code> to show one or more execution plans for the SQL statement identified by SQL handle.

See Also:

- "About the DBMS_SPM Evolve Functions" describes the functions related to SQL plan evolution.
- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS SPM and DBMS XPLAN packages

Basic Tasks in SQL Plan Management

This topic explains the basic tasks in using SQL plan management to prevent performance regressions and enable the optimizer to consider new execution plans.

The tasks are as follows:

 Set initialization parameters to control whether the database captures and uses SQL plan baselines, and whether it evolves new plans.

See "Configuring SQL Plan Management".

Display plans in a SQL plan baseline.

See "Displaying Plans in a SQL Plan Baseline".

Manually load plans into SQL plan baselines.

Load plans from AWR, SQL tuning sets, the shared SQL area, a staging table, or stored outlines.

See "Loading SQL Plan Baselines".

Manually evolve plans into SQL plan baselines.

Use PL/SQL to verify the performance of specified plans and add them to plan baselines.

See "Evolving SQL Plan Baselines Manually".

Drop all or some plans in SQL plan baselines.

See "Dropping SQL Plan Baselines".

Manage the SMB.

Alter disk space limits and change the length of the plan retention policy.

See "Managing the SQL Management Base".

Migrate stored outlines to SQL plan baselines.

See "Migrating Stored Outlines to SQL Plan Baselines".

Configuring the Capture and Use of SQL Plan Baselines

You control SQL plan management with the initialization parameters
OPTIMIZER CAPTURE SQL PLAN BASELINES and OPTIMIZER USE SQL PLAN BASELINES.

The default values are as follows:

• OPTIMIZER CAPTURE SQL PLAN BASELINES=false

For any repeatable SQL statement that does not already exist in the plan history, the database does *not* automatically create an initial SQL plan baseline for the statement.

If OPTIMIZER_CAPTURE_SQL_PLAN_BASELINES=true, then you can use the DBMS_SPM.CONFIGURE procedure to configure filters that determine which statements are eligible for plan capture. By default, no filters are configured, which means that all repeatable statements are eligible for plan capture.

• OPTIMIZER USE SQL PLAN BASELINES=true

For any SQL statement that has an existing SQL plan baseline, the database automatically adds new plans to the SQL plan baseline as unaccepted plans.



The settings of the preceding parameters are independent of each other. For example, if <code>OPTIMIZER_CAPTURE_SQL_PLAN_BASELINES</code> is true, then the database creates initial plan baselines for new statements even if <code>OPTIMIZER_USE_SQL_PLAN_BASELINES</code> is false.

If the default behavior is what you intend, then skip this section.

The following sections explain how to change the default parameter settings from the command line. If you use Cloud Control, then set these parameters in the SQL Plan Baseline subpage.



See Also:

- "Figure 29-1"
- "Automatic Initial Plan Capture"
- "Plan Selection"

Enabling Automatic Initial Plan Capture for SQL Plan Management

Setting the OPTIMIZER CAPTURE SQL PLAN BASELINES initialization parameter to true is all that is necessary for the database to automatically create an initial SQL plan baseline for any eligible SQL statement not already in the plan history.

By default, the database considers all repeatable SQL statements as eligible for capture, with the following exceptions:

- CREATE TABLE when the AS SELECT clause is not specified
- DROP TABLE
- INSERT INTO ... VALUES

Caution:

By default, when automatic baseline capture is enabled, the database creates a SQL plan baseline for every eligible repeatable statement, including all recursive SQL and monitoring SQL. Thus, automatic capture may result in the creation of an extremely large number of plan baselines. To limit the statements that are eligible for plan baselines, configure filters using the DBMS SPM.CONFIGURE procedure.

The OPTIMIZER CAPTURE SQL PLAN BASELINES parameter does not control the automatic addition of newly discovered plans to a previously created SQL plan baseline.

To enable automatic initial plan capture for SQL plan management:

- In SQL*Plus, log in to the database with the necessary privileges.
- Show the current settings for SQL plan management.

For example, connect SQL*Plus to the database with administrator privileges and execute the following command (sample output included):

```
SHOW PARAMETER SQL PLAN
```

The following sample output shows that automatic initial plan capture is disabled:

NAME	TYPE	VALUE
<pre>optimizer_capture_sql_plan_baselines</pre>	boolean	FALSE
optimizer use sql plan baselines	boolean	TRUE



If the parameters are set as you intend, then skip the remaining steps.

3. To enable the automatic recognition of repeatable SQL statements and the generation of SQL plan baselines for these statements, enter the following statement:

ALTER SYSTEM SET OPTIMIZER CAPTURE SQL PLAN BASELINES=true;



Oracle Database Reference to learn more about OPTIMIZER CAPTURE SQL PLAN BASELINES

Configuring Filters for Automatic Plan Capture

If OPTIMIZER_CAPTURE_SQL_PLAN_BASELINES=true, then you can use the DBMS_SPM.CONFIGURE procedure to create an automatic capture filter for repeatable statements.

An automatic filter enables you to capture only statements that you want, and exclude noncritical statements. This technique saves space in the SYSAUX tablespace.

The following table describes the relevant parameters of the DBMS SPM.CONFIGURE procedure.

Table 29-2 DBMS_SPM.CONFIGURE Parameters

	I
Parameter	Description
parameter_name	The type of filter for automatic capture.
	Possible values are AUTO_CAPTURE_SQL_TEXT, AUTO_CAPTURE_PARSING_SCHEMA_NAME, AUTO_CAPTURE_MODULE, and AUTO_CAPTURE_ACTION.
parameter_value	The search criteria for the automatic capture filter.
	When parameter_name is set to AUTO_CAPTURE_SQL_TEXT, the search pattern depends on the allow setting:
	• LIKE
	The parameter uses this pattern when allow=>true.
	• NOT LIKE
	The parameter uses this pattern when allow=>false.
	For all other non-null parameter_name values, the search pattern depends on the allow setting:
	• =
	The parameter uses this pattern when allow=>true.
	•
	The parameter uses this pattern when allow=>false.
	A null value removes the filter for parameter_name entirely.
allow	Whether to include (true) or exclude (false) matching SQL statements and plans. If null, then the procedure ignores the specified parameter.

You can configure multiple parameters of different types. Also, you can specify multiple values for the same parameter in separate statements, which the database combines. The settings

are additive: one parameter setting does not override a previous setting. For example, the following filter captures SQL in the parsing schema SYS or SYSTEM:

```
EXEC DBMS_SPM.CONFIGURE('AUTO_CAPTURE_PARSING_SCHEMA_NAME','SYS',true);
EXEC DBMS SPM.CONFIGURE('AUTO CAPTURE PARSING SCHEMA NAME','SYSTEM',true);
```

However, you cannot configure multiple values for the same parameter in the same procedure. For example, you cannot specify multiple SQL text strings for AUTO CAPTURE SQL TEXT.

The DBA SQL MANAGEMENT CONFIG view shows the current parameter values.

This tutorial assumes the following:

- The OPTIMIZER CAPTURE SQL PLAN BASELINES initialization parameter is set to true.
- You want to include only statements parsed in the sh schema to be eligible for baselines.
- You want to exclude statements that contain the text TEST ONLY.

To filter out all statements except those parsed in the sh schema:

- Connect SQL*Plus to the database with the appropriate privileges.
- To remove any existing filters for parsing schema and SQL text, execute the following PL/SQL programs:

```
EXEC DBMS_SPM.CONFIGURE('AUTO_CAPTURE_PARSING_SCHEMA_NAME', null, true);
EXEC DBMS_SPM.CONFIGURE('AUTO_CAPTURE_SQL_TEXT', null, true);
```

3. Include only statements parsed in the sh schema for consideration for automatic capture:

```
EXEC DBMS_SPM.CONFIGURE('AUTO_CAPTURE_PARSING_SCHEMA_NAME','sh',true);
```

4. Exclude any statement that contains the text TEST_ONLY from consideration for automatic capture:

```
EXEC DBMS SPM.CONFIGURE ('AUTO CAPTURE SQL TEXT', '%TEST ONLY%', false);
```

5. Optionally, to confirm the filters, query DBA SQL MANAGEMENT CONFIG.

For example, use the following query (sample output included):

```
COL PARAMETER_NAME FORMAT a32

COL PARAMETER_VALUE FORMAT a32

SELECT PARAMETER_NAME, PARAMETER_VALUE
FROM DBA_SQL_MANAGEMENT_CONFIG
WHERE PARAMETER_NAME LIKE '%AUTO%';

PARAMETER_NAME PARAMETER_VALUE

AUTO_CAPTURE_PARSING_SCHEMA_NAME parsing_schema IN (SH)
AUTO_CAPTURE_MODULE
AUTO_CAPTURE_ACTION
AUTO_CAPTURE_SQL_TEXT (sql_text_NOT_LIKE %TEST_ONLY%)
```



See Also:

- "Automatic Initial Plan Capture"
- Oracle Database PL/SQL Packages and Types Reference to learn more about the DBMS SPM.CONFIGURE procedure
- Oracle Database Reference to learn more about the DBA_SQL_MANAGEMENT_CONFIG view

Disabling All SQL Plan Baselines

When you set the OPTIMIZER_USE_SQL_PLAN_BASELINES initialization parameter to false, the database does not use *any* plan baselines in the database.

Typically, you might want to disable one or two plan baselines, but not all of them. A possible use case might be testing the benefits of SQL plan management.

To disable all SQL plan baselines in the database:

 Connect SQL*Plus to the database with the appropriate privileges, and then show the current settings for SQL plan management.

For example, connect SQL*Plus to the database with administrator privileges and execute the following command (sample output included):

SQL> SHOW PARAMETER SQL PLAN

NAME	TYPE	VALUE
optimizer_capture_sql_plan_baselines	boolean	FALSE
optimizer use sql plan baselines	boolean	TRUE

If the parameters are set as you intend, then skip the remaining steps.

2. To ignore all existing plan baselines enter the following statement:

SQL> ALTER SYSTEM SET OPTIMIZER USE SQL PLAN BASELINES=false



Oracle Database Reference to learn about the SQL plan baseline initialization parameters

Configuring SQL Plan Management

You can configure the capture and use of SQL plan baselines, and the SPM Evolve Advisor task.

Managing the SPM Evolve Advisor Task

SPM Evolve Advisor is a SQL advisor that evolves plans that have recently been added to the SQL plan baseline. The advisor simplifies plan evolution by eliminating the requirement to do it manually.

About the SPM Evolve Advisor Task

By default, SYS_AUTO_SPM_EVOLVE_TASK runs daily in the scheduled maintenance window

The SPM Evolve Advisor task perform the following operations:

- Locates unaccepted plans
- 2. Ranks all unaccepted plans
- Performs test executions of as many plans as possible during the maintenance window
- 4. Selects the lowest-cost plan to compare against each unaccepted plan
- Accepts automatically any unaccepted plan that performs sufficiently better, using a costbased algorithm, than the existing accepted plan



The SPM Evolve Advisor task can accept more than one plan for a single statement.

SPM Evolve Advisor locates unaccepted plans that the optimizer added to the SMB during earlier hard parses. In some cases, more optimal plans may reside in other locations. By using the <code>DBMS_SPM.SET_EVOLVE_TASK_PARAMETER</code> procedure, you can configure the automatic task to search the shared SQL area, AWR repository, or SQL tuning sets for plans that are not yet in the SMB plan history. If the advisor finds eligible plans in alternative locations, then it includes them along with the other unaccepted plans. By default, the automatic task does not search alternative locations.



Oracle Database PL/SQL Packages and Types Reference to learn more about the DBMS_SPM.SET_EVOLVE_TASK_PARAMETER procedure

Enabling and Disabling the Automatic SPM Evolve Advisor Task

No separate scheduler client exists for the Automatic SPM Evolve Advisor task.

One client controls both Automatic SQL Tuning Advisor and Automatic SPM Evolve Advisor. Thus, the same task enables or disables both. You can also disable it using DBMS_SPM.SET_EVOLVE_TASK_PARAMETER.

To disable the Automatic SPM Evolve Advisor task:

Log in to the database with the appropriate privileges.

2. Set the ALTERNATE PLAN BASELINE parameter to null:

```
BEGIN
   DBMS_SPM.SET_EVOLVE_TASK_PARAMETER(
        task_name => 'SYS_AUTO_SPM_EVOLVE_TASK',
        parameter => 'ALTERNATE_PLAN_BASELINE',
        value => '');
END;
//
```

3. Set the ALTERNATE_PLAN_SOURCE parameter to an empty string:

```
BEGIN
   DBMS_SPM.SET_EVOLVE_TASK_PARAMETER(
        task_name => 'SYS_AUTO_SPM_EVOLVE_TASK',
        parameter => 'ALTERNATE_PLAN_SOURCE',
        value => '');
END;
//
```

✓ See Also:

"Enabling and Disabling the Automatic SQL Tuning Task" to learn how to enable and disable Automatic SPM Evolve Advisor

Configuring the Automatic SPM Evolve Advisor Task

Configure automatic plan evolution by using the DBMS SPM package.

Overview of the Automatic SPM Evolve Advisor Task

Specify the automatic task parameters using the <code>SET_EVOLVE_TASK_PARAMETER</code> procedure. The following table describes some procedure parameters.

Table 29-3 DBMS_SPM.SET_EVOLVE_TASK_PARAMETER Parameters

Parameter	Description	Default
alternate_plan_source	Determines which sources to search for additional plans:	The default depends on whether the SPM Evolve Advisor task is automated
	 AUTO (the database selects the source automatically) AUTOMATIC_WORKLOAD_REPOSITORY CURSOR_CACHE SQL_TUNING_SET You can combine multiple values with the plus sign (+). 	or manual: If automated, the default is AUTO. If manual, the default is CURSOR_CACHE+AUTOMATIC_WORK LOAD_REPOSITORY.

Table 29-3 (Cont.) DBMS_SPM.SET_EVOLVE_TASK_PARAMETER Parameters

Parameter	Description	Default
alternate_plan_baseline	Determines which alternative plans should be loaded: • AUTO lets Autonomous Database choose whether to load plans for statements with or without baselines. • EXISTING loads alternate plans with for statements with existing baselines. • NEW loads alternative plans for statements without a baseline, in which case a new baseline is created. You can combine multiple values with the plus sign (+), as in EXISTING+NEW.	EXISTING
alternate_plan_limit	Specifies the maximum number of plans to load in total (that is, not the limit for each SQL statement).	The default depends on whether the SPM Evolve Advisor task is automated or manual: If automated, the default is UNLIMITED. If manual, the default is 10.
accept_plans	Specifies whether to accept recommended plans automatically. When ACCEPT_PLANS is true, SQL plan management automatically accepts all plans recommended by the task. When ACCEPT_PLANS is false, the task verifies the plans and generates a report of its findings, but does not evolve the plans automatically. You can use a report to identify new SQL plan baselines and accept them manually.	true (regardless of whether the advisor is run automatically or manually)
time_limit	Global time limit in seconds. This is the total time allowed for the task.	The default depends on whether the SPM Evolve Advisor task is automated or manual: If automated, the default is 3600. If manual, the default is 2147483646.

Assumptions

The tutorial in this section assumes the following:

- You can log in to the database as SYS. Because the SYS_AUTO_SPM_EVOLVE_TASK task is owned by SYS, only SYS can set task parameters.
- You want the database to accept plans automatically.
- You want the task to time out after 1200 seconds per execution.
- You want the evolve task to look for up to a maximum of 500 plans in the shared SQL area and AWR repository

To set automatic evolution task parameters:

1. Start SQL*Plus, and then log in to the database as SYS.

2. Query the current parameter settings for SYS AUTO SPM EVOLVE TASK.

For example, connect SQL*Plus to the database with administrator privileges and execute the following query:

Sample output appears as follows:

3. Set parameters using PL/SQL code of the following form:

```
BEGIN
   DBMS_SPM.SET_EVOLVE_TASK_PARAMETER(
     task_name => 'SYS_AUTO_SPM_EVOLVE_TASK'
,   parameter => parameter_name
,   value => value
);
END;
//
```

For example, the following PL/SQL block configures the SYS_AUTO_SPM_EVOLVE_TASK task to automatically accept plans, seek up a maximum of 500 plans in the shared SQL area and AWR repository, and time out after 20 minutes:

```
BEGIN
  DBMS_SPM.SET_EVOLVE_TASK_PARAMETER(
    task_name => 'SYS_AUTO_SPM_EVOLVE_TASK'
,  parameter => 'TIME_LIMIT'
,  value => '1200'
);

  DBMS_SPM.SET_EVOLVE_TASK_PARAMETER(
    task_name => 'SYS_AUTO_SPM_EVOLVE_TASK'
,  parameter => 'ACCEPT_PLANS'
,  value => 'true'
);

  DBMS_SPM.SET_EVOLVE_TASK_PARAMETER(
    task_name => 'SYS_AUTO_SPM_EVOLVE_TASK'
,  parameter => 'ALTERNATE_PLAN_LIMIT'
,  value => '500'
);
```

```
END;
```

 Optionally, confirm your changes by querying the current parameter settings for SYS AUTO SPM EVOLVE TASK.

For example, execute the following query:

Sample output appears as follows:

See Also:

- Oracle Database PL/SQL Packages and Types Reference for complete reference information for DBMS SPM.SET EVOLVE TASK PARAMETER
- Oracle Database Reference to learn more about the DBA_ADVISOR_PARAMETERS

Automatic SQL Plan Management

A SQL plan baseline prevents performance regressions caused by suboptimal plans.

The High Frequency SPM Evolve Advisor Task can be configured to identify execution plan changes and compare the performance of the new plan with plans previously captured in the automatic SQL tuning set.

If a SQL statement does not have a SQL plan baseline, then the High Frequency SPM Evolve Advisor Task can sometimes resolve such performance regressions automatically. The advisor compares all available plans and chooses the best-performing plan as the baseline.

There are two modes of plan performance verification. In background-verification mode, the high frequency SPM evolve task does the verification. In real-time mode, plan performance verification is done by the foreground process immediately.

The SPM configuration parameter AUTO_SPM_EVOLVE_TASK controls Automatic SQL plan management. The setting ON is for background verification (using the task) and AUTO is for Real-Time SQL plan management.

The following figure shows the workflow for the High Frequency SPM Evolve Advisor Task in background verification mode:

AWR Look for alternative Check top plans SQL Automatic SQL **Tuning Set** Automatic SQL **Tuning Set** Compare performance Test execute Add alternative plans to and accept best plan alternative plans plan history **SQL Plan History** GB GB

Figure 29-2 High Frequency SPM Evolve Advisor Task

Whenever it runs, the High Frequency SPM Evolve Advisor Task performs the following tasks:

- Searches AWR and ASTS for resource-intensive SQL statements.
- 2. Looks for alternative plans in the automatic SQL tuning set.
- 3. If multiple plans are identified, they are added to the SQL plan history.
- **4.** Test executes and measures the performance of the plans alternatives.

The database test executes the statement and records the performance statistics.

- 5. Performs either of the following actions, depending on whether the alternative plan performs better than the current plan:
 - If performance is better, then High Frequency SPM Evolve Advisor accepts the plan. The alternative plan is now in the baseline.
 - If performance is worse, then the plan remains in the statement history, but not the baseline.

See Also:

- The Automatic SQL Tuning Set, the system-maintained record of SQL execution
 plans and SQL statement performance metrics seen by the database. The High
 Frequency SPM Evolve Advisor Task compares the performance of a new plan
 with plans previously captured in the automatic SQL tuning set.
- Oracle Database PL/SQL Packages and Types Reference to learn more about the DBMS SPM.SET EVOLVE TASK PARAMETER procedure

Configuring Automatic SQL Plan Management

The High-Frequency SPM Evolve Advisor task complements the standard Automatic SPM Evolve Advisor task and is a component of automatic SQL Plan Management.

This task runs periodically throughout the day.

Both the standard Automatic SPM Evolve Advisor task and high-frequency task have the same name: SYS_AUTO_SPM_EVOLVE_TASK. In DBA_ADVISOR_EXECUTIONS, the two tasks are distinguished by execution name. The name of the standard task execution has the form EXEC_number, whereas the name of the high-frequency execution has the form SYS_SPM_timestamp.

Either real-time SPM or SPM (with background-verification) is enabled by DBMS_SPM.CONFIGURE to set AUTO SPM EVOLVE TASK to ON (background) or AUTO (real-time).

DBMS_SPM.CONFIGURE enables the high-frequency task, but has no dependency on the SPM Evolve Advisor. The standard task and high-frequency task are independent and are scheduled through two different frameworks.

See Also:

- #unique_994
- See DBMS_SPM.CONFIGURE in the *Oracle Database PL/SQL Packages and Types Reference* to learn more about the DBMS_SPM.CONFIGURE procedure.

Enabling Automatic SQL Plan Management: Tutorial

To enable and disable automatic SQL plan management (automatic SPM), use the $\mbox{DBMS_SPM.CONFIGURE}$ procedure.

You can set auto_spm_evolve_task to any of the following values:

- ON Enables automatic SPM and the High-Frequency SPM Evolve Advisor Task with background-verified Automatic SPM.
- OFF Disables automatic SPM and the High-Frequency SPM Evolve Advisor task. This is the default.
- AUTO Enables Automatic SPM in real-time mode (running in the foreground).

Note that the task interval and run-time are fixed and cannot be adjusted by the user.



To enable the high-frequency SPM Evolve Advisor task:

Query the current setting for DBMS SPM.CONFIGURE (sample output included):

```
COL PARAMETER_NAME FORMAT a32
COL PARAMETER_VALUE FORMAT a32
SELECT PARAMETER_NAME, PARAMETER_VALUE
FROM DBA_SQL_MANAGEMENT_CONFIG
WHERE PARAMETER_NAME LIKE '%SPM%';
```

PARAMETER_NAME
PARAMETER_VALUE

AUTO_SPM_EVOLVE_TASK OFF
AUTO_SPM_EVOLVE_TASK_INTERVAL 3600

AUTO_SPM_EVOLVE_TASK_MAX_RUNTIME 1800

2. Enable the task.

Execute the following PL/SQL code:

```
EXEC DBMS SPM.CONFIGURE('AUTO SPM EVOLVE TASK', 'ON');
```

To confirm that the task is enabled, query the current setting for AUTO_SPM_EVOLVE_TASK (sample output included):

```
COL PARAMETER_NAME FORMAT a32

COL PARAMETER_VALUE FORMAT a32

SELECT PARAMETER_NAME, PARAMETER_VALUE

FROM DBA_SQL_MANAGEMENT_CONFIG

WHERE PARAMETER_NAME = 'AUTO_SPM_EVOLVE_TASK';

PARAMETER_NAME PARAMETER_VALUE

AUTO SPM_EVOLVE_TASK ON
```

4. Optionally, wait a few hours, and then query the status of the task executions:

```
SET LINESIZE 150

COL TASK_NAME FORMAT a30

COL EXECUTION_NAME FORMAT a30

SELECT TASK_NAME, EXECUTION_NAME, STATUS

FROM DBA_ADVISOR_EXECUTIONS

WHERE TASK_NAME LIKE '%SPM%'

AND (EXECUTION_NAME LIKE 'SYS_SPM%' OR EXECUTION_NAME LIKE 'EXEC_%')

ORDER BY EXECUTION_END;
```

TASK_NAME	EXECUTION_NAME	STATUS
SYS_AUTO_SPM_EVOLVE_TASK	SYS_SPM_2019-06-03/13:15:26	COMPLETED
SYS_AUTO_SPM_EVOLVE_TASK	SYS_SPM_2019-06-03/14:16:04	COMPLETED
SYS_AUTO_SPM_EVOLVE_TASK	EXEC_6	COMPLETED
SYS_AUTO_SPM_EVOLVE_TASK	SYS_SPM_2019-06-03/15:16:32	COMPLETED
SYS_AUTO_SPM_EVOLVE_TASK	SYS_SPM_2019-06-03/16:17:00	COMPLETED



In the preceding output, $EXEC_6$ is the execution name of the standard SPM Automatic Advisor task. The other executions are of the high-frequency task.

Monitoring and Reporting in Automatic SQL Plan Management

SQL plan baselines are shown in the view DBA_SQL_PLAN_BASELINES. The DBA SQL PLAN BASELINES.ORIGIN column in this view displays FOREGROUND-CAPTURE.

Displaying Plans in a SQL Plan Baseline

To view the plans stored in the SQL plan baseline for a specific statement, use the DBMS XPLAN.DISPLAY SQL PLAN BASELINE function.

This function uses plan information stored in the plan history to display the plans. The following table describes the relevant parameters for DBMS XPLAN.DISPLAY SQL PLAN BASELINE.

Table 29-4 DBMS_XPLAN.DISPLAY_SQL_PLAN_BASELINE Parameters

Function Parameter	Description
sql_handle	SQL handle of the statement. Retrieve the SQL handle by joining the V\$SQL.SQL_PLAN_BASELINE and DBA_SQL_PLAN_BASELINES views on the PLAN_NAME columns.
plan_name	Name of the plan for the statement.

This section explains how to show plans in a baseline from the command line. If you use Cloud Control, then display plan baselines from the SQL Plan Baseline subpage shown in Figure 29-1.

To display plans in a SQL plan baselines:

 Connect SQL*Plus to the database with the appropriate privileges, and then obtain the SQL ID of the query whose plan you want to display.

For example, assume that a SQL plan baseline exists for a SELECT statement with the SQL ID 31d96zzzpcys9.

Query the plan by SQL ID.

The following query displays execution plans for the statement with the SQL ID 31d96zzzpcys9:

The sample query results are as follows:

The results show that the plan for SQL ID 31d96zzzpcys is named SQL PLAN 52gvzja8jfysuc0e983c6 and was captured automatically.

See Also:

- "SQL Management Base"
- Oracle Database PL/SQL Packages and Types Reference to learn about additional parameters used by the DISPLAY SQL PLAN BASELINE function

Loading SQL Plan Baselines

Using <code>DBMS_SPM</code>, you can initiate the bulk load of a set of existing plans into a SQL plan baseline.

About Loading SQL Plan Baselines

The DBMS SPM package enables you to load plans from multiple sources.

The goal of this task is to load plans from the following sources:

AWR

Load plans from Automatic Workload Repository (AWR) snapshots. You must specify the beginning and ending of the snapshot range. Optionally, you can apply a filter to load only plan that meet specified criteria. By default, the optimizer uses the loaded plans the next time that the database executes the SQL statements.

Shared SQL area

Load plans for statements directly from the shared SQL area, which is in the shared pool of the SGA. By applying a filter on the module name, the schema, or the SQL ID you identify the SQL statement or set of SQL statements to capture. The optimizer uses the plans the next time that the database executes the SQL statements.

Loading plans directly from the shared SQL area is useful when application SQL has been hand-tuned using hints. Because you probably cannot change the SQL to include the hint, populating the SQL plan baseline ensures that the application SQL uses optimal plans.

SQL tuning set (STS)

Capture the plans for a SQL workload into an STS, and then load the plans into the SQL plan baselines. The optimizer uses the plans the next time that the database executes the SQL statements. Bulk loading execution plans from an STS is an effective way to prevent plan regressions after a database upgrade.

Staging table

Use the DBMS_SPM package to define a staging table, DBMS_SPM.PACK_STGTAB_BASELINE to copy the baselines into a staging table, and Oracle Data Pump to transfer the table to another database. On the destination database, use DBMS_SPM.UNPACK_STGTAB_BASELINE to unpack the plans from the staging table and put the baselines into the SMB.

A use case is the introduction of new SQL statements into the database from a new application module. A vendor can ship application software with SQL plan baselines for the new SQL. In this way, the new SQL uses plans that are known to give optimal performance under a standard test configuration. Alternatively, if you develop or test an application inhouse, export the correct plans from the test database and import them into the production database.

Stored outline

Migrate stored outlines to SQL plan baselines. After the migration, you maintain the same plan stability that you had using stored outlines while being able to use the more advanced features provided by SQL Plan Management, such as plan evolution. See .

See Also:

- "Migrating Stored Outlines to SQL Plan Baselines"
- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS_SPM.PACK_STGTAB_BASELINE Function

Loading Plans from AWR

This topic explains how to load plans from AWR using PL/SQL.

Load plans with the LOAD_PLANS_FROM_AWR function of the DBMS_SPM package. The following table describes some function parameters.

Table 29-5 LOAD_PLANS_FROM_AWR Parameters

Function Parameter	Description
begin_snap	Number of the beginning snapshot in the range. Required.



Table 29-5 (Cont.) LOAD_PLANS_FROM_AWR Parameters

Function Parameter	Description	
end_snap Number of the ending snapshot in the range. Required.		
basic_filter	A filter applied to AWR to select only qualifying plans to be loaded. The default null means that all plans in AWR are selected. The filter can take the form of any WHERE clause predicate that can be specified against the column DBA_HIST_SQLTEXT.SQL_TEXT. An example is basic_filter => 'sql_text like ''SELECT /*LOAD_STS*/%'''.	
fixed	Default NO means the loaded plans are used as nonfixed plans. YES means the loaded plans are fixed plans. "Plan Selection" explains that the optimizer chooses a fixed plan in the plan baseline over a nonfixed plan.	

This section explains how to load plans using the command line. In Cloud Control, go to the SQL Plan Baseline subpage (shown in Figure 29-1) and click **Load** to load plan baselines from AWR.

This tutorial assumes the following:

You want to load plans for the following query into the SMB:

```
SELECT /*LOAD_AWR*/ *
FROM sh.sales
WHERE quantity_sold > 40
ORDER BY prod id;
```

- You want the loaded plans to be nonfixed.
- The user sh has privileges to query DBA_HIST_SNAPSHOT and DBA_SQL_PLAN_BASELINES, execute DBMS_WORKLOAD_REPOSITORY.CREATE_SNAPSHOT, and execute DBMS SPM.LOAD PLANS FROM AWR.

To load plans from the shared SQL area:

 Log in to the database with the appropriate privileges, and then query the most recent 3 AWR snapshots.

For example, query DBA_HIST_SNAPSHOT as follows:

Query sh.sales, using the LOAD_AWR tag to identify the SQL statement.

For example, use the following query:

```
SELECT /*LOAD_AWR*/ *
FROM sh.sales
WHERE quantity_sold > 40
ORDER BY prod id;
```

3. Take a new AWR snapshot.

For example, use the following program:

```
EXEC DBMS WORKLOAD REPOSITORY.CREATE SNAPSHOT;
```

4. Query the most recent 3 AWR snapshots to confirm that a new snapshot was taken.

For example, query DBA HIST SNAPSHOT as follows:

5. Load the plans for the most recent 2 snapshots from AWR.

For example, execute the <code>LOAD_PLANS_FROM_AWR</code> function in SQL*Plus to load the plans from snapshot 212 to 213:

```
VARIABLE v_plan_cnt NUMBER
EXEC :v_plan_cnt := DBMS_SPM.LOAD_PLANS_FROM_AWR(begin_snap => 212,
end snap => 213);
```

In the preceding example, the variable v_plan_cnt contains the number of plans that were loaded

6. Query the data dictionary to ensure that the plans were loaded into the baseline for the LOAD_AWR statement.

The following statement queries DBA SQL PLAN BASELINES (sample output included):

```
COL SQL_HANDLE FORMAT a20
COL SQL_TEXT FORMAT a20
COL PLAN_NAME FORMAT a30
COL ORIGIN FORMAT a20

SELECT SQL_HANDLE, SQL_TEXT, PLAN_NAME,
ORIGIN, ENABLED, ACCEPTED
FROM DBA_SQL_PLAN_BASELINES
WHERE SQL_TEXT_LIKE '%LOAD AWR%';
```

```
SQL_HANDLE SQL_TEXT PLAN_NAME ORIGIN ENA ACC

SQL_495d29c5f4612cda SELECT /*LOAD_AWR SQL_PLAN_4kr99sru62b6u54bc MANUAL-LOAD- YES YES

*/ * FROM 8843 FROM-AWR

sh.sales WHERE
quantity_sold
> 40
ORDER BY prod_id
```

The output shows that the plan is accepted, which means that it is in the plan baseline for the statement. Also, the origin is MANUAL-LOAD-FROM-AWR, which means that the statement was loaded manually from AWR rather than automatically captured.

See Also:

- "Fixed Plans"
- Oracle Database PL/SQL Packages and Types Reference to learn how to use the DBMS SPM.LOAD PLANS FROM AWR function
- Oracle Database Reference to learn more about the DBA_SQL_PLAN_BASELINES view

Loading Plans from the Shared SQL Area

This topic explains how to load plans from the shared SQL area, also called the cursor cache, using PL/SQL.

Load plans with the LOAD_PLANS_FROM_CURSOR_CACHE function of the DBMS_SPM package. The following table describes some function parameters.

Table 29-6 LOAD_PLANS_FROM_CURSOR_CACHE Parameters

Function Parameter	Description
sql_id	SQL statement identifier. Identifies a SQL statement in the shared SQL area.
fixed	Default NO means the loaded plans are used as nonfixed plans. YES means the loaded plans are fixed plans. "Plan Selection" explains that the optimizer chooses a fixed plan in the plan baseline over a nonfixed plan.

This section explains how to load plans using the command line. In Cloud Control, go to the SQL Plan Baseline subpage (shown in Figure 29-1) and click **Load** to load plan baselines from the shared SQL area.

This tutorial assumes the following:

You have executed the following query:

```
SELECT /*LOAD_CC*/ *
FROM sh.sales
WHERE quantity_sold > 40
ORDER BY prod id;
```

You want the loaded plans to be nonfixed.

To load plans from the shared SQL area:

 Connect SQL*Plus to the database with the appropriate privileges, and then determine the SQL IDs of the relevant statements in the shared SQL area.

For example, query V\$SQL for the SQL ID of the sh.sales query (sample output included):

The preceding output shows that the SQL ID of the statement is 27m0sdw9snw59.

2. Load the plans for the specified statements into the SQL plan baseline.

For example, execute the LOAD_PLANS_FROM_CURSOR_CACHE function in SQL*Plus to load the plan for the statement with the SQL ID 27m0sdw9snw59:

```
VARIABLE v_plan_cnt NUMBER
BEGIN
   :v_plan_cnt:=DBMS_SPM.LOAD_PLANS_FROM_CURSOR_CACHE(
        sql_id => '27m0sdw9snw59');
END;
```

In the preceding example, the variable v_{plan_cnt} contains the number of plans that were loaded.

Query the data dictionary to ensure that the plans were loaded into the baseline for the statement.

The following statement queries DBA SQL PLAN BASELINES (sample output included):

The output shows that the plan is accepted, which means that it is in the plan baseline for the statement. Also, the origin is MANUAL-LOAD-FROM-CC, which means that the statement was loaded manually from the shared SQL area rather than automatically captured.

- "Fixed Plans"
- Oracle Database PL/SQL Packages and Types Reference to learn how to use the DBMS SPM.LOAD PLANS FROM CURSOR CACHE function
- Oracle Database Reference to learn more about the DBA_SQL_PLAN_BASELINES view

Loading Plans from a SQL Tuning Set

A **SQL tuning set** (STS) is a database object that includes one or more SQL statements, execution statistics, and execution context. This topic explains how to load plans from an STS.

Load plans with the <code>DBMS_SPM.LOAD_PLANS_FROM_SQLSET</code> function or using Cloud Control. The following table describes some function parameters.

Table 29-7 LOAD_PLANS_FROM_SQLSET Parameters

Function Parameter	Description
sqlset_name	Name of the STS from which the plans are loaded into SQL plan baselines.
basic_filter	A filter applied to the STS to select only qualifying plans to be loaded. The filter can take the form of any WHERE clause predicate that can be specified against the view DBA_SQLSET_STATEMENTS. An example is basic_filter => 'sql_text like ''SELECT /*LOAD_STS*/%'''.
fixed	Default ${\tt NO}$ means the loaded plans are used as nonfixed plans. YES means the loaded plans are fixed plans. "Plan Selection" explains that the optimizer chooses a fixed plan in the plan baseline over a nonfixed plan.

This section explains how to load plans from the command line. In Cloud Control, go to the SQL Plan Baseline subpage (shown in Figure 29-1) and click **Load** to load plan baselines from SQL tuning sets.

Assumptions

This tutorial assumes the following:

- You want the loaded plans to be nonfixed.
- You have executed the following query:

```
SELECT /*LOAD_STS*/ *
FROM sh.sales
WHERE quantity_sold > 40
ORDER BY prod_id;
```

- You have loaded the plan from the shared SQL area into the SQL tuning set named SPM_STS, which is owned by user SPM.
- After the operation, you want to drop the STS using <code>DBMS_SQLTUNE.DROP_SQLSET</code> rather than the equivalent <code>DBMS_SQLSET.DROP_SQLSET</code>.

To load plans from a SQL tuning set:

1. Connect SQL*Plus to the database with the appropriate privileges, and then verify which plans are in the SQL tuning set.

For example, query DBA SQLSET STATEMENTS for the STS name (sample output included):

The output shows that the plan for the select /*LOAD STS*/ statement is in the STS.

2. Load the plan from the STS into the SQL plan baseline.

For example, in SQL*Plus execute the function as follows:

The <code>basic_filter</code> parameter specifies a <code>WHERE</code> clause that loads only the plans for the queries of interest. The variable <code>v_plan_cnt</code> stores the number of plans loaded from the STS.

3. Query the data dictionary to ensure that the plan was loaded into the baseline for the statement.

The following statement queries the DBA_SQL_PLAN_BASELINES view (sample output included).

```
SQL> SELECT SQL_HANDLE, SQL_TEXT, PLAN_NAME,

2 ORIGIN, ENABLED, ACCEPTED

3 FROM DBA_SQL_PLAN_BASELINES;

SQL_HANDLE SQL_TEXT PLAN_NAME ORIGIN ENA ACC

SQL_a8632bd857a4a25e SELECT SQL_PLAN_ahstb wanual-Load-FROM-STS YES YES

/*Load_STS*/* v1bu98ky1694fc6b

FROM sh.sales
WHERE
quantity_sold
> 40 ORDER BY
prod_id
```

The output shows that the plan is accepted, which means that it is in the plan baseline. Also, the origin is MANUAL-LOAD-FROM-STS, which means that the plan was loaded manually from a SQL tuning set rather than automatically captured.

4. Optionally, drop the STS.

For example, execute <code>DBMS_SQLTUNE.DROP_SQLSET</code> to drop the <code>SPM_STS</code> tuning set as follows:

See Also:

- "Command-Line Interface to SQL Tuning Sets"
- Oracle Database Reference to learn about the DBA_SQL_PLAN_BASELINES view
- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS SPM.LOAD PLANS FROM SQLSET function

Loading Plans from a Staging Table

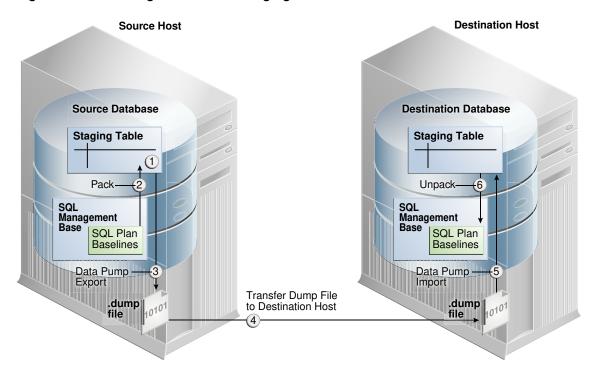
You may want to transfer optimal plans from a source database to a different destination database.

For example, you may have investigated a set of plans on a test database and confirmed that they have performed well. You may then want to load these plans into a production database.

A staging table is a table that, for the duration of its existence, stores plans so that the plans do not disappear from the table while you are unpacking them. Use the <code>DBMS_SPM.CREATE_STGTAB_BASELINE</code> procedure to create a staging table. To pack (insert row into) and unpack (extract rows from) the staging table, use the <code>PACK_STGTAB_BASELINE</code> and <code>UNPACK_STGTAB_BASELINE</code> functions of the <code>DBMS_SPM</code> package. Oracle <code>Data Pump Import</code> and <code>Export</code> enable you to copy the staging table to a different database.



Figure 29-3 Loading Plans from a Staging Table



Export plans with the PACK_STGTAB_BASELINE function of the DBMS_SPM package, and then import them with UNPACK_STGTAB_BASELINE. The following table describes some function parameters.

Table 29-8 PACK_STGTAB_BASELINE and UNPACK_STGTAB_BASELINE Parameters

Function Parameter	Description
table_name	Specifies the table to be imported or exported.
origin	Origin of SQL plan baseline. These procedures accept all possible values of DBA_SQL_PLAN_BASELINES.ORIGIN as the origin argument. For example, the origin parameter permits MANUAL-LOAD-FROM-STS, MANUAL-LOAD-FROM-AWR, and MANUAL-LOAD-FROM-CC.

This tutorial assumes the following:

- You want to create a staging table named stage1 in the source database.
- You want to load all plans owned by user spm into the staging table.
- You want to transfer the staging table to a destination database.
- You want to load the plans in stage1 as fixed plans.

To transfer a set of SQL plan baselines from one database to another:

1. Connect SQL*Plus to the source database with the appropriate privileges, and then create a staging table using the CREATE STGTAB BASELINE procedure.



The following example creates a staging table named stage1:

```
BEGIN
   DBMS_SPM.CREATE_STGTAB_BASELINE (
    table_name => 'stage1');
END;
/
```

On the source database, pack the SQL plan baselines you want to export from the SQL management base into the staging table.

The following example packs enabled plan baselines created by user spm into staging table stage1. Select SQL plan baselines using the plan name (plan_name), SQL handle (sql handle), or any other plan criteria. The table name parameter is mandatory.

```
DECLARE
  v_plan_cnt NUMBER;
BEGIN
  v_plan_cnt := DBMS_SPM.PACK_STGTAB_BASELINE (
    table_name => 'stage1'
,    enabled => 'yes'
,    creator => 'spm'
);
END;
//
```

- 3. Export the staging table stage1 into a dump file using Oracle Data Pump Export.
- Transfer the dump file to the host of the destination database.
- 5. On the destination database, import the staging table stage1 from the dump file using the Oracle Data Pump Import utility.
- 6. On the destination database, unpack the SQL plan baselines from the staging table into the SQL management base.

The following example unpacks all fixed plan baselines stored in the staging table stage1:

```
DECLARE
  v_plan_cnt NUMBER;
BEGIN
  v_plan_cnt := DBMS_SPM.UNPACK_STGTAB_BASELINE (
    table_name => 'stage1'
, fixed => 'yes'
);
END;
//
```

See Also:

- Oracle Database PL/SQL Packages and Types Reference for more information about using the DBMS SPM package
- Oracle Database Reference to learn more about the DBA_SQL_PLAN_BASELINES view
- Oracle Database Utilities for detailed information about using the Data Pump Export and Import utilities

Evolving SQL Plan Baselines Manually

You can use PL/SQL or Cloud Control to manually evolve an unaccepted plan to determine whether it performs better than any plan currently in the plan baseline.



"Managing the SPM Evolve Advisor Task"

About the DBMS_SPM Evolve Functions

This topic describes the most relevant DBMS_SPM functions for managing plan evolution. Execute evolution tasks manually or schedule them to run automatically.

Table 29-9 DBMS_SPM Functions and Procedures for Managing Plan Evolution Tasks

Procedure or Function	Description
ACCEPT_SQL_PLAN_BASELINE	This function accepts one recommendation to evolve a single plan into a SQL plan baseline.
CREATE_EVOLVE_TASK	This function creates an advisor task to prepare the plan evolution of one or more plans for a specified SQL statement. The input parameters can be a SQL handle, plan name or a list of plan names, time limit, task name, and description.
EXECUTE_EVOLVE_TASK	This function executes an evolution task. The input parameters can be the task name, execution name, and execution description. If not specified, the advisor generates the name, which is returned by the function.
IMPLEMENT_EVOLVE_TASK	This function implements all recommendations for an evolve task. Essentially, this function is equivalent to using ACCEPT_SQL_PLAN_BASELINE for all recommended plans. Input parameters include task name, plan name, owner name, and execution name.
REPORT_EVOLVE_TASK	This function displays the results of an evolve task as a CLOB. Input parameters include the task name and section of the report to include.

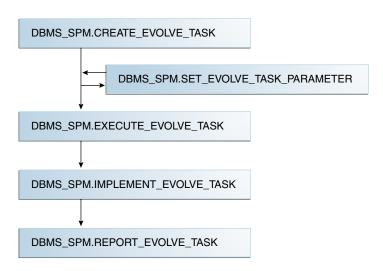


Table 29-9 (Cont.) DBMS_SPM Functions and Procedures for Managing Plan Evolution Tasks

Procedure or Function	Description	
SET_EVOLVE_TASK_PARAMETER	This function updates the value of an evolve task parameter.	

Oracle recommends that you configure SPM Evolve Advisor to run automatically. You can also evolve SQL plan baselines manually. The following graphic shows the basic workflow for managing SQL plan management tasks.

Figure 29-4 Evolving SQL Plan Baselines



Typically, you manage SQL plan evolution tasks in the following sequence:

- 1. Create an evolve task
- 2. Optionally, set evolve task parameters
- 3. Execute the evolve task
- Implement the recommendations in the task
- 5. Report on the task outcome

See Also:

- "Configuring the Automatic SPM Evolve Advisor Task" to learn about SET_EVOLVE_TASK_PARAMETER
- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS_SPM package

Managing an Evolve Task

This topic describes a typical use case in which you create and execute a task, and then implement its recommendations.

The following table describes some parameters of the CREATE EVOLVE TASK function.

Table 29-10 DBMS_SPM.CREATE_EVOLVE_TASK Parameters

Function Parameter	Description
sql_handle	SQL handle of the statement. The default ${\tt NULL}$ considers all SQL statements with unaccepted plans.
plan_name	Plan identifier. The default NULL means consider all unaccepted plans of the specified SQL handle or all SQL statements if the SQL handle is NULL.
time_limit	Time limit in number of minutes. The time limit for first unaccepted plan equals the input value. The time limit for the second unaccepted plan equals the input value minus the time spent in first plan verification, and so on. The default <code>DBMS_SPM.AUTO_LIMIT</code> means let the system choose an appropriate time limit based on the number of plan verifications required to be done.
task_name	User-specified name of the evolution task.

This section explains how to evolve plan baselines from the command line. In Cloud Control, from the SQL Plan Baseline subpage, select a plan, and then click **Evolve**.

This tutorial assumes the following:

- You do not have the automatic evolve task enabled.
- You want to create a SQL plan baseline for the following query:

```
SELECT /* q1_group_by */ prod_name, sum(quantity_sold)
FROM products p, sales s
WHERE p.prod_id = s.prod_id
AND p.prod_category_id =203
GROUP BY prod name;
```

You want to create two indexes to improve the query performance, and then evolve the
plan that uses these indexes if it performs better than the plan currently in the plan
baseline.

To evolve a specified plan:

- Perform the initial setup as follows:
 - a. Connect SQL*Plus to the database with administrator privileges, and then prepare for the tutorial by flushing the shared pool and the buffer cache:

```
ALTER SYSTEM FLUSH SHARED_POOL;
ALTER SYSTEM FLUSH BUFFER CACHE;
```

b. Enable the automatic capture of SQL plan baselines.

For example, enter the following statement:

```
ALTER SYSTEM SET OPTIMIZER CAPTURE SQL PLAN BASELINES=true;
```

c. Connect to the database as user sh, and then set SQL*Plus display parameters:

```
CONNECT sh
-- enter password

SET PAGES 10000 LINES 140

SET SERVEROUTPUT ON

COL SQL_TEXT FORMAT A20

COL SQL_HANDLE FORMAT A30

COL ORIGIN FORMAT A12

SET LONGC 60535

SET LONG 60535

SET ECHO ON
```

- Execute the SELECT statements so that SQL plan management captures them:
 - a. Execute the SELECT /* q1 group by */ statement for the first time.

Because the database only captures plans for repeatable statements, the plan baseline for this statement is empty.

b. Query the data dictionary to confirm that no plans exist in the plan baseline.

For example, execute the following query (sample output included):

```
SELECT SQL_HANDLE, SQL_TEXT, PLAN_NAME, ORIGIN, ENABLED,
ACCEPTED, FIXED, AUTOPURGE

FROM DBA_SQL_PLAN_BASELINES
WHERE SQL_TEXT LIKE '%q1_group%';

no rows selected
```

SQL plan management only captures repeatable statements, so this result is expected.

- c. Execute the SELECT /* q1 group by */ statement for the second time.
- Query the data dictionary to ensure that the plans were loaded into the plan baseline for the statement.

The following statement queries DBA SQL PLAN BASELINES (sample output included):



```
sales s WHERE
p.prod_id =
s.prod_id AND
p.prod_category
_id =203 GROUP
BY prod_name
```

The output shows that the plan is accepted, which means that it is in the plan baseline for the statement. Also, the origin is AUTO-CAPTURE, which means that the statement was automatically captured and not manually loaded.

4. Explain the plan for the statement and verify that the optimizer is using this plan.

For example, explain the plan as follows, and then display it:

```
EXPLAIN PLAN FOR
  SELECT /* q1_group_by */ prod_name, sum(quantity_sold)
  FROM   products p, sales s
  WHERE  p.prod_id = s.prod_id
  AND   p.prod_category_id =203
  GROUP BY prod_name;

SELECT * FROM TABLE(DBMS XPLAN.DISPLAY(null, null, 'basic +note'));
```

Sample output appears below:

```
Plan hash value: 1117033222
```

I	Id	1	Operation		Name	I
	U		SELECT STATEMENT			
	1		HASH GROUP BY			
	2		HASH JOIN			
	3		TABLE ACCESS FULL		PRODUCTS	
	4		PARTITION RANGE ALL			
	5		TABLE ACCESS FULL		SALES	

Note

```
- SQL plan baseline "SQL_PLAN_0gwbcfvzskcu242949306" used for this statement
```

The note indicates that the optimizer is using the plan shown with the plan name listed in the previous step.

Create two indexes to improve the performance of the SELECT /* q1_group_by */ statement.

For example, use the following statements:

```
CREATE INDEX ind_prod_cat_name
  ON products(prod category id, prod name, prod id);
```

```
CREATE INDEX ind_sales_prod_qty_sold
  ON sales(prod id, quantity sold);
```

6. Execute the select /* q1 group by */ statement again.

Because automatic capture is enabled, the plan baseline is populated with the new plan for this statement.

Query the data dictionary to ensure that the plan was loaded into the SQL plan baseline for the statement.

The following statement queries DBA SQL PLAN BASELINES (sample output included).

```
SELECT SQL HANDLE, SQL TEXT, PLAN NAME, ORIGIN, ENABLED, ACCEPTED
FROM DBA SQL PLAN BASELINES
WHERE SQL HANDLE IN ('SQL 07f16c76ff893342')
ORDER BY SQL HANDLE, ACCEPTED;
SQL_HANDLE SQL_TEXT PLAN_NAME
                                                            ORIGIN ENA ACC
SQL 07f16c76ff893342 SELECT /* q1 group b SQL PLAN 0qwbcfvzskcu2 AUTO-CAPTURE YES NO
                   y */ prod name, sum( 0135fd6c
                   quantity sold)
                   FROM products p, s
                   ales s
                   WHERE p.prod_id = s
                    .prod id
                   AND p.prod catego
                   ry id = 203
                   GROUP BY prod name
SQL 07f16c76ff893342 SELECT /* q1 group b SQL PLAN 0qwbcfvzskcu2 AUTO-CAPTURE YES YES
                   y */ prod name, sum( 42949306
                   quantity sold)
                   FROM products p, s
                   WHERE p.prod_id = s
                    .prod id
                   AND
                        p.prod catego
                   ry id =203
                   GROUP BY prod name
```

The output shows that the new plan is unaccepted, which means that it is in the statement history but not the SQL plan baseline.

8. Explain the plan for the statement and verify that the optimizer is using the original unindexed plan.

For example, explain the plan as follows, and then display it:

```
EXPLAIN PLAN FOR
  SELECT /* q1_group_by */ prod_name, sum(quantity_sold)
  FROM    products p, sales s
  WHERE    p.prod_id = s.prod_id
  AND        p.prod_category_id =203
    GROUP BY prod_name;
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY(null, null, 'basic +note'));
```

Sample output appears below:

Plan hash value: 1117033222

1	Id		Operation	- 	Name	
	0		SELECT STATEMENT			
	1		HASH GROUP BY			
	2		HASH JOIN			
	3		TABLE ACCESS FULL		PRODUCTS	
	4		PARTITION RANGE ALL			
	5		TABLE ACCESS FULL		SALES	

Note

- SQL plan baseline "SQL_PLAN_0gwbcfvzskcu242949306" used for this statement

The note indicates that the optimizer is using the plan shown with the plan name listed in Step 3.

Connect as an administrator, and then create an evolve task that considers all SQL statements with unaccepted plans.

For example, execute the <code>DBMS_SPM.CREATE_EVOLVE_TASK</code> function and then obtain the name of the task:

```
CONNECT / AS SYSDBA
VARIABLE cnt NUMBER
VARIABLE tk_name VARCHAR2(50)
VARIABLE exe_name VARCHAR2(50)
VARIABLE evol_out CLOB

EXECUTE :tk_name := DBMS_SPM.CREATE_EVOLVE_TASK(
   sql_handle => 'SQL_07f16c76ff893342',
   plan_name => 'SQL_PLAN_0gwbcfvzskcu20135fd6c');

SELECT :tk name FROM DUAL;
```

The following sample output shows the name of the task:

```
:EVOL_OUT
-----TASK_11
```

Now that the task has been created and has a unique name, execute the task.

10. Execute the task.

For example, execute the <code>DBMS_SPM.EXECUTE_EVOLVE_TASK</code> function (sample output included):

```
EXECUTE :exe_name :=DBMS_SPM.EXECUTE_EVOLVE_TASK(task_name=>:tk_name);
SELECT :exe_name FROM DUAL;
```

```
:EXE_NAME
-----EXEC 1
```

11. View the report.

For example, execute the <code>DBMS_SPM.REPORT_EVOLVE_TASK</code> function (sample output included):

```
EXECUTE :evol_out := DBMS_SPM.REPORT_EVOLVE_TASK( task_name=>:tk_name,
execution_name=>:exe_name );
SELECT :evol_out FROM DUAL;
GENERAL INFORMATION SECTION
```

Task Information:

Task Name : TASK_11
Task Owner : SYS
Execution Name : EXEC_1
Execution Type : SPM EVOLVE
Scope : COMPREHENSIVE
Status : COMPLETED

 Started
 : 01/09/2012 12:21:27

 Finished
 : 01/09/2012 12:21:29

 Last Updated
 : 01/09/2012 12:21:29

Global Time Limit : 2147483646
Per-Plan Time Limit : UNUSED
Number of Errors : 0

SUMMARY SECTION

Number of plans processed : 1
Number of findings : 1
Number of recommendations : 1
Number of errors : 0

DETAILS SECTION

Object ID : 2

Test Plan Name : SQL_PLAN_0gwbcfvzskcu20135fd6c
Base Plan Name : SQL_PLAN_0gwbcfvzskcu242949306

SQL Handle : SQL 07f16c76ff893342

Parsing Schema : SH Test Plan Creator : SH

SQL Text : SELECT /*ql group by*/ prod name,

sum(quantity_sold)
FROM products p, sales s

WHERE p.prod id=s.prod id AND p.prod category id=203

GROUP BY prod name

Execution Statistics:

	Base Plan	Test Plan
Elapsed Time (s):	.044336	.012649
CPU Time (s):	.044003	.012445
Buffer Gets:	360	99
Optimizer Cost:	924	891
Disk Reads:	341	82
Direct Writes:	0	0
Rows Processed:	4	2
Executions:	5	9

FINDINGS SECTION

Findings (1):

1. The plan was verified in 2.18 seconds. It passed the benefit criterion because its verified performance was 2.01 times better than that of the baseline plan.

Recommendation:

```
Consider accepting the plan. Execute
dbms_spm.accept_sql_plan_baseline(task_name => 'TASK_11',
object_id => 2, task_owner => 'SYS');
```

EXPLAIN PLANS SECTION

Baseline Plan

Plan Id : 1

Plan Hash Value : 1117033222

Id Operation	Name	Rows Bytes Cost	Time
0 SELECT STATEMENT		21 861 924	00:00:12
1 HASH GROUP BY		21 861 924	00:00:12
*2 HASH JOIN		267996 10987836 742	00:00:09
*3 TABLE ACCESS FULL	PRODUCT	S 21 714 2	00:00:01
4 PARTITION RANGE ALL		918843 6431901 662	00:00:08
5 TABLE ACCESS FULL	SALES	918843 6431901 662	00:00:08

Predicate Information (identified by operation id):

```
* 2 - access("P"."PROD ID"="S"."PROD ID")
```

Test Plan

Plan Id : 2

Plan Hash Value : 20315500

^{* 3 -} filter("P"."PROD CATEGORY ID"=203)

This report indicates that the new execution plan, which uses the two new indexes, performs better than the original plan.

12. Implement the recommendations of the evolve task.

For example, execute the IMPLEMENT EVOLVE TASK function:

```
BEGIN
  :cnt := DBMS_SPM.IMPLEMENT_EVOLVE_TASK(
    task_name=>:tk_name, execution_name=>:exe_name);
END;
```

13. Query the data dictionary to ensure that the new plan is accepted.

The query provides the following sample output:

```
SELECT SQL HANDLE, SQL TEXT, PLAN NAME, ORIGIN, ENABLED, ACCEPTED
FROM DBA SQL PLAN BASELINES
WHERE SQL HANDLE IN ('SQL 07f16c76ff893342')
ORDER BY SQL HANDLE, ACCEPTED;
           SQL_TEXT PLAN NAME
SQL HANDLE
SQL 07f16c76ff893342 SELECT /* q1 group b SQL PLAN 0gwbcfvzskcu2 AUTO-CAPTURE YES YES
                    y */ prod name, sum ( 0135fd6c
                    quantity sold)
                    FROM products p, s
                    ales s
                    WHERE p.prod id = s
                    .prod id
                    AND
                         p.prod catego
                    ry id = 203
                    GROUP BY prod name
SQL 07f16c76ff893342 SELECT /* q1 group b SQL PLAN 0gwbcfvzskcu2 AUTO-CAPTURE YES YES
                    y */ prod name, sum( 42949306
                    quantity sold)
                    FROM products p, s
                    ales s
                    WHERE p.prod id = s
                    .prod id
```

```
AND p.prod_catego
ry_id =203
GROUP BY prod_name
```

The output shows that the new plan is accepted.

14. Clean up after the example.

For example, enter the following statements:

```
EXEC :cnt := DBMS_SPM.DROP_SQL_PLAN_BASELINE('SQL_07f16c76ff893342');
EXEC :cnt := DBMS_SPM.DROP_SQL_PLAN_BASELINE('SQL_9049245213a986b3');
EXEC :cnt := DBMS_SPM.DROP_SQL_PLAN_BASELINE('SQL_bb77077f5f90a36b');
EXEC :cnt := DBMS_SPM.DROP_SQL_PLAN_BASELINE('SQL_02a86218930bbb20');
DELETE FROM SQLLOG$;
CONNECT sh
-- enter password
DROP INDEX IND_SALES_PROD_QTY_SOLD;
DROP INDEX IND_PROD_CAT_NAME;
```

See Also:

- "Managing the SPM Evolve Advisor Task"
- Oracle Database PL/SQL Packages and Types Reference to learn more about the DBMS SPM evolve functions

Adding SQL Plans to a Baseline

You can use DBMS SPM.ADD VERIFIED SQL PLAN BASELINE to add SQL plans to a baseline.

This function fetches all existing plans for a given SQL ID. It chooses the best performing plans and accepts them. The sources may include the Cursor Cache, Automatic Workload Repository, and Auto SQL Tuning Set (SYS_AUTO_STS).

These are the steps that DBMS SPM.ADD VERIFIED SQL PLAN BASELINE executes:

- Loads plans from all sources into SQL plan management SQL plan history in a nonaccepted state.
- Uses the SQL Plan Management Evolve Advisor internally to identify the best-performing execution plans. The best-performing plans are accepted.

When the function has completed, it will have accepted the best plans among the plans available from the Cursor Cache, AWR, and the Automatic SQL Tuning Set.

Syntax

```
DBMS_SPM.ADD_VERIFIED_SQL_PLAN_BASELINE (
     sql_id IN VARCHAR
);
```



Usage

```
var report clob
exec :report := dbms_spm.add_verified_sql_plan_baseline('fazrk33ng71km');
select :report report from dual;
```



As of Oracle DB 23ai, for SELECTS that do not require table access, FROM DUAL is now implicit when there is no FROM clause. FROM DUAL is supported but is no longer required. SELECT $\langle expr_list \rangle$; is sufficient.

Required Privileges

- Administer SQL Management Object
- Advisor
- Select on dba sql plan baselines

Table 29-11 ADD_VERIFIED_SQL_PLAN_BASELINE Function Parameters

Parameter	Description
sql_id	The SQL statement identifier, which is used to identify the plans in different sources, such as the cursor cache.



The Database Packages and Types Reference provides further details and usage examples for ADD_VERIFIED _SQL_PLAN_BASELINE.

Dropping SQL Plan Baselines

You can remove some or all plans from a SQL plan baseline. This technique is sometimes useful when testing SQL plan management.

Drop plans with the DBMS_SPM.DROP_SQL_PLAN_BASELINE function. This function returns the number of dropped plans. The following table describes input parameters.

Table 29-12 DROP_SQL_PLAN_BASELINE Parameters

Function Parameter	Description
sql_handle	SQL statement identifier.
plan_name	Name of a specific plan. Default NULL drops all plans associated with the SQL statement identified by sql_handle.

This section explains how to drop baselines from the command line. In Cloud Control, from the SQL Plan Baseline subpage, select a plan, and then click **Drop**.

This tutorial assumes that you want to drop all plans for the following SQL statement, effectively dropping the SQL plan baseline:

```
SELECT /* repeatable sql */ COUNT(*) FROM hr.jobs;
```

To drop a SQL plan baseline:

 Connect SQL*Plus to the database with the appropriate privileges, and then query the data dictionary for the plan baseline.

The following statement queries DBA_SQL_PLAN_BASELINES (sample output included):

```
SQL> SELECT SQL_HANDLE, SQL_TEXT, PLAN_NAME, ORIGIN,

2 ENABLED, ACCEPTED

3 FROM DBA_SQL_PLAN_BASELINES

4 WHERE SQL_TEXT LIKE 'SELECT /* repeatable_sql%';

SQL_HANDLE SQL_TEXT PLAN_NAME ORIGIN ENA ACC

SQL_b6b0d1c71cd1807b SELECT /* repeatable SQL_PLAN_bdc6jswfd303v AUTO-CAPTURE YES YES

_sql */ count(*) 2fle9c20
from hr.jobs
```

2. Drop the SQL plan baseline for the statement.

The following example drops the plan baseline with the SQL handle $SQL_b6b0d1c71cd1807b$, and returns the number of dropped plans. Specify plan baselines using the plan name (plan_name), SQL handle (sql_handle), or any other plan criteria. The table name parameter is mandatory.

```
DECLARE
   v_dropped_plans number;
BEGIN
   v_dropped_plans := DBMS_SPM.DROP_SQL_PLAN_BASELINE (
        sql_handle => 'SQL_b6b0d1c71cd1807b'
);
   DBMS_OUTPUT.PUT_LINE('dropped ' || v_dropped_plans || ' plans');
END;
//
```

Confirm that the plans were dropped.

For example, execute the following query:

```
SELECT SQL HANDLE, SQL_TEXT, PLAN_NAME, ORIGIN,
ENABLED, ACCEPTED

FROM DBA_SQL_PLAN_BASELINES
WHERE SQL_TEXT_LIKE 'SELECT /* repeatable_sql%';
no rows selected
```

See Also:

Oracle Database PL/SQL Packages and Types Reference to learn about the ${\tt DROP_SQL_PLAN_BASELINE}$ function

Managing the SQL Management Base

The SQL management base is a part of the data dictionary that resides in the SYSAUX tablespace. It stores statement logs, plan histories, SQL plan baselines, and SQL profiles.

About Managing the SMB

Use the <code>DBMS_SPM.CONFIGURE</code> procedure to set configuration options for the SMB and the maintenance of SQL plan baselines.

The DBA_SQL_MANAGEMENT_CONFIG view shows the current configuration settings for the SMB. The following table describes the parameters in the PARAMETER NAME column.

Table 29-13 Parameters in DBA_SQL_MANAGEMENT_CONFIG.PARAMETER_NAME

Parameter	Description
SPACE_BUDGET_PERCENT	Maximum percent of SYSAUX space that the SQL management base can use. The default is 10. The allowable range for this limit is between 1% and 50%.
PLAN_RETENTION_WEEKS	Number of weeks to retain unused plans before they are purged. The default is 53.
AUTO_CAPTURE_PARSING_SCHEMA_ NAME	A list of the form ($%$ LIKE a OR $%$ LIKE b) AND ($%$ NOT LIKE c AND $%$ NOT LIKE d), which is the internal representation of the parsing schema name filter. If no parsing schema filters exist, then one side of the outer conjunction will be absent.
AUTO_CAPTURE_MODULE	A list of the form ($%$ LIKE a OR $%$ LIKE b) AND ($%$ NOT LIKE c AND $%$ NOT LIKE d), which is the internal representation of the module filter. If no module filters exist, then one side of the outer conjunction will be absent.
AUTO_CAPTURE_ACTION	A list of the form (% LIKE a OR % LIKE b) AND (% NOT LIKE c AND % NOT LIKE d), which is the internal representation of the action filter. If no action filters exist, then one side of the outer conjunction will be absent.
AUTO_CAPTURE_SQL_TEXT	A list of the form ($%$ LIKE a OR $%$ LIKE b) AND ($%$ NOT LIKE c AND $%$ NOT LIKE d), which is the internal representation of the SQL text filter. If no SQL text filters exist, then one side of the outer conjunction will be absent.



See Also:

- "Eligibility for Automatic Initial Plan Capture"
- Oracle Database Reference to learn more about DBA SQL MANAGEMENT CONFIG
- Oracle Database PL/SQL Packages and Types Reference to learn more about DBMS SPM.CONFIGURE

Changing the Disk Space Limit for the SMB

A weekly background process measures the total space occupied by the SMB.

When the defined limit is exceeded, the process writes a warning to the alert log. The database generates alerts weekly until either the SMB space limit is increased, the size of the SYSAUX tablespace is increased, or the disk space used by the SMB is decreased by purging SQL management objects (SQL plan baselines or SQL profiles). This task explains how to change the limit with the DBMS SPM.CONFIGURE procedure.

Assumptions

This tutorial assumes the following:

- The current SMB space limit is the default of 10%.
- You want to change the percentage limit to 30%

To change the percentage limit of the SMB:

 Connect SQL*Plus to the database with the appropriate privileges, and then query the data dictionary to see the current space budget percent.

For example, execute the following query (sample output included):

2. Change the percentage setting.

For example, execute the following command to change the setting to 30%:

```
EXECUTE DBMS SPM.CONFIGURE('space budget percent', 30);
```

3. Query the data dictionary to confirm the change.

For example, execute the following join (sample output included):

See Also:

Oracle Database PL/SQL Packages and Types Reference to learn more about the DBMS SPM.CONFIGURE procedure

Changing the Plan Retention Policy in the SMB

A weekly scheduled purging task manages disk space used by SQL plan management.

The task runs as an automated task in the maintenance window. The database purges plans that have not been used for longer than the plan retention period, as identified by the LAST_EXECUTED timestamp stored in the SMB for that plan. The default retention period is 53 weeks. The period can range between 5 and 523 weeks.

This task explains how to change the plan retention period with the DBMS_SPM.CONFIGURE procedure. In Cloud Control, set the plan retention policy in the SQL Plan Baseline subpage (shown in Figure 29-1).

To change the plan retention period for the SMB:

 Connect SQL*Plus to the database with the appropriate privileges, and then query the data dictionary to see the current plan retention period.

For example, execute the following query (sample output included):

```
SQL> SELECT PARAMETER_NAME, PARAMETER_VALUE

2 FROM DBA_SQL_MANAGEMENT_CONFIG

3 WHERE PARAMETER_NAME = 'PLAN_RETENTION_WEEKS';

PARAMETER_NAME PARAMETER_VALUE

PLAN RETENTION WEEKS 53
```

2. Change the retention period.

For example, execute the CONFIGURE procedure to change the period to 105 weeks:

```
EXECUTE DBMS SPM.CONFIGURE('plan retention weeks',105);
```

3. Query the data dictionary to confirm the change.

For example, execute the following query:

See Also:

Oracle Database PL/SQL Packages and Types Reference to learn more about the ${\tt CONFIGURE}$ procedure

