# Explaining and Displaying Execution Plans

Knowledge of how to explain a statement and display its plan is essential to SQL tuning.

## Introduction to Execution Plans

An **execution plan** is the sequence of operations that the database performs to run a SQL statement.

## Contents of an Execution Plan

Plan hash value: 1219589317

The execution plan operation alone cannot differentiate between well-tuned statements and those that perform suboptimally.

The plan consists of a series of steps. Every step either retrieves rows of data physically from the database or prepares them for the user issuing the statement. The following plan shows a join of the <code>employees</code> and <code>departments</code> tables:

| 2 | TABLE ACCESS BY INDEX ROWID| EMPLOYEES | 5 | 110 | 2 (0) | 00:00:01 | | \* 3 | INDEX RANGE SCAN | EMP\_NAME\_IX | 5 | 1 (0) | 00:00:01 | | \* 4 | TABLE ACCESS FULL | DEPARTMENTS | 1 | 16 | 1 (0) | 00:00:01 |

Predicate Information (identified by operation id):

```
3 - access("LAST_NAME" LIKE 'T%')
    filter("LAST_NAME" LIKE 'T%')
4 - filter("E"."DEPARTMENT ID"="D"."DEPARTMENT ID")
```

The row source tree is the core of the execution plan. The tree shows the following information:

The join order of the tables referenced by the statement
 In the preceding plan, employees is the outer row source and departments is the inner row source.

An access path for each table mentioned in the statement

In the preceding plan, the optimizer chooses to access employees using an index scan and departments using a full scan.

A join method for tables affected by join operations in the statement
 In the preceding plan, the optimizer chooses a nested loops join.

· Data operations like filter, sort, or aggregation

In the preceding plan, the optimizer filters on last names that begin with  $\ensuremath{\mathtt{T}}$  and matches on department id.

In addition to the row source tree, the plan table contains information about the following:

- Optimization, such as the cost and cardinality of each operation
- Partitioning, such as the set of accessed partitions
- Parallel execution, such as the distribution method of join inputs

# Why Execution Plans Change

Execution plans can and do change as the underlying optimizer inputs change.



To avoid possible SQL performance regression that may result from execution plan changes, consider using SQL plan management.

## See Also:

- "Overview of SQL Plan Management"
- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS SPM package

### Different Schemas

Schemas can differ for various reasons.

Principal reasons include the following:

- The execution and explain plan occur on different databases.
- The user explaining the statement is different from the user running the statement. Two
  users might be pointing to different objects in the same database, resulting in different
  execution plans.
- Schema changes (often changes in indexes) between the two operations.

### Different Costs

Even if the schemas are the same, the optimizer can choose different execution plans when the costs are different.

Some factors that affect the costs include the following:

- Data volume and statistics
- Bind variable types and values
- Initialization parameters set globally or at session level

# Generating Plan Output Using the EXPLAIN PLAN Statement

The EXPLAIN PLAN statement enables you to examine the execution plan that the optimizer chose for a SQL statement.

## About the EXPLAIN PLAN Statement

The EXPLAIN PLAN statement displays execution plans that the optimizer chooses for SELECT, UPDATE, INSERT, and DELETE statements.

EXPLAIN PLAN output shows how the database would have run the SQL statement when the statement was explained. Because of differences in the execution environment and explain plan environment, the explained plan can differ from the actual plan used during statement execution.

When the EXPLAIN PLAN statement is issued, the optimizer chooses an execution plan and then inserts a row describing each step of the execution plan into a specified plan table. You can also issue the EXPLAIN PLAN statement as part of the SQL trace facility.

The EXPLAIN PLAN statement is a DML statement rather than a DDL statement. Therefore, Oracle Database does not implicitly commit the changes made by an EXPLAIN PLAN statement.

### See Also:

- "SOL Row Source Generation"
- Oracle Database SQL Language Reference to learn about the EXPLAIN PLAN statement

## About PLAN TABLE

PLAN\_TABLE is the default sample output table into which the EXPLAIN PLAN statement inserts rows describing execution plans.

Oracle Database automatically creates a global temporary table <code>PLAN\_TABLE\$</code> in the <code>SYS</code> schema, and creates <code>PLAN\_TABLE</code> as a synonym. All necessary privileges to <code>PLAN\_TABLE</code> are granted to <code>PUBLIC</code>. Consequently, every session gets its own private copy of <code>PLAN\_TABLE</code> in its temporary tablespace.

You can use the SQL script <code>catplan.sql</code> to manually create the global temporary table and the <code>PLAN\_TABLE</code> synonym. The name and location of this script depends on your operating system. On UNIX and Linux, the script is located in the <code>\$ORACLE HOME/rdbms/admin</code> directory. For



example, start a SQL\*Plus session, connect with SYSDBA privileges, and run the script as follows:

@\$ORACLE\_HOME/rdbms/admin/catplan.sql

The definition of a sample output table PLAN\_TABLE is available in a SQL script on your distribution media. Your output table must have the same column names and data types as this table. The common name of this script is utlxplan.sql. The exact name and location depend on your operating system.

### See Also

Oracle Database SQL Language Reference for a complete description of EXPLAIN PLAN syntax.

### **EXPLAIN PLAN Restrictions**

Oracle Database does not support EXPLAIN PLAN for statements performing implicit type conversion of date bind variables.

With bind variables in general, the EXPLAIN PLAN output might not represent the real execution plan.

From the text of a SQL statement, TKPROF cannot determine the types of the bind variables. It assumes that the type is VARCHAR, and gives an error message otherwise. You can avoid this limitation by putting appropriate type conversions in the SQL statement.

## See Also:

- "Performing Application Tracing"
- "Guideline for Avoiding the Argument Trap"
- Oracle Database SQL Language Reference to learn more about SQL data types

# Explaining a SQL Statement: Basic Steps

Use EXPLAIN PLAN to store the plan for a SQL statement in PLAN TABLE.

#### **Prerequisites**

This task assumes that a sample output table named PLAN\_TABLE exists in your schema. If this table does not exist, then run the SQL script catplan.sql.

To execute EXPLAIN PLAN, you must have the following privileges:

- You must have the privileges necessary to insert rows into an existing output table that you specify to hold the execution plan
- You must also have the privileges necessary to execute the SQL statement for which you
  are determining the execution plan. If the SQL statement accesses a view, then you must



have privileges to access any tables and views on which the view is based. If the view is based on another view that is based on a table, then you must have privileges to access both the other view and its underlying table.

To examine the execution plan produced by an EXPLAIN PLAN statement, you must have the privileges necessary to query the output table.

#### To explain a statement:

- Start SQL\*Plus or SQL Developer, and log in to the database as a user with the requisite permissions.
- 2. Include the EXPLAIN PLAN FOR clause immediately before the SQL statement.

The following example explains the plan for a query of the employees table:

```
EXPLAIN PLAN FOR
  SELECT e.last_name, d.department_name, e.salary
  FROM employees e, departments d
  WHERE salary < 3000
  AND     e.department_id = d.department_id
  ORDER BY salary DESC;</pre>
```

3. After issuing the EXPLAIN PLAN statement, use a script or package provided by Oracle Database to display the most recent plan table output.

The following example uses the DBMS XPLAN.DISPLAY function:

```
SELECT * FROM TABLE(DBMS XPLAN.DISPLAY(format => 'ALL'));
```

4. Review the plan output.

For example, the following plan shows a hash join:

```
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY(format => 'ALL'));
Plan hash value: 3556827125
```

Id	Operation	N	Name	Ro	ws		Bytes	Cc	st	(%CPU)	Time	
0     1    * 2    * 3     4	TABLE ACCESS	FULL  E		     		     	124 124 60	 	5 4	(20)   (0)   (0)	00:00:01 00:00:01 00:00:01 00:00:01 00:00:01	

Query Block Name  $\ /\$  Object Alias (identified by operation id):

```
1 - SEL$1
3 - SEL$1 / E@SEL$1
4 - SEL$1 / D@SEL$1

Predicate Information (identified by operation id):
```

2 - access("E"."DEPARTMENT ID"="D"."DEPARTMENT ID")



Plan operations request data from their children. The execution order in EXPLAIN PLAN output is as follows:

- 1. Execution starts at the first operation with no children, which in the example above is the full scan of EMPLOYEES (Id 3).
- 2. EMPLOYEES returns its data to the parent (Id 2).
- 3. Execution then proceeds to next child of the hash join and does a full scan of DEPARTMENTS (Id 4).
- 4. DEPARTMENTS has no children and so returns data to the parent (Id 2).
- 5. The hash join combines the rows from the two tables and passes them up to the SORT ORDER BY (Id 1)
- 6. Finally the SELECT returns the data to the client.

#### **Note:**

If this example included more operations such as additional joins, execution would continue from Step 5 following the same pattern for each operation down to the end of the plan, the final step where the SELECT returns data to the client.

The steps in the EXPLAIN PLAN output as described here may be different on some of your databases. This is because the optimizer may choose a different EXECUTION PLAN, depending on the database configuration.



### See Also:

- "About PLAN\_TABLE"
- "About the Display of PLAN\_TABLE Output"
- Oracle Database SQL Language Reference for the syntax and semantics of EXPLAIN PLAN
- How to Read an Execution Plan. This Oracle blog post describes how to read an EXECUTION PLAN, but the same order of execution applies to an EXPLAIN PLAN, so it may give you a better understanding of the process in both types of plan.

# Specifying a Statement ID in EXPLAIN PLAN: Example

With multiple statements, you can specify a statement identifier and use that to identify your specific execution plan.

Before using SET STATEMENT ID, remove any existing rows for that statement ID. In the following example, st1 is specified as the statement identifier.

#### Example 6-1 Using EXPLAIN PLAN with the STATEMENT ID Clause

```
EXPLAIN PLAN
   SET STATEMENT_ID = 'st1' FOR
   SELECT last name FROM employees;
```

# Specifying a Different Location for EXPLAIN PLAN Output: Example

The INTO clause of EXPLAIN PLAN specifies a different table in which to store the output.

If you do not want to use the name  $PLAN_TABLE$ , create a new synonym after running the catplan.sql script. For example:

```
CREATE OR REPLACE PUBLIC SYNONYM my plan table for plan table$
```

The following statement directs output to my plan table:

```
EXPLAIN PLAN
  INTO my_plan_table FOR
  SELECT last_name FROM employees;
```

You can specify a statement ID when using the INTO clause, as in the following statement:

```
EXPLAIN PLAN
   SET STATEMENT_ID = 'st1'
   INTO my_plan_table FOR
   SELECT last_name FROM employees;
```



### See Also:

- "PLAN\_TABLE Columns" for a description of the columns in PLAN\_TABLE
- Oracle Database SQL Language Reference to learn about CREATE SYNONYM

# EXPLAIN PLAN Output for a CONTAINERS Query: Example

The CONTAINERS clause can be used to query both user-created and Oracle-supplied tables and views. It enables you to query these tables and views across all containers.

The following example illustrates the output of an EXPLAIN PLAN for a query using the CONTAINERS clause.

SQL> explain plan for select con\_id, count(\*) from containers(sys.dba\_tables) where con id < 10 group by con id order by con id;

Explained.

SQL> @?/rdbms/admin/utlxpls

PLAN TABLE OUTPUT

\_\_\_\_\_\_

-----

Plan hash value: 891225627

| Name | Rows | Bytes | Cost | Id | Operation (%CPU) | Time | Pstart | Pstop | | 0 | SELECT STATEMENT | 234K| 2970K| 145 (100) | 00:00:01 | | 1 | PX COORDINATOR (100) | 00:00:01 | | | 3 | SORT GROUP BY | 234K| 2970K| (100) | 00:00:01 | | 4 | PX RECEIVE | 234K| 2970K| 145 5 | PX SEND RANGE 0) | 00:00:01 (100) | 00:00:01 | |:TQ10000 | 234K| 2970K| 145 (100) | 00:00:01 | | 234K| 2970K| 6 | HASH GROUP BY 145 (100) | 00:00:01 | 7 | PX PARTITION LIST ITERATOR| | 234K| 2970K| 139 (100) | 00:00:01 | 1 | 9 | 8 | CONTAINERS FULL | DBA TABLES | 234K| 2970K| 139 (100) | 00:00:01 | |

\_\_\_\_\_



```
15 rows selected.
```

At Row 8 of this plan, CONTAINERS is shown in the Operation column as the value CONTAINERS FULL. The Name column in the same row shows the argument to CONTAINERS.

#### **Default Partitioning**

A query using the CONTAINERS clause is partitioned by default. At Row 7 in the plan, the PX PARTITION LIST ITERATOR in the Operation column indicates that the query is partitioned. Iteration over containers is implemented in this partition iterator. On the same row, the Pstart and Pstop values 1 and 9 are derived from the con id < 10 predicate in the query.

#### **Default Parallelism**

A query using the CONTAINERS clause uses parallel execution servers by default. In Row 1 of the plan above, PX COORDINATOR in the Operation column indicates that parallel execution servers will be used. Each container is assigned to a parallel execution process (P00\*). When the parallel execution process executes the part of the query EXECUTION PLAN that corresponds to CONTAINERS FULL, then the process switches into the container it has been assigned to work on. It retrieves rows from the base object by executing a recursive SQL statement.

# **Displaying Execution Plans**

The easiest way to display execution plans is to use DBMS XPLAN display functions or V\$ views.

# About the Display of PLAN\_TABLE Output

To display the plan table output, you can use either SQL scripts or the DBMS XPLAN package.

After you have explained the plan, use the following SQL scripts or PL/SQL package provided by Oracle Database to display the most recent plan table output:

DBMS XPLAN.DISPLAY table function

This function accepts options for displaying the plan table output. You can specify:

- A plan table name if you are using a table different than PLAN TABLE
- A statement ID if you have set a statement ID with the EXPLAIN PLAN
- A format option that determines the level of detail: BASIC, SERIAL, TYPICAL, and ALL

Examples of using DBMS XPLAN to display PLAN TABLE output are:

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE(DBMS_XPLAN.DISPLAY());

SELECT PLAN_TABLE_OUTPUT
  FROM TABLE(DBMS_XPLAN.DISPLAY('MY_PLAN_TABLE', 'st1','TYPICAL'));
```

• utlxpls.sql

This script displays the plan table output for serial processing

• utlxplp.sql

This script displays the plan table output including parallel execution columns.

See Also:

Oracle Database PL/SQL Packages and Types Reference for more information about the  $\tt DBMS$  XPLAN package

# DBMS\_XPLAN Display Functions

You can use the DBMS\_XPLAN display functions to show plans.

The display functions accept options for displaying the plan table output. You can specify:

- A plan table name if you are using a table different from PLAN TABLE
- A statement ID if you have set a statement ID with the EXPLAIN PLAN
- A format option that determines the level of detail: BASIC, SERIAL, TYPICAL, ALL, and in some cases ADAPTIVE

Table 6-1 DBMS\_XPLAN Display Functions

Display Functions	Notes					
DISPLAY	This table function displays the contents of the plan table.					
	In addition, you can use this table function to display any plan (with or without statistics) stored in a table as long as the columns of this table are named the same as columns of the plan table (or V\$SQL_PLAN_STATISTICS_ALL if statistics are included). You can apply a predicate on the specified table to select rows of the plan to display.					
	The format parameter controls the level of the plan. It accepts the values BASIC, TYPICAL, SERIAL, and ALL.					
DISPLAY_AWR	This table function displays the contents of an execution plan stored in AWR.					
	The format parameter controls the level of the plan. It accepts the values BASIC, TYPICAL, SERIAL, and ALL.					
DISPLAY_CURSOR	This table function displays the explain plan of any cursor loaded in the cursor cache. In addition to the explain plan, various plan statistics (such as. I/O, memory and timing) can be reported (based on the V\$SQL_PLAN_STATISTICS_ALL_VIEWS).					
	The format parameter controls the level of the plan. It accepts the values BASIC, TYPICAL, SERIAL, ALL, and ADAPTIVE. When you specify ADAPTIVE, the output includes:					
	<ul> <li>The final plan. If the execution has not completed, then the output shows the current plan. This section also includes notes about run-time optimizations that affect the plan.</li> </ul>					
	<ul> <li>Recommended plan. In reporting mode, the output includes the plan that would be chosen based on execution statistics.</li> </ul>					
	<ul> <li>Dynamic plan. The output summarizes the portions of the plan that differ from the default plan chosen by the optimizer.</li> </ul>					
	<ul> <li>Reoptimization. The output displays the plan that would be chosen on a subsequent execution because of reoptimization.</li> </ul>					



Table 6-1 (Cont.) DBMS\_XPLAN Display Functions

Display Functions	Notes
DISPLAY_PLAN	This table function displays the contents of the plan table in a variety of formats with CLOB output type.  The format parameter controls the level of the plan. It accepts the values BASIC, TYPICAL, SERIAL, ALL, and ADAPTIVE. When you specify ADAPTIVE, the output includes the default plan. For each dynamic subplan, the plan shows a list of the row sources from the original that may be replaced, and
	the row sources that would replace them.  If the format argument specifies the outline display, then the function displays the hints for each option in the dynamic subplan. If the plan is not an adaptive query plan, then the function displays the default plan. When you do not specify ADAPTIVE, the plan is shown as-is, but with additional comments in the Note section that show any row sources that are dynamic.
DISPLAY_SQL_PLAN_BA SELINE	This table function displays one or more execution plans for the specified SQL handle of a SQL plan baseline.
	This function uses plan information stored in the plan baseline to explain and display the plans. The $plan\_id$ stored in the SQL management base may not match the $plan\_id$ of the generated plan. A mismatch between the stored $plan\_id$ and generated $plan\_id$ means that it is a non-reproducible plan. Such a plan is deemed invalid and is bypassed by the optimizer during SQL compilation.
DISPLAY_SQLSET	This table function displays the execution plan of a given statement stored in a SQL tuning set.
	The format parameter controls the level of the plan. It accepts the values BASIC, TYPICAL, SERIAL, and ALL.

## See Also:

Oracle Database PL/SQL Packages and Types Reference to learn more about  $\tt DBMS \ XPLAN \ display \ functions$ 



# Plan-Related Views

You can obtain information about execution plans by querying dynamic performance and data dictionary views.

Table 6-2 Execution Plan Views

View	Description
V\$SQL	Lists statistics for cursors and contains one row for each child of the original SQL text entered.
	Starting in Oracle Database 19c, V\$SQL.QUARANTINED indicates whether a statement has been terminated by the Resource Manager because the statement consumed too many resources. Oracle Database records and marks the quarantined plans and prevents the execution of statements using these plans from executing. The AVOIDED_EXECUTIONS column indicates the number of executions attempted but prevented because of the quarantined statement.
V\$SQL_SHARED_CURSOR	Explains why a particular child cursor is not shared with existing child cursors. Each column identifies a specific reason why the cursor cannot be shared.
	The USE_FEEDBACK_STATS column shows whether a child cursor fails to match because of reoptimization.
V\$SQL_PLAN	Contains the plan for every statement stored in the shared SQL area.
	The view definition is similar to PLAN_TABLE. The view includes a superset of all rows appearing in all final plans. PLAN_LINE_ID is consecutively numbered, but for a single final plan, the IDs may not be consecutive.
	As an alternative to EXPLAIN PLAN, you can display the plan by querying V\$SQL_PLAN. The advantage of V\$SQL_PLAN over EXPLAIN PLAN is that you do not need to know the compilation environment that was used to execute a particular statement. For EXPLAIN PLAN, you would need to set up an identical environment to get the same plan when executing the statement.
V\$SQL_PLAN_STATISTICS	Provides the actual execution statistics for every operation in the plan, such as the number of output rows and elapsed time. All statistics, except the number of output rows, are cumulative. For example, the statistics for a join operation also includes the statistics for its two inputs. The statistics in V\$SQL_PLAN_STATISTICS are available for cursors that have been compiled with the STATISTICS_LEVEL initialization parameter set to ALL.
V\$SQL_PLAN_STATISTICS_ALL	Contains memory usage statistics for row sources that use SQL memory (sort or hash join). This view concatenates information in V\$SQL_PLAN with execution statistics from V\$SQL_PLAN_STATISTICS and V\$SQL_WORKAREA.
	V\$SQL_PLAN_STATISTICS_ALL enables side-by-side comparisons of the estimates that the optimizer provides for the number of rows and elapsed time. This view combines both V\$SQL_PLAN and V\$SQL_PLAN_STATISTICS information for every cursor.



### See Also:

- "PLAN\_TABLE Columns"
- "Monitoring Database Operations " for information about the V\$SQL PLAN MONITOR view
- Oracle Database Reference for more information about V\$SQL PLAN views
- Oracle Database Reference for information about the STATISTICS\_LEVEL initialization parameter

# Displaying Execution Plans: Basic Steps

The DBMS XPLAN.DISPLAY function is a simple way to display an explained plan.

By default, the DISPLAY function uses the format setting of TYPICAL. In this case, the plan the most relevant information in the plan: operation id, name and option, rows, bytes and optimizer cost. Pruning, parallel and predicate information are only displayed when applicable.

#### To display an execution plan:

- 1. Start SQL\*Plus or SQL Developer and log in to the session in which you explained the plan.
- 2. Explain a plan.
- 3. Query PLAN TABLE using DBMS XPLAN.DISPLAY.

Specify the query as follows:

```
SELECT PLAN TABLE OUTPUT FROM TABLE (DBMS XPLAN.DISPLAY);
```

Alternatively, specify the statement ID using the statement id parameter:

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE(DBMS_XPLAN.DISPLAY(statement_id =>
'statement_id));
```

#### Example 6-2 EXPLAIN PLAN for Statement ID ex\_plan1

This example explains a query of employees that uses the statement ID  $ex_plan1$ , and then queries PLAN TABLE:

```
EXPLAIN PLAN
  SET statement_id = 'ex_plan1' FOR
  SELECT phone_number
  FROM employees
  WHERE phone_number LIKE '650%';

SELECT PLAN_TABLE_OUTPUT
  FROM TABLE(DBMS XPLAN.DISPLAY(statement id => 'ex plan1'));
```



#### Sample output appears below:

Plan hash value: 1445457117

Id   Operation		Name	Ro	WS	١	Bytes		Cost	(%CPU)	Time	
0  SELECT STATEMENT  * 1  TABLE ACCESS FULI											
Predicate Information (identified by operation id):											
1 - filter("PHONE NUMBER" LIKE '650%')											

### Example 6-3 EXPLAIN PLAN for Statement ID ex\_plan2

This example explains a query of employees that uses the statement ID  $ex_plan2$ , and then displays the plan using the BASIC format:

```
EXPLAIN PLAN
  SET statement_id = 'ex_plan2' FOR
  SELECT last_name
  FROM employees
  WHERE last_name LIKE 'Pe%';

SELECT PLAN_TABLE_OUTPUT
  FROM TABLE(DBMS XPLAN.DISPLAY(NULL, 'ex plan2', 'BASIC'));
```

#### Sample output appears below:

```
| Id | Operation | Name |
| 0 | SELECT STATEMENT | |
| 1 | INDEX RANGE SCAN | EMP_NAME_IX |
```

#### See Also:

Oracle Database PL/SQL Packages and Types Reference for more information about the  $\tt DBMS\_XPLAN$  package

# Displaying Adaptive Query Plans: Tutorial

The **adaptive optimizer** is a feature of the optimizer that enables it to adapt plans based on run-time statistics. All adaptive mechanisms can execute a final plan for a statement that differs from the default plan.

An adaptive query plan chooses among subplans *during* the current statement execution. In contrast, automatic reoptimization changes a plan only on executions that occur *after* the current statement execution.

You can determine whether the database used adaptive query optimization for a SQL statement based on the comments in the Notes section of plan. The comments indicate whether row sources are dynamic, or whether automatic reoptimization adapted a plan.

#### **Assumptions**

This tutorial assumes the following:

- The STATISTICS LEVEL initialization parameter is set to ALL.
- The database uses the default settings for adaptive execution.
- As user oe, you want to issue the following separate queries:

- Before executing each query, you want to query DBMS\_XPLAN.DISPLAY\_PLAN to see the
  default plan, that is, the plan that the optimizer chose before applying its adaptive
  mechanism.
- After executing each query, you want to query DBMS\_XPLAN.DISPLAY\_CURSOR to see the final
  plan and adaptive query plan.
- SYS has granted oe the following privileges:

```
- GRANT SELECT ON V_$SESSION TO OE
- GRANT SELECT ON V_$SQL TO OE
- GRANT SELECT ON V_$SQL_PLAN TO OE
- GRANT SELECT ON V $SQL PLAN STATISTICS ALL TO OE
```

#### To see the results of adaptive optimization:

1. Start SQL\*Plus, and then connect to the database as user oe.

2. Query orders.

For example, use the following statement:

3. View the plan in the cursor.

For example, run the following commands:

```
SET LINESIZE 165
SET PAGESIZE 0
SELECT * FROM TABLE(DBMS XPLAN.DISPLAY CURSOR(FORMAT=>'+ALLSTATS'));
```

The following sample output has been reformatted to fit on the page. In this plan, the optimizer chooses a nested loops join. The original optimizer estimates are shown in the E-Rows column, whereas the actual statistics gathered during execution are shown in the A-Rows column. In the MERGE JOIN operation, the difference between the estimated and actual number of rows is significant.

- 4. Run the same guery of orders that you ran in Step 2.
- 5. View the execution plan in the cursor by using the same SELECT statement that you ran in Step 3.

The following example shows that the optimizer has chosen a different plan, using a hash join. The Note section shows that the optimizer used statistics feedback to adjust its cost estimates for the second execution of the query, thus illustrating automatic reoptimization.



```
|*5| INDEX UNIQUE SCAN |ORDER_PK |269| 1|269|00:00:00.01|21|0| | | |

Predicate Information (identified by operation id):

2 - access("P"."PRODUCT_ID"="O"."PRODUCT_ID")

3 - filter(("MIN_PRICE"<40 AND "LIST_PRICE"<50))

5 - access("O"."ORDER_ID"="ORDER_ID")

Note
----
```

- statistics feedback used for this statement

6. Query V\$SQL to verify the performance improvement.

The following query shows the performance of the two statements (sample output included).

```
SELECT CHILD_NUMBER, CPU_TIME, ELAPSED_TIME, BUFFER_GETS
FROM V$SQL
WHERE SQL_ID = 'gm2npz344xqn8';

CHILD_NUMBER CPU_TIME ELAPSED_TIME BUFFER_GETS

0 92006 131485 1831
1 12000 24156 60
```

The second statement executed, which is child number 1, used statistics feedback. CPU time, elapsed time, and buffer gets are all significantly lower.

7. Explain the plan for the query of order\_items.

For example, use the following statement:

```
EXPLAIN PLAN FOR
  SELECT product_name
  FROM    order_items o, product_information p
  WHERE    o.unit_price = 15
  AND     quantity > 1
  AND    p.product_id = o.product_id
```

8. View the plan in the plan table.

For example, run the following statement:

```
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

Sample output appears below:

Id  Operation	Name	Rows Bytes Cost (%CPU) Time
0  SELECT STATEMENT		4 128 7 (0) 00:00:01
1  NESTED LOOPS		
2  NESTED LOOPS		4 128 7 (0) 00:00:01
*3  TABLE ACCESS FULL	ORDER_ITEMS	4 48  3 (0) 00:00:01
*4  INDEX UNIQUE SCAN	PRODUCT_INFORM	MATION_PK 1   0 (0) 00:00:01
5  TABLE ACCESS BY INDEX	ROWID   PRODUCT_INFORM	MATION  1 20  1 (0) 00:00:01



\_\_\_\_\_\_

```
Predicate Information (identified by operation id):

3 - filter("O"."UNIT_PRICE"=15 AND "QUANTITY">1)
4 - access("P"."PRODUCT_ID"="O"."PRODUCT_ID")
```

In this plan, the optimizer chooses a nested loops join.

9. Run the query that you previously explained.

For example, use the following statement:

```
SELECT product_name
FROM order_items o, product_information p
WHERE o.unit_price = 15
AND quantity > 1
AND p.product id = o.product id
```

**10.** View the plan in the cursor.

For example, run the following commands:

```
SET LINESIZE 165
SET PAGESIZE 0
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY(FORMAT=>'+ADAPTIVE'));
```

Sample output appears below. Based on statistics collected at run time (Step 4), the optimizer chose a hash join rather than the nested loops join. The dashes (-) indicate the steps in the nested loops plan that the optimizer considered but do not ultimately choose. The switch illustrates the adaptive guery plan feature.

```
|Id | Operation
                     | Name
                           |Rows|Bytes|Cost(%CPU)|Time
______
| 0| SELECT STATEMENT
                                   |4|128|7(0)|00:00:01|
                   | |
| *1| HASH JOIN
                                   |4|128|7(0)|00:00:01|
|- 2| NESTED LOOPS
                                  |- 3| NESTED LOOPS
                                  | |128|7(0)|00:00:01| | | |
|- 7| TABLE ACCESS BY INDEX ROWID| PRODUCT INFORMATION |1| 20|1(0)|00:00:01|
| 8| TABLE ACCESS FULL | PRODUCT INFORMATION |1| 20|1(0)|00:00:01|
```

```
Predicate Information (identified by operation \operatorname{id}):
```

```
1 - access("P"."PRODUCT_ID"="O"."PRODUCT_ID")
5 - filter("O"."UNIT_PRICE"=15 AND "QUANTITY">1)
6 - access("P"."PRODUCT_ID"="O"."PRODUCT_ID")
```

Note

----

- this is an adaptive plan (rows marked '-' are inactive)

### See Also:

- "Adaptive Query Plans"
- "Table 6-1"
- "Controlling Adaptive Optimization"
- Oracle Database Reference to learn about the STATISTICS\_LEVEL initialization parameter
- Oracle Database PL/SQL Packages and Types Reference to learn more about DBMS\_XPLAN

# Display Execution Plans: Examples

These examples show different ways of displaying execution plans.

## Customizing PLAN\_TABLE Output

If you have specified a statement identifier, then you can write your own script to query the  ${\tt PLAN\_TABLE}.$ 

#### For example:

- Start with ID = 0 and given STATEMENT ID.
- Use the CONNECT BY clause to walk the tree from parent to child, the join keys being STATEMENT ID = PRIOR STATMENT ID and PARENT ID = PRIOR ID.
- Use the pseudo-column LEVEL (associated with CONNECT BY) to indent the children.

The  $\mathtt{NULL}$  in the  $\mathtt{Rows}$  column indicates that the optimizer does not have any statistics on the table. Analyzing the table shows the following:

```
Rows Plan
```



16957 SELECT STATEMENT
16957 TABLE ACCESS FULL EMPLOYEES

You can also select the COST. This is useful for comparing execution plans or for understanding why the optimizer chooses one execution plan over another.



These simplified examples are not valid for recursive SQL.

## Displaying Parallel Execution Plans: Example

Plans for parallel queries differ in important ways from plans for serial queries.

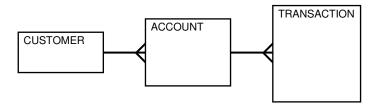
## About EXPLAIN PLAN and Parallel Queries

Tuning a parallel query begins much like a non-parallel query tuning exercise by choosing the driving table. However, the rules governing the choice are different.

In the serial case, the best driving table produces the fewest numbers of rows after applying limiting conditions. The database joins a small number of rows to larger tables using non-unique indexes.

For example, consider a table hierarchy consisting of customer, account, and transaction.

Figure 6-1 A Table Hierarchy



In this example, customer is the smallest table, whereas transaction is the largest table. A typical OLTP query retrieves transaction information about a specific customer account. The query drives from the customer table. The goal is to minimize logical I/O, which typically minimizes other critical resources including physical I/O and CPU time.

For parallel queries, the driving table is usually the *largest* table. It would not be efficient to use parallel query in this case because only a few rows from each table are accessed. However, what if it were necessary to identify all customers who had transactions of a certain type last month? It would be more efficient to drive from the transaction table because no limiting conditions exist on the customer table. The database would join rows from the transaction table to the account table, and then finally join the result set to the customer table. In this case, the used on the account and customer table are probably highly selective primary key or unique indexes rather than the non-unique indexes used in the first query. Because the transaction table is large and the column is not selective, it would be beneficial to use parallel query driving from the transaction table.

Parallel operations include the following:



- PARALLEL TO PARALLEL
- PARALLEL TO SERIAL

A PARALLEL\_TO\_SERIAL operation is always the step that occurs when the query coordinator consumes rows from a parallel operation. Another type of operation that does not occur in this query is a SERIAL operation. If these types of operations occur, then consider making them parallel operations to improve performance because they too are potential bottlenecks.

- PARALLEL FROM SERIAL
- PARALLEL TO PARALLEL

If the workloads in each step are relatively equivalent, then the PARALLEL\_TO\_PARALLEL operations generally produce the best performance.

- PARALLEL\_COMBINED\_WITH\_CHILD
- PARALLEL COMBINED WITH PARENT

A PARALLEL\_COMBINED\_WITH\_PARENT operation occurs when the database performs the step simultaneously with the parent step.

If a parallel step produces many rows, then the QC may not be able to consume the rows as fast as they are produced. Little can be done to improve this situation.

### See Also:

The OTHER TAG column in "PLAN\_TABLE Columns"

## Viewing Parallel Queries with EXPLAIN PLAN: Example

When using EXPLAIN PLAN with parallel queries, the database compiles and executes one parallel plan. This plan is derived from the serial plan by allocating row sources specific to the parallel support in the QC plan.

The table queue row sources (PX Send and PX Receive), the granule iterator, and buffer sorts, required by the two parallel execution server set PQ model, are directly inserted into the parallel plan. This plan is the same plan for all parallel execution servers when executed in parallel or for the QC when executed serially.

#### Example 6-4 Parallel Query Explain Plan

The following simple example illustrates an EXPLAIN PLAN for a parallel query:

```
CREATE TABLE emp2 AS SELECT * FROM employees;

ALTER TABLE emp2 PARALLEL 2;

EXPLAIN PLAN FOR
SELECT SUM(salary)
FROM emp2
GROUP BY department_id;

SELECT PLAN_TABLE_OUTPUT FROM TABLE(DBMS_XPLAN.DISPLAY());
```



Id	Operation	Name	Rows	Bytes	10	Cost	%CPU	TQ	:	IN-OUT	PQ	Distrib	)
0	SELECT STATEMENT		107	2782	3	(34)							
1	PX COORDINATOR PX SEND QC (RANDOM)	  :TQ10001	  107	2782 I	3	(34)	1 0	1 . 01		P->S		(RAND)	1
3	HASH GROUP BY			2782						PCWP		(IdlivD)	i
4	PX RECEIVE		107	2782	3	(34)	I Q	1,01		PCWP			
5	PX SEND HASH	:TQ10000								P->P	HAS	H	
6	HASH GROUP BY			2782		, ,	ΙQ	1,00		PCWP			
7	PX BLOCK ITERATOR			2782		, ,		1,00		PCWP			
8	TABLE ACCESS FULL	EMP2 	107  	2782	2 	(0)	Q	1,00 		PCWP	 		

One set of parallel execution servers scans EMP2 in parallel, while the second set performs the aggregation for the GROUP BY operation. The PX BLOCK ITERATOR row source represents the splitting up of the table EMP2 into pieces to divide the scan workload between the parallel execution servers. The PX SEND and PX RECEIVE row sources represent the pipe that connects the two sets of parallel execution servers as rows flow up from the parallel scan, get repartitioned through the HASH table queue, and then read by and aggregated on the top set. The PX SEND QC row source represents the aggregated values being sent to the QC in random (RAND) order. The PX COORDINATOR row source represents the QC or Query Coordinator which controls and schedules the parallel plan appearing below it in the plan tree.

## Displaying Bitmap Index Plans: Example

Index row sources using bitmap indexes appear in the EXPLAIN PLAN output with the word BITMAP indicating the type of the index.

#### Example 6-5 EXPLAIN PLAN with Bitmap Indexes

In this example, the predicate c1=2 yields a bitmap from which a subtraction can take place. From this bitmap, the bits in the bitmap for c2=6 are subtracted. Also, the bits in the bitmap for c2 IS NULL are subtracted, explaining why there are two MINUS row sources in the plan. The NULL subtraction is necessary for semantic correctness unless the column has a NOT NULL constraint. The TO ROWIDS option generates the rowids necessary for the table access.



Queries using bitmap join index indicate the bitmap join index access path. The operation for bitmap join index is the same as bitmap index.

```
EXPLAIN PLAN FOR SELECT *

FROM t

WHERE c1 = 2

AND c2 <> 6

OR c3 BETWEEN 10 AND 20;

SELECT STATEMENT

TABLE ACCESS T BY INDEX ROWID

BITMAP CONVERSION TO ROWID

BITMAP OR

BITMAP MINUS
```

BITMAP MINUS

BITMAP INDEX C1\_IND SINGLE VALUE
BITMAP INDEX C2\_IND SINGLE VALUE
BITMAP INDEX C2\_IND SINGLE VALUE
BITMAP MERGE
BITMAP INDEX C3 IND RANGE SCAN

## Displaying Result Cache Plans: Example

When your query contains the result\_cache hint, the ResultCache operator is inserted into the execution plan.

Starting in Oracle Database 21c, the <code>result\_cache</code> hint accepts a new option:  $result_cache (TEMP=\{TRUE \mid FALSE\})$ . A value of <code>TRUE</code> enables the query to spill to disk, whereas <code>FALSE</code> prevents a Temp object from being formed. Instead, the Result object will enter the 'Bypass' status.

For example, you might explain a guery as follows:

```
EXPLAIN PLAN FOR
SELECT /*+ result_cache(TEMP=TRUE) */ department_id, AVG(salary)
FROM employees
GROUP BY department_id;
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY(format => 'ALL'));
```

The EXPLAIN PLAN output for this query includes a Result Cache Information section, and should look similar to the following:

```
PLAN TABLE OUTPUT
                -----
Plan hash value: 1192169904
                 | Name
| Id | Operation
                                                | Rows |Bytes|Cost (%CPU)| Time|
  0 | SELECT STATEMENT |
                                               | 11 | 77 | 4 (25) | 00:00:01 |
| 1 | RESULT CACHE | ch5r45jxt05rk0xc1brct197fp | 11 | 77 | 4 (25) | 00:00:01 | 2 | HASH GROUP BY | 11 | 77 | 4 (25) | 00:00:01 |
       TABLE ACCESS FULL | EMPLOYEES
                                                | 107 | 749 | 3 (0) | 00:00:01 |
Query Block Name / Object Alias (identified by operation id):
______
  1 - SEL$1
  3 - SEL$1 / "EMPLOYEES"@"SEL$1"
Column Projection Information (identified by operation id):
______
  1 - "DEPARTMENT ID" [NUMBER, 22], SUM("SALARY")/COUNT("SALARY")[22]
  2 - (#keys=1) "DEPARTMENT_ID"[NUMBER,22], COUNT("SALARY")[22], SUM("SALARY")[22]
  3 - (rowset=256) "SALARY"[NUMBER, 22], "DEPARTMENT ID"[NUMBER, 22]
Result Cache Information (identified by operation id):
```



1 - column-count=2; dependencies=(HR.EMPLOYEES);

```
name="SELECT /*+ result_cache(TEMP=TRUE) */ department_id, AVG(salary)
FROM employees
GROUP BY department id"
```

In this plan, the RESULT CACHE operations is identified by its cache ID, which is ch5r45jxt05rk0xc1brct197fp. You can query the V\$RESULT\_CACHE\_OBJECTS view by using this CACHE ID, as shown in the following example (sample output included):

```
SELECT SUBCACHE_ID, TYPE, STATUS, BLOCK_COUNT,
ROW_COUNT, INVALIDATIONS

FROM V$RESULT_CACHE_OBJECTS
WHERE CACHE_ID = 'ch5r45jxt05rk0xc1brct197fp';

SUBCACHE_ID TYPE STATUS BLOCK_COUNT ROW_COUNT INVALIDATIONS

0 Result Published 1 12 0
```

## Displaying Plans for Partitioned Objects: Example

Use  $\tt EXPLAIN\ PLAN\ to\ determine\ how\ Oracle\ Database\ accesses\ partitioned\ objects\ for\ specific\ queries.$ 

Partitions accessed after pruning are shown in the PARTITION START and PARTITION STOP columns. The row source name for the range partition is PARTITION RANGE. For hash partitions, the row source name is PARTITION HASH.

A join is implemented using partial partition-wise join if the DISTRIBUTION column of the plan table of one of the joined tables contains PARTITION (KEY). Partial partition-wise join is possible if one of the joined tables is partitioned on its join column and the table is parallelized.

A join is implemented using full partition-wise join if the partition row source appears before the join row source in the EXPLAIN PLAN output. Full partition-wise joins are possible only if both joined tables are equipartitioned on their respective join columns. Examples of execution plans for several types of partitioning follow.

## Displaying Range and Hash Partitioning with EXPLAIN PLAN: Examples

This example illustrates pruning by using the <code>emp\_range</code> table, which partitioned by range on hire date.

Assume that the tables employees and departments from the Oracle Database sample schema exist.

```
CREATE TABLE emp_range

PARTITION BY RANGE(hire_date)
(

PARTITION emp_p1 VALUES LESS THAN (TO_DATE('1-JAN-1992','DD-MON-YYYY')),

PARTITION emp_p2 VALUES LESS THAN (TO_DATE('1-JAN-1994','DD-MON-YYYY')),

PARTITION emp_p3 VALUES LESS THAN (TO_DATE('1-JAN-1996','DD-MON-YYYY')),

PARTITION emp_p4 VALUES LESS THAN (TO_DATE('1-JAN-1998','DD-MON-YYYY')),

PARTITION emp_p5 VALUES LESS THAN (TO_DATE('1-JAN-2001','DD-MON-YYYY'))

AS SELECT * FROM employees;
```

#### For the first example, consider the following statement:

```
EXPLAIN PLAN FOR SELECT * FROM emp range;
```

#### Oracle Database displays something similar to the following:

Id  Operation	Name	Rows  Bytes Cost Pstart Pstop
0  SELECT STATEMENT	1	105  13965   2
1  PARTITION RANGE AL	L	105  13965   2   1   5
2  TABLE ACCESS FULL	EMP_RANGE	GE   105  13965   2   1   5

The database creates a partition row source on top of the table access row source. It iterates over the set of partitions to be accessed. In this example, the partition iterator covers all partitions (option ALL), because a predicate was not used for pruning. The PARTITION\_START and PARTITION STOP columns of the PLAN TABLE show access to all partitions from 1 to 5.

For the next example, consider the following statement:

In the previous example, the partition row source iterates from partition 4 to 5 because the database prunes the other partitions using a predicate on hire\_date.

Finally, consider the following statement:



In the previous example, only partition 1 is accessed and known at compile time; thus, there is no need for a partition row source.



EXPLAIN PLAN FOR

Oracle Database displays the same information for hash partitioned objects, except the partition row source name is PARTITION HASH instead of PARTITION RANGE. Also, with hash partitioning, pruning is only possible using equality or IN-list predicates.

## Pruning Information with Composite Partitioned Objects: Examples

To illustrate how Oracle Database displays pruning information for composite partitioned objects, consider the table <code>emp\_comp</code>. It is range-partitioned on <code>hiredate</code> and subpartitioned by hash on <code>deptno</code>.

```
CREATE TABLE emp_comp PARTITION BY RANGE(hire_date)

SUBPARTITION BY HASH(department_id) SUBPARTITIONS 3

(

PARTITION emp_p1 VALUES LESS THAN (TO_DATE('1-JAN-1992','DD-MON-YYYY')),

PARTITION emp_p2 VALUES LESS THAN (TO_DATE('1-JAN-1994','DD-MON-YYYY')),

PARTITION emp_p3 VALUES LESS THAN (TO_DATE('1-JAN-1996','DD-MON-YYYY')),

PARTITION emp_p4 VALUES LESS THAN (TO_DATE('1-JAN-1998','DD-MON-YYYY')),

PARTITION emp_p5 VALUES LESS THAN (TO_DATE('1-JAN-2001','DD-MON-YYYY'))

AS SELECT * FROM employees;
```

For the first example, consider the following statement:

This example shows the plan when Oracle Database accesses all subpartitions of all partitions of a composite object. The database uses two partition row sources for this purpose: a range partition row source to iterate over the partitions, and a hash partition row source to iterate over the subpartitions of each accessed partition.

In the following example, the range partition row source iterates from partition 1 to 5, because the database performs no pruning. Within each partition, the hash partition row source iterates over subpartitions 1 to 3 of the current partition. As a result, the table access row source

accesses subpartitions 1 to 15. In other words, the database accesses all subpartitions of the composite object.

In the previous example, only the last partition, partition 5, is accessed. This partition is known at compile time, so the database does not need to show it in the plan. The hash partition row source shows accessing of all subpartitions within that partition; that is, subpartitions 1 to 3, which translates into subpartitions 13 to 15 of the emp comp table.

Now consider the following statement:

In the previous example, the predicate <code>deptno=20</code> enables pruning on the hash dimension within each partition. Therefore, Oracle Database only needs to access a single subpartition. The number of this subpartition is known at compile time, so the hash partition row source is not needed.

Finally, consider the following statement:

```
VARIABLE dno NUMBER;

EXPLAIN PLAN FOR

SELECT *

FROM emp_comp
WHERE department_id = :dno;

Id| Operation | Name | Rows| Bytes | Cost| Pstart| Pstop|

0 | SELECT STATEMENT | | 101| 13433 | 78 | | |
```



```
| 1 | PARTITION RANGE ALL | | 101 | 13433 | 78 | 1 | 5 | | 2 | PARTITION HASH SINGLE | | 101 | 13433 | 78 | KEY | KEY | | *3 | TABLE ACCESS FULL | EMP_COMP | 101 | 13433 | 78 | | |
```

The last two examples are the same, except that department\_id = :dno replaces deptno=20. In this last case, the subpartition number is unknown at compile time, and a hash partition row source is allocated. The option is SINGLE for this row source because Oracle Database accesses only one subpartition within each partition. In Step 2, both PARTITION\_START and PARTITION\_STOP are set to KEY. This value means that Oracle Database determines the number of subpartitions at run time.

### **Examples of Partial Partition-Wise Joins**

In these examples, the PQ\_DISTRIBUTE hint explicitly forces a partial partition-wise join because the query optimizer could have chosen a different plan based on cost in this query.

#### Example 6-6 Partial Partition-Wise Join with Range Partition

In the following example, the database joins <code>emp\_range\_did</code> on the partitioning column <code>department\_id</code> and parallelizes it. The database can use a partial partition-wise join because the <code>dept2</code> table is not partitioned. Oracle Database dynamically partitions the <code>dept2</code> table before the join.

```
CREATE TABLE dept2 AS SELECT * FROM departments;
ALTER TABLE dept2 PARALLEL 2;
CREATE TABLE emp range did PARTITION BY RANGE (department id)
   (PARTITION emp p1 VALUES LESS THAN (150),
   PARTITION emp p5 VALUES LESS THAN (MAXVALUE) )
 AS SELECT * FROM employees;
ALTER TABLE emp range did PARALLEL 2;
EXPLAIN PLAN FOR
 SELECT /*+ PQ DISTRIBUTE(d NONE PARTITION) ORDERED */ e.last name,
        d.department name
 FROM emp range did e, dept2 d
 WHERE e.department id = d.department id;
                              |Name |Row|Byte|Cost|Pstart|Pstop|TQ|IN-OUT|PQ Distrib|
|Id| Operation

    HASH JOIN
    |
    | 284 | 16188 | 6 | | Q1,01 | PCWP |

    PX PARTITION RANGE ALL
    | 284 | 7668 | 2 | 1 | 2 | Q1,01 | PCWC |

|*3| HASH JOIN
| 4|
TABLE ACCESS FULL | EMP_RANGE_DID|284 |7668 |2|1 |2| Q1,01 |PCWP|
        PX SEND PARTITION (KEY)|:TQ10000 | 21 | 630 |2| | | | |S->P|PART (KEY)|

TABLE ACCESS FULL |DEPT2 | 21 | 630 |2| | | | |
```

The execution plan shows that the table dept2 is scanned serially and all rows with the same partitioning column value of  $emp\_range\_did$  ( $department\_id$ ) are sent through a PART (KEY), or partition key, table queue to the same parallel execution server doing the partial partitionwise join.

#### **Example 6-7 Partial Partition-Wise Join with Composite Partition**

In the following example, <code>emp\_comp</code> is joined on the partitioning column and is parallelized, enabling use of a partial partition-wise join because <code>dept2</code> is not partitioned. The database dynamically partitions <code>dept2</code> before the join.

The plan shows that the optimizer selects partial partition-wise join from one of two columns. The PX SEND node type is PARTITION (KEY) and the PQ Distrib column contains the text PART (KEY), or partition key. This implies that the table dept2 is re-partitioned based on the join column department\_id to be sent to the parallel execution servers executing the scan of EMP COMP and the join.

## Example of Full Partition-Wise Join

In this example, <code>emp\_comp</code> and <code>dept\_hash</code> are joined on their hash partitioning columns, enabling use of a full partition-wise join.

The Partition hash row source appears on top of the join row source in the plan table output.

```
CREATE TABLE dept_hash

PARTITION BY HASH(department_id)

PARTITIONS 3

PARALLEL 2

AS SELECT * FROM departments;
```



Id	Operation	Name  Ro	VS.	Byte	es	Cost	Psta	rt F	st	op TQ	II	N-OUT	PQ I	)istrib
0	SELECT STATEMENT			106		2544	8							
1	PX COORDINATOR													1
2	PX SEND QC (RANDOM)	:TQ10000		106		2544	8			Q1,00		P->S	QC	(RAND)
3	PX PARTITION HASH ALL			106		2544	8 1	3		Q1,00		PCWC		1
* 4	HASH JOIN			106		2544	8			Q1,00		PCWP		1
5	PX PARTITION RANGE AL:	L		107		1070	3 1	5	-	Q1,00		PCWC		1
6	TABLE ACCESS FULL	EMP COMP		107		1070	3 1	15	-	Q1,00		PCWP		
7	TABLE ACCESS FULL	DEPT_HASH		27		378	4 1	3		Q1,00		PCWP		1

The PX PARTITION HASH row source appears on top of the join row source in the plan table output while the PX PARTITION RANGE row source appears over the scan of emp\_comp. Each parallel execution server performs the join of an entire hash partition of emp\_comp with an entire partition of dept hash.

### Examples of INLIST ITERATOR and EXPLAIN PLAN

An INLIST ITERATOR operation appears in the EXPLAIN PLAN output if an index implements an IN-list predicate.

#### Consider the following statement:

```
SELECT * FROM emp WHERE empno IN (7876, 7900, 7902);
```

#### The EXPLAIN PLAN output appears as follows:

OPERATION	OPTIONS	OBJECT_NAME
SELECT STATEMENT		
INLIST ITERATOR		
TABLE ACCESS	BY ROWID	EMP
INDEX	RANGE SCAN	EMP_EMPNO

The INLIST ITERATOR operation iterates over the next operation in the plan for each value in the IN-list predicate. The following sections describe the three possible types of IN-list columns for partitioned tables and indexes.

### When the IN-List Column is an Index Column: Example

If the IN-list column empno is an index column but not a partition column, then the IN-list operator appears before the table operation but after the partition operation in the plan.

OPERATION	OPTIONS	OBJECT	_NAME	PARTIT_	_START	PARTITI_	_STOP



SELECT STATEMENT				
PARTITION RANGE	ALL		KEY(INLIST)	KEY(INLIST)
INLIST ITERATOR				
TABLE ACCESS	BY LOCAL INDEX ROWID	EMP	KEY(INLIST)	KEY(INLIST)
INDEX	RANGE SCAN	EMP_EMPNO	KEY(INLIST)	KEY(INLIST)

The KEY (INLIST) designation for the partition start and stop keys specifies that an IN-list predicate appears on the index start and stop keys.

### When the IN-List Column is an Index and a Partition Column: Example

If empno is an indexed and a partition column, then the plan contains an INLIST ITERATOR operation before the partition operation.

OPERATION	OPTIONS	OBJECT_NAME	PARTITION_START	PARTITION_STOP
SELECT STATEMENT				
INLIST ITERATOR				
PARTITION RANGE	ITERATOR		KEY(INLIST)	KEY(INLIST)
TABLE ACCESS	BY LOCAL INDEX ROWID	EMP	KEY(INLIST)	KEY(INLIST)
INDEX	RANGE SCAN	EMP_EMPNO	KEY(INLIST)	KEY(INLIST)

### When the IN-List Column is a Partition Column: Example

If empno is a partition column and no indexes exist, then no INLIST ITERATOR operation is allocated.

OPERATION	OPTIONS	OBJECT_NAME	PARTITION_START	PARTITION_STOP
SELECT STATEMENT				
PARTITION RANGE	INLIST		KEY(INLIST)	KEY(INLIST)
TABLE ACCESS	FULL	EMP	KEY(INLIST)	KEY(INLIST)

If emp empno is a bitmap index, then the plan is as follows:

OPERATION	OPTIONS	OBJECT_NAME
SELECT STATEMENT		
INLIST ITERATOR		
TABLE ACCESS	BY INDEX ROWID	EMP
BITMAP CONVERSION	TO ROWIDS	
BITMAP INDEX	SINGLE VALUE	EMP_EMPNO

## Example of Domain Indexes and EXPLAIN PLAN

You can use EXPLAIN PLAN to derive user-defined CPU and I/O costs for domain indexes.

EXPLAIN PLAN displays domain index statistics in the OTHER column of PLAN\_TABLE. For example, assume table emp has user-defined operator CONTAINS with a domain index emp\_resume on the resume column, and the index type of emp\_resume supports the operator CONTAINS. You explain the plan for the following query:

```
SELECT * FROM emp WHERE CONTAINS (resume, 'Oracle') = 1
```

#### The database could display the following plan:

OPERATION	OPTIONS	OBJECT_NAME	OTHER
SELECT STATEMENT			
TABLE ACCESS	BY ROWID	EMP	
DOMAIN INDEX		EMP RESUME	CPU: 300, I/O: 4

# **Comparing Execution Plans**

The plan comparison tool takes a reference plan and an arbitrary list of test plans and highlights the differences between them. The plan comparison is logical rather than line by line.

# Purpose of Plan Comparison

The plan comparison report identifies the source of differences, which helps users triage plan reproducibility issues.

The plan comparison report is particularly useful in the following scenarios:

- You want to compare the current plan of a query whose performance is regressing with an old plan captured in AWR.
- A SQL plan baseline fails to reproduce the originally intended plan, and you want to determine the difference between the new plan and the intended plan.
- You want to determine how adding a hint, changing a parameter, or creating an index will affect a plan.
- You want to determine how a plan generated based on a SQL profile or by SQL Performance Analyzer differs from the original plan.

# User Interface for Plan Comparison

You can use DBMS XPLAN.COMPARE PLANS to generate a report in text, XML, or HTML format.

#### **Compare Plans Report Format**

The report begins with a summary. The COMPARE PLANS REPORT section includes information such as the user who ran the report and the number of plans compared, as shown in the following example:

```
COMPARE PLANS REPORT

Current user : SH

Total number of plans : 2

Number of findings : 1
```

The COMPARISON DETAILS section of the report contains the following information:

Plan information

The information includes the plan number, the plan source, plan attributes (which differ depending on the source), parsing schema, and SQL text.

Plans

This section displays the plan rows, including the predicates and notes.

Comparison results

This section summarizes the comparison findings, highlighting logical differences such as join order, join methods, access paths, and parallel distribution method. The findings start at number 1. For findings that relate to a particular query block, the text starts with the name of the block. For findings that relate to a particular object alias, the text starts with the name of the query block and the object alias. The following

```
Comparison Results (1):
```

 Query block SEL\$1, Alias PRODUCTS@SEL\$1: Some columns (OPERATION, OPTIONS, OBJECT\_NAME) do not match between the reference plan (id: 2) and the current plan (id: 2).

#### DBMS\_XPLAN.PLAN\_OBJECT\_LIST Table Type

The plan\_object\_list type allows for a list of generic objects as input to the DBMS XPLAN.COMPARE PLANS function. The syntax is as follows:

TYPE plan\_object\_list IS TABLE OF generic\_plan\_object;

The generic object abstracts the common attributes of plans from all plan sources. Every plan source is a subclass of the <code>generic\_plan\_object</code> superclass. The following table summarizes the different plan sources. Note that when an optional parameter is null, it can correspond to multiple objects. For example, if you do not specify a child number for <code>cursor\_cache\_object</code>, then it matches all cursor cache statements with the specified SQL ID.

Table 6-3 Plan Sources for PLAN\_OBJECT\_LIST

Plan Source	Specification	Description
Plan table	<pre>plan_table_object(owner, plan_table_name, statement_id, plan_id)</pre>	The parameters are as follows:  owner—The owner of the plan table  plan_table_name—The name of the plan table  statement_id—The ID of the statement (optional)  plan_id—The ID of the plan (optional)
Cursor cache	<pre>cursor_cache_object(sql_id, child_number)</pre>	<ul> <li>The parameters are as follows:</li> <li>sql_id—The SQL ID of the plan</li> <li>child_number—The child number of the plan in the cursor cache (optional)</li> </ul>
AWR	<pre>awr_object(sql_id, dbid, con_dbid, plan_hash_value)</pre>	The parameters are as follows:  sql_id—The SQL ID of the plan  dbid—The database ID (optional)  con_dbid—The CDB ID (optional)  plan_hash_value—The hash value of the plan (optional)



Table 6-3 (Cont.) Plan Sources for PLAN\_OBJECT\_LIST

Plan Source	Specification	Description
SQL tuning set	<pre>sqlset_object (sqlset_owner, sqlset_name, sql_id, plan_hash_value)</pre>	The parameters are as follows:  sqlset_owner—The owner of the SQL tuning set  sqlset_name—The name of the SQL tuning set  sql_id—The SQL ID of the plan plan_hash_value—The hash value of the plan (optional)
SQL plan management	<pre>spm_object (sql_handle, plan_name)</pre>	<ul> <li>The parameters are as follows:</li> <li>sql_handle—The SQL handle of plans protected by SQL plan management</li> <li>plan_name—The name of the SQL plan baseline (optional)</li> </ul>
SQL profile	<pre>sql_profile_object (profile_name)</pre>	The profile_name parameter specifies the name of the SQL profile.
Advisor	<pre>advisor_object (task_name, execution_name, sql_id, plan_id)</pre>	The parameters are as follows:  task_name—The name of the advisor task  execution_name—The name of the task execution  sql_id—The SQL ID of the plan  plan_id—The advisor plan ID (optional)

### DBMS\_XPLAN.COMPARE\_PLANS Function

The interface for the compare plan tools is the following function:

The following table describes the parameters that specify that plans to be compared.

Table 6-4 Parameters for the COMPARE\_PLANS Function

Parameter	Description
reference_plan	Specifies a single plan of type generic_plan_object.
compare_plan_list	Specifies a list of plan objects. An object might correspond to one or more plans.



#### **Example 6-8 Comparing Plans from Child Cursors**

This example compares the plan of child cursor number 2 for the SQL ID 8mkxm7ur07za0 with the plan for child cursor number 4 for the same SQL ID.

```
VAR v_report CLOB;

BEGIN
   :v_report := DBMS_XPLAN.COMPARE_PLANS(
     reference_plan => CURSOR_CACHE_OBJECT('8mkxm7ur07za0', 2),
     compare_plan_list =>
PLAN_OBJECT_LIST(CURSOR_CACHE_OBJECT('8mkxm7ur07za0', 4)));
END;
/
PRINT v report
```

#### Example 6-9 Comparing Plan from Child Cursor with Plan from SQL Plan Baseline

This example compares the plan of child cursor number 2 for the SQL ID 8mkxm7ur07za0 with the plan from the SQL plan baseline. The baseline query has a SQL handle of SQL 024d0f7d21351f5d and a plan name of SQL PLAN sdfjkd.

```
VAR v_report CLOB;
BEGIN
   :v_report := DBMS_XPLAN.COMPARE_PLANS( -
      reference_plan => CURSOR_CACHE_OBJECT('8mkxm7ur07za0', 2),
      compare_plan_list => PLAN_OBJECT_LIST(SPM_OBJECT('SQL_024d0f7d21351f5d',
   'SQL_PLAN_sdfjkd')));
END;

PRINT v report
```

#### Example 6-10 Comparing a Plan with Plans from Multiple Sources

This example prints the summary section only. The program compares the plan of child cursor number 2 for the SQL ID <code>8mkxm7ur07za0</code> with every plan in the following list:

- All plans in the shared SQL area that are generated for the SQL ID 8mkxm7ur07za0
- All plans generated in the SQL tuning set SH. SQLT\_WORKLOAD for the SQL ID 6vfqvav0rgyad
- All plans in AWR that are captured for database ID 5 and SQL ID 6vfqvav0rgyad
- The plan baseline for the query with handle SQL\_024d0f7d21351f5d with name SQL\_PLAN\_sdfjkd
- The plan stored in sh.plan table identified by plan id=38
- The plan identified by the SQL profile name pe3r3ejsfd
- All plans stored in SQL advisor identified by task name TASK\_1228, execution name EXEC 1928, and SQL ID 8mkxm7ur07za0

```
VAR v_report CLOB
BEGIN
:v report := DBMS XPLAN.COMPARE PLANS(
```

```
=> CURSOR CACHE OBJECT('8mkxm7ur07za0', 2),
    reference plan
    compare_plan_list => plan_object_list(
         cursor cache object('8mkxm7ur07za0'),
         sqlset_object('SH', 'SQLT_WORKLOAD', '6vfqvav0rgyad'),
         awr object('6vfqvav0rgyad', 5),
         spm object('SQL 024d0f7d21351f5d', 'SQL PLAN sdfjkd'),
         plan table object('SH', 'plan table', 38),
         sql profile object('pe3r3ejsfd'),
         advisor object('TASK 1228', 'EXEC 1928', '8mkxm7ur07za0')),
                     => 'XML',
    type
    level
                      => 'ALL',
    section => 'SUMMARY');
END;
PRINT v report
```

### ✓ Note:

Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS XPLAN package

# Comparing Execution Plans: Tutorial

To compare plans, use the DBMS XPLAN.COMPARE PLANS function.

In this tutorial, you compare two distinct queries. The compare plans report shows that the optimizer was able to use a join elimination transformation in one query but not the other.

#### **Assumptions**

This tutorial assumes that user sh issued the following queries:

```
select count(*)
from products p, sales s
where p.prod_id = s.prod_id
and p.prod_min_price > 200;
select count(*)
from products p, sales s
where p.prod_id = s.prod_id
and s.quantity sold = 43;
```

## To compare execution plans:



As of Oracle Database 23ai, for SELECTS that do not require table access, FROM DUAL is now implicit when there is no FROM clause. FROM DUAL is supported but is no longer required. SELECT <<a href="mailto:sexpr\_list">sexpr\_list</a>; is sufficient. Also note that these operations still appear in the execution plan as FAST DUAL.

- Start SQL\*Plus, and log in to the database with administrative privileges.
- 2. Query V\$SQL to determine the SQL IDs of the two queries.

The following query queries V\$SQL for queries that contain the string products:

```
SET LINESIZE 120

COL SQL_ID FORMAT a20

COL SQL_TEXT FORMAT a60

SELECT SQL_ID, SQL_TEXT

FROM V$SQL

WHERE SQL_TEXT LIKE '%products%'

AND SQL_TEXT NOT LIKE '%SQL_TEXT%'

ORDER BY SQL_ID;

SQL_ID SQL_TEXT

Ohxmvnfkasg6q select count(*) from products p, sales s where p.prod_id = s.prod_id and s.quantity_sold = 43

10dqxjph6bwum select count(*) from products p, sales s where p.prod id = s.prod id and p.prod min price > 200
```

- 3. Log in to the database as user sh.
- 4. Execute the DBMS\_XPLAN.COMPARE\_PLANS function, specifying the SQL IDs obtained in the previous step.

For example, execute the following program:

```
VARIABLE v_rep CLOB

BEGIN
   :v_rep := DBMS_XPLAN.COMPARE_PLANS(
    reference_plan => cursor_cache_object('Ohxmvnfkasg6q', NULL),
    compare_plan_list =>
plan_object_list(cursor_cache_object('10dqxjph6bwum', NULL)),
    type => 'TEXT',
    level => 'TYPICAL',
    section => 'ALL');

END;
//
```

Print the report.

# For example, run the following query.

SET PAGESIZE 50000 SET LONG 100000 SET LINESIZE 210

COLUMN report FORMAT a200 SELECT :v\_rep REPORT FROM DUAL;

The Comparison Results section of the following sample report shows that only the first query used a join elimination transformation:

REPORT

\_\_\_\_\_\_

#### COMPARE PLANS REPORT

\_\_\_\_\_\_

Current user : SH
Total number of plans : 2
Number of findings : 1

#### COMPARISON DETAILS

-----

Plan Number : 1 (Reference Plan)

Plan Found : Yes

Plan Source : Cursor Cache
SQL ID : Ohxmvnfkasg6q

Child Number : 0

Plan Database Version : 19.0.0.0 Parsing Schema : "SH"

SQL Text : select count(\*) from products p, sales s where

p.prod id = s.prod id and s.quantity sold = 43

Plan

-----

Plan Hash Value : 3519235612

-	 I	d 	   	Operation	   	Name	   	Rows	 	Bytes	   	Cost	   	 Time	 
		0		SELECT STATEMENT SORT AGGREGATE						3					
   	*			PARTITION RANGE ALL										00:00:01 00:00:01	

Predicate Information (identified by operation id):

\* 3 - filter("S"."QUANTITY SOLD"=43)

\_\_\_\_\_\_

Plan Number : 2 Plan Found : Yes

Plan Source : Cursor Cache

SQL ID : 10dqxjph6bwum

Child Number : 0

Plan Database Version : 19.0.0.0
Parsing Schema : "SH"

SQL Text : select count(\*) from products p, sales s where p.prod id = s.prod id and p.prod min price > 200

Plan

\_\_\_\_\_

Plan Hash Value : 3037679890

Id  Operation	Name	Rows	Bytes	Cost  Time
0  SELECT STATEMENT   1  SORT AGGREGATE  *2  HASH JOIN  *3  TABLE ACCESS FULL   4  PARTITION RANGE ALL   5  BITMAP CONVERSION TO ROWIDS   6  BITMAP INDEX FAST FULL SCAN	     PRODUCTS       SALES_PROD_BIX	61  918843  918843	549 3675372	34

Predicate Information (identified by operation id):

-----

#### Notes

----

- This is an adaptive plan

#### Comparison Results (1):

\_\_\_\_\_

 Query block SEL\$1: Transformation JOIN REMOVED FROM QUERY BLOCK occurred only in the reference plan (result query block: SEL\$A43D1678).



Oracle Database PL/SQL Packages and Types Reference for more information about the  $\tt DBMS$  XPLAN package

# Comparing Execution Plans: Examples

These examples demonstrate how to generate compare plans reports for queries of tables in the  ${\tt sh}$  schema.

<sup>\* 2 -</sup> access("P"."PROD\_ID"="S"."PROD\_ID")

<sup>\* 3 -</sup> filter("P"."PROD MIN PRICE">200)

## Example 6-11 Comparing an Explained Plan with a Plan in a Cursor

This example explains a plan for a query of tables in the sh schema, and then executes the query:

```
EXPLAIN PLAN

SET STATEMENT_ID='TEST' FOR

SELECT c.cust_city, SUM(s.quantity_sold)

FROM customers c, sales s, products p

WHERE c.cust_id=s.cust_id

AND p.prod_id=s.prod_id

AND prod_min_price>100

GROUP BY c.cust_city;

SELECT c.cust_city, SUM(s.quantity_sold)

FROM customers c, sales s, products p

WHERE c.cust_id=s.cust_id

AND p.prod_id=s.prod_id

AND prod_min_price>100

GROUP BY c.cust city;
```

Assume that the SQL ID of the executed query is 9mp7z6qq83k5y. The following PL/SQL program compares the plan in PLAN\_TABLE and the plan in the shared SQL area:

```
BEGIN
  :v_rep := DBMS_XPLAN.COMPARE_PLANS(
    reference_plan => plan_table_object('SH', 'PLAN_TABLE', 'TEST', NULL),
    compare_plan_list =>
plan_object_list(cursor_cache_object('9mp7z6qq83k5y')),
    type => 'TEXT',
    level => 'TYPICAL',
    section => 'ALL');
END;
/
PRINT v rep
```

The following sample report shows that the plans are the same:

```
Compare plans report

Current user : SH
Total number of plans : 2
Number of findings : 1

Comparison Details

Plan Number : 1 (Reference Plan)
Plan Found : Yes
Plan Source : Plan Table
Plan Table Owner : SH
Plan Table Name : PLAN_TABLE
Statement ID : TEST
Plan ID : 52
```



Plan Database Version : 19.0.0.0 Parsing Schema : "SH"

SQL Text : No SQL Text

Plan

Plan Hash Value : 3473931970

Id	Operation	Name	Rows	Bytes	Cost	Time	
0    1   * 2    3   * 4	SELECT STATEMENT  HASH GROUP BY  HASH JOIN  TABLE ACCESS FULL  HASH JOIN  TABLE ACCESS FULL	        CUSTOMERS    PRODUCTS	160348     55500   160348	22320   5772528   832500	1213  1209  414  472	00:00:01 00:00:01 00:00:01 00:00:01 00:00:01	       
6    7	PARTITION RANGE ALL TABLE ACCESS FULL	  SALES	918843	11026116	467		

## Predicate Information (identified by operation id):

```
_____
* 2 - access("C"."CUST ID"="S"."CUST ID")
```

#### Notes

- This is an adaptive plan

: 2 : Yes Plan Number Plan Found

: Cursor Cache Plan Source SQL ID : 9mp7z6qq83k5y

Child Number : 0 Plan Database Version : 19.0.0.0 : "SH" Parsing Schema

SQL Text : select c.cust city, sum(s.quantity sold) from

customers c, sales s, products p where

c.cust\_id=s.cust\_id and p.prod id=s.prod id and

prod\_min\_price>100 group by c.cust\_city

Plan

-----

Plan Hash Value : 3473931970

Id	Operation		Name	Rows	Bytes   Cost	: Time	
0	SELECT STATEMENT	1		1 1	1213		
1	HASH GROUP BY	1		620	22320 1213	00:00:01	
* 2	HASH JOIN	1		160348	5772528 1209	00:00:01	
3	TABLE ACCESS	FULL	CUSTOMERS	55500	832500  414	00:00:01	
* 4	HASH JOIN	1		160348	3367308  472	00:00:01	



<sup>\* 4 -</sup> access("P"."PROD ID"="S"."PROD ID")

<sup>\* 5 -</sup> filter("PROD MIN PRICE">100)

```
1* 51
                                                 117| 2 |00:00:01 |
        TABLE ACCESS FULL | PRODUCTS | 13|
        PARTITION RANGE ALL | | 918843|11026116| 467 |00:00:01
1 61
          TABLE ACCESS FULL |SALES |918843|11026116| 467 |00:00:01 |
Predicate Information (identified by operation id):
_____
* 2 - access("C"."CUST ID"="S"."CUST ID")
* 4 - access("P"."PROD ID"="S"."PROD ID")
* 5 - filter("PROD MIN PRICE">100)
Notes
____
- This is an adaptive plan
Comparison Results (1):
 1. The plans are the same.
```

## Example 6-12 Comparing Plans in a Baseline and SQL Tuning Set

Assume that you want to compare the plans for the following queries, which differ only in the NO MERGE hint contained in the subquery:

The plan for the first query is captured in a SQL plan management baseline with SQL handle  $SQL\_c522f5888cc4613e$ . The plan for the second query is stored in a SQL tuning set named MYSTS1 and has a SQL ID of d07p7qmrm13nc. You run the following PL/SQL program to compare the plans:

```
VAR v_rep CLOB

BEGIN

v_rep := DBMS_XPLAN.COMPARE_PLANS(
  reference_plan => spm_object('SQL_c522f5888cc4613e'),
  compare_plan_list => plan_object_list(sqlset_object('SH', 'MYSTS1',
'd07p7qmrm13nc', null)),
  type => 'TEXT',
  level => 'TYPICAL',
  section => 'ALL');

END;
```

```
PRINT v rep
```

The following output shows that the only the reference plan, which corresponds to the guery without the hint, used a view merge:

\_\_\_\_\_\_

COMPARE PLANS REPORT

\_\_\_\_\_\_

Current user : SH Total number of plans : 2 Number of findings

#### COMPARISON DETAILS

: 1 (Reference Plan) Plan Number

Plan Found : Yes

Plan Source : SQL Plan Baseline
SQL Handle : SQL\_c522f5888cc4613e
Plan Name : SQL\_PLAN\_ca8rpj26c8s9y7c2279c4

Plan Database Version : 19.0.0.0 : "SH" Parsing Schema

SQL Text : select c.cust city, sum(s.quantity sold) from

> customers c, sales s, (select prod id from products where prod min price>100) p where c.cust id=s.cust id and p.prod id=s.prod id

group by c.cust city

Plan

Plan Hash Value : 2082634180

	Id		Operation		Name	I	Rows	Bytes	C	ost		Time	
	0	I	SELECT STATEMENT			1				22			I
	1		HASH GROUP BY				300	11400		22		00:00:01	
	2		HASH JOIN				718	27284		21		00:00:01	
	3		TABLE ACCESS FULL		CUSTOMERS		630	9450		5		00:00:01	
	4		HASH JOIN				718	16514		15		00:00:01	
	5		TABLE ACCESS FULL		PRODUCTS		573	5730		9		00:00:01	
	6		PARTITION RANGE ALL				960	12480		5		00:00:01	
	7		TABLE ACCESS FULL		SALES		960	12480		5		00:00:01	

\_\_\_\_\_\_

Plan Number Plan Found : Yes

Plan Source : SQL Tuning Set

SQL Tuning Set Owner : SH SQL Tuning Set Name : MYSTS1

SQL ID : d07p7qmrm13nc



Plan Hash Value : 655891922
Plan Database Version : 19.0.0.0
Parsing Schema : "SH"

SQL Text : select c.cust\_city, sum(s.quantity\_sold) from

customers c, sales s, (select /\*+ NO\_MERGE \*/
prod\_id from products where prod\_min\_price>100)

p where c.cust\_id=s.cust\_id and

p.prod\_id=s.prod\_id group by c.cust\_city

Plan

-----

Plan Hash Value : 655891922

Id   Operation   Name   R			1
2   HASH JOIN	300   718   718   718   573   573   960   960   630	23694   21  00:00:01   12924   15  00:00:01   2865   9  00:00:01   5730   9  00:00:01   12480   5  00:00:01   12480   5  00:00:01	

Notes

\_\_\_\_

- This is an adaptive plan

## Comparison Results (1):

-----

1. Query block SEL\$1: Transformation VIEW MERGE occurred only in the reference plan (result query block: SEL\$F5BB74E1).

# Example 6-13 Comparing Plans Before and After Adding an Index

In this example, you test the effect of an index on a query plan:

```
EXPLAIN PLAN

SET STATEMENT_ID='TST1' FOR

SELECT COUNT(*) FROM products WHERE prod_min_price>100;

CREATE INDEX newprodidx ON products(prod_min_price);

EXPLAIN PLAN

SET STATEMENT_ID='TST2' FOR

SELECT COUNT(*) FROM products WHERE prod min price>100;
```

You execute the following PL/SQL program to generate the report:

VAR v\_rep CLOB

BEGIN



```
:v rep := DBMS XPLAN.COMPARE PLANS(
   reference plan => plan table object('SH','PLAN TABLE','TST1',NULL),
   compare plan list => plan object list(plan table object('SH','PLAN TABLE','TST2',NULL)),
              => 'TEXT',
              => 'TYPICAL',
  level
            => 'ALL');
  section
END;
PRINT v_rep
          The following report indicates that the operations in the two plans are different:
          COMPARE PLANS REPORT
           Current user
           Total number of plans : 2
           Number of findings : 1
          COMPARISON DETAILS
          ______
                      : 1 (Reference Plan)
           Plan Number
           Plan Found
                          : Yes
                          : Plan Table
           Plan Source
           Plan Table Owner : SH
Plan Table Name : PLAN_TABLE
                          : TST1
           Statement ID
           Plan ID
                           : 56
           Plan Database Version : 19.0.0.0
           Parsing Schema : "SH"
           SQL Text
                          : No SQL Text
          Plan
          _____
           Plan Hash Value : 3421487369
          ______
                         | Name
          | Id | Operation
                                     | Rows | Bytes | Cost | Time
          ______
          | * 2 | TABLE ACCESS FULL | PRODUCTS | 13 | 65 | 2 | 00:00:01 |
          Predicate Information (identified by operation id):
           * 2 - filter("PROD MIN PRICE">100)
          ______
                          : 2
           Plan Number
           Plan Found
                          : Yes
           Plan Source
                          : Plan Table
           Plan Table Owner
                          : SH
           Plan Table Name : PLAN_TABLE
```

```
Statement ID : TST2
Plan ID : 57
Plan Database Version : 19.0.0.0
Parsing Schema : "SH"
SQL Text : No SQL Text
```

Plan

Plan Hash Value : 2694011010

Id	Operation	Name	ı	Rows		Bytes		Cost		Time	
1	SELECT STATEMENT   SORT AGGREGATE	İ	İ	1	İ	5	İ		İ		İ
* 2	INDEX RANGE SCAN	NEWPRODIDX		13		65		1		00:00:01	

Predicate Information (identified by operation id):
----\* 2 - access("PROD MIN PRICE">100)

## Comparison Results (1):

-----

 Query block SEL\$1, Alias PRODUCTS@SEL\$1: Some columns (OPERATION, OPTIONS, OBJECT\_NAME) do not match between the reference plan (id: 2) and the current plan (id: 2).

## Example 6-14 Comparing Plans with Visible and Invisible Indexes

In this example, an application executes the following query:

```
select count(*)
  from products p, sales s
where p.prod_id = s.prod_id
  and p.prod status = 'obsolete';
```

The plan for this query uses two indexes: <code>sales\_prod\_bix</code> and <code>products\_prod\_status\_bix</code>. The database generates four plans, using all combinations of visible and invisible for both indexes. Assume that SQL plan management accepts the following plans in the baseline for the query:

- sales prod bix visible and products prod status bix visible
- sales prod bix visible and products prod status bix invisible
- sales\_prod\_bix invisible and products\_prod\_status\_bix visible

You make both indexes invisible, and then execute the query again. The optimizer, unable to use the invisible indexes, generates a new plan. The three baseline plans, all of which rely on at least one index being visible, fail to reproduce. Therefore, the optimizer uses the new plan and adds it to the SQL plan baseline for the query. To compare the plan currently in the shared SQL area, which is the reference plan, with all four plans in the baseline, you execute the following PL/SQL code:

VAR v rep CLOB

```
BEGIN
  :v rep := DBMS XPLAN.COMPARE PLANS(
    reference plan => cursor cache object('45ns3tzutg0ds'),
    compare plan list => plan object list(spm object('SQL aec814b0d452da8a')),
                     => 'TEXT',
   TYPE
   level
                     => 'TYPICAL',
   section
                    => 'ALL');
END;
PRINT v_rep
```

### The following report compares all five plans:

#### COMPARE PLANS REPORT

Current user : SH Total number of plans : 5 Number of findings : 19

#### COMPARISON DETAILS

\_\_\_\_\_\_

Plan Number : 1 (Reference Plan) : 1 (F

Plan Found

Plan Source : Cursor Cache : 45ns3tzutq0ds SQL ID

Child Number : 0

Plan Database Version : 19.0.0.0 Parsing Schema : "SH"

SQL Text : select count(\*) from products p, sales s where p.prod id

= s.prod id and p.prod status = 'obsolete'

Plan

\_\_\_\_\_

Plan Hash Value : 1136711713

Id	Operation	   	Name		Rows	 	Bytes	   	Cost		Time	
0   1   * 2   3   * 4	SELECT STATEMENT   SORT AGGREGATE   HASH JOIN   JOIN FILTER CREATE   TABLE ACCESS FULL   JOIN FILTER USE		:BF0000 PRODUCTS :BF0000		1 320 255 255 960		30 9600 6375 6375 4800	   	9 9	 	00:00:01 00:00:01 00:00:01 00:00:01	 
6   * 7	PARTITION RANGE ALL   TABLE ACCESS FULL		SALES		960 960		4800 4800		5 5		00:00:01 00:00:01	

\_\_\_\_\_\_

Predicate Information (identified by operation id):

```
_____
* 2 - access("P"."PROD ID"="S"."PROD ID")
```

<sup>\* 7 -</sup> filter(SYS OP BLOOM FILTER(:BF0000, "S". "PROD ID"))



<sup>\* 4 -</sup> filter("P"."PROD STATUS"='obsolete')

## Notes

\_\_\_\_

- baseline\_repro\_fail = yes

\_\_\_\_\_\_

Plan Number : 2 Plan Found : Yes

Plan Source : SQL Plan Baseline
SQL Handle : SQL aec814b0d452da8a

Plan Name : SQL\_PLAN\_axkOnq3a55qna6e039463

Plan Database Version : 19.0.0.0
Parsing Schema : "SH"

SQL Text : select count(\*) from products p, sales s where p.prod id =

s.prod\_id and p.prod\_status = 'obsolete'

Plan

-----

Plan Hash Value : 1845728355

Id	Operation		Name	Row	 s Byte	es Co	st  Time
0    1    *2    3    *4    *5    6	SELECT STATEMENT SORT AGGREGATE HASH JOIN JOIN FILTER CREATE VIEW HASH JOIN BITMAP CONVERSION TO ROWIDS BITMAP INDEX SINGLE VALUE		<pre>index\$_join\$_001  PRODUCTS_PROD_STATUS_BIX</pre>	255  255  255    255	30   9600   6375   6375     6375	  11   5   5     1	00:00:01
8    9    10	INDEX FAST FULL SCAN JOIN FILTER USE PARTITION RANGE ALL		:BF0000	960  960	6375   4800   4800	5   5	00:00:01   00:00:01   00:00:01
*11	TABLE ACCESS FULL		SALES	1960	4800	5	00:00:01

# Predicate Information (identified by operation id):

```
* 2 - access("P"."PROD_ID"="S"."PROD_ID")
* 4 - filter("P"."PROD_STATUS"='obsolete')
* 5 - access(ROWID=ROWID)
* 7 - access("P"."PROD_STATUS"='obsolete')
* 11 - filter(SYS OP BLOOM FILTER(:BF0000,"S"."PROD_ID"))
```

# Comparison Results (4):

-----

- 1. Query block SEL\$1, Alias P@SEL\$1: Some lines (id: 4) in the reference plan are missing in the current plan.
- 2. Query block SEL\$1, Alias S@SEL\$1: Some columns (ID) do not match between the reference plan (id: 5) and the current plan (id: 9).
- 3. Query block SEL\$1, Alias S@SEL\$1: Some columns (ID, PARENT\_ID, PARTITION\_ID) do not match between the reference plan (id: 6) and the current plan (id: 10).
- 4. Query block SEL\$1, Alias S@SEL\$1: Some columns (ID, PARENT\_ID, PARTITION\_ID) do not match between the reference plan (id: 7) and the current plan (id: 11).

\_\_\_\_\_

Plan Number : 3
Plan Found : Yes

Plan Source : SQL Plan Baseline
SQL Handle : SQL aec814b0d452da8a

Plan Name : SQL PLAN axk0nq3a55qna43c0d821

Plan Database Version : 19.0.0.0
Parsing Schema : "SH"

SQL Text : select count(\*) from products p, sales s where p.prod\_id =

s.prod\_id and

p.prod status = 'obsolete'

Plan

Plan Hash Value : 1136711713

	I	d 		Operation	   	Name		Rows		Bytes		Cost	    -	Time	- 
-       		0 1 2 3 4 5	   	SELECT STATEMENT SORT AGGREGATE HASH JOIN JOIN FILTER CREATE TABLE ACCESS FULL JOIN FILTER USE	       	:BF0000 PRODUCTS :BF0000	       	1 1 320 255 255 960		30 30 9600 6375 6375 4800	   	15 15 9 9		00:00:01   00:00:01   00:00:01   00:00:01   00:00:01	 
	*	6 7	İ	PARTITION RANGE ALL TABLE ACCESS FULL	İ	SALES	İ	960 960	   	4800 4800	İ	5 5		00:00:01   00:00:01	ĺ

# Predicate Information (identified by operation id):

- \* 2 access("P"."PROD\_ID"="S"."PROD\_ID")
  \* 4 filter("P"."PROD\_STATUS"='obsolete')
- \* 7 filter(SYS OP BLOOM FILTER(:BF0000, "S". "PROD ID"))

#### Comparison Results (1):

-----

1. The plans are the same.

\_\_\_\_\_\_

Plan Number : 4
Plan Found : Yes

Plan Source : SQL Plan Baseline SQL Handle : SQL aec814b0d452da8a

Plan Name : SQL PLAN axkOnq3a55qna1b7aea6c

Plan Database Version : 19.0.0.0 Parsing Schema : "SH"

SQL Text : select count(\*) from products p, sales s where p.prod\_id =

s.prod\_id and

p.prod status = 'obsolete'

Plan

\_\_\_\_\_

Plan Hash Value : 461040236

------

Id   Operation	Name	Rows Bytes   Cost   Time
0   SELECT STATEMENT   1   SORT AGGREGATE   2   NESTED LOOPS	   	1   30   10   00:00:01     1   30
* 3   TABLE ACCESS FULL   4   PARTITION RANGE ALL	PRODUCTS	255   6375   9   00:00:01      1   5   10   00:00:01
5   BITMAP CONVERSION COUNT  * 6   BITMAP INDEX SINGLE VALUE	   SALES_PROD_BIX	1   5   10   00:00:01

Predicate Information (identified by operation id):

-----

- \* 3 filter("P"."PROD\_STATUS"='obsolete')
  \* 6 access("P"."PROD ID"="S"."PROD ID")
- Comparison Results (7):

\_\_\_\_\_

- 1. Query block SEL\$1, Alias P@SEL\$1: Some lines (id: 3) in the reference plan are missing in the current plan.
- 2. Query block SEL\$1, Alias S@SEL\$1: Some lines (id: 5) in the reference plan are missing in the current plan.
- 3. Query block SEL\$1, Alias S@SEL\$1: Some lines (id: 7) in the reference plan are missing in the current plan.
- 4. Query block SEL\$1, Alias S@SEL\$1: Some lines (id: 5,6) in the current plan are missing in the reference plan.
- 5. Query block SEL\$1, Alias P@SEL\$1: Some columns (OPERATION) do not match between the reference plan (id: 2) and the current plan (id: 2).
- 6. Query block SEL\$1, Alias P@SEL\$1: Some columns (ID, PARENT\_ID, DEPTH) do not match between the reference plan (id: 4) and the current plan (id: 3).
- 7. Query block SEL\$1, Alias S@SEL\$1: Some columns (ID, PARENT\_ID, DEPTH, POSITION, PARTITION ID) do not match between the reference plan (id: 6) and the current plan (id: 4).

\_\_\_\_\_\_

Plan Number : 5
Plan Found : Yes

Plan Source : SQL Plan Baseline SQL Handle : SQL aec814b0d452da8a

Plan Name : SQL PLAN axk0nq3a55qna0628afbd

Plan Database Version : 19.0.0.0 Parsing Schema : "SH"

SQL Text : select count(\*) from products p, sales s where p.prod id =

s.prod id and

p.prod status = 'obsolete'

Plan

-----

Plan Hash Value : 103329725

Id  Operation	Name	Rows Bytes Cost Time
0  SELECT STATEMENT   1  SORT AGGREGATE   2  NESTED LOOPS   3  VIEW	       index\$_join\$_001	



4	HASH JOIN			I		1			1
5	BITMAP CONVERSION TO ROWIDS			25	5	6375		1	00:00:01
6	BITMAP INDEX SINGLE VALUE		PRODUCTS_PROD_STATUS_BIX						1
7	INDEX FAST FULL SCAN		PRODUCTS_PK	25	5	6375		4	00:00:01
8	PARTITION RANGE ALL				1	5		5	00:00:01
9	BITMAP CONVERSION TO ROWIDS				1	5		5	00:00:01
10	BITMAP INDEX SINGLE VALUE		SALES_PROD_BIX						1

#### Comparison Results (7):

\_\_\_\_\_

- 1. Query block SEL\$1, Alias P@SEL\$1: Some lines (id: 3) in the reference plan are missing in the current plan.
- 2. Query block SEL\$1, Alias P@SEL\$1: Some lines (id: 4) in the reference plan are missing in the current plan.
- 3. Query block SEL\$1, Alias S@SEL\$1: Some lines (id: 5) in the reference plan are missing in the current plan.
- 4. Query block SEL\$1, Alias S@SEL\$1: Some lines (id: 7) in the reference plan are missing in the current plan.
- 5. Query block SEL\$1, Alias S@SEL\$1: Some lines (id: 9,10) in the current plan are missing in the reference plan.
- 6. Query block SEL\$1, Alias P@SEL\$1: Some columns (OPERATION) do not match between the reference plan (id: 2) and the current plan (id: 2).
- 7. Query block SEL\$1, Alias S@SEL\$1: Some columns (ID, PARENT\_ID, DEPTH, POSITION, PARTITION ID) do not match between the reference plan (id: 6) and the current plan (id: 8).

The preceding report shows the following:

- Plan 1 is the reference plan from the shared SQL area. The plan does not use the indexes, which are both invisible, and does not reproduce a baseline plan.
- Plan 2 is in the baseline and assumes sales\_prod\_bix is invisible and products prod status bix is visible.
- Plan 3 is in the baseline and assumes both indexes are invisible. Plan 1 and Plan 3 are the same.
- Plan 4 is in the baseline and assumes sales\_prod\_bix is visible and products prod status bix is invisible.
- Plan 5 is in the baseline and assumes that both indexes are visible.

The comparison report shows that Plan 1 could not reproduce a plan from that baseline. The reason is that the plan in the cursor (Plan 1) was added to the baseline because no baseline plan was available at the time of execution, so the database performed a soft parse of the statement and generated the no-index plan. If the current cursor were to be invalidated, and if the query were to be executed again, then a comparison report would show that the cursor plan did reproduce a baseline plan.

# See Also:

Oracle Database PL/SQL Packages and Types Reference for more information about the  $\tt DBMS\_XPLAN$  package



# Example 6-15 Comparing a Baseline That Fails to Reproduce

One use case is to compare a cost-based plan with a SQL plan baseline. In this example, you create a unique index. The database captures a plan baseline that uses this index. You then make the index invisible and execute the query again. The baseline plan fails to reproduce because the index is not visible, forcing the optimizer to choose a different plan. A compare plans report between the baseline plan and the cost-based plan shows the difference in the access path between the two plans.

1. Log in to the database as user hr, and then create a plan table:

```
CREATE TABLE PLAN TABLE (
STATEMENT ID
                              VARCHAR2 (30),
PLAN ID
                              NUMBER,
TIMESTAMP
                              DATE,
REMARKS
                              VARCHAR2 (4000),
OPERATION
                              VARCHAR2 (30),
OPTIONS
                             VARCHAR2 (255),
OBJECT NODE
                             VARCHAR2 (128),
OBJECT OWNER
                              VARCHAR2(30),
OBJECT NAME
                              VARCHAR2(30),
OBJECT ALIAS
                              VARCHAR2 (65),
OBJECT INSTANCE
                              NUMBER (38),
OBJECT TYPE
                              VARCHAR2 (30),
OPTIMIZER
                              VARCHAR2 (255),
SEARCH COLUMNS
                              NUMBER,
                              NUMBER (38),
PARENT ID
                              NUMBER (38),
DEPTH
                              NUMBER (38),
POSITION
                              NUMBER (38),
COST
                              NUMBER (38),
                              NUMBER (38),
CARDINALITY
BYTES
                              NUMBER (38),
OTHER TAG
                              VARCHAR2 (255),
PARTITION START
                              VARCHAR2 (255),
PARTITION STOP
                              VARCHAR2 (255),
PARTITION ID
                              NUMBER (38),
OTHER
                              LONG,
DISTRIBUTION
                              VARCHAR2 (30),
CPU COST
                              NUMBER (38),
IO COST
                              NUMBER (38),
TEMP SPACE
                              NUMBER (38),
ACCESS PREDICATES
                              VARCHAR2 (4000),
FILTER PREDICATES
                              VARCHAR2 (4000),
PROJECTION
                              VARCHAR2 (4000),
TIME
                              NUMBER (38),
                              VARCHAR2 (30),
QBLOCK NAME
OTHER XML
                              CLOB);
```

2. Execute the following DDL statements, which create a table named staff and an index on the staff.employee id column:

```
CREATE TABLE staff AS (SELECT * FROM employees);
CREATE UNIQUE INDEX staff employee id ON staff (employee id);
```



3. Execute the following statements to place a query of staff under the protection of SQL plan management, and then make the index invisible:

```
ALTER SESSION SET optimizer_capture_sql_plan_baselines = TRUE;
SELECT COUNT(*) FROM staff WHERE employee_id = 20;
-- execute query a second time to create a baseline
SELECT COUNT(*) FROM staff WHERE employee_id = 20;
ALTER SESSION SET optimizer_capture_sql_plan_baselines = FALSE;
ALTER INDEX staff employee id INVISIBLE;
```

4. Explain the plan, and then query the plan table (sample output included):

- dynamic statistics used: dynamic sampling (level=2)
- Failed to use SQL plan baseline for this statement

As the preceding output shows, the optimizer chooses a full table scan because the index is invisible. Because the SQL plan baseline uses an index, the optimizer cannot reproduce the plan.

5. In a separate session, log in as SYS and query the handle and plan name of the SQL plan baseline (sample output included):

```
SET LINESIZE 120

COL SQL_HANDLE FORMAT a25

COL PLAN_NAME FORMAT a35

SELECT DISTINCT SQL_HANDLE, PLAN_NAME, ACCEPTED FROM DBA_SQL_PLAN_BASELINES

WHERE PARSING_SCHEMA_NAME = 'HR';
```



SQL_HANDLE	PLAN_NAME	ACC
SQL_3fa3b23c5ba1bf60	SQL_PLAN_3z8xk7jdu3gv0b7aa092a	YES

**6.** Compare the plans, specifying the SQL handle and plan baseline name obtained from the previous step:

```
VAR v_report CLOB

BEGIN

:v_report := DBMS_XPLAN.COMPARE_PLANS(
   reference_plan => plan_table_object('HR', 'PLAN_TABLE', 'STAFF'),
   compare_plan_list => plan_object_list

(SPM_OBJECT('SQL_3fa3b23c5ba1bf60','SQL_PLAN_3z8xk7jdu3gv0b7aa092a')),
   type => 'TEXT',
   level => 'ALL',
   section => 'ALL');

END;
//
```

7. Query the compare plans report (sample output included):

```
SET LONG 1000000
SET PAGESIZE 50000
SET LINESIZE 200
SELECT : v report rep FROM DUAL;
REP
COMPARE PLANS REPORT
                     : SYS
 Current user
 Total number of plans : 2
 Number of findings : 1
COMPARISON DETAILS
Plan Number : 1 (Reference Plan)
Plan Found
                    : Yes
Plan Source
                    : Plan Table
Plan Table Owner : HR
Plan Table Name : PLAN_TABLE
Statement ID
                    : STAFF
                     : 72
Plan ID
Plan Database Version : 19.0.0.0
Parsing Schema : "HR"
SQL Text
                     : No SQL Text
Plan
_____
Plan Hash Value : 1766070819
```

| Name |Rows| Bytes | Cost | Time

| Id | Operation

```
0| SELECT STATEMENT
                                  13 | 2 | 00:00:01 |
| 1| SORT AGGREGATE
                                         | * 2| TABLE ACCESS FULL | STAFF | 1 | 13 | 2 | 00:00:01 |
Predicate Information (identified by operation id):
_____
* 2 - filter("EMPLOYEE ID"=20)
Notes
- Dynamic sampling used for this statement ( level = 2 )
- baseline repro fail = yes
Plan Number : 2
Plan Found
                : Yes
Plan Source
                 : SQL Plan Baseline
SQL Handle : SQL_3fa3b23c5ba1bf60
Plan Name : SQL_PLAN_3z8xk7jdu3gv0b7aa092a
Plan Database Version : 19.0.0.0
Parsing Schema : "HR"

SQL Text : SELECT COUNT(*) FROM staff WHERE employee_id = 20
Plan
_____
Plan Hash Value : 3081373994
_____
|Id| Operation
              | Name
                                |Rows|Bytes |Cost |Time
|*2| INDEX UNIQUE SCAN | STAFF EMPLOYEE ID | 1 | 13 | 0 |00:00:01|
Predicate Information (identified by operation id):
_____
* 2 - access("EMPLOYEE ID"=20)
Comparison Results (1):
_____
1. Query block SEL$1, Alias "STAFF"@"SEL$1": Some columns (OPERATION,
OPTIONS, OBJECT NAME) do not match between the reference plan (id: 2)
and the current plan (id: 2)
```

