Managing Extended Statistics

DBMS_STATS enables you to collect **extended statistics**, which are statistics that can improve cardinality estimates when multiple predicates exist on different columns of a table, or when predicates use expressions.

An **extension** is either a column group or an expression. Column group statistics can improve cardinality estimates when multiple columns from the same table occur together in a SQL statement. Expression statistics improves optimizer estimates when predicates use expressions, for example, built-in or user-defined functions.



You cannot create extended statistics on virtual columns.

See Also:

Oracle Database SQL Language Reference for a list of restrictions on virtual columns

Managing Column Group Statistics

A column group is a set of columns that is treated as a unit.

Essentially, a column group is a virtual column. By gathering statistics on a column group, the optimizer can more accurately determine the cardinality estimate when a query groups these columns together.

The following sections provide an overview of column group statistics, and explain how to manage them manually.

See Also:

Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS_STATS package

About Statistics on Column Groups

Individual column statistics are useful for determining the selectivity of a single predicate in a WHERE clause.

When the WHERE clause includes multiple predicates on different columns from the same table, individual column statistics do not show the relationship between the columns. This is the problem solved by a column group.

The optimizer calculates the selectivity of the predicates independently, and then combines them. However, if a correlation between the individual columns exists, then the optimizer cannot take it into account when determining a cardinality estimate, which it creates by multiplying the selectivity of each table predicate by the number of rows.

The following graphic contrasts two ways of gathering statistics on the <code>cust_state_province</code> and <code>country_id</code> columns of the <code>sh.customers</code> table. The diagram shows <code>DBMS_STATS</code> collecting statistics on each column individually and on the group. The column group has a system-denerated name.

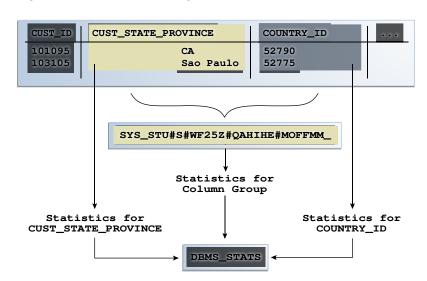


Figure 14-1 Column Group Statistics



The optimizer uses column group statistics for equality predicates, inlist predicates, and for estimating the $\tt GROUP$ BY cardinality.

Why Column Group Statistics Are Needed: Example

This example demonstrates how column group statistics enable the optimizer to give a more accurate cardinality estimate.

The following query of the DBA_TAB_COL_STATISTICS table shows information about statistics that have been gathered on the columns cust_state_province and country_id from the sh.customers table:

```
COL COLUMN_NAME FORMAT a20
COL NDV FORMAT 999

SELECT COLUMN_NAME, NUM_DISTINCT AS "NDV", HISTOGRAM
FROM DBA_TAB_COL_STATISTICS
WHERE OWNER = 'SH'
AND TABLE_NAME = 'CUSTOMERS'
AND COLUMN NAME IN ('CUST STATE PROVINCE', 'COUNTRY ID');
```



Sample output is as follows:

```
COLUMN_NAME NDV HISTOGRAM

CUST_STATE_PROVINCE 145 FREQUENCY
COUNTRY ID 19 FREQUENCY
```

As shown in the following query, 3341 customers reside in California:

Consider an explain plan for a query of customers in the state CA and in the country with ID 52790 (USA):

```
EXPLAIN PLAN FOR
 SELECT *
 FROM sh.customers
 WHERE cust_state_province = 'CA'
    country id=52790;
Explained.
sys@PROD> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN TABLE OUTPUT
______
Plan hash value: 1683234692
| Id | Operation
               | Name | Rows | Bytes |Cost (%CPU)| Time |
 0 | SELECT STATEMENT | | 128 | 24192 | 442 (7) | 00:00:06 |
|* 1 | TABLE ACCESS FULL| CUSTOMERS | 128 | 24192 | 442 (7) | 00:00:06 |
Predicate Information (identified by operation id):
______
PLAN TABLE OUTPUT
______
  1 - filter("CUST STATE PROVINCE"='CA' AND "COUNTRY ID"=52790)
13 rows selected.
```

Based on the single-column statistics for the <code>country_id</code> and <code>cust_state_province</code> columns, the optimizer estimates that the query of California customers in the USA will return 128 rows. In fact, 3341 customers reside in California, but the optimizer does not know that the state of

California is in the country of the USA, and so greatly underestimates cardinality by assuming that both predicates reduce the number of returned rows.

You can make the optimizer aware of the real-world relationship between values in <code>country_id</code> and <code>cust_state_province</code> by gathering column group statistics. These statistics enable the optimizer to give a more accurate cardinality estimate.

See Also:

- "Detecting Useful Column Groups for a Specific Workload"
- "Creating Column Groups Detected During Workload Monitoring"
- "Creating and Gathering Statistics on Column Groups Manually"

Automatic and Manual Column Group Statistics

Oracle Database can create column group statistics either automatically or manually.

The optimizer can use SQL plan directives to generate a more optimal plan. If the <code>DBMS_STATS</code> preference <code>AUTO_STAT_EXTENSIONS</code> is set to <code>ON</code> (by default it is <code>OFF</code>), then a SQL plan directive can automatically trigger the creation of column group statistics based on usage of predicates in the workload. You can set <code>AUTO_STAT_EXTENSIONS</code> with the <code>SET_TABLE_PREFS</code>, <code>SET_GLOBAL_PREFS</code>, or <code>SET_SCHEMA_PREFS</code> procedures.

When you want to manage column group statistics manually, then use DBMS STATS as follows:

- Detect column groups
- Create previously detected column groups
- Create column groups manually and gather column group statistics

See Also:

- "Detecting Useful Column Groups for a Specific Workload"
- "Creating Column Groups Detected During Workload Monitoring"
- "Creating and Gathering Statistics on Column Groups Manually"
- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS STATS procedures for setting optimizer statistics

User Interface for Column Group Statistics

Several DBMS STATS program units have preferences that are relevant for column groups.



Table 14-1 DBMS_STATS APIs Relevant for Column Groups

Program Unit or Preference	Description			
SEED_COL_USAGE Procedure	Iterates over the SQL statements in the specified workload, compiles them, and then seeds column usage information for the columns that appear in these statements.			
	To determine the appropriate column groups, the database must observe a representative workload. You do not need to run the queries themselves during the monitoring period. Instead, you can run EXPLAIN PLAN for some longer-running queries in your workload to ensure that the database is recording column group information for these queries.			
REPORT_COL_USAGE Function	Generates a report that lists the columns that were seen in filter predicates, join predicates, and GROUP BY clauses in the workload.			
	You can use this function to review column usage information recorded for a specific table.			
CREATE_EXTENDED_STATS Function	Creates extensions, which are either column groups or expressions. The database gathers statistics for the extension when either a user-generated or automatic statistics gathering job gathers statistics for the table.			
AUTO_STAT_EXTENSIONS Preference	Controls the automatic creation of extensions, including column groups, when optimizer statistics are gathered. Set this preference using SET_TABLE_PREFS, SET_SCHEMA_PREFS, or SET_GLOBAL_PREFS. When AUTO_STAT_EXTENSIONS is set to OFF (default), the database does not create column group statistics automatically. To create extensions, you must execute the CREATE_EXTENDED_STATS function or specify extended statistics explicitly in the METHOD_OPT parameter in the DBMS_STATS API.			
	When set to ON, a SQL plan directive can trigger the creation of column group statistics automatically based on usage of columns in the predicates in the workload.			

See Also:

- "Setting Artificial Optimizer Statistics for a Table"
- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS STATS package

Detecting Useful Column Groups for a Specific Workload

You can use <code>DBMS_STATS.SEED_COL_USAGE</code> and <code>REPORT_COL_USAGE</code> to determine which column groups are required for a table based on a specified workload.

This technique is useful when you do not know which extended statistics to create. This technique does not work for expression statistics.

Assumptions

This tutorial assumes the following:

- Cardinality estimates have been incorrect for queries of the sh.customers_test table (created from the customers table) that use predicates referencing the columns country id and cust state province.
- You want the database to monitor your workload for 5 minutes (300 seconds).
- You want the database to determine which column groups are needed automatically.

To detect column groups:

- 1. Start SQL*Plus or SQL Developer, and log in to the database as user sh.
- 2. Create the customers test table and gather statistics for it:

```
DROP TABLE customers_test;
CREATE TABLE customers_test AS SELECT * FROM customer;
EXEC DBMS STATS.GATHER TABLE STATS(user, 'customers test');
```

3. Enable workload monitoring.

In a different SQL*Plus session, connect as SYS and run the following PL/SQL program to enable monitoring for 300 seconds:

```
BEGIN
    DBMS_STATS.SEED_COL_USAGE(null,null,300);
END;
/
```

4. As user sh, run explain plans for two queries in the workload.

The following examples show the explain plans for two queries on the <code>customers_test</code> table:

```
EXPLAIN PLAN FOR

SELECT *

FROM customers_test

WHERE cust_city = 'Los Angeles'

AND cust_state_province = 'CA'

AND country_id = 52790;

SELECT PLAN_TABLE_OUTPUT

FROM TABLE(DBMS_XPLAN.DISPLAY('plan_table', null, 'basic rows'));

EXPLAIN PLAN FOR

SELECT country_id, cust_state_province, count(cust_city)

FROM customers_test

GROUP BY country_id, cust_state_province;

SELECT PLAN_TABLE_OUTPUT

FROM TABLE(DBMS_XPLAN.DISPLAY('plan_table', null, 'basic rows'));
```

Sample output appears below:

Id Operation Na	ame	 F	Rows	
0 SELECT STATEMENT 1 TABLE ACCESS FULL CU				
8 rows selected.				
PLAN_TABLE_OUTPUT				
Plan hash value: 3050654408				
Id Operation 1	Name	 	Rows	
0 SELECT STATEMENT 1 HASH GROUP BY 2 TABLE ACCESS FULL (İ	1949 1949 55500	·

9 rows selected.

The first plan shows a cardinality of 1 row for a query that returns 932 rows. The second plan shows a cardinality of 1949 rows for a query that returns 145 rows.

5. Optionally, review the column usage information recorded for the table.

Call the DBMS STATS.REPORT COL USAGE function to generate a report:

```
SET LONG 100000
SET LINES 120
SET PAGES 0
SELECT DBMS_STATS.REPORT_COL_USAGE(user, 'customers_test')
FROM DUAL;
```

The report appears below:

In the preceding report, the first three columns were used in equality predicates in the first monitored query:

```
WHERE cust_city = 'Los Angeles'
AND cust_state_province = 'CA'
AND country id = 52790;
```

All three columns appeared in the same where clause, so the report shows them as a group filter. In the second query, two columns appeared in the <code>GROUP BY</code> clause, so the report labels them as $\texttt{GROUP_BY}$. The sets of columns in the <code>FILTER</code> and $\texttt{GROUP_BY}$ report are candidates for column groups.

See Also:

- "Capturing Workloads in SQL Tuning Sets"
- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS_STATS package

Creating Column Groups Detected During Workload Monitoring

You can use the <code>DBMS_STATS.CREATE_EXTENDED_STATS</code> function to create column groups that were detected previously by executing <code>DBMS_STATS.SEED_COL_USAGE</code>.

Assumptions

This tutorial assumes that you have performed the steps in "Detecting Useful Column Groups for a Specific Workload".

To create column groups:

1. Create column groups for the <code>customers_test</code> table based on the usage information captured during the monitoring window.

For example, run the following query:

```
SELECT DBMS STATS.CREATE EXTENDED STATS(user, 'customers test') FROM DUAL;
```

Note that FROM DUAL is no longer required for queries that do not access tables.

Sample output appears below:

The database created two column groups for <code>customers_test</code>: one column group for the filter predicate and one group for the <code>GROUP BY</code> operation.

2. Regather table statistics.

```
Run GATHER_TABLE_STATS to regather the statistics for customers_test:
EXEC DBMS STATS.GATHER TABLE STATS(user,'customers test');
```

3. As user sh, run explain plans for two queries in the workload.

Check the $\tt USER_TAB_COL_STATISTICS$ view to determine which additional statistics were created by the database:

```
SELECT COLUMN_NAME, NUM_DISTINCT, HISTOGRAM
FROM USER_TAB_COL_STATISTICS
WHERE TABLE_NAME = 'CUSTOMERS_TEST'
ORDER BY 1;
```

Partial sample output appears below:

```
CUST_CITY 620 HEIGHT BALANCED
...

SYS_STU#S#WF25Z#QAHIHE#MOFFMM 145 NONE
SYS_STUMZ$C3AIHLPBROI#SKA58H N 620 HEIGHT BALANCED
```

This example shows the two column group names returned from the DBMS_STATS.CREATE_EXTENDED_STATS function. The column group created on CUST_CITY, CUST_STATE_PROVINCE, and COUNTRY_ID has a height-balanced histogram.

4. Explain the plans again.

The following examples show the explain plans for two queries on the <code>customers_test</code> table:

```
EXPLAIN PLAN FOR
   SELECT *
   FROM    customers_test
   WHERE   cust_city = 'Los Angeles'
   AND     cust_state_province = 'CA'
   AND     country_id = 52790;

SELECT PLAN_TABLE_OUTPUT
FROM   TABLE(DBMS_XPLAN.DISPLAY('plan_table', null,'basic rows'));

EXPLAIN PLAN FOR
```

```
SELECT country_id, cust_state_province, count(cust_city)
FROM customers_test
GROUP BY country_id, cust_state_province;

SELECT PLAN_TABLE_OUTPUT
FROM TABLE(DBMS XPLAN.DISPLAY('plan table', null, 'basic rows'));
```

The new plans show more accurate cardinality estimates:

Id Operation	Name	Rows
0 SELECT STATEMENT 1 TABLE ACCESS FULL	CUSTOMERS_TEST	1093 1093
8 rows selected.		
Plan hash value: 3050654408	}	
Id Operation	Name	Rows
0 SELECT STATEMENT 1 HASH GROUP BY 2 TABLE ACCESS FULI	 CUSTOMERS_TEST	145 145 55500
9 rows selected.		

See Also:

Oracle Database PL/SQL Packages and Types Reference to learn about the ${\tt DBMS_STATS}$ package

Creating and Gathering Statistics on Column Groups Manually

In some cases, you may know the column group that you want to create.

The METHOD_OPT argument of the DBMS_STATS.GATHER_TABLE_STATS function can create and gather statistics on a column group automatically. You can create a new column group by specifying the group of columns using FOR COLUMNS.

Assumptions

This tutorial assumes the following:

- You want to create a column group for the <code>cust_state_province</code> and <code>country_id</code> columns in the <code>customers</code> table in sh schema.
- You want to gather statistics (including histograms) on the entire table and the new column group.

To create a column group and gather statistics for this group:

- In SQL*Plus, log in to the database as the sh user.
- 2. Create the column group and gather statistics.

For example, execute the following PL/SQL program:

```
See Also:
```

Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS STATS.GATHER TABLE STATS procedure

Displaying Column Group Information

To obtain the name of a column group, use the <code>DBMS_STATS.SHOW_EXTENDED_STATS_NAME</code> function or a database view.

You can also use views to obtain information such as the number of distinct values, and whether the column group has a histogram.

Assumptions

This tutorial assumes the following:

- You created a column group for the cust_state_province and country_id columns in the customers table in sh schema.
- You want to determine the column group name, the number of distinct values, and whether a histogram has been created for a column group.

To monitor a column group:

- Start SQL*Plus and connect to the database as the sh user.
- 2. To determine the column group name, do one of the following.
 - Execute the SHOW EXTENDED STATS NAME function.

For example, run the following PL/SQL program:



The output is similar to the following:

Query the USER STAT EXTENSIONS view.

For example, run the following query:

```
SELECT EXTENSION_NAME, EXTENSION

FROM USER_STAT_EXTENSIONS
WHERE TABLE_NAME='CUSTOMERS';

EXTENSION_NAME EXTENSION

SYS STU#S#WF25Z#QAHIHE#MOFFMM_ ("CUST_STATE_PROVINCE", "COUNTRY_ID")
```

Query the number of distinct values and find whether a histogram has been created for a column group.

For example, run the following query:



Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS_STATS.SHOW_EXTENDED_STATS_NAME function

Dropping a Column Group

Use the DBMS STATS.DROP EXTENDED STATS function to delete a column group from a table.

Assumptions

This tutorial assumes the following:

- You created a column group for the <code>cust_state_province</code> and <code>country_id</code> columns in the <code>customers</code> table in sh schema.
- You want to drop the column group.

To drop a column group:

Start SQL*Plus and connect to the database as the sh user.

2. Drop the column group.

For example, the following PL/SQL program deletes a column group from the customers table:

See Also:

Oracle Database PL/SQL Packages and Types Reference to learn about the $\tt DBMS_STATS.DROP_EXTENDED_STATS$ function

Managing Expression Statistics

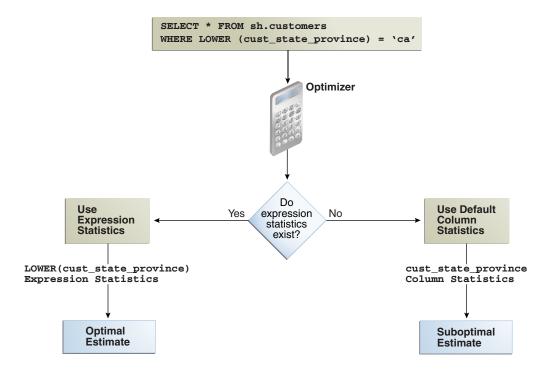
The type of extended statistics known as **expression statistics** improve optimizer estimates when a WHERE clause has predicates that use expressions.

About Expression Statistics

For an **expression** in the form (function(col) = constant) applied to a WHERE clause column, the optimizer does not know how this function affects predicate cardinality unless a function-based index exists. However, you can gather expression statistics on the expression (function(col)) itself.

The following graphic shows the optimizer using statistics to generate a plan for a query that uses a function. The top shows the optimizer checking statistics for the column. The bottom shows the optimizer checking statistics corresponding to the expression used in the query. The expression statistics yield more accurate estimates.

Figure 14-2 Expression Statistics



As shown in Figure 14-2, when expression statistics are not available, the optimizer can produce suboptimal plans.



Oracle Database SQL Language Reference to learn about SQL functions

When Expression Statistics Are Useful: Example

The following query of the sh.customers table shows that 3341 customers are in the state of California:

```
sys@PROD> SELECT COUNT(*) FROM sh.customers WHERE cust_state_province='CA';
COUNT(*)
------
3341
```

Consider the plan for the same query with the LOWER () function applied:

```
sys@PROD> EXPLAIN PLAN FOR
  2 SELECT * FROM sh.customers WHERE LOWER(cust_state_province)='ca';
Explained.
sys@PROD> select * from table(dbms xplan.display);
```

Because no expression statistics exist for <code>LOWER(cust_state_province)='ca'</code>, the optimizer estimate is significantly off. You can use <code>DBMS STATS</code> procedures to correct these estimates.

Creating Expression Statistics

You can use DBMS STATS to create statistics for a user-specified expression.

You can use either of the following program units:

- GATHER TABLE STATS procedure
- CREATE EXTENDED STATISTICS function followed by the GATHER TABLE STATS procedure

Assumptions

This tutorial assumes the following:

- Selectivity estimates are inaccurate for queries of sh.customers that use the UPPER (cust state province) function.
- You want to gather statistics on the UPPER (cust state province) expression.

To create expression statistics:

- 1. Start SQL*Plus and connect to the database as the sh user.
- 2. Gather table statistics.

For example, run the following command, specifying the function in the $method_opt$ argument:

```
See Also:
```

Oracle Database PL/SQL Packages and Types Reference to learn about the $\tt DBMS_STATS.GATHER_TABLE_STATS$ procedure

Displaying Expression Statistics

To obtain information about expression statistics, use the database view <code>DBA_STAT_EXTENSIONS</code> and the <code>DBMS_STATS.SHOW_EXTENDED_STATS_NAME</code> function.

You can also use views to obtain information such as the number of distinct values, and whether the column group has a histogram.

Assumptions

This tutorial assumes the following:

- You created extended statistics for the LOWER (cust state province) expression.
- You want to determine the column group name, the number of distinct values, and whether a histogram has been created for a column group.

To monitor expression statistics:

- 1. Start SQL*Plus and connect to the database as the sh user.
- Query the name and definition of the statistics extension.

For example, run the following query:

```
COL EXTENSION_NAME FORMAT a30
COL EXTENSION FORMAT a35

SELECT EXTENSION_NAME, EXTENSION
FROM USER_STAT_EXTENSIONS
WHERE TABLE NAME='CUSTOMERS';
```

Sample output appears as follows:

```
EXTENSION_NAME EXTENSION

SYS STUBPHJSBRKOIK902YV3W8HOUE (LOWER("CUST STATE PROVINCE"))
```

Query the number of distinct values and find whether a histogram has been created for the expression.

For example, run the following query:

```
SELECT e.EXTENSION expression, t.NUM_DISTINCT, t.HISTOGRAM

FROM USER_STAT_EXTENSIONS e, USER_TAB_COL_STATISTICS t

WHERE e.EXTENSION_NAME=t.COLUMN_NAME

AND e.TABLE_NAME=t.TABLE_NAME

AND t.TABLE_NAME='CUSTOMERS';

EXPRESSION NUM DISTINCT HISTOGRAM
```

```
(LOWER ("CUST_STATE_PROVINCE")) 145 FREQUENCY
```

See Also:

- Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS STATS.SHOW EXTENDED STATS NAME procedure
- Oracle Database Reference to learn about the DBA STAT EXTENSIONS view

Dropping Expression Statistics

To delete a column group from a table, use the DBMS STATS.DROP EXTENDED STATS function.

Assumptions

This tutorial assumes the following:

- You created extended statistics for the LOWER (cust_state_province) expression.
- You want to drop the expression statistics.

To drop expression statistics:

- Start SQL*Plus and connect to the database as the sh user.
- 2. Drop the column group.

For example, the following PL/SQL program deletes a column group from the customers table:

```
BEGIN
   DBMS_STATS.DROP_EXTENDED_STATS(
        'sh'
,   'customers'
,   '(LOWER(cust_state_province))'
);
END;
/
```

See Also:

Oracle Database PL/SQL Packages and Types Reference to learn about the DBMS STATS.DROP EXTENDED STATS procedure