XML Schema Evolution

You can use XML schema evolution to update your XML schema after you have registered it with Oracle XML DB.

Oracle XML DB supports the W3C XML Schema recommendation. XML instance documents that conform to an XML schema can be stored and retrieved using SQL and protocols such as FTP, HTTP(S), and WebDAV. In addition to specifying the structure of XML documents, XML schemas determine the mapping between XML and object-relational storage.

Overview of XML Schema Evolution

A major challenge for developers using an XML schema with Oracle XML DB is how to deal with changes in the content or structure of XML documents. In some environments, the need for changes may be frequent or extensive, arising from new regulations, internal needs, or external opportunities.

Copy-Based Schema Evolution

You perform copy-based XML schema evolution using PL/SQL procedure <code>DBMS_XMLSCHEMA.copyEvolve</code>. This backs up existing instance documents to temporary <code>XMLType</code> tables, drops the old version of the XML schema (which also deletes the associated instance documents), registers the new version, and copies the backed-up instance documents to new <code>XMLType</code> tables.

In-Place XML Schema Evolution

In-place XML schema evolution makes changes to an XML schema without requiring that existing data be copied, deleted, and reinserted. In-place evolution is thus much faster than copy-based evolution. However, in-place evolution also has several restrictions that do not apply to copy-based evolution.

Related Topics

XML Schema Storage and Query: Basic

XML Schema is a standard for describing the content and structure of XML documents. You can register, update, and delete an XML schema used with Oracle XML DB. You can define storage structures to use for your XML schema-based data and map XML Schema data types to SQL data types.

Overview of XML Schema Evolution

A major challenge for developers using an XML schema with Oracle XML DB is how to deal with changes in the content or structure of XML documents. In some environments, the need for changes may be frequent or extensive, arising from new regulations, internal needs, or external opportunities.

For example, you might need to add new elements or attributes to an XML schema definition, modify a data type, or relax or tighten certain minimum and maximum occurrence requirements.

In such cases, you need to "evolve" the XML schema so that new requirements are accommodated, while any existing instance documents (the data) remain valid (or can be made valid), and existing applications can continue to run.

If you do not care about any existing documents, you can of course simply drop the XMLType tables that are dependent on the XML schema, delete the old XML schema, and register the new XML schema at the same URL. In most cases, however, you need to keep the existing documents, possibly transforming them to accommodate the new XML schema.

Oracle XML DB supports two kinds of schema evolution:

- Copy-based schema evolution, in which all instance documents that conform to the schema are copied to a temporary location in the database, the old schema is deleted, the modified schema is registered, and the instance documents are inserted into their new locations from the temporary area
- **In-place schema evolution**, which does not require copying, deleting, and inserting existing data and thus is much faster than copy-based evolution, but which has restrictions that do not apply to copy-based evolution

In general, in-place evolution is permitted if you are not changing the storage model and if the changes do not invalidate existing documents (that is, if existing documents are conformant with the new schema or can be made conformant with it). A more detailed explanation of restrictions and guidelines is presented in In-Place XML Schema Evolution.

Each approach has its own PL/SQL procedure: DBMS_XMLSCHEMA.copyEvolve for copy-based evolution, DBMS_XMLSCHEMA.inPlaceEvolve for in-place evolution. This chapter explains the use of each procedure and presents guidelines for using its associated approach to schema evolution.

Copy-Based Schema Evolution

You perform copy-based XML schema evolution using PL/SQL procedure <code>DBMS_XMLSCHEMA.copyEvolve</code>. This backs up existing instance documents to temporary <code>XMLType</code> tables, drops the old version of the XML schema (which also deletes the associated instance documents), registers the new version, and copies the backed-up instance documents to new <code>XMLType</code> tables.

In case of a problem, the backup copies are restored — see Rollback When Procedure DBMS_XMLSCHEMA.COPYEVOLVE Raises an Error.

Using procedure <code>copyEvolve</code>, you can evolve your registered XML schema in such a way that existing XML instance documents continue to be valid.

- Scenario for Copy-Based Evolution
 An evolved version of a purchase-order XML schema is shown. It is used in examples that illustrate the use of copy-based XML schema evolution.
- COPYEVOLVE Parameters and Errors
 The parameters of PL/SQL procedure DBMS_XMLSCHEMA.copyEvolve are described, as are the errors associated with this procedure.
- Limitations of Procedure COPYEVOLVE

 The use of PL/SQL procedure DBMS_XMLSCHEMA.copyEvolve involves certain limitations.
- Guidelines for Using Procedure COPYEVOLVE
 General guidelines for using PL/SQL procedure DBMS_XMLSCHEMA.copyEvolveare
 presented, as well as guidelines that are specific to particular contexts.
- Update of Existing XML Instance Documents Using an XSLT Stylesheet
 After you modify a registered XML schema, you must update any existing XML instance
 documents that use the schema. You do this by applying an XSLT stylesheet to each of the
 instance documents. The stylesheet represents the difference between the old and new
 XML schemas.

Examples of Using Procedure COPYEVOLVE

Several examples are presented of using PL/SQL procedureDBMS_XMLSCHEMA.copyEvolve to update an XML schema. (Be sure to back up all registered XML schemas and XML documents that reference them, before using the procedure.)

Scenario for Copy-Based Evolution

An evolved version of a purchase-order XML schema is shown. It is used in examples that illustrate the use of copy-based XML schema evolution.

Example 20-1 shows a *partial* listing of a revised version of the purchase-order XML schema of Example A-2. See Example A-3 for the *complete* revised schema listing. Text that is in **bold** here is new or different from that in the original schema.

Example 20-1 Revised Purchase-Order XML Schema

```
<xs:schema xmlns:xs="http://www.w3.org/2001/XMLSchema"</pre>
           xmlns:xdb="http://xmlns.oracle.com/xdb"
           version="1.0">
  <xs:element.</pre>
   name="PurchaseOrder" type="PurchaseOrderType"
   xdb:defaultTable="PURCHASEORDER"
   xdb:columnProps=
      "CONSTRAINT purchaseorder pkey PRIMARY KEY (XMLDATA.reference),
       CONSTRAINT valid_email_address FOREIGN KEY (XMLDATA.userid)
         REFERENCES hr.employees (EMAIL) "
    xdb:tableProps=
      "VARRAY XMLDATA.ACTIONS.ACTION STORE AS TABLE ACTION TABLE
        ((CONSTRAINT action pkey PRIMARY KEY (NESTED TABLE ID, SYS NC ARRAY INDEX$)))
       VARRAY XMLDATA.LINEITEMS.LINEITEM STORE AS TABLE LINEITEM TABLE
        ((constraint LINEITEM PKEY primary key (NESTED TABLE ID, SYS NC ARRAY INDEX$)))
       lob (XMLDATA.NOTES) STORE AS (ENABLE STORAGE IN ROW STORAGE (INITIAL 4K NEXT 32K))"/>
  <xs:complexType name="PurchaseOrderType" xdb:SQLType="PURCHASEORDER T">
    <xs:sequence>
      <xs:element name="Actions" type="ActionsType" xdb:SQLName="ACTIONS"/>
      <xs:element name="Reject" type="RejectionType" minOccurs="0" xdb:SQLName="REJECTION"/>
      <xs:element name="Requestor" type="RequestorType" xdb:SQLName="REQUESTOR"/>
      <xs:element name="User" type="UserType" xdb:SQLName="USERID"/>
      <xs:element name="CostCenter" type="CostCenterType" xdb:SQLName="COST_CENTER"/>
      <xs:element name="BillingAddress" type="AddressType" minOccurs="0"</pre>
                  xdb:SQLName="BILLING ADDRESS"/>
      <xs:element name="ShippingInstructions" type="ShippingInstructionsType"</pre>
                 xdb:SQLName="SHIPPING INSTRUCTIONS"/
      <xs:element name="SpecialInstructions" type="SpecialInstructionsType"</pre>
                 xdb:SQLName="SPECIAL INSTRUCTIONS"/>
      <xs:element name="LineItems" type="LineItemsType" xdb:SQLName="LINEITEMS"/>
      <xs:element name="Notes" type="NotesType" minOccurs="0" xdb:SQLType="CLOB"</pre>
                  xdb:SQLName="NOTES"/>
    </xs:sequence>
    <xs:attribute name="Reference" type="ReferenceType" use="required" xdb:SQLName="REFERENCE"/>
    <xs:attribute name="DateCreated" type="xs:dateTime" use="required"</pre>
                  xdb:SQLType="TIMESTAMP WITH TIME ZONE"/>
  </xs:complexType>
  <xs:complexType name="LineItemsType" xdb:SQLType="LINEITEMS T">
    <xs:sequence>
      <xs:element name="LineItem" type="LineItemType" maxOccurs="unbounded" xdb:SQLName="LINEITEM"</pre>
                 xdb:SQLCollType="LINEITEM V"/>
    </xs:sequence>
  </xs:complexType>
  <xs:complexType name="LineItemType" xdb:SQLType="LINEITEM T">
    <xs:sequence>
      <xs:element name="Part" type="PartType" xdb:SQLName="PART"/>
      <xs:element name="Quantity" type="quantityType"/>
    </xs:sequence>
    <xs:attribute name="ItemNumber" type="xs:integer" xdb:SQLName="ITEMNUMBER"</pre>
                  xdb:SQLType="NUMBER"/>
  </xs:complexType>
```

```
<xs:complexType name="PartType" xdb:SQLType="PART T">
  <xs:simpleContent>
    <xs:extension base="UPCCodeType">
      <xs:attribute name="Description" type="DescriptionType" use="required"</pre>
                    xdb:SQLName="DESCRIPTION"/>
      <xs:attribute name="UnitCost" type="moneyType" use="required"/>
    </xs:extension>
  </xs:simpleContent>
</xs:complexType>
<xs:simpleType name="ReferenceType">
  <xs:restriction base="xs:string">
    <xs:minLength value="18"/>
    <xs:maxLength value="30"/>
  </xs:restriction>
</xs:simpleType>
<xs:complexType name="RejectionType" xdb:SQLType="REJECTION T">
    <xs:element name="User" type="UserType" minOccurs="0" xdb:SQLName="REJECTED_BY"/>
    <xs:element name="Date" type="DateType" minOccurs="0" xdb:SQLName="DATE REJECTED"/>
    <xs:element name="Comments" type="CommentsType" minOccurs="0" xdb:SQLName="REASON REJECTED"/>
  </xs:all>
</xs:complexType>
<xs:complexType name="ShippingInstructionsType" xdb:SQLType="SHIPPING INSTRUCTIONS T">
  <xs:sequence>
    <xs:element name="name" type="NameType" minOccurs="0" xdb:SQLName="SHIP TO NAME"/>
      <xs:element name="address" type="AddressType" minOccurs="0"/>
      <xs:element name="fullAddress" type="FullAddressType" minOccurs="0"</pre>
                  xdb:SQLName="SHIP_TO_ADDRESS"/>
    </xs:choice>
    <xs:element name="telephone" type="TelephoneType" minOccurs="0" xdb:SQLName="SHIP TO PHONE"/>
  </xs:sequence>
</xs:complexType>
<xs:simpleType name="NameType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="20"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="FullAddressType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="256"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="DescriptionType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="256"/>
  </xs:restriction>
</xs:simpleType>
<xs:complexType name="AddressType" xdb:SQLType="ADDRESS T">
  <xs:sequence>
    <xs:element name="StreetLine1" type="StreetType"/>
    <xs:element name="StreetLine2" type="StreetType" minOccurs="0"/>
    <xs:element name="City" type="CityType"/>
    <xs:choice>
      <xs:sequence>
        <xs:element name="State" type="StateType"/>
        <xs:element name="ZipCode" type="ZipCodeType"/>
```

```
</xs:sequence>
      <xs:sequence>
        <xs:element name="Province" type="ProvinceType"/>
        <xs:element name="PostCode" type="PostCodeType"/>
      </xs:sequence>
      <xs:sequence>
        <xs:element name="County" type="CountyType"/>
        <xs:element name="Postcode" type="PostCodeType"/>
      </xs:sequence>
    </xs:choice>
    <xs:element name="Country" type="CountryType"/>
  </xs:sequence>
</xs:complexType>
<xs:simpleType name="StreetType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="128"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="CityType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="64"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="StateType">
  <xs:restriction base="xs:string">
    <xs:minLength value="2"/>
    <xs:maxLength value="2"/>
    <xs:enumeration value="AK"/>
    <xs:enumeration value="AL"/>
    <xs:enumeration value="AR"/>
. . . -- A value for each US state abbreviation
    <xs:enumeration value="WY"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="ZipCodeType">
  <xs:restriction base="xs:string">
    <xs:pattern value="\d{5}"/>
    <xs:pattern value="\d{5}-\d{4}"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="CountryType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="64"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="CountyType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="32"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="PostCodeType">
  <xs:restriction base="xs:string">
    <xs:minLength value="1"/>
    <xs:maxLength value="12"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="ProvinceType">
  <xs:restriction base="xs:string">
    <xs:minLength value="2"/>
    <xs:maxLength value="2"/>
  </xs:restriction>
</xs:simpleType>
<xs:simpleType name="NotesType">
  <xs:restriction base="xs:string">
```

```
<xs:maxLength value="32767"/>
    </xs:restriction>
</xs:simpleType>
<xs:simpleType name="UPCCodeType">
    <xs:restriction base="xs:string">
        <xs:minLength value="11"/>
        <xs:maxLength value="14"/>
        <xs:pattern value="\d{11}"/>
        <xs:pattern value="\d{12}"/>
        <xs:pattern value="\d{13}"/>
        <xs:pattern value="\d{14}"/>
        <xs:pattern value="\d{14}"/>
        <xs:pattern value="\d{14}"/>
        <xs:pattern value="\d{14}"/>
        </xs:restriction>
        </xs:simpleType>
</xs:schema>
```

COPYEVOLVE Parameters and Errors

The parameters of PL/SQL procedure <code>DBMS_XMLSCHEMA.copyEvolve</code> are described, as are the errors associated with this procedure.

This is the signature of procedure DBMS XMLSCHEMA.copyEvolve:

Table 20-1 describes the individual parameters. Table 20-2 describes the errors associated with the procedure.

Table 20-1 Parameters of Procedure DBMS_XMLSCHEMA.COPYEVOLVE

Parameter	Description
schemaURLs	Varray of URLs of XML schemas to be evolved (varray of VARCHAR2 (4000). This should include the dependent schemas as well. Unless the force parameter is TRUE, the URLs should be in the dependency order, that is, if URL A comes before URL B in the varray, then schema A should not be dependent on schema B but schema B may be dependent on schema A.
newSchemas	Varray of new XML schema documents (XMLType instances). Specify this in exactly the same order as the corresponding URLs. If no change is necessary in an XML schema, provide the unchanged schema.
transforms	Varray of XSL documents (XMLType instances) that are applied to XML schema based documents to make them conform to the new schemas. Specify these in exactly the same order as the corresponding URLs. If no transformations are required, this parameter need not be specified.
preserveOldDocs	If this is TRUE, then the temporary tables holding old data are not dropped at the end of schema evolution. See also Guidelines for Using Procedure COPYEVOLVE.
mapTabName	Specifies the name of table that maps old XMLType table or column names to names of corresponding temporary tables.

Table 20-1 (Cont.) Parameters of Procedure DBMS_XMLSCHEMA.COPYEVOLVE

Parameter	Description
generateTables	By default this parameter is TRUE. If FALSE then XMLType tables or columns are not generated after registering new XML schemas. If FALSE, preserveOldDocs must be TRUE and mapTabName must not be NULL.
force	If this is TRUE, then errors during the registration of new schemas are ignored. If there are circular dependencies among the schemas, set this flag to TRUE to ensure that each schema is stored even though there may be errors in registration.
schemaOwners	Varray of names of schema owners. Specify these in exactly the same order as the corresponding URLs.
parallelDegree	Specifies the degree of parallelism to be used in a PARALLEL hint during the data-copy stage. If this is 0 (default value), a PARALLEL hint is absent from the data-copy statements.
options	Miscellaneous options. The only option is <code>COPYEVOLVE_BINARY_XML</code> , which means to register the new XML schemas for binary XML data and create the new tables or columns with binary XML as the storage model.

Table 20-2 Errors Associated with Procedure DBMS_XMLSCHEMA.COPYEVOLVE

Error Number and Message	Cause	Action
30942 XML Schema Evolution error for schema ' <schema_url>' table "<owner_name>.<table_nam e="">" column '<column_name>'</column_name></table_nam></owner_name></schema_url>	The given XMLType table or column that conforms to the given XML schema had errors during evolution. In the case of a table, the column name is empty. See also the more specific error that follows this.	Based on the schema, table, and column information in this error and the more specific error that follows, take corrective action.
30943 XML Schema ' <schema_url>' is dependent on XML schema '<schema_url>'</schema_url></schema_url>	Not all dependent XML schemas were specified or the schemas were not specified in dependency order, that is, if schema S1 is dependent on schema S, S must appear before S1.	Include the previously unspecified schema in the list of schemas or correct the order in which the schemas are specified. Then retry the operation.
30944 Error during rollback for XML schema ' <schema_url>' table "<owner_name>.<table_nam e>" column '<column_name>'</column_name></table_nam </owner_name></schema_url>	The given XMLType table or column that conforms to the given XML schema had errors during a rollback of XML schema evolution. For a table, the column name is empty. See also the more specific error that follows this.	Based on the schema, table, and column information in this error and the more specific error that follows, take corrective action.
30945 Could not create mapping table ' <table_name>'</table_name>	A mapping table could not be created during XML schema evolution. See also the more specific error that follows this.	Ensure that a table with the given name does not exist and retry the operation.
30946 XML Schema Evolution warning: temporary tables not cleaned up	An error occurred after the schema was evolved while cleaning up temporary tables. The schema evolution was successful.	If you need to remove the temporary tables, use the mapping table to get the temporary table names and drop them.



Limitations of Procedure COPYEVOLVE

The use of PL/SQL procedure DBMS XMLSCHEMA.copyEvolve involves certain limitations.

- Indexes, triggers, constraints, row-level security (RLS) policies, and other metadata related
 to the XMLType tables that are dependent on the schemas are not preserved. These must
 be re-created after evolution.
- If top-level element names are changed, additional steps are required after copyEvolve finishes executing. See Top-Level Element Name Changes.
- Copy-based evolution cannot be used if there is a table with an object-type column that
 has an XMLType attribute that is dependent on any of the schemas to be evolved. For
 example, consider this table:

```
CREATE TYPE t1 AS OBJECT (n NUMBER, x XMLType);
CREATE TABLE tab1 (e NUMBER, o t1) XMLType
COLUMN o.x XMLSchema "s1.xsd" ELEMENT "Employee";
```

This assumes that an XML schema with a top-level element <code>Employee</code> has been registered under URL ${\tt s1.xsd.}$ It is not possible to evolve this XML schema, because table <code>tab1</code> with column o with <code>XMLType</code> attribute <code>x</code> is dependent on the XML schema. Although <code>copyEvolve</code> does not handle <code>XMLType</code> object attributes, it does raise an error in such cases.

Guidelines for Using Procedure COPYEVOLVE

General guidelines for using PL/SQL procedure DBMS_XMLSCHEMA.copyEvolveare presented, as well as guidelines that are specific to particular contexts.

The following general guideline applies to using <code>copyEvolve</code>. The rest of this section describes specific guidelines that can also be appropriate in particular contexts.

1. Turn off the recycle bin, to prevent dropped tables from being copied to it:

```
ALTER SESSION SET RECYCLEBIN=off;
```

2. Identify the XML schemas that are dependent on the XML schema that is to be evolved. You can acquire the URLs of the dependent XML schemas using the following query, where schema to be evolved is the schema to be evolved, and

```
owner of schema to be evolved is its owner (database user).
```

```
SELECT dxs.SCHEMA_URL, dxs.OWNER

FROM DBA_DEPENDENCIES dd, DBA_XML_SCHEMAS dxs

WHERE dd.REFERENCED_NAME = (SELECT INT_OBJNAME

FROM DBA_XML_SCHEMAS

WHERE SCHEMA_URL = schema_to_be_evolved

AND OWNER = owner_of_schema_to_be_evolved)

AND dxs.INT_OBJNAME = dd.NAME;
```

In many cases, no changes are needed in the dependent XML schemas. But if the dependent XML schemas need to be changed, then you must also prepare new versions of those XML schemas.

3. If the existing instance documents do not conform to the new XML schema, then you must provide an XSL stylesheet that, when applied to an instance document, transforms it to conform to the new schema. You must do this for each XML schema identified in Step 2.

The transformation must handle documents that conform to all top-level elements in the new XML schema.

4. Call procedure DBMS_XMLSCHEMA.copyEvolve, specifying the XML schema URLs, new schemas, and transformation stylesheet.

Top-Level Element Name Changes

PL/SQL procedure DBMS_XMLSCHEMA.copyEvolve assumes that top-level elements have not been dropped in new schemas and that their names have not been changed. If there are such changes then call procedure <code>copyEvolve</code> with parameter <code>generateTables</code> set to <code>FALSE</code> and parameter <code>preserveOldDocs</code> set to <code>TRUE</code>, so that backup document copies are available.

- User-Created Virtual Columns of Tables Other Than Default Tables

 For tables that are not default tables, any virtual columns that you create are not re-created during copy-based evolution. If such columns are needed then set <code>copyEvolve</code> parameter <code>preserveOldDocs</code> to <code>TRUE</code>, create the tables, and copy the old documents after procedure <code>copyEvolve</code> has finished.
- Ensure That the XML Schema and Dependents Are Not Used by Concurrent Sessions
 Ensure that the XML schema to be evolved, as well as its dependents, are not used by any
 concurrent session during the XML schema evolution process.
- Rollback When Procedure DBMS_XMLSCHEMA.COPYEVOLVE Raises an Error Procedure DBMS_XMLSCHEMA.copyEvolve either completely succeeds or it raises an error. If it raises an error then it tries to roll back as much of the operation as possible.
- Failed Rollback From Insufficient Privileges
 In certain cases you cannot roll back a copy-based evolution operation. For example, if table creation fails due to reasons not related to the new XML schema, then there is no way to roll back.
- Privileges Needed for XML Schema Evolution
 There are several database privileges that you might need, in order to perform copy-based XML schema evolution.

Top-Level Element Name Changes

PL/SQL procedure <code>DBMS_XMLSCHEMA.copyEvolve</code> assumes that top-level elements have not been dropped in new schemas and that their names have not been changed. If there are such changes then call procedure <code>copyEvolve</code> with parameter <code>generateTables</code> set to <code>FALSE</code> and parameter <code>preserveOldDocs</code> set to <code>TRUE</code>, so that backup document copies are available.

With those parameter values, new tables are not generated, and the temporary tables holding the old documents (backup copies) are not dropped at the end of the procedure. You can then store the old documents in whatever form is appropriate and drop the temporary tables. See COPYEVOLVE Parameters and Errors for more details on using these parameters.

User-Created Virtual Columns of Tables Other Than Default Tables

For tables that are not default tables, any virtual columns that you create are not re-created during copy-based evolution. If such columns are needed then set <code>copyEvolve</code> parameter <code>preserveOldDocs</code> to <code>TRUE</code>, create the tables, and copy the old documents after procedure <code>copyEvolve</code> has finished.



Ensure That the XML Schema and Dependents Are Not Used by Concurrent Sessions

Ensure that the XML schema to be evolved, as well as its dependents, are not used by any concurrent session during the XML schema evolution process.

If other, concurrent sessions have shared locks on this schema at the beginning of the evolution process, then procedure <code>DBMS_XMLSCHEMA.copyEvolve</code> waits for these sessions to release the locks so that it can acquire an exclusive lock. However, this lock is released immediately to allow the rest of the process to continue.

Rollback When Procedure DBMS XMLSCHEMA.COPYEVOLVE Raises an Error

Procedure DBMS_XMLSCHEMA.copyEvolve either completely succeeds or it raises an error. If it raises an error then it tries to roll back as much of the operation as possible.

Evolving an XML schema involves many database DDL statements. When an error occurs, compensating DDL statements are executed to undo the effect of all steps executed to that point. If the old tables or schemas have been dropped, they are re-created, but any table, column, and storage properties and any auxiliary structures (such as indexes, triggers, constraints, and RLS policies) associated with the tables and columns are lost.

Failed Rollback From Insufficient Privileges

In certain cases you cannot roll back a copy-based evolution operation. For example, if table creation fails due to reasons not related to the new XML schema, then there is no way to roll back.

An example is failure due to insufficient privileges. The temporary tables are not deleted even if preserveOldDocs is FALSE, so the data can be recovered. If the mapTabName parameter is null, the mapping table name is XDB\$MAPTAB followed by a sequence number. The exact table name can be found using a query such as the following:

SELECT TABLE_NAME FROM USER_TABLES WHERE TABLE_NAME LIKE 'XDB\$MAPTAB%';

Privileges Needed for XML Schema Evolution

There are several database privileges that you might need, in order to perform copy-based XML schema evolution.

Copy-based XML schema evolution can involve dropping or creating SQL data types, so you need type-related privileges such as DROP TYPE, CREATE TYPE, and ALTER TYPE. You need privileges to delete and register the XML schemas involved in the evolution. You need all privileges on XMLType tables that conform to the schemas being evolved. For XMLType columns, the ALTER TABLE privilege is needed on corresponding tables. If there are schema-based XMLType tables or columns in other database schemas, you need privileges such as the following:

- CREATE ANY TABLE
- CREATE ANY INDEX
- SELECT ANY TABLE
- READ ANY TABLE



- UPDATE ANY TABLE
- INSERT ANY TABLE
- DELETE ANY TABLE
- DROP ANY TABLE
- ALTER ANY TABLE
- DROP ANY INDEX

To avoid needing to grant all these privileges to the database- schema owner, Oracle recommends that a database administrator perform the evolution if there are XML schema-based XMLType table or columns belonging to other database schemas.

Update of Existing XML Instance Documents Using an XSLT Stylesheet

After you modify a registered XML schema, you must update any existing XML instance documents that use the schema. You do this by applying an XSLT stylesheet to each of the instance documents. The stylesheet represents the difference between the old and new XML schemas.

Example 20-2 shows an XSLT stylesheet, in file <code>evolvePurchaseOrder.xsl</code>, that transforms existing purchase-order documents that use the old XML schema, so they use the new XML schema instead.

Example 20-2 evolvePurchaseOrder.xsl: XSLT Stylesheet to Update Instance Documents

```
<?xml version="1.0" encoding="UTF-8"?>
<xsl:stylesheet</pre>
 version="1.0"
 xmlns:xsl="http://www.w3.org/1999/XSL/Transform"
 xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
 <xsl:output method="xml" encoding="UTF-8"/>
  <xsl:template match="/PurchaseOrder">
    <PurchaseOrder>
      <xsl:attribute name="xsi:noNamespaceSchemaLocation">
       http://localhost:8080/source/schemas/poSource/xsd/purchaseOrder.xsd
      </xsl:attribute>
      <xsl:for-each select="Reference">
        <xsl:attribute name="Reference">
          <xsl:value-of select="."/>
        </xsl:attribute>
      </xsl:for-each>
      <xsl:variable name="V264 394" select="'2004-01-01T12:00:00.000000-08:00'"/>
      <xsl:attribute name="DateCreated">
        <xsl:value-of select="$V264 394"/>
      </xsl:attribute>
      <xsl:for-each select="Actions">
        <Actions>
          <xsl:for-each select="Action">
            <Action>
              <xsl:for-each select="User">
                <User>
                  <xsl:value-of select="."/>
                </User>
              </xsl:for-each>
              <xsl:for-each select="Date">
                <Date>
                  <xsl:value-of select="."/>
                </Date>
              </xsl:for-each>
            </Action>
          </xsl:for-each>
        </Actions>
```



```
</xsl:for-each>
<xsl:for-each select="Reject">
  <Reject>
    <xsl:for-each select="User">
      <User>
        <xsl:value-of select="."/>
      </User>
    </xsl:for-each>
    <xsl:for-each select="Date">
      <Date>
        <xsl:value-of select="."/>
      </Date>
    </xsl:for-each>
    <xsl:for-each select="Comments">
        <xsl:value-of select="."/>
      </Comments>
    </xsl:for-each>
  </Reject>
</xsl:for-each>
<xsl:for-each select="Requestor">
  <Requestor>
    <xsl:value-of select="."/>
  </Requestor>
</xsl:for-each>
<xsl:for-each select="User">
  <User>
    <xsl:value-of select="."/>
  </User>
</xsl:for-each>
<xsl:for-each select="CostCenter">
  <CostCenter>
    <xsl:value-of select="."/>
  </CostCenter>
</xsl:for-each>
<ShippingInstructions>
  <xsl:for-each select="ShippingInstructions">
    <xsl:for-each select="name">
      <name>
        <xsl:value-of select="."/>
      </name>
    </xsl:for-each>
  </xsl:for-each>
  <xsl:for-each select="ShippingInstructions">
    <xsl:for-each select="address">
      <fullAddress>
        <xsl:value-of select="."/>
      </fullAddress>
    </xsl:for-each>
  </xsl:for-each>
  <xsl:for-each select="ShippingInstructions">
    <xsl:for-each select="telephone">
      <telephone>
        <xsl:value-of select="."/>
      </telephone>
    </xsl:for-each>
  </xsl:for-each>
</ShippingInstructions>
<xsl:for-each select="SpecialInstructions">
  <SpecialInstructions>
    <xsl:value-of select="."/>
  </SpecialInstructions>
</xsl:for-each>
<xsl:for-each select="LineItems">
  <LineItems>
    <xsl:for-each select="LineItem">
      <xsl:variable name="V22" select="."/>
      <LineItem>
        <xsl:for-each select="@ItemNumber">
          <xsl:attribute name="ItemNumber">
```



```
<xsl:value-of select="."/>
                </xsl:attribute>
              </xsl:for-each>
              <xsl:for-each select="$V22/Part">
                <xsl:variable name="V24" select="."/>
                <xsl:for-each select="@Id">
                  <Part>
                    <xsl:for-each select="$V22/Description">
                      <xsl:attribute name="Description">
                        <xsl:value-of select="."/>
                     </xsl:attribute>
                    </xsl:for-each>
                    <xsl:for-each select="$V24/@UnitPrice">
                      <xsl:attribute name="UnitCost">
                        <xsl:value-of select="."/>
                      </xsl:attribute>
                    </xsl:for-each>
                    <xsl:value-of select="."/>
                  </Part>
                </xsl:for-each>
              </xsl:for-each>
              <xsl:for-each select="$V22/Part">
                <xsl:for-each select="@Quantity">
                    <xsl:value-of select="."/>
                  </Quantity>
                </xsl:for-each>
              </xsl:for-each>
            </LineItem>
          </xsl:for-each>
        </LineItems>
      </xsl:for-each>
   </PurchaseOrder>
  </xsl:template>
</xsl:stylesheet>
```

Examples of Using Procedure COPYEVOLVE

Several examples are presented of using PL/SQL procedureDBMS XMLSCHEMA.copyEvolve to update an XML schema. (Be sure to back up all registered XML schemas and XML documents that reference them, before using the procedure.)

Example 20-3 loads a revised XML schema and evolution XSL stylesheet into Oracle XML DB Repository.

Example 20-4 shows how to use procedure DBMS XMLSCHEMA.copyEvolve to evolve the XML schema purchaseOrder.xsd to revisedPurchaseOrder.xsd using the XSLT stylesheet evolvePurchaseOrder.xsl.

Procedure DBMS XMLSCHEMA.copyEvolve evolves registered XML schemas in such a way that existing instance documents continue to remain valid.



Caution:

Before executing procedure DBMS XMLSCHEMA.copyEvolve, always back up all registered XML schemas and all XML documents that conform to them. Procedure copyEvolve deletes all documents that conform to registered XML schemas.

First, procedure copyEvolve copies the data in XML schema-based XMLType tables and columns to temporary tables. It then drops the original tables and columns, and deletes the old XML schemas. After registering the new XML schemas, it creates XMLType tables and columns and populates them with data (unless parameter GENTABLES is FALSE) but it does not create any auxiliary structures such as indexes, constraints, triggers, and row-level security (RLS) policies. Procedure copyEvolve creates the tables and columns as follows:

- It creates default tables while registering the new schemas.
- It creates tables that are not default tables using a statement of the following form:

```
CREATE TABLE table_name OF XMLType OID 'oid' XMLSCHEMA schema url ELEMENT element name
```

where OID is the original OID of the table, before it was dropped.

It adds XMLType columns using a statement of the following form:

```
ALTER TABLE table_name ADD (column_name XMLType) XMLType COLUMN column name XMLSCHEMA schema url ELEMENT element name
```

When a new XML schema is registered, types are generated if the registration of the corresponding old schema had generated types. If an XML schema was global before the evolution, then it is also global after the evolution. Similarly, if an XML schema was local before the evolution, then it is also local (owned by the same user) after the evolution. You have the option to preserve the temporary tables that contain the old documents, by setting parameter preserveOldDocs to TRUE. All temporary tables are created in the database schema of the current user. For XMLType tables, the temporary table has the columns shown in Table 20-3.

Table 20-3 XML Schema Evolution: XMLType Table Temporary Table Columns

Name	Туре	Comment
Data	CLOB	XML document from the old table, in CLOB format.
OID	RAW(16)	OID of the corresponding row in the old table.
ACLOID	RAW(16)	This column is present only if the old table is hierarchy-enabled. ACLOID of corresponding row in old table.
OWNERID	RAW(16)	This column is present only if old table is hierarchyenabled. OWNERID of corresponding row in old table.

For XMLType columns, the temporary table has the columns shown in Table 20-4.

Table 20-4 XML Schema Evolution: XMLType Column Temporary Table Columns

Name	Туре	Comment
Data	CLOB	XML document from the old column, in CLOB format.
RID	ROWID	ROWID of the corresponding row in the table containing this column.

Procedure <code>copyEvolve</code> stores information about the mapping from the old table or column name to the corresponding temporary table name in a separate table specified by parameter <code>mapTabName</code>. If <code>preserveOldDocs</code> is <code>TRUE</code>, then the <code>mapTabName</code> parameter must not be <code>NULL</code>, and it must not be the name of any existing table in the current database schema. Each row in the mapping table has information about one of the old tables/columns. Table 20-5 shows the mapping table columns.



Table 20-5 Procedure COPYEVOLVE Mapping Table

Column Name	Column Type	Comment
SCHEMA_URL	VARCHAR2(700)	URL of the schema to which this table or column conforms.
SCHEMA_OWNER	VARCHAR2(30)	Owner of the schema.
ELEMENT_NAME	VARCHAR2(256)	Element to which this table or column conforms.
TABLE_NAME	VARCHAR2(65)	Qualified name of the table (<owner_name>.<table_name>).</table_name></owner_name>
TABLE_OID	RAW(16)	OID of table.
COLUMN_NAME	VARCHAR2(4000)	Name of the column (NULL for XMLType tables).
TEMP_TABNAME	VARCHAR2(30)	Name of temporary table that holds the data for this table or column.

You can avoid generating any tables or columns after registering the new XML schema by setting parameter GENTABLES to FALSE. If GENTABLES is FALSE, parameter PRESERVEOLDDOCS must be TRUE and parameter MAPTABNAME must not be NULL. This ensures that the data in the old tables is not lost. This is useful if you do not want the tables to be created by the procedure, as described in section COPYEVOLVE Parameters and Errors.

By default, it is assumed that all XML schemas are owned by the current user. If this is not true, then you must specify the owner of each XML schema in the schemaOwners parameter.



Oracle Database SQL Language Reference for the complete description of ${\tt ALTER}$ ${\tt TABLE}$

Example 20-3 Loading Revised XML Schema and XSLT Stylesheet

Example 20-4 Updating an XML Schema Using DBMS_XMLSCHEMA.COPYEVOLVE

```
BEGIN

DBMS_XMLSCHEMA.copyEvolve(
   xdb$string_list_t('http://localhost:8080/source/schemas/poSource/xsd/purchaseOrder.xsd'),
   XMLSequenceType(XDBURIType('/source/schemas/poSource/revisedPurchaseOrder.xsd').getXML()),
```

The same query would have produced the following result before the schema evolution:

In-Place XML Schema Evolution

In-place XML schema evolution makes changes to an XML schema without requiring that existing data be copied, deleted, and reinserted. In-place evolution is thus much faster than copy-based evolution. However, in-place evolution also has several restrictions that do not apply to copy-based evolution.

You use procedure <code>DBMS_XMLSCHEMA.inPlaceEvolve</code> to perform in-place evolution. Using this procedure, you identify the changes to be made to an existing XML schema by specifying an XML schema-differences document, and you optionally specify flags to be applied to the evolution process.

In-place evolution constructs a new version of an XML schema by applying changes specified in a diffxML document, validates that new XML schema (against the XML schema for XML schemas), constructs DDL statements to evolve the disk structures used to store the XML instance documents associated with the XML schema, executes these DDL statements, and replaces the old version of the XML schema with the new, in that order. If the new version of the XML schema is not a valid schema, then in-place evolution fails.

- Restrictions for In-Place XML Schema Evolution
 - Because in-place XML schema evolution avoids copying data, it does not permit arbitrary changes to an XML schema. The primary restriction on using in-place evolution can be stated generally as a requirement that a given XML schema can be evolved in place in only a backward-compatible way.
- Supported Operations for In-Place XML Schema Evolution
 Some of the operations that are supported for in-place schema evolution are described.
 Some of these are not permitted in specific contexts, which are specified.
- Guidelines for Using In-Place XML Schema Evolution
 Guidelines that apply to in-place XML-schema evolution are presented. (Be sure to back up your data before performing in-place schema evolution.)

inPlaceEvolve Parameters

The parameters of PL/SQL procedure <code>DBMS_XMLSCHEMA.inPlaceEvolve</code> are described, as are the errors associated with this procedure.

The diffXML Parameter Document

The value of parameter <code>diffxml</code> of procedure <code>DBMS_XMLSCHEMA.inPlaceEvolve</code> is an <code>XMLType</code> instance that specifies the changes to be applied to an XML schema for in-place evolution. This <code>diffxml</code> document contains a sequence of operations that describe the changes between the old XML schema and the new (the intended evolution result).

Restrictions for In-Place XML Schema Evolution

Because in-place XML schema evolution avoids copying data, it does not permit arbitrary changes to an XML schema. The primary restriction on using in-place evolution can be stated generally as a requirement that a given XML schema can be evolved in place in only a backward-compatible way.

For the complete list of changes supported by in-place evolution, see Supported Operations for In-Place XML Schema Evolution.

Backward-compatible here means that any possible instance document that would validate against a given XML schema must also validate against a later (evolved) version of that XML schema. This applies to *all possible* conforming instance documents, not only to *existing* instance documents For XML data that is stored as binary XML, backward compatibility also means that any XML schema annotations that affect binary XML treatment must not change during evolution.

In addition to this general backward-compatibility restriction, there are some other restrictions for in-place evolution.

Backward-Compatibility Restrictions

Restrictions on in-place XML schema evolution are described that ensure backward compatibility of an evolved schema, so that any possible instance documents that satisfy the old XML schema also satisfy the new one.

• Other Restrictions on In-Place Evolution

Some restrictions on in-place XML schema evolution are necessary for reasons other than backward compatibility of the evolved XML schema.

Backward-Compatibility Restrictions

Restrictions on in-place XML schema evolution are described that ensure backward compatibility of an evolved schema, so that any possible instance documents that satisfy the old XML schema also satisfy the new one.

Changes in Data Layout on Disk

Certain changes to an XML schema alter the layout of the associated instance documents on disk, and are therefore not permitted. This situation is more common when the storage layer is tightly integrated with information derived from the XML schema, as is the case for object-relational storage.

Reorder of XML Schema Constructs

You cannot use in-place evolution to reorder XML schema elements in a way that affects the DOM fidelity of XML instance documents. For example, you cannot change the order of elements within a <sequence> element in a complex type definition.

Changes from a Collection to a Non-Collection

You cannot use in-place evolution to change a collection to a non-collection. An example would be changing from maxOccurs greater than one to maxOccurs equal to one. In-place

evolution thus cannot delete an element from a complex type if the deletion requires that a collection be evolved to a non-collection.

Model Changes within a complexType Element

A model is a group, choice, sequence, or all element. Within a complexType element you cannot use in-place evolution to either add a new model or replace an existing model with a model of another type (for example, replace a choice element with a sequence element).

Changes in Data Layout on Disk

Certain changes to an XML schema alter the layout of the associated instance documents on disk, and are therefore not permitted. This situation is more common when the storage layer is tightly integrated with information derived from the XML schema, as is the case for object-relational storage.

One such example is an XML schema, registered for object-relational storage mapping, that is evolved by splitting a complex type into two complex types. In Example 20-5, complex type ShippingInstructionsType is split into two complex types, Person-Name and Contact-Info, and the ShippingInstructionsType complex type is deleted.

Even if this XML schema has no associated instance documents, and therefore no data copy is required, a change in the layout of existing tables is required to accommodate future instance documents.

Example 20-5 Splitting a Complex Type into Two Complex Types

These code excerpts show the definitions of the original ShippingInstructionsType type and the new Person-Name and Contact-Info types.

```
<complexType name="ShippingInstructionsType">
    <sequence>
        <element name="name"</pre>
                               type="NameType" minOccurs="0"/>
        <element name="address" type="AddressType" minOccurs="0"/>
        <element name="telephone" type="TelephoneType" minOccurs="0"/>
   </sequence>
</complexType>
<complexType name="Person-Name">
    <sequence>
        <element name="name" type="NameType" minOccurs="0"/>
    </sequence>
</complexType>
<complexType name="Contact-Info">
    <sequence>
        <element name="address" type="AddressType" minOccurs="0"/>
        <element name="telephone" type="TelephoneType" minOccurs="0"/>
    </sequence>
</complexType>
```

Reorder of XML Schema Constructs

You cannot use in-place evolution to reorder XML schema elements in a way that affects the DOM fidelity of XML instance documents. For example, you cannot change the order of elements within a <sequence> element in a complex type definition.

As an example, if a complex type named <code>ShippingInstructionsType</code> requires that its child elements <code>name</code>, <code>address</code>, and <code>telephone</code> be in that order, you cannot use in-place evolution to change the order to <code>name</code>, <code>telephone</code>, and <code>address</code>.

Changes from a Collection to a Non-Collection

You cannot use in-place evolution to change a collection to a non-collection. An example would be changing from maxOccurs greater than one to maxOccurs equal to one. In-place evolution thus cannot delete an element from a complex type if the deletion requires that a collection be evolved to a non-collection.

Model Changes within a complexType Element

A model is a group, choice, sequence, or all element. Within a complexType element you cannot use in-place evolution to either add a new model or replace an existing model with a model of another type (for example, replace a choice element with a sequence element).

You can, however, add a global group element, that is, add a group element outside of a complexType element.

Other Restrictions on In-Place Evolution

Some restrictions on in-place XML schema evolution are necessary for reasons other than backward compatibility of the evolved XML schema.

- Changes to Attributes in Namespace xdb
 Except for attribute xdb:defaultTable, you cannot use in-place evolution to modify any attributes in namespace http://xmlns.oracle.com/xdb (which has the predefined prefix xdb).
- Changes from a Non-Collection to a Collection

 For object-relational XMLType data, you cannot use in-place evolution to change a noncollection object type to a collection object type. An example would be adding an element
 to a complex type if the element is already present in the type (or in a type related through
 inheritance).

Changes to Attributes in Namespace xdb

Except for attribute xdb:defaultTable, you cannot use in-place evolution to modify any attributes in namespace http://xmlns.oracle.com/xdb (which has the predefined prefix xdb).

Changes from a Non-Collection to a Collection

For object-relational XMLType data, you cannot use in-place evolution to change a non-collection object type to a collection object type. An example would be adding an element to a complex type if the element is already present in the type (or in a type related through inheritance).

Supported Operations for In-Place XML Schema Evolution

Some of the operations that are supported for in-place schema evolution are described. Some of these are not permitted in specific contexts, which are specified.

 Add an optional element to a complex type or group: Always permitted. An example is the addition of the optional element shipmethod in the following complex type definition:

Add an optional attribute to a complex type or attribute group: Always permitted. An
example is the addition of the optional attribute shipbydate in the following complex type
definition:

- Convert an element from simple type to complex type with simple content: Supported only
 if the storage model is binary XML.
- Modify the value attribute of an existing maxLength element: Always permitted. The value can only be increased, not decreased.
- Add an enumeration value: You can add a new enumeration value only to the end of an enumeration list.
- Add a global element: Always permitted. An example is the addition of the global element PurchaseOrderComment in the following schema definition:

Add a global attribute: Always permitted.

 Add or delete a global complex type: Always permitted. An example is the addition of the global complex type ComplexAddressType in the following schema definition:

- Add or delete a global simple type: Always permitted.
- Change the minOccurs attribute value: The value of minOccurs can only be decreased.
- Change the maxOccurs attribute value: The value of maxOccurs can only be increased, and this is only possible for data stored as binary XML. That is, you cannot make any change to the maxOccurs attribute for data stored object-relationally.
- Add or delete a global group or attributeGroup: Always permitted. An example is the addition of an Instructions group in the following type definition:

- Change the xdb:defaultTable attribute value: Always permitted. Changes are *not* permitted to any other attributes in the xdb namespace.
- Add, modify, or delete a comment or processing instruction: Always permitted.

Guidelines for Using In-Place XML Schema Evolution

Guidelines that apply to in-place XML-schema evolution are presented. (Be sure to back up your data before performing in-place schema evolution.)

- Before you perform an in-place XML-schema evolution:
 - Back up all existing data (instance documents) for the XML schema to be evolved.

Caution:

Make sure that you back up your data before performing in-place XML schema evolution, in case the result is not what you intended. There is no rollback possible after an in-place evolution. If any errors occur during evolution, or if you make a major mistake and need to redo the entire operation, you must be able to go back to the backup copy of your original

Perform a dry run using trace only, that is, without actually evolving the XML schema or updating any instance documents, produce a trace of the update operations that would be performed during evolution. To do this, set the flag parameter value to only INPLACE TRACE. Do not also use INPLACE EVOLVE.

After performing the dry run, examine the trace file, verifying that the listed DDL operations are in fact those that you intend.

After you perform an in-place XML-schema evolution:

If you are accessing the database using a client that caches data, or if you are not sure whether this is the case, then restart your client. Otherwise, the pre-evolution version of the XML schema might continue to be used locally, with unpredictable results.



Oracle Database Administrator's Guide for information about using trace files

inPlaceEvolve Parameters

The parameters of PL/SQL procedure DBMS XMLSCHEMA.inPlaceEvolve are described, as are the errors associated with this procedure.

This is the signature of procedure DBMS XMLSCHEMA.inPlaceEvolve:

```
procedure inPlaceEvolve(schemaURL IN VARCHAR2,
                      diffXML IN XMLType,
                       flags IN NUMBER);
```

Table 20-6 describes the individual parameters.

Table 20-6 Parameters of Procedure DBMS XMLSCHEMA.INPLACEEVOLVE

Parameter	Description
schemaURL	URL of the XML schema to be evolved (VARCHAR2).
diffXML	XML document (XMLType instance) that conforms to the xdiff XML schema, and that specifies the changes to apply and the locations in the XML schema where the changes are to be applied. For information about how to create the document for this parameter, see The diffXML Parameter Document.

Table 20-6 (Cont.) Parameters of Procedure DBMS_XMLSCHEMA.INPLACEEVOLVE

Parameter Description

flags

A bit mask that controls the behavior of the procedure. You can set the following bit values in this mask independently, summing them to define the overall effect. The default flags value is 1 (bit 1 on, bit 2 off), meaning that in-place evolution is performed and no trace is written.

- INPLACE_EVOLVE (value 1, meaning that bit 1 is on) Perform in-place XML schema evolution. Construct a new XML schema and validate it (against the XML schema for XML schemas). Construct the DDL statements needed to evolve the instance-document disk structures. Execute the DDL statements. Replace the old XML schema with the new.
- INPLACE_TRACE (value 2, meaning that bit 2 is on) Perform all steps necessary
 for in-place evolution, except executing the DDL statements and overwriting the old
 XML schema with the new, then write both the DDL statements and the new XML
 schema to a trace file.

That is, each of the bits constructs the new XML schema, validates it, and determines the steps needed to evolve the disk structures underlying the instance documents. In addition:

- Bit INPLACE_EVOLVE carries out those evolution steps and replaces the old XML schema with the new.
- Bit INPLACE_TRACE saves the evolution steps and the new XML schema in a trace file (it does not carry out the evolution steps).

Procedure DBMS XMLSCHEMA.inPlaceEvolve raises an error in the following cases:

- An XPath expression is invalid, or is syntactically correct but does not target a node in the XML schema.
- The diffXML document does not conform to the xdiff XML schema.
- The change makes the XML schema invalid or not well formed.
- A generated DDL statement (CREATE TYPE, ALTER TYPE, and so on) causes a problem when it is executed.
- An index object associated with an XMLType table is in an unsafe state, which could be caused by partition management operations.

The diffXML Parameter Document

The value of parameter <code>diffxml</code> of procedure <code>DBMS_XMLSCHEMA.inPlaceEvolve</code> is an <code>XMLType</code> instance that specifies the changes to be applied to an XML schema for in-place evolution. This <code>diffxml</code> document contains a sequence of operations that describe the changes between the old XML schema and the new (the intended evolution result).

The changes specified by the ${\tt diffxml}$ document are applied in order.

You must create the XML document to be used for the diffXML parameter You can do this in any of the following ways:

- The XMLDiff JavaBean (oracle.xml.differ.XMLDiff)
- The xmldiff command-line utility
- SQL function XMLDiff

The diffXML parameter document must conform to the xdiff XML schema.

The rest of this section presents examples of some operations in a document that conforms to the xdiff XML schema.

diffXML Operations and Examples

Operations that can be specified in the diffXML document supplied to procedure DBMS_XMLSCHEMA.inPlaceEvolve are described. An example XML document that conforms to the xdiff XML schema is shown.

Related Topics

xdiff.xsd: XML Schema for Comparing Schemas for In-Place Evolution
 A full listing is presented of xdiff.xsd, the Oracle XML DB-supplied XML schema to which
 the document specified as the diffXML parameter to procedure
 DBMS XMLSCHEMA.inPlaceEvolve must conform.

See Also:

- Oracle XML Developer's Kit Programmer's Guide for information on using the XMLDiff JavaBean
- Oracle XML Developer's Kit Programmer's Guide for information on commandline utility xmldiff
- Oracle Database SQL Language Reference for information on SQL function
 XMLDiff

diffXML Operations and Examples

Operations that can be specified in the diffXML document supplied to procedure DBMS_XMLSCHEMA.inPlaceEvolve are described. An example XML document that conforms to the xdiff XML schema is shown.

The <append-node> element is used for most of the supported changes, such as adding a new attribute to a complex type or appending a new element to a group.

The <insert-node-before> element specifies that a node of the given type should be inserted before the specified node. The xpath attribute specifies the location of the specified node and the node-type attribute specifies the type of node to be inserted. The node to be inserted is specified by the <content> child element. The <insert-node-before> element is mainly used for inserting comments and processing instructions, and for changing and adding add annotation elements.

The <delete-node> element specifies that the node with the given XPath (specified by the xpath attribute) should be deleted along with all its children. For example, you can use this element to delete comments and annotation elements. You can also use this element, in conjunction with <append-node> or <insert-node-before>, to make changes to an existing node.

Example 20-6 shows an XML document for the diffXML parameter that specifies the following changes:

- Delete complex type PartType.
- Add complex type PartType with a maximum length of 28.
- Add a comment before element ShippingInstructions.



Add a required element shipmethod to element ShippingInstructions.

Example 20-6 diffXML Parameter Document

```
<xd:xdiff xmlns="http://www.w3c.org/2001/XMLSchema"</pre>
           xmlns:xd="http://xmlns.oracle.com/xdb/xdiff.xsd"
           xmlns:xsi="http://www.w3c.org/2001/XMLSchema-Instance"
           xsi:schemaLocation="http://xmlns.oracle.com/xdb/xdiff.xsd
           http://xmlns.oracle.com/xdb/xdiff.xsd">
 <xd:delete-node xpath="/schema/complexType[@name=&quote;PartType&quote;]//maxLength/>
 <xd:append-node
 parent-xpath = "/schema/complexType[@name=&quote;PartType&quote;]//restriction"
 node-type = "element">
 <xd:content>
    <xs:maxLength value = "28"/>
 </xd:content>
</xd:append-node>
 <xd:insert-node-before
 xpath="/schema/complexType[@name =&quote;ShippingInstructionsType&quote;]/sequence"
 node-type="comment">
 <xd:content>
   <!-- A type representing instructions for shipping -->
 </xd:content>
</xd:insert-node-before>
 <xd:append-node
 parent-xpath="/schema/complexType[@name=&quote;ShippingInstructionsType&quote;]/sequence"
 node-type="element">
 <xd:content>
  <xs:element name = "shipmethod" type = "xs:string" minOccurs = "1"/>
 </xd:content>
</xd:append-node>
</xd:xdiff>
```

